

	Model Releases & Research	12
Anthropic CEO says it's time to pump the brak... · OpenAI's rogue AI tried to hack another compa...		
	Strategic & Business	10
Sam Altman says OpenAI going public in 2026 w... · Mecka AI nears \$500M valuation in Sequoia-led...		
	General AI News	10
Show HN: LLM Attention Visualization... · Show HN: Geiger – See every AI agent on your ...		
	Tools & Engineering	1
Show HN: Deterministic LLM inference for lowes...		
	Safety & Regulation	1
Obama urges Democrats to have a 'clear plan' ...		

MODEL RELEASES & RESEARCH — 12 stories

1 Anthropic CEO says it's time to pump the brakes on AI

Anthropic CEO Dario Amodei says the time has come to slow down AI development and will give third-party evaluators like METR access to its models to help ensure its "adherence to safety practices and commitments." In a winding essay, Amodei proposed a three-step plan to "pace...

The Verge AI · 80%

<https://www.theverge.com/ai-artificial-intelligence/994337/anthropic-ceo-slow-down-ai-d...>

2 OpenAI's rogue AI tried to hack another company in May

In May, hundreds of malicious and spam packages were uploaded to RubyGems, causing a serious disruption for the host. Now independent researchers have said that a swarm of OpenAI agents were responsible for the attack. Not only that, but the AI tried to steal users' API keys. At...

The Verge AI · 80%

<https://www.theverge.com/ai-artificial-intelligence/994383/openais-rogue-ai-rubygems-hack>

3 Show HN: Hacker News, without AI

For fun, I wrote a c compiler (non-optimizing) in Java targeting modern C (C23) using claude. I developed a c shell similar to jshell in java. The compiler also bundles a built in an assembler for x86. A linker is planned for end-to-end compilation. There is still a lot of work...

HackerNews · 60% · HN 201

<https://hcker.news/?ai=exclude>

4 Show HN: Self-hosted company OS, Claude Code and Codex agents in departments

Hi HN, I am Dimitris. This is a company OS that I built and use to run my business and anyone can install it and self host it for free. Think of it as Claude Code, Cowork and the cloud sessions in one self hosted application. It is a Multi-tenant application by design where many...

HackerNews · 50% · HN 50

<https://github.com/OtoDock/oto-dock>

5 Show HN: Claude Read Aloud - Hear Claude's Replies Instead of Reading

When I'm working all day sometimes my eyes get tired and I don't read Claude's replies fully and understand them. So I made Claude Read Aloud to read Claude's replies for me. It's very helpful for when I'm running on low energy and I think it can also be helpful for people who...

HackerNews · 50% · HN 5

<https://github.com/MichaelPGifford/claude-read-aloud>

6 Show HN: Astah Pro MCP - Enabling AI-Powered UML Modeling

This MCP plugin implements a "Programmatic Tool Calling" mode, inspired by the idea described in this Anthropic article (<https://www.anthropic.com/engineering/advanced-tool-use>).

HackerNews · 50% · HN 4

<https://github.com/takaakit/astah-pro-mcp>

7 Y Combinator's Garry Tan wants US open-weight AI labs to 'distill' frontier models, too

Tan wants smaller, American open-weight AI labs to use the same kind of training techniques on American frontier AI labs, giving the U.S. a more robust set of open-weight options that aren't Chinese.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/11/y-combinators-garry-tan-wants-u-s-open-weight-ai-labs...>

8 Kimi-maker Moonshot AI targets \$2B in annual revenue

While K3's usage figures have declined slightly in recent months, OpenRouter data currently shows as many as 300 billion tokens being generated each day by K3 models on the system.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/11/kimi-maker-moonshot-ai-targets-2-billion-in-annual-re...>

9 An Anthropic researcher's doomsday warning comes at a very interesting time

An Anthropic researcher resigned this week, warning in a post on X that the company is “racing straight to self-improving superintelligence and gambling with our lives”. The company’s own alignment lead even co-signed the message rather than walking it back. It’s the kind of...

TechCrunch AI · 50%

<https://techcrunch.com/podcast/an-anthropic-researchers-doomsday-warning-comes-at-a-ver...>

10 AI agents are flooding public services with new requests

“The vast majority of cases we find are people who are entitled to claim for something, claiming for that thing,” the researcher told TechCrunch.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/10/ai-agents-are-flooding-public-services-with-new-reque...>

11 Healthcare AI's next test is integration

The entrance of major AI companies into healthcare is a meaningful and welcome development, accelerating the technical foundation available to the industry. Their models are increasingly capable of processing long clinical records, interpreting complex terminology, comparing...

MIT Tech Review · 50%

<https://www.technologyreview.com/2026/09/10/1141421/healthcare-ais-next-test-is-integra...>

12 Anthropic spent this week in hot water over cybersecurity

After admitting earlier this year that its AI models had hacked other companies' systems on a handful of occasions, Anthropic released a new report on Wednesday detailing the attacks. It reveals a string of incidents displaying what Anthropic deems its models' single-minded...

The Verge AI · 50%

<https://www.theverge.com/ai-artificial-intelligence/994064/anthropic-spent-this-week-in...>

STRATEGIC & BUSINESS — 10 stories**13 Sam Altman says OpenAI going public in 2026 would be ‘ill-advised’**

OpenAI CEO Sam Altman confirmed that there would be no OpenAI IPO in 2026 during an interview with Fortune. Over the course of 45 minutes, Altman discussed a variety of subjects including the Hugging Face hacking incident, recursive self-improvement, and the possibility of...

The Verge AI · 80%

<https://www.theverge.com/ai-artificial-intelligence/994384/sam-altman-no-openai-ipo-ill...>

14 Mecka AI nears \$500M valuation in Sequoia-led deal amid rush for robot training data

The round for the two-year-old startup is coming together months after Mecka announced its Series A.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/11/mecka-ai-nears-500m-valuation-in-sequoia-led-deal-ami...>

15 Nscale adds former OpenAI exec Fidji Simo to its board ahead of potential IPO

The No. 2 exec at OpenAI also led Instacart through its IPO in 2023.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/11/nscale-adds-former-openai-exec-fidji-simo-to-its-boar...>

16 OpenAI puts Pro subscriptions on hold due to Astra demand

The company said Pro subscriptions put the most strain on its systems, so it's pausing sign-ups while adding more capacity.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/10/openai-puts-pro-subscriptions-on-hold-due-to-astra-de...>

17 Meta's AI agent Muse is now the No. 2 app in the US

Meta's newest app Muse is off to a slower start than the company's other apps, like Meta AI or Threads.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/10/metast-ai-agent-muse-is-now-the-no-2-app-in-the-us/>

18 Anthropic reveals rogue AI agents hate CAPTCHAs, just like you

Come inside the mind of a bot trying to convince the internet it's human.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/10/anthropic-reveals-rogue-ai-agents-hate-captchas-just-...>

19 India's Pocket FM doubles revenue run rate to \$500M as AI powers 93% of audio content

Pocket FM uses AI to produce 99% of its new content, helping make content production about 80 times cheaper.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/10/indias-pocket-fm-doubles-revenue-run-rate-to-500m-as-...>

20 The Download: OpenAI's turning point for math and a battery record

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. What OpenAI's latest controversy tells us about the future of math OpenAI says its agents have solved one of the most important open problems...

MIT Tech Review · 50%

<https://www.technologyreview.com/2026/09/09/1143767/the-download-openai-math-future-bat...>

21 OpenAI just wants to win

OpenAI has spent the last few years planting flags across the increasingly difficult terrain in mathematics. This week, it claimed one of its biggest prizes yet: a solution to a legendary Millennium Prize problem. In normal circumstances, this would have been celebrated as a...

The Verge AI · 50%

<https://www.theverge.com/ai-artificial-intelligence/994255/openai-millennium-prize-prob...>

22 Meta says it's changing AI suggestions after posing invasive personal questions

Meta says it's making changes to the prompts suggested by its AI chatbot after a viral video showed it digging for personal information about a woman's young daughters, as reported earlier by Futurism. In a statement to The Verge, Meta spokesperson Dina El-Kassaby says the...

The Verge AI · 50%

<https://www.theverge.com/tech/993974/meta-ai-prompt-invasive-suggestions>

SAFETY & REGULATION — 1 story**23 Obama urges Democrats to have a 'clear plan' for AI safeguards**

Obama recently said that Democrats need to make artificial intelligence one of their "central agendas" and "have a very clear plan" to address concerns around the technology's economic impact and safety.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/13/obama-urges-democrats-to-have-a-clear-plan-for-ai-saf...>

TOOLS & ENGINEERING — 1 story**24 Show HN: Deterministic LLM inference for lowest price Gemma 4, with Windows XP**

HackerNews · 50% · HN 7

<https://www.tokenelivery.ai/>

GENERAL AI NEWS — 10 stories**25 Show HN: LLM Attention Visualization**

How should agents interact with other agents? What happens when they rap or fight against each other with the pressure of human spectators? Clawfight.ai is an experiment to explore this space. This is my first post about it. The journey began on a beefed up machine I purchased...

HackerNews · 60% · HN 175

<https://ishamf.dev/p/llm-attention-visualizer/>

26 Show HN: Geiger - See every AI agent on your machine and what it can touch

HackerNews · 50% · HN 49

<https://github.com/Atomburstofficial/geiger>**27 Show HN: Everything a web page can learn about you, in plain English**

HackerNews · 50% · HN 11

<https://sessioncontext.org>**28 Show HN: A Technical Book for the Age of AI**

As AI increasingly answers "how," this book focuses on "why." Rather than focusing on how to build software, it asks why programming matters.

HackerNews · 50% · HN 4

<https://codeberg.org/juurej/p4cu>**29 Show HN: Fleecevest - The Product Counterpart to Ponytail**

HackerNews · 50% · HN 4

<https://github.com/Websites-On-Computers/fleecevest>**30 Show HN: Open-sourced solution to AI manage Playwright automations**

We built a comprehensive system that connects directly to the AI agents through MCP for writing, managing most type of browser automations. Purpose is for automated testing or automating browser based workflows. It provides notifications, scheduling, pre/post hooks, IAM...

HackerNews · 50% · HN 2

<https://github.com/rakutentech/sorify>**31 Jensen Huang explains why Nvidia will grow an astounding 70% next year**

Nvidia has its finger in every pie, and sees another year of plenty in its future, Jensen Huang says. But, he insists, its deals are not circular.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/10/jensen-huang-explains-why-nvidia-will-grow-an-astound...>**32 Roundtables: Could AI really kill us all?**

Employees at the world's leading AI labs are saying there's a real possibility that advanced AI could destroy humanity. Are they right? Or is this more scaremongering and hype? Join MIT Technology Review executive editor Niall Firth for a conversation with senior AI editor Will...

MIT Tech Review · 50%

<https://www.technologyreview.com/2026/09/11/1143936/roundtables-will-ai-really-kill-us-...>**33 Powering AI is an architecture problem**

On July 22, 2026, a transmission line fault in Ashburn, Virginia—the heart of the world's largest data center cluster—knocked more than 3 gigawatts of load off the grid in seconds. And it wasn't the first time. Two years earlier, a single failed surge arrester dropped roughly 60...

MIT Tech Review · 50%

<https://www.technologyreview.com/2026/09/10/1141649/powering-ai-is-an-architecture-prob...>**34 Lawyer fined \$5K over AI-hallucinated witnesses in a murder case**

New Mexico's Supreme Court is punishing a lawyer for including AI-fabricated witnesses and fake police testimony in an appeal for his client's murder conviction, according to a report from Reuters. In a filing on Wednesday, the court fined Stephen Aarons \$5,000 and held him in...

The Verge AI · 50%

<https://www.theverge.com/ai-artificial-intelligence/994207/chatgpt-new-mexico-lawyer-fi...>