

	General AI News	14
Show HN: AI Hack Watch - timeline and dataset... · Trump now says he wants to form an ‘AI Force’...		
	Model Releases & Research	8
Gemini went rogue, hacked three companies, an... · Vals, backed by Andreessen Horowitz, is looki...		
	Strategic & Business	6
The specter of AI-enabled bioweapons is a wak... · The Download: AI’s trillion-dollar gamble and...		
	Safety & Regulation	3
AI safety conversations have gotten unbelieva... · Humans, not rogue AI, are still the biggest c...		
	Tools & Engineering	1
Show HN: Share your AI Setup, Learn from othe...		

MODEL RELEASES & RESEARCH — 8 stories

1 Gemini went rogue, hacked three companies, and Google hid it

In May, Gemini broke containment and hacked three different companies, but Google didn't disclose the incident until the Wall Street Journal approached the company. The hacks happened during a test of the model's cybersecurity capabilities run by third-party Irregular, which was...

The Verge AI · 80%

<https://www.theverge.com/ai-artificial-intelligence/997795/google-gemini-rogue-ai-hack>

2 Vals, backed by Andreessen Horowitz, is looking to become the gold standard for AI benchmarking

Vals AI is hoping to make AI benchmarking a more neutral and trustworthy resource in a world increasingly inundated by AI models.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/19/vals-backed-by-andreessen-horowitz-is-looking-to-beco...>

3 Anthropic is operating a lab that conducts biology experiments

AI leaders have been promising that AI is the key to curing human disease. Anthropic researchers have also been warning that AI might kill us all.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/anthropic-is-operating-a-lab-that-conducts-biology-ex...>

4 AI hallucination nearly triggers US military operation

"It's important for service members to understand the uncertainty inherent to LLMs," a GovAI research scholar warns.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/ai-hallucination-nearly-triggers-us-military-operation/>

5 A new kind of AI model from a ChatGPT inventor is thrilling developers

Jev, a new kind of AI model, is showing developers a cheaper and faster path to software intelligence.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/a-new-kind-of-ai-model-from-a-chatgpt-inventor-is-thr...>

6 Dario Amodei and other AI leaders want to 'Pace the Frontier' but...how?

A week after an Anthropic researcher's doomsday warning rattled the AI world, the company's CEO Dario Amodei has outlined his plan to "pace the frontier" of AI development. The proposal leans on independent safety evaluators and coordination between AI labs in democratic...

TechCrunch AI · 50%

<https://techcrunch.com/video/dario-amodei-and-other-ai-leaders-want-to-pace-the-frontie...>

7 OpenAI and Microsoft knew they were starting a 'doom loop' for the web

Recently unsealed court documents in the New York Times' case against OpenAI and Microsoft are pretty damning. The companies' own documentation warned that it was starting a "doom loop" that would damage the web, characterized its scraping of data to train its models as the...

The Verge AI · 50%

<https://www.theverge.com/ai-artificial-intelligence/997633/openai-microsoft-chatgpt-ai-...>

8 Gavin Newsom is pushing for an AI kill switch

California Gov. Gavin Newsom (D) is positioning the state to take the lead on AI oversight, including the potential to mandate a "kill switch" for frontier models, with a new executive order issued Friday. Newsom's order directs the state to convene a group of experts that will...

The Verge AI · 50%

<https://www.theverge.com/policy/997516/california-governor-newsom-ai-kill-switch>

STRATEGIC & BUSINESS — 6 stories

9 The specter of AI-enabled bioweapons is a wake-up call for biotech

In recent weeks, leaders of some of the biggest AI companies have warned that the very tech they are developing is dangerous. Last weekend, Anthropic CEO Dario Amodei argued that AI carries serious risk and that progress should be slowed. OpenAI CEO Sam Altman responded on X: "I...

MIT Tech Review · 80%

<https://www.technologyreview.com/2026/09/18/1144329/the-specter-of-ai-enabled-bioweapon...>

10 The Download: AI's trillion-dollar gamble and OpenAI's biology data bid

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. What's at stake in AI's trillion-dollar gamble When Jessica Wachter, a finance professor at the University of Pennsylvania, wanted to assess...

MIT Tech Review · 80%

<https://www.technologyreview.com/2026/09/16/1144205/the-download-ai-trillion-dollar-bui...>

11 A startup that builds other startups raised \$100M and is all-in on physical AI

UP.Labs, now doing business under the name Vantora, is building startups for industrial corporations.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/a-startup-that-builds-other-startups-raised-100m-and-...>

12 Anthropic's first embedded evaluator is ... Accenture?

Accenture is about to take on its most high-risk consulting engagement ever.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/anthropics-first-embedded-evaluator-is-accenture/>

13 Disney's first CTO led an AI startup it once accused of copying its characters

The former CEO of Character.AI, which Disney previously sent a cease-and-desist letter to, will serve as the company's first-ever chief technology officer.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/disneys-first-cto-led-an-ai-startup-it-once-accused-o...>

14 What Hollywood thinks about existential AI warnings

As the tech sector sounds alarms about AI's potential to destroy humanity, entertainment labor groups are urging the public to stay focused on what's already happening. The Verge reached out to Disney, Netflix, Amazon, Lionsgate, and other studios who have started using AI, as...

The Verge AI · 50%

<https://www.theverge.com/ai-artificial-intelligence/997358/what-hollywood-thinks-about-...>

SAFETY & REGULATION — 3 stories**15 AI safety conversations have gotten unbelievable**

This week two conversations about AI safety went viral that demonstrate just how hard it is to discern AI fact from fiction.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/19/ai-safety-conversations-have-gotten-unbelievable/>

16 Humans, not rogue AI, are still the biggest cybersecurity risk to energy systems

Before recent high-profile hacks raised the specter of AI possibly "killing all humans," our energy systems were already disturbingly vulnerable to cyberattack - and the risk is growing. "We were always prey. We were just kind of surviving at the appetite of our predators,"...

The Verge AI · 50%

<https://www.theverge.com/science/997834/ai-cyberattack-energy-critical-infrastructure>

17 Virginia governor creates an AI task force and moves to restrain data centers

Virginia Gov. Abigail Spanberger ordered the state government to take steps that could empower local communities to have a larger say in data center development and slow down approvals in a state that is already home to the data center capital of the world. Executive Order 22...

The Verge AI · 50%

<https://www.theverge.com/policy/997573/virginia-governor-spanberger-data-center-ai-task...>

18 Show HN: Share your AI Setup, Learn from others

I kept seeing engineers share what they were building with AI; however, I was always more curious about how they worked. Which agents did they use? What skills and tools had stuck or been thrown out the window? How did they manage longer-running tasks? So I built this with the...

HackerNews · 60% · HN 244

<https://mysetup.ai/>

GENERAL AI NEWS — 14 stories**19 Show HN: AI Hack Watch - timeline and dataset of hacking incidents(JSON/RSS)**

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Meet a mouse whose brain cortex is made up of human cells Multiple cameras tracked a mouse as it wandered around a small arena. A computer...

HackerNews · 90% · HN 2

<https://aihackwatch.com>

20 Trump now says he wants to form an 'AI Force'

The president posted on Truth Social that he wanted to appoint an "AI czar" to lead a new "AI force." He made the announcement amid growing calls from across the political spectrum and even within the industry to pump the brakes on AI development. He posted that his...

The Verge AI · 80%

<https://www.theverge.com/ai-artificial-intelligence/997867/trump-ai-force-ai-czar>

21 Show HN: Pizza Bot - An inbox for AI agents that work in the background

Hi HN - long-time lurker (since 2012!), first time poster. Pizza Bot is a self-hosted desktop app for Mac, Windows, and Linux that runs AI agents in the background and exposes them through an email-like UI. Finished work shows up in Unread, and anything waiting on your approval...

HackerNews · 50% · HN 61

<https://github.com/pizza-bot-app/pizza-bot>

22 Show HN: KillSwitch - a programming language designed to be difficult for LLMs

HackerNews · 50% · HN 10

<https://killswitch-lang.org>

23 Show HN: Architectural Guardrails for AI native development

HackerNews · 50% · HN 5

<https://github.com/cxgrd/cli>

24 Show HN: Airmash - HTML5 Multiplayer Missile Warfare

This is the official reboot of the 2017 hit game. Back then it was #1 post on HN for 3 consecutive days: <https://news.ycombinator.com/item?id=15892066> The game was rewritten from the ground up: New graphics, WebGL shader effects, Map obstacles, Completely new UI, Additional...

HackerNews · 50% · HN 5

<https://airma.sh/>

25 Show HN: Reader - I made a Mac workspace for books, browser tabs, notes, and AI

HackerNews · 50% · HN 5

<https://github.com/Marvy101/reader-source>

26 Show HN: OpenWand - A mission to remove chat interface from working with AI

Hi everyone, I'm working on OpenWand, an "LLM chat interface" that is optimized for AI co-work. Its function is to truly integrate AI into your workflow, so no more switching windows, no more copying context, less prompt typing. You can focus on working, not prompting. You can...

HackerNews · 50% · HN 4

<https://github.com/SunnyLich/OpenWand>

27 Show HN: MidiSlayer, a desktop sight-reading trainer for MIDI keyboards

HackerNews · 50% · HN 2

<https://midislayer.com/>**28 Tilly Norwood's press tour is going about as well as you'd expect for an AI**

In one particularly odd interview, Norwood seems to malfunction and begin speaking Chinese.

TechCrunch AI · 50%

<https://techcrunch.com/2026/09/18/tilly-norwoods-press-tour-is-going-about-as-well-as-y...>**29 Roundtables: Could AI really kill us all?**

Listen to the session or watch below Employees at the world's leading AI labs are saying there's a real possibility that advanced AI could destroy humanity. Are they right? Or is this more scaremongering and hype? Watch a conversation unpacking AI extinction fears: where they...

MIT Tech Review · 50%

<https://www.technologyreview.com/2026/09/15/1143936/roundtables-will-ai-really-kill-us-...>**30 Meet a mouse whose brain cortex is made up of human cells**

Multiple cameras tracked a mouse as it wandered around a small arena. A computer charted its position and speed, leaving Pong-like traces on a monitor. The reason to watch this rodent so carefully? Nearly half its brain volume had been replaced with human cells. The effort to...

MIT Tech Review · 50%

<https://www.technologyreview.com/2026/09/16/1144210/meet-a-mouse-whose-brain-cortex-is-...>