



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	13	Media
TechCrunch AI	9	Media
The Verge AI	7	Media
HackerNews	5	Community
MIT Tech Review	2	Media

VERIFIED INTELLIGENCE — 5 stories

VERI

Show HN: Needle: We Distilled Gemini Tool Calling into a 26M Model

HackerNews

+1 source

100%

HN 390

Hey HN, Henry here from Cactus. We open-sourced Needle, a 26M parameter function-calling (tool use) model. It runs at 6000 tok/s prefill and 1200 tok/s decode on consumer devices. We were always frustrated by the little effort made towards building agentic models that run on...

<https://github.com/cactus-compute/needle>

VERI

Explaining Graph Neural Networks for Node Similarity on Graphs

ArXiv CS.AI

+1 source

80%

arXiv:2407.07639v2 Announce Type: replace-cross Abstract: Similarity search is a fundamental task for exploiting information in various applications dealing with graph data, such as citation networks or knowledge graphs. While this task has been intensively approached from...

<https://arxiv.org/abs/2407.07639>

VERI

Rivian's AI-powered voice assistant is ready to roll

The Verge AI

+1 source

80%

Rivian's AI-powered voice assistant is rolling out today to the company's vehicle fleet. The assistant will be available through a software update to all compatible Rivian Gen 1 and Gen 2 vehicle owners who subscribe to the company's Connect Plus cellular service, which costs...

<https://www.theverge.com/transportation/928651/rivian-ai-voice-assistant-r1-r2>

VERI

The Download: a Nobel winner on AI, and the case for fixing everything

MIT Tech Review

+1 source

80%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Three things in AI to watch, according to a Nobel-winning economist A few months before he won the Nobel Prize in economics in 2024, Daron...

<https://www.technologyreview.com/2026/05/12/1137103/the-download-nobel-winner-ai-mainte...>

VERI

Sam Altman takes the stand in trial against Elon Musk

The Verge AI

+1 source

80%

OpenAI CEO Sam Altman has begun his testimony against Elon Musk in a high-profile jury trial in a California federal courtroom. Altman, alongside OpenAI president Greg Brockman, is a primary defendant in the trial brought by Musk. Altman, Brockman, and Musk were all part of the...

<https://www.theverge.com/ai-artificial-intelligence/916975/altman-takes-stand-elon-musk...>

CONFIRMED SIGNALS — 26 stories

CONF

Show HN: Statewright - Visual state machines that make AI agents reliable

HackerNews

50%

HN 90

Hi HN, I'm Namanyay from Gigacatalyst (link: <https://gigacatalyst.com/>). Gigacatalyst allows sales, CS, and users to build one-off features, so your SaaS can support long-tail customer workflows and engineers aren't pulled away from the roadmap. When you sell software to large...

<https://github.com/statewright/statewright>

CONF

Show HN: Profine - Profile and rewrite your PyTorch training loop on real GPUs

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/ProfineAI/profine-cli>

CONF

Show HN: Rapunzel - a tree-style tab terminal emulator for Codex Claude Gemini

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/salmanjavid/rapunzel/tree/main>

CONF

Show HN: OpenClaw OS - OSS Claude Cowork Built on Top of OpenClaw

HackerNews 50% HN 3

Hi HN, We made OSS Claude Cowork, built as an OpenClaw plugin. It lets you create live artifacts (like Claude) that connect to datasources instead of datasets. (eg: fetching Stripe data automatically) Other tools(Paperclip, Multica) focus on task management but our vision is to...

<https://github.com/thesysdev/openclaw-os>

0 CONF

Musk mulled handing OpenAI to his children, Altman testifies

TechCrunch AI 50%

Altman said that Musk's focus on controlling the initial for-profit gave him pause because OpenAI was dedicated to keeping advanced AI out of the hands of a single person, and Altman, with his experience running the prominent startup accelerator Y Combinator, knew "founders who..."

<https://techcrunch.com/2026/05/12/musk-mulled-handing-openai-to-his-children-altman-tes...>

1 CONF

Anthropic warns investors against secondary platforms offering access to its shares

TechCrunch AI 50%

"Any sale or transfer of Anthropic stock, or any interest in Anthropic stock, offered by these firms is void and will not be recognized on our books and records," the company's support page reads.

<https://techcrunch.com/2026/05/12/anthropic-warns-investors-against-secondary-platforms...>

2 CONF

Google adds Gemini-powered dictation to Gboard, which could be bad news for dictation startups

TechCrunch AI 50%

Google's transcription feature will initially launch with Samsung Galaxy and Google Pixel phones.

<https://techcrunch.com/2026/05/12/google-adds-gemini-powered-dictation-to-gboard-which-...>

3 CONF

The AI legal services industry is heating up — Anthropic is getting in on the action

TechCrunch AI 50%

Anthropic's new tools are designed to help law firms automate specific clerical functions — things like document search and review, case law resources, deposition prep, document drafting, and other related areas.

<https://techcrunch.com/2026/05/12/the-ai-legal-services-industry-is-heating-up-anthropi...>

4 CONF

Google brings agentic AI and vibe-coded widgets to Android

TechCrunch AI 50%

Gemini Intelligence will also include Gboard-based dictation and form-filling capabilities.

<https://techcrunch.com/2026/05/12/google-brings-agentic-ai-and-vibe-coded-widgets-to-an...>

5 CONF

Dessn raises \$6M for its production-focused design tool

TechCrunch AI 50%

A new startup called Dessn has raised \$6 million to build AI-powered design tools that work directly with production codebases.

<https://techcrunch.com/2026/05/12/dessn-raises-6m-for-its-production-focused-design-tool/>

6 CONF

AI voice startup Vapi hits \$500M valuation after winning Amazon Ring over 40 rivals

TechCrunch AI 50%

Vapi says its enterprise business has grown tenfold since early 2025 as companies shift customer support and sales calls to AI agents.

<https://techcrunch.com/2026/05/12/vapi-hits-500m-valuation-as-amazon-ring-chose-its-ai-...>

7 CONF

World Models: 10 Things That Matter in AI Right Now

MIT Tech Review 50%

World models recently made our list of 10 Things That Matter in AI Right Now. Watch executive editor Niall Firth explain why this emerging area of AI is gaining so much attention. Join MIT Technology Review editors and reporters for a subscriber-only Roundtables discussion, "Can..."

<https://www.technologyreview.com/2026/05/12/1137134/world-models-10-things-that-matter-...>

8 CONF

Meta won't let you block its AI account on Threads

The Verge AI 50%

Meta announced on Tuesday that it's testing a Threads feature that lets users tag a Meta AI account to get answers to questions or context about a conversation on the platform. If you've spent any time looking at replies on X as of late, this new feature sounds a lot like Meta's...

<https://www.theverge.com/tech/929091/meta-ai-threads-account-block>

9 CONF

Sam Altman says Elon Musk's mind games were damaging OpenAI

The Verge AI 50%

OpenAI CEO Sam Altman says Elon Musk did "huge damage" to the culture of the AI startup. During testimony as part of Musk's lawsuit against OpenAI, Altman said Musk required OpenAI president Greg Brockman and former chief scientist Ilya Sutskever to rank researchers by their...

<https://www.theverge.com/ai-artificial-intelligence/928861/openai-sam-altman-elon-musk-...>

0 CONF

Gemini's latest updates are all about controlling your phone

The Verge AI 50%

It is, once again, Gemini season. Google is announcing a host of new Gemini features during its pre-I/O Android showcase, many of which aim to help use your phone for you. You'll find Gemini in more places, like Chrome on Android, in your autofill suggestions, and all up in your...

<https://www.theverge.com/tech/928724/gemini-intelligence-android-io-autofill>

1 CONF

Parents say ChatGPT got their son killed with bad advice on party drugs

The Verge AI 50%

The family of a 19-year-old college student is suing OpenAI over claims that his conversations with ChatGPT led to an accidental overdose. In the lawsuit filed on Tuesday, Sam Nelson's parents allege ChatGPT "encouraged" the teen to "consume a combination of substances that any..."

<https://www.theverge.com/ai-artificial-intelligence/928691/openai-chatgpt-wrongful-deat...>

- 2

CONF

George Clooney, Tom Hanks, and Meryl Streep back new 'Human Consent Standard' for AI licensing

The Verge AI

50%

Hollywood actors and producers are standing behind a new AI licensing standard that will tell AI systems whether they'll need to pay to use a person's likeness, creative work, characters, and designs. With the Human Consent Standard, people can set terms for the use of their...

<https://www.theverge.com/ai-artificial-intelligence/928534/rsl-media-human-consent-stan...>
- 3

CONF

DeepRefine: Agent-Compiled Knowledge Refinement via Reinforcement Learning

ArXiv CS.AI

50%

arXiv:2605.10488v1 Announce Type: cross Abstract: Agent-compiled knowledge bases provide persistent external knowledge for large language model (LLM) agents in open-ended, knowledge-intensive downstream tasks. Yet their quality is systematically limited by \emph{incompleteness},...

<https://arxiv.org/abs/2605.10488>
- 4

CONF

The Echo Amplifies the Knowledge: Somatic Marker Analogues in Language Models via Emotion Vector Re-Injection

ArXiv CS.AI

50%

arXiv:2605.08611v1 Announce Type: new Abstract: Current language model memory systems store what happened but not how it felt. This distinction -- between semantic memory (knowing about a past event) and episodic memory (re-experiencing it) -- was identified by Tulving as the...

<https://arxiv.org/abs/2605.08611>
- 5

CONF

How Much is Brain Data Worth for Machine Learning?

ArXiv CS.AI

50%

arXiv:2605.09243v1 Announce Type: new Abstract: If a person can solve a task, can measuring their brain make it easier to train a model to solve that task too? Recent NeuroAI work suggests that supplementing task training with neural recordings can modestly improve model...

<https://arxiv.org/abs/2605.09243>
- 6

CONF

Control-Augmented Autoregressive Diffusion for Data Assimilation

ArXiv CS.AI

50%

arXiv:2510.06637v3 Announce Type: replace-cross Abstract: Despite advances in test-time scaling and diffusion finetuning, guidance for Auto-Regressive Diffusion Models (ARDMs) remains underexplored. We introduce an amortized framework that augments a pretrained ARDM with an...

<https://arxiv.org/abs/2510.06637>
- 7

CONF

The Pok\emon Theorem and other Fairness Impossibility Results

ArXiv CS.AI

50%

arXiv:2605.09221v1 Announce Type: cross Abstract: Fairness impossibility results often look like distinct scalar incompatibility statements. We show that several share one RKHS geometry: fairness criteria are linear constraints on conditional mean embeddings, and unequal base...

<https://arxiv.org/abs/2605.09221>

8

CONF

LEVI: Stronger Search Architectures Can Substitute for Larger LLMs in Evolutionary Search

ArXiv CS.AI 50%

arXiv:2605.09764v1 Announce Type: cross Abstract: LLM-guided evolutionary methods such as AlphaEvolve have proven effective in domains like math, systems research, and algorithmic discovery, but their reliance on frontier models makes each run expensive. We argue this is largely...

<https://arxiv.org/abs/2605.09764>

9

CONF

SoK: A Systematic Bidirectional Literature Review of AI & DLT Convergence

ArXiv CS.AI 50%

arXiv:2605.10515v1 Announce Type: cross Abstract: The integration of Artificial Intelligence (AI) with Distributed Ledger Technology (DLT) has become a growing research area, yet contributions tend to cluster around specific application domains or examine only one direction of...

<https://arxiv.org/abs/2605.10515>

0

CONF

AdaRubric: Task-Adaptive Rubrics for Reliable LLM Agent Evaluation and Reward Learning

ArXiv CS.AI 50%

arXiv:2603.21362v3 Announce Type: replace Abstract: Evaluating LLM agent trajectories is fundamentally task-specific: a code-debugging agent should be judged on Correctness and Error Handling, not on Fluency or Safety. Yet the dominant paradigm -- LLM-as-Judge with a fixed...

<https://arxiv.org/abs/2603.21362>

1

CONF

OASIS: A Multilingual and Multimodal Dataset for Culturally Grounded Spoken Visual QA

ArXiv CS.AI 50%

arXiv:2510.06371v3 Announce Type: replace-cross Abstract: Large-scale multimodal models achieve strong results on tasks like Visual Question Answering (VQA), but they are often limited when queries require cultural and visual information, everyday knowledge, particularly in...

<https://arxiv.org/abs/2510.06371>