



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	146	Media
TechCrunch AI	9	Media
HackerNews	6	Community
MIT Tech Review	4	Media
The Verge AI	3	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

OpenAI's Codex is now in the ChatGPT mobile app

The Verge AI

+2 sources

100%

OpenAI is going to let users access Codex, its desktop AI tool that can write code and use apps on your computer, from the ChatGPT app on your phone. Following the surge in popularity for Anthropic's Claude Code, OpenAI has been working quickly to try and catch up, including by...

<https://www.theverge.com/ai-artificial-intelligence/930763/openai-codex-chatgpt-ios-and...>

VERI

Holistic Evaluation and Failure Diagnosis of AI Agents

ArXiv CS.AI

+1 source

80%

arXiv:2605.14865v1 Announce Type: new Abstract: AI agents execute complex multi-step processes, but current evaluation falls short: outcome metrics report success or failure without explaining why, and process-level approaches struggle to connect failure types to their precise...

<https://arxiv.org/abs/2605.14865>

VERI

Data readiness for agentic AI in financial services

MIT Tech Review

+1 source

80%

Financial services companies have unique needs when it comes to business AI. They operate in one of the most highly regulated sectors while responding to external events that are updated by the second. As a result, the success of agentic AI in financial services depends less on...

<https://www.technologyreview.com/2026/05/14/1137034/data-readiness-for-agentic-ai-in-fi...>

VERI

GhostCite: A Large-Scale Analysis of Citation Validity in the Age of Large Language Models

ArXiv CS.AI

+1 source

80%

arXiv:2602.06718v2 Announce Type: replace-cross Abstract: Citations provide the basis for trusting scientific claims; when they are invalid or fabricated, this trust collapses. With the advent of Large Language Models (LLMs), this risk has intensified: LLMs are increasingly used...

<https://arxiv.org/abs/2602.06718>

CONFIRMED SIGNALS — 159 stories

CONF

Show HN: Torrix, self hosted, LLM Observability,(no Postgres, no Redis)

HackerNews

50%

HN 45

I work as a SAP Integration consultant and built this as a side project. Friction point: Most self hosted LLM observability tools require Postgres, Redis and non trivial infrastructure. Teams just want to see what their agents are actually doing in Production, that set up cost...

<https://github.com/torrix-ai/install>

CONF

Show HN: Claude-pee: use Claude -p without the programmatic usage credit pool

HackerNews

50%

HN 6

Anthropic announced today that starting June 15, paid Claude plans get a separate monthly credit pool for programmatic usage (claude -p, Agent SDK). Seems OK at first glance but it turns out the monthly credit pool is charged at API rates, which effectively kills any serious...

<https://github.com/sbhattach/claude-pee/tree/main>

CONF

Show HN: Halgorithm - Catching AI Hallucinations Using Trees, No AI in Pipeline

HackerNews 50% HN 5

No summary — see link for full article.

<https://github.com/TangibleResearch/Halgorithm>

CONF

Show HN: Parse LLM Markdown streams incrementally on the server or client

HackerNews 50% HN 3

Most AI chat applications (such as ChatGPT or Claude) stream their responses to the client as markdown text. As each new chunk of text arrives, the front end typically re-parses the entire markdown document to render the updated message. This works, but it can quickly slow down...

<https://github.com/nimeshnayaju/markdown-parser>

CONF

Show HN: 3D-Agent - AI that edits Blender scenes through the Python API

HackerNews 50% HN 3

No summary — see link for full article.

<https://3d-agent.com/blender-ai/plugin>

0 CONF

Show HN: Visualizing Tiny LLMs from OpenAI's Parameter Golf

HackerNews 50% HN 2

The two from parameter golf (one I trained, one was the baseline) are just 16MB each! They produce barely plausible English

<https://leebutterman.com/2026/05/01/visualizing-tiny-llms-in-parameter-golf.html>

1 CONF

Elon Musk's SpaceXAI has been bleeding staff since its merger

TechCrunch AI 50%

More than 50 employees have reportedly left Elon Musk's newly merged SpaceXAI since February, raising questions about burnout, leadership changes, talent poaching, and whether liquidity events weakened retention incentives.

[https://techcrunch.com/2026/05/14/elon-musks-spacexai-has-been-bleeding-staff-since-its...](https://techcrunch.com/2026/05/14/elon-musks-spacexai-has-been-bleeding-staff-since-its-...)

2 CONF

What happens when AI starts building itself?

TechCrunch AI 50%

Richard Socher's new \$650 million startup wants to build an AI that can research and improve itself indefinitely — and he insists it will actually ship products.

<https://techcrunch.com/2026/05/14/what-happens-when-ai-starts-building-itself/>

3 CONF

OpenAI is reportedly preparing legal action against Apple; it wouldn't be the first partner to feel burned

TechCrunch AI 50%

OpenAI is so frustrated with Apple over a ChatGPT integration that failed to deliver the subscribers and prominence it expected that the company is now actively exploring legal action against the iPhone maker.

[https://techcrunch.com/2026/05/14/openai-is-reportedly-preparing-legal-action-against-a...](https://techcrunch.com/2026/05/14/openai-is-reportedly-preparing-legal-action-against-a-...)

4 CONF

ClawdMeter turns your Claude Code usage stats into a tiny desktop dashboard

TechCrunch AI 50%

A new open source gadget called ClawdMeter turns Claude Code usage stats into a tiny desktop dashboard for AI coding power users.

[https://techcrunch.com/2026/05/14/clawd-meter-turns-your-claude-code-usage-stats-into-a...](https://techcrunch.com/2026/05/14/clawd-meter-turns-your-claude-code-usage-stats-into-a-...)

5 CONF

Cerebras raises \$5.5B, then stock pops \$108%, in the first huge tech IPO of 2026

TechCrunch AI 50%

A year ago, it looked like this day would never happen for Cerebras.

<https://techcrunch.com/2026/05/14/cerebras-raises-5-5b-kicking-off-2026s-ipo-season-wit...>

6 CONF

Cisco cuts nearly 4,000 jobs to spend more on AI, reports 'record quarterly revenue'

TechCrunch AI 50%

This is Cisco's latest layoff in recent years, while the company's chief executive touts record revenue and growth.

<https://techcrunch.com/2026/05/14/cisco-cuts-nearly-4000-jobs-to-spend-more-on-ai-repor...>

7 CONF

Wirestock raises \$23M to supply creative multimodal data to AI labs

TechCrunch AI 50%

The company pivoted to being a data provider in 2023 and now supplies datasets of images, videos, design assets, and gaming and 3D content to AI labs.

<https://techcrunch.com/2026/05/14/wirestock-raises-23m-to-supply-multi-modal-data-to-ai...>

8 CONF

The Download: deepfake porn's stolen bodies and AI sharing private numbers

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. The shock of seeing your body used in deepfake porn When Jennifer got a research job in 2023, she ran her new professional headshot through a...

<https://www.technologyreview.com/2026/05/14/1137257/the-download-deepfake-porn-bodies-a...>

9 CONF

Microsoft starts canceling Claude Code licenses

The Verge AI 50%

Microsoft first started opening up access to Claude Code in December, inviting thousands of its own developers to use Anthropic's AI coding tool daily. It was part of an effort to get project managers, designers, and other employees to experiment with coding for the first time,...

<https://www.theverge.com/tech/930447/microsoft-claude-code-discontinued-notepad>

0 CONF

Americans do not want AI data centers in their backyards

The Verge AI 50%

Over 70 percent of Americans oppose AI data center construction in their area, according to a new Gallup survey. Just 7 percent said they were "strongly" in favor of new data centers. According to Gallup, data centers are so strongly disliked that Americans would prefer to live...

<https://www.theverge.com/ai-artificial-intelligence/930477/ai-data-centers-gallup-surve...>

1 CONF

A Two-Dimensional Framework for AI Agent Design Patterns: Cognitive Function and Execution Topology

ArXiv CS.AI 50%

arXiv:2605.13850v1 Announce Type: new Abstract: Existing frameworks for LLM-based agent architectures describe systems from a single perspective: industry guides (Anthropic, Google, LangChain) focus on execution topology -- how data flows -- while cognitive science surveys focus...

<https://arxiv.org/abs/2605.13850>

2 CONF

Invisible Orchestrators Suppress Protective Behavior and Dissociate Power-Holders: Safety Risks in Multi-Agent LLM Systems

ArXiv CS.AI 50%

arXiv:2605.13851v1 Announce Type: new Abstract: Multi-agent orchestration -- in which a hidden coordinator manages specialized worker agents -- is becoming the default architecture for enterprise AI deployment, yet the safety implications of orchestrator invisibility have never...

<https://arxiv.org/abs/2605.13851>

3 CONF

NeuroAtlas: Benchmarking Foundation Models for Clinical EEG and Brain-Computer Interfaces

ArXiv CS.AI 50%

arXiv:2605.14698v1 Announce Type: cross Abstract: Foundation models (FM) promise to extract unified representations that generalize across downstream tasks. They have emerged across fields, including electroencephalography (EEG), but it is less clear how effective they are in...

<https://arxiv.org/abs/2605.14698>

4 CONF

Sheaf-Theoretic Transport and Obstruction for Detecting Scientific Theory Shift in AI Agents

ArXiv CS.AI 50%

arXiv:2605.14033v1 Announce Type: new Abstract: Scientific theory shift in AI agents requires more than fitting equations to data. An artificial scientific agent must detect whether an existing representational framework remains transportable into a new regime, or whether its...

<https://arxiv.org/abs/2605.14033>

5 CONF

From Descriptive to Prescriptive: Uncover the Social Value Alignment of LLM-based Agents

ArXiv CS.AI 50%

arXiv:2605.14034v1 Announce Type: new Abstract: Wide applications of LLM-based agents require strong alignment with human social values. However, current works still exhibit deficiencies in self-cognition and dilemma decision, as well as self-emotions. To remedy this, we propose...

<https://arxiv.org/abs/2605.14034>

6 CONF

Case-Based Calibration of Adaptive Reasoning and Execution for LLM Tool Use

ArXiv CS.AI 50%

arXiv:2605.15041v1 Announce Type: new Abstract: Tool use extends large language models beyond parametric knowledge, but reliable execution requires balancing appropriate reasoning depth with strict structural validity. We approach this problem from a case-based perspective to...

<https://arxiv.org/abs/2605.15041>

7 CONF

Network-Aware Bilinear Tokenization for Brain Functional Connectivity Representation Learning

ArXiv CS.AI 50%

arXiv:2605.14048v1 Announce Type: new Abstract: Masked autoencoders (MAEs) have recently shown promise for self-supervised representation learning of resting-state brain functional connectivity (FC). However, a fundamental question remains unresolved: how should FC matrices be...

<https://arxiv.org/abs/2605.14048>

8 CONF

Bridging Legal Interpretation and Formal Logic: Faithfulness, Assumption, and the Future of AI Legal Reasoning

ArXiv CS.AI 50%

arXiv:2605.14049v1 Announce Type: new Abstract: The growing adoption of large language models in legal practice brings both significant promise and serious risk. Legal professionals stand to benefit from AI that can reason over contracts, draft documents, and analyze sources at...

<https://arxiv.org/abs/2605.14049>

9 CONF

SPIN: Structural LLM Planning via Iterative Navigation for Industrial Tasks

ArXiv CS.AI 50%

arXiv:2605.14051v1 Announce Type: new Abstract: Industrial LLM agent systems often separate planning from execution, yet LLM planners frequently produce structurally invalid or unnecessarily long workflows, leading to brittle failures and avoidable tool and API cost. We propose...

<https://arxiv.org/abs/2605.14051>

0 CONF

MathAtlas: A Benchmark for Autoformalization in the Wild

ArXiv CS.AI 50%

arXiv:2605.14061v1 Announce Type: new Abstract: Current autoformalization benchmarks are largely focused on olympiad or undergraduate mathematics, while graduate and research-level mathematics remains underexplored. In this paper, we introduce MathAtlas, the first large-scale...

<https://arxiv.org/abs/2605.14061>

1 CONF

Know When To Fold 'Em: Token-Efficient LLM Synthetic Data Generation via Multi-Stage In-Flight Rejection

ArXiv CS.AI 50%

arXiv:2605.14062v1 Announce Type: new Abstract: While synthetic data generation with large language models (LLMs) is widely used in post-training pipelines, existing approaches typically generate full outputs before applying quality filters, leading to substantial token waste on...

<https://arxiv.org/abs/2605.14062>

2 CONF

CRANE: Constrained Reasoning Injection for Code Agents via Nullspace Editing

ArXiv CS.AI 50%

arXiv:2605.14084v1 Announce Type: cross Abstract: Code agents must both reason over long-horizon repository state and obey strict tool-use protocols. In paired Instruct/Thinking checkpoints, these capabilities are complementary but misaligned. The Instruct model is concise and...

<https://arxiv.org/abs/2605.14084>

3 CONF

Optimizing PyTorch Inference with LLM-Based Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2511.16964v2 Announce Type: replace-cross Abstract: Maximizing performance on available GPU hardware is an ongoing challenge for modern AI inference systems. Traditional approaches include writing custom GPU kernels and using specialized model compilers to tune high-level...

<https://arxiv.org/abs/2511.16964>

4 **CONF****Grounded Continuation: A Linear-Time Runtime Verifier for LLM Conversations****ArXiv CS.AI** **50%**

arXiv:2605.14175v1 Announce Type: new Abstract: In long conversations, an LLM can produce a next utterance that sounds plausible but rests on premises the conversation has already abandoned. Context-manipulation attacks against deployed agents now actively exploit this gap. We...

<https://arxiv.org/abs/2605.14175>

5 **CONF****Proximal Action Replacement for Behavior Cloning Actor-Critic in Offline Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2602.07441v2 Announce Type: replace-cross Abstract: Offline reinforcement learning (RL), which optimizes policies using a previously collected static dataset, is an important branch of RL. A popular and promising approach is to regularize actor-critic methods with behavior...

<https://arxiv.org/abs/2602.07441>

6 **CONF****GenCircuit-RL: Reinforcement Learning from Hierarchical Verification for Genetic Circuit Design****ArXiv CS.AI** **50%**

arXiv:2605.14215v1 Announce Type: new Abstract: Genetic circuit design remains a laborious, expert-driven process despite decades of progress in synthetic biology. We study this problem through code generation: models produce Python code in pysbol3 to construct genetic circuits...

<https://arxiv.org/abs/2605.14215>

7 **CONF****Fusion-fission forecasts when AI will shift to undesirable behavior****ArXiv CS.AI** **50%**

arXiv:2605.14218v1 Announce Type: new Abstract: The key problem facing ChatGPT-like AI's use across society is that its behavior can shift, unnoticed, from desirable to undesirable -- encouraging self-harm, extremist acts, financial losses, or costly medical and military...

<https://arxiv.org/abs/2605.14218>

8 **CONF****Asymmetric Generative Recommendation via Multi-Expert Projection and Multi-Faceted Hierarchical Quantization****ArXiv CS.AI** **50%**

arXiv:2605.14512v1 Announce Type: cross Abstract: Generative Recommendation (GenRec) models reformulate recommendation as a sequence generation task, representing items as discrete Semantic IDs used symmetrically as both inputs and prediction targets. We identify a critical...

<https://arxiv.org/abs/2605.14512>

9 **CONF****Agentic AI Ecosystems in Higher Education: A Perspective on AI Agents to Emerging Inclusive, Agentic Multi-Agent AI Framework for Learning, Teaching and Institutional Intelligence****ArXiv CS.AI** **50%**

arXiv:2605.14266v1 Announce Type: new Abstract: Integration of artificial intelligent (AI) agents in higher education is transforming teaching, learning and administrative processes. Although existing AI agents effectively support individual tasks, their implementation remains...

<https://arxiv.org/abs/2605.14266>

0 **CONF****Precise Verification of Transformers through ReLU-Catalyzed Abstraction Refinement**

ArXiv CS.AI 50%

arXiv:2605.14294v1 Announce Type: new Abstract: Formal verification of transformers has become increasingly important due to their widespread deployment in safety-critical applications. Compared to classic neural networks, the inferences of transformers involve highly complex...

<https://arxiv.org/abs/2605.14294>

1 **CONF****Semantic Feature Segmentation for Interpretable Predictive Maintenance in Complex Systems**

ArXiv CS.AI 50%

arXiv:2605.14318v1 Announce Type: new Abstract: Predictive maintenance in complex systems is often complicated by the heterogeneity and redundancy of monitored variables, which can obscure fault-relevant information and reduce model interpretability. This work proposes a semantic...

<https://arxiv.org/abs/2605.14318>

2 **CONF****Metis AI: The Overlooked Middle Zone Between AI-Native and World-Movers**

ArXiv CS.AI 50%

arXiv:2605.14407v1 Announce Type: new Abstract: The dominant discourse on AI limitations frames the boundary of AI capability as a divide between digital tasks (where AI excels) and physical tasks (where embodiment is required). We argue this framing misses the most...

<https://arxiv.org/abs/2605.14407>

3 **CONF****DVMap: Fine-Grained Pluralistic Value Alignment via High-Consensus Demographic-Value Mapping**

ArXiv CS.AI 50%

arXiv:2605.14420v1 Announce Type: new Abstract: Current Large Language Models (LLMs) typically rely on coarse-grained national labels for pluralistic value alignment. However, such macro-level supervision often obscures intra-country value heterogeneity, yielding a loose...

<https://arxiv.org/abs/2605.14420>

4 **CONF****Prompting Policies for Multi-step Reasoning and Tool-Use in Black-box LLMs with Iterative Distillation of Experience**

ArXiv CS.AI 50%

arXiv:2605.14443v1 Announce Type: new Abstract: The shift toward interacting with frozen, "black-box" Large Language Models (LLMs) has transformed prompt engineering from a heuristic exercise into a critical optimization challenge. We propose a Reinforcement Learning (RL)...

<https://arxiv.org/abs/2605.14443>

5 **CONF****A plug-and-play generative framework for multi-satellite precipitation estimation**

ArXiv CS.AI 50%

arXiv:2605.14426v1 Announce Type: cross Abstract: Reliable precipitation monitoring is essential for disaster risk reduction, water resources management, and agricultural decision-making. Multi-source satellite observations, particularly the combination of geostationary infrared...

<https://arxiv.org/abs/2605.14426>

6 CONF

OmniDrop: Layer-wise Token Pruning for Omni-modal LLMs via Query-Guidance

ArXiv CS.AI 50%

arXiv:2605.14458v1 Announce Type: new Abstract: Omni-modal large language models have demonstrated remarkable potential in holistic multimodal understanding; however, the token explosion caused by high-resolution audio and video inputs remains a critical bottleneck for real-time...

<https://arxiv.org/abs/2605.14458>

7 CONF

Why Retrieval-Augmented Generation Fails: A Graph Perspective

ArXiv CS.AI 50%

arXiv:2605.14192v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) has become a powerful and widely used approach for improving large language models by grounding generation in retrieved evidence. However, RAG systems still produce incorrect answers in many...

<https://arxiv.org/abs/2605.14192>

8 CONF

Learning Scenario Reduction for Two-Stage Robust Optimization with Discrete Uncertainty

ArXiv CS.AI 50%

arXiv:2605.14494v1 Announce Type: new Abstract: Two-Stage Robust Optimization (2RO) with discrete uncertainty is challenging, often rendering exact solutions prohibitive. Scenario reduction alleviates this issue by selecting a small, representative subset of scenarios to enable...

<https://arxiv.org/abs/2605.14494>

9 CONF

When Robots Do the Chores: A Benchmark and Agent for Long-Horizon Household Task Execution

ArXiv CS.AI 50%

arXiv:2605.14504v1 Announce Type: new Abstract: Long-horizon household tasks demand robust high-level planning and sustained reasoning capabilities, which are largely overlooked by existing embodied AI benchmarks that emphasize short-horizon navigation or manipulation and rely...

<https://arxiv.org/abs/2605.14504>

0 CONF

ExploitBench: A Capability Ladder Benchmark for LLM Cybersecurity Agents

ArXiv CS.AI 50%

arXiv:2605.14153v1 Announce Type: cross Abstract: Exploitation is not a binary event. It is a ladder of acquiring progressive capabilities, from executing a single buggy line of code to taking full control of the target. However, existing LLM security benchmarks treat a crash as...

<https://arxiv.org/abs/2605.14153>

1 CONF

Complacent, Not Sycophantic: Reframing Large Language Models and Designing AI Literacy for Complacent Machines

ArXiv CS.AI 50%

arXiv:2605.14544v1 Announce Type: new Abstract: Large language models are often described as sycophantic, in the sense that they appear to flatter users or mirror their beliefs. We argue that this label is conceptually misleading: sycophancy implies motives and strategic intent,...

<https://arxiv.org/abs/2605.14544>

2 **CONF****TERMS-Bench: Diagnosing LLM Negotiation Agents Beyond Deal Rate**

ArXiv CS.AI 50%

arXiv:2605.13909v1 Announce Type: cross Abstract: Negotiation is a central mechanism of economic exchange, shaping markets, procurement, labor agreements, and resource allocation. It is also a canonical testbed for agentic language models, requiring multi-turn interaction under...

<https://arxiv.org/abs/2605.13909>
3 **CONF****Prompt Segmentation and Annotation Optimisation: Controlling LLM Behaviour via Optimised Segment-Level Annotations**

ArXiv CS.AI 50%

arXiv:2605.14561v1 Announce Type: new Abstract: Prompt engineering is crucial for effective interaction with generative artificial intelligence systems, yet existing optimisation methods often operate over an unstructured and vast prompt space, leading to high computational...

<https://arxiv.org/abs/2605.14561>
4 **CONF****Sycophancy is an Educational Safety Risk: Why LLM Tutors Need Sycophancy Benchmarks**

ArXiv CS.AI 50%

arXiv:2605.14604v1 Announce Type: new Abstract: This position paper argues that effective tutoring requires corrective friction: surfacing misconceptions and challenging them supportively to drive conceptual change. Yet preference-aligned LLMs can trade epistemic rigor for...

<https://arxiv.org/abs/2605.14604>
5 **CONF****COTCAgent: Preventive Consultation via Probabilistic Chain-of-Thought Completion**

ArXiv CS.AI 50%

arXiv:2605.15016v1 Announce Type: cross Abstract: As large language models empower healthcare, intelligent clinical decision support has developed rapidly. Longitudinal electronic health records (EHR) provide essential temporal evidence for accurate clinical diagnosis and...

<https://arxiv.org/abs/2605.15016>
6 **CONF****MindGap: A Conversational AI Framework for Upstream Neuroplastic Intervention in Post-Traumatic Stress Disorder**

ArXiv CS.AI 50%

arXiv:2605.14660v1 Announce Type: new Abstract: Post-Traumatic Stress Disorder (PTSD) is fundamentally a neuroplastic problem traumatic contact events encode over-reactive neural pathways through Hebbian long-term potentiation, producing hair-triggered amygdala-HPA stress...

<https://arxiv.org/abs/2605.14660>
7 **CONF****Falkor-IRAC: Graph-Constrained Generation for Verified Legal Reasoning in Indian Judicial AI**

ArXiv CS.AI 50%

arXiv:2605.14665v1 Announce Type: new Abstract: Legal reasoning is not semantic similarity search. A court judgment encodes constrained symbolic reasoning: precedent propagation, procedural state transitions, and statute-bound inference. These are properties that vector-based...

<https://arxiv.org/abs/2605.14665>

8 CONF

How Sensitive Are Radiomic AI Models to Acquisition Parameters?

ArXiv CS.AI 50%

arXiv:2605.14667v1 Announce Type: new Abstract: A main barrier for the deployment of AI radiomic systems in clinical routine is their drop in performance under heterogeneous multicentre acquisition protocols. This work presents a performance-oriented framework for quantifying...

<https://arxiv.org/abs/2605.14667>

9 CONF

Agentifying Patient Dynamics within LLMs through Interacting with Clinical World Model

ArXiv CS.AI 50%

arXiv:2605.14723v1 Announce Type: new Abstract: Sepsis management in the ICU requires sequential treatment decisions under rapidly evolving patient physiology. Although large language models (LLMs) encode broad clinical knowledge and can reason over guidelines, they are not...

<https://arxiv.org/abs/2605.14723>

0 CONF

Does RAG Know When Retrieval Is Wrong? Diagnosing Context Compliance under Knowledge Conflict

ArXiv CS.AI 50%

arXiv:2605.14473v1 Announce Type: cross Abstract: The Context-Compliance Regime in Retrieval-Augmented Generation (RAG) occurs when retrieved context dominates the final answer even when it conflicts with the model's parametric knowledge. Accuracy alone does not reveal how...

<https://arxiv.org/abs/2605.14473>

1 CONF

Self-Distilled Agentic Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.15155v1 Announce Type: cross Abstract: Reinforcement learning (RL) has emerged as a central paradigm for post-training LLM agents, yet its trajectory-level reward signal provides only coarse supervision for long-horizon interaction. On-Policy Self-Distillation (OPSD)...

<https://arxiv.org/abs/2605.15155>

2 CONF

AI Outperforms Humans in Personalized Image Aesthetics Assessment via LLM-Based Interviews and Semantic Feature Extraction

ArXiv CS.AI 50%

arXiv:2605.14761v1 Announce Type: new Abstract: Accurately predicting individual aesthetic evaluation for images is a fundamental challenge for AI. Various deep learning (DL)-based models have been proposed for this task, training on image evaluation data to extract objective...

<https://arxiv.org/abs/2605.14761>

3 CONF

Identifying Culprits Through Deep Deterministic Policy Gradient Deep Learning Investigation

ArXiv CS.AI 50%

arXiv:2605.14774v1 Announce Type: new Abstract: In the world of AI and advanced technologies investigation aspects identification of a crime or criminal plays a major problem. In this research we focus on a Conventional ways of implicating criminal investigations usually rely on...

<https://arxiv.org/abs/2605.14774>

- 4 **CONF**
- A Heterogeneous Temporal Memory Governance Framework for Long-Term LLM Persona Consistency**
- ArXiv CS.AI 50%
- arXiv:2605.14802v1 Announce Type: new Abstract: Large language models often suffer from fact loss, timeline confusion, persona drift, and reduced stability during long-range interaction, especially under high-noise knowledge bases, context clearing, and cross-model transfer. To...
- <https://arxiv.org/abs/2605.14802>
-
- 5 **CONF**
- BiFedKD: Bidirectional Federated Knowledge Distillation Framework for Non-IID and Long-Tailed ECG Monitoring**
- ArXiv CS.AI 50%
- arXiv:2605.14886v1 Announce Type: new Abstract: Electrocardiogram (ECG) monitoring in Internet of Medical Things (IoMT) networks is constrained by strict data-sharing regulations and privacy concerns. Federated learning (FL) enables collaborative learning by keeping raw ECG data...
- <https://arxiv.org/abs/2605.14886>
-
- 6 **CONF**
- Beyond Individual Intelligence: Surveying Collaboration, Failure Attribution, and Self-Evolution in LLM-based Multi-Agent Systems**
- ArXiv CS.AI 50%
- arXiv:2605.14892v1 Announce Type: new Abstract: LLM-based autonomous agents have demonstrated strong capabilities in reasoning, planning, and tool use, yet remain limited when tasks require sustained coordination across roles, tools, and environments. Multi-agent systems address...
- <https://arxiv.org/abs/2605.14892>
-
- 7 **CONF**
- From Sycophantic Consensus to Pluralistic Repair: Why AI Alignment Must Surface Disagreement**
- ArXiv CS.AI 50%
- arXiv:2605.14912v1 Announce Type: new Abstract: Pluralistic alignment is typically operationalised as preference aggregation: producing responses that span (Overton), steer toward (Steerable), or proportionally represent (Distributional) diverse human values. We argue that...
- <https://arxiv.org/abs/2605.14912>
-
- 8 **CONF**
- GraphFlow: An Architecture for Formally Verifiable Visual Workflows Enabling Reliable Agentic AI Automation**
- ArXiv CS.AI 50%
- arXiv:2605.14968v1 Announce Type: new Abstract: GraphFlow is a visual workflow system designed to improve the reliability of agentic AI automation in multi-step, mission-critical processes. In these workflows, small errors compound rapidly: under an idealized model of...
- <https://arxiv.org/abs/2605.14968>
-
- 9 **CONF**
- Explainable Detection of Depression Status Shifts from User Digital Traces**
- ArXiv CS.AI 50%
- arXiv:2605.14995v1 Announce Type: new Abstract: Every day, users generate digital traces (e.g., social media posts, chats, and online interactions) that are inherently timestamped and may reflect aspects of their mental state. These traces can be organized into temporal...
- <https://arxiv.org/abs/2605.14995>

0 CONF

Addressing Terminal Constraints in Data-Driven Demand Response Scheduling

ArXiv CS.AI 50%

arXiv:2605.14741v1 Announce Type: cross Abstract: Electrified chemical processes are incentivized by exposure to time-varying electricity markets to operate flexibly, but participating in demand response schemes can require satisfying terminal constraints over long horizons....

<https://arxiv.org/abs/2605.14741>

1 CONF

Why Neighborhoods Matter: Traversal Context and Provenance in Agentic GraphRAG

ArXiv CS.AI 50%

arXiv:2605.15109v1 Announce Type: new Abstract: Retrieval-Augmented Generation can improve factuality by grounding answers in external evidence, but Agentic GraphRAG complicates what it means for citations to be faithful. In these systems, an agent explores a knowledge graph...

<https://arxiv.org/abs/2605.15109>

2 CONF

Query-Conditioned Test-Time Self-Training for Large Language Models

ArXiv CS.AI 50%

arXiv:2605.13369v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are typically deployed with fixed parameters, and their performance is often improved by allocating more computation at inference time. While such test-time scaling can be effective, it cannot...

<https://arxiv.org/abs/2605.13369>

3 CONF

WARD: Adversarially Robust Defense of Web Agents Against Prompt Injections

ArXiv CS.AI 50%

arXiv:2605.15030v1 Announce Type: cross Abstract: Web agents can autonomously complete online tasks by interacting with websites, but their exposure to open web environments makes them vulnerable to prompt injection attacks embedded in HTML content or visual interfaces. Existing...

<https://arxiv.org/abs/2605.15030>

4 CONF

BiSpikCLM: A Spiking Language Model integrating Softmax-Free Spiking Attention and Spike-Aware Alignment Distillation

ArXiv CS.AI 50%

arXiv:2605.13859v1 Announce Type: cross Abstract: Spiking Neural Networks (SNNs) offer promising energy-efficient alternatives to large language models (LLMs) due to their event-driven nature and ultra-low power consumption. However, to preserve capacity, most existing spiking...

<https://arxiv.org/abs/2605.13859>

5 CONF

Elastic Spiking Transformers for Efficient Gesture Understanding

ArXiv CS.AI 50%

arXiv:2605.13869v1 Announce Type: cross Abstract: Spiking Neural Networks (SNNs), particularly Spiking Transformers, offer energy-efficient processing of event-based sensor data for healthcare applications. Yet current architectures are rigid: they are trained and deployed as...

<https://arxiv.org/abs/2605.13869>

6 CONF

S-AI-Recursive: A Bio-Inspired and Temporal Sparse AI Architecture for Iterative, Introspective, and Energy-Frugal Reasoning

ArXiv CS.AI 50%

arXiv:2605.13872v1 Announce Type: cross Abstract: This article introduces S-AI-Recursive, a bio-inspired Sparse Artificial Intelligence architecture in which reasoning is operationalized as a hormonal closed-loop iteration rather than a single feed-forward pass. Building upon...

<https://arxiv.org/abs/2605.13872>

7 CONF

Large Language Models for Web Accessibility: A Systematic Literature Review

ArXiv CS.AI 50%

arXiv:2605.13873v1 Announce Type: cross Abstract: Web accessibility aims to ensure that web content and services are usable by people with diverse abilities. In recent years, Large Language Models (LLMs) have been increasingly explored to support accessibility-related tasks on...

<https://arxiv.org/abs/2605.13873>

8 CONF

ARES-LSHADE: Autoresearch-Enhanced LSHADE with Memetic Polish for the GNBG Benchmark

ArXiv CS.AI 50%

arXiv:2605.13877v1 Announce Type: cross Abstract: We present ARES-LSHADE, a memetic differential-evolution variant submitted to the GECCO 2026 competition on LLM-designed evolutionary algorithms for the Generalized Numerical Benchmark Generator (GNBG). The algorithm builds on...

<https://arxiv.org/abs/2605.13877>

9 CONF

NeuroTrain: Surveying Local Learning Rules for Spiking Neural Networks with an Open Benchmarking Framework

ArXiv CS.AI 50%

arXiv:2605.15058v1 Announce Type: cross Abstract: The rapid expansion of spiking neural networks (SNNs) has led to a proliferation of training algorithms that differ widely in biological inspiration, computational structure, and hardware suitability. Despite this progress, the...

<https://arxiv.org/abs/2605.15058>

10 CONF

A Non-Destructive Methodological Framework for Modernizing Legacy Clinical Reporting Systems for AI-Driven Pharmacoinformatics: A SAS Case Study

ArXiv CS.AI 50%

arXiv:2605.13905v1 Announce Type: cross Abstract: Drug development and pharmacovigilance are frequently bottlenecked by legacy clinical reporting pipelines. These monolithic systems encode regulatory-grade logic but resist AI integration by producing opaque output with no...

<https://arxiv.org/abs/2605.13905>

11 CONF

Multi-Scale Dequant: Eliminating Dequantization Bottleneck via Activation Decomposition for Efficient LLM Inference

ArXiv CS.AI 50%

arXiv:2605.13915v1 Announce Type: cross Abstract: Quantization is essential for efficient large language model (LLM) inference, yet the dequantization step-converting low-bit weights back to high-precision for matrix multiplication has become a critical bottleneck on modern AI...

<https://arxiv.org/abs/2605.13915>

2 **CONF****Towards the Next Frontier of LLMs, Training on Private Data: A Cross-Domain Benchmark for Federated Fine-Tuning****ArXiv CS.AI** **50%**

arXiv:2605.13936v1 Announce Type: cross Abstract: The recent success of large language models (LLMs) has been largely driven by vast public datasets. However, the next frontier for LLM development lies beyond public data. Much of the world's most valuable information is private,...

<https://arxiv.org/abs/2605.13936>

3 **CONF****AgentTrap: Measuring Runtime Trust Failures in Third-Party Agent Skills****ArXiv CS.AI** **50%**

arXiv:2605.13940v1 Announce Type: cross Abstract: Third-party skills are becoming the package ecosystem for LLM agents. They package natural-language instructions, helper scripts, templates, documents, and service configuration into reusable workflows. This makes skills useful,...

<https://arxiv.org/abs/2605.13940>

4 **CONF****EvolveMem: Self-Evolving Memory Architecture via AutoResearch for LLM Agents****ArXiv CS.AI** **50%**

arXiv:2605.13941v1 Announce Type: cross Abstract: Long-term memory is essential for LLM agents that operate across multiple sessions, yet existing memory systems treat retrieval infrastructure as fixed: stored content evolves while scoring functions, fusion strategies, and...

<https://arxiv.org/abs/2605.13941>

5 **CONF****WarmPrior: Straightening Flow-Matching Policies with Temporal Priors****ArXiv CS.AI** **50%**

arXiv:2605.13959v1 Announce Type: cross Abstract: Generative policies based on diffusion and flow matching have become a dominant paradigm for visuomotor robotic control. We show that replacing the standard Gaussian source distribution with WarmPrior, a simple temporally...

<https://arxiv.org/abs/2605.13959>

6 **CONF****Few Channels Draw The Whole Picture: Revealing Massive Activations in Diffusion Transformers****ArXiv CS.AI** **50%**

arXiv:2605.13974v1 Announce Type: cross Abstract: Diffusion Transformers (DiTs) and related flow-based architectures are now among the strongest text-to-image generators, yet the internal mechanisms through which prompts shape image semantics remain poorly understood. In this...

<https://arxiv.org/abs/2605.13974>

7 **CONF****Towards Resource-Efficient LLMs: End-to-End Energy Accounting of Distillation Pipelines****ArXiv CS.AI** **50%**

arXiv:2605.13981v1 Announce Type: cross Abstract: The rise in deployment of large language models has driven a surge in GPU demand and datacenter scaling, raising concerns about electricity use, grid stress, and the impacts of modern AI workloads. Distillation is often promoted...

<https://arxiv.org/abs/2605.13981>

8 CONF

HodgeCover: Higher-Order Topological Coverage Drives Compression of Sparse Mixture-of-Experts

ArXiv CS.AI 50%

arXiv:2605.13997v1 Announce Type: cross Abstract: Sparse Mixture-of-Experts (MoE) layers route tokens through a handful of experts, and learning-free compression of these layers reduces inference cost without retraining. A subtle obstruction blocks every existing compressor in...

<https://arxiv.org/abs/2605.13997>

9 CONF

Measuring Google AI Overviews: Activation, Source Quality, Claim Fidelity, and Publisher Impact

ArXiv CS.AI 50%

arXiv:2605.14021v1 Announce Type: cross Abstract: Google AI Overviews (AIOs) are arguably the most widely encountered deployment of generative AI, reaching over 2 billion users who may not realize the answers they see are AI-generated. Where search engines have traditionally...

<https://arxiv.org/abs/2605.14021>

0 CONF

A Benchmark for Early-stage Parkinson's Disease Detection from Speech

ArXiv CS.AI 50%

arXiv:2605.14066v1 Announce Type: cross Abstract: Early-stage Parkinson's disease (EarlyPD) detection from speech is clinically meaningful yet underexplored, and published results are hard to compare because studies differ in datasets, languages, tasks, evaluation protocols, and...

<https://arxiv.org/abs/2605.14066>

1 CONF

Reinforcement Learning for Tool-Calling Agents in Fast Healthcare Interoperability Resources (FHIR)

ArXiv CS.AI 50%

arXiv:2605.14126v1 Announce Type: cross Abstract: Fast Healthcare Interoperability Resources (FHIR) is the dominant standard for interoperable exchange of healthcare data. In FHIR, electronic health records form a directed graph of resources. Answering clinically meaningful...

<https://arxiv.org/abs/2605.14126>

2 CONF

Thinking Ahead: Prospection-Guided Retrieval of Memory with Language Models

ArXiv CS.AI 50%

arXiv:2605.14177v1 Announce Type: cross Abstract: Long-horizon personalization requires dialogue assistants to retrieve user-specific facts from extended interaction histories. In practice, many relevant facts often have low semantic similarity to the query under dense retrieval....

<https://arxiv.org/abs/2605.14177>

3 CONF

LLM-Based Robustness Testing of Microservice Applications: An Empirical Study

ArXiv CS.AI 50%

arXiv:2605.14202v1 Announce Type: cross Abstract: Malformed, missing, or boundary-value inputs in microservice APIs can cascade across dependent services, threatening reliability. Robustness testing systematically exercises such inputs to expose server-side failures, but...

<https://arxiv.org/abs/2605.14202>

4 **CONF****Towards Fine-Grained and Verifiable Concept Bottleneck Models****ArXiv CS.AI** **50%**

arXiv:2605.14210v1 Announce Type: cross Abstract: Concept Bottleneck Models (CBMs) offer interpretable alternatives to black-box predictors by introducing human-relatable concepts before the final output. However, existing CBMs struggle to verify whether predicted concepts...

<https://arxiv.org/abs/2605.14210>

5 **CONF****RQ-MoE: Residual Quantization via Mixture of Experts for Efficient Input-Dependent Vector Compression****ArXiv CS.AI** **50%**

arXiv:2605.14359v1 Announce Type: cross Abstract: Vector quantization is a fundamental tool for compressing high-dimensional embeddings, yet existing multi-codebook methods rely on static codebooks that limit expressiveness under heterogeneous data geometry. While recent dynamic...

<https://arxiv.org/abs/2605.14359>

6 **CONF****Artificial Intelligence-Assistant Cardiocography: Unified Model for Signal Reconstruction, Fetal Heart Rate Analysis, and Variability Assessment****ArXiv CS.AI** **50%**

arXiv:2605.14242v1 Announce Type: cross Abstract: The monitoring of fetal heart rate (FHR) and the assessment of its variability are crucial for preventing fetal compromise and adverse outcomes. However, traditional methods encounter limitations arising from equipment...

<https://arxiv.org/abs/2605.14242>

7 **CONF****Dynamics of the Transformer Residual Stream: Coupling Spectral Geometry to Network Topology****ArXiv CS.AI** **50%**

arXiv:2605.14258v1 Announce Type: cross Abstract: Large language models are remarkably capable, yet how computation propagates through their layers remains poorly understood. A growing line of work treats depth as discrete time and the residual stream as a dynamical system,...

<https://arxiv.org/abs/2605.14258>

8 **CONF****Image Restoration via Diffusion Models with Dynamic Resolution****ArXiv CS.AI** **50%**

arXiv:2605.14267v1 Announce Type: cross Abstract: Diffusion models (DMs) have exhibited remarkable efficacy in various image restoration tasks. However, existing approaches typically operate within the high-dimensional pixel space, resulting in high computational overhead. While...

<https://arxiv.org/abs/2605.14267>

9 **CONF****Beyond Binary: Reframing GUI Critique as Continuous Semantic Alignment****ArXiv CS.AI** **50%**

arXiv:2605.14311v1 Announce Type: cross Abstract: Test-Time Scaling (TTS), which samples multiple candidate actions and ranks them via a Critic Model, has emerged as a promising paradigm for generalist GUI agents. Its efficacy thus hinges on the critic's fine-grained ranking...

<https://arxiv.org/abs/2605.14311>

00 CONF

AIM-DDI: A Model-Agnostic Multimodal Integration Module for Drug-Drug Interaction Prediction

ArXiv CS.AI 50%

arXiv:2605.14327v1 Announce Type: cross Abstract: Drug-drug interaction (DDI) prediction is a critical task in computational biomedicine, as adverse interactions between co-administered drugs can cause severe side effects and clinical risks. A key challenge is unseen-drug...

<https://arxiv.org/abs/2605.14327>

01 CONF

Analog RF Computing: A New Paradigm for Energy-Efficient Edge AI Over MU-MIMO Systems

ArXiv CS.AI 50%

arXiv:2605.14331v1 Announce Type: cross Abstract: Modern edge devices increasingly rely on neural networks for intelligent applications. However, conventional digital computing-based edge inference requires substantial memory and energy consumption. In analog radio frequency...

<https://arxiv.org/abs/2605.14331>

02 CONF

Correctness-Aware Repository Filtering Under Maximum Effective Context Window Constraints

ArXiv CS.AI 50%

arXiv:2605.14362v1 Announce Type: cross Abstract: Context window efficiency is a practical constraint in large language model (LLM)-based developer tools. Paulsen [12] shows that all tested models degrade in accuracy well before their advertised context limits the Maximum...

<https://arxiv.org/abs/2605.14362>

03 CONF

LoMETab: Beyond Rank-1 Ensembles for Tabular Deep Learning

ArXiv CS.AI 50%

arXiv:2605.14365v1 Announce Type: cross Abstract: Recent tabular learning benchmarks increasingly show a tight performance cluster rather than a clear hierarchy among leading methods, spanning gradient boosted decision trees, attention-based architectures, and implicit ensembles...

<https://arxiv.org/abs/2605.14365>

04 CONF

(How) Do Large Language Models Understand High-Level Message Sequence Charts?

ArXiv CS.AI 50%

arXiv:2605.13773v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are being employed widely to automate tasks across the software development life-cycle. It is, however, unclear whether these tasks are performed consistently with respect to the semantics of...

<https://arxiv.org/abs/2605.13773>

05 CONF

Darwin Family: MRI-Trust-Weighted Evolutionary Merging for Training-Free Scaling of Language-Model Reasoning

ArXiv CS.AI 50%

arXiv:2605.14386v1 Announce Type: cross Abstract: We present Darwin Family, a framework for training-free evolutionary merging of large language models via gradient-free weight-space recombination. We ask whether frontier-level reasoning performance can be improved without...

<https://arxiv.org/abs/2605.14386>

06 CONF

Workspace-Bench 1.0: Benchmarking AI Agents on Workspace Tasks with Large-Scale File Dependencies

ArXiv CS.AI 50%

arXiv:2605.03596v4 Announce Type: replace Abstract: Workspace learning requires AI agents to identify, reason over, exploit, and update explicit and implicit dependencies among heterogeneous files in a worker's workspace, enabling them to complete both routine and advanced tasks...

<https://arxiv.org/abs/2605.03596>

07 CONF

The Great Pretender: A Stochasticity Problem in LLM Jailbreak

ArXiv CS.AI 50%

arXiv:2605.14418v1 Announce Type: cross Abstract: "Oh-Oh, yes, I'm the great pretender. Pretending that I'm doing well. My need is such, I pretend too much..." summarizes the state in the area of jailbreak creation and evaluation. You find this method to generate adversarial...

<https://arxiv.org/abs/2605.14418>

08 CONF

MemLineage: Lineage-Guided Enforcement for LLM Agent Memory

ArXiv CS.AI 50%

arXiv:2605.14421v1 Announce Type: cross Abstract: We introduce MemLineage, a defense for LLM agent memory that attaches both cryptographic provenance and LLM-mediated derivation lineage to every entry. Recent and concurrent work shows that untrusted content can be written into...

<https://arxiv.org/abs/2605.14421>

09 CONF

Collaborative Yet Personalized Policy Training: Single-Timescale Federated Actor-Critic

ArXiv CS.AI 50%

arXiv:2605.14423v1 Announce Type: cross Abstract: Despite the popularity of the actor-critic method and the practical needs of collaborative policy training, existing works typically either overlook environmental heterogeneity or give up personalization altogether by training a...

<https://arxiv.org/abs/2605.14423>

10 CONF

Efficient Generative Retrieval for E-commerce Search with Semantic Cluster IDs and Expert-Guided RL

ArXiv CS.AI 50%

arXiv:2605.14434v1 Announce Type: cross Abstract: Generative retrieval offers a promising alternative by unifying the fragmented multi-stage retrieval process into a single end-to-end model. However, its practical adoption in industrial e-commerce search remains challenging,...

<https://arxiv.org/abs/2605.14434>

11 CONF

GEAR: Granularity-Adaptive Advantage Reweighting for LLM Agents via Self-Distillation

ArXiv CS.AI 50%

arXiv:2605.11853v2 Announce Type: replace-cross Abstract: Reinforcement learning has become a widely used post-training approach for LLM agents, where training commonly relies on outcome-level rewards that provide only coarse supervision. While finer-grained credit assignment is...

<https://arxiv.org/abs/2605.11853>

12 CONF

Quantifying Cyber-Vulnerability in Power Electronics Systems via an Impedance-Based Attack Reachable Domain

ArXiv CS.AI 50%

arXiv:2605.14502v1 Announce Type: cross Abstract: Power electronics systems are increasingly exposed to cyber threats due to their integration with digital controllers and communication networks. However, an attacker-oriented metric is still lacking to quantify the extent to...

<https://arxiv.org/abs/2605.14502>

13 CONF

HASTE: Training-Free Video Diffusion Acceleration via Head-Wise Adaptive Sparse Attention

ArXiv CS.AI 50%

arXiv:2605.14513v1 Announce Type: cross Abstract: Diffusion-based video generation has advanced substantially in visual fidelity and temporal coherence, but practical deployment remains limited by the quadratic complexity of full attention. Training-free sparse attention is...

<https://arxiv.org/abs/2605.14513>

14 CONF

VLRS-Bench: A Vision-Language Reasoning Benchmark for Remote Sensing

ArXiv CS.AI 50%

arXiv:2602.07045v2 Announce Type: replace-cross Abstract: Recent advancements in Multimodal Large Language Models (MLLMs) have enabled complex reasoning. However, existing remote sensing (RS) benchmarks remain heavily biased toward perception tasks, such as object recognition...

<https://arxiv.org/abs/2602.07045>

15 CONF

One Step to the Side: Why Defenses Against Malicious Finetuning Fail Under Adaptive Adversaries

ArXiv CS.AI 50%

arXiv:2605.14605v1 Announce Type: cross Abstract: Model providers increasingly release open weights or allow users to fine-tune foundation models through APIs. Although these models are safety-aligned before release, their safeguards can often be removed by fine-tuning on...

<https://arxiv.org/abs/2605.14605>

16 CONF

Action-Inspired Generative Models

ArXiv CS.AI 50%

arXiv:2605.14631v1 Announce Type: cross Abstract: We introduce Action-Inspired Generative Models (AGMs), a dual-network generative framework motivated by the observation that existing bridge-matching methods assign uniform regression weight to every stochastic transition in the...

<https://arxiv.org/abs/2605.14631>

17 CONF

MultiEmo-Bench: Multi-label Visual Emotion Analysis for Multi-modal Large Language Models

ArXiv CS.AI 50%

arXiv:2605.14635v1 Announce Type: cross Abstract: This paper introduces a multi-label visual emotion analysis benchmark dataset for comprehensively evaluating the ability of multimodal large language models (MLLMs) to predict the emotions evoked by images. Recent user studies...

<https://arxiv.org/abs/2605.14635>

18 CONF

AI-assisted cultural heritage dissemination: Comparing NMT and glossary-augmented LLM translation in rock art documents

ArXiv CS.AI 50%

arXiv:2605.14679v1 Announce Type: cross Abstract: Cultural heritage institutions increasingly disseminate research and interpretive materials globally, but multilingual dissemination is constrained by limited budgets and staffing. In terminology-dense domains such as rock art,...

<https://arxiv.org/abs/2605.14679>

19 CONF

Mechanical Enforcement for LLM Governance:Evidence of Governance-Task Decoupling in Financial Decision Systems

ArXiv CS.AI 50%

arXiv:2605.14744v1 Announce Type: cross Abstract: Large language models in regulated financial workflows are governed by natural-language policies that the same model interprets, creating a principal--agent failure: outputs can appear compliant without being compliant. Existing...

<https://arxiv.org/abs/2605.14744>

20 CONF

Video2GUI: Synthesizing Large-Scale Interaction Trajectories for Generalized GUI Agent Pretraining

ArXiv CS.AI 50%

arXiv:2605.14747v1 Announce Type: cross Abstract: Recent advances in multimodal large language models have driven growing interest in graphical user interface (GUI) agents, yet their generalization remains constrained by the scarcity of large-scale training data spanning diverse...

<https://arxiv.org/abs/2605.14747>

21 CONF

Streaming Speech-to-Text Translation with a SpeechLLM

ArXiv CS.AI 50%

arXiv:2605.14766v1 Announce Type: cross Abstract: Normally, a system that translates speech into text consists of separate modules for speech recognition and text-to-text translation. Combining those tasks into a SpeechLLM promises to exploit paralinguistic information in the...

<https://arxiv.org/abs/2605.14766>

22 CONF

Beyond GRPO and On-Policy Distillation: An Empirical Sparse-to-Dense Reward Principle for Language-Model Post-Training

ArXiv CS.AI 50%

arXiv:2605.12483v2 Announce Type: replace-cross Abstract: When labeled verifiable training data is scarce, each checked example should be used where it has the most value. A common approach is to train the deployment student model directly with sparse RL methods such as GRPO. We...

<https://arxiv.org/abs/2605.12483>

23 CONF

Beyond AI as Assistants: Toward Autonomous Discovery in Cosmology

ArXiv CS.AI 50%

arXiv:2605.14791v1 Announce Type: cross Abstract: Recent advances in artificial intelligence (AI) agents are pushing AI beyond tools toward autonomous scientific discovery. We discuss two complementary agentic systems for cosmology: \texttt{CMBEvolve}, which targets tasks with...

<https://arxiv.org/abs/2605.14791>

24 **CONF****GPart: End-to-End Isometric Fine-Tuning via Global Parameter Partitioning**

ArXiv CS.AI 50%

arXiv:2605.14841v1 Announce Type: cross Abstract: Low-rank adaptation (LoRA) has become the dominant paradigm for parameter-efficient fine-tuning (PEFT) of large language models (LLMs). However, its bilinear structure introduces a critical limitation: the mapping from trainable...

<https://arxiv.org/abs/2605.14841>
25 **CONF****CUICurate: A GraphRAG-based Framework for Automated Clinical Concept Curation for NLP applications**

ArXiv CS.AI 50%

arXiv:2602.17949v2 Announce Type: replace-cross Abstract: Background: Clinical named entity recognition tools commonly map free text to Unified Medical Language System (UMLS) Concept Unique Identifiers (CUIs). For many downstream tasks, however, the clinically meaningful unit is...

<https://arxiv.org/abs/2602.17949>
26 **CONF****Exploitation of Hidden Context in Dynamic Movement Forecasting: A Neural Network Journey from Recurrent to Graph Neural Networks and General Purpose Transformers**

ArXiv CS.AI 50%

arXiv:2605.14855v1 Announce Type: cross Abstract: Forecasting within signal processing pipelines is crucial for mitigating delays, particularly in predicting the dynamic movements of objects such as NBA players. This task poses significant challenges due to the inherently...

<https://arxiv.org/abs/2605.14855>
27 **CONF****Your CLIP has 164 dimensions of noise: Exploring the embeddings covariance eigenspectrum of contrastively pretrained vision-language transformers**

ArXiv CS.AI 50%

arXiv:2605.14893v1 Announce Type: cross Abstract: Contrastively pre-trained Vision-Language Models (VLMs) serve as powerful feature extractors. Yet, their shared latent spaces are prone to structural anomalies and act as repositories for non-semantic, multi-modal noise. To...

<https://arxiv.org/abs/2605.14893>
28 **CONF****Critic-Driven Voronoi-Quantization for Distilling Deep RL Policies to Explainable Models**

ArXiv CS.AI 50%

arXiv:2605.14897v1 Announce Type: cross Abstract: Despite many successful attempts at explaining Deep Reinforcement Learning policies using distillation, it remains difficult to balance the performance-interpretability trade-off and select a fitting surrogate model. In addition...

<https://arxiv.org/abs/2605.14897>
29 **CONF****Predicting Response to Neoadjuvant Chemotherapy in Ovarian Cancer from CT Baseline Using Multi-Loss Deep Learning**

ArXiv CS.AI 50%

arXiv:2605.14991v1 Announce Type: cross Abstract: Ovarian cancer is the most lethal gynecologic malignancy: around 60% of patients are diagnosed at an advanced stage, with an associated 5-year survival rate of about 30%. Early identification of non-responders to neoadjuvant...

<https://arxiv.org/abs/2605.14991>

30 **CONF****Towards Gaze-Informed AI Disclosure Interfaces: Eye-Tracking Attentional and Cognitive Load While Reading AI-Assisted News****ArXiv CS.AI** **50%**

arXiv:2605.14999v1 Announce Type: cross Abstract: As generative AI becomes increasingly integrated into journalism, designing effective AI-use disclosures that inform readers without imposing unnecessary burden is a key challenge. While prior research has primarily focused on...

<https://arxiv.org/abs/2605.14999>

31 **CONF****SpeakerLLM: A Speaker-Specialized Audio-LLM for Speaker Understanding and Verification Reasoning****ArXiv CS.AI** **50%**

arXiv:2605.15044v1 Announce Type: cross Abstract: As audio-first agents become increasingly common in physical AI, conversational robots, and screenless wearables, audio large language models (audio-LLMs) must integrate speaker-specific understanding to support user...

<https://arxiv.org/abs/2605.15044>

32 **CONF****TFGN: Task-Free, Replay-Free Continual Pre-Training Without Catastrophic Forgetting at LLM Scale****ArXiv CS.AI** **50%**

arXiv:2605.15053v1 Announce Type: cross Abstract: Continually pre-training a large language model on heterogeneous text domains, without replay or task labels, has remained an unsolved architectural problem at LLM scale. Existing methods rely on replay buffers, task identifiers,...

<https://arxiv.org/abs/2605.15053>

33 **CONF****Concurrency without Model Changes: Future-based Asynchronous Function Calling for LLMs****ArXiv CS.AI** **50%**

arXiv:2605.15077v1 Announce Type: cross Abstract: Function calling, also known as tool use, is a core capability of modern LLM agents but is typically constrained by synchronous execution semantics. Under these semantics, LLM decoding is blocked until each function call...

<https://arxiv.org/abs/2605.15077>

34 **CONF****Understanding How International Students in the U.S. Are Using Conversational AI to Support Cross-Cultural Adaptation****ArXiv CS.AI** **50%**

arXiv:2605.15127v1 Announce Type: cross Abstract: Moving to a new culture and adapting to a new life, as an international student, can be a stressful experience. In the US, international students face unique overlapping challenges, yet the current support ecosystem, including...

<https://arxiv.org/abs/2605.15127>

35 **CONF****Position: Behavioural Assurance Cannot Verify the Safety Claims Governance Now Demands****ArXiv CS.AI** **50%**

arXiv:2605.15164v1 Announce Type: cross Abstract: This position paper argues that behavioural assurance, even when carefully designed, is being asked to carry safety claims it cannot verify. AI governance frameworks enacted between 2019 and early 2026 require reviewable evidence...

<https://arxiv.org/abs/2605.15164>

36 CONF

Text Knows What, Tables Know When: Clinical Timeline Reconstruction via Retrieval-Augmented Multimodal Alignment

ArXiv CS.AI 50%

arXiv:2605.15168v1 Announce Type: cross Abstract: Reconstructing precise clinical timelines is essential for modeling patient trajectories and forecasting risk in complex, heterogeneous conditions like sepsis. While unstructured clinical narratives offer semantically rich and...

<https://arxiv.org/abs/2605.15168>

37 CONF

Eradicating Negative Transfer in Multi-Physics Foundation Models via Sparse Mixture-of-Experts Routing

ArXiv CS.AI 50%

arXiv:2605.15179v1 Announce Type: cross Abstract: Scaling Scientific Machine Learning (SciML) toward universal foundation models is bottlenecked by negative transfer: the simultaneous co-training of disparate partial differential equation (PDE) regimes can induce gradient...

<https://arxiv.org/abs/2605.15179>

38 CONF

AgenticEval: Toward Agentic and Self-Evolving Safety Evaluation of Large Language Models

ArXiv CS.AI 50%

arXiv:2509.26100v2 Announce Type: replace Abstract: The rapid integration of Large Language Models (LLMs) into high-stakes domains necessitates reliable safety and compliance evaluation. However, existing static benchmarks are ill-equipped to address the dynamic nature of AI...

<https://arxiv.org/abs/2509.26100>

39 CONF

CausalReasoningBenchmark: A Real-World Benchmark for Disentangled Evaluation of Causal Identification and Estimation

ArXiv CS.AI 50%

arXiv:2602.20571v2 Announce Type: replace Abstract: Many benchmarks for automated causal inference evaluate a system's performance based on a single numerical output, such as an Average Treatment Effect (ATE). This approach conflates two distinct steps in causal analysis:...

<https://arxiv.org/abs/2602.20571>

40 CONF

LLMs learn scientific taste from institutional traces across the social sciences

ArXiv CS.AI 50%

arXiv:2603.16659v3 Announce Type: replace Abstract: Reinforcement-learned reasoning has powered recent AI leaps on verifiable tasks, including mathematics, code, and structure prediction. The harder bottleneck is evaluative judgment in low-verifiability domains, where no oracle...

<https://arxiv.org/abs/2603.16659>

41 CONF

What Do EEG Foundation Models Capture from Human Brain Signals?

ArXiv CS.AI 50%

arXiv:2605.11410v2 Announce Type: replace Abstract: Clinical electroencephalogram (EEG) analysis rests on a hand-crafted feature catalog refined over decades, \emph{e.g.,} band power, connectivity, complexity, and more. Modern EEG foundation models bypass this catalog, learn...

<https://arxiv.org/abs/2605.11410>

42 CONF

CuSearch: Curriculum Rollout Sampling via Search Depth for Agentic RAG

ArXiv CS.AI 50%

arXiv:2605.11611v2 Announce Type: replace Abstract: Reinforcement Learning with Verifiable Rewards (RLVR) has emerged as a promising paradigm for training agentic retrieval-augmented generation (RAG) systems from outcome-only supervision. Most existing methods optimize policies...

<https://arxiv.org/abs/2605.11611>

43 CONF

D-VLA: A High-Concurrency Distributed Asynchronous Reinforcement Learning Framework for Vision-Language-Action Models

ArXiv CS.AI 50%

arXiv:2605.13276v2 Announce Type: replace Abstract: The rapid evolution of Embodied AI has enabled Vision-Language-Action (VLA) models to excel in multimodal perception and task execution. However, applying Reinforcement Learning (RL) to these massive models in large-scale...

<https://arxiv.org/abs/2605.13276>

44 CONF

MMSkills: Towards Multimodal Skills for General Visual Agents

ArXiv CS.AI 50%

arXiv:2605.13527v2 Announce Type: replace Abstract: Reusable skills have become a core substrate for improving agent capabilities, yet most existing skill packages encode reusable behavior primarily as textual prompts, executable code, or learned routines. For visual agents,...

<https://arxiv.org/abs/2605.13527>

45 CONF

The Impact of Generative AI on Collaborative Open-Source Software Development: Evidence from GitHub Copilot

ArXiv CS.AI 50%

arXiv:2410.02091v3 Announce Type: replace-cross Abstract: Generative artificial intelligence (AI) facilitates content production and enhances ideation capabilities, which can significantly influence developer productivity and participation in software development. To explore its...

<https://arxiv.org/abs/2410.02091>

46 CONF

DUET: Optimizing Training Data Mixtures via Feedback from Unseen Evaluation Tasks

ArXiv CS.AI 50%

arXiv:2502.00270v3 Announce Type: replace-cross Abstract: The performance of an LLM depends heavily on the relevance of its training data to the downstream evaluation task. However, in practice, the data involved in an unseen evaluation task is often unknown (e.g., conversations...

<https://arxiv.org/abs/2502.00270>

47 CONF

Pro-DG: Procedural Diffusion Guidance for Architectural Facade Generation

ArXiv CS.AI 50%

arXiv:2504.01571v2 Announce Type: replace-cross Abstract: We use hierarchical procedural rules for the generation of control maps within the stable diffusion framework to produce photo-realistic architectural facade images. Starting from a single input image and its...

<https://arxiv.org/abs/2504.01571>

48 **CONF****Dual Ascent Diffusion for Inverse Problems**

ArXiv CS.AI 50%

arXiv:2505.17353v2 Announce Type: replace-cross Abstract: Ill-posed inverse problems are fundamental in many domains, ranging from astrophysics to medical imaging. Emerging diffusion models provide a powerful prior for solving these problems. Existing maximum-a-posteriori (MAP)...

<https://arxiv.org/abs/2505.17353>
49 **CONF****VER: Vision Expert Transformer for Robot Learning via Foundation Distillation and Dynamic Routing**

ArXiv CS.AI 50%

arXiv:2510.05213v2 Announce Type: replace-cross Abstract: Pretrained vision foundation models (VFMs) advance robotic learning via rich visual representations, yet individual VFMs typically excel only in specific domains, limiting generality across tasks. Distilling multiple VFMs...

<https://arxiv.org/abs/2510.05213>
50 **CONF****Seeing Through the Brain: New Insights from Decoding Visual Stimuli with fMRI**

ArXiv CS.AI 50%

arXiv:2510.16196v2 Announce Type: replace-cross Abstract: Understanding how the brain encodes visual information is a central challenge in neuroscience and machine learning. A promising approach is to reconstruct visual stimuli, essentially images, from functional Magnetic...

<https://arxiv.org/abs/2510.16196>
51 **CONF****From Ranking to Reasoning: Explainable Web API Recommendation via Semantic Reasoning**

ArXiv CS.AI 50%

arXiv:2511.05820v2 Announce Type: replace-cross Abstract: The rapid growth of Web APIs has made automated Web API recommendation essential for efficient mashup development. However, existing approaches suffer from two major limitations: 1) they rely on fixed top-N recommendation...

<https://arxiv.org/abs/2511.05820>
52 **CONF****AI-Driven Optimization under Uncertainty for Mineral Processing Operations**

ArXiv CS.AI 50%

arXiv:2512.01977v2 Announce Type: replace-cross Abstract: The global capacity for mineral processing must expand rapidly to meet the demand for critical minerals, which are essential for building the clean energy technologies necessary to mitigate climate change. However, the...

<https://arxiv.org/abs/2512.01977>
53 **CONF****Teaching and Evaluating LLMs to Reason About Polymer Design Related Tasks**

ArXiv CS.AI 50%

arXiv:2601.16312v2 Announce Type: replace-cross Abstract: Research in AI4Science has shown promise in many science applications, including polymer design. However, current LLMs are ineffective in this problem space because: (i) most models lack polymer-specific knowledge, and...

<https://arxiv.org/abs/2601.16312>

54 CONF

OPT-Engine: Benchmarking the Limits of LLMs in Optimization Modeling via Complexity Scaling

ArXiv CS.AI 50%

arXiv:2601.19924v2 Announce Type: replace-cross Abstract: We investigate the capabilities and scalability of Large Language Models (LLMs) in optimization modeling, a domain requiring structured reasoning and precise formulation. To this end, we introduce OPT-ENGINE, an...

<https://arxiv.org/abs/2601.19924>

55 CONF

Artificial Intelligence Specialization in the European Union: Underexplored Role of the Periphery at NUTS-3 Level

ArXiv CS.AI 50%

arXiv:2602.15249v2 Announce Type: replace-cross Abstract: This study examines the distribution of Artificial Intelligence (AI) research across European NUTS-3 regions during the period 2015-2024. Using bibliometric data from Clarivate InCites and the Citation Topics...

<https://arxiv.org/abs/2602.15249>

56 CONF

M^{\$^2}RNN: Non-Linear RNNs with Matrix-Valued States for Scalable Language Modeling

ArXiv CS.AI 50%

arXiv:2603.14360v2 Announce Type: replace-cross Abstract: Transformers are highly parallel but are limited to computations in the TC^{\$^0\$} complexity class, excluding tasks such as entity tracking and code execution that provably require greater expressive power. Motivated by this...

<https://arxiv.org/abs/2603.14360>

57 CONF

A Security Analysis of the OpenClaw AI Agent Framework

ArXiv CS.AI 50%

arXiv:2603.27517v3 Announce Type: replace-cross Abstract: AI agent frameworks connecting large language model (LLM) reasoning to host execution surfaces -- shell, filesystem, containers, and messaging -- introduce security challenges structurally distinct from conventional...

<https://arxiv.org/abs/2603.27517>

58 CONF

Predict-then-Diffuse: Adaptive Response Length for Compute-Budgeted Inference in Diffusion LLMs

ArXiv CS.AI 50%

arXiv:2605.04215v2 Announce Type: replace-cross Abstract: Diffusion-based Large Language Models (D-LLMs) represent a promising frontier in generative AI, offering fully parallel token generation that can lead to significant throughput advantages and superior GPU utilization over...

<https://arxiv.org/abs/2605.04215>

59 CONF

Predictive Maps of Multi-Agent Reasoning: A Successor-Representation Spectrum for LLM Communication Topologies

ArXiv CS.AI 50%

arXiv:2605.11453v2 Announce Type: replace-cross Abstract: Practitioners deploying multi-agent large language model (LLM) systems must currently choose between communication topologies such as chain, star, mesh, and richer variants without any pre-inference diagnostic for which...

<https://arxiv.org/abs/2605.11453>

60 CONF

Overcoming Dynamics-Blindness: Training-Free Pace-and-Path Correction for VLA Models

ArXiv CS.AI 50%

arXiv:2605.11459v2 Announce Type: replace-cross Abstract: Vision-Language-Action (VLA) models achieve remarkable flexibility and generalization beyond classical control paradigms. However, most prevailing VLAs are trained under a single-frame observation paradigm, which leaves...

<https://arxiv.org/abs/2605.11459>

61 CONF

Boosting Omni-Modal Language Models: Staged Post-Training with Visually Debiased Evaluation

ArXiv CS.AI 50%

arXiv:2605.12034v2 Announce Type: replace-cross Abstract: Omni-modal language models are intended to jointly understand audio, visual inputs, and language, but benchmark gains can be inflated when visual evidence alone is enough to answer a query. We study whether current...

<https://arxiv.org/abs/2605.12034>

62 CONF

Learning, Fast and Slow: Towards LLMs That Adapt Continually

ArXiv CS.AI 50%

arXiv:2605.12484v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are trained for downstream tasks by updating their parameters (e.g., via RL). However, updating parameters forces them to absorb task-specific information, which can result in catastrophic...

<https://arxiv.org/abs/2605.12484>

63 CONF

Does language matter for spoken word classification? A multilingual generative meta-learning approach

ArXiv CS.AI 50%

arXiv:2605.13084v2 Announce Type: replace-cross Abstract: Meta-learning has been shown to have better performance than supervised learning for few-shot monolingual spoken word classification. However, the meta-learning approach remains under-explored in multilingual spoken word...

<https://arxiv.org/abs/2605.13084>