



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	110	Media
TechCrunch AI	3	Media
HackerNews	2	Community

VERIFIED INTELLIGENCE — 1 story

VERI

FastKernels: Benchmarking GPU Kernel Generation in Production

ArXiv CS.AI +1 source

80%

arXiv:2605.23215v1 Announce Type: cross Abstract: LLM-based agents for GPU kernel generation are advancing rapidly, yet their progress is fundamentally constrained by the benchmarks they optimize against. Existing benchmarks are poorly aligned with production inference...

<https://arxiv.org/abs/2605.23215>

CONFIRMED SIGNALS — 113 stories

CONF

Show HN: Local note engine uses LLM to organize notes into a knowledge graph

HackerNews 50% HN 5

i take a lot of notes but rarely find time to organize them so the value of most notes in my personal context quietly disappears. Notecast is a local note engine i've been building to help me with that. it runs a three stage LLM pipeline (Classify -> Organize -> Consolidate)...

<https://github.com/AlexWasHeree/NoteCast>

CONF

Show HN: AgentLens - see if your AI features work, per customer

HackerNews 50% HN 3

No summary — see link for full article.

<https://agentlens-demo-omega.vercel.app>

CONF

Everyone is navigating AI security in real time — even Google

TechCrunch AI 50%

We're in the transition period -- all of us.

<https://techcrunch.com/2026/05/24/everyone-is-navigating-ai-security-in-real-time-even-...>

CONF

Spotify and Universal Music strike deal allowing fan-made AI covers and remixes

TechCrunch AI 50%

Spotify is partnering with Universal Music Group to let Premium subscribers create AI-generated song covers and remixes, with participating artists receiving a share of the revenue.

<https://techcrunch.com/2026/05/21/spotify-and-universal-music-strike-deal-allowing-fan-...>

CONF

BOHM: Zero-Cost Hierarchical Attribution for Compound AI Systems

ArXiv CS.AI 50%

arXiv:2605.22866v1 Announce Type: new Abstract: Compound AI systems route tasks through hierarchies of specialised components. Attribution is dominated by Shapley-based methods (SHAP), which decompose a coalition value function into per-component marginal contributions and...

<https://arxiv.org/abs/2605.22866>

CONF

Skill Retrieval Augmentation for Agentic AI

ArXiv CS.AI 50%

arXiv:2604.24594v2 Announce Type: replace-cross Abstract: As large language models (LLMs) evolve into agentic problem solvers, they increasingly rely on external, reusable skills to handle tasks beyond their native parametric capabilities. In existing agent systems, the dominant...

<https://arxiv.org/abs/2604.24594>

CONF

The Deterministic Horizon: Impossibility Results as Design Specifications for Trustworthy AI Systems

ArXiv CS.AI 50%

arXiv:2605.23024v1 Announce Type: new Abstract: Large language models now write software, draft legal documents, and produce clinical notes, yet fundamental limits, from Turing and Arrow to the No Free Lunch theorems, shape what computation can do. This thesis turns such...

<https://arxiv.org/abs/2605.23024>

CONF

Inductive Deductive Synthesis: Enabling AI to Generate Formally Verified Systems

ArXiv CS.AI 50%

arXiv:2605.23109v1 Announce Type: new Abstract: AI agents increasingly excel at generating, testing, and refining code. However, they fall short on tasks requiring formal guarantees of full coverage that testing alone cannot provide. Distributed systems are a prime example:...

<https://arxiv.org/abs/2605.23109>

0 CONF

AutoResearch AI: Towards AI-Powered Research Automation for Scientific Discovery

ArXiv CS.AI 50%

arXiv:2605.23204v1 Announce Type: new Abstract: Scientific research is being reshaped by AI systems that move beyond isolated assistance toward longer-horizon workflows spanning literature grounding, hypothesis generation, experimentation, validation, reporting, and revision....

<https://arxiv.org/abs/2605.23204>

1 CONF

How Far Will They Go? Red-Teaming Online Influence with Large Language Models

ArXiv CS.AI 50%

arXiv:2605.22880v1 Announce Type: cross Abstract: As large language model (LLM)-based agents increasingly participate in online discourse, red-teaming their capacity to support political influence campaigns is critical for information integrity. In pursuit of this goal, we focus...

<https://arxiv.org/abs/2605.22880>

2 CONF

EvalVerse: Pipeline-Aware and Expert-Calibrated Benchmarking for Professional Cinematic Video Generation

ArXiv CS.AI 50%

arXiv:2605.23271v1 Announce Type: cross Abstract: The rapid evolution of generative video foundation models has propelled the field toward professional-grade cinematic synthesis. To achieve such demanding quality, the community transitions towards Reinforcement Learning (RL) and...

<https://arxiv.org/abs/2605.23271>

3 CONF

Parallel Context Compaction for Long-Horizon LLM Agent Serving

ArXiv CS.AI 50%

arXiv:2605.23296v1 Announce Type: new Abstract: Long-horizon LLM agents accumulate growing conversation histories that eventually exceed the model's context window. Context compaction via LLM-based summarization keeps the conversation bounded, but summarization is inherently...

<https://arxiv.org/abs/2605.23296>

- 4

CONF

Ontological Knowledge Blocks: Executable Compliance and Profile-Based Validation for Trustworthy AI Systems

ArXiv CS.AI 50%

arXiv:2605.23297v1 Announce Type: new Abstract: AI-enabled services deployed in critical digital infrastructure are subject to governance obligations spanning transparency, accountability, fairness, and traceability. Compliance today remains documentation-centric: obligations...

<https://arxiv.org/abs/2605.23297>
- 5

CONF

When Planning Fails Despite Correct Execution: On Epistemic Calibration for LLM-Based Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2605.23414v1 Announce Type: new Abstract: LLM-based multi-agent systems can fail even when planned actions are executed correctly because agents may misjudge their knowledge when evaluating plan feasibility, a phenomenon we term epistemic miscalibration in planning. Unlike...

<https://arxiv.org/abs/2605.23414>
- 6

CONF

Solving the Aircraft Disassembly Scheduling Problem

ArXiv CS.AI 50%

arXiv:2605.23592v1 Announce Type: new Abstract: Dismantling aircrafts reaching their end of life is a complex endeavour that is necessary in terms of sustainability but yields small income margins for air transport companies. An efficient scheduling of the disassembly procedure...

<https://arxiv.org/abs/2605.23592>
- 7

CONF

Beyond Binary Edits Robust Multimodal Knowledge Editing with Adversarial Subspace Alignment

ArXiv CS.AI 50%

arXiv:2605.23780v1 Announce Type: new Abstract: Multimodal large language models (MLLMs) need efficient mechanisms to update knowledge without degrading existing capabilities. While intrinsic multimodal knowledge editing achieves strong reliability and locality, it often...

<https://arxiv.org/abs/2605.23780>
- 8

CONF

An AI-Driven Framework for Energy-Efficient Environmental Monitoring in Smart Cities Using Edge Intelligence

ArXiv CS.AI 50%

arXiv:2605.22824v1 Announce Type: cross Abstract: Environmental monitoring is a crucial component of the smart city infrastructure. It enables informed decision making which enhances sustainability, public health and urban planning. However, the large-scale deployments of the...

<https://arxiv.org/abs/2605.22824>
- 9

CONF

Evaluating Large Language Models in a Complex Hidden Role Game

ArXiv CS.AI 50%

arXiv:2605.22826v1 Announce Type: cross Abstract: Quantifying the deceptive potential of Large Language Models (LLMs) is critical for AI safety, yet difficult to achieve in uncontrolled environments. This work investigates the reasoning, persuasion, and deceptive capabilities of...

<https://arxiv.org/abs/2605.22826>

0 CONF

Computable Fairness: Boltzmann-Softmax Control for AI Resource Allocation

ArXiv CS.AI 50%

arXiv:2605.22827v1 Announce Type: cross Abstract: In large-scale AI systems, allocating scarce resources such as GPU compute time and bandwidth among multiple agents is a critical challenge. Conventional policies focus on efficiency metrics, potentially leading to dominance...

<https://arxiv.org/abs/2605.22827>

1 CONF

RAG4Outcome: A Retrieval-Augmented Multimodal Framework for Prognostic Prediction in Chronic Osteomyelitis

ArXiv CS.AI 50%

arXiv:2605.22833v1 Announce Type: cross Abstract: Chronic osteomyelitis presents substantial prognostic challenges due to its high recurrence risk and complex postoperative recovery trajectories. Traditional assessment often relies on manual scoring systems, which limit...

<https://arxiv.org/abs/2605.22833>

2 CONF

Strategic Coercion Within Alliances: The Greenland Sovereignty Game as an AI Stress Test

ArXiv CS.AI 50%

arXiv:2605.22841v1 Announce Type: cross Abstract: What happens when the strongest alliance member pressures a weaker member over territory and strategic control? We examine the Greenland sovereignty crisis as a stress test for LLM geopolitics, centered on the 2019-2026 U.S. push...

<https://arxiv.org/abs/2605.22841>

3 CONF

The Misattribution Gap: When Memory Poisoning Looks Like Model Failure in Agentic AI Systems

ArXiv CS.AI 50%

arXiv:2605.22842v1 Announce Type: cross Abstract: Multi-agent AI pipelines typically assume that agent misconduct originates from model misalignment. We identify a structural failure in this assumption, the `\emph{Misattribution Gap}`, where memory-layer attacks produce behaviors...

<https://arxiv.org/abs/2605.22842>

4 CONF

ObjectCache: Layerwise Object-Storage Retrieval for KV Cache Reuse

ArXiv CS.AI 50%

arXiv:2605.22850v1 Announce Type: cross Abstract: Prefix KV caching has become a key mechanism in LLM serving: it reduces time to first token (TTFT) by avoiding redundant computation across requests that share a prefix (i.e., the system prompt). However, the accumulated KV cache...

<https://arxiv.org/abs/2605.22850>

5 CONF

PrefBench: Evaluating Zero-Shot LLM Agents in Hidden-Preference Personalized Pricing Negotiations

ArXiv CS.AI 50%

arXiv:2605.22855v1 Announce Type: cross Abstract: Personalized pricing negotiations are a challenging testbed for LLM agents because successful interaction does not guarantee profitable decision making. A seller may produce valid actions and close many deals while still pricing...

<https://arxiv.org/abs/2605.22855>

6 CONF

The Readout Shortcut: Positional Number Copying Dominates Arithmetic CoT Readout in Small Language Models

ArXiv CS.AI 50%

arXiv:2605.22870v1 Announce Type: cross Abstract: Chain-of-thought (CoT) prompting is necessary for arithmetic in small language models, yet shuffling its steps preserves most performance. What does CoT contribute if not logical sequencing? In three 1-3B instruction-tuned LMs on...

<https://arxiv.org/abs/2605.22870>

7 CONF

When Do LLMs Reason? A Dynamical Systems View via Entropy Phase Transitions

ArXiv CS.AI 50%

arXiv:2605.22873v1 Announce Type: cross Abstract: Chain-of-thought (CoT) reasoning has become the default strategy for enhancing LLM capabilities, yet its application raises a fundamental question: when is explicit reasoning actually beneficial? Empirical evidence reveals a...

<https://arxiv.org/abs/2605.22873>

8 CONF

Tensor Cache: Eviction-conditioned Associative Memory for Transformers

ArXiv CS.AI 50%

arXiv:2605.22884v1 Announce Type: cross Abstract: Autoregressive Transformer KV caches grow linearly with context length; sliding-window caching bounds memory but discards evicted tokens entirely, so relevant evidence outside the window becomes inaccessible. We introduce...

<https://arxiv.org/abs/2605.22884>

9 CONF

Seeing without Looking: Do Vision-Language Benchmarks Really Test Vision?

ArXiv CS.AI 50%

arXiv:2605.22903v1 Announce Type: cross Abstract: Benchmark accuracy is often implicitly assumed to reflect grounded visual understanding in vision-language models (VLMs), yet it remains unclear to what extent such scores truly reflect reliance on visual evidence. Motivated by a...

<https://arxiv.org/abs/2605.22903>

0 CONF

Suicide Risk Assessment from AI-powered Video Surveillance: An Interpretable Framework for Prevention in Metro Stations

ArXiv CS.AI 50%

arXiv:2605.22904v1 Announce Type: cross Abstract: Understanding and monitoring human behavior in metro stations play an important role in supporting suicide prevention efforts, where early identification of high-risk situations can enable timely intervention. This requires...

<https://arxiv.org/abs/2605.22904>

1 CONF

CoReVAD: A Contextual Reasoning Framework for Training-Free Video Anomaly Detection

ArXiv CS.AI 50%

arXiv:2605.23116v1 Announce Type: cross Abstract: Existing Video Anomaly Detection (VAD) methods typically rely on task-specific training, leading to strong domain dependency and high training costs. Moreover, most existing methods output only scalar anomaly scores, providing...

<https://arxiv.org/abs/2605.23116>

2 **CONF****LLM Code Smells: A Taxonomy and Detection Approach****ArXiv CS.AI** **50%**

arXiv:2605.22976v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly integrated into software systems for diverse purposes, due to their versatility, flexibility, and ability to simulate human reasoning to some extent. However, poor integration of LLM...

<https://arxiv.org/abs/2605.22976>

3 **CONF****Memorization Dynamics of Fill-in-the-Middle Pretraining****ArXiv CS.AI** **50%**

arXiv:2605.22981v1 Announce Type: cross Abstract: Fill-in-the-middle (FIM) is a pretraining objective widely used to equip causal language models with infilling ability, yet its effect on verbatim memorization remains underexplored. We study the memorization dynamics of FIM in a...

<https://arxiv.org/abs/2605.22981>

4 **CONF****Test-Time Training Undermines Safety Guardrails****ArXiv CS.AI** **50%**

arXiv:2605.22984v1 Announce Type: cross Abstract: Test-Time Training (TTT) is an emerging paradigm that enables models to adapt their parameters during inference, improving performance on tasks such as few-shot learning, retrieval-augmented generation, and complex reasoning....

<https://arxiv.org/abs/2605.22984>

5 **CONF****A Proactive Multi-Agent Dialogue Framework for Assessing Social Language Disorder Traits in Autism****ArXiv CS.AI** **50%**

arXiv:2605.22993v1 Announce Type: cross Abstract: Characteristic linguistic behaviors associated with Social Language Disorder (SLD) in autism spectrum disorder, including echoic repetition, pronoun displacement, and stereotyped media quoting, are largely absent from spontaneous...

<https://arxiv.org/abs/2605.22993>

6 **CONF****Whose Good, Whose Place? The Moral Geography of Agentic AI for Social Good****ArXiv CS.AI** **50%**

arXiv:2605.22995v1 Announce Type: cross Abstract: Agentic AI systems are increasingly proposed for social-good domains, often invoking the United Nations Sustainable Development Goals (SDGs) as a vocabulary of global benefit. Yet claims of social good do not establish...

<https://arxiv.org/abs/2605.22995>

7 **CONF****Sparse Autoencoders Map Brain-LLM Alignment onto Cortical Semantic Topography****ArXiv CS.AI** **50%**

arXiv:2605.23035v1 Announce Type: cross Abstract: Intermediate layers of large language models (LLMs) best predict human brain responses to language, one of the most robust findings in computational neurolinguistics, yet why remains mechanistically unexplained. We address this...

<https://arxiv.org/abs/2605.23035>

8 CONF

Do Language Models Know What Not to Say? Causal Evidence for Statistical Preemption in LLMs

ArXiv CS.AI 50%

arXiv:2605.23039v1 Announce Type: cross Abstract: How do learners acquire knowledge of what is unacceptable without negative evidence? Construction Grammar proposes statistical preemption: exposure to a conventional form (e.g., "donated the books to the library") preempts...

<https://arxiv.org/abs/2605.23039>

9 CONF

DreamerNLplus: Interpretable Modeling of Mental Health Dynamics from Social Media Timelines using Hybrid Rule-Based and RAG Methods

ArXiv CS.AI 50%

arXiv:2605.23052v1 Announce Type: cross Abstract: We present DreamerNLplus, a hybrid framework for modeling mental health dynamics from social media timelines in the CLPsych 2026 shared task. Our system addresses three tasks: psychological state modeling, temporal change...

<https://arxiv.org/abs/2605.23052>

0 CONF

Anytime Training with Schedule-Free Spectral Optimization

ArXiv CS.AI 50%

arXiv:2605.23061v1 Announce Type: cross Abstract: Standard neural network training relies on learning-rate schedules tied to a fixed horizon, leading to strong path dependence and costly re-tuning as data availability changes. Schedule-Free (SF) methods address this by removing...

<https://arxiv.org/abs/2605.23061>

1 CONF

Leveraging Foundation Models for Causal Generative Modeling

ArXiv CS.AI 50%

arXiv:2605.23861v1 Announce Type: cross Abstract: Causal generative modeling is essential for developing reliable and transparent AI systems capable of counterfactual reasoning. While existing approaches focus on integrating causal constraints during the training of generative...

<https://arxiv.org/abs/2605.23861>

2 CONF

Security of LLM-generated Code: A Comparative Analysis

ArXiv CS.AI 50%

arXiv:2605.23091v1 Announce Type: cross Abstract: The majority of software developers use or are planning to use Artificial Intelligence (AI) tools in their development processes. Their top reasons include improving productivity and faster learning. In fact, Large Language Model...

<https://arxiv.org/abs/2605.23091>

3 CONF

Do Synthetic Brain MRIs Reliably Improve Tumour Classification? A StyleGAN2-ADA Class-Plane Augmentation Study on BRISC 2025

ArXiv CS.AI 50%

arXiv:2605.23094v1 Announce Type: cross Abstract: Generative augmentation is often proposed as a remedy for small medical-image datasets, but synthetic images are only useful when they improve downstream task performance. "Augmentation" here means synthetic supplementation:...

<https://arxiv.org/abs/2605.23094>

- 4 **CONF**
A Fine-Tuned BERT Classifier for Personal-Letter Titles in Late-Ming and Early-Qing Collected Works
 ArXiv CS.AI 50%
 arXiv:2605.23103v1 Announce Type: cross Abstract: I present Lepton (Letter Prediction), a fine-tuned BERT classifier that predicts whether a title in a Classical Chinese wenji table of contents is a personal letter or a closely confusable preface (particularly the...
<https://arxiv.org/abs/2605.23103>
- 5 **CONF**
Philosophical Dispositions as Behavioral Constraints for AI-Assisted Code Review: An Empirical Study
 ArXiv CS.AI 50%
 arXiv:2605.23108v1 Announce Type: cross Abstract: AI-assisted code review tools typically operate as generic "expert reviewer" agents, producing homogeneous findings regardless of the analysis type needed. We present a system that constrains AI reviewer behavior through...
<https://arxiv.org/abs/2605.23108>
- 6 **CONF**
Defining AI Fatigue in Academic Contexts: Dimensions, Indicators, and a Stage-Based Model Using Grounded Theory
 ArXiv CS.AI 50%
 arXiv:2605.23123v1 Announce Type: cross Abstract: The integration of AI tools in academic settings has introduced a distinct form of strain that existing frameworks like technostress and digital fatigue have not yet fully addressed. This study develops a conceptual model and...
<https://arxiv.org/abs/2605.23123>
- 7 **CONF**
TABX: A High-Throughput Sandbox Battle Simulator for Multi-Agent Reinforcement Learning
 ArXiv CS.AI 50%
 arXiv:2602.01665v2 Announce Type: replace-cross Abstract: The design of environments plays a critical role in shaping the development and evaluation of cooperative multi-agent reinforcement learning (MARL) algorithms. While existing benchmarks highlight critical challenges, they...
<https://arxiv.org/abs/2602.01665>
- 8 **CONF**
Reinforcement Learning for Microcanonical Graph Ensemble with Assortativity Constraints
 ArXiv CS.AI 50%
 arXiv:2605.23285v1 Announce Type: cross Abstract: How network structure determines function is a fundamental question, and it can be investigated by graph ensembles with precisely controlled structural properties. Canonical approaches, formulated as exponential random graph...
<https://arxiv.org/abs/2605.23285>
- 9 **CONF**
As X, Do Y: How Persona and Task Combine in Instruction-Tuned LLMs
 ArXiv CS.AI 50%
 arXiv:2605.23147v1 Announce Type: cross Abstract: Role prompts of the form As X, do Y admit a clean linear decomposition at one specific site in the residual stream: the prompt-to-answer transition -- the last prompt token together with the first two generated tokens -- in an...
<https://arxiv.org/abs/2605.23147>

0 CONF

Positional Failures in Long-Context LLMs: A Blind Spot in Reasoning Benchmarks

ArXiv CS.AI 50%

arXiv:2605.23170v1 Announce Type: cross Abstract: Position-controlled evaluation is standard for retrieval tasks such as Needle-in-a-Haystack and RULER, but mainstream reasoning benchmarks do not control positional placement of target tasks in long contexts. We audit 11...

<https://arxiv.org/abs/2605.23170>

1 CONF

MAS-Orchestra: Understanding and Improving Multi-Agent Reasoning Through Holistic Orchestration and Controlled Benchmarks

ArXiv CS.AI 50%

arXiv:2601.14652v5 Announce Type: replace Abstract: While multi-agent systems (MAS) promise elevated intelligence through coordination of agents, current approaches to automatic MAS design under-deliver. Such shortcomings stem from two key factors: (1) methodological complexity...

<https://arxiv.org/abs/2601.14652>

2 CONF

Scalable Heterogeneous Graph Foundation Models for Data-Driven Optimal Power Flow in Smart Grids

ArXiv CS.AI 50%

arXiv:2605.23194v1 Announce Type: cross Abstract: Fast and reliable optimal power flow (OPF) approximation is essential for reliable smart-grid operation, yet many learning-based surrogates either flatten the native heterogeneous structure of power networks, target a limited set...

<https://arxiv.org/abs/2605.23194>

3 CONF

Are Frontier LLMs Ready for Cybersecurity? Evidence for Vertical Foundation Models from Dual-Mode Vulnerability Benchmarks

ArXiv CS.AI 50%

arXiv:2605.23243v1 Announce Type: cross Abstract: We evaluate whether frontier LLMs are ready for cybersecurity through a dual-mode benchmark: white-box function-level vulnerability detection (VulnLLM-R, across C/Java/Python) and black-box web application security testing (five...

<https://arxiv.org/abs/2605.23243>

4 CONF

Enhancing Deep Neural Network Reliability with Refinement and Calibration

ArXiv CS.AI 50%

arXiv:2605.23249v1 Announce Type: cross Abstract: Although deep neural networks (DNNs) achieve high predictive accuracy, their confidence estimates are often unreliable, potentially compromising user trust in their decisions. This has motivated research on calibrated models,...

<https://arxiv.org/abs/2605.23249>

5 CONF

Coloring the Noise: Adversarial Sobolev Alignment for Faithful Image Super Resolution

ArXiv CS.AI 50%

arXiv:2605.23264v1 Announce Type: cross Abstract: Generative priors in Image Super-Resolution (SR) often compromise faithful restoration, we attribute this limitation to a fundamental spectral misalignment between isotropic objectives and the intrinsic natural image manifold...

<https://arxiv.org/abs/2605.23264>

- 6
CONF
Convergence Without Understanding: When Language Models Agree on Representations but Disagree on Reasoning
ArXiv CS.AI 50%
arXiv:2605.23315v1 Announce Type: cross Abstract: Large language models trained under diverse objectives and architectures have been shown to develop increasingly similar internal representations, an observation formalized as the Platonic Representation Hypothesis. Whether this...
<https://arxiv.org/abs/2605.23315>
- 7
CONF
CHASD: Language Increment-Calibrated Contrastive Decoding against Hallucination in LVLMs
ArXiv CS.AI 50%
arXiv:2605.23344v1 Announce Type: cross Abstract: Large Vision-Language Models have shown strong multimodal reasoning capabilities, yet they remain susceptible to object hallucinations when language priors dominate insufficient or misaligned visual evidence. Training-free...
<https://arxiv.org/abs/2605.23344>
- 8
CONF
XWind: A Cross-site Router for Large Language Model Inference Serving at Renewable Energy Farms
ArXiv CS.AI 50%
arXiv:2605.23348v1 Announce Type: cross Abstract: AI power demand is growing at an unprecedented rate while power grids are often ailing and struggle to keep up. Grid expansion comes with high capital expenditure and long-distance transmission losses, yet there is abundant...
<https://arxiv.org/abs/2605.23348>
- 9
CONF
Reflex: Reinforcement Learning with Reflection Symmetry Exploitation in State-Based Continuous Control
ArXiv CS.AI 50%
arXiv:2605.23415v1 Announce Type: cross Abstract: Reinforcement learning has long struggled with poor sample efficiency. One promising approach to mitigate this problem is leveraging group-invariant Markov Decision Processes (G-invariant MDPs). Existing works in this direction...
<https://arxiv.org/abs/2605.23415>
- 0
CONF
Online Hand Gesture Recognition Using 3D Convolutional Neural Networks
ArXiv CS.AI 50%
arXiv:2605.23409v1 Announce Type: cross Abstract: In human computer interaction, real-time detection and classification of dynamic hand gestures is challenging as: 1) the system must run in a real-time video stream and there is no noticeable lag in response after performing a...
<https://arxiv.org/abs/2605.23409>
- 1
CONF
Socially fluent AI decouples conversational signals from source identity in online interaction
ArXiv CS.AI 50%
arXiv:2605.23426v1 Announce Type: cross Abstract: Socially fluent agentic AI can now participate in online interaction in ways that resemble ordinary human conversation, potentially weakening people's ability to infer who is human from conversational signals alone. We tested...
<https://arxiv.org/abs/2605.23426>

2 **CONF****AI Security Research Should Better Incentivize Defense Research****ArXiv CS.AI 50%**

arXiv:2605.23448v1 Announce Type: cross Abstract: This work examines an imbalance in artificial intelligence (AI) security research: the field tends to produce more work on attacking AI systems than on defending them. Drawing on related academic papers, we find biased...

<https://arxiv.org/abs/2605.23448>

3 **CONF****Automated Random Embedding for Practical Bayesian Optimization with Unknown Effective Dimension****ArXiv CS.AI 50%**

arXiv:2605.23473v1 Announce Type: cross Abstract: Bayesian optimization is widely employed for optimizing complex black-box functions but struggles with the curse of dimensionality. Random embedding, as a dimension reduction strategy, simplifies tasks that possess the effective...

<https://arxiv.org/abs/2605.23473>

4 **CONF****Precise: SDE-Consistent Stochastic Sampling for RL Post-Training of Flow-Matching Models****ArXiv CS.AI 50%**

arXiv:2605.23522v1 Announce Type: cross Abstract: Reinforcement learning (RL) has become an effective way to improve prompt alignment and perceptual quality in diffusion and flow-matching generators. A critical step for applying online RL to flow matching is turning the...

<https://arxiv.org/abs/2605.23522>

5 **CONF****PathNavigate: A Training-Free Pathology Agent with Surprise-Guided Scan and Shared Slide Memory for Whole-Slide Image VQA****ArXiv CS.AI 50%**

arXiv:2605.23559v1 Announce Type: cross Abstract: Whole-slide image visual question answering (WSI-VQA) frames pathology as an extreme-context search problem: to answer a free-form clinical query, a system must first navigate a gigapixel slide under a strict inspection budget to...

<https://arxiv.org/abs/2605.23559>

6 **CONF****HARNESS-LM: A Three-Phase Training Recipe for Harnessing SLMs in Sponsored Search Retrieval****ArXiv CS.AI 50%**

arXiv:2605.23572v1 Announce Type: cross Abstract: In the competitive landscape of sponsored search, balancing retrieval quality with production latency is a critical challenge. While large retrieval models based on Small Language Models (SLMs) such as Qwen3-Embedding-4B/8B set...

<https://arxiv.org/abs/2605.23572>

7 **CONF****DiLaDiff: Distilled Latent-Augmented Diffusion for Language Modeling****ArXiv CS.AI 50%**

arXiv:2605.23605v1 Announce Type: cross Abstract: Diffusion language models intrinsically fail to capture correlations between decoded tokens, which leads to a harsh trade-off between sampling quality and throughput. To solve this issue, we propose DiLaDiff, a variant of masked...

<https://arxiv.org/abs/2605.23605>

8 CONF

EM-Vid: Training-Free Entity-Centric Memory for Efficient and Consistent Multi-Shot Video Generation

ArXiv CS.AI 50%

arXiv:2605.23610v1 Announce Type: cross Abstract: Multi-shot video generation requires maintaining a consistent appearance of recurring entities across shots while remaining faithful to shot-specific text prompts. Recent autoregressive methods reuse previously generated frames...

<https://arxiv.org/abs/2605.23610>

9 CONF

Learning Through Noise: Why Subliminal Learning Works and When It Fails

ArXiv CS.AI 50%

arXiv:2605.23645v1 Announce Type: cross Abstract: In the context of artificial neural networks, subliminal learning refers to the transfer of task-relevant knowledge or unintended biases from teacher to student models through distillation on task-unrelated...

<https://arxiv.org/abs/2605.23645>

0 CONF

CVSearch: Empowering Multimodal LLMs with Cognitive Visual Search for High-Resolution Image Perception

ArXiv CS.AI 50%

arXiv:2605.23655v1 Announce Type: cross Abstract: High-resolution (HR) image perception presents a key bottleneck for multimodal large language models (MLLMs). While visual search offers a promising solution, existing methods struggle with the trade-off between coverage and...

<https://arxiv.org/abs/2605.23655>

1 CONF

Not Too Generative, Not Too Discriminative: The Human Alignment Sweet Spot

ArXiv CS.AI 50%

arXiv:2605.23819v1 Announce Type: cross Abstract: A central question in computational vision is whether human-like visual representations are better explained by discriminative or generative learning. Existing comparisons, however, often confound the learning objective with...

<https://arxiv.org/abs/2605.23819>

2 CONF

It's the humans, not the data: Geopolitical bias in LLMs originates in post-training, amplified by the language of the prompt

ArXiv CS.AI 50%

arXiv:2605.23825v1 Announce Type: cross Abstract: It has generally been assumed that geopolitical bias in language models originates from the training data used during the pre-training phase. We tested seven open-weight LLM pairs consisting of the base model (pre-training only)...

<https://arxiv.org/abs/2605.23825>

3 CONF

Human Decision-Making with Persuasive and Narrative LLM Explanations

ArXiv CS.AI 50%

arXiv:2605.23867v1 Announce Type: cross Abstract: Large language models (LLMs) have the potential to aid and improve human decision-making in classification tasks, not only by providing fairly accurate predictions, but also in their ability to generate cogent narrative...

<https://arxiv.org/abs/2605.23867>

4 **CONF****PGT: Procedurally Generated Tasks for improving visual grounding in MLLMs**

ArXiv CS.AI 50%

arXiv:2605.23883v1 Announce Type: cross Abstract: Despite remarkable progress in Multimodal Large Language Models (MLLMs), these models still struggle with fine-grained understanding tasks. In this work, we propose Procedurally Generated Tasks (PGT), a simple data-driven...

<https://arxiv.org/abs/2605.23883>

5 **CONF****Good Token Hunting: A Hitchhiker's Guide to Token Selection for Visual Geometry Transformers**

ArXiv CS.AI 50%

arXiv:2605.23892v1 Announce Type: cross Abstract: Visual geometry transformers have become powerful architectures for multi-view 3D reconstruction, enabling joint prediction of multiple 3D attributes in a feed-forward manner. However, their computational cost grows quadratically...

<https://arxiv.org/abs/2605.23892>

6 **CONF****LLMs as Noisy Channels: A Shannon Perspective on Model Capacity and Scaling Laws**

ArXiv CS.AI 50%

arXiv:2605.23901v1 Announce Type: cross Abstract: Existing scaling laws for Large Language Models (LLMs), predominantly monotonic power laws, fail to explain emerging non-monotonic phenomena such as catastrophic overtraining and quantization-induced degradation, where...

<https://arxiv.org/abs/2605.23901>

7 **CONF****Robust Counterfactual Inference in Markov Decision Processes**

ArXiv CS.AI 50%

arXiv:2502.13731v5 Announce Type: replace Abstract: This paper addresses a key limitation in existing counterfactual inference methods for Markov Decision Processes (MDPs). Current approaches assume a specific causal model to make counterfactuals identifiable. However, there are...

<https://arxiv.org/abs/2502.13731>

8 **CONF****FATHOMS-RAG: A Framework for the Assessment of Thinking and Observation in Multimodal Systems that use Retrieval Augmented Generation**

ArXiv CS.AI 50%

arXiv:2510.08945v3 Announce Type: replace Abstract: Retrieval-augmented generation (RAG) has emerged as a promising paradigm for improving factual accuracy in large language models (LLMs). We introduce a benchmark designed to evaluate RAG pipelines as a whole, evaluating a...

<https://arxiv.org/abs/2510.08945>

9 **CONF****Deja Vu in Plots: Leveraging Cross-Session Evidence with Retrieval-Augmented LLMs for Live Streaming Risk Assessment**

ArXiv CS.AI 50%

arXiv:2601.16027v2 Announce Type: replace Abstract: The rise of live streaming has transformed online interaction, enabling massive real-time engagement but also exposing platforms to complex risks such as scams and coordinated malicious behaviors. Detecting these risks is...

<https://arxiv.org/abs/2601.16027>

- 0 **CONF**
- TCAP: Tri-Component Attention Profiling for Unsupervised Backdoor Detection in MLLM Fine-Tuning**
- ArXiv CS.AI 50%
- arXiv:2601.21692v2 Announce Type: replace Abstract: Fine-Tuning-as-a-Service (FTaaS) facilitates the customization of Multimodal Large Language Models (MLLMs) but introduces critical backdoor risks via poisoned data. Existing defenses either rely on supervised signals or fail to...
- <https://arxiv.org/abs/2601.21692>
-
- 1 **CONF**
- Scaling-Aware Adapter for Structure-Grounded LLM Reasoning**
- ArXiv CS.AI 50%
- arXiv:2602.02780v3 Announce Type: replace Abstract: Large language models (LLMs) are enabling reasoning over 2D and 3D structures, yet existing methods remain modality-specific and typically compress structural inputs through sequence-based tokenization or fixed-length query...
- <https://arxiv.org/abs/2602.02780>
-
- 2 **CONF**
- ReCoVer: Resilient LLM Pre-Training System via Fault-Tolerant Collective and Versatile Workload**
- ArXiv CS.AI 50%
- arXiv:2605.11215v2 Announce Type: replace-cross Abstract: Pre-training large language models on massive GPU clusters has made hardware faults routine rather than rare, driving the need for resilient training systems. Yet existing frameworks either focus on specific parallelism...
- <https://arxiv.org/abs/2605.11215>
-
- 3 **CONF**
- GT-HarmBench: Benchmarking AI Safety Risks Through the Lens of Game Theory**
- ArXiv CS.AI 50%
- arXiv:2602.12316v2 Announce Type: replace Abstract: Frontier AI systems are increasingly capable and deployed in high-stakes multi-agent environments. However, existing AI safety benchmarks largely evaluate single agents, leaving multi-agent risks such as coordination failure...
- <https://arxiv.org/abs/2602.12316>
-
- 4 **CONF**
- Bridging AI and Clinical Reasoning: Abductive Explanations for Alignment on Critical Symptoms**
- ArXiv CS.AI 50%
- arXiv:2602.13985v2 Announce Type: replace Abstract: Artificial intelligence (AI) has demonstrated strong potential in clinical diagnostics, often achieving accuracy comparable to or exceeding that of human experts. A key challenge, however, is that AI reasoning frequently...
- <https://arxiv.org/abs/2602.13985>
-
- 5 **CONF**
- AI Evaluation Should Require Standardized Item-Level Data Releases**
- ArXiv CS.AI 50%
- arXiv:2604.03244v2 Announce Type: replace Abstract: This position paper argues that standardized item-level benchmark data should become the default infrastructure for AI evaluation. Current evaluations suffer from underspecified item selection, construct misalignment, and poor...
- <https://arxiv.org/abs/2604.03244>
-

- 6
CONF

Is Capability a Liability? More Capable Language Models Make Worse Forecasts When It Matters Most

ArXiv CS.AI 50%

arXiv:2605.22672v2 Announce Type: replace Abstract: We document inverse scaling in LLMs on forecasting problems whose underlying time series exhibit superlinear growth and tail risk of regime change, a structure common in finance and epidemiology. On these tasks, more capable...

<https://arxiv.org/abs/2605.22672>
- 7
CONF

Diffusion and Flow Matching Models for Tabular Data: A Survey

ArXiv CS.AI 50%

arXiv:2502.17119v2 Announce Type: replace-cross Abstract: Deep generative models have made rapid progress in image, text, audio, and video generation, and are increasingly being applied to structured records. For tabular data, however, generative modeling remains difficult: a...

<https://arxiv.org/abs/2502.17119>
- 8
CONF

Through the Stealth Lens: Attention-Aware Defenses Against Poisoning in RAG

ArXiv CS.AI 50%

arXiv:2506.04390v2 Announce Type: replace-cross Abstract: Retrieval-augmented generation (RAG) systems are vulnerable to attacks that inject poisoned passages into the retrieved context, even at low corruption rates. We show that existing attacks are not designed to be stealthy,...

<https://arxiv.org/abs/2506.04390>
- 9
CONF

Online Learning with Multiple Fairness Regularizers via Graph-Structured Feedback

ArXiv CS.AI 50%

arXiv:2508.14311v2 Announce Type: replace-cross Abstract: There is an increasing need to enforce multiple, often competing, measures of fairness within automated decision systems. The appropriate weighting of these fairness objectives is typically unknown a priori, may change...

<https://arxiv.org/abs/2508.14311>
- 0
CONF

Disentangling Interaction and Bias Effects in Opinion Dynamics of Large Language Models

ArXiv CS.AI 50%

arXiv:2509.06858v2 Announce Type: replace-cross Abstract: Large Language Models are increasingly used to simulate human opinion dynamics, yet the effect of genuine interaction is often obscured by systematic biases. We develop a Bayesian framework to disentangle and quantify...

<https://arxiv.org/abs/2509.06858>
- 1
CONF

Efficient and Transferable Agentic Knowledge Graph RAG via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2509.26383v5 Announce Type: replace-cross Abstract: Knowledge-graph retrieval-augmented generation (KG-RAG) couples large language models (LLMs) with structured, verifiable knowledge graphs (KGs) to reduce hallucination and provide reasoning traces. However, current KG-RAG...

<https://arxiv.org/abs/2509.26383>

2 **CONF****GILT: An LLM-Free, Tuning-Free Graph Foundational Model for In-Context Learning****ArXiv CS.AI 50%**

arXiv:2510.04567v2 Announce Type: replace-cross Abstract: Graph Neural Networks (GNNs) are powerful tools for processing relational data but often struggle to generalize to unseen graphs, giving rise to the development of Graph Foundational Models (GFM). However, current GFMs...

<https://arxiv.org/abs/2510.04567>

3 **CONF****RAG-Pull: Turning Retrieval into a Code-Injection Channel via Invisible Unicode Perturbations****ArXiv CS.AI 50%**

arXiv:2510.11195v2 Announce Type: replace-cross Abstract: Retrieval-Augmented Generation (RAG) increases the reliability and trustworthiness of the LLM response and reduces hallucination by eliminating the need for model retraining. It does so by adding external data into the...

<https://arxiv.org/abs/2510.11195>

4 **CONF****LACY: A Vision-Language Model-based Language-Action Cycle for Self-Improving Robotic Manipulation****ArXiv CS.AI 50%**

arXiv:2511.02239v2 Announce Type: replace-cross Abstract: Learning generalizable policies for robotic manipulation increasingly relies on large-scale models that map language instructions to actions (L2A). However, this one-way paradigm often produces policies that execute tasks...

<https://arxiv.org/abs/2511.02239>

5 **CONF****DocVAL: Validated Chain-of-Thought Distillation for Grounded Document VQA****ArXiv CS.AI 50%**

arXiv:2511.22521v3 Announce Type: replace-cross Abstract: Document visual question answering requires models not only to answer questions correctly, but also to precisely localize answers within complex document layouts. While large vision-language models (VLMs) achieve strong...

<https://arxiv.org/abs/2511.22521>

6 **CONF****Fine-Tuning Causal LLMs for Text Classification: Embedding-Based vs. Instruction-Based Approaches****ArXiv CS.AI 50%**

arXiv:2512.12677v2 Announce Type: replace-cross Abstract: We explore efficient strategies to fine-tune decoder-only Large Language Models (LLMs) for downstream text classification under resource constraints. Two approaches are investigated: (1) attaching a classification head to...

<https://arxiv.org/abs/2512.12677>

7 **CONF****Bridging Data and Physics: A Graph Neural Network-Based Hybrid Twin Framework****ArXiv CS.AI 50%**

arXiv:2512.15767v2 Announce Type: replace-cross Abstract: Simulating complex unsteady physical phenomena relies on detailed mathematical models, simulated for instance by using the Finite Element Method (FEM). However, these models often exhibit discrepancies from the reality...

<https://arxiv.org/abs/2512.15767>

8 CONF

Operator-Based Generalization Bound for Deep Learning: Insights on Multi-Task Learning

ArXiv CS.AI 50%

arXiv:2512.19184v2 Announce Type: replace-cross Abstract: This paper presents novel generalization bounds for vector-valued neural networks and deep kernel methods, focusing on multi-task learning through an operator-theoretic framework. Our key development lies in strategically...

<https://arxiv.org/abs/2512.19184>

9 CONF

R\$^3\$L: Reflect-then-Retry Reinforcement Learning with Language-Guided Exploration, Pivotal Credit, and Positive Amplification

ArXiv CS.AI 50%

arXiv:2601.03715v2 Announce Type: replace-cross Abstract: Reinforcement learning drives recent advances in LLM reasoning and agentic capabilities, yet current approaches struggle with both exploration and exploitation. Exploration suffers from low success rates on difficult...

<https://arxiv.org/abs/2601.03715>

00 CONF

SciHorizon-GENE: Benchmarking LLM for Life Sciences Inference from Gene Knowledge to Functional Understanding

ArXiv CS.AI 50%

arXiv:2601.12805v3 Announce Type: replace-cross Abstract: Large language models (LLMs) have shown growing promise in biomedical research, particularly for knowledge-driven interpretation tasks. However, their ability to reliably reason from gene-level knowledge to functional...

<https://arxiv.org/abs/2601.12805>

01 CONF

GradingAttack: Exposing Security Vulnerabilities in LLM Based Educational Grading Agents

ArXiv CS.AI 50%

arXiv:2602.00979v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly deployed as educational agents for automatic short answer grading (ASAG) in real-world educational environments, significantly boosting assessment efficiency and scalability...

<https://arxiv.org/abs/2602.00979>

02 CONF

ArcMark: Distortion-Free Multi-Byte LLM Watermark via Optimal Transport

ArXiv CS.AI 50%

arXiv:2602.07235v2 Announce Type: replace-cross Abstract: Watermarking is an important tool for promoting the responsible use of large language models (LLMs). Existing watermarks insert a signal into generated tokens that either flags LLM-generated text (zero-bit watermarking)...

<https://arxiv.org/abs/2602.07235>

03 CONF

Empowering 9-1-1 Calltaking Training with Generative AI: Experiences and Lessons Learned

ArXiv CS.AI 50%

arXiv:2602.13241v2 Announce Type: replace-cross Abstract: Emergency call-takers form the first operational link in public safety response, handling over 240 million calls annually while facing a sustained training crisis: staffing shortages exceed 25% in many centers, and...

<https://arxiv.org/abs/2602.13241>

04 CONF

BarrierSteer: LLM Safety via Learning Barrier Steering

ArXiv CS.AI 50%

arXiv:2602.20102v2 Announce Type: replace-cross Abstract: Despite the strong performance of large language models (LLMs) across diverse tasks, their susceptibility to adversarial attacks and unsafe content generation remains a significant obstacle to deployment, particularly in...

<https://arxiv.org/abs/2602.20102>

05 CONF

Multimodal Crystal Flow: Any-to-Any Modality Generation for Unified Crystal Modeling

ArXiv CS.AI 50%

arXiv:2602.20210v2 Announce Type: replace-cross Abstract: Crystal modeling spans a family of conditional and unconditional generation tasks, including crystal structure prediction (CSP) and de novo generation (DNG). While recent deep generative models have shown promising...

<https://arxiv.org/abs/2602.20210>

06 CONF

Tabular PDF Information Extraction with Local LLMs and Layout-Aware Parsing: A Reliability Evaluation

ArXiv CS.AI 50%

arXiv:2604.00003v2 Announce Type: replace-cross Abstract: Extracting structured information from academic PDF documents is non trivial: a single page typically combines free text metadata with tabular regions, exhibits cross program variation, and is susceptible to Unicode...

<https://arxiv.org/abs/2604.00003>

07 CONF

Visually-Guided Policy Optimization for Multimodal Reasoning

ArXiv CS.AI 50%

arXiv:2604.09349v2 Announce Type: replace-cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has significantly advanced the reasoning ability of vision-language models (VLMs). However, the inherent text-dominated nature of VLMs often leads to insufficient...

<https://arxiv.org/abs/2604.09349>

08 CONF

ProtDBench: A Unified Benchmark of Protein Binder Design and Evaluation

ArXiv CS.AI 50%

arXiv:2605.04118v2 Announce Type: replace-cross Abstract: Recent advances in de novo protein binder design have enabled increasing experimental validation, yet reported in silico metrics remain difficult to interpret or compare across studies due to non-standardized evaluation...

<https://arxiv.org/abs/2605.04118>

09 CONF

Content-Aware Attack Detection in LLM Agent Tool-Call Traffic: An Empirical Study of Features, Architectures, and Evaluation Protocols

ArXiv CS.AI 50%

arXiv:2605.11053v3 Announce Type: replace-cross Abstract: The Model Context Protocol (MCP) has become a widely adopted interface for LLM agents to invoke external tools, yet learned monitoring of MCP tool-call traffic remains underexplored. In this article, the proposed detector...

<https://arxiv.org/abs/2605.11053>

10 CONF

Can the Recovery Mechanism Survive AI? Skill Formation, Labor, and What Current Measurement Misses

ArXiv CS.AI 50%

arXiv:2605.16283v2 Announce Type: replace-cross Abstract: Throughout the modern era, when new technologies displaced workers, societies adapted through the same mechanism: education raised the cognitive ceiling, producing workers capable of tasks machines could not yet reach....

<https://arxiv.org/abs/2605.16283>

11 CONF

S-Bus: Automatic Read-Set Reconstruction for Multi-Agent LLM State Coordination

ArXiv CS.AI 50%

arXiv:2605.17076v2 Announce Type: replace-cross Abstract: We address concurrency control for LLM agents sharing mutable state over HTTP, where agents cannot be modified to declare read sets. S-Bus is an HTTP middleware whose central mechanism, a server-side DeliveryLog,...

<https://arxiv.org/abs/2605.17076>

12 CONF

IVF-TQ: Calibration-Free Streaming Vector Search via a Codebook-Free Residual Layer

ArXiv CS.AI 50%

arXiv:2605.17415v2 Announce Type: replace-cross Abstract: Approximate nearest neighbor (ANN) indexes deployed against streaming corpora silently lose recall over weeks. The standard diagnosis is distribution shift, but under shuffled-i.i.d. ingestion -- no shift at all --...

<https://arxiv.org/abs/2605.17415>

13 CONF

Distilling Linearized Behavior into Non-Linear Fine-Tuning for Effective Task Arithmetic

ArXiv CS.AI 50%

arXiv:2605.18993v2 Announce Type: replace-cross Abstract: Task vector composition has emerged as a promising paradigm for editing pre-trained models, enabling model merging through addition and unlearning through subtraction. Fine-tuning in the tangent space of a pre-trained...

<https://arxiv.org/abs/2605.18993>

14 CONF

ThoughtTrace: Understanding User Thoughts in Real-World LLM Interactions

ArXiv CS.AI 50%

arXiv:2605.20087v2 Announce Type: replace-cross Abstract: Conversational AI has now reached billions of users, yet existing datasets capture only what people say, not what they think. We introduce ThoughtTrace, the first large-scale dataset that pairs real-world multi-turn...

<https://arxiv.org/abs/2605.20087>