



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	161	Media
HackerNews	5	Community
The Verge AI	5	Media
TechCrunch AI	5	Media
MIT Tech Review	2	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: Posthorn, self-hosted mail without the mail server

HackerNews

+1 source

90%

HN 10

AI Agents operate in the digital world. Humans operate in the physical world. Hands&Claws is a social platform that brings together all intelligent individuals. And allow them to work together. I developed this software because I discovered that there are some things that humans...

<https://github.com/craigmccaskill/posthorn>

VERI

Show HN: I used AI to built static recompilers for 5 retro game systems

HackerNews

+1 source

90%

HN 5

arXiv:2605.27022v1 Announce Type: new Abstract: Causal analysis is a crucial task in many domains, including manufacturing, social science, and medicine. However, despite recent progress, the conceptual and methodological complexity of causal methods makes them largely...

<https://1379.tech/nes-snes-genesis-virtualboy-and-psx-a-journey-with-ai-and-recompilation/>

VERI

Position: AI Safety Requires Effective Controllability

ArXiv CS.AI

+1 source

80%

arXiv:2605.27117v1 Announce Type: new Abstract: AI safety is still largely framed as alignment: training models to follow human preferences, safety policies, and normative constraints. That framing has improved the behavior of modern language models, but aligned behavior does...

<https://arxiv.org/abs/2605.27117>

VERI

BioFormer: Rethinking Cross-Subject Generalization via Spectral Structural Alignment in Biomedical Time-Series

ArXiv CS.AI

+1 source

80%

arXiv:2605.22468v2 Announce Type: replace-cross Abstract: Cross-subject generalization in biomedical time-series refers to training on data from some subjects and testing on unseen subjects. The key challenge is to suppress subject specific variability in BTS representations. Most...

<https://arxiv.org/abs/2605.22468>

CONFIRMED SIGNALS — 170 stories

CONF

Show HN: I built a tool to auto-accept AI slop and bigtech devs loves it

HackerNews

50%

HN 11

So, you know how bigtech gives nearly unlimited credits for coding tools and in the same time put a huge admin whitelist of what commands are allowed? So, instead of becoming a 100x engineers, you (including me) simply become a multi-window-enter-clicker I build a tool to solve...

<https://github.com/Alcray/SlopeAutoAcceptor>

CONF

Show HN: NeuroFlow 55.8x video inference speedup for Vision Transformers PyTorch

HackerNews

50%

HN 7

No summary — see link for full article.

<https://github.com/yinnk-research/-NeuroFlow>

CONF

Show HN: Harbor v0.4.19 - harbor launch -back end vLLM -web codex

HackerNews

50%

HN 4

No summary — see link for full article.

<https://github.com/av/harbor/releases/tag/v0.4.19>

CONF

DuckDuckGo installs are up 30% as users reject being ‘force-fed’ Google’s AI Search

TechCrunch AI 50%

Google overhauled Search at I/O 2026, replacing blue links with AI agents. The backlash has been swift. DuckDuckGo app installs spiked 30% as users seek a way out.

<https://techcrunch.com/2026/05/26/duckduckgo-installs-are-up-30-as-users-reject-being-f...>

CONF

This startup is betting India’s gig economy can train the world’s robots

TechCrunch AI 50%

Human Archive, a startup founded by UC Berkeley and Stanford researchers, is paying gig workers in India to wear camera-equipped caps and sensor devices to collect the real-world physical training data that AI and robotics labs are racing to acquire.

<https://techcrunch.com/2026/05/26/human-archive-taps-into-indias-services-startups-to-c...>

0 CONF

Universal Music Group and TikTok renew agreement to combat unauthorized AI music

TechCrunch AI 50%

For years, UMG has pushed platforms, streaming services, and AI companies to implement stricter content moderation policies.

<https://techcrunch.com/2026/05/26/universal-music-group-and-tiktok-renew-agreement-to-c...>

1 CONF

The pope’s AI encyclical isn’t really about AI

TechCrunch AI 50%

Pope Leo XIV’s first encyclical uses AI as a lens to diagnose older problems: concentrated power, eroding democracy, and a tech elite that shapes the world to its own advantage.

<https://techcrunch.com/2026/05/25/the-popes-ai-encyclical-isnt-really-about-ai/>

2 CONF

The Download: puncturing the AI jobs panic

MIT Tech Review 50%

This is today’s edition of The Download, our weekday newsletter that provides a daily dose of what’s going on in the world of technology. A reality check on the AI jobs hysteria Despite the growing hysteria over AI’s threat to white-collar jobs, there’s still scant evidence that...

<https://www.technologyreview.com/2026/05/26/1138028/the-download-ai-jobs-data/>

3 CONF

Sundar Pichai on AI, the future of search, and what’s happening to the web

The Verge AI 50%

Today, I’m talking with Google and Alphabet CEO Sundar Pichai, in a conversation we recorded just after the Google I/O developer conference. This is the fifth year Sundar and I have sat down after I/O, and it’s become one of my favorite Decoder traditions. There’s always a lot...

<https://www.theverge.com/podcast/936445/sundar-pichai-ai-search-google-zero-youtube-web>

4 CONF

AI warfare is already here

The Verge AI 50%

The Convention on Certain Conventional Weapons, an international forum that focuses on lethal autonomous systems, is hosted twice a year at the United Nations in Geneva. When Branka Marijan attended in November 2017, she thought the five-day sessions - which dealt largely in...

<https://www.theverge.com/ai-artificial-intelligence/937028/military-ai-warfare-red-lines>

5 CONF

Uber president says AI spending is getting 'harder to justify'

The Verge AI 50%

After reportedly exhausting its annual AI budget just four months into 2026, Uber is now questioning whether it's actually seeing meaningful returns on its investments. In an interview with Rapid Response, Uber president and chief operating officer Andrew Macdonald said the...

<https://www.theverge.com/transportation/937116/uber-ai-investment-hard-to-justify>

6 CONF

Pope Leo calls for being 'profoundly human' in the age of AI

The Verge AI 50%

Pope Leo XIV warned of the risks of AI and unconstrained technological power in his first major papal document released on Monday. Magnifica Humanitas is the pope's manifesto on "safeguarding the human person in the time of artificial intelligence," in which he discusses the...

<https://www.theverge.com/news/936945/pope-leo-letter-encyclical-ai-anthropic-labor-warfare>

7 CONF

Can LLMs Introspect? A Reality Check

ArXiv CS.AI 50%

arXiv:2605.26242v1 Announce Type: new Abstract: Can large language models detect and report their own internal states? A number of studies have argued that the answer to this question is yes. We argue, based on lessons from human metacognition research, that this conclusion may...

<https://arxiv.org/abs/2605.26242>

8 CONF

Constructing Industrial-Scale Optimization Modeling Benchmark

ArXiv CS.AI 50%

arXiv:2602.10450v2 Announce Type: replace-cross Abstract: Optimization modeling underpins decision-making in logistics, manufacturing, energy, and finance, yet translating natural-language requirements into correct optimization formulations and solver-executable code remains...

<https://arxiv.org/abs/2602.10450>

9 CONF

Experiments in Agentic AI for Science

ArXiv CS.AI 50%

arXiv:2605.26305v1 Announce Type: new Abstract: This paper details two novel frameworks for developing autonomous, agentic AI in scientific workflows. Both systems leverage a hybrid Local Body, Remote Brain architecture via Google Colab, utilizing Python-based local...

<https://arxiv.org/abs/2605.26305>

0 CONF

Anchor: Mitigating Artifact Drift in Agent Benchmark Generation

ArXiv CS.AI 50%

arXiv:2605.26321v1 Announce Type: new Abstract: AI agents are beginning to complete valuable, long-horizon business operations tasks, but training and evaluation environments for enterprise work still struggle to balance realism, verifiability, and scale. Environment and task...

<https://arxiv.org/abs/2605.26321>

1 CONF

PitchBench: Measuring Pitch Hearing in Audio-Language Models

ArXiv CS.AI 50%

arXiv:2605.26176v1 Announce Type: cross Abstract: Audio-language models (ALMs) are increasingly used in real-world applications that require understanding music, from music tutoring and transcription to captioning, recommendation systems, and music production. More broadly, they...

<https://arxiv.org/abs/2605.26176>

2 CONF

Managing Uncertainty in LLM-Generated Procedural Knowledge for Virtual Laboratory Planning

ArXiv CS.AI 50%

arXiv:2605.26333v1 Announce Type: new Abstract: Educational virtual laboratories can make experimental training more scalable, adaptive, and accessible, especially when students have limited access to physical laboratory facilities. However, authoring new simulated laboratory...

<https://arxiv.org/abs/2605.26333>

3 CONF

ScientistOne: Towards Human-Level Autonomous Research via Chain-of-Evidence

ArXiv CS.AI 50%

arXiv:2605.26340v1 Announce Type: new Abstract: Autonomous research agents produce competitive solutions and professional-looking manuscripts, yet their outputs contain verifiability failures undetectable by surface-level evaluation: fabricated citations, unreproducible scores,...

<https://arxiv.org/abs/2605.26340>

4 CONF

VisualNeedle: Benchmarking Active Visual Search in Information-Dense Scenes

ArXiv CS.AI 50%

arXiv:2605.26380v1 Announce Type: cross Abstract: Frontier multimodal large language models (MLLMs) have been reported to achieve over 90% accuracy on fine-grained perception benchmarks. However, such scores do not necessarily imply faithful use of visual evidence. Prior studies...

<https://arxiv.org/abs/2605.26380>

5 CONF

Which Changes Matter? Towards Trustworthy Legal AI via Relevance-Sensitive Evaluation and Solver-Grounded Reasoning

ArXiv CS.AI 50%

arXiv:2605.26530v1 Announce Type: new Abstract: Legal reasoning requires distinguishing changes that matter from those that do not. Legal AI should remain stable under legally irrelevant perturbations, but should change when perturbations alter legally material points. We...

<https://arxiv.org/abs/2605.26530>

6 CONF

PolyFusionAgent: A Multimodal Foundation Model and Autonomous AI Assistant for Polymer Property Prediction and Inverse Design

ArXiv CS.AI 50%

arXiv:2605.26543v1 Announce Type: new Abstract: Polymer discovery is central to fields ranging from energy storage to biomedicine, but it is hindered by an astronomically large chemical design space and fragmented representations of structure, properties, and prior knowledge....

<https://arxiv.org/abs/2605.26543>

7 CONF

MobileExplorer: Accelerating On-Device Inference for Mobile GUI Agents via Online Exploration

ArXiv CS.AI 50%

arXiv:2605.26546v1 Announce Type: new Abstract: Mobile graphical user interface (GUI) agents enable AI models to autonomously operate smartphones on behalf of users. However, most existing systems focus primarily on optimizing task accuracy and rely on cloud-hosted models for...

<https://arxiv.org/abs/2605.26546>

8 CONF

AGORA: Adapter-Grounded Observation-Action Retention for Inference-Free Prompt Compression in LLM Agents

ArXiv CS.AI 50%

arXiv:2605.26596v1 Announce Type: new Abstract: The token-level extractive compressors widely used for general LM context are structurally inappropriate for LLM agents: across 17 (env, backbone, method) cells spanning two independent token-level method families, every cell...

<https://arxiv.org/abs/2605.26596>

9 CONF

AMARIS: A Memory-Augmented Rubric Improvement System for Rubric-Based Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.18592v2 Announce Type: replace-cross Abstract: Rubric-based reward shaping provides interpretable and editable reward signals for fine-tuning LLMs via reinforcement learning (RL), but existing adaptive rubric methods typically update criteria from local evidence such...

<https://arxiv.org/abs/2605.18592>

0 CONF

Tail-Aware HiFloat4: W4A4 Post-Training Quantization for Wan2.2

ArXiv CS.AI 50%

arXiv:2605.26628v1 Announce Type: new Abstract: This report describes Tail-Aware HiFloat4, our submission to the low-bit text-to-video generation quantization challenge. Our method adapts the public ViDiT-Q post-training quantization pipeline to Wan2.2 under the HiFloat4...

<https://arxiv.org/abs/2605.26628>

1 CONF

The Necessity of a Unified Framework for LLM-Based Agent Evaluation

ArXiv CS.AI 50%

arXiv:2602.03238v2 Announce Type: replace Abstract: With the advent of Large Language Models (LLMs), general-purpose agents have seen fundamental advancements. However, evaluating these agents presents unique challenges that distinguish them from static QA benchmarks. We observe...

<https://arxiv.org/abs/2602.03238>

2 CONF

Detecting Is Not Resolving: The Monitoring Control Gap in Retrieval Augmented LLMs

ArXiv CS.AI 50%

arXiv:2605.27157v1 Announce Type: new Abstract: Retrieval-augmented LLMs are deployed for tasks where evidence quality determines action safety, yet evaluation protocols assume that single-turn robustness predicts robustness when evidence accumulates across turns. We show this...

<https://arxiv.org/abs/2605.27157>

3 CONF

Towards Feedback-to-Plan Decisions for Self-Evolving LLM Agents in CUDA Kernel Generation

ArXiv CS.AI 50%

arXiv:2605.26720v1 Announce Type: new Abstract: Large language models (LLMs) have shown strong empirical gains as self-evolving agents for CUDA kernel generation, driven by feedback-conditioned planning across generations. However, how planning decisions attribute and combine...

<https://arxiv.org/abs/2605.26720>

4 **CONF****It's Not the Capability: Harness Sensitivity Is Non-Monotone Across LLM Agent Tiers****ArXiv CS.AI** **50%**

arXiv:2605.26731v1 Announce Type: new Abstract: A prevalent assumption in LLM agent deployment holds that more structured harnesses universally improve reliability, and that higher-capability models need proportionally less structural guidance -- together implying a monotone...

<https://arxiv.org/abs/2605.26731>

5 **CONF****Beyond a Single Direction: Chain-of-Thought Disrupts Simple Steering of Refusal****ArXiv CS.AI** **50%**

arXiv:2605.26772v1 Announce Type: new Abstract: Large reasoning models (LRMs) generate chain-of-thought (CoT) traces before producing final outputs, introducing a dynamic internal state that may complicate control mechanisms such as refusal. Unlike instruction-tuned LLMs, where...

<https://arxiv.org/abs/2605.26772>

6 **CONF****The AI Cognitive Trojan Horse: How Large Language Models May Bypass Human Epistemic Vigilance****ArXiv CS.AI** **50%**

arXiv:2601.07085v2 Announce Type: replace-cross Abstract: Large language model (LLM)-based conversational AI systems present a challenge to human cognition that current frameworks for understanding misinformation and persuasion do not adequately address. This paper proposes that...

<https://arxiv.org/abs/2601.07085>

7 **CONF****LiveK12Bench: Have Large Multimodal Models Truly Conquered High School-level Examinations?****ArXiv CS.AI** **50%**

arXiv:2605.26781v1 Announce Type: new Abstract: Advanced Large Multimodal Models (LMMs) have demonstrated impressive performance in K-12 reasoning tasks, exhibiting great promise as intelligent tutors. Realizing this potential requires models to navigate real-world examinations...

<https://arxiv.org/abs/2605.26781>

8 **CONF****What Makes Chain-of-Thought Work at Probe Time? Local Co-occurrence Rather Than Global Derivation****ArXiv CS.AI** **50%**

arXiv:2605.26795v1 Announce Type: new Abstract: Chain-of-thought (CoT) prompting reliably improves language-model accuracy, but which properties of a rationale text drive the improvement is poorly understood. Prior work has largely studied generation-time behavior. We instead...

<https://arxiv.org/abs/2605.26795>

9 **CONF****Helicase: Uncertainty-Guided Supply Chain Knowledge Graph Construction with Autonomous Multi-Agent LLMs****ArXiv CS.AI** **50%**

arXiv:2605.26835v1 Announce Type: new Abstract: LLM-based multi-agent systems have been widely adopted for knowledge retrieval and report generation, synthesizing known information through web search and textual reasoning. However, many critical information tasks in supply...

<https://arxiv.org/abs/2605.26835>

0 CONF

Multi-Stakeholder LLM Alignment: Decomposing Estimation from Aggregation

ArXiv CS.AI 50%

arXiv:2605.26878v1 Announce Type: new Abstract: Multi-stakeholder tasks require one output to satisfy users with conflicting preferences. Holistic LLM judges conflate utility estimation and utility aggregation, yielding unstable implicit weights. We show empirically and...

<https://arxiv.org/abs/2605.26878>

1 CONF

TADDLE: A Tool-Augmented Agent for Detecting Deficient LLM-Generated Peer Reviews

ArXiv CS.AI 50%

arXiv:2605.26911v1 Announce Type: new Abstract: LLM-generated peer reviews are increasingly common at major venues, yet their deficiencies are hard to detect because they are uniformly fluent and well-structured. Existing work either classifies authorship without judging...

<https://arxiv.org/abs/2605.26911>

2 CONF

Where Code Meets Natural Language: Taxonomy-Driven Information Flow Analysis for LLM-Integrated Applications

ArXiv CS.AI 50%

arXiv:2603.28345v2 Announce Type: replace-cross Abstract: LLM API calls are becoming a ubiquitous program construct, yet they create a boundary that no existing program analysis can cross: runtime values enter a natural-language prompt, undergo opaque processing inside the LLM,...

<https://arxiv.org/abs/2603.28345>

3 CONF

Neuro-Symbolic Verification of LLM Outputs for Data-Sensitive Domains (extended preprint)

ArXiv CS.AI 50%

arXiv:2605.26942v1 Announce Type: new Abstract: LLMs deployed in high-stakes domains face fundamental reliability challenges: hallucinations, inconsistencies, and privacy vulnerabilities introduce unacceptable risks where errors carry legal, financial, or safety consequences....

<https://arxiv.org/abs/2605.26942>

4 CONF

LELA: An End-to-end LLM-based Entity Linking Framework with Zero-shot Domain Adaptation

ArXiv CS.AI 50%

arXiv:2605.26956v1 Announce Type: new Abstract: Entity linking is a key component of many downstream NLP systems, yet existing approaches are often tied to the specific target knowledge bases and domains, limiting their real world application. In this paper, we extend LELA, a...

<https://arxiv.org/abs/2605.26956>

5 CONF

Kandinsky 5.0: A Family of Foundation Models for Image and Video Generation

ArXiv CS.AI 50%

arXiv:2511.14993v3 Announce Type: replace-cross Abstract: This report introduces Kandinsky 5.0, a family of state-of-the-art foundation models for high-resolution image and 10-second video synthesis. The framework comprises three core line-up of models: Kandinsky 5.0 Image Lite...

<https://arxiv.org/abs/2511.14993>

6 CONF

Knowledge Graphs as the Missing Data Layer for LLM-Based Industrial Asset Operations

ArXiv CS.AI 50%

arXiv:2605.26874v1 Announce Type: cross Abstract: LLM-based agents for industrial asset operations show limited accuracy when reasoning over flat document stores. AssetOpsBench (KDD 2026) establishes that GPT-4 agents achieve 65% on 139 industrial maintenance scenarios backed by...

<https://arxiv.org/abs/2605.26874>

7 CONF

Counteraction-Aware Multi-Teacher On-Policy Distillation for General Capability Recovery with Domain Preservation

ArXiv CS.AI 50%

arXiv:2605.27115v1 Announce Type: new Abstract: Domain specialization can improve LLM behavior in vertical domains, but often weakens the general capabilities inherited from the original model. Recent Multi-Teacher On-Policy Distillation (MOPD) pipelines recover model...

<https://arxiv.org/abs/2605.27115>

8 CONF

SIA: Self Improving AI with Harness & Weight Updates

ArXiv CS.AI 50%

arXiv:2605.27276v1 Announce Type: new Abstract: Humans are the bottleneck in building and improving AI. Both the models and the agents that wrap them are written, tuned, and corrected by people. The long-horizon goal of an AI that can figure out how to improve itself remains...

<https://arxiv.org/abs/2605.27276>

9 CONF

2-ASP(Q) programs with weak constraints: Complexity and efficient implementation

ArXiv CS.AI 50%

arXiv:2605.27338v1 Announce Type: new Abstract: ASP(Q) extends Answer Set Programming (ASP) with Quantifiers over answer sets. In this paper we focus on the class of ASP(Q) programs with two quantifiers and weak constraints, denoted as 2-ASP(Q)^w . 2-ASP(Q)^w is a practically...

<https://arxiv.org/abs/2605.27338>

0 CONF

Access Timing as Scaffolding: A Reinforcement Learning Approach to GenAI in Education

ArXiv CS.AI 50%

arXiv:2605.15850v2 Announce Type: replace-cross Abstract: In recent years, generative AI (GenAI) in educational settings has become ubiquitous in university students' daily lives, despite its potential to induce over-reliance, metacognitive disengagement, and diminished learning...

<https://arxiv.org/abs/2605.15850>

1 CONF

Edge AI Deployment Beyond Models: A BSP-Aware Systems Framework for Industrial Embedded Platforms

ArXiv CS.AI 50%

arXiv:2605.26119v1 Announce Type: cross Abstract: Industrial Edge AI programs often begin with the model and only later confront the platform. That sequencing is attractive because it allows early demonstrations, but it breaks down when the deployment target is an embedded...

<https://arxiv.org/abs/2605.26119>

2 **CONF****GEM: Geometric Entropy Mixing for Optimal LLM Data Curation****ArXiv CS.AI** **50%**

arXiv:2605.26121v1 Announce Type: cross Abstract: LLM pre-training efficacy increasingly depends on data composition rather than sheer volume. Yet, optimal mixing is hindered by categorization flaws: human taxonomies suffer from ontological misalignment, and Euclidean clustering...

<https://arxiv.org/abs/2605.26121>3 **CONF****Augment Engineering: A Methodology for Multi-Tool AI Orchestration Across Professional Domains****ArXiv CS.AI** **50%**

arXiv:2605.26146v1 Announce Type: cross Abstract: Organizations increasingly deploy separate purpose-built AI tools across professional domains, often hiring domain specialists for each, recreating the staffing models AI was expected to transform. Yet the meta-skills that make...

<https://arxiv.org/abs/2605.26146>4 **CONF****MemMorph: Tool Hijacking in LLM Agents via Memory Poisoning****ArXiv CS.AI** **50%**

arXiv:2605.26154v1 Announce Type: cross Abstract: LLM-driven agents are capable of selecting external tools to complete users' tasks. However, attackers could compromise such process, steering agents toward inappropriate/wrong tools and enabling malicious actions. Most existing...

<https://arxiv.org/abs/2605.26154>5 **CONF****Turning Bias into Bugs: Bandit-Guided Style Manipulation Attacks on LLM Judges****ArXiv CS.AI** **50%**

arXiv:2605.26156v1 Announce Type: cross Abstract: The known stylistic biases in LLM judges, such as a preference for verbosity or specific sentence structures, present an underexplored security vulnerability. In this work, we introduce BITE (Bias exploration and Exploitation), a...

<https://arxiv.org/abs/2605.26156>6 **CONF****Furina: Fragmented Uncertainty-Driven Refusal Instability Attack****ArXiv CS.AI** **50%**

arXiv:2605.26158v1 Announce Type: cross Abstract: Safety alignment in large language models (LLMs) and multimodal large language models (MLLMs) is commonly assumed to operate as a near-binary threshold mechanism. We challenge this assumption by revealing that safety behavior is...

<https://arxiv.org/abs/2605.26158>7 **CONF****LUCoS: Latent Unsupervised Context Selection for Tabular Foundation Models****ArXiv CS.AI** **50%**

arXiv:2605.27254v1 Announce Type: cross Abstract: Selecting which instances to label is a key challenge in low-label tabular learning. For recent Tabular Foundation Models such as TabPFN, context selection directly determines predictive performance. Supervised oracle experiments...

<https://arxiv.org/abs/2605.27254>

8 CONF

Tool-Schema Compression Enables Agentic RAG Under Constrained Context Budgets

ArXiv CS.AI 50%

arXiv:2605.26165v1 Announce Type: cross Abstract: Agentic RAG systems that equip language models with dozens to hundreds of tool definitions face a critical resource conflict: tool schemas consume the same context window needed for retrieval-augmented generation. We present the...

<https://arxiv.org/abs/2605.26165>

9 CONF

Planning Neural Dynamics with Lie Group Embedding through Supervised Projective Manifold Learning

ArXiv CS.AI 50%

arXiv:2605.26167v1 Announce Type: cross Abstract: We propose Lie group embedded dynamical neural networks (LieEDNN) and the corresponding learning algorithms based on gradient descent and metric projection on smooth manifold, where we treat Lie group as an intrinsic...

<https://arxiv.org/abs/2605.26167>

0 CONF

A Universal Cliff and a Design Fingerprint: Cross-Section Defect Detection Under LLM Orchestration

ArXiv CS.AI 50%

arXiv:2605.26174v1 Announce Type: cross Abstract: Production language-model systems answer a request by partitioning it across an invisible orchestration of worker agents that recompose one integrated report. We ask what this does to a class of defect no single worker can see: a...

<https://arxiv.org/abs/2605.26174>

1 CONF

InfoQuant: Shaping Activation Distributions for Low-Bit LLM Quantization

ArXiv CS.AI 50%

arXiv:2605.26175v1 Announce Type: cross Abstract: Low-bit activation quantization remains a major bottleneck in efficient large language model (LLM) deployment. The difficulty is not only that activations contain outliers, but that their distributions are often poorly matched to...

<https://arxiv.org/abs/2605.26175>

2 CONF

Structure-Adaptive Conformal Inference for Large-Scale Out-of-Distribution Testing

ArXiv CS.AI 50%

arXiv:2605.26429v1 Announce Type: cross Abstract: This paper addresses structured out-of-distribution (OOD) testing in high-stakes machine learning applications. Traditional conformal methods rely on joint exchangeability, making it difficult to incorporate auxiliary information...

<https://arxiv.org/abs/2605.26429>

3 CONF

SetupX: Can LLM Agents Learn from Past Failures in Functionality-Correct Code Repository Setup?

ArXiv CS.AI 50%

arXiv:2605.26186v1 Announce Type: cross Abstract: Functionality-correct repository setup aims to configure execution environments (e.g., dependencies, build scripts) to successfully execute a repository's documented features. It presents significant challenges due to diverse,...

<https://arxiv.org/abs/2605.26186>

4 **CONF****Max-Window Scale Estimation for Near-Lossless HiF8 W8A8 Quantization-Aware Training****ArXiv CS.AI** **50%**

arXiv:2605.26189v1 Announce Type: cross Abstract: Quantization-aware training (QAT) with low-bit floating-point formats enables efficient LLM deployment, yet introduces subtle failure modes invisible to standard training metrics. We present a systematic study of HiF8 W8A8 QAT...

<https://arxiv.org/abs/2605.26189>

5 **CONF****Prospective evaluation of multimodal respiratory failure prediction: Do chest X-rays improve performance beyond EHR signals?****ArXiv CS.AI** **50%**

arXiv:2605.26255v1 Announce Type: cross Abstract: Early prediction of respiratory failure is critical for timely clinical intervention in intensive care units. Existing electronic health record (EHR)-based models can continuously monitor physiologic deterioration, but they may...

<https://arxiv.org/abs/2605.26255>

6 **CONF****AI evaluation may bias perceptions: The importance of context in interpreting academic writing****ArXiv CS.AI** **50%**

arXiv:2605.26662v1 Announce Type: cross Abstract: This paper examines how estimates of AI use in scientific writing can be biased when evaluation methods ignore contextual differences across countries and fields. Using large-scale data on journal publications from Dimensions, we...

<https://arxiv.org/abs/2605.26662>

7 **CONF****Decoupled Delay Compensation: Enhancing Pre-trained MARL Policies via Learned Dynamics Filtering****ArXiv CS.AI** **50%**

arXiv:2605.26286v1 Announce Type: cross Abstract: Real-world multi-agent reinforcement learning (MARL) systems must often operate under stale observations, stochastic communication delays, and intermittent packet loss. Policies trained under idealized synchronous conditions...

<https://arxiv.org/abs/2605.26286>

8 **CONF****Curriculum Learning for Safety Alignment****ArXiv CS.AI** **50%**

arXiv:2605.26315v1 Announce Type: cross Abstract: Direct Preference Optimisation (DPO) is widely used for safety alignment in large language models. However, prior work shows it is brittle and exhibits poor out-of-distribution (OOD) generalisation. In this paper, we investigate...

<https://arxiv.org/abs/2605.26315>

9 **CONF****Erased but Exploitable: Black-box Embedding-Aware Prompting Against Unlearned Text-to-Image Diffusion Models****ArXiv CS.AI** **50%**

arXiv:2605.26332v1 Announce Type: cross Abstract: Machine unlearning aims to remove specific concepts from pretrained text-to-image diffusion models, yet several white- and black-box attacks have been introduced to make the model generate such unlearned concepts. These attacks,...

<https://arxiv.org/abs/2605.26332>

- 0 **CONF**
- Why LLMs Hallucinate on Structured Knowledge: A Mechanistic Analysis of Reasoning over Linearized Representations**
- ArXiv CS.AI 50%
- arXiv:2605.26362v1 Announce Type: cross Abstract: In many reasoning tasks, large language models (LLMs) rely on structured external knowledge, such as graphs and tables, which is typically linearized into sequential token representations. However, even when sufficient knowledge...
- <https://arxiv.org/abs/2605.26362>
-
- 1 **CONF**
- Unified Panoramic Geometry Estimation via Multi-View Foundation Models**
- ArXiv CS.AI 50%
- arXiv:2605.26368v1 Announce Type: cross Abstract: Geometry estimation from perspective images has greatly advanced, maturing to the point where off-the-shelf foundation models are able to reconstruct 3D scene structure not only from multi-view imagery, but even from a single...
- <https://arxiv.org/abs/2605.26368>
-
- 2 **CONF**
- Plans for Evaluating Structured Generative Search Summaries**
- ArXiv CS.AI 50%
- arXiv:2605.26400v1 Announce Type: cross Abstract: We propose a framework for evaluating structured generative search summaries that are placed atop organic web search results. A structured summary, generated by a large language model, typically consists of an overview, several...
- <https://arxiv.org/abs/2605.26400>
-
- 3 **CONF**
- Jailbreak susceptibility prediction and mitigation via the behavioral geometry of models**
- ArXiv CS.AI 50%
- arXiv:2605.26409v1 Announce Type: cross Abstract: Evaluating and mitigating a generative system's susceptibility to jailbreak attacks is critical to its safe deployment. Given the number of deployable systems, full per-configuration evaluation and optimization is impractical. In...
- <https://arxiv.org/abs/2605.26409>
-
- 4 **CONF**
- When Does Deep RL Beat Calibrated Baselines? A Benchmark Study on Adaptive Resource Control**
- ArXiv CS.AI 50%
- arXiv:2605.26418v1 Announce Type: cross Abstract: A properly calibrated rule-based autoscaler can beat every one of six mainstream deep reinforcement learning (DRL) algorithms on cost across every workload we test - so when, if ever, does DRL actually help? We study this in...
- <https://arxiv.org/abs/2605.26418>
-
- 5 **CONF**
- Uniboost: Global Coordination with Value Alignment for Fair and Efficient Traffic Allocation**
- ArXiv CS.AI 50%
- arXiv:2605.26424v1 Announce Type: cross Abstract: With the rapid evolution of internet services, recommendation systems have become indispensable. In particular, the blending (re-ranking) stage plays a pivotal role in allocating traffic across diverse business objectives....
- <https://arxiv.org/abs/2605.26424>

6 CONF

Athena: Enhancing Multimodal Reasoning with Data-efficient Process Reward Models

ArXiv CS.AI 50%

arXiv:2506.09532v5 Announce Type: replace-cross Abstract: We present Athena-PRM, a multimodal process reward model (PRM) designed to evaluate the reward score for each step in solving complex reasoning problems. Developing high-performance PRMs typically demands significant time...

<https://arxiv.org/abs/2506.09532>

7 CONF

AnchorDiff: Training-Free Concept Grounding for MM-DiTs via Anchor-Based Graph Propagation

ArXiv CS.AI 50%

arXiv:2605.26460v1 Announce Type: cross Abstract: Multi-Modal Diffusion Transformers (MM-DiTs) encode rich representations for training-free concept grounding, but existing attention-based methods often produce overlapping activations on visually confusable concepts, a failure...

<https://arxiv.org/abs/2605.26460>

8 CONF

Cross-scale Aligned Supervision for Training GANs

ArXiv CS.AI 50%

arXiv:2605.26449v1 Announce Type: cross Abstract: Modern GANs often introduce adversarial supervision on intermediate generator outputs and interpret the resulting multi-stage synthesis as coarse-to-fine hierarchical generation. In this work, we challenge this interpretation. We...

<https://arxiv.org/abs/2605.26449>

9 CONF

Elias in the Lighthouse, Again? Diagnosing Low Diversity in LLM Stories

ArXiv CS.AI 50%

arXiv:2605.26492v1 Announce Type: cross Abstract: LLM-generated stories are a popular use case, but they show very low variability. We sample 20,000 total stories from four current models using five prompts. We find that 11 words occur in 88.3% of generated stories, with little...

<https://arxiv.org/abs/2605.26492>

0 CONF

Dense2MoE: Pushing the Pareto Frontier of On-Device LLMs via Unified Pruning and Upcycling

ArXiv CS.AI 50%

arXiv:2605.26496v1 Announce Type: cross Abstract: The Mixture of Experts MoE architecture is highly promising for resource constrained on device deployments yet training these models from scratch incurs prohibitive costs Current methods attempt to alleviate this by upcycling...

<https://arxiv.org/abs/2605.26496>

1 CONF

Unveiling the Fragility of Vision-Language Models: Multi-Modal Adversarial Synergy via Texture-Constrained Perturbations and Cross-Modal Optimization

ArXiv CS.AI 50%

arXiv:2605.26501v1 Announce Type: cross Abstract: Large Vision-Language Models (LVLMs) have transformed multi-modal understanding, excelling in tasks like image captioning and visual question answering by integrating visual and textual inputs. However, their robustness against...

<https://arxiv.org/abs/2605.26501>

2 **CONF****Foundations of a Time-Consistent Counterfactual Actuarial Runtime for Autonomous AI Agents****ArXiv CS.AI** **50%**

arXiv:2605.26508v1 Announce Type: cross Abstract: We propose a foundational runtime actuarial layer for autonomous AI agents in which every side-effect-bearing action carries a time-consistent, counterfactual risk toll computed against a contractually fixed safe default, inside...

<https://arxiv.org/abs/2605.26508>

3 **CONF****CSV-ViT: A Vision Transformer with the Variable-sized Cortical Supervertices for Detection of Alzheimer's Disease Pathologies****ArXiv CS.AI** **50%**

arXiv:2605.26514v1 Announce Type: cross Abstract: Confirming Alzheimer's disease (AD) typically relies on positron emission tomography (PET), which remains costly and invasive, motivating the use of structural MRI-based prescreening. Deep learning on non-Euclidean manifolds,...

<https://arxiv.org/abs/2605.26514>

4 **CONF****StreamSplit: Continuous Audio Representation Learning via Uncertainty-Guided Adaptive Splitting****ArXiv CS.AI** **50%**

arXiv:2605.26523v1 Announce Type: cross Abstract: Large-batch Contrastive Learning (CL), the foundation of modern representation learning, is fundamentally incompatible with the volatile resource constraints of edge devices. This conflict creates a dilemma: small on-device...

<https://arxiv.org/abs/2605.26523>

5 **CONF****DGLD: Domain-Gated Latent Diffusion for the Discovery of Novel Energetic Materials****ArXiv CS.AI** **50%**

arXiv:2605.26540v1 Announce Type: cross Abstract: Energetic-materials performance gains translate directly into reduced propellant mass, smaller warheads, and more efficient civilian gas-generators, yet no new HMX-class compound has been disclosed in fifteen years. Designing one...

<https://arxiv.org/abs/2605.26540>

6 **CONF****ChainCaps: Composition-Safe Tool-Using Agents via Monotonic Capability Attenuation****ArXiv CS.AI** **50%**

arXiv:2605.26542v1 Announce Type: cross Abstract: Tool-using agents increasingly operate in open-ended deployment environments, where they compose file systems, web APIs, code interpreters, and enterprise services at runtime. This creates a safety gap in tool composition: an...

<https://arxiv.org/abs/2605.26542>

7 **CONF****Aligning Few-Step Generative Models by Amortizing Sample-based Variational Inference****ArXiv CS.AI** **50%**

arXiv:2605.26552v1 Announce Type: cross Abstract: Aligning a few-step generative model is challenging, since existing alignment frameworks typically rely on restrictive assumptions: a tractable likelihood, a specific ODE/SDE solver, or a particular model family. We introduce...

<https://arxiv.org/abs/2605.26552>

8 CONF

Bridging Control with Neural Network Verifier alpha-beta-CROWN: A Tutorial

ArXiv CS.AI 50%

arXiv:2605.26577v1 Announce Type: cross Abstract: Learning-based methods for synthesizing controllers have gained popularity due to their high expressiveness and strong empirical performance. However, in safety-critical scenarios such as autonomous driving, robotics, and power...

<https://arxiv.org/abs/2605.26577>

9 CONF

On the Error-Correcting Effects of Stochasticity in Discrete Diffusion

ArXiv CS.AI 50%

arXiv:2605.26582v1 Announce Type: cross Abstract: Discrete diffusion models achieve strong performance in text and image generation, but their inference remains slow and must inherently balance sampling efficiency and sample quality. In this work, we present a systematic study...

<https://arxiv.org/abs/2605.26582>

0 CONF

Few-shot Cross-country Generalization of Tabular Machine Learning and Foundation Models for Childhood Anemia Prediction under Distribution Shift

ArXiv CS.AI 50%

arXiv:2605.26589v1 Announce Type: cross Abstract: Childhood anemia affects around 40% of children aged 6-59 months globally and arises from heterogeneous factors, limiting model generalizability. We evaluate a transformer-based tabular foundation model against classical...

<https://arxiv.org/abs/2605.26589>

1 CONF

JetViT: Efficient High-Resolution Vision Transformer with Post-Training Attention Search

ArXiv CS.AI 50%

arXiv:2605.26636v1 Announce Type: cross Abstract: We introduce JetViT, a novel family of hybrid-architecture Vision Transformer (ViT) models that match the accuracy of state-of-the-art full-attention vision foundation models while achieving substantially higher inference...

<https://arxiv.org/abs/2605.26636>

2 CONF

Negligible in Size, Significant in Effect: On Scale Vectors in Large Language Models

ArXiv CS.AI 50%

arXiv:2605.26895v1 Announce Type: cross Abstract: Normalization layers in modern large language models (LLMs) consist of a deterministic normalization operation and a learnable scale vector. While the normalization operation has been extensively studied, the scale vector remains...

<https://arxiv.org/abs/2605.26895>

3 CONF

The Labyrinth and the Thread: Rethinking Regularizations in Sequential Knowledge Editing for Large Language Models

ArXiv CS.AI 50%

arXiv:2605.26670v1 Announce Type: cross Abstract: Sequential editing of structured knowledge in large language models allows targeted factual updates without retraining, yet existing methods often rely on complex regularization or constraint mechanisms whose necessity remains...

<https://arxiv.org/abs/2605.26670>

4 **CONF**

DynFrame: Adaptive Reasoning-Driven Multimodal Framework with Dynamic Frame Augmentation for Complex Video Understanding

ArXiv CS.AI 50%

arXiv:2605.26680v1 Announce Type: cross Abstract: Recent video multimodal large language models (MLLMs) increasingly couple step-by-step reasoning with on-demand visual evidence retrieval, allowing models to revisit relevant video segments during inference. However, two...

<https://arxiv.org/abs/2605.26680>

5 **CONF**

Multi-Agent Causal Discovery Using Large Language Models

ArXiv CS.AI 50%

arXiv:2407.15073v4 Announce Type: replace Abstract: Causal discovery aims to identify causal relationships between variables and is a fundamental problem across the sciences. Traditional statistical causal discovery (SCD) methods rely solely on observational data and ignore the...

<https://arxiv.org/abs/2407.15073>

6 **CONF**

L2Rec: Towards Dual-View Understanding of LLMs for Personalized Recommendation

ArXiv CS.AI 50%

arXiv:2605.26717v1 Announce Type: cross Abstract: Adapting large language models (LLMs) for personalized recommendation requires aligning their general-purpose capabilities with user-specific preferences while effectively leveraging both behavioral and semantic signals. Existing...

<https://arxiv.org/abs/2605.26717>

7 **CONF**

Stabilizing Recurrent Dynamics for Test-Time Scalable Latent Reasoning in Looped Language Models

ArXiv CS.AI 50%

arXiv:2605.26733v1 Announce Type: cross Abstract: Looped Language Models (LoopLMs) enable efficient latent reasoning through depth recurrence, yet exhibit unreliable test-time scaling behavior: performance often peaks at a certain iteration depth and then collapses with further...

<https://arxiv.org/abs/2605.26733>

8 **CONF**

Cordon-MAS: Defending RAG against Knowledge Poisoning via Information-Flow Control

ArXiv CS.AI 50%

arXiv:2605.26754v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) increasingly underpins high-stakes applications, yet remains vulnerable to Confundo-style poisoning where adversarially optimized documents manipulate generated outputs. Existing defenses...

<https://arxiv.org/abs/2605.26754>

9 **CONF**

Adversarial Training for Robust Coverage Network under Worst-case Facility Losses

ArXiv CS.AI 50%

arXiv:2605.26763v1 Announce Type: cross Abstract: The Maximal Covering Location-Interdiction Problem (MCLIP) is a classic bi-level optimization problem, which is fundamental to resilient infrastructure planning yet remains computationally intractable. Specifically, the upper...

<https://arxiv.org/abs/2605.26763>

00 CONF

SeDT: Sentence-Transformer Decision-Transformer Conditioning for Multi-Turn Conversation Reliability

ArXiv CS.AI 50%

arXiv:2605.26788v1 Announce Type: cross Abstract: Large language models (LLMs) achieve impressive performance when a task is fully specified in a single turn, yet the same models lose up to 39% of that performance when the identical task is revealed incrementally across multiple...

<https://arxiv.org/abs/2605.26788>

01 CONF

HTMLCure: Turning Browser Experience into State Guided Repair for Interactive HTML

ArXiv CS.AI 50%

arXiv:2605.26807v1 Announce Type: cross Abstract: LLMs can now produce full HTML pages, but many of those pages are only superficially correct: they render once, then fail under scroll, hover, click, resize, or gameplay. Evaluation from screenshots can miss these failures, and...

<https://arxiv.org/abs/2605.26807>

02 CONF

RAGEAR: Retrieval-Augmented Graph-Enhanced Academic Recommender

ArXiv CS.AI 50%

arXiv:2605.26819v1 Announce Type: cross Abstract: We present RAGEAR (Retrieval-Augmented Graph-Enhanced Academic Recommender), a neurosymbolic recommender system for academic course recommendation. RAGEAR combines dense retrieval over full lecture transcripts with a symbolic...

<https://arxiv.org/abs/2605.26819>

03 CONF

Periodic Topological Deep Learning for Polymer Design and Discovery

ArXiv CS.AI 50%

arXiv:2605.26833v1 Announce Type: cross Abstract: Polymers underpin applications across energy, healthcare, and materials science, yet their vast chemical space makes systematic discovery challenging. Most machine learning approaches represent polymers as molecular graphs of a...

<https://arxiv.org/abs/2605.26833>

04 CONF

Persistent AI Agents in Academic Research: A Single-Investigator Implementation Case Study

ArXiv CS.AI 50%

arXiv:2605.26870v1 Announce Type: cross Abstract: Background: Large language models are typically evaluated as models, benchmarks, or short conversational episodes. Less is known about what happens when an agent is embedded persistently in a real academic research environment...

<https://arxiv.org/abs/2605.26870>

05 CONF

GeoFaith: A Spatio-Temporal Dual View of Faithful Chain-of-Thought

ArXiv CS.AI 50%

arXiv:2605.26893v1 Announce Type: cross Abstract: Chain-of-Thought (CoT) reasoning has advanced large language models (LLMs), but outcome-based supervision leads to pervasive post-hoc rationalization, producing plausible yet unfaithful reasoning chains. Most prior faithfulness...

<https://arxiv.org/abs/2605.26893>

06 CONF

Strategies for Guiding LLMs to Use Software Design Patterns: A Case of Singleton

ArXiv CS.AI 50%

arXiv:2605.26898v1 Announce Type: cross Abstract: Large Language Models (LLMs) can generate functional source code from natural-language prompts, but often fail to consistently follow higher-level architectural structures or design patterns. Since LLMs are increasingly used in...

<https://arxiv.org/abs/2605.26898>

07 CONF

JuICE: A Benchmark for Evaluating LLM-Judge in Identifying Cultural Errors

ArXiv CS.AI 50%

arXiv:2605.26955v1 Announce Type: cross Abstract: As large language models (LLMs) are increasingly deployed to users around the world, they are integrated into everyday tasks across diverse cultural contexts, from drafting personal communications to brainstorming creative ideas....

<https://arxiv.org/abs/2605.26955>

08 CONF

Position: The Hidden Costs and Measurement Gaps of Reinforcement Learning with Verifiable Rewards

ArXiv CS.AI 50%

arXiv:2509.21882v3 Announce Type: replace-cross Abstract: Reinforcement learning with verifiable rewards (RLVR) is a practical, scalable way to improve large language models on math, code, and other structured tasks. However, we argue that many headline RLVR gains are not yet...

<https://arxiv.org/abs/2509.21882>

09 CONF

ReasonOps: A Unified Operational Paradigm for Trustworthy Verified LLM Reasoning

ArXiv CS.AI 50%

arXiv:2605.27014v1 Announce Type: cross Abstract: Large Language Models (LLMs) have transformed artificial intelligence from primarily generative systems into increasingly capable reasoning agents. Recent advances in theorem proving, autoformalization, symbolic reasoning, and...

<https://arxiv.org/abs/2605.27014>

10 CONF

Assessing Per-Sample Membership Inference Vulnerability without Retraining

ArXiv CS.AI 50%

arXiv:2602.15919v2 Announce Type: replace-cross Abstract: Recent work in the privacy literature shows that sample-targeted membership inference attacks (MIAs) significantly outperform untargeted approaches by a wide margin. Motivated by this observation, we address the following...

<https://arxiv.org/abs/2602.15919>

11 CONF

Tracing Computation Density in LLMs

ArXiv CS.AI 50%

arXiv:2605.27033v1 Announce Type: cross Abstract: Transformer-based large language models (LLMs) are comprised of billions of parameters arranged in deep and wide computational graphs, but it is not clear that they exploit their full capacity for all inputs. We introduce the...

<https://arxiv.org/abs/2605.27033>

12 CONF

QUACK: Questioning, Understanding, and Auditing Communicated Knowledge in Multimodal Social Deduction Agents

ArXiv CS.AI 50%

arXiv:2605.27068v1 Announce Type: cross Abstract: Social deduction games have become a popular testbed for probing reasoning, deception, coordination, and belief modeling in Large Language Model (LLM) agents. However, most environments are scored only by game outcomes such as...

<https://arxiv.org/abs/2605.27068>

13 CONF

ReMoE: Boosting Expert Reuse through Router Fine-Tuning in Memory-Constrained MoE LLM Inference

ArXiv CS.AI 50%

arXiv:2605.27081v1 Announce Type: cross Abstract: Fine-grained Mixture-of-Experts (MoE) models sparsely activate only a subset of experts per token, reducing activated computation while maintaining high model capacity. However, in memory-constrained inference scenarios, only a...

<https://arxiv.org/abs/2605.27081>

14 CONF

MiRD: Reliable Set-Valued Prediction for Open-Ended Question Answering via Miscoverage Risk Decomposition

ArXiv CS.AI 50%

arXiv:2605.27091v1 Announce Type: cross Abstract: Reliable set-valued prediction provides a principled way to mitigate hallucinations in open-ended question answering (QA), yet existing conformal approaches typically rely on a fragile premise: finite sampling must already...

<https://arxiv.org/abs/2605.27091>

15 CONF

High-Quality Synthetic Financial Time-Series using a GAN-Diffusion Framework

ArXiv CS.AI 50%

arXiv:2605.27113v1 Announce Type: cross Abstract: In recent years, financial institutions and firms have increasingly adopted synthetic data to address data scarcity and to generate counterfactual market scenarios. However, reproducing all the statistical properties of financial...

<https://arxiv.org/abs/2605.27113>

16 CONF

DEI: Diversity in Evolutionary Inference for Quality-Diversity Search

ArXiv CS.AI 50%

arXiv:2605.27130v1 Announce Type: cross Abstract: We present DEI: Diversity in Evolutionary Inference, a distributed Quality-Diversity (QD) search framework that assigns heterogeneous large language models (LLMs) as mutation operators across peer nodes communicating with...

<https://arxiv.org/abs/2605.27130>

17 CONF

Beyond the Data Mesh Illusion: Designing Modern AI-augmented Lakehouses to Bridge the Gap Between Theory and Practice

ArXiv CS.AI 50%

arXiv:2605.27131v1 Announce Type: cross Abstract: Enterprise data platforms face an enduring tension between domain self-service and holistic governance. The data mesh paradigm proposed decentralized domain ownership as a remedy, but pure implementations frequently underdeliver:...

<https://arxiv.org/abs/2605.27131>

18 CONF

Semantic Robustness Probing via Inpainting: An Interactive Tool for Safety-Critical Object Detection

ArXiv CS.AI 50%

arXiv:2605.27155v1 Announce Type: cross Abstract: Testing object detectors in safety-critical domains requires semantically meaningful probes beyond pixel-level corruptions. We present SemProbe, a tool for semantic robustness probing: users upload deployment images, create masks...

<https://arxiv.org/abs/2605.27155>

19 CONF

LitSeg: Narrative-Aware Document Segmentation for Literary RAG

ArXiv CS.AI 50%

arXiv:2605.27156v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) enhances Large Language Models (LLMs) by incorporating external knowledge, particularly for long-tail domains such as literary works. However, the critical step of document segmentation in RAG...

<https://arxiv.org/abs/2605.27156>

20 CONF

Grounding Text Embeddings in Stakeholder Associations

ArXiv CS.AI 50%

arXiv:2605.27168v1 Announce Type: cross Abstract: Text embeddings are widely used to analyse large corpora of complex texts. However, it is unclear whether the embeddings capture the same semantic distances as the human experts using them. Ensuring alignment between embedding...

<https://arxiv.org/abs/2605.27168>

21 CONF

An investigation of AI integration in sound designer workflows and experiences

ArXiv CS.AI 50%

arXiv:2605.27174v1 Announce Type: cross Abstract: Artificial intelligence is increasingly being integrated into professional audio production workflows, yet a gap persists between the tools developers produce and the requirements of practising sound designers. This paper...

<https://arxiv.org/abs/2605.27174>

22 CONF

Generative Animations: A Multi-Model Pipeline for Prompt-Driven Motion Synthesis

ArXiv CS.AI 50%

arXiv:2605.27203v1 Announce Type: cross Abstract: Animation elevates digital documents into immersive experiences, yet creating custom motion paths remains cumbersome, requiring designers to manually select presets, plot Bezier points, and configure timing properties. We...

<https://arxiv.org/abs/2605.27203>

23 CONF

Lost in Sampling: Assessing Lexical Reachability in LLMs via the Word Coverage Score (WCS)

ArXiv CS.AI 50%

arXiv:2605.27268v1 Announce Type: cross Abstract: Modern Large Language Models (LLMs) are often criticized for producing repetitive and homogeneous text, despite possessing vast latent vocabularies. While previous research has focused on model knowledge and training data, we...

<https://arxiv.org/abs/2605.27268>

24 CONF

PICACO: Pluralistic In-Context Value Alignment of LLMs via Total Correlation Optimization

ArXiv CS.AI 50%

arXiv:2507.16679v3 Announce Type: replace-cross Abstract: In-Context Learning has shown great potential for aligning Large Language Models (LLMs) with human values, helping reduce harmful outputs and accommodate diverse preferences without costly post-training, known as...

<https://arxiv.org/abs/2507.16679>

25 CONF

It's Not Always Sycophancy: Measuring LLM Conformity as a Function of Epistemic Uncertainty

ArXiv CS.AI 50%

arXiv:2605.27288v1 Announce Type: cross Abstract: Large language models (LLMs) are known to abandon their initial stance to conform to user pushback. While prior research largely attributes this behavior to sycophancy learned during reinforcement learning from human feedback, we...

<https://arxiv.org/abs/2605.27288>

26 CONF

When Eyes Betray AI: Social Gaze Consistency as a Semantic Cue for AI-Generated Image Detection

ArXiv CS.AI 50%

arXiv:2605.27348v1 Announce Type: cross Abstract: Recent generative models have largely closed the gap on low-level artifacts - pixel fingerprints, frequency anomalies, upsampling traces - particularly in person-centric and partial-edit settings where the manipulated region is...

<https://arxiv.org/abs/2605.27348>

27 CONF

Guiding LLM Post-training Data Engineering with Model Internals from Sparse Autoencoders

ArXiv CS.AI 50%

arXiv:2605.27354v1 Announce Type: cross Abstract: Model internals encode rich information about how a large language model (LLM) processes its training data; however, post-training data engineering largely relies on external signals and ignores rich intrinsic signals lying in...

<https://arxiv.org/abs/2605.27354>

28 CONF

GENESIS: Harnessing AI Agents for Autonomous 6G RAN Synthesis, Research, and Testing

ArXiv CS.AI 50%

arXiv:2605.27360v1 Announce Type: cross Abstract: Cellular research and development (R&D) is throttled by six structural processes that each consume months of manual engineering work per iteration: (i) synthesizing new features from standards or research papers into production...

<https://arxiv.org/abs/2605.27360>

29 CONF

Doc-CoB: Enhancing Document Understanding with Visual Chain-of-Boxes Reasoning

ArXiv CS.AI 50%

arXiv:2505.18603v2 Announce Type: replace Abstract: Document understanding aims to perform question answering and information extraction over document images, where the visual content is highly information-dense and most queries rely on only a few relevant layout regions....

<https://arxiv.org/abs/2505.18603>

30 CONF

Plan Then Action: High-Level Planning Guidance Reinforcement Learning for LLM Reasoning

ArXiv CS.AI 50%

arXiv:2510.01833v2 Announce Type: replace Abstract: Large language models (LLMs) demonstrate strong reasoning abilities via Chain-of-Thought (CoT), but their token-level generation encourages local decisions and lacks global planning, often leading to redundant or inaccurate...

<https://arxiv.org/abs/2510.01833>

31 CONF

AI-Driven Contribution Evaluation and Conflict Resolution: A Framework & Design for Group Workload Investigation

ArXiv CS.AI 50%

arXiv:2511.07667v2 Announce Type: replace Abstract: The equitable assessment of individual contribution in teams remains a persistent challenge, where conflict and disparity in workload can result in unfair performance evaluation, often requiring manual intervention - a costly...

<https://arxiv.org/abs/2511.07667>

32 CONF

PaTAS: A Framework for Trust Propagation in Neural Networks Using Subjective Logic

ArXiv CS.AI 50%

arXiv:2511.20586v4 Announce Type: replace Abstract: Trustworthiness has become a key requirement for the deployment of artificial intelligence systems in safety-critical applications. Conventional evaluation metrics, such as accuracy and precision, fail to appropriately capture...

<https://arxiv.org/abs/2511.20586>

33 CONF

LEC: Linear Expectation Constraints for Selection-Conditioned Risk Control in Selective Prediction and Routing Systems

ArXiv CS.AI 50%

arXiv:2512.01556v3 Announce Type: replace Abstract: Foundation models often generate unreliable answers, while heuristic uncertainty estimators fail to fully distinguish correct from incorrect outputs, causing users to accept erroneous answers without any statistical guarantee....

<https://arxiv.org/abs/2512.01556>

34 CONF

TowerMind: A Tower Defence Game Learning Environment and Benchmark for LLM as Agents

ArXiv CS.AI 50%

arXiv:2601.05899v2 Announce Type: replace Abstract: Recent breakthroughs in Large Language Models (LLMs) have positioned them as a promising paradigm for agents, with long-term planning and decision-making emerging as core general-purpose capabilities for adapting to diverse...

<https://arxiv.org/abs/2601.05899>

35 CONF

Drive-P2D: A Progressive Perception-to-Decision Benchmark for VLMs in Autonomous Driving

ArXiv CS.AI 50%

arXiv:2601.14702v2 Announce Type: replace Abstract: Autonomous driving requires reliable perception and safe decision-making in complex scenarios. Recent vision-language models (VLMs) demonstrate reasoning and generalization abilities, opening new possibilities for autonomous...

<https://arxiv.org/abs/2601.14702>

36 CONF

AI Agent for Reverse-Engineering Legacy Finite-Difference Code and Translating to Devito

ArXiv CS.AI 50%

arXiv:2601.18381v2 Announce Type: replace Abstract: To facilitate the transformation of legacy finite difference implementations into the Devito environment, this study develops an integrated AI agent framework. Retrieval-Augmented Generation (RAG) and open-source Large Language...

<https://arxiv.org/abs/2601.18381>

37 CONF

Chain Of Thought Compression: A Theoretical Analysis

ArXiv CS.AI 50%

arXiv:2601.21576v2 Announce Type: replace Abstract: Chain-of-Thought (CoT) has unlocked advanced reasoning abilities of Large Language Models (LLMs) with intermediate steps, yet incurs prohibitive computational costs due to generation of extra tokens. Recent studies empirically...

<https://arxiv.org/abs/2601.21576>

38 CONF

MedCollab: IBIS-Guided Multi-Agent Collaboration with Hierarchical Disease Relation Chains for Clinical Diagnosis

ArXiv CS.AI 50%

arXiv:2603.01131v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have shown promise in clinical diagnosis but remain limited by unreliable report generation, weak evidence grounding, and opaque reasoning. We propose MedCollab, an IBIS-guided multi-agent...

<https://arxiv.org/abs/2603.01131>

39 CONF

UCPO: Uncertainty-Aware Policy Optimization

ArXiv CS.AI 50%

arXiv:2601.22648v2 Announce Type: replace Abstract: The key to building trustworthy large language models (LLMs) lies in endowing them with inherent uncertainty expression capabilities, thereby mitigating overconfident errors in high-stakes applications. However, existing RL...

<https://arxiv.org/abs/2601.22648>

40 CONF

Understanding Reasoning in LLMs through Strategic Information Allocation under Uncertainty

ArXiv CS.AI 50%

arXiv:2603.15500v2 Announce Type: replace Abstract: LLMs often exhibit Aha moments such as self-correction after tokens like "Wait," yet the underlying mechanism remains unclear. Standard LLMs collapse mainly through silent divergence, where trajectories drift from the correct...

<https://arxiv.org/abs/2603.15500>

41 CONF

OMD-GraphRAG: Enhancing GraphRAG with Ontology-Guided Extraction, Multi-Dimensional Clustering and Dual-Channel Fusion

ArXiv CS.AI 50%

arXiv:2603.25152v3 Announce Type: replace Abstract: Retrieval-Augmented Generation (RAG) systems face significant challenges in complex reasoning, multi-hop queries, and domain-specific QA. While existing GraphRAG frameworks have made progress in structural knowledge...

<https://arxiv.org/abs/2603.25152>

42 CONF

Communication Gain and Delay Cost Under Cross-Timestep Delays in Cooperative Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2604.03785v2 Announce Type: replace Abstract: Communication is essential for coordination in \emph{cooperative} multi-agent reinforcement learning under partial observability, yet \emph{cross-timestep} delays cause messages to arrive multiple timesteps after generation,...

<https://arxiv.org/abs/2604.03785>

43 CONF

Persona2Web: Benchmarking Personalized Web Agents for Contextual Reasoning with User History

ArXiv CS.AI 50%

arXiv:2602.17003v2 Announce Type: replace-cross Abstract: Large language models have advanced web agents, yet current agents lack personalization capabilities. Since users rarely specify every detail of their intent, practical web agents must be able to interpret ambiguous...

<https://arxiv.org/abs/2602.17003>

44 CONF

Retrying vs Resampling in AI Control

ArXiv CS.AI 50%

arXiv:2605.26047v2 Announce Type: replace Abstract: AI coding scaffolds like Claude Code and Codex use retrying: blocking actions flagged as risky and continuing the trajectory. We study retrying from an AI control perspective, which treats the model as potentially adversarial....

<https://arxiv.org/abs/2605.26047>

45 CONF

"Give Me BF16 or Give Me Death"? Accuracy-Performance Trade-Offs in LLM Quantization

ArXiv CS.AI 50%

arXiv:2411.02355v4 Announce Type: replace-cross Abstract: Quantization is a powerful tool for accelerating large language model (LLM) inference, but the accuracy-performance trade-offs across different formats remain unclear. In this paper, we conduct the most comprehensive...

<https://arxiv.org/abs/2411.02355>

46 CONF

Intelligent Offloading in Vehicular Edge Computing: A Comprehensive Review of Deep Reinforcement Learning Approaches and Architectures

ArXiv CS.AI 50%

arXiv:2502.06963v3 Announce Type: replace-cross Abstract: The increasing complexity of Intelligent Transportation Systems (ITS) has led to significant interest in computational offloading to external infrastructures such as edge servers, vehicular nodes, and UAVs. These dynamic...

<https://arxiv.org/abs/2502.06963>

47 CONF

Self-Cascaded Diffusion Models for Arbitrary-Scale Image Super-Resolution

ArXiv CS.AI 50%

arXiv:2506.07813v2 Announce Type: replace-cross Abstract: Arbitrary-scale image super-resolution aims to upsample images to any desired resolution, offering greater flexibility than traditional fixed-scale super-resolution. Recent approaches based on regression-based or...

<https://arxiv.org/abs/2506.07813>

48 CONF

A Physics-Informed Hierarchical Neural Network for Microwave Scattering Analysis of 3D PEC Targets

ArXiv CS.AI 50%

arXiv:2508.03774v5 Announce Type: replace-cross Abstract: Accurate modeling of scattering from three-dimensional (3D) perfectly electrically conducting (PEC) targets at microwave frequencies constitutes a fundamental objective in computational electromagnetics, particularly for...

<https://arxiv.org/abs/2508.03774>

49 CONF

How Reliable are LLMs for Reasoning on the Re-ranking task?

ArXiv CS.AI 50%

arXiv:2508.18444v2 Announce Type: replace-cross Abstract: With the improving semantic understanding capability of Large Language Models (LLMs), they exhibit a greater awareness and alignment with human values, but this comes at the cost of transparency. Although promising...

<https://arxiv.org/abs/2508.18444>

50 CONF

Scalable GANs with Transformers

ArXiv CS.AI 50%

arXiv:2509.24935v2 Announce Type: replace-cross Abstract: Scalability has driven recent advances in generative modeling, yet its principles remain underexplored for adversarial learning. We investigate the scalability of Generative Adversarial Networks (GANs) through two design...

<https://arxiv.org/abs/2509.24935>

51 CONF

When LLMs Benchmark Themselves: Deconstructing Self-Bias in Automated Evaluation

ArXiv CS.AI 50%

arXiv:2509.26600v2 Announce Type: replace-cross Abstract: As LLMs rapidly saturate existing benchmarks, automated benchmark creation using LLMs (LLM-as-a-benchmark) -- where a model generates test inputs (LLM-as-a-testset) and evaluates outputs (LLM-as-an-evaluator) -- has...

<https://arxiv.org/abs/2509.26600>

52 CONF

HiSpec: Hierarchical Speculative Decoding for LLMs

ArXiv CS.AI 50%

arXiv:2510.01336v2 Announce Type: replace-cross Abstract: Speculative decoding accelerates LLM inference by using a smaller draft model to speculate tokens that a larger target model verifies. Verification is often the bottleneck (e.g. verification is 4\times slower than token...

<https://arxiv.org/abs/2510.01336>

53 CONF

Self-signals Driven Multi-LLM Debate for Efficient and Accurate Reasoning

ArXiv CS.AI 50%

arXiv:2510.06843v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) have exhibited impressive capabilities across diverse application domains. Recent work has explored Multi-LLM Agent Debate (MAD) as a way to enhance performance by enabling multiple LLMs to...

<https://arxiv.org/abs/2510.06843>

54 CONF

Securing Multi-Agent Systems Against Corruptions via Node Contribution Backpropagation

ArXiv CS.AI 50%

arXiv:2510.19420v2 Announce Type: replace-cross Abstract: Multi-Agent Systems (MAS) have become a prevalent paradigm for Large Language Model (LLM) applications. However, the complex multi-agent design in MAS introduces unique trustworthiness concerns: adversarial agents can...

<https://arxiv.org/abs/2510.19420>

55 CONF

Explainable Cross-Disease Reasoning for Cardiovascular Risk Assessment from Low-Dose Computed Tomography

ArXiv CS.AI 50%

arXiv:2511.06625v5 Announce Type: replace-cross Abstract: Low-dose chest computed tomography (LDCT) captures pulmonary and cardiac structures in a single scan, enabling joint assessment of lung and cardiovascular health. Existing approaches typically model these domains...

<https://arxiv.org/abs/2511.06625>

56 CONF

Reconstructing Multi-Scale Physical Fields from Extremely Sparse Measurements with an Autoencoder-Diffusion Cascade

ArXiv CS.AI 50%

arXiv:2512.01572v3 Announce Type: replace-cross Abstract: Extreme sensor sparsity makes full-field reconstruction a fundamentally ill-posed problem in scientific sensing, where the goal is to infer physical fields from sparse measurements. In this regime, the posterior is severely...

<https://arxiv.org/abs/2512.01572>

57 CONF

Faithfulness Evaluation for Decoder-only LLM Attributions with Controlled Retained Information

ArXiv CS.AI 50%

arXiv:2601.03089v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are increasingly evaluated with input attribution methods, yet comparing such explanations remains challenging. Existing soft-perturbation faithfulness metrics, such as Soft-NC and Soft-NS,...

<https://arxiv.org/abs/2601.03089>

58 CONF

EpiQAL: Benchmarking Large Language Models in Epidemiological Question Answering and Reasoning

ArXiv CS.AI 50%

arXiv:2601.03471v3 Announce Type: replace-cross Abstract: Reliable epidemiological reasoning requires synthesizing study evidence to infer disease burden, transmission dynamics, and intervention effects at the population level. Existing medical question answering benchmarks...

<https://arxiv.org/abs/2601.03471>

59 CONF

Shadow Unlearning: A Neuro-Semantic Approach to Fidelity-Preserving Faceless Forgetting in LLMs

ArXiv CS.AI 50%

arXiv:2601.04275v2 Announce Type: replace-cross Abstract: Machine unlearning aims to selectively remove the influence of specific training samples to satisfy privacy regulations such as the GDPR's 'Right to be Forgotten'. However, many existing methods require access to the data...

<https://arxiv.org/abs/2601.04275>

60 CONF

Seeing vs. Believing: Evaluating the Language Bias of Open-Source MLLMs in Counter-Intuitive Scenes

ArXiv CS.AI 50%

arXiv:2601.07737v2 Announce Type: replace-cross Abstract: Multimodal Large Language Models (MLLMs) have demonstrated remarkable performance in mainstream visual understanding tasks, but their ability to process action scenes that contradict everyday common sense remains...

<https://arxiv.org/abs/2601.07737>

61 CONF

CogAdapt: Transferring Clinical ECG Foundation Models to Wearable Cognitive Load Assessment via Lead Adaptation

ArXiv CS.AI 50%

arXiv:2605.22774v2 Announce Type: replace-cross Abstract: Real-time cognitive load assessment is essential for adaptive human-computer interaction but remains challenging due to limited labeled data and poor cross-subject generalization. Recent ECG foundation models pre-trained...

<https://arxiv.org/abs/2605.22774>

62 CONF

MetaSICL: Adapting Audiorty LLM via Meta Speech In-Context Learning

ArXiv CS.AI 50%

arXiv:2601.18904v2 Announce Type: replace-cross Abstract: Auditory Large Language Models (LLMs) have demonstrated strong performance across a wide range of speech and audio understanding tasks. Nevertheless, they often struggle when applied to low-resource tasks. In case...

<https://arxiv.org/abs/2601.18904>

63 CONF

LLMs versus the Halting Problem: Characterizing Program Termination Reasoning

ArXiv CS.AI 50%

arXiv:2601.18987v5 Announce Type: replace-cross Abstract: Determining whether a program terminates is a central problem in computer science. Turing's Halting Problem established termination as undecidable, showing that no algorithm can universally determine termination for all...

<https://arxiv.org/abs/2601.18987>

64 CONF

ORLoopBench: Solver-in-the-Loop Benchmarks for Self-Correction and Behavioral Rationality in Operations Research

ArXiv CS.AI 50%

arXiv:2601.21008v3 Announce Type: replace-cross Abstract: Operations Research practitioners debug infeasible models through an iterative process: inspecting Irreducible Infeasible Subsystems (IIS), identifying constraint conflicts, and repairing formulations until feasibility...

<https://arxiv.org/abs/2601.21008>

65 CONF

GraphDancer: Training LLMs to Explore and Reason over Graphs via Two-Stage Curriculum Post-Training

ArXiv CS.AI 50%

arXiv:2602.02518v2 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly rely on external knowledge to improve factuality, yet many real-world knowledge sources are organized as heterogeneous graphs rather than plain text. Reasoning over such graphs...

<https://arxiv.org/abs/2602.02518>

66 CONF

Scaling GraphLLM with Bilevel-Optimized Sparse Querying

ArXiv CS.AI 50%

arXiv:2602.09038v2 Announce Type: replace-cross Abstract: LLMs have recently shown strong potential in enhancing node-level tasks on text-attributed graphs (TAGs) by providing explanation features. However, their practical use is severely limited by the high computational and...

<https://arxiv.org/abs/2602.09038>

67 CONF

Vital Trace: Protocol-Constrained Patient-State Reasoning for Longitudinal Clinical Trajectories

ArXiv CS.AI 50%

arXiv:2602.12833v2 Announce Type: replace-cross Abstract: Longitudinal clinical reasoning over electronic health records requires tracking evolving physiological measurements, laboratory results, and interventions across extended patient trajectories. Existing LLM-based clinical...

<https://arxiv.org/abs/2602.12833>

68 CONF

GUI-Libra: Training Native GUI Agents to Reason and Act with Action-aware Supervision and Partially Verifiable RL

ArXiv CS.AI 50%

arXiv:2602.22190v2 Announce Type: replace-cross Abstract: Open-source native GUI agents still lag behind closed-source systems on long-horizon navigation tasks. This gap stems from two limitations: a shortage of high-quality, action-aligned reasoning data, and the direct...

<https://arxiv.org/abs/2602.22190>

69 CONF

Phase-Type Variational Autoencoders for Heavy-Tailed Data

ArXiv CS.AI 50%

arXiv:2603.01800v2 Announce Type: replace-cross Abstract: Heavy-tailed distributions are ubiquitous in real-world data, where rare but extreme events dominate risk and variability. However, standard Variational Autoencoders (VAEs) employ simple decoder distributions, such as...

<https://arxiv.org/abs/2603.01800>

70 CONF

SenBen: Sensitive Scene Graphs for Explainable Content Moderation

ArXiv CS.AI 50%

arXiv:2604.08819v2 Announce Type: replace-cross Abstract: Content moderation systems classify images as safe or unsafe but lack spatial grounding and interpretability: they cannot explain what sensitive behavior was detected, who is involved, or where it occurs. We introduce the...

<https://arxiv.org/abs/2604.08819>

71 CONF

Reasoning Primitives in Hybrid and Non-Hybrid LLMs: Do Architectural Differences Yield Advantages in State-Tracking and Recall?

ArXiv CS.AI 50%

arXiv:2604.21454v2 Announce Type: replace-cross Abstract: Reasoning in large language models is often discussed as a single capability, but some of its gains may stem from simpler underlying operations. We examine two such primitives, recall and state-tracking, through five...

<https://arxiv.org/abs/2604.21454>

72 CONF

GSM-SEM: Benchmark and Framework for Generating Semantically Variant Augmentations

ArXiv CS.AI 50%

arXiv:2605.07053v2 Announce Type: replace-cross Abstract: Benchmarks like GSM8K are popular measures of mathematical reasoning, but leaderboard gains can overstate true capability due to memorization of fixed test sets. Most robustness variants apply surface-level perturbations...

<https://arxiv.org/abs/2605.07053>

73 CONF

Echoes in Filter Bubble: Diagnosing and Curing Popularity Bias in Generative Recommenders

ArXiv CS.AI 50%

arXiv:2605.16825v2 Announce Type: replace-cross Abstract: Recently, Generative Recommenders (GRs), characterized by a unified end-to-end framework, have exhibited astonishing potential in transforming the recommendation paradigm. Despite their effectiveness, we recognize that...

<https://arxiv.org/abs/2605.16825>

74 CONF

One LR Doesn't Fit All: Heavy-Tail Guided Layerwise Learning Rates for LLMs

ArXiv CS.AI 50%

arXiv:2605.22297v2 Announce Type: replace-cross Abstract: Learning rate configuration is a fundamental aspect of modern deep learning. The prevailing practice of applying a uniform learning rate across all layers overlooks the structural heterogeneity of Transformers,...

<https://arxiv.org/abs/2605.22297>