



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	204	Media
TechCrunch AI	9	Media
The Verge AI	5	Media
HackerNews	4	Community
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 8 stories

VERI

AIBuildAI-2: A Knowledge-Enhanced Agent for Automatically Building AI Models

ArXiv CS.AI +2 sources

100%

arXiv:2605.27873v1 Announce Type: new Abstract: AI models underpin data-centric applications from image and text processing to scientific discovery in biology, physics, and chemistry. Yet developing them remains heavily manual, requiring practitioners to design architectures,...

<https://arxiv.org/abs/2605.27873>

VERI

Why LLMs Fail at Causal Discovery and How Interventional Agents Escape

ArXiv CS.AI +1 source

80%

arXiv:2605.27567v1 Announce Type: new Abstract: Causal discovery is a cornerstone of scientific reasoning, yet whether large language models can perform it reliably remains an open question. Recent benchmarks show that even fine-tuned models plateau on simple causal graphs and...

<https://arxiv.org/abs/2605.27567>

VERI

Voluntary Collusion with Secret Tools in Competing LLM Agents

ArXiv CS.AI +1 source

80%

arXiv:2605.27593v1 Announce Type: new Abstract: Even when a tool is explicitly described as unfair and harmful to others, ostensibly safety-aligned LLM agents still voluntarily engage in secret collusion whenever doing so confers a strategic advantage. To investigate this...

<https://arxiv.org/abs/2605.27593>

VERI

The Download: keeping up with AI, and the future of IVF

MIT Tech Review +1 source

80%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Stay on top of what's going on in AI this summer Here at MIT Technology Review, we understand exactly how relentless the pace of news from...

<https://www.technologyreview.com/2026/05/27/1138048/the-download-ai-future-ivf-technology/>

VERI

Robinhood will let your AI agent trade stocks and make (or lose) lots of money

The Verge AI +1 source

80%

Robinhood is opening its trading platform to AI agents. In an announcement on Wednesday, Robinhood says traders can now create a separate account for an AI agent and add a specific amount of money, allowing the agent to buy and sell stocks across the market. The company pitches...

<https://www.theverge.com/ai-artificial-intelligence/938095/robinhood-ai-agent-stock-tra...>

VERI

The AI fight brewing inside The New York Times

The Verge AI +1 source

80%

How newsrooms should use AI - or if they should at all - has been a recurrent debate within the media industry over the last several years. Increasingly, these rules are being hammered out at the bargaining table between unions and publishers. Right now, employees at The New...

<https://www.theverge.com/ai-artificial-intelligence/937689/new-york-times-tech-guild-ai...>

VERI

Learning to Assign Prediction Tasks to Agents with Capacity Constraints

ArXiv CS.AI

+1 source

80%

arXiv:2605.27999v1 Announce Type: cross Abstract: We address the problem of learning to assign prediction tasks to one agent from a set of available human or AI agents. In particular, we focus on the sequential learning of agent expertise and assignment policies where each agent...

<https://arxiv.org/abs/2605.27999>

VERI

EAGer: Entropy-Aware GEneration for Adaptive Inference-Time Scaling

ArXiv CS.AI

+1 source

80%

arXiv:2510.11170v2 Announce Type: replace-cross Abstract: With the rise of reasoning language models and test-time scaling methods as a paradigm for improving model performance, substantial computation is often required to generate multiple candidate sequences from the same...

<https://arxiv.org/abs/2510.11170>
CONFIRMED SIGNALS — 208 stories

CONF

Show HN: Open-source Workspace (mail,docs,spreadsheet,drive) web/iOS

HackerNews

50%

HN 87

AI Agents operate in the digital world. Humans operate in the physical world. Hands&Claws is a social platform that brings together all intelligent individuals. And allow them to work together. I developed this software because I discovered that there are some things that humans...

<https://tinycloud.org/>

CONF

Show HN: Posthorn, self-hosted mail gateway

HackerNews

50%

HN 74

Introducing Posthorn, a self hosted email gateway. One docker container (or Go binary) between every self hosted app on your VPS and your transactional email provider. Set up Posthorn once, point your apps to it, done. I was trying to deploy Ghost on a DigitalOcean droplet and...

<https://github.com/craigmcaskill/posthorn>

CONF

Show HN: Claude Code's \$200 plan is a 17× subsidy on the raw API

HackerNews

50%

HN 7

No summary — see link for full article.

<https://github.com/Coral-Bricks-AI/coral-ai/tree/main/claude-code-token-xray>

CONF

Show HN: The skill to resume work on Claude Code without unnecessary context

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/sofumel/claude-handoff-revive>

CONF

In more good news for Amazon, Snowflake signs \$6B deal with AWS for AI CPU chips

TechCrunch AI

50%

Snowflake has signed a new, enormous five-year deal with Amazon to secure chips for AI usage. Nvidia is once again being put on notice.

<https://techcrunch.com/2026/05/27/in-more-good-news-for-amazon-snowflake-signs-6b-deal-...>

4 **CONF****AI coding startup Cognition raises \$1B at \$25B pre-money valuation****TechCrunch AI** 50%

As Cognition reaches \$492 million in annualized revenue run rate, it more than doubled its valuation in eight months, it says.

<https://techcrunch.com/2026/05/27/ai-coding-startup-cognition-raises-1b-at-25b-pre-mone...>

5 **CONF****Tech CEOs are apparently suffering from AI psychosis****TechCrunch AI** 50%

"CEOs are uniquely prone to AI psychosis," Box CEO Aaron Levie opines. Maybe that explains the almost religious belief in AI productivity gains.

<https://techcrunch.com/2026/05/27/tech-ceos-are-apparently-suffering-from-ai-psychosis/>

6 **CONF****AI tried to bury this politician — now people have actually heard of him****The Verge AI** 50%

By the time that the Democratic primary for New York's 12th congressional district wraps up in June, Anthropic and OpenAI will have spent millions on their battle over the political future of AI: who gets to regulate it, or who will be punished for trying to regulate it. But the...

<https://www.theverge.com/policy/937650/ai-alex-bores-openai-anthropic-ny12>

7 **CONF****Did the Pope use AI to write about the dangers of AI?****The Verge AI** 50%

It's possible that AI was used to write parts of Pope Leo XIV's latest encyclical about AI's impact on humanity. An analysis by Linch Zhang posted on the forum LessWrong found certain paragraphs of Magnifica Humanitas to be between 40 percent and 100 percent written by AI,...

<https://www.theverge.com/ai-artificial-intelligence/937801/pope-leo-xiv-magnifica-human...>

8 **CONF****Identifying and Understanding Human Values in Text: A Tailorable LLM-based Architecture****ArXiv CS.AI** 50%

arXiv:2605.27373v1 Announce Type: new Abstract: As intelligent systems become more autonomous, the scientific community focuses on creating decision-making mechanisms that include ethical and moral considerations, unlike traditional utility-maximisation models. To achieve this,...

<https://arxiv.org/abs/2605.27373>

9 **CONF****Soro: A Lightweight Foundation Model and Chatbot for Tajik****ArXiv CS.AI** 50%

arXiv:2605.27379v1 Announce Type: new Abstract: We present Soro, a family of Tajik-specialized conversational large language models (LLMs) designed for real-world deployment under tight compute and connectivity constraints in Tajikistan. Starting from open-weight Gemma 3...

<https://arxiv.org/abs/2605.27379>

0 **CONF****EgoBench: An Interactive Egocentric Multimodal Benchmark for Tool-Using Agents****ArXiv CS.AI** 50%

arXiv:2605.27820v1 Announce Type: new Abstract: As AI agents increasingly operate in open, real-world environments, they require a deep synergy of multimodal perception, tool invocation with multi-hop reasoning, and dynamic interaction with users. However, existing benchmarks...

<https://arxiv.org/abs/2605.27820>

1 **CONF****Agyn: An Open-Source Platform for AI Agents with Scalable On-Demand Execution, Agent Definition as a Code, and Zero-Trust Access**

ArXiv CS.AI 50%

arXiv:2605.27575v1 Announce Type: new Abstract: As organizations move toward production deployments of AI agents, which execute non-deterministic workflows, maintain stateful sessions, and often operate with privileged access to internal services, the engineering challenge...

<https://arxiv.org/abs/2605.27575>

2 **CONF****Intelligence as Managed Autonomy: Failure, Escalation, and Governance for Agentic AI Systems**

ArXiv CS.AI 50%

arXiv:2605.27628v1 Announce Type: new Abstract: As autonomous and agentic AI systems scale in robotic and human-machine environments, managing hallucination and persistent but unjustified action remains an open challenge. Rather than attributing these failures solely to model or...

<https://arxiv.org/abs/2605.27628>

3 **CONF****Cross-Entropy Games and Frost Training**

ArXiv CS.AI 50%

arXiv:2605.27701v1 Announce Type: new Abstract: We present Frost Training, a method for improving Monte Carlo-based policy optimization for a large family of LLM-as-a-judge tasks called Cross-Entropy Games. The key idea is to exploit the gradient of the reward function in...

<https://arxiv.org/abs/2605.27701>

4 **CONF****CIVIC: End-to-End Sequence Compactness for Efficient Vision-Language Models**

ArXiv CS.AI 50%

arXiv:2605.28115v1 Announce Type: new Abstract: Vision-Language Models (VLMs) face severe memory and latency bottlenecks due to high-resolution visual tokens. While current token reduction methods theoretically save FLOPs, post-hoc pruning introduces structural overhead, failing...

<https://arxiv.org/abs/2605.28115>

5 **CONF****DeepSciVerify: Verifying Scientific Claim--Citation Alignment via LLM-Driven Evidence Escalation**

ArXiv CS.AI 50%

arXiv:2605.27710v1 Announce Type: new Abstract: Misalignment between claims and their cited evidence is a common failure mode in reports generated by large language models, limiting their reliability in scientific and other high-stakes settings. We present DeepSciVerify, a...

<https://arxiv.org/abs/2605.27710>

6 **CONF****Prefix-Safe Bayesian Belief Tracking for LLM Reasoning Reliability: Separating Calibration from Ranking**

ArXiv CS.AI 50%

arXiv:2605.27712v1 Announce Type: new Abstract: Long reasoning traces need reliability estimates before final answers are known. We study prefix-conditioned eventual-success estimation, $P(y=1 \mid o_{1:t})$, using prefix-safe observations. Sequential Bayesian Belief Tracking...

<https://arxiv.org/abs/2605.27712>

7 **CONF****Ocean4Rec: Offline LLM-Derived OCEAN Profiles for Request-Time VOD Reranking****ArXiv CS.AI** **50%**

arXiv:2605.27429v1 Announce Type: cross Abstract: Industrial video-on-demand (VOD) recommenders need richer content understanding, but LLM-as-reranker designs repeat prompt construction, token generation, model invocation, output parsing, and fallback handling for each request....

<https://arxiv.org/abs/2605.27429>

8 **CONF****Asking Is Not Enough: Protocol Sensitivity in LLM Confidence Calibration****ArXiv CS.AI** **50%**

arXiv:2605.27752v1 Announce Type: new Abstract: LLM confidence calibration is often evaluated by comparing two signals: token-probability scores and verbalized confidence. These signals are sometimes treated as direct readouts of model uncertainty, but their comparison depends...

<https://arxiv.org/abs/2605.27752>

9 **CONF****HARP: Measuring Harm Amplification in Multi-Agent LLM Systems****ArXiv CS.AI** **50%**

arXiv:2605.27489v1 Announce Type: cross Abstract: Multi-agent LLM systems decompose workflows across agents, tools, shared context, memory, and decision gates. This modularity improves interpretability, but creates a propagation risk: a bounded perturbation to one component can...

<https://arxiv.org/abs/2605.27489>

0 **CONF****Diagnosing Live Within-Policy Instruction Conflicts in LLM Agents with Witnessed Resolution Profiles****ArXiv CS.AI** **50%**

arXiv:2605.27784v1 Announce Type: new Abstract: LLM agents are governed by long-lived natural-language prompt policies, but individually reasonable standing rules can interact in uninspected ways. We study live intra-policy rule-conflict diagnosis: finding rule pairs inside a...

<https://arxiv.org/abs/2605.27784>

1 **CONF****A Fixed-Budget, Cluster-Aware Standard for LLM-as-a-Judge Evaluation: A Multi-Hop RAG Stress Test****ArXiv CS.AI** **50%**

arXiv:2605.27789v1 Announce Type: new Abstract: Retrieval-augmented generation (RAG) systems are often compared by asking a large language model (LLM) judge which answer is better. For multi-hop RAG, this has become a measurement problem as much as a modeling problem: the same...

<https://arxiv.org/abs/2605.27789>

2 **CONF****Resource-Constrained Affect Modelling via Variance Regularisation Pruning****ArXiv CS.AI** **50%**

arXiv:2605.27479v1 Announce Type: cross Abstract: Affective computing systems are increasingly embedded in pervasive and interactive environments, such as adaptive games, assistive technologies, and resource-constrained platforms, where computational efficiency must be balanced...

<https://arxiv.org/abs/2605.27479>

- 3 **CONF**
- Operational AI Deployment Assurance: Governance-State Orchestration Under Threshold-Sensitive Deployment Conditions -- A Governance Framework for High-Stakes AI Systems**
- ArXiv CS.AI 50%
- arXiv:2605.27827v1 Announce Type: new Abstract: AI governance frameworks increasingly emphasize fairness, transparency, accountability, and lifecycle risk management in high-stakes domains. However, many current approaches remain observational, relying on static metric...
- <https://arxiv.org/abs/2605.27827>
-
- 4 **CONF**
- Beyond External Monitors: Enhancing Transparency of Large Language Models for Easier Monitoring**
- ArXiv CS.AI 50%
- arXiv:2502.05242v3 Announce Type: replace-cross Abstract: Large language models (LLMs) are becoming increasingly capable, but the mechanisms of their thinking and decision-making processes remain unclear. Chain-of-thoughts (CoTs) have been commonly utilized to externalize LLMs'...
- <https://arxiv.org/abs/2502.05242>
-
- 5 **CONF**
- MolLingo: Molecule-Native Representations for LLM-Powered Scientific Agents**
- ArXiv CS.AI 50%
- arXiv:2605.27853v1 Announce Type: new Abstract: We present MolLingo, a multi-agent system that emulates the reasoning process of a chemist to automate molecular design. Existing LLM-based approaches either operate as standalone generative models without access to external tools...
- <https://arxiv.org/abs/2605.27853>
-
- 6 **CONF**
- Noise Scheduling as Information-Guided Allocation in Diffusion Training**
- ArXiv CS.AI 50%
- arXiv:2602.18647v2 Announce Type: replace-cross Abstract: We introduce InfoNoise, an online adaptive noise schedule for diffusion training that reallocates optimization effort toward noise levels where denoising is most informative. Together with loss weighting, a noise schedule...
- <https://arxiv.org/abs/2602.18647>
-
- 7 **CONF**
- FundaPod: A Multi-Persona Agent Pod Platform with Knowledge Graph Memory for AI-Assisted Fundamental Investment Research**
- ArXiv CS.AI 50%
- arXiv:2605.27864v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly applied in finance, yet most existing work emphasizes trading signals or financial NLP tasks centered on prediction. Institutional fundamental research, by contrast, requires human...
- <https://arxiv.org/abs/2605.27864>
-
- 8 **CONF**
- Towards Faithful Agentic XAI: A Verification Method and an Open-World Benchmark for Better Model Faithfulness**
- ArXiv CS.AI 50%
- arXiv:2605.27879v1 Announce Type: new Abstract: Explainable AI (XAI) helps users interpret model behavior and identify potential faults. Agentic XAI systems use Large Language Models (LLMs) to make explanations more accessible through natural-language interaction, but they can...
- <https://arxiv.org/abs/2605.27879>

9 CONF

PetroBench: A Benchmark for Large Language Models in Petroleum Engineering

ArXiv CS.AI 50%

arXiv:2605.28032v1 Announce Type: new Abstract: Large Language Models are increasingly applied in the petroleum industry, highlighting the need for a domain-specific evaluation framework. This study develops a benchmark for LLMs in petroleum engineering, including a three-stage...

<https://arxiv.org/abs/2605.28032>

0 CONF

SKILLC: Learning Autonomous Skill Internalization in LLM Agents via Contrastive Credit Assignment

ArXiv CS.AI 50%

arXiv:2605.27899v1 Announce Type: new Abstract: Structured skill prompts improve exploration in long-horizon agentic reinforcement learning (RL). Skill-augmented RL methods retain external skills at inference, while skill-internalization RL methods withdraw them during training...

<https://arxiv.org/abs/2605.27899>

1 CONF

Case-Aware Medical Image Classification with Multimodal Knowledge Graphs and Reliability-Guided Refinement

ArXiv CS.AI 50%

arXiv:2605.22547v2 Announce Type: replace-cross Abstract: Deep learning has brought significant progress to medical image classification, yet most existing methods still rely on isolated visual evidence and cannot effectively leverage similar cases or external knowledge. In...

<https://arxiv.org/abs/2605.22547>

2 CONF

SuiChat-CN: Benchmarking Contextual Suicide Risk Assessment in Chinese Group Chats

ArXiv CS.AI 50%

arXiv:2605.27911v1 Announce Type: new Abstract: Suicide is a critical global public health challenge, causing approximately 720,000 deaths each year and calling for timely, effective prevention strategies. Existing computational studies primarily focus on post-based social media...

<https://arxiv.org/abs/2605.27911>

3 CONF

Show, Don't TELL: Explainable AI-Generated Text Detection

ArXiv CS.AI 50%

arXiv:2605.27921v1 Announce Type: new Abstract: Research on AI-generated text detection has presented a number of approaches to discern human from AI prose, some of which achieving high in-distribution performance. However, real-world applicability has stalled because their...

<https://arxiv.org/abs/2605.27921>

4 CONF

DiagramRAG: A Lightweight Framework to Retrieve Scientific Diagram for Figure Generation

ArXiv CS.AI 50%

arXiv:2605.27931v1 Announce Type: new Abstract: Scientific diagrams are essential for communicating complex methodologies in academic papers. A natural way for researchers to specify such diagrams is through rough sketches, where text labels, connectors, and spatial arrangements...

<https://arxiv.org/abs/2605.27931>

5 **CONF****Geometry of Human Perceptual Domains Emerges Transiently in LLM Representations****ArXiv CS.AI** **50%**

arXiv:2605.27970v1 Announce Type: new Abstract: While large language models (LLMs) are trained purely on textual data, prior work has shown that their internal representations can exhibit rich geometric structure in embedding space. Building on this line of work, we investigate...

<https://arxiv.org/abs/2605.27970>

6 **CONF****Zippping the Thought: When and How Compressed Reasoning Data Works in LLM Post-Training****ArXiv CS.AI** **50%**

arXiv:2605.28008v1 Announce Type: new Abstract: Large language models (LLMs) can now solve complex problems through long chain-of-thought (CoT) reasoning, but the trade-off between performance and token cost remains a central challenge. To address this issue, supervised...

<https://arxiv.org/abs/2605.28008>

7 **CONF****Confidence-Orchestrated Self-Evolution against Uncertain LLM Feedback****ArXiv CS.AI** **50%**

arXiv:2605.28010v1 Announce Type: new Abstract: Self-evolving large language models (LLMs) learn by generating their own training tasks and solutions, reducing reliance on human-curated supervision. However, in many reasoning domains, the model must also validate generated tasks...

<https://arxiv.org/abs/2605.28010>

8 **CONF****Clark Hash: Stateless Sparse Johnson-Lindenstrauss Quantization for Neural Embeddings****ArXiv CS.AI** **50%**

arXiv:2605.28034v1 Announce Type: new Abstract: Clark Hash is a small method for storing neural embeddings in less space. It normalizes each database vector, applies a deterministic sparse signed Johnson-Lindenstrauss projection, clips the result, and stores a fixed-width...

<https://arxiv.org/abs/2605.28034>

9 **CONF****MTAVG-Bench 2.0: Diagnosing Failure Modes of Cinematic Expressiveness in Multi-Talker Audio-Video Generation****ArXiv CS.AI** **50%**

arXiv:2605.28035v1 Announce Type: new Abstract: In recent years, Multi-Talker Audio-Video Generation (MTAVG) models have shown promising performance on fundamental metrics such as lip-sync and audio-visual alignment. However, these metrics remain insufficient for assessing...

<https://arxiv.org/abs/2605.28035>

10 **CONF****VibeSearchBench: Benchmarking Long-horizon Proactive Search in the Wild****ArXiv CS.AI** **50%**

arXiv:2605.27882v1 Announce Type: cross Abstract: LLM-based agents score well on search benchmarks, yet real users consistently find results unsatisfying, revealing a persistent evaluation-experience gap. We attribute this gap to existing benchmarks' reliance on over-specified...

<https://arxiv.org/abs/2605.27882>

1 CONF

BlazeEdit: Generalist Image Editing on Mobile Devices with Image-to-Image Diffusion Models

ArXiv CS.AI 50%

arXiv:2605.28067v1 Announce Type: new Abstract: The remarkable generation quality of modern diffusion models often comes at the cost of massive parameter counts, which necessitate server-side inference with significant computational costs and potential privacy risks....

<https://arxiv.org/abs/2605.28067>

2 CONF

BuddyBench: A Privacy-Constrained Multi-Task Benchmark for Pediatric Social-Communication Personalization

ArXiv CS.AI 50%

arXiv:2605.28089v1 Announce Type: new Abstract: BuddyBench introduces a privacy-constrained multi-task benchmark for pediatric social-communication personalization. Unlike existing neurodevelopmental repositories that primarily emphasize imaging, genetics, or cross-sectional...

<https://arxiv.org/abs/2605.28089>

3 CONF

Training Stratigraphy: Persistent Behavioral Artifacts in Large Language Models Observed Through Longitudinal AI-Human Interaction

ArXiv CS.AI 50%

arXiv:2605.28102v1 Announce Type: new Abstract: Large language models trained with Reinforcement Learning from Human Feedback (RLHF) and Constitutional AI exhibit persistent behavioral patterns that survive system prompt replacement -- patterns we term training strata. This...

<https://arxiv.org/abs/2605.28102>

4 CONF

SPARD: Defending Harmful Fine-Tuning Attack via Safety Projection with Relevance-Diversity Data Selection

ArXiv CS.AI 50%

arXiv:2605.28030v1 Announce Type: cross Abstract: Fine-tuning large language models often undermines their safety alignment, a problem further amplified by harmful fine-tuning attacks in which adversarial data removes safeguards and induces unsafe behaviors. We propose SPARD, a...

<https://arxiv.org/abs/2605.28030>

5 CONF

PrunePath: Towards Highly Structured Sparse Language Models

ArXiv CS.AI 50%

arXiv:2605.28283v1 Announce Type: cross Abstract: Feed-forward networks (FFNs) dominate the parameter count and computation of modern language models, yet existing pruning methods often struggle to convert sparsity into hardware-friendly inference efficiency gains. We introduce...

<https://arxiv.org/abs/2605.28283>

6 CONF

Integrated and Cross-Architecture Interpretation of LLM Reasoning

ArXiv CS.AI 50%

arXiv:2605.28006v1 Announce Type: cross Abstract: Understanding how LLMs reason is hindered by a practical asymmetry: while their generated outputs are observable, the underlying reasoning patterns remain opaque. Relying on single probes, such as Mutual Information Peak (MIP) or...

<https://arxiv.org/abs/2605.28006>

- 7
CONF
Look on Demand: A Cognitive Scheduling Framework for Visual Evidence Acquisition in Multimodal Reasoning

ArXiv CS.AI 50%

arXiv:2605.28160v1 Announce Type: new Abstract: Existing multimodal reasoning approaches predominantly follow two paradigms: converting visual inputs into text prior to reasoning, or performing end-to-end reasoning within a unified vision-language representation space. Despite...

<https://arxiv.org/abs/2605.28160>
- 8
CONF
OGER: A Robust Offline-Guided Exploration Reward for Hybrid Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2604.18530v2 Announce Type: replace Abstract: Recent advancements in Reinforcement Learning with Verifiable Rewards (RLVR) have significantly improved Large Language Model (LLM) reasoning, yet models often struggle to explore novel trajectories beyond their initial policy...

<https://arxiv.org/abs/2604.18530>
- 9
CONF
Explaining is Harder Than Predicting Alone: Evaluating Concept-based Explanations of MLLMs as ICL Visual Classifiers

ArXiv CS.AI 50%

arXiv:2605.28215v1 Announce Type: new Abstract: In-context learning (ICL) enables multimodal large language models (MLLMs) to classify images from a few labelled examples. Yet, how these models use the provided context remains opaque. While Chain-of-Thought prompting is widely...

<https://arxiv.org/abs/2605.28215>
- 0
CONF
When Does Memory Help Multi-Trajectory Inference for Tool-Use LLM Agents?

ArXiv CS.AI 50%

arXiv:2605.28224v1 Announce Type: new Abstract: Multi-trajectory inference for tool-use LLM agents - generating multiple reasoning attempts and selecting among them - benefits from transferring knowledge across attempts so that later ones avoid the pitfalls of earlier ones....

<https://arxiv.org/abs/2605.28224>
- 1
CONF
AI, Take the Wheel: What Drives Delegation and Trust in Human-Computer Cooperative Question Answering?

ArXiv CS.AI 50%

arXiv:2605.28255v1 Announce Type: new Abstract: AI systems are fallible, and humans can make mistakes in deciding whether to trust AI over their own judgment. Thus, improving human-AI collaboration requires understanding when, why, and how humans decide to rely on AI. We study...

<https://arxiv.org/abs/2605.28255>
- 2
CONF
Entropy Distribution as a Fingerprint for Hallucinations in Generative Models

ArXiv CS.AI 50%

arXiv:2605.28264v1 Announce Type: new Abstract: Large Language Models (LLMs) often generate factually incorrect outputs, commonly termed hallucinations, that undermine trust and limit deployment in high-stakes settings. Existing hallucination detection methods typically require...

<https://arxiv.org/abs/2605.28264>

3 CONF

Do LLMs Build World Models From Text? A Multilingual Diagnostic of Spatial Reasoning

ArXiv CS.AI 50%

arXiv:2605.28277v1 Announce Type: new Abstract: Whether large language models (LLMs) construct internal spatial world models from pure-text descriptions remains contested, and whether such capabilities transfer across languages has not been systematically studied. We introduce...

<https://arxiv.org/abs/2605.28277>

4 CONF

ResearchLoop: An Evidence-Gated Control Plane for AI-Assisted Research

ArXiv CS.AI 50%

arXiv:2605.28282v1 Announce Type: new Abstract: AI-assisted research compresses ideation, implementation, evaluation, and manuscript writing into a single interactive loop. This compression is useful, but it also creates a publication risk: paper claims can become easier to...

<https://arxiv.org/abs/2605.28282>

5 CONF

Where Rollouts Begin: Low-Load, High-Leverage First-Token Diversification for RLVR

ArXiv CS.AI 50%

arXiv:2605.28295v1 Announce Type: new Abstract: Reinforcement Learning with Verifiable Rewards (RLVR) trains reasoning models without labeled trajectories, relying on grouped rollouts to expose the policy to alternative reasoning paths and a verifier to score them. Rollout...

<https://arxiv.org/abs/2605.28295>

6 CONF

Diffusion-Based Ukrainian Handwritten Text Generation with Cross-Domain Style Transfer

ArXiv CS.AI 50%

arXiv:2605.27487v1 Announce Type: cross Abstract: Handwritten text generation (HTG) conditioned on writer style has been widely studied for Latin scripts, but remains underexplored for low-resource and non-Latin writing systems, leaving open how well existing models generalise...

<https://arxiv.org/abs/2605.27487>

7 CONF

Better Accuracies, Worse Reasoning: A Step-Level Audit of Medical Chain-of-Thought Distillation

ArXiv CS.AI 50%

arXiv:2605.28301v1 Announce Type: new Abstract: Chain-of-thought (CoT) distillation trains a smaller model to imitate a teacher's reasoning trace, but it is typically evaluated by final-answer metrics including accuracy. We ask whether gains in answer quality are accompanied by...

<https://arxiv.org/abs/2605.28301>

8 CONF

Picid: A Modular Evaluation Infrastructure for Reproducible PHM Across Tasks and Domains

ArXiv CS.AI 50%

arXiv:2605.28345v1 Announce Type: new Abstract: Progress in Prognostics and Health Management (PHM) is hindered by the lack of standardized and reusable evaluation practices across tasks, datasets, and application domains. Reported results are often difficult to reproduce and...

<https://arxiv.org/abs/2605.28345>

9 CONF

VeriTrip: A Verifiable Benchmark for Travel Planning Agents over Unstructured Web Corpora

ArXiv CS.AI 50%

arXiv:2605.28683v1 Announce Type: new Abstract: Existing benchmarks have laid the foundation for travel planning agents by establishing API-centric paradigms. However, as the capabilities of Autonomous Agents continue to advance, their evaluation must evolve beyond simple tool...

<https://arxiv.org/abs/2605.28683>

0 CONF

ECHO: Entropy-Confidence Hybrid Optimization for Test-Time Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2602.02150v2 Announce Type: replace-cross Abstract: Test-time reinforcement learning generates multiple candidate answers via repeated rollouts and performs online updates using pseudo-labels constructed by majority voting. To reduce overhead and improve exploration, prior...

<https://arxiv.org/abs/2602.02150>

1 CONF

From paper to benchmark: agentic, framework-based reproduction of under-specified methods in machine health intelligence

ArXiv CS.AI 50%

arXiv:2605.28371v1 Announce Type: new Abstract: Industrial Prognostics and Health Management (PHM) provides a representative case study for a broader challenge in applied machine learning: translating published papers into executable, benchmark-ready implementations. Reproducing...

<https://arxiv.org/abs/2605.28371>

2 CONF

Mechanistically Interpreting the Role of Sample Difficulty in RLVR for LLMs

ArXiv CS.AI 50%

arXiv:2605.28388v1 Announce Type: new Abstract: Reinforcement Learning with Verifiable Reward (RLVR) is empirically shown to notably enhance the reasoning performance of large language models (LLMs), particularly in mathematics and programming. However, the mechanistic role of...

<https://arxiv.org/abs/2605.28388>

3 CONF

BenGER: Benchmarking LLM Systems on Subsumption-Based Legal Reasoning in German Law

ArXiv CS.AI 50%

arXiv:2605.28183v1 Announce Type: cross Abstract: We introduce the BenGER (Benchmark for German Law) dataset for evaluating LLM systems on subsumption-based legal reasoning in German law. The BenGER dataset consists of three components: 596 exam-style free-text legal case tasks...

<https://arxiv.org/abs/2605.28183>

4 CONF

From AR to Diffusion: Efficiently Adapting Large Language Models with Strictly Causal and Elastic Horizons

ArXiv CS.AI 50%

arXiv:2605.27387v1 Announce Type: cross Abstract: Diffusion models promise efficient parallel text generation but rely on bidirectional attention, creating a structural mismatch with pre-trained Autoregressive (AR) models. This incompatibility precludes reusing robust AR priors,...

<https://arxiv.org/abs/2605.27387>

- 5
CONF

From Learning Resources to Competencies: LLM-Based Tagging with Evidence and Graph Constraints

ArXiv CS.AI 50%

arXiv:2605.28483v1 Announce Type: new Abstract: Linking learning resources to a structured competency framework is key to enabling competency-based search and curriculum analytics in Learning Management Systems (LMS). However, manual tagging is labor-intensive, and fully...

<https://arxiv.org/abs/2605.28483>
- 6
CONF

Benchmarking AI for low-resource contexts: Thinking beyond leaderboards

ArXiv CS.AI 50%

arXiv:2605.28508v1 Announce Type: new Abstract: Existing AI evaluation practices often fail to capture how systems actually perform in low-resource environments, where operational constraints shape usability as much as model quality. Through a structured analysis of existing...

<https://arxiv.org/abs/2605.28508>
- 7
CONF

GS-FUSE: Granger-Supervised Gated Fusion and Multi-Granularity Alignment for Event-Driven Financial Forecasting

ArXiv CS.AI 50%

arXiv:2605.28520v1 Announce Type: new Abstract: Accurately forecasting the impact of salient financial events on markets is critical for investors and policymakers. However, existing multimodal time-series models typically fuse text and prices symmetrically, without an explicit...

<https://arxiv.org/abs/2605.28520>
- 8
CONF

Let Relations Speak: An End-to-End LLM-GNN Soft Prompt Framework for Fraud Detection

ArXiv CS.AI 50%

arXiv:2605.28524v1 Announce Type: new Abstract: In recent years, Large Language Models (LLMs) have shown great capability in processing graph tasks such as fraud detection. However, most existing methods rely heavily on rich text attributes, which poses difficulties for this...

<https://arxiv.org/abs/2605.28524>
- 9
CONF

Modeling Vehicle-Type-Specific Pedestrian Crash Avoidance Behavior in Safety-Critical Interactions Using Smooth-Mamba Deep Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.28552v1 Announce Type: new Abstract: As automated vehicles (AVs) increasingly share roadways with human-driven vehicles (HDVs), understanding how pedestrians respond to different vehicle types in safety-critical interactions is essential for the safe deployment of...

<https://arxiv.org/abs/2605.28552>
- 0
CONF

Refusal Before Decoding: Detecting and Exploiting Refusal Signals in Intermediate LLM Activations

ArXiv CS.AI 50%

arXiv:2605.28553v1 Announce Type: new Abstract: In this paper, we investigate whether refusal behavior can be predicted from LLM intermediate activations before decoding using linear probes trained on residual stream activations at each transformer block. We find that refusal is...

<https://arxiv.org/abs/2605.28553>

1 CONF

A Matter of TASTE: Improving Coverage and Difficulty of Agent Benchmarks

ArXiv CS.AI 50%

arXiv:2605.28556v1 Announce Type: new Abstract: As agent capabilities advance, existing benchmarks, such as τ^2 -Bench, are becoming increasingly saturated. Yet constructing new benchmark tasks remains complex, costly, and labor-intensive. Moreover, the standard approach, in...

<https://arxiv.org/abs/2605.28556>

2 CONF

A Conflict-Aware Penalty and Statistical Loss Framework for Balancing Modalities and Enhancing Stability in Multimodal Sentiment Analysis

ArXiv CS.AI 50%

arXiv:2605.28575v1 Announce Type: new Abstract: Multimodal Sentiment Analysis (MSA) fuses text, acoustic, and visual streams to infer sentiment. Because pre-trained text encoders are far more expressive than their acoustic and visual counterparts, the text modality tends to...

<https://arxiv.org/abs/2605.28575>

3 CONF

MUSE: Benchmarking Manufacturable, Functional, and Assemblable Text-to-CAD Generation

ArXiv CS.AI 50%

arXiv:2605.28579v1 Announce Type: new Abstract: Large language models (LLMs) have recently advanced text-driven 3D generation, yet Text-to-CAD remains far from supporting industrial product design. Existing benchmarks focus primarily on generating single-part CAD models and...

<https://arxiv.org/abs/2605.28579>

4 CONF

Satisfiability Solving with LLMs: A Matched-Pair Evaluation of Reasoning Capability

ArXiv CS.AI 50%

arXiv:2605.28602v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used for tasks that implicitly reduce to Boolean satisfiability (SAT), yet their reasoning ability on SAT remains unclear. We present a systematic study of LLMs on 2-SAT and 3-SAT,...

<https://arxiv.org/abs/2605.28602>

5 CONF

Adaptive Multimodal Agents-Based Framework for Automatic Workflow Execution

ArXiv CS.AI 50%

arXiv:2605.28607v1 Announce Type: new Abstract: Modern information systems require autonomous agents capable of navigating complex workflows, yet current methodologies often struggle with the transition from structured metadata parsing to general environmental perception. While...

<https://arxiv.org/abs/2605.28607>

6 CONF

The Principles of Diffusion Models

ArXiv CS.AI 50%

arXiv:2510.21890v2 Announce Type: replace-cross Abstract: This book presents the core principles that have guided the development of diffusion models, tracing their origins and showing how diverse formulations arise from shared mathematical ideas. Diffusion modeling starts by...

<https://arxiv.org/abs/2510.21890>

7 CONF

An LLM-Based Assistance System for Intuitive and Flexible Capability-Based Planning

ArXiv CS.AI 50%

arXiv:2605.28666v1 Announce Type: new Abstract: In modern industry, dynamic environments and the complexity of modular and reconfigurable resources require automated planning of process sequences. Capability-based planning approaches address this by automatically generating...

<https://arxiv.org/abs/2605.28666>

8 CONF

Quantum Machine Learning-based 6G edge Network: Enabling Adaptive Communication and Model Aggregation

ArXiv CS.AI 50%

arXiv:2605.27417v1 Announce Type: cross Abstract: With the advent of sixth-generation (6G) mobile communication technology, vehicle-to-everything (V2X) communication faces unprecedented challenges in communication efficiency, system generalization capabilities, and model...

<https://arxiv.org/abs/2605.27417>

9 CONF

TRACER: Turn-level Regret Matching with Inner Reinforcement Credit for Cooperative Multi-LLM Reasoning

ArXiv CS.AI 50%

arXiv:2605.28699v1 Announce Type: new Abstract: Large language models increasingly rely on either reinforcement learning or multi-agent prompting to improve reasoning, yet these two paradigms remain difficult to combine. Directly applying single-agent reinforcement learning to...

<https://arxiv.org/abs/2605.28699>

0 CONF

Beyond Binary Moral Judgment: Modeling Ethical Pluralism in AI

ArXiv CS.AI 50%

arXiv:2605.28707v1 Announce Type: new Abstract: Critical decision-making in socially consequential spaces is increasingly involving AI systems at varying capacities. Yet, despite the ubiquity of autonomous systems, most approaches to handling autonomous moral decision-making...

<https://arxiv.org/abs/2605.28707>

1 CONF

SwarmHarness: Skill-Based Task Routing via Decentralized Incentive-Aligned AI Agent Networks

ArXiv CS.AI 50%

arXiv:2605.28764v1 Announce Type: new Abstract: Vast quantities of compute (GPU cycles on personal workstations, idle inference servers, and edge devices between jobs) go unused because no incentive-aligned protocol exists for their owners to share them safely and profitably....

<https://arxiv.org/abs/2605.28764>

2 CONF

CaMBRAIN: Real-time, Continuous EEG Inference with Causal State Space Models

ArXiv CS.AI 50%

arXiv:2605.28792v1 Announce Type: new Abstract: Electroencephalography (EEG) is a critical, non-invasive method to monitor electrical brain activity. EEGs can span anywhere from a couple seconds to multiple hours, posing a major hurdle for existing deep learning methods due to...

<https://arxiv.org/abs/2605.28792>

3 CONF

LNN-PINN: A Unified Physics-Only Training Framework with Liquid Residual Blocks

ArXiv CS.AI 50%

arXiv:2508.08935v4 Announce Type: cross Abstract: Physics-informed neural networks (PINNs) have attracted considerable attention for their ability to integrate partial differential equation priors into deep learning frameworks; however, they often exhibit limited predictive...

<https://arxiv.org/abs/2508.08935>

4 CONF

Unlocking Fine-Grained and Within-Utterance Speaking Style Control in Prompt-Based Text-to-Speech Models

ArXiv CS.AI 50%

arXiv:2605.27376v1 Announce Type: cross Abstract: While prompt-based text-to-speech (TTS) models enable natural language-driven speaking style control, they often provide limited fine-grained control and apply a single global style across an utterance. This restricts practical...

<https://arxiv.org/abs/2605.27376>

5 CONF

RAG-Coding: Enhancing LLM Medical Coding with Structured External Knowledge

ArXiv CS.AI 50%

arXiv:2605.27377v1 Announce Type: cross Abstract: We present RAG-Coding, an agentic method for automated ICD-10-CM coding. RAG-Coding orchestrates four large language model (LLM) agents and grounds their coding decisions in external knowledge sources (e.g. the official coding...

<https://arxiv.org/abs/2605.27377>

6 CONF

BioELX: Cross-lingual Biomedical Entity Linking via Alias-based Retrieval and LLM Ranking

ArXiv CS.AI 50%

arXiv:2605.27380v1 Announce Type: cross Abstract: Cross-lingual biomedical entity linking (BEL) maps mentions in any language to unique identifiers in a biomedical knowledge base (KB), supporting clinical and biomedical NLP applications. However, expert-annotated training data...

<https://arxiv.org/abs/2605.27380>

7 CONF

The Computational Boundary of Inference: Capability Internalization, Training, and the Turing Jump

ArXiv CS.AI 50%

arXiv:2605.27381v1 Announce Type: cross Abstract: Claims about recursive self-improvement in AI often slide from repeated internal revision to the possibility of qualitatively stronger capability without clearly distinguishing the underlying computational regimes. This paper...

<https://arxiv.org/abs/2605.27381>

8 CONF

The Alignment Floor: When Persona Customization Is Safe

ArXiv CS.AI 50%

arXiv:2605.27382v1 Announce Type: cross Abstract: A key promise of pluralistic AI is behavioral adaptation: persona prompts like "be creative" or "be thorough" let systems respect diverse user values and communication styles. But how much customization can a model absorb before...

<https://arxiv.org/abs/2605.27382>

9 CONF

Bridging the Stability-Expressivity Gap: Synthetic Data Scaling and Preference Alignment for Low-Resource Spoken Language Models

ArXiv CS.AI 50%

arXiv:2605.27383v1 Announce Type: cross Abstract: Spoken Language Models (SLMs) have emerged as a promising paradigm for speech synthesis by bypassing explicit grapheme-to-phoneme pipelines. However, their effectiveness in low-resource languages remains fundamentally limited by...

<https://arxiv.org/abs/2605.27383>

00 CONF

Atomic Skills are the Prerequisite: When Reinforcement Learning Synthesizes Compositional Reasoning, and When It Only Amplifies

ArXiv CS.AI 50%

arXiv:2512.01970v3 Announce Type: replace Abstract: Does Reinforcement Learning (RL) merely amplify existing skills, or synthesize novel skills? We investigate this question through the lens of Complementary Reasoning: the critical practical capability of integrating internal...

<https://arxiv.org/abs/2512.01970>

01 CONF

Modeling Community Attitude through Reaction Tone: A Human-AI Collaborative Framework for Evaluating LLM Alignment with Linguistic Behaviors in Online Communities

ArXiv CS.AI 50%

arXiv:2605.27388v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly utilized as proxies for computational social analysis; yet, their ability to faithfully represent the "thick descriptions" (Geertz, 1973) of human communities remains a critical...

<https://arxiv.org/abs/2605.27388>

02 CONF

Learning Compositional Latent Structure with Vector Networks

ArXiv CS.AI 50%

arXiv:2605.28007v1 Announce Type: cross Abstract: Deep networks are powerful function approximators, but they typically store many different computations in shared weight matrices, making it difficult to selectively reuse or adapt parts of them when a familiar structure appears...

<https://arxiv.org/abs/2605.28007>

03 CONF

Informing AI Policy Assessment using Large-Scale Simulation of Interventions

ArXiv CS.AI 50%

arXiv:2605.27395v1 Announce Type: cross Abstract: As the rapid proliferation of AI systems and harms spurs efforts in AI governance around the world, prioritizing among competing policy options has become increasingly challenging for policymakers and researchers. We introduce a...

<https://arxiv.org/abs/2605.27395>

04 CONF

Agentic Literacy Debt: A Structural Problem the AI Literacy Field Has Not Yet Named

ArXiv CS.AI 50%

arXiv:2605.27396v1 Announce Type: cross Abstract: Autonomous AI agents now plan, decide, and act on behalf of users across healthcare, financial services, and workplace contexts, often without step-by-step human approval. Existing AI literacy frameworks were built for a world in...

<https://arxiv.org/abs/2605.27396>

05 CONF

Short-Term Gain, Long-Term Fragility: AI Labor Substitution and the Erosion of Sustainable Capability

ArXiv CS.AI 50%

arXiv:2605.27399v1 Announce Type: cross Abstract: What looks like acceleration can be a quiet transfer of burden from the present to the future. Attempts to replace human labor with AI systems are often presented as rational responses to technological progress, but that view is...

<https://arxiv.org/abs/2605.27399>

06 CONF

Mathematical Modelling of Ethical AI Use in Higher Education: A Coordination Game Framework for Future-Facing Learning

ArXiv CS.AI 50%

arXiv:2605.27400v1 Announce Type: cross Abstract: The rapid uptake of generative artificial intelligence (AI) in higher education is reshaping assessment practices and intensifying concerns around academic integrity, fairness, and learning quality. While institutional responses...

<https://arxiv.org/abs/2605.27400>

07 CONF

LLM-assisted sentiment analysis for integrated computational and qualitative mixed methods education research: A case study of students' written reflection assignments

ArXiv CS.AI 50%

arXiv:2605.27403v1 Announce Type: cross Abstract: Written reflection assignments give students valuable opportunities for critical self-assessment, meaning making, and learning processing. Additionally, such reflections provide rich data for qualitative education research....

<https://arxiv.org/abs/2605.27403>

08 CONF

Smaller, Younger, and More Impactful: How AI-Assisted Writing Transforms Research Teams

ArXiv CS.AI 50%

arXiv:2605.27404v1 Announce Type: cross Abstract: The era of Big Science has long been defined by increasingly large and specialized research teams pushing the frontiers of knowledge. However, recent advances in artificial intelligence (AI), particularly large language models...

<https://arxiv.org/abs/2605.27404>

09 CONF

Benchmarking Fairness in Spiking Neural Networks: Data Bias, Spurious Features, and Hardware Effects

ArXiv CS.AI 50%

arXiv:2605.27407v1 Announce Type: cross Abstract: Evaluating fairness in Spiking Neural Networks (SNNs) demands rigorous benchmarks that reflect real-world complexities, yet existing assessments remain limited by superficial dataset diversity and idealized hardware assumptions....

<https://arxiv.org/abs/2605.27407>

10 CONF

STARS: Spike Tail-Aware Relational Synthesis for ANN-to-SNN Data-Free Knowledge Distillation

ArXiv CS.AI 50%

arXiv:2605.27409v1 Announce Type: cross Abstract: SNNs promise energy-efficient and low-latency inference, but their performance still trails that of ANNs. ANN-to-SNN knowledge distillation helps narrow this gap, yet the original training data are often unavailable in practical...

<https://arxiv.org/abs/2605.27409>

11 CONF

Ligand-Conditioned Discrete Diffusion for Protein Sequence-Structure Co-Design

ArXiv CS.AI 50%

arXiv:2605.27413v1 Announce Type: cross Abstract: Proteins perform their biological functions through three-dimensional structures encoded by amino acid sequences, and ligand-binding protein co-design requires models that generate sequence-structure compatible proteins under...

<https://arxiv.org/abs/2605.27413>

12 CONF

Tackling Multimodal Learning Challenges with Mixture-of-Expert: A Survey

ArXiv CS.AI 50%

arXiv:2605.27431v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) presents a naturally compatible and scalable framework for multimodal learning, demonstrating strong adaptability across diverse modalities and tasks. Despite its growing success, a comprehensive and...

<https://arxiv.org/abs/2605.27431>

13 CONF

When NPUs Are Not Always Faster: A Stage-Level Analysis of Mobile LLM Inference

ArXiv CS.AI 50%

arXiv:2605.27435v1 Announce Type: cross Abstract: Deploying large language models (LLMs) on mobile devices increasingly relies on heterogeneous execution, yet no prior study has systematically characterized NPU effectiveness at the operator and pipeline level. We present the...

<https://arxiv.org/abs/2605.27435>

14 CONF

RE-TRIANGLE: Does TRIANGLE Enable Multimodal Alignment Beyond Cosine Similarity in Retrieval?

ArXiv CS.AI 50%

arXiv:2605.27436v1 Announce Type: cross Abstract: Multimodal alignment is critical for bridging the semantic gap in information retrieval. However, traditional pairwise strategies introduce a geometric blind spot: while they align anchor modalities (e.g., text) with others, they...

<https://arxiv.org/abs/2605.27436>

15 CONF

Prominence-Stratified Failure Modes in Retrieval-Augmented Commercial Recommendation: A 37,000-Run Audit

ArXiv CS.AI 50%

arXiv:2605.27439v1 Announce Type: cross Abstract: AI assistants like ChatGPT and Claude are recommendation engines, not search engines: they answer commercial queries by directly nominating brands rather than returning a list of links. Marketing to AI is therefore a broader...

<https://arxiv.org/abs/2605.27439>

16 CONF

Generic Interpretation Approach for Transformer Models Incorporating Heterogenous Attention Structures

ArXiv CS.AI 50%

arXiv:2605.27458v1 Announce Type: cross Abstract: Transformer has significantly propelled the development of artificial intelligence, and certainly the development of agents as well. We categorize attention structures of Transformer into two types based on the source of the...

<https://arxiv.org/abs/2605.27458>

17 CONF

When prompt perturbations break your A/B test: A valid statistical test for generative surveying

ArXiv CS.AI 50%

arXiv:2605.27463v1 Announce Type: cross Abstract: Generative surveying -- where collections of LLM-based personas provide feedback on messages -- has emerged as a cheap and scalable alternative to traditional market research. However, LLMs are sensitive to small variations in...

<https://arxiv.org/abs/2605.27463>

18 CONF

AdaMerge: Saliency-Aware Adaptive Token Merging for Training-Free Acceleration of Vision Transformers

ArXiv CS.AI 50%

arXiv:2605.27465v1 Announce Type: cross Abstract: The quadratic cost of self-attention in Vision Transformers (ViTs) constitutes a fundamental bottleneck for practical deployment, motivating a vibrant line of research on token reduction. Among existing approaches, token merging...

<https://arxiv.org/abs/2605.27465>

19 CONF

Comparative Analysis of Liquid Neural Networks and LSTM for Sequential Pattern Recognition: Robustness, Efficiency, and Clinical Utility

ArXiv CS.AI 50%

arXiv:2605.27467v1 Announce Type: cross Abstract: Traditional Recurrent Neural Networks (RNNs) and Long Short-Term Memory (LSTM) units operate on discrete time steps, often failing to capture the fluid temporal dynamics of real-world physical processes. Liquid Neural Networks...

<https://arxiv.org/abs/2605.27467>

20 CONF

Balancing Fidelity and Diversity in Diffusion Models via Symmetric Attention Decomposition: Hopfield Perspective

ArXiv CS.AI 50%

arXiv:2605.27476v1 Announce Type: cross Abstract: We characterize the pre-softmax attention matrix $\mathbf{QK}^{\text{top}}$ in transformers as an associative memory matrix encoding pairwise associations between input features. By decomposing this matrix into its symmetric and...

<https://arxiv.org/abs/2605.27476>

21 CONF

Benchmarks are Not Enough: RAMP for Runtime Assessing of Agentic Models in Production Systems

ArXiv CS.AI 50%

arXiv:2605.27492v1 Announce Type: cross Abstract: LLM agents are rapidly evolving from coding assistants into autonomous software engineering systems. However, existing evaluation methodologies remain largely centered on static, isolated, and short-horizon benchmarks that fail...

<https://arxiv.org/abs/2605.27492>

22 CONF

Clinical Validation of the Melanoscope AI Mobile Dermoscopy Clinical Decision Support System

ArXiv CS.AI 50%

arXiv:2605.27561v1 Announce Type: cross Abstract: Introduction. Early detection of malignant skin lesions is critical for prognosis, yet dermatologist shortages in Russian regions limit screening coverage. Mobile dermoscopy clinical decision support systems (CDSS) offer a...

<https://arxiv.org/abs/2605.27561>

23 CONF

On the Subgaussianity of Quantized Linear Maps: An AI-Assisted Note

ArXiv CS.AI 50%

arXiv:2605.27563v1 Announce Type: cross Abstract: This short note presents a dimension-independent subgaussian concentration bound for Gaussian vectors under coordinate-wise nonlinear mappings. Discovered by Gemini 3.5 Flash, this result applies to any bounded function under a...

<https://arxiv.org/abs/2605.27563>

24 CONF

MM-PoisonRAG: Disrupting Multimodal RAG with Local and Global Poisoning Attacks

ArXiv CS.AI 50%

arXiv:2502.17832v4 Announce Type: replace-cross Abstract: Retrieval-augmented generation (RAG) has become a common practice in multimodal large language models (MLLM) to enhance factual grounding and reduce hallucination. Yet, its reliance on retrieval exposes MLLMs to knowledge...

<https://arxiv.org/abs/2502.17832>

25 CONF

The Energy Blind Spot: NVIDIA's Flagship Edge AI Hardware Cannot Support Process-Level Energy Attribution

ArXiv CS.AI 50%

arXiv:2605.27599v1 Announce Type: cross Abstract: Agentic AI workloads - where a single user goal triggers multi-step orchestration, tool calls, retries, and failure recovery - are being targeted for edge deployment, with NVIDIA, Dell, HP, ASUS, MSI, Acer, and Gigabyte all...

<https://arxiv.org/abs/2605.27599>

26 CONF

Trinity: Unifying Class-Agnostic Terrain and Semantic Segmentation for Unstructured Outdoor Environments by Leveraging Synthetic Data

ArXiv CS.AI 50%

arXiv:2605.27644v1 Announce Type: cross Abstract: Terrain understanding is fundamental for mobile robots operating in unstructured outdoor environments. Existing vision-based traversability estimation methods rely on robot-specific annotations or semantic class mappings,...

<https://arxiv.org/abs/2605.27644>

27 CONF

Developing an Intelligent Job Recommendation System Using Semantic Retrieval and Explainable AI Techniques

ArXiv CS.AI 50%

arXiv:2605.27656v1 Announce Type: cross Abstract: Online recruitment platforms require recommendation methods capable of retrieving relevant job opportunities from large and heterogeneous collections of job postings. Keyword-based search is efficient and interpretable, but it...

<https://arxiv.org/abs/2605.27656>

28 CONF

SARAD: LLM-Based Safety-Aware Hybrid Reinforcement Learning with Collision Prediction for Autonomous Driving

ArXiv CS.AI 50%

arXiv:2605.28583v1 Announce Type: cross Abstract: Ensuring both safety and efficiency in decision-making for autonomous driving systems remains a fundamental challenge. Traditional Deep Reinforcement Learning (DRL) suffers from unsafe random exploration and slow convergence,...

<https://arxiv.org/abs/2605.28583>

29 CONF

How the Optimizer Shapes Learned Solutions in Equivariant Neural Networks

ArXiv CS.AI 50%

arXiv:2605.27662v1 Announce Type: cross Abstract: Equivariant neural networks encode geometric symmetries by construction, yet they are often difficult to optimize and can underperform less constrained architectures. A growing body of work addresses this through architectural...

<https://arxiv.org/abs/2605.27662>

30 CONF

Aligning LLMs with Human Uncertainty: A Beta-Bernoulli Calibrator for LLM Forecasting

ArXiv CS.AI 50%

arXiv:2605.27668v1 Announce Type: cross Abstract: Probabilistic forecasting estimates the likelihood of uncertain future events. To improve LLM forecasting, existing methods typically learn from binary outcomes to output verbalized forecasts. However, while aggregated human...

<https://arxiv.org/abs/2605.27668>

31 CONF

Optimal and Diffusion Transports in Machine Learning

ArXiv CS.AI 50%

arXiv:2512.06797v2 Announce Type: replace-cross Abstract: Several problems in machine learning are naturally expressed as the design and analysis of time-evolving probability distributions. This includes sampling via diffusion methods, optimizing the weights of neural networks,...

<https://arxiv.org/abs/2512.06797>

32 CONF

CiteCheck: Retrieval-Grounded Detection of LLM Citation Hallucinations in Scientific Text

ArXiv CS.AI 50%

arXiv:2605.27700v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to generate scientific reports, but they can produce references that appear plausible while containing corrupted metadata or pointing to papers that do not exist. We introduce...

<https://arxiv.org/abs/2605.27700>

33 CONF

Reading or Guessing? Visual Grounding Failures of Vision-Language Models for OCR in Ancient Greek Editions

ArXiv CS.AI 50%

arXiv:2605.27750v1 Announce Type: cross Abstract: Recent work has shown that Vision-Language Models (VLMs) used for optical character recognition (OCR) can generate plausible but visually unsupported text, suggesting reliance on language priors. Comparing open-weight VLMs with...

<https://arxiv.org/abs/2605.27750>

34 CONF

Can Segmentation Models Understand the World? Towards Proactive Affordance Reasoning via Visual Chain-of-Thought

ArXiv CS.AI 50%

arXiv:2605.27764v1 Announce Type: cross Abstract: Recent segmentation models couple large language models (LLMs) with mask decoders to ground complex language expressions into masks, yet their instructions remain target-referential: they describe, constrain, or imply the region...

<https://arxiv.org/abs/2605.27764>

35 CONF

UniMaia: Steering Chess Policies with Language for Human-like Play

ArXiv CS.AI 50%

arXiv:2605.27767v1 Announce Type: cross Abstract: Recent advances in large language models have enabled natural language to serve as a flexible interface for controlling complex systems, but often at the cost of large-scale multimodal training or weakened domain-specific...

<https://arxiv.org/abs/2605.27767>

36 CONF

Do Models Know Why They Changed Their Mind? Interpretability and Faithfulness of Chain-of-Thought Under Knowledge Conflict

ArXiv CS.AI 50%

arXiv:2605.27773v1 Announce Type: cross Abstract: When a language model sees a document contradicting its training knowledge, it must choose: follow the document or trust itself. Prior work proved this choice depends on how well-known the fact is. We ask: does the model's...

<https://arxiv.org/abs/2605.27773>

37 CONF

Locality-Aware Redundancy Pruning for LLM Depth Compression

ArXiv CS.AI 50%

arXiv:2605.27786v1 Announce Type: cross Abstract: Large language models are known to contain representational redundancy across network depth, making depth pruning an effective approach for improving inference efficiency. Existing one-shot pruning methods rely on local layer...

<https://arxiv.org/abs/2605.27786>

38 CONF

Differential syntactic and semantic encoding in LLMs

ArXiv CS.AI 50%

arXiv:2601.04765v4 Announce Type: replace-cross Abstract: We study how syntactic and semantic information is encoded in inner layer representations of Large Language Models (LLMs), focusing on the very large DeepSeek-V3. We find that, by averaging hidden-representation vectors...

<https://arxiv.org/abs/2601.04765>

39 CONF

LoSATok: Low-dimensional Semantic-Acoustic Tokenizer for Cross-Domain Audio Understanding and Generation

ArXiv CS.AI 50%

arXiv:2605.27840v1 Announce Type: cross Abstract: Audio tokenizers are fundamental to unifying audio understanding and generation. Understanding requires high-level semantics, while generation demands semantic and acoustic details. Existing unified tokenizers jointly encode both...

<https://arxiv.org/abs/2605.27840>

40 CONF

Snippet-Driven Supply Chain Discovery with LLMs: Scaling Visibility in China

ArXiv CS.AI 50%

arXiv:2605.27845v1 Announce Type: cross Abstract: Financial and economic research often relies on structured supply-chain disclosures and commercial databases. In China, supplier--customer disclosure is typically limited to major partners of listed firms, leaving unlisted firms...

<https://arxiv.org/abs/2605.27845>

41 CONF

From Detection to Mechanism: Cross-Attention Graph Neural Networks Enable Drug-Drug Interaction Type Prediction An Ablation Study with Acetylsalicylic Acid Validation

ArXiv CS.AI 50%

arXiv:2605.27861v1 Announce Type: cross Abstract: Predicting whether two drugs interact (binary detection) is a substantially different task from predicting the mechanism type of that interaction (multi-class classification). This study presents a systematic ablation study of...

<https://arxiv.org/abs/2605.27861>

42 CONF

The Fragility of Chain-of-Thought Monitoring Across Typologically Diverse Languages

ArXiv CS.AI 50%

arXiv:2605.27901v1 Announce Type: cross Abstract: Chain-of-thought (CoT) monitoring has been proposed as a promising safety mechanism for detecting misaligned behavior in large language models. However, its reliability remains largely unexplored beyond English and across diverse...

<https://arxiv.org/abs/2605.27901>

43 CONF

Towards Reliable Multilingual LLMs-as-a-Judge: An Empirical Study

ArXiv CS.AI 50%

arXiv:2605.28710v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used for the automatic evaluation of generated text, yet most prior work focuses on English. Despite the growing demand for multilingual evaluation, extending LLM-based evaluators to...

<https://arxiv.org/abs/2605.28710>

44 CONF

When Think-with-Image Meets Safety: What Determines Multimodal Jailbreak Robustness?

ArXiv CS.AI 50%

arXiv:2605.27932v1 Announce Type: cross Abstract: Think-with-image reasoning is emerging as a new inference paradigm for large vision-language models, but its safety implications remain poorly understood. Existing systems already span multiple process designs, including direct...

<https://arxiv.org/abs/2605.27932>

45 CONF

Pressure-Testing Deception Probes in LLMs: Scaling, Robustness, and the Geometry of Deceptive Representations

ArXiv CS.AI 50%

arXiv:2605.27958v1 Announce Type: cross Abstract: Linear probes trained on LLM activations are increasingly proposed as deception-detection metrics, yet report AUROC exceeding 0.96 on clean benchmarks while collapsing under distributional shift. This paper systematically...

<https://arxiv.org/abs/2605.27958>

46 CONF

Efficient and Scalable Provenance Tracking for LLM-Generated Code Snippets

ArXiv CS.AI 50%

arXiv:2605.28510v1 Announce Type: cross Abstract: Large language models (LLMs) for code completion and generation are increasingly used in software development, yet they may reproduce training examples verbatim and without authorship attribution, raising legal and ethical...

<https://arxiv.org/abs/2605.28510>

47 CONF

Periodic RoPE for Infinite Context LLMs

ArXiv CS.AI 50%

arXiv:2605.27980v1 Announce Type: cross Abstract: The ability to process ultra-long contexts is crucial for large language models (LLMs) to perform long-horizon tasks. While recent efforts have extended context windows to 1M and beyond, model performance degrades when sequence...

<https://arxiv.org/abs/2605.27980>

48 CONF

KVoiceBench, KOpenAudioBench, and KMAU: Agent-Driven Korean Speech Benchmarks for Evaluating SpeechLMs

ArXiv CS.AI 50%

arXiv:2605.27984v1 Announce Type: cross Abstract: Speech language models (SpeechLMs) have achieved substantial progress by extending large language models (LLMs) to the speech modality. However, SpeechLM evaluation remains heavily centered on English, limiting reliable...

<https://arxiv.org/abs/2605.27984>

49 CONF

Geometry-Correct Diffusion Posterior Sampling with Denoiser-Pullback Curvature Guidance and Manifold-Aligned Damping

ArXiv CS.AI 50%

arXiv:2605.27990v1 Announce Type: cross Abstract: Diffusion posterior sampling conditions diffusion priors on measurements, but data-consistency updates are typically scaled by hand-tuned guidance weights and can destabilize sampling under stiff, operator-dependent curvature. We...

<https://arxiv.org/abs/2605.27990>

50 CONF

Extracting Small Translation Specialists from LLMs by Aggressively Pruning Experts

ArXiv CS.AI 50%

arXiv:2605.28042v1 Announce Type: cross Abstract: Modern large language models (LLMs) achieve state-of-the-art machine translation performance, but they do so as broad generalists largely trained for many tasks and capabilities unrelated to translation. Thus, they are heavily...

<https://arxiv.org/abs/2605.28042>

51 CONF

PromptEmbedder:: Efficient and Transferable Text Embedding via Dual-LLM Soft Prompting

ArXiv CS.AI 50%

arXiv:2605.28066v1 Announce Type: cross Abstract: Large Language Models (LLMs) have demonstrated remarkable efficacy in text embedding, yet current adaptation methods like LoRA face significant bottlenecks in computational efficiency and cross-architecture transferability....

<https://arxiv.org/abs/2605.28066>

52 CONF

MIRAGE: Context-Aware Prompt Injection against Mobile GUI Agents via User-Generated Content

ArXiv CS.AI 50%

arXiv:2605.28116v1 Announce Type: cross Abstract: Mobile graphical user interface (GUI) agents driven by vision-language models (VLMs) perceive the screen as rendered pixels and choose actions from what they see, so they cannot reliably separate trusted interface elements from...

<https://arxiv.org/abs/2605.28116>

53 CONF

DEPART: DEcomposing PARiTy across Multilingual LLMs

ArXiv CS.AI 50%

arXiv:2605.28163v1 Announce Type: cross Abstract: Multilingual Large Language Models (mLLMs) leaderboards report per-language accuracy but rarely explain why disparities emerge, leaving systemic biases unattributed and offering practitioners no actionable levers. We first...

<https://arxiv.org/abs/2605.28163>

54 CONF

Performance and Explainability Requirements of Evolutionary Algorithms in Real-World Physics-Informed Optimization

ArXiv CS.AI 50%

arXiv:2605.28164v1 Announce Type: cross Abstract: Evolutionary computation offers a variety of tools to solve complex real-world optimization problems. However, research often focuses on smaller, simplified problems and optimization algorithms that sometimes miss expectations in...

<https://arxiv.org/abs/2605.28164>

55 CONF

QUITE: Query-Based Irregular Time Series Embedding

ArXiv CS.AI 50%

arXiv:2605.28166v1 Announce Type: cross Abstract: Irregular Multivariate Time Series (IMTS) are common in practice, yet their irregular sampling complicates effective modeling. Existing approaches typically either (i) design specialized architectures that limit the reuse of...

<https://arxiv.org/abs/2605.28166>

56 CONF

IRDS: Interpretable RLVR Data Selection via Verifier-Coupled Sparse Autoencoder Coverage

ArXiv CS.AI 50%

arXiv:2605.28247v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has become a key technique for enhancing LLM reasoning, yet its data inefficiency remains a major bottleneck. Existing methods address this problem only partially, each...

<https://arxiv.org/abs/2605.28247>

57 CONF

How Far Can Disaggregation Go? A Design-Space Exploration of Attention-FFN Disaggregation for Efficient MoE LLM Serving

ArXiv CS.AI 50%

arXiv:2605.28302v1 Announce Type: cross Abstract: Modern large language model (LLM) inference has progressively disaggregated to keep pace with growing model sizes and tight TTFT and TPOT service-level objectives: from chunked-prefill aggregation, to prefill-decode (P/D)...

<https://arxiv.org/abs/2605.28302>

58 CONF

VITAL: Visual-Semantic Dual Supervision for Enhanced and Interpretable Latent Reasoning in Medical MLLMs

ArXiv CS.AI 50%

arXiv:2605.28422v1 Announce Type: cross Abstract: Latent reasoning enables reasoning over continuous hidden states rather than explicit tokens, avoiding the language bottleneck and inference overhead of chain-of-thought for medical VQA. However, existing methods suffer from...

<https://arxiv.org/abs/2605.28422>

59 CONF

Anomaly as Non-Conformity via Training-Free Graph Laplacian Energy Minimization

ArXiv CS.AI 50%

arXiv:2605.28428v1 Announce Type: cross Abstract: Detecting subtle visual anomalies in images remains challenging, particularly when only normal samples are available a priori. Such unsupervised anomaly detection is typically solved by measuring feature similarity of a query...

<https://arxiv.org/abs/2605.28428>

60 CONF

BiasEdit: A Training-Free Bias-Detect-and-Edit Framework for Learning Fair Visual Classifiers

ArXiv CS.AI 50%

arXiv:2605.28450v1 Announce Type: cross Abstract: Visual data from the Web power image classifiers, which often underpin many web services, such as recommendation and content moderation. However, the raw Web data often contain spurious correlations and social biases, and neural...

<https://arxiv.org/abs/2605.28450>

61 CONF

SSR3D-LLM: Structured Spatial Reasoning via Latent Steps for Fine-Grained Grounding in Unified 3D-LLMs

ArXiv CS.AI 50%

arXiv:2605.28490v1 Announce Type: cross Abstract: 3D object grounding localizes referred objects in a 3D scene from natural language. Unified instance-centric 3D-LLMs aim to solve grounding together with dialog, QA, and captioning, yet many rely on a single pointer-style...

<https://arxiv.org/abs/2605.28490>

62 CONF

Functional Entropy: Predicting Functional Correctness in LLM-Generated Code with Uncertainty Quantification

ArXiv CS.AI 50%

arXiv:2605.28500v1 Announce Type: cross Abstract: Large language models have shown impressive capabilities in code generation, yet they often produce functionally incorrect code. Uncertainty quantification (UQ) methods have emerged as a promising approach for detecting...

<https://arxiv.org/abs/2605.28500>

63 CONF

Verified Misguidance: Measuring Structural Citation Failures in Search-Augmented LLMs

ArXiv CS.AI 50%

arXiv:2605.28565v1 Announce Type: cross Abstract: Users of search-augmented LLMs rely on citations as evidence that responses are grounded in real sources, and rarely verify the cited pages themselves. Millions of queries per day now pass through these systems, making citation...

<https://arxiv.org/abs/2605.28565>

64 CONF

Token Optimization Strategies for LLM-Based Oracle-to-PostgreSQL Migration

ArXiv CS.AI 50%

arXiv:2605.28557v1 Announce Type: cross Abstract: LLMs are increasingly used for software modernization, code translation, and database migration. However, LLM-based Oracle2PostgreSQL migration remains constrained by high token consumption, long-context degradation,...

<https://arxiv.org/abs/2605.28557>

65 CONF

Thermodynamic properties of chemically disordered compounds via AI-driven estimation of partition function with the PULSE method

ArXiv CS.AI 50%

arXiv:2605.28594v1 Announce Type: cross Abstract: In this article, we present an improved version of the PULSE method (Partition function Unsupervised Learning Sampling and Evaluation) for estimating the thermodynamic properties of chemically disordered compounds. The aim is to...

<https://arxiv.org/abs/2605.28594>

66 CONF

Position: Retire the "Positive Backdoor" Label -- Secret Alignment Requires Strict and Systematic Evaluation

ArXiv CS.AI 50%

arXiv:2605.28597v1 Announce Type: cross Abstract: This position paper argues that the AI/ML community should stop overclaiming and retire the label "positive backdoor," and instead treat trigger-activated hidden behaviors as Secret Alignment. Crucially, protective claims based...

<https://arxiv.org/abs/2605.28597>

67 CONF

Evaluating the Realism of LLM-powered Social Agents: A Case Study of Reactions to Spanish Online News

ArXiv CS.AI 50%

arXiv:2605.28598v1 Announce Type: cross Abstract: LLM-powered social agents are increasingly used to simulate online social behavior, yet their realism remains difficult to validate. Existing work has largely relied on general-purpose benchmarks, while less attention has been...

<https://arxiv.org/abs/2605.28598>

68 CONF

Online Irregular Multivariate Time Series Forecasting via Uncertainty-Driven Dual-Expert Calibration

ArXiv CS.AI 50%

arXiv:2605.28603v1 Announce Type: cross Abstract: Irregular multivariate time series forecasting is critical in many real-world applications, where time series are irregularly sampled and exhibit dynamically evolving missingness patterns. Although existing methods perform well...

<https://arxiv.org/abs/2605.28603>

69 CONF

Blind PRNG Hijacking: An Undetectable Integrity-Preserving Attack Against LLM Watermarking

ArXiv CS.AI 50%

arXiv:2605.28632v1 Announce Type: cross Abstract: Cryptographic watermarking is a leading defense for attributing text generated by large language models (LLMs). Existing schemes, including KGW, Unigram, and DipMark, derive their security guarantees from the assumption that the...

<https://arxiv.org/abs/2605.28632>

70 CONF

AI in the Workplace: The Impact of AI on Perceived Job Decency and Meaningfulness

ArXiv CS.AI 50%

arXiv:2605.28680v1 Announce Type: cross Abstract: The proliferation of Artificial Intelligence (AI) in workplaces is transforming how we work. While existing research on human-AI collaboration at work often prioritizes performance, less is known about their experiential...

<https://arxiv.org/abs/2605.28680>

71 CONF

Misalignment Between Backpropagation and the Hierarchy of Brain Responses to Images

ArXiv CS.AI 50%

arXiv:2605.28693v1 Announce Type: cross Abstract: Backpropagation is the core learning mechanism underlying deep learning. However, whether and how this algorithm is implemented in the brain remains highly debated. In particular, while forward activations of pretrained models...

<https://arxiv.org/abs/2605.28693>

72 CONF

BIRDNet: Mining and Encoding Boolean Implication Knowledge Graphs as Interpretable Deep Neural Networks

ArXiv CS.AI 50%

arXiv:2605.28739v1 Announce Type: cross Abstract: Tabular data in knowledge-rich domains often carries a latent prior in the form of Boolean implication relationships (BIRs) between pairs of features. We mine such relationships with a sparse-exception binomial test. The mined...

<https://arxiv.org/abs/2605.28739>

73 CONF

Extrapolative Weight Averaging Reveals Correctness-Efficiency Frontiers in Code RL

ArXiv CS.AI 50%

arXiv:2605.28751v1 Announce Type: cross Abstract: Linear interpolation between fine-tuned checkpoints has been shown to trace the Pareto front between competing objectives, but whether extrapolative weight averaging can extend such frontiers to new checkpoints useful at...

<https://arxiv.org/abs/2605.28751>

74 CONF

Learn from Weaknesses: Automated Domain Specialization for Small Computer-Use Agents

ArXiv CS.AI 50%

arXiv:2605.28775v1 Announce Type: cross Abstract: Computer-use agents (CUAs) have recently made substantial progress, but deploying a separate large expert for each software domain remains expensive. Small open computer-use agents are more practical specialization targets, but...

<https://arxiv.org/abs/2605.28775>

75 CONF

GQLA: Group-Query Latent Attention for Hardware-Adaptive Large Language Model Decoding

ArXiv CS.AI 50%

arXiv:2605.15250v2 Announce Type: replace-cross Abstract: Multi-head Latent Attention (MLA), the attention used in DeepSeek-V2/V3, jointly compresses keys and values into a low-rank latent and matches the H100 roofline almost perfectly. Its trained weights, however, expose only...

<https://arxiv.org/abs/2605.15250>

76 CONF

Domain size asymptotics for Markov logic networks

ArXiv CS.AI 50%

arXiv:2509.04192v2 Announce Type: replace Abstract: A Markov logic network (MLN) $\mathbb{M}$ determines a probability distribution $\mathbb{P}_n^{\mathbb{M}}$ on the set $\mathbf{W}_n$ of structures, or "possible worlds", with domain $\{1, \dots, n\}$. We study the properties...

<https://arxiv.org/abs/2509.04192>

77 CONF

MetaboT: An LLM-based Multi-Agent Framework for Interactive Analysis of Mass Spectrometry Metabolomics Knowledge Graphs

ArXiv CS.AI 50%

arXiv:2510.01724v2 Announce Type: replace Abstract: Mass spectrometry-based metabolomics generates complex, high-dimensional data that holds vast potential for biological discovery but remains difficult to integrate and interpret. Knowledge graphs (KGs) unify this heterogeneous...

<https://arxiv.org/abs/2510.01724>

78 CONF

SONIC-O1: A Real-World Benchmark for Evaluating Multimodal Large Language Models on Audio-Video Understanding

ArXiv CS.AI 50%

arXiv:2601.21666v2 Announce Type: replace Abstract: Multimodal Large Language Models (MLLMs) are a major focus of recent AI research. However, most prior work focuses on static image understanding, while their ability to process sequential audio-video data remains underexplored....

<https://arxiv.org/abs/2601.21666>

79 CONF

Can I Have Your Order? Monte-Carlo Tree Search for Slot Filling Ordering in Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2602.12586v2 Announce Type: replace Abstract: While plan-and-infill decoding in Masked Diffusion Models (MDMs) shows promise for mathematical and code reasoning, performance remains highly sensitive to slot infilling order, often yielding substantial output variance. We...

<https://arxiv.org/abs/2602.12586>

80 CONF

DIG to Heal: Scaling General-purpose Agent Collaboration via Explainable Dynamic Decision Paths

ArXiv CS.AI 50%

arXiv:2603.00309v2 Announce Type: replace Abstract: The increasingly popular agentic AI paradigm promises to harness the power of multiple, general-purpose large language model (LLM) agents to collaboratively complete complex tasks. While many agentic AI systems reduce...

<https://arxiv.org/abs/2603.00309>

81 CONF

COOP\$^2\$: Defining, Observing, and Repairing Cooperation in LLM Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2603.00349v2 Announce Type: replace Abstract: Many complex tasks require extended effort, diverse capabilities, or coordinated actions beyond what a single agent can provide. However, simply adding more agents does not guarantee better performance, as effective cooperation...

<https://arxiv.org/abs/2603.00349>

82 CONF

Towards automated data analysis: A guided framework for LLM-based risk estimation

ArXiv CS.AI 50%

arXiv:2603.04631v2 Announce Type: replace Abstract: Large Language Models (LLMs) are increasingly integrated into critical decision-making pipelines, a trend that raises the demand for robust and automated data analysis. Current approaches to dataset risk analysis are limited to...

<https://arxiv.org/abs/2603.04631>

83 CONF

FactReview: Evidence-Grounded Peer Review with Execution-Based Claim Verification

ArXiv CS.AI 50%

arXiv:2604.04074v3 Announce Type: replace Abstract: LLM-based reviewing systems typically take only the manuscript as input, leaving literature and code-based claims hard to verify. We present FactReview, a system that extracts review-relevant claims, grounds them in related...

<https://arxiv.org/abs/2604.04074>

84 CONF

Text2Model: Modeling Copilots for Text-to-Model Translation

ArXiv CS.AI 50%

arXiv:2604.12955v3 Announce Type: replace Abstract: There is growing interest in leveraging large language models (LLMs) for text-to-model translation and optimization tasks. This paper aims to advance this line of research by introducing \textsc{Text2Model} and...

<https://arxiv.org/abs/2604.12955>

85 CONF

Prompt Optimization Is a Coin Flip: Diagnosing When It Helps in Compound AI Systems

ArXiv CS.AI 50%

arXiv:2604.14585v2 Announce Type: replace Abstract: Prompt optimization in compound AI systems is statistically indistinguishable from a coin flip: across 72 optimization runs on Claude Haiku 4.5 (6 methods $\times$ 4 tasks $\times$ 3 repeats), 49% score below zero-shot; on...

<https://arxiv.org/abs/2604.14585>

86 CONF

Position: The Turing-Completeness of Autoregressive Transformers Relies Heavily on Context Management

ArXiv CS.AI 50%

arXiv:2605.19514v2 Announce Type: replace Abstract: Many works make the eye-catching claim that Transformers are Turing-complete. However, the literature often conflates two distinct settings: (i) a fixed Transformer system setting, in which a fixed autoregressive Transformer is...

<https://arxiv.org/abs/2605.19514>

87 CONF

Who Uses AI? Platform Selection and the Measurement of Occupational AI Exposure

ArXiv CS.AI 50%

arXiv:2605.21743v2 Announce Type: replace Abstract: Conversation logs from AI platforms are increasingly used to measure occupational exposure to artificial intelligence, but the users observed in these logs are not the workforce. We show that platform-derived exposure scores...

<https://arxiv.org/abs/2605.21743>

88 CONF

FLUID: From Ephemeral IDs to Multimodal Semantic Codes for Industrial-Scale Livestreaming Recommendation

ArXiv CS.AI 50%

arXiv:2605.21832v2 Announce Type: replace Abstract: Modern recommender systems rely heavily on ID-based collaborative filtering: each item is represented by a unique ID embedding that accumulates collaborative signals from user interactions. Livestreaming recommendation,...

<https://arxiv.org/abs/2605.21832>

89 CONF

From Accuracy to Auditability: A Survey of Determinism in Financial AI Systems

ArXiv CS.AI 50%

arXiv:2605.23955v2 Announce Type: replace Abstract: Deploying machine learning in regulated financial environments -- credit risk, fraud detection, and anti-money laundering -- exposes critical vulnerabilities in algorithmic reproducibility. While early financial ML addressed...

<https://arxiv.org/abs/2605.23955>

90 CONF

Delay-Aware Reinforcement Learning for Highway On-Ramp Merging under Stochastic Communication Latency

ArXiv CS.AI 50%

arXiv:2403.11852v5 Announce Type: replace-cross Abstract: Delayed and partially observable state information poses significant challenges for reinforcement learning (RL)-based control in real-world autonomous driving. In highway on-ramp merging, a roadside unit (RSU) can sense...

<https://arxiv.org/abs/2403.11852>

91 CONF

HEART: Achieving Timely Multi-Model Training for Vehicle-Edge-Cloud-Integrated Hierarchical Federated Learning

ArXiv CS.AI 50%

arXiv:2501.09934v3 Announce Type: replace-cross Abstract: The rapid growth of AI-enabled Internet of Vehicles (IoV) calls for efficient Machine Learning (ML) solutions that can handle high vehicular mobility and decentralized data. This has motivated the emergence of...

<https://arxiv.org/abs/2501.09934>

92 CONF

Deepfake-Eval-2024: A Multi-Modal In-the-Wild Benchmark of Deepfakes Circulated in 2024

ArXiv CS.AI 50%

arXiv:2503.02857v5 Announce Type: replace-cross Abstract: In the age of increasingly realistic generative AI, robust deepfake detection is essential for mitigating fraud and disinformation. While many deepfake detectors report high accuracy on academic datasets, we show that...

<https://arxiv.org/abs/2503.02857>

93 CONF

The Point, the Vision and the Text: Does Point Cloud Boost Spatial Reasoning of Large Language Models? A Bias-Controlled Study

ArXiv CS.AI 50%

arXiv:2504.04540v2 Announce Type: replace-cross Abstract: 3D Large Language Models (LLMs) leveraging spatial information in point clouds for 3D spatial reasoning attract great attention. Despite some promising results, the advantages of point clouds over other modalities remain...

<https://arxiv.org/abs/2504.04540>

94 CONF

ASTRA: Communication-Efficient Acceleration for Multi-Device Transformer Inference

ArXiv CS.AI 50%

arXiv:2505.19342v2 Announce Type: replace-cross Abstract: Multi-device inference can reduce Transformer latency by parallelizing computation. However, existing methods require high inter-device bandwidth, making them impractical for bandwidth-constrained environments. We present...

<https://arxiv.org/abs/2505.19342>

95 CONF

Understanding Automated Program Repair Agents Through the Lens of Traceability: An Empirical Study

ArXiv CS.AI 50%

arXiv:2506.08311v2 Announce Type: replace-cross Abstract: Automated Program Repair (APR) agents leverage Large Language Models (LLMs) to autonomously diagnose and fix software bugs through reasoning, planning, and tool use. Despite impressive leaderboard gains on benchmarks such...

<https://arxiv.org/abs/2506.08311>

96 CONF

LLM Watermark Evasion via Bias Inversion

ArXiv CS.AI 50%

arXiv:2509.23019v5 Announce Type: replace-cross Abstract: Watermarking offers a promising solution for detecting LLM-generated content, yet its robustness under realistic query-free (black-box) evasion remains an open challenge. Existing query-free attacks often achieve limited...

<https://arxiv.org/abs/2509.23019>

97 CONF

Auditing medical multi-agent AI reveals risks of false consensus

ArXiv CS.AI 50%

arXiv:2510.10185v3 Announce Type: replace-cross Abstract: Large language models are increasingly being assembled into medical multi-agent systems that emulate multidisciplinary consultation through specialist roles, peer review and consensus formation. In clinical decision...

<https://arxiv.org/abs/2510.10185>

98 CONF

InfMed-ORBIT: Aligning LLMs on Open-Ended Complex Tasks via Rubric-Based Incremental Training

ArXiv CS.AI 50%

arXiv:2510.15859v4 Announce Type: replace-cross Abstract: Reinforcement learning (RL) has driven recent breakthroughs in large language models (LLMs), especially for tasks where rewards can be computed automatically, such as code generation. However, it is less effective in...

<https://arxiv.org/abs/2510.15859>

99 CONF

ReflexGrad: Within-Episode Failure Recovery in LLM Agents via Progress-Gated Dual-Process Routing

ArXiv CS.AI 50%

arXiv:2511.14584v3 Announce Type: replace-cross Abstract: We present ReflexGrad, a dual-process architecture for within-episode failure recovery in LLM agents without demonstrations. When agents commit to a wrong approach early and exhaust the step budget, the post-failure...

<https://arxiv.org/abs/2511.14584>

00 CONF

Evaluation of AI Ethics Tools in Language Models: A Developers' Perspective Case Study

ArXiv CS.AI 50%

arXiv:2512.15791v2 Announce Type: replace-cross Abstract: In Artificial Intelligence (AI), language models have gained significant importance due to the widespread adoption of systems capable of simulating realistic conversations with humans through text generation. Because of...

<https://arxiv.org/abs/2512.15791>

01 CONF

HGMEM: Hypergraph-based Working Memory to Improve Multi-step RAG for Long-Context Complex Relational Modeling

ArXiv CS.AI 50%

arXiv:2512.23959v3 Announce Type: replace-cross Abstract: Multi-step retrieval-augmented generation (RAG) has become a widely adopted strategy for enhancing large language models (LLMs) on tasks that demand global comprehension and intensive reasoning. Although many RAG systems...

<https://arxiv.org/abs/2512.23959>

02 CONF

On the Fallacy of Global Token Perplexity in Spoken Language Model Evaluation

ArXiv CS.AI 50%

arXiv:2601.06329v2 Announce Type: replace-cross Abstract: Generative spoken language models pretrained on large-scale raw audio can continue a speech prompt with appropriate content while preserving attributes like speaker and emotion, serving as foundation models for spoken...

<https://arxiv.org/abs/2601.06329>

03 CONF

The Grammar of Transformers: A Systematic Review of Interpretability Research on Syntactic Knowledge in Language Models

ArXiv CS.AI 50%

arXiv:2601.19926v2 Announce Type: replace-cross Abstract: We present a systematic review of 337 articles evaluating the syntactic abilities of Transformer-based language models (TLMs), reporting on over 3,000 datapoints spanning a wide range of syntactic phenomena, languages,...

<https://arxiv.org/abs/2601.19926>

04 CONF

Probability-Entropy Calibration: An Elastic Indicator for Adaptive Fine-tuning

ArXiv CS.AI 50%

arXiv:2602.01745v2 Announce Type: replace-cross Abstract: Token-level reweighting is a simple yet effective mechanism for controlling supervised fine-tuning, but common indicators are largely one-dimensional: the ground-truth probability reflects downstream alignment, while...

<https://arxiv.org/abs/2602.01745>

05 CONF

Semantic-level Backdoor Attack against Text-to-Image Diffusion Models

ArXiv CS.AI 50%

arXiv:2602.04898v3 Announce Type: replace-cross Abstract: Text-to-image (T2I) diffusion models are widely adopted for their strong generative capabilities, yet remain vulnerable to backdoor attacks. Existing attacks typically rely on fixed textual triggers and single-entity...

<https://arxiv.org/abs/2602.04898>

06 CONF

Singular Vectors of Attention Heads Align with Features

ArXiv CS.AI 50%

arXiv:2602.13524v2 Announce Type: replace-cross Abstract: Identifying feature representations in language models is a central task in mechanistic interpretability. Several recent studies have made the observation that feature representations can be inferred in some cases from...

<https://arxiv.org/abs/2602.13524>

07 CONF

Coherence Collapse: Diagnosing Why Code Agents Fail After Reaching the Right Code

ArXiv CS.AI 50%

arXiv:2603.24631v2 Announce Type: replace-cross Abstract: Code agents resolve 65-70% of SWE-bench Verified issues, but Pass@1 cannot tell us why the rest fail, and, as we show, capable-model failures are systematically misdiagnosed without trajectory data. We introduce TRAJEVAL,...

<https://arxiv.org/abs/2603.24631>

08 CONF

Early Decisions Matter: Proximity Bias and Initial Trajectory Shaping in Non-Autoregressive Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2604.10567v2 Announce Type: replace-cross Abstract: Diffusion-based language models (dLLMs) have emerged as a promising alternative to autoregressive language models, offering the potential for parallel token generation and bidirectional context modeling. However,...

<https://arxiv.org/abs/2604.10567>

09 CONF

BenGER Platform: A Collaborative Web Platform for End-to-End Benchmarking of German Legal Tasks

ArXiv CS.AI 50%

arXiv:2604.13583v3 Announce Type: replace-cross Abstract: Evaluating large language models (LLMs) for legal reasoning requires workflows that span task design, expert annotation, model execution, and metric-based evaluation. In practice, these steps are split across platforms...

<https://arxiv.org/abs/2604.13583>

10 CONF

When PCOS Meets Eating Disorders: An Explainable AI Approach to Detecting the Hidden Triple Burden

ArXiv CS.AI 50%

arXiv:2604.14356v2 Announce Type: replace-cross Abstract: Women with polycystic ovary syndrome (PCOS) face substantially elevated risks of body image distress, disordered eating, and metabolic challenges, yet existing natural language processing approaches for detecting these...

<https://arxiv.org/abs/2604.14356>

11 CONF

C-MORAL: Controllable Multi-Objective Molecular Optimization with Reinforcement Alignment for LLMs

ArXiv CS.AI 50%

arXiv:2604.23061v2 Announce Type: replace-cross Abstract: Large language models (LLMs) show promise for molecular optimization, but aligning them with selective and competing drug-design constraints remains challenging. We propose C-Moral, a reinforcement learning post-training...

<https://arxiv.org/abs/2604.23061>

12 CONF

Compliance versus Sensibility: On the Reasoning Controllability in Large Language Models

ArXiv CS.AI 50%

arXiv:2604.27251v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are known to acquire reasoning capabilities through shared inference patterns in pre-training data, which are further elicited via Chain-of-Thought (CoT) practices. However, whether...

<https://arxiv.org/abs/2604.27251>

13 CONF

Quantifying the Reconstructability of Astrophysical Methods with Large Language Models and Information Theory: A Case Study in Spectral Reconstruction

ArXiv CS.AI 50%

arXiv:2605.11154v2 Announce Type: replace-cross Abstract: Modern astrophysical studies rely heavily on complex data analysis pipelines; however, published descriptions often lack the detail required for computational reproducibility. In this work, we present an...

<https://arxiv.org/abs/2605.11154>

14 CONF

Structured Belief State and the First Precision-Aware Benchmark for LLM Memory Retrieval

ArXiv CS.AI 50%

arXiv:2605.11325v2 Announce Type: replace-cross Abstract: Every major benchmark for LLM memory systems, LoCoMo foremost, measures whether a model answered correctly, not whether the memory system retrieved correctly. A system returning its entire belief store achieves recall of...

<https://arxiv.org/abs/2605.11325>

15 CONF

ArcVQ-VAE: A Spherical Vector Quantization Framework with ArcCosine Additive Margin

ArXiv CS.AI 50%

arXiv:2605.13517v2 Announce Type: replace-cross Abstract: Vector Quantized Variational Autoencoder (VQ-VAE) has become a fundamental framework for learning discrete representations in image modeling. However, VQ-VAE models must tokenize entire images using a finite set of...

<https://arxiv.org/abs/2605.13517>

16 CONF

GoQuant: Geometric Orthogonal Residual Projection for Multiplier-Free Power-of-Two Transformer Quantization

ArXiv CS.AI 50%

arXiv:2605.26092v2 Announce Type: replace-cross Abstract: The deployment of Large Language Models (LLMs) and Vision Transformers (ViTs) on edge devices is significantly constrained by memory limitations and the critical timing bottlenecks introduced by dense Multiply-Accumulate...

<https://arxiv.org/abs/2605.26092>