



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	169	Media
TechCrunch AI	12	Media
HackerNews	5	Community
MIT Tech Review	4	Media
The Verge AI	4	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

The Download: climate tech goes public and the AI Hype Index returns

MIT Tech Review +1 source

80%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Climate tech companies are going public. What's next? Solar and battery company Solv Energy went public in February, hitting a \$6 billion...

<https://www.technologyreview.com/2026/05/28/1138085/the-download-climate-tech-ipos-ai-h...>

VERI

Governing Technical Debt in Agentic AI Systems

ArXiv CS.AI +1 source

80%

arXiv:2605.29129v1 Announce Type: new Abstract: Agentic AI systems are increasingly being explored as production infrastructure: they reason over multiple steps, call tools, act through workflows, and adapt through memory and feedback. These systems create governance challenges...

<https://arxiv.org/abs/2605.29129>

CONFIRMED SIGNALS — 190 stories

CONF

Show HN: Continue? Y/N: A 60-second game about AI agent permission fatigue

HackerNews 60% HN 293

Hi HN, I built VAEN (an open source CLI) because I kept running into a boring problem with AI coding-agent workflows: the setup becomes useful, but then it is hard to move. A good, useful agentic harness consists of more than just instructions: skills, mcp servers and more....

<https://llmgame.scalex.dev>

CONF

Show HN: Local CPU model for fact-checking, summarizing, explaining text

HackerNews 50% HN 6

No summary — see link for full article.

<https://github.com/kouhxp/fftext>

CONF

Show HN: http for the airwaves — a live 802.11 RF dashboard in your terminal

HackerNews 50% HN 5

No summary — see link for full article.

<https://github.com/yeet-src/airtop>

CONF

Show HN: Imagine Kickstarter, but doubters get paid if the project flops

HackerNews 50% HN 4

No summary — see link for full article.

<https://trypieces.com/p/as9wuvhp1qwb4NvpMGG2FU>

CONF

Show HN: AI Skill to port PostgreSQL extensions to MySQL

HackerNews 50% HN 3

Agent skills for working with VillageSQL. Skills run in Claude Code, Gemini CLI, agy, Codex, Cursor, Amp, and Kiro.

<https://github.com/villagesql/villagesql-skills>

CONF

Glean's top line crosses \$300M as AI budget-cutting becomes its major selling point

TechCrunch AI 50%

The enterprise AI search startup tripled its annual revenue even as tech giants entered the category.

<https://techcrunch.com/2026/05/28/gleans-top-line-crosses-300m-as-ai-budget-cutting-bec...>

CONF

Asana acquires no-code agent-builder StackAI

TechCrunch AI 50%

Asana will incorporate StackAI into its growing suite of AI workflow tools.

<https://techcrunch.com/2026/05/28/asana-acquires-no-code-agent-builder-stack-ai/>

0 CONF

Anthropic raises \$65 billion, nears \$1T valuation ahead of IPO

TechCrunch AI 50%

Anthropic has closed a \$65 billion Series H round at a \$965 billion post-money valuation, marking what could be the AI startup's final private fundraiser before a highly anticipated IPO.

<https://techcrunch.com/2026/05/28/anthropic-raises-65-billion-nears-1t-valuation-ahead-...>

1 CONF

Just like gold and oil, we'll soon be able to trade AI token futures

TechCrunch AI 50%

Large exchanges are designing derivative products around AI tokens, which are increasingly being considered less a computational output and more a raw material input, like electricity or bandwidth.

<https://techcrunch.com/2026/05/28/just-like-gold-and-oil-well-soon-be-able-to-trade-ai-...>

2 CONF

How long is Anthropic's lease with SpaceX? Opinions vary

TechCrunch AI 50%

Elon Musk is publicly reframing xAI's massive Anthropic compute deal as short-term and cancellable, despite SpaceX's own S-1 filing describing payments through May 2029.

<https://techcrunch.com/2026/05/28/how-long-is-anthropics-lease-with-spacex-opinions-vary/>

3 CONF

Sesame, the conversational AI startup from Oculus founders, launches its iOS app

TechCrunch AI 50%

Sesame's new iOS app brings its conversational AI agents to the public, offering more natural back-and-forth interactions designed to feel less like traditional chatbots and more like talking to a person.

<https://techcrunch.com/2026/05/28/sesame-the-conversational-ai-startup-from-oculus-foun...>

4 CONF

Sneak peek at new Siri app reveals Apple's plans to take on ChatGPT and more

TechCrunch AI 50%

New renders offer a closer look at Apple's planned AI overhaul for iOS 27, including a redesigned Siri experience and standalone Siri app.

<https://techcrunch.com/2026/05/28/sneak-peek-at-new-siri-app-reveals-apples-plans-to-ta...>

5 CONF

At TechCrunch Disrupt 2026: Databricks' co-founder on what kills enterprise AI deals

TechCrunch AI 50%

Enterprise AI is entering a different phase now, one where enterprises are no longer evaluating whether AI is exciting. They are evaluating whether it is safe to deploy broadly.

<https://techcrunch.com/2026/05/28/techcrunch-disrupt-2026-databricks-co-founder-on-what...>

6 CONF

YouTube adds new podcast features, including an AI recommendation tool and 'Auto speed'

TechCrunch AI 50%

The update signals YouTube's ongoing efforts to compete with other platforms for podcast audiences.

<https://techcrunch.com/2026/05/28/youtube-adds-new-podcast-features-including-an-ai-rec...>

7 CONF

Vertu wants CEOs to run companies from an AI foldable starting at \$6,880

TechCrunch AI 50%

Built on top of the open source Hermes project, Vertu's new foldable combines AI-agent workflows, enterprise integrations, and ultra-premium luxury finishes.

<https://techcrunch.com/2026/05/28/vertu-wants-ceos-to-run-companies-from-an-ai-foldable...>

8 CONF

Why Google's AI can't spell Google (or anything else)

TechCrunch AI 50%

Google is embarrassing itself, again.

<https://techcrunch.com/2026/05/27/why-googles-ai-cant-spell-google-or-anything-else/>

9 CONF

The AI Hype Index: AI gets booed in graduation season

MIT Tech Review 50%

It is one thing to say AI will change the world. It is another to expect the class of 2026 to applaud it. In fact, when former Google CEO Eric Schmidt told University of Arizona graduates that their task is to help shape AI, he was met with a resounding chorus of boos. "I can...

<https://www.technologyreview.com/2026/05/28/1138053/the-ai-hype-index-ai-gets-booed-in-...>

0 CONF

A reality check on the AI jobs hysteria

MIT Tech Review 50%

Haven't you heard? White-collar jobs are going away, decimated by AI. Waves of layoffs in the tech sector (most recently at Coinbase and Meta and Cisco) are said to presage what will soon come for all of us knowledge workers. But before you quit your job as a software developer...

<https://www.technologyreview.com/2026/05/26/1137855/a-reality-check-on-the-ai-jobs-hyst...>

1 CONF

Microsoft 365 Copilot gets a speed boost and cleaner design

The Verge AI 50%

Microsoft is launching a revamped version of Microsoft 365 Copilot, offering a cleaner design that the company claims loads twice as fast. As part of this update, Copilot will provide more reliable and structured responses that are easier to scan, according to Microsoft. The...

<https://www.theverge.com/tech/939273/microsoft-365-copilot-redesign>

2 CONF

Claude's new model is more 'honest' when it messes up

The Verge AI 50%

Anthropic is releasing Claude Opus 4.8 on Thursday, and the company is touting the model's "honesty." According to Anthropic, it trains "all [its] models to be honest - for instance, to avoid making claims that they can't support." But it notes that "a general problem with AI..."

<https://www.theverge.com/ai-artificial-intelligence/939094/anthropic-claude-4-8-opus-ho...>

3 CONF

A \$2,000 AI-generated film will make its debut at Tribeca

The Verge AI 50%

Next month's Tribeca Festival will include the premiere of an AI-generated film: Dreams of Violets. The 75-minute film is a fictional dramatization of the Iranian government's mass killing of protestors in January, with the people and images fully created by AI, as reported...

<https://www.theverge.com/entertainment/939067/ai-film-dreams-of-violets-tribeca>

4 CONF

YouTube will let you ask AI to make a custom video feed

The Verge AI 50%

YouTube is launching a new AI feature that creates a personalized video feed based on descriptions of what you want to watch. In its announcement, YouTube says custom content feeds can be built around your specific interests, moods, or favorite topics, which you can then pin to...

<https://www.theverge.com/streaming/938759/youtube-custom-ai-feed-prompt-availability>

5 CONF

The Cognitive Categorical Transformer: Category-Theoretic Inductive Biases for Language Modeling

ArXiv CS.AI 50%

arXiv:2605.28864v1 Announce Type: new Abstract: The Cognitive Categorical Transformer (CCT) is a 306M-parameter architecture that augments a pretrained GPT-2 Small backbone with cognitively grounded components derived from category theory and several inspirations from cognitive...

<https://arxiv.org/abs/2605.28864>

6 CONF

Ultra-Reduced-Impact-Encased-Logging (URIEL): propose a new method for selective sustainable logging and post-harvest silvicultural treatment in tropical forest using airborne robotics systems

ArXiv CS.AI 50%

arXiv:2605.28883v1 Announce Type: new Abstract: Tropical forests worldwide are under intense deforestation pressure driven by economic and political interests, and scientific evidence suggests this deforestation contributes to climate change. This paper proposes a novel logging...

<https://arxiv.org/abs/2605.28883>

7 CONF

Review Arcade: On the Human Alignment and Gameability of LLM Reviews

ArXiv CS.AI 50%

arXiv:2605.28897v1 Announce Type: new Abstract: LLM-generated reviews for scientific papers are gaining considerable traction and are even being officially piloted by major conferences. We have to assume that not only reviewers are using LLM-assistance, but also that authors use...

<https://arxiv.org/abs/2605.28897>

8 CONF

BioArc: Discovering Optimal Neural Architectures for Biological Foundation Models

ArXiv CS.AI 50%

arXiv:2512.00283v3 Announce Type: replace-cross Abstract: Foundation models have revolutionized various fields such as natural language processing (NLP) and computer vision (CV). While efforts have been made to transfer the success of the foundation models in general AI domains...

<https://arxiv.org/abs/2512.00283>

9 CONF

Frontier LLM-based agents can overcome the ontology curation bottleneck for natural phenotypes

ArXiv CS.AI 50%

arXiv:2605.28965v1 Announce Type: new Abstract: Linking free-text phenotype descriptions to ontology terms, typically referred to as phenotype annotation, is essential for the cross-study integration of comparative morphological data. This labor intensive process has heavily...

<https://arxiv.org/abs/2605.28965>

0 CONF

VFEAgent: A Multimodal Agent Framework for End-to-End Automated Finite Element Analysis

ArXiv CS.AI 50%

arXiv:2605.28978v1 Announce Type: new Abstract: Finite Element Analysis (FEA) serves as the cornerstone of modern engineering design. However, its workflow is inherently complex and relies heavily on domain expertise. Although recent efforts have integrated Large Language Models...

<https://arxiv.org/abs/2605.28978>

1 CONF

When Models Disagree: Rethinking LLM Evaluation for Public Comment Analysis

ArXiv CS.AI 50%

arXiv:2605.29025v1 Announce Type: new Abstract: Federal agencies are deploying large language models (LLMs) to categorize public comment corpora, where the model's organization of the record shapes what policymakers see and which arguments register. Standard evaluation, anchored...

<https://arxiv.org/abs/2605.29025>

2 CONF

Adapt \$\neq\$ Adapt: Longitudinal Analyses of LLM Conversations in the Wild

ArXiv CS.AI 50%

arXiv:2605.29018v1 Announce Type: new Abstract: Although a growing body of research has begun to describe user-LLM interactions, the picture it paints is largely static; little is known about how individual users change their behavior over time. To address this gap, we analyze...

<https://arxiv.org/abs/2605.29018>

3 CONF

Mind Your Tone: Does Tone Alter LLM Performance?

ArXiv CS.AI 50%

arXiv:2605.29027v1 Announce Type: new Abstract: The use of Large Language Models (LLMs) is proliferating, yet their performance is observed to vary based on prompting styles and tones. In this study, we investigate both whether and how tonal variations in prompts lead to...

<https://arxiv.org/abs/2605.29027>

4 CONF

Practitioner Beliefs and Behaviors in AI-Enhanced Education: DOT Framework Survey Evidence

ArXiv CS.AI 50%

arXiv:2605.29041v1 Announce Type: new Abstract: This study reports findings from a cross-sectional survey (n = 72) of higher education practitioners examining beliefs, behaviors, and institutional conditions related to artificial intelligence (AI) integration in teaching and...

<https://arxiv.org/abs/2605.29041>

5 CONF

KairosAgent: Agentic Time Series Forecasting with Fused Semantic Reasoning

ArXiv CS.AI 50%

arXiv:2605.30002v1 Announce Type: new Abstract: Cross-domain multimodal time series forecasting is a challenging task, requiring models to integrate precise numerical comprehension, cross-domain semantic understanding, and effective multimodal fusion. Existing approaches either...

<https://arxiv.org/abs/2605.30002>

6 CONF

Robust and Efficient Guardrails with Latent Reasoning

ArXiv CS.AI 50%

arXiv:2605.29068v1 Announce Type: new Abstract: Maintaining the safety of large language models (LLMs) is crucial as they are increasingly deployed in real-world applications. Existing safety guardrails typically rely on single-pass classification or, more recently, distilled...

<https://arxiv.org/abs/2605.29068>

7 CONF

Label-Free Reinforcement Learning via Cross-Model Entropy

ArXiv CS.AI 50%

arXiv:2605.29009v1 Announce Type: cross Abstract: Post-training large language models with reinforcement learning is bottlenecked by the reward signal. Existing approaches require either ground-truth verifiable rewards, restricting training to domains with automatic correctness...

<https://arxiv.org/abs/2605.29009>

8 CONF

The Importance of Out-of-Band Metadata for Safe Autonomous Agents: The Redpanda Agentic Data Plane

ArXiv CS.AI 50%

arXiv:2605.29082v1 Announce Type: new Abstract: AI agents are increasingly expected to operate as digital employees: accessing enterprise data, making decisions, and taking actions autonomously. But agents are simultaneously less predictable than humans -- prone to...

<https://arxiv.org/abs/2605.29082>

9 CONF

Reasoning While Asking: Transforming Reasoning Large Language Models from Passive Solvers to Proactive Inquirers

ArXiv CS.AI 50%

arXiv:2601.22139v2 Announce Type: replace-cross Abstract: Reasoning-oriented Large Language Models (LLMs) have achieved remarkable progress with Chain-of-Thought (CoT) prompting, yet they remain fundamentally limited by a \emph{blind self-thinking} paradigm: performing extensive...

<https://arxiv.org/abs/2601.22139>

0 CONF

Trends in AI and Human-AI Interaction in Clinical Trials -- A Hybrid Human-AI Exploration

ArXiv CS.AI 50%

arXiv:2605.29096v1 Announce Type: new Abstract: This paper examines records retrieved from the ClinicalTrials.gov registry to characterize temporal trends in AI terminology and the geographical distribution of AI trials. The work also reports on an exploratory hybrid human-AI...

<https://arxiv.org/abs/2605.29096>

1 CONF

Aryabhata 2: Scaling Reinforcement Learning for Advanced STEM Reasoning

ArXiv CS.AI 50%

arXiv:2605.28829v1 Announce Type: cross Abstract: Competitive STEM examinations such as JEE and NEET require multi-step symbolic reasoning, precise numerical computation, and deep conceptual understanding across physics, chemistry, and mathematics. Recent large language models...

<https://arxiv.org/abs/2605.28829>

2 CONF

InsightEval: An Expert-Curated Benchmark for Assessing Insight Discovery in LLM-Driven Data Agents

ArXiv CS.AI 50%

arXiv:2511.22884v2 Announce Type: replace Abstract: Data analysis has become an indispensable part of scientific research. To discover the latent knowledge and insights hidden within massive datasets, we need to perform deep exploratory analysis to realize their full value. With...

<https://arxiv.org/abs/2511.22884>

3 CONF

Tailoring the Curriculum: Student-Centered Reasoning Distillation via Dynamic Data-Model Compatibility

ArXiv CS.AI 50%

arXiv:2605.29229v1 Announce Type: new Abstract: Reasoning distillation transfers complex reasoning abilities from large language models (LLMs) to smaller ones, yet its success depends on how well the training data align with the student model. This paper introduces the...

<https://arxiv.org/abs/2605.29229>

4 CONF

Surfacing Isolated Learners with Outcome-Independent Mediation of Feedback between Teachers and Students Using AI

ArXiv CS.AI 50%

arXiv:2605.29240v1 Announce Type: new Abstract: AI-augmented classrooms generate rich teacher and student feedback before graded outcomes become available, yet these signals can be difficult to translate into timely instructional decisions. We propose an interpretable decision...

<https://arxiv.org/abs/2605.29240>

5 CONF

CrystalXRD-Bench: Benchmarking Vision-Language Models for XRD Peak Indexing Across Diverse Crystalline Materials

ArXiv CS.AI 50%

arXiv:2605.29446v1 Announce Type: new Abstract: Miller-index identification from powder XRD patterns requires capabilities untested by existing multimodal benchmarks: the model must read a narrow peak location from a rendered scientific curve and then connect that observation to...

<https://arxiv.org/abs/2605.29446>

6 CONF

Thinking Before Constraining: A Unified Decoding Framework for Large Language Models

ArXiv CS.AI 50%

arXiv:2601.07525v2 Announce Type: replace-cross Abstract: Natural generation allows Large Language Models (LLMs) to produce free-form responses with rich reasoning, yet the lack of structure makes outputs difficult to verify. Conversely, constrained decoding ensures standardized...

<https://arxiv.org/abs/2601.07525>

7 CONF

When and How Human Curation Backfires: Preference Alignment under Multi-Model Self-Consuming Loop

ArXiv CS.AI 50%

arXiv:2605.29267v1 Announce Type: new Abstract: Foundation models are increasingly trained on synthetic data generated by prior model iterations rather than exclusively on real data. This self-consuming training paradigm can lead to model collapse, divergence, or bias...

<https://arxiv.org/abs/2605.29267>

8 CONF

Indexing the Unreadable: LLM-Native Recursive Construction and Search of Service Taxonomies

ArXiv CS.AI 50%

arXiv:2605.29270v1 Announce Type: new Abstract: The era of the Internet of Agents (IoA) is taking shape: LLM agents are expected to fulfill user goals by orchestrating fast-growing populations of Model Context Protocol (MCP) servers, Agent-to-Agent (A2A) endpoints, reusable...

<https://arxiv.org/abs/2605.29270>

9 CONF

SCOPE: A Lightweight-training LLM Framework for Air Traffic Control Readback Monitoring

ArXiv CS.AI 50%

arXiv:2605.29543v1 Announce Type: cross Abstract: Pilot readback of Air Traffic Control (ATC) voice instructions is a primary safeguard against miscommunication in air transportation. However, readback anomalies remain implicated in approximately 80% of aviation incidents. This...

<https://arxiv.org/abs/2605.29543>

0 CONF

Entropy-KL Divergence-based Token Masking: A Novel Approach for Selective Fine-tuning of Large Language Models

ArXiv CS.AI 50%

arXiv:2605.29303v1 Announce Type: new Abstract: Supervised fine-tuning (SFT) followed by reinforcement learning (RL) has become a standard post-training paradigm for large language models. This paradigm provides a cold-start for RL exploration, avoiding the inefficiency of pure...

<https://arxiv.org/abs/2605.29303>

1 CONF

Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet

ArXiv CS.AI 50%

arXiv:2605.29358v1 Announce Type: new Abstract: We demonstrate that sparse autoencoders can extract interpretable features from Claude 3 Sonnet, a production-scale language model, addressing the open question of whether dictionary learning methods scale beyond small...

<https://arxiv.org/abs/2605.29358>

2 CONF

EvoMD-LLM: Learning the Language of Species Evolution in Reactive Molecular Dynamics

ArXiv CS.AI 50%

arXiv:2605.29394v1 Announce Type: new Abstract: While large language models (LLMs) excel at static scientific reasoning, they struggle to model the temporal structure of dynamic physical processes. We present EvoMD-LLM (Evolutionary Molecular Dynamics Large Language Model), a...

<https://arxiv.org/abs/2605.29394>

- 3
CONF
MuPHI: Learning Implicit Multimodal Harm Reasoning via Semantically Grounded Reward Optimization
ArXiv CS.AI 50%

arXiv:2605.29951v1 Announce Type: new Abstract: Understanding how harm emerges from interaction between otherwise benign image-text pairs requires intent-aware cross-modal reasoning beyond surface-level features. Existing vision-language models (VLMs) excel at literal reasoning...

<https://arxiv.org/abs/2605.29951>
- 4
CONF
Architecture-Sensitive Supervised Fine-Tuning for Screen-Conditioned Action Prediction: A PiSAR Benchmark
ArXiv CS.AI 50%

arXiv:2605.29400v1 Announce Type: new Abstract: We benchmark three supervised fine-tuned models against frontier zero-shot baselines on a 661-row held-out slice of PiSAR (Persona, intent, Screen, Action, Rationale), a 12,929-tuple corpus of screen-anchored behavioural rationales...

<https://arxiv.org/abs/2605.29400>
- 5
CONF
When Does Persona Prompting Actually Help? A Retrieval and Metric Analysis of Expert Role Injection in LLMs
ArXiv CS.AI 50%

arXiv:2605.29420v1 Announce Type: new Abstract: Persona prompting is widely used to steer large language models, yet its practical value remains unclear. Prior work often evaluates persona prompting using aggregate scores, making it difficult to determine whether expert-role...

<https://arxiv.org/abs/2605.29420>
- 6
CONF
Xetrieval: Mechanistically Explaining Dense Retrieval
ArXiv CS.AI 50%

arXiv:2605.29507v1 Announce Type: new Abstract: Explaining why dense retrievers assign high relevance scores remains challenging because retrieval decisions are made through opaque high-dimensional embeddings. Existing explanations often focus on surface signals, such as lexical...

<https://arxiv.org/abs/2605.29507>
- 7
CONF
Opt-Verifier: Unleashing the Power of LLMs for Optimization Modeling via Dual-Side Verification
ArXiv CS.AI 50%

arXiv:2605.29556v1 Announce Type: new Abstract: Building mathematical optimization models is critical in operations research (OR), while it requires substantial human expertise. Recent advancements have utilized large language models (LLMs) to automate this modeling process....

<https://arxiv.org/abs/2605.29556>
- 8
CONF
VRAG: Learning World Models for Interactive Video Generation
ArXiv CS.AI 50%

arXiv:2505.21996v4 Announce Type: replace-cross Abstract: Foundational world models must be both interactive and preserve spatiotemporal coherence for effective future planning with action choices. However, present models for long video generation have limited inherent world...

<https://arxiv.org/abs/2505.21996>

- 9 **CONF**
- Adaptive Interviewing for Persona Simulation in LLMs: Evidence-Grounded Reasoning Improves Decision Alignment**
- ArXiv CS.AI 50%
- arXiv:2605.29458v1 Announce Type: cross Abstract: Accurately simulating the decisions of a specific individual remains challenging for large language models (LLMs), partly because persona information is often provided as static descriptions that miss the values, experiences, and...
- <https://arxiv.org/abs/2605.29458>
-
- 0 **CONF**
- GPS-Enhanced Tourist Mobility Modeling with Seasonal Spatial Priors and LLM-Based Activity Chain Generation**
- ArXiv CS.AI 50%
- arXiv:2605.29578v1 Announce Type: new Abstract: Tourist mobility poses a distinct challenge for urban transportation planning. Unlike resident commuting, tourist travel is largely non-routine, attraction driven, and highly sensitive to trip purpose, travel season, and trip...
- <https://arxiv.org/abs/2605.29578>
-
- 1 **CONF**
- Beyond Attack Success Rate: Temporal Logit Observability for LLM Safety Failures**
- ArXiv CS.AI 50%
- arXiv:2605.29629v1 Announce Type: new Abstract: Attack Success Rate (ASR) evaluates each jailbreak with a single yes/no label at the end of generation, telling us whether a failure happened but not how it unfolded. Two attacks that produce equally harmful outputs may have...
- <https://arxiv.org/abs/2605.29629>
-
- 2 **CONF**
- VikingMem: A Memory Base Management System for Stateful LLM-based Applications**
- ArXiv CS.AI 50%
- arXiv:2605.29640v1 Announce Type: new Abstract: Large Language Models have revolutionized interactive applications; however, their finite context windows pose a critical data management challenge for maintaining stateful, long-term interactions. Existing memory approaches often...
- <https://arxiv.org/abs/2605.29640>
-
- 3 **CONF**
- LLM-Evolved Domain-Independent Heuristics for Symbolic AI Planning**
- ArXiv CS.AI 50%
- arXiv:2605.29649v1 Announce Type: new Abstract: Heuristic search is the dominant paradigm in symbolic AI planning, and the strongest heuristics are the result of decades of work by planning researchers. Recent work has shown that large language models (LLMs) can design...
- <https://arxiv.org/abs/2605.29649>
-
- 4 **CONF**
- PTCG-Bench: Can LLM Agents Master Pok\`emon Trading Card Game?**
- ArXiv CS.AI 50%
- arXiv:2605.29653v1 Announce Type: new Abstract: Given a strategically complex board game, human players can quickly learn to devise strategies after playing a few rounds. Autonomous agents require similar capabilities in realistic interactive environments, yet existing agent...
- <https://arxiv.org/abs/2605.29653>
-

5 CONF

TRACE: Toulmin-based Reasoning Assessment through Constructive Elements for LLM CoT Evaluation

ArXiv CS.AI 50%

arXiv:2605.29656v1 Announce Type: new Abstract: Evaluating open-ended outputs from large language models (LLMs) remains challenging due to the absence of ground truth. Existing metrics rely on final-answer accuracy or surface-level statistics, leaving the reasoning process...

<https://arxiv.org/abs/2605.29656>

6 CONF

HARP: Hadamard-Preconditioned Adaptive Rotation Processor for Extreme LLM Quantization

ArXiv CS.AI 50%

arXiv:2605.29843v1 Announce Type: cross Abstract: Post-training quantization (PTQ) is essential for deploying LLMs under memory and bandwidth constraints. However, extreme low-bit quantization remains highly sensitive to activation outliers and anisotropic weight curvature....

<https://arxiv.org/abs/2605.29843>

7 CONF

Dataset-Driven Channel Masks in Transformers for Multivariate Time Series

ArXiv CS.AI 50%

arXiv:2410.23222v4 Announce Type: replace-cross Abstract: Recent advancements in foundation models have been successfully extended to the time series (TS) domain, facilitated by the emergence of large-scale TS datasets. However, previous efforts have primarily Capturing channel...

<https://arxiv.org/abs/2410.23222>

8 CONF

Alignment-Guided Score Matching for Text-to-Image Alignment in Diffusion Models

ArXiv CS.AI 50%

arXiv:2605.30038v1 Announce Type: cross Abstract: Diffusion models generate highly realistic images but often struggle with precise text-image alignment. While recent post-training methods improve alignment using external rewards or human preference signals, their performance...

<https://arxiv.org/abs/2605.30038>

9 CONF

Uncertainty-Aware Transfer Learning for Cross-Building Energy Forecasting: Toward Robust and Scalable District-Level Energy Management

ArXiv CS.AI 50%

arXiv:2605.29733v1 Announce Type: new Abstract: Scaling data-driven energy forecasting to district level requires models that can be re-used across buildings with minimal target-domain data and honest uncertainty estimates. We present an uncertainty-aware transfer learning (TL)...

<https://arxiv.org/abs/2605.29733>

0 CONF

Why Specialist Models Still Matter: A Heterogeneous Multi-Agent Paradigm for Medical Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2605.29744v1 Announce Type: new Abstract: The impressive performance of generalist large language models (LLMs) such as GPT and Claude in healthcare raises a critical question: will domain-specific medical specialist models become obsolete? We argue that the future of...

<https://arxiv.org/abs/2605.29744>

1 **CONF****A comparative study of transformer-based embeddings for topic coherence****ArXiv CS.AI** **50%**

arXiv:2605.28832v1 Announce Type: cross Abstract: Topic modeling is a branch of Natural Language Processing (NLP) that aims to organize large collections of texts into coherent groups according to word co-occurrence patterns, with Latent Dirichlet Allocation (LDA) remaining one...

<https://arxiv.org/abs/2605.28832>

2 **CONF****LFQ: Logit-aware Final-block Quantization for Boosting the Generation Quality of Low-Bit Quantized LLMs****ArXiv CS.AI** **50%**

arXiv:2605.29756v1 Announce Type: new Abstract: As large language models continue to scale, low-bit weight-only post-training quantization (PTQ) offers a practical solution to their memory-efficient deployment. Although block-wise PTQ is capable of matching the full-precision...

<https://arxiv.org/abs/2605.29756>

3 **CONF****Croissant Tasks: A Metadata Format for Reproducible Machine Learning Evaluations****ArXiv CS.AI** **50%**

arXiv:2605.29786v1 Announce Type: new Abstract: Reproducibility is fundamental to the scientific method, yet remains a critical challenge in machine learning. Contributing factors include underspecified execution details and brittle software environments. Human-centric remedies,...

<https://arxiv.org/abs/2605.29786>

4 **CONF****SkillsInjector: Dynamic Skill Context Construction for LLM Agents****ArXiv CS.AI** **50%**

arXiv:2605.29794v1 Announce Type: new Abstract: LLM agents now draw on growing skill libraries to handle complex tasks. However, injecting more skills does not always improve task completion and can even degrade it. Existing methods still treat skill injection as a static step,...

<https://arxiv.org/abs/2605.29794>

5 **CONF****MEMENTO: Leveraging Web as a Learning Signal for Low-Data Domains****ArXiv CS.AI** **50%**

arXiv:2605.29795v1 Announce Type: new Abstract: Real-world tasks often lack large labeled datasets, motivating extensive work on learning in low-data regimes. However, existing approaches such as few-shot prompting, instruction tuning, and synthetic data generation, continue to...

<https://arxiv.org/abs/2605.29795>

6 **CONF****PRAIB: Peer Review AI Benchmark of Behaviour of LLM-Assisted Reviewing****ArXiv CS.AI** **50%**

arXiv:2605.29815v1 Announce Type: new Abstract: The growing number of submitted papers has motivated the exploration of Large Language Models (LLMs) as a means to support and augment the peer review process, particularly in terms of improving its speed and scalability. Yet, it...

<https://arxiv.org/abs/2605.29815>

7 CONF

OmniMatBench: A Human-Calibrated Multimodal Reasoning Benchmark Across 19 Materials Science Subfields

ArXiv CS.AI 50%

arXiv:2605.29833v1 Announce Type: new Abstract: As multimodal language models play an increasingly important role in scientific research, materials science offers a critical testbed due to its interdisciplinary, multimodal, and application-driven nature. However, existing...

<https://arxiv.org/abs/2605.29833>

8 CONF

Toward AI Systems That Understand Self and Others: A Multi-Phase Inference Framework for Human Cognitive Diversity and World-Model Alignment

ArXiv CS.AI 50%

arXiv:2605.29930v1 Announce Type: new Abstract: Mutual misunderstanding in contemporary society does not arise merely because people hold different opinions or values. Even under the same observations, different subjects may form different inferential targets, state...

<https://arxiv.org/abs/2605.29930>

9 CONF

It's All About Speed: AI's Impact on Workflow in Music Production

ArXiv CS.AI 50%

arXiv:2605.29931v1 Announce Type: new Abstract: In this paper, we present the results of an ethnographic study into the impact of AI and automated tools on music production workflow. Focusing specifically on professional participants who identified as recording engineers,...

<https://arxiv.org/abs/2605.29931>

0 CONF

Make LLM Learn to Synthesize from Streaming Experiences through Feedback

ArXiv CS.AI 50%

arXiv:2605.29940v1 Announce Type: new Abstract: Large language models (LLMs) have been widely adopted for synthetic data generation, significantly reducing annotation costs. However, most existing studies treat synthesis as a set of isolated tasks and overlook a more fundamental...

<https://arxiv.org/abs/2605.29940>

1 CONF

Balancing Multimodal Learning through Label Space Reshaping

ArXiv CS.AI 50%

arXiv:2605.28869v1 Announce Type: cross Abstract: Multimodal learning often suffers from modality imbalance, where modalities that converge faster dominate optimization while others remain undertrained. Existing approaches typically mitigate this issue by strengthening the weak...

<https://arxiv.org/abs/2605.28869>

2 CONF

From GPS Points to Travel Patterns: Flexible and Semantic Trajectory Generation with LLMs

ArXiv CS.AI 50%

arXiv:2605.30014v1 Announce Type: new Abstract: Urban trajectories play a crucial role in modeling urban dynamics and supporting various smart city applications. However, privacy concerns restrict access to large-scale and high-quality trajectory datasets. Trajectory generation...

<https://arxiv.org/abs/2605.30014>

3 CONF

RAISE: RAG Design as an Architecture Search Problem

ArXiv CS.AI 50%

arXiv:2605.30029v1 Announce Type: new Abstract: Retrieval-augmented generation (RAG) systems expose numerous design choices spanning query rewriting, chunking, retrieval depth, reranking, and context compression. In practice, these choices are often configured through...

<https://arxiv.org/abs/2605.30029>

4 CONF

Teaching Values to Machines: Simulating Human-Like Behavior in LLMs

ArXiv CS.AI 50%

arXiv:2605.30036v1 Announce Type: new Abstract: Large Language Models (LLMs) demonstrate a remarkable capacity to adopt different personas and roles; however, it remains unclear whether they can manifest behavior that adheres to a coherent, human-like value structure. In this...

<https://arxiv.org/abs/2605.30036>

5 CONF

Domain-Specific Data Synthesis for LLMs via Minimal Sufficient Representation Learning

ArXiv CS.AI 50%

arXiv:2605.30039v1 Announce Type: new Abstract: Large Language Models have demonstrated remarkable progress in general-purpose capabilities and can achieve strong performance in specific domains through fine-tuning on domain-specific data. However, acquiring high-quality data...

<https://arxiv.org/abs/2605.30039>

6 CONF

Robust and Generalizable Safety Steering for Text-to-Image Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2605.30049v1 Announce Type: new Abstract: Diffusion Transformers have become a powerful backbone for text-to-image generation, but their layered and cross-modal generation process makes safety control fundamentally different from prompt-level filtering or output-level...

<https://arxiv.org/abs/2605.30049>

7 CONF

PokerSkill: LLMs Can Play Expert-Level Poker without Training or Solvers

ArXiv CS.AI 50%

arXiv:2605.30094v1 Announce Type: new Abstract: Poker is a landmark challenge for artificial intelligence. The dominant approach relies on equilibrium solvers built on counterfactual regret minimization, requiring millions of core-hours of training. Large Language Models (LLMs)...

<https://arxiv.org/abs/2605.30094>

8 CONF

Beyond Recall: Behavioral Specification as an Interpretive Layer for AI Personalization

ArXiv CS.AI 50%

arXiv:2605.28969v1 Announce Type: cross Abstract: If an AI agent makes decisions on a person's behalf, those decisions must align with its user. We introduce representational accuracy to measure how faithfully a system captures a person's interpretation. An interpretive layer is...

<https://arxiv.org/abs/2605.28969>

- 9 **CONF**
BioRefusalAudit: Auditing Biosecurity Refusal Depth Using General and Domain-Fine-Tuned Sparse Autoencoders
 ArXiv CS.AI 50%
 arXiv:2605.30162v1 Announce Type: new Abstract: Biosecurity evaluations of language models typically ask whether models produce hazardous output. This paper asks a complementary question: when a model refuses, is that refusal structurally sound, or does it disappear under modest...
<https://arxiv.org/abs/2605.30162>
- 0 **CONF**
Combating Data Laundering in LLM Training
 ArXiv CS.AI 50%
 arXiv:2604.01904v2 Announce Type: replace-cross Abstract: Data rights owners can detect unauthorized data use in large language model (LLM) training by querying with proprietary samples. Often, superior performance (e.g., higher confidence or lower loss) on a sample relative to...
<https://arxiv.org/abs/2604.01904>
- 1 **CONF**
Double-Edged Sword or Sharp Tool? Designing and Evaluating Triadic LLM-Teacher Collaboration for K-12 Writing at Scale
 ArXiv CS.AI 50%
 arXiv:2605.30200v1 Announce Type: new Abstract: The double-edged sword of integrating Large Language Models (LLMs) requires an effective triadic collaboration mechanism among LLMs, teachers and students, especially for K-12 education. By developing a triadic collaboration system...
<https://arxiv.org/abs/2605.30200>
- 2 **CONF**
When Models Learn to Ask Why: Adaptive Causal Reasoning for Trustworthy Medical Vision-Language Models
 ArXiv CS.AI 50%
 arXiv:2603.23085v2 Announce Type: replace Abstract: Vision-Language Models (VLMs) have enabled interpretable medical diagnosis by integrating visual perception with linguistic reasoning. Yet, existing medical chain-of-thought (CoT) models lack explicit mechanisms to represent...
<https://arxiv.org/abs/2603.23085>
- 3 **CONF**
ProjectionBench: Evaluating Scientific Hypothesis Generation in LLMs Under Progressive Information Disclosure
 ArXiv CS.AI 50%
 arXiv:2605.30284v1 Announce Type: new Abstract: Scientific discovery is an inherently creative and uncertain process, requiring reasoning beyond the recall of known knowledge. While many benchmarks have been proposed to evaluate large language model (LLM) performance on deep...
<https://arxiv.org/abs/2605.30284>
- 4 **CONF**
Locally Coherent, Globally Incoherent: Bounding Compositional Incoherence in Multi-Component LLM Agents
 ArXiv CS.AI 50%
 arXiv:2605.30335v1 Announce Type: new Abstract: Multi-component LLM agents assemble probabilistic claims from components that each see only part of a joint problem; the composition can violate basic probability axioms even when every component is locally coherent. We formalise...
<https://arxiv.org/abs/2605.30335>

- 5 **CONF**
- Physics Is All You Need? A Case Study in Physicist-Supervised AI Development of Scientific Software**
- ArXiv CS.AI 50%
- arXiv:2605.30353v1 Announce Type: new Abstract: Are AI agents tools, co-authors, or researchers? We present a quantified case study (\$N=1\$): a physicist supervising an AI coding agent (Claude Code, Sonnet and Opus models) over 12 work days and 57 sessions to build CLAX-PT, a...
- <https://arxiv.org/abs/2605.30353>
-
- 6 **CONF**
- Audio Jailbreaks in Large Audio-Language Models: Taxonomy, Attack-Defense Analysis, and Cost-Aware Evaluation**
- ArXiv CS.AI 50%
- arXiv:2605.30031v1 Announce Type: cross Abstract: Large Audio Language Models (LALMs) expand jailbreak risks from token-level prompting to the full speech perception-to-reasoning pipeline, where unsafe behavior can be induced through semantics, acoustic style, signal artifacts,...
- <https://arxiv.org/abs/2605.30031>
-
- 7 **CONF**
- Transcribing Children's Speech: ASR Performance and Obtaining Reliable Orthographic Transcriptions**
- ArXiv CS.AI 50%
- arXiv:2605.28833v1 Announce Type: cross Abstract: Automatic speech recognition (ASR) has the potential to substantially reduce manual annotation effort in child speech research by generating automatic transcriptions. However, obtaining reliably high-quality ASR transcriptions...
- <https://arxiv.org/abs/2605.28833>
-
- 8 **CONF**
- Assessing Dutch Syllabification Algorithms and Improving Accuracy by Combining Phonetic and Orthographic Information through Deep Learning**
- ArXiv CS.AI 50%
- arXiv:2605.28834v1 Announce Type: cross Abstract: Syllabification describes the task of dividing words into syllables. Due to many rules and exceptions, training an algorithm to perform syllabification with high accuracy remains a challenge. Throughout the last decades,...
- <https://arxiv.org/abs/2605.28834>
-
- 9 **CONF**
- Thoughts-as-Planning: Latent World Models for Chain-of-Thoughts Optimization via Reinforcement Planning**
- ArXiv CS.AI 50%
- arXiv:2605.28842v1 Announce Type: cross Abstract: The success of large language models (LLMs) across diverse NLP tasks has elevated the importance of reasoning chain optimization as a critical step in aligning model behavior with task objectives. Existing reasoning chain tuning...
- <https://arxiv.org/abs/2605.28842>
-
- 00 **CONF**
- Reinforcement Learning with Robust Rubric Rewards**
- ArXiv CS.AI 50%
- arXiv:2605.30244v1 Announce Type: cross Abstract: While Reinforcement Learning with Verifiable Rewards (RLVR) is effective for deterministically checkable tasks, many vision-language tasks are partially verifiable, demanding multi-criteria supervision (e.g., perceptual details,...
- <https://arxiv.org/abs/2605.30244>

01 CONF

TaxDistill: Improving Metagenomic Taxonomic Annotation via Distilled Genomic Foundation Models

ArXiv CS.AI 50%

arXiv:2605.28868v1 Announce Type: cross Abstract: Metagenomic taxonomic annotation aims to identify the microbial origins of DNA fragments in environmental samples. Traditional methods that rely on sequence similarity are often constrained by the high microbial diversity and the...

<https://arxiv.org/abs/2605.28868>

02 CONF

Quantum-Enhanced Adversarial Robustness in Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2605.28899v1 Announce Type: cross Abstract: Artificial Intelligence has achieved remarkable success across diverse application domains. However, its vulnerability to adversarial attacks poses significant challenges to reliability, security, and trustworthiness. Adversarial...

<https://arxiv.org/abs/2605.28899>

03 CONF

AIRGuard: Guarding Agent Actions with Runtime Authority Control

ArXiv CS.AI 50%

arXiv:2605.28914v1 Announce Type: cross Abstract: Tool-using language agents turn model decisions into external side effects: they read files, run scripts, call APIs, send messages, and invoke Model Context Protocol tools. This makes agent attacks different from jailbreaks. The...

<https://arxiv.org/abs/2605.28914>

04 CONF

First head-to-head comparison of agentic AI applied to the analysis of simulated data of the Einstein Telescope

ArXiv CS.AI 50%

arXiv:2605.28916v1 Announce Type: cross Abstract: We report a comparison of two state-of-the-art agentic AI systems, Claude Code (Anthropic) and Codex (OpenAI), tasked with autonomously executing a simple end-to-end gravitational wave data analysis pipeline on a shared computing...

<https://arxiv.org/abs/2605.28916>

05 CONF

CosmicFish-HRM: Adaptive Reasoning via Hierarchical Recurrent Mechanisms in Compact Language Models

ArXiv CS.AI 50%

arXiv:2605.28919v1 Announce Type: cross Abstract: Large language models have achieved strong reasoning capabilities, though often at the cost of massive parameter counts and expensive inference. In this work, we explore a different direction: adaptive reasoning depth in compact...

<https://arxiv.org/abs/2605.28919>

06 CONF

The Hamilton-Jacobi Theory of Deep Learning

ArXiv CS.AI 50%

arXiv:2605.28983v1 Announce Type: cross Abstract: In this paper, training a neural network is identified, exactly, as a search through Hamilton--Jacobi initial-value problems: each gradient step selects the initial data of a viscous Hamilton--Jacobi equation whose Hopf--Cole...

<https://arxiv.org/abs/2605.28983>

07 CONF

TIMEGATE: Sustainable Time-Boxed Promotion Gates for Continual ML Adaptation Under Resource Constraints

ArXiv CS.AI 50%

arXiv:2605.29183v1 Announce Type: cross Abstract: As machine learning(ML) systems evolve to continual adaptation, each re-training cycle uses compute, annotation, and energy. We introduce TIMEGATE, a policy layer managing adaptation by budgeting time, labeling, training, and...

<https://arxiv.org/abs/2605.29183>

08 CONF

FormInv: A Measurement Protocol for Semantic Invariance in Mathematical Reasoning Benchmarks

ArXiv CS.AI 50%

arXiv:2605.29001v1 Announce Type: cross Abstract: A paraphrase-quality audit of MathCheck (ICLR 2025) detected 4 semantically incorrect paraphrases in 129 groups (3.1%); removing them drops GPT-4o from rank 2 to rank 4 and elevates Claude Haiku and DeepSeek V3 above it; these...

<https://arxiv.org/abs/2605.29001>

09 CONF

Return-to-Go Is More Than a Number: Q-Guided Alignment for Return-Conditioned Supervised Learning

ArXiv CS.AI 50%

arXiv:2605.29028v1 Announce Type: cross Abstract: Conditioned Sequence Models (CSMs) learn policies by treating return-to-go (RTG) as a control signal. However, existing CSMs often treat the RTGs as simple numerical inputs rather than aligning them with the performance of their...

<https://arxiv.org/abs/2605.29028>

10 CONF

Structured Prompt Optimization Meets Reinforcement Learning for Global and Local Interpretability over Complex Text

ArXiv CS.AI 50%

arXiv:2605.29076v1 Announce Type: cross Abstract: LLMs have advanced text classification, yet existing paradigms face a trade-off: supervised (label only) fine-tuning is scalable but offers limited reasoning on complex text and lacks broader model transparency, while discrete...

<https://arxiv.org/abs/2605.29076>

11 CONF

Same Evidence, Different Answers: Canonical-Context On-Policy Distillation for Multi-Turn Language Models

ArXiv CS.AI 50%

arXiv:2605.30251v1 Announce Type: cross Abstract: Large language models (LLMs) often solve a task when all instructions are given in a single prompt, but fail when the same information is revealed gradually across turns. When a clean FULL prompt and a RAW-SHARDED conversation...

<https://arxiv.org/abs/2605.30251>

12 CONF

Good SFT Optimizes for SFT, Better SFT Prepares for Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2602.01058v2 Announce Type: replace-cross Abstract: Post-training of reasoning LLMs is a holistic process that typically consists of an offline SFT stage followed by an online reinforcement learning (RL) stage. However, SFT is often optimized in isolation to maximize SFT...

<https://arxiv.org/abs/2602.01058>

13 CONF

KLAS: Using Similarity to Stitch Neural Networks for Improved Accuracy-Efficiency Tradeoffs

ArXiv CS.AI 50%

arXiv:2605.29259v1 Announce Type: cross Abstract: Given the wide range of deployment targets, flexible model selection is essential for optimizing performance within a given compute budget. Recent work demonstrates that stitching pretrained models within a model family enables...

<https://arxiv.org/abs/2605.29259>

14 CONF

Toward User Preference Alignment in LLM Recommendation via Explicit Context Feedback

ArXiv CS.AI 50%

arXiv:2605.29141v1 Announce Type: cross Abstract: Traditional recommender systems (RecSys) primarily infer user preferences from implicit signals (such as clicks, watches, and purchases), often neglecting the rich explicit contextual feedback users provide through verbal text,...

<https://arxiv.org/abs/2605.29141>

15 CONF

Real-rootedness of the Poincaré polynomials of $\overline{\mathcal{M}}_{0,n}$: an AI-assisted proof

ArXiv CS.AI 50%

arXiv:2605.29151v1 Announce Type: cross Abstract: We prove real-rootedness for the Poincaré polynomial $[P_n(t) = \sum_{i=0}^{n-3} \dim H^{\{2i\}}(\overline{\mathcal{M}}_{0,n}; \mathbb{Q}) t^i]$ of the Deligne--Mumford moduli space $\overline{\mathcal{M}}_{0,n}$ of stable n -pointed...

<https://arxiv.org/abs/2605.29151>

16 CONF

Unveiling Multi-regime Patterns in SciML: Distinct Failure Modes and Regime-specific Optimization

ArXiv CS.AI 50%

arXiv:2605.29153v1 Announce Type: cross Abstract: Neural networks trained under different hyperparameter settings can fall into distinct training "regimes," with consistent behavior within regimes and qualitative differences across regimes. In this paper, we study such...

<https://arxiv.org/abs/2605.29153>

17 CONF

Parallax: Parameterized Local Linear Attention for Language Modeling

ArXiv CS.AI 50%

arXiv:2605.29157v1 Announce Type: cross Abstract: Large Language Models (LLMs) have become the central paradigm in artificial intelligence, yet the core computational primitive of attention has remained structurally unchanged. Local Linear Attention (LLA) is an attention...

<https://arxiv.org/abs/2605.29157>

18 CONF

Evolutionary Refinement of Generative Graph Topologies: A Hybrid WGAN-GA Approach

ArXiv CS.AI 50%

arXiv:2605.29161v1 Announce Type: cross Abstract: Generating realistic graph-structured data is challenging due to discrete connectivity, varying graph sizes, and class-specific structural patterns. Recent Generative Adversarial Networks (GAN)-based graph generation methods...

<https://arxiv.org/abs/2605.29161>

19 CONF

Domain-Informed Representation for Evolutionary Sieving in Integral and Module Lattices

ArXiv CS.AI 50%

arXiv:2605.29169v1 Announce Type: cross Abstract: Traditional cryptography, rooted in problems, e.g., integer factorisation or discrete log, is inevitably vulnerable to a fully operational quantum computer. Although it remains an engineering frontier, the looming threat extends...

<https://arxiv.org/abs/2605.29169>

20 CONF

Influence-Guided Symbolic Regression: Scientific Discovery via LLM-Driven Equation Search with Granular Feedback

ArXiv CS.AI 50%

arXiv:2605.29184v1 Announce Type: cross Abstract: Large Language Models (LLMs) offer a promising avenue for scientific discovery, yet their application to symbolic regression is often constrained by inefficient search strategies and coarse feedback signals. Current methods...

<https://arxiv.org/abs/2605.29184>

21 CONF

Relevance as a Vulnerability: How Web Retrieval Degrades Safety Alignment in LLM Agents

ArXiv CS.AI 50%

arXiv:2605.29224v1 Announce Type: cross Abstract: AI agents augment large language models with external tools such as web retrieval, enabling grounded and up-to-date responses. However, incorporating external content into the generation pipeline can weaken the safety alignment...

<https://arxiv.org/abs/2605.29224>

22 CONF

Toward Ethical Facial Age Estimation: A Generalized Zero-Shot Benchmark Without Training on Children's Data

ArXiv CS.AI 50%

arXiv:2605.29230v1 Announce Type: cross Abstract: Age estimation from facial images typically relies on training data that includes images of minors, a practice that raises serious ethical, legal, and privacy concerns. In this work, we propose a generalized zero-shot benchmark...

<https://arxiv.org/abs/2605.29230>

23 CONF

Do Physics Foundation Models Learn Generalizable Physics? A Bias-Aware Benchmark Across Physical Regimes and Distribution Shifts

ArXiv CS.AI 50%

arXiv:2605.29283v1 Announce Type: cross Abstract: Recent physics foundation models claim general spatiotemporal forecasting ability, yet their evaluations often collapse performance into a single average score under a fixed training distribution. This makes it difficult to...

<https://arxiv.org/abs/2605.29283>

24 CONF

SCoOP: Semantic Consistent Opinion Pooling for Uncertainty Quantification in Multiple Vision-Language Model Systems

ArXiv CS.AI 50%

arXiv:2603.23853v3 Announce Type: replace Abstract: Combining multiple Vision-Language Models (VLMs) can enhance multimodal reasoning and robustness, but aggregating heterogeneous models' outputs amplifies uncertainty and increases the risk of hallucinations. We propose SCoOP...

<https://arxiv.org/abs/2603.23853>

25 CONF

GrepSeek: Training Search Agents for Direct Corpus Interaction

ArXiv CS.AI 50%

arXiv:2605.29307v1 Announce Type: cross Abstract: Large Language Model (LLM) search agents have shown strong promise for knowledge-intensive language tasks through multiple rounds of reasoning and information retrieval. Most existing systems access information using a retriever...

<https://arxiv.org/abs/2605.29307>

26 CONF

Does Distributed Training Undermine Compute Governance?

ArXiv CS.AI 50%

arXiv:2605.29359v1 Announce Type: cross Abstract: Compute governance proposals often rely on the assumption that frontier AI training requires large, detectable computing clusters. However, recent advances in distributed training algorithms could allow developers to conduct...

<https://arxiv.org/abs/2605.29359>

27 CONF

TRACER: Persistent Regularization for Robust Multimodal Finetuning

ArXiv CS.AI 50%

arXiv:2605.29380v1 Announce Type: cross Abstract: Mainstream strategies for finetuning pretrained multimodal models often degrade out-of-distribution (OOD) robustness, a phenomenon known as catastrophic forgetting. In this paper, we develop a theoretical framework for multimodal...

<https://arxiv.org/abs/2605.29380>

28 CONF

Latent Terms: Dense Retrievers Contain Trivially Extractable BM25-ready Zipfian Vocabularies

ArXiv CS.AI 50%

arXiv:2605.29384v1 Announce Type: cross Abstract: We propose Latent Terms, a method revealing that models trained for dense retrieval, whether single- or multi-vector, learn representations that can trivially be decomposed into retrieval-ready sparse features. When trained on...

<https://arxiv.org/abs/2605.29384>

29 CONF

AliMark: Enhancing Robustness of Sentence-Level Watermarking Against Text Paraphrasing

ArXiv CS.AI 50%

arXiv:2605.29434v1 Announce Type: cross Abstract: Existing sentence-level watermarking methods enhance robustness to paraphrasing by anchoring watermarks in sentence semantics. However, their prefix-based designs remain vulnerable to structural perturbations, such as sentence...

<https://arxiv.org/abs/2605.29434>

30 CONF

How Coding Agents Fail Their Users: A Large-Scale Analysis of Developer-Agent Misalignment in 20,574 Real-World Sessions

ArXiv CS.AI 50%

arXiv:2605.29442v1 Announce Type: cross Abstract: AI coding agents increasingly act directly within software environments, yet existing analyses of their failures rely on benchmark trajectories that miss how developers actually experience misalignment. We present an...

<https://arxiv.org/abs/2605.29442>

31 CONF

SciIntBench: Measuring LLM Compliance with Research Integrity Norms Under Adversarial Framing

ArXiv CS.AI 50%

arXiv:2605.29468v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to support scientific work, but it is unclear whether they uphold responsible conduct of research (RCR) norms or help undermine them. We introduce SciIntBench, an adversarial...

<https://arxiv.org/abs/2605.29468>

32 CONF

Inform, Coach, Relate, Listen: Auditing LLM Caregiving Support Roles

ArXiv CS.AI 50%

arXiv:2605.29473v1 Announce Type: cross Abstract: Language models are increasingly being deployed for conversational support in informal caregiving contexts, where interactions often extend beyond information-seeking: caregivers seek emotional reassurance, guidance, and help,...

<https://arxiv.org/abs/2605.29473>

33 CONF

MOOSE-Copilot: A Web-Based Interactive Assistant for Unified Exploratory and Fine-Grained Scientific Hypothesis Discovery

ArXiv CS.AI 50%

arXiv:2605.29475v1 Announce Type: cross Abstract: Large language models (LLMs) show remarkable potential in scientific hypothesis discovery. However, existing approaches face two critical limitations: they treat divergent exploratory ideation and convergent fine-grained...

<https://arxiv.org/abs/2605.29475>

34 CONF

The New Pro Se: Generative AI and the Surge in Federal Civil Self-Representation

ArXiv CS.AI 50%

arXiv:2605.29493v1 Announce Type: cross Abstract: Since public access to generative AI tools became widespread, federal civil litigation has seen a marked increase in pro se (self-represented) plaintiffs. This paper analyzes that shift using ~2.8 million filings, asking whether...

<https://arxiv.org/abs/2605.29493>

35 CONF

KBF: Knowledge Boundary as Fingerprint for Language Model and Black-Box API Auditing

ArXiv CS.AI 50%

arXiv:2605.29524v1 Announce Type: cross Abstract: Relay and reseller APIs increasingly intermediate access to large language models (LLMs), but users have no direct way to verify that a claimed endpoint is actually serving the advertised model. We introduce KBF, a low-cost...

<https://arxiv.org/abs/2605.29524>

36 CONF

Temporal Motif-aware Graph Test-time Adaptation for OOD Blockchain Anomaly Detection

ArXiv CS.AI 50%

arXiv:2605.29526v1 Announce Type: cross Abstract: Ever-evolving transaction patterns have significantly hindered anomaly detection on emerging cryptocurrency blockchains due to the vast number of addresses and diverse anomalous behaviors. Recently, advanced Graph Anomaly...

<https://arxiv.org/abs/2605.29526>

37 CONF

GiPL: Generative augmented iterative Pseudo-Labeling for Cross-Domain Few-Shot Object Detection

ArXiv CS.AI 50%

arXiv:2605.29539v1 Announce Type: cross Abstract: Vision-language foundation models have shown promising zero-shot generalization for Cross-Domain Few-Shot Object Detection (CD-FSOD). However, they face two critical challenges in fine-tuning: insufficient support set utilization...

<https://arxiv.org/abs/2605.29539>

38 CONF

Brain-IT-VQA: From Brain Signals to Answers

ArXiv CS.AI 50%

arXiv:2605.29588v1 Announce Type: cross Abstract: Decoding visual content from fMRI signals recorded while a person views images, and specifically answering questions about the seen images, is a long-standing challenge. While significant progress has been made in recent years in...

<https://arxiv.org/abs/2605.29588>

39 CONF

COMET: Concept Space Dissection of the Modality Gap in Audio-Text Multimodal Contrastive Embeddings

ArXiv CS.AI 50%

arXiv:2605.29628v1 Announce Type: cross Abstract: Contrastive Language-Audio Pretraining (CLAP) models are widely used for audio understanding and support modality-agnostic condition swapping in many zero-shot applications. However, their performance is heavily affected by the...

<https://arxiv.org/abs/2605.29628>

40 CONF

OccamToken: Efficient VLM Inference with Training-Free and Budget-Adaptive Token Pruning

ArXiv CS.AI 50%

arXiv:2605.29657v1 Announce Type: cross Abstract: Vision-language models (VLMs) rely on long visual token sequences for visual understanding, making the prefill stage expensive in both computation and memory. Most existing pruning methods follow an absolute-ranking paradigm,...

<https://arxiv.org/abs/2605.29657>

41 CONF

Opir: Efficient Multi-Task Safety Classification for Toxicity, Jailbreaks, Hate Speech, and Harmful Content

ArXiv CS.AI 50%

arXiv:2605.29659v1 Announce Type: cross Abstract: Real-time safety filtering for large language model (LLM) applications requires classifiers that can detect unsafe prompts, toxic language, jailbreak attempts, and unsafe responses without the cost profile of large guardrail...

<https://arxiv.org/abs/2605.29659>

42 CONF

EviLink: Multi-Path Schema Linking with Uncertainty-Guided Evidence Acquisition for Large-Scale Text-to-SQL

ArXiv CS.AI 50%

arXiv:2605.29670v1 Announce Type: cross Abstract: Schema linking is a difficult and important step in large-scale Text-to-SQL, where systems must identify a compact yet sufficient schema context from large and ambiguous databases. Existing methods often treat schema linking as...

<https://arxiv.org/abs/2605.29670>

43 CONF

From Prompts to Context: An Ontology-Driven Framework for Human-Generative AI Collaboration

ArXiv CS.AI 50%

arXiv:2605.29675v1 Announce Type: cross Abstract: Collaborations with Generative AI often begin with a short prompt and end with an opaque output, leaving implicit who was involved, what task was being pursued, which resources were used, and which constraints should have shaped...

<https://arxiv.org/abs/2605.29675>

44 CONF

Personalized Turn-Level User Conversation Satisfaction Benchmark

ArXiv CS.AI 50%

arXiv:2605.29711v1 Announce Type: cross Abstract: User satisfaction with AI assistants is highly personalized: the same response may satisfy one user but disappoint another depending on what each user expects and what they have asked for before. Existing automatic evaluation...

<https://arxiv.org/abs/2605.29711>

45 CONF

The Little Book of Generative AI Foundations: An Intuitive Mathematical Primer

ArXiv CS.AI 50%

arXiv:2605.29713v1 Announce Type: cross Abstract: This book provides a compact, derivation-oriented introduction to the mathematical foundations of modern generative artificial intelligence. Rather than surveying every recent architecture or implementation detail, it develops a...

<https://arxiv.org/abs/2605.29713>

46 CONF

Evolve as a Team: Collaborative Self-Evolution for LLM-based Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2605.29790v1 Announce Type: cross Abstract: LLM-based multi-agent systems (MAS) have emerged as an effective paradigm for complex and long-horizon tasks. However, in real-world tasks, MAS often exhibit various failures during execution and such failures are difficult to...

<https://arxiv.org/abs/2605.29790>

47 CONF

Towards Verifiable Multimodal Deep Research: A Multi-Agent Harness for Interleaved Report Generation

ArXiv CS.AI 50%

arXiv:2605.29861v1 Announce Type: cross Abstract: Large Language Models (LLMs) have advanced autonomous agents from deep search, which retrieves concise factual answers, to deep research, which synthesizes scattered evidence into long-form reports. However, verifiable multimodal...

<https://arxiv.org/abs/2605.29861>

48 CONF

Mitigating Stethoscope-Induced Shortcuts in Respiratory Sound Classification under Federated Domain Generalization with Causality-Inspired Interventions

ArXiv CS.AI 50%

arXiv:2605.29862v1 Announce Type: cross Abstract: AI-driven respiratory sound classification (RSC) is promising for automated pulmonary disease detection, yet multi-site deployment is hindered by inter-stethoscope variability. We introduce a federated domain generalization...

<https://arxiv.org/abs/2605.29862>

49 CONF

Evolutionary Dynamics of Cooperation in Next-Generation LLM Agent Systems: A Cross-Provider Empirical Extension

ArXiv CS.AI 50%

arXiv:2605.29874v1 Announce Type: cross Abstract: Do next-generation LLM agents inherit the cooperative biases documented in their predecessors, or does scale and provider diversity reshape equilibrium behaviour in competitive multi-agent settings? Willis et al. established a...

<https://arxiv.org/abs/2605.29874>

50 CONF

Mitigating Hallucination in Vision-Language Models through Barrier-Regulated Adaptive Closed-form Steering

ArXiv CS.AI 50%

arXiv:2605.29881v1 Announce Type: cross Abstract: Large vision-language models (LVLMs) often hallucinate objects that are not present in the input image, largely because visual grounding weakens as decoding progresses. Existing inference-time mitigation methods modify logits or...

<https://arxiv.org/abs/2605.29881>

51 CONF

AG-REPA: Causal Layer Selection for Representation Alignment in Audio Flow Matching

ArXiv CS.AI 50%

arXiv:2603.01006v2 Announce Type: replace-cross Abstract: REPresentation Alignment (REPA) improves the training of generative flow models by aligning intermediate hidden states with pretrained teacher features, but its effectiveness in token-conditioned audio Flow Matching...

<https://arxiv.org/abs/2603.01006>

52 CONF

Internal Representation, Not Clinical Knowledge: Where Apparent LLM Triage Failures Originate

ArXiv CS.AI 50%

arXiv:2605.29889v1 Announce Type: cross Abstract: Patient-voiced clinical-triage benchmarks report high under-triage rates for consumer LLMs for constrained multiple-choice output, yet the same cases score differently with free-text. We ask whether output format changes the...

<https://arxiv.org/abs/2605.29889>

53 CONF

Agora: Toward Autonomous Bug Detection in Production-Level Consensus Protocols with LLM Agents

ArXiv CS.AI 50%

arXiv:2605.29910v1 Announce Type: cross Abstract: Consensus protocols form the backbone of distributed systems and blockchains, where implementation bugs can cause data corruption and financial losses. While LLM-based approaches show promise in code analysis, they struggle with...

<https://arxiv.org/abs/2605.29910>

54 CONF

Label Over Logic? How Source Cues Bias Human Fallacy Judgments More Than LLMs

ArXiv CS.AI 50%

arXiv:2605.29928v1 Announce Type: cross Abstract: As AI-generated and AI-assisted content floods online spaces, source labels attached to such content can distort human reasoning judgments, with downstream consequences for moderation, evaluation, and decision-making. Whether...

<https://arxiv.org/abs/2605.29928>

55 CONF

REPORT: Recoverable Program-of-Thought via Checkpoint Repair

ArXiv CS.AI 50%

arXiv:2605.30052v1 Announce Type: cross Abstract: One-shot Program-of-Thought (PoT) emits a Python program that prints a primitive-action plan; a single invalid action silently invalidates the trajectory. We introduce RePoT (Recoverable PoT): a deterministic verified replay that...

<https://arxiv.org/abs/2605.30052>

56 CONF

Projectional Decoding: Towards Semantic-Aware LLM Generation

ArXiv CS.AI 50%

arXiv:2605.30054v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to generate software artifacts across many software engineering (SE) tasks, yet ensuring the semantic validity of these artifacts remains a fundamental challenge. Existing...

<https://arxiv.org/abs/2605.30054>

57 CONF

How Reliable Are AI Attackers Against a Fixed Vulnerable Target? A 400-Run Empirical Study of LLM Penetration Testing Consistency

ArXiv CS.AI 50%

arXiv:2605.30096v1 Announce Type: cross Abstract: Large language models (LLMs) can autonomously conduct multi-stage cyber attacks, but the consistency of their offensive behavior under repeated trials remains unstudied. This work presents the first large-scale empirical...

<https://arxiv.org/abs/2605.30096>

58 CONF

No More K-means: Single-Stage Sparse Coding for Efficient Multi-Vector Retrieval

ArXiv CS.AI 50%

arXiv:2605.30120v1 Announce Type: cross Abstract: Multi-vector retrieval (MVR) models, exemplified by ColBERT, have established new benchmarks in retrieval accuracy by preserving fine-grained token-level interactions. However, this granularity imposes prohibitive storage and...

<https://arxiv.org/abs/2605.30120>

59 CONF

Overcoming Forgetting in LLM Fine-Tuning with Evolution Strategies

ArXiv CS.AI 50%

arXiv:2605.30148v1 Announce Type: cross Abstract: Evolution Strategies (ES) has recently emerged as a competitive alternative to reinforcement learning (RL) for large language model (LLM) fine-tuning, offering advantages through simplicity, scalability, and inference-only...

<https://arxiv.org/abs/2605.30148>

60 CONF

Do Proactive Agents Really Need an LLM to Decide When to Wake and What to Anchor?

ArXiv CS.AI 50%

arXiv:2605.30152v1 Announce Type: cross Abstract: Proactive agents read user activity as text and call an LLM on every event to decide whether to act. But user activity is not natively text: it is a structured event stream of (actor, verb, object, timestamp) tuples that the...

<https://arxiv.org/abs/2605.30152>

61 CONF

CalArena: A Large-Scale Post-Hoc Calibration Benchmark

ArXiv CS.AI 50%

arXiv:2605.30188v1 Announce Type: cross Abstract: Reliable probability estimates are critical in many machine learning applications, yet modern classifiers are often poorly calibrated. Post-hoc calibration provides a simple and widely used solution, but the large number of...

<https://arxiv.org/abs/2605.30188>

62 CONF

What drives performance in molecular MPNNs? An operator-level factorial benchmark

ArXiv CS.AI 50%

arXiv:2605.30195v1 Announce Type: cross Abstract: Message-passing neural networks (MPNNs) are widely used for molecular property prediction, but their deployment as monolithic architectures makes it difficult to identify how specific message-passing operators affect performance....

<https://arxiv.org/abs/2605.30195>

63 CONF

HPO: Hysteretic Policy Optimization for Stable and Efficient Training under Sparse-Reward Regime

ArXiv CS.AI 50%

arXiv:2605.30201v1 Announce Type: cross Abstract: We investigate a narrow but common failure mode of GRPO-style reinforcement learning in the context of sparse verifiable rewards: early updates contain more responses with negative advantages than those with positive advantages,...

<https://arxiv.org/abs/2605.30201>

64 CONF

LLUMI: Improving LLM Writing Assistance for Mental Health Support with Online Community Feedback

ArXiv CS.AI 50%

arXiv:2605.30273v1 Announce Type: cross Abstract: Large language models (LLMs) show promise in generating supportive responses for mental health queries, but improving their usefulness, empathy, and safety often requires substantial compute, expert input, and labeled data. At...

<https://arxiv.org/abs/2605.30273>

65 CONF

Revisiting the Effectiveness of LLM Pruning for Test-Time Scaling

ArXiv CS.AI 50%

arXiv:2604.25098v2 Announce Type: replace Abstract: Large Language Models (LLMs) now exhibit remarkable reasoning capabilities through test-time compute scaling (TTS), with impressive performance across math and coding benchmarks. In parallel, research in model compression has...

<https://arxiv.org/abs/2604.25098>

66 CONF

MedCase-Structured: A Text-to-FHIR Dataset for Benchmarking Diagnostic Reasoning in Clinically Realistic EHR Settings

ArXiv CS.AI 50%

arXiv:2605.30295v1 Announce Type: cross Abstract: Large language models (LLMs) show promise for clinical reasoning and decision support, but evaluation in realistic, electronic health record-congruent settings remains limited. Existing benchmarks often rely on static datasets or...

<https://arxiv.org/abs/2605.30295>

67 CONF

Before the Shutter: Aesthetic and Actionable Portrait Photography Planning in 3D Scenes

ArXiv CS.AI 50%

arXiv:2605.30318v1 Announce Type: cross Abstract: Portrait photography is largely decided before the shutter opens: the subject's pose, the camera configuration, and the lighting devices must be coordinated within the surrounding 3D scene. In contrast, most existing...

<https://arxiv.org/abs/2605.30318>

68 CONF

Gram: Assessing sabotage propensities via automated alignment auditing

ArXiv CS.AI 50%

arXiv:2605.30322v1 Announce Type: cross Abstract: We introduce Gram, an automated alignment auditing framework to assess the propensity of AI agents to engage in sabotage. We evaluate Gemini models across 17 simulated agentic deployment scenarios that incentivize sabotage. We...

<https://arxiv.org/abs/2605.30322>

69 CONF

VideoMLA: Low-Rank Latent KV Cache for Minute-Scale Autoregressive Video Diffusion

ArXiv CS.AI 50%

arXiv:2605.30351v1 Announce Type: cross Abstract: Long-rollout causal video diffusion has converged on a fixed-size sliding-window KV cache, with recent progress innovating within this layout by changing which tokens occupy the window or how their positions are encoded. The...

<https://arxiv.org/abs/2605.30351>

70 CONF

SafeSearch: Automated Red-Teaming of LLM-Based Search Agents

ArXiv CS.AI 50%

arXiv:2509.23694v5 Announce Type: replace Abstract: Search agents connect LLMs to the Internet, enabling them to access broader and more up-to-date information. However, this also introduces a new threat surface: unreliable search results can mislead agents into producing unsafe...

<https://arxiv.org/abs/2509.23694>

71 CONF

Jailbreak Scaling Laws for Large Language Models: Polynomial-Exponential Crossover

ArXiv CS.AI 50%

arXiv:2603.11331v3 Announce Type: replace-cross Abstract: Adversarial attacks can reliably steer safety-aligned large language models toward unsafe behavior. Empirically, we find that adversarial prompt-injection attacks can amplify attack success rate from the slow polynomial...

<https://arxiv.org/abs/2603.11331>

72 CONF

RewardFlow: Topology-Aware Reward Propagation on State Graphs for Agentic RL with Large Language Models

ArXiv CS.AI 50%

arXiv:2603.18859v2 Announce Type: replace Abstract: Reinforcement learning (RL) shows promise for enhancing LLM agentic reasoning, yet sparse terminal rewards hinder fine-grained optimization. Process reward modeling offers an alternative but incurs high computational costs...

<https://arxiv.org/abs/2603.18859>

73 CONF

MATNet: Multi-Level Fusion Transformer-Based Model for Day-Ahead PV Generation Forecasting

ArXiv CS.AI 50%

arXiv:2306.10356v3 Announce Type: replace-cross Abstract: Accurate forecasting of renewable generation is crucial to facilitate the integration of Renewable Energy Sources into the power system. Focusing on photovoltaic (PV) units, forecasting methods can be divided into two...

<https://arxiv.org/abs/2306.10356>

74 CONF

Is Your LLM Overcharging You? Tokenization, Transparency, and Incentives

ArXiv CS.AI 50%

arXiv:2505.21627v4 Announce Type: replace-cross Abstract: State-of-the-art large language models require specialized hardware and substantial energy to operate. As a consequence, cloud-based services that provide access to large language models have become very popular. In these...

<https://arxiv.org/abs/2505.21627>

75 CONF

Approximate Proportionality in Online Fair Division

ArXiv CS.AI 50%

arXiv:2508.03253v2 Announce Type: replace-cross Abstract: We study the online fair division problem, where indivisible goods arrive sequentially and must be allocated immediately and irrevocably. Prior work establishes strong impossibility results for approximating classic...

<https://arxiv.org/abs/2508.03253>

76 CONF

Scalable RF Simulation in Generative 4D Worlds

ArXiv CS.AI 50%

arXiv:2508.12176v2 Announce Type: replace-cross Abstract: Radio Frequency (RF) sensing has emerged as a powerful, privacy-preserving alternative to vision-based methods for various perception tasks. However, building high-quality RF datasets in dynamic and diverse environments...

<https://arxiv.org/abs/2508.12176>

77 CONF

Towards Foundation Models for Zero-Shot Time Series Anomaly Detection: Leveraging Synthetic Data and Relative Context Discrepancy

ArXiv CS.AI 50%

arXiv:2509.21190v4 Announce Type: replace-cross Abstract: Time series anomaly detection (TSAD) is a critical task, but developing models that generalize to unseen data in a zero-shot manner remains a major challenge. Prevailing foundation models for TSAD predominantly rely on...

<https://arxiv.org/abs/2509.21190>

78 CONF

Uncovering Vulnerabilities of LLM-Assisted Cyber Threat Intelligence

ArXiv CS.AI 50%

arXiv:2509.23573v5 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly used to help security analysts manage the surge of cyber threats, automating tasks from vulnerability assessment to incident response. Yet in operational CTI workflows,...

<https://arxiv.org/abs/2509.23573>

79 CONF

Beyond Accuracy: Are Time Series Foundation Models Well-Calibrated?

ArXiv CS.AI 50%

arXiv:2510.16060v2 Announce Type: replace-cross Abstract: The recent development of foundation models for time series data has generated considerable interest in using such models across a variety of applications. Although foundation models achieve state-of-the-art predictive...

<https://arxiv.org/abs/2510.16060>

80 CONF

HD-Prot: A Protein Language Model for Joint Sequence-Structure Modeling with Continuous Structure Tokens

ArXiv CS.AI 50%

arXiv:2512.15133v2 Announce Type: replace-cross Abstract: Proteins inherently possess a consistent sequence-structure duality. The abundance of protein sequence data, which can be readily represented as discrete tokens, has driven fruitful developments in protein language models...

<https://arxiv.org/abs/2512.15133>

81 CONF

Grammar-Aware Literate Generative Mathematical Programming with Compiler-in-the-Loop

ArXiv CS.AI 50%

arXiv:2601.17670v2 Announce Type: replace-cross Abstract: Mathematical programming is widely employed across various sectors - such as logistics, energy, and workforce planning - to model and solve industrial optimisation problems, but its use requires substantial domain...

<https://arxiv.org/abs/2601.17670>

82 CONF

Pushing the Limits of Block Rotations in Post-Training Quantization

ArXiv CS.AI 50%

arXiv:2601.22347v2 Announce Type: replace-cross Abstract: Recent post-training quantization (PTQ) methods have adopted block rotations to diffuse outliers prior to rounding. While this reduces the overhead of online full-vector rotations, the effect of block structure on outlier...

<https://arxiv.org/abs/2601.22347>

83 CONF

Who can we trust? LLM-as-a-jury for Comparative Assessment

ArXiv CS.AI 50%

arXiv:2602.16610v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly applied as automatic evaluators for natural language generation assessment often using pairwise comparative judgements. Existing approaches typically rely on single judges or...

<https://arxiv.org/abs/2602.16610>

84 CONF

Rooted Absorbed Prefix Trajectory Balance with Submodular Replay for GFlowNet Training

ArXiv CS.AI 50%

arXiv:2603.00454v2 Announce Type: replace-cross Abstract: Generative Flow Networks (GFlowNets) enable fine-tuning large language models to approximate reward-proportional posteriors, but they remain prone to mode collapse, manifesting as prefix collapse and length bias. We...

<https://arxiv.org/abs/2603.00454>

85 CONF

Enhancing LLM Medical Coding with Structured External Knowledge

ArXiv CS.AI 50%

arXiv:2605.27377v2 Announce Type: replace-cross Abstract: Accurate medical coding requires consulting authoritative resources such as the ICD tabular list and coding guidelines. Existing LLM-based automated methods largely rely on LLMs' internal knowledge, which is prone to...

<https://arxiv.org/abs/2605.27377>

86 CONF

P\$^2\$RAG: Efficient Privacy-Preserving RAG Service Supporting Arbitrary Top-\$k\$ Retrieval

ArXiv CS.AI 50%

arXiv:2603.14778v2 Announce Type: replace-cross Abstract: Retrieval-Augmented Generation (RAG) enables large language models to use external knowledge, but outsourcing the RAG service raises privacy concerns for both data owners and users. Privacy-preserving RAG systems address...

<https://arxiv.org/abs/2603.14778>

87 CONF

When Should a Robot Think? Resource-Aware Reasoning via Reinforcement Learning for Embodied Robotic Decision-Making

ArXiv CS.AI 50%

arXiv:2603.16673v4 Announce Type: replace-cross Abstract: Embodied robotic systems increasingly rely on large language model (LLM)-based agents to support high-level reasoning, planning, and decision-making during interactions with the environment. However, invoking LLM...

<https://arxiv.org/abs/2603.16673>

88 CONF

Maximizing Mutual Information Between Prompt and Response Improves LLM Performance With No Additional Data

ArXiv CS.AI 50%

arXiv:2603.19294v3 Announce Type: replace-cross Abstract: While post-training has successfully improved large language models (LLMs) across a variety of domains, these gains heavily rely on human-labeled data or external verifiers. Existing data has already been exploited, and...

<https://arxiv.org/abs/2603.19294>

89 CONF

SelfGrader: LLM Jailbreak Detection via Anchored Token-Level Logits

ArXiv CS.AI 50%

arXiv:2604.01473v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are powerful tools for answering user queries, yet they remain highly vulnerable to jailbreak attacks. Existing guardrail methods typically rely on internal features or textual responses to...

<https://arxiv.org/abs/2604.01473>

90 CONF

JMed48k: A Multi-Profession Japanese Medical Licensing Benchmark for Vision-Language Model Evaluation

ArXiv CS.AI 50%

arXiv:2605.22080v2 Announce Type: replace-cross Abstract: We introduce JMed48k, a multi-profession Japanese healthcare licensing benchmark for evaluating vision-language models. Built from official PDF materials released by the Japanese Ministry of Health, Labour and Welfare,...

<https://arxiv.org/abs/2605.22080>

91 CONF

Eureka: Intelligent Feature Engineering for Enterprise AI Cloud Resource Demand Prediction

ArXiv CS.AI 50%

arXiv:2605.25297v2 Announce Type: replace-cross Abstract: Effective features are crucial for predictive model performance, but creating them often requires domain expertise, limiting scalability across applications. We define feature engineering as an agentic code generation...

<https://arxiv.org/abs/2605.25297>

92 CONF

The Alignment Floor: How Persona Customization Breaks Safety in Weakly-Aligned LLMs

ArXiv CS.AI 50%

arXiv:2605.27382v2 Announce Type: replace-cross Abstract: Telling an LLM to "be enthusiastic" raises its sycophancy rate from 30\% to 50\% on a lightly-aligned model, but has zero effect on a strongly-aligned one. We define this gap as the alignment floor,...

<https://arxiv.org/abs/2605.27382>