



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	137	Media
HackerNews	2	Community
TechCrunch AI	2	Media
The Verge AI	1	Media

CONFIRMED SIGNALS — 142 stories

CONF

Show HN: Lite-Harness - Self-Hosted Cursor Agents (Use Claude Code/OpenCode)

HackerNews 50% HN 6

We built this Dockerfile because we wanted a simple harness server to run our agents and get memory, durable sessions, cron scheduling, and a vault, out of the box.

<https://github.com/LiteLLM-Labs/lite-harness>

CONF

Show HN: FuguUX - science-backed, AI user testing to identify key web UX issues

HackerNews 50% HN 5

No summary — see link for full article.

<https://www.fuguux.com>

CONF

Erin Brockovich takes aim at data center secrecy

TechCrunch AI 50%

Environmental activist Erin Brockovich has a new mission.

<https://techcrunch.com/2026/05/31/erin-brockovich-takes-aim-at-data-center-secrecy/>

CONF

Making sense of the debate over AI psychosis

TechCrunch AI 50%

On the latest episode of Equity, we debate whether tech CEOs are "uniquely prone to AI psychosis."

<https://techcrunch.com/2026/05/31/making-sense-of-the-debate-over-ai-psychosis/>

CONF

I went looking for the AI weed vape that gives you Bitcoin for smoking

The Verge AI 50%

The crypto weed vape found me on 4/20, the high holiday of cannabis enthusiasts everywhere. It arrived over Slack with the thumbnail of a man exhaling a plume of vapor, the words "every hit delivers Bitcoin" emblazoned across it. It claimed to be advertising a device called...

<https://www.theverge.com/ai-artificial-intelligence/933916/ai-powered-crypto-cannabis-vape>

CONF

Physically Viable World Models: A Case for Query-Conditioned Embodied AI

ArXiv CS.AI 50%

arXiv:2605.30542v1 Announce Type: new Abstract: World models for embodied AI must be physically viable: constructed to answer intervention queries by representing the physical structure governing action outcomes, rather than merely predicting future observations. Existing...

<https://arxiv.org/abs/2605.30542>

CONF

The Terminal Representation in Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.31289v1 Announce Type: cross Abstract: Representation learning is a powerful tool for spatio-temporal abstraction within reinforcement learning (RL). Two well established approaches are through the successor representation (SR) and the default representation (DR). The...

<https://arxiv.org/abs/2605.31289>

CONF

Counterfactual Evaluation Reveals Hidden Capability Profiles in Clinical LLMs and Agents

ArXiv CS.AI 50%

arXiv:2605.30590v1 Announce Type: cross Abstract: Two clinical AI systems can score nearly identically on coverage-based rubrics yet behave radically differently when their patient inputs change: one updates its recommendations to match the new clinical signal, while the other...

<https://arxiv.org/abs/2605.30590>

CONF

Inferring Events from Time Series using Language Models

ArXiv CS.AI 50%

arXiv:2503.14190v3 Announce Type: replace Abstract: A common goal in analyzing time series data is to understand how events cause observed variations. We study whether Large Language Models (LLMs) can infer natural language events associated with time series data. We introduce...

<https://arxiv.org/abs/2503.14190>

0 CONF

PRemISE: Policy Rubrics as Measurement Specifications for LLM Judges

ArXiv CS.AI 50%

arXiv:2605.30803v1 Announce Type: new Abstract: LLM judges are increasingly used to evaluate open-ended responses, but their scores depend strongly on the rubrics that condition them. A vague rubric asking for a response to be ``helpful and factual" can reward polished answers...

<https://arxiv.org/abs/2605.30803>

1 CONF

Inverse Reinforcement Learning without an Optimal Demonstrator: A Feasible Reward Set Approach

ArXiv CS.AI 50%

arXiv:2605.30903v1 Announce Type: cross Abstract: Inverse reinforcement learning (IRL) typically assumes demonstrations from a single optimal demonstrator, but in many applications data come from multiple imperfect demonstrators with heterogeneous suboptimality levels. We study...

<https://arxiv.org/abs/2605.30903>

2 CONF

COMPASS: Cognitive MCTS-Guided Process Alignment for Safe Search Agents

ArXiv CS.AI 50%

arXiv:2605.30838v1 Announce Type: new Abstract: LLM-powered search agents enable multi-step reasoning and tool use. However, these capabilities introduce retrieval-induced safety degradation, as harmful intents may decompose into seemingly innocuous sub-queries that lead to...

<https://arxiv.org/abs/2605.30838>

3 CONF

Distilling LLM Feedback for Lean Theorem Proving

ArXiv CS.AI 50%

arXiv:2605.30861v1 Announce Type: new Abstract: Post-training for reasoning models typically combines supervised fine-tuning with reinforcement learning from verifiable rewards, most commonly with GRPO. However, this algorithm suffers from sparse rewards, limited exploration,...

<https://arxiv.org/abs/2605.30861>

- 4

CONF

UniScale: Adaptive Unified Inference Scaling via Online Joint Optimization of Model Routing and Test-Time Scaling

ArXiv CS.AI 50%

arXiv:2605.30898v1 Announce Type: new Abstract: In real-world deployments of large language models (LLMs), balancing inference quality and computational cost has become a central challenge. Existing approaches tackle this trade-off along two largely independent dimensions: model...
<https://arxiv.org/abs/2605.30898>
- 5

CONF

BilliardPhys-Bench: Benchmarking Physical Reasoning and Visual Dynamics of Multimodal LLMs

ArXiv CS.AI 50%

arXiv:2605.30900v1 Announce Type: new Abstract: Current multimodal models handle static image recognition well, but intuitive physical reasoning remains a weakness. Predicting how objects will move and interact from a single image is still difficult for these systems. We present...
<https://arxiv.org/abs/2605.30900>
- 6

CONF

Personalized to Persuade: The Effects of Contextualization and Warmth on Trust and Reliance in Conversational AI

ArXiv CS.AI 50%

arXiv:2605.31275v1 Announce Type: cross Abstract: Artificial Intelligence (AI) agents personalize their responses by tailoring explanations to users' backgrounds, interests, and prior interactions, referred to as contextualization. Personalization has been identified as a...
<https://arxiv.org/abs/2605.31275>
- 7

CONF

TunerDiT: Training-free Progressive Steering of Diffusion Transformer for Multi-Event Video Generation

ArXiv CS.AI 50%

arXiv:2605.31590v1 Announce Type: cross Abstract: Text-to-video (T2V) generation faces challenging questions when generating videos with long horizons containing multiple events. Inspired by the intrinsics of the diffusion process, we probe video diffusion transformers (DiTs)...
<https://arxiv.org/abs/2605.31590>
- 8

CONF

GraphARC: A Comprehensive Benchmark for Graph-Based Abstract Reasoning

ArXiv CS.AI 50%

arXiv:2605.31031v1 Announce Type: new Abstract: Relational reasoning lies at the heart of intelligence, but existing benchmarks are typically confined to formats such as grids or text. We introduce GraphARC, a benchmark for abstract reasoning on graph-structured data. GraphARC...
<https://arxiv.org/abs/2605.31031>
- 9

CONF

Vector Linking via Cross-Model Local Isometric Consistency

ArXiv CS.AI 50%

arXiv:2605.31100v1 Announce Type: new Abstract: We study Vector Linking: given two embedding clouds produced by different black-box encoders over partially overlapping datasets, recover cross-model object correspondences using only vectors. Empirically and theoretically, we show...
<https://arxiv.org/abs/2605.31100>

- 0 **CONF**
LLM-FACETS: A Privacy-Preserving Framework for Evaluating LLM Transparency and Accountability
 ArXiv CS.AI 50%
 arXiv:2605.31167v1 Announce Type: new Abstract: Assessing whether Large Language Models outputs are factually grounded, epistemically calibrated, and methodologically reproducible is a prerequisite for responsible AI deployment. Yet auditing LLMs remains inaccessible to...
<https://arxiv.org/abs/2605.31167>
-
- 1 **CONF**
COLLEAGUE.SKILL: Automated AI Skill Generation via Expert Knowledge Distillation
 ArXiv CS.AI 50%
 arXiv:2605.31264v1 Announce Type: new Abstract: LLM agents are increasingly expected not only to complete isolated tasks, but also to carry bounded representations of human expertise, judgment, and interaction style. Building such person-grounded agents remains difficult because...
<https://arxiv.org/abs/2605.31264>
-
- 2 **CONF**
Industrializing Prediction-Powered Inference: The GLIDE Library for Reliable GenAI and Agentic Systems Evaluation
 ArXiv CS.AI 50%
 arXiv:2605.31278v1 Announce Type: new Abstract: Reliable evaluation of agentic systems requires unbiased estimates with valid uncertainty, but standard practice navigates between costly human annotation and biased LLM-as-judge proxies. Prediction-powered inference (PPI) combines...
<https://arxiv.org/abs/2605.31278>
-
- 3 **CONF**
Diagnosing Failure Modes of Shared-State Collaboration in Resource-Constrained Visual Agents
 ArXiv CS.AI 50%
 arXiv:2605.31354v1 Announce Type: new Abstract: Modular visual reasoning systems increasingly rely on shared working memory for multi-step collaboration, yet the failure dynamics of intermediate state evolution in low-capacity regimes remain underexplored. We study failure modes...
<https://arxiv.org/abs/2605.31354>
-
- 4 **CONF**
Reliable Self-Improvement Training by Verifying Reasoning, Not Just Answers
 ArXiv CS.AI 50%
 arXiv:2603.21558v2 Announce Type: replace Abstract: Self-improvement training, where models learn from self-generated solutions, promises sustained capability gains but suffers from a pervasive failure mode: across multiple rounds, compounding reasoning errors cause accuracy to...
<https://arxiv.org/abs/2603.21558>
-
- 5 **CONF**
When LLM Reward Design Fails: Diagnostic-Driven Refinement for Sparse Structured RL
 ArXiv CS.AI 50%
 arXiv:2605.28918v1 Announce Type: cross Abstract: For sparse, structured reinforcement-learning tasks with semantic reward-function interfaces, LLM-generated reward shaping is better framed as debugging than one-shot generation. We study PPO-trained agents using MiniGrid as core...
<https://arxiv.org/abs/2605.28918>

6 CONF

Gradient-Free Training of Spiking Neural Networks via Low-Rank Evolution Strategies

ArXiv CS.AI 50%

arXiv:2605.30361v1 Announce Type: cross Abstract: Spiking Neural Networks (SNNs) offer compelling energy efficiency on neuromorphic hardware, yet their training remains challenging because the discrete spike threshold is non-differentiable. Surrogate-gradient methods sidestep...

<https://arxiv.org/abs/2605.30361>

7 CONF

XOResNet: Exclusive-OR Meta-Residuals Facilitate Deep Spiking Neural Networks Learning

ArXiv CS.AI 50%

arXiv:2605.30362v1 Announce Type: cross Abstract: Spiking neural networks (SNNs) hold promise for demonstrating superior learning and representation capabilities in deep models. Given the tremendous success of ResNet in deep learning, it would naturally follow to train deep SNNs...

<https://arxiv.org/abs/2605.30362>

8 CONF

Updating the standard neuron model in artificial neural networks

ArXiv CS.AI 50%

arXiv:2605.30370v1 Announce Type: cross Abstract: From their inception in the 1950s, artificial neural networks (ANNs) started using the so-called point neuron model then prevalent in neuroscience, hoping that this analogy would allow for a better emulation of brain function....

<https://arxiv.org/abs/2605.30370>

9 CONF

Full-field prediction for engineering-scale three-dimensional aircraft with multigrid-hierarchical learning

ArXiv CS.AI 50%

arXiv:2605.30375v1 Announce Type: cross Abstract: High-fidelity computational fluid dynamics is essential for aerospace design, but engineering-scale simulations of practical three-dimensional aircraft remain computationally expensive. Learning-based flow-field initialization...

<https://arxiv.org/abs/2605.30375>

0 CONF

Benchmarking and Enhancing Text-to-Image Models for Generating Visual Representations in Early Arithmetic Education

ArXiv CS.AI 50%

arXiv:2605.31212v1 Announce Type: cross Abstract: AI systems are increasingly used to support educational content creation, yet it remains unclear whether they can generate outputs that faithfully represent the pedagogical concepts they are intended to teach. Thus, we introduce...

<https://arxiv.org/abs/2605.31212>

1 CONF

LLMs Without Deep Neural Networks: New Architecture, Benefits and Case Study

ArXiv CS.AI 50%

arXiv:2605.30385v1 Announce Type: cross Abstract: The purpose of this article is to provide validation to my deep neural network alternative in the context of LLMs. Very recently, there has been a significant interest by Chinese researchers in a model called RBF network, as a...

<https://arxiv.org/abs/2605.30385>

2 CONF

Functional MRI Time Series Generation via Wavelet-Based Image Transform and Spectral Flow Matching for Brain Disorder Identification

ArXiv CS.AI 50%

arXiv:2605.30387v1 Announce Type: cross Abstract: Functional Magnetic Resonance Imaging (fMRI) provides non-invasive access to dynamic brain activity by measuring blood oxygen level-dependent (BOLD) signals over time. However, the resource-intensive nature of fMRI acquisition...

<https://arxiv.org/abs/2605.30387>

3 CONF

CodeGolf Bench: A Multi-Language Benchmark for Evaluating Concise Code Generation Capabilities of Large Language Models

ArXiv CS.AI 50%

arXiv:2605.30394v1 Announce Type: cross Abstract: This paper introduces Code Bench, a benchmark capable of evaluating Large Language Models (LLMs) concise code generation abilities in 60 programming languages. Based on code golf, a recreational programming competition focused on...

<https://arxiv.org/abs/2605.30394>

4 CONF

Automatically Attacking Software Reverse Engineering AI Agents

ArXiv CS.AI 50%

arXiv:2605.30667v1 Announce Type: cross Abstract: Software tools for reverse engineering executable binary files, such as Ghidra, enable malware analysts to safely conduct robust static analysis without having access to original source code. Coupled with the analytic power of...

<https://arxiv.org/abs/2605.30667>

5 CONF

SANA-Streaming: Real-time Streaming Video Editing with Hybrid Diffusion Transformer

ArXiv CS.AI 50%

arXiv:2605.30409v1 Announce Type: cross Abstract: Real-time streaming video-to-video editing (V2V) is critical for interactive applications such as live broadcasting and gaming, yet it remains a formidable challenge due to the stringent requirements for temporal consistency and...

<https://arxiv.org/abs/2605.30409>

6 CONF

Domain Adaptation and Reasoning Frameworks in Language Models: A Controlled Experiment with Historical Cosmology

ArXiv CS.AI 50%

arXiv:2605.30415v1 Announce Type: cross Abstract: We investigate how domain adaptation reshapes explanatory behavior in language models using historical cosmology as a controlled setting. In Phase 1, we train a small language model from scratch on a pre-Copernican corpus from...

<https://arxiv.org/abs/2605.30415>

7 CONF

LongDS-Bench: On the Failure of Long-Horizon Agentic Data Analysis

ArXiv CS.AI 50%

arXiv:2605.30434v1 Announce Type: cross Abstract: Real-world data analysis is inherently iterative, yet existing benchmarks mostly evaluate isolated or short interactive tasks, leaving agents' ability to track evolving analytical context over long horizons untested. We introduce...

<https://arxiv.org/abs/2605.30434>

8 **CONF****idSCD: Identifying Training Datasets through Semantic Correlation Descriptors**

ArXiv CS.AI 50%

arXiv:2605.30462v1 Announce Type: cross Abstract: Can a dataset be recognized from the spurious correlations it induces during training? We argue that datasets leave dataset-specific traces in a model's learned semantic correlation structure: incidental regularities that are...

<https://arxiv.org/abs/2605.30462>
9 **CONF****A Novel Global Context-aware Deep Neural Network for Enhanced Brain Tumor Segmentation using Magnetic Resonance Images**

ArXiv CS.AI 50%

arXiv:2605.30510v1 Announce Type: cross Abstract: Brain cancer's severity necessitates precise brain tumor segmentation, which is crucial for effective brain tumor diagnosis. Manual identification, burdened by high costs, labor, and error risks, highlights the need for automated...

<https://arxiv.org/abs/2605.30510>
10 **CONF****Generalistic or Specific Embeddings, Which is Better? An Empirical Study on Search for Clinical Coding in Non-English Languages**

ArXiv CS.AI 50%

arXiv:2605.30529v1 Announce Type: cross Abstract: Sentence-embedding models for semantic search are overwhelmingly developed and evaluated on English corpora. When applied to clinical retrieval in other languages -- particularly retrieval of ICD-10-CM / CIE-10 codes -- recall...

<https://arxiv.org/abs/2605.30529>
1 **CONF****Simple Token-Efficient Vision-Language Model for Case-level Pathology Synoptic Report Generation**

ArXiv CS.AI 50%

arXiv:2605.30716v1 Announce Type: cross Abstract: Generating clinically useful pathology reports for pathology cases from whole-slide images (WSIs) is challenging due to gigapixel resolution, long visual-token sequences, and the complexity of case-level reasoning, where a single...

<https://arxiv.org/abs/2605.30716>
2 **CONF****Memory-Bound but Not Bandwidth-Limited: The Physical AI Inference Gap in Batch-1 LLM Decode**

ArXiv CS.AI 50%

arXiv:2605.30571v1 Announce Type: cross Abstract: Physical AI systems, including robots, autonomous vehicles, embodied agents and edge copilots, often run a different inference workload from cloud LLM serving: single-stream, batch-1 autoregressive decode, where one robot, camera...

<https://arxiv.org/abs/2605.30571>
3 **CONF****Prior Availability in Industrial Visual Sim-to-Real: A Review of CAD-Guided and CAD-Unavailable Regimes**

ArXiv CS.AI 50%

arXiv:2605.30581v1 Announce Type: cross Abstract: Industrial visual sim-to-real is often described as transferring from synthetic images to real images, but industrial deployment usually involves a broader mismatch between available evidence and required decisions. A system may...

<https://arxiv.org/abs/2605.30581>

- 4

CONF

Benchmarking Machine Learning Uncertainty Quantification Methodologies for Predicting Turbine Gas Temperature Degradation

ArXiv CS.AI

50%

arXiv:2605.30585v1 Announce Type: cross Abstract: Effective prognostics and health management of modern engines relies on accurate turbine gas temperature predictions and robust uncertainty quantification to ensure reliability and safety. This paper investigates five major...

<https://arxiv.org/abs/2605.30585>

- 5

CONF

Organizational Adaptation to Generative AI in Cybersecurity

ArXiv CS.AI

50%

arXiv:2506.12060v2 Announce Type: replace-cross Abstract: Cybersecurity organizations are adapting to GenAI integration through modified frameworks and hybrid operational processes, with success influenced by existing security maturity, regulatory requirements, and investments...

<https://arxiv.org/abs/2506.12060>

- 6

CONF

The Architecture of Errors: From Universal Impossibility to Patch-Local LLM Reliability

ArXiv CS.AI

50%

arXiv:2605.30628v1 Announce Type: cross Abstract: Universal LLM reliability is not a finite-library problem: across all possible tasks, tools, schemas, knowledge sources, and evaluator expectations, new intervention-distinguishable failure modes can appear without bound, so no...

<https://arxiv.org/abs/2605.30628>

- 7

CONF

Controllable Lung Nodule Synthesis via Histogram-Regularized Latent Diffusion Models

ArXiv CS.AI

50%

arXiv:2605.30631v1 Announce Type: cross Abstract: While automated diagnosis systems have achieved remarkable success in computed tomography (CT)-based lung cancer screening, their development remains limited by the scarcity of diverse, annotated pulmonary nodule datasets....

<https://arxiv.org/abs/2605.30631>

- 8

CONF

Boundary-Guided Policy Optimization for Memory-efficient RL of Diffusion Large Language Models

ArXiv CS.AI

50%

arXiv:2510.11683v3 Announce Type: replace-cross Abstract: A key challenge in applying reinforcement learning (RL) to diffusion large language models (dLLMs) is the intractability of their likelihood functions, which are essential for the RL objective, necessitating corresponding...

<https://arxiv.org/abs/2510.11683>

- 9

CONF

Same Patient, Different Words, Different Diagnosis? Evaluating Semantic Stability in Clinical LLMs

ArXiv CS.AI

50%

arXiv:2605.30646v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used in clinical applications. However, their behavior remains highly sensitive to subtle linguistic variations, such as rephrasing or syntactic variation. This sensitivity poses...

<https://arxiv.org/abs/2605.30646>

0 CONF

EUDAIMONIA: Evaluating Undesirable Dynamics in AI

ArXiv CS.AI 50%

arXiv:2605.30654v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used as conversational partners for companionship, emotional disclosure, and interpersonal advice, but the social dynamics of these interactions can create harms that are not captured...

<https://arxiv.org/abs/2605.30654>

1 CONF

Human-Alignment and Calibration of Inference-Time Uncertainty in Large Language Models

ArXiv CS.AI 50%

arXiv:2508.08204v2 Announce Type: replace-cross Abstract: There has been much recent interest in evaluating large language models for uncertainty calibration to facilitate model control and modulate user trust. Inference time uncertainty, which may provide a real-time signal to...

<https://arxiv.org/abs/2508.08204>

2 CONF

Investigating Detection and Obfuscation of Prompt Injection Attacks Against Software Reverse Engineering AI Agents

ArXiv CS.AI 50%

arXiv:2605.30677v1 Announce Type: cross Abstract: Agentic software reverse engineering systems are vulnerable to prompt injection attacks placed into the source code of executable binary files. This research demonstrates defensive tactics for detecting the presences of prompt...

<https://arxiv.org/abs/2605.30677>

3 CONF

How Early Adopters Used Generative AI Worldwide: Variation by Country Income and Language

ArXiv CS.AI 50%

arXiv:2605.30685v1 Announce Type: cross Abstract: AI is being used by people globally, but not everyone is using it in the same ways. Using a large-scale dataset of anonymized, de-identified, and privacy-scrubbed interactions with a widely available and free AI chatbot, we...

<https://arxiv.org/abs/2605.30685>

4 CONF

SAGE: A Novelty Gate for Efficient Memory Evolution in Agentic LLMs

ArXiv CS.AI 50%

arXiv:2605.30711v1 Announce Type: cross Abstract: Agentic LLMs must continuously decide whether newly extracted facts should be added, merged with existing memories, or ignored, yet prior work has focused more on retrieval and storage than on principled write-side control. We...

<https://arxiv.org/abs/2605.30711>

5 CONF

When are LLMs Sufficient Policy Optimizers for Sequential RL Tasks?

ArXiv CS.AI 50%

arXiv:2605.30719v1 Announce Type: cross Abstract: We study when large language models (LLMs) can serve as effective black-box policy optimizers for reinforcement learning (RL) tasks, i.e., when can we replace classical RL algorithms with an LLM? We explore this question by...

<https://arxiv.org/abs/2605.30719>

6 **CONF****OrcaRouter: A Production-Oriented LLM Router with Hybrid Offline-Online Learning**

ArXiv CS.AI 50%

arXiv:2605.30736v1 Announce Type: cross Abstract: The rapid development of large language models, each with distinct capabilities and inference costs, raises a practical deployment question: given an incoming request, which model should handle it? We present OrcaRouter, a...

<https://arxiv.org/abs/2605.30736>
7 **CONF****Chatterbox-Flash: Prior-Calibrated Block Diffusion for Streaming Zero-Shot TTS**

ArXiv CS.AI 50%

arXiv:2605.30748v1 Announce Type: cross Abstract: We present Chatterbox-Flash, a zero-shot text-to-speech model obtained by fine-tuning a pretrained autoregressive TTS decoder into a block-diffusion decoder, enabling parallel token generation within each block while retaining...

<https://arxiv.org/abs/2605.30748>
8 **CONF****Learning Cardiac Latent Representations in Vectorcardiogram Space**

ArXiv CS.AI 50%

arXiv:2605.31249v1 Announce Type: cross Abstract: Electrocardiography (ECG) is a cornerstone of cardiac assessment, making the learning of informative ECG representations fundamental to tasks ranging from disease diagnosis to clinical report generation. However, existing methods...

<https://arxiv.org/abs/2605.31249>
9 **CONF****FBHM: Functional Benchmarking and Steering of VLMs for Hateful Meme Detection**

ArXiv CS.AI 50%

arXiv:2605.31349v1 Announce Type: cross Abstract: Hateful meme detection remains a formidable challenge for vision-language models, as existing benchmarks are structurally observational - confounding rhetorical hate mechanisms with target community features and preventing causal...

<https://arxiv.org/abs/2605.31349>
10 **CONF****Design and Evaluation of Multi-Agent AI Oracle Systems for Prediction Market Resolution**

ArXiv CS.AI 50%

arXiv:2605.30802v1 Announce Type: cross Abstract: Prediction markets aggregate collective intelligence to forecast uncertain events, but their utility depends on reliable outcome resolution. Existing oracle systems tradeoff fast but brittle automation against accurate but costly...

<https://arxiv.org/abs/2605.30802>
1 **CONF****Unlearning in Diffusion Models: A Unified Framework with KL Divergence and Likelihood Constraints**

ArXiv CS.AI 50%

arXiv:2605.30825v1 Announce Type: cross Abstract: Unlearning in diffusion models aims to remove undesirable data or concepts while preserving the utility of pretrained models -- two fundamentally conflicting objectives. We propose a principled constrained optimization framework...

<https://arxiv.org/abs/2605.30825>

2 **CONF****Hide-and-Seek in Trajectories: Discovering Failure Signals for VLA Runtime Monitoring****ArXiv CS.AI** **50%**

arXiv:2605.30834v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models enable robots to follow natural language instructions and generalize across diverse tasks, but they remain vulnerable to execution failures that compromise reliability in real-world deployment....

<https://arxiv.org/abs/2605.30834>

3 **CONF****Shared Doubt: Zero-shot Cross-Lingual Confidence Estimation for Language Models****ArXiv CS.AI** **50%**

arXiv:2605.31220v1 Announce Type: cross Abstract: Confidence estimation (CE), i.e. quantifying the reliability of a model's prediction, has attracted great interest in the context of large language models (LLMs). However, most studies focus on English, ignoring the multilingual...

<https://arxiv.org/abs/2605.31220>

4 **CONF****DARTS: Distribution-Aware Active Rollout Trajectory Shaping for Accelerating LLM Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2605.30859v1 Announce Type: cross Abstract: Reinforcement Learning (RL) has become pivotal for improving model capabilities yet suffers from rollout efficiency bottlenecks due to the long-tail response length distribution. While existing works mitigate the impact of long...

<https://arxiv.org/abs/2605.30859>

5 **CONF****Federated Variational Preference Alignment with Gumbel-Softmax Prior for Personalized User Preferences****ArXiv CS.AI** **50%**

arXiv:2605.30873v1 Announce Type: cross Abstract: Federated Learning (FL) offers a privacy-preserving pathway for aligning Large Language Models (LLMs); however, existing frameworks typically enforce a monolithic reward model, inevitably averaging out inherently conflicting user...

<https://arxiv.org/abs/2605.30873>

6 **CONF****BlueFin: Benchmarking LLM Agents on Financial Spreadsheets****ArXiv CS.AI** **50%**

arXiv:2605.30907v1 Announce Type: cross Abstract: We present BlueFin, a benchmark that tasks large language model (LLM) agents with synthesis, manipulation, and comprehension tasks over spreadsheet workbooks in the professional finance domain. Though estimates of the global...

<https://arxiv.org/abs/2605.30907>

7 **CONF****Toxic HallucinAltions: Perturbing Prompts and Tracing LLM Circuits****ArXiv CS.AI** **50%**

arXiv:2605.30913v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in conversational settings where user tone ranges from polite to adversarial or toxic, yet less is known about whether toxic language in otherwise semantically equivalent...

<https://arxiv.org/abs/2605.30913>

8 CONF

TUX: Measuring Human--AI Tacit Understanding

ArXiv CS.AI 50%

arXiv:2605.30930v1 Announce Type: cross Abstract: As large language models (LLMs) increasingly act as collaborative partners, human--AI alignment is often evaluated through explicit task success, accuracy, or reward optimization. Yet many collaborative settings depend on tacit...

<https://arxiv.org/abs/2605.30930>

9 CONF

Do Large Language Models Encode Institutional Experience? Evidence from Cross-Linguistic Moral Reasoning Under Ambiguity

ArXiv CS.AI 50%

arXiv:2605.30934v1 Announce Type: cross Abstract: Large language models (LLMs) exhibit systematic differences in moral reasoning across languages, yet the source of this variation remains unclear. We test the hypothesis that languages encode aspects of the institutional...

<https://arxiv.org/abs/2605.30934>

0 CONF

Reasoning-Aware Multimodal Fusion for Hateful Video Detection

ArXiv CS.AI 50%

arXiv:2512.02743v2 Announce Type: replace-cross Abstract: Hate speech in online videos is posing an increasingly serious threat to digital platforms, especially as video content becomes increasingly multimodal and context-dependent. Existing methods often struggle to effectively...

<https://arxiv.org/abs/2512.02743>

1 CONF

Reading Between the Citations: A Typed Claim Network for Scientific Literature

ArXiv CS.AI 50%

arXiv:2605.30966v1 Announce Type: cross Abstract: Knowledge graphs over corpora of inter-referencing documents - scholarly papers, legal opinions, policy briefs - encode the topology of reference but not its stance. The standard representation collapses a rich evaluative...

<https://arxiv.org/abs/2605.30966>

2 CONF

From Prompt Injection to Persistent Control: Defending Agentic Harness Against Trojan Backdoors

ArXiv CS.AI 50%

arXiv:2605.31042v1 Announce Type: cross Abstract: LLM agents are evolving from conversational chatbots to operational tools in real-world workspaces. In local agentic harnesses, an LLM can read and write files, call tools, and reuse workspace state across sessions. While such...

<https://arxiv.org/abs/2605.31042>

3 CONF

Routing on the Stiefel Manifold: When Does Adaptive Subspace Selection Help for Cross-Domain EEG Decoding?

ArXiv CS.AI 50%

arXiv:2605.31043v1 Announce Type: cross Abstract: Cross-domain EEG decoding remains challenging despite advances in Riemannian deep learning: covariance matrices from different subjects occupy systematically distinct regions of the SPD manifold, yet existing domain adaptation...

<https://arxiv.org/abs/2605.31043>

4 **CONF****No More K-means: Single-Stage Sparse Coding for Efficient Multi-Vector Retrieval****ArXiv CS.AI** **50%**

arXiv:2605.30120v2 Announce Type: replace-cross Abstract: Multi-vector retrieval (MVR) models, exemplified by ColBERT, have established new benchmarks in retrieval accuracy by preserving fine-grained token-level interactions. However, this granularity imposes prohibitive storage...

<https://arxiv.org/abs/2605.30120>

5 **CONF****A Pilot Study on Curator-Guided Multilingual Art Description for Blind and Low-Vision Audiences with Small Vision-Language Models****ArXiv CS.AI** **50%**

arXiv:2605.31080v1 Announce Type: cross Abstract: Blind and low-vision (BLV) audiences remain underserved by visual art descriptions, particularly across languages and in museum settings where privacy and intellectual-property constraints may favour small on-premise...

<https://arxiv.org/abs/2605.31080>

6 **CONF****SpecDB: LLM-Generated Customized Databases via Feature-Oriented Decomposition****ArXiv CS.AI** **50%**

arXiv:2605.31097v1 Announce Type: cross Abstract: Mainstream relational databases ship a uniform feature set across deployments, although individual workloads exercise only a fraction of the available subsystems. We investigate whether a database can instead be generated on...

<https://arxiv.org/abs/2605.31097>

7 **CONF****KnowledgeGain: Evaluating and Optimizing Science News Generation for Reader Learning****ArXiv CS.AI** **50%**

arXiv:2605.31099v1 Announce Type: cross Abstract: Science news is an important medium to communicate discoveries between the research communities and the public. Yet, most metrics for generated or summarized text evaluate semantic similarity and factual consistency, but do not...

<https://arxiv.org/abs/2605.31099>

8 **CONF****From Evidence to Design: Developing an AI-Augmented UX Research Point of View for Digital Wellbeing in Emergency and Public Safety Contexts****ArXiv CS.AI** **50%**

arXiv:2605.31146v1 Announce Type: cross Abstract: This paper investigates how User Experience Research (UXR) methods can be combined with AI-supported analysis to develop clearer design direction for digital wellbeing interventions targeting Emergency and Public Safety Personnel...

<https://arxiv.org/abs/2605.31146>

9 **CONF****On the Robustness of Multilingual Text Embedding Rankings Across Learning Tasks, Languages, and Benchmark Datasets****ArXiv CS.AI** **50%**

arXiv:2605.31142v1 Announce Type: cross Abstract: Large-scale multilingual text embedding models play crucial role in both research and industry, yet their behavior in language-specific, multi-task settings remains insufficiently understood. Although benchmarking platforms such...

<https://arxiv.org/abs/2605.31142>

0 **CONF****Extending the UXR Point of View Pyramid: A Generative AI-Augmented Methodology for Human-Centred AI Systems****ArXiv CS.AI** **50%**

arXiv:2605.31143v1 Announce Type: cross Abstract: Rising household debt and cost-of-living pressures in the United Kingdom have intensified the role of AI-driven financial technologies in mediating credit assessment, repayment structuring, and debt support services. These...

<https://arxiv.org/abs/2605.31143>

1 **CONF****FOCUS: Forcing In-Context Object Localization through Visual Support Constraints and Policy Optimization****ArXiv CS.AI** **50%**

arXiv:2605.31145v1 Announce Type: cross Abstract: In-context localization (ICL) seeks to localize a target object specified by a small set of support examples in a query image, operating on the fly without training or parameter updates. Despite rapid advances in vision-language...

<https://arxiv.org/abs/2605.31145>

2 **CONF****D³: Dynamic Directional Graph-Constrained Data Scheduling for LLM Training****ArXiv CS.AI** **50%**

arXiv:2605.31164v1 Announce Type: cross Abstract: Training data plays a central role in large language models (LLMs) optimization, motivating extensive research on data scheduling strategies. Most existing approaches concentrate on adjusting the overall data distribution but...

<https://arxiv.org/abs/2605.31164>

3 **CONF****Emergent Languages in Populations of Language Model Agents: From Token Efficiency to Oversight Evasion****ArXiv CS.AI** **50%**

arXiv:2605.31170v1 Announce Type: cross Abstract: Monitoring autonomous language model agents currently relies mostly on surface behavior. But what happens when agent populations invent new languages with the goal of avoiding human oversight. Here, we study the emergent...

<https://arxiv.org/abs/2605.31170>

4 **CONF****MindVoice: Reconstructing Intelligible Speech from Non-invasive Neural Signals with Pretrained Priors****ArXiv CS.AI** **50%**

arXiv:2605.31173v1 Announce Type: cross Abstract: Reconstructing continuous speech from non-invasive neural recordings is a fundamental problem for probing human auditory perception and building safe, scalable speech brain-computer interfaces. Despite recent progress,...

<https://arxiv.org/abs/2605.31173>

5 **CONF****Steering LLMs? Actually, Sparse Autoencoders can outperform simple baselines****ArXiv CS.AI** **50%**

arXiv:2605.31183v1 Announce Type: cross Abstract: Sparse Autoencoders (SAEs) have been seen as a promising avenue for exploring the internals of Large Language Models (LLMs) and for steering model output generation. When AxBench - a model steering benchmark - was introduced in...

<https://arxiv.org/abs/2605.31183>

6 CONF

Comparing LLM-Based Conversational and Graphical Interfaces for Industrial Decision Tasks: An Exploratory Mixed-Methods Study

ArXiv CS.AI 50%

arXiv:2605.31224v1 Announce Type: cross Abstract: The use of Generative AI Conversational User Interfaces (CUI) as a new way to access and analyze data is growing in all sectors, and the industrial one is no exception. There, large amounts of data produced by IoT devices are...

<https://arxiv.org/abs/2605.31224>

7 CONF

Neural Network Verification using Partial Multi-Neuron Relaxation

ArXiv CS.AI 50%

arXiv:2605.30155v2 Announce Type: replace-cross Abstract: The increasing integration of deep neural networks in critical systems has spawned a theoretical and practical interest in formally guaranteeing safety properties about their behavior. To achieve this, contemporary...

<https://arxiv.org/abs/2605.30155>

8 CONF

Correcting Split Selection in Online Decision Trees via Anytime-Valid Inference

ArXiv CS.AI 50%

arXiv:2605.31239v1 Announce Type: cross Abstract: Bagging-based ensembles, most notably Adaptive Random Forests, are among the strongest performers for learning from data streams. A common denominator across these methods is their reliance on Hoeffding Trees as base learners,...

<https://arxiv.org/abs/2605.31239>

9 CONF

Entropic Projection Alignment: Estimating, Explaining, and Improving Model Performance Under Distribution Shift

ArXiv CS.AI 50%

arXiv:2605.31250v1 Announce Type: cross Abstract: We propose a unified framework for addressing three key challenges of distribution shift: (1) estimating a model's performance on an unlabeled target domain, (2) explaining the shift by identifying the features responsible, and...

<https://arxiv.org/abs/2605.31250>

0 CONF

Practical Cross-Band Channel Prediction for AI-RAN via Physics-Guided Deep Unfolding

ArXiv CS.AI 50%

arXiv:2605.31279v1 Announce Type: cross Abstract: To make cross-band channel prediction practical for AI-native RAN, algorithms must generalize across diverse environments and support real-time inference. Existing approaches achieve one but not both. To bridge this gap, we...

<https://arxiv.org/abs/2605.31279>

1 CONF

DeMaVLA: A Vision-Language-Action Foundation Model for Generalizable Deformable Manipulation

ArXiv CS.AI 50%

arXiv:2605.31286v1 Announce Type: cross Abstract: Real-world household robots require Vision-Language-Action (VLA) foundation models that can acquire reusable manipulation skills across diverse objects, task conditions, and household environments. Deformable-object folding is a...

<https://arxiv.org/abs/2605.31286>

2 CONF

Appropriateness of Empathy in AI: A Signal-Cost Perspective

ArXiv CS.AI 50%

arXiv:2605.31340v1 Announce Type: cross Abstract: The appropriateness of empathy in AI has emerged as a critical concern, as excessive empathy risks seeming manipulative while insufficient empathy appears dismissive. While prior research has explored how to quantify empathy in...

<https://arxiv.org/abs/2605.31340>

3 CONF

Dreaming Of Others: Latent Teammate Modeling In World Models For Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.31361v1 Announce Type: cross Abstract: In cooperative multi-agent reinforcement learning (MARL), agents must coordinate with partners whose internal policies and intentions are not directly observable. While world models such as Dreamer have demonstrated strong...

<https://arxiv.org/abs/2605.31361>

4 CONF

The Sword, Shield, and Achilles' Heel: Characterizing the Linguistic Inductive Bias of Large Language Models for Spatial Reasoning in Navigation Planning

ArXiv CS.AI 50%

arXiv:2605.31404v1 Announce Type: cross Abstract: Large Language Model (LLM)-based navigation systems commonly construct explicit spatial representations (e.g., topological graphs, semantic raster maps) and translate them into textual descriptions as LLMs' inputs. However, the...

<https://arxiv.org/abs/2605.31404>

5 CONF

Neuro-symbolic Syntactic Parsing: Shaping a Neural Network with the CYK Algorithm

ArXiv CS.AI 50%

arXiv:2605.31421v1 Announce Type: cross Abstract: In this paper, we show the possibility of a direct injection of algorithms into neural network architecture. We focus on a complex algorithm, that is, Cocke-Younger-Kasami (CYK) for parsing context-free grammars in Chomsky Normal...

<https://arxiv.org/abs/2605.31421>

6 CONF

DOA: Training-Free Decoder-Only Attention Policy for Long-Form Simultaneous Translation with SpeechLLMs

ArXiv CS.AI 50%

arXiv:2605.31432v1 Announce Type: cross Abstract: Simultaneous speech-to-text translation (SimulST) generates translations while speech is still unfolding, requiring a streaming policy that decides when to read and when to write. State-of-the-art approaches rely on...

<https://arxiv.org/abs/2605.31432>

7 CONF

Used Car Salesbots? Honesty and Credulity of LLMs as Bargaining Agents under Partial Information

ArXiv CS.AI 50%

arXiv:2605.31445v1 Announce Type: cross Abstract: In this work we study agents in simulated bargaining scenarios, where a buyer and a seller communicate through a text channel and attempt to negotiate mutually beneficial trades, under different information regimes (complete...

<https://arxiv.org/abs/2605.31445>

8 CONF

Fine-grained Verification via Diagnostic Reasoning Supervision for Aspect Sentiment Triplet Extraction

ArXiv CS.AI 50%

arXiv:2605.31446v1 Announce Type: cross Abstract: Aspect Sentiment Triplet Extraction (ASTE) aims to identify aspect terms, opinion terms, and sentiment polarities as structured triplets, providing essential inputs for downstream information system applications such as opinion...

<https://arxiv.org/abs/2605.31446>

9 CONF

PithTrain: A Compact and Agent-Native MoE Training System

ArXiv CS.AI 50%

arXiv:2605.31463v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) has become the dominant architecture for frontier language models. To meet this demand, production frameworks have built optimized MoE training stacks over years of engineering effort. Yet evolving these...

<https://arxiv.org/abs/2605.31463>

00 CONF

GPU Forecasters: Language Models as Selective Surrogates for Kernel Runtime Optimization

ArXiv CS.AI 50%

arXiv:2605.31464v1 Announce Type: cross Abstract: GPU kernels are the workhorse of modern deep learning, and optimizing them (via evolutionary search or coding agents) usually requires repeated measurement on target hardware. While these measurements provide the ground-truth...

<https://arxiv.org/abs/2605.31464>

01 CONF

If LLMs Have Human-Like Attributes, Then So Does Age of Empires II

ArXiv CS.AI 50%

arXiv:2605.31514v1 Announce Type: cross Abstract: Much research has been carried out on large language models (LLMs) and LLM-powered agentic workflows. However, many works within the field state emergence of, ascribe to, or assume, generalised anthropomorphic attributes to them...

<https://arxiv.org/abs/2605.31514>

02 CONF

What Gets Unmasked First? Trajectory Analysis of Diffusion Models for Graph-to-Text Generation

ArXiv CS.AI 50%

arXiv:2605.31564v1 Announce Type: cross Abstract: We present the first systematic study of masked diffusion language models (MDLMs) for graph-to-text generation. We analyze MDLM generation trajectories -- the order in which tokens are unmasked during iterative decoding -- and...

<https://arxiv.org/abs/2605.31564>

03 CONF

Language Models Learn Constructional Semantics, Not To Mention Syntax: Investigating LM Understanding of Paired-Focus Constructions

ArXiv CS.AI 50%

arXiv:2605.31586v1 Announce Type: cross Abstract: Grasping the semantics of rare constructions (form-meaning pairings) has been shown to be a challenging problem that has currently only been solved by the largest LLMs. It remains an open question if open-source models have...

<https://arxiv.org/abs/2605.31586>

04 CONF

Chain-of-Thought Reasoning In The Wild Is Not Always Faithful

ArXiv CS.AI 50%

arXiv:2503.08679v5 Announce Type: replace Abstract: Recent studies indicate that when faced with explicit biases in prompts, models often omit mentioning these biases in their Chain-of-Thought (CoT) output, revealing that verbalized reasoning can give an incorrect picture of how...

<https://arxiv.org/abs/2503.08679>

05 CONF

SEMA-RAG: A Self-Evolving Multi-Agent Retrieval-Augmented Generation Framework for Medical Reasoning

ArXiv CS.AI 50%

arXiv:2605.17101v2 Announce Type: replace-cross Abstract: Retrieval-Augmented Generation (RAG) is widely employed to mitigate risks such as hallucinations and knowledge obsolescence in medical question answering, yet its predominantly single-round, static retrieval paradigm...

<https://arxiv.org/abs/2605.17101>

06 CONF

ReTabAD: A Benchmark for Restoring Semantic Context in Tabular Anomaly Detection

ArXiv CS.AI 50%

arXiv:2510.02060v2 Announce Type: replace Abstract: In tabular anomaly detection (AD), textual semantics often carry critical signals, as the definition of an anomaly is closely tied to domain-specific context. However, existing benchmarks provide only raw data points without...

<https://arxiv.org/abs/2510.02060>

07 CONF

Post-Training LLMs as Better Decision-Making Agents: A Regret-Minimization Approach

ArXiv CS.AI 50%

arXiv:2511.04393v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly deployed as "agents" for decision-making (DM) in interactive and dynamic environments. Yet, since they were not originally designed for DM, recent studies show that LLMs can...

<https://arxiv.org/abs/2511.04393>

08 CONF

Agentic Physical AI toward a Domain-Specific Foundation Model for Energy Systems: A Case Study on Nuclear Reactor Control

ArXiv CS.AI 50%

arXiv:2512.23292v4 Announce Type: replace Abstract: The prevailing paradigm in AI for physical systems: scaling general-purpose foundation models toward universal multimodal reasoning, confronts a barrier at the control interface. Frontier vision-language models achieve only...

<https://arxiv.org/abs/2512.23292>

09 CONF

ConSensus: Multi-Agent Collaboration for Multimodal Sensing

ArXiv CS.AI 50%

arXiv:2601.06453v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly grounded in sensor data to perceive and reason about human physiology and the physical world. However, accurately interpreting heterogeneous multimodal sensor data remains a...

<https://arxiv.org/abs/2601.06453>

10 CONF

Diagnosing the Reliability of LLM-as-a-Judge via Item Response Theory

ArXiv CS.AI 50%

arXiv:2602.00521v2 Announce Type: replace Abstract: While LLM-as-a-Judge is widely used in automated evaluation, existing validation practices primarily operate at the level of observed outputs, offering limited insight into whether LLM judges themselves function as stable and...

<https://arxiv.org/abs/2602.00521>

11 CONF

MedCoG: Maximizing LLM Inference Density in Medical Reasoning via Meta-Cognitive Regulation

ArXiv CS.AI 50%

arXiv:2602.07905v2 Announce Type: replace Abstract: Large Language Models (LLMs) have shown strong potential in complex medical reasoning yet face diminishing gains under inference scaling laws. While existing studies augment LLMs with various knowledge types, it remains unclear...

<https://arxiv.org/abs/2602.07905>

12 CONF

Discovering Differences in Strategic Behavior Between Humans and LLMs

ArXiv CS.AI 50%

arXiv:2602.10324v2 Announce Type: replace Abstract: As Large Language Models (LLMs) are increasingly deployed in social and strategic scenarios, it becomes critical to understand where and why their behavior diverges from that of humans. While behavioral game theory (BGT)...

<https://arxiv.org/abs/2602.10324>

13 CONF

Counterfactual Trace Auditing of LLM Agent Skills

ArXiv CS.AI 50%

arXiv:2605.11946v2 Announce Type: replace Abstract: Large Language Model agents are increasingly augmented with agent skills. Current evaluation methods for skills remain limited. Most deployed benchmarks report only pass rate before and after a skill is attached, treating the...

<https://arxiv.org/abs/2605.11946>

14 CONF

BoxLitE: A Faithful Knowledge Base Embedding Based on Convex Optimization

ArXiv CS.AI 50%

arXiv:2605.23937v2 Announce Type: replace Abstract: Knowledge base (KB) embeddings aim at combining the capability of classical knowledge graph embeddings to generalize the information present in facts, the ABox, with conceptual knowledge represented in an ontology language, the...

<https://arxiv.org/abs/2605.23937>

15 CONF

MIRA: Mid-training Rubric Anchoring for Source-Aware Data Selection

ArXiv CS.AI 50%

arXiv:2605.30288v2 Announce Type: replace Abstract: Mid-training has become an important stage in modern LLM development, using large-scale curated mixtures to strengthen capabilities before final post-training. Its data selection problem is distinct: the data are optimized...

<https://arxiv.org/abs/2605.30288>

16 CONF

Gap-K%: Measuring Top-1 Prediction Gap for Detecting Pretraining Data

ArXiv CS.AI 50%

arXiv:2601.19936v2 Announce Type: replace-cross Abstract: The opacity of massive pretraining corpora in Large Language Models (LLMs) raises significant privacy and copyright concerns, making pretraining data detection a critical challenge. Existing state-of-the-art methods...

<https://arxiv.org/abs/2601.19936>

17 CONF

EMCEE: Improving Multilingual Capability of LLMs via Bridging Knowledge and Reasoning with Extracted Synthetic Multilingual Context

ArXiv CS.AI 50%

arXiv:2503.05846v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) have achieved impressive progress across a wide range of tasks, yet their heavy reliance on English-centric training data leads to significant performance degradation in non-English languages....

<https://arxiv.org/abs/2503.05846>

18 CONF

Who Gets Credit or Blame? Attributing Accountability in Modern AI Systems

ArXiv CS.AI 50%

arXiv:2506.00175v5 Announce Type: replace-cross Abstract: Modern AI systems are typically developed through multiple stages-pretraining, fine-tuning rounds, and subsequent adaptation or alignment, where each stage builds on the previous ones and updates the model in distinct...

<https://arxiv.org/abs/2506.00175>

19 CONF

PictSure: Pretraining Embeddings Matters for In-Context Learning Image Classifiers

ArXiv CS.AI 50%

arXiv:2506.14842v2 Announce Type: replace-cross Abstract: Building image classification models remains cumbersome in data-scarce domains, where collecting large labeled datasets is impractical. In-context learning (ICL) is a promising paradigm for few-shot image classification...

<https://arxiv.org/abs/2506.14842>

20 CONF

Human Psychometric Questionnaires Mischaracterize LLM Behavior

ArXiv CS.AI 50%

arXiv:2509.10078v4 Announce Type: replace-cross Abstract: We examine whether human psychometric questionnaires can serve as reliable tools for characterizing and predicting LLM behavior in everyday user interactions. We analyze eight open-source LLMs by comparing their value and...

<https://arxiv.org/abs/2509.10078>

21 CONF

LLMs Lean on Priors, Not Programming Language Semantics

ArXiv CS.AI 50%

arXiv:2510.03415v3 Announce Type: replace-cross Abstract: Recent work asks whether large language models (LLMs) condition their reasoning on explicit rules rather than statistical regularities from pretraining. Program execution provides a canonical instance: formal semantics...

<https://arxiv.org/abs/2510.03415>

22 CONF

OBCache: Optimal Brain KV Cache Pruning for Efficient Long-Context LLM Inference

ArXiv CS.AI 50%

arXiv:2510.07651v2 Announce Type: replace-cross Abstract: Large language models (LLMs) with extended context windows enable powerful applications but impose significant memory overhead, as caching all key-value (KV) states scales linearly with sequence length and batch size....

<https://arxiv.org/abs/2510.07651>

23 CONF

Scaling Multi-Agent Environment Co-Design with Diffusion Models

ArXiv CS.AI 50%

arXiv:2511.03100v2 Announce Type: replace-cross Abstract: The agent-environment co-design paradigm jointly optimises agent policies and environment configurations in search of improved system performance. With application domains ranging from warehouse logistics to windfarm...

<https://arxiv.org/abs/2511.03100>

24 CONF

Conditional Coverage Diagnostics for Conformal Prediction

ArXiv CS.AI 50%

arXiv:2512.11779v2 Announce Type: replace-cross Abstract: Evaluating conditional coverage remains one of the most persistent challenges in assessing the reliability of predictive systems. Although conformal methods can give guarantees on marginal coverage, no method can...

<https://arxiv.org/abs/2512.11779>

25 CONF

Bottom-up Policy Optimization: Your Language Model Policy Secretly Contains Internal Policies

ArXiv CS.AI 50%

arXiv:2512.19673v3 Announce Type: replace-cross Abstract: Existing reinforcement learning (RL) approaches treat large language models (LLMs) as a unified policy, overlooking their internal mechanisms. In this paper, we decompose the LLM-based policy into Internal Layer Policies...

<https://arxiv.org/abs/2512.19673>

26 CONF

Rethinking Multimodal Few-Shot 3D Point Cloud Segmentation: From Fused Refinement to Decoupled Arbitration

ArXiv CS.AI 50%

arXiv:2601.01456v2 Announce Type: replace-cross Abstract: In this paper, we revisit multimodal few-shot 3D point cloud semantic segmentation (FS-PCS), identifying a conflict in "Fuse-then-Refine" paradigms: the "Plasticity-Stability Dilemma." In addition, CLIP's inter-class...

<https://arxiv.org/abs/2601.01456>

27 CONF

Plain Transformers are Surprisingly Powerful Link Predictors

ArXiv CS.AI 50%

arXiv:2602.01553v2 Announce Type: replace-cross Abstract: Link prediction is a core challenge in graph machine learning, demanding models that capture rich and complex topological dependencies. While Graph Neural Networks (GNNs) are the standard solution, state-of-the-art...

<https://arxiv.org/abs/2602.01553>

28 CONF

Stop the Flip-Flop: Context-Preserving Verification for Fast Revocable Diffusion Decoding

ArXiv CS.AI 50%

arXiv:2602.06161v2 Announce Type: replace-cross Abstract: Parallel diffusion decoding can accelerate diffusion language model inference by unmasking multiple tokens per step, but aggressive parallelism often harms quality. Revocable decoding mitigates this by rechecking earlier...

<https://arxiv.org/abs/2602.06161>

29 CONF

Pull Requests as a Training Signal for Repo-Level Code Editing

ArXiv CS.AI 50%

arXiv:2602.07457v2 Announce Type: replace-cross Abstract: Repository-level code editing requires models to understand complex dependencies and execute precise multi-file modifications across a large codebase. While recent gains on SWE-bench rely heavily on complex agent...

<https://arxiv.org/abs/2602.07457>

30 CONF

Effective Reasoning Chains Reduce Intrinsic Dimensionality

ArXiv CS.AI 50%

arXiv:2602.09276v2 Announce Type: replace-cross Abstract: Chain-of-thought (CoT) reasoning and its variants have substantially improved the performance of language models on complex reasoning tasks, yet the precise mechanisms by which different strategies facilitate...

<https://arxiv.org/abs/2602.09276>

31 CONF

Biases in the Blind Spot: Detecting What LLMs Fail to Mention

ArXiv CS.AI 50%

arXiv:2602.10117v5 Announce Type: replace-cross Abstract: Large Language Models (LLMs) often provide chain-of-thought (CoT) reasoning traces that appear plausible, but may hide internal biases. We call these un verbalized biases. Monitoring models via their stated reasoning is...

<https://arxiv.org/abs/2602.10117>

32 CONF

Less is Enough: Synthesizing Diverse Data in LLM Feature Space with Sparse Autoencoders

ArXiv CS.AI 50%

arXiv:2602.10388v4 Announce Type: replace-cross Abstract: The diversity of post-training data is critical for effective downstream performance in large language models (LLMs). Many existing approaches to constructing post-training data quantify diversity using text-based metrics...

<https://arxiv.org/abs/2602.10388>

33 CONF

SCOPE: Selective Conformal Optimized Pairwise LLM Judging

ArXiv CS.AI 50%

arXiv:2602.13110v3 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly used as scalable judges in pairwise evaluation, but they remain prone to miscalibration and biases. We propose SCOPE (Selective Conformal Optimized Pairwise Evaluation), a...

<https://arxiv.org/abs/2602.13110>

34 CONF

The Global Landscape of Environmental AI Regulation: From the Cost of Reasoning to a Right to Green AI

ArXiv CS.AI 50%

arXiv:2603.00068v2 Announce Type: replace-cross Abstract: Artificial intelligence (AI) systems impose substantial and growing environmental costs, yet transparency about these impacts has declined even as their deployment has accelerated. This paper makes three contributions....

<https://arxiv.org/abs/2603.00068>

35 CONF

Circuit-Inspired High-Order Neural Networks with Unified Neural Dynamics Modeling for PDE Solving and Visual Perception

ArXiv CS.AI 50%

arXiv:2603.23977v2 Announce Type: replace-cross Abstract: Deep networks often rely on architectural heuristics to shape representation evolution, limiting their ability to model data governed by intrinsic dynamics. We present the Circuit-inspired High-Order Neural Network...

<https://arxiv.org/abs/2603.23977>

36 CONF

Beyond Static Uncertainty: Modeling Temporal Uncertainty Dynamics for Probabilistic Time Series Forecasting

ArXiv CS.AI 50%

arXiv:2603.24254v3 Announce Type: replace-cross Abstract: Real-world time series exhibit temporally structured uncertainty: volatility clusters in turbulent regimes, dissipates in stable periods, and shifts abruptly around structural breaks. Yet many probabilistic forecasting...

<https://arxiv.org/abs/2603.24254>

37 CONF

Symmetry Reveals Layerwise Dynamics: How Transformers Perform In-Context Classification

ArXiv CS.AI 50%

arXiv:2604.11613v3 Announce Type: replace-cross Abstract: Transformers can perform in-context classification from a few labeled examples, yet the inference-time algorithm remains opaque. We study multi-class linear classification in the hard no-margin regime and make the...

<https://arxiv.org/abs/2604.11613>

38 CONF

Aligning Dense Retrievers with LLM Utility via Distillation

ArXiv CS.AI 50%

arXiv:2604.22722v2 Announce Type: replace-cross Abstract: Dense vector retrieval is the practical backbone of Retrieval-Augmented Generation (RAG), but similarity search can suffer from precision limitations. Conversely, utility-based approaches leveraging LLM re-ranking often...

<https://arxiv.org/abs/2604.22722>

39 CONF

Self-Captioning Multimodal Interaction Tuning: Amplifying Exploitable Redundancies for Robust Vision Language Models

ArXiv CS.AI 50%

arXiv:2605.08145v2 Announce Type: replace-cross Abstract: Current vision language models face hallucination and robustness issues against ambiguous or corrupted modalities. We hypothesize that these issues can be addressed by exploiting the shared information between modalities...

<https://arxiv.org/abs/2605.08145>

40 CONF

Spurious Correlation Learning in Preference Optimization: Mechanisms, Consequences, and Mitigation via Tie Training

ArXiv CS.AI 50%

arXiv:2605.11134v2 Announce Type: replace-cross Abstract: Preference learning methods like Direct Preference Optimization (DPO) are known to induce reliance on spurious correlations, leading to sycophancy and length bias in today's language models and potentially severe goal...

<https://arxiv.org/abs/2605.11134>

41 CONF

Pair-In, Pair-Out: Latent Multi-Token Prediction for Efficient LLMs

ArXiv CS.AI 50%

arXiv:2605.27255v2 Announce Type: replace-cross Abstract: Long chain-of-thought reasoning has made autoregressive decoding the dominant inference cost of modern large language models. Existing methods target either the input side (latent compression) or the output side...

<https://arxiv.org/abs/2605.27255>

42 CONF

Pocket-Dentist: On-Device Dental Image Understanding via Efficient Multimodal Large Language Models

ArXiv CS.AI 50%

arXiv:2605.29299v2 Announce Type: replace-cross Abstract: Evaluations of dental vision-language models remain fragmented across datasets, task definitions and metrics, and often ignore their computational cost. This limits their widespread deployment for dental screening outside...

<https://arxiv.org/abs/2605.29299>