



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

## Source Breakdown

| Source          | Articles | Category  |
|-----------------|----------|-----------|
| ArXiv CS.AI     | 273      | Media     |
| TechCrunch AI   | 7        | Media     |
| The Verge AI    | 6        | Media     |
| HackerNews      | 2        | Community |
| MIT Tech Review | 2        | Media     |

## VERIFIED INTELLIGENCE — 5 stories

VERI

**Principled Synthetic Data Enables the First Scaling Laws for LLMs in Recommendation**

ArXiv CS.AI +1 source

80%

arXiv:2602.07298v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) represent a promising frontier for recommender systems, yet their development has been impeded by the absence of predictable scaling laws, which are crucial for guiding research and optimizing...

<https://arxiv.org/abs/2602.07298>

VERI

**Anthropic has officially filed to go public**

The Verge AI +1 source

80%

After months of speculation about whether OpenAI or Anthropic would be first in their race to IPO, Anthropic on Monday reached a key milestone: filing to kick off the process with the U.S. Securities and Exchange Commission. The filing sets the stage for what's sure to be a...

<https://www.theverge.com/ai-artificial-intelligence/941016/anthropic-has-officially-fil...>

VERI

**This AI weather startup is out-forecasting government agencies**

TechCrunch AI +1 source

80%

WindBorne benefits from its unique combination of model-building and data collection. The company now has about 400 balloons in flight gathering sensor readings at any given time, launched from 15 sites around the globe. The advances in its current model come from improvements...

<https://techcrunch.com/2026/06/01/this-ai-weather-startup-is-out-forecasting-government...>

VERI

**Characterizing Web Search in The Age of Generative AI**

ArXiv CS.AI +1 source

80%

arXiv:2510.11560v2 Announce Type: replace-cross Abstract: The advent of LLMs has given rise to generative search, a new search paradigm in which LLMs retrieve information from the web related to a query and synthesize it into a single, coherent response. This paradigm differs...

<https://arxiv.org/abs/2510.11560>

VERI

**Gemini's new AI agent is about as good as Google's demo**

The Verge AI +1 source

80%

Google's new "24/7" AI agent, Gemini Spark, can be shockingly good at doing things on your behalf. But I'm not sure it's worth the financial cost and potential privacy tradeoffs. The company gave me access to Spark last week. Google advertises Spark as an AI agent that can take...

<https://www.theverge.com/tech/941138/google-gemini-spark-ai-agent-hands-on>

## CONFIRMED SIGNALS — 280 stories

CONF

**Show HN: Viveka: filter LLM output against a Lean-verified Advaita Vedanta model**

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/SpecStudio-net/Viveka>

CONF

**Show HN: Agentpack - isolated config layers for Claude Code, Codex, and OpenCode**

HackerNews 50% HN 4

No summary — see link for full article.

<https://nexo.sh/posts/agentpack/>

CONF

**Alphabet plans to raise \$80B to pay for AI buildout**

TechCrunch AI 50%

"The company is experiencing strong demand for its AI solutions and services from enterprises and consumers, at levels that are exceeding the company's available supply," Alphabet said in its statement.

<https://techcrunch.com/2026/06/01/alphabet-plans-to-raise-80-billion-to-pay-for-ai-buil...>

CONF

**Nvidia chases \$200B CPU market with AI agent PCs from Microsoft, Dell, and HP**

TechCrunch AI 50%

If Nvidia has cracked a way to bring AI agents easily, safely, and usefully to the masses, it could — and should — be big.

<https://techcrunch.com/2026/06/01/nvidia-chases-200b-cpu-market-with-ai-agent-pcs-from-...>

0

CONF

**DuckDuckGo makes its 'no-AI' search engine easier to access as its traffic booms**

TechCrunch AI 50%

Alternative search engine DuckDuckGo launches 'no AI' web extensions for Chrome and Firefox users.

<https://techcrunch.com/2026/06/01/duckduckgo-makes-its-no-ai-search-engine-easier-to-ac...>

1

CONF

**The Download: China's brain implant ambitions**

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. China has approved the world's first invasive brain-computer chip—here's what's next Sitting in the courtyard of his house in China's Henan...

<https://www.technologyreview.com/2026/06/01/1138207/the-download-china-bci-brain-implan...>

2

CONF

**China has approved the world's first invasive brain-computer chip—here's what's next**

MIT Tech Review 50%

One day last October, sitting in the courtyard of his house in China's Henan province, Dong Hui decided to see if he could hold a pen to write. Dong, 39, had sustained spinal cord injuries in a car accident six years earlier that left him paralyzed from the neck down. Slowly but...

<https://www.technologyreview.com/2026/06/01/1138133/china-world-first-brain-chip/>

3

CONF

**Meta's own AI was exploited to hijack Instagram accounts**

The Verge AI 50%

Meta's AI support chatbot helped hackers hijack Instagram accounts, as reported earlier by 404 Media. In a video shared on Telegram, a hacker shows how they could take over an account by asking Meta's chatbot to switch the email associated with someone else's profile and then...

<https://www.theverge.com/tech/941179/meta-instagram-ai-support-chatbot-exploit-hacked>

4

CONF

**Microsoft to unveil new AI models and Windows improvements at Build**

The Verge AI 50%

Microsoft is heading to San Francisco this week in a bid to win back developers at its Build conference. I've been attending Build since the days when Microsoft called it the Professional Developers Conference, and I can't remember a more pivotal moment. As Microsoft continues...

<https://www.theverge.com/report/940861/microsoft-build-ai-models-windows-dev-mode-what-...>

5 **CONF****AI is blowing up music. How should the Grammys handle it?****The Verge AI** 50%

Today I'm talking with Harvey Mason Jr., who is CEO of the Recording Academy — that's the outfit that puts on the Grammy Awards. I last talked to Harvey in 2024, when it was obvious that generative AI would upend the music industry, but still not exactly clear how that would...

<https://www.theverge.com/podcast/940831/ai-grammys-music-recording-harvey-mason>
6 **CONF****Strava blames zero-code AI apps and scrapers as it tightens API access****The Verge AI** 50%

The popular fitness-tracking platform, Strava, is restricting access to its API as part of efforts to clamp down on AI scraping, as reported earlier by TechCrunch. Developers who want to build an app using Strava's data now need to pay for a flat \$11.99 / month subscription. In...

<https://www.theverge.com/gadgets/940854/strava-restricts-api-access-ai-apps>
7 **CONF****Efficient LLM Moderation with Multi-Layer Latent Prototypes****ArXiv CS.AI** 50%

arXiv:2502.16174v4 Announce Type: replace-cross Abstract: Although modern LLMs are aligned with human values during post-training, robust moderation remains essential to prevent harmful outputs at deployment time. Existing approaches suffer from performance-efficiency trade-offs...

<https://arxiv.org/abs/2502.16174>
8 **CONF****MindGames Arena Generalization Track: In2AI Solution with Delayed Per-Step Reward Attribution****ArXiv CS.AI** 50%

arXiv:2606.00017v1 Announce Type: new Abstract: Training language model agents for multi-agent strategic interaction presents a core difficulty: the quality of any action may depend on future events that never materialize, on moves that violate game rules, or on decisions made...

<https://arxiv.org/abs/2606.00017>
9 **CONF****Universal Quantum Transformer****ArXiv CS.AI** 50%

arXiv:2606.00045v1 Announce Type: new Abstract: Classical continuous-space neural networks fundamentally struggle to lock into exact mathematical symmetries, such as modular arithmetic and non-commutative algebra. To approximate these discrete logical rules, they often rely on...

<https://arxiv.org/abs/2606.00045>
0 **CONF****Relational Intervention During Functional Collapse in Large Language Models: A Lexical-Statistical Ablation and a Structure x Register Factorial****ArXiv CS.AI** 50%

arXiv:2606.00935v1 Announce Type: new Abstract: We test whether a relational-style intervention delivered during functional collapse in a small language model produces post-collapse behavior distinguishable from technical feedback, from a lexically-matched scrambled control, and...

<https://arxiv.org/abs/2606.00935>

1 **CONF****A Multi-AI-agent Framework Enabling End-to-end Finite Element Analysis for Solid Mechanics Problems****ArXiv CS.AI** **50%**

arXiv:2606.00138v1 Announce Type: new Abstract: Finite element analysis (FEA) is the most important numerical approach for solid mechanics. Challenges of FEA include a steep learning curve for entry-level users and potential false simulations due to incorrect definitions of key...

<https://arxiv.org/abs/2606.00138>

2 **CONF****TIGER: Traceable Inference with Graph-Based Evidence Routing for Mitigating Hallucinations in Multimodal Generation****ArXiv CS.AI** **50%**

arXiv:2606.00232v1 Announce Type: new Abstract: We study fact-level repair for multimodal generation, where a fluent output may contain specific facts that are not supported by the input. Existing inference-time repair methods often generate feedback by jointly conditioning on...

<https://arxiv.org/abs/2606.00232>

3 **CONF****Capability Self-Assessment: Teaching LLMs to Know Their Limits****ArXiv CS.AI** **50%**

arXiv:2606.00251v1 Announce Type: new Abstract: The ability to recognize one's own limitations and decide whether to solve a problem or delegate is fundamental for reliable intelligent systems. Yet we show that modern large language models systematically lack this ability:...

<https://arxiv.org/abs/2606.00251>

4 **CONF****PR2: Predictive Routing Replay for MoE-Based LLM Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2606.00395v1 Announce Type: cross Abstract: Mixture of Experts (MoE) Large Language Models (LLMs) achieve strong performance at scale. However, reinforcement learning (RL) on MoE-based LLMs often suffers from training instability. A root cause is router drift, i.e., expert...

<https://arxiv.org/abs/2606.00395>

5 **CONF****Coupling Language Models with Physics-based Simulation for Synthesis of Inorganic Materials****ArXiv CS.AI** **50%**

arXiv:2606.00315v1 Announce Type: new Abstract: Modern generative machine learning (ML) models can propose novel inorganic crystalline materials with targeted properties; however, synthesis planning of these materials remains difficult due to the complexity of the associated...

<https://arxiv.org/abs/2606.00315>

6 **CONF****From Noise to Order: Learning to Rank via Denoising Diffusion****ArXiv CS.AI** **50%**

arXiv:2602.11453v2 Announce Type: replace-cross Abstract: In information retrieval (IR), learning-to-rank (LTR) methods have traditionally limited themselves to discriminative machine learning approaches that model the probability of the document being relevant to the query...

<https://arxiv.org/abs/2602.11453>

7 CONF

**The Deterministic Horizon: When Extended Reasoning Fails and Tool Delegation Becomes Necessary**

ArXiv CS.AI 50%

arXiv:2606.00376v1 Announce Type: new Abstract: Extended chain-of-thought reasoning can degrade performance on deterministic state-tracking tasks, not due to preference biases, but limits rooted in the information-theoretic capacity of decoder-only attention. We establish: (1)...

<https://arxiv.org/abs/2606.00376>

8 CONF

**Doing What They Say, Not What They Reason: Locating the Faithfulness Gap in LLM Agents**

ArXiv CS.AI 50%

arXiv:2606.00476v1 Announce Type: new Abstract: Do LLM agents act on the reasoning they state? This question of process fidelity is central to using LLMs in social simulation, yet it is hard to measure where no reference for correct behavior exists. We study it in a controlled...

<https://arxiv.org/abs/2606.00476>

9 CONF

**BudgetDraft: Acceptance-Aware Multi-View Training for Sparse-KV Speculative Decoding**

ArXiv CS.AI 50%

arXiv:2606.00144v1 Announce Type: cross Abstract: Speculative decoding speeds up autoregressive decoding by using a drafter to propose multiple tokens that a verifier validates in parallel. In resource-constrained deployments, the drafter uses a sparse KV cache to limit peak GPU...

<https://arxiv.org/abs/2606.00144>

0 CONF

**EnergyMamba: An Uncertainty-Aware Graph-Enhanced Selective State Space Model for Energy Consumption Prediction**

ArXiv CS.AI 50%

arXiv:2606.00506v1 Announce Type: new Abstract: Energy consumption prediction is essential for efficient grid management, demand-side optimization, and sustainable energy planning. Although advanced machine learning methods have been employed for better prediction performance,...

<https://arxiv.org/abs/2606.00506>

1 CONF

**Threshold-Based Exclusive Batching for LLM Inference**

ArXiv CS.AI 50%

arXiv:2606.00516v1 Announce Type: new Abstract: Mixed batching (MB)--interleaving prefill and decode in a single batch--has become the standard scheduling strategy for large language model (LLM) inference due to its efficiency in maximizing compute and memory utilization....

<https://arxiv.org/abs/2606.00516>

2 CONF

**AgentRedBench: Dynamic Redteaming and Integration-Aware Defense for LLM Agents over SaaS Integrations**

ArXiv CS.AI 50%

arXiv:2606.02240v1 Announce Type: cross Abstract: Indirect prompt injection in tool-use agents is a concrete production threat: LLM agents read from integrations (third-party services such as Gmail, Salesforce, or Jira accessed through tool calls) whose response content the user...

<https://arxiv.org/abs/2606.02240>

3 CONF

**Probe Before You Edit: Probing-Guided Molecular Optimization for LLM Agents in Structure-Based Drug Design**

ArXiv CS.AI 50%

arXiv:2606.00555v1 Announce Type: new Abstract: Structure-based drug design increasingly employs LLM agents to iteratively refine ligands against a target pocket, yet a viable ligand must satisfy two often-conflicting objectives -- binding affinity and druggability -- which...

<https://arxiv.org/abs/2606.00555>

4 CONF

**Fair Finetuning Mitigates Distribution Inference Attacks**

ArXiv CS.AI 50%

arXiv:2606.01719v1 Announce Type: cross Abstract: Machine learning models trained on sensitive data can inadvertently leak population-level information about their training distributions -- a threat known as distribution inference attack (DIA). An adversary with black-box access...

<https://arxiv.org/abs/2606.01719>

5 CONF

**Beyond One-shot: AI Agents for Learning in Field Experiments**

ArXiv CS.AI 50%

arXiv:2606.02458v1 Announce Type: new Abstract: Organizations routinely run experiments for A/B testing, yet the data generated from one experiment is underutilized to inform subsequent intervention design. Significant barriers exist to extracting actionable knowledge from prior...

<https://arxiv.org/abs/2606.02458>

6 CONF

**MOSAIC: Modular Orchestration for Structured Agentic Intelligence and Composition**

ArXiv CS.AI 50%

arXiv:2606.00708v1 Announce Type: new Abstract: Automated data science is a structured model-selection problem. A solution must choose data transformations, feature representations, architecture, training procedure, evaluation protocol, and refinement strategy for a task. AutoML...

<https://arxiv.org/abs/2606.00708>

7 CONF

**LLM-Driven Co-Evolutionary Automated Heuristic Design for Bi-Component Coupled Combinatorial Optimization**

ArXiv CS.AI 50%

arXiv:2606.00718v1 Announce Type: new Abstract: While Large Language Models (LLMs) have recently shown promise in Automated Heuristic Design (AHD), existing methods typically generate and evolve heuristics as a single operator or search strategy, limiting their ability to model...

<https://arxiv.org/abs/2606.00718>

8 CONF

**Plausibility Is Not Prediction: Contrastive Evidence for LLM-Based Cellular Perturbation Reasoning**

ArXiv CS.AI 50%

arXiv:2606.01042v1 Announce Type: cross Abstract: Perturbation experiments are central to understanding cellular mechanisms, but remain costly and sparse, motivating prediction of gene expression responses for unobserved conditions. A promising recent direction leverages large...

<https://arxiv.org/abs/2606.01042>



9 CONF

**AI Sovereignty as National Learning Capacity: A Human-Centered Learning Mechanics Viewpoint on France, the United States, and China**

ArXiv CS.AI 50%

arXiv:2606.00729v1 Announce Type: new Abstract: Artificial Intelligence is often discussed in France in terms of investment, compute capacity, regulation, employment, sovereignty, and education. These dimensions are usually treated separately. This viewpoint paper proposes a...

<https://arxiv.org/abs/2606.00729>

0 CONF

**COMAP: Co-Evolving World Models and Agent Policies for LLM Agents**

ArXiv CS.AI 50%

arXiv:2606.02372v1 Announce Type: new Abstract: Equipping language agents with world models enables them to anticipate environment dynamics and evaluate candidate actions before execution. However, existing textual world models are typically fixed after training, preventing them...

<https://arxiv.org/abs/2606.02372>

1 CONF

**FALAT: Tracing Failures in LLM Agent Trajectories via Dependency-Guided Search**

ArXiv CS.AI 50%

arXiv:2606.00765v1 Announce Type: new Abstract: LLM-based agents increasingly solve complex tasks through long trajectories involving reasoning steps, tool calls, and inter-agent communication. However, when these agents fail, it is often unclear which agent caused the failure...

<https://arxiv.org/abs/2606.00765>

2 CONF

**Interaction-Centered Intelligence: Toward Interaction as the Primary Unit of Analysis in Co-Creative AI and Human-AI Systems**

ArXiv CS.AI 50%

arXiv:2606.00807v1 Announce Type: new Abstract: Traditional artificial intelligence has largely conceptualized intelligence as isolated computation occurring within bounded agents. Across classical AI, machine learning, and many generative systems, the dominant unit of analysis...

<https://arxiv.org/abs/2606.00807>

3 CONF

**Make a Video Call with LLM: A Measurement Campaign over Six Mainstream Apps**

ArXiv CS.AI 50%

arXiv:2510.00481v2 Announce Type: replace-cross Abstract: In 2025, Large Language Model (LLM) services have launched a new feature -- AI video chat -- allowing users to interact with AI agents via real-time video communication (RTC), just like chatting with real people. Despite...

<https://arxiv.org/abs/2510.00481>

4 CONF

**Prospect-Theory Behavior from Bellman Optimality in MDPs with Catastrophic States**

ArXiv CS.AI 50%

arXiv:2606.00970v1 Announce Type: new Abstract: We study risk-neutral control in Markov decision processes with an absorbing catastrophic state. Even though rewards are linear and the agent has no utility curvature, probability weighting, or framing dependence, standard Bellman...

<https://arxiv.org/abs/2606.00970>

5 CONF

**Large Language Models in Transportation Systems Management and Operations: From Text Reasoning to Multi-modal Decision Support**

ArXiv CS.AI 50%

arXiv:2606.00991v1 Announce Type: new Abstract: Transportation systems management and operations (TSMO) increasingly depends on timely interpretation of heterogeneous data, from various sensor streams, incident reports, traveler feedback, and visual observations. Large language...

<https://arxiv.org/abs/2606.00991>

6 CONF

**Constitutional Black-Box Monitoring for Scheming in LLM Agents**

ArXiv CS.AI 50%

arXiv:2603.00829v2 Announce Type: replace-cross Abstract: Safe deployment of Large Language Model (LLM) agents in autonomous settings requires reliable oversight mechanisms. A central challenge is detecting scheming, where agents covertly pursue misaligned goals. One approach to...

<https://arxiv.org/abs/2603.00829>

7 CONF

**Property Prediction of Stacked Bilayer Materials: A Multimodal Learning Approach**

ArXiv CS.AI 50%

arXiv:2606.01012v1 Announce Type: new Abstract: AI for materials science is a critical topic within AI for science, aiming to accelerate materials discovery and produce accurate property predictions. Bilayer 2D material stacking is essential for exploring new materials with...

<https://arxiv.org/abs/2606.01012>

8 CONF

**Can AI Review Improve Paper Drafting? An Empirical Study on 20 Computer Architecture Submissions**

ArXiv CS.AI 50%

arXiv:2606.01013v1 Announce Type: new Abstract: Research is advancing faster than ever with artificial intelligence (AI); and so are the corresponding research papers. The exploding volume of AI-generated papers have put a strain to peer review, leading to the usage of...

<https://arxiv.org/abs/2606.01013>

9 CONF

**RoboBenchMart: Benchmarking Robots in Retail Environment**

ArXiv CS.AI 50%

arXiv:2511.10276v2 Announce Type: replace-cross Abstract: Most existing robotic manipulation benchmarks focus on tabletop or household scenarios. While these setups have driven impressive progress, it remains unclear whether generalist VLAs that excel there can truly generalize...

<https://arxiv.org/abs/2511.10276>

10 CONF

**Diagnosing LLM Arbitration Behavior over Pre-evidence Epistemic States in RAG-based Fact-Checking**

ArXiv CS.AI 50%

arXiv:2606.01120v1 Announce Type: new Abstract: In RAG-based fact-checking, LLMs are increasingly used as verifiers to check given claims against retrieved evidence. Their parametric knowledge can induce pre-evidence tendencies that may conflict with the retrieved context, yet...

<https://arxiv.org/abs/2606.01120>

1 CONF

**SkillRevise: Improving LLM-Authored Agent Skills via Trace-Conditioned Skill Revision**

ArXiv CS.AI 50%

arXiv:2606.01139v1 Announce Type: new Abstract: Agent skills are procedural artifacts that enable LLM agents to execute workflows, verify constraints, and recover from failures. Existing self-evolving methods refine skills using accumulated trajectories. However, they struggle...

<https://arxiv.org/abs/2606.01139>

2 CONF

**Expected Value Alignment for Generative Reward Modeling in Formal Mathematics Verification**

ArXiv CS.AI 50%

arXiv:2606.01160v1 Announce Type: new Abstract: Large Language Models (LLMs) are increasingly used with formal interactive theorem provers such as Lean 4. Scaling these systems with reinforcement learning or search methods requires process reward models (PRMs) that can evaluate...

<https://arxiv.org/abs/2606.01160>

3 CONF

**Can LLM Agents Sustain Long-Horizon Organizational Dynamics?**

ArXiv CS.AI 50%

arXiv:2606.01199v1 Announce Type: new Abstract: Large language agents are increasingly used for social simulation, yet it remains unclear whether they can sustain coherent behavior in structured organizations, where goals must propagate through hierarchy, tasks depend on prior...

<https://arxiv.org/abs/2606.01199>

4 CONF

**Policy and World Modeling Co-Training for Language Agents**

ArXiv CS.AI 50%

arXiv:2606.02388v1 Announce Type: cross Abstract: Reinforcement learning (RL) improves large language model (LLM) agents by teaching them which actions lead to high rewards, but provides little supervision on what those actions do to the environment. World modeling (WM) can fill...

<https://arxiv.org/abs/2606.02388>

5 CONF

**Advanced Mathematics Learning Behavior Prediction and Academic Early Warning Model Based on Multimodal Data Analysis**

ArXiv CS.AI 50%

arXiv:2606.01224v1 Announce Type: new Abstract: Early detection of at-risk students and timely academic intervention pose major challenges in advanced mathematics education, where complex conceptual hierarchies and nonlinear learning trajectories often hold back students'...

<https://arxiv.org/abs/2606.01224>

6 CONF

**MuFeRL: Enhancing Reinforcement Learning with Verbal Feedback in a Multi-turn Loop**

ArXiv CS.AI 50%

arXiv:2601.22900v2 Announce Type: replace Abstract: Reinforcement Learning with Verifiable Rewards (RLVR) is widely used to improve reasoning across domains, but outcome-only scalar rewards are often sparse and uninformative. This limitation is especially severe for failed...

<https://arxiv.org/abs/2601.22900>

## 7 **CONF** **Brain-Atlas-Guided Generative Counterfactual Attention for Explainable Cognitive Decline Diagnosis Using Multimodal Connectomes**

ArXiv CS.AI 50%

arXiv:2606.01237v1 Announce Type: new Abstract: Mild cognitive impairment (MCI) and subjective cognitive decline (SCD) are closely associated with the early Alzheimer's disease continuum, where accurate and explainable diagnosis is important for early risk assessment and...

<https://arxiv.org/abs/2606.01237>

## 8 **CONF** **Emergent Ordinal Geometry in Transformers Trained on Local Comparisons**

ArXiv CS.AI 50%

arXiv:2606.01269v1 Announce Type: new Abstract: Transitive inference is the challenge of inferring that  $A < C$  from knowing only adjacent relations ( $A < B$ ,  $B < C$ ). It is solved by humans and animals not through logical chaining but via an analogue mental number line, whose...

<https://arxiv.org/abs/2606.01269>

## 9 **CONF** **DataShield: Safety-degrading Data Filtering for LLM Benign Instruction Fine-Tuning**

ArXiv CS.AI 50%

arXiv:2606.00160v1 Announce Type: cross Abstract: Large language models (LLMs) suffer from degraded safety capabilities even when fine-tuned with benign datasets. However, existing methods for identifying safety-degrading samples in benign datasets suffer from high computational...

<https://arxiv.org/abs/2606.00160>

## 0 **CONF** **DiffuSent: Towards a Unified Diffusion Framework for Aspect-Based Sentiment Analysis**

ArXiv CS.AI 50%

arXiv:2606.01323v1 Announce Type: cross Abstract: Aspect-Based Sentiment Analysis (ABSA) encompasses seven distinct subtasks, each focusing on different extracted elements. Despite the proven success of generative models in unified aspect sentiment analysis, existing approaches...

<https://arxiv.org/abs/2606.01323>

## 1 **CONF** **Recognize Your Orchestrator: An Entropy Dynamics Perspective for LLM Multi-Agent Systems**

ArXiv CS.AI 50%

arXiv:2606.01351v1 Announce Type: new Abstract: The transition from single-turn models to Multi-Agent Systems (MAS) promises enhanced problem-solving capabilities, yet the centralized orchestration topology remains a critical point of fragility. To analyze this, we propose a...

<https://arxiv.org/abs/2606.01351>

## 2 **CONF** **FlowTime: Towards Continuous Generative Watch Time Prediction via Flow-based Personalized Priors**

ArXiv CS.AI 50%

arXiv:2606.01352v1 Announce Type: new Abstract: Watch time has emerged as a pivotal metric for optimizing deep user engagement in short-video recommender systems. However, current methods of watch time prediction (WTP) suffer from inherent paradigm-specific limitations. Direct...

<https://arxiv.org/abs/2606.01352>

3 **CONF****How Generation Architecture Shapes Code Complexity in Multi-Agent LLM Systems: A Paired Study on HumanEval**

ArXiv CS.AI 50%

arXiv:2606.00308v1 Announce Type: cross Abstract: Large-language-model code generation has shifted from single-shot prompting to multi-agent orchestrations - analyst, coder, tester, and debugger pipelines - and is evaluated almost exclusively on functional correctness. Whether...

<https://arxiv.org/abs/2606.00308>
4 **CONF****Herculean: An Agentic Benchmark for Financial Intelligence**

ArXiv CS.AI 50%

arXiv:2605.14355v3 Announce Type: replace Abstract: As AI agents improve, the central question is no longer whether they can solve isolated well-defined financial tasks, but whether they can reliably carry out financial professional work. Existing financial benchmarks offer only...

<https://arxiv.org/abs/2605.14355>
5 **CONF****GovAI-Pipe: A Layered AI Governance Pipeline for Citizen-Facing AI in Turkey's e-Government Gateway**

ArXiv CS.AI 50%

arXiv:2606.01417v1 Announce Type: new Abstract: Turkey's e-Government Gateway (e-Devlet) serves over 68 million registered users with more than 9,200 government services, and is increasingly integrating artificial intelligence into citizen-facing applications such as chatbot...

<https://arxiv.org/abs/2606.01417>
6 **CONF****Don't Ask the LLM to Track Freshness: A Deterministic Recipe for Memory Conflict Resolution**

ArXiv CS.AI 50%

arXiv:2606.01435v1 Announce Type: new Abstract: LLM-based memory systems increasingly maintain facts that evolve over time, where a recurring failure is conflict resolution: when a fact has multiple contradictory values, which should the agent return? MemoryAgentBench (MAB; Hu...

<https://arxiv.org/abs/2606.01435>
7 **CONF****Dive into Ambiguity: A\*-Inspired Multi-Agents Commonsense Obfuscation Attack on LLM Prompts**

ArXiv CS.AI 50%

arXiv:2606.01441v1 Announce Type: new Abstract: Large language models (LLMs) excel in reasoning and knowledge-intensive tasks but remain vulnerable to prompt-level adversarial attacks that preserve intent while triggering commonsense hallucinations. This vulnerability is urgent,...

<https://arxiv.org/abs/2606.01441>
8 **CONF****A Minimalist Brain-Computer Musical Interface for Real-Time Emotion-Driven Sonification: System Design and Preliminary Evaluation**

ArXiv CS.AI 50%

arXiv:2606.01473v1 Announce Type: new Abstract: This paper presents a minimalist brain-computer Musical Interface (BCMI) that functions as a real-time affective sonification system, translating prefrontal EEG activity into adaptive music. Emotional valence is estimated from...

<https://arxiv.org/abs/2606.01473>

9 CONF

**SpikeWFM: Spiking-Aided Wireless Foundation Model for Robust Channel Prediction**

ArXiv CS.AI 50%

arXiv:2606.00120v1 Announce Type: cross Abstract: This paper proposes SpikeWFM, a novel hybrid architecture that integrates spiking neural networks (SNNs) with conventional artificial neural network (ANN)-based transformers for wireless foundation models (WFMs). Inspired by the...

<https://arxiv.org/abs/2606.00120>

0 CONF

**Understanding LLM Behavior in Multi-Target Cross-Lingual Summarization**

ArXiv CS.AI 50%

arXiv:2606.01252v1 Announce Type: cross Abstract: Multi-target cross-lingual text summarization (MTXLS), which summarizes a source document into multiple target languages, is increasingly important as users consume content in diverse languages, but remains underexplored. To...

<https://arxiv.org/abs/2606.01252>

1 CONF

**EvoBrain: Continual Learning of EEG Foundation Models Across Heterogeneous BCI Tasks**

ArXiv CS.AI 50%

arXiv:2606.01767v1 Announce Type: new Abstract: Electroencephalography (EEG) is the cornerstone of non-invasive brain-computer interfaces (BCIs), yet conventional decoding relies on fragmented, task-specific architectures that severely limit cross-task scalability. While EEG...

<https://arxiv.org/abs/2606.01767>

2 CONF

**On the Evaluation of Spiking Neural Network Configurations for Network Intrusion Detection**

ArXiv CS.AI 50%

arXiv:2606.01442v1 Announce Type: cross Abstract: Network intrusion detection is a core component of modern cybersecurity infrastructure, yet the deep learning models that dominate the field are computationally demanding, motivating interest in lightweight alternatives suited to...

<https://arxiv.org/abs/2606.01442>

3 CONF

**Does Compression Preserve Uncertainty? A Unified Benchmark for Quantized and Sparse LLMs via Conformal Prediction**

ArXiv CS.AI 50%

arXiv:2606.01850v1 Announce Type: new Abstract: Model compression techniques such as quantization and pruning are widely used to reduce the deployment cost of large language models (LLMs), with existing evaluations focusing almost exclusively on accuracy preservation. However,...

<https://arxiv.org/abs/2606.01850>

4 CONF

**3DCodeBench: Benchmarking Agentic Procedural 3D Modeling Via Code**

ArXiv CS.AI 50%

arXiv:2606.01057v1 Announce Type: cross Abstract: Procedural 3D modeling through code is emerging as a versatile paradigm, offering deterministic, engine-ready, and precisely editable assets that neural 3D generators inherently lack. Authoring such procedural content, however,...

<https://arxiv.org/abs/2606.01057>

5 **CONF****Physically-Constrained Mamba-SDE for Remaining Useful Life Prediction under Irregular Observations****ArXiv CS.AI** **50%**

arXiv:2606.01894v1 Announce Type: new Abstract: Accurate Remaining Useful Life prediction is critical for industrial predictive maintenance. However, real-world deployment is challenging due to the irregular nature of sensor observations, characterized by asynchronous sampling,...

<https://arxiv.org/abs/2606.01894>

6 **CONF****VET: A Framework for Analyzing AI Discourse****ArXiv CS.AI** **50%**

arXiv:2606.01929v1 Announce Type: new Abstract: Public discourse on AI has become polarized; exaggerated positions on AI in traditional and social media threaten the development of AI Literacy among the general public. In this article, I introduce the VET Framework, a method for...

<https://arxiv.org/abs/2606.01929>

7 **CONF****Domain-Shift-Aware Conformal Prediction for Large Language Models****ArXiv CS.AI** **50%**

arXiv:2510.05566v2 Announce Type: replace-cross Abstract: Large language models have achieved impressive performance across diverse tasks. However, their tendency to produce overconfident and factually incorrect outputs, known as hallucinations, poses risks in real-world...

<https://arxiv.org/abs/2510.05566>

8 **CONF****Algorithmic algorithm development with LLMs: A Case Study on LLM-Usage for Contraction Order Optimization in Tensor Networks****ArXiv CS.AI** **50%**

arXiv:2606.01975v1 Announce Type: new Abstract: We consider LLM-based algorithm development through a case study on contractionorder optimisation for tensor networks with OpenEvolve. We pay particular attention to the choice of the LLM as well as design choices such as...

<https://arxiv.org/abs/2606.01975>

9 **CONF****An NLP-Driven Framework for Curriculum-Labor Market Alignment: Schema-Constrained LLM Extraction, ESCO-Anchored Semantic Matching, and Multi-Dimensional Gap Quantification****ArXiv CS.AI** **50%**

arXiv:2606.01982v1 Announce Type: new Abstract: Schema-constrained information extraction from diverse educational and labor-market corpora remains an open challenge in natural language processing because existing pipelines rely primarily on lexical-surface methods that cannot...

<https://arxiv.org/abs/2606.01982>

0 **CONF****A Lightweight Deep Learning-based Model for Ranking Influential Nodes in Complex Networks****ArXiv CS.AI** **50%**

arXiv:2507.19702v1 Announce Type: cross Abstract: Identifying influential nodes in complex networks is a critical task with a wide range of applications across different domains. However, existing approaches often face trade-offs between accuracy and computational efficiency. To...

<https://arxiv.org/abs/2507.19702>

1 **CONF****Privacy Policy Enforcement Guardrails for Data-Sensitive Retrieval-Augmented Generation****ArXiv CS.AI** **50%**

arXiv:2605.17034v2 Announce Type: replace-cross Abstract: Standard PII filters often miss contextual data leakage in RAG systems, such as non-regulated attribute clusters that collectively identify individuals. We introduce a Privacy Policy Enforcement (PPE) framework using dual...

<https://arxiv.org/abs/2605.17034>

2 **CONF****S3TS: Stochastic Scenario-Structured Tree Search for Advanced Planning Under Uncertainty****ArXiv CS.AI** **50%**

arXiv:2606.02151v1 Announce Type: new Abstract: Effective scheduling in the energy sector is essential to ensure the reliable operation of electrical grids and their connected assets by, for instance, optimizing the dispatch of generation units and storage systems. An effective...

<https://arxiv.org/abs/2606.02151>

3 **CONF****Multi-Objective Reinforcement Learning for Tactical Decision Making for Trucks in Highway Traffic****ArXiv CS.AI** **50%**

arXiv:2601.18783v2 Announce Type: replace-cross Abstract: Balancing safety, efficiency, and operational costs in highway driving poses a challenging decision-making problem for heavy-duty vehicles. A central difficulty is that conventional scalar reward formulations, obtained by...

<https://arxiv.org/abs/2601.18783>

4 **CONF****Spatial Representation Learning Beyond Pixels: Unifying Raster Data and Vector Semantics for Human-Centric Geospatial Foundation Models****ArXiv CS.AI** **50%**

arXiv:2606.02374v1 Announce Type: new Abstract: Earth Observation (EO) has fundamentally transformed the monitoring of environmental processes and human activities up to planetary scale. Recent advances in self-supervised learning have given rise to Earth Observation Foundation...

<https://arxiv.org/abs/2606.02374>

5 **CONF****LLMs Need Encoders for Semantic IDs Too****ArXiv CS.AI** **50%**

arXiv:2606.00324v1 Announce Type: cross Abstract: Multimodal LLMs use dedicated encoders to bridge non-language modalities (vision encoders for images, depth models for audio codec tokens) because raw token embeddings alone cannot capture modality-specific structure. We argue...

<https://arxiv.org/abs/2606.00324>

6 **CONF****Food Noise & False Safety: A Systematic Evaluation of How LLMs Fail to Adapt to Eating Disorder Queries with Clinician Feedback****ArXiv CS.AI** **50%**

arXiv:2606.02444v1 Announce Type: new Abstract: Recent evidence shows that people with eating disorders (EDs) are increasingly seeking guidance, advice, and emotional support from Large Language Model (LLM)-based chat systems. Although these systems are not designed to provide...

<https://arxiv.org/abs/2606.02444>



7 CONF

**Med-Scout: Curing MLLMs' Geometric Blindness in Medical Perception via Geometry-Aware RL Post-Training**

ArXiv CS.AI 50%

arXiv:2601.23220v2 Announce Type: replace-cross Abstract: Despite recent Multimodal Large Language Models (MLLMs)' linguistic prowess in medical diagnosis, we find even state-of-the-art MLLMs suffer from a critical perceptual deficit: geometric blindness. This failure to ground...

<https://arxiv.org/abs/2601.23220>

8 CONF

**SafeSteer: Localized On-Policy Distillation for Efficient Safety Alignment**

ArXiv CS.AI 50%

arXiv:2606.02530v1 Announce Type: new Abstract: Aligning Large Language Models (LLMs) with human values often degrades their general capabilities, termed the alignment tax. Existing methods mitigate this by balancing dual objectives, which heavily rely on massive general-purpose...

<https://arxiv.org/abs/2606.02530>

9 CONF

**A Novel Data Augmentation Strategy for Robust Deep Learning Classification of Biomedical Time-Series Data: Application to ECG and EEG Analysis**

ArXiv CS.AI 50%

arXiv:2507.12645v1 Announce Type: cross Abstract: The increasing need for accurate and unified analysis of diverse biological signals, such as ECG and EEG, is paramount for comprehensive patient assessment, especially in synchronous monitoring. Despite advances in multi-sensor...

<https://arxiv.org/abs/2507.12645>

0 CONF

**Business Utility of Large Language Models as Exploratory Data Analysis Agents**

ArXiv CS.AI 50%

arXiv:2606.00051v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used in analytical workflows, but their suitability as exploratory data analysis (EDA) agents in business settings remains uncertain. In practice, a deployable EDA agent must provide...

<https://arxiv.org/abs/2606.00051>

1 CONF

**Empathic and agentic artificial intelligence in nursing: perspectives on a human-centered framework for cancer care navigation in the United States**

ArXiv CS.AI 50%

arXiv:2606.00010v1 Announce Type: cross Abstract: For patients experiencing cancer, nurse navigation can ease the burden of complex care by enhancing coordination of health services and patient outcomes. However, in under-resourced areas, trained nurse navigators may be limited...

<https://arxiv.org/abs/2606.00010>

2 CONF

**RuleEdit: Failure-Guided Human-AI Model Editing with Prospective Impact Preview**

ArXiv CS.AI 50%

arXiv:2606.00011v1 Announce Type: cross Abstract: Despite the promise of AI to assist complex decisions, practitioners still lack ways to detect likely failures and inspect the consequences of model edits before committing them. We present RuleEdit, an interactive, rule-guided...

<https://arxiv.org/abs/2606.00011>

3 **CONF****DenseMLLM: Standard Multimodal LLMs for Dense Prediction**

ArXiv CS.AI 50%

arXiv:2602.14134v2 Announce Type: replace-cross Abstract: Multimodal Large Language Models (MLLMs) have demonstrated exceptional capabilities in high-level visual understanding. However, extending these models to fine-grained dense prediction tasks, such as semantic segmentation...

<https://arxiv.org/abs/2602.14134>

4 **CONF****A phenomenon of AI-conformity: how algorithms change human moral decision-making**

ArXiv CS.AI 50%

arXiv:2606.00013v1 Announce Type: cross Abstract: Social conformity is a well-documented phenomenon in which individuals shift their opinions towards those of a social majority. As artificial intelligence (AI) becomes increasingly integrated into everyday life it may also create...

<https://arxiv.org/abs/2606.00013>

5 **CONF****Toward Robust In-Context Learning: Leveraging Out-of-distribution Proxies for Target Inaccessible Demonstration Retrieval**

ArXiv CS.AI 50%

arXiv:2606.00014v1 Announce Type: cross Abstract: Although studies have demonstrated that Large Language Models (LLMs) can perform well on Out-of-Distribution (OOD) tasks, their advantage tends to diminish as the distribution shift becomes more severe. Consequently, researchers...

<https://arxiv.org/abs/2606.00014>

6 **CONF****Consistency Training while Mitigating Obfuscation via Rate Matching**

ArXiv CS.AI 50%

arXiv:2606.02211v1 Announce Type: cross Abstract: Large language models are often influenced by extraneous input features, such as cues revealing a user's preferred answer. Consistency training reduces this influence by training models to behave similarly across inputs with and...

<https://arxiv.org/abs/2606.02211>

7 **CONF****Understanding Stigmatizing Language in Clinical Documentation: A Paired Comparison of Ambient AI Drafts and Clinician Finalized Notes**

ArXiv CS.AI 50%

arXiv:2606.00019v1 Announce Type: cross Abstract: Ambient artificial intelligence (AI) documentation tools are increasingly deployed to reduce clinician documentation burden, but their implications for biased language in clinical notes remain unclear. We conducted a large-scale...

<https://arxiv.org/abs/2606.00019>

8 **CONF****CSRP: Chain-of-Thought Reasoning for Chinese Text Correction via Reinforcement Learning with Efficiency-Aware Rewards**

ArXiv CS.AI 50%

arXiv:2606.00020v1 Announce Type: cross Abstract: Large Language Model (LLM) based Chinese Grammatical Error Correction (CGEC) systems face two critical challenges: general-purpose models lack specialized linguistic priors for subtle grammatical distinctions, and Supervised...

<https://arxiv.org/abs/2606.00020>

9 CONF

**SENSE: Semantic Embedding Navigation with Soft-gated Evaluation for Retrieval-based Speculative Decoding**

ArXiv CS.AI 50%

arXiv:2606.00021v1 Announce Type: cross Abstract: Speculative Decoding (SD) accelerates Large Language Model (LLM) inference by employing a lightweight draft model to propose candidate tokens, which are verified in parallel by the target model, without compromising generation...

<https://arxiv.org/abs/2606.00021>

00 CONF

**Imfaoooo at SemEval-2026 Task 1: Humor Is an Audience. Preference Modeling for Constrained Humor Generation**

ArXiv CS.AI 50%

arXiv:2606.00022v1 Announce Type: cross Abstract: Humor generation remains difficult not only because producing fluent, novel jokes is hard, but because "funny" is audience-dependent and supervision is noisy -- preferences vary with audience, context, and culture, and annotator...

<https://arxiv.org/abs/2606.00022>

01 CONF

**Update Opacity: Epistemic Accessibility and Governance Under AI System Change**

ArXiv CS.AI 50%

arXiv:2606.00037v1 Announce Type: cross Abstract: Machine learning models embedded in deployed AI systems are routinely updated to maintain correct functioning over time. Yet such updates can generate update opacity: users may not be able to understand why the same input now...

<https://arxiv.org/abs/2606.00037>

02 CONF

**Beyond Tool Adoption: A Practical Five-Stage Developmental Continuum for AI Literacy in Higher Education**

ArXiv CS.AI 50%

arXiv:2606.00038v1 Announce Type: cross Abstract: Artificial intelligence (AI) literacy is increasingly recognized as a foundational competency for all university graduates. Yet students' engagement with AI tools often clusters at two problematic extremes: avoidance driven by...

<https://arxiv.org/abs/2606.00038>

03 CONF

**Beyond Categories of Caste: Examining Caste Bias and Morality in Text-to-Image AI Models**

ArXiv CS.AI 50%

arXiv:2606.00039v1 Announce Type: cross Abstract: Text-to-Image (T2I) models have shown promising utility across various domains. However, such models are also amplifying harmful societal biases in their outputs. In the context of South Asia, recent work has shown caste biases...

<https://arxiv.org/abs/2606.00039>

04 CONF

**Tracing GenAI Literacy: Uncovering Student-AI Interaction Patterns in Academic Writing through Epistemic Network Analysis**

ArXiv CS.AI 50%

arXiv:2606.00040v1 Announce Type: cross Abstract: As Generative AI (GenAI) becomes integral to education, fostering GenAI literacy is critical. However, current assessments largely rely on self-reported scales, lacking insights into how literacy manifests in actual learning...

<https://arxiv.org/abs/2606.00040>

05 CONF

**Lookahead Sample Reward Guidance for Test-Time Scaling of Diffusion Models**

ArXiv CS.AI 50%

arXiv:2602.03211v2 Announce Type: replace-cross Abstract: Diffusion models have demonstrated strong generative performance; however, generated samples often fail to fully align with human intent. This paper studies an efficient test-time scaling method for sampling from regions...

<https://arxiv.org/abs/2602.03211>

06 CONF

**Physics-Informed Deep Learning for Entropy Prediction in Heterogeneous Systems: Thermodynamic and Information-Theoretic Case Studies**

ArXiv CS.AI 50%

arXiv:2606.01179v1 Announce Type: cross Abstract: Entropy production governs irreversibility and uncertainty in both physical and information-theoretic systems. While Physics-Informed Neural Networks (PINNs) successfully solve differential equations, current architectures remain...

<https://arxiv.org/abs/2606.01179>

07 CONF

**Beyond Text and Tables: Vision-Language Model Integration in ComProScanner for Extracting Materials Data from Scientific Figures with High Accuracy**

ArXiv CS.AI 50%

arXiv:2606.00065v1 Announce Type: cross Abstract: Automated extraction of materials composition-property data from scientific literature has advanced considerably with the development of large language model-based pipelines; however, existing frameworks remain limited to textual...

<https://arxiv.org/abs/2606.00065>

08 CONF

**Flow-Based Generative Modeling for Optimizing Sampling Policies in Compressed Sensing Applications**

ArXiv CS.AI 50%

arXiv:2606.00078v1 Announce Type: cross Abstract: Numerous modern applications in signal processing and medical imaging necessitate acquiring high-dimensional signals under tight resource constraints. Traditional sampling theory suggests that accurate signal reconstruction...

<https://arxiv.org/abs/2606.00078>

09 CONF

**BitsMoE: Efficient Spectral Energy-Guided Bit Allocation for MoE LLM Quantization**

ArXiv CS.AI 50%

arXiv:2606.00079v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) large language models reduce per-token computation through sparse expert activation, but their deployment remains memory-intensive because all expert weights must be kept resident in memory. Existing MoE...

<https://arxiv.org/abs/2606.00079>

10 CONF

**Pause and Think: A Dataset and Benchmark for Video-Grounded Assistive Action Suggestion**

ArXiv CS.AI 50%

arXiv:2606.00616v1 Announce Type: cross Abstract: Recent Vision-Language Models (VLMs) struggle with grounded reasoning, temporal consistency, and context aware planning in videos. We introduce pause-and-think-T, a reasoning-centric training dataset that encourages models to...

<https://arxiv.org/abs/2606.00616>

11 CONF

**DASStatFormer: A Hybrid Multibranch Transformer with Statistical Feature Integration for DAS-Based Pattern Recognitions**

ArXiv CS.AI 50%

arXiv:2606.00081v1 Announce Type: cross Abstract: Distributed Acoustic Sensing (DAS) enables large-scale monitoring through optical fibers, but its high dimensionality and complex spatio-temporal patterns make event classification demanding. Existing deep learning...

<https://arxiv.org/abs/2606.00081>

12 CONF

**Structured Visual Evidence Decomposition for Evidence-Grounded Multimodal Screening of Obstructive Sleep Apnea-Hypopnea Syndrome**

ArXiv CS.AI 50%

arXiv:2606.00087v1 Announce Type: cross Abstract: Effective pre-polysomnography screening for obstructive sleep apnea-hypopnea syndrome (OSAHS) requires combining clinical risk factors with visible craniofacial and neck cues. Directly prompting general-purpose multimodal...

<https://arxiv.org/abs/2606.00087>

13 CONF

**SECUREVENT: Hybrid AI/ML Security Monitoring for Distributed Event-Based Systems**

ArXiv CS.AI 50%

arXiv:2606.01741v1 Announce Type: cross Abstract: Distributed event-based systems have become a common substrate for Internet-scale publish/subscribe services, IoT telemetry, cloud-native microservices, and security operations pipelines. Their loose coupling and asynchronous...

<https://arxiv.org/abs/2606.01741>

14 CONF

**Diffusion Image Generation with Explicit Modeling of Data Manifold Geometry**

ArXiv CS.AI 50%

arXiv:2606.00094v1 Announce Type: cross Abstract: Image generative models aim to sample data points from the underlying data manifold, a task that requires learning and decoding a dense, low-dimensional, and compact parameterization space. To achieve this, we propose the Data...

<https://arxiv.org/abs/2606.00094>

15 CONF

**Diversity Over Frequency: Rethinking Tool Use in Visual Chain-of-Thought Agents**

ArXiv CS.AI 50%

arXiv:2606.00096v1 Announce Type: cross Abstract: Visual agents employ external visual tools within visual chains of thought to incorporate fine-grained evidence. While prior work has mainly studied these tools in visual search tasks, their role in more complex visual reasoning...

<https://arxiv.org/abs/2606.00096>

16 CONF

**CoCoVideo: The High-Quality Commercial-Model-Based Contrastive Benchmark for AI-Generated Video Detection**

ArXiv CS.AI 50%

arXiv:2606.00101v1 Announce Type: cross Abstract: With the rapid advancement of artificial intelligence generated content (AIGC) technologies, video forgery has become increasingly prevalent, posing new challenges to public discourse and societal security. Despite remarkable...

<https://arxiv.org/abs/2606.00101>

17 **CONF****A Methodological Framework for Explicit Control of the Speed-Accuracy Trade-off in Brain-Computer Interfaces**

ArXiv CS.AI 50%

arXiv:2606.00106v1 Announce Type: cross Abstract: Brain-computer interfaces (BCIs) are limited by low signal-to-noise ratio in modalities such as electroencephalography, which requires multiple trials to reliably decode user intentions. This induces a speed-accuracy trade-off,...

<https://arxiv.org/abs/2606.00106>

18 **CONF****VDSB-GWSyn: Diffusion Schrödinger Bridge for Controllable and Anatomically Feasible Guidewire Synthesis in Coronary Angiography**

ArXiv CS.AI 50%

arXiv:2606.00109v1 Announce Type: cross Abstract: Coronary guidewire endpoint localization is a fundamental capability for computer-assisted PCI, and its importance increases as robot-assisted PCI is progressively adopted to reduce operator radiation exposure. However, the...

<https://arxiv.org/abs/2606.00109>

19 **CONF****Versatile Framework with Semantic and Structural guidance for Image Reconstruction from Brain Activity**

ArXiv CS.AI 50%

arXiv:2606.00121v1 Announce Type: cross Abstract: Reconstructing visual stimuli from brain recordings has been a meaningful and challenging task in brain decoding. Especially, the achievement of precise and controllable image reconstruction bears great significance in propelling...

<https://arxiv.org/abs/2606.00121>

20 **CONF****CardioLens: Revealing the Clinical Reality Gap of MLLMs via Multi-Sequence Cardiac MRI Evaluations**

ArXiv CS.AI 50%

arXiv:2606.00123v1 Announce Type: cross Abstract: Multimodal Large Language Models (MLLMs) have shown strong performance on public medical benchmarks, yet existing evaluations often remain weak proxies for clinical use, relying on isolated inputs and simplified recognition-style...

<https://arxiv.org/abs/2606.00123>

21 **CONF****Towards a Physics Foundation Model**

ArXiv CS.AI 50%

arXiv:2509.13805v4 Announce Type: replace-cross Abstract: Foundation models have revolutionized natural language processing through a "train once, deploy anywhere" paradigm, where a single pre-trained model adapts to countless downstream tasks without retraining. Access to a...

<https://arxiv.org/abs/2509.13805>

22 **CONF****A Shared Valence Axis Across Modern LLMs and Human EEG: The Saturation Regularity**

ArXiv CS.AI 50%

arXiv:2606.00129v1 Announce Type: cross Abstract: Large language models (LLMs) have emerged as powerful representation learners whose internal features increasingly align with human cognition. We study whether modern LLMs can serve as a lens for understanding neural...

<https://arxiv.org/abs/2606.00129>

23 CONF

**Automatically Differentiable Nonlinear Tensor Networks (ADNTNs) for Exponential Compression of Deep Neural Networks**

ArXiv CS.AI 50%

arXiv:2606.00130v1 Announce Type: cross Abstract: We study Automatically Differentiable Nonlinear Tensor Networks (ADNTNs), a family of structured weight generators whose compact core tensors are trained end-to-end by reverse-mode automatic differentiation (AD). The approach can...

<https://arxiv.org/abs/2606.00130>

24 CONF

**AI-PROPELLER: Warehouse-Scale Interprocedural Code Layout Optimization with AlphaEvolve**

ArXiv CS.AI 50%

arXiv:2606.00131v1 Announce Type: cross Abstract: Post-link optimizers (PLOs) such as Propeller and BOLT have demonstrated that precise, profile-guided code layout can extract significant performance gains from heavily optimized binaries. However, these systems are currently...

<https://arxiv.org/abs/2606.00131>

25 CONF

**On Effectiveness and Efficiency of Agentic Tool-calling and RL Training**

ArXiv CS.AI 50%

arXiv:2606.00135v1 Announce Type: cross Abstract: Tool-calling is a central component of modern large language model (LLM) agents, equipping them with skills beyond their parametric knowledge. This paper studies tool-calling along two complementary axes: effectiveness, i.e., how...

<https://arxiv.org/abs/2606.00135>

26 CONF

**Generative AI and Digital Ecosystem Resilience: A Proactive Lifecycle-Based Survey**

ArXiv CS.AI 50%

arXiv:2606.00136v1 Announce Type: cross Abstract: The proliferation of adversarial synthetic content, accelerated by Generative AI (GenAI) is rendering traditional reactive detection methods ineffective. This survey synthesizes emerging research to demonstrate a paradigm shift...

<https://arxiv.org/abs/2606.00136>

27 CONF

**Geodesics with Unified Tangent-constrained Priors and Curvature Regularization**

ArXiv CS.AI 50%

arXiv:2606.00139v1 Announce Type: cross Abstract: Curvature-penalized geodesic models have proven their effectiveness in image segmentation by computing globally optimal curves. Unfortunately, these models remain susceptible to shortcuts when delineating objects with complex...

<https://arxiv.org/abs/2606.00139>

28 CONF

**StemBind: When MLLMs Get Lost Between Rules and Instances in Abstract Visual Reasoning**

ArXiv CS.AI 50%

arXiv:2606.00148v1 Announce Type: cross Abstract: Multimodal large language models (MLLMs) often know the rule but pick the wrong answer: on abstract visual reasoning (AVR) tasks, a model can describe what it sees and name the underlying pattern, yet still fail to choose the...

<https://arxiv.org/abs/2606.00148>

29 CONF

**Persona Attack: Incremental Memory Injection Jailbreak Attack against Large Language Models**

ArXiv CS.AI 50%

arXiv:2606.00150v1 Announce Type: cross Abstract: As Large Language Models evolve for user convenience, vulnerability to jailbreak attacks continues to be reported despite ongoing efforts in safety training. Traditional jailbreak techniques typically focus on a single prompt...

<https://arxiv.org/abs/2606.00150>

30 CONF

**PrivacyPeek: Auditing What LLM-Based Agents Acquire, Not Just What They Say**

ArXiv CS.AI 50%

arXiv:2606.00152v1 Announce Type: cross Abstract: LLM-based agents are rapidly advancing, autonomously invoking external tools to complete multi-step tasks for users. However, agents often acquire more sensitive information than the task requires. Existing privacy benchmarks...

<https://arxiv.org/abs/2606.00152>

31 CONF

**DiffCrossGait: Trajectory-Level Alignment for 2D-3D Cross-Modal Gait Recognition via Latent Diffusion**

ArXiv CS.AI 50%

arXiv:2606.00153v1 Announce Type: cross Abstract: Cross-modal 2D-3D gait recognition is impeded by inherent domain discrepancies between 2D silhouette and 3D LiDAR range-view representations. While prior methods align only final embeddings, we propose DiffCrossGait, which...

<https://arxiv.org/abs/2606.00153>

32 CONF

**A Protocol-Language Model for Network Intrusion (Without Deep Packet Inspection)**

ArXiv CS.AI 50%

arXiv:2606.00155v1 Announce Type: cross Abstract: Modern network intrusion detection systems (NIDS) are caught in a structural contradiction: the protocols carrying the highest threat intelligence are precisely those encrypted under TLS 1.3 and QUIC, where payload inspection...

<https://arxiv.org/abs/2606.00155>

33 CONF

**InfoAtlas: A Foundation Model for Zero-Shot Statistical Dependence Estimate**

ArXiv CS.AI 50%

arXiv:2606.00241v1 Announce Type: cross Abstract: Measuring statistical dependency between high-dimensional random variables is a fundamental task in data science and machine learning. Neural mutual information (MI) estimators offer a promising avenue, but they typically require...

<https://arxiv.org/abs/2606.00241>

34 CONF

**ChurnNet: A Optimized Modern AI for Churn Prediction**

ArXiv CS.AI 50%

arXiv:2606.00169v1 Announce Type: cross Abstract: Increased competition and the growing similarity of products and services offered by retailers have lowered the barriers for customers to switch to competitors. Accurate churn prediction can be a valuable tool for driving...

<https://arxiv.org/abs/2606.00169>



35 CONF

**The New Social Image: How AI Competency and AI Proactivity Influence Self- and Peer-Perceptions in the Workplace**

ArXiv CS.AI 50%

arXiv:2606.00182v1 Announce Type: cross Abstract: Human-AI collaboration is considered the most promising way to incorporate AI in the workplace. What remains unexplored are the experiential consequences of this teaming. More specifically, in a team with AI, how humans perceive...

<https://arxiv.org/abs/2606.00182>

36 CONF

**BAGEN: Are LLM Agents Budget-Aware?**

ArXiv CS.AI 50%

arXiv:2606.00198v1 Announce Type: cross Abstract: While agents are increasingly spending more resources, today agent cost is mostly measured only after execution. A Budget-Aware Agent (BAGEN) should treat budget as an active control signal, rather than a passive cost metric. We...

<https://arxiv.org/abs/2606.00198>

37 CONF

**Effects of Varying LLM Access on Essay Writing Behavior**

ArXiv CS.AI 50%

arXiv:2606.00250v1 Announce Type: cross Abstract: Investigating the degree to which large language models (LLMs) affect teaching and learning in universities can help identify strategies for integrating LLMs in a way that supports, rather than undermines, student learning...

<https://arxiv.org/abs/2606.00250>

38 CONF

**When Softmax Fails at the Top: Extreme Value Corrections for InfoNCE**

ArXiv CS.AI 50%

arXiv:2606.00262v1 Announce Type: cross Abstract: InfoNCE is the standard contrastive learning objective, but its softmax form is not only a computational convenience: it also encodes a statistical assumption about how the top-scoring example is selected. Using extreme value...

<https://arxiv.org/abs/2606.00262>

39 CONF

**Real2SAM2Real: Generative 3D Caches as Complementary Context for Video Diffusion**

ArXiv CS.AI 50%

arXiv:2606.00299v1 Announce Type: cross Abstract: While Video Diffusion Models (VDMs) excel at synthesizing high-fidelity videos, enabling precise camera and scene control remains challenging. Existing methods predominantly rely on implicit diffusion priors to generate...

<https://arxiv.org/abs/2606.00299>

40 CONF

**Beyond Access: Guided LLM Scaffolding for Independent Learning in Undergraduate Statistics**

ArXiv CS.AI 50%

arXiv:2606.01375v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly entering students' learning practices, but their educational value depends on whether they support reasoning or enable task completion without engagement. This study examines guided...

<https://arxiv.org/abs/2606.01375>

41 CONF

**c-TPE: Tree-structured Parzen Estimator with Inequality Constraints for Expensive Hyperparameter Optimization**

ArXiv CS.AI 50%

arXiv:2211.14411v5 Announce Type: replace-cross Abstract: Hyperparameter optimization (HPO) is crucial for strong performance of deep learning algorithms and real-world applications often impose some constraints, such as on memory usage or latency, on top of the performance...

<https://arxiv.org/abs/2211.14411>

42 CONF

**RadAgent: A tool-using AI agent for stepwise interpretation of chest computed tomography**

ArXiv CS.AI 50%

arXiv:2604.15231v2 Announce Type: replace Abstract: Vision-language models (VLM) have markedly advanced AI-driven interpretation and reporting of complex medical imaging, such as computed tomography (CT). Yet, existing methods largely relegate clinicians to passive observers of...

<https://arxiv.org/abs/2604.15231>

43 CONF

**On the Limits of LLM Adaptability: Impact of Model-Internalized Priors on Annotation Task Performance**

ArXiv CS.AI 50%

arXiv:2606.00467v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used for zero-shot annotation and LLM-as-a-judge tasks, yet their reliability hinges on how model-internalized priors interact with user-provided instructions. We investigate three...

<https://arxiv.org/abs/2606.00467>

44 CONF

**CodeCytos: AI-assisted spatial molecular imaging analysis via code-augmented agent action space**

ArXiv CS.AI 50%

arXiv:2606.00472v1 Announce Type: cross Abstract: Conventional tissue image analysis software provides foundational capabilities for cellular analysis, including segmentation, basic morphological feature extraction, and spatial organization analysis. However, these tools often...

<https://arxiv.org/abs/2606.00472>

45 CONF

**V-LynX: Token Interface Alignment for Video+X LLMs**

ArXiv CS.AI 50%

arXiv:2606.00508v1 Announce Type: cross Abstract: This study introduces an intriguing phenomenon in Video LLMs: rather than merely translating frames into textual embeddings, Video LLMs establish a continuous manifold, token interface, allowing visual tokens to operate as...

<https://arxiv.org/abs/2606.00508>

46 CONF

**Improving Diffusion Planners by Self-Supervised Action Gating with Energies**

ArXiv CS.AI 50%

arXiv:2603.02650v2 Announce Type: replace-cross Abstract: Diffusion planners are a strong approach for offline reinforcement learning, but they can fail when value-guided selection favours trajectories that score well yet are locally inconsistent with the environment dynamics,...

<https://arxiv.org/abs/2603.02650>

47 CONF

**MemGraphRAG: Memory-based Multi-Agent System for Graph Retrieval-Augmented Generation**

ArXiv CS.AI 50%

arXiv:2606.00610v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) has become an essential method for mitigating hallucinations in Large Language Models (LLMs) by leveraging external knowledge. Although effective for simple queries, traditional RAG struggles...

<https://arxiv.org/abs/2606.00610>

48 CONF

**Linguistics-Aware Non-Distortionary LLM Watermarking**

ArXiv CS.AI 50%

arXiv:2606.00613v1 Announce Type: cross Abstract: Watermarking should identify language-model output without degrading quality or limiting verification to the model provider. Multilingual deployment makes this harder because morphology, segmentation, and script change where...

<https://arxiv.org/abs/2606.00613>

49 CONF

**LinguUTics at PsyDefDetect: Iterative Imbalance-Aware Fine-tuning of Qwen3-8B for Psychological Defense Mechanism Classification**

ArXiv CS.AI 50%

arXiv:2606.00647v1 Announce Type: cross Abstract: Detecting psychological defense mechanisms in conversational text remains a challenging clinical NLP problem. For the PsyDefDetect 2026 shared task (nine-class utterance classification evaluated via macro F1), our team...

<https://arxiv.org/abs/2606.00647>

50 CONF

**Scaling Behavior of Single LLM-Driven Multi-Agent Systems**

ArXiv CS.AI 50%

arXiv:2606.00655v1 Announce Type: cross Abstract: The burgeoning field of LLM-based Multi-Agent Systems (MAS) promises to tackle complex tasks through collaborative intelligence, yet fundamental questions regarding their scaling behavior and intrinsic collective dynamics remain...

<https://arxiv.org/abs/2606.00655>

51 CONF

**Demystifying the Optimal Fair Classifier in Multi-Class Classification**

ArXiv CS.AI 50%

arXiv:2606.00656v1 Announce Type: cross Abstract: Ensuring fair and equitable treatment across diverse groups, particularly in multi-class classification tasks, poses a significant challenge due to the persistent biases inherent in machine learning models. Most existing bias...

<https://arxiv.org/abs/2606.00656>

52 CONF

**Beyond the Mouth: Upper-Face Affective Cues in Audiovisual Sentence Recognition under Acoustic Uncertainty**

ArXiv CS.AI 50%

arXiv:2606.00670v1 Announce Type: cross Abstract: Face-to-face speech comprehension is inherently multimodal, integrating acoustic signals with visible articulation, facial expression, head motion, and other socially relevant cues. While audiovisual speech systems typically...

<https://arxiv.org/abs/2606.00670>

53 CONF

**The Paradox of Outcome Optimization: A Causal Information-Theoretic Bound on Reasoning Shortcuts in LLMs**

ArXiv CS.AI 50%

arXiv:2606.00674v1 Announce Type: cross Abstract: Large Language Models (LLMs) aligned via outcome-based Reinforcement Learning (RL) frequently exhibit a critical failure mode: they achieve high performance on in-distribution benchmarks while demonstrating brittle reasoning...

<https://arxiv.org/abs/2606.00674>

54 CONF

**COPF: An Online Framework for Deployment-Stable Counterfactual Fairness in Evolving Graphs**

ArXiv CS.AI 50%

arXiv:2606.00700v1 Announce Type: cross Abstract: Online link recommendation on evolving graphs is performative: by choosing which candidate links to show users, the system changes which links form and what feedback it later observes. Consequently, fairness estimates from logged...

<https://arxiv.org/abs/2606.00700>

55 CONF

**Information-Theoretic Lower Bounds for Bit-Constrained Stochastic Optimization via a Reduction to Compressed Gaussian Mean Estimation**

ArXiv CS.AI 50%

arXiv:2606.00703v1 Announce Type: cross Abstract: Low-precision pretraining (FP8, MXFP4, NVFP4) is now standard for frontier language models, yet the literature is almost entirely achievability -- algorithms and empirical scaling laws -- with no matching characterization of what...

<https://arxiv.org/abs/2606.00703>

56 CONF

**WaveFilter: Enhancing the Long-Context Capability of Diffusion LLMs via Wavelet-Guided KV Cache Filtering**

ArXiv CS.AI 50%

arXiv:2606.00724v1 Announce Type: cross Abstract: Diffusion Large Language Models (DLMs) have demonstrated significant advantages across various tasks. However, constrained by their multi-step iterative inference mechanism, their computational overhead and inference latency in...

<https://arxiv.org/abs/2606.00724>

57 CONF

**Quantum Tunneling-Aware Machine Learning: Physics-Derived Noise Models for Robust Deployment**

ArXiv CS.AI 50%

arXiv:2606.00741v1 Announce Type: cross Abstract: Transistor scaling is approaching a quantum-mechanical limit, as thin gate oxides induce electron leakage through quantum tunneling. Unlike conventional digital systems, AI inference can tolerate such errors provided their...

<https://arxiv.org/abs/2606.00741>

58 CONF

**Bayesian Inference of Nonlinear Malaria Dynamics in Ghana via an Ensemble Markov Chain Monte Carlo Sampler**

ArXiv CS.AI 50%

arXiv:2606.00783v1 Announce Type: cross Abstract: Reliable quantification of malaria dynamics in sub-Saharan Africa is hindered by short, noisy, and spatially heterogeneous surveillance records. In Ghana, health-facility data from 2014 to 2023 reveal non-linear and age-specific...

<https://arxiv.org/abs/2606.00783>

59 CONF

**Certificates without Electrons? Theory and Evidence on Impacts from AI-Driven Power Demand**

ArXiv CS.AI 50%

arXiv:2606.00811v1 Announce Type: cross Abstract: Data centers now account for 4.4% of United States electricity demand, yet the grid-level effectiveness of the renewable energy certificates (RECs) and power purchase agreements (PPAs) hyperscalers use to claim carbon neutrality...

<https://arxiv.org/abs/2606.00811>

60 CONF

**LERD: Latent Event-Relational Dynamics for Neurodegenerative Classification**

ArXiv CS.AI 50%

arXiv:2602.18195v2 Announce Type: replace-cross Abstract: Alzheimer's disease (AD) alters brain electrophysiology and disrupts multichannel EEG dynamics, making accurate and clinically useful EEG-based diagnosis increasingly important for screening and disease monitoring....

<https://arxiv.org/abs/2602.18195>

61 CONF

**GenPT: Beyond Self-Report for Reliable LLM Psychometrics via Generative Projective Testing**

ArXiv CS.AI 50%

arXiv:2606.00860v1 Announce Type: cross Abstract: Self-report questionnaires remain the prevailing tool for probing the psychological states of persona-conditioned agents (PC-Agents). However, classical instruments inherit two well-known threats: contamination from training...

<https://arxiv.org/abs/2606.00860>

62 CONF

**Benchmarks for Vision-Language Models in Urban Perception Should Be Reliability-Aware and Negotiated**

ArXiv CS.AI 50%

arXiv:2606.00871v1 Announce Type: cross Abstract: Vision-language models (VLMs) are increasingly used to generate structured descriptions of street-level imagery for tasks such as streetscape auditing, mapping, and public consultation. These uses combine observable attributes...

<https://arxiv.org/abs/2606.00871>

63 CONF

**Dive into Waves: Morlet Spectral Transformer for Cross-Subject Emotion Decoding from EEG**

ArXiv CS.AI 50%

arXiv:2606.00884v1 Announce Type: cross Abstract: We study cross-subject emotion recognition from EEG, a practically important yet challenging problem in brain-computer interfaces. Unlike tasks with clear waveform signatures, emotion-related EEG signals are primarily encoded in...

<https://arxiv.org/abs/2606.00884>

64 CONF

**Memory-Efficient LLM Training with Dynamic Sparsity: From Stability to Practical Scaling**

ArXiv CS.AI 50%

arXiv:2606.00888v1 Announce Type: cross Abstract: Dynamic Sparse Training (DST) offers a promising paradigm for improving the training and inference efficiency of deep neural networks; however, we find that in large language model training, DST can suffer from optimization...

<https://arxiv.org/abs/2606.00888>

65 CONF

**Accuracy, Stability, and Repeated-Run Reliability of Large Language Models on Deterministic Programming Tasks**

ArXiv CS.AI 50%

arXiv:2606.00920v1 Announce Type: cross Abstract: Run-level pass rate overstates retry-free coverage by up to 17.8 percentage points -- and the gap is largest precisely for mid-performing systems. We investigate this accuracy--stability relationship in large language model (LLM)...

<https://arxiv.org/abs/2606.00920>

66 CONF

**CV-Arena: An Open Benchmark for Instructional Computer Vision Problem Solving with Human-AI Collaborative Preferences**

ArXiv CS.AI 50%

arXiv:2606.00931v1 Announce Type: cross Abstract: Instruction-guided image editing is becoming a general interface for visual work, yet existing benchmarks still focus largely on narrow appearance edits and do not fully capture the diversity of real-image tasks in professional...

<https://arxiv.org/abs/2606.00931>

67 CONF

**Lodestar: An Online-Learning LLM Inference Router**

ArXiv CS.AI 50%

arXiv:2606.00946v1 Announce Type: cross Abstract: Efficiently serving large language model (LLM) inference tasks is crucial both for user-perceived latency such as time-to-first-token (TTFT) and for GPU utilization. However, LLM request routing, that is, assigning each inference...

<https://arxiv.org/abs/2606.00946>

68 CONF

**Network Distributed Multi-Agent Reinforcement Learning for Consensus Control of Quadcopters**

ArXiv CS.AI 50%

arXiv:2606.02107v1 Announce Type: cross Abstract: This paper proposes a Network Distributed Multi-Agent Reinforcement Learning (ND-MARL) framework for quadcopter consensus control. Compared to conventional multi-agent MARL formulations that rely on centralized planning or fully...

<https://arxiv.org/abs/2606.02107>

69 CONF

**Beyond Task-Agnostic: Task-Aware Grouping for Communication-Efficient Multi-Task MoE Inference**

ArXiv CS.AI 50%

arXiv:2606.01007v1 Announce Type: cross Abstract: Sparsely activated Mixture-of-Experts (MoE) models scale capacity via conditional computation, but distributed inference suffers from cross-GPU expert communication and routing-induced load imbalance. Existing placement methods...

<https://arxiv.org/abs/2606.01007>

70 CONF

**AI-IoT-Robotics Integration: Survey of Frameworks, Emerging Trends, and the Path Toward Connected Robotics**

ArXiv CS.AI 50%

arXiv:2606.01015v1 Announce Type: cross Abstract: The convergence of Artificial Intelligence, the Internet of Things, and Robotics is no longer a futuristic vision; it is rapidly becoming the foundation of real-time, intelligent, and context-aware systems. AI enables perception...

<https://arxiv.org/abs/2606.01015>

71 CONF

**PolySpeech-100: A Large-Scale Benchmark for Speech Understanding Across 100+ Languages and Dialects**

ArXiv CS.AI 50%

arXiv:2606.01016v1 Announce Type: cross Abstract: While End-to-End (E2E) Speech-Large Language Models (Speech-LLMs) are rapidly evolving, their evaluation methodologies remain limited to the era of simple transcription. Existing benchmarks suffer from three critical limitations:...

<https://arxiv.org/abs/2606.01016>

72 CONF

**Data Collection for Training Quality-Control AI in Carpet Manufacturing**

ArXiv CS.AI 50%

arXiv:2606.01023v1 Announce Type: cross Abstract: Visual inspection remains the dominant quality-control practice in woven and tufted carpet production, yet it is slow, subjective, and inconsistent at the line speeds and widths of modern looms. We present a design proposal for...

<https://arxiv.org/abs/2606.01023>

73 CONF

**DSL-LLaDA: Scaling Continuous Denoising to 8B Masked Diffusion LMs**

ArXiv CS.AI 50%

arXiv:2606.01024v1 Announce Type: cross Abstract: Discrete Masked diffusion language models generate text by iterative parallel decoding, but few-step decoding suffers from a tradeoff between length and quality: with a fixed step budget, standard methods can generate a short,...

<https://arxiv.org/abs/2606.01024>

74 CONF

**From Segments to Scenes: Temporal Understanding in Autonomous Driving via Vision-Language Model**

ArXiv CS.AI 50%

arXiv:2512.05277v3 Announce Type: replace-cross Abstract: Vision-Language Models (VLMs) are increasingly deployed as the perception and reasoning backbone of autonomous agents acting in the wild, with autonomous driving (AD) being one of the most safety-critical instances....

<https://arxiv.org/abs/2512.05277>

75 CONF

**Test-Time Training for Zero-Resource Dense Retrieval Reranking**

ArXiv CS.AI 50%

arXiv:2606.01070v1 Announce Type: cross Abstract: Dense retrievers excel at first-stage candidate generation but lack effective reranking in zero-resource settings. Existing approaches face a fundamental dilemma: cross-encoders deliver strong reranking quality but require costly...

<https://arxiv.org/abs/2606.01070>

76 CONF

**HASTE: Hardware-Aware Dynamic Sparse Training for Large Output Spaces**

ArXiv CS.AI 50%

arXiv:2606.01117v1 Announce Type: cross Abstract: Extreme multi-label classification (XMC) involves learning models over large output spaces with millions of labels, making the output layer a memory-compute bottleneck. While sparsity-based methods reduce arithmetic complexity,...

<https://arxiv.org/abs/2606.01117>

77 CONF

**AI From the Margins (AIM): Rethinking Participatory AI Design Through the Lived Experience of Minoritized Communities**

ArXiv CS.AI 50%

arXiv:2606.01171v1 Announce Type: cross Abstract: Artificial intelligence (AI) can reproduce and amplify the structural inequities faced by minoritized communities. Participatory AI has been proposed as a response, but participation typically starts after problem definitions and...

<https://arxiv.org/abs/2606.01171>

78 CONF

**Low-Resource Safety Failures Are Action Failures, Not Representation Failures**

ArXiv CS.AI 50%

arXiv:2606.01196v1 Announce Type: cross Abstract: Safety alignment learned in high-resource languages transfers poorly to low-resource languages. Models refuse harmful prompts in English but fail to refuse when the same prompts are translated into Swahili or Burmese. Adaptive...

<https://arxiv.org/abs/2606.01196>

79 CONF

**Implicit Geographic Inference in LLM Medical Triage: Language-Driven Disparities in Emergency Recommendations**

ArXiv CS.AI 50%

arXiv:2606.01204v1 Announce Type: cross Abstract: We investigate whether large language models produce different medical triage recommendations for identical symptoms based solely on the language of the patient prompt. Using Gemini 3.5 Flash, we evaluate a neurological symptom...

<https://arxiv.org/abs/2606.01204>

80 CONF

**Distilling Neuro-Symbolic Programs into 3D Multi-modal LLMs**

ArXiv CS.AI 50%

arXiv:2606.01215v1 Announce Type: cross Abstract: Current 3D spatial reasoning methods face a fundamental trade-off: neuro-symbolic 3D (NS3D) concept learners achieve interpretable reasoning through compositional programs but are constrained to closed-set concept vocabularies...

<https://arxiv.org/abs/2606.01215>

81 CONF

**IndoBias: A Dual Track Culturally Grounded Benchmark for LLMs Bias Evaluation in Indonesian Languages**

ArXiv CS.AI 50%

arXiv:2606.01260v1 Announce Type: cross Abstract: Despite being home to more than 1300 ethnic groups and 700 indigenous languages, bias in Large Language Models has not been fully studied in Indonesia, thus leaving a critical gap in evaluating representational fairness and...

<https://arxiv.org/abs/2606.01260>

82 CONF

**Quantum Algorithm for Distributed Reduction of Entanglements (QADR): A Trainable and Simulation-Efficient QML Framework**

ArXiv CS.AI 50%

arXiv:2606.01291v1 Announce Type: cross Abstract: Training Variational Quantum Circuits (VQCs) under Noisy Intermediate-Scale Quantum (NISQ) constraints introduces severe computational limitations: classical statevector simulation memory scales exponentially...

<https://arxiv.org/abs/2606.01291>



83 CONF

**A Lightweight Context-Driven Training-Free Network for Scene Text Segmentation and Recognition**

ArXiv CS.AI 50%

arXiv:2503.15639v2 Announce Type: replace-cross Abstract: Modern scene text recognition systems often depend on large end-to-end architectures that require extensive training and are prohibitively expensive for real-time scenarios. In such cases, the deployment of heavy models...

<https://arxiv.org/abs/2503.15639>

84 CONF

**A Communication-Centric 6G-LLM Architecture for Scalable Tactical Autonomous Defense Vehicle Networks**

ArXiv CS.AI 50%

arXiv:2606.01312v1 Announce Type: cross Abstract: The integration of Artificial Intelligence (AI) and emerging 6G networks introduces new opportunities for scalable coordination in tactical autonomous vehicle systems. This paper proposes a communication-centric hierarchical...

<https://arxiv.org/abs/2606.01312>

85 CONF

**Needles at Scale: LLM-Assisted Target Selection for Windows Vulnerability Research**

ArXiv CS.AI 50%

arXiv:2606.01364v1 Announce Type: cross Abstract: The attack surface of a modern operating system is a haystack: thousands of signed binaries and millions of functions, almost none relevant to any given vulnerability. A human analyst or an LLM agent must pick the function worth...

<https://arxiv.org/abs/2606.01364>

86 CONF

**Consistent and Distinctive: LLM Benchmark Efficiency via Maximum Independent Set Prompt Selection on Similarity Graphs**

ArXiv CS.AI 50%

arXiv:2606.01400v1 Announce Type: cross Abstract: Evaluating large language models (LLMs) across comprehensive benchmarks is expensive and time-consuming. We propose a graph-based prompt selection framework that models each benchmark as a similarity graph -- nodes are prompts...

<https://arxiv.org/abs/2606.01400>

87 CONF

**Neural Network Compression by Approximate Differential Equivalence**

ArXiv CS.AI 50%

arXiv:2606.01402v1 Announce Type: cross Abstract: Neural network compression is commonly achieved by pruning parameters based on local importance scores, e.g., magnitude-based pruning. We propose a complementary approach that compresses models by aggregating neurons with similar...

<https://arxiv.org/abs/2606.01402>

88 CONF

**UR-JEPA: Uniform Rectifiability as a Regularizer for Joint-Embedding Predictive Architectures**

ArXiv CS.AI 50%

arXiv:2606.01443v1 Announce Type: cross Abstract: A central difficulty in training Joint-Embedding Predictive Architectures (JEPAs) is preventing representation collapse. LeJEPa addresses this by enforcing an isotropic Gaussian target on the embeddings via Sketched Isotropic...

<https://arxiv.org/abs/2606.01443>

89 CONF

**Hierarchical Online Prompt Mutation with Dual-Loop Feedback for Guardrailed Evidence Document Generation: A Production-Evaluation Case Study**

ArXiv CS.AI 50%

arXiv:2606.01472v1 Announce Type: cross Abstract: High-stakes production document-generation systems require language models to be adaptive, evidence-grounded, and auditable. We present HOPM, a hierarchical online prompt mutation framework evaluated on a real marketplace...

<https://arxiv.org/abs/2606.01472>

90 CONF

**LLM Consortium for Software Design Refinement: A Controlled Experiment on Multi-Agent Collaboration Topologies**

ArXiv CS.AI 50%

arXiv:2606.01490v1 Announce Type: cross Abstract: We present a controlled experiment evaluating 12 multi-agent LLM collaboration topologies for software architecture design. Using a  $2 \times 2 \times 2$  factorial design (Authority  $\times$  Roles  $\times$  Dynamics), we conducted 520...

<https://arxiv.org/abs/2606.01490>

91 CONF

**FedS2R: One-Shot Federated Domain Generalization for Synthetic-to-Real Semantic Segmentation in Autonomous Driving**

ArXiv CS.AI 50%

arXiv:2507.19881v2 Announce Type: replace-cross Abstract: Federated domain generalization has shown promising progress in image classification by enabling collaborative training across multiple clients without sharing raw data. However, its potential in the semantic segmentation...

<https://arxiv.org/abs/2507.19881>

92 CONF

**Self-Conditioned Positional HNSW for Overlap-Aware Retrieval in Chunked-Document RAG Systems: Method and Industrial Evidence-Quality Audit**

ArXiv CS.AI 50%

arXiv:2606.01542v1 Announce Type: cross Abstract: Chunked-document retrieval is a common component of retrieval-augmented generation (RAG) systems. Documents are split into overlapping chunks, embedded, and indexed with approximate nearest-neighbor search such as hierarchical...

<https://arxiv.org/abs/2606.01542>

93 CONF

**GJNet: Robust Graph Neural Networks via Joint Disentangled Learning Against Adversarial Attacks**

ArXiv CS.AI 50%

arXiv:2606.01560v1 Announce Type: cross Abstract: Graph Neural Networks (GNNs) are vulnerable to adversarial attacks, which inherently invert connectivity patterns by introducing disassortative edges in assortative graphs and assortative edges in disassortative graphs. This...

<https://arxiv.org/abs/2606.01560>

94 CONF

**Identifying High-Confidence Social Biases in LLMs for Trustworthy Conversational Tutoring Agents**

ArXiv CS.AI 50%

arXiv:2606.01584v1 Announce Type: cross Abstract: Conversational tutoring agents have been shown to improve learning engagement and student outcomes, and large language models (LLMs) are increasingly used in these systems to provide scalable, personalized feedback. However, LLMs...

<https://arxiv.org/abs/2606.01584>

95 CONF

**Demystifying Multimodal Biomolecular Co-design With Intrinsic Geodesic Coupling**

ArXiv CS.AI 50%

arXiv:2606.01628v1 Announce Type: cross Abstract: Biomolecules such as proteins and small-molecule ligands play a central role in biological systems, arising from the tight interplay between sequence and three-dimensional structure. Recent generative models for biomolecular...

<https://arxiv.org/abs/2606.01628>

96 CONF

**E4GEN: Event-level Explainable Extreme-Enhanced Time-series Generation**

ArXiv CS.AI 50%

arXiv:2606.01634v1 Announce Type: cross Abstract: Generating realistic time series is essential for scientific research and real-world applications. However, existing methods often emphasize overall distributional fidelity while failing to faithfully capture extreme events. To...

<https://arxiv.org/abs/2606.01634>

97 CONF

**AlphaToken: Decoupling Adaptation and Stability for Path-Aware Response Token Valuation in LLM Post-Training**

ArXiv CS.AI 50%

arXiv:2606.01635v1 Announce Type: cross Abstract: Token selection is pivotal for effective LLM post-training. However, existing methods mostly rely on local heuristics and rarely formulate token selection as a principled valuation of individual response tokens. We introduce...

<https://arxiv.org/abs/2606.01635>

98 CONF

**Easier to Mislead Than to Correct: Harmful and Beneficial Revision in LLM Conformity**

ArXiv CS.AI 50%

arXiv:2606.01637v1 Announce Type: cross Abstract: Large language models are increasingly used in multi-agent systems, where they see and respond to other agents' answers. A key risk is conformity: a model may abandon its own answer simply because others agree on a different one....

<https://arxiv.org/abs/2606.01637>

99 CONF

**Off-the-Shelf LLMs as Process Scorers: Training-Free Alternative to PRMs for Mathematical Reasoning**

ArXiv CS.AI 50%

arXiv:2606.01682v1 Announce Type: cross Abstract: Selecting the best response from multiple small-model samples using a stronger scorer is a simple inference-time strategy, but fails when the small model has already committed to incorrect reasoning paths. PRM guided search...

<https://arxiv.org/abs/2606.01682>

00 CONF

**Argument Collapse: LLMs Flatten Long-Form Public Debate**

ArXiv CS.AI 50%

arXiv:2606.01736v1 Announce Type: cross Abstract: As LLMs are increasingly used to draft public-facing arguments, they may flatten public debate by repeatedly introducing the same polished, plausible arguments. We study argument collapse, the tendency of essays generated by...

<https://arxiv.org/abs/2606.01736>

01 CONF

**Adaptive Auto-Harness: Sustained Self-Improvement for Agentic System Deployment on Open-Ended Task Streams**

ArXiv CS.AI 50%

arXiv:2606.01770v1 Announce Type: cross Abstract: Auto-harness systems such as A-Evolve, GEPA, and Meta-Harness improve LLM agents by optimizing prompts, skills, tools, memories, and supporting infrastructure from execution feedback, but they are typically evaluated on fixed...

<https://arxiv.org/abs/2606.01770>

02 CONF

**LayerRoute: Input-Conditioned Adaptive Layer Skipping via LoRA Fine-Tuning for Agentic Language Models**

ArXiv CS.AI 50%

arXiv:2606.01838v1 Announce Type: cross Abstract: Agentic language model systems alternate between two structurally distinct step types: structured tool calls (short, deterministic, low perplexity) and open-ended planning/reasoning steps (long, complex, high perplexity). Despite...

<https://arxiv.org/abs/2606.01838>

03 CONF

**ProbeScale: Probing Analysis to Optimize Neural Scaling Laws for Efficient Small Language Model Inference**

ArXiv CS.AI 50%

arXiv:2606.01806v1 Announce Type: cross Abstract: Small Language Models (SLMs) offer a balance between capability and computational feasibility. Neural scaling laws inform their optimal training, suggesting that they possess rich internal representations that scale with their...

<https://arxiv.org/abs/2606.01806>

04 CONF

**Train, Test, Re-evaluate: Schedule-Sensitive Evaluation of Generative Data for Hand Detection**

ArXiv CS.AI 50%

arXiv:2606.01896v1 Announce Type: cross Abstract: Generated (or synthetic) image data is increasingly used to augment or replace real training datasets when target imagery is scarce, expensive, or biased. For hand detection, particularly in occupational safety settings, public...

<https://arxiv.org/abs/2606.01896>

05 CONF

**RA-LWLM: Retrieval-Augmented In-Context Localization with Wireless Foundation Models**

ArXiv CS.AI 50%

arXiv:2606.01899v1 Announce Type: cross Abstract: Wireless localization is a fundamental capability of sixth-generation (6G) networks. Conventional model-based methods require accurate modeling of the propagation environment and degrade in complex multipath and non-line-of-sight...

<https://arxiv.org/abs/2606.01899>

06 CONF

**The Image Reconstruction Game: Drawing Common Ground Through Iterative Multimodal Dialogue**

ArXiv CS.AI 50%

arXiv:2606.01901v1 Announce Type: cross Abstract: We introduce the Image Reconstruction Game, a fully automated benchmark in which a vision-language model issues corrective instructions to an image generator across multiple turns, making accumulated common ground directly...

<https://arxiv.org/abs/2606.01901>

07 CONF

**Echo: A Joint-Embedding Predictive Architecture for Speaker Diarization and Speech Recognition in a Shared Latent Space**

ArXiv CS.AI 50%

arXiv:2606.01909v1 Announce Type: cross Abstract: We present Echo, a proof-of-concept audio system built around a single 25 M-parameter ViT encoder. The encoder is pretrained with a JEPA objective and then specialised by stages to carry speaker identity, phonetic content, and...

<https://arxiv.org/abs/2606.01909>

08 CONF

**Collaborative and Efficient Fine-tuning: Leveraging Task Similarity**

ArXiv CS.AI 50%

arXiv:2602.07218v2 Announce Type: replace-cross Abstract: Adaptability has been regarded as a central feature in the foundation models, enabling them to effectively acclimate to unseen downstream tasks. Parameter-efficient fine-tuning methods such as celebrated LoRA facilitate...

<https://arxiv.org/abs/2602.07218>

09 CONF

**Towards 3D-Aware Video Diffusion Models: Render-Free Human Motion Control with Mesh Tokenization**

ArXiv CS.AI 50%

arXiv:2606.02000v1 Announce Type: cross Abstract: Diffusion models have shown remarkable success in video generation. However, whether such models are truly aware of the 3D structure underlying visual observations, rather than simply reproducing plausible 2D projections, remains...

<https://arxiv.org/abs/2606.02000>

10 CONF

**Agentic-J: An AI Agent for Biological Microscopy Image Analysis**

ArXiv CS.AI 50%

arXiv:2606.02080v1 Announce Type: cross Abstract: Biological image analysis increasingly demands integration across heterogeneous tools, programming environments, and domain knowledge that few researchers can command simultaneously. We present Agentic-J, a containerised,...

<https://arxiv.org/abs/2606.02080>

11 CONF

**From Features to Actions: Explainability in Traditional and Agentic AI Systems**

ArXiv CS.AI 50%

arXiv:2602.06841v4 Announce Type: replace Abstract: Over the last decade, Explainable AI has primarily focused on interpreting individual model predictions, producing post-hoc explanations that relate inputs to outputs under a fixed decision structure. Recent advances in large...

<https://arxiv.org/abs/2602.06841>

12 CONF

**A Primer in Post-Training Reasoning Data: What We Know About How It Works**

ArXiv CS.AI 50%

arXiv:2606.02113v1 Announce Type: cross Abstract: Post-training has become a primary driver of recent progress in large reasoning models, and reasoning data are often the key variable determining whether this stage succeeds. Work on post-training reasoning data has grown...

<https://arxiv.org/abs/2606.02113>

13 CONF

**Rethinking Evaluation Paradigms in IBP-based Certified Training**

ArXiv CS.AI 50%

arXiv:2606.02134v1 Announce Type: cross Abstract: Deep neural networks achieve strong performance on many supervised learning tasks but remain vulnerable to adversarial perturbations. Neural network verification provides mathematically rigorous robustness guarantees, yet at...

<https://arxiv.org/abs/2606.02134>

14 CONF

**Multimodal Approaches for Visually-Rich Document Type Classification: A Comparative Analysis**

ArXiv CS.AI 50%

arXiv:2606.02162v1 Announce Type: cross Abstract: Document type classification in visually rich documents remains challenging, as relevant information is distributed across textual, visual, and layout modalities. To capture this complexity, current approaches rely on diverse...

<https://arxiv.org/abs/2606.02162>

15 CONF

**Stop Wandering, Find the Keys: LLMs Discriminate Key States for Efficient Multi-Agent Exploration**

ArXiv CS.AI 50%

arXiv:2410.02511v2 Announce Type: replace Abstract: With expansive state-action spaces, efficient multi-agent exploration remains a longstanding challenge in reinforcement learning. Although pursuing novelty, diversity, or uncertainty attracts increasing attention, redundant...

<https://arxiv.org/abs/2410.02511>

16 CONF

**Faster Synchronous On-Policy RL via Straggler-Aware Group Sizing**

ArXiv CS.AI 50%

arXiv:2606.02218v1 Announce Type: cross Abstract: Synchronous reinforcement learning methods such as Group Relative Policy Optimization (GRPO) provide stable and reproducible on-policy training, but they are highly vulnerable to stragglers, a single unusually long rollout can...

<https://arxiv.org/abs/2606.02218>

17 CONF

**Who Annotates in NLP? A Large-scale Assessment of Human Annotation Reporting between 2018 and 2025**

ArXiv CS.AI 50%

arXiv:2606.02255v1 Announce Type: cross Abstract: Human annotation is the empirical foundation of much NLP research, from dataset construction to model evaluation, but papers often leave unclear who produced the annotations and how the annotation process was controlled. We...

<https://arxiv.org/abs/2606.02255>

18 CONF

**Repurposing Adversarial Perturbations for Continual Learning: From Defense to Active Alignment**

ArXiv CS.AI 50%

arXiv:2606.02322v1 Announce Type: cross Abstract: In dynamic environments, large language models need to keep adapting to new tasks, but continual learning often suffers from forgetting, limited transfer, and vulnerability to adversarial perturbations. To address this, we...

<https://arxiv.org/abs/2606.02322>

19 CONF

**Do Multimodal Agents Really Benefit from Tool Use? A Systematic Study of Capability Gains**

ArXiv CS.AI 50%

arXiv:2606.02357v1 Announce Type: cross Abstract: Tool-augmented multimodal agents show strong benchmark gains, often taken as evidence that agents have learned to use tools. We argue that this interpretation can be premature: a tool-call trace alone does not show whether the...

<https://arxiv.org/abs/2606.02357>

20 CONF

**Evolutionary Discovery of Bivariate Bicycle Codes with LLM-Guided Search**

ArXiv CS.AI 50%

arXiv:2606.02418v1 Announce Type: cross Abstract: Quantum LDPC code discovery requires searching large algebraic design spaces while reliably certifying the parameters and equivalence classes of any candidates found. We introduce an LLM-guided evolutionary workflow in which...

<https://arxiv.org/abs/2606.02418>

21 CONF

**ODTQA-FoRe: An Open-Domain Tabular Question Answering Dataset for Future Data Forecasting and Reasoning**

ArXiv CS.AI 50%

arXiv:2606.02433v1 Announce Type: cross Abstract: The rapid development of LLMs has significantly advanced tabular question answering, but most systems cannot perform future-oriented numerical prediction. To address this gap, we introduce a novel task, Open-Domain Tabular...

<https://arxiv.org/abs/2606.02433>

22 CONF

**ReasonBENCH: Benchmarking the (In)Stability of LLM Reasoning**

ArXiv CS.AI 50%

arXiv:2512.07795v2 Announce Type: replace Abstract: Benchmark scores for LLM reasoning systems are reported as single numbers, yet the same model, strategy, and task can produce meaningfully different answers and costs across repeated executions, even under greedy decoding (T =...

<https://arxiv.org/abs/2512.07795>

23 CONF

**Why Not Hyperparameter-Friendly Optimisation? A Monotonic Adaptive Norm Rescaling Approach For Long-Tailed Recognition**

ArXiv CS.AI 50%

arXiv:2606.02526v1 Announce Type: cross Abstract: Long-tailed recognition poses a significant challenge for deep learning. The two-stage decoupling paradigm, which separates representation learning from classifier retraining, offers a promising solution. During the classifier...

<https://arxiv.org/abs/2606.02526>

24 CONF

**From Layers to Submodules: Rethinking Granularity in Replacement-Based LLM Compression**

ArXiv CS.AI 50%

arXiv:2606.02559v1 Announce Type: cross Abstract: Post-training compression of Large Language Models (LLMs) removes entire architectural components, either deleting them or replacing them with fitted modules. Existing replacement-based methods share two design constraints:...

<https://arxiv.org/abs/2606.02559>

25 CONF

**Permissive Safety Through Trusted Inference: Verifiable Belief-Space Neural Safety Filters for Assured Interactive Robotics**

ArXiv CS.AI 50%

arXiv:2606.02562v1 Announce Type: cross Abstract: Autonomous robots that interact with people must make safe and efficient decisions under human-induced uncertainty, such as their preferences, goals, competency, and willingness to cooperate. Safety filters are a popular approach...

<https://arxiv.org/abs/2606.02562>

26 CONF

**AdaCodec: A Predictive Visual Code for Video MLLMs**

ArXiv CS.AI 50%

arXiv:2606.02569v1 Announce Type: cross Abstract: Video is temporally redundant: adjacent frames usually share most objects, background, and layout. Yet existing video multimodal large language models (video MLLMs) usually encode each sampled frame as an independent RGB image,...

<https://arxiv.org/abs/2606.02569>

27 CONF

**Mitigating Perceptual Judgment Bias in Multimodal LLM-as-a-Judge via Perceptual Perturbation and Reward Modeling**

ArXiv CS.AI 50%

arXiv:2606.02578v1 Announce Type: cross Abstract: Recent multimodal large language models have demonstrated strong reasoning ability, yet their reliability as automated evaluators remains limited by a critical weakness: when visual evidence conflicts with textual cues, MLLM...

<https://arxiv.org/abs/2606.02578>

28 CONF

**Algebraic anti-unification**

ArXiv CS.AI 50%

arXiv:2407.15510v2 Announce Type: replace Abstract: Abstraction is key to human and artificial intelligence as it allows one to identify common structure in otherwise distinct objects or situations. Anti-unification (or generalization) is the branch of theoretical computer...

<https://arxiv.org/abs/2407.15510>

29 CONF

**Explainable AI Through a Democratic Lens: DhondtXAI for D'Hondt-Projected Feature Attribution**

ArXiv CS.AI 50%

arXiv:2411.05196v3 Announce Type: replace Abstract: This study presents DhondtXAI as a SHAP-independent, D'Hondt-based attribution framework for tabular XAI. Instead of model-native feature importance or SHAP values, DhondtXAI computes background-interventional removal effects,...

<https://arxiv.org/abs/2411.05196>

30 CONF

**Safety Mirage: How Spurious Correlations Undermine VLM Safety Fine-Tuning and Can Be Mitigated by Machine Unlearning**

ArXiv CS.AI 50%

arXiv:2503.11832v5 Announce Type: replace Abstract: Recent vision language models (VLMs) have made remarkable strides in generative modeling with multimodal inputs, particularly text and images. However, their susceptibility to generating harmful content when exposed to unsafe...

<https://arxiv.org/abs/2503.11832>



31 CONF

**Finding the Minimal Parameter Budget for Implicit Reasoning: A Data Complexity Driven Scaling Law for Language Models**

ArXiv CS.AI 50%

arXiv:2504.03635v4 Announce Type: replace Abstract: Reasoning is a core capability of language models (LMs), yet it remains unclear how much model capacity is necessary to support reasoning during pretraining. In this work, we study the minimal parameter budget required for...

<https://arxiv.org/abs/2504.03635>

32 CONF

**EMoE: Training-Free Expert Disagreement for Uncertainty-Aware Text-to-Image Diffusion**

ArXiv CS.AI 50%

arXiv:2505.13273v2 Announce Type: replace Abstract: Large text-to-image diffusion models rarely expose reliable signals of when a prompt is likely to produce a poorly aligned generation, especially when training data is undisclosed. We study whether expert disagreement inside...

<https://arxiv.org/abs/2505.13273>

33 CONF

**InPhyRe Discovers: Large Multimodal Models Struggle in Inductive Physical Reasoning**

ArXiv CS.AI 50%

arXiv:2509.12263v3 Announce Type: replace Abstract: Large multimodal models (LMMs) encode physical laws observed during training, such as momentum conservation, as parametric knowledge. It allows LMMs to answer physical reasoning queries, such as the outcome of a potential...

<https://arxiv.org/abs/2509.12263>

34 CONF

**Query Circuits: Explaining How Language Models Answer User Prompts**

ArXiv CS.AI 50%

arXiv:2509.24808v2 Announce Type: replace Abstract: Explaining why a language model produces a particular output requires local, input-level explanations. Existing methods uncover global capability circuits (e.g., indirect object identification), but not why the model answers a...

<https://arxiv.org/abs/2509.24808>

35 CONF

**Safety Alignment of LMs via Non-cooperative Games**

ArXiv CS.AI 50%

arXiv:2512.20806v3 Announce Type: replace Abstract: Ensuring the safety of language models (LMs) while maintaining their usefulness remains a critical challenge in AI alignment. Current approaches rely on sequential adversarial training: generating adversarial prompts and...

<https://arxiv.org/abs/2512.20806>

36 CONF

**PolarMem: A Training-Free Polarized Latent Graph Memory for Verifiable Vision-Language Models**

ArXiv CS.AI 50%

arXiv:2602.00415v2 Announce Type: replace Abstract: Memory is not merely a storage mechanism for intelligent systems, but a structure for organizing evidence and constraining belief. This is especially important for multimodal reasoning, where retrieved evidence must be both...

<https://arxiv.org/abs/2602.00415>

37 CONF

**Structure Enables Effective Self-Localization of Errors in LLMs**

ArXiv CS.AI 50%

arXiv:2602.02416v2 Announce Type: replace Abstract: Self-correction in language models remains elusive. In this work, we explore whether language models can explicitly localize errors in incorrect reasoning, as a path toward building AI systems that can effectively correct...

<https://arxiv.org/abs/2602.02416>

38 CONF

**REAL: Resolving Knowledge Conflicts in Knowledge-Intensive Visual Question Answering via Reasoning-Pivot Alignment**

ArXiv CS.AI 50%

arXiv:2602.14065v2 Announce Type: replace Abstract: Knowledge-intensive Visual Question Answering (KI-VQA) frequently suffers from severe knowledge conflicts caused by the inherent limitations of open-domain retrieval. However, existing paradigms face critical limitations due to...

<https://arxiv.org/abs/2602.14065>

39 CONF

**LLM-WikiRace Benchmark: How Far Can LLMs Plan over Real-World Knowledge Graphs?**

ArXiv CS.AI 50%

arXiv:2602.16902v4 Announce Type: replace Abstract: We introduce LLM-Wikirace, a benchmark for evaluating planning, reasoning, and world knowledge in large language models (LLMs). In LLM-Wikirace, models must efficiently navigate Wikipedia hyperlinks step by step to reach a...

<https://arxiv.org/abs/2602.16902>

40 CONF

**Retrieval-aligned Tabular Foundation Models Enable Robust Clinical Risk Prediction in Electronic Health Records Under Real-world Constraints**

ArXiv CS.AI 50%

arXiv:2604.01841v2 Announce Type: replace Abstract: Clinical prediction from structured electronic health records (EHRs) is challenging due to high dimensionality, heterogeneity, class imbalance, and distribution shift. While tabular in-context learning (TICL) and...

<https://arxiv.org/abs/2604.01841>

41 CONF

**Perspective on Bias in Biomedical AI: Preventing Downstream Healthcare Disparities**

ArXiv CS.AI 50%

arXiv:2604.14514v2 Announce Type: replace Abstract: Healthcare disparities persist across socioeconomic boundaries, often attributed to unequal access to screening, diagnostics, and therapeutics. However, this perspective highlights that critical biases can emerge much earlier,...

<https://arxiv.org/abs/2604.14514>

42 CONF

**Trustworthy AI Suffers from Invariance Conflicts and Causality is The Solution**

ArXiv CS.AI 50%

arXiv:2605.02640v2 Announce Type: replace Abstract: As artificial intelligence (AI), including machine learning (ML) models and foundation models (FMs), are increasingly deployed in high-stakes domains, ensuring their trustworthiness has become a central challenge. However, the...

<https://arxiv.org/abs/2605.02640>

43 CONF

**Ethical Hyper-Velocity (EHV): A Hardware-Rooted Zero-Trust Runtime Enforcement Architecture for Agentic AI Systems**

ArXiv CS.AI 50%

arXiv:2605.17909v2 Announce Type: replace Abstract: As autonomous agentic systems scale across regulated critical infrastructures, the lack of mechanistic, hardware-rooted enforcement for high-frequency policy updates presents a fundamental safety gap. We present Ethical...

<https://arxiv.org/abs/2605.17909>

44 CONF

**Perturbation Effects on Accuracy and Fairness among Similar Individuals**

ArXiv CS.AI 50%

arXiv:2404.01356v3 Announce Type: replace-cross Abstract: Deep neural networks are vulnerable to adversarial perturbations that can simultaneously degrade prediction robustness and individual fairness across diverse application settings. However, existing evaluation protocols...

<https://arxiv.org/abs/2404.01356>

45 CONF

**Simulating Macroeconomic Expectations in Survey Experiments with LLM-based Economic Agents**

ArXiv CS.AI 50%

arXiv:2505.17648v5 Announce Type: replace-cross Abstract: We introduce a framework for simulating macroeconomic expectations in survey experiments using LLM-based economic agents (LLM Agents). We construct LLM Agents equipped with several functional modules that retrieve...

<https://arxiv.org/abs/2505.17648>

46 CONF

**Can LLMs Reason Structurally? Benchmarking via the Lens of Data Structures**

ArXiv CS.AI 50%

arXiv:2505.24069v4 Announce Type: replace-cross Abstract: Large language models (LLMs) are deployed on increasingly complex tasks that require multi-step decision-making. Understanding their algorithmic reasoning abilities is therefore crucial. However, we lack a diagnostic...

<https://arxiv.org/abs/2505.24069>

47 CONF

**GFlowGR: Fine-tuning Generative Recommendation Frameworks with Generative Flow Networks**

ArXiv CS.AI 50%

arXiv:2506.16114v3 Announce Type: replace-cross Abstract: Generative recommendations (GR), which usually include item tokenizers and generative Large Language Models (LLMs), have demonstrated remarkable success across a wide range of scenarios. The majority of existing research...

<https://arxiv.org/abs/2506.16114>

48 CONF

**Truth, Trust, and Trouble: Medical AI on the Edge**

ArXiv CS.AI 50%

arXiv:2507.02983v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) hold significant promise for transforming digital health by enabling automated medical question answering. However, ensuring these models meet critical industry standards for factual accuracy,...

<https://arxiv.org/abs/2507.02983>

49 CONF

**Toward accurate RUL and SoH estimation using reinforced graph-based physics-informed neural networks enhanced with dynamic weights**

ArXiv CS.AI 50%

arXiv:2507.09766v2 Announce Type: replace-cross Abstract: Accurate estimation of Remaining Useful Life (RUL) and State of Health (SoH) is essential for reliable Prognostics and Health Management (PHM), supporting timely maintenance and dependable industrial operation. However,...

<https://arxiv.org/abs/2507.09766>

50 CONF

**Beyond Model Base Retrieval: Weaving Knowledge to Master Fine-grained Neural Network Design**

ArXiv CS.AI 50%

arXiv:2507.15336v3 Announce Type: replace-cross Abstract: Designing high-performance neural networks for new tasks requires balancing optimization quality with search efficiency. Current methods fail to achieve this balance: neural architectural search is computationally...

<https://arxiv.org/abs/2507.15336>

51 CONF

**Graph is a Natural Regularization: Revisiting Vector Quantization for Graph Representation Learning**

ArXiv CS.AI 50%

arXiv:2508.06588v3 Announce Type: replace-cross Abstract: Vector Quantization (VQ) has recently emerged as a promising approach for learning compressed and discrete representations for graph-structured data. However, a fundamental challenge, i.e., codebook collapse, remains...

<https://arxiv.org/abs/2508.06588>

52 CONF

**Position: Beyond Sensitive Attributes, ML Fairness Should Quantify Structural Injustice via Social Determinants**

ArXiv CS.AI 50%

arXiv:2508.08337v3 Announce Type: replace-cross Abstract: Algorithmic fairness research has largely framed unfairness as discrimination along sensitive attributes. However, this approach limits visibility into unfairness as structural injustice instantiated through social...

<https://arxiv.org/abs/2508.08337>

53 CONF

**Language-Native Materials Processing Design by Lightly Structured Text Database and Reasoning Large Language Model**

ArXiv CS.AI 50%

arXiv:2509.06093v4 Announce Type: replace-cross Abstract: Materials synthesis procedures are predominantly documented as narrative text in papers, protocols, and laboratory records, placing them beyond the reach of conventional data-driven optimization frameworks. This...

<https://arxiv.org/abs/2509.06093>

54 CONF

**Deep Learning as the Disciplined Construction of Tame Objects**

ArXiv CS.AI 50%

arXiv:2509.18025v2 Announce Type: replace-cross Abstract: One can see deep-learning models as compositions of functions within the so-called tame geometry. In this expository note, we give an overview of some topics at the interface of tame geometry (also known as o-minimality),...

<https://arxiv.org/abs/2509.18025>

55 CONF

**Simultaneous Multi-objective Alignment Across Verifiable and Non-verifiable Rewards**

ArXiv CS.AI 50%

arXiv:2510.01167v2 Announce Type: replace-cross Abstract: Aligning large language models to human preferences is inherently multidimensional, yet most pipelines collapse heterogeneous signals into a single objective. We seek to answer what it would take to simultaneously align a...

<https://arxiv.org/abs/2510.01167>

56 CONF

**HRTFformer: A Spatially-Aware Transformer for Individual HRTF Upsampling in Immersive Audio Rendering**

ArXiv CS.AI 50%

arXiv:2510.01891v2 Announce Type: replace-cross Abstract: Individual Head-Related Transfer Functions (HRTFs) are starting to be introduced in many commercial immersive audio applications and are crucial for realistic spatial audio rendering. However, one of the main hesitations...

<https://arxiv.org/abs/2510.01891>

57 CONF

**SHERLOCK: Towards Dynamic Knowledge Adaptation in LLM-enhanced E-commerce Risk Management**

ArXiv CS.AI 50%

arXiv:2510.08948v4 Announce Type: replace-cross Abstract: Effective e-commerce risk management requires in-depth case investigations to identify emerging fraud patterns in highly adversarial environments. However, manual investigation typically requires analyzing the...

<https://arxiv.org/abs/2510.08948>

58 CONF

**Video Reasoning without Training**

ArXiv CS.AI 50%

arXiv:2510.17045v2 Announce Type: replace-cross Abstract: Video reasoning using Large Multimodal Models (LMMs) relies on costly reinforcement learning (RL) and verbose chain-of-thought, resulting in substantial computational overhead during both training and inference. Moreover,...

<https://arxiv.org/abs/2510.17045>

59 CONF

**Who Evaluates AI's Social Impacts? Mapping Coverage and Gaps in First and Third Party Evaluations**

ArXiv CS.AI 50%

arXiv:2511.05613v2 Announce Type: replace-cross Abstract: Foundation models are increasingly central to high-stakes AI systems, and governance frameworks now depend on evaluations to assess their risks and capabilities. Although general capability evaluations are widespread,...

<https://arxiv.org/abs/2511.05613>

60 CONF

**Evaluating the Performance of Deep Learning Models in Whole-body Dynamic 3D Posture Prediction During Load-reaching Activities**

ArXiv CS.AI 50%

arXiv:2511.20615v2 Announce Type: replace-cross Abstract: This study aimed to explore the application of deep neural networks for whole-body human posture prediction during dynamic load-reaching activities. Two time-series models were trained using bidirectional long short-term...

<https://arxiv.org/abs/2511.20615>

61 CONF

**SpeedAug: Policy Acceleration via Tempo-Enriched Policy and RL Fine-Tuning**

ArXiv CS.AI 50%

arXiv:2512.00062v2 Announce Type: replace-cross Abstract: Robotic policy learning for complex real-world manipulation tasks has seen rapid recent progress, enabled in large part by the ability to collect demonstrations through human operation. However, policies trained from such...

<https://arxiv.org/abs/2512.00062>

62 CONF

**InFerActive: Interactive Tree-Based Exploration of LLM Sampling for Safety Evaluation**

ArXiv CS.AI 50%

arXiv:2512.10234v2 Announce Type: replace-cross Abstract: Even LLMs that appear safe during evaluation can still produce harmful responses in deployment. Because stochastic sampling yields different responses to the same prompt, low-probability harmful outputs can still reach...

<https://arxiv.org/abs/2512.10234>

63 CONF

**Calibrating Uncertainty for Zero-Shot Adversarial CLIP**

ArXiv CS.AI 50%

arXiv:2512.12997v2 Announce Type: replace-cross Abstract: CLIP delivers strong zero-shot classification but remains highly vulnerable to adversarial attacks. Prior adversarial fine-tuning work primarily matches predicted logits between clean and adversarial examples, which...

<https://arxiv.org/abs/2512.12997>

64 CONF

**MGRegBench: A Novel Benchmark Dataset with Anatomical Landmarks for Mammography Image Registration**

ArXiv CS.AI 50%

arXiv:2512.17605v2 Announce Type: replace-cross Abstract: Robust mammography registration is essential for clinically relevant applications like tracking disease progression in breast tissue. However, progress has been limited by the absence of transparent public datasets and...

<https://arxiv.org/abs/2512.17605>

65 CONF

**naPINN: Noise-Adaptive Physics-Informed Neural Networks for Recovering Physics from Corrupted Measurement**

ArXiv CS.AI 50%

arXiv:2602.02547v2 Announce Type: replace-cross Abstract: Physics-Informed Neural Networks (PINNs) are effective methods for solving inverse problems and discovering governing equations from observational data. However, their performance degrades significantly under complex...

<https://arxiv.org/abs/2602.02547>

66 CONF

**The Alignment Curse: Modality Alignment Supercharges Audio Attacks via Text Transfer**

ArXiv CS.AI 50%

arXiv:2602.02557v2 Announce Type: replace-cross Abstract: Recent advances in end-to-end trained omni-models have substantially improved audio capabilities by strengthening text-audio modality alignment. However, whether such alignment inadvertently facilitates the transfer of...

<https://arxiv.org/abs/2602.02557>

67 CONF

**When Single Answer Is Not Enough: Rethinking Single-Step Retrosynthesis Benchmarks for LLMs**

ArXiv CS.AI 50%

arXiv:2602.03554v2 Announce Type: replace-cross Abstract: Recent progress has expanded the use of large language models (LLMs) in drug discovery, including synthesis planning. However, objective evaluation of retrosynthesis performance remains limited. Existing benchmarks and...

<https://arxiv.org/abs/2602.03554>

68 CONF

**Evaluating Reliability Asymmetries in Chinese Factual Search and AI Answers**

ArXiv CS.AI 50%

arXiv:2602.22221v2 Announce Type: replace-cross Abstract: Search engines and AI-powered systems increasingly mediate access to factual information, yet their reliability remains difficult to evaluate in realistic information-seeking settings. We study this problem in the Chinese...

<https://arxiv.org/abs/2602.22221>

69 CONF

**Optimal Bayesian Stopping for Efficient Inference of Consistent LLM Answers**

ArXiv CS.AI 50%

arXiv:2602.05395v2 Announce Type: replace-cross Abstract: A simple strategy for improving LLM accuracy, especially in math and reasoning problems, is to sample multiple responses and submit the answer most consistently reached. In this paper we leverage Bayesian prior...

<https://arxiv.org/abs/2602.05395>

70 CONF

**SoK: DARPA's AI Cyber Challenge (AixCC): Competition Design, Architectures, and Lessons Learned**

ArXiv CS.AI 50%

arXiv:2602.07666v4 Announce Type: replace-cross Abstract: DARPA's AI Cyber Challenge (AixCC, 2023--2025) is the largest competition to date for building fully autonomous cyber reasoning systems (CRSs) that leverage recent advances in AI -- particularly large language models...

<https://arxiv.org/abs/2602.07666>

71 CONF

**Knowing Isn't Understanding: Re-grounding Generative Proactivity with Epistemic and Behavioral Insight**

ArXiv CS.AI 50%

arXiv:2602.15259v2 Announce Type: replace-cross Abstract: Generative AI agents equate understanding with resolving explicit queries, an assumption that confines interaction to what users can articulate. This assumption breaks down when users themselves lack awareness of what is...

<https://arxiv.org/abs/2602.15259>

72 CONF

**Scaling Search Relevance: Augmenting App Store Ranking with LLM-Generated Judgments**

ArXiv CS.AI 50%

arXiv:2602.23234v4 Announce Type: replace-cross Abstract: Large-scale commercial search systems optimize for relevance to drive successful sessions that help users find what they are looking for. To maximize relevance, we leverage two complementary objectives: behavioral...

<https://arxiv.org/abs/2602.23234>

73 CONF

**Interpretable Multimodal Gesture Recognition for Drone and Mobile Robot Teleoperation via Log-Likelihood Ratio Fusion**

ArXiv CS.AI 50%

arXiv:2602.23694v3 Announce Type: replace-cross Abstract: Human operators are still frequently exposed to hazardous environments such as disaster zones and industrial facilities, where intuitive and reliable teleoperation of mobile robots and Unmanned Aerial Vehicles (UAVs) is...

<https://arxiv.org/abs/2602.23694>

74 CONF

**Ethical Fairness in Ubiquitous Health Sensing without Known Attributes**

ArXiv CS.AI 50%

arXiv:2603.13373v4 Announce Type: replace-cross Abstract: In ubiquitous and mobile health systems, computational models infer human states from wearable, behavioral, and physiological sensing data. In these settings, high accuracy alone is insufficient; models must act ethically...

<https://arxiv.org/abs/2603.13373>

75 CONF

**ES-Merging: Biological MLLM Merging via Embedding Space Signals**

ArXiv CS.AI 50%

arXiv:2603.14405v2 Announce Type: replace-cross Abstract: Biological multimodal large language models (MLLMs) have emerged as powerful foundation models for scientific discovery. However, existing models are specialized to a single modality, limiting their ability to solve...

<https://arxiv.org/abs/2603.14405>

76 CONF

**Large Language Model Guided Incentive Aware Reward Design for Cooperative Multi-Agent Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2603.24324v4 Announce Type: replace-cross Abstract: Designing effective auxiliary rewards for cooperative multi-agent systems remains challenging, as misaligned incentives can induce suboptimal coordination, particularly when sparse task rewards provide insufficient...

<https://arxiv.org/abs/2603.24324>

77 CONF

**Limits of Spatial Imagery Reasoning in Frontier LLM Models**

ArXiv CS.AI 50%

arXiv:2603.26779v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) have demonstrated impressive reasoning capabilities, yet they struggle with spatial tasks that require mental simulation, such as mental rotation. This paper investigates whether equipping an...

<https://arxiv.org/abs/2603.26779>

78 CONF

**Gradients with Respect to Semantics Preserving Embeddings Tell the Uncertainty of Large Language Models**

ArXiv CS.AI 50%

arXiv:2605.04638v2 Announce Type: replace-cross Abstract: Uncertainty quantification (UQ) is an important technique for ensuring the trustworthiness of LLMs, given their tendency to hallucinate. Existing state-of-the-art UQ approaches for free-form generation rely heavily on...

<https://arxiv.org/abs/2605.04638>



79 CONF

**OGLS-SD: On-Policy Self-Distillation with Outcome-Guided Logit Steering for LLM Reasoning**

ArXiv CS.AI 50%

arXiv:2605.12400v2 Announce Type: replace-cross Abstract: We study on-policy self-distillation (OPSD), where a language model improves its reasoning ability by distilling privileged teacher distributions along its own on-policy trajectories. Despite its promise, OPSD can suffer...

<https://arxiv.org/abs/2605.12400>

80 CONF

**RISED: A Pre-Deployment Evaluation Framework for High-Stakes AI Decision-Support Systems, with Application to Healthcare**

ArXiv CS.AI 50%

arXiv:2605.12895v2 Announce Type: replace-cross Abstract: Clinical decision-support systems are expert systems whose recommendations clinicians act on directly, yet they are usually cleared on one aggregate accuracy number from a held-out test set. That number says nothing about...

<https://arxiv.org/abs/2605.12895>

81 CONF

**Avoiding Structural Failure Modes in Tabular Fair SSL: Online Primal-Dual Allocation under Confidence Gating**

ArXiv CS.AI 50%

arXiv:2605.16446v2 Announce Type: replace-cross Abstract: Semi-supervised learning (SSL) enables prediction with limited labels, but high-stakes tabular applications (medical, credit, recidivism) require statistical fairness guarantees. We identify a structural conflict in...

<https://arxiv.org/abs/2605.16446>

82 CONF

**Can Vision Models Truly Forget? Mirage: Representation-Level Certification of Visual Unlearning**

ArXiv CS.AI 50%

arXiv:2605.20282v2 Announce Type: replace-cross Abstract: Machine unlearning in Vertical Federated Learning (VFL) has attracted growing interest, yet existing methods certify forgetting solely using output-level metrics. We challenge these claims by introducing Mirage, a...

<https://arxiv.org/abs/2605.20282>

83 CONF

**Algorithmic Fragility and Persona Bias in LLM-Generated Autistic Communication**

ArXiv CS.AI 50%

arXiv:2605.26397v2 Announce Type: replace-cross Abstract: Safety alignment reduces explicitly harmful outputs but inadvertently encodes a sanitized, neuronormative representation of marginalized communication. We investigate this encoding using a dual-persona rewrite paradigm,...

<https://arxiv.org/abs/2605.26397>

84 CONF

**Targeted Remasking: Replacing Token Editing with Token-to-Mask Refinement in Discrete Diffusion Language Models**

ArXiv CS.AI 50%

arXiv:2605.26436v2 Announce Type: replace-cross Abstract: Discrete masked diffusion language models such as LLaDA generate text through iterative denoising, where mask tokens are progressively replaced with predicted tokens. LLaDA2.1 introduced a Token-to-Token (T2T) editing...

<https://arxiv.org/abs/2605.26436>

85 CONF

**Dissociative Identity: Language Model Agents Lack Grounding for Reputation Mechanisms**

ArXiv CS.AI 50%

arXiv:2605.30169v2 Announce Type: replace-cross Abstract: As autonomous language model agents proliferate, forming an emerging agentic web with real-world consequences, what credibility signals can you use to decide whether to trust an unfamiliar agent in the wild and delegate...

<https://arxiv.org/abs/2605.30169>