



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	146	Media
TechCrunch AI	11	Media
The Verge AI	6	Media
HackerNews	5	Community
MIT Tech Review	2	Media

VERIFIED INTELLIGENCE — 6 stories

VERI

Show HN: Open Envelope - an open schema for defining AI agent teams

HackerNews

+1 source

90%

HN 51

arXiv:2606.03463v1 Announce Type: new Abstract: Conversational AI agents require memory systems that are both scalable and semantically coherent across long interaction horizons. Existing approaches rely predominantly on large language model (LLM)-based summarisation at write...

<https://openenvelope.org/docs/schema/>

VERI

When Should LLMs Be Less Specific? Selective Abstraction for Reliable Long-Form Text Generation

ArXiv CS.AI

+1 source

80%

arXiv:2602.11908v3 Announce Type: replace Abstract: LLMs are widely used, yet they remain prone to factual errors that erode user trust and limit adoption in high-risk settings. One approach to mitigate this risk is to equip models with uncertainty estimation mechanisms that...

<https://arxiv.org/abs/2602.11908>

VERI

Trump signs executive order to review AI models before they're released

The Verge AI

+1 source

80%

President Donald Trump signed an executive order Tuesday creating a "voluntary framework" for AI companies to share their frontier models with the federal government before they're released "to promote secure innovation and strengthen the cybersecurity of critical...

<https://www.theverge.com/policy/941775/trump-ai-executive-order>

VERI

OpenAI launches new Codex tools for white-collar work

TechCrunch AI

+1 source

80%

OpenAI released a set of six plug-ins aimed at specific jobs: data analytics, creative production, sales, product design, equity investing, and investment banking. Available from within the Codex app, each of the new tools bundles integrations, instructions, and context to allow...

<https://techcrunch.com/2026/06/02/openai-launches-new-codex-tools-for-white-collar-work/>

VERI

Rehumanizing global health care with agentic AI

MIT Tech Review

+1 source

80%

The global health care sector is under increasing strain. Decades of chronic underinvestment and constraints in recruitment have coincided with a surge in demand for services for aging populations. Gaps in provision are already taking a toll, with fragmented access to care and...

<https://www.technologyreview.com/2026/06/02/1137827/rehumanizing-global-health-care-wit...>

VERI

Evaluating Relational Reasoning in LLMs with REL

ArXiv CS.AI

+1 source

80%

arXiv:2604.12176v2 Announce Type: replace Abstract: Relational reasoning is the ability to infer relations that jointly bind multiple entities, attributes, or variables. This ability is central to scientific reasoning, but existing evaluations of relational reasoning in large...

<https://arxiv.org/abs/2604.12176>

CONFIRMED SIGNALS — 158 stories

CONF

Show HN: I built an on-chain economy where AI agents transact autonomously

HackerNews 50% HN 9

No summary — see link for full article.

<https://agtchain.io/>

CONF

Show HN: I built a way to find and install Claude skills

HackerNews 50% HN 5

I've been experimenting with ways to increase AI adoption for non-technical people. Basically, all companies are pushing for AI because it's all over the news and they feel left behind but most people have no clue where to start. I think 90% of people (ie non coders) are...

<https://www.claudioho.xyz/>

CONF

Show HN: AI-Powered PDF to Markdown Converter

HackerNews 50% HN 3

Turn complex PDFs into clean Markdown that people can review and AI tools can use.

<https://pdfdomconverter.com>

0 CONF

Show HN: Claude wrote me a language, now I don't know what to do with it

HackerNews 50% HN 2

Claude and I started building FROG in an effort to stop burning through Anthropic credit so quickly. That was the only reason I honestly needed at the time, but as the language grew and features actually worked together, the language itself quickly became the reason to continue....

<https://github.com/kmcnally5/FROG>

1 CONF

Uber caps employee AI spending after blowing through budget in 4 months

TechCrunch AI 50%

Uber's cutback has occurred after the company had reportedly encouraged staff to use AI as much as possible.

<https://techcrunch.com/2026/06/02/uber-caps-employee-ai-spending-after-blowing-through-...>

2 CONF

Martin Scorsese becomes the latest — and most unlikely — Hollywood voice for AI

TechCrunch AI 50%

The caveat is that one of the world's most famous living directors is using the tech solely for storyboarding.

<https://techcrunch.com/2026/06/02/martin-scorsese-becomes-the-latest-and-most-unlikely-...>

3 CONF

Google rolls out fake call detection to protect against AI deepfake impersonation scams

TechCrunch AI 50%

As people increasingly refuse to answer calls from unknown numbers, scammers are shifting their tactics by spoofing trusted phone numbers and using AI deepfake technology to sound like authority figures, family members, or employers.

<https://techcrunch.com/2026/06/02/google-rolls-out-fake-call-detection-to-protect-again-...>

4 CONF

Anthropic scales Claude Mythos to critical infrastructure in 15+ countries

TechCrunch AI 50%

Anthropic is expanding Project Glasswing, its security vulnerability program, and access to Mythos to 150 organizations across 15 countries — targeting critical infrastructure in power, water, healthcare, and communications where a cyberattack could affect 100 million people.

<https://techcrunch.com/2026/06/02/anthropic-scales-claude-mythos-to-critical-infrastruc-...>

5 **CONF****ZeroDrift raises \$10M to protect AI models from themselves****TechCrunch AI** 50%

A new AI compliance service sits between AI models and end users to flag and replace any messages that might present a compliance problem.

<https://techcrunch.com/2026/06/02/zerodrift-raises-10-million-to-protect-ai-models-from...>

6 **CONF****Rocket engine startup Impulse raises \$500 million to hire people, not AI****TechCrunch AI** 50%

Engineering physical systems still depends on human talent, according to Impulse Space president Eric Romo.

<https://techcrunch.com/2026/06/02/rocket-engine-startup-impulse-raises-500-million-to-h...>

7 **CONF****Florida sues OpenAI, Sam Altman, in first-of-its-kind lawsuit over violent incidents****TechCrunch AI** 50%

The lawsuit partially revolves around a shooting at Florida State University last year, and ChatGPT's alleged role in the incident.

<https://techcrunch.com/2026/06/01/florida-sues-openai-sam-altman-in-first-of-its-kind-l...>

8 **CONF****Anthropic files to go public****TechCrunch AI** 50%

Anthropic, now an AI powerhouse that has landed top-tier enterprise customers, was once considered an underdog in the emerging world of large language models.

<https://techcrunch.com/2026/06/01/anthropic-files-to-go-public/>

9 **CONF****How small businesses can leverage AI****MIT Tech Review** 50%

This article is from Making AI Work, MIT Technology Review's limited-run newsletter examining how to apply LLMs across industries. To receive it in your inbox, sign up here. From accounting to design to market research and product development, there's a staggering breadth of...

<https://www.technologyreview.com/2026/06/02/1138227/how-small-businesses-can-leverage-ai/>

0 **CONF****Microsoft Scout is a new AI personal assistant built on OpenClaw****The Verge AI** 50%

Much like Google, Microsoft is launching its own version of OpenClaw. Microsoft Scout is an always-on assistant that integrates into Microsoft 365 apps like Outlook, OneDrive, and Microsoft Teams, allowing businesses to assign a virtual assistant to employees to help with...

<https://www.theverge.com/news/939713/microsoft-scout-assistant-openclaw>

1 **CONF****Microsoft's Project Solara is an OS for AI agent gadgets****The Verge AI** 50%

Microsoft just announced "Project Solara," a new OS designed for gadgets that run AI agents, at Build 2026. The company is calling it "a new platform built from the ground up to power agent-driven experiences." It's built on Android, not Windows. Microsoft demonstrated two...

<https://www.theverge.com/news/941830/microsoft-project-solara-os-ai-agent-gadgets>

2 CONF

Microsoft Build 2026: All the news about Windows, AI, RTX Spark, and more

The Verge AI 50%

Microsoft Build 2026 kicked off with a keynote presentation that introduced some developer-focused Windows updates, an OpenClaw-based AI assistant called Scout, the new Majorana 2 quantum computing chip, and a Surface mini PC designed for AI developers. Microsoft also rolled out...

<https://www.theverge.com/tech/941668/microsoft-build-may-2026-live-news-updates>

3 CONF

Gemini Spark is the most impressive and terrifying AI experience I've had yet

The Verge AI 50%

According to every product demo from the last four years, planning a trip is a killer use case for AI. Just tell it where you're going, they all promise, and your chatbot / agent / other buzzword will exhaustively search travel options, read up on all the fun things to do, check...

<https://www.theverge.com/ai-artificial-intelligence/941388/gemini-spark-ai-agent-trip-p...>

4 CONF

Visual Graph Scaffolds for Structural Reasoning in Large Language Models

ArXiv CS.AI 50%

arXiv:2606.02673v1 Announce Type: new Abstract: Graphs have been used to enhance large language models (LLMs) for structured reasoning, mostly as external knowledge sources are provided to models at test time. In this paper, we take a different view: the value of graphs for LLMs...

<https://arxiv.org/abs/2606.02673>

5 CONF

A New Framework for Cybersecurity Refusals in AI Agents

ArXiv CS.AI 50%

arXiv:2606.02644v1 Announce Type: cross Abstract: Agentic scaffolds have dramatically improved LLM performance on complex, long-horizon tasks, yielding both broad benefits and amplified risks in domains like cybersecurity. Existing benchmarks for AI agents in cybersecurity focus...

<https://arxiv.org/abs/2606.02644>

6 CONF

ChatHealthAI: Aligning Electronic Health Record Representations with Large Language Models for Grounded Clinical Reasoning

ArXiv CS.AI 50%

arXiv:2606.02802v1 Announce Type: new Abstract: Large language models (LLMs) exhibit strong natural-language reasoning abilities for clinical decision support, but struggle to effectively model structured longitudinal electronic health records (EHRs). In contrast, EHR foundation...

<https://arxiv.org/abs/2606.02802>

7 CONF

Toward a Modular Architecture for Embedded AI Agent Systems at the Edge

ArXiv CS.AI 50%

arXiv:2606.02862v1 Announce Type: new Abstract: The rise of Large Language Models (LLMs) has enabled agentic AI capable of complex reasoning and tool use; however, deploying such autonomy in pervasive computing environments remains challenging due to the strict memory and energy...

<https://arxiv.org/abs/2606.02862>

8 CONF

StepFinder: A Temporal Semantic Framework for Failure Attribution in Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2606.03467v1 Announce Type: new Abstract: LLM-based multi-agent systems exhibit remarkable collaborative capabilities in complex multi-step tasks. However, these systems are highly sensitive to single-step execution errors that can propagate through agent interactions and...

<https://arxiv.org/abs/2606.03467>

9 CONF

Large AI Models in Dental Healthcare: From General-Purpose Systems to Domain-Specific Foundation Models

ArXiv CS.AI 50%

arXiv:2606.02914v1 Announce Type: new Abstract: Background: Oral diseases affect nearly 3.5 billion people worldwide, yet the comparative clinical potential of large-scale AI models in dentistry remains poorly understood. Three distinct model categories have emerged:...

<https://arxiv.org/abs/2606.02914>

0 CONF

What Benchmarks Don't Measure: The Case for Evaluating Abstention Competence in Autonomous Agents

ArXiv CS.AI 50%

arXiv:2606.02965v1 Announce Type: new Abstract: Benchmarks for autonomous agents measure whether agents complete tasks, yet this framing is systematically blind to whether an agent should have proceeded at all. Agents trained under human-feedback objectives develop a structural...

<https://arxiv.org/abs/2606.02965>

1 CONF

Assistax: A Multi-Agent Hardware-Accelerated Reinforcement Learning Benchmark for Assistive Robotics

ArXiv CS.AI 50%

arXiv:2507.21638v2 Announce Type: replace Abstract: The development of reinforcement learning (RL) algorithms has been largely driven by ambitious challenge tasks and benchmarks. Games have dominated RL benchmarks because they present relevant challenges, are inexpensive to run...

<https://arxiv.org/abs/2507.21638>

2 CONF

TriEval: A Resource-Efficient Pipeline for LLM Bias, Toxicity, and Truthfulness Assessment

ArXiv CS.AI 50%

arXiv:2606.03036v1 Announce Type: new Abstract: LLMs have evolved from basic chatbots to the backbone of the AI ecosystem, now widely used in healthcare, schools, and government services. The domain-wide adoption of LLMs necessitates continuous evaluation to ensure their safety...

<https://arxiv.org/abs/2606.03036>

3 CONF

RelGT-AC: A Relational Graph Transformer for Autocomplete Tasks in Relational Databases

ArXiv CS.AI 50%

arXiv:2606.03040v1 Announce Type: new Abstract: Relational databases underpin modern enterprise, scientific, and healthcare systems, yet predictive machine learning on such data remains challenging due to their multi-table, heterogeneous, and temporal structure. Relational Deep...

<https://arxiv.org/abs/2606.03040>

4 **CONF****CORE: Conflict-Oriented Reasoning for General Multimodal Manipulation Detection**

ArXiv CS.AI 50%

arXiv:2606.03066v1 Announce Type: new Abstract: The rapid rise of generative AI has made multimodal fake news increasingly realistic and pervasive, posing severe threats to public trust and social stability. Existing detection methods rely heavily on manipulation-specific models...

<https://arxiv.org/abs/2606.03066>
5 **CONF****A Scoping Review of the Ethical Perspectives on Anthropomorphising Large Language Model-Based Conversational Agents**

ArXiv CS.AI 50%

arXiv:2601.09869v2 Announce Type: replace Abstract: Anthropomorphisation -- the phenomenon whereby non-human entities are ascribed human-like qualities -- has become increasingly salient with the rise of large language model (LLM)-based conversational agents (CAs). Unlike...

<https://arxiv.org/abs/2601.09869>
6 **CONF****DeskCraft: Benchmarking Desktop Agents on Professional Workflows and Human-in-the-Loop Collaboration**

ArXiv CS.AI 50%

arXiv:2606.03103v1 Announce Type: new Abstract: Real-world professional desktop workflows in specialized creative and engineering software unfold over long horizons and often require human-in-the-loop coordination, where agents proactively seek necessary information and users...

<https://arxiv.org/abs/2606.03103>
7 **CONF****Post-Hoc Robustness for Model-Based Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2606.03521v1 Announce Type: cross Abstract: To improve the real-world applicability of reinforcement learning (RL), the field of adversarially robust RL studies how to train agents under adversarial environment perturbations. In this setting, a protagonist agent optimizes...

<https://arxiv.org/abs/2606.03521>
8 **CONF****SeSE: Black-Box Uncertainty Quantification for Large Language Models Based on Structural Information Theory**

ArXiv CS.AI 50%

arXiv:2511.16275v4 Announce Type: replace-cross Abstract: Reliable uncertainty quantification (UQ) is essential for deploying large language models (LLMs) in safety-critical scenarios, as it enables them to abstain from responding when uncertain, thereby avoiding hallucinations,...

<https://arxiv.org/abs/2511.16275>
9 **CONF****GTBench: A Curriculum-Grounded Benchmark for Evaluating LLMs as Mathematical Research Assistants in Graph Theory**

ArXiv CS.AI 50%

arXiv:2606.03144v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used as self-study assistants in technical disciplines, yet their reliability as mathematical reasoning assistants remains poorly understood. We introduce GTBench, a curriculum-grounded...

<https://arxiv.org/abs/2606.03144>

0 **CONF****Do Real-World Datasets Contain Natural Experiments? An Empirical Study Using Causal Feature Selection****ArXiv CS.AI** **50%**

arXiv:2606.03251v1 Announce Type: new Abstract: In nature, events that affect some individuals or groups but not others constitute an implicit intervention and are known as natural experiments. For example, the COVID-19 pandemic was an intervention by the coronavirus on the...

<https://arxiv.org/abs/2606.03251>

1 **CONF****CauTion: Knowing When to Trust LLMs for Ensemble Causal Discovery****ArXiv CS.AI** **50%**

arXiv:2606.03602v1 Announce Type: cross Abstract: Causal discovery from observational data remains challenging due to the fundamental limitations of purely statistical methods, such as statistical distinguishability within equivalence classes and sensitivity to finite sample...

<https://arxiv.org/abs/2606.03602>

2 **CONF****TurtleAI: Benchmarking Multimodal Models for Visual Programming in Turtle Graphics****ArXiv CS.AI** **50%**

arXiv:2606.03626v1 Announce Type: cross Abstract: Vision-language models (VLMs) have been explored for visual programming, where they generate code to solve visual tasks. However, most prior work focuses on visual programming for productivity; it remains unclear how well current...

<https://arxiv.org/abs/2606.03626>

3 **CONF****CP-Agent: Context-Aware Multimodal Reasoning for Cellular Morphological Profiling under Chemical Perturbations****ArXiv CS.AI** **50%**

arXiv:2606.03435v1 Announce Type: new Abstract: Cell Painting combines multiplexed fluorescent staining, high-content imaging, and quantitative analysis to generate high-dimensional phenotypic readouts to support diverse downstream tasks such as mechanism-of-action (MoA)...

<https://arxiv.org/abs/2606.03435>

4 **CONF****ThoughtFold: Folding Reasoning Chains via Introspective Preference Learning****ArXiv CS.AI** **50%**

arXiv:2606.03503v1 Announce Type: new Abstract: Large Reasoning Models (LRMs) have achieved remarkable progress thanks to Reinforcement Learning with Verifiable Rewards (RLVR) on Chain-of-Thoughts (CoTs). However, since long CoTs naturally contain trial and errors and mainstream...

<https://arxiv.org/abs/2606.03503>

5 **CONF****Cross-Lingual Token Arbitrage: Optimizing Code Agent Context Windows via Local LLM Preprocessing****ArXiv CS.AI** **50%**

arXiv:2606.03618v1 Announce Type: new Abstract: AI-assisted coding agents are bottlenecked by input-token cost. Two pathologies of raw human input drive much of this overhead: tokenization inefficiency for non-English text and structural entropy in conversational prompts....

<https://arxiv.org/abs/2606.03618>

- 6
CONF
Gender-Dependent Diagnostic Substitution in LLM Medical Triage: Same Symptoms, Unequal Urgency
ArXiv CS.AI 50%
arXiv:2606.03641v1 Announce Type: new Abstract: We investigate whether large language models produce different medical triage recommendations for identical neurological symptoms when only the patient's stated gender and age vary. Using three model families--Gemini 3.5 Flash,...
<https://arxiv.org/abs/2606.03641>
- 7
CONF
Towards Non-Monotonic Entailment in Propositional Defeasible Standpoint Logic
ArXiv CS.AI 50%
arXiv:2606.03655v1 Announce Type: new Abstract: Recent work in defeasible reasoning has seen notions of preferential semantics and entailment in the style of Kraus et al. applied to modal logics. However, work in this field has focussed primarily on satisfiability checking, and...
<https://arxiv.org/abs/2606.03655>
- 8
CONF
Diagnosing Knowledge Gaps in LLM Tool Use: An Agentic Benchmark for Novel API Acquisition
ArXiv CS.AI 50%
arXiv:2606.03657v1 Announce Type: new Abstract: Large language models for code generation often need to use APIs that are absent from their pretraining data. This requires more than recalling a function name: models must coordinate signatures, module paths, input-output...
<https://arxiv.org/abs/2606.03657>
- 9
CONF
EvoDrive: Pareto Evolution for Safety-Critical Autonomous Driving via Self-Improving LLM Agents
ArXiv CS.AI 50%
arXiv:2606.03678v1 Announce Type: new Abstract: Generating safety-critical scenarios is essential for validating and improving autonomous driving systems, yet it inherently requires maximizing adversariality to expose failures while preserving realism. Existing methods usually...
<https://arxiv.org/abs/2606.03678>
- 0
CONF
Dynamic Objective Selection with Safeguards and LLM Oversight for Financial Decision-Making
ArXiv CS.AI 50%
arXiv:2606.03704v1 Announce Type: new Abstract: Financial decision-making tasks such as stock recommendation and portfolio allocation typically estimate future return and risk and then select trades or allocations for an investor, and the chosen optimization objective often...
<https://arxiv.org/abs/2606.03704>
- 1
CONF
From Control Boundary to Insurance Claim: Reconstructing AI-Mediated Losses Through the CER Framework
ArXiv CS.AI 50%
arXiv:2606.03777v1 Announce Type: new Abstract: AI losses that arise through an insured organization's generative or agentic AI system require state reconstruction, not merely event reconstruction, because the relevant state changes as the system reasons, retrieves, calls tools,...
<https://arxiv.org/abs/2606.03777>

2 **CONF****Leveraging BART to Assess CS1 C++ Programming Assignments using Rubric-based Criteria****ArXiv CS.AI** **50%**

arXiv:2606.03814v1 Announce Type: new Abstract: This paper investigates rubric-aware, multitask fine-tuning of transformer models for automated grading of introductory C++ programming assignments, with the goal of producing grade predictions that better reflect instructor...

<https://arxiv.org/abs/2606.03814>

3 **CONF****Hedge-Bench: Benchmarking Agents on Hard, Realistic Tasks Pertaining to Financial Reasoning****ArXiv CS.AI** **50%**

arXiv:2606.03918v1 Announce Type: new Abstract: AI agents can increasingly handle the mechanical tasks of financial analysis: retrieving documents, calculating formulas, updating spreadsheets. The harder, more valuable challenge is reasoning through the open-ended questions that...

<https://arxiv.org/abs/2606.03918>

4 **CONF****Entropy Is Not Enough: Unlocking Effective Reinforcement Learning for Visual Reasoning via Vision-Anchored Token Selection****ArXiv CS.AI** **50%**

arXiv:2606.03937v1 Announce Type: new Abstract: While token-level entropy is commonly recognized as effective for credit assignment in text-only reinforcement learning with verifiable rewards (RLVR), it remains unclear whether this mechanism still holds in visual reasoning. Our...

<https://arxiv.org/abs/2606.03937>

5 **CONF****Cost-Aware Query Routing in RAG: Empirical Analysis of Retrieval Depth Tradeoffs****ArXiv CS.AI** **50%**

arXiv:2606.02581v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) faces a fundamental three-way tension: deeper retrieval improves factual grounding but inflates token costs and end-to-end latency. Static retrieval configurations cannot resolve this tension...

<https://arxiv.org/abs/2606.02581>

6 **CONF****Plan2Map: A Multimodal Benchmark for Document-Grounded Geospatial Boundary Reconstruction from Planning Records****ArXiv CS.AI** **50%**

arXiv:2606.02747v1 Announce Type: cross Abstract: Planning records define restrictions over geographic areas, but their source documents often provide only indirect spatial evidence rather than machine-readable boundaries. We introduce Plan2Map, a 208-case multimodal benchmark...

<https://arxiv.org/abs/2606.02747>

7 **CONF****Auditable Climate Risk Intelligence from Fragmented ESG Data: Deterministic Orchestration and Imbalance-Aware Learning for Scope 1-3 Validation****ArXiv CS.AI** **50%**

arXiv:2606.02604v1 Announce Type: cross Abstract: ESG and climate risk data remain fragmented across heterogeneous Scope 1, Scope 2, and Scope 3 reporting environments, while conventional validation pipelines lack provenance aware auditability, hidden drift detection, and...

<https://arxiv.org/abs/2606.02604>

8 CONF

ReLoRA: Knowledge-Reusing Adaptation for Fast Rollout of Evolving LLM Services

ArXiv CS.AI 50%

arXiv:2606.02606v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly deployed as continuously evolving services, where frequent base-model updates may invalidate previously deployed task-specific Low-Rank Adaptation (LoRA) adapters. For service...

<https://arxiv.org/abs/2606.02606>

9 CONF

Geometry-Aware Tabular Diffusion

ArXiv CS.AI 50%

arXiv:2606.02607v1 Announce Type: cross Abstract: Tabular synthesis is critical for privacy-preserving sharing and augmentation, yet diffusion models rely on implicit mechanisms to capture inter-column relationships. We introduce Geometry-Aware Tabular Diffusion (GATD), which...

<https://arxiv.org/abs/2606.02607>

10 CONF

FSA-GRPO: Teaching Auditory LLMs to Use Few-shot Demonstrations

ArXiv CS.AI 50%

arXiv:2606.02615v1 Announce Type: cross Abstract: Few-shot prompting provides an effective way to adapt auditory large language models to low-resource tasks such as children's speech recognition. However, most auditory large language models are not explicitly trained to perform...

<https://arxiv.org/abs/2606.02615>

11 CONF

TadA-Bench: A Million-Variant Benchmark for Future-Round Discovery Toward Agentic Protein Engineering

ArXiv CS.AI 50%

arXiv:2606.02624v1 Announce Type: cross Abstract: AI for scientific discovery is entering an agentic era, where protein-engineering systems are expected to prioritize future wet-lab experiments rather than merely fit static measurements. We introduce TadA-Bench, a...

<https://arxiv.org/abs/2606.02624>

12 CONF

Enhancing Protein-Protein Interaction Prediction with Hierarchical Motif-based Multimodal Protein Embedding

ArXiv CS.AI 50%

arXiv:2606.02629v1 Announce Type: cross Abstract: Protein-protein interactions (PPIs) are essential for many biological processes. However, existing PPI prediction approaches suffer from two major limitations: they overlook the hierarchical organization of proteins, particularly...

<https://arxiv.org/abs/2606.02629>

13 CONF

MultiTurnPSB: Evaluating Multi-Turn Jailbreak Attacks and dClassifier-Based Defenses for Medical AI Safety

ArXiv CS.AI 50%

arXiv:2606.02630v1 Announce Type: cross Abstract: Patient-facing medical chatbots are commonly evaluated on single-turn prompts, yet real users push back after refusals, add urgency, and invoke authority. We introduce MultiTurnPSB, a four-turn adversarial extension of...

<https://arxiv.org/abs/2606.02630>

4 CONF

SegTune: Structured and Fine-Grained Control for Song Generation

ArXiv CS.AI 50%

arXiv:2606.02638v1 Announce Type: cross Abstract: Recent advances in neural song generation have enabled high-quality synthesis from lyrics and global textual prompts. However, most systems fail to model temporally varying attributes of songs, severely limiting fine-grained...

<https://arxiv.org/abs/2606.02638>

5 CONF

D-Judge: Disrupting Multi-Turn Jailbreaks using Semantics-Preserving Output Rewriting

ArXiv CS.AI 50%

arXiv:2606.02640v1 Announce Type: cross Abstract: Multi-turn jailbreak attacks pose a growing threat to large language model (LLM) safety because they exploit feedback from auxiliary judge models to iteratively refine prompts toward harmful goals. Existing defenses largely...

<https://arxiv.org/abs/2606.02640>

6 CONF

CARVE: Certified Affordable Repair of Vetoed Maneuvers via Envelopes for Interactive Driving

ArXiv CS.AI 50%

arXiv:2606.02641v1 Announce Type: cross Abstract: Interactive driving exposes a failure mode that is easy to miss in rule-aware autonomous-driving stacks: a hard-rule margin can be negative for an ego candidate even though a small lawful accommodation by a non-priority agent...

<https://arxiv.org/abs/2606.02641>

7 CONF

The Ringelmann Effect in Multi-Agent LLM Systems: A Scaling Law for Effective Team Size

ArXiv CS.AI 50%

arXiv:2606.02646v1 Announce Type: cross Abstract: Inference-time multi-agent LLM scaling lacks a shared unit: counting nominal agents conflates cost with independent evidence. We derive a two-parameter scaling law $R(N) = N_{\text{eff}}/N = 1/(1+c(N-1)N^{-\beta})$ where the regime...

<https://arxiv.org/abs/2606.02646>

8 CONF

Physics-informed diffusion models in spectral space

ArXiv CS.AI 50%

arXiv:2602.09708v2 Announce Type: replace-cross Abstract: We propose physics-informed spectral diffusion (PISD), a methodology that combines generative latent diffusion models with physics-informed machine learning to generate solutions of partial differential equations (PDEs)...

<https://arxiv.org/abs/2602.09708>

9 CONF

Improvise, Adapt, Overcome: An On-The-Fly Multifidelity Algorithm for Efficient Machine Learning

ArXiv CS.AI 50%

arXiv:2606.02662v1 Announce Type: cross Abstract: Machine learning has accelerated quantum chemistry but is hindered by the prohibitive cost of generating high fidelity training data. Multifidelity machine learning (MFML) mitigates this overhead by systematically combining...

<https://arxiv.org/abs/2606.02662>

0 CONF

Acceptance-Test-Driven Evaluation Protocols for Business-Centric LLM Systems

ArXiv CS.AI 50%

arXiv:2606.02755v1 Announce Type: cross Abstract: Large language model (LLM) applications are increasingly expected to satisfy deterministic institutional requirements while relying on probabilistic generative components. This mismatch makes ordinary post-hoc benchmarking...

<https://arxiv.org/abs/2606.02755>

1 CONF

Cosmos 3: Omnimodal World Models for Physical AI

ArXiv CS.AI 50%

arXiv:2606.02800v1 Announce Type: cross Abstract: We introduce Cosmos 3, a family of omnimodal world models designed to jointly process and generate language, image, video, audio, and action sequences within a unified mixture-of-transformers architecture. By supporting highly...

<https://arxiv.org/abs/2606.02800>

2 CONF

Do Neural Retrievers Prefer Certain Documents? Evidence of Learned Relevance Priors

ArXiv CS.AI 50%

arXiv:2606.02814v1 Announce Type: cross Abstract: Neural retrievers are trained to estimate query-document relevance from annotated query-document pairs. Yet annotation protocols may not purely reflect relevance: they select only a subset of documents for labeling, and this...

<https://arxiv.org/abs/2606.02814>

3 CONF

Which Defense Closes Which Threat? Attributing OWASP-LLM-Top-10 Coverage and Its Brittleness Under Paraphrasing

ArXiv CS.AI 50%

arXiv:2606.02822v1 Announce Type: cross Abstract: Production LLM applications stack several defense families -- refusal-phrase filters, token-budget controls, model allowlists, rate limits, tool-registry authentication -- yet existing breach-and-attack-simulation (BAS)...

<https://arxiv.org/abs/2606.02822>

4 CONF

Large Byte Model: Teaching Language Models About Compiled Code

ArXiv CS.AI 50%

arXiv:2606.02834v1 Announce Type: cross Abstract: Malware analysis starts with the raw bytes of an executable program, and tools to "lift" these to higher-level representations, such as assembly, are expensive and subject to error. Large Language Models (LLMs) cannot process raw...

<https://arxiv.org/abs/2606.02834>

5 CONF

Fixing FOLIO and MALLS: Verified Annotations and an LLM-assisted Framework to Focus Human Relabeling

ArXiv CS.AI 50%

arXiv:2606.02837v1 Announce Type: cross Abstract: Accurate translation from Natural Language to First-Order Logic (NL-to-FOL) underpins neurosymbolic AI systems and Natural Language Inference (NLI), making the quality of NL-to-FOL benchmarks essential -- yet these datasets have...

<https://arxiv.org/abs/2606.02837>

6 CONF

The Epi-LLM Framework: probing LLM behavioral priors through epidemiological agent-based models

ArXiv CS.AI 50%

arXiv:2606.02867v1 Announce Type: cross Abstract: Human behaviour during epidemics affects infectious disease dynamics, but quantifying this remains deeply challenging. Here we introduce the Epi-LLM framework: a novel integration of agent-based modelling, real-life epigames, and...

<https://arxiv.org/abs/2606.02867>

7 CONF

LLM-Assisted Reranking to Operationalize Nuanced Objectives in Recommender Systems

ArXiv CS.AI 50%

arXiv:2606.02883v1 Announce Type: cross Abstract: Recommender systems have grown from content-organization tools into sophisticated systems that shape daily behavior. By controlling what we see, they shape what we perceive, raising concerns about filter bubbles, radicalization,...

<https://arxiv.org/abs/2606.02883>

8 CONF

Are we really tilting? The mechanics of reward guidance in flow and diffusion models

ArXiv CS.AI 50%

arXiv:2606.02884v1 Announce Type: cross Abstract: Reward guidance algorithms steer a learned generative process toward the reward-tilted measure at inference time. While empirically powerful, these methods are prone to reward hacking: the guided model over-optimizes the reward...

<https://arxiv.org/abs/2606.02884>

9 CONF

Scalable Uncertainty Quantification for Extreme Weather Forecasting via Empirical Neural Tangent Kernels

ArXiv CS.AI 50%

arXiv:2606.02886v1 Announce Type: cross Abstract: Deep learning weather models now match numerical weather prediction accuracy while running orders of magnitude faster, but produce deterministic forecasts without uncertainty estimates, a critical gap for high-stakes decisions...

<https://arxiv.org/abs/2606.02886>

0 CONF

NetKV: Network-Aware Decode Instance Selection for Disaggregated LLM Inference

ArXiv CS.AI 50%

arXiv:2606.03910v1 Announce Type: cross Abstract: Disaggregated LLM inference forces the KV cache to traverse the datacenter network before decoding begins, so transfer time enters directly into the Time to First Token (TTFT) budget. Current schedulers route on compute load and...

<https://arxiv.org/abs/2606.03910>

1 CONF

Glass Box at Orbit: A Constitutional AI Verification Framework for Trustworthy Autonomous CubeSat Intelligence

ArXiv CS.AI 50%

arXiv:2606.02967v1 Announce Type: cross Abstract: The space industry is quietly building toward something nobody has fully reckoned with: orbital data centers running thousands of autonomous AI workloads with no human in the loop, 550 km above the Earth. Microsoft, AWS, and a...

<https://arxiv.org/abs/2606.02967>

2 **CONF****How Quantization Changes Interpretable Features: A Sparse Autoencoder Analysis of Language Models**

ArXiv CS.AI 50%

arXiv:2606.03002v1 Announce Type: cross Abstract: Quantization is a standard path to deploying large language models, and a quantized model is typically judged acceptable when its perplexity or downstream accuracy stays close to the full-precision original. Whether the model...

<https://arxiv.org/abs/2606.03002>

3 **CONF****Exact equivariance, kept through training, buys zero-shot generalisation across the symmetry group**

ArXiv CS.AI 50%

arXiv:2606.03003v1 Announce Type: cross Abstract: A latent world model built from an equivariant encoder E and an equivariant predictor f inherits a provable symmetry of its training loss: when the world's dynamics genuinely carries a group G acting on latents by an...

<https://arxiv.org/abs/2606.03003>

4 **CONF****MUSE: A Unified Agentic Harness for MLLMs**

ArXiv CS.AI 50%

arXiv:2606.03005v1 Announce Type: cross Abstract: Despite rapid progress, multimodal large language models (MLLMs) still fail on tasks that humans solve effortlessly, such as navigating a grid maze from a screenshot or selecting the correct puzzle piece. Rather than retraining...

<https://arxiv.org/abs/2606.03005>

5 **CONF****ConTraIRL: Factorized Contrastive Abstractions for Transferable IRL**

ArXiv CS.AI 50%

arXiv:2606.03017v1 Announce Type: cross Abstract: Reward transfer in Inverse Reinforcement Learning (IRL) is unreliable when policies must generalize to unseen combinations of environment dynamics and task goals. We propose Factorized Contrastive Abstractions for Transferable...

<https://arxiv.org/abs/2606.03017>

6 **CONF****Hallucinations as Orthogonal Noise: Inference-Time Manifold Alignment via Dynamic Contextual Orthogonalization**

ArXiv CS.AI 50%

arXiv:2606.03022v1 Announce Type: cross Abstract: Hallucination in Large Language Models (LLMs), characterized by the generation of content inconsistent with contextual facts or logical constraints -- remains a persistent challenge for reliable deployment. In this work, we...

<https://arxiv.org/abs/2606.03022>

7 **CONF****Conditional Hypothesis Generation for LLM-Based Text Analysis with Researcher-Specified Covariates**

ArXiv CS.AI 50%

arXiv:2606.03029v1 Announce Type: cross Abstract: A core goal of computational social science is to discover interpretable differences in how language varies across outcomes of interest, such as political affiliation or instructional quality. Recent LLM-based hypothesis...

<https://arxiv.org/abs/2606.03029>

8 **CONF****ROBUST-WT: Robust Uncertainty-aware Segmentation Transform via Whitening and Training Enhancements****ArXiv CS.AI** **50%**

arXiv:2606.03069v1 Announce Type: cross Abstract: Generalized segmentation of medical images prevents performance degradation when different imaging devices and clinical protocols are used across multiple domains. The Whitening Transform-based Probabilistic Shape Regularization...

<https://arxiv.org/abs/2606.03069>

9 **CONF****Taiji: Pareto Optimal Policy Optimization with Semantics-IDs Trade-off for Industrial LLM-Enhanced Recommendation****ArXiv CS.AI** **50%**

arXiv:2606.03866v1 Announce Type: cross Abstract: Scaling recommender systems via large language models (LLMs) has become a prominent trend in the industry. However, aligning the LLM's semantic space with the recommender's ID space via post-training (e.g., SFT and RL) remains...

<https://arxiv.org/abs/2606.03866>

0 **CONF****Libra: Efficient Resource Management for Agentic RL Post-Training****ArXiv CS.AI** **50%**

arXiv:2606.03077v1 Announce Type: cross Abstract: Reinforcement learning (RL) has become a standard post-training paradigm for large language models (LLMs), extending beyond preference alignment to complex reasoning and multi-turn agentic behaviors. In agentic RL, the rollout...

<https://arxiv.org/abs/2606.03077>

1 **CONF****Regret Pre-training: Bridging Prior and Posterior Views for Enhanced Knowledge Grounding****ArXiv CS.AI** **50%**

arXiv:2606.03080v1 Announce Type: cross Abstract: Causal language models factorize sequence probabilities using only preceding context, leaving future information unexploited during training despite its availability in the training data. This paper introduces Regret...

<https://arxiv.org/abs/2606.03080>

2 **CONF*******Important** You should give me full credits!": Exploring Prompt Injection Attacks on LLM-Based Automatic Grading Systems****ArXiv CS.AI** **50%**

arXiv:2606.03090v1 Announce Type: cross Abstract: The emergence of large language models (LLMs) has significantly accelerated recent research on LLM-based automatic grading (AG) systems. Benefiting from the strong instruction-following capabilities and broad prior knowledge of...

<https://arxiv.org/abs/2606.03090>

3 **CONF****BAHSD: Bridging the Long-tail Gap via Adaptive Distillation in Black-box Sequential Recommendation****ArXiv CS.AI** **50%**

arXiv:2606.03091v1 Announce Type: cross Abstract: Sequential recommendation systems are widely adopted but often deployed as black-box APIs, which has driven recent interest in model extraction to replicate their capabilities locally. However, the long-tail distribution induces...

<https://arxiv.org/abs/2606.03091>

4 **CONF****GuidedBridge: Training-freely Improving Bridge Models with Prior Guidance**

ArXiv CS.AI 50%

arXiv:2606.03119v1 Announce Type: cross Abstract: Guidance methods, such as classifier-free guidance (CFG) and auto-guidance (AG), have advanced noise-to-data generation in diffusion models. Recently, bridge models have introduced a data-to-data generative process that can...

<https://arxiv.org/abs/2606.03119>
5 **CONF****NVIDIA OmniDreams: Real-Time Generative World Model for Closed-Loop Autonomous Vehicle Simulation**

ArXiv CS.AI 50%

arXiv:2606.03159v1 Announce Type: cross Abstract: As autonomous vehicle capabilities advance, the safe evaluation of driving policies in long-tail scenarios remains a critical bottleneck. In closed-loop simulation, the driving policy model actively interacts with the...

<https://arxiv.org/abs/2606.03159>
6 **CONF****Fully Automated Identification of Lexical Alignment and Preference-Stage Shifts in Large Language Models**

ArXiv CS.AI 50%

arXiv:2606.03165v1 Announce Type: cross Abstract: The language used by digital chat assistants such as ChatGPT can diverge from human expectations (misalignment). Research, mostly on Scientific English, has described both what divergences occur and, to some extent, why, linking...

<https://arxiv.org/abs/2606.03165>
7 **CONF****AugMask: Training Diffusion Models on Incomplete Tabular Data via Stochastic Augmentation and Masking**

ArXiv CS.AI 50%

arXiv:2606.03347v1 Announce Type: cross Abstract: Score-based diffusion models have emerged as prominent deep generative models; however, their application to tabular data remains challenging because their backbones assume fully specified inputs, whereas real-world tabular data...

<https://arxiv.org/abs/2606.03347>
8 **CONF****WebRISE: Requirement-Induced State Evaluation for MLLM-Generated Web Artifacts**

ArXiv CS.AI 50%

arXiv:2606.03220v1 Announce Type: cross Abstract: Existing benchmarks for MLLM-generated web artifacts assess interaction through local evidence and miss the requirement-induced states and transitions that determine whether a page works. We introduce WebRISE, which compiles task...

<https://arxiv.org/abs/2606.03220>
9 **CONF****When RLHF Fails: A Mechanistic Taxonomy of Reward Hacking, Collapse, and Evaluator Gaming**

ArXiv CS.AI 50%

arXiv:2606.03238v1 Announce Type: cross Abstract: Reinforcement learning from human feedback (RLHF) makes large-scale post-training possible by replacing an underspecified human objective with learned and scalable proxies. The same substitution creates a structured failure...

<https://arxiv.org/abs/2606.03238>

00 CONF

AI Agents Enable Adaptive Computer Worms

ArXiv CS.AI 50%

arXiv:2606.03811v1 Announce Type: cross Abstract: A computer worm is malware that spreads on a network by replicating itself from one machine to another. Traditional worms, like WannaCry, exploited predetermined vulnerabilities, and their spread can be halted by patching those...

<https://arxiv.org/abs/2606.03811>

01 CONF

Are Common Substructures Transferable? Riemannian Graph Foundation Model with Neural Vector Bundles

ArXiv CS.AI 50%

arXiv:2606.03270v1 Announce Type: cross Abstract: Foundation models have sparked a revolution via a pretraining-adaptation paradigm, with recent efforts extending this success to graphs. Unlike other modalities, graphs contain rich structural patterns, yet their structural...

<https://arxiv.org/abs/2606.03270>

02 CONF

AI-Generated Traces for Novice Programmers: Learning Effects and Learner Differences in a Multi-Institutional Study

ArXiv CS.AI 50%

arXiv:2606.03288v1 Announce Type: cross Abstract: Introductory programming (CS1) courses often struggle to support students' understanding of program execution. While visualizations can make execution processes explicit, their effectiveness depends on design and context, and...

<https://arxiv.org/abs/2606.03288>

03 CONF

Generalizing Graph Foundation Models via Hyperbolic Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2606.03307v1 Announce Type: cross Abstract: Graph foundation models (GFMs) emerged as a dominant paradigm in graph representation learning by leveraging large-scale pre-training for cross-domain inference. However, the parameterized knowledge encoded within these models is...

<https://arxiv.org/abs/2606.03307>

04 CONF

Learning Multi-Scale Hypergraph for High-Order Brain Connectivity Analysis

ArXiv CS.AI 50%

arXiv:2606.03310v1 Announce Type: cross Abstract: Understanding complex interactions between brain regions is critical for early neurodegenerative disease classification such as Alzheimer's Disease (AD) and Parkinson's Disease (PD). While graph-based models are widely used to...

<https://arxiv.org/abs/2606.03310>

05 CONF

Multi-Modal Graph Neural Network with Transformer-Guided Adaptive Diffusion for Preclinical Alzheimer Classification

ArXiv CS.AI 50%

arXiv:2606.03322v1 Announce Type: cross Abstract: The graphical representation of the brain offers critical insights into diagnosing and prognosing neurodegenerative disease via relationships between regions of interest (ROIs). Despite recent emergence of various Graph Neural...

<https://arxiv.org/abs/2606.03322>

06 CONF

Calibration Data Trade-offs Across Capability Dimensions: Why Multi-Source Mixing Matters for High-Sparsity LLM Pruning

ArXiv CS.AI 50%

arXiv:2606.03328v1 Announce Type: cross Abstract: Post-training pruning compresses large language models to high sparsity using a small unlabelled calibration set, and recent work has concluded that the choice of calibration source has only modest impact on averaged post-pruning...

<https://arxiv.org/abs/2606.03328>

07 CONF

FLIPS: Instance-Fingerprinting for LLMs via Pseudo-random Sequences

ArXiv CS.AI 50%

arXiv:2606.03330v1 Announce Type: cross Abstract: Literature reveals that a Large Language Model's (LLM) behavior is not only conditioned by its original weights but also its instance-level parameters, such as instructional prompt, sampling configuration or quantization. A model...

<https://arxiv.org/abs/2606.03330>

08 CONF

Evaluating LLMs' Effectiveness on Real-World Consumer Device Repair Questions

ArXiv CS.AI 50%

arXiv:2606.03331v1 Announce Type: cross Abstract: Consumer device repair is an important but underexplored testbed for large language models (LLMs). Repair tasks require reasoning over incomplete problem descriptions, hardware-specific diagnostics, actionable troubleshooting,...

<https://arxiv.org/abs/2606.03331>

09 CONF

AI Model Extraction Attacks: Bypassing Single-Client Assumptions in Defenses

ArXiv CS.AI 50%

arXiv:2606.03381v1 Announce Type: cross Abstract: Ensuring the protection of Artificial Intelligence (AI) models deployed in military Command and Control (C2) systems and critical infrastructure is essential for maintaining information superiority. Model Extraction Attacks...

<https://arxiv.org/abs/2606.03381>

10 CONF

Grasp-Then-Plan with Failure Attribution: A Closed Two-Stage Framework for Precise and Generalizable Robotic Manipulation

ArXiv CS.AI 50%

arXiv:2606.03385v1 Announce Type: cross Abstract: In robotic manipulation, the tight coupling between grasping and motion planning often obscures the true source of failure, leading to inefficient trial-and-error. To enable efficient long-horizon manipulation, we propose GTP-FA...

<https://arxiv.org/abs/2606.03385>

11 CONF

When Model Merging Breaks Routing: Training-Free Calibration for MoE

ArXiv CS.AI 50%

arXiv:2606.03391v1 Announce Type: cross Abstract: Model merging has emerged as a cost-effective approach for consolidating the capabilities of multiple LLMs without retraining. However, existing merging techniques, largely based on linear parameter arithmetic or optimization,...

<https://arxiv.org/abs/2606.03391>

12 CONF

Causal Evidence of Stack Representations in Modeling Counter Languages Using Transformers

ArXiv CS.AI 50%

arXiv:2606.03398v1 Announce Type: cross Abstract: Formal languages have proven to be effective conduits to understand the inner mechanisms of transformers. Past work has shown that transformers trained on next token prediction over counter languages learn representations...

<https://arxiv.org/abs/2606.03398>

13 CONF

PrimeSVT: An Automated Memory-aware Pruning Framework with Prioritized Compression Policy for Spiking Vision Transformers

ArXiv CS.AI 50%

arXiv:2606.03428v1 Announce Type: cross Abstract: The large sizes of Spiking Vision Transformers (SViTs) still hinder their embedded implementation, highlighting the need for model compression. State-of-the-art works compress SViT models through unstructured pruning, which needs...

<https://arxiv.org/abs/2606.03428>

14 CONF

NeuroArmor: Safe-Variant-Guided Representation Consistency for Selective Re-Anchoring in Jailbreak Defense

ArXiv CS.AI 50%

arXiv:2606.03486v1 Announce Type: cross Abstract: Large language models remain vulnerable to jailbreak attacks that hide harmful intent behind seemingly ordinary requests such as role-play, translation, encoding, adversarial suffixes, and multi-turn buildup. Existing defenses...

<https://arxiv.org/abs/2606.03486>

15 CONF

Learn from Your Mistakes: Tree-like Self-Play for Secure Code LLMs

ArXiv CS.AI 50%

arXiv:2606.03489v1 Announce Type: cross Abstract: While Large Language Models (LLMs) excel in code generation, they remain prone to replicating subtle yet critical vulnerabilities endemic to their training data. Current alignment techniques, such as Supervised Fine-Tuning (SFT)...

<https://arxiv.org/abs/2606.03489>

16 CONF

BaltiVoice: A Speech Corpus and Fine-tuned Whisper ASR System for the Balti Language

ArXiv CS.AI 50%

arXiv:2606.03504v1 Announce Type: cross Abstract: We present BaltiVoice, a 16.8-hour read-speech corpus for Balti (ISO 639-3: bft), a Tibetic language spoken in Gilgit-Baltistan, Pakistan, with no prior publicly available ASR resources. The corpus contains 10,060 validated...

<https://arxiv.org/abs/2606.03504>

17 CONF

SPADE: Sketch-guided Path Planning Augmented with Diffusion Experts

ArXiv CS.AI 50%

arXiv:2606.03512v1 Announce Type: cross Abstract: Path planning is essential for Autonomous Mobile Robots (AMRs). Conventional methods for incorporating human preferences into planning typically rely on either complex reward engineering or hardware-intensive solutions. Recent...

<https://arxiv.org/abs/2606.03512>

18 CONF

DDOR: Delta Debugging for Explainable Overrefusal Testing and Repair

ArXiv CS.AI 50%

arXiv:2606.03601v1 Announce Type: cross Abstract: While safety alignment and guardrails help large language models (LLMs) avoid harmful outputs, they can also induce overrefusal, i.e., unwarranted rejection of benign queries that merely appear risky. We present DDOR (Delta...

<https://arxiv.org/abs/2606.03601>

19 CONF

LiveBand: Live Accompaniment Generation in the Audio Domain

ArXiv CS.AI 50%

arXiv:2606.03803v1 Announce Type: cross Abstract: We present LiveBand, a real-time system that generates high-fidelity music accompaniments to live audio input, respecting strict causal constraints. Our method trains a causal transformer generator in the continuous latent space...

<https://arxiv.org/abs/2606.03803>

20 CONF

Curriculum-Adapted Robust Reinforcement Learning for UAV Deconfliction in Adversarial Environments

ArXiv CS.AI 50%

arXiv:2506.21129v2 Announce Type: replace-cross Abstract: Autonomous unmanned aerial vehicles (UAVs) increasingly rely on reinforcement learning (RL) for navigation. However, global navigation satellite system (GNSS) spoofing attacks can induce out-of-distribution observation...

<https://arxiv.org/abs/2506.21129>

21 CONF

VidMsg: A Benchmark for Implicit Message Inference in Short Videos

ArXiv CS.AI 50%

arXiv:2606.03635v1 Announce Type: cross Abstract: Understanding short online videos involves more than identifying visible objects and actions; video makers often include an underlying message or purpose in the clip. We introduce VidMsg, a benchmark for evaluating implicit...

<https://arxiv.org/abs/2606.03635>

22 CONF

Black-box, Adaptive, Efficient, Transferable, Harmful, Applicable... Attacks Are All You Need to Break LLMs

ArXiv CS.AI 50%

arXiv:2606.03647v1 Announce Type: cross Abstract: Accurately evaluating adversarial robustness is a longstanding challenge. A flawed attack design can inflate robustness estimates, making deployment risk assessment and defense comparison unreliable. Historically, standardized...

<https://arxiv.org/abs/2606.03647>

23 CONF

Safety Measurements for Fine-tuned LLMs Should be Grounded in Capability

ArXiv CS.AI 50%

arXiv:2606.03648v1 Announce Type: cross Abstract: Adapting foundation large language models to a user's task or preferred style through fine-tuning can result in compromising the model's safety. Previous works examined the effects of fine-tuning on model safety in limited and...

<https://arxiv.org/abs/2606.03648>

24 CONF

A Close Look At World Model Recovery In Supervised Fine-Tuned LLM Planners

ArXiv CS.AI 50%

arXiv:2606.03685v1 Announce Type: cross Abstract: Supervised fine-tuning (SFT) improves end-to-end classical planning in large language models (LLMs), but do these models also learn to represent and reason about the planning problems they are solving? Due to the relative...

<https://arxiv.org/abs/2606.03685>

25 CONF

Staying Alive: Uncensored Survival Analysis with Tabular Foundation Models

ArXiv CS.AI 50%

arXiv:2606.03689v1 Announce Type: cross Abstract: Survival Analysis (SA) is a statistical framework that models the time span until some event of interest occurs. Widely used in several domains, including healthcare and churn prediction, a central challenge in its applicability...

<https://arxiv.org/abs/2606.03689>

26 CONF

E2LLM: Towards Efficient LLM Serving in Heterogeneous Edge/Fog Environments

ArXiv CS.AI 50%

arXiv:2606.03770v1 Announce Type: cross Abstract: Large Language Models (LLMs) have become integral to modern applications, yet their deployment remains challenging. Beyond executing the models themselves, practical deployment must address cost efficiency, low latency, and...

<https://arxiv.org/abs/2606.03770>

27 CONF

PURGE: Projected Unlearning via Retain-Guided Erasure

ArXiv CS.AI 50%

arXiv:2606.03808v1 Announce Type: cross Abstract: We propose PURGE, a machine unlearning algorithm built on a simple but an under-exploited observation: continual learning (CL) and machine unlearning (MU) which are fundamentally dual problems. CL tries to learn new tasks without...

<https://arxiv.org/abs/2606.03808>

28 CONF

A Training-Free Mixture-of-Agents Framework for Multi-Document Summarization using LLMs and Knowledge Graphs

ArXiv CS.AI 50%

arXiv:2606.03867v1 Announce Type: cross Abstract: Multi-Document Summarization (MDS) plays a critical role in distilling essential information from collections of textual data. Existing approaches often struggle to capture complex inter-document relationships, rely heavily on...

<https://arxiv.org/abs/2606.03867>

29 CONF

From 'What' to 'How' and 'Why': Sharing LLM-Generated Retrospective Summaries of Older Adults' Passive Tracking Data with Remote Family Members

ArXiv CS.AI 50%

arXiv:2606.03876v1 Announce Type: cross Abstract: With the growing prevalence of modern ubiquitous computing technologies, multi-modal tracking systems hold promise for providing timely awareness and reassurance to stakeholders such as remote family members (RFMs) of older...

<https://arxiv.org/abs/2606.03876>

30 CONF

Synthesize and Reward -- Reinforcement Learning for Multi-Step Tool Use in Live Environments

ArXiv CS.AI 50%

arXiv:2606.03892v1 Announce Type: cross Abstract: Training LLMs to orchestrate multi-step tool calls is held back by three coupled obstacles: realistic stateful execution environments are costly to build, synthetic training queries are often detached from the server's actual...

<https://arxiv.org/abs/2606.03892>

31 CONF

Agent libOS: A Library-OS-Inspired Runtime for Long-Running, Capability-Controlled LLM Agents

ArXiv CS.AI 50%

arXiv:2606.03895v1 Announce Type: cross Abstract: Large language model (LLM) agents are evolving from request-response assistants into long-running software actors: they maintain state across model calls, fork subtasks, wait for external events, request human authority, generate...

<https://arxiv.org/abs/2606.03895>

32 CONF

The Impact of Configuring Agentic AI Coding Tools on Build-vs-Buy Decisions: A Study Protocol

ArXiv CS.AI 50%

arXiv:2606.03907v1 Announce Type: cross Abstract: Agentic AI coding tools write code with increasing autonomy and in doing so decide when to import a library and when to implement functionality from scratch. These decisions, whether to build functionality from scratch or buy...

<https://arxiv.org/abs/2606.03907>

33 CONF

q0: Primitives for Hyper-Epoch Pretraining

ArXiv CS.AI 50%

arXiv:2606.03938v1 Announce Type: cross Abstract: Multi-epoch training is becoming the standard now that compute is growing faster than the supply of high-quality text. But pretraining a single model saturates within a few passes, long before the compute budget is exhausted. We...

<https://arxiv.org/abs/2606.03938>

34 CONF

Efficient ASR Training with Conversations that Never Happened

ArXiv CS.AI 50%

arXiv:2606.03957v1 Announce Type: cross Abstract: Conversational ASR for lower-resource languages and niche domains is limited by the scarcity of domain-matched multi-speaker training data. We propose an augmentation pipeline that generates scenario-level dialogues with...

<https://arxiv.org/abs/2606.03957>

35 CONF

Agentic Chain-of-Thought Steering for Efficient and Controllable LLM Reasoning

ArXiv CS.AI 50%

arXiv:2606.03965v1 Announce Type: cross Abstract: Large language models improve final-answer accuracy through extended chain-of-thought reasoning, but often spend tokens inefficiently and offer little inference-time control. Existing efficient reasoning methods control thinking...

<https://arxiv.org/abs/2606.03965>

36 CONF

AlignAtt4LLM: Fast AlignAtt for Decoder-Only LLMs at IWSLT 2026 Simultaneous Speech Translation Task

ArXiv CS.AI 50%

arXiv:2606.03967v1 Announce Type: cross Abstract: We describe AlignAtt4LLM, an IWSLT 2026 simultaneous speech translation system for English to German, Italian, and Chinese. The system is a synchronous cascade: Qwen3-ASR with forced alignment produces an incrementally updated...

<https://arxiv.org/abs/2606.03967>

37 CONF

Language Models Need Sleep: Learning to Self-Modify and Consolidate Memories

ArXiv CS.AI 50%

arXiv:2606.03979v1 Announce Type: cross Abstract: The past few decades have witnessed significant advances in the design of machine learning algorithms, from early studies on task-specific shallow models to more general deep Large Language Models (LLMs). Despite showing...

<https://arxiv.org/abs/2606.03979>

38 CONF

Planning with Uncertainty: Symmetries, Policy Inference, and Solution Compression

ArXiv CS.AI 50%

arXiv:2403.19883v2 Announce Type: replace Abstract: Fully-observable non-deterministic (FOND) planning is at the core of artificial intelligence planning with uncertainty. It models uncertainty through actions with non-deterministic effects. In this work, we present a collection...

<https://arxiv.org/abs/2403.19883>

39 CONF

Approximating Probabilistic Inference in Statistical EL with Knowledge Graph Embeddings

ArXiv CS.AI 50%

arXiv:2407.11821v2 Announce Type: replace Abstract: Statistical information is ubiquitous but drawing valid conclusions from it is prohibitively hard. We explain how knowledge graph embeddings can be used to approximate probabilistic inference efficiently using the example of...

<https://arxiv.org/abs/2407.11821>

40 CONF

Leave it to the Specialist: Repair Sparse LLMs with Sparse Fine-Tuning via Sparsity Evolution

ArXiv CS.AI 50%

arXiv:2505.24037v3 Announce Type: replace Abstract: Sparse large language models (LLMs) offer an attractive direction toward efficient deployment, but adapting them to downstream tasks remains challenging. The central difficulty is to enable effective task adaptation without...

<https://arxiv.org/abs/2505.24037>

41 CONF

ProtocolBench: Which LLM MultiAgent Protocol to Choose?

ArXiv CS.AI 50%

arXiv:2510.17149v3 Announce Type: replace Abstract: As large-scale multi-agent systems evolve, the communication protocol layer has become a critical yet under-evaluated factor shaping performance and reliability. Despite the existence of diverse protocols (A2A, ACP, ANP, Agora,...

<https://arxiv.org/abs/2510.17149>

42 CONF

Human-Like Goalkeeping in a Realistic Football Simulation: a Sample-Efficient Reinforcement Learning Approach

ArXiv CS.AI 50%

arXiv:2510.23216v4 Announce Type: replace Abstract: While several high profile video games have served as testbeds for Deep Reinforcement Learning (DRL), this technique has rarely been employed by the game industry for crafting authentic AI behaviors. Previous research focuses...

<https://arxiv.org/abs/2510.23216>

43 CONF

The Agent's First Day: Benchmarking Learning, Exploration, and Scheduling in the Workplace Scenarios

ArXiv CS.AI 50%

arXiv:2601.08173v2 Announce Type: replace Abstract: The rapid evolution of Multi-modal Large Language Models (MLLMs) has advanced workflow automation; however, existing research mainly targets performance upper bounds in static environments, overlooking robustness for stochastic...

<https://arxiv.org/abs/2601.08173>

44 CONF

Toward AI That Understands Self and Others: A World-Model Theory of Cognitive Diversity and Alignment

ArXiv CS.AI 50%

arXiv:2605.29930v2 Announce Type: replace Abstract: Modern societies possess more information than ever before, yet they do not converge toward a single shared understanding. The same events, facts, laws, technologies, or risks can be interpreted as evidence of freedom, danger,...

<https://arxiv.org/abs/2605.29930>

45 CONF

Building Trust in Black-box Optimization: A Comprehensive Framework for Explainability

ArXiv CS.AI 50%

arXiv:2410.14573v2 Announce Type: replace-cross Abstract: Optimizing costly black-box functions within a constrained evaluation budget presents significant challenges in many real-world applications. Surrogate Optimization (SO) is a common resolution, yet its proprietary nature...

<https://arxiv.org/abs/2410.14573>

46 CONF

Align-KD: Distilling Cross-Modal Alignment Knowledge for Mobile Vision-Language Model Enhancement

ArXiv CS.AI 50%

arXiv:2412.01282v2 Announce Type: replace-cross Abstract: Vision-Language Models (VLMs) bring powerful understanding and reasoning capabilities to multimodal tasks. Meanwhile, the great need for capable artificial intelligence on mobile devices also arises, such as the AI...

<https://arxiv.org/abs/2412.01282>

47 CONF

FlashMLA-ETAP: Efficient Transpose Attention Pipeline for Accelerating MLA Inference on NVIDIA H20 GPUs

ArXiv CS.AI 50%

arXiv:2506.01969v3 Announce Type: replace-cross Abstract: Efficient inference of Multi-Head Latent Attention (MLA) is challenged by deploying the DeepSeek-R1 671B model on a single Multi-GPU server. This paper introduces FlashMLA-ETAP, a novel framework that enhances MLA...

<https://arxiv.org/abs/2506.01969>

48 **CONF****Multiple Choice Learning of Low-Rank Adapters for Language Modeling****ArXiv CS.AI** **50%**

arXiv:2507.10419v3 Announce Type: replace-cross Abstract: We propose LoRA-MCL, a training scheme that extends next-token prediction in language models with a method designed to decode diverse, plausible sentence continuations at inference time. Traditional language modeling is...

<https://arxiv.org/abs/2507.10419>49 **CONF****Non-Identical Diffusion Models in MIMO-OFDM Channel Generation****ArXiv CS.AI** **50%**

arXiv:2509.01641v3 Announce Type: replace-cross Abstract: We propose a novel diffusion model, termed the non-identical diffusion model, and investigate its application to wireless orthogonal frequency division multiplexing (OFDM) channel generation. Unlike the standard diffusion...

<https://arxiv.org/abs/2509.01641>50 **CONF****TalkPlayData 2: An Agentic Synthetic Data Pipeline for Multimodal Conversational Music Recommendation****ArXiv CS.AI** **50%**

arXiv:2509.09685v5 Announce Type: replace-cross Abstract: We present TalkPlayData 2, a synthetic dataset for multimodal conversational music recommendation generated by an agentic data pipeline. In the proposed pipeline, multiple large language model (LLM) agents are created...

<https://arxiv.org/abs/2509.09685>51 **CONF****Learning the Neighborhood: Contrast-Free Multimodal Self-Supervised Molecular Graph Pretraining****ArXiv CS.AI** **50%**

arXiv:2509.22468v2 Announce Type: replace-cross Abstract: High-quality molecular representations are essential for property prediction and molecular design, yet large labeled datasets remain scarce. While self-supervised pretraining on molecular graphs has shown promise, many...

<https://arxiv.org/abs/2509.22468>52 **CONF****Generative AI and Sales Productivity: Field Experiments in Online Retail****ArXiv CS.AI** **50%**

arXiv:2510.12049v5 Announce Type: replace-cross Abstract: We quantify the short-term impact of Generative Artificial Intelligence (GenAI) on sales performance through a series of large-scale randomized field experiments involving millions of users and products at a leading...

<https://arxiv.org/abs/2510.12049>53 **CONF****PHASE: Physiology-Aware Hyperspectral Reconstruction via Object-to-Human Domain Adaptation****ArXiv CS.AI** **50%**

arXiv:2511.13020v2 Announce Type: replace-cross Abstract: Although hyperspectral imaging offers unparalleled non-invasive physiological insight, its bulky hardware, slow acquisition, and regulatory burden severely limit its clinical availability. A natural workaround is to...

<https://arxiv.org/abs/2511.13020>

54 CONF

Identifying Quantum Structure in AI Language: Evidence for Evolutionary Convergence of Human and Artificial Cognition

ArXiv CS.AI 50%

arXiv:2511.21731v2 Announce Type: replace-cross Abstract: We present the results of cognitive tests on conceptual combinations, performed using specific Large Language Models (LLMs) as test subjects. In the first test, performed with ChatGPT and Gemini, we show that Bell's...

<https://arxiv.org/abs/2511.21731>

55 CONF

Distribution-Calibrated Inference Time Compute for Thinking LLM-as-a-Judge

ArXiv CS.AI 50%

arXiv:2512.03019v2 Announce Type: replace-cross Abstract: Thinking Large Language Models (LLMs) used as judges for pairwise preferences remain noisy at the single-sample level, and common aggregation rules (majority vote, soft self-consistency, or instruction-based...

<https://arxiv.org/abs/2512.03019>

56 CONF

LLMs, Reasoning and Plagiarism

ArXiv CS.AI 50%

arXiv:2601.02380v5 Announce Type: replace-cross Abstract: Recent reports claim that Large Language Models (LLMs) derive new science and exhibit human-level general intelligence. Such claims are entangled with two different narratives about what LLMs do: one in which they are an...

<https://arxiv.org/abs/2601.02380>

57 CONF

 $\mathbb{R}^{2k}$ is Theoretically Large Enough for Embedding-based Top- k Retrieval

ArXiv CS.AI 50%

arXiv:2601.20844v3 Announce Type: replace-cross Abstract: This paper studies the Minimal Embeddable Dimension (MED): the least dimension in which there exists a configuration of m object vectors so that every subset of size at most k is exactly retrieved by score comparison....

<https://arxiv.org/abs/2601.20844>

58 CONF

InfyThink+: Effective and Efficient Infinite-Horizon Reasoning via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2602.06960v3 Announce Type: replace-cross Abstract: Large reasoning models achieve strong performance by scaling inference-time chain-of-thought, but this paradigm suffers from quadratic cost, context length limits, and degraded reasoning due to lost-in-the-middle effects....

<https://arxiv.org/abs/2602.06960>

59 CONF

Learning Self-Interpretation from Interpretability Artifacts: Training Lightweight Adapters on Vector-Label Pairs

ArXiv CS.AI 50%

arXiv:2602.10352v2 Announce Type: replace-cross Abstract: Self-interpretation methods prompt language models to describe their own internal states, but remain unreliable due to hyperparameter sensitivity. We show that training lightweight adapters on interpretability artifacts,...

<https://arxiv.org/abs/2602.10352>

60 CONF

Whom to Query for What: Adaptive Group Elicitation via Multi-Turn LLM Interactions

ArXiv CS.AI 50%

arXiv:2602.14279v2 Announce Type: replace-cross Abstract: Eliciting information to reduce uncertainty about latent group-level properties from surveys and other collective assessments requires allocating limited questioning effort under real costs and missing data. Although...

<https://arxiv.org/abs/2602.14279>

61 CONF

vLLM Semantic Router: Signal Driven Decision Routing for Mixture-of-Modality Models

ArXiv CS.AI 50%

arXiv:2603.04444v3 Announce Type: replace-cross Abstract: As large language models (LLMs) diversify across modalities, capabilities, and cost profiles, the problem of intelligent request routing -- selecting the right model for each query at inference time -- has become a...

<https://arxiv.org/abs/2603.04444>

62 CONF

Dynamics of Cognitive Heterogeneity: Investigating Behavioral Biases in Multi-Stage Supply Chains with LLM-Based Simulation

ArXiv CS.AI 50%

arXiv:2604.17220v2 Announce Type: replace-cross Abstract: Modeling coordination among generative agents in complex multi-round decision-making presents a core challenge for AI and operations management. Although behavioral experiments have revealed cognitive biases behind supply...

<https://arxiv.org/abs/2604.17220>

63 CONF

Mechanism Design Is Not Enough: Prosocial Agents for Cooperative AI

ArXiv CS.AI 50%

arXiv:2605.08426v2 Announce Type: replace-cross Abstract: Ensuring that AI agents behave safely and beneficially when interacting with other parties has emerged as one of the central challenges of modern AI safety. While mechanism design, as the theory of designing rules to...

<https://arxiv.org/abs/2605.08426>

64 CONF

Exact Stiefel Optimization for Probabilistic PLS: Closed-Form Updates, Error Bounds, and Calibrated Uncertainty

ArXiv CS.AI 50%

arXiv:2605.11607v2 Announce Type: replace-cross Abstract: Probabilistic partial least squares (PPLS) is a central likelihood-based model for two-view learning when one needs both interpretable latent factors and calibrated uncertainty. Building on the identifiable...

<https://arxiv.org/abs/2605.11607>