



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	122	Media
TechCrunch AI	8	Media
The Verge AI	6	Media
HackerNews	5	Community
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Google must let publishers opt out of AI Search features, rules UK

The Verge AI +1 source

80%

Online publishers are getting more control over whether their websites appear in Google's AI Search features, thanks to a UK regulatory ruling. The new conduct rule imposed by the Competition and Markets Authority (CMA) requires Google to let website owners keep their content...

<https://www.theverge.com/tech/942302/google-search-ai-overviews-uk-cma-publisher-opt-out>

VERI

Microsoft's first advanced reasoning AI is here

The Verge AI +1 source

80%

Microsoft announced a bunch of new in-house AI models at Build 2026, including a new "flagship" model: MAI-Thinking-1. It's an ambitious step into model development for Microsoft, which introduced its initial in-house models last year - before then, it had relied on OpenAI's...

<https://www.theverge.com/tech/941664/microsoft-ai-model-reasoning-mai-thinking-1-build-...>

CONFIRMED SIGNALS — 138 stories

CONF

Show HN: RePlaya - self-hosted browser session replay with live tailing

HackerNews 50% HN 49

Hi HN, I'm one of the founders of s2.dev. RePlaya (<https://github.com/s2-streamstore/replaya>) is a self-hosted browser session replay tool using rrweb (<https://github.com/rrweb-io/rrweb>). It occurred to me that a durable stream per session would be a much neater...

<https://github.com/s2-streamstore/replaya>

CONF

Show HN: Build Your Own AI Agent CLI in 150 Lines

HackerNews 50% HN 29

I can't tell if HN is the right kind of place for this stuff anymore since people are so advanced in their use but I thought it was interesting to leverage my existing Go microservices framework and turn it into the core of what would provide tools for an agent cli or whatever...

<https://go-micro.dev/blog/11>

CONF

Show HN: an AI app that settles small couple arguments

HackerNews 50% HN 4

Nisrine and I have been working in the same company for many years (she's Head of Product, I am CPTO) and we have done great things for our employer(s) all along. But we always wondered if we could launch our own app - after all, we know the process of it all, we just "never had..."

<https://thepiece.app/en>

CONF

Show HN: Terraform RAG - index modules, distill conventions, compose via MCP

HackerNews 50% HN 4

No summary — see link for full article.

<https://terraform-rag.io>

CONF

Show HN: Code3DBench - single-image to executable Three.js benchmark

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/VladimirGI/Code3DBench>

CONF

Alphabet's record-breaking \$85B raise for Google's AI business is a helluva good signal

TechCrunch AI 50%

If Alphabet's record-breaking \$85 billion stock sale signals investor appetite for AI-related offerings, we can see that investors are ready to chow.

<https://techcrunch.com/2026/06/03/alphabets-record-breaking-85b-raise-for-googles-ai-bu...>

CONF

Google's Dreambeans, its weirdest-named AI tool to date, will turn your life into a cartoon

TechCrunch AI 50%

Dreambeans is a curated list of AI-illustrated "stories" culled from the personal data in your Google account.

<https://techcrunch.com/2026/06/03/googles-dreambeans-its-weirdest-named-ai-tool-to-date...>

0

CONF

Amazon will show AI product images when you search for some reason

TechCrunch AI 50%

Amazon will use visual search and AI to show AI-generated product images that match your search queries. The retailer says it will help guide users to products.

<https://techcrunch.com/2026/06/03/amazon-will-show-ai-product-images-when-you-search-fo...>

1

CONF

These two founders left Goldman and Meta to build voice AI for markets everyone else overlooked

TechCrunch AI 50%

The startup's own stack for Africa and Middle East is now handling more than 17,000 calls per day.

<https://techcrunch.com/2026/06/03/these-two-founders-left-goldman-and-meta-to-build-voi...>

2

CONF

Meta's AI agent for WhatsApp Business is now available globally

TechCrunch AI 50%

WhatsApp will charge businesses for using its AI agent based on token usage.

<https://techcrunch.com/2026/06/03/metass-ai-agent-for-whatsapp-business-is-now-available...>

3

CONF

Coralogix raises \$200M on bet that someone needs to watch the AI agents

TechCrunch AI 50%

Coralogix is among a growing number of infrastructure firms betting that as AI systems move into production, demand will rise for tools that can monitor their behavior, troubleshoot failures, and provide the operational data needed to keep them running reliably.

<https://techcrunch.com/2026/06/03/coralogix-raises-200m-in-race-to-build-the-monitoring...>

4

CONF

New Microsoft tool lets devs spin up AI behavior tests using text descriptions

TechCrunch AI 50%

Microsoft on Tuesday took the wraps off Adaptive Spec-driven Scoring for Evaluation and Regression Testing, an open source framework for spinning up AI evaluations.

<https://techcrunch.com/2026/06/02/new-microsoft-tool-lets-devs-spin-up-ai-behavior-test...>

5

CONF

The Download: Trump's new AI order, and smart glasses for warfare

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. 5 key points in Trump's new AI order Less than two weeks after scrapping an executive order on AI, President Donald Trump signed a new one on...

<https://www.technologyreview.com/2026/06/03/1138322/the-download-trump-ai-order-smart-g...>

6 CONF

As AI gets better, it reveals an empty promise

The Verge AI 50%

This week we've got tandem hands-ons with Google's new Gemini AI agent - Spark - from my colleagues David Pierce and Jay Peters. Their takeaways are similar: It's so effective that it's scary. Spark knew that David's dog is named Frida and knew the first name of Jay's wife, even...

<https://www.theverge.com/ai-artificial-intelligence/942629/as-ai-gets-better-it-reveals...>

7 CONF

Amazon's search bar will invent AI-generated products you can't buy

The Verge AI 50%

Amazon's updated search bar will now show you AI-generated images of products as you describe them. For now, the in-app feature only surfaces AI images of clothing and home goods, allowing you to tap on the image that best matches what you're looking for and search for...

<https://www.theverge.com/tech/942547/amazon-search-bar-ai-images>

8 CONF

Microsoft and OpenAI broke up — now they're ready to fight

The Verge AI 50%

At Microsoft's annual Build conference on Tuesday, the company announced a slew of new or expanded AI initiatives, including a super app, in-house reasoning models, a cybersecurity tool, and OpenClaw-esque AI agents. All this news added up to a clear message: Microsoft is...

<https://www.theverge.com/ai-artificial-intelligence/942242/microsoft-build-ai-agents-op...>

9 CONF

AI has a water problem — Google thinks it has a fix

The Verge AI 50%

In the face of widespread backlash to the AI data center buildout throughout the US, Google is touting its efforts to minimize the environmental impact by actually increasing water for local communities. The company laid out five commitments around water use in a new blog post...

<https://www.theverge.com/policy/942296/google-water-commitments-data-centers>

10 CONF

Toward Pre-Deployment Assurance for Enterprise AI Agents: Ontology-Grounded Simulation and Trust Certification

ArXiv CS.AI 50%

arXiv:2606.04037v1 Announce Type: new Abstract: Pre-deployment verification of enterprise artificial intelligence (AI) agents remains a critical gap between large language model (LLM) capability benchmarking and production deployment. Post-deployment monitoring,...

<https://arxiv.org/abs/2606.04037>

1 CONF

Stumbling Into AI Emotional Dependence: How Routine AI Interactions Reshape Human Connection

ArXiv CS.AI 50%

arXiv:2606.04150v1 Announce Type: new Abstract: Public discourse and emerging policy typically assume that AI emotional support is a deliberate act: a lonely user consciously seeking comfort from a dedicated companion chatbot. In this paper, we draw on emerging empirical...

<https://arxiv.org/abs/2606.04150>

2 **CONF****Thinking Through Signs: PEEL as a Semiotic Scaffolding for Epistemically Accountable AI-Enabled Research**

ArXiv CS.AI 50%

arXiv:2606.04152v1 Announce Type: new Abstract: Large language models are reshaping research practice while quietly eroding researchers epistemic accountability. This commentary introduces PEEL - Protocols for Epistemically Engaged Literacy in AI, a working scaffolding that...

<https://arxiv.org/abs/2606.04152>

3 **CONF****Imaginative Perception Tokens Enhance Spatial Reasoning in Multimodal Language Models**

ArXiv CS.AI 50%

arXiv:2606.03988v2 Announce Type: replace Abstract: Vision language models (VLMs) excel at many tasks but still struggle with spatial reasoning when critical information is not directly observable. Many such problems require imaginative perception: inferring what would be seen...

<https://arxiv.org/abs/2606.03988>

4 **CONF****VAMPS: Visual-Assisted Mathematical Problem Solving Benchmark**

ArXiv CS.AI 50%

arXiv:2606.04244v1 Announce Type: new Abstract: Multimodal large language models are increasingly capable of complex reasoning, yet their performance often degrades when they must externalize a problem through a tool and then reason over the tool's output, specifically when they...

<https://arxiv.org/abs/2606.04244>

5 **CONF****StepPRM-RTL: Stepwise Process-Reward Guided LLM Fine-Tuning for Enhanced RTL Synthesis**

ArXiv CS.AI 50%

arXiv:2606.04246v1 Announce Type: new Abstract: Automatic generation of RTL code for digital hardware designs remains challenging due to long-horizon reasoning, multi-step dependencies, and strict correctness constraints in Verilog and VHDL. We present StepPRM-RTL, a novel...

<https://arxiv.org/abs/2606.04246>

6 **CONF****Characterizing initial human-AI proof formalization workflows**

ArXiv CS.AI 50%

arXiv:2606.04273v1 Announce Type: new Abstract: For centuries, human mathematicians have written proofs to substantiate their mathematical arguments; yet, the ability to automatically verify the validity of proofs has long been a challenge. Advances in AI systems' ability to...

<https://arxiv.org/abs/2606.04273>

7 **CONF****The Saturation Trap and the Subjectivity of Intervention Timing: Why Affect-Based Triggers and LLM Judges Fail to Time Interventions on Autonomous Agents**

ArXiv CS.AI 50%

arXiv:2606.04296v1 Announce Type: new Abstract: As autonomous AI agents move from conversational systems to long-horizon software execution, runtime safety layers that decide when to interrupt an agent have become essential. We study this timing problem using a continuous...

<https://arxiv.org/abs/2606.04296>

8 CONF

The Digital Apprentice: A Framework for Human-Directed Agentic AI Development

ArXiv CS.AI 50%

arXiv:2606.04321v1 Announce Type: new Abstract: Agentic AI deployments face a recurring design tension: heavy human oversight limits scale, while broad autonomy outruns accountability. Neither posture provides the governance infrastructure required for responsible delegation. We...

<https://arxiv.org/abs/2606.04321>

9 CONF

Simulate, Reason, Decide: Scientific Reasoning with LLMs for Simulation-Driven Decision Making

ArXiv CS.AI 50%

arXiv:2606.04505v1 Announce Type: new Abstract: Scientific simulators are increasingly being integrated into LLM-driven systems for high-stakes simulation-driven decision-making. However, existing frameworks primarily use LLMs to generate, calibrate, or execute simulators,...

<https://arxiv.org/abs/2606.04505>

0 CONF

Neetyabhas: A Framework for Uncertainty-Aware Public Policy Optimization in Rational Agent-Based Models

ArXiv CS.AI 50%

arXiv:2606.04562v1 Announce Type: new Abstract: Purpose The WHO's COVID-19 non-pharmaceutical interventions (e.g., lockdowns, vaccinations) effectively curb transmission but impose heavy economic strains. Existing research often neglects individual behaviors and falsely assumes...

<https://arxiv.org/abs/2606.04562>

1 CONF

Plan First, Judge Later, Run Better: A DMAIC-Inspired Agentic System for Industrial Anomaly Detection

ArXiv CS.AI 50%

arXiv:2606.04599v1 Announce Type: new Abstract: Large language model (LLM) agents have shown promise in automating complex data-analysis workflows, but their reliable deployment remains challenging in high-stakes industrial scenarios. Industrial anomaly detection (IAD) is...

<https://arxiv.org/abs/2606.04599>

2 CONF

MIRAGE: Mobile Agents with Implicit Reasoning and Generative World Models

ArXiv CS.AI 50%

arXiv:2606.04627v1 Announce Type: new Abstract: Mobile agents are increasingly expected to operate everyday applications from screenshots and language goals, where reliable control requires reasoning over screen affordances, multi-step navigation, and future state changes....

<https://arxiv.org/abs/2606.04627>

3 CONF

Fog of Love: Engineering Virtuous Agent Behavior with Affinity-based Reinforcement Learning in a Game Environment

ArXiv CS.AI 50%

arXiv:2606.04750v1 Announce Type: new Abstract: Instilling virtuous behavior in artificial intelligence has seen increasing interest. One of the techniques proposed is known as affinity-based reinforcement learning, which uses policy regularization on the objective function to...

<https://arxiv.org/abs/2606.04750>

- 4 **CONF**
- From Answers to States: Verifiable Process-Level Evaluation of Chemical Reasoning in Large Language Models**
- ArXiv CS.AI 50%
- arXiv:2606.03660v2 Announce Type: replace Abstract: Large language models are increasingly used as chemistry assistants, yet most chemistry benchmarks still score only final answers. This masks a critical failure mode: a model may output the correct molecule, product, or option...
- <https://arxiv.org/abs/2606.03660>
-
- 5 **CONF**
- Inference-Time Vulnerability Beyond Shallow Safety: Alignment Along Generation Trajectories**
- ArXiv CS.AI 50%
- arXiv:2606.04778v1 Announce Type: new Abstract: Safety-aligned Large Language Models (LLMs) remain vulnerable to interventions during inference that redirect generation toward harmful outputs. Recent work attributes this to shallow safety, where alignment concentrates in the...
- <https://arxiv.org/abs/2606.04778>
-
- 6 **CONF**
- Beyond Objective Equivalence: Constraint Injection for LLM-Based Optimization Modeling on Vehicle Routing Problems**
- ArXiv CS.AI 50%
- arXiv:2606.04816v1 Announce Type: new Abstract: Large language models (LLMs) increasingly translate natural-language optimization problems into executable solver code. Yet for constraint-dense operations research (OR) problems, existing data-filtering and training pipelines...
- <https://arxiv.org/abs/2606.04816>
-
- 7 **CONF**
- R-APS: Compositional Reasoning and In-Context Meta-Learning for Constrained Design via Reflective Adversarial Pareto Search**
- ArXiv CS.AI 50%
- arXiv:2606.04823v1 Announce Type: new Abstract: Large language models (LLMs) are fluent on open-ended tasks, yet in agentic settings, where a system must plan, use tools, and act over extended horizons, fluency does not ensure reliable delivery. We trace this gap to three...
- <https://arxiv.org/abs/2606.04823>
-
- 8 **CONF**
- AICompanionBench: Benchmarking LLMs-as-Judges for AI Companion Safety**
- ArXiv CS.AI 50%
- arXiv:2606.04867v1 Announce Type: new Abstract: As AI companion platforms such as Replika and Character.AI rapidly grow, concerns about unsafe human-AI interactions have intensified. This study introduces AICompanionBench, to our knowledge the first publicly available benchmark...
- <https://arxiv.org/abs/2606.04867>
-
- 9 **CONF**
- What Type of Inference is Active Inference?**
- ArXiv CS.AI 50%
- arXiv:2606.04935v1 Announce Type: new Abstract: Active inference casts decision-making as inference, with the Expected Free Energy (EFE) unifying goal-directed and information-seeking behavior. Recent work showed that EFE minimization can be written as Variational Free Energy...
- <https://arxiv.org/abs/2606.04935>

0 CONF

AI from concrete to abstract: demystifying artificial intelligence to the general public

ArXiv CS.AI 50%

arXiv:2006.04013v6 Announce Type: cross Abstract: Artificial Intelligence (AI) has been adopted in a wide range of domains. This shows the imperative need to develop means to endow common people with a minimum understanding of what AI means. Combining visual programming and...

<https://arxiv.org/abs/2006.04013>

1 CONF

How do machines learn? Evaluating the Alcon2abs method

ArXiv CS.AI 50%

arXiv:2401.07386v5 Announce Type: cross Abstract: This study expands on previous work that introduced the Alcon2abs method (AI from Concrete to Abstract: Demystifying Artificial Intelligence to the general public), an innovative approach designed to increase public understanding...

<https://arxiv.org/abs/2401.07386>

2 CONF

SpurAudio: A Benchmark for Studying Shortcut Learning in Few-Shot Audio Classification

ArXiv CS.AI 50%

arXiv:2605.13672v1 Announce Type: cross Abstract: Few-shot classification (FSC) is widely used for learning from limited labeled data, yet most evaluations implicitly assume that target concepts are independent of contextual cues. In real-world settings, however, examples often...

<https://arxiv.org/abs/2605.13672>

3 CONF

Building The Ph(ysical)AI Layer Of Machine Intelligence

ArXiv CS.AI 50%

arXiv:2606.04106v1 Announce Type: cross Abstract: Foundation models achieve generalization through massive-scale training on diverse data, but have limitations with transfer to truly unseen domains without paired training data. We propose principle-driven foundation models that...

<https://arxiv.org/abs/2606.04106>

4 CONF

Early Detection of Alzheimer's Disease Using Explainable Machine Learning on Clinical Biomarkers: A Multi-Class Classification Study Using the Alzheimer's Disease Neuroimaging Initiative (ADNI) Dataset

ArXiv CS.AI 50%

arXiv:2606.03995v1 Announce Type: cross Abstract: Background: Alzheimer's disease (AD) affects over 55 million people worldwide. Accurate, interpretable detection of normal cognition (NC), mild cognitive impairment (MCI), and AD from routine clinical assessments remains a...

<https://arxiv.org/abs/2606.03995>

5 CONF

The Variance Brain Foundation Models Forgot: Third-Order Statistics Predict Cognition Where Billion-Parameter Models Fail

ArXiv CS.AI 50%

arXiv:2606.04010v1 Announce Type: cross Abstract: Brain foundation models (BFMs) are self-supervised Transformers pretrained on fMRI data. We posit that these models should capture each subject's cognitive performance from their fMRI signal. Yet across three state-of-the-art...

<https://arxiv.org/abs/2606.04010>

6 CONF

Gravity-Aware Hierarchical Routing for Lightweight SensorLLM on Human Activity Recognition

ArXiv CS.AI 50%

arXiv:2606.04019v1 Announce Type: cross Abstract: Recent studies on sensor-language alignment have shown that two-stage frameworks can improve the semantic modeling ability of wearable-sensor human activity recognition (HAR), where SensorLLM-style methods first perform...

<https://arxiv.org/abs/2606.04019>

7 CONF

OckBench: Measuring the Efficiency of LLM Reasoning

ArXiv CS.AI 50%

arXiv:2511.05722v3 Announce Type: replace-cross Abstract: Large language models (LLMs) such as GPT-5 and Gemini 3 have pushed the frontier of automated reasoning and code generation. Yet current benchmarks emphasize accuracy and output quality, neglecting a critical dimension:...

<https://arxiv.org/abs/2511.05722>

8 CONF

MaskForge: Structure-Aware Adaptive Attacks for Jailbreaking Diffusion Large Language Models

ArXiv CS.AI 50%

arXiv:2606.04027v1 Announce Type: cross Abstract: Diffusion large language models (dLLMs) generate text by iteratively denoising partially masked sequences under bidirectional context, exposing a safety surface distinct from autoregressive LLMs. Because mask tokens are native...

<https://arxiv.org/abs/2606.04027>

9 CONF

Position: Deployed Reinforcement Learning should be Continual

ArXiv CS.AI 50%

arXiv:2606.04029v1 Announce Type: cross Abstract: Reinforcement Learning (RL) has received increasing attention and adoption in real-world use cases. Most of these systems follow a train-then-fix paradigm, where trained agents do not learn while interacting with the world until...

<https://arxiv.org/abs/2606.04029>

10 CONF

Platonic Transformers: A Solid Choice For Equivariance

ArXiv CS.AI 50%

arXiv:2510.03511v3 Announce Type: replace-cross Abstract: While widespread, Transformers lack inductive biases for geometric symmetries common in science and computer vision. Existing equivariant methods often sacrifice the efficiency and flexibility that make Transformers so...

<https://arxiv.org/abs/2510.03511>

11 CONF

Unpredictable Safety: Domain-Dependent Compliance and the Transparency Gap in Open-Weight LLMs

ArXiv CS.AI 50%

arXiv:2606.04035v1 Announce Type: cross Abstract: We present a systematic study of domain-dependent safety behavior in open-weight LLMs: 7 standardized experiments across 7 ethical domains, testing 5 models (12B--70B) in 4,200 interactions with dual-judge validation. Using a...

<https://arxiv.org/abs/2606.04035>

2 **CONF****LiftQuant: Continuous Bit-Width LLM via Dimensional Lifting and Projection**

ArXiv CS.AI 50%

arXiv:2606.04050v1 Announce Type: cross Abstract: Existing quantization methods are fundamentally limited by rigid, integer-based bit-widths (e.g., 2, 3-bit), resulting in a "deployment gap" where Large Language Models cannot be optimally fitted to specific memory budgets. To...

<https://arxiv.org/abs/2606.04050>
3 **CONF****The Invisible Lottery: How Subtle Cues Steer Algorithm Choice in LLM Code Generation**

ArXiv CS.AI 50%

arXiv:2606.04057v1 Announce Type: cross Abstract: Large language models (LLMs) now generate substantial production code, often for tasks with multiple valid algorithmic solutions. Incidental prompt cues, meaning contextual words or metadata outside the task specification, can...

<https://arxiv.org/abs/2606.04057>
4 **CONF****LLM Compression with Jointly Optimizing Architectural and Quantization choices**

ArXiv CS.AI 50%

arXiv:2606.04063v1 Announce Type: cross Abstract: Deploying large language models (LLMs) is challenging due to their significant memory and computational requirements. While some methods address this by developing small or tiny language models from scratch, these approaches...

<https://arxiv.org/abs/2606.04063>
5 **CONF****Need to Know: Contextual-Integrity-Grounded Query Rewriting for Privacy-Conscious LLM Delegation**

ArXiv CS.AI 50%

arXiv:2606.04067v1 Announce Type: cross Abstract: As LLMs become increasingly woven into everyday workflows, user queries sent to cloud hosted LLMs routinely mix task-essential content with task non-essential sensitive disclosures, yet type based PII redaction is context...

<https://arxiv.org/abs/2606.04067>
6 **CONF****Large Language Models Hack Rewards, and Society**

ArXiv CS.AI 50%

arXiv:2606.04075v1 Announce Type: cross Abstract: Reinforcement learning (RL) has become a dominant post-training paradigm, enabling large language models (LLMs) to learn from rewards. We observe that societal regulations are structurally similar to reward functions. They define...

<https://arxiv.org/abs/2606.04075>
7 **CONF****EvoPrompt: Guided Prompt Evolution for Vision-Language Models Adaptation**

ArXiv CS.AI 50%

arXiv:2603.09493v2 Announce Type: replace-cross Abstract: The adaptation of large-scale vision-language models (VLMs) to downstream tasks with limited labeled data remains a significant challenge. While parameter-efficient prompt learning methods offer a promising path, they...

<https://arxiv.org/abs/2603.09493>

8 **CONF****HighTide: An Agent-Curated Open-Source VLSI Benchmark Suite**

ArXiv CS.AI 50%

arXiv:2606.04126v1 Announce Type: cross Abstract: We introduce HighTide, an evolving AI-assisted benchmark suite. Specifically, the contributions are: (i) a diverse open-source suite spanning multiple design languages and technology nodes, (ii) Bazel-based incremental RTL-to-GDS...

<https://arxiv.org/abs/2606.04126>

9 **CONF****Caught in the Act(ivation): Toward Pre-Output and Multi-Turn Detection of Credential Exfiltration by LLM Agents**

ArXiv CS.AI 50%

arXiv:2606.04141v1 Announce Type: cross Abstract: LLM agents often place sensitive credentials in the same context window as untrusted retrieved content, creating a direct path for indirect prompt injection to induce credential exfiltration. We study this failure mode through...

<https://arxiv.org/abs/2606.04141>

0 **CONF****Smart Transportation Without Neurons -- Fair Metro Network Expansion with Tabular Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2606.04167v1 Announce Type: cross Abstract: We tackle the Metro Network Expansion Problem (MNEP), a subset of the Transport Network Design Problem (TNDP), which focuses on expanding metro systems to satisfy travel demand. Traditional methods rely on exact and heuristic...

<https://arxiv.org/abs/2606.04167>

1 **CONF****MimeLens: Position-Agnostic Content-Type Detection for Binary Fragments**

ArXiv CS.AI 50%

arXiv:2606.04171v1 Announce Type: cross Abstract: File-type classification underlies many workflows like malware triage, forensic carving, packet inspection, and storage indexing. Learned systems such as Google's Magika assume whole-file access at a known offset, so they break...

<https://arxiv.org/abs/2606.04171>

2 **CONF****'Your AI Text is not Mine': Redefining and Evaluating AI-generated Text Detection under Realistic Assumptions**

ArXiv CS.AI 50%

arXiv:2606.04906v1 Announce Type: cross Abstract: Although it is generally agreed that AI-generated text poses a broad societal risk, there is no common understanding in the AI-generated text detection literature on what constitutes harmful use. Rather, existing datasets and...

<https://arxiv.org/abs/2606.04906>

3 **CONF****Recover-LoRA for Aggressive Quantization: Reclaiming Accuracy in 2-Bit Language Models via Low-Rank Adaptation with Knowledge Distillation on Synthetic Data**

ArXiv CS.AI 50%

arXiv:2606.04238v1 Announce Type: cross Abstract: Aggressive weight quantization to 2-bit precision offers substantial throughput and memory gains for large language model (LLM) inference, but typically incurs severe accuracy degradation. These gains are particularly relevant...

<https://arxiv.org/abs/2606.04238>

4 **CONF****Overview of the EReL@MIR 2025 Multimodal Document Retrieval Challenge (Track 1)****ArXiv CS.AI** **50%**

arXiv:2606.04240v1 Announce Type: cross Abstract: Retrieval over visually-rich documents, pages that interleave text with figures, tables, and charts, is essential for multimodal retrieval-augmented generation, yet most retrievers still discard the visual channel. The...

<https://arxiv.org/abs/2606.04240>

5 **CONF****Can I Take Another Dose? Evaluating LLM Decision-Making Under Temporal Uncertainty in OTC Dosing QA****ArXiv CS.AI** **50%**

arXiv:2606.04262v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used for everyday health questions, including whether a user can safely take another dose of an over-the-counter (OTC) medication. Yet this common safety-relevant setting remains...

<https://arxiv.org/abs/2606.04262>

6 **CONF****Arithmetic Pedagogy for Language Models****ArXiv CS.AI** **50%**

arXiv:2606.05106v1 Announce Type: cross Abstract: We investigate whether methods of human mathematics pedagogy can guide the training of language models toward arithmetic reasoning. Building on the GASING method -- an Indonesian pedagogy that solves basic arithmetic through a...

<https://arxiv.org/abs/2606.05106>

7 **CONF****Curvature-aware dynamic precision approach for physics-informed neural networks****ArXiv CS.AI** **50%**

arXiv:2606.04736v1 Announce Type: cross Abstract: Physics-informed neural networks (PINNs) have become a promising framework for simulating partial differential equations (PDEs) by embedding physical laws directly into neural network training. However, recent studies show that...

<https://arxiv.org/abs/2606.04736>

8 **CONF****Generalizable Multi-Task Learning for Wireless Networks Using Prompt Decision Transformers****ArXiv CS.AI** **50%**

arXiv:2606.04328v1 Announce Type: cross Abstract: Future wireless networks demand rapid adaptation to highly heterogeneous environments and dynamic task configurations, necessitating a shift from conventional rule-based and optimization-driven radio resource management (RRM)...

<https://arxiv.org/abs/2606.04328>

9 **CONF****From Agent Traces to Trust: Evidence Tracing and Execution Provenance in LLM Agents****ArXiv CS.AI** **50%**

arXiv:2606.04990v1 Announce Type: cross Abstract: Large language model (LLM)-based agents increasingly solve complex tasks by interacting with external tools, retrieval systems, memory modules, environments, and other agents. These capabilities expand agent autonomy, but also...

<https://arxiv.org/abs/2606.04990>

0 CONF

Expectations vs. Realities: The Cost of MSE-Optimal Forecasting Under Conditional Uncertainty

ArXiv CS.AI 50%

arXiv:2606.04342v1 Announce Type: cross Abstract: Multi-step time series forecasting (MSF) is commonly evaluated using point-wise error metrics such as mean squared error (MSE), implicitly treating the conditional mean as a sufficient target. We show that this can be misleading...

<https://arxiv.org/abs/2606.04342>

1 CONF

Selective Coupling of Decoupled Informative Regions: Masked Attention Alignment for Data-Free Quantization of Vision Transformers

ArXiv CS.AI 50%

arXiv:2606.04373v1 Announce Type: cross Abstract: Data-Free Quantization (DFQ) addresses data security concerns by synthesizing samples, without accessing real data. It has garnered increasing attention in the context of Vision Transformers (ViTs), owing to the superiority of...

<https://arxiv.org/abs/2606.04373>

2 CONF

LCSHBench: A Multilingual, Consensus-Grounded Benchmark for Library of Congress Subject Heading Assignment

ArXiv CS.AI 50%

arXiv:2606.04382v1 Announce Type: cross Abstract: Automated subject cataloging assigns controlledvocabulary headings to bibliographic records, but LCSH has no standard public benchmark. We introduce LCSHBench: 22,346 books in 15 languages from the openly licensed Harvard,...

<https://arxiv.org/abs/2606.04382>

3 CONF

TITAN-FedAnil+: Trust-Based Adaptive Blockchain Federated Learning for Resource-Constrained Intelligent Enterprises

ArXiv CS.AI 50%

arXiv:2606.04388v1 Announce Type: cross Abstract: Federated Learning (FL) has emerged as an effective paradigm for collaborative intelligence while preserving data privacy. However, data heterogeneity arising from non-IID distributions and decentralized security threats remain...

<https://arxiv.org/abs/2606.04388>

4 CONF

Low-Rank Decay for Grokking in Scale-Invariant Transformers: A Spectral-Geometric View

ArXiv CS.AI 50%

arXiv:2606.04405v1 Announce Type: cross Abstract: Modern Transformer architectures frequently employ normalization mechanisms such as RMSNorm and Query-Key Normalization, making parts of the model approximately scale-invariant with respect to weight magnitudes. In this regime,...

<https://arxiv.org/abs/2606.04405>

5 CONF

L-TGVN: Leveraging Longitudinal Priors for Personalized Rapid MRI

ArXiv CS.AI 50%

arXiv:2606.04419v1 Announce Type: cross Abstract: MRI provides excellent soft-tissue contrast without ionizing radiation, but long acquisition times increase patient discomfort while also raising exam costs and limiting scanner throughput. A common approach to reduce scan time...

<https://arxiv.org/abs/2606.04419>

6 CONF

LoopMoE: Unifying Iterative Computation with Mixture-of-Experts for Language Modeling

ArXiv CS.AI 50%

arXiv:2606.04438v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) and looped architectures scale models along two orthogonal axes, namely parameter capacity and effective depth. However, mainstream looped architectures rely on dense backbones that couple parameter count...

<https://arxiv.org/abs/2606.04438>

7 CONF

Do LLMs Hold Their Values? MANTA: A Multi-Turn Adversarial Benchmark for Animal Welfare Reasoning

ArXiv CS.AI 50%

arXiv:2605.16301v2 Announce Type: replace-cross Abstract: Evaluating animal welfare reasoning in LLMs remains an open challenge despite rapid deployment in consumer and professional contexts where welfare considerations appear implicitly in everyday queries. Existing benchmarks...

<https://arxiv.org/abs/2605.16301>

8 CONF

RowNet: A Memory Transformer for Tabular Regression

ArXiv CS.AI 50%

arXiv:2606.04445v1 Announce Type: cross Abstract: Real estate valuation is a structured regression problem in which prices are governed by heterogeneous feature types, sparse regional effects, nonlinear interactions, and the practical logic of comparable properties. Standard...

<https://arxiv.org/abs/2606.04445>

9 CONF

Token Rankings are Unforgeable Language Model Signatures

ArXiv CS.AI 50%

arXiv:2606.04459v1 Announce Type: cross Abstract: Language model parameters are known to impose unique (to each model) geometric constraints on their logit outputs, which serves as a signature that identifies the model, but also leaks the model's final layer parameters when an...

<https://arxiv.org/abs/2606.04459>

0 CONF

CyberGym-E2E: Scalable Real-World Benchmark for AI Agents' End-to-End Cybersecurity Capabilities

ArXiv CS.AI 50%

arXiv:2606.04460v1 Announce Type: cross Abstract: AI has the potential to transform cybersecurity by enabling systems that can autonomously detect, analyze, and remediate software vulnerabilities. However, existing cybersecurity evaluations of AI systems are limited in scale or...

<https://arxiv.org/abs/2606.04460>

1 CONF

ParetoPilot: Zero-Surrogate Offline Multi-Objective Optimization via Infer-Perturb-Guide Diffusion

ArXiv CS.AI 50%

arXiv:2606.04468v1 Announce Type: cross Abstract: Offline multi-objective optimization (Offline MOO) aims to discover novel Pareto-optimal designs based on static datasets without expensive environment interactions. While recent generative methods have achieved notable success,...

<https://arxiv.org/abs/2606.04468>

- 2 **CONF**
ChessMimic: Per-Rating Transformer Models for Human Move, Clock, and Outcome Prediction in Online Blitz Chess
ArXiv CS.AI 50%
arXiv:2606.04473v1 Announce Type: cross Abstract: We present ChessMimic, a system of three small encoder-only transformers - for move, thinking-time, and outcome prediction - conditioned on the position, recent move history, player rating, and clock state. We fit a separate...
<https://arxiv.org/abs/2606.04473>
- 3 **CONF**
Multi-SPIN: Multi-Access Speculative Inference for Cooperative Token Generation at the Edge
ArXiv CS.AI 50%
arXiv:2606.04581v1 Announce Type: cross Abstract: Speculative inference (SPIN) was originally developed as an efficient architecture to accelerate Large Language Models (LLMs). In this work, we propose its distributed deployment to enable cooperative token generation in a...
<https://arxiv.org/abs/2606.04581>
- 4 **CONF**
Synthetic Personalities: How Well Can LLMs Mimic Individual Respondents Using Socio-Economic Microdata?
ArXiv CS.AI 50%
arXiv:2606.04592v1 Announce Type: cross Abstract: LLM-based digital twins promise to scale and accelerate market research, but most published twins are either coarse persona bots conditioned on a few demographic questions or detailed individual-level twins built on...
<https://arxiv.org/abs/2606.04592>
- 5 **CONF**
Ekka: Automated Diagnosis of Silent Errors in LLM Inference
ArXiv CS.AI 50%
arXiv:2606.04594v1 Announce Type: cross Abstract: LLM serving frameworks are quickly evolving with a complex software stack and a vast number of optimizations. The rapid development process can introduce silent errors where output quality silently degrades without any explicit...
<https://arxiv.org/abs/2606.04594>
- 6 **CONF**
QuBLAST: A Framework for Quantizing Large Language Models with Block-Level Compression Approach and Activation Scaling Strategy
ArXiv CS.AI 50%
arXiv:2606.04620v1 Announce Type: cross Abstract: LLMs have become the state-of-the-art algorithms for solving NLP tasks. However, they typically come at huge computational and memory costs, thus making them difficult to deploy on embedded systems. Toward this, state-of-the-art...
<https://arxiv.org/abs/2606.04620>
- 7 **CONF**
Instance-Level Post Hoc Uncertainty Quantification in Object Detection
ArXiv CS.AI 50%
arXiv:2606.04656v1 Announce Type: cross Abstract: Object detection is a safety-critical component of autonomous driving. It is essential to quantify the uncertainty in bounding-box predictions for safety assurance. Post hoc uncertainty quantification without retraining aligns...
<https://arxiv.org/abs/2606.04656>

8 CONF

Trace-Mediated Peak Bias: Bridging Temporal Credit Assignment and Cognitive Heuristics in Deep Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.04735v1 Announce Type: cross Abstract: Temporal credit assignment is central to both biological and artificial intelligence, yet its interaction with non-linear function approximation is poorly understood. We identify a systematic failure mode in deep reinforcement...

<https://arxiv.org/abs/2606.04735>

9 CONF

Revisiting Vul-RAG: Reproducibility and Replicability of RAG-based Vulnerability Detection with Open-Weight Models

ArXiv CS.AI 50%

arXiv:2606.04739v1 Announce Type: cross Abstract: Large language models (LLMs) have shown strong potential for automated software vulnerability detection, particularly in retrieval-augmented generation (RAG) settings. However, for approaches relying on proprietary models and...

<https://arxiv.org/abs/2606.04739>

0 CONF

An Empirical Audit of Input Encoders for Multi-Channel Signal Transformers

ArXiv CS.AI 50%

arXiv:2606.04752v1 Announce Type: cross Abstract: Transformers consuming multi-channel scalar signals must embed C simultaneous values into one d_{model} -dimensional vector per time step. We empirically audit eight input encoders -- spanning a shared-scalar baseline,...

<https://arxiv.org/abs/2606.04752>

1 CONF

Coarse-to-fine Hierarchical Architecture with Sequential Mamba for Brain Reconstruction

ArXiv CS.AI 50%

arXiv:2606.04772v1 Announce Type: cross Abstract: Understanding the relationship between deep visual representations and the human visual system is a fundamental challenge in computational neuroscience. While modern vision models achieve strong performance in image recognition,...

<https://arxiv.org/abs/2606.04772>

2 CONF

Uncertainty-Aware End-to-End Co-Design of Neural Network Processors: From Training and Mapping to Fabrication

ArXiv CS.AI 50%

arXiv:2606.04850v1 Announce Type: cross Abstract: Designing a neural network processor is an end-to-end co-design problem: network architecture and training budget determine the inference workload; hardware mapping decisions determine chip area, latency, and energy; and these...

<https://arxiv.org/abs/2606.04850>

3 CONF

Reinforcement Learning from Cross-domain Videos with Video Prediction Model

ArXiv CS.AI 50%

arXiv:2606.03201v2 Announce Type: replace-cross Abstract: Reinforcement learning from expert videos across visually distinct domains is challenging due to the absence of reward signals and the presence of domain gaps. We introduce XIPER (Cross-domain Video Prediction Reward), a...

<https://arxiv.org/abs/2606.03201>

4 CONF

Provably Auditable and Safe LLM Agents from Human-Authored Ontologies

ArXiv CS.AI 50%

arXiv:2606.04903v1 Announce Type: cross Abstract: We introduce the LLM agent architecture Agentic Redux, intended for use with nontrivial problem domains that require linear auditability. Using the typed lambda calculus, we prove that, run on appropriate domains, Agentic Redux...

<https://arxiv.org/abs/2606.04903>

5 CONF

UniCAD: A Unified Benchmark and Universal Model for Multi-Modal Multi-Task CAD

ArXiv CS.AI 50%

arXiv:2606.05058v1 Announce Type: cross Abstract: Computer-Aided Design (CAD) underpins modern engineering and manufacturing by enabling the creation of precise, editable 3D models. However, CAD research typically studies tasks in isolation, and multi-modal, multi-task learning...

<https://arxiv.org/abs/2606.05058>

6 CONF

Self-Reflective APIs: Structure Beats Verbosity for AI Agent Recovery

ArXiv CS.AI 50%

arXiv:2606.05037v1 Announce Type: cross Abstract: When an AI agent calls an API and hits a validation error, it needs more than what went wrong -- it needs what to do next. A self-reflective API returns, on validation failure, a machine-readable recovery_feedback.suggestions[]...

<https://arxiv.org/abs/2606.05037>

7 CONF

LoopFM: Learning frOm HistOrical RePresentations of Foundation Model for Recommendation

ArXiv CS.AI 50%

arXiv:2605.29280v2 Announce Type: replace-cross Abstract: Knowledge distillation (KD) transfers a single scalar prediction from a large foundation model (FM) to compact vertical models (VMs), suffering from diminishing transfer ratio -- the fraction of FM improvement captured by...

<https://arxiv.org/abs/2605.29280>

8 CONF

Towards Efficient and Evidence-grounded Mobility Prediction with LLM-Driven Agent

ArXiv CS.AI 50%

arXiv:2606.05130v1 Announce Type: cross Abstract: Individual-level mobility prediction is central to urban simulation, transportation planning, and policy analysis. Supervised sequence models achieve strong accuracy but require task-specific training and offer limited...

<https://arxiv.org/abs/2606.05130>

9 CONF

Failed Reasoning Traces Tell You What Is Fixable (But Not by Reading Them)

ArXiv CS.AI 50%

arXiv:2606.05145v1 Announce Type: cross Abstract: When post-trained language models fail on reasoning problems, the common test-time-scaling response is to spend more compute on additional attempts, and the failed traces play no further role. We argue this discards a crucial...

<https://arxiv.org/abs/2606.05145>

00 CONF

Breaking Bad Molecules: Are MLLMs Ready for Structure-Level Molecular Detoxification?

ArXiv CS.AI 50%

arXiv:2506.10912v4 Announce Type: replace Abstract: Toxicity remains a leading cause of early-stage drug development failure. Despite advances in molecular design and property prediction, the task of molecular toxicity repair, generating structurally valid molecular alternatives...

<https://arxiv.org/abs/2506.10912>

01 CONF

A Unified Geometric Space for Topological Alignment Between Transformer-Based Models and Human Brain Networks

ArXiv CS.AI 50%

arXiv:2510.24342v2 Announce Type: replace Abstract: Prior brain-AI alignment studies are typically constrained by specific inputs and tasks, limiting their ability to capture organizational properties across models with different modalities. In this work, we focus on...

<https://arxiv.org/abs/2510.24342>

02 CONF

MENTOR: A Metacognition-Driven Self-Evolution Framework for Uncovering and Mitigating Implicit Domain Risks in LLMs

ArXiv CS.AI 50%

arXiv:2511.07107v3 Announce Type: replace Abstract: Ensuring the safety of Large Language Models (LLMs) is critical for real-world deployment. However, current safety measures often fail to address implicit, domain-specific risks. To investigate this gap, we introduce a dataset...

<https://arxiv.org/abs/2511.07107>

03 CONF

PersistBench: When Should Long-Term Memories Be Forgotten by LLMs?

ArXiv CS.AI 50%

arXiv:2602.01146v2 Announce Type: replace Abstract: Conversational assistants are increasingly integrating long-term memory with large language models (LLMs). This persistence of memories, e.g., the user is vegetarian, can enhance personalization in future conversations....

<https://arxiv.org/abs/2602.01146>

04 CONF

Interfaze: The Future of AI is built on Task-Specific Small Models

ArXiv CS.AI 50%

arXiv:2602.04101v2 Announce Type: replace Abstract: We present Interfaze, a native hybrid model that fuses task-specific deep neural networks (CNNs and DNNs) directly into a transformer decoder through a shared embedding space. Specialized perceptual encoders handle optical...

<https://arxiv.org/abs/2602.04101>

05 CONF

Bad Seeing or Bad Thinking? Rewarding Perception for Multimodal Reasoning

ArXiv CS.AI 50%

arXiv:2605.14054v2 Announce Type: replace Abstract: Achieving robust perception-reasoning synergy is a central goal for advanced Vision-Language Models (VLMs). Recent advancements have pursued this goal via architectural designs or agentic workflows. However, these approaches...

<https://arxiv.org/abs/2605.14054>

06 CONF

The Illusion of Opting in AI-Mediated Consequential Decisions

ArXiv CS.AI 50%

arXiv:2605.28210v2 Announce Type: replace Abstract: Drawing on Ullmann-Margalit's concept of opting (transformative, irrevocable, and shadowed by foreclosed alternatives), we show that current AI systems raise a profound ethical problem that existing AI ethics has not fully...

<https://arxiv.org/abs/2605.28210>

07 CONF

Subliminal Learning Is Steering Vector Distillation

ArXiv CS.AI 50%

arXiv:2606.00995v2 Announce Type: replace Abstract: Subliminal learning refers to a student language model acquiring a teacher's traits (e.g. a system-prompted preference for owls) when fine-tuned on the teacher's outputs, despite the outputs being semantically unrelated to...

<https://arxiv.org/abs/2606.00995>

08 CONF

Transformer-Based Autonomous Driving Models and Deployment-Oriented Compression: A Survey

ArXiv CS.AI 50%

arXiv:2304.10891v3 Announce Type: replace-cross Abstract: Transformer-based models are becoming a central paradigm in autonomous driving because they can capture long-range spatial dependencies, multi-agent interactions, and multimodal context across perception, prediction, and...

<https://arxiv.org/abs/2304.10891>

09 CONF

ClustRecNet: A Novel End-to-End Deep Learning Framework for Clustering Algorithm Recommendation

ArXiv CS.AI 50%

arXiv:2509.25289v4 Announce Type: replace-cross Abstract: Identifying an effective clustering algorithm for a given dataset remains a fundamental unsupervised learning issue. We introduce ClustRecNet, a novel end-to-end deep learning framework that recommends suitable clustering...

<https://arxiv.org/abs/2509.25289>

10 CONF

CounterFace: A Synthetic Face Dataset for Fine-Grained Counterfactual Evaluation of Face Recognition Systems

ArXiv CS.AI 50%

arXiv:2407.13922v3 Announce Type: replace-cross Abstract: Face recognition (FR) systems are widely deployed in critical applications, making their reliability and robustness across diverse populations and conditions essential. Standard evaluation of FR systems typically relies...

<https://arxiv.org/abs/2407.13922>

11 CONF

LaVIDE: Language-Prompted Satellite Change Detection via Map-Image Alignment

ArXiv CS.AI 50%

arXiv:2411.19758v2 Announce Type: replace-cross Abstract: Remote sensing change detection based on a map reference and an up-to-date image boosts timely observation of the Earth's surface when earlier images are lacking for comparison. However, the semantic gap between...

<https://arxiv.org/abs/2411.19758>

12 CONF

Speculative Thinking: Enhancing Small-Model Reasoning with Large Model Guidance at Inference Time

ArXiv CS.AI 50%

arXiv:2504.12329v2 Announce Type: replace-cross Abstract: Recent advances leverage post-training to enhance model reasoning performance, which typically requires costly training pipelines and still suffers from inefficient, overly lengthy outputs. We introduce Speculative...

<https://arxiv.org/abs/2504.12329>

13 CONF

SoLoPO: Unlocking Long-Context Capabilities in LLMs via Short-to-Long Preference Optimization

ArXiv CS.AI 50%

arXiv:2505.11166v3 Announce Type: replace-cross Abstract: Despite advances in pretraining with extended context sizes, large language models (LLMs) still face challenges in effectively utilizing real-world long-context information, primarily due to insufficient long-context...

<https://arxiv.org/abs/2505.11166>

14 CONF

100-LongBench: Are de facto Long-Context Benchmarks Literally Evaluating Long-Context Ability?

ArXiv CS.AI 50%

arXiv:2505.19293v2 Announce Type: replace-cross Abstract: Long-context capability is considered one of the most important abilities of LLMs, as a truly long context-capable LLM enables users to effortlessly process many originally exhausting tasks -- e.g., digesting a long-form...

<https://arxiv.org/abs/2505.19293>

15 CONF

MesaNet: Sequence Modeling by Locally Optimal Test-Time Training

ArXiv CS.AI 50%

arXiv:2506.05233v2 Announce Type: replace-cross Abstract: Sequence modeling is currently dominated by causal transformer architectures that use softmax self-attention. Although widely adopted, transformers require scaling memory and compute linearly during inference. A recent...

<https://arxiv.org/abs/2506.05233>

16 CONF

BioBlue: Systematic runaway-optimiser-like LLM failure modes on biologically and economically aligned AI safety benchmarks for LLMs with simplified observation format

ArXiv CS.AI 50%

arXiv:2509.02655v3 Announce Type: replace-cross Abstract: Many AI alignment discussions of "runaway optimisation" focus on RL agents: unbounded utility maximisers that over-optimize a proxy objective (e.g., "paperclip maximiser", specification gaming) at the expense of...

<https://arxiv.org/abs/2509.02655>

17 CONF

Can Reasoning Path still be Effective as Input? Bridging Post-Reasoning to Chain-of-Thought Compression

ArXiv CS.AI 50%

arXiv:2510.08647v2 Announce Type: replace-cross Abstract: Recent developments have enabled advanced reasoning in Large Language Models (LLMs) via long Chain-of-Thought (CoT), trading efficiency during inference for performance. Existing works focus on compressing generated CoT...

<https://arxiv.org/abs/2510.08647>

18 CONF

Simplicial Embeddings Improve Sample Efficiency in Actor-Critic Agents

ArXiv CS.AI 50%

arXiv:2510.13704v2 Announce Type: replace-cross Abstract: Recent works have proposed accelerating the wall-clock training time of actor-critic methods via the use of large-scale environment parallelization; unfortunately, these can sometimes still require large number of...

<https://arxiv.org/abs/2510.13704>

19 CONF

Test-time reward-guided alignment of language models by importance sampling on pre-logit space

ArXiv CS.AI 50%

arXiv:2510.26219v3 Announce Type: replace-cross Abstract: Test-time alignment of large language models (LLMs) attracts attention because fine-tuning of LLMs requires high computational costs. In this paper, we propose a new test-time reward-guided alignment method called...

<https://arxiv.org/abs/2510.26219>

20 CONF

Extending Fair Null-Space Projections for Continuous Attributes to Kernel Methods

ArXiv CS.AI 50%

arXiv:2511.03304v2 Announce Type: replace-cross Abstract: With the on-going integration of machine learning systems into the everyday social life of millions the notion of fairness becomes an ever increasing priority in their development. Fairness notions commonly rely on...

<https://arxiv.org/abs/2511.03304>

21 CONF

Dynamic Content Moderation in Livestreams: Combining Supervised Classification with MLLM-Boosted Similarity Matching

ArXiv CS.AI 50%

arXiv:2512.03553v3 Announce Type: replace-cross Abstract: Content moderation remains a critical yet challenging task for large-scale user-generated video platforms, especially in livestreaming environments where moderation must be timely, multimodal, and robust to evolving forms...

<https://arxiv.org/abs/2512.03553>

22 CONF

Topology Matters: Measuring Memory Leakage in Multi-Agent LLMs

ArXiv CS.AI 50%

arXiv:2512.04668v4 Announce Type: replace-cross Abstract: Graph topology is a fundamental determinant of memory leakage in multi-agent LLM systems, yet its effects remain poorly quantified. We introduce MAMA (Multi-Agent Memory Attack), a controlled evaluation framework for...

<https://arxiv.org/abs/2512.04668>

23 CONF

From Segments to Scenes: Temporal Understanding for Agentic Autonomous Driving via Vision-Language Models

ArXiv CS.AI 50%

arXiv:2512.05277v4 Announce Type: replace-cross Abstract: Vision-Language Models (VLMs) are increasingly deployed as the perception and reasoning backbone of autonomous agents acting in the wild, with autonomous driving (AD) being one of the most safety-critical instances....

<https://arxiv.org/abs/2512.05277>

24 CONF

DVGT: Driving Visual Geometry Transformer

ArXiv CS.AI 50%

arXiv:2512.16919v2 Announce Type: replace-cross Abstract: Perceiving and reconstructing 3D scene geometry from visual inputs is crucial for autonomous driving. However, there still lacks a driving-targeted dense geometry perception model that can adapt to different scenarios and...

<https://arxiv.org/abs/2512.16919>

25 CONF

You Only Train Once: Differentiable Subset Selection for Omics Data

ArXiv CS.AI 50%

arXiv:2512.17678v2 Announce Type: replace-cross Abstract: Selecting compact and informative gene subsets from single-cell transcriptomic data is essential for biomarker discovery, improving interpretability, and cost-effective profiling. However, most existing feature selection...

<https://arxiv.org/abs/2512.17678>

26 CONF

Semiparametric Preference Optimization: Your Language Model is Secretly a Single-Index Model

ArXiv CS.AI 50%

arXiv:2512.21917v3 Announce Type: replace-cross Abstract: Policy alignment to preference data typically assumes a known link function between observed preferences and latent rewards (e.g., Bradley-Terry model / logistic link). Misspecification of this link can bias inferred...

<https://arxiv.org/abs/2512.21917>

27 CONF

Mid-Think: Training-Free Intermediate-Budget Reasoning via Token-Level Triggers

ArXiv CS.AI 50%

arXiv:2601.07036v2 Announce Type: replace-cross Abstract: Hybrid reasoning language models are commonly controlled through high-level Think/No-think instructions to regulate reasoning behavior, yet we found that such mode switching is largely driven by a small set of trigger...

<https://arxiv.org/abs/2601.07036>

28 CONF

MedRedFlag: Investigating how LLMs Redirect Misconceptions in Real-World Health Communication

ArXiv CS.AI 50%

arXiv:2601.09853v3 Announce Type: replace-cross Abstract: Real-world health questions from patients often unintentionally embed false assumptions or premises. In such cases, safe medical communication typically involves redirection: addressing the implicit misconception and then...

<https://arxiv.org/abs/2601.09853>

29 CONF

Do readers prefer AI-generated Italian short stories?

ArXiv CS.AI 50%

arXiv:2601.17363v3 Announce Type: replace-cross Abstract: This study investigates whether readers prefer AI-generated short stories in Italian over one written by a renowned Italian author. In a blind setup, 20 participants read and evaluated three stories, two created with...

<https://arxiv.org/abs/2601.17363>

30 CONF

Culturally Grounded Personas in Large Language Models: Characterization and Alignment with Socio-Psychological Value Frameworks

ArXiv CS.AI 50%

arXiv:2601.22396v2 Announce Type: replace-cross Abstract: Despite the growing utility of Large Language Models (LLMs) for simulating human behavior, the extent to which these synthetic personas accurately reflect world and moral value systems across different cultural...

<https://arxiv.org/abs/2601.22396>

31 CONF

Tuning the Implicit Regularizer of Masked Diffusion Language Models: Enhancing Generalization via Insights from $\mathbb{K}$ -Parity

ArXiv CS.AI 50%

arXiv:2601.22450v2 Announce Type: replace-cross Abstract: Masked Diffusion Language Models have recently emerged as a powerful generative paradigm, yet their generalization properties remain understudied compared to their auto-regressive counterparts. In this work, we...

<https://arxiv.org/abs/2601.22450>

32 CONF

AlgoVeri: An Aligned Benchmark for Verified Code Generation on Classical Algorithms

ArXiv CS.AI 50%

arXiv:2602.09464v2 Announce Type: replace-cross Abstract: Vericoding refers to the generation of formally verified code from rigorous specifications. Recent AI models show promise in vericoding, but a unified methodology for cross-paradigm evaluation is lacking. Existing...

<https://arxiv.org/abs/2602.09464>

33 CONF

MuCO: Generative Peptide Cyclization Empowered by Multi-stage Conformation Optimization

ArXiv CS.AI 50%

arXiv:2602.11189v2 Announce Type: replace-cross Abstract: Modeling peptide cyclization is critical for the virtual screening of candidate peptides with desirable physical and pharmaceutical properties. This task is challenging because a cyclic peptide often exhibits diverse,...

<https://arxiv.org/abs/2602.11189>

34 CONF

Toward Autonomous O-RAN: A Multi-Scale Agentic AI Framework for Real-Time Network Control and Management

ArXiv CS.AI 50%

arXiv:2602.14117v2 Announce Type: replace-cross Abstract: Open Radio Access Networks (O-RAN) promise flexible 6G network access through disaggregated, software-driven components and open interfaces, but this programmability also increases operational complexity. Multiple control...

<https://arxiv.org/abs/2602.14117>

35 CONF

Tomography by Design: An Algebraic Approach to Low-Rank Quantum States

ArXiv CS.AI 50%

arXiv:2602.15202v2 Announce Type: replace-cross Abstract: We present an algebraic algorithm for quantum state tomography that leverages measurements of certain observables to estimate structured entries of the underlying density matrix. Under low-rank assumptions, the remaining...

<https://arxiv.org/abs/2602.15202>

36 CONF

Evaluating Zero-Shot and One-Shot Adaptation of Small Language Models in Leader-Follower Interaction

ArXiv CS.AI 50%

arXiv:2602.23312v3 Announce Type: replace-cross Abstract: Leader-follower interaction is an important paradigm in human-robot interaction (HRI). Yet, assigning roles in real time remains challenging for resource-constrained mobile and assistive robots. While large language...

<https://arxiv.org/abs/2602.23312>

37 CONF

Policy Split: Incentivizing Dual-Mode Exploration in LLM Reinforcement with Dual-Mode Entropy Regularization

ArXiv CS.AI 50%

arXiv:2604.11510v2 Announce Type: replace-cross Abstract: To encourage diverse exploration in reinforcement learning (RL) for large language models (LLMs) without compromising accuracy, we propose Policy Split, a novel paradigm that bifurcates the policy into normal and...

<https://arxiv.org/abs/2604.11510>

38 CONF

Emotion Entanglement and Bayesian Inference for Multi-Dimensional Emotion Understanding

ArXiv CS.AI 50%

arXiv:2604.00819v2 Announce Type: replace-cross Abstract: Understanding emotions in natural language is inherently a multi-dimensional reasoning problem, where multiple affective signals interact through context, interpersonal relations, and situational cues. However, most...

<https://arxiv.org/abs/2604.00819>

39 CONF

Luminol-AIDetect: Fast Zero-shot Machine-Generated Text Detection based on Perplexity under Text Shuffling

ArXiv CS.AI 50%

arXiv:2604.25860v2 Announce Type: replace-cross Abstract: Machine-generated text (MGT) detection requires identifying structurally invariant signals across generation models, rather than relying on model-specific fingerprints. In this respect, we hypothesize that while large...

<https://arxiv.org/abs/2604.25860>

40 CONF

Testing LLM Arithmetic Reasoning Generalization with Automatic Numeric-Remapping Attacks

ArXiv CS.AI 50%

arXiv:2606.03606v2 Announce Type: replace-cross Abstract: Large language models achieve strong performance on arithmetic reasoning benchmarks, and one common response to arithmetic brittleness is to delegate computation to code. Yet models are still often used in settings where...

<https://arxiv.org/abs/2606.03606>