



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	152	Media
TechCrunch AI	3	Media
HackerNews	3	Community
The Verge AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 1 story

VERI

MAViS: A Multimodal Conversational Assistant For Avian Species

ArXiv CS.AI +1 source

80%

arXiv:2603.07294v2 Announce Type: replace-cross Abstract: Fine-grained understanding and species-specific multimodal question answering are vital for advancing biodiversity conservation and ecological monitoring. However, existing multimodal large language models face challenges...

<https://arxiv.org/abs/2603.07294>

CONFIRMED SIGNALS — 159 stories

CONF

Show HN: Lowfat - pluggable CLI filter that saved 91.8% of my LLM tokens

HackerNews 60% HN 126

Finally made something I've always wanted, using the model we built. • SOTA omni embedding model, fully local, indexes text, PDF, image, audio, and video • Swift-native app UI + mlx-swift-transformer core. No Python. • Tested on M3 Pro 18G / M3 Ultra 512G / M4 Pro 48G. All work...

<https://github.com/zdk/lowfat>

CONF

Show HN: Nanocode-CLI - A lightweight terminal-based AI coding assistant

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/hit9/nanocode>

CONF

Show HN: SnapToCode - Screenshot any UI and get clean Tailwind code

HackerNews 50% HN 2

No summary — see link for full article.

<https://chromewebstore.google.com/detail/snaptocode/jpchamlmjfoccmkdoiaibbpgkidapcnk>

CONF

The token bill comes due: Inside the industry scramble to manage AI's runaway costs

TechCrunch AI 50%

"The whole conversation shifted from tokenmaxxing and 'go fast' to 'we need guardrails, how do we control this?'"

<https://techcrunch.com/2026/06/05/the-token-bill-comes-due-inside-the-industry-scramble...>

CONF

AirTrunk commits \$30B to build 5GW of AI data centers in India

TechCrunch AI 50%

The Australian data center operator plans to set up 5GW of capacity in India.

<https://techcrunch.com/2026/06/05/airtrunk-commits-30b-to-build-5gw-of-ai-data-centers-...>

CONF

The Meta hack shows there's more to AI security than Mythos

MIT Tech Review 50%

On June 5, 404 Media reported that attackers had been using Meta's AI customer support agent to steal Instagram accounts. Their approach was simple: They asked the agent to link the accounts to email addresses that they controlled, and the agent complied. One attacker broke into...

<https://www.technologyreview.com/2026/06/05/1138437/the-meta-hack-shows-theres-more-to-...>

CONF

This is your laptop... on AI

The Verge AI 50%

We're now deep into developer conference season, and one of the themes so far is the relentless conviction from Big Tech companies that AI is going to change everything about how we do everything. Nvidia's Jensen Huang made that clearer than anyone this week, when he described a...

<https://www.theverge.com/podcast/944058/ai-laptop-nvidia-build-gemini-spark-vergecast>

CONF

This AI startup says it can tell if a script will make a hit film

The Verge AI 50%

When Quilty hit the industry trades earlier this year, the AI startup promised that its tool could accurately predict a film's success just by reading the script. When people actually got a chance to experiment with Quilty's product, though, they were left skeptical. Even with...

<https://www.theverge.com/ai-artificial-intelligence/943531/ai-script-quilty-simon-horsm...>

0

CONF

How Far Did They Go? The Persuasive Tactics of Covert LLM Agents in a Discontinued Field Experiment

ArXiv CS.AI 50%

arXiv:2606.05256v1 Announce Type: new Abstract: This study analyzes a publicly released dataset from a discontinued field experiment on Reddit's r/ChangeMyView. The intervention, conducted by unknown, external researchers and halted following ethical backlash, involved...

<https://arxiv.org/abs/2606.05256>

1

CONF

GITCO: Gated Inference-Time Context Optimization in TSFMs

ArXiv CS.AI 50%

arXiv:2606.05332v1 Announce Type: new Abstract: Patch-based Time Series Foundation Models (TSFMs) suffer from context poisoning: structurally anomalous patches capture disproportionate attention and silently degrade zero-shot forecast quality. We propose improving TSFM accuracy...

<https://arxiv.org/abs/2606.05332>

2

CONF

Uncertainty Aware Functional Behavior Prediction and Material Fatigue Assessment for Circular Factory

ArXiv CS.AI 50%

arXiv:2606.05334v1 Announce Type: new Abstract: Returned products in circular factories re-enter production with heterogeneous degradation states, usage histories, and remaining capability. Reuse cannot be decided from the current inspection alone, because future function...

<https://arxiv.org/abs/2606.05334>

3

CONF

SubtleMemory: A Benchmark for Fine-Grained Relational Memory Discrimination in Long-Horizon AI Agents

ArXiv CS.AI 50%

arXiv:2606.05761v1 Announce Type: new Abstract: Persistent AI assistants, such as OpenClaw, accumulate large collections of related memories over long-term interactions. As these memories grow, they may reinforce one another, diverge across contexts, or directly conflict, making...

<https://arxiv.org/abs/2606.05761>

- 4 **CONF**
- An interpretable and trustworthy AI framework for large-scale longitudinal structure-pain association studies using data from the Osteoarthritis Initiative (OAI)**
- ArXiv CS.AI 50%
- arXiv:2606.05357v1 Announce Type: new Abstract: Purpose: To develop an interpretable and trustworthy AI framework that combines deep learning based MRI Osteoarthritis Knee Score (MOAKS) prediction with interpretable statistical modeling to study structure-pain relationships at...
- <https://arxiv.org/abs/2606.05357>
-
- 5 **CONF**
- Stability vs. Manipulability: Evaluating Robustness Under Post-Decision Interaction in LLM Judges**
- ArXiv CS.AI 50%
- arXiv:2606.05384v1 Announce Type: new Abstract: LLM-as-judge evaluation is widely used in benchmarking pipelines, where model outputs are compared and ranked using automated evaluators. These pipelines typically assume that judgments are stable properties of fixed inputs. We...
- <https://arxiv.org/abs/2606.05384>
-
- 6 **CONF**
- Evaluation of LLMs for Mathematical Formalization in Lean**
- ArXiv CS.AI 50%
- arXiv:2606.05632v1 Announce Type: new Abstract: Within the past few years, the ability of Large Language Models (LLMs) to generate formal mathematical proofs has improved drastically. We provide a comparison of various LLMs' effectiveness in producing formal proofs in Lean 4...
- <https://arxiv.org/abs/2606.05632>
-
- 7 **CONF**
- Reformulating Neural Operators in $d+1$ Dimensions for Embedding Evolution**
- ArXiv CS.AI 50%
- arXiv:2505.11766v4 Announce Type: replace-cross Abstract: Neural Operators (NOs) are powerful architectures for learning mappings between function spaces. While most advances focus on refining kernel parameterizations over the d -dimensional physical domain, the evolution of...
- <https://arxiv.org/abs/2505.11766>
-
- 8 **CONF**
- Channel-Wise Mixed-Precision Quantization for Large Language Models**
- ArXiv CS.AI 50%
- arXiv:2410.13056v4 Announce Type: replace-cross Abstract: Large Language Models (LLMs) have demonstrated remarkable success across a wide range of language tasks, but their deployment on edge devices remains challenging due to the substantial memory requirements imposed by their...
- <https://arxiv.org/abs/2410.13056>
-
- 9 **CONF**
- Cross-Epoch Adaptive Rollout Optimization for RL Post-Training**
- ArXiv CS.AI 50%
- arXiv:2606.05606v1 Announce Type: cross Abstract: LLM post-training often relies on reinforcement learning methods that sample multiple rollouts per prompt, yet most existing approaches use a fixed rollout budget for every prompt, despite large differences in the training signal...
- <https://arxiv.org/abs/2606.05606>
-

0 CONF

Ten Headache Specialists versus Artificial Intelligence for Clinical Literature Summarization: A Critical Evaluation and Comparison

ArXiv CS.AI 50%

arXiv:2606.05436v1 Announce Type: new Abstract: Summarizing the latest medical literature to guide clinical decision-making is essential for evidence-based medicine and high-quality patient care. Yet clinicians face increasing challenges due to limited time with patients and a...

<https://arxiv.org/abs/2606.05436>

1 CONF

Brick-Composer: Using MLLMs for Assembly with Diverse Bricks

ArXiv CS.AI 50%

arXiv:2606.05445v1 Announce Type: new Abstract: We dream of AI agents that can read arbitrary designs and construct real-world objects from reusable building blocks. As a first step toward this vision, we study whether multimodal large language models (MLLMs) possess the visual...

<https://arxiv.org/abs/2606.05445>

2 CONF

Output Type Before Quality: A Standards-Derived XAI Admissibility Rubric for Autonomous-Driving Safety

ArXiv CS.AI 50%

arXiv:2606.05461v1 Announce Type: new Abstract: Safety standards for ML-based autonomous driving specify the kind of evidence an assurance case must contain (directed cause-and-effect chains, quantified interventional effects, named root-cause variables), yet the XAI literature...

<https://arxiv.org/abs/2606.05461>

3 CONF

TAPO: Tool-Aware Policy Optimization via Credit Transfer for Multimodal Search Agents

ArXiv CS.AI 50%

arXiv:2606.05784v1 Announce Type: new Abstract: We identify and formally characterize credit misassignment as a systematic failure mode of GRPO in tool-augmented multimodal search agents: its uniform broadcast of trajectory-level advantages to all tokens causes valuable tool-use...

<https://arxiv.org/abs/2606.05784>

4 CONF

Where Should Knowledge Enter? A Layered Framework for Knowledge Infusion in Multimodal Iterative Generative Mo

ArXiv CS.AI 50%

arXiv:2606.06356v1 Announce Type: new Abstract: Multimodal generative models produce fluent outputs but remain unreliable when generation must respect structured, domain-specific, or safety-critical knowledge. Existing methods incorporate knowledge through mechanisms such as...

<https://arxiv.org/abs/2606.06356>

5 CONF

Individual Gain, Collective Loss: Metacognitive Adaptation in AI-Assisted Creativity

ArXiv CS.AI 50%

arXiv:2606.05532v1 Announce Type: new Abstract: Recent studies reveal a paradox: AI enhances individual creative outputs while reducing collective diversity. Current explanations -- cognitive offloading and over-reliance -- identify symptoms but not mechanisms. We propose...

<https://arxiv.org/abs/2606.05532>

6 CONF

Statistically Reliable LLM-Based Ranking Evaluation via Prediction-Powered Inference

ArXiv CS.AI 50%

arXiv:2606.05308v1 Announce Type: cross Abstract: With PRECISE, we extended Prediction-Powered Inference to produce bias-corrected estimates of ranking evaluation metrics by combining a small human-labeled set with a large LLM-judged set. PPI is provably unbiased regardless of...

<https://arxiv.org/abs/2606.05308>

7 CONF

DAST: A VLM-LLM Framework for Cross-Interface Anomaly Detection in O-RAN

ArXiv CS.AI 50%

arXiv:2606.06261v1 Announce Type: cross Abstract: O-RAN enables a disaggregated baseband stack with programmable functions that communicate over standardized open interfaces. The same openness that enables multi-vendor composition also expands the attack surface across logically...

<https://arxiv.org/abs/2606.06261>

8 CONF

Fix the Mind, Not the Move: Interpretable AI Assistance via Knowledge-Gap Localization

ArXiv CS.AI 50%

arXiv:2606.05602v1 Announce Type: new Abstract: AI assistants in human-AI collaboration often correct suboptimal human actions through behavioral feedback (e.g., alerts or steering-wheel nudges in assistive driving). Such interventions can mitigate immediate errors, but...

<https://arxiv.org/abs/2606.05602>

9 CONF

Multilingual Fine-Tuning via Localized Gradient Conflict Resolution

ArXiv CS.AI 50%

arXiv:2606.05613v1 Announce Type: new Abstract: The rapid evolution of Large Language Models (LLMs) has established cross-lingual versatility as a defining feature of modern systems. However, fine-tuning these models frequently induces negative interference across languages. To...

<https://arxiv.org/abs/2606.05613>

0 CONF

Safety Paradox: How Enhanced Safety Awareness Leaves LLMs Vulnerable to Posterior Attack

ArXiv CS.AI 50%

arXiv:2606.05614v1 Announce Type: new Abstract: Large language models (LLMs) are rigorously aligned to refuse harmful requests, a process that inherently cultivates a latent capacity to evaluate and recognize unsafe content. In this work, we reveal that this advanced safety...

<https://arxiv.org/abs/2606.05614>

1 CONF

A Study of LLMs' Preferences for Libraries and Programming Languages

ArXiv CS.AI 50%

arXiv:2503.17181v4 Announce Type: replace-cross Abstract: Despite the rapid progress of large language models (LLMs) in code generation, existing evaluations focus on functional correctness or syntactic validity, overlooking how LLMs make critical design choices such as which...

<https://arxiv.org/abs/2503.17181>

2 CONF

HypRAG: Hyperbolic Dense Retrieval for Retrieval Augmented Generation

ArXiv CS.AI 50%

arXiv:2602.07739v2 Announce Type: replace-cross Abstract: Embedding geometry plays a fundamental role in retrieval quality, yet dense retrievers for retrieval-augmented generation (RAG) remain largely confined to Euclidean space. However, natural language exhibits hierarchical...

<https://arxiv.org/abs/2602.07739>

3 **CONF****Coding with "Enemy": Can Human Developers Detect AI Agent Sabotage?**

ArXiv CS.AI 50%

arXiv:2606.05647v1 Announce Type: new Abstract: AI coding agents are increasingly embedded in real-world software development, collaborating with human developers while gaining broader access to codebases and tools. This creates a new attack surface: an agent can exploit human...

<https://arxiv.org/abs/2606.05647>
4 **CONF****Continual Learning Bench: Evaluating Frontier AI Systems in Real-World Stateful Environments**

ArXiv CS.AI 50%

arXiv:2606.05661v1 Announce Type: new Abstract: Continual learning, the ability of AI systems to improve through sequential experience, has attracted substantial interest, but no high-quality benchmark exists to evaluate it. We introduce Continual Learning Bench (CL-Bench), the...

<https://arxiv.org/abs/2606.05661>
5 **CONF****Memory is Reconstructed, Not Retrieved: Graph Memory for LLM Agents**

ArXiv CS.AI 50%

arXiv:2606.06036v1 Announce Type: new Abstract: Despite recent progress, LLM agents still struggle with reasoning over long interaction histories. While current memory-augmented agents rely on a static retrieve-then-reason paradigm, this rigid pipeline design prevents them from...

<https://arxiv.org/abs/2606.06036>
6 **CONF****Beyond Output Matching: Preserving Internal Geometry in NVFP4 LLM Distillation**

ArXiv CS.AI 50%

arXiv:2606.05682v1 Announce Type: new Abstract: Demand for low-precision inference, including NVFP4-based approaches, has grown as large language models are increasingly deployed in latency and cost constrained production environments. Quantization-aware distillation (QAD) helps...

<https://arxiv.org/abs/2606.05682>
7 **CONF****PerceptUI: LLM Agents as Human-Aligned Synthetic Users for UI/UX Evaluation**

ArXiv CS.AI 50%

arXiv:2606.05697v1 Announce Type: new Abstract: User interface (UI) and user experience (UX) evaluation is central to product development, yet reliable feedback still relies on recruiting human participants or running online A/B tests, making early-stage iteration slow and...

<https://arxiv.org/abs/2606.05697>
8 **CONF****Self-Augmenting Retrieval for Diffusion Language Models**

ArXiv CS.AI 50%

arXiv:2606.06474v1 Announce Type: cross Abstract: Discrete diffusion language models generate text by iteratively denoising an entire response in parallel. At each step, they predict tentative tokens for every masked position, committing the confident predictions to the output...

<https://arxiv.org/abs/2606.06474>

9 CONF

DiG-Plan: Mitigating Early Commitment for Tool-Graph Planning via Diffusion Guidance

ArXiv CS.AI 50%

arXiv:2606.05728v1 Announce Type: new Abstract: Generating executable tool plans requires selecting appropriate subsets from tool libraries, a combinatorial search problem with an exponentially large solution space. However, we identify a critical misalignment in predominant...

<https://arxiv.org/abs/2606.05728>

0 CONF

When AI Says It Feels

ArXiv CS.AI 50%

arXiv:2606.05734v1 Announce Type: new Abstract: Large language models (LLMs) are generally constrained from expressing feelings through human-preference alignment in post-training processes. This policy is designed using a top-down approach and may conflict with the goal of...

<https://arxiv.org/abs/2606.05734>

1 CONF

Can LLMs Write Correct TLA+ Specifications? Evaluating Natural-Language-to-TLA+ Generation

ArXiv CS.AI 50%

arXiv:2606.05792v1 Announce Type: new Abstract: TLA+ has supported industrial verification at companies such as Amazon and Microsoft, yet writing correct TLA+ specifications from natural language still requires time and expertise, which limits adoption. LLMs show promise, but no...

<https://arxiv.org/abs/2606.05792>

2 CONF

From Risk Classification to Action Plan Remediation: A Guardrail Feedback Driven Framework for LLM Agents

ArXiv CS.AI 50%

arXiv:2606.05805v1 Announce Type: new Abstract: LLM-based guardrails typically safeguard agents by evaluating proposed actions or inputs before execution, producing safety signals such as binary allow/deny decisions, risk categories, and/or explanatory rationales about potential...

<https://arxiv.org/abs/2606.05805>

3 CONF

Entropy-Based Evaluation of AI Agents: A Lightweight Framework for Measuring Behavioral Patterns

ArXiv CS.AI 50%

arXiv:2606.05872v1 Announce Type: new Abstract: AI agents are commonly evaluated using task success, reward, latency, and cost. These metrics are useful, but they often miss important aspects of agent behavior: whether an agent explores too much, repeats itself too rigidly, uses...

<https://arxiv.org/abs/2606.05872>

4 CONF

PAMF: Prior-Aware Multimodal Fusion for Incomplete Time Series Data

ArXiv CS.AI 50%

arXiv:2606.06328v1 Announce Type: cross Abstract: In healthcare, multimodal time series tasks often operate on incomplete observations in practice, for example when ECG segments are lost because electrodes detach or an entire respiratory channel is unavailable during overnight...

<https://arxiv.org/abs/2606.06328>

5 CONF

Retrospective Harness Optimization: Improving LLM Agents via Self-Preference over Trajectory Rollouts

ArXiv CS.AI 50%

arXiv:2606.05922v1 Announce Type: new Abstract: AI agents rely on a harness of skills, tools, and workflows to solve complex problems. Continually improving this harness is essential for adapting to new tasks. However, existing optimization methods typically require ground-truth...

<https://arxiv.org/abs/2606.05922>

6 CONF

SUPERNOVA: Eliciting General Reasoning in LLMs with Reinforcement Learning on Natural Instructions

ArXiv CS.AI 50%

arXiv:2604.08477v2 Announce Type: replace Abstract: Reinforcement Learning with Verifiable Rewards (RLVR) has substantially improved reasoning in formal domains such as mathematics and code, but extending these gains beyond STEM remains challenging. Extending RLVR beyond STEM is...

<https://arxiv.org/abs/2604.08477>

7 CONF

The Self-Correction Illusion: LLMs Correct Others but Not Themselves

ArXiv CS.AI 50%

arXiv:2606.05976v1 Announce Type: new Abstract: Recent work shows that LLM agents struggle to correct errors in their own reasoning traces yet show markedly higher correction rates when identical claims appear under external sources. We ask whether this asymmetry reflects a...

<https://arxiv.org/abs/2606.05976>

8 CONF

Framing, Judging, Steering: An Assessable Competency Model for Teaching Students to Reason With Generative AI

ArXiv CS.AI 50%

arXiv:2606.05983v1 Announce Type: new Abstract: Generative AI makes answers easy and understanding hard, and uncritical use invites cognitive offloading. Schools still measure unaided performance, yet the real task is to produce good work with AI: framing an ill-defined task,...

<https://arxiv.org/abs/2606.05983>

9 CONF

Beyond Similarity: Trustworthy Memory Search for Personal AI Agents

ArXiv CS.AI 50%

arXiv:2606.06054v1 Announce Type: new Abstract: Personal AI agents increasingly rely on long-term memory to provide persistent personalization across sessions. However, existing memory pipelines are largely driven by semantic similarity: memory data close to the current query is...

<https://arxiv.org/abs/2606.06054>

0 CONF

RedditPersona: A Modular Framework for Community-Conditioned LLM Adaptation from Reddit

ArXiv CS.AI 50%

arXiv:2606.06027v1 Announce Type: new Abstract: Community-conditioned language model adaptation requires choices about data collection, community definition, and evaluation that are currently made independently in each study, making it hard to compare assumptions or reuse...

<https://arxiv.org/abs/2606.06027>

1 CONF

A Framework for Measuring Appropriate Reliance on Set-Valued AI Advice

ArXiv CS.AI 50%

arXiv:2606.06081v1 Announce Type: new Abstract: Appropriate reliance on AI advice has become a central research theme in human-AI collaboration. Existing frameworks have focused exclusively on point predictions as AI advice. However, set-valued AI advice (e.g., discrete sets or...

<https://arxiv.org/abs/2606.06081>

2 CONF

Step-adaptive multimodal fusion network with multi-scale cloud feature learning for ultra-short-term solar irradiance forecasting

ArXiv CS.AI 50%

arXiv:2606.06102v1 Announce Type: new Abstract: Ultra-short-term solar irradiance prediction is critical for photovoltaic system dispatch and power grid stability. Existing approaches suffer from three key shortcomings: single time-series models cannot capture the spatial...

<https://arxiv.org/abs/2606.06102>

3 CONF

TRACE: A Temporal Conditional Estimation for Multimodal Time Series Foundation Models

ArXiv CS.AI 50%

arXiv:2606.06285v1 Announce Type: new Abstract: Time series foundation models (TS-FMs) aim to learn generalizable temporal representations that can be adapted to a wide range of downstream tasks. In real-world multimodal settings, time series are frequently affected by temporal...

<https://arxiv.org/abs/2606.06285>

4 CONF

Where does Absolute Position come from in decoder-only Transformers?

ArXiv CS.AI 50%

arXiv:2606.06160v1 Announce Type: new Abstract: RoPE-trained transformers distinguish absolute position in their attention patterns, even though RoPE encodes only relative offsets in the inner product. We trace this leakage to two architectural components, The causal mask is...

<https://arxiv.org/abs/2606.06160>

5 CONF

Learning to replenish: A hybrid deep reinforcement learning for dynamic inventory management in the pharmaceutical supply chains

ArXiv CS.AI 50%

arXiv:2606.06201v1 Announce Type: new Abstract: Pharmaceutical supply chains (PSCs) struggle with inventory management (IM) due to unpredictable demand patterns and variable lead times associated with restocking. This complexity is further compounded by the finite shelf lives of...

<https://arxiv.org/abs/2606.06201>

6 CONF

Evaluating Agentic Configuration Repair for Computer Networks

ArXiv CS.AI 50%

arXiv:2606.06212v1 Announce Type: new Abstract: Misconfigurations in computer networks remain a major source of critical Internet outages. Research is turning to Large Language Models (LLMs) to automate the complex, error-prone task of network configuration. However, even...

<https://arxiv.org/abs/2606.06212>

7 CONF

RedKnot: Efficient Long-Context LLM Serving with Head-Aware KV Reuse and SegPagedAttention

ArXiv CS.AI 50%

arXiv:2606.06256v1 Announce Type: new Abstract: As the input length of large language model (LLM) serving continues to grow, the KV cache has become a dominant bottleneck in AI infrastructure. It limits GPU memory capacity, serving concurrency, cache reuse, and distributed...

<https://arxiv.org/abs/2606.06256>

8 CONF

ToolChoiceConfusion: Causal Minimal Tool Filtering for Reliable LLM Agents

ArXiv CS.AI 50%

arXiv:2606.06284v1 Announce Type: new Abstract: Large language model agents increasingly rely on external tools, but larger tool menus can reduce reliability and efficiency by increasing wrong-tool calls, premature actions, and token cost. Existing tool-selection methods often...

<https://arxiv.org/abs/2606.06284>

9 CONF

Multi-ResNets for Subspace Preconditioning in Constrained Optimization

ArXiv CS.AI 50%

arXiv:2606.06300v1 Announce Type: new Abstract: We propose MResOpt, a staged residual neural network architecture for constrained optimization problems. Our architecture fits within predict-complete-correct pipelines and decomposes constraint satisfaction by priority through...

<https://arxiv.org/abs/2606.06300>

0 CONF

Separation Power of Equivariant Neural Networks

ArXiv CS.AI 50%

arXiv:2406.08966v3 Announce Type: replace-cross Abstract: The separation power of a machine learning model refers to its ability to distinguish between different inputs and is often used as a proxy for its expressivity. Indeed, knowing the separation power of a family of models...

<https://arxiv.org/abs/2406.08966>

1 CONF

LLM Self-Recognition: Steering and Retrieving Activation Signatures

ArXiv CS.AI 50%

arXiv:2606.06315v1 Announce Type: new Abstract: Recent advances in interpretability suggest that large language models (LLMs) implicitly encode signals in their generated text that enable self-recognition of their outputs. We demonstrate that this capability is reliable, even in...

<https://arxiv.org/abs/2606.06315>

2 CONF

CTIConnect: A Benchmark for Retrieval-Augmented LLMs over Heterogeneous Cyber Threat Intelligence

ArXiv CS.AI 50%

arXiv:2510.11974v2 Announce Type: replace-cross Abstract: Cyber Threat Intelligence (CTI) is foundational to modern cybersecurity, enabling organizations to proactively defend against evolving threats. However, the sheer volume and heterogeneity of CTI data, spanning structured...

<https://arxiv.org/abs/2510.11974>

3 CONF

TokenMizer: Graph-Structured Session Memory for Long-Horizon LLM Context Management

ArXiv CS.AI 50%

arXiv:2606.06337v1 Announce Type: new Abstract: Large language model (LLM) deployments for long-horizon tasks face a fundamental constraint: context windows are finite while productive work sessions are not. When history exceeds the Maximum Effective Context Window (MECW),...

<https://arxiv.org/abs/2606.06337>

4 CONF

Boosting Brain-to-Image Decoding with TRIBE v2 Data Augmentation

ArXiv CS.AI 50%

arXiv:2606.06345v1 Announce Type: new Abstract: Brain decoding is limited by the availability of labeled neural data, and remains challenging in low-data regimes. To address this issue, we investigate whether and when brain decoding can be boosted by augmenting small fMRI...

<https://arxiv.org/abs/2606.06345>

5 CONF

In-Training Defenses against Emergent Misalignment in Language Models

ArXiv CS.AI 50%

arXiv:2508.06249v3 Announce Type: replace-cross Abstract: Fine-tuning lets practitioners repurpose aligned large language models (LLMs) for new domains, yet recent work reveals emergent misalignment (EM): Even a small, domain-specific fine-tune can induce harmful behaviors far...

<https://arxiv.org/abs/2508.06249>

6 CONF

Risk Assessment of Autonomous Driving: Integrating Technical Failures, Ethical Dilemmas, and Policy Frameworks

ArXiv CS.AI 50%

arXiv:2606.06396v1 Announce Type: new Abstract: Autonomous driving technology has the potential to reduce the large number of road traffic accidents caused by human error each year, but it also brings new types of risks that need to be evaluated from the aspects of technology,...

<https://arxiv.org/abs/2606.06396>

7 CONF

VASO: Formally Verifiable Self-Evolving Skills for Physical AI Agents

ArXiv CS.AI 50%

arXiv:2606.05395v1 Announce Type: cross Abstract: Reusable robot skills are becoming the basic units through which embodied agents turn open-ended instructions into long-horizon physical behavior. We argue that, while foundation models have collapsed the cost of creating these...

<https://arxiv.org/abs/2606.05395>

8 CONF

Benchmarks in Leipzig

ArXiv CS.AI 50%

arXiv:2606.05818v1 Announce Type: cross Abstract: Between April 1 and May 15, 2026, a group of 49 mathematicians compiled a dataset of research-level mathematics questions with known answers. Most of the work was done during the 3-day workshop *Benchmarks in Leipzig* with 35...

<https://arxiv.org/abs/2606.05818>

9 CONF

MLEvolve: A Self-Evolving Framework for Automated Machine Learning Algorithm Discovery

ArXiv CS.AI 50%

arXiv:2606.06473v1 Announce Type: new Abstract: Large language model (LLM) agents are increasingly applied to long-horizon tasks such as scientific discovery and machine learning engineering (MLE), where sustained self-evolution becomes a key capability. However, existing MLE...

<https://arxiv.org/abs/2606.06473>

0 CONF

RAINO: Anchoring Agents in Reality, A Systematic Review and Conceptual Framework for Realism in Agent-Based Modelling

ArXiv CS.AI 50%

arXiv:2606.05167v1 Announce Type: cross Abstract: Realism is a central yet seemingly under-theorized concept in Agent-Based Modelling. This paper presents a Systematic Literature Review, aiming to identify how realism is currently operationalized and demonstrated. The results...

<https://arxiv.org/abs/2606.05167>

1 CONF

Improving Answer Extraction in Context-based Question Answering Systems Using LLMs

ArXiv CS.AI 50%

arXiv:2606.06197v1 Announce Type: cross Abstract: Question answering (QA) systems have achieved notable progress with the advent of large language models (LLMs). However, they still face challenges in accurately extracting and generating precise answers from given contexts,...

<https://arxiv.org/abs/2606.06197>

2 CONF

The Virtual Roundtable: Multi-Agent Personas Simulating the Dynamics of Human Brainstorming

ArXiv CS.AI 50%

arXiv:2606.05178v1 Announce Type: cross Abstract: As AI-driven product development accelerates, the bottleneck is shifting from how we build to what we build. Traditional human brainstorming faces challenges including groupthink, echo chambers, and limited diversity. To address...

<https://arxiv.org/abs/2606.05178>

3 CONF

From Scoring to Explanations: Evaluating SHAP and LLM Rationales for Rubric-based Teaching Quality Assessment

ArXiv CS.AI 50%

arXiv:2606.05180v1 Announce Type: cross Abstract: Automated scoring models are increasingly used to assign rubric-based quality ratings to complex language performances, including classroom transcripts, yet they typically provide little insight into why a particular score is...

<https://arxiv.org/abs/2606.05180>

4 CONF

SAGE: Scalable AI Governance & Evaluation

ArXiv CS.AI 50%

arXiv:2602.07840v3 Announce Type: replace-cross Abstract: Evaluating relevance in large-scale search systems is fundamentally constrained by the governance gap between nuanced, resource-constrained human oversight and the high-throughput requirements of production systems. While...

<https://arxiv.org/abs/2602.07840>

- 5 **CONF**
- A Cartography of Open Collaboration in Open Source AI: Mapping Practices, Motivations, and Governance in 14 Open Large Language Model Projects**
- ArXiv CS.AI 50%
- arXiv:2509.25397v2 Announce Type: replace-cross Abstract: The proliferation of open large language models (LLMs) is fostering a vibrant ecosystem in artificial intelligence (AI). However, the methods of collaboration used to develop open LLMs, both before and after their public...
- <https://arxiv.org/abs/2509.25397>
-
- 6 **CONF**
- Finite Element-Based Material Learning via Automatic Differentiation: Learning constitutive neural network models from full-field deformation data**
- ArXiv CS.AI 50%
- arXiv:2606.05199v1 Announce Type: cross Abstract: The identification of constitutive neural network models from heterogeneous full-field deformation data provides a robust alternative to traditional calibration methods based on homogeneous stress-strain experiments, particularly...
- <https://arxiv.org/abs/2606.05199>
-
- 7 **CONF**
- Ontology-constrained multi-LLM scoring of hypothesis support in the predictive processing literature**
- ArXiv CS.AI 50%
- arXiv:2606.05206v1 Announce Type: cross Abstract: Fragmentation is common in interdisciplinary fields with diverse methods and theoretical commitments. Predictive coding neuroscience is a clear example: its literature spans computational theory, electrophysiology, imaging,...
- <https://arxiv.org/abs/2606.05206>
-
- 8 **CONF**
- The Invisible Hand of Physics: When Video Diffusion Models Know More Than They Show**
- ArXiv CS.AI 50%
- arXiv:2606.05328v1 Announce Type: cross Abstract: Modern video diffusion models generate increasingly realistic and temporally coherent videos, motivating their use as candidate world simulators. Yet it remains unclear whether these models internally encode physical structure,...
- <https://arxiv.org/abs/2606.05328>
-
- 9 **CONF**
- Where's the Structure? A Systematic Literature Review of Empirical Research on Human-AI Collaboration and Hybrid Intelligence for Learning**
- ArXiv CS.AI 50%
- arXiv:2606.05222v1 Announce Type: cross Abstract: Artificial intelligence (AI) has been applied across educational contexts to support learning. One approach to such support is "human-AI collaboration" (also termed "hybrid intelligence"), where human(s) and AI components...
- <https://arxiv.org/abs/2606.05222>
-
- 0 **CONF**
- Domain-Conditioned Safety in Frontier Computer-Using Agents: A 793-Episode Browser Benchmark, a Coding-Domain Cross-Reference, and a Reproducibility Audit of Recent Red-Teaming**
- ArXiv CS.AI 50%
- arXiv:2606.05233v1 Announce Type: cross Abstract: Recent computer-using-agent (CUA) red-teaming papers report prompt-injection attack success rates (ASR) of 42-98%, but these headline numbers cluster on retired models and on the most-vulnerable model in each paper's panel. We...
- <https://arxiv.org/abs/2606.05233>

1 CONF

Search-Time Contamination in Deep Research Agents: Measuring Performance Inflation in Public Benchmark Evaluation

ArXiv CS.AI 50%

arXiv:2606.05241v1 Announce Type: cross Abstract: Public benchmarks enable fair and reproducible evaluation of LLM reasoning, but they become fragile for deep research agents that actively search the web during inference. Such agents may retrieve public benchmark metadata,...

<https://arxiv.org/abs/2606.05241>

2 CONF

Personal AI Agent for Camera Roll VQA

ArXiv CS.AI 50%

arXiv:2606.05275v1 Announce Type: cross Abstract: We study the personal camera roll visual question answering setting. In this setting, a conversational AI assistant can access a user's personal camera roll and retrieve relevant photos to answer queries, ranging from simple...

<https://arxiv.org/abs/2606.05275>

3 CONF

Pattern Selectivity is Not Task-Causal Structure: A Cross-Architecture Mechanistic Study of Composed-Task Circuits in 1B-Class Language Models

ArXiv CS.AI 50%

arXiv:2606.05378v1 Announce Type: cross Abstract: We test whether a single screen-and-ablate recipe -- identify attention-head circuits by task-pattern selectivity, then verify by causal ablation against a matched-random null -- produces consistent mechanistic claims across...

<https://arxiv.org/abs/2606.05378>

4 CONF

Can AI Refute Economic Theory? Evidence from Beyond the Knowledge Cutoff

ArXiv CS.AI 50%

arXiv:2606.05383v1 Announce Type: cross Abstract: Can artificial intelligence (AI) refute economic theory? I document experiments in which I asked several AI models (Gemini, Refine, Claude, and ChatGPT) to check the correctness of four published papers in economic theory, each...

<https://arxiv.org/abs/2606.05383>

5 CONF

Willing but Unable: Separating Refusal from Capability in Code LLMs via Abliteration

ArXiv CS.AI 50%

arXiv:2606.05396v1 Announce Type: cross Abstract: Producing a labeled vulnerable code at scale is a recurring obstacle for learning-based vulnerability detection: mined corpora carry substantial label noise, and existing LLM-based augmentation propagates these inaccuracies...

<https://arxiv.org/abs/2606.05396>

6 CONF

A2RAG: Adaptive Agentic Graph Retrieval for Cost-Aware and Reliable Reasoning

ArXiv CS.AI 50%

arXiv:2601.21162v2 Announce Type: replace-cross Abstract: Graph Retrieval-Augmented Generation (Graph-RAG) enhances multihop question answering by organizing corpora into knowledge graphs and routing evidence through relational structure. However, practical deployments face two...

<https://arxiv.org/abs/2601.21162>

7 CONF

When Evidence is Sparse: Weakly Supervised Early Failure Alerting in Dialogs and LLM-Agent Trajectories

ArXiv CS.AI 50%

arXiv:2606.05414v1 Announce Type: cross Abstract: Early failure alerting requires deciding, while a dialog or agent trajectory is still unfolding, whether to flag it as likely to fail. This is challenging because supervision is typically available only as a trajectory-level...

<https://arxiv.org/abs/2606.05414>

8 CONF

Selective-Advantage Entropy-Adaptive Horizon GRPO: Asymmetric Token-Level Discounting for Efficient Reinforcement Learning of Language Models

ArXiv CS.AI 50%

arXiv:2606.05434v1 Announce Type: cross Abstract: Group Relative Policy Optimisation (GRPO) has emerged as an effective reinforcement-learning algorithm for aligning language models on reasoning tasks, but it treats every token position and every sampled rollout symmetrically....

<https://arxiv.org/abs/2606.05434>

9 CONF

GOTabPFN: From Feature Ordering to Compact Tokenization for Tabular Foundation Models on High-Dimensional Data

ArXiv CS.AI 50%

arXiv:2606.05441v1 Announce Type: cross Abstract: We investigate how to make small tabular foundation models effective for High-Dimensional, Low-Sample Size (HDLSS) tabular prediction without retraining large backbones. We introduce Graph-guided Ordering with Local Refinement...

<https://arxiv.org/abs/2606.05441>

0 CONF

Towards Unified and Data-Efficient Prognostics and Health Management with Tabular Foundation Models

ArXiv CS.AI 50%

arXiv:2606.05481v1 Announce Type: cross Abstract: Data-driven Prognostics and Health Management (PHM) uses time-varying condition-monitoring data to diagnose system states and estimate remaining useful life in engineered assets. These tasks are central to maintenance planning,...

<https://arxiv.org/abs/2606.05481>

1 CONF

The Role of Instructional Guidance in Generative AI-Assisted Learning: Empirical Evidence from Construction Engineering Education

ArXiv CS.AI 50%

arXiv:2606.05509v1 Announce Type: cross Abstract: Generative artificial intelligence (AI) is increasingly used to support self-directed learning, yet student interaction with such systems often remains unstructured, limiting engagement in deeper cognitive processes. This study...

<https://arxiv.org/abs/2606.05509>

2 CONF

Exploring LLMs for South Asian Music Understanding and Generation

ArXiv CS.AI 50%

arXiv:2606.05522v1 Announce Type: cross Abstract: Recent advancements in Large Language Models (LLMs) have shown promising results in music understanding and generation tasks. However, existing works remain confined to Western tonal traditions, offering little insight into...

<https://arxiv.org/abs/2606.05522>

3 **CONF****ADK Arena: Evaluating Agent Development Kits via LLM-as-a-Developer****ArXiv CS.AI 50%**

arXiv:2606.05548v1 Announce Type: cross Abstract: The rapid proliferation of Agent Development Kits (ADKs), SDK-level frameworks for building LLM-powered autonomous agents, has outpaced any empirical understanding of how framework choice affects agent performance. We propose...

<https://arxiv.org/abs/2606.05548>

4 **CONF****Representation Learning Enables Scalable Multitask Deep Reinforcement Learning****ArXiv CS.AI 50%**

arXiv:2606.05555v1 Announce Type: cross Abstract: Scaling reinforcement learning (RL) to diverse multitask settings remains a central challenge. While recent advances in model-based RL achieve strong performance, they rely on planning and complex training pipelines, making it...

<https://arxiv.org/abs/2606.05555>

5 **CONF****HDST-GNN: Heterogeneous Dynamic Spatiotemporal Graph Neural Networks for Multi-Object Tracking in UAV Aerial Imagery****ArXiv CS.AI 50%**

arXiv:2606.05587v1 Announce Type: cross Abstract: Multi-object tracking (MOT) from UAV imagery presents unique challenges: altitude varies across sequences, objects are small and densely packed, and frequent occlusion causes identity switches. Existing graph-based trackers...

<https://arxiv.org/abs/2606.05587>

6 **CONF****The End of Software Engineering: How AI Agents Are Fundamentally Restructuring the Software Paradigm****ArXiv CS.AI 50%**

arXiv:2606.05608v1 Announce Type: cross Abstract: For over half a century, software engineering has operated on a foundational premise: human engineers decompose problems, encode decision logic into static code, and manually adapt that code as requirements evolve. This paper...

<https://arxiv.org/abs/2606.05608>

7 **CONF****When Surface Form Changes Moderation Decisions: A Paired Study of Code-Mixed Workflow Instability****ArXiv CS.AI 50%**

arXiv:2606.05654v1 Announce Type: cross Abstract: Hate moderation is often evaluated as classification on clean English inputs, but deployed systems must route content to actions such as ALLOW, FLAG, or REVIEW. We study how this workflow changes under code-mixed inputs using a...

<https://arxiv.org/abs/2606.05654>

8 **CONF****Agent-Orchestrated Adaptive RAG: A Comparative Study on Structured and Multi-Hop Retrieval****ArXiv CS.AI 50%**

arXiv:2606.05658v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) enhances Large Language Models (LLMs) by grounding their responses in external knowledge, but conventional pipelines rely on static, single-step retrieval that limits performance on complex...

<https://arxiv.org/abs/2606.05658>

9 CONF

Safe Embodied AI for Long-horizon Tasks: A Cross-layer Analysis of Robotic Manipulation

ArXiv CS.AI 50%

arXiv:2606.05660v1 Announce Type: cross Abstract: Embodied AI systems are increasingly expected to reason and act over extended horizons in physical environments. This growing capability brings safety to the foreground, because failures in the physical world can harm people,...

<https://arxiv.org/abs/2606.05660>

00 CONF

Value-and-Structure Alignment for Routing-Consistent Quantization of Mixture-of-Experts Models

ArXiv CS.AI 50%

arXiv:2606.05688v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) models scale foundation models efficiently by activating only a subset of experts for each token, but their large number of expert parameters still makes quantization essential for practical deployment....

<https://arxiv.org/abs/2606.05688>

01 CONF

Cognitive Threat Intelligence and Explainable Federated Security Analytics for distributed Infrastructure Systems

ArXiv CS.AI 50%

arXiv:2606.05701v1 Announce Type: cross Abstract: The increasing adoption of distributed infrastructure systems, cloud computing, Internet of Things (IoT) technologies, and edge-based architectures has significantly expanded the cybersecurity attack surface and introduced...

<https://arxiv.org/abs/2606.05701>

02 CONF

Explainable AI-Driven Cyber Risk Analytics and Model Reliability Assessment for Intelligent Governance of U.S. Critical Infrastructure: An XGBoost and SHAP-Based Intrusion Detection Framework

ArXiv CS.AI 50%

arXiv:2606.05710v1 Announce Type: cross Abstract: The increasing penetrations of the critical infrastructure sector in the United States with intelligent digital technologies have greatly increased exposure to advanced cyber adversaries and operational vulnerabilities....

<https://arxiv.org/abs/2606.05710>

03 CONF

ViCuR: Visual Cues as Recoverable Privilege for Multimodal On-Policy Distillation

ArXiv CS.AI 50%

arXiv:2606.05718v1 Announce Type: cross Abstract: On-policy distillation (OPD) improves reasoning by training a student on trajectories sampled from its own policy under supervision from a teacher. In multimodal reasoning, a common extension is to use a privileged teacher that...

<https://arxiv.org/abs/2606.05718>

04 CONF

UNIVID: Unified Vision-Language Model for Video Moderation

ArXiv CS.AI 50%

arXiv:2606.05748v1 Announce Type: cross Abstract: Global-scale video moderation faces a dual challenge: the need for fine-grained multi-modal reasoning and the demand for interpretable outputs to support downstream enforcement. Traditional moderation systems often rely on...

<https://arxiv.org/abs/2606.05748>

05 CONF

SagnacAssisted Enhanced OTDR for Distributed Acoustic Sensing: A Standardized Benchmark and Engineering Evaluation Framework

ArXiv CS.AI 50%

arXiv:2606.05754v1 Announce Type: cross Abstract: Phase-sensitive optical time-domain reflectometry (ϕ -OTDR) is widely used in large-scale distributed acoustic sensing (DAS) because it provides distributed spatiotemporal monitoring over long sensing distances. Its field...

<https://arxiv.org/abs/2606.05754>

06 CONF

Beyond Soft Masks: Hard-Perturbation Mixup Explainer for Robust GNN Explainability

ArXiv CS.AI 50%

arXiv:2606.05756v1 Announce Type: cross Abstract: Graph Neural Networks (GNNs) have demonstrated remarkable performance across a range of applications involving graph-structured data, particularly in high-stakes domains. However, the opaque nature of their decision-making...

<https://arxiv.org/abs/2606.05756>

07 CONF

Human Oversight and Overload: Two Hidden and Costly Burdens of AI-Assisted Software Engineering

ArXiv CS.AI 50%

arXiv:2606.05770v1 Announce Type: cross Abstract: AI is changing how software engineers work, but it often comes with hidden burdens and costs. In this paper, we characterize two such often-overlooked burdens: (1) the constant need for human oversight and inspection of...

<https://arxiv.org/abs/2606.05770>

08 CONF

CollabBench: Benchmarking and Unleashing Collaborative Ability of LLMs with Diverse Players via Proactive Engagement

ArXiv CS.AI 50%

arXiv:2606.05793v1 Announce Type: cross Abstract: While LLM-based agents excel at individual tasks, effective collaboration with realistic human partners remains challenging. Most of the existing conversation-level collaborative studies lack grounded interaction and behavioral...

<https://arxiv.org/abs/2606.05793>

09 CONF

Emotion-Aware Image Generation from Korean Diary Text via LLM-based Prompt Translation and LoRA Fine-Tuning

ArXiv CS.AI 50%

arXiv:2606.05816v1 Announce Type: cross Abstract: T2I models cannot effectively capture sentiment from various types of text, including diaries, as they primarily focus on visual object-related patterns rather than contextual emotional understanding. This paper proposes an...

<https://arxiv.org/abs/2606.05816>

10 CONF

Consistency Training Along the Transformer Stack

ArXiv CS.AI 50%

arXiv:2606.05817v1 Announce Type: cross Abstract: Consistency training encourages models to behave similarly across different contexts, and has shown promise for reducing misalignment. We broaden the scope of consistency training in two ways. First, we introduce two new internal...

<https://arxiv.org/abs/2606.05817>

11 CONF

Mechanistic Insights into Functional Sparsity in Multimodal LLMs via CoRe Heads

ArXiv CS.AI 50%

arXiv:2606.05843v1 Announce Type: cross Abstract: While Multimodal Large Language Models (MLLMs) demonstrate remarkable proficiency on complex vision-language tasks, the mechanisms by which they extract query-relevant visual features from complex, noisy contexts remain opaque....

<https://arxiv.org/abs/2606.05843>

12 CONF

Deciphering Two Training Clocks in Grokking via Deep Linear Network Theory with Conditional ReLU Reduction

ArXiv CS.AI 50%

arXiv:2606.05863v1 Announce Type: cross Abstract: Grokking suggests that fitting the training data and learning a simple underlying rule may occur on different time scales. We formalize this phenomenon by separating the fast decay of the classification loss from the slower...

<https://arxiv.org/abs/2606.05863>

13 CONF

Staying with the Uncertainty: Uncertainty-Scaffolding Strategies for Artificial Moral Advisors in LLM-to-LLM Simulated Conversations

ArXiv CS.AI 50%

arXiv:2606.05890v1 Announce Type: cross Abstract: LLMs are increasingly deployed as Artificial Moral Advisors (AMA) in a variety of contexts: what kind of conversational patterns should they display? In this paper, we study how AMA can help their interlocutors "stay with the..."

<https://arxiv.org/abs/2606.05890>

14 CONF

Better Literary Translation: A Multi-Aspect Data Generation and LLM Training Approach

ArXiv CS.AI 50%

arXiv:2606.05924v1 Announce Type: cross Abstract: Literary translation poses unique challenges due to the scarcity of high-quality annotated data and the need to balance expression fluency with literary effect. We present a multi-aspect iterative refinement framework that...

<https://arxiv.org/abs/2606.05924>

15 CONF

To Be Multimodal or Not to Be: Query-Adaptive Audio-Visual Person Retrieval via Active Modality Detection

ArXiv CS.AI 50%

arXiv:2606.05931v1 Announce Type: cross Abstract: When retrieving a person from a video archive by voice and face, should the system be multimodal or not? In real-world broadcast archives, unlike curated benchmarks, a target may be heard but unseen, seen but unheard, or both....

<https://arxiv.org/abs/2606.05931>

16 CONF

Causal Scaffolding for Physical Reasoning: A Benchmark for Causally-Informed Physical World Understanding in VLMs

ArXiv CS.AI 50%

arXiv:2606.05966v1 Announce Type: cross Abstract: Understanding and reasoning about the physical world is the foundation of intelligent behavior, yet state-of-the-art vision-language models (VLMs) still fail at causal physical reasoning, often producing plausible but incorrect...

<https://arxiv.org/abs/2606.05966>

17 CONF

Measuring the sensitivity of LLM-based structured extraction to prompt, model, and schema choices in clinical discharge summaries

ArXiv CS.AI 50%

arXiv:2606.05970v1 Announce Type: cross Abstract: Large language models are increasingly used for structured extraction from clinical free-text notes, but the sensitivity of their output to upstream configuration choices is less understood than their accuracy on fixed...

<https://arxiv.org/abs/2606.05970>

18 CONF

ATT-CR: Adaptive Triangular Transformer for Cloud Removal

ArXiv CS.AI 50%

arXiv:2606.05999v1 Announce Type: cross Abstract: Cloud removal aims to accurately reconstruct the ground objects obscured by clouds in remote sensing images. Existing Transformer-based methods utilizing self-attention have shown impressive results by effectively modeling...

<https://arxiv.org/abs/2606.05999>

19 CONF

MDP-GRPO: Stabilized Group Relative Policy Optimization for Multi-Constraint Instruction Following

ArXiv CS.AI 50%

arXiv:2606.06058v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards is ideal for multi-constraint instruction following, yet standard group-relative policy optimization (GRPO) becomes unstable under discrete, low-dispersion rewards, where...

<https://arxiv.org/abs/2606.06058>

20 CONF

LatentSkill: From In-Context Textual Skills to In-Weight Latent Skills for LLM Agents

ArXiv CS.AI 50%

arXiv:2606.06087v1 Announce Type: cross Abstract: Agent systems increasingly use textual skills to encode reusable task procedures, but injecting these skills into the prompt at every step incurs substantial context overhead and exposes skill content as plaintext. We present...

<https://arxiv.org/abs/2606.06087>

21 CONF

ITP-STDP: An Intrinsic-Timing Power-of-Two Learning Engine for On-Chip SNN Training

ArXiv CS.AI 50%

arXiv:2606.06159v1 Announce Type: cross Abstract: Spiking neural networks (SNNs) have the potential to emerge as the third generation of neural networks and have attracted increasing attention across a wide range of applications. However, the large number of synaptic connections...

<https://arxiv.org/abs/2606.06159>

22 CONF

Learning to Route LLMs from Implicit Cost-Performance Preferences via Meta-Learning

ArXiv CS.AI 50%

arXiv:2606.06178v1 Announce Type: cross Abstract: Large language models (LLMs) present a trade-off between performance and cost, where more powerful models incur greater expense. LLM routing aims to mitigate expenses while maintaining performance by sending queries to the most...

<https://arxiv.org/abs/2606.06178>

23 CONF

Dense Contexts Are Hard Contexts: Lexical Density Limits Effective Context in LLMs

ArXiv CS.AI 50%

arXiv:2606.06203v1 Announce Type: cross Abstract: Input length and the position of relevant information are widely cited as the primary causes of degraded LLM long-context performance. Here, we study lexical density -- the rate at which a context introduces distinct information...

<https://arxiv.org/abs/2606.06203>

24 CONF

Towards the Readability of LLM-Generated Codes through Multitask Representation Engineering

ArXiv CS.AI 50%

arXiv:2606.06214v1 Announce Type: cross Abstract: Correctness and readability are key measures of code quality, respectively ensuring functional fidelity and ease of comprehension. While most existing research focuses on improving the correctness of large language models~(LLMs)...

<https://arxiv.org/abs/2606.06214>

25 CONF

Design a Reliable LLM-Integrated Interface for Mortality Forecasting

ArXiv CS.AI 50%

arXiv:2606.06235v1 Announce Type: cross Abstract: Mortality forecasting plays an important role in actuarial and policy decision-making, but its implementation remains technically complex and inaccessible to non-expert users. This project proposes a reliable large language model...

<https://arxiv.org/abs/2606.06235>

26 CONF

TOKI: A Bitemporal Operator Algebra for Contradiction Resolution in LLM-Agent Persistent Memory

ArXiv CS.AI 50%

arXiv:2606.06240v1 Announce Type: cross Abstract: Persistent memory for an LLM agent is a write-heavy substrate: every belief update is a versioned write, and a new claim may contradict a stored one. Production systems use four resolution heuristics (last-writer-wins,...

<https://arxiv.org/abs/2606.06240>

27 CONF

Benchmarking Open-Source Layout Detection Models for Data Snapshot Extraction from Institutional Documents

ArXiv CS.AI 50%

arXiv:2606.06242v1 Announce Type: cross Abstract: Institutional documents contain substantial amounts of operational and analytical information embedded within figures and tables. Current approaches for extracting visual content from documents are largely built around generic...

<https://arxiv.org/abs/2606.06242>

28 CONF

LLMs Can Leak Training Data But Do They Want To? A Propensity-Aware Evaluation of Memorization in LLMs

ArXiv CS.AI 50%

arXiv:2606.06286v1 Announce Type: cross Abstract: Large language models can reproduce training data, but existing memorization evaluations mostly measure whether models can be forced to do so, rather than whether they do so under ordinary use. We introduce PropMe, a...

<https://arxiv.org/abs/2606.06286>

29 **CONF****Plug-and-Play Guidance for Discrete Diffusion Models via Gradient-Informed Logit Correction**

ArXiv CS.AI 50%

arXiv:2606.06303v1 Announce Type: cross Abstract: Controllable generation with discrete diffusion models is often hindered by high computational overhead or the need for retraining. In this paper, we present $\nabla_{\textbf{G}}$ -informed...

<https://arxiv.org/abs/2606.06303>
30 **CONF****Learning What to Forget: Improving LLM Unlearning via Learned Token-Level Importance**

ArXiv CS.AI 50%

arXiv:2606.06320v1 Announce Type: cross Abstract: Machine unlearning aims to remove targeted knowledge from a trained model while preserving its general capabilities. For autoregressive language models, not all tokens in a forget sample are equally relevant to forgetting....

<https://arxiv.org/abs/2606.06320>
31 **CONF****Bridging Domain Expertise and Generalization for Performance Estimation**

ArXiv CS.AI 50%

arXiv:2606.06335v1 Announce Type: cross Abstract: Performance estimation under distribution shift aims to predict how a model behaves on an unlabeled test set whose distribution differs from the training data, a scenario that requires reliable indicators that can faithfully...

<https://arxiv.org/abs/2606.06335>
32 **CONF****LatentWave: JEPA Pretraining for Wireless Foundation Models**

ArXiv CS.AI 50%

arXiv:2606.06373v1 Announce Type: cross Abstract: Wireless foundation models have emerged as a promising alternative to building separate models for each wireless task. However, existing approaches rely on masked input reconstruction, which can bias representations toward...

<https://arxiv.org/abs/2606.06373>
33 **CONF****EasyLens: A Training-Free Plug-and-Play Subtle-Lesion Representation Amplifier for Medical Vision-Language Models**

ArXiv CS.AI 50%

arXiv:2606.06379v1 Announce Type: cross Abstract: Medical vision-language models (VLMs) have shown increasing potential for clinical image interpretation, including lesion detection and report generation. However, their practical utility remains limited by insufficient...

<https://arxiv.org/abs/2606.06379>
34 **CONF****Emergent Language as an Approach to Conscious AI**

ArXiv CS.AI 50%

arXiv:2606.06380v1 Announce Type: cross Abstract: The question of whether artificial systems can be conscious remains open, in part because existing approaches either evaluate systems against theory-derived checklists (discriminative) or engineer consciousness-inspired modules...

<https://arxiv.org/abs/2606.06380>

35 CONF

Will the Agent Recuse Itself? Measuring LLM-Agent Compliance with In-Band Access-Deny Signals

ArXiv CS.AI 50%

arXiv:2606.06460v1 Announce Type: cross Abstract: As autonomous LLM agents increasingly hold real credentials and operate infrastructure without a human in the loop, operators have no standard way to tell an agent that a resource is off-limits. Access controls either let the...

<https://arxiv.org/abs/2606.06460>

36 CONF

PC Layer: Polynomial Weight Preconditioning for Improving LLM Pre-Training

ArXiv CS.AI 50%

arXiv:2606.06470v1 Announce Type: cross Abstract: We propose a preconditioning (PC) layer, a weight parameterization via polynomial preconditioner that ensures stable weight conditioning throughout LLM training. The PC module reshapes the singular-value spectrum of weight...

<https://arxiv.org/abs/2606.06470>

37 CONF

Operation-Guided Progressive Human-to-AI Text Transformation Benchmark for Multi-Granularity AI-Text Detection

ArXiv CS.AI 50%

arXiv:2606.06481v1 Announce Type: cross Abstract: As AI writing assistants become increasingly integrated into real-world drafting and revision workflows, many documents are no longer purely human-written or AI-generated, but instead result from progressive human-AI co-editing....

<https://arxiv.org/abs/2606.06481>

38 CONF

Code2LoRA: Hypernetwork-Generated Adapters for Code Language Models under Software Evolution

ArXiv CS.AI 50%

arXiv:2606.06492v1 Announce Type: cross Abstract: Code language models need repository-level context to resolve imports, APIs, and project conventions. Existing methods inject this knowledge as long inputs (retrieved through RAG or dependency analysis) or through per-repository...

<https://arxiv.org/abs/2606.06492>

39 CONF

Towards AI epidemiology: a measurement standardisation framework for prospective risk detection

ArXiv CS.AI 50%

arXiv:2512.15783v3 Announce Type: replace Abstract: This paper proposes a measurement standardisation framework that compresses expert-AI interactions into structured, comparable fields for prospective risk detection in deployed AI systems, without access to model internals. The...

<https://arxiv.org/abs/2512.15783>

40 CONF

DPBench: Structural Determinants of Multi-Agent LLM Coordination Under Simultaneous Resource Contention

ArXiv CS.AI 50%

arXiv:2602.13255v2 Announce Type: replace Abstract: We present DPBench, a benchmark for evaluating coordination in multi-agent systems built from large language models. Existing benchmarks measure task-level success under a fixed protocol; the structural conditions under which...

<https://arxiv.org/abs/2602.13255>

41 CONF

CuTeGen: An LLM-Based Agentic Framework for Generation and Optimization of High-Performance GPU Kernels using CuTe

ArXiv CS.AI 50%

arXiv:2604.01489v2 Announce Type: replace-cross Abstract: High-performance GPU kernels are critical to modern machine learning systems, yet developing them remains a manual, expert-driven process. Recent work has explored using LLMs to automate kernel generation, but generated...

<https://arxiv.org/abs/2604.01489>

42 CONF

Semantic Partial Grounding via LLMs

ArXiv CS.AI 50%

arXiv:2602.22067v2 Announce Type: replace Abstract: Grounding is a critical step in classical planning, yet it often becomes a computational bottleneck due to the exponential growth in grounded actions and atoms as task size increases. Recent advances in partial grounding have...

<https://arxiv.org/abs/2602.22067>

43 CONF

Binary Gaussian Copula Synthesis: an LLM-powered data augmentation framework for early dialysis prediction in chronic kidney disease

ArXiv CS.AI 50%

arXiv:2403.00965v2 Announce Type: replace-cross Abstract: Only a small fraction of patients with chronic kidney disease (CKD) progress to dialysis, creating severe class imbalance that limits the performance of machine learning models for early dialysis prediction. This...

<https://arxiv.org/abs/2403.00965>

44 CONF

Comprehensive and Reliable Feature Attribution for Diverse Modalities and Models via Frequency-Domain Insights

ArXiv CS.AI 50%

arXiv:2411.18343v3 Announce Type: replace-cross Abstract: Personalized Federal learning(PFL) allows clients to cooperatively train a personalized model without disclosing their private dataset. However, PFL suffers from Non-IID, heterogeneous devices, lack of fairness, and...

<https://arxiv.org/abs/2411.18343>

45 CONF

The Mirage of Performance Gains: Why Contrastive Decoding Fails to Mitigate Object Hallucinations in MLLMs?

ArXiv CS.AI 50%

arXiv:2504.10020v4 Announce Type: replace-cross Abstract: Contrastive decoding strategies are widely used to reduce object hallucinations in multimodal large language models (MLLMs). These methods work by constructing contrastive samples to induce hallucinations and then...

<https://arxiv.org/abs/2504.10020>

46 CONF

CLASH: Evaluating Language Models on Judging High-Stakes Dilemmas from Multiple Perspectives

ArXiv CS.AI 50%

arXiv:2504.10823v4 Announce Type: replace-cross Abstract: Navigating dilemmas involving conflicting values is challenging even for humans in high-stakes domains, let alone for AI, yet prior work has been limited to everyday scenarios. To close this gap, we introduce CLASH...

<https://arxiv.org/abs/2504.10823>

47 CONF

RAG Security and Privacy: Formalizing the Threat Model and Attack Surface

ArXiv CS.AI 50%

arXiv:2509.20324v2 Announce Type: replace-cross Abstract: Retrieval-Augmented Generation (RAG) is an emerging approach in natural language processing that combines large language models (LLMs) with external document retrieval to produce more accurate and grounded responses....

<https://arxiv.org/abs/2509.20324>

48 CONF

Correcting Prompt Dependence in LLM Benchmarks: A Bayesian Hierarchical Model with Embedding-Space Clustering

ArXiv CS.AI 50%

arXiv:2510.05709v2 Announce Type: replace-cross Abstract: LLM benchmarking metrics often misstate performance and uncertainty as they rely on two assumptions that frequently do not hold in practice: (i) a sufficient number of evaluations are available for classical inference,...

<https://arxiv.org/abs/2510.05709>

49 CONF

Using street view images and visual LLMs to predict heritage values for governance support: Risks, ethics, and policy implications

ArXiv CS.AI 50%

arXiv:2601.06056v2 Announce Type: replace-cross Abstract: During 2025 and 2026, the Energy Performance of Buildings Directive is being implemented in the European Union member states, requiring all member states to have National Building Renovation Plans. In Sweden, there is no...

<https://arxiv.org/abs/2601.06056>

50 CONF

Aligning Tree-Search Policies with Fixed Token Budgets in Test-Time Scaling of LLMs

ArXiv CS.AI 50%

arXiv:2602.09574v2 Announce Type: replace-cross Abstract: Tree-search decoding is an effective form of test-time scaling for large language models (LLMs), but real-world deployment often imposes a fixed per-query token budget that varies across settings. Existing tree-search...

<https://arxiv.org/abs/2602.09574>

51 CONF

ContactExplorer: Contact Coverage-Guided Exploration for General-Purpose Dexterous Manipulation

ArXiv CS.AI 50%

arXiv:2603.10971v3 Announce Type: replace-cross Abstract: Reinforcement learning has achieved remarkable success in domains such as Atari games, navigation, and locomotion, where exploration can often be guided by novelty over states or dynamics. In contrast, dexterous...

<https://arxiv.org/abs/2603.10971>

52 CONF

LeWorldModel: Stable End-to-End Joint-Embedding Predictive Architecture from Pixels

ArXiv CS.AI 50%

arXiv:2603.19312v3 Announce Type: replace-cross Abstract: Joint Embedding Predictive Architectures (JEPAs) offer a compelling framework for learning world models in compact latent spaces, yet existing methods remain fragile, relying on complex multi-term losses, exponential...

<https://arxiv.org/abs/2603.19312>

53 CONF

Brain-CLIPLM: Semantic Compression for EEG-to-Text Decoding

ArXiv CS.AI 50%

arXiv:2604.16370v2 Announce Type: replace-cross Abstract: Decoding natural language from non-invasive electroencephalography (EEG) remains constrained by low signal-to-noise ratio and limited information bandwidth. This raises a central question: can sentence-level language be...

<https://arxiv.org/abs/2604.16370>

54 CONF

Frontier Lag: A Bibliometric Audit of Capability Misrepresentation in Academic AI Evaluation

ArXiv CS.AI 50%

arXiv:2605.04135v2 Announce Type: replace-cross Abstract: Readers of applied-domain LLM capability evaluations want to know what AI systems can currently do. That literature answers a related, but consequentially different, question: what older, cheaper, less-elicited models...

<https://arxiv.org/abs/2605.04135>

55 CONF

Path-Coupled Bellman Flows for Distributional Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.08253v2 Announce Type: replace-cross Abstract: Distributional reinforcement learning (DRL) models the full return distribution, but existing finite-support or quantile-based methods rely on projections, while recent flow-based approaches can suffer from boundary...

<https://arxiv.org/abs/2605.08253>

56 CONF

Macro: Enhancing Multilingual Counterfactual Explanations through Alignment-as-Preference Optimization

ArXiv CS.AI 50%

arXiv:2605.11632v2 Announce Type: replace-cross Abstract: Self-generated counterfactual explanations (SCEs) are minimally modified inputs (minimality) generated by large language models (LLMs) that flip their own predictions (validity), offering a causally grounded approach to...

<https://arxiv.org/abs/2605.11632>

57 CONF

Spectral Probe-Circuits: A Three-Step Recipe for Identifying Attention-Head Circuits in Pretrained Transformers

ArXiv CS.AI 50%

arXiv:2605.24059v2 Announce Type: replace-cross Abstract: We present a three-step recipe for identifying attention-head circuits in pretrained transformers. A per-head spectral signal -- the time-integrated participation ratio of each head's attention output -- ranks heads doing...

<https://arxiv.org/abs/2605.24059>

58 CONF

CoEval: Ranking Language Models for Custom Tasks Without Labeled Data or Trustworthy Benchmarks

ArXiv CS.AI 50%

arXiv:2606.03650v2 Announce Type: replace-cross Abstract: Selecting a pretrained language model, or evaluating a fine-tuned one, for a specific application is a high-value decision, yet the public benchmarks used to make it are poorly suited: a generic benchmark need not reflect...

<https://arxiv.org/abs/2606.03650>

59 CONF

Do Transformers Need Three Projections? Systematic Study of QKV Variants

ArXiv CS.AI 50%

arXiv:2606.04032v2 Announce Type: replace-cross Abstract: Transformers have become the standard solution for various AI tasks, with the query, key, and value (QKV) attention formulation playing a central role. However, the individual contribution of these three projections and...

<https://arxiv.org/abs/2606.04032>

60 CONF

VISTA: Vision-Grounded and Physics-Validated Adaptation of UMI data for VLA Training

ArXiv CS.AI 50%

arXiv:2606.04708v2 Announce Type: replace-cross Abstract: Universal Manipulation Interface (UMI) enables scalable real-world robot data collection without hardware-specific teleoperation, yet leveraging UMI data to train large-scale Vision-Language-Action (VLA) models remains...

<https://arxiv.org/abs/2606.04708>