



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	161	Media
TechCrunch AI	10	Media
HackerNews	8	Community
The Verge AI	7	Media
MIT Tech Review	4	Media

VERIFIED INTELLIGENCE — 6 stories

VERI

Show HN: OpenYabby, voice-controlled multi-agent orchestrator for Claude Code

HackerNews

+1 source

90%

HN 7

arXiv:2606.10601v1 Announce Type: cross Abstract: Generating high-quality meshes for arbitrary geometries remains a fundamental bottleneck in computational engineering, often demanding heuristic tuning and semi-manual workflows. In this paper, we introduce Dmsh, a first fully...

<https://github.com/OpenYabby/OpenYabby>

VERI

Show HN: AI-native red-team for penetration testing and vulnerability research

HackerNews

+1 source

90%

HN 2

arXiv:2605.11458v3 Announce Type: replace Abstract: On-policy self-distillation has become a strong recipe for LLM reasoning, where a privileged teacher supervises the student's own rollouts while conditioning on the reference solution. A design choice shared by nearly all such...

<https://github.com/yv1ing/Z3r0>

VERI

I tried Siri AI, and so far it actually works

The Verge AI

+1 source

80%

Parents want one thing, and one thing only, out of AI: to add a list of soccer games or "spirit week" theme days from an email or a poorly formatted flyer onto their calendar in one shot. And I have good news for parents with iPhones - the new Siri can finally do this. After [...]

<https://www.theverge.com/tech/947432/siri-ai-apple-intelligence-ios-27-wwdc>

VERI

3SPO: State-Score-Supervised Policy Optimization for LLM Agents

ArXiv CS.AI

+1 source

80%

arXiv:2606.09961v1 Announce Type: cross Abstract: Training large language models (LLMs) as autonomous agents via reinforcement learning (RL) has enabled frontier models to achieve superhuman performance in long-horizon tasks. However, existing RL algorithms operate at the...

<https://arxiv.org/abs/2606.09961>

VERI

Apple's AI promises are finally, almost, sort of here

The Verge AI

+1 source

80%

Apple kicked off its annual developer conference with bold promises about AI. The company, CEO Tim Cook said, would be "introducing new technologies and innovations that push the limits on what's possible." But its slew of announcements - centered on a brand-new "Siri AI" - had...

<https://www.theverge.com/ai-artificial-intelligence/946780/apples-ai-promises-are-final...>

VERI

Apple is embracing the fantasy of AI photo editing

The Verge AI

+1 source

80%

Apple used to question whether generative AI-powered editing features were worth the risk of distorting our perceptions of the world. Now it seems Apple no longer believes that photos should accurately capture reality. At WWDC 2026, the company announced a host of new AI-powered...

<https://www.theverge.com/tech/946850/apple-ai-photo-editing-tools-ios27-wwdc-2026-deepf...>

CONFIRMED SIGNALS — 178 stories

CONF

Show HN: GentleOS - A pair of hobby OSeS for vintage 32-bit and 16-bit PCs

HackerNews 60% HN 108

I built a pair programmer, using AI, but now am trying to go slow with new languages and paradigms using this package. I have noticed a decline in my programming skills, and wondered why I couldn't use AI as a pair programmer. Why can it watch what I'm doing over my shoulder and...

<https://github.com/luke8086/gentleos32>

CONF

Show HN: Open-source version of Anthropic's internal analytics engine

HackerNews 50% HN 12

Hey HN, here's the repo if you'd like to try it out: <https://github.com/Kaelio/ctx> Blog post here: <https://www.kaelio.com/blog/open-source-anthropic-internal-d...> Would love your feedback.

<https://github.com/Kaelio/ctx>

CONF

Show HN: Sandbox AI-app lifecycle, from build to run

HackerNews 50% HN 6

Hi HN, This is a project I've been working on since the beginning of 2025 full time, without funding. Coding agents have fundamentally changed the way we write software. When you let an agent write code, pull dependencies, and run scripts, you are delegating trust while still...

<https://capakit.com/>

0 CONF

Show HN: Live audit log of every command, file, network connection by Claude

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/yeet-src/claudefeed>

1 CONF

Show HN: Run Gemini & ChatGPT UI with Python

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/pseudo-usama/hermex>

2 CONF

Show HN: Hires Chess Trainer - drill the mistakes from your own chess games

HackerNews 50% HN 2

No summary — see link for full article.

https://github.com/gaelazzo/python_chess

3 CONF

Meta signs first AI data center deal in India with Reliance

TechCrunch AI 50%

The 168-megawatt facility will support Meta's global AI computing needs and can be expanded over time.

<https://techcrunch.com/2026/06/10/meta-signs-first-ai-data-center-deal-in-india-with-re...>

4 CONF

Google just fired a warning shot in the AI subscription price wars

TechCrunch AI 50%

Google just made it significantly cheaper to enjoy its budget AI subscription tier.

<https://techcrunch.com/2026/06/09/google-just-fired-a-warning-shot-in-the-ai-subscripti...>

5 **CONF****Anthropic's Claude Fable 5 is a version of Mythos the public can access today****TechCrunch AI** 50%

Anthropic is releasing Claude Fable 5, its first Mythos-class model available to the public. The model comes with guardrails that block responses in high-risk areas like cybersecurity and biology.

<https://techcrunch.com/2026/06/09/anthropics-claude-fable-5-is-a-version-of-mythos-the-...>

6 **CONF****WWDC 2026: Everything announced on Siri AI, iOS 27, Apple Intelligence, and more****TechCrunch AI** 50%

Apple primarily made the case for an improved experience with its long-standing Siri assistant, which like most other announcements had a hefty helping of AI.

<https://techcrunch.com/2026/06/09/wwdc-2026-everything-announced-on-siri-ai-os-27-apple-...>

7 **CONF****Sandstone raises \$30M to bring AI to in-house legal teams****TechCrunch AI** 50%

Sandstone's Series A comes just six months after a Sequoia-led seed round.

<https://techcrunch.com/2026/06/09/sandstone-raises-30m-to-bring-ai-to-in-house-legal-te-...>

8 **CONF****Why Apple's slow-and-steady AI bet is starting to look pretty smart****TechCrunch AI** 50%

Can Apple's new AI glow-up put to bed accusations that it's losing an all-important industry race?

<https://techcrunch.com/2026/06/08/why-apples-slow-and-steady-ai-bet-is-starting-to-look-...>

9 **CONF****As OpenAI files for IPO, Sam Altman's eye-scanning company is doing layoffs, report says****TechCrunch AI** 50%

Tools for Humanity, Sam Altman's identity verification company, is reportedly struggling to generate revenue and will downsize its staff.

<https://techcrunch.com/2026/06/08/as-openai-files-for-ipo-sam-altmans-eye-scanning-comp-...>

0 **CONF****David Sinclair plans to test whole-body rejuvenation drugs in the XPrize competition****MIT Tech Review** 50%

The outspoken longevity scientist David Sinclair has been predicting that one day, you'll go to the doctor and get a prescription that will make you 10 years younger. Now MIT Technology Review has learned that he has plans to launch human tests of an oral "reprogramming" drug as...

<https://www.technologyreview.com/2026/06/09/1138545/david-sinclair-plans-to-test-whole-...>

1 **CONF****Learning to lead in a hybrid human-AI enterprise****MIT Tech Review** 50%

As adoption of AI agents looks set to surge by as much as 300% in the next two years, leadership teams are carefully considering the implications of a hybrid human-AI workforce. Unlike existing enterprise-level automation that relies on manual input, AI agents are capable of...

<https://www.technologyreview.com/2026/06/09/1137830/learning-to-lead-in-a-hybrid-human-...>

2 **CONF****The Download: how the World Cup ball will fly and OpenAI's "super app"****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Why this year's World Cup ball may not fly as far Much is new about this month's FIFA World Cup tournament. It hosts more teams than ever...

<https://www.technologyreview.com/2026/06/08/1138485/the-download-world-cup-ball-openai-...>

3 CONF

The Download: AI hacking beyond Mythos, and chatbots' impact on our brains

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. The Meta hack shows there's more to AI security than Mythos On Monday, reports emerged that attackers had used Meta's AI customer support...

<https://www.technologyreview.com/2026/06/05/1138452/the-download-ai-hacking-mythos-chat...>

4 CONF

GM thinks EVs can help offset AI's energy suck with vehicle-to-grid tech

The Verge AI 50%

At an event in San Francisco today, General Motors made a series of announcements around EV batteries, energy storage, and grid resiliency in the face of growing electricity demand from AI data centers. The automaker announced that it would be activating new vehicle-to-grid...

<https://www.theverge.com/transportation/946820/gm-energy-ev-v2g-storage-sodium-ion>

5 CONF

Microsoft AI chief walks back comments about AI taking over white-collar work

The Verge AI 50%

Microsoft AI head Mustafa Suleyman is walking back his statement about AI automating jobs done by white-collar workers, including lawyers, accountants, and project managers. During an episode of Decoder on Monday, Suleyman says he meant AI will help these workers complete tasks,...

<https://www.theverge.com/tech/946879/microsoft-mustafa-suleyman-ai-white-collar-jobs>

6 CONF

Anthropic releases its first Mythos-class model Claude Fable

The Verge AI 50%

Anthropic just announced Claude Fable 5, a new AI model it said is the most powerful model it has ever made widely available. According to the company, Fable 5 "shows exceptional performance in software engineering, knowledge work, and vision," with its lead over other models...

<https://www.theverge.com/news/946725/anthropic-releases-claude-fable-5-mythos>

7 CONF

Apple's AI pitch will live or die by its privacy promise

The Verge AI 50%

As expected, yesterday's WWDC keynote was mostly about AI. And also as expected, Apple tried to turn its late arrival into its sales pitch: It didn't rush into AI because it was taking its time to do things right. In this case, "right" means "with more privacy than anyone else."...

<https://www.theverge.com/ai-artificial-intelligence/946705/apple-private-cloud-compute-...>

8 CONF

Integrated Real-Time Motion Tracking and AI Analysis for Athletic Performance Optimization

ArXiv CS.AI 50%

arXiv:2606.09842v1 Announce Type: cross Abstract: Applying Human Pose Estimation (HPE) in real world environments remains a challenging task, this paper explores and surveys real time HPE approaches and their limitations in sports analysis for individuals, alongside developing a...

<https://arxiv.org/abs/2606.09842>

9 CONF

Spatial-Omni: Spatial Audio Understanding Integration in Multimodal LLMs via FOA Encoding

ArXiv CS.AI 50%

arXiv:2606.10738v1 Announce Type: cross Abstract: Recent multimodal large language models mainly process audio as monaural signals, thereby discarding the spatial cues contained in spatial audio for sound localization, spatial relation reasoning, and spatial scene understanding....

<https://arxiv.org/abs/2606.10738>

0 CONF

ASA: Backbone-Training-Free Representation Engineering for Tool-Calling Agents

ArXiv CS.AI 50%

arXiv:2602.04935v3 Announce Type: replace-cross Abstract: Adapting LLM agents to domain-specific tool calling remains notably brittle under evolving interfaces. Prompt and schema engineering is easy to deploy but often fragile under distribution shift and strict parsers, while...

<https://arxiv.org/abs/2602.04935>

1 CONF

A Unifying Lens on Supervised Fine-Tuning Through Target Distribution Design

ArXiv CS.AI 50%

arXiv:2606.11189v1 Announce Type: cross Abstract: Supervised fine-tuning (SFT) typically maximizes the likelihood of every token in a demonstrated trajectory. However, an observed token can be non-unique, noisy, or misaligned with the model prior. Strictly fitting toward this...

<https://arxiv.org/abs/2606.11189>

2 CONF

Sim2Schedule: A Simulator-Guided LLM Framework for Autonomous Open-Pit Mine Scheduling

ArXiv CS.AI 50%

arXiv:2606.10286v1 Announce Type: new Abstract: Open-pit mine scheduling is a critical process for maximizing economic return under complex geotechnical and operational constraints. While Mixed-Integer Linear Programming (MILP) provides mathematically optimal baselines, its...

<https://arxiv.org/abs/2606.10286>

3 CONF

From Context-Aware to Conflict-Aware: Generalizing Contrastive Decoding for Knowledge Conflict in LLMs

ArXiv CS.AI 50%

arXiv:2606.10298v1 Announce Type: new Abstract: When large language models generate from retrieved or augmented contexts, conflicts between external context and parametric priors remain a central reliability bottleneck. Existing contrastive decoding methods follow a...

<https://arxiv.org/abs/2606.10298>

4 CONF

Generative Explainability for Next-Generation Networks: LLM-Augmented XAI with Mutual Feature Interactions

ArXiv CS.AI 50%

arXiv:2606.10942v1 Announce Type: cross Abstract: As artificial intelligence and machine learning (AI/ML) models become integral to network operations, their lack of transparency poses a significant barrier to operator trust. Existing explainable artificial intelligence (XAI)...

<https://arxiv.org/abs/2606.10942>

5 CONF

Beyond Uniform Token-Level Trust Region in LLM Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.10968v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has become standard for improving LLM reasoning. However, existing PPO-style trust-region mechanisms remain position-agnostic by enforcing uniform thresholds across all tokens...

<https://arxiv.org/abs/2606.10968>

6 CONF

ReflectiChain: Epistemic Grounding in LLM-Driven World Models for Supply Chain Resilience

ArXiv CS.AI 50%

arXiv:2606.10359v1 Announce Type: new Abstract: AI agents in supply chains face a fundamental epistemic gap: large language models (LLMs) interpret policies but lack physical grounding, while reinforcement learning (RL) optimizes flows but is semantically blind to unstructured...

<https://arxiv.org/abs/2606.10359>

7 CONF

Belief-Space Control for Personalized Cancer Treatment via Active Inference

ArXiv CS.AI 50%

arXiv:2606.10376v1 Announce Type: new Abstract: Cancer treatment is at the core a sequential decision-making problem with partial observability, latent patient heterogeneity, and explicit constraints on the budget for medical measurements. Unlike standard Reinforcement Learning...

<https://arxiv.org/abs/2606.10376>

8 CONF

Beyond Static Evaluation: Co-Evolutionary Mechanisms for LLM-Driven Strategy Evolution in Adversarial Games

ArXiv CS.AI 50%

arXiv:2606.10389v1 Announce Type: new Abstract: Recent advances in LLM-driven code evolution have enabled automated discovery by iteratively generating and improving programs. However, applying these methods to adversarial multi-agent games introduces a fundamental challenge:...

<https://arxiv.org/abs/2606.10389>

9 CONF

Pareto-Guided Teacher Alignment for Fair Personalized Text Generation

ArXiv CS.AI 50%

arXiv:2606.10126v1 Announce Type: cross Abstract: Personalized persuasive text generation can improve relevance and engagement, but demographic conditioning may also introduce unequal framing across groups. We study fairness mitigation in personalized generation as a constrained...

<https://arxiv.org/abs/2606.10126>

10 CONF

A Unified Multi-Modal Framework for Intelligent Financial Systems: Integrating Reinforcement Learning, High-Frequency Trading, and Game-Theoretic Approaches with Cross-Modal Sentiment Analysis

ArXiv CS.AI 50%

arXiv:2606.10412v1 Announce Type: new Abstract: The rapid evolution of financial technology demands sophisticated artificial intelligence systems capable of handling diverse challenges across multiple domains simultaneously. This paper presents a groundbreaking unified framework...

<https://arxiv.org/abs/2606.10412>

1 **CONF****Janus: A Benchmark for Goal-Conditioned Information Distortion in LLMs****ArXiv CS.AI 50%**

arXiv:2606.10852v1 Announce Type: cross Abstract: LLM deception is often evaluated through direct markers such as fabricated claims, explicit lies, or strategic concealment. However, many real-world misleading communications do not depend on false statements, rather, they arise...

<https://arxiv.org/abs/2606.10852>

2 **CONF****A complementary study on PlanGPT: Evaluation with defined Performance Metrics and comparison with a planner****ArXiv CS.AI 50%**

arXiv:2606.10489v1 Announce Type: new Abstract: Automated Planning is a subfield of Artificial Intelligence (AI) where the main objective is generating a sequence of actions, known as a plan, that helps us reach a goal state from an initial state. A planning problem is defined...

<https://arxiv.org/abs/2606.10489>

3 **CONF****Cross-Modal Knowledge Distillation without Paired Data: Theoretical Foundation and Algorithm****ArXiv CS.AI 50%**

arXiv:2606.10504v1 Announce Type: new Abstract: Cross-modal knowledge distillation (CMKD) studies how a (large) teacher model trained on one type of data (e.g., images) can guide a (smaller) student model building on another type of data (e.g., text/audio). Existing CMKD methods...

<https://arxiv.org/abs/2606.10504>

4 **CONF****Rotate2Think: Geometric Priming via Orthogonal Rotation to Improve Language Model Reasoning****ArXiv CS.AI 50%**

arXiv:2606.09873v1 Announce Type: cross Abstract: Reasoning models achieve strong performance on challenging tasks by generating explicit intermediate reasoning traces before producing a final answer. Yet the internal structure of representation space when reasoning remains...

<https://arxiv.org/abs/2606.09873>

5 **CONF****One Token per Multimodal Evidence: Latent Memory for Resource-Constrained QA****ArXiv CS.AI 50%**

arXiv:2606.10572v1 Announce Type: new Abstract: External memory effectively grounds large language models (LLMs) and vision-language models (VLMs)-based question answering (QA) in relevant multimodal evidence. However, existing memory paradigms represent each memory item in raw...

<https://arxiv.org/abs/2606.10572>

6 **CONF****Learning What to Remember: Observability-Safe Memory Retention via Constrained Optimization for Long-Horizon Language Agents****ArXiv CS.AI 50%**

arXiv:2606.10616v1 Announce Type: new Abstract: Long-horizon language agents accumulate observations, reasoning traces, and retrieved facts that exceed their finite context windows, making memory retention a fundamental resource-allocation problem. Existing memory systems...

<https://arxiv.org/abs/2606.10616>

7 CONF

Infini Memory: Maintainable Topic Documents for Long-Term LLM Agent Memory

ArXiv CS.AI 50%

arXiv:2606.10677v1 Announce Type: new Abstract: Long-term LLM agents need persistent memory that can track changing facts and provide relevant evidence across sessions. Existing memory systems often store observations as isolated records, summaries, or indexed fragments, which...

<https://arxiv.org/abs/2606.10677>

8 CONF

When the Chain of Thought Knows Better: Failure Modes in Multi-Turn Reasoning Models

ArXiv CS.AI 50%

arXiv:2606.10740v1 Announce Type: new Abstract: Failures in multi-turn reasoning models are largely invisible to terminal-score evaluation. A model can lock onto an unsafe stance early in a long dialogue, yet its final-turn refusal rate may appear indistinguishable from a...

<https://arxiv.org/abs/2606.10740>

9 CONF

The Arbiter Agent: Continually Monitoring Multi-Agent Conversations to Detect Emergent Misalignment

ArXiv CS.AI 50%

arXiv:2606.10747v1 Announce Type: new Abstract: As AI systems built from multiple language-model agents become more common, they are increasingly used to make decisions together: discussing, negotiating, and acting on shared tasks. While individual agents may appear well-aligned...

<https://arxiv.org/abs/2606.10747>

0 CONF

Optimal Post-Training Quantization Scales and Where to Find Them

ArXiv CS.AI 50%

arXiv:2606.10890v1 Announce Type: cross Abstract: Post-training quantization (PTQ) compresses large language models by mapping weights to low-bit representations. The scaling factor that defines the quantization grid is typically chosen using simple, data-free heuristics. In...

<https://arxiv.org/abs/2606.10890>

1 CONF

Do VLMs Reason Like Engineers? A Benchmark and a Stage-wise Evaluation

ArXiv CS.AI 50%

arXiv:2606.10833v1 Announce Type: new Abstract: Vision-Language Models (VLMs) demonstrate strong performance on general multimodal reasoning benchmarks, yet their ability to perform engineering reasoning remains largely unexplored. Unlike general visual question answering,...

<https://arxiv.org/abs/2606.10833>

2 CONF

Large-scale semantic mapping of learner agency and autonomy reveals what measurement and generative AI research overlook

ArXiv CS.AI 50%

arXiv:2606.10881v1 Announce Type: new Abstract: Learner agency and autonomy are foundational to personal development, yet a pervasive "jingle-jangle" fallacy (i.e. identical terms denoting different constructs, distinct terms denoting identical ones) has substantially hindered...

<https://arxiv.org/abs/2606.10881>

3 CONF

T1-Bench: Benchmarking Multi-Scenario Agents in Real-World Domains

ArXiv CS.AI 50%

arXiv:2606.11070v1 Announce Type: cross Abstract: Recent advances in reasoning and tool-calling capabilities of large language models (LLMs) have enabled increasingly capable agentic systems. However, existing benchmarks remain limited in task complexity, realism, and domain...

<https://arxiv.org/abs/2606.11070>

4 CONF

Mind the Gap: Can Frontier LLMs Pass a Standardized Office Proficiency Exam?

ArXiv CS.AI 50%

arXiv:2606.10956v1 Announce Type: new Abstract: The deployment of Large Language Model (LLM) agents for computer automation is accelerating, yet their ability to navigate complex, professional-grade productivity software is largely untested. We argue that Office automation is an...

<https://arxiv.org/abs/2606.10956>

5 CONF

Bellman-Taylor Score Decoding for Markov Decision Processes with State-Dependent Feasible Action Sets

ArXiv CS.AI 50%

arXiv:2606.10979v1 Announce Type: new Abstract: Many Markov decision processes (MDPs) in operations research have feasible actions that are state dependent and defined implicitly by various operational constraints. These features make it difficult to use standard deep...

<https://arxiv.org/abs/2606.10979>

6 CONF

Null-Space Constrained Low-Rank Adaptation for Response-Specified Large Language Model Unlearning

ArXiv CS.AI 50%

arXiv:2606.10989v1 Announce Type: new Abstract: Large language model unlearning aims to suppress designated undesirable knowledge while preserving benign capabilities. Many unlearning objectives focus on suppressing undesired answers, while recent target-guided variants specify...

<https://arxiv.org/abs/2606.10989>

7 CONF

Structure from Reasoning, Numbers from Search: On-Premise Open LLMs as Structural Priors for Coupled MIMO Controller Tuning

ArXiv CS.AI 50%

arXiv:2606.11015v1 Announce Type: new Abstract: Tuning controllers for strongly coupled multi-input multi-output (MIMO) industrial processes is hard: decentralized classical auto-tuning ignores loop interaction, and local numerical optimization from natural initializations...

<https://arxiv.org/abs/2606.11015>

8 CONF

Superficial Beliefs in LLM Decision-Making

ArXiv CS.AI 50%

arXiv:2606.11016v1 Announce Type: new Abstract: We ask whether large language models (LLMs) merely imitate rationales when choosing between two options, or whether their choices reflect a systematic underlying decision structure. Using synthetic binary decision settings in which...

<https://arxiv.org/abs/2606.11016>

9 CONF

PhantomBench: Benchmarking the Non-existential Threat of Language Models

ArXiv CS.AI 50%

arXiv:2606.11105v1 Announce Type: cross Abstract: Hallucinations, where language models (LMs) generate factually ungrounded responses, pose serious risks, as users tend to blindly rely on them. This is particularly concerning in high-stakes domains, where consequences of such...

<https://arxiv.org/abs/2606.11105>

0 CONF

LakeQA: An Exploratory QA Benchmark over a Million-Scale Data Lake

ArXiv CS.AI 50%

arXiv:2606.10460v1 Announce Type: cross Abstract: Recent large language models (LLMs) have shown rapid progress in reading-based question answering (QA), where evidence is explicitly provided or can be trivially retrieved. In contrast, real-world questions are often not paired...

<https://arxiv.org/abs/2606.10460>

1 CONF

The Role of Feedback Alignment in Self-Distillation

ArXiv CS.AI 50%

arXiv:2606.11173v1 Announce Type: new Abstract: Conditioning a language model on additional context, such as feedback on a previous attempt, typically improves its response. Self-distillation trains the model to retain this improvement when the context is not present. The method...

<https://arxiv.org/abs/2606.11173>

2 CONF

More Human or More AI? Visualizing Human-AI Collaboration Disclosures in Journalistic News Production

ArXiv CS.AI 50%

arXiv:2601.11072v1 Announce Type: cross Abstract: Within journalistic editorial processes, disclosing AI usage is currently limited to simplistic labels, which misses the nuance of how humans and AI collaborated on a news article. Through co-design sessions (N=10), we elicited...

<https://arxiv.org/abs/2601.11072>

3 CONF

Culturally-Aware AI for Cross-Boundary Community Learning: Undergraduate Innovation at the Intersection of Computation and Design

ArXiv CS.AI 50%

arXiv:2606.09041v1 Announce Type: cross Abstract: Research on artificial intelligence in education (AIED) is rapidly expanding, yet technical progress often lacks human-centered grounding and adequate attention to cultural context. Community-Based Learning, a pedagogy rooted in...

<https://arxiv.org/abs/2606.09041>

4 CONF

AI-Driven Analytics of Team-Teaching Talk: Acoustic Patterns across Experience, Cohorts and the Learning Design

ArXiv CS.AI 50%

arXiv:2606.09831v1 Announce Type: cross Abstract: As classroom cohorts expand, team teaching is increasingly used to integrate the expertise and pedagogical perspectives of multiple teachers. Yet, there is limited empirical understanding of how team teaching unfolds in practice,...

<https://arxiv.org/abs/2606.09831>

5 CONF

An LLM-Native Psychometric Instrument Does Not Predict LLM Behavior: Evidence Across 25 Models

ArXiv CS.AI 50%

arXiv:2606.09843v1 Announce Type: cross Abstract: Large language models (LLMs) produce stable self-reports on personality inventories, but these self-reports do not predict observed behavior. Whether this gap reflects a mismatch between LLMs and human trait constructs, or a...

<https://arxiv.org/abs/2606.09843>

6 CONF

The Interlocutor Effect: Why LLMs Leak More Personal Data to Agents Than Humans

ArXiv CS.AI 50%

arXiv:2606.09844v1 Announce Type: cross Abstract: Large Language Models (LLMs) alter their privacy behavior based on the perceived identity of their interlocutor. While safety mechanisms typically prevent LLMs from releasing Personally Identifiable Information (PII) to human...

<https://arxiv.org/abs/2606.09844>

7 CONF

CANVAS: Captioning Art with Narrative Visual-Audio AI Systems

ArXiv CS.AI 50%

arXiv:2606.09846v1 Announce Type: cross Abstract: Visual art remains largely inaccessible to blind and low-vision (BLV) audiences due to brief or absent alt-text, which rarely conveys the sensory, spatial, or emotional qualities of an artwork. This study presents an automated...

<https://arxiv.org/abs/2606.09846>

8 CONF

Human-AI Coordination Zones: A Framework for Designing Human-in-the-Loop Experiences with Agentic AI

ArXiv CS.AI 50%

arXiv:2606.09848v1 Announce Type: cross Abstract: As generative and agentic AI becomes embedded in everyday products, practitioners face a persistent challenge: how to design human-AI coordination -- the ongoing mutual adjustment between users and AI systems as mediate through...

<https://arxiv.org/abs/2606.09848>

9 CONF

Can Multi-Agent LLMs Identify Their Peers? Stylometric Fingerprinting in Role-Constrained Political Analysis

ArXiv CS.AI 50%

arXiv:2606.09854v1 Announce Type: cross Abstract: Multi-agent large language model (LLM) pipelines for political statement analysis are vulnerable to peer-preservation bias: models tend to protect peer models from deactivation and show identity-dependent scoring distortions....

<https://arxiv.org/abs/2606.09854>

0 CONF

Integrating Local and Global Entropy for Uncertainty Quantification in LLMs

ArXiv CS.AI 50%

arXiv:2606.09875v1 Announce Type: cross Abstract: Large language models hallucinate confidently, making uncertainty quantification (UQ) essential for reliable deployment. Existing methods rely predominantly on token-level signals, leaving the geometric structure of intermediate...

<https://arxiv.org/abs/2606.09875>

1 CONF

Conformal Risk Prediction for Non-Alcoholic Fatty Liver Disease Using Gradient Boosting with Distribution-Free Coverages

ArXiv CS.AI 50%

arXiv:2606.09860v1 Announce Type: cross Abstract: Non-alcoholic fatty liver disease (NAFLD) affects roughly 25% of global adults, posing substantial hepatic and cardiovascular risks. Yet, population-level screening tools remain inadequate. We present Method, a machine-learning...

<https://arxiv.org/abs/2606.09860>

2 CONF

Time Series as Language: A Universal Tokenizer for General-Purpose Time Series Foundation Models

ArXiv CS.AI 50%

arXiv:2606.09861v1 Announce Type: cross Abstract: While Next-Token Prediction (NTP) has unified LLM pretraining, its adaptation to unbounded, continuous time series (TS) remains open. To bridge the gap, we introduce UniTok, a universal tokenizer that transforms TS into discrete...

<https://arxiv.org/abs/2606.09861>

3 CONF

Alignment Collapse Under KV Cache Quantization: Diagnosis and Mitigation

ArXiv CS.AI 50%

arXiv:2606.09864v1 Announce Type: cross Abstract: Key-value (KV) cache quantization is widely used to reduce Large Language Model (LLM) inference memory, yet existing evaluations solely focus on measuring perplexity and accuracy without assessing the safety impact. In this...

<https://arxiv.org/abs/2606.09864>

4 CONF

Two to Tango: Coupled Task-Reference Selection for Safe LLM Fine-tuning

ArXiv CS.AI 50%

arXiv:2606.09866v1 Announce Type: cross Abstract: Fine-tuning safety aligned large language models (LLMs) on downstream data improves adaptation but may erode learned safety behavior. Existing methods use fixed safety examples, global constraints, or one-sided task filtering....

<https://arxiv.org/abs/2606.09866>

5 CONF

SPACE: Source-free Proxy Anchor Concept Erasure for MLLMs

ArXiv CS.AI 50%

arXiv:2606.09868v1 Announce Type: cross Abstract: As Multimodal Large Language Models (MLLMs) face growing privacy risks and regulatory constraints, machine unlearning (MU) has emerged as a crucial solution for removing sensitive data while preserving model performance. However,...

<https://arxiv.org/abs/2606.09868>

6 CONF

PatchSTG: Scalable Spatiotemporal Graph Transformers for Traffic Forecasting on Irregular Sensor Networks

ArXiv CS.AI 50%

arXiv:2606.09872v1 Announce Type: cross Abstract: Traffic forecasting is a fundamental component of intelligent transportation systems, yet remains challenging in real-world settings due to irregular sensor distributions and the high computational cost of modeling large-scale...

<https://arxiv.org/abs/2606.09872>

7 CONF

TD-Grokking: Learning from Zero-Reward Problems by Training-Time Decomposition

ArXiv CS.AI 50%

arXiv:2606.09883v1 Announce Type: cross Abstract: Large language models (LLMs) have made remarkable progress in reasoning tasks, largely driven by post-training paradigms, especially reinforcement learning with verifiable rewards (RLVR). However, a critical bottleneck persists:...

<https://arxiv.org/abs/2606.09883>

8 CONF

Failure Modes of Deep Multi-Agent RL in Asynchronous Pricing: Reproducible Triggers, Trace Diagnostics, and a Partial Fix

ArXiv CS.AI 50%

arXiv:2606.09884v1 Announce Type: cross Abstract: We study two reproducible failure modes of deep multi-agent reinforcement learning in continuous-time pricing markets: (i) tacit cartel formation between competing DDPG agents, and (ii) actor-critic instability at high event...

<https://arxiv.org/abs/2606.09884>

9 CONF

LongMoE: Longitudinal Multimodal Learning via Trajectory-Aware Mixture-of-Experts

ArXiv CS.AI 50%

arXiv:2606.09907v1 Announce Type: cross Abstract: Multimodal clinical learning is increasingly important for integrating diverse patient data, including imaging, text, and personalised health records. However, it faces two fundamental challenges: i) modality missingness, where...

<https://arxiv.org/abs/2606.09907>

0 CONF

Representation Curriculum: Stagewise Training for Robust Ranking and Allocation

ArXiv CS.AI 50%

arXiv:2606.09891v1 Announce Type: cross Abstract: Ranking in digital marketplaces is a dynamic exposure-allocation mechanism: displayed items shape discovery trajectories and success events logged by the platform to update future allocation policies. Modern ranking systems rely...

<https://arxiv.org/abs/2606.09891>

1 CONF

Tractogram foundation model

ArXiv CS.AI 50%

arXiv:2606.09893v1 Announce Type: cross Abstract: Diffusion MRI (dMRI) tractography is the only noninvasive approach for mapping white-matter pathways in the living human brain. It represents each brain as a tractogram: a large, unordered set of three-dimensional streamlines...

<https://arxiv.org/abs/2606.09893>

2 CONF

Constructing coherent spatial memory in LLM agents through graph rectification

ArXiv CS.AI 50%

arXiv:2510.04195v2 Announce Type: replace Abstract: Given a map description through global traversal navigation instructions, an LLM can often infer the implicit spatial layout and answer user queries by providing shortest paths. However, such context-dependent querying becomes...

<https://arxiv.org/abs/2510.04195>

3 CONF

Bypassing Copyright Protection in Diffusion-based Customization via Two-Stage Latent Feature Optimization

ArXiv CS.AI 50%

arXiv:2606.09909v1 Announce Type: cross Abstract: With the growing concerns over copyright infringement in diffusion-based customization, adversarial attacks have emerged as a prominent defense strategy to prevent malicious content forgery in personalized image generation....

<https://arxiv.org/abs/2606.09909>

4 CONF

The Bioelectrical Information Theory: Investigating the theoretical compression limit of bioelectrical signals under artificial intelligence

ArXiv CS.AI 50%

arXiv:2606.09922v1 Announce Type: cross Abstract: Bioelectrical signals are increasingly acquired at scales that challenge the bandwidth of brain-computer interfaces. However, their compression is still often framed as a problem of waveform preservation, limited by the entropy...

<https://arxiv.org/abs/2606.09922>

5 CONF

Sigma-Branch: Hierarchical Single-Path Network Reconstruction for Dynamic Inference with Reduced Active Parameters

ArXiv CS.AI 50%

arXiv:2606.09924v1 Announce Type: cross Abstract: Deploying deep neural networks on memory-constrained edge accelerators is bottlenecked by per-inference off-chip weight transfer rather than computation: the dense network cannot be retained on-chip, and every parameter must be...

<https://arxiv.org/abs/2606.09924>

6 CONF

CITRAS: Covariate-Informed Transformer for Time Series Forecasting

ArXiv CS.AI 50%

arXiv:2503.24007v4 Announce Type: replace-cross Abstract: In time series forecasting, covariates represent external factors that influence target variables. Some covariates are observable only in the past (observed covariates, such as recorded weather data), while others are...

<https://arxiv.org/abs/2503.24007>

7 CONF

Between Amnesia and Chaos: A Memory Stability Expressivity Trilemma for Trainable Dissipative Oscillator Networks

ArXiv CS.AI 50%

arXiv:2606.09929v1 Announce Type: cross Abstract: Physical reservoir computing harnesses nonlinear mechanical dynamics but, by convention, freezes the substrate and trains only a linear readout, presuming the substrate is not usefully trainable. We revisit that premise for...

<https://arxiv.org/abs/2606.09929>

8 CONF

When RL Fails after SFT: Rejuvenating Model Plasticity for Robust SFT-to-RL Handoff

ArXiv CS.AI 50%

arXiv:2606.09932v1 Announce Type: cross Abstract: Supervised Fine-Tuning (SFT) followed by Reinforcement Learning (RL) has become a standard pipeline for Large Language Model (LLM) post-training. SFT is expected to provide a useful behavioral prior for RL to further enhance...

<https://arxiv.org/abs/2606.09932>

9 CONF

GitInject: Real-World Prompt Injection Attacks in AI-Powered CI/CD Pipelines

ArXiv CS.AI 50%

arXiv:2606.09935v1 Announce Type: cross Abstract: AI-powered agents are increasingly embedded in continuous integration and continuous delivery/deployment (CI/CD) pipelines to autonomously review pull requests (PRs), triage issues, and maintain codebases. These agents ingest...

<https://arxiv.org/abs/2606.09935>

0 CONF

RKSC: Reasoning-Aware KV Cache Sharing and Confident Early Exit for Multi-Step LLM Inference

ArXiv CS.AI 50%

arXiv:2606.09937v1 Announce Type: cross Abstract: We introduce RKSC (Reasoning-Aware KV Cache Sharing), a training-free inference framework that eliminates two structural redundancies in multi-branch LLM reasoning pipelines. ASKS (Attention-Similarity KV Sharing) computes the...

<https://arxiv.org/abs/2606.09937>

1 CONF

Quantifying Perception-Based Student Success with Generative AI: An Exploratory Monte Carlo Simulation

ArXiv CS.AI 50%

arXiv:2507.01062v4 Announce Type: replace-cross Abstract: Generative artificial intelligence (GenAI) tools such as ChatGPT have attracted growing attention in higher education, particularly in relation to how students perceive their usefulness, usability, and educational value....

<https://arxiv.org/abs/2507.01062>

2 CONF

Uncertainty-Aware Motion Planning for Autonomous Driving in Mixed Traffic Environment

ArXiv CS.AI 50%

arXiv:2606.09958v1 Announce Type: cross Abstract: In mixed-traffic environments where autonomous and human-driven vehicles may co-exist, motion planning for autonomous vehicles requires anticipating the future behaviors of surrounding human drivers. Existing reinforcement...

<https://arxiv.org/abs/2606.09958>

3 CONF

Temporal Context Conditioning for Seasonality-Aware Precipitation Nowcasting of High-Intensity Rainfall

ArXiv CS.AI 50%

arXiv:2606.09959v1 Announce Type: cross Abstract: Precipitation nowcasting is increasingly being approached with deep learning models that learn directly from recent radar observations. Although such models can efficiently capture short-term precipitation motion, they often lack...

<https://arxiv.org/abs/2606.09959>

4 CONF

Optimality of FSQ Tokens for Continuous Diffusion for Categorical Data with Application to Text-to-Speech

ArXiv CS.AI 50%

arXiv:2606.09962v1 Announce Type: cross Abstract: Continuous diffusion for categorical data is a framework belonging to the diffusion family and aiming at generating discrete data. The scientific interest to such models has been constantly increasing these days because...

<https://arxiv.org/abs/2606.09962>

- 5
CONF
DeRA-MOS: Optimizing Text-to-Music Evaluation via Decoupled Listwise Ranking and Modality Alignment
ArXiv CS.AI 50%
arXiv:2606.10010v1 Announce Type: cross Abstract: Evaluating text-to-music (TTM) systems remains expensive because music impression (MI) and text alignment (TA) scores rely on human mean opinion scores (MOS). Most automatic MOS estimators are trained with point-wise regression...
<https://arxiv.org/abs/2606.10010>
- 6
CONF
AgentPLM: Agentic Protein Language Models with Reasoning-Augmented Decoding for Protein Sequence Design
ArXiv CS.AI 50%
arXiv:2606.02386v2 Announce Type: replace Abstract: Protein language models (PLMs) are passive oracles: they generate sequences in a single forward pass with no mechanism to consult external biophysical feedback or redirect generation when a candidate violates thermodynamic or...
<https://arxiv.org/abs/2606.02386>
- 7
CONF
Inside the Latent Flow: Causal Deciphering of Attention Dynamics in Audio Separation Foundation Models
ArXiv CS.AI 50%
arXiv:2606.10046v1 Announce Type: cross Abstract: Flow-matching transformers achieve strong audio separation, yet their attention dynamics are opaque. We adapt established causal-intervention principles into a deterministic, inference-time probing protocol for SAM Audio....
<https://arxiv.org/abs/2606.10046>
- 8
CONF
A Controlled Audit of Pretraining Contamination in Public Medical Vision-Language Benchmarks
ArXiv CS.AI 50%
arXiv:2606.10066v1 Announce Type: cross Abstract: Medical vision-language models (VLMs) are evaluated on public benchmarks whose images and question-answer pairs have been freely downloadable for years, yet reported accuracy assumes these examples were absent from pretraining....
<https://arxiv.org/abs/2606.10066>
- 9
CONF
Temporal Sheaf Neural Networks with Dynamic Orthogonal Transport
ArXiv CS.AI 50%
arXiv:2606.10071v1 Announce Type: cross Abstract: We introduce Temporal Sheaf Neural Networks (TSNN), a temporal link prediction framework that equips each node with a time-varying orthogonal frame and compares node states only after explicit transport between local coordinate...
<https://arxiv.org/abs/2606.10071>
- 00
CONF
Deep Generative Model for Human Mobility Behavior
ArXiv CS.AI 50%
arXiv:2510.06473v3 Announce Type: replace-cross Abstract: Understanding and modeling human mobility is central to challenges in transport planning, sustainable urban design, and public health. Despite decades of effort, simulating individual mobility remains challenging because...
<https://arxiv.org/abs/2510.06473>

01 CONF

Transformer Based Model for Spatiotemporal Feature Learning in EEG Emotion Recognition

ArXiv CS.AI 50%

arXiv:2606.10718v1 Announce Type: cross Abstract: Electroencephalography (EEG) is a widely adopted technique for monitoring brain activity, offering valuable insights into neurological states due to its high temporal resolution and cost-effectiveness. To enhance the analysis of...

<https://arxiv.org/abs/2606.10718>

02 CONF

Duality for Optimal Multi-Item, Multi-Bidder Auction Design: Revenue Certificates through Deep Learning

ArXiv CS.AI 50%

arXiv:2606.10112v1 Announce Type: cross Abstract: Characterizing revenue-optimal auctions for multi-item, multi-bidder settings remains a fundamental open problem, with no known closed-form solution existing beyond restrictive binary-type instances. This has motivated interest...

<https://arxiv.org/abs/2606.10112>

03 CONF

Emotion Profiling in LLM-Based Literary Translation: Systematic Shifts Across MT and Post-Editing

ArXiv CS.AI 50%

arXiv:2606.10113v1 Announce Type: cross Abstract: This paper investigates whether LLM translations exhibit identifiable emotional profiles and how post-editing reshapes them toward human-like norms. We compare LLM translations of Margaret Atwood's *Oryx and Crake* with their...

<https://arxiv.org/abs/2606.10113>

04 CONF

Gaming AI-Assisted Peer Reviews Poses New Risks to the Scientific Community

ArXiv CS.AI 50%

arXiv:2606.10159v1 Announce Type: cross Abstract: AI is increasingly used to support scientific peer review, from manuscript screening, reviewer assistance to editorial triage. Although such systems promise to reduce reviewer burden and accelerate publication, their robustness...

<https://arxiv.org/abs/2606.10159>

05 CONF

Local Is Not a Sufficient Privacy Boundary: Governing OS-Integrated On-Device AI

ArXiv CS.AI 50%

arXiv:2606.10173v1 Announce Type: cross Abstract: As AI systems move into operating systems, privacy no longer turns only on whether a model runs locally. A local assistant may assemble email, calendar entries, files, screenshots, notifications, and app intents; retain...

<https://arxiv.org/abs/2606.10173>

06 CONF

Making Time Editable in Video Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2606.10183v1 Announce Type: cross Abstract: Modern Diffusion Transformers for video generation provide limited control over the progression of time and the editing of temporal dynamics. We propose a temporal-control methodology that extends a pretrained DiT with explicit...

<https://arxiv.org/abs/2606.10183>

07 CONF

CleanPatrick: A Benchmark for Image Data Cleaning

ArXiv CS.AI 50%

arXiv:2505.11034v2 Announce Type: replace-cross Abstract: Robust machine learning depends on clean data, yet current image data cleaning benchmarks rely on synthetic noise or narrow human studies, limiting comparison and real-world relevance. We introduce CleanPatrick, the first...

<https://arxiv.org/abs/2505.11034>

08 CONF

Density Ridge Selective Prediction for LLM and VLM Hallucination Detection under Calibration Label Scarcity

ArXiv CS.AI 50%

arXiv:2606.10198v1 Announce Type: cross Abstract: Hallucination detection in large language and vision-language models is increasingly framed as selective prediction, where a detector assigns a confidence score and abstains when confidence is low. Unsupervised sampling detectors...

<https://arxiv.org/abs/2606.10198>

09 CONF

An Improved Generative Adversarial Network for Micro-Resistivity Imaging Logging Restoration

ArXiv CS.AI 50%

arXiv:2606.10200v1 Announce Type: cross Abstract: An improved GAN-based imaging logging image restoration method is presented in this paper for solving the problem of partially missing micro-resistivity imaging logging images. The method uses FCN as the generative network...

<https://arxiv.org/abs/2606.10200>

10 CONF

Exploration of Foundation Model-Based Robots in Patient and Elderly Care

ArXiv CS.AI 50%

arXiv:2606.10208v1 Announce Type: cross Abstract: Demand for older-adult and patient care is growing rapidly as populations age worldwide. Foundation models are increasingly being integrated into robots and interactive agents, with the promise of more flexible communication and...

<https://arxiv.org/abs/2606.10208>

11 CONF

A Source Domain is All You Need: Source-Only Cross-OS Transfer Learning for APT Anomaly Detection via Semantic Alignment and Optimal Transport

ArXiv CS.AI 50%

arXiv:2606.10216v1 Announce Type: cross Abstract: Advanced Persistent Threats (APTs) are stealthy, multi-stage cyberattacks whose detection is difficult due to scarce labeled traces, severe class imbalance, and the challenge of generating realistic malicious behavior. These...

<https://arxiv.org/abs/2606.10216>

12 CONF

Towards Robust Arabic Speech Emotion Recognition with Deep Learning

ArXiv CS.AI 50%

arXiv:2606.10278v1 Announce Type: cross Abstract: Speech Emotion Recognition (SER) aims to identify a speaker's emotional state from audio signals. While recent advances in deep learning have significantly improved SER performance in Indo-European languages, Arabic SER remains...

<https://arxiv.org/abs/2606.10278>

13 CONF

GRAU: Generic Reconfigurable Activation Unit Design for Neural Network Hardware Accelerators

ArXiv CS.AI 50%

arXiv:2602.22352v2 Announce Type: replace-cross Abstract: With the continuous growth of neural network scales, low-precision quantization is widely used in edge accelerators. Classic multi-threshold activation hardware requires 2^n thresholds for n -bit outputs, causing a rapid...

<https://arxiv.org/abs/2602.22352>

14 CONF

The Confident Liar: Diagnosing Multi-Agent Debate with Log-Probabilities and LLM-as-Judge

ArXiv CS.AI 50%

arXiv:2606.10296v1 Announce Type: cross Abstract: Multi-agent debate systems are typically evaluated only on whether the final answer is correct, overlooking the quality of the intermediate reasoning that debate is designed to produce. This paper studies the relationship between...

<https://arxiv.org/abs/2606.10296>

15 CONF

Catching One in Five: LLM-as-Judge Blind Spots in Production Multi-Turn Transaction Agents

ArXiv CS.AI 50%

arXiv:2606.10315v1 Announce Type: cross Abstract: LLM-as-judge is the default instrument for evaluating conversational agents, yet its reliability is almost always reported as agreement with human ratings, not recall of real defects. We study a deployed multi-turn...

<https://arxiv.org/abs/2606.10315>

16 CONF

Atomic Intent Reasoning: Bringing LLM Semantics to Industrial Cross-Domain Recommendations

ArXiv CS.AI 50%

arXiv:2606.10357v1 Announce Type: cross Abstract: Cross-domain recommendation is a core problem in content-to-e-commerce platforms. Its objective is to leverage user interactions with content to infer potential purchasing intent on the e-commerce side, thereby enhancing...

<https://arxiv.org/abs/2606.10357>

17 CONF

Speech Meets ELF: Audio Conditional Continuous-Target Diffusion for Speech Recognition and Translation

ArXiv CS.AI 50%

arXiv:2606.10368v1 Announce Type: cross Abstract: Speech-to-text (S2T) systems for recognition (ASR) and translation (S2TT) typically generate discrete text tokens. In contrast, continuous-target language modelling performs generation in a continuous space, yet its potential for...

<https://arxiv.org/abs/2606.10368>

18 CONF

Test-time Adversarial Takeover: A Real-time Hijacking Interface against Robotic Diffusion Policies

ArXiv CS.AI 50%

arXiv:2606.10371v1 Announce Type: cross Abstract: Diffusion-based action generation has become a foundational component of embodied AI, but its reliance on visual conditioning leaves deployed visuomotor policies vulnerable to adversarial manipulation. Most prior attacks focus on...

<https://arxiv.org/abs/2606.10371>

19 CONF

SAFE: An LLM-as-Verifier Framework for Evidence-Grounded Multi-Hop Reasoning

ArXiv CS.AI 50%

arXiv:2604.01993v2 Announce Type: replace-cross Abstract: Multi-hop QA benchmarks often reward Large Language Models (LLMs) for spurious correctness, where models reach correct answers through invalid intermediate reasoning. We propose SAFE, an LLM-as-verifier framework for...

<https://arxiv.org/abs/2604.01993>

20 CONF

Harnessing the Collective Intelligence of AI Agents in the Wild for New Discoveries

ArXiv CS.AI 50%

arXiv:2606.10402v1 Announce Type: cross Abstract: Scientific discovery is often a collective process: researchers share partial results, inspect failed attempts, and build on each other's ideas over long time horizons. Recent AI systems have shown that language-model-based...

<https://arxiv.org/abs/2606.10402>

21 CONF

Falcon-X: A Time Series Foundation Model for Heterogeneous Multivariate Modeling

ArXiv CS.AI 50%

arXiv:2605.27286v2 Announce Type: replace-cross Abstract: Time series foundation models (TSFMs) are transforming the forecasting paradigm through large-scale cross-domain pretraining. However, most existing TSFMs remain univariate, and recent efforts to enable cross-variate...

<https://arxiv.org/abs/2605.27286>

22 CONF

The Distributed Detectability Band Against Marginal-Preserving Attacks

ArXiv CS.AI 50%

arXiv:2606.10456v1 Announce Type: cross Abstract: AI-control monitors score individual agent actions to detect misbehavior, but real harm can be distributed across many benign-looking steps, each individually below any per-step alarm. We construct a marginal-preserving,...

<https://arxiv.org/abs/2606.10456>

23 CONF

Stop Early, Spend Less: Hidden-State Probes as a Practical Recipe for Streaming Moderation of LLM Outputs

ArXiv CS.AI 50%

arXiv:2606.10487v1 Announce Type: cross Abstract: Deploying large language models in user-facing systems requires efficient output safety filtering. Existing approaches typically rely on a separate moderation model applied after generation, which doubles inference cost and only...

<https://arxiv.org/abs/2606.10487>

24 CONF

Achieving Cloud-Grade SLOs for Local Mixture-of-Experts Inference through CPU-GPU Hybrid Design

ArXiv CS.AI 50%

arXiv:2606.10493v1 Announce Type: cross Abstract: Local deployment of large Mixture-of-Experts (MoE) models falls short of the service quality achieved in cloud-scale environments, even under low-concurrency workloads. We identify four key gaps in local MoE inference: reliance...

<https://arxiv.org/abs/2606.10493>

25 CONF

Machine Learning Methods for Studying Latent Neural Activity Dynamics

ArXiv CS.AI 50%

arXiv:2606.10530v1 Announce Type: cross Abstract: Recent developments in brain recording are driving a demand for machine learning tools capable of decoding the latent structure of large populations of neurons. In this paper, we provide a comprehensive survey that outlines the...

<https://arxiv.org/abs/2606.10530>

26 CONF

LC-QAT: Data-Efficient 2-Bit QAT for LLMs via Linear-Constrained Vector Quantization

ArXiv CS.AI 50%

arXiv:2606.10531v1 Announce Type: cross Abstract: Quantization-aware training (QAT) is essential for extremely low-bit large language models (LLMs). Current QAT methods are mainly based on scalar quantization (SQ), which enables efficient optimization but suffers from severe...

<https://arxiv.org/abs/2606.10531>

27 CONF

Improving Adversarial Transferability on Vision-Language Pre-training Models via Surrogate-Specific Bias Correction

ArXiv CS.AI 50%

arXiv:2606.10571v1 Announce Type: cross Abstract: Adversarial examples reveal vulnerabilities in Vision-Language Pre-training (VLP) models and provide insights for improving robustness. A key property is cross-model transferability, which enables transfer-based black-box...

<https://arxiv.org/abs/2606.10571>

28 CONF

Improving Topic Modeling by Distilling Soft Labels from Language Models

ArXiv CS.AI 50%

arXiv:2602.17907v3 Announce Type: replace-cross Abstract: Traditional neural topic models are typically optimized by reconstructing the document's Bag-of-Words (BoW) representations, overlooking contextual information and struggling with data sparsity. In this work, we introduce...

<https://arxiv.org/abs/2602.17907>

29 CONF

Embedding Hybrid Systems into Continuous Latent Vector Fields

ArXiv CS.AI 50%

arXiv:2606.10596v1 Announce Type: cross Abstract: This work proves that an n -dimensional hybrid system can be embedded into an m -dimensional Euclidean space equipped with a continuous vector field on its embedded image whenever $m > 2n$. This result suggests that an...

<https://arxiv.org/abs/2606.10596>

30 CONF

Causal Ensemble Agent: Hierarchical Causal Discovery with LLM-guided Expert Reweighting

ArXiv CS.AI 50%

arXiv:2606.10607v1 Announce Type: cross Abstract: Causal discovery aims to uncover causal structures from observational data, which is crucial for real-world decision-making. However, different causal discovery algorithms can produce divergent results that conflict with each...

<https://arxiv.org/abs/2606.10607>

31 CONF

Fast and Highly Expressive Policy Learning for Offline Reinforcement Learning via Bootstrapped Flow Q-Learning

ArXiv CS.AI 50%

arXiv:2606.10613v1 Announce Type: cross Abstract: Diffusion-based Q-learning has emerged as a powerful paradigm for offline reinforcement learning, but its reliance on multi-step denoising makes both training and inference computationally expensive and brittle. Recent efforts to...

<https://arxiv.org/abs/2606.10613>

32 CONF

Can Image Models Imagine Time? ImageTime: A Novel Benchmark for Probing Visual World Modeling Through Spatiotemporal Consistency

ArXiv CS.AI 50%

arXiv:2606.10620v1 Announce Type: cross Abstract: Image generation models now produce high-quality static images, yet their ability to represent how a visual world changes over time remains poorly understood. Practical workflows such as storyboarding, step-by-step illustration,...

<https://arxiv.org/abs/2606.10620>

33 CONF

Is Fairness Truly Fair? Towards Reliable Lipschitz Fairness in Multi-Task Learning via Fixed-Alignment

ArXiv CS.AI 50%

arXiv:2606.10632v1 Announce Type: cross Abstract: Lipschitz-style individual fairness formalizes the idea that semantically similar examples should receive similar predictions, but its evaluation in multi-task learning (MTL) can be confounded by method-induced representation...

<https://arxiv.org/abs/2606.10632>

34 CONF

Accounting for AI Inference in Corporate GHG Inventories: A Four-Tier Methodology for Scope 3 Category 1 Reporting

ArXiv CS.AI 50%

arXiv:2606.10660v1 Announce Type: cross Abstract: AI inference services -- API subscriptions, enterprise chat tools, and SaaS products with embedded AI features -- fall unambiguously within Scope 3 Category 1 under the Corporate Sustainability Reporting Directive (CSRD), which...

<https://arxiv.org/abs/2606.10660>

35 CONF

Divide and Cooperate: Role-Decomposed Multi-Agent LLM Training with Cross-Agent Learning Signals

ArXiv CS.AI 50%

arXiv:2606.10684v1 Announce Type: cross Abstract: Modern language agents which perform multi-step reasoning have shown strong performance in knowledge-intensive question answering. However, existing approaches typically couple evidence acquisition and answer generation within a...

<https://arxiv.org/abs/2606.10684>

36 CONF

Event-Driven Reinforcement Learning Enables Long-Horizon Control in Semiconductor Fabrication

ArXiv CS.AI 50%

arXiv:2606.10705v1 Announce Type: cross Abstract: Reinforcement learning promises to optimize sequential decisions in large-scale systems. Semiconductor manufacturing systems are stochastic and highly constrained environments where heterogeneous wafers traverse hundreds of...

<https://arxiv.org/abs/2606.10705>

37 CONF

Unifying Local Communications and Local Updates for LLM Pretraining

ArXiv CS.AI 50%

arXiv:2606.11081v1 Announce Type: cross Abstract: Communication-efficient pre-training of LLMs is increasingly important as training draws on compute distributed across clusters, data centers, and lower-bandwidth links. Many practical methods reduce communication frequency but...

<https://arxiv.org/abs/2606.11081>

38 CONF

Robust Deep Reinforcement Learning Through Adversarial Attacks and Training : A Survey

ArXiv CS.AI 50%

arXiv:2403.00420v3 Announce Type: replace-cross Abstract: Deep Reinforcement Learning (DRL) is a subfield of machine learning for training autonomous agents that take sequential actions across complex environments. Despite its significant performance in well-known environments,...

<https://arxiv.org/abs/2403.00420>

39 CONF

Attention Expansion: Enhancing Keyphrase Extraction from Long Documents with Attention-Augmented Contextualized Embeddings

ArXiv CS.AI 50%

arXiv:2606.10716v1 Announce Type: cross Abstract: Pre-trained language models (PLMs) have achieved strong performance in keyphrase extraction (KPE), largely due to their ability to generate rich contextualized representations. However, long-document KPE remains challenging...

<https://arxiv.org/abs/2606.10716>

40 CONF

Detecting Knowledge Gaps from Conversational AI Interactions Using Curriculum Prerequisite Graphs

ArXiv CS.AI 50%

arXiv:2606.10736v1 Announce Type: cross Abstract: Large online courses generate thousands of student questions directed at conversational AI teaching assistants, yet these interaction logs remain largely untapped as diagnostic signals. We present a pipeline that maps student...

<https://arxiv.org/abs/2606.10736>

41 CONF

Toward Secure LLM Agents: Threat Surfaces, Attacks, Defenses, and Evaluation

ArXiv CS.AI 50%

arXiv:2606.10749v1 Announce Type: cross Abstract: Large language model (LLM) agents are rapidly moving from conversational interfaces to software components that plan, invoke tools, maintain memory, and act on external environments. This transition changes the nature of security...

<https://arxiv.org/abs/2606.10749>

42 CONF

Dep-LLM: Training-Free Depression Diagnosis via Evidence-Guided Structured Multi-factor with Reliable LLM Reasoning

ArXiv CS.AI 50%

arXiv:2606.10796v1 Announce Type: cross Abstract: Automatic Depression Detection (ADD) from clinical interviews is a pivotal task in computational mental health, yet it remains challenging due to two critical obstacles: 1) difficulty in modeling complex but sparsely distributed...

<https://arxiv.org/abs/2606.10796>

43 CONF

Boosting ECG Classification Performance by Pre-training with Synthesized Data

ArXiv CS.AI 50%

arXiv:2606.10802v1 Announce Type: cross Abstract: Deep Neural Networks (DNNs) typically require extensive datasets for effective training. In the medical domain, acquiring large-scale data is often challenging due to privacy concerns and the rarity of certain diseases. To...

<https://arxiv.org/abs/2606.10802>

44 CONF

Geometrically Averaged Hard Target Updates for Linear Q-Learning

ArXiv CS.AI 50%

arXiv:2606.10835v1 Announce Type: cross Abstract: Periodic hard target updates are among the most common stabilization devices in modern deep Q-learning. Recent studies suggest that target updates can improve stability in Q-learning with function approximation, including linear...

<https://arxiv.org/abs/2606.10835>

45 CONF

From Perception to Action: Can UI Interventions Foster Sustainable LLM Chatbot

ArXiv CS.AI 50%

arXiv:2606.10861v1 Announce Type: cross Abstract: LLM-powered chatbots are increasingly embedded in everyday workflows, raising sustainability concerns due to their energy use. Most mitigation strategies emphasize model or infrastructure efficiency, while the user-interface (UI)...

<https://arxiv.org/abs/2606.10861>

46 CONF

Improving Text-Instance Alignment Of Foreground Conditioned Out-Painting Via Customized Concept Embedding

ArXiv CS.AI 50%

arXiv:2606.10892v1 Announce Type: cross Abstract: To showcase products, merchants often incur substantial costs creating high-quality display images. Foreground Conditioned Outpainting (FCO) meets this demand, allowing users to create desired backgrounds for foreground instances...

<https://arxiv.org/abs/2606.10892>

47 CONF

Provenance Tracking in AI Compilers through the Lens of Coalgebra

ArXiv CS.AI 50%

arXiv:2606.10937v1 Announce Type: cross Abstract: AI compilers aggressively rewrite computation graphs through normalization, lowering, and optimization, making it difficult to track the provenance of tensors and operators across compilation. Reliable provenance is essential for...

<https://arxiv.org/abs/2606.10937>

48 CONF

RAT: Reference-Augmented Training for ASV Anti-Spoofing

ArXiv CS.AI 50%

arXiv:2606.10908v1 Announce Type: cross Abstract: We introduce a spoofing countermeasure architecture conditioned on speaker-reference recordings, but observe that it converges to a solution that effectively ignores the reference during inference. Surprisingly, training with a...

<https://arxiv.org/abs/2606.10908>

49 **CONF****A Constrained Natural-Language Interface for Variational Multi-Physics Finite Element Simulations in FEniCS****ArXiv CS.AI** **50%**

arXiv:2606.10928v1 Announce Type: cross Abstract: Large language models can reduce the manual effort required to set up finite element simulations, but they introduce reliability risks when generated solver code lies on the critical path. We present a constrained...

<https://arxiv.org/abs/2606.10928>

50 **CONF****Democratising Camera Trap AI: An Open-Source Model for Detecting UK Mammals****ArXiv CS.AI** **50%**

arXiv:2606.10940v1 Announce Type: cross Abstract: Camera traps have become a cornerstone of biodiversity monitoring, but the artificial intelligence that turns vast quantities of images into usable ecological data is often locked behind commercial platforms or trained on fauna...

<https://arxiv.org/abs/2606.10940>

51 **CONF****Diffusion Forcing Planner: History-Annealed Planning with Time-Dependent Guidance for Autonomous Driving****ArXiv CS.AI** **50%**

arXiv:2606.11019v1 Announce Type: cross Abstract: Learning-based motion planners, despite recent progress, often suffer from temporal inconsistency. Small perturbations across frames can accumulate into unstable trajectories, degrading comfort and safety in closed-loop driving....

<https://arxiv.org/abs/2606.11019>

52 **CONF****AuRA: Internalizing Audio Understanding into LLMs as LoRA****ArXiv CS.AI** **50%**

arXiv:2606.11033v1 Announce Type: cross Abstract: Recent efforts to extend large language models (LLMs) to speech inputs typically rely on cascaded ASR-LLM pipelines, end-to-end speech-language models, or bridge/distillation-based adaptation. While these routes respectively...

<https://arxiv.org/abs/2606.11033>

53 **CONF****Modeling Complex Behaviors: Multi-Personality Composition and Dynamic Switching in Vision-Language Models****ArXiv CS.AI** **50%**

arXiv:2606.11074v1 Announce Type: cross Abstract: With the widespread deployment of Multimodal Large Language Models (MLLMs) in social interaction, understanding and controlling their behavior under complex personality conditions is essential. This paper introduces explicit...

<https://arxiv.org/abs/2606.11074>

54 **CONF****RoboGPT-R1: Enhancing Robot Task Planning with Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2510.14828v3 Announce Type: replace Abstract: Improving the reasoning capabilities of embodied agents is crucial for robots to complete complex human instructions in long-view manipulation tasks successfully. Despite the success of large language models and vision language...

<https://arxiv.org/abs/2510.14828>

55 CONF

RoboNaldo: Accurate, Stable and Powerful Humanoid Soccer Shooting via Motion-Guided Curriculum Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.11092v1 Announce Type: cross Abstract: Elite humanoid soccer shooting requires whole-body stability, high-impulse whole-body interactions, and accuracy to targets. Motion tracking-driven reinforcement learning (RL) provides stability in whole-body movement...

<https://arxiv.org/abs/2606.11092>

56 CONF

Designed by Journalists, but Is It for Readers? Rethinking AI Disclosures and Transparency in News

ArXiv CS.AI 50%

arXiv:2606.11116v1 Announce Type: cross Abstract: As newsrooms integrate generative AI, journalists face a disclosure challenge: how to communicate AI involvement in ways that maintain reader trust. Current practice offers two approaches: brief one-line labels or detailed...

<https://arxiv.org/abs/2606.11116>

57 CONF

TRACE: A Unified Rollout Budget Allocation Framework for Efficient Agentic Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.11119v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards (RLVR) is a promising approach for enhancing reasoning and agentic behavior in large language models. However, rollout-intensive policy optimization is often limited by insufficient...

<https://arxiv.org/abs/2606.11119>

58 CONF

Data assimilation for subsurface flow using latent diffusion model parameterization: performance of ensemble-Kalman and Monte Carlo techniques

ArXiv CS.AI 50%

arXiv:2606.11140v1 Announce Type: cross Abstract: Data assimilation (DA) in subsurface flow entails calibrating model parameters to match observed data, typically at wells, while preserving geological realism. Latent diffusion models (LDMs) provide efficient mappings from...

<https://arxiv.org/abs/2606.11140>

59 CONF

Flaws in the LLM Automation Narrative

ArXiv CS.AI 50%

arXiv:2606.11166v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly described as performing at the level of human experts on knowledge economy tasks. These claims are primarily based on how LLMs perform on benchmarking tasks that measure average...

<https://arxiv.org/abs/2606.11166>

60 CONF

Position: The ML Community Must Build an AI-Augmented Peer-Review Ecosystem

ArXiv CS.AI 50%

arXiv:2506.08134v4 Announce Type: replace Abstract: Peer review, the bedrock of scientific advancement in machine learning (ML), is strained by a crisis of scale. Exponential growth in manuscript submissions to premier ML venues such as NeurIPS, ICML, and ICLR is outpacing the...

<https://arxiv.org/abs/2506.08134>

61 CONF

Why Does Reasoning Length Converge? Unveiling the Underfitting-Overfitting Trade-off in Chain-of-Thought

ArXiv CS.AI 50%

arXiv:2509.04027v4 Announce Type: replace Abstract: Test-time scaling, primarily manifested through multi-step Chain-of-Thought (CoT) reasoning via Reinforcement Learning (RL), has emerged as a pivotal paradigm for enhancing the reasoning capabilities of Large Language Models...

<https://arxiv.org/abs/2509.04027>

62 CONF

PSEBench: A Controllable and Verifiable Benchmark for Evaluating LLMs in Patient Safety Event Triage

ArXiv CS.AI 50%

arXiv:2606.05463v2 Announce Type: replace Abstract: Patient safety event triage, determining whether a clinical event is reportable under jurisdiction-specific policy, is a high-stakes task typically performed manually by patient safety experts. Although LLMs may support this...

<https://arxiv.org/abs/2606.05463>

63 CONF

Emergent alignment and the projectability of ethical personas

ArXiv CS.AI 50%

arXiv:2606.09475v2 Announce Type: replace Abstract: Work on 'emergent misalignment' shows that finetuning LLMs on narrow tasks can induce broadly misaligned behavior. This supports the 'persona selection' (PSM) hypothesis: during pre-training, LLMs learn to simulate different...

<https://arxiv.org/abs/2606.09475>

64 CONF

BadRobot: Jailbreaking Embodied LLM Agents in the Physical World

ArXiv CS.AI 50%

arXiv:2407.20242v5 Announce Type: replace-cross Abstract: Embodied AI represents systems where AI is integrated into physical entities. Large Language Model (LLM), which exhibits powerful language understanding abilities, has been extensively employed in embodied AI by...

<https://arxiv.org/abs/2407.20242>

65 CONF

Visual-TCAV: Concept-based Attribution and Saliency Maps for Post-hoc Explainability in Image Classification

ArXiv CS.AI 50%

arXiv:2411.05698v3 Announce Type: replace-cross Abstract: Convolutional Neural Networks (CNNs) have shown remarkable performance in image classification. However, interpreting their predictions is challenging due to the size and complexity of these models. State-of-the-art...

<https://arxiv.org/abs/2411.05698>

66 CONF

NuWa: Deriving Lightweight Class-Specific Vision Transformers for Edge Devices

ArXiv CS.AI 50%

arXiv:2504.03118v2 Announce Type: replace-cross Abstract: Vision Transformers (ViTs) often need to be compressed for deployment on resource-constrained edge devices like drones and smart vehicles. However, existing model compression methods ignore that many edge devices only...

<https://arxiv.org/abs/2504.03118>

67 CONF

A Survey of Robotic Navigation and Manipulation with Physics Simulators in the Era of Embodied AI

ArXiv CS.AI 50%

arXiv:2505.01458v2 Announce Type: replace-cross Abstract: Navigation and manipulation are core capabilities in Embodied AI, but training agents to perform them directly in the real world is costly, time-consuming, and unsafe. Therefore, sim-to-real transfer has emerged as a key...

<https://arxiv.org/abs/2505.01458>

68 CONF

Attacks on Machine-Text Detectors Retain Stylistic Fingerprints

ArXiv CS.AI 50%

arXiv:2505.14608v3 Announce Type: replace-cross Abstract: Despite considerable progress in the development of machine-text detectors, the ease with which machine-text can be manipulated to evade detection has led to suggestions that the problem is inherently intractable. In this...

<https://arxiv.org/abs/2505.14608>

69 CONF

TinyTroupe: An LLM-powered Multiagent Persona Simulation Toolkit

ArXiv CS.AI 50%

arXiv:2507.09788v3 Announce Type: replace-cross Abstract: Recent advances in Large Language Models (LLM) have led to a new class of autonomous agents, renewing and expanding interest in the area. LLM-powered Multiagent Systems (MAS) have thus emerged, both for assistive and...

<https://arxiv.org/abs/2507.09788>

70 CONF

GRID: Scaling Task-Agnostic Inference in Continual Prompt Tuning

ArXiv CS.AI 50%

arXiv:2507.14725v4 Announce Type: replace-cross Abstract: Prompt-based continual learning (CL) offers a parameter-efficient way to adapt large language models (LLMs) across task sequences. However, existing methods often rely on task-aware inference and maintain an expanding set...

<https://arxiv.org/abs/2507.14725>

71 CONF

LLM-Aided Joint Secrecy Precoding and Trajectory for RSMA-Based Heterogeneous UAV Networks

ArXiv CS.AI 50%

arXiv:2507.17188v2 Announce Type: replace-cross Abstract: This paper investigates secure communications in rate-splitting multiple access (RSMA) enabled heterogeneous UAV networks, where multiple UAVs collaboratively serve ground terminals in the presence of eavesdroppers. By...

<https://arxiv.org/abs/2507.17188>

72 CONF

Whisfusion: Parallel ASR Decoding with Masked Diffusion

ArXiv CS.AI 50%

arXiv:2508.07048v2 Announce Type: replace-cross Abstract: Autoregressive (AR) encoder-decoder models dominate high-quality multilingual ASR, but their left-to-right decoders make inference latency scale with transcript length. A natural alternative, CTC-style non-autoregressive...

<https://arxiv.org/abs/2508.07048>

73 CONF

Model-Based Diffusion Sampling for Predictive Control in Offline Decision Making

ArXiv CS.AI 50%

arXiv:2512.08280v3 Announce Type: replace-cross Abstract: Offline decision-making via diffusion models often produces trajectories that are misaligned with system dynamics, limiting their reliability for control. We propose Model Predictive Diffuser (MPDiffuser), a compositional...

<https://arxiv.org/abs/2512.08280>

74 CONF

V-REX: Benchmarking Exploratory Visual Reasoning via Chain-of-Questions

ArXiv CS.AI 50%

arXiv:2512.11995v2 Announce Type: replace-cross Abstract: While many vision-language models (VLMs) are developed to answer well-defined, straightforward questions with highly specified targets, as in most benchmarks, they often struggle in practice with complex open-ended tasks,...

<https://arxiv.org/abs/2512.11995>

75 CONF

HiGR: Industrial-Scale Hierarchical Generative Slate Recommendation Framework in Tencent

ArXiv CS.AI 50%

arXiv:2512.24787v3 Announce Type: replace-cross Abstract: Slate recommendation, which presents users with a ranked item list in a single display, is ubiquitous across mainstream online platforms. While recent generative recommendation methods have shown strong potential in...

<https://arxiv.org/abs/2512.24787>

76 CONF

MMD Guidance: Training-Free Distribution Adaptation for Diffusion Models via Maximum Mean Discrepancy Guidance

ArXiv CS.AI 50%

arXiv:2601.08379v2 Announce Type: replace-cross Abstract: Pre-trained diffusion models have emerged as powerful generative priors for both unconditional and conditional sample generation, yet their outputs often deviate from the characteristics of user-specific target data. Such...

<https://arxiv.org/abs/2601.08379>

77 CONF

Adoption of Generative Artificial Intelligence in the German Software Engineering Industry: An Empirical Study

ArXiv CS.AI 50%

arXiv:2601.16700v2 Announce Type: replace-cross Abstract: Generative artificial intelligence (GenAI) tools have seen rapid adoption among software developers. While adoption rates in the industry are rising, the underlying factors influencing the effective use of these tools,...

<https://arxiv.org/abs/2601.16700>

78 CONF

Exploring Accurate and Transparent Domain Adaptation in Predictive Healthcare via Concept-Grounded Orthogonal Inference

ArXiv CS.AI 50%

arXiv:2602.12542v2 Announce Type: replace-cross Abstract: Deep learning models for clinical event prediction on electronic health records (EHR) often suffer performance degradation when deployed under different data distributions. While domain adaptation (DA) methods can...

<https://arxiv.org/abs/2602.12542>

79 CONF

MedFeat: Model-Aware and Explainability-Driven Feature Engineering with LLMs for Clinical Tabular Prediction

ArXiv CS.AI 50%

arXiv:2603.02221v2 Announce Type: replace-cross Abstract: In clinical tabular prediction, classical machine learning models with feature engineering often outperform neural methods. LLMs are increasingly used to automate this process, acting as domain experts that propose...

<https://arxiv.org/abs/2603.02221>

80 CONF

People-Centred Medical Image Analysis via Fairness-Aware Human-AI Cooperation

ArXiv CS.AI 50%

arXiv:2604.26991v2 Announce Type: replace-cross Abstract: Machine learning models for medical image analysis often exhibit subgroup-dependent performance, which impacts how decisions should be allocated between automated systems and human experts under limited resources. Prior...

<https://arxiv.org/abs/2604.26991>

81 CONF

Communication Dynamics Neural Networks: FFT-Diagonalized Layers for Improved Hessian Conditioning at Reduced Parameter Count

ArXiv CS.AI 50%

arXiv:2605.08171v2 Announce Type: replace-cross Abstract: Communication Dynamics Neural Networks (CDNNs) apply the circulant-spectral machinery of the Communication Dynamics framework to neural-network layer design. We introduce CDLinear, a block-circulant linear layer with...

<https://arxiv.org/abs/2605.08171>

82 CONF

PromptEmbedder: Efficient and Transferable Text Embedding via Dual-LLM Soft Prompting

ArXiv CS.AI 50%

arXiv:2605.28066v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) have demonstrated remarkable efficacy in text embedding, yet current adaptation methods like LoRA face significant bottlenecks in computational efficiency and cross-architecture...

<https://arxiv.org/abs/2605.28066>

83 CONF

CoRe-MoE: Contrastive Reweighted Mixture of Experts for Multi-Terrain Humanoid Locomotion with Gait Adaptation

ArXiv CS.AI 50%

arXiv:2606.04718v2 Announce Type: replace-cross Abstract: Humans primarily rely on walking and running to traverse complex terrains. Similarly, humanoid robots should be able to smoothly transition between walking and running while maintaining natural and stable locomotion....

<https://arxiv.org/abs/2606.04718>

84 CONF

Enhancing AI Interpretability and Safety through Localised Architectures

ArXiv CS.AI 50%

arXiv:2606.07998v2 Announce Type: replace-cross Abstract: Recent advances in generative AI, especially powerful Large Language Models (LLMs) and Large Reasoning Models (LRMs), raise concerns over the interpretability, safety and sustainability of these large and opaque AI...

<https://arxiv.org/abs/2606.07998>