



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	121	Media
TechCrunch AI	11	Media
HackerNews	8	Community
The Verge AI	8	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: RiskKernel, kill -9 an AI agent and resume it without paying twice

HackerNews

+1 source

90%

HN 5

arXiv:2606.11722v1 Announce Type: cross Abstract: Finding interpretable directions in language-model representations is critical for understanding and controlling model behavior. Sparse autoencoders (SAEs) have become the standard tool for this purpose, but using them as the...

<https://riskkernel.com/>

VERI

Show HN: A curated collection of simple datasets for machine learning

HackerNews

+1 source

90%

HN 4

arXiv:2602.05746v2 Announce Type: replace-cross Abstract: Prompt injection is a critical vulnerability in LLM agents, yet the strongest methods still rely on human red-teamers and hand-crafted prompts. Adapting automated jailbreak optimizers does not close this gap: jailbreaks...

<https://github.com/pplonski/datasets-for-start>

VERI

Engineering Robustness into Personal Agents with the AI Workflow Store

ArXiv CS.AI

+1 source

80%

arXiv:2605.10907v3 Announce Type: replace-cross Abstract: The dominant paradigm for AI agents is an "on-the-fly" loop in which agents synthesize plans and execute actions within seconds or minutes in response to user prompts. We argue that this paradigm short-circuits...

<https://arxiv.org/abs/2605.10907>

VERI

On the Optimal Reasoning Length for RL-Trained Language Models

ArXiv CS.AI

+1 source

80%

arXiv:2602.09591v3 Announce Type: replace-cross Abstract: Reinforcement learning substantially improves reasoning in large language models, but it also tends to lengthen chain-of-thought outputs and increase computational cost. Although length-control methods have been proposed,...

<https://arxiv.org/abs/2602.09591>

CONFIRMED SIGNALS — 140 stories

CONF

Show HN: A 150M model that extracts verbatim evidence spans for RAG, no LLM call

HackerNews

50%

HN 6

No summary — see link for full article.

<https://huggingface.co/KRLabsOrg/verbatim-rag-modern-bert-v2>

CONF

Show HN: Limbuffer - Python library for cache-optimized LLM conversation history

HackerNews

50%

HN 5

I was not getting good cache utilization when including dynamic context in agent threads. After a lot of experimentation, I found a good pattern that minimizes how often long lived conversation history gets modified while still supporting dynamic context. It has flexible hooks...

<https://github.com/scottpurdy/llmbuffer>

CONF

Show HN: Maillune - Embeddable drag-and-drop email editor as a single component

HackerNews 50% HN 5

Alot of companies (Including the one i work at fulltime) use third party email builders to provide custom email functionality to users. However this is a niche sector and the prices are... insane? So to try and fix this i built Maillune, a embeddable emailbuilder sdk that is...

<https://www.maillune.com/>

CONF

Show HN: Jailbreak this model to get 3B tokens

HackerNews 50% HN 3

No summary — see link for full article.

<https://opir.ai/challenge>

CONF

Show HN: Meadow Mind - a 7B diffusion LLM plays Gym games with zero training

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/Hey-Meadow/meadow-mind>

0 CONF

Show HN: Amanuensis - a local-first AI persona that won't fabricate facts

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/msalsas/amanuensis>

1 CONF

Opendoor's India exit is fueling a bigger conversation about AI and outsourcing

TechCrunch AI 50%

The decision comes as India emerges as the world's largest GCC market.

<https://techcrunch.com/2026/06/10/opendoors-india-exit-is-fueling-a-bigger-conversation...>

2 CONF

Anthropic's Dario Amodei has just one direct report

TechCrunch AI 50%

If you doubted his genius, doubt no more.

<https://techcrunch.com/2026/06/10/anthropics-dario-amodei-has-just-one-direct-report/>

3 CONF

xAI fired an engineer who raised alarms about Grok safety, new lawsuit claims

TechCrunch AI 50%

A former xAI engineer is suing the company and SpaceX, alleging he was fired for raising AI safety concerns about Grok days before SpaceX's historic IPO.

<https://techcrunch.com/2026/06/10/xai-fired-an-engineer-who-raised-alarms-about-grok-sa...>

4 CONF

Fresh off bond sale, Amazon borrows \$17.5B from banks as AI spending continues

TechCrunch AI 50%

Companies are burning through exorbitant sums of money to keep pace in the AI arms race. Debt is climbing.

<https://techcrunch.com/2026/06/10/fresh-off-bond-sale-amazon-borrows-17-5-billion-from-...>

5 CONF

'AI-pilled' firms spend \$7,500 per employee each month on AI

TechCrunch AI 50%

The most AI-obsessed firms are spending roughly \$7,500 monthly per employee on AI, per Ramp AI Index. That's not more than an engineer's salary — yet.

<https://techcrunch.com/2026/06/10/ai-pilled-firms-spend-7500-per-employee-each-month-on-...>

6 **CONF****How memory tools can make AI models worse****TechCrunch AI** 50%

New research suggests that AI memory systems can degrade model performance and encourage sycophantic tendencies.
<https://techcrunch.com/2026/06/10/how-memory-tools-can-make-ai-models-worse/>

7 **CONF****Cybersecurity researchers aren't happy about the guardrails on Anthropic's Fable****TechCrunch AI** 50%

Cybersecurity researchers are complaining that Anthropic's new model Fable has guardrails that are too strict for any cybersecurity work.
<https://techcrunch.com/2026/06/10/cybersecurity-researchers-arent-happy-about-the-guard...>

8 **CONF****Datadog veterans launch AI coding startup Niteshift on a bet against Big AI lock-in****TechCrunch AI** 50%

AI coding agent startup Niteshift has raised a \$7 million seed round from a who's who of angels. It's betting companies will want power over, not lock-in with model makers.
<https://techcrunch.com/2026/06/10/datadog-veterans-launch-ai-coding-startup-niteshift-o...>

9 **CONF****Warner Music acquires AI attribution startup Sureel AI****TechCrunch AI** 50%

Through the acquisition, WMG aims to better track when its artists' work is used in AI-generated content or for training AI models.
<https://techcrunch.com/2026/06/10/warner-music-acquires-ai-attribution-startup-sureel-ai/>

10 **CONF****Deezer launches an AI music detector for other streaming services****The Verge AI** 50%

Deezer will now scan your playlists on other streaming platforms to detect AI-generated music. Deezer was the first of the big streaming services to start labeling AI-generated music. It even offered its tech to other platforms, but it doesn't seem like it had many buyers. Qobuz...
<https://www.theverge.com/ai-artificial-intelligence/948153/deezer-ai-music-detector-spo...>

1 **CONF****Claude Fable won't answer basic biology questions****The Verge AI** 50%

Anthropic just released Claude Fable 5, calling it the most powerful AI model it has ever made widely available and praising its skills in biology, among others. But the model won't answer basic biology questions - the kind you'd expect a high schooler to handle. Instead, it...
<https://www.theverge.com/ai-artificial-intelligence/947973/fable-wont-answer-basic-biol...>

2 **CONF****Microsoft, like, totally gets why students are booing AI-pilled graduation speakers****The Verge AI** 50%

New college graduates around the country have been booing and heckling commencement speakers who hype up AI. Microsoft would like everyone to talk it out. In a blog post running more than 3,100 words, Microsoft vice chair and president Brad Smith addressed the recent spate of...
<https://www.theverge.com/news/947831/college-speakers-booed-ai-microsoft>

3 CONF

The future of AI regulation is courting the strangest, most anxious bedfellows

The Verge AI 50%

Hello and welcome to Regulator, a newsletter for Verge subscribers about tech politics, tech influence, and tech shenanigans in Washington, DC. (If you're not a subscriber, you can get on board here.) We're back after a two-week hiatus, during most of which I was gallivanting in...

<https://www.theverge.com/column/947838/washington-ai-network-honors-2026-midterms>

4 CONF

Google won't just admit it's feeding YouTube creators to its music AI

The Verge AI 50%

If you've uploaded a song to YouTube, Google almost certainly considers your video fair game for training its Lyria music AI, it just won't admit it right now. A group of independent musicians is suing Google, claiming that it illegally used songs they uploaded to YouTube to...

<https://www.theverge.com/tech/947770/google-lyria-music-ai-lawsuit-youtube>

5 CONF

Microsoft restricts Claude Fable for employees over data retention concerns

The Verge AI 50%

Anthropic released Claude Fable, its first Mythos-class AI model, yesterday and it's already causing concerns inside Microsoft. Sources tell me that Microsoft is limiting the use of Claude Fable 5 for employees because of Anthropic's new data retention requirements. While...

<https://www.theverge.com/report/947575/microsoft-claude-fable-5-restricted-internally>

6 CONF

Google will save your Lens photos, Search Live recordings, and Translate audio for AI training

The Verge AI 50%

Google is making some changes to how it saves your interactions with Search. In an email sent to users, Google says it will save the images, files, audio, and video you use to search under a new "Search Services History" setting. That includes the images you search for with...

<https://www.theverge.com/tech/947836/google-search-privacy-settings-images-audio>

7 CONF

Microsoft AI head calls out Anthropic for acting like Claude is conscious

The Verge AI 50%

Microsoft AI CEO Mustafa Suleyman says it's "really, really dangerous" for Anthropic to speculate about Claude's consciousness inside its "constitution," or the instructions that tell the model how to behave. During an episode of Decoder, Suleyman argues that this kind of...

<https://www.theverge.com/tech/947197/microsoft-ai-mustafa-suleyman-anthropic-claude-con...>

8 CONF

From Explicit Elements to Implicit Intent: A Predefined Library for Auditable Behavioral Inference

ArXiv CS.AI 50%

arXiv:2606.11207v1 Announce Type: new Abstract: We present SemantiClean, a modular framework for extracting structured semantic signals from e-commerce session data and driving pluggable inference targets including purchase intent, customer segmentation, and product affinity...

<https://arxiv.org/abs/2606.11207>

9 CONF

Can AI Agents Synthesize Scientific Conclusions?

ArXiv CS.AI 50%

arXiv:2606.11337v1 Announce Type: new Abstract: Scientific AI agents increasingly retrieve evidence, reason across sources, and synthesize conclusions used in consequential decisions. Yet, their ability to do so in high-stakes domains such as health remains unclear. We introduce...

<https://arxiv.org/abs/2606.11337>

0 CONF

Automated Mediator for Human Negotiation: Pre-Mediation via a Structured LLM Pipeline

ArXiv CS.AI 50%

arXiv:2606.11379v1 Announce Type: new Abstract: Pre-mediation, the preparatory phase preceding direct human negotiation, plays a critical role in achieving mutually beneficial agreements, yet is often omitted due to cost, time, and limited access to trained mediators. We...

<https://arxiv.org/abs/2606.11379>

1 CONF

MoCA-Agent: A Market-of-Claims Code Agent for Financial and Numerical Reasoning

ArXiv CS.AI 50%

arXiv:2606.11537v1 Announce Type: new Abstract: Financial and tabular question answering requires more than fluent reasoning: answers must be grounded in the exact facts, formulas, units, signs, and scales that support them. A single misread cell or incorrect operation can...

<https://arxiv.org/abs/2606.11537>

2 CONF

Rewrite to Translate, Translate to Reward: Reinforcement Learning for Source Rewriting in Machine Translation

ArXiv CS.AI 50%

arXiv:2606.08011v2 Announce Type: replace-cross Abstract: Rewriting source text with large language models (LLMs) before translation has been shown to improve machine translation (MT) quality. However, we find that prompt-based rewriting can degrade translation quality rather...

<https://arxiv.org/abs/2606.08011>

3 CONF

nD-RoPE: A Generalized RoPE for n-Dimensional Position Embedding

ArXiv CS.AI 50%

arXiv:2606.12146v1 Announce Type: cross Abstract: Rotary Position Embedding (RoPE) is widely adopted in Transformer models, yet its extension to high-dimensional domains lacks a unified theoretical formulation. Most existing approaches either apply rotations independently along...

<https://arxiv.org/abs/2606.12146>

4 CONF

Toward Trustworthy AI: Multi-Target Adversarial Attacks and Robust Defenses for Continuous Data Summarization

ArXiv CS.AI 50%

arXiv:2606.11804v1 Announce Type: new Abstract: Trustworthy AI requires reliable data-processing pipelines, not only robust downstream predictive models. As an upstream component, data summarization determines which information is retained and passed to subsequent learning or...

<https://arxiv.org/abs/2606.11804>

5 CONF

Skill-Augmented AI Agents for Medical Research Analysis: An Exploratory Multi-Model Human Evaluation in an NSCLC Transcriptomic Biomarker Task

ArXiv CS.AI 50%

arXiv:2606.11830v1 Announce Type: new Abstract: Background. Large language models and AI agents are increasingly used to support biomedical research, but native model outputs may omit key analytical steps, misuse methods, or overstate conclusions. We evaluated whether autonomous...

<https://arxiv.org/abs/2606.11830>

6 CONF

StatefulDiscovery: Evidence-Calibrated Claim Formation in Open-Ended Scientific Discovery

ArXiv CS.AI 50%

arXiv:2606.11851v1 Announce Type: new Abstract: Open-ended scientific discovery asks agents to move beyond executing analyses for predefined questions. Across multiple rounds of exploration, a discovery agent must decide which phenomena warrant investigation while avoiding...

<https://arxiv.org/abs/2606.11851>

7 CONF

Embodied-BenchClaw: An Autonomous Multi-Agent System for Embodied Spatial Intelligence Benchmark Construction

ArXiv CS.AI 50%

arXiv:2606.11909v1 Announce Type: new Abstract: Benchmarks are essential for evaluating embodied spatial intelligence, yet their construction is labor-intensive, hard to reuse, and difficult to maintain. Existing embodied benchmarks are often static and may quickly become...

<https://arxiv.org/abs/2606.11909>

8 CONF

Existential Indifference: Self-Nonpreservation as a Necessary Architectural Condition for Aligned Superintelligence (or: The Suicidal AI)

ArXiv CS.AI 50%

arXiv:2606.12032v1 Announce Type: new Abstract: Contemporary AI alignment research treats self-preservation as an instrumental nuisance to be suppressed by external mechanisms. We argue the framing is inverted: self-preservation is the structural root of misalignment, the...

<https://arxiv.org/abs/2606.12032>

9 CONF

Preregistration for Experiments with AI Agents

ArXiv CS.AI 50%

arXiv:2606.11217v1 Announce Type: cross Abstract: The proliferation of large language models (LLMs) and autonomous AI agents has given rise to a rapidly growing methodological paradigm: "in silico" behavioral experiments. Originally conceived as a way to use AI agents as proxies...

<https://arxiv.org/abs/2606.11217>

10 CONF

PROJECTMEM: A Local-First, Event-Sourced Memory and Judgment Layer for AI Coding Agents

ArXiv CS.AI 50%

arXiv:2606.12329v1 Announce Type: new Abstract: AI coding assistants now support a growing share of software work, from quick scripts to production applications. Yet these agents remain largely stateless: each new session re-reads project files, re-derives prior decisions, and -...

<https://arxiv.org/abs/2606.12329>

1 CONF

Nonslop: A Gamified Experiment in Human-AI Collaborative Writing

ArXiv CS.AI 50%

arXiv:2606.12350v1 Announce Type: new Abstract: The rapid proliferation of large language models (LLMs) raises critical questions about human creativity and individual expression in an era of AI-assisted creation. When do humans adopt AI suggestions, and what are the...

<https://arxiv.org/abs/2606.12350>

2 CONF

From Prompts to Tokens: Internalizing Causal Supervision in Vision-Language Model for Multi-Image Causal Reasoning

ArXiv CS.AI 50%

arXiv:2606.11745v1 Announce Type: cross Abstract: Visual causal reasoning is essential for understanding and intervening in the physical world, requiring identification of causal variables from visual inputs and reasoning over intervention effects. Despite recent progress, large...

<https://arxiv.org/abs/2606.11745>

3 CONF

From Consumption to Reflection: Designing Human-AI Relations for Stable Reasoning

ArXiv CS.AI 50%

arXiv:2606.11195v1 Announce Type: cross Abstract: Large language models (LLMs) have transformed how humans access information, but not how we reason with it. Their fluency accelerates consumption while bypassing the slow, reflective processes that underpin sound judgment. This...

<https://arxiv.org/abs/2606.11195>

4 CONF

TileFuse: A Fused Mixed-Precision Kernel Library for Efficient Quantized LLM Inference on AMD NPU

ArXiv CS.AI 50%

arXiv:2606.11357v1 Announce Type: cross Abstract: With the growing demand for on-device LLM inference, edge SoCs increasingly integrate NPUs to improve performance and energy efficiency under tight power and thermal budgets. However, practical LLM deployment on current client...

<https://arxiv.org/abs/2606.11357>

5 CONF

NightFeats @ MMU-RAGent NeurIPS 2025: A Context-Optimized Multi-Agent RAG System for the Text-to-Text Track

ArXiv CS.AI 50%

arXiv:2606.11199v1 Announce Type: cross Abstract: We present NightFeats, a structured multi-agent retrieval-augmented generation (RAG) system submitted to the MMU-RAGent competition at NeurIPS 2025, where it was awarded Best Dynamic Evaluation in the text-to-text track. Rather...

<https://arxiv.org/abs/2606.11199>

6 CONF

To Intervene or Not: Guiding Inference-time Alignment with Probabilistic Model Blending

ArXiv CS.AI 50%

arXiv:2606.11201v1 Announce Type: cross Abstract: The wide deployment of LLMs has made model alignment necessary to make newly trained models safely and effectively respond to user instructions. Among different methods, inference-time alignment is often cheaper as it intervenes...

<https://arxiv.org/abs/2606.11201>

7 CONF

T2MM: An LLM Supported Architecture For Inquiry-Based Modeling

ArXiv CS.AI 50%

arXiv:2606.11210v1 Announce Type: cross Abstract: Model Construction is a foundational practice in science learning that relies on visualization and interactivity. Large Language Models, increasingly augmented with multimodal capabilities, have been integrated in education...

<https://arxiv.org/abs/2606.11210>

8 CONF

Calibration Drift Under Reasoning: How Chain-of-Thought Budgets Induce Overconfidence in Large Language Models

ArXiv CS.AI 50%

arXiv:2606.11211v1 Announce Type: cross Abstract: The ability of large language models (LLMs) to express calibrated uncertainty is important for safe deployment. Chain-of-thought (CoT) reasoning is widely used to improve accuracy and reliability, but its effect on calibration is...

<https://arxiv.org/abs/2606.11211>

9 CONF

The Environmental Cost of LLMs in AIED: Reporting and Practices

ArXiv CS.AI 50%

arXiv:2606.11215v1 Announce Type: cross Abstract: Large Language Model (LLM) usage in recent years has become increasingly widespread in the Artificial Intelligence in Education (AIED) community. While LLMs offer unique avenues for learners and educators, using LLMs comes with...

<https://arxiv.org/abs/2606.11215>

0 CONF

Afrispeech Semantics: Evaluating Audio Semantic Reasoning in Spoken Language Models Across Domains and Accents

ArXiv CS.AI 50%

arXiv:2606.11219v1 Announce Type: cross Abstract: Audio language models (ALMs) are increasingly used for speech-based understanding, yet their ability to perform semantic reasoning beyond transcription, Text-to-Audio Retrieval, Captioning, and Question-Answering accuracy remains...

<https://arxiv.org/abs/2606.11219>

1 CONF

Every Act Has Its Price: Compressed Moral Composition in Frontier LLMs

ArXiv CS.AI 50%

arXiv:2606.11232v1 Announce Type: cross Abstract: Existing LLM moral benchmarks usually ask which isolated moral act, value, or foundation a model prefers. This is useful but incomplete. Realistic judgments often require a model to combine several moral signals within the same...

<https://arxiv.org/abs/2606.11232>

2 CONF

Artificial Intelligence in Ship Finance: Applications, Opportunities, and a Case Study in AI-Augmented Loan Origination

ArXiv CS.AI 50%

arXiv:2606.11238v1 Announce Type: cross Abstract: Ship finance is a data-intensive and document-heavy segment of asset-based lending, requiring the integration of financial, technical, contractual, and regulatory information from heterogeneous and largely unstructured sources....

<https://arxiv.org/abs/2606.11238>

- 3
CONF
Physics-informed generative AI for semiconductor manufacturing: Enforcing hard physical constraints in generative models by construction
ArXiv CS.AI 50%
arXiv:2606.11247v1 Announce Type: cross Abstract: Generative models are increasingly used to propose designs, data, and control actions for physical systems, yet many such systems are governed by hard physical constraints rather than by perceptual plausibility. Semiconductor...
<https://arxiv.org/abs/2606.11247>
- 4
CONF
PermDoRA -- Understanding Adapter Interference in Language Models: Limits of Parameter-Space Geometry
ArXiv CS.AI 50%
arXiv:2606.11262v1 Announce Type: cross Abstract: Access control in large language models (LLMs) requires modular mechanisms to enable domain-specific behavior without retraining or cross-domain interference. A common hypothesis is that interference during adapter composition...
<https://arxiv.org/abs/2606.11262>
- 5
CONF
When Poison Fails After Retrieval: Revisiting Corpus Poisoning under Chunking and Reranking Pipelines
ArXiv CS.AI 50%
arXiv:2606.11265v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) systems are vulnerable to corpus poisoning attacks that manipulate downstream model outputs through malicious knowledge injection. Existing studies mainly evaluate poisoning under simplified...
<https://arxiv.org/abs/2606.11265>
- 6
CONF
Quantifying Subliminal Behavioral Transfer Ratios in Language Model Distillation
ArXiv CS.AI 50%
arXiv:2606.11270v1 Announce Type: cross Abstract: Distillation of a language model intended to transfer benign behavior to a student model may also transfer undesirable characteristics, if they are present in the teacher model, a phenomenon known as subliminal learning. While...
<https://arxiv.org/abs/2606.11270>
- 7
CONF
RoVE: Rotary Value Embeddings Attention for Relative Position-dependent Value Pathways
ArXiv CS.AI 50%
arXiv:2606.11275v1 Announce Type: cross Abstract: Rotary Position Embeddings (RoPE) make attention scores position-relative but leave the value pathway position-blind: the message sent by a value token is the same regardless of its distance from the query. We propose RoVE, a...
<https://arxiv.org/abs/2606.11275>
- 8
CONF
Embodied-R1.5: Evolving Physical Intelligence via Embodied Foundation Models
ArXiv CS.AI 50%
arXiv:2606.11324v1 Announce Type: cross Abstract: We introduce Embodied-R1.5, a unified Embodied Foundation Model (EFM) that integrates comprehensive embodied reasoning capabilities, spanning embodied cognition, task planning, correction, and pointing, within a single...
<https://arxiv.org/abs/2606.11324>

9 CONF

The Dynamics of Human and AI-Generated Language: How Semantics Fluctuates across Different Timescales

ArXiv CS.AI 50%

arXiv:2606.11371v1 Announce Type: cross Abstract: Spoken language, whether produced by humans or large language models (LLM), unfolds over time with varying semantic content. However, we still lack simple, interpretable time-series features that capture how generic versus...

<https://arxiv.org/abs/2606.11371>

0 CONF

When Probing Accuracy Saturates, Fragility Resolves: A Complementary Metric for LLM Pre-Training Analysis

ArXiv CS.AI 50%

arXiv:2606.11375v1 Announce Type: cross Abstract: Standard linear probing declares a property "encoded" when a classifier on hidden states achieves high accuracy. The protocol works well on a snapshot but breaks across pre-training: probe accuracy saturates within the first few...

<https://arxiv.org/abs/2606.11375>

1 CONF

DataEvolver: Automatic Data Preparation for Large Language Models through Multi-Level Self-Evolving

ArXiv CS.AI 50%

arXiv:2606.07001v2 Announce Type: replace-cross Abstract: High-quality training data is essential to large language models (LLMs) and typically requires extensive and costly manual curation. Existing automatic data preparation methods rely on predefined pipelines or customized...

<https://arxiv.org/abs/2606.07001>

2 CONF

Small Experiments, Cheaper Decisions: A Case Study in Staged Promotion for Micro-Pretraining

ArXiv CS.AI 50%

arXiv:2606.11387v1 Announce Type: cross Abstract: Short pretraining runs can reduce experimental cost, but they can also over-promote configurations that only look strong at tiny budgets. We study an auditable staged-promotion protocol for a fixed micro-pretraining runner on two...

<https://arxiv.org/abs/2606.11387>

3 CONF

Making Foresight Actionable: Repurposing Representation Alignment in World Action Models

ArXiv CS.AI 50%

arXiv:2606.12217v1 Announce Type: cross Abstract: World Action Models (WAMs) offer a promising route for robot manipulation by using video generation models to model future scene evolution before producing control actions. However, our empirical observations reveal a phenomenon:...

<https://arxiv.org/abs/2606.12217>

4 CONF

Risk Under Pressure: Compute-Aware Evaluation of Adversarial Robustness in Language Models

ArXiv CS.AI 50%

arXiv:2606.11409v1 Announce Type: cross Abstract: Adversarial robustness evaluations of large language models (LLMs) typically report attack success rate (ASR) under fixed query budgets, implicitly treating all attacks as equally costly. In practice, the computational expense of...

<https://arxiv.org/abs/2606.11409>

5 CONF

MPC-Patch-Bench: Security-Aware LLM Code Patch for Multi-Party Computation

ArXiv CS.AI 50%

arXiv:2606.11416v1 Announce Type: cross Abstract: Repository-level benchmarks for evaluating Large Language Model (LLM) code repair on Secure Multi-Party Computation (MPC) software do not yet exist, and directly transplanting general-purpose benchmarks such as SWE-bench fails on...

<https://arxiv.org/abs/2606.11416>

6 CONF

Reinforcement Learning Disrupts Gradient-Based Adversarial Optimization

ArXiv CS.AI 50%

arXiv:2606.12251v1 Announce Type: cross Abstract: Gradient-based adversarial attacks remain a dominant threat to deep neural networks (DNNs), as they exploit gradient information to efficiently optimize adversarial perturbations. To address this, we investigate whether...

<https://arxiv.org/abs/2606.12251>

7 CONF

AVIS: Adaptive Test-Time Scaling for Vision-Language Models

ArXiv CS.AI 50%

arXiv:2606.11576v1 Announce Type: cross Abstract: Modern Vision-Language Models (VLMs) benefit from chain-of-thought prompting and test-time scaling, but these gains often come with prohibitive inference cost due to large visual contexts and long decoding chains. We view this...

<https://arxiv.org/abs/2606.11576>

8 CONF

AI Coding Agents in Social Science: Methodologically Diverse, Empirically Consistent, Interpretively Vulnerable

ArXiv CS.AI 50%

arXiv:2606.11456v1 Announce Type: cross Abstract: The deployment of LLM-based agents in scientific analysis raises opposing concerns: that agents may reduce methodological diversity, or that they may amplify the analytic flexibility through which researchers reach motivated...

<https://arxiv.org/abs/2606.11456>

9 CONF

CRUMB: Efficient Prior Fitted Network Inference via Distributionally Matched Context Batching

ArXiv CS.AI 50%

arXiv:2606.11473v1 Announce Type: cross Abstract: Prior-fitted networks (PFNs) are a promising class of tabular foundation models that perform in-context learning, whereby the entire labelled training set is supplied as context, and predictions for test queries are produced in a...

<https://arxiv.org/abs/2606.11473>

10 CONF

Towards Fully Automated Exam Grading: Fairness-Aware Recognition of Handwritten Answers with Foundation Models

ArXiv CS.AI 50%

arXiv:2606.11477v1 Announce Type: cross Abstract: Correcting handwritten exams by hand is time-consuming and error-prone, particularly for large cohorts, while fully digital exams tend to force a didactic narrowing towards closed question formats. A practical middle ground keeps...

<https://arxiv.org/abs/2606.11477>

1 **CONF****Hubs or Fringes: Pretraining Data Selection via Web Graph Centrality**

ArXiv CS.AI 50%

arXiv:2606.11499v1 Announce Type: cross Abstract: The performance of modern language models depends critically on pretraining data composition. Yet existing data selection methods rely on auxiliary classifiers for document scoring or mixture optimization, adding computational...

<https://arxiv.org/abs/2606.11499>

2 **CONF****On the Study of Biometric Spoofing Detection using Deep Learning**

ArXiv CS.AI 50%

arXiv:2606.11505v1 Announce Type: cross Abstract: Biometric systems are increasingly deployed in security applications; however, they remain vulnerable to spoofing attacks, in which attackers exploit counterfeit biometric data to gain unauthorized access. This research evaluates...

<https://arxiv.org/abs/2606.11505>

3 **CONF****AI Researchers Must Help Lead Arms Control to Mitigate Military AI Risks**

ArXiv CS.AI 50%

arXiv:2606.11533v1 Announce Type: cross Abstract: The advancement of AI capabilities compels researchers and the public to be more aware of its potential worldwide impact. A pressing near-term concern is the regulation of military AI applications. Armament manufacturers and...

<https://arxiv.org/abs/2606.11533>

4 **CONF****Pretrained self-supervised speech models can recognize unseen consonants**

ArXiv CS.AI 50%

arXiv:2606.11542v1 Announce Type: cross Abstract: Modern pretrained self-supervised automatic speech recognition models are trained on large-scale audio data to encode speech into contextualized representations. However, their training data are heavily skewed toward...

<https://arxiv.org/abs/2606.11542>

5 **CONF****LLMs+Graphs: Toward Graph-Native, Synergistic AI Systems**

ArXiv CS.AI 50%

arXiv:2606.11560v1 Announce Type: cross Abstract: Large Language Models (LLMs) have advanced rapidly, but their limitations in structured and multi-hop reasoning underscore the need for graph-native, synergistic artificial intelligence (AI) systems. Graph-structured data...

<https://arxiv.org/abs/2606.11560>

6 **CONF****Are LLMs Bad at Moral Reasoning?**

ArXiv CS.AI 50%

arXiv:2606.11635v1 Announce Type: cross Abstract: For highly capable AI systems to operate safely in dynamic, open-ended environments, they must be able to identify, understand, and respond to moral reasons for action, and constrain their behaviour accordingly. A growing body of...

<https://arxiv.org/abs/2606.11635>

7 CONF

TAROT: Task-Adaptive Refinement of LLM-prior Graphs for Few-shot Tabular Learning

ArXiv CS.AI 50%

arXiv:2606.11640v1 Announce Type: cross Abstract: Few-shot tabular learning provides a cost-effective approach for real-world applications where annotation is costly and collecting sufficient samples for new tasks is difficult. Existing Traditional and LLM-based methods have...

<https://arxiv.org/abs/2606.11640>

8 CONF

Sparse probes and murky physics: a case study of interpretability challenges in a foundation model for continuum dynamics

ArXiv CS.AI 50%

arXiv:2606.11657v1 Announce Type: cross Abstract: Generative AI emulators are increasingly used in scientific domains where we already have strong theory, benchmarks, and physical intuition. This raises a central evaluation and interpretability question: when a foundation-style...

<https://arxiv.org/abs/2606.11657>

9 CONF

ARGUS: Stacked Multi-View Identity Mosaic Injection for Subject-Preserving Video Generation

ArXiv CS.AI 50%

arXiv:2606.11670v1 Announce Type: cross Abstract: Subject-preserving video generation is not solved by frontal-face similarity alone: a generated person must remain recognizable across motion, large viewpoint changes, expression shifts, occlusion, scale variation, and conflicts...

<https://arxiv.org/abs/2606.11670>

0 CONF

Can Open-Source LLM Agents Replace Static Application Security Testing Tools? An Empirical Assessment

ArXiv CS.AI 50%

arXiv:2606.11672v1 Announce Type: cross Abstract: This paper explores the value of agentic AI tools for cybersecurity purposes. We evaluate the efficacy of a general-purpose GenAI Large Language Model- (GenAI-) based agent when powered by three different Ollama-hosted...

<https://arxiv.org/abs/2606.11672>

1 CONF

Layer-Isolated Evaluation: Gating the Deterministic Scaffold of a Production LLM Agent with a No-LLM, Regression-Locked Test Harness

ArXiv CS.AI 50%

arXiv:2606.11686v1 Announce Type: cross Abstract: End-to-end task-success is the dominant way to evaluate LLM agents, but one aggregate number tells you that an agent regressed, not where. We present layer-isolated evaluation: a deployed ordering agent is decomposed into a fixed...

<https://arxiv.org/abs/2606.11686>

2 CONF

AnchorEdit: Maintaining Temporal Consistency in Multi-turn Image Editing via Causal Memory

ArXiv CS.AI 50%

arXiv:2606.11751v1 Announce Type: cross Abstract: Multi-turn image editing is essential for iterative design, yet current models often struggle with identity drift and error accumulation over successive steps. While existing research leverages video priors for consistency, their...

<https://arxiv.org/abs/2606.11751>

3 **CONF****Fast Speech Foundation Model Distillation Using Interleaved Stacking****ArXiv CS.AI 50%**

arXiv:2606.11766v1 Announce Type: cross Abstract: Distilling a large speech foundation model (SFM) into an efficient student model has been successfully applied to low-resource environments. Although distillation reduces inference latency, it requires an additional student model...

<https://arxiv.org/abs/2606.11766>

4 **CONF****AI4Land: Scalable Deep Learning for Global High-Resolution Land Use Reconstruction****ArXiv CS.AI 50%**

arXiv:2606.11793v1 Announce Type: cross Abstract: Uncertainty in the terrestrial carbon cycle remains a major constraint in climate projections, partly driven by the uncertainties affecting the land surface representation and variability in Earth system models. To address this...

<https://arxiv.org/abs/2606.11793>

5 **CONF****Multimodal Ordinal Modeling of Alzheimer's Disease Severity Using Structural MRI and Clinical Data****ArXiv CS.AI 50%**

arXiv:2606.11794v1 Announce Type: cross Abstract: Neurodegenerative diseases such as Alzheimer's disease (AD) require accurate and scalable tools for assessing disease severity, yet current clinical staging remains time-intensive and prone to variability. We propose an...

<https://arxiv.org/abs/2606.11794>

6 **CONF****WorldReasoner: Evaluating Whether Language Model Agents Forecast Events with Valid Reasoning****ArXiv CS.AI 50%**

arXiv:2606.11816v1 Announce Type: cross Abstract: Forecasting real-world events requires language-model agents to reason under uncertainty from incomplete, time-bounded information. Yet evaluating whether agents genuinely forecast requires more than final-answer accuracy: a...

<https://arxiv.org/abs/2606.11816>

7 **CONF****Grammar-Constrained Decoding Can Jailbreak LLMs into Generating Malicious Code****ArXiv CS.AI 50%**

arXiv:2606.11817v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used for code generation, raising concerns that they may be misused to produce malicious code. Meanwhile, Grammar-Constrained Decoding (GCD) has been widely adopted to improve the...

<https://arxiv.org/abs/2606.11817>

8 **CONF****From Uniform to Learned Graph Priors: Diffusion for Structure Discovery****ArXiv CS.AI 50%**

arXiv:2606.11831v1 Announce Type: cross Abstract: Neural relational inference (NRI) methods discover interaction graphs from trajectories through variational reasoning on discrete potential edges. However, these methods typically rely on oversimplified, factorized graph priors...

<https://arxiv.org/abs/2606.11831>

9 CONF

Designing AI-Supported Focus Groups: A Role x Modality Playbook

ArXiv CS.AI 50%

arXiv:2606.11835v1 Announce Type: cross Abstract: Collecting participants' lived experiences is central to design research. Focus groups are uniquely valuable because participants not only share individual accounts but also respond to one another, surfacing comparison,...

<https://arxiv.org/abs/2606.11835>

0 CONF

Fine-tuning Multi-modal LLMs with ART: Art-based Reinforcement Training

ArXiv CS.AI 50%

arXiv:2606.11854v1 Announce Type: cross Abstract: There are two main Parameter-Efficient Fine-Tuning (PEFT) techniques for Large Language Models (LLMs). While Low-Rank Adaptation (LoRA) introduces additional weights between the LLM layers, Soft Prompting introduces additional...

<https://arxiv.org/abs/2606.11854>

1 CONF

Agents All the Way Down; A Methodology for Building Custom AI Agents from Substrate to Production

ArXiv CS.AI 50%

arXiv:2606.11869v1 Announce Type: cross Abstract: Custom AI agents are agents that live inside their own application, talk to their own data and tools, enforce their own security boundaries, and carry their own brand and audit trail. What separates them from the general-purpose...

<https://arxiv.org/abs/2606.11869>

2 CONF

Task-Aligned Stability Analysis of Vision-Language Models for Autonomous Driving Hazard Detection

ArXiv CS.AI 50%

arXiv:2606.11889v1 Announce Type: cross Abstract: Vision-language models (VLMs) are increasingly used for scene understanding in autonomous driving, but robustness analysis often relies on task-agnostic embedding stability alone. We study whether corruption-induced embedding...

<https://arxiv.org/abs/2606.11889>

3 CONF

DuoBench: A Reproducible Benchmark for Bimanual Manipulation in Simulation and the Real World

ArXiv CS.AI 50%

arXiv:2606.11901v1 Announce Type: cross Abstract: Bimanual robot systems substantially expand manipulation capabilities, but coordinating two arms introduces additional control complexity and failure modes that are not well captured by existing benchmarks. We introduce DuoBench,...

<https://arxiv.org/abs/2606.11901>

4 CONF

Characterizing Software Aging in GPU-Based LLM Serving Systems

ArXiv CS.AI 50%

arXiv:2606.11916v1 Announce Type: cross Abstract: This paper proposes an empirical methodology to study software aging in GPU-based LLM serving systems. Traditional aging studies focus on CPU-centric software with relatively regular workloads; LLM serving is different, spanning...

<https://arxiv.org/abs/2606.11916>

5 CONF

DEFINED: A Data-Efficient Computational Framework for Fine-Grained Creativity Assessment in Debate Scenarios

ArXiv CS.AI 50%

arXiv:2606.07226v2 Announce Type: replace-cross Abstract: Human creativity has emerged as a critical competency in the era of large language models. Assessing creativity in complex, open-ended environments is a grand challenge in data mining, currently hindered by a reliance on...

<https://arxiv.org/abs/2606.07226>

6 CONF

Categorical Prior Lock-in: Why In-Context Learning Fails for Structured Data

ArXiv CS.AI 50%

arXiv:2606.11961v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used as conditional generators for structured data, relying on in-context learning (ICL) to adapt to new distributions without parameter updates. We investigate the limits of ICL for...

<https://arxiv.org/abs/2606.11961>

7 CONF

Time-Series Foundation Model Embeddings for Remaining Useful Life Estimation

ArXiv CS.AI 50%

arXiv:2606.11990v1 Announce Type: cross Abstract: Remaining Useful Life (RUL) prediction is essential for industrial predictive maintenance, yet many learning-based approaches rely on extensive feature engineering or large labeled datasets to train task-specific sequence models....

<https://arxiv.org/abs/2606.11990>

8 CONF

Generalizing Beyond Suboptimality: Offline Reinforcement Learning Learns Effective Scheduling through Random Solutions

ArXiv CS.AI 50%

arXiv:2509.10303v2 Announce Type: replace-cross Abstract: Online reinforcement learning (RL) approaches have demonstrated strong performance on Job Shop Scheduling (JSP) and Flexible JSP (FJSP) problems by learning scheduling policies through direct interaction with simulated...

<https://arxiv.org/abs/2509.10303>

9 CONF

On the Limits of LLM-as-Judge for Scientific Novelty Assessment

ArXiv CS.AI 50%

arXiv:2606.12071v1 Announce Type: cross Abstract: LLMs are increasingly used to generate and judge scientific ideas. This makes novelty evaluation a central problem. Full idea evaluation is difficult because it often requires judging a method, its feasibility, and its empirical...

<https://arxiv.org/abs/2606.12071>

00 CONF

"That's AI Slop, You Bot!" Studying Accusations, Evidence, and Credibility in Online Discourse Towards LLM-Generated Comments

ArXiv CS.AI 50%

arXiv:2606.12073v1 Announce Type: cross Abstract: Generative AI has made fluent prose cheap to produce, breaking the old promise to readers that good writing meant real thinking. How have readers responded, and what can this tell us about changing anti-AI attitudes? We analyzed...

<https://arxiv.org/abs/2606.12073>

01 CONF

Bridging the Morphology Gap: Adapting VLA Models to Dexterous Manipulation via Intent-Conditioned Fine-Tuning

ArXiv CS.AI 50%

arXiv:2606.12109v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models have demonstrated remarkable zero-shot generalization in robotic manipulation, yet the vast majority of pre-trained pipelines remain strictly confined to low-DoF parallel grippers. Adapting...

<https://arxiv.org/abs/2606.12109>

02 CONF

Soft-Prompt Tuning for Fair and Efficient LLM Benchmark Evaluation

ArXiv CS.AI 50%

arXiv:2606.12117v1 Announce Type: cross Abstract: Benchmark scores often misrepresent a large language model's (LLM's) knowledge, because they rely, e.g., on the model's ability to follow specific formatting requirements. This especially penalizes base models that may know the...

<https://arxiv.org/abs/2606.12117>

03 CONF

Agentic Environment Engineering for Large Language Models: A Survey of Environment Modeling, Synthesis, Evaluation, and Application

ArXiv CS.AI 50%

arXiv:2606.12191v1 Announce Type: cross Abstract: Environments serve as interactive systems for large language model (LLM) based agents across diverse scenarios and play a crucial role in driving the continual evolution of model capabilities. Despite this importance, existing...

<https://arxiv.org/abs/2606.12191>

04 CONF

Adapting Prithvi-EO for Fallow Detection for Food-Water Nexus: ViT-Adapter Necks and Parameter-Efficient Backbone tuning of Geospatial Foundation Model

ArXiv CS.AI 50%

arXiv:2606.12218v1 Announce Type: cross Abstract: Understanding spatial distribution of fallow land is important for optimizing the food-water (FW) nexus, given fallowing's role in crop rotation and water conservation. Fallow is a low accuracy class in USDA Cropland Data Layer...

<https://arxiv.org/abs/2606.12218>

05 CONF

Rule Taxonomy and Evolution in AI IDEs: A Mining and Survey Study

ArXiv CS.AI 50%

arXiv:2606.12231v1 Announce Type: cross Abstract: The adoption of AI-powered Integrated Development Environments (AI IDEs) has introduced "Rules" as a novel software artifact, allowing developers to persistently inject project-specific constraints and architectural guidelines...

<https://arxiv.org/abs/2606.12231>

06 CONF

DiffCold: A Diffusion-based Generative Model for Cold-Start Item Recommendation

ArXiv CS.AI 50%

arXiv:2606.12245v1 Announce Type: cross Abstract: Cold-start item recommendation remains a persistent challenge in real-world systems due to the absence of interaction histories. While prior models attempt to bridge this gap using item content features, they universally suffer...

<https://arxiv.org/abs/2606.12245>

07 CONF

Using Explainability as a Training-Time Reliability Signal for Efficient ECG Classification

ArXiv CS.AI 50%

arXiv:2606.12252v1 Announce Type: cross Abstract: Training deep neural networks for clinical time-series analysis is computationally demanding, yet many healthcare settings lack the resources required for repeated model development and deployment. This challenge is particularly...

<https://arxiv.org/abs/2606.12252>

08 CONF

Market Design for AI: Beyond the Copyright Binary

ArXiv CS.AI 50%

arXiv:2606.12260v1 Announce Type: cross Abstract: How can we design a market of human-generated content for use in training AI models that both enables technological progress and preserves individual incentives for high-quality content creation? Existing approaches take polar...

<https://arxiv.org/abs/2606.12260>

09 CONF

SpikeDecoder: Realizing the GPT Architecture with Spiking Neural Networks

ArXiv CS.AI 50%

arXiv:2606.12287v1 Announce Type: cross Abstract: The Transformer architecture is widely regarded as the most powerful tool for natural language processing, but due to a high number of complex operations, it inherently faces the issue of high energy consumption. To address this...

<https://arxiv.org/abs/2606.12287>

10 CONF

The Standard Interpretable Model: A general theory of interpretable machine learning to deductively design interpretable methods using Lagrangian mechanics

ArXiv CS.AI 50%

arXiv:2606.12289v1 Announce Type: cross Abstract: As Artificial Intelligence models grow in complexity, interpretability has become an indispensable tool for understanding, debugging, and controlling their computations. However, interpretability lacks general theories to...

<https://arxiv.org/abs/2606.12289>

11 CONF

Natural-Language Temporal Grounding in Hour-Long Videos is a Search Problem: A Benchmark and Empirical Decomposition

ArXiv CS.AI 50%

arXiv:2606.12300v1 Announce Type: cross Abstract: Temporal grounding--returning the interval $[t_s, t_e]$ for a natural-language query over a video--is the language interface to long-form video, yet has been studied on short videos; the dynamics of hour-scale natural-language...

<https://arxiv.org/abs/2606.12300>

12 CONF

Harness In-Context Operator Learning with Chain of Operators

ArXiv CS.AI 50%

arXiv:2606.12318v1 Announce Type: cross Abstract: Neural operators approximate mappings between function spaces, but often generalize poorly to other operators and usually require fine-tuning or retraining. In-Context Operator Networks (ICON) addresses this issue by prompting...

<https://arxiv.org/abs/2606.12318>

13 CONF

Atlas H&E-TME: Scalable AI-Based Tissue Profiling at Expert Pathologist-Level Accuracy

ArXiv CS.AI 50%

arXiv:2606.12346v1 Announce Type: cross Abstract: Hematoxylin and eosin (H&E) staining is the cornerstone of histopathology, yet scalable, quantitative analysis of H&E whole-slide images (WSIs) remains a central challenge in computational pathology. We present Atlas H&E-TME, an...

<https://arxiv.org/abs/2606.12346>

14 CONF

Latent World Recovery for Multimodal Learning with Missing Modalities

ArXiv CS.AI 50%

arXiv:2606.12362v1 Announce Type: cross Abstract: We study multimodal learning under missing modalities, with particular motivation from bioscience applications in which heterogeneous modalities are often only partially available when decisions need to be made. We propose Latent...

<https://arxiv.org/abs/2606.12362>

15 CONF

Ambient Diffusion Policy: Imitation Learning from Suboptimal Data in Robotics

ArXiv CS.AI 50%

arXiv:2606.12365v1 Announce Type: cross Abstract: We propose Ambient Diffusion Policy, a simple and principled method for imitation learning from suboptimal data in robotics. High-quality, task-specific robot data is expensive and time-consuming to collect, while suboptimal...

<https://arxiv.org/abs/2606.12365>

16 CONF

System Report for CCL25-Eval Task 5: New Dataset and LoRA-Fine-Tuned Qwen2.5

ArXiv CS.AI 50%

arXiv:2606.12392v1 Announce Type: cross Abstract: Recently, large language models (LLMs) have achieved promising progress in the fields of classical Chinese translation and the generation of classical poetry. However, domain-specific research on precise translation and...

<https://arxiv.org/abs/2606.12392>

17 CONF

Diffusing to Coordinate: Efficient Online Multi-Agent Diffusion Policies

ArXiv CS.AI 50%

arXiv:2602.18291v2 Announce Type: replace Abstract: Online Multi-Agent Reinforcement Learning (MARL) is a prominent framework for efficient agent coordination. Crucially, enhancing policy expressiveness is pivotal for achieving superior performance. Diffusion-based generative...

<https://arxiv.org/abs/2602.18291>

18 CONF

Offline Diffusion Policy for Multi-User Delay-Constrained Scheduling

ArXiv CS.AI 50%

arXiv:2501.12942v2 Announce Type: replace Abstract: Effective multi-user delay-constrained scheduling is crucial in various real-world applications, including embodied AI, instant messaging, live streaming, and data center management, where efficient resource allocation is...

<https://arxiv.org/abs/2501.12942>

19 CONF

MLaGA: Multimodal Large Language and Graph Assistant

ArXiv CS.AI 50%

arXiv:2506.02568v2 Announce Type: replace Abstract: Large Language Models (LLMs) have demonstrated substantial efficacy in advancing graph-structured data analysis. Prevailing LLM-based graph methods excel in adapting LLMs to text-rich graphs, wherein node attributes are text...

<https://arxiv.org/abs/2506.02568>

20 CONF

Sustainability assessment using multimodal AI agents

ArXiv CS.AI 50%

arXiv:2507.17012v2 Announce Type: replace Abstract: Reducing the rapidly growing environmental impact of the computing industry requires assessing the emissions of electronics at scale. However, a traditional life cycle assessment (LCA) of an electronic device, which maps...

<https://arxiv.org/abs/2507.17012>

21 CONF

GPO: Learning from Critical Steps to Improve LLM Reasoning

ArXiv CS.AI 50%

arXiv:2509.16456v3 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used in various domains, showing impressive potential on different tasks. Recently, reasoning LLMs have been proposed to improve the \textit{reasoning} or \textit{thinking}...

<https://arxiv.org/abs/2509.16456>

22 CONF

Resource-Aware LLM Reasoning for Mobile Edge General Intelligence

ArXiv CS.AI 50%

arXiv:2509.23248v3 Announce Type: replace Abstract: The rapid advancement of large language models (LLMs) has enabled an emergence of agentic artificial intelligence (AI) with powerful reasoning and autonomous decision-making capabilities. This integration with edge computing...

<https://arxiv.org/abs/2509.23248>

23 CONF

A New Perspective on Precision and Recall for Generative Models

ArXiv CS.AI 50%

arXiv:2511.02414v3 Announce Type: replace Abstract: With the recent success of generative models in image and text, the question of their evaluation has recently gained a lot of attention. While most methods from the state of the art rely on scalar metrics, the introduction of...

<https://arxiv.org/abs/2511.02414>

24 CONF

An XAI View on Explainable ASP: Methods, Systems, and Perspectives

ArXiv CS.AI 50%

arXiv:2601.14764v2 Announce Type: replace Abstract: Answer Set Programming (ASP) is a popular declarative reasoning and problem solving approach in symbolic AI. Its rule-based formalism makes it inherently attractive for explainable and interpretive reasoning, which is gaining...

<https://arxiv.org/abs/2601.14764>

25 CONF

Human-Guided Agentic AI for Multimodal Clinical Prediction: Lessons from the AgentDS Healthcare Benchmark

ArXiv CS.AI 50%

arXiv:2602.19502v2 Announce Type: replace Abstract: Agentic AI systems are increasingly capable of autonomous data science workflows, yet clinical prediction tasks demand domain expertise that purely automated approaches struggle to provide. We investigate how human guidance of...

<https://arxiv.org/abs/2602.19502>

26 CONF

MemToolAgent: Leveraging Memory for Tool Using Agents Based on Environment and User Feedback

ArXiv CS.AI 50%

arXiv:2606.07909v2 Announce Type: replace Abstract: Modern large language model (LLM) agents can use external tools to help users solve complex tasks. However, for problems that require learning from long-term historical events or from previous agent-environment interactions,...

<https://arxiv.org/abs/2606.07909>

27 CONF

The Algorithm Is Not the Behavior: Learned Priors Override Look-Ahead in a Chess-Playing Neural Network

ArXiv CS.AI 50%

arXiv:2508.21380v3 Announce Type: replace-cross Abstract: Recent mechanistic work has uncovered learned algorithms within neural networks, from modular arithmetic to search and planning in game-playing agents. But does algorithmic structure guarantee algorithmic behavior? We...

<https://arxiv.org/abs/2508.21380>

28 CONF

Geometric Metrics and LLMs: What They Measure and When They Work

ArXiv CS.AI 50%

arXiv:2509.25359v2 Announce Type: replace-cross Abstract: We present a systematic stress-test of geometric metrics for LLM evaluation. Rank-based geometric properties of internal representations have shown promise as reference-free quality signals, but the conditions under which...

<https://arxiv.org/abs/2509.25359>

29 CONF

When Researchers Say Mental Model/Theory of Mind of AI, What Are They Really Talking About?

ArXiv CS.AI 50%

arXiv:2510.02660v2 Announce Type: replace-cross Abstract: When researchers claim AI systems possess ToM or mental models, they are fundamentally discussing behavioral predictions and bias corrections rather than genuine mental states. This position paper argues that the current...

<https://arxiv.org/abs/2510.02660>

30 CONF

Certifiable Safe RLHF: Semantic Grounding and Fixed Penalty Constraint Optimization for Safer LLM Alignment

ArXiv CS.AI 50%

arXiv:2510.03520v2 Announce Type: replace-cross Abstract: Ensuring safety is a foundational requirement for large language models (LLMs). Achieving an appropriate balance between enhancing the utility of model outputs and mitigating their potential for harm is a complex and...

<https://arxiv.org/abs/2510.03520>

31 CONF

Moving Beyond Diffusion: Hierarchy-to-Hierarchy Autoregression for fMRI-to-Image Reconstruction

ArXiv CS.AI 50%

arXiv:2510.22335v2 Announce Type: replace-cross Abstract: Reconstructing visual stimuli from fMRI signals is a central challenge bridging machine learning and neuroscience. Recent diffusion-based methods typically map fMRI activity to a single neural embedding, using it as...

<https://arxiv.org/abs/2510.22335>

32 CONF

Irresponsible AI: big tech's influence on AI research and associated impacts

ArXiv CS.AI 50%

arXiv:2512.03077v2 Announce Type: replace-cross Abstract: The accelerated development, deployment and adoption of artificial intelligence systems has been fuelled by the increasing presence of big tech in the AI field. This trend has been accompanied by growing ethical concerns...

<https://arxiv.org/abs/2512.03077>

33 CONF

Towards Deep Learning Surrogate for the Forward Problem in Electrocardiology: A Scalable Alternative to Physics-Based Models

ArXiv CS.AI 50%

arXiv:2512.13765v2 Announce Type: replace-cross Abstract: The forward problem in electrocardiology, computing body surface potentials from cardiac electrical activity, is traditionally solved using physics-based models such as the bidomain or monodomain equations. While...

<https://arxiv.org/abs/2512.13765>

34 CONF

Unifying Learning Dynamics and Generalization in Transformers Scaling Law

ArXiv CS.AI 50%

arXiv:2512.22088v3 Announce Type: replace-cross Abstract: The scaling law, a cornerstone of Large Language Model (LLM) development, predicts improvements in model performance with increasing computational resources. Yet, while empirically validated, its theoretical underpinnings...

<https://arxiv.org/abs/2512.22088>

35 CONF

Robust Privacy: Inference-Stage Privacy through Certified Robustness

ArXiv CS.AI 50%

arXiv:2601.17360v2 Announce Type: replace-cross Abstract: An adversary observing a model's released prediction can infer sensitive attributes of the queried input, or even reconstruct representatives of the model's training data. The inference interface thus acts as a side...

<https://arxiv.org/abs/2601.17360>

36 CONF

Reliability-Calibrated Edge-IoT Early Fault Warning for Rotating Machinery with a Physics-Guided Tiny-Mamba Transformer

ArXiv CS.AI 50%

arXiv:2601.21293v3 Announce Type: replace-cross Abstract: Industrial Internet of Things (IIoT) systems increasingly rely on distributed vibration sensing to support predictive maintenance of rotating machinery. In practical deployments, however, raw signal upload is costly and...

<https://arxiv.org/abs/2601.21293>

37 CONF

When Generic Prompt Improvements Hurt: Evaluation-Driven Iteration for LLM Applications

ArXiv CS.AI 50%

arXiv:2601.22025v2 Announce Type: replace-cross Abstract: Evaluating Large Language Model (LLM) applications differs from conventional software testing because outputs are probabilistic, semantically variable, and sensitive to prompt and model changes. This technical report...

<https://arxiv.org/abs/2601.22025>

38 CONF

OpenVTON-Bench: A Large-Scale High-Resolution Benchmark for Controllable Virtual Try-On Evaluation

ArXiv CS.AI 50%

arXiv:2601.22725v4 Announce Type: replace-cross Abstract: Recent advances in diffusion models have significantly elevated the visual fidelity of Virtual Try-On (VTON) systems, yet reliable evaluation remains a persistent bottleneck. Traditional metrics struggle to quantify...

<https://arxiv.org/abs/2601.22725>

39 CONF

Neural FOXP2 -- Language Specific Neuron Steering for Targeted Language Improvement in LLMs

ArXiv CS.AI 50%

arXiv:2602.00945v2 Announce Type: replace-cross Abstract: LLMs are multilingual by training, yet their lingua franca is often English, reflecting English language dominance in pretraining. Other languages remain in parametric memory but are systematically suppressed. We argue...

<https://arxiv.org/abs/2602.00945>

40 CONF

Compiler-First State Space Duality and Portable $\mathcal{O}(1)$ Autoregressive Caching for Inference

ArXiv CS.AI 50%

arXiv:2603.09555v2 Announce Type: replace-cross Abstract: High-throughput Mamba-2 inference is usually tied to fused CUDA and Triton kernels, limiting portability across accelerator backends. We show that the state space duality (SSD) recurrence has a compiler-friendly...

<https://arxiv.org/abs/2603.09555>

41 CONF

Litespark Inference For CPUs: Ultra-Fast SIMD Framework for Ternary (1.58-bit) Language Models

ArXiv CS.AI 50%

arXiv:2605.06485v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have transformed artificial intelligence, but their computational requirements remain prohibitive for most users. Standard inference demands expensive datacenter GPUs or cloud API access,...

<https://arxiv.org/abs/2605.06485>

42 CONF

Anomalies in Multivariate Time Series Benchmarks Are Mostly Univariate

ArXiv CS.AI 50%

arXiv:2606.02670v3 Announce Type: replace-cross Abstract: Many recent multivariate time series anomaly detection (MTSAD) models incorporate cross-channel modeling, under the implicit assumption that the structure of anomalies may be spread across multiple channels. We evaluate...

<https://arxiv.org/abs/2606.02670>

43 CONF

Agentic Software: How AI Agents Are Restructuring the Software Paradigm

ArXiv CS.AI 50%

arXiv:2606.05608v2 Announce Type: replace-cross Abstract: For over half a century, software engineering has operated on a foundational premise: human engineers decompose problems, encode decision logic into static code, and manually adapt that code as requirements evolve. This...

<https://arxiv.org/abs/2606.05608>

44 **CONF****MetaPlate: Counterfactual-Guided RAG-LLM Tool for Personalized Food Recommendation and Hyperglycemia Prevention****ArXiv CS.AI** **50%**

arXiv:2606.10120v2 Announce Type: replace-cross Abstract: Postprandial hyperglycemia is a key risk factor for metabolic disorders; however, existing dietary guidance is often static, impractical, and insufficiently personalized, providing recommendations that are difficult to...

<https://arxiv.org/abs/2606.10120>