



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	120	Media
TechCrunch AI	6	Media
HackerNews	4	Community
MIT Tech Review	2	Media
VentureBeat AI	1	Media
The Verge AI	1	Media

VERIFIED INTELLIGENCE — 1 story

VERI

The Containment Gap: How Deployed Agentic AI Frameworks Fail Public-Facing Safety Requirements

ArXiv CS.AI +1 source

80%

arXiv:2606.12797v1 Announce Type: new Abstract: Agentic large language model systems that autonomously invoke tools, maintain persistent memory, and execute multi-step plans are increasingly deployed in public-facing domains, including government services, healthcare triage, and...

<https://arxiv.org/abs/2606.12797>

CONFIRMED SIGNALS — 132 stories

CONF

Show HN: HelixDB - A graph database built on object storage

HackerNews 60% HN 148

Hey HN, it's been just over a year since we launched HelixDB (<https://news.ycombinator.com/item?id=43975423>), a project a friend and I started in college. It's an OLTP graph database built on object-storage, with native vector search and full-text search (FTS). Why graph,...

<https://github.com/HelixDB/helix-db/tree/main>

CONF

Show HN: TunnelMind - reputation API for IPs, ASNs, and ad-tech supply chains

HackerNews 50% HN 9

I'm a network engineer that likes to think about the future of the internet and this is what I've built over many nights and weekends. One reputation graph over IPs, ASNs, domains, and entities, exposed as a JSON API. Try it: [curl https://api.tunnelmind.ai/v1/check/1.1.1.1](https://api.tunnelmind.ai/v1/check/1.1.1.1) Every...

<https://tunnelmind.ai/>

CONF

Show HN: Claumon - forecasting Claude Code usage limits with a Gamma process

HackerNews 50% HN 6

Hey HN, I built this a side project because I'm a soccer fan that has been vibing and tokenmaxxing with Claude Code maybe too much. So, the World Cup is here and it was the perfect excuse to build and ship something from 0 to 1. Enter Claudinho, a CLI and MCP that puts World Cup...

<https://github.com/fabioconcina/claumon>

CONF

Show HN: Brooks-Lint - AI code reviews grounded in 12 classic engineering books

HackerNews 50% HN 6

I once managed a small team and hired an intern to develop the backend services for a software system. He used AI tools to write the entire project; while it ran correctly and produced the right results, the directory structure was atrocious upon review. Typically, a FastAPI...

<https://github.com/hyhmrright/brooks-lint>

CONF

Cheaper, faster, and culturally aware, Avataar's video AI is built for India's scale

TechCrunch AI 50%

Avataar AI's distilled video model is priced at \$0.005 for every second of generation

<https://techcrunch.com/2026/06/11/cheaper-faster-and-culturally-aware-avataars-video-ai...>

CONF

Theker just raised \$85M to build the factory robot that doesn't specialize in anything

TechCrunch AI 50%

Unlike humanoid robots designed around a fixed form — think Boston Dynamics — Theker's machines are built to be reconfigured.

<https://techcrunch.com/2026/06/11/theker-just-raised-85m-to-build-the-factory-robot-tha...>

CONF

Jeff Bezos's Prometheus raises \$12B to build an 'artificial general engineer' for the physical world

TechCrunch AI 50%

The new round values the physical AI startup that aims to automate heavy engineering and drug design at \$41 billion.

<https://techcrunch.com/2026/06/11/jeff-bezoss-prometheus-raises-12b-to-build-an-artific...>

CONF

Deezer's new tool can identify AI music from Spotify, Apple Music, and others

TechCrunch AI 50%

Deezer introduced a tool that scans playlists from Spotify, Apple Music, and other platforms to identify AI music.

<https://techcrunch.com/2026/06/11/deezers-new-tool-can-identify-ai-music-from-spotify-a...>

0 CONF

DoorDash's new AI chatbot lets you order with prompts and photos

TechCrunch AI 50%

The new chatbot, called Ask DoorDash, allows users to search the app for what they're looking for in their own words instead of having to scroll through restaurants and stores to build a cart.

<https://techcrunch.com/2026/06/11/doordashs-new-ai-chatbot-lets-you-order-with-prompts-...>

1 CONF

Jedify raises \$24M to help companies arm AI agents with context on their business

TechCrunch AI 50%

The funding round was led by Norwest, with participation from S Capital VC, Cerca Partners, and Oceans Ventures. Snowflake Ventures also participated as a strategic investor.

<https://techcrunch.com/2026/06/10/jedify-raises-24m-to-help-companies-arm-ai-agents-wit...>

2 CONF

The Download: soccer's data renaissance and China's big nuclear plans

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Inside soccer's data renaissance Imagine tuning in to the opening kickoff of a World Cup match and seeing a player intentionally kick the...

<https://www.technologyreview.com/2026/06/11/1138809/the-download-soccer-football-data-a...>

3 CONF

Google DeepMind is worried about what happens when millions of agents start to interact

MIT Tech Review 50%

Google DeepMind is funding research into the potential dangers of situations where millions of different AI agents interact with each other online. According to Rohin Shah, who directs the company's AGI safety and alignment research, the mass-market arrival of agents that can...

<https://www.technologyreview.com/2026/06/11/1138794/google-deepmind-is-worried-about-wh...>

4 CONF

Anthropic apologizes for invisible Claude Fable guardrails

The Verge AI 50%

Anthropic has apologized for stealthily throttling its new AI model, Claude Fable 5, with hidden guardrails that undermine both researchers and rivals using it to develop competing systems. The company says it is reversing course and will be more transparent about when the...

<https://www.theverge.com/ai-artificial-intelligence/948280/anthropic-claude-fable-invis...>

5 CONF

Definitional alignment before capability alignment: a Design-Science framework for adjudicating claims about AGI

ArXiv CS.AI 50%

arXiv:2606.12713v1 Announce Type: new Abstract: Claims that artificial general intelligence has already arrived and claims that it remains decades away are often defended from overlapping evidence. "AGI" lacks a single shared and stable referent and competing operationalizations...

<https://arxiv.org/abs/2606.12713>

6 CONF

Arbor: Tree Search as a Cognition Layer for Autonomous Agents

ArXiv CS.AI 50%

arXiv:2606.12563v1 Announce Type: new Abstract: Arbor is a multi-agent framework that introduces structured tree search as a cognition layer for autonomous agents operating in large, stateful action spaces. Prior autonomous optimization systems operate on isolated targets with...

<https://arxiv.org/abs/2606.12563>

7 CONF

Strategic Decision Support for AI Agents

ArXiv CS.AI 50%

arXiv:2606.12587v1 Announce Type: new Abstract: Traditionally, decision support studies how humans use machine learning models to make better decisions. In modern agentic systems, this division of roles is increasingly reversed: AI agents act on behalf of users, while humans and...

<https://arxiv.org/abs/2606.12587>

8 CONF

PhononBench:A Large-Scale Phonon-Based Benchmark for Dynamical Stability in Crystal Generation

ArXiv CS.AI 50%

arXiv:2512.21227v3 Announce Type: replace-cross Abstract: In recent years, generative artificial intelligence has made significant advances in the design of crystalline materials, giving rise to approaches based on graph neural networks, diffusion models, and large language...

<https://arxiv.org/abs/2512.21227>

9 CONF

Evoflux: Inference-Time Evolution of Executable Tool Workflows for Compact Agents

ArXiv CS.AI 50%

arXiv:2606.12674v1 Announce Type: new Abstract: Compact language models (LMs) reduce cost, latency, and deployment risk for tool agents. Yet MCP-style tool use requires more than isolated function calling: an agent must discover tools from live catalogs, satisfy schemas,...

<https://arxiv.org/abs/2606.12674>

0 CONF

Deployment-Centered Evaluation: Predicting Query-Level Rejection Risk in a Clinical LLM System

ArXiv CS.AI 50%

arXiv:2606.12702v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly integrated into clinical systems, making it essential to evaluate the real-world utility of these systems. However, static benchmarks tend to measure correctness rather than user...

<https://arxiv.org/abs/2606.12702>

1 **CONF****Rethinking Psychometric Evaluation of LLMs: When and Why Self-Reports Predict Behavior****ArXiv CS.AI** **50%**

arXiv:2606.12730v1 Announce Type: new Abstract: Anticipating LLM behavioral tendencies from low-cost psychometric probes is critical for safe deployment, but only if self-reports (SR) reliably predict behavior. Recent work documented substantial SR-behavior dissociation in LLMs,...

<https://arxiv.org/abs/2606.12730>

2 **CONF****SciR: A Controllable Benchmark for Scientific Reasoning in LLMs****ArXiv CS.AI** **50%**

arXiv:2606.13020v1 Announce Type: new Abstract: Three paradigmatic forms of inference recur across scientific reasoning: deduction, induction, and causal abduction. Reliably evaluating LLMs on these in scientific settings is currently out of reach: scientific benchmarks built on...

<https://arxiv.org/abs/2606.13020>

3 **CONF****Agentic Large Language Models for Automated Structural Analysis of 3D Frame Systems****ArXiv CS.AI** **50%**

arXiv:2606.06525v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have emerged as powerful foundation models with strong reasoning capabilities across domains. Beyond reactive text generation, agentic LLMs enable autonomous workflow execution through modular...

<https://arxiv.org/abs/2606.06525>

4 **CONF****The Emergence of Autonomous Penetration Capabilities in Large Language Model-Powered AI Systems****ArXiv CS.AI** **50%**

arXiv:2606.13079v1 Announce Type: cross Abstract: Nowadays, the autonomous execution of cyberattacks capable of causing substantial real-world harm is widely regarded as one of the critical red lines that frontier AI systems must not cross. Within this broader red-line scenario,...

<https://arxiv.org/abs/2606.13079>

5 **CONF****Constructing Evaluation Datasets for Procedural Reasoning: Balancing Naturalness, Grounding, and Multi-Hop Coverage****ArXiv CS.AI** **50%**

arXiv:2606.12767v1 Announce Type: new Abstract: Evaluating procedural reasoning in AI-supported learning systems requires question-answer datasets that are both learner-like and grounded in the instructional knowledge the system is expected to use. We study how TMK-based...

<https://arxiv.org/abs/2606.12767>

6 **CONF****A Tutorial on World Models and Physical AI****ArXiv CS.AI** **50%**

arXiv:2606.12783v1 Announce Type: new Abstract: World modeling is emerging as a central principle for building intelligent systems capable of prediction, reasoning, and decision making. A central distinction can be drawn between explicit world models, which learn structured...

<https://arxiv.org/abs/2606.12783>

7 CONF

SupraBench: A Benchmark for Supramolecular Chemistry

ArXiv CS.AI 50%

arXiv:2606.13477v1 Announce Type: cross Abstract: Supramolecular chemistry, which includes the study of non-covalent host-guest assemblies, has advanced various applications. However, designing host-guest systems remains time-consuming, requiring days of dry-lab verification per...

<https://arxiv.org/abs/2606.13477>

8 CONF

Emerging Flexible Designs for Geospatial Multimodal Foundation Models

ArXiv CS.AI 50%

arXiv:2606.12595v1 Announce Type: cross Abstract: Foundation models are rapidly transforming Earth observation by enabling scalable pretraining across diverse unlabeled geospatial modalities. However, their architectural diversity ranging from encoder-only to encoder-decoder and...

<https://arxiv.org/abs/2606.12595>

9 CONF

Topical Phase Transitions in Artificial Intelligence Research: Large-Scale Evidence and an Early-Warning Signature for Emerging Topics

ArXiv CS.AI 50%

arXiv:2606.12828v1 Announce Type: new Abstract: Do research topics in artificial intelligence grow gradually, or do they advance through abrupt, detectable jumps? Analyzing 80,814 accepted main-track papers from five premier AI conferences (ACL, CVPR, ICLR, ICML, NeurIPS)...

<https://arxiv.org/abs/2606.12828>

0 CONF

(Human) Attention Is (Still) All You Need: Human oversight makes AI-assisted social science reliable

ArXiv CS.AI 50%

arXiv:2606.12848v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used for tasks once reserved for trained researchers, including hypothesis generation, specification choice, and drafting conclusions. We argue that the reliability of AI-assisted...

<https://arxiv.org/abs/2606.12848>

1 CONF

DailyReport: An Open-ended Benchmark for Evaluating Search Agents on Daily Search Tasks

ArXiv CS.AI 50%

arXiv:2606.12871v1 Announce Type: new Abstract: Search Agents (SAs) typically leverage large language models (LLMs) to support complex information-seeking tasks by autonomously exploring web sources and synthesizing information into comprehensive responses. For SAs evaluation,...

<https://arxiv.org/abs/2606.12871>

2 CONF

Zero-source LLM Hallucination Detection with Human-like Criteria Probing

ArXiv CS.AI 50%

arXiv:2606.12900v1 Announce Type: new Abstract: Large language models (LLMs) often hallucinate by generating factually incorrect or unfaithful content, posing significant risks to their safe use. Detecting such hallucinations is particularly challenging under the zero-source...

<https://arxiv.org/abs/2606.12900>

3 CONF

MARS: Margin-Adversarial Risk-controlled Stopping for Parallel LLM Test-time Scaling

ArXiv CS.AI 50%

arXiv:2606.12935v1 Announce Type: new Abstract: Parallel test-time scaling samples many reasoning traces and majority-votes their answers, improving LLM accuracy but requiring traces to run to completion, incurring substantial computational overhead. We observe that probing...

<https://arxiv.org/abs/2606.12935>

4 CONF

PRISMR: Overcoming Parse Collapse in Multimodal Listwise Ranking via Parameterized Representation Internalization

ArXiv CS.AI 50%

arXiv:2606.12942v1 Announce Type: new Abstract: Generative listwise ranking with Large Multimodal Models (LMMs) aims to capture global list context in a single forward pass, but its effectiveness degrades in long-context multimodal scenarios. We identify a recurring failure...

<https://arxiv.org/abs/2606.12942>

5 CONF

The Safety-Aware Denoiser for Text Diffusion Models

ArXiv CS.AI 50%

arXiv:2605.08116v2 Announce Type: replace-cross Abstract: Recent work on text diffusion models offers a promising alternative to autoregressive generation, but controlling their safety remains underexplored. Existing safety approaches are geared toward autoregressive models and...

<https://arxiv.org/abs/2605.08116>

6 CONF

A Mathematical Forum Platform for Collaborative Problem Solving and Dataset Generation for AI Reasoning

ArXiv CS.AI 50%

arXiv:2606.12976v1 Announce Type: new Abstract: Sharing mathematical content in online forums remains a significant friction point for students and educators: writing raw LATEX is error-prone, standalone optical character recognition tools require platform switching, and current...

<https://arxiv.org/abs/2606.12976>

7 CONF

Structured Testbench Generation for LLM-Driven HDL Design and Verification-Oriented Data Curation

ArXiv CS.AI 50%

arXiv:2606.12983v1 Announce Type: new Abstract: Automated testbench generation has become a critical bottleneck in large language model (LLM)-driven Register Transfer Level (RTL) workflows, where large numbers of candidate designs must be verified rapidly and reliably. Existing...

<https://arxiv.org/abs/2606.12983>

8 CONF

Otters++: A Time-to-first-spike Based Energy Efficient Optical Spiking Transformer

ArXiv CS.AI 50%

arXiv:2606.13016v1 Announce Type: new Abstract: Spiking neural networks (SNNs) are promising for energy-efficient inference, and time-to-first-spike (TTFS) coding is especially attractive because each neuron fires at most once. In practice, however, this benefit is often reduced...

<https://arxiv.org/abs/2606.13016>

9 CONF

Mental-R1: Aligning LLM Reasoning for Mental Health Assessment

ArXiv CS.AI 50%

arXiv:2606.13176v1 Announce Type: new Abstract: Mental health problems such as anxiety, depression, and suicide remain urgent global challenges, where timely and accurate assessment is critical for effective intervention. Recently, large language models have been explored for...

<https://arxiv.org/abs/2606.13176>

0 CONF

Reasoning for Mobile User Experience with Multimodal LLMs: Task, Benchmark, and Approach

ArXiv CS.AI 50%

arXiv:2606.13192v1 Announce Type: new Abstract: User experience (UX) centered on usability, perceived consistency, and functional clarity is fundamental to real-world user interfaces (UI). The application of multimodal large language models (MLLMs) in the field of user...

<https://arxiv.org/abs/2606.13192>

1 CONF

ARMOR-MAD: Adaptive Routing for Heterogeneous Multi-Agent Debate in Large Language Model Reasoning

ArXiv CS.AI 50%

arXiv:2606.13197v1 Announce Type: new Abstract: Multi-agent debate (MAD) can improve large language model reasoning, but fixed debate pipelines often waste computation and can amplify correlated errors among similar agents. We propose ARMOR-MAD, a training-free heterogeneous MAD...

<https://arxiv.org/abs/2606.13197>

2 CONF

Hallucination in Medical Imaging AI: A Cross-Modality Analytical Framework for Taxonomy, Detection, and Mitigation under Regulatory Constraints

ArXiv CS.AI 50%

arXiv:2606.13211v1 Announce Type: new Abstract: AI systems are being deployed across medical imaging faster than their failure modes are understood. At this point in time, the failure of greatest clinical concern is hallucination: clinically plausible but factually incorrect...

<https://arxiv.org/abs/2606.13211>

3 CONF

LLM-as-an-Investigator: Evidence-First Reasoning for Robust Interactive Problem Diagnosis

ArXiv CS.AI 50%

arXiv:2606.13220v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used as interactive assistants for technical problem solving. However, when users provide incomplete descriptions or plausible but unverified explanations, LLMs may prematurely align...

<https://arxiv.org/abs/2606.13220>

4 CONF

MOSAIC: Modality-Specific Adaptation for Incremental Continual Learning in Parkinson's Disease Gait Assessment

ArXiv CS.AI 50%

arXiv:2606.13258v1 Announce Type: new Abstract: Gait-based Parkinson's disease assessment increasingly relies on heterogeneous sensors, but clinical systems rarely collect all modalities simultaneously. New sensors may arrive through device upgrades, protocol changes, or...

<https://arxiv.org/abs/2606.13258>

5 CONF

From Verdict to Process: Agentic Reinforcement Learning for Multi-Stage Fact Verification

ArXiv CS.AI 50%

arXiv:2606.13262v1 Announce Type: new Abstract: Recent approaches combining Large Language Models (LLMs) with retrieval-augmented reasoning have shown promise for automated fact verification. To process complex claims, these verification pipelines typically execute multi-stage...

<https://arxiv.org/abs/2606.13262>

6 CONF

ERTS: Adversarial Robustness Testing of Ethical AI via Semantic Perturbation in a Bounded Consequence Space

ArXiv CS.AI 50%

arXiv:2606.13282v1 Announce Type: new Abstract: As AI systems are deployed in high-stakes ethical contexts such as healthcare triage, autonomous vehicle control, and employment screening, formal methods for evaluating their robustness against adversarial manipulation of ethical...

<https://arxiv.org/abs/2606.13282>

7 CONF

A Quantitative Experimental Repeated Measures Study of Training Dynamics in a Small Llama Style Language Model Under a Compute-Aware Token Budget

ArXiv CS.AI 50%

arXiv:2606.13370v1 Announce Type: new Abstract: This study examines training dynamics in a small Llama-style language model trained under a fixed, compute-constrained token budget. Rather than evaluating efficiency solely through endpoint performance, the study uses a...

<https://arxiv.org/abs/2606.13370>

8 CONF

MoReBench: Evaluating Procedural and Pluralistic Moral Reasoning in Language Models, More than Outcomes

ArXiv CS.AI 50%

arXiv:2510.16380v2 Announce Type: replace-cross Abstract: As AI systems progress, we rely more on them to make decisions with us and for us. To ensure that such decisions are aligned with human values, it is imperative for us to understand not only what decisions they make but...

<https://arxiv.org/abs/2510.16380>

9 CONF

CloudCons: A Comprehensive End-to-End Benchmark for Cloud Resource Consolidation

ArXiv CS.AI 50%

arXiv:2606.13513v1 Announce Type: new Abstract: Driven by conservative over-provisioning to guarantee service reliability, resource utilization in cloud data centers remains at low levels. To mitigate this, the forecast-then-optimize paradigm has emerged to optimize...

<https://arxiv.org/abs/2606.13513>

0 CONF

Uncertainty-Aware Hybrid Retrieval for Long-Document RAG

ArXiv CS.AI 50%

arXiv:2606.13550v1 Announce Type: new Abstract: Retrieval augmented generation (RAG) depends critically on the quality and granularity of retrieved evidence. Large retrieval units preserve context but often introduce irrelevant content, which can dilute answer bearing evidence...

<https://arxiv.org/abs/2606.13550>

1 **CONF****Is It You or Your Environment? A Bayesian Inference Framework for Genomically-Anchored Personalized Physiological Interpretation**

ArXiv CS.AI 50%

arXiv:2606.13556v1 Announce Type: new Abstract: Personalized health AI systems face a fundamental cold-start problem: machine learning models for physiological interpretation require weeks of individual behavioral data before they can distinguish constitutional variation from...

<https://arxiv.org/abs/2606.13556>

2 **CONF****A Three-Layer Framework for AI in Scientific Discovery**

ArXiv CS.AI 50%

arXiv:2606.13566v1 Announce Type: new Abstract: Current discussions of AI in scientific discovery are often dominated by two visible capabilities: search over existing knowledge and execution through optimization, simulation, and automation. Both are important, but neither fully...

<https://arxiv.org/abs/2606.13566>

3 **CONF****Reasoning as Pattern Matching: Shared Mechanisms in Human and LLM Everyday Reasoning**

ArXiv CS.AI 50%

arXiv:2606.13607v1 Announce Type: new Abstract: When large language models (LLMs) fail to generalize or make haphazard errors in reasoning, it is often taken as evidence that LLMs are not truly reasoning, but rather performing a kind of pattern matching. The implication is that...

<https://arxiv.org/abs/2606.13607>

4 **CONF****Before You Think: System 0, AI-Mediated Cognition and Cognitive Colonization**

ArXiv CS.AI 50%

arXiv:2606.13658v1 Announce Type: new Abstract: This paper examines three recent frameworks for understanding the cognitive and epistemic consequences of artificial intelligence: Tri-System Theory, Thinkframes, and System 0. It argues that while the first two capture important...

<https://arxiv.org/abs/2606.13658>

5 **CONF****The AI Legal Specialist: A Juridically Autonomous Professional Profile for AI Governance**

ArXiv CS.AI 50%

arXiv:2606.12415v1 Announce Type: cross Abstract: The rapid global expansion of artificial intelligence regulation has generated, across multiple jurisdictions, a demand for legal expertise dedicated to AI that the market has addressed in a fragmented manner. Data protection...

<https://arxiv.org/abs/2606.12415>

6 **CONF****Divination by Prompt: LLM-Mediated Xuanxue on Chinese Social Media**

ArXiv CS.AI 50%

arXiv:2606.12418v1 Announce Type: cross Abstract: The rapid proliferation of large language models (LLMs) has produced a striking cultural practice: using conversational AI for divination. This paper offers one of the first systematic studies of LLM-mediated divination in the...

<https://arxiv.org/abs/2606.12418>

- 7 **CONF**
GeoDial: A Multimodal Conversational Tutoring Dataset for Geometry Problem-Solving with Visual Tutor Turns
 ArXiv CS.AI 50%
 arXiv:2606.12419v1 Announce Type: cross Abstract: Several educational domains rely heavily on diagrams and visual cues, yet most existing tutoring datasets are limited to text-only interactions. This limits the development of AI tutors that can teach in visually grounded ways...
<https://arxiv.org/abs/2606.12419>
- 8 **CONF**
Creating and Evaluating K-12 GenAI Assessment Graders Through Context Engineering
 ArXiv CS.AI 50%
 arXiv:2606.12422v1 Announce Type: cross Abstract: The integration of large language models (LLMs) into educational assessment represents a transformative shift in classroom grading practices. While automated scoring systems and machine learning techniques have existed for...
<https://arxiv.org/abs/2606.12422>
- 9 **CONF**
The Challenges of Balancing AI Compliance and Technological Innovations in Critical Sectors: A Systematic Literature Review
 ArXiv CS.AI 50%
 arXiv:2606.12423v1 Announce Type: cross Abstract: The rapid integration of artificial intelligence (AI) into critical infrastructure including healthcare, finance, energy, and defense, offers transformative benefits but also conflicts with evolving regulatory and governance...
<https://arxiv.org/abs/2606.12423>
- 0 **CONF**
AI-Automation Tooling in Computer Engineering Education: Mixed-Methods TAM/UTAUT Evidence for a General Acceptance Attitude
 ArXiv CS.AI 50%
 arXiv:2606.12424v1 Announce Type: cross Abstract: As generative AI and low-code workflow platforms become routine in software practice, a key educational question is whether the next generation of computer engineers will accept these tools as useful, usable, and worthy of...
<https://arxiv.org/abs/2606.12424>
- 1 **CONF**
An Explainable AI Assistant for Introductory Programming Education: Improving Feedback Reliability with Instructor-AI Collaboration
 ArXiv CS.AI 50%
 arXiv:2606.12425v1 Announce Type: cross Abstract: Active learning is widely recognized as an effective approach for improving learning outcomes in introductory programming courses. However, insufficient instructional support often limits students' access to timely, personalized...
<https://arxiv.org/abs/2606.12425>
- 2 **CONF**
Mapping AI Programs in the U.S: A Status Report from Early 2026 and an Analysis of AI Majors and Minors
 ArXiv CS.AI 50%
 arXiv:2606.12428v1 Announce Type: cross Abstract: We present a report on the status of undergraduate Artificial Intelligence (AI) programs in the United States in Spring 2026. In so doing, we 1) describe our scraping and mapping tools, which dynamically update to track the state...
<https://arxiv.org/abs/2606.12428>

3 CONF

Will AI Agents Free Us From Meaningless Work? A Human-Centered Analysis

ArXiv CS.AI 50%

arXiv:2606.12430v1 Announce Type: cross Abstract: Some claim that AI agents will free workers from the boring parts of their jobs, yet little is known about how workers themselves identify which tasks should be automated. Prior research focuses on occupations, overlooking that...

<https://arxiv.org/abs/2606.12430>

4 CONF

Position: Generative Engine Optimization Creates Underexamined Risks, Governance Must Target Concentration, Disclosure, and Academic Blind Spots

ArXiv CS.AI 50%

arXiv:2606.12439v1 Announce Type: cross Abstract: Large language model (LLM) answer engines are increasingly used for information seeking, shifting visibility from ranked lists to synthesized answers. This enables Generative Engine Optimization (GEO), which targets LLM answer...

<https://arxiv.org/abs/2606.12439>

5 CONF

Generativism: Toward a Learning Theory for the Age of Generative Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2606.12441v1 Announce Type: cross Abstract: The four dominant learning theories of behaviorism, cognitivism, constructivism, and connectivism show significant conceptual limitations as generative artificial intelligence (AI) proliferates in educational settings. These...

<https://arxiv.org/abs/2606.12441>

6 CONF

Reframing AI Loss of Control: What It Is, How to Have It, How to Lose It

ArXiv CS.AI 50%

arXiv:2606.12442v1 Announce Type: cross Abstract: At present, loss of control risks have gained much prominence in public discussion, particularly in relation to AI, with extensive discourse present among academics, frontier labs, and even governments. However, in the existing...

<https://arxiv.org/abs/2606.12442>

7 CONF

The Internet of Agentic AI: Communication, Coordination, and Collective Intelligence at Scale

ArXiv CS.AI 50%

arXiv:2606.12835v1 Announce Type: cross Abstract: The rapid emergence of autonomous AI agents is transforming artificial intelligence from isolated model inference into distributed systems of reasoning, communication, and action. This paper develops the vision of the Internet of...

<https://arxiv.org/abs/2606.12835>

8 CONF

ReCal: Reward Calibration for RL-based LLM Routing

ArXiv CS.AI 50%

arXiv:2606.12479v1 Announce Type: cross Abstract: Large language model (LLM) routing has emerged as an effective paradigm for leveraging the complementary strengths of multiple LLMs through dynamic model and reasoning-strategy selection. Recent reinforcement learning (RL)-based...

<https://arxiv.org/abs/2606.12479>

9 CONF

Improving Crash Frequency Prediction from Simulated Traffic Conflicts Using Machine Learning Based Microsimulation

ArXiv CS.AI 50%

arXiv:2606.12500v1 Announce Type: cross Abstract: Traffic microsimulation combined with surrogate safety measures has increasingly been used as a proactive alternative to historical crash data for predicting crash frequency for current or planned road infrastructure designs....

<https://arxiv.org/abs/2606.12500>

0 CONF

Graph Reduction in Multirelational Networks: A Spreading-Oriented Reduction Benchmark

ArXiv CS.AI 50%

arXiv:2606.12581v1 Announce Type: cross Abstract: Real-world networks are inherently incomplete, noisy, and dynamically evolving, making it difficult to capture all actors and their relationships. Their scale often renders direct analysis computationally demanding. While...

<https://arxiv.org/abs/2606.12581>

1 CONF

Analyzing and Improving Fine-grained Preference Optimization in Medical LVLMs

ArXiv CS.AI 50%

arXiv:2606.12590v1 Announce Type: cross Abstract: Large Vision-Language Models (LVLMs) have achieved strong performance across medical imaging tasks, yet they remain prone to factual inconsistencies, poor visual grounding, and misalignment with clinically meaningful feedback....

<https://arxiv.org/abs/2606.12590>

2 CONF

From Imitation to Alignment: Human-Preference Flow Policies for Long-Horizon Sidewalk Navigation

ArXiv CS.AI 50%

arXiv:2606.12603v1 Announce Type: cross Abstract: Autonomous long-horizon sidewalk navigation is essential for micro-mobility applications such as robotic food delivery and assistive electronic wheelchairs. Unlike autonomous driving on the road, long-horizon sidewalk navigation...

<https://arxiv.org/abs/2606.12603>

3 CONF

HybridCodeAuthorship: A Benchmark Dataset for Line-Level Code Authorship Detection

ArXiv CS.AI 50%

arXiv:2606.12620v1 Announce Type: cross Abstract: Thanks to the rapid adoption of AI code assistants powered by large language models (LLMs), industry codebases are, increasingly, a hybrid of AI- and human-authored code. For risk management and productivity analysis purposes, it...

<https://arxiv.org/abs/2606.12620>

4 CONF

Bag of Dims: Training-Free Mechanistic Interpretability via Dimension-Level Sign Patterns

ArXiv CS.AI 50%

arXiv:2606.12629v1 Announce Type: cross Abstract: We show that the standard basis of transformer hidden states already provides a training-free, architecture-general feature basis. Individual dimensions encode semantic content via their signs and confidence via their magnitudes,...

<https://arxiv.org/abs/2606.12629>

5 **CONF****Token Complexity Theory for AI-Augmented Computing**

ArXiv CS.AI 50%

arXiv:2606.12647v1 Announce Type: cross Abstract: AI-augmented computing delegates natural language queries, code generation requests, and other open-ended tasks to a cluster of AI models that processes queries and generates responses. This paradigm introduces a resource...

<https://arxiv.org/abs/2606.12647>

6 **CONF****M*: A Modular, Extensible, Serving System for Multimodal Models**

ArXiv CS.AI 50%

arXiv:2606.12688v1 Announce Type: cross Abstract: We are entering a new era of composite model architectures that integrate diverse components such as vision encoders, language backbones, diffusion and flow heads, audio codecs, action generators, and world-model predictors. Such...

<https://arxiv.org/abs/2606.12688>

7 **CONF****LLM-Powered Personalized Glycemic Assessment in Type 2 Diabetes with Wearable Sensor Data**

ArXiv CS.AI 50%

arXiv:2606.12699v1 Announce Type: cross Abstract: Type 2 Diabetes (T2D) poses an increasing global health threat, demanding effective glycemic assessment to support personalized and improved diabetes care. Wearable sensors such as continuous glucose monitors (CGM) and fitness...

<https://arxiv.org/abs/2606.12699>

8 **CONF****SMR: Certified Defence Against Runtime Memory Poisoning in Persistent LLM Agent Systems**

ArXiv CS.AI 50%

arXiv:2606.12703v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) agents increasingly run with persistent memory that accumulates across user sessions. This creates a new attack surface: an adversary interacting only through normal channels can inject...

<https://arxiv.org/abs/2606.12703>

9 **CONF****LLMs Can Better Capture Human Judgments--With the Right Prompts**

ArXiv CS.AI 50%

arXiv:2606.12754v1 Announce Type: cross Abstract: Are large language models (LLMs) bad at capturing human judgment? Two commonly stated limitations are that LLMs fail to capture full distributions of responses, and that their judgments are unstable across wording variations. We...

<https://arxiv.org/abs/2606.12754>

0 **CONF****Exploring How Agent Voice Accents Shape Human-AI Collaboration in K-12 Group Learning**

ArXiv CS.AI 50%

arXiv:2606.12805v1 Announce Type: cross Abstract: Collaboration is widely recognized as a cornerstone of 21st-century education, yet teachers still encounter persistent challenges in fostering productive peer interaction. LLM conversational peer agents introduce new...

<https://arxiv.org/abs/2606.12805>

1 CONF

Stubborn: A Streamlined and Unified Reinforcement Learning Framework for Robust Motion Tracking and Fall Recovery for Humanoids

ArXiv CS.AI 50%

arXiv:2606.12814v1 Announce Type: cross Abstract: Recent reinforcement learning approaches have shown great promise in improving humanoid motion tracking performance and achieving fall recovery under disturbances. However, most existing works treat motion tracking and fall...

<https://arxiv.org/abs/2606.12814>

2 CONF

Acquisition state behaves as a structured, measurable variable governing lung-nodule AI: kernel-driven measurement instability and noise-driven detection fragility, invisible to DICOM metadata

ArXiv CS.AI 50%

arXiv:2606.12824v1 Announce Type: cross Abstract: AI governance for medical imaging is formalizing: the 2026 ACR-SIIM Practice Parameter recommends local acceptance testing and ongoing drift monitoring, and the ACR Assess-AI registry monitors AI outputs using DICOM metadata for...

<https://arxiv.org/abs/2606.12824>

3 CONF

Beyond Problem Solving: UOJ-Bench for Evaluating Code Generation, Hacking, and Repair in Competitive Programming

ArXiv CS.AI 50%

arXiv:2606.12864v1 Announce Type: cross Abstract: Despite strong performance in competitive programming, the role of Large Language Models (LLMs) in supporting human learning in the same setting remains largely unexplored. In this work, we introduce UOJ-Bench, a benchmark...

<https://arxiv.org/abs/2606.12864>

4 CONF

Efficient, Robust, and Anti-Collusion Fingerprinting of Image Diffusion Models

ArXiv CS.AI 50%

arXiv:2606.12977v1 Announce Type: cross Abstract: Model fingerprinting, embedding user-specific identifiers (fingerprints) into generated outputs, has recently emerged as a popular solution to protect the intellectual property rights (IPR) of generative text-to-image (T2I)...

<https://arxiv.org/abs/2606.12977>

5 CONF

A Machine Learning Framework for Real-Time Personalized Ergonomic Pose Analysis

ArXiv CS.AI 50%

arXiv:2606.12988v1 Announce Type: cross Abstract: This paper introduces a new methodology for real-time prediction of ergonomic and non-ergonomic human poses using volumetric video data in three dimensions. Although the methodology was designed for ergonomic assessments, it can...

<https://arxiv.org/abs/2606.12988>

6 CONF

scLLM-DSC: LLM-Knowledge Enhanced Cross-Modal Deep Structural Clustering for Single-Cell RNA Sequencing

ArXiv CS.AI 50%

arXiv:2606.13007v1 Announce Type: cross Abstract: Clustering is fundamental to scRNA-seq analysis, serving as a cornerstone for identifying cell populations and resolving tissue heterogeneity. However, existing methods focus on mining numerical statistical patterns, suffering...

<https://arxiv.org/abs/2606.13007>

7 **CONF****CausalMoE: A Billion-Scale Multimodal Foundation Model for Granger Causal Discovery with Pattern-Routed Heterogeneous Experts**

ArXiv CS.AI 50%

arXiv:2606.13024v1 Announce Type: cross Abstract: Granger Causal Discovery (GCD) is fundamental for analyzing temporal dependencies in complex systems. However, existing neural GCD methods predominantly rely on a "one-size-fits-all" paradigm, struggling to capture distribution...

<https://arxiv.org/abs/2606.13024>

8 **CONF****TetherCache: Stabilizing Autoregressive Long-Form Video Generation with Gated Recall and Trusted Alignment**

ArXiv CS.AI 50%

arXiv:2606.13035v1 Announce Type: cross Abstract: Autoregressive video diffusion models provide a natural formulation for streaming and variable-length video generation by conditioning newly generated frames on previously generated content. However, extending these models to...

<https://arxiv.org/abs/2606.13035>

9 **CONF****Fault Lines: Navigating Ethics and Responsible AI Where National Policy Meets Local Practice in Public Sector Transformation**

ArXiv CS.AI 50%

arXiv:2606.13039v1 Announce Type: cross Abstract: The UK government has adopted a pro-AI stance to help transform public service delivery in the face of severe financial pressures, but the path to translate this vision into responsible AI practice remains ill-defined. While UK...

<https://arxiv.org/abs/2606.13039>

0 **CONF****TWLA: Achieving Ternary Weights and Low-Bit Activations for LLMs via Post-Training Quantization**

ArXiv CS.AI 50%

arXiv:2606.13054v1 Announce Type: cross Abstract: Large language models (LLMs) exhibit exceptional general language processing capabilities, but their memory and compute costs hinder deployment. Ternarization has emerged as a promising compression technique, offering significant...

<https://arxiv.org/abs/2606.13054>

1 **CONF****MP3: Multi-Period Pattern Pre-training for Spatio-Temporal Forecasting**

ArXiv CS.AI 50%

arXiv:2606.13119v1 Announce Type: cross Abstract: Spatio-Temporal forecasting is crucial in diverse fields, such as transportation, climate, and energy. Urban spatio-temporal data exhibits temporal mirage: similar short-window inputs have divergent future trends, and vice versa....

<https://arxiv.org/abs/2606.13119>

2 **CONF****Select and Improve: Understanding the Mechanics of Post-Training for Reasoning**

ArXiv CS.AI 50%

arXiv:2606.13125v1 Announce Type: cross Abstract: Reinforcement learning has rapidly emerged as a key component in the training of reasoning and coding models, yet it remains poorly understood from a mechanistic perspective. We study how and through what underlying processes...

<https://arxiv.org/abs/2606.13125>

3 **CONF****Iterative Visual Thinking: Teaching Vision-Language Models Spatial Self-Correction through Visual Feedback****ArXiv CS.AI** **50%**

arXiv:2606.13156v1 Announce Type: cross Abstract: Vision-language models (VLMs) achieve strong singleshot spatial grounding, yet lack any mechanism to observe and correct their own predictions. We find that naively prompting a VLM to iterate over rendered visualizations of its...

<https://arxiv.org/abs/2606.13156>

4 **CONF****NTS-CoT: Mitigating Hallucinations in LLM-based News Timeline Summarization with Chain-of-Thought Reasoning****ArXiv CS.AI** **50%**

arXiv:2606.13171v1 Announce Type: cross Abstract: The rapid updates of online news make tracking event developments challenging, highlighting the need for timeline summarization (TLS). Hallucinations, where LLM-generated content deviates from source news, still remain a critical...

<https://arxiv.org/abs/2606.13171>

5 **CONF****Contrast-Informed Augmentation and Domain-Adversarial Training for Adult-to-Neonatal MR Reconstruction Generalization****ArXiv CS.AI** **50%**

arXiv:2606.13562v1 Announce Type: cross Abstract: Purpose: To investigate whether contrast-informed data augmentation and domain-adversarial training improve the adult-to-neonatal generalization of the E2E-VarNet. Methods: Three training regimes were investigated: (1) adult-only...

<https://arxiv.org/abs/2606.13562>

6 **CONF****Towards More General Control of Diffusion Models Using Jeffrey Guidance****ArXiv CS.AI** **50%**

arXiv:2606.13240v1 Announce Type: cross Abstract: A key strength of diffusion models lies in their flexibility, since their outputs can be controlled at sampling time through guidance. However, beyond simple cases such as conditional sampling, the target distribution is often...

<https://arxiv.org/abs/2606.13240>

7 **CONF****Humor Style Drives Laughter, Topic Shapes Acceptability: Evaluating Bilingual Personal and Political Robot-Delivered AI Jokes****ArXiv CS.AI** **50%**

arXiv:2606.13256v1 Announce Type: cross Abstract: Humor plays a central role in human social relationships, and recent advances in computational humor create new opportunities for integrating humor into human-robot interaction (HRI). While large language models (LLMs) can...

<https://arxiv.org/abs/2606.13256>

8 **CONF****Different Layers, Different Manifolds: Module-Wise Weight-Space Geometry in Transformer Optimization****ArXiv CS.AI** **50%**

arXiv:2606.13276v1 Announce Type: cross Abstract: Weight-space geometry plays a central role in neural network optimization, yet manifold constraints are often applied uniformly across all weight matrices. In this work, we ask whether different transformer modules prefer...

<https://arxiv.org/abs/2606.13276>

9 CONF

HYDRA-X: Native Unified Multimodal Models with Holistic Visual Tokenizers

ArXiv CS.AI 50%

arXiv:2606.13289v1 Announce Type: cross Abstract: Holistic visual tokenizers are fundamental to unified multimodal models (UMMs) as they map diverse visual inputs into a unified representation space. In this paper, we present HYDRA-X, the first UMM that unifies image and video...

<https://arxiv.org/abs/2606.13289>

00 CONF

Mining Architectural Quality Under Agentic AI Adoption: A Causal Study of Java Repositories

ArXiv CS.AI 50%

arXiv:2606.13298v1 Announce Type: cross Abstract: AI coding tools are now used by a majority of developers, and agentic use of these tools has popularized the practice colloquially called "vibe coding". Yet causal evidence on their effect on software architecture is scarce....

<https://arxiv.org/abs/2606.13298>

01 CONF

An LLM System for Autonomous Variational Quantum Circuit Design

ArXiv CS.AI 50%

arXiv:2606.13380v1 Announce Type: cross Abstract: The design of high performing quantum circuits remains largely dependent on human expertise. We introduce an autonomous agentic framework that employs large language models (LLMs) to conduct iterative quantum circuit designs...

<https://arxiv.org/abs/2606.13380>

02 CONF

Who Pays the Price? Stakeholder-Centric Prompt Injection Benchmarking for Real-world Web Agents

ArXiv CS.AI 50%

arXiv:2606.13385v1 Announce Type: cross Abstract: Web agents driven by large language models (LLMs) are increasingly deployed in real-world environments, where they operate over untrusted web content and execute actions with direct consequences. This makes them vulnerable to...

<https://arxiv.org/abs/2606.13385>

03 CONF

Mod-Guide: An LLM-based Content Moderation Feedback System to Address Insensitive Speech toward Indigenous Ethnic and Religious Minority Communities

ArXiv CS.AI 50%

arXiv:2606.13397v1 Announce Type: cross Abstract: Language operates as a mechanism of both marginalization and resistance, especially for minority communities navigating insensitive and harmful speech online. As content moderation increasingly depends on large language models...

<https://arxiv.org/abs/2606.13397>

04 CONF

PolyFlow: Safe and Efficient Polytope-Constrained Flow Matching with Constraint Embedding and Projection-free Update

ArXiv CS.AI 50%

arXiv:2606.13400v1 Announce Type: cross Abstract: While flow-based generative models have demonstrated strong performance across a wide range of domains, deploying them in safety-critical physical systems remains challenging due to strict constraint requirements. Existing...

<https://arxiv.org/abs/2606.13400>

05 CONF

OmniDirector: General Multi-Shot Camera Cloning without Cross-Paired Data

ArXiv CS.AI 50%

arXiv:2606.13432v1 Announce Type: cross Abstract: Cloning camera motion from reference videos is an important task in video generation, as videos provide intuitive and precise control. Existing methods either directly use parametric representations that fail to handle multi-shot...

<https://arxiv.org/abs/2606.13432>

06 CONF

Understanding the Rejection of Fixes Generated by Agentic Pull Requests -- Insights from the AIDev Dataset

ArXiv CS.AI 50%

arXiv:2606.13468v1 Announce Type: cross Abstract: AI coding agents are increasingly used to generate pull requests (PRs) that propose code fixes in software projects. From a first exploration of the AIDev dataset, we find that 46.41\% of the fixes proposed by the agents Copilot,...

<https://arxiv.org/abs/2606.13468>

07 CONF

MaxProof: Scaling Mathematical Proof with Generative-Verifier RL and Population-Level Test-Time Scaling

ArXiv CS.AI 50%

arXiv:2606.13473v1 Announce Type: cross Abstract: We present MaxProof, a population-level test-time scaling framework for competition-level mathematical proof in the MiniMax-M3 series. M3 first trains three proof-oriented capabilities -- proof generation, proof verification, and...

<https://arxiv.org/abs/2606.13473>

08 CONF

Beyond the Commitment Boundary: Probing Epiphenomenal Chain-of-Thought in Large Reasoning Models

ArXiv CS.AI 50%

arXiv:2606.13603v1 Announce Type: cross Abstract: Chain-of-thought (CoT) reasoning is the dominant paradigm for inference-time scaling in language models, yet the causal influence of individual steps on the final answer poorly understood. We estimate each step's causal...

<https://arxiv.org/abs/2606.13603>

09 CONF

One Polluted Page Is Enough: Evaluating Web Content Pollution in Generative Recommenders

ArXiv CS.AI 50%

arXiv:2606.13610v1 Announce Type: cross Abstract: Search-augmented LLMs increasingly mediate everyday consumer recommendations by retrieving live web content. This creates a new risk: generative recommenders may consume polluted web content, such as fake reviews and promotional...

<https://arxiv.org/abs/2606.13610>

10 CONF

Transformer Field Theory: A Response-Theoretic Approach to Mechanistic Interpretability

ArXiv CS.AI 50%

arXiv:2605.25225v2 Announce Type: replace-cross Abstract: Mechanistic interpretability often studies Transformer behavior by intervening on internal activations through activation patching, causal tracing, path patching, and steering directions. This paper develops Transformer...

<https://arxiv.org/abs/2605.25225>

11 CONF

SkMTEB: Slovak Massive Text Embedding Benchmark and Model Adaptation

ArXiv CS.AI 50%

arXiv:2606.13647v1 Announce Type: cross Abstract: We introduce SkMTEB, the first comprehensive MTEB-style text embedding benchmark for Slovak, a low-resource West Slavic language, comprising 31 datasets across 7 task types -- nearly 4\times\$ the depth of existing multilingual...

<https://arxiv.org/abs/2606.13647>

12 CONF

Learning to Reason by Analogy via Retrieval-Augmented Reinforcement Fine-Tuning

ArXiv CS.AI 50%

arXiv:2606.13680v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) has become a standard mechanism for grounding language models in external knowledge, yet conventional retrieval based on lexical or semantic similarity is poorly suited for complex reasoning...

<https://arxiv.org/abs/2606.13680>

13 CONF

Interaction-Centered Intelligence: Toward an Interaction-Based Theory of Human-AI Co-Creation

ArXiv CS.AI 50%

arXiv:2606.00807v2 Announce Type: replace Abstract: Traditional artificial intelligence has largely conceptualized intelligence as isolated computation occurring within bounded agents. Across classical AI, machine learning, and many generative systems, the dominant unit of...

<https://arxiv.org/abs/2606.00807>

14 CONF

On Pitfalls of \$\textit{RemOve-And-Retrain}\$: Data Processing Inequality Perspective

ArXiv CS.AI 50%

arXiv:2304.13836v5 Announce Type: replace-cross Abstract: The RemOve-And-Retrain (ROAR) benchmark is widely used to evaluate feature attribution methods, yet its validity remains underexplored from an information-theoretic perspective. We show that model- and data-agnostic...

<https://arxiv.org/abs/2304.13836>

15 CONF

Competition and Diversity in Generative AI

ArXiv CS.AI 50%

arXiv:2412.08610v3 Announce Type: replace-cross Abstract: Recent evidence, both in the lab and in the wild, suggests that the use of generative artificial intelligence reduces the diversity of content produced. The use of the same or similar AI models appears to lead to more...

<https://arxiv.org/abs/2412.08610>

16 CONF

Lightweight and Interpretable Transformer via Mixed Graph Algorithm Unrolling for Traffic Forecast

ArXiv CS.AI 50%

arXiv:2505.13102v4 Announce Type: replace-cross Abstract: Unlike conventional "black-box" transformers with classical self-attention mechanism, we build a lightweight and interpretable transformer-like neural net by unrolling a mixed-graph-based optimization algorithm to...

<https://arxiv.org/abs/2505.13102>

17 CONF

Meta-Learning Transformers to Improve In-Context Generalization

ArXiv CS.AI 50%

arXiv:2507.05019v2 Announce Type: replace-cross Abstract: In-context learning enables transformer models to generalize to new tasks based solely on input prompts, without any need for weight updates. However, existing training paradigms typically rely on large, unstructured...

<https://arxiv.org/abs/2507.05019>

18 CONF

Decoding the Multimodal Maze: A Systematic Review on the Adoption of Explainability in Multimodal Attention-based Models

ArXiv CS.AI 50%

arXiv:2508.04427v2 Announce Type: replace-cross Abstract: Multimodal learning has witnessed remarkable advancements in recent years, particularly with the integration of attention-based models, leading to significant performance gains across a variety of tasks. Parallel to this...

<https://arxiv.org/abs/2508.04427>

19 CONF

Structuring The Future: Diffusion LLM Speculative Decoding via Calibrated Draft Graphs

ArXiv CS.AI 50%

arXiv:2509.18085v4 Announce Type: replace-cross Abstract: Diffusion LLMs (dLLMs) have recently emerged as a powerful alternative to autoregressive LLMs (AR-LLMs) with the potential to operate at significantly higher token-generation rates. To unlock this potential, we present...

<https://arxiv.org/abs/2509.18085>

20 CONF

Improving Pre-trained Adult Glioma Segmentation Models Using only Post-processing Techniques

ArXiv CS.AI 50%

arXiv:2512.14937v2 Announce Type: replace-cross Abstract: Gliomas are the most common malignant brain tumors in adults and are among the most lethal. Despite aggressive treatment, the median survival rate is less than 15 months. Accurate multiparametric MRI (mpMRI) tumor...

<https://arxiv.org/abs/2512.14937>

21 CONF

CuMA: Aligning LLMs with Sparse Cultural Values via Demographic-Aware Mixture of Adapters

ArXiv CS.AI 50%

arXiv:2601.04885v2 Announce Type: replace-cross Abstract: As Large Language Models (LLMs) serve a global audience, alignment must transition from enforcing universal consensus to respecting cultural pluralism. We demonstrate that dense models, when forced to fit conflicting...

<https://arxiv.org/abs/2601.04885>

22 CONF

When Smaller Wins: Dual-Stage Distillation and Pareto-Guided Compression of Liquid Neural Networks for Edge Battery Prognostics

ArXiv CS.AI 50%

arXiv:2601.06227v3 Announce Type: replace-cross Abstract: Battery management systems increasingly require accurate battery health prognostics under strict on-device constraints. This paper presents DLNet, a practical framework with dual-stage distillation of liquid neural...

<https://arxiv.org/abs/2601.06227>

23 CONF

Language Model Circuits Are Sparse in the Neuron Basis

ArXiv CS.AI 50%

arXiv:2601.22594v2 Announce Type: replace-cross Abstract: The high-level concepts that a neural network uses to perform computation need not be aligned to individual neurons (Smolensky, 1986). Language model interpretability research has thus turned to techniques which decompose...

<https://arxiv.org/abs/2601.22594>

24 CONF

Fin-RATE: A Real-world Financial Analytics and Tracking Evaluation Benchmark for LLMs on SEC Filings

ArXiv CS.AI 50%

arXiv:2602.07294v4 Announce Type: replace-cross Abstract: With the increasing deployment of Large Language Models (LLMs) in the finance domain, LLMs are increasingly expected to parse complex regulatory disclosures. However, existing benchmarks often focus on isolated details,...

<https://arxiv.org/abs/2602.07294>

25 CONF

Unsafer in Many Turns: Benchmarking and Defending Multi-Turn Safety Risks in Tool-Using Agents

ArXiv CS.AI 50%

arXiv:2602.13379v2 Announce Type: replace-cross Abstract: LLM-based agents are becoming increasingly capable, yet their safety lags behind. This creates a gap between what agents can do and should do. This gap widens as agents engage in multi-turn interactions and employ diverse...

<https://arxiv.org/abs/2602.13379>

26 CONF

FENCE: A Financial and Multimodal Jailbreak Detection Dataset

ArXiv CS.AI 50%

arXiv:2602.18154v2 Announce Type: replace-cross Abstract: Jailbreaking poses a significant risk to the deployment of Large Language Models (LLMs) and Vision Language Models (VLMs). VLMs are particularly vulnerable because they process both text and images, creating broader...

<https://arxiv.org/abs/2602.18154>

27 CONF

CMI-RewardBench: Evaluating Music Reward Models with Compositional Multimodal Instruction

ArXiv CS.AI 50%

arXiv:2603.00610v3 Announce Type: replace-cross Abstract: While music generation models have evolved to handle complex multimodal inputs mixing text, lyrics, and reference audio, evaluation mechanisms have lagged behind. In this paper, we bridge this critical gap by establishing...

<https://arxiv.org/abs/2603.00610>

28 CONF

Fusion Learning from Dynamic Functional Connectivity: Combining the Amplitude and Phase of fMRI Signals to Identify Brain Disorders

ArXiv CS.AI 50%

arXiv:2603.24603v2 Announce Type: replace-cross Abstract: Dynamic functional connectivity (dFC) derived from resting-state functional magnetic resonance imaging (fMRI) has been extensively utilized in brain science research. The sliding window correlation (SWC) method is a...

<https://arxiv.org/abs/2603.24603>

29 CONF

A Survey on Long-Term Memory Security in LLM Agents: Attacks, Defenses, and Governance Across the Memory Lifecycle

ArXiv CS.AI 50%

arXiv:2604.16548v2 Announce Type: replace-cross Abstract: The emergence of writable, cross-session persistent memory in LLM agents introduces a qualitatively different threat landscape from conventional input-centric security concerns, characterized by three properties:...

<https://arxiv.org/abs/2604.16548>

30 CONF

Versioned Late Materialization for Ultra-Long Sequence Training in Recommendation Systems at Scale

ArXiv CS.AI 50%

arXiv:2604.24806v2 Announce Type: replace-cross Abstract: Modern Deep Learning Recommendation Models (DLRMs) follow scaling laws with sequence length, driving the frontier toward ultra-long User Interaction History (UIH). However, the industry-standard "Fat Row" paradigm, which...

<https://arxiv.org/abs/2604.24806>

31 CONF

Benchmarking Counterfactual Prediction in Epidemic Time Series with Time-Varying Interventions

ArXiv CS.AI 50%

arXiv:2606.05692v2 Announce Type: replace-cross Abstract: Deep learning has enabled significant advances in time-series causal inference, yet progress remains constrained by the lack of realistic benchmarks with observable counterfactual outcomes. Existing datasets either rely...

<https://arxiv.org/abs/2606.05692>

32 CONF

Towards Data-free and Training-free Compression for Speech Foundation Models Using Parameter Clustering

ArXiv CS.AI 50%

arXiv:2606.11836v2 Announce Type: replace-cross Abstract: This paper presents a novel data-free and training-free compression approach for speech foundation models using channelwise clustering via k-means. More fine-grained, mixed sparsity pruning by layer-level varying number...

<https://arxiv.org/abs/2606.11836>

33 CONF

Frozen Multimodal Embeddings for AI-Assisted Interview Assessment of Personality and Cognitive Ability

ArXiv CS.AI 50%

arXiv:2606.11930v2 Announce Type: replace-cross Abstract: Predicting psychological traits from asynchronous video interviews (AVIs) is a challenging problem in AI-assisted interview assessment because labeled datasets are limited while each response contains high-dimensional...

<https://arxiv.org/abs/2606.11930>