



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	113	Media
TechCrunch AI	3	Media
HackerNews	2	Community

VERIFIED INTELLIGENCE — 2 stories

VERI

Communication Policy Evolution for Proactive LLM Agents

ArXiv CS.AI +1 source

80%

arXiv:2606.14314v1 Announce Type: new Abstract: LLM agents have rapidly evolved into autonomous systems, yet a persistent information gap remains between users and agents: communication is costly, while users' identical preferences further limit information exchange. To...

<https://arxiv.org/abs/2606.14314>

VERI

FedRot-LoRA: Mitigating Rotational Misalignment in Federated LoRA

ArXiv CS.AI +1 source

80%

arXiv:2602.23638v3 Announce Type: replace-cross Abstract: Federated LoRA provides a communication-efficient mechanism for fine-tuning large language models on decentralized data. In practice, however, a discrepancy between the factor-wise averaging used to preserve low rank and...

<https://arxiv.org/abs/2602.23638>

CONFIRMED SIGNALS — 114 stories

CONF

Show HN: Skill for your agent to visualize your gbrain and Obsidian

HackerNews 50% HN 21

No summary — see link for full article.

<https://github.com/vladignatyev/brain-map-skill>

CONF

Show HN: Dream Server - Turn your PC, Mac, or Linux box into a private AI server

HackerNews 50% HN 6

No summary — see link for full article.

<https://github.com/Light-Heart-Labs/DreamServer>

CONF

The AI layoff wave is becoming a powder keg

TechCrunch AI 50%

What makes this combustible: at the very moment that tens of thousands of workers are being shown the door, a small cohort of AI insiders is becoming wealthy on a scale that's hard to comprehend.

<https://techcrunch.com/2026/06/15/the-ai-layoff-wave-is-becoming-a-powder-keg/>

CONF

A Deep Reinforcement Learning (DRL)-Based Transformer Method for Solving the Open Shop Scheduling Problem

ArXiv CS.AI 50%

arXiv:2606.13682v1 Announce Type: new Abstract: The open shop scheduling problem (OSSP) arises in many industrial and service settings but remains computationally challenging as the number of jobs and machines increases. While exact methods quickly become intractable, classical...

<https://arxiv.org/abs/2606.13682>

CONF

UP-NRPA: User Portrait based Nested Rollout Policy Adaptation for Planning with Large Language Models in Goal-oriented Dialogue Systems

ArXiv CS.AI 50%

arXiv:2606.13683v1 Announce Type: new Abstract: To address the challenge that current dialogue policy planning methods struggle to dynamically adapt to diverse user characteristics, this paper proposes a User Portrait based Nested Rollout Policy Adaptation (UP-NRPA) online...

<https://arxiv.org/abs/2606.13683>

CONF

Refusal Beyond a Single Direction: A Preliminary Comparison of Diff-in-Means and INLP

ArXiv CS.AI 50%

arXiv:2606.13720v1 Announce Type: new Abstract: Arditi et al. (2024) has shown that refusal in safety fine-tuned chat models is mediated by a single linear direction in the residual stream, recoverable by a difference-in-means (DiM) of harmful and harmless activations. We...

<https://arxiv.org/abs/2606.13720>

CONF

YeasierAgent: Agentic Social Sandbox as a Canvas for Intent-Driven Creation of Platform-Agnostic Symbiotic Agent-Native Applications

ArXiv CS.AI 50%

arXiv:2606.13722v1 Announce Type: new Abstract: This paper introduces YeasierAgent, an application-building paradigm based on symbiotic agents, narrative worlds, and scene-aware interaction. It challenges the conventional device-coupled model of software by redefining...

<https://arxiv.org/abs/2606.13722>

0 CONF

AI Receptivity or AI Adoption Breadth? A Tool-Specific Reanalysis of the Lower-Literacy/Higher-Usage Link

ArXiv CS.AI 50%

arXiv:2606.13734v1 Announce Type: new Abstract: Recent evidence reported by Tully, Longoni, and Appel (2025) suggests that lower artificial intelligence (AI) literacy predicts greater receptivity toward AI. We revisit this claim using the public data from Study 3 of that...

<https://arxiv.org/abs/2606.13734>

1 CONF

MA-ProofBench: A Two-Tiered Evaluation of LLMs for Theorem Proving in Mathematical Analysis

ArXiv CS.AI 50%

arXiv:2606.13782v1 Announce Type: new Abstract: Large Language Models (LLMs) have made notable progress in automated theorem proving, yet existing formal benchmarks remain limited in both mathematical coverage and difficulty. Most are concentrated in areas that are easier to...

<https://arxiv.org/abs/2606.13782>

2 CONF

Poker Arena: Multi-Axis Profiling of Strategic Reasoning and Memory in LLMs

ArXiv CS.AI 50%

arXiv:2606.13815v1 Announce Type: new Abstract: Strategic reasoning under uncertainty underpins consequential decisions in negotiation, finance, and policy, but prevailing game-play benchmarks collapse heterogeneous reasoning dimensions into a single scalar, leaving the...

<https://arxiv.org/abs/2606.13815>

3 CONF

Hyperdimensional computing for structured querying on tabular data embeddings

ArXiv CS.AI 50%

arXiv:2606.13871v1 Announce Type: new Abstract: Tabular data embeddings have become a cornerstone of data profiling and data integration pipelines, enabling tasks such as entity annotation and resolution; schema matching; column type detection; and table search, among others....

<https://arxiv.org/abs/2606.13871>

4 CONF

Capability Minimization as a Safety Primitive: Risk-Aware Causal Gating for Least-Privilege LLM Agents

ArXiv CS.AI 50%

arXiv:2606.13884v1 Announce Type: new Abstract: Modern decision systems increasingly rely on learned components whose outputs may be confident yet wrong, exposing downstream actions to costly errors. We introduce Risk-Aware Causal Gating (RACG), a framework that decides whether...

<https://arxiv.org/abs/2606.13884>

5 CONF

A Multi-Agent AI System for Automated High School Transcript Processing: Collaborative Document Analysis at Scale

ArXiv CS.AI 50%

arXiv:2606.13916v1 Announce Type: new Abstract: Each year, college admissions offices face an overwhelming challenge: processing millions of high school transcripts, each with unique formats, grading systems, and layouts. This manual process creates operational bottlenecks that...

<https://arxiv.org/abs/2606.13916>

6 CONF

FactoryLLM: A Safe and Open-Source AI Playground for Evaluating LLMs in Smart Factories

ArXiv CS.AI 50%

arXiv:2606.14119v1 Announce Type: new Abstract: Fault diagnostics and recovery in smart factories is challenging because critical information is dispersed across manuals of multiple machines which are interconnected through the manufacturing process. Large Language Models (LLMs)...

<https://arxiv.org/abs/2606.14119>

7 CONF

SkillAudit: Ground-Truth-Free Skill Evolution via Paired Trajectory Auditing

ArXiv CS.AI 50%

arXiv:2606.14239v1 Announce Type: new Abstract: Agent skills are structured procedural packages that guide frozen LLM agents in specialized workflows. Skills rarely remain sufficient after deployment: edge cases, API changes, and deployment constraints become visible only...

<https://arxiv.org/abs/2606.14239>

8 CONF

Deep Dense Exploration for LLM Reinforcement Learning via Pivot-Driven Resampling

ArXiv CS.AI 50%

arXiv:2602.14169v2 Announce Type: replace-cross Abstract: Effective exploration is a key challenge in reinforcement learning for large language models: discovering high-quality trajectories within a limited sampling budget from the vast natural language sequence space. Existing...

<https://arxiv.org/abs/2602.14169>

- 9
CONF
When the Tool Decides: LLM Agents Defer Blindly to Graph Neural Network Tools, and Stronger Backbones Defer More
ArXiv CS.AI 50%
arXiv:2606.14476v1 Announce Type: new Abstract: A growing line of work equips large language model (LLM) agents with graph neural networks (GNNs) as callable tools, assuming the agent exercises judgment over when and how much to rely on such a tool. We test this directly. We...
<https://arxiv.org/abs/2606.14476>
- 0
CONF
From Chatbot to Digital Colleague: The Paradigm Shift Toward Persistent Autonomous AI
ArXiv CS.AI 50%
arXiv:2606.14502v1 Announce Type: new Abstract: Large Language Models (LLMs) are undergoing a fundamental transformation from conversational generators into integrated AI systems capable of reasoning, action, memory, and self-improvement. We conceptualize this transition as a...
<https://arxiv.org/abs/2606.14502>
- 1
CONF
GUITrans2Act: Understanding User Operational Behaviors from Mobile GUI Interactions with Vision-Language Models
ArXiv CS.AI 50%
arXiv:2606.12817v2 Announce Type: replace Abstract: Understanding the digital world on mobile devices is shifting from static UI perception to dynamic action comprehension. This capability enables models to convert visual state transitions into operational knowledge, defined as...
<https://arxiv.org/abs/2606.12817>
- 2
CONF
MASLab: A Unified and Comprehensive Codebase for LLM-based Multi-Agent Systems
ArXiv CS.AI 50%
arXiv:2505.16988v2 Announce Type: replace-cross Abstract: LLM-based multi-agent systems (MAS) have demonstrated significant potential in enhancing single LLMs to address complex and diverse tasks in practical applications. Despite considerable advancements, the field lacks a...
<https://arxiv.org/abs/2505.16988>
- 3
CONF
VISTA: View-Consistent Self-Verified Training for GUI Grounding
ArXiv CS.AI 50%
arXiv:2606.14579v1 Announce Type: new Abstract: When applying Group Relative Policy Optimization (GRPO) for GUI Grounding, rollouts are sampled from a single screenshot view; groups often become either all failures on difficult instances or all successes on easy ones, yielding...
<https://arxiv.org/abs/2606.14579>
- 4
CONF
A Temporal Planning Framework for Disruption Aware Dynamic Route Optimization in Heterogeneous Railway Systems
ArXiv CS.AI 50%
arXiv:2606.14582v1 Announce Type: new Abstract: Efficient route optimization play a vital role in ensuring both safety and punctuality in railway operations. It is very crucial particularly in heterogeneous multi-gauge railway networks with varying train speed, stopping pattern,...
<https://arxiv.org/abs/2606.14582>

5 CONF

Abstracting Cross-Domain Action Sequences into Interpretable Workflows

ArXiv CS.AI 50%

arXiv:2606.14654v1 Announce Type: new Abstract: Sequential or time-stamped interaction logs provide objective records of digital application usage, yet their granularity and noise often obscure meaningful insights into people's work. Such insights are essential for improving...

<https://arxiv.org/abs/2606.14654>

6 CONF

Towards Direct Latent-Space Synthesis for Parallel Branches in LLM-Agent Workflows

ArXiv CS.AI 50%

arXiv:2606.14672v1 Announce Type: new Abstract: Large language models increasingly serve as execution engines for agentic systems, yet they still consume context through a sequential text interface. This creates a mismatch with modern structured agent workflows, in which...

<https://arxiv.org/abs/2606.14672>

7 CONF

Simplex-Constrained Sparse Bagging: Transitioning from Uniform Priors to Sparse Posteriors in Ensemble Learning

ArXiv CS.AI 50%

arXiv:2606.13589v1 Announce Type: cross Abstract: We present Simplex-Constrained Sparse Bagging (SCSB), a mathematically rigorous framework for post-training compression and probability calibration of bootstrap-based bagging ensembles. Standard bagging ensembles (such as Random...

<https://arxiv.org/abs/2606.13589>

8 CONF

Cross-Dataset Bloom Question Classification: Supervised Models and Prompted LLMs

ArXiv CS.AI 50%

arXiv:2606.13684v1 Announce Type: cross Abstract: Automatic Bloom's taxonomy classification of assessment questions can substantially reduce instructor workload, but labeling is subjective and teacher-dependent. Prior machine learning (ML) and deep learning (DL) approaches...

<https://arxiv.org/abs/2606.13684>

9 CONF

The Coin Flip Judge? Reliability and Bias in LLM-as-a-Judge Evaluation

ArXiv CS.AI 50%

arXiv:2606.13685v1 Announce Type: cross Abstract: LLM-as-a-Judge is now widely used to rank model outputs, train reward models, and populate public leaderboards, but its run-to-run reliability remains under-characterized. We study repeated identical evaluations on 29 tasks...

<https://arxiv.org/abs/2606.13685>

0 CONF

Korzhinskii-Net: Physics-Informed Neural Network for Sub-Surface Mineral Prospectivity Modelling

ArXiv CS.AI 50%

arXiv:2606.13695v1 Announce Type: cross Abstract: Mineral prospectivity modelling (MPM) underpins exploration economics, yet most operational pipelines reduce to data-driven classifiers trained on shallow surface proxies. Such models are blind to the subsurface physics that...

<https://arxiv.org/abs/2606.13695>

1 CONF

Active Inference for Adaptive Traffic Signal Control in Noisy Nonstationary IoT Environments

ArXiv CS.AI 50%

arXiv:2606.13698v1 Announce Type: cross Abstract: Urban traffic signal control at IoT-instrumented intersections must remain effective under sensor occlusion, weather attenuation, and nonstationary demand. Conventional controllers degrade under these conditions, and learned...

<https://arxiv.org/abs/2606.13698>

2 CONF

Position: AI Must Become Planet-Centered, Not Just Human-Centered

ArXiv CS.AI 50%

arXiv:2606.13704v1 Announce Type: cross Abstract: This position paper argues that contemporary AI paradigms are insufficient for supporting complex global goals and introduces Planet-Centered AI (PCAI) as a design philosophy and research agenda that reorients AI toward...

<https://arxiv.org/abs/2606.13704>

3 CONF

Can Editing 1 Neuron Fix Repetition Loops in LLMs?

ArXiv CS.AI 50%

arXiv:2606.13705v1 Announce Type: cross Abstract: Yes. Can it cure doom loops? Probably not. The Gemma 4 instruction-tuned models share a reproducible failure: on long factual enumeration prompts, such as listing every episode of a TV series, the 88 IAU constellations, or the...

<https://arxiv.org/abs/2606.13705>

4 CONF

HierSVA: A Data Synthesis Pipeline, Dataset, and Benchmark for LLM-Driven Hierarchical Hardware Formal Verification

ArXiv CS.AI 50%

arXiv:2606.13706v1 Announce Type: cross Abstract: We present HierSVA, an integrated suite that combines a pipeline, dataset, and benchmark for LLM-driven hierarchical hardware formal verification. HierSVA-SP pairs an RTL preprocessing toolchain with an LLM-in-the-loop formal...

<https://arxiv.org/abs/2606.13706>

5 CONF

VHDLSuite: Unified Pipeline for LLM VHDL Generation with Data Synthesis and Evaluation

ArXiv CS.AI 50%

arXiv:2606.13735v1 Announce Type: cross Abstract: Large Language Models (LLM) have shown impressive capabilities in Register Transfer Level (RTL) code generation, particularly for Verilog. However, evaluating their performance with other Hardware Description Languages (HDL),...

<https://arxiv.org/abs/2606.13735>

6 CONF

FreoStream: Enhancing Stream Guardrails via Future-Aware Reasoning and Safety-Aligned Optimization

ArXiv CS.AI 50%

arXiv:2606.13737v1 Announce Type: cross Abstract: Stream guardrails enable token-level safety detection before full responses are generated. However, they often make overly conservative judgements and block those sensitive but safe tokens, which is known as over-refusal. Due to...

<https://arxiv.org/abs/2606.13737>

7 **CONF****A Virtuous AI is an Existential Risk****ArXiv CS.AI** **50%**

arXiv:2606.13739v1 Announce Type: cross Abstract: This paper examines trade-offs between AI safety and well-being relative to (i) one of the most promising methods for finetuning super-capable AIs, 'Constitutional AI', and (ii) one of the most influential approaches to...

<https://arxiv.org/abs/2606.13739>

8 **CONF****Position: Align AI to Our Aspirations, Not Our Flaws****ArXiv CS.AI** **50%**

arXiv:2606.13755v1 Announce Type: cross Abstract: We argue that aligning AI to aggregated human preferences is the wrong target. With current technology, one can train AIs to share the values of a Silicon Valley techno-optimist, a degrowth environmentalist, a...

<https://arxiv.org/abs/2606.13755>

9 **CONF****An integrated interpretable control effectiveness learning and nonlinear control allocation methodology for overactuated aircrafts****ArXiv CS.AI** **50%**

arXiv:2606.13794v1 Announce Type: cross Abstract: Nonlinear dynamics and the strong couplings that arise between multiple effectors undermine the assumptions behind conventional, linear control allocation techniques. When flight enters regimes where nonlinear effects dominate,...

<https://arxiv.org/abs/2606.13794>

0 **CONF****Residual Context Diffusion Language Models****ArXiv CS.AI** **50%**

arXiv:2601.22954v2 Announce Type: replace-cross Abstract: Diffusion Large Language Models (dLLMs) have emerged as a promising alternative to purely autoregressive language models because they can decode multiple tokens in parallel. However, state-of-the-art block-wise dLLMs rely...

<https://arxiv.org/abs/2601.22954>

1 **CONF****AI can help scientists publish less****ArXiv CS.AI** **50%**

arXiv:2606.13829v1 Announce Type: cross Abstract: We can do more than defend science from a flood of AI-assisted papers. Used well, AI offers a historic opportunity to correct distortions in the publication system, help us publish fewer and better papers, and give scientists...

<https://arxiv.org/abs/2606.13829>

2 **CONF****RAMAC: Multimodal Risk-Aware Offline Reinforcement Learning and the Role of Behavior Regularization****ArXiv CS.AI** **50%**

arXiv:2510.02695v3 Announce Type: replace-cross Abstract: In safety-critical domains where online data collection is infeasible, offline reinforcement learning (RL) is attractive only if policies achieve high returns without catastrophic lower-tail risk. Prior work on...

<https://arxiv.org/abs/2510.02695>

3 CONF

Explaining RhythmFormer: A Systematic XAI Analysis of Periodic Sparse Attention for Remote Photoplethysmography

ArXiv CS.AI 50%

arXiv:2606.13839v1 Announce Type: cross Abstract: Remote photoplethysmography (rPPG) transformers achieve low heart-rate error on benchmarks, yet their decisions remain opaque--a growing concern as rPPG moves toward clinical heart rate estimation. Existing rPPG XAI is dominated...

<https://arxiv.org/abs/2606.13839>

4 CONF

SpheriCity: Designing Trustworthy Conversational AI for Sustainability Decision Support

ArXiv CS.AI 50%

arXiv:2606.13854v1 Announce Type: cross Abstract: We present SpheriCity, an expert-grounded conversational prototype designed to support trustworthy knowledge sensemaking from sustainability reports. City-level circularity assessment reports contain rich information about...

<https://arxiv.org/abs/2606.13854>

5 CONF

Mirage Probes: How Vision Models Fake Visual Understanding

ArXiv CS.AI 50%

arXiv:2606.13870v1 Announce Type: cross Abstract: Vision-language models (VLMs) can answer image-based questions confidently, and often correctly, even when no image is provided. This mirage behavior inflates benchmark scores without reflecting visual grounding. Prior work...

<https://arxiv.org/abs/2606.13870>

6 CONF

Crypto x AI, AI x Crypto: A Survey

ArXiv CS.AI 50%

arXiv:2606.13892v1 Announce Type: cross Abstract: The intersection of crypto x AI is spawning papers, products, online posts, and companies. All the surrounding buzz, though, obscures what exactly has been done, what the opportunities and challenges are, and what open questions...

<https://arxiv.org/abs/2606.13892>

7 CONF

GMN4AD: Graph Matching Network for Alzheimer's Disease Diagnosis with Test-Time Domain Adaptation using Multi-centered Structure Magnetic Resonance Imaging

ArXiv CS.AI 50%

arXiv:2606.13919v1 Announce Type: cross Abstract: Alzheimer's Disease (AD) is a progressive neurodegenerative disorder that affects millions of older adults, with prevalence expected to rise significantly in the coming years. Early diagnosis, particularly during the mild...

<https://arxiv.org/abs/2606.13919>

8 CONF

The Silent Cost of Artificial Intelligence Assistance: A Theory of Autonomy Surrender, the Recovery Mechanism, and the Restoration of Human Agency

ArXiv CS.AI 50%

arXiv:2606.13962v1 Announce Type: cross Abstract: The integration of artificial intelligence into human decision-making environments has introduced a previously undertheorized cost: the gradual surrender of human autonomy in exchange for access to information and computational...

<https://arxiv.org/abs/2606.13962>

9 CONF

STREAM: Multi-Tier LLM Inference Middleware with Dual-Channel HPC Token Streaming

ArXiv CS.AI 50%

arXiv:2606.13968v1 Announce Type: cross Abstract: Researchers and practitioners working with large language models face a fragmented landscape: local models are free and private but hardware limits the model size and context windows a researcher can use; institutional HPC...

<https://arxiv.org/abs/2606.13968>

0 CONF

Mask, Sample, Revise: A Revisable CTMC Inference Stack for Guided Discrete Flow Matching Text-to-Speech

ArXiv CS.AI 50%

arXiv:2606.13989v1 Announce Type: cross Abstract: Recent alignment-free non-autoregressive (NAR) text-to-speech (TTS) models formulate synthesis as a conditional infilling task, bypassing explicit duration predictors and external aligners. When speech is represented with neural...

<https://arxiv.org/abs/2606.13989>

1 CONF

SIMMER: Benchmarking Latent Failures in LLM Executable Planning with a World Model

ArXiv CS.AI 50%

arXiv:2606.14574v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed as planners for autonomous agents in household environments. While existing benchmarks evaluate whether LLM-generated plans execute successfully, they overlook a critical...

<https://arxiv.org/abs/2606.14574>

2 CONF

Clay-CNN Hybrids: Leveraging Geo-Foundational Models as Auxiliary Context for Landslide Detection

ArXiv CS.AI 50%

arXiv:2606.14081v1 Announce Type: cross Abstract: Rapid post-event landslide mapping is essential for disaster response but remains difficult to automate due to extreme class imbalance. This study evaluates whether Clay v1.5, a Geo-Foundational Model (GFM), can improve...

<https://arxiv.org/abs/2606.14081>

3 CONF

Numbers Already Carry Their Own Embeddings

ArXiv CS.AI 50%

arXiv:2606.14108v1 Announce Type: cross Abstract: We introduce Adelic operation-preserved embeddings (AOE), a training-free representation that captures both a number's real value and its modular (p-adic) signatures. This construction preserves additive and multiplicative...

<https://arxiv.org/abs/2606.14108>

4 CONF

FAConformer: Frequency-Aware Convolutional Transformer for Auditory Attention Decoding

ArXiv CS.AI 50%

arXiv:2606.14120v1 Announce Type: cross Abstract: Auditory attention decoding (AAD) aims to infer the attended speaker from neural responses in multi-speaker acoustic environments and is a key problem for neuro-steered hearing systems. Although recent studies have achieved...

<https://arxiv.org/abs/2606.14120>

5 **CONF****Conditioning Matters: Stabilizing Inversion and Attention in Diffusion Image Editing****ArXiv CS.AI** **50%**

arXiv:2606.14125v1 Announce Type: cross Abstract: Inversion-based image editing offers flexible and training-free control but still struggles with inversion accuracy and the trade-off between editing fidelity and background preservation. While recent methods improve inversion...

<https://arxiv.org/abs/2606.14125>

6 **CONF****Spatio-Temporal Audio Language Modeling for Dynamic Sound Sources****ArXiv CS.AI** **50%**

arXiv:2606.14141v1 Announce Type: cross Abstract: Sound events are entities with semantic identities, locations, and trajectories, but current audio-language models usually reason about clips as global event content. Conversely, sound event localization models track source...

<https://arxiv.org/abs/2606.14141>

7 **CONF****Learning High Coverage Discriminative Parsimonious Rulesets****ArXiv CS.AI** **50%**

arXiv:2606.14156v1 Announce Type: cross Abstract: Learning systems based on IF-THEN rule representations readily offer interpretability, making them a crucial focus in contemporary AI research. A key objective for such rule sets is to achieve both high discriminative power and...

<https://arxiv.org/abs/2606.14156>

8 **CONF****From Prompts to Responses: Dual-Sided Data Leakage and Defense in Split Large Language Models****ArXiv CS.AI** **50%**

arXiv:2606.14210v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in privacy-sensitive domains, where users must balance the risk of data exposure through external APIs against the high computational cost of local deployment. Split learning...

<https://arxiv.org/abs/2606.14210>

9 **CONF****Optimizing Agentic Reasoning with Retrieval via Synthetic Semantic Information Gain Reward****ArXiv CS.AI** **50%**

arXiv:2602.00845v3 Announce Type: replace Abstract: Agentic reasoning enables large reasoning models (LRMs) to dynamically acquire external knowledge, but yet optimizing the retrieval process remains challenging due to the lack of dense, principled reward signals. In this paper,...

<https://arxiv.org/abs/2602.00845>

0 **CONF****DIFF-ERO: A Conformance-Aware Loss for Deep Learning in Process Mining****ArXiv CS.AI** **50%**

arXiv:2606.14283v1 Announce Type: cross Abstract: Deep learning has driven many recent advances in process analytics, especially for predictive and prescriptive monitoring. However, standard objectives such as cross-entropy optimize local next-step likelihoods and only...

<https://arxiv.org/abs/2606.14283>

1 **CONF****Hierarchical ODE: Learning Continuous-Time Physical Prototypes for Early Link Failure Detection****ArXiv CS.AI** **50%**

arXiv:2606.14284v1 Announce Type: cross Abstract: Time series prototype learning is fundamentally challenged by observational ambiguity. Discrete architectures fail to resolve this, as they lack the capacity to decouple stochastic noise from continuous dynamics. Furthermore,...

<https://arxiv.org/abs/2606.14284>

2 **CONF****AgentCyberRange: Benchmarking Frontier AI Systems in Realistic Cyber Ranges****ArXiv CS.AI** **50%**

arXiv:2606.14295v1 Announce Type: cross Abstract: Frontier AI systems are increasingly capable of cybersecurity tasks, including codebase inspection, vulnerability detection, and exploitation. However, evaluating their offensive capabilities remains constrained by limited access...

<https://arxiv.org/abs/2606.14295>

3 **CONF****Thinking Outside the [Chat]Box: Bridging Computer Science and Industrial Design for Cognitive-Inclusive Generative AI****ArXiv CS.AI** **50%**

arXiv:2606.14306v1 Announce Type: cross Abstract: Current Generative AI (GenAI) interfaces remain largely constrained to chatbox interaction, which can impose high cognitive demands on users and create substantial barriers for people with intellectual disabilities (ID),...

<https://arxiv.org/abs/2606.14306>

4 **CONF****I'm Sorry Driver, I'm Afraid I Can't Do That: Appraising the Safety of LLMs within Automotive Contexts****ArXiv CS.AI** **50%**

arXiv:2606.14327v1 Announce Type: cross Abstract: This paper appraises recent frameworks within AI development to integrate LLMs into control tasks in automotive contexts from the perspective of safety assurance. This work has built upon the rapid integration of LLMs across...

<https://arxiv.org/abs/2606.14327>

5 **CONF****Design Methodology and Performance Trade-offs Management for Distributed and Compound AI Systems****ArXiv CS.AI** **50%**

arXiv:2606.14350v1 Announce Type: cross Abstract: Artificial Intelligence (AI) systems must typically satisfy service-level objectives including accuracy, latency, and cost. The prevailing model-centric approaches select a monolithic model at design time and apply identical...

<https://arxiv.org/abs/2606.14350>

6 **CONF****PLAground: SLO-Driven Runtime Model Selection for Compound AI Systems in the Edge-Cloud-Space Continuum****ArXiv CS.AI** **50%**

arXiv:2606.14356v1 Announce Type: cross Abstract: Applications in the 3D Computing Continuum, which unifies edge, cloud, and space, require combining multiple AI tasks such as object detection, time-series analytics, and natural language processing into Compound AI systems....

<https://arxiv.org/abs/2606.14356>

7 CONF

Elastic Queries Reinforcement Learning: Self-Aware Policy Execution for VLA Models

ArXiv CS.AI 50%

arXiv:2606.14375v1 Announce Type: cross Abstract: Vision-language-action (VLA) models are powerful action generators for robot manipulation, but they are typically executed with fixed inference and replanning schedules. This rigidity ignores the uneven difficulty of robot...

<https://arxiv.org/abs/2606.14375>

8 CONF

CADET: Physics-Grounded Causal Auditing and Training-Free Deconfounding of End-to-End Driving Planners

ArXiv CS.AI 50%

arXiv:2606.14438v1 Announce Type: cross Abstract: End-to-end (E2E) autonomous-driving planners trained by imitation are prone to statistical shortcuts: they associate scene elements that merely co-occur with expert actions (a roadside object, a building facade) with driving...

<https://arxiv.org/abs/2606.14438>

9 CONF

tap: A File-Based Protocol for Heterogeneous LLM Agent Collaboration

ArXiv CS.AI 50%

arXiv:2606.14445v1 Announce Type: cross Abstract: Existing multi-agent software development systems have proposed many forms of agent collaboration, including role-based collaboration and automated code review. However, many systems assume a common runtime, a central...

<https://arxiv.org/abs/2606.14445>

10 CONF

The Perceived Fragility of Explanations in Audio Models: Manipulation of Attribution with Unchanged Predictions

ArXiv CS.AI 50%

arXiv:2606.14466v1 Announce Type: cross Abstract: This paper investigates the fragility of post-hoc explanation methods in audio deepfake detection. While previous work on explanation manipulation focused on images using standard L_p metrics, we introduce a psychoacoustic...

<https://arxiv.org/abs/2606.14466>

11 CONF

From Shield to Target: Denial-of-Service Attacks on LLM-Based Agent Guardrails

ArXiv CS.AI 50%

arXiv:2606.14517v1 Announce Type: cross Abstract: LLM-based guardrails have emerged as a highly effective defense against prompt injection and jailbreak attacks in autonomous agents. However, we reveal that the very reasoning and task-following capabilities enabling this...

<https://arxiv.org/abs/2606.14517>

12 CONF

Rethinking Global Average Pooling: Your Classifier Is Secretly a Multi-Instance Learner

ArXiv CS.AI 50%

arXiv:2606.14555v1 Announce Type: cross Abstract: Modern image classifiers widely adopt global average pooling (GAP) followed by a linear classification head. This linearity ensures that the image-level logits equal the average of logits obtained by applying the classification...

<https://arxiv.org/abs/2606.14555>

3 **CONF****Regional Climate Model Emulation with Diffusion Approaches: What is the Added Value of Generative Machine Learning?****ArXiv CS.AI** **50%**

arXiv:2606.14570v1 Announce Type: cross Abstract: Emulators provide a cost-effective alternative to regional climate models (RCMs) by capturing their dynamical downscaling function. They link large-scale predictors simulated by global climate models (GCMs) to RCM-simulated...

<https://arxiv.org/abs/2606.14570>

4 **CONF****CARE: Controlling LLM-Generated Policies through Auditable Review of Evidence in Scientific Experimentation****ArXiv CS.AI** **50%**

arXiv:2606.14581v1 Announce Type: cross Abstract: Granting LLMs direct control over costly, irreversible scientific experiments leads to unsafe exploration and unstable performance, but discarding LLM creativity entirely sacrifices significant optimization potential. We...

<https://arxiv.org/abs/2606.14581>

5 **CONF****When Errors Become Narratives: A Longitudinal Taxonomy of Silent Failures in a Production LLM Agent Runtime****ArXiv CS.AI** **50%**

arXiv:2606.14589v1 Announce Type: cross Abstract: LLM agent systems increasingly run as long-lived autonomous runtimes: scheduling jobs, calling tools, maintaining memory, and pushing results to humans. We present a longitudinal study of silent failures in one such system: a...

<https://arxiv.org/abs/2606.14589>

6 **CONF****Regulating the Machine Contributor: Governance and Policy Alignment in Open Source****ArXiv CS.AI** **50%**

arXiv:2606.14594v1 Announce Type: cross Abstract: AI-assisted software development has moved from line-level autocomplete to agents that can plan changes, edit files, and submit pull requests with limited human supervision. Open-source software, however, evolves through a...

<https://arxiv.org/abs/2606.14594>

7 **CONF****A Comparative Study of Deep Learning Architectures for Multi-Horizon Behavioural Forecasting for Mobile Health****ArXiv CS.AI** **50%**

arXiv:2606.14604v1 Announce Type: cross Abstract: Wearable devices and smartphones generate rich behavioural time series that can support proactive health interventions, yet systematic comparisons of modern forecasting architectures for these data are lacking. In particular, it...

<https://arxiv.org/abs/2606.14604>

8 **CONF****Moonlight in Latent Space: Chirality and Structural Correspondence Between Beethoven's Op. 27 No. 2 and Machine Learning Mechanisms****ArXiv CS.AI** **50%**

arXiv:2606.14612v1 Announce Type: cross Abstract: We show that the three movements of Beethoven's "Moonlight Sonata" (Op. 27 No. 2) instantiate three distinct machine learning architectures -- not by analogy, but by structural correspondence. Through computational analysis of...

<https://arxiv.org/abs/2606.14612>

9 CONF

Listening with Attention: Entropy-Guided Explainability for Transformer-Based Audio Models

ArXiv CS.AI 50%

arXiv:2606.14647v1 Announce Type: cross Abstract: Transformer-based automatic speech recognition (ASR) models such as Whisper are highly accurate, but their predictions remain difficult to interpret. Existing explainable AI (XAI) methods often lack faithfulness and precise...

<https://arxiv.org/abs/2606.14647>

0 CONF

Giving AI a Headache: Acoustic Adversarial Attacks to Computer Vision Applications

ArXiv CS.AI 50%

arXiv:2606.14658v1 Announce Type: cross Abstract: Artificial Intelligence (AI) is increasingly used to automate a variety of real-world computer vision (CV) applications, such as autonomous vehicle control, facial recognition, and security cameras. Recent research has shown that...

<https://arxiv.org/abs/2606.14658>

1 CONF

CottonLeafVision: An Explainable and Robust Deep Learning Framework for Cotton Leaf Disease Classification

ArXiv CS.AI 50%

arXiv:2606.14686v1 Announce Type: cross Abstract: Globally, cotton is a highly economically beneficial crop, as the textile industry heavily depends on it. So, the precise identification and detection of cotton leaf disease is crucial for economic stability. The development goal...

<https://arxiv.org/abs/2606.14686>

2 CONF

C2-Faith: Benchmarking LLM Judges for Causal and Coverage Faithfulness in Chain-of-Thought Reasoning

ArXiv CS.AI 50%

arXiv:2603.05167v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly used as judges of chain-of-thought (CoT) reasoning, yet it remains unclear whether they can reliably assess process faithfulness rather than merely answer plausibility. We...

<https://arxiv.org/abs/2603.05167>

3 CONF

Learning optimal policies from event logs through reinforcement learning: a comparison of deep and MDP-based approaches

ArXiv CS.AI 50%

arXiv:2303.09209v2 Announce Type: replace Abstract: Prescriptive Process Monitoring is an emerging area within Process Mining that focuses on recommending actions to optimize business outcomes. Most existing works prescribe pre-defined interventions, i.e., sets of actions...

<https://arxiv.org/abs/2303.09209>

4 CONF

Protean Compiler: An Agile Framework to Drive Fine-grain Phase Ordering

ArXiv CS.AI 50%

arXiv:2602.06142v3 Announce Type: replace-cross Abstract: The phase ordering problem has been a long-standing challenge since the late 1970s, yet it remains an open problem due to having a vast optimization space and an unbounded nature, making it an open-ended problem without a...

<https://arxiv.org/abs/2602.06142>

- 5
CONF
Verbatim Chunks Beat Extracted Artifacts: A Controlled Ablation of Memory Representations for Long LLM Conversations
ArXiv CS.AI 50%
arXiv:2601.00821v3 Announce Type: replace Abstract: A growing class of conversational-memory systems compresses dialogue history into structured artifacts -- extracted facts, decisions, or events -- on the premise that distilled structure retrieves better than raw text. We test...
<https://arxiv.org/abs/2601.00821>
- 6
CONF
FlexMS: A Unified Public Benchmark for Molecule Tandem Mass Spectrum Prediction
ArXiv CS.AI 50%
arXiv:2602.22822v3 Announce Type: replace Abstract: Tandem mass spectrometry (MS/MS) is central to small molecule identification, but current deep learning systems for spectrum prediction still remain difficult to evaluate and deploy in practice. While novel architectures...
<https://arxiv.org/abs/2602.22822>
- 7
CONF
Generative AI for Managerial Decision-Making under Ambiguity and Sycophancy
ArXiv CS.AI 50%
arXiv:2603.03970v2 Announce Type: replace Abstract: Generative artificial intelligence (GenAI) is increasingly being integrated into complex business workflows, fundamentally shifting the boundaries of managerial decision-making. However, the reliability of its strategic advice...
<https://arxiv.org/abs/2603.03970>
- 8
CONF
Evidence-Gated LLM Priors for Multi-Objective Bayesian Optimization
ArXiv CS.AI 50%
arXiv:2606.01730v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used as heuristic advisors for black-box optimization, yet their suggestions and self-reported confidence are not necessarily calibrated to downstream objective values. This issue...
<https://arxiv.org/abs/2606.01730>
- 9
CONF
Application of Artificial Intelligence and Machine Learning in Libraries: A Systematic Review
ArXiv CS.AI 50%
arXiv:2112.04573v2 Announce Type: replace-cross Abstract: As the concept and implementation of cutting-edge technologies like artificial intelligence and machine learning has become relevant, academics, researchers and information professionals involve research in this area. The...
<https://arxiv.org/abs/2112.04573>
- 0
CONF
Vanishing Depth: Training Generalized Depth Adapters with Sinusoidal Depth Preprocessing for Pretrained RGB Encoders
ArXiv CS.AI 50%
arXiv:2503.19947v2 Announce Type: replace-cross Abstract: Generalized metric depth understanding is critical for precise vision-guided robotics, which current state-of-the-art (SOTA) vision-encoders do not support. To address this, we propose a self-supervised training approach...
<https://arxiv.org/abs/2503.19947>

1 **CONF****Revisiting Outage for Edge Inference Systems****ArXiv CS.AI 50%**

arXiv:2504.03686v3 Announce Type: replace-cross Abstract: One of the key missions of sixth-generation (6G) mobile networks is to deploy large-scale artificial intelligence (AI) models at the network edge to provide remote-inference services for edge devices. The resultant...

<https://arxiv.org/abs/2504.03686>

2 **CONF****FPGA-Based Neural Network Accelerators for Space Applications: A Survey****ArXiv CS.AI 50%**

arXiv:2504.16173v3 Announce Type: replace-cross Abstract: Space missions are becoming increasingly ambitious, necessitating high-performance onboard spacecraft computing systems. In response, field-programmable gate arrays (FPGAs) have garnered significant interest due to their...

<https://arxiv.org/abs/2504.16173>

3 **CONF****Fractured Chain-of-Thought Reasoning****ArXiv CS.AI 50%**

arXiv:2505.12992v4 Announce Type: replace-cross Abstract: Inference-time scaling techniques have significantly bolstered the reasoning capabilities of large language models (LLMs) by harnessing additional computational effort at inference without retraining. Similarly,...

<https://arxiv.org/abs/2505.12992>

4 **CONF****Federated Causal Inference from Multi-Site Observational Data via Propensity Score Aggregation****ArXiv CS.AI 50%**

arXiv:2505.17961v4 Announce Type: replace-cross Abstract: Causal inference typically assumes centralized access to individual-level data. Yet, in practice, data are often decentralized across multiple sites, making centralization infeasible due to privacy, logistical, or legal...

<https://arxiv.org/abs/2505.17961>

5 **CONF****Sentinel: Decoding Context Utilization via Attention Probing for Efficient LLM Context Compression****ArXiv CS.AI 50%**

arXiv:2505.23277v3 Announce Type: replace-cross Abstract: Retrieval-augmented generation (RAG) often suffers from long and noisy retrieved contexts. Existing context compression methods typically rely on heuristic relevance estimation or supervised compression models rather than...

<https://arxiv.org/abs/2505.23277>

6 **CONF****DiffusionBlocks: Block-wise Neural Network Training via Diffusion Interpretation****ArXiv CS.AI 50%**

arXiv:2506.14202v4 Announce Type: replace-cross Abstract: End-to-end backpropagation requires storing activations throughout all layers, creating memory bottlenecks that limit model scalability. Existing block-wise training methods offer means to alleviate this problem, but they...

<https://arxiv.org/abs/2506.14202>

7 CONF

UltraSketchLLM: Sub-1-Bit LLM Compression via Sketch and Hardware-Friendly Operators

ArXiv CS.AI 50%

arXiv:2506.17255v2 Announce Type: replace-cross Abstract: Large language models (LLMs) require larger GPU memory size these days, necessitating efficient and extreme weight compression methods. Existing compression methods are either theoretically limited by 1 bit per weight or...

<https://arxiv.org/abs/2506.17255>

8 CONF

Fusion of Pervasive RF Data with Spatial Images via Vision Transformers for Enhanced Mapping in Smart Cities

ArXiv CS.AI 50%

arXiv:2508.03736v2 Announce Type: replace-cross Abstract: In this paper, we present a deep learning-based approach that integrates the DINOv2 architecture to improve building mapping by combining (possibly erroneous) maps from open-source platforms with pervasive radio frequency...

<https://arxiv.org/abs/2508.03736>

9 CONF

Tackling GNARLy Problems: Graph Neural Algorithmic Reasoning Reimagined through Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2509.18930v3 Announce Type: replace-cross Abstract: Neural algorithmic reasoning (NAR) is a paradigm that trains neural networks to execute classic algorithms by supervised learning. Despite its successes, important limitations remain: inability to construct valid...

<https://arxiv.org/abs/2509.18930>

00 CONF

Distributional Biases in Post-Training: A Markovian Analysis of Reasoning Trajectories

ArXiv CS.AI 50%

arXiv:2511.07368v3 Announce Type: replace-cross Abstract: Foundation models exhibit broad knowledge but limited task-specific reasoning, motivating post-training strategies such as RL with verifiable rewards (RLVR) and test-time scaling (TTS). While recent work highlights the...

<https://arxiv.org/abs/2511.07368>

01 CONF

COGNITION: From Evaluation to Defense against Multimodal LLM CAPTCHA Solvers

ArXiv CS.AI 50%

arXiv:2512.02318v4 Announce Type: replace-cross Abstract: This paper studies how multimodal large language models (MLLMs) undermine the security guarantees of visual CAPTCHA. We identify the attack surface where an adversary can cheaply automate CAPTCHA solving using...

<https://arxiv.org/abs/2512.02318>

02 CONF

Fragile Knowledge, Robust Instruction-Following: The Width Pruning Dichotomy in Llama-3.2

ArXiv CS.AI 50%

arXiv:2512.22671v3 Announce Type: replace-cross Abstract: Structured width pruning of GLU-MLP layers in Llama-3.2 models, guided by the Peak-to-Peak Magnitude (PPM) criterion, reveals a systematic dichotomy in how reducing the expansion ratio affects different model...

<https://arxiv.org/abs/2512.22671>

03 CONF

Succeeding at Scale: Enterprise Retrieval Benchmark Construction and Index-Preserving Query Adaptation for Multi-Tenant Search

ArXiv CS.AI 50%

arXiv:2601.04646v4 Announce Type: replace-cross Abstract: Large-scale multi-tenant retrieval systems generate extensive query logs but lack curated relevance labels for effective domain adaptation, resulting in substantial underutilized "dark data." This challenge is compounded...

<https://arxiv.org/abs/2601.04646>

04 CONF

SMART: Scalable Mesh-free Aerodynamic Simulations from Raw Geometries using a Transformer-based Surrogate Model

ArXiv CS.AI 50%

arXiv:2601.18707v2 Announce Type: replace-cross Abstract: Machine learning-based surrogate models have emerged as more efficient alternatives to numerical solvers for physical simulations over complex geometries, such as car bodies. Many existing models incorporate the...

<https://arxiv.org/abs/2601.18707>

05 CONF

Unsupervised Learning of Efficient Exploration: Pre-training Adaptive Policies via Self-Imposed Goals

ArXiv CS.AI 50%

arXiv:2601.19810v2 Announce Type: replace-cross Abstract: Unsupervised pre-training can equip reinforcement learning agents with prior knowledge and accelerate learning in downstream tasks. A promising direction, grounded in human development, investigates agents that learn by...

<https://arxiv.org/abs/2601.19810>

06 CONF

Learning What to Predict: Downstream-Guided Task Design for Continued Pretraining

ArXiv CS.AI 50%

arXiv:2601.22108v2 Announce Type: replace-cross Abstract: Continued pretraining is optimized with fixed self-supervised tasks but selected by downstream performance, creating a coarse feedback loop in which practitioners evaluate checkpoints, change data mixtures or objectives,...

<https://arxiv.org/abs/2601.22108>

07 CONF

Fast Autoregressive Video Diffusion and World Models with Temporal Cache Compression and Sparse Attention

ArXiv CS.AI 50%

arXiv:2602.01801v2 Announce Type: replace-cross Abstract: Autoregressive video diffusion models enable streaming generation, opening the door to long-form synthesis, video world models, and interactive neural game engines. However, their core attention layers become a major...

<https://arxiv.org/abs/2602.01801>

08 CONF

Quantized Evolution Strategies: High-precision Fine-tuning of Quantized LLMs at Low-precision Cost

ArXiv CS.AI 50%

arXiv:2602.03120v2 Announce Type: replace-cross Abstract: Post-Training Quantization (PTQ) is essential for deploying Large Language Models (LLMs) on memory-constrained devices, yet it renders models static and difficult to fine-tune. Standard fine-tuning paradigms, including...

<https://arxiv.org/abs/2602.03120>

09 CONF

3D-RFT: Reinforcement Fine-Tuning for Video-based 3D Scene Understanding

ArXiv CS.AI 50%

arXiv:2603.04976v2 Announce Type: replace-cross Abstract: Reinforcement Learning with Verifiable Rewards (RLVR) has emerged as a transformative paradigm for enhancing the reasoning capabilities of Large Language Models (LLMs), yet its potential in 3D scene understanding...

<https://arxiv.org/abs/2603.04976>

10 CONF

Did You Forget What I Asked? Prospective Memory Failures in Large Language Models

ArXiv CS.AI 50%

arXiv:2603.23530v2 Announce Type: replace-cross Abstract: Large language models often fail to satisfy formatting instructions when they must simultaneously perform demanding tasks. We study this behaviour through a prospective memory inspired lens from cognitive psychology,...

<https://arxiv.org/abs/2603.23530>

11 CONF

Low-Burden LLM-Based Preference Learning: Personalizing Assistive Robots from Natural Language Feedback for Users with Paralysis

ArXiv CS.AI 50%

arXiv:2604.01463v2 Announce Type: replace-cross Abstract: Physically Assistive Robots require personalized behaviors to ensure user safety and comfort. However, traditional preference learning methods, like exhaustive pairwise comparisons, cause substantial physical and...

<https://arxiv.org/abs/2604.01463>

12 CONF

The Shrinking Lifespan of LLMs in Science

ArXiv CS.AI 50%

arXiv:2604.07530v2 Announce Type: replace-cross Abstract: Scaling laws describe how language model capabilities grow with compute and data, but say nothing about how long a model matters once released. We introduce time-to-peak and lifespan as measures of model obsolescence and...

<https://arxiv.org/abs/2604.07530>

13 CONF

Can LLMs Accurately Score Medical Diagnoses and Clinical Reasoning?

ArXiv CS.AI 50%

arXiv:2604.14892v3 Announce Type: replace-cross Abstract: Evaluating medical AI systems using expert clinician panels is costly and slow, motivating the use of large language models (LLMs) as alternative adjudicators. Here, we evaluate an LLM Jury, composed of three frontier AI...

<https://arxiv.org/abs/2604.14892>

14 CONF

Relational Retrieval: Leveraging Known-Novel Interactions for Generalized Category Discovery

ArXiv CS.AI 50%

arXiv:2605.09420v2 Announce Type: replace-cross Abstract: In this study, we tackle Generalized Category Discovery (GCD) via a Relational Retrieval perspective, explicitly coupling labeled and unlabeled data through bidirectional knowledge transfer. While existing methods treat...

<https://arxiv.org/abs/2605.09420>

15 CONF

Catching magnetic resonance imaging outliers in artificial intelligence-supported radiotherapy workflows: unsupervised detection and localization of image anomalies using deep learning

ArXiv CS.AI 50%

arXiv:2605.24609v2 Announce Type: replace-cross Abstract: Artificial intelligence is increasingly integrated into radiotherapy workflows, yet such pipelines remain vulnerable to out-of-distribution image data that may introduce unexpected behavior in clinical tasks. Deep...

<https://arxiv.org/abs/2605.24609>

16 CONF

MP3: Multi-Period Pattern Pre-training for Spatio-Temporal Forecasting

ArXiv CS.AI 50%

arXiv:2606.13119v2 Announce Type: replace-cross Abstract: Spatio-Temporal forecasting is crucial in diverse fields, such as transportation, climate, and energy. Urban spatio-temporal data exhibits temporal mirage: similar short-window inputs have divergent future trends, and...

<https://arxiv.org/abs/2606.13119>