



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	125	Media
TechCrunch AI	8	Media
HackerNews	5	Community
MIT Tech Review	1	Media
The Verge AI	1	Media

CONFIRMED SIGNALS — 140 stories

CONF

Show HN: Veterinarian turned founder, AI lawn diagnosis

HackerNews 50% HN 74

I know, it's kind of weird. What is a veterinarian doing creating an analysis tool for lawn problems? Frankly, the idea was born of my own lawn care struggles. Endless lawn care company fees without any actual improvement. Googling problems and finding generic solutions without...

<https://grassdx.com/>

CONF

Show HN: Memento - Self-hosted agentic search and LLM wiki over your email

HackerNews 50% HN 8

Our email inboxes carry multiple decades of messages (100K-500K). This is a good proxy for all the important things that happened in your life, the projects you have done and the people that you have connected with. With the chronological view of messages in the inbox, these...

<https://news.ycombinator.com/item?id=48557937>

CONF

Show HN: The Dictionary Game (Fictionary/Balderdash) as a Daily Puzzle

HackerNews 50% HN 5

I love the dictionary game: someone takes an unfamiliar word from the dictionary and writes down the definition. Then the other players come up with their own, fake, definition, and guess which one is real. I built a Wordle-style daily game out of this concept. You get shown a...

<https://dictionarypuzzle.com/>

CONF

Show HN: Brainfuck but with Turtle Graphics

HackerNews 50% HN 5

I made a Brainfuck interpreter (<https://en.wikipedia.org/wiki/Brainfuck>) with additional commands for turtle graphics. Check it out: <https://czterycztery.pl/programy/turtlefreeze/>

<https://czterycztery.pl/programy/turtlefreeze/>

CONF

Show HN: Pair your iPhone to your own Ollama over Tailscale with a QR scan

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/TeamDzX/myllm-connect>

CONF

Anthropic's latest feud with the Trump admin may actually help it, sales data suggests

TechCrunch AI 50%

Anthropic's popularity with business users is growing so well that the latest beef with the government might actually boost it, data from Ramp suggests.

<https://techcrunch.com/2026/06/16/anthropics-latest-feud-with-the-trump-admin-may-actuaa...>

CONF

Android 17 launches with new multitasking tools as Google expands Gemini features

TechCrunch AI 50%

Google has released Android 17 and Wear OS 7, introducing new multitasking features, parental controls, security tools, and smartwatch upgrades. The launch is also accompanied by a Pixel Drop that brings Google's latest AI models to its devices.

<https://techcrunch.com/2026/06/16/android-17-launches-with-new-multitasking-tools-as-go...>

CONF

Sixty percent of US consumers say 'AI' in brand messaging is a turnoff, survey finds

TechCrunch AI 50%

WordPress VIP's latest survey suggests consumers are wary of AI-generated answers even as companies increasingly view AI search as an important referral channel.

<https://techcrunch.com/2026/06/16/sixty-percent-of-u-s-consumers-say-ai-in-brand-messag...>

CONF

DOJ claims xAI's unpermitted gas turbines are a matter of 'national, economic, and energy security'

TechCrunch AI 50%

The Justice department says the Pentagon needs xAI to keep using its unpermitted gas turbines.

<https://techcrunch.com/2026/06/16/doj-claims-xais-unpermitted-gas-turbines-are-a-matter...>

0 CONF

Plaud says its software business topped \$100M in ARR after shipping over 2M AI notetakers

TechCrunch AI 50%

Plaud is trying to make a mark in a crowded market full of AI-powered meeting notetakers.

<https://techcrunch.com/2026/06/16/plaud-says-its-software-business-topped-100m-in-arr-a...>

1 CONF

Robinhood's note on 10% layoffs shows blaming AI isn't cutting it

TechCrunch AI 50%

Unlike many of his tech industry peers who have cut thousands of jobs citing the need to restructure to make the most of AI, Robinhood's CEO Vlad Tenev conspicuously made no mention of AI in his note about layoffs.

<https://techcrunch.com/2026/06/16/robinhoods-note-on-10-layoffs-shows-blaming-ai-isnt-c...>

2 CONF

Probably raises \$9M to build a more reliable kind of AI

TechCrunch AI 50%

Probably wants to prevent hallucinations and factual errors from reaching users, and achieve accuracy on par with deterministic systems.

<https://techcrunch.com/2026/06/16/probably-raises-9m-to-build-a-more-reliable-kind-of-ai/>

3 CONF

ChatGPT's market share slips below 50% for first time

TechCrunch AI 50%

The chatbot still remains the most popular AI assistant worldwide with over 1.1 billion monthly users, followed by Gemini with 662 million and Claude with 245 million.

<https://techcrunch.com/2026/06/16/chatgpts-market-share-slips-below-50-for-first-time/>

4 CONF

Exclusive eBook: How AI is becoming the next military advisor

MIT Tech Review 50%

A collection of stories about how militaries are using AI models to make decisions. This subscriber-only eBook is a package of six stories that were originally published in MIT Technology Review between April 11, 2025, and April 21, 2026, and have been updated to reflect recent...

<https://www.technologyreview.com/2026/06/16/1138905/exclusive-ebook-how-ai-is-becoming-...>

5 CONF

Apple 2027 rumors: AirPods with cameras for AI and the second folding iPhone

The Verge AI 50%

Now that we're clear of WWDC and all of the new AI-powered features coming to Apple's platforms, Bloomberg reporter Mark Gurman has more details about rumored new hardware, like the camera-equipped AirPods he'd previously written about. He says they are currently on schedule for...

<https://www.theverge.com/tech/950826/apple-airpod-camera-ai-foldable-iphone-rumor>

6 CONF

Can LLMs Be CEOs? Benchmarking Strategic Resource Reallocation with Multi-Role Agent Simulation

ArXiv CS.AI 50%

arXiv:2606.17459v1 Announce Type: new Abstract: Evaluating the decision-making capabilities of large language models (LLMs) is a growing research priority, yet existing benchmarks focus on isolated cognitive tasks such as reasoning, knowledge retrieval, and economic rationality...

<https://arxiv.org/abs/2606.17459>

7 CONF

Mordal: Automated Pretrained Model Selection for Vision Language Models

ArXiv CS.AI 50%

arXiv:2502.00241v2 Announce Type: replace-cross Abstract: Incorporating multiple modalities into large language models (LLMs) is a powerful way to enhance their understanding of non-textual data, enabling them to perform multimodal tasks. Vision language models (VLMs) form the...

<https://arxiv.org/abs/2502.00241>

8 CONF

Reinforcing Dual-Path Reasoning in Spatial Vision Language Models

ArXiv CS.AI 50%

arXiv:2606.17539v1 Announce Type: cross Abstract: Spatial VLMs have made substantial progress in geometric perception, yet complex spatial reasoning requiring multi-step inference over depth, distance, and scene relations remains challenging. Moreover, different spatial queries...

<https://arxiv.org/abs/2606.17539>

9 CONF

Quantifying Consistency in LLM Logical Reasoning via Structural Uncertainty

ArXiv CS.AI 50%

arXiv:2606.17312v1 Announce Type: new Abstract: Large language models can arrive at the same answer through reasoning paths that are unstable, contradictory, or difficult to rank consistently -- a failure mode especially prevalent in multi-step deductive reasoning. Existing...

<https://arxiv.org/abs/2606.17312>

0 CONF

SpeechDx: A Multi-Task Benchmark for Clinical Speech AI

ArXiv CS.AI 50%

arXiv:2606.17339v1 Announce Type: new Abstract: Speech offers a uniquely informative window into health by simultaneously engaging neurological, motor, respiratory, and vocal systems. Current clinical speech AI methods have largely progressed through isolated condition-specific...

<https://arxiv.org/abs/2606.17339>

1 CONF

Treatment Response Optimized Clinical Decision Support AI System via Digital Twin Simulation

ArXiv CS.AI 50%

arXiv:2606.17405v1 Announce Type: new Abstract: Clinical decision support AI systems (CDSASs) must adapt to evolving patient conditions in real-time while adhering to strict safety constraints. We present an online adaptive framework that integrates Treatment Effect (TE)...

<https://arxiv.org/abs/2606.17405>

- 2 **CONF**
- Incumbent Advantage: Brand Bias and Cognitive Manipulation Dynamics in LLM Recommendation Systems**
- ArXiv CS.AI 50%
- arXiv:2606.17443v1 Announce Type: new Abstract: Large language models (LLMs) are becoming a major way for consumers to find products, but we do not yet understand how brands compete in this new channel. We study brand dynamics in LLM recommendations using skincare products -- a...
- <https://arxiv.org/abs/2606.17443>
-
- 3 **CONF**
- MapSatisfyBench: Benchmarking Satisfaction-Aware Map Agents through Behavior-Grounded Implicit Decision Factors**
- ArXiv CS.AI 50%
- arXiv:2606.17453v1 Announce Type: new Abstract: Large language model agents are increasingly integrated into map services. Since map services are embedded in everyday-life scenarios rather than professional task settings, users often express their needs informally, resulting in...
- <https://arxiv.org/abs/2606.17453>
-
- 4 **CONF**
- LLM-as-Judge in Education: A Curriculum-Grounded Marking Pipeline**
- ArXiv CS.AI 50%
- arXiv:2606.17507v1 Announce Type: new Abstract: Generative AI and large language models (LLMs) are increasingly applied to question generation and automated assessment. However, deploying LLMs in preparation for high-stakes exams requires more than prompt engineering; it demands...
- <https://arxiv.org/abs/2606.17507>
-
- 5 **CONF**
- An Evaluation of Data Leakage Risks in Tool-Using LLM Agents in Realistic Scenarios**
- ArXiv CS.AI 50%
- arXiv:2606.17114v1 Announce Type: cross Abstract: AI agents are increasingly being adopted in enterprise and personal settings with access to emails, databases, documents, and other tools where they can read, update, and disseminate sensitive information. Much of prior research...
- <https://arxiv.org/abs/2606.17114>
-
- 6 **CONF**
- DeepInsight: A Unified Evaluation Infrastructure Across the Physical AI Stack**
- ArXiv CS.AI 50%
- arXiv:2606.17574v1 Announce Type: new Abstract: Evaluating a Physical AI stack spans operators that differ by more than three orders of magnitude -- from a single foundation-model decoding step to thousands of physics ticks of whole-body control -- varying orthogonally in...
- <https://arxiv.org/abs/2606.17574>
-
- 7 **CONF**
- Geometric Consistency Protocol for Foundation Model Features in Multi-View Satellite Imagery**
- ArXiv CS.AI 50%
- arXiv:2606.17564v1 Announce Type: cross Abstract: Standardized evaluation protocols are indispensable for robust benchmarking in remote sensing, particularly as foundation features are increasingly transferred across diverse sensors and complex imaging geometries. In satellite...
- <https://arxiv.org/abs/2606.17564>

8 CONF

Closing the Feedback Loop: From Experience Extraction to Insight Governance in Verbal Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.17591v1 Announce Type: new Abstract: Training-free verbal reinforcement learning enables LLM agents to learn from world feedback -- objective signals such as dynamic task outcomes, market returns, or demand forecasts -- by extracting verbal rules from experience and...

<https://arxiv.org/abs/2606.17591>

9 CONF

FinAcumen: Financial Multimodal Reasoning via Self-Evolving Experience Memory Harness

ArXiv CS.AI 50%

arXiv:2606.17642v1 Announce Type: new Abstract: Financial multimodal reasoning requires agents to coordinate numerical computation, retrieval, visual interpretation, and temporal grounding across heterogeneous evidence sources. Existing tool-augmented agents improve execution...

<https://arxiv.org/abs/2606.17642>

10 CONF

Beyond Domains: Reusing Web Skills via Transferable Interaction Patterns

ArXiv CS.AI 50%

arXiv:2606.17645v1 Announce Type: new Abstract: Large language model (LLM) web agents are usually deployed as tool callers: each turn, the model reads a fresh page observation and emits one structured tool action. When every action is a low-level primitive, horizons grow quickly...

<https://arxiv.org/abs/2606.17645>

1 CONF

From Brewing to Resolution: Tracing the Internal Lifecycle of Code Reasoning in LLMs

ArXiv CS.AI 50%

arXiv:2606.17648v1 Announce Type: new Abstract: Standard accuracy metrics cannot explain why LLMs handle variable tracking but fail on semantically equivalent loops. We study an internal lifecycle of code reasoning in which models first brew the answer, making it linearly...

<https://arxiv.org/abs/2606.17648>

2 CONF

FilumaOne: A Code-Native Multimodal CAD Dataset with Executable Programs and Kernel-Validated Feature Histories

ArXiv CS.AI 50%

arXiv:2606.17696v1 Announce Type: new Abstract: Parametric computer-aided design records both final geometry and the ordered construction history that determines how a part can be edited. Datasets for editable CAD research should therefore expose modeling operations, parameters,...

<https://arxiv.org/abs/2606.17696>

3 CONF

EComAgentBench: Benchmarking Shopping Agents on Long-Horizon Tasks with Distributed Hidden Intent

ArXiv CS.AI 50%

arXiv:2606.17698v1 Announce Type: new Abstract: As LLM-based shopping agents enter production, existing benchmarks fail to capture how a shopper's requirements arrive: stated implicitly in the query, recorded in a profile, or revealed only when the right question is asked....

<https://arxiv.org/abs/2606.17698>

4 CONF

DecoSearch: Complexity-Aware Routing and Plan-Level Repair for Text-to-SQL

ArXiv CS.AI 50%

arXiv:2606.17821v1 Announce Type: new Abstract: Large Language Models (LLMs) have demonstrated remarkable capabilities in translating natural language to SQL, yet existing methods still falter on complex queries requiring multi-step, data-aware reasoning. We introduce...

<https://arxiv.org/abs/2606.17821>

5 CONF

A homotopy-type-theoretic generalization of neurosymbolic inference

ArXiv CS.AI 50%

arXiv:2606.17851v1 Announce Type: new Abstract: A wide range of neurosymbolic (NeSy) systems compute one functional: a belief-weighted sum of a logical quantity over a space of Σ -structures, of which weighted model counting, fuzzy logic, and probabilistic logic are...

<https://arxiv.org/abs/2606.17851>

6 CONF

FlowRAG: Synergizing Explicit Reasoning via Frequency-Aware Multi-Granularity Graph Flow

ArXiv CS.AI 50%

arXiv:2606.17856v1 Announce Type: new Abstract: Graph-based retrieval-augmented generation (GraphRAG) is effective for knowledge-intensive and multi-hop query tasks; however, many existing methods primarily seed entity-based graphs and rely on implicit semantic relevance...

<https://arxiv.org/abs/2606.17856>

7 CONF

Structural Preservation and the Logical Expressiveness of Graph Neural Networks

ArXiv CS.AI 50%

arXiv:2606.17882v1 Announce Type: new Abstract: Bridges between graph neural networks (GNNs) and logical formalisms have been established by fixing architectural choices, such as the types of aggregation, combination, and activation functions. These choices define restricted...

<https://arxiv.org/abs/2606.17882>

8 CONF

MathVis-Fine: Aligning Visual Supervision with Necessity via Progressive Dependency-Guided Training for Multimodal Mathematical Reasoning

ArXiv CS.AI 50%

arXiv:2606.17888v1 Announce Type: new Abstract: Chain-of-Thought (CoT) reasoning has extended from purely linguistic domains to multimodal scenarios; however, existing approaches often treat visual inputs as homogeneous or auxiliary signals, failing to capture the intricate and...

<https://arxiv.org/abs/2606.17888>

9 CONF

Learn to Quantify Social Interaction with Constraints for Pedestrian Walking

ArXiv CS.AI 50%

arXiv:2606.17897v1 Announce Type: new Abstract: Long-term human path forecasting in crowds is critical for autonomous moving platforms (like autonomous driving cars and social robots) to avoid collision and make high-quality planning. Although the current research take into...

<https://arxiv.org/abs/2606.17897>

- 0 CONF**
- DiagFlowBench: Evaluating How Language Models Handle Off-Procedure Inputs in Grounded Diagnostic Dialogue**
- ArXiv CS.AI 50%
- arXiv:2606.17904v1 Announce Type: new Abstract: Language models increasingly serve as advisory systems in maintenance operations. To prevent hallucination, recent systems ground these models in procedural documentation to constrain them to approved steps. In practice, however,...
- <https://arxiv.org/abs/2606.17904>
-
- 1 CONF**
- How Inference Compute Shapes Frontier LLM Evaluation**
- ArXiv CS.AI 50%
- arXiv:2606.17930v1 Announce Type: new Abstract: AI evaluations are shifting toward harder tasks that benefit from longer trajectories involving tool use and iterative problem solving. As a result, performance is increasingly sensitive to the amount and allocation of compute...
- <https://arxiv.org/abs/2606.17930>
-
- 2 CONF**
- MoCo-AIS: A Contrastive Learning Framework for Similarity Computation of Vessel Trajectories**
- ArXiv CS.AI 50%
- arXiv:2606.17978v1 Announce Type: new Abstract: Trajectory similarity is a fundamental task in analyzing mobility patterns, essential for applications such as route pattern extraction, mobility prediction, and anomaly detection. Traditional distance-based measures for computing...
- <https://arxiv.org/abs/2606.17978>
-
- 3 CONF**
- Detail++: Training-Free Detail Enhancer for Text-to-Image Diffusion Models**
- ArXiv CS.AI 50%
- arXiv:2507.17853v2 Announce Type: replace-cross Abstract: Recent advances in text-to-image (T2I) generation have led to impressive visual results. However, these models still face significant challenges when handling complex prompt, particularly those involving multiple subjects...
- <https://arxiv.org/abs/2507.17853>
-
- 4 CONF**
- LLM Consumer Behavior Theory: Foundations of a Novel Research Field**
- ArXiv CS.AI 50%
- arXiv:2606.18005v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly deployed as autonomous agents that make consumption decisions on behalf of users. This shift raises fundamental questions for consumer theory, which has traditionally modeled humans as...
- <https://arxiv.org/abs/2606.18005>
-
- 5 CONF**
- LegalHalluLens: Typed Hallucination Auditing and Calibrated Multi-Agent Debate for Trustworthy Legal AI**
- ArXiv CS.AI 50%
- arXiv:2606.18021v1 Announce Type: new Abstract: AI systems deployed in legal workflows hallucinate at rates that aggregate metrics report at ~52%, but this average conceals where errors concentrate and in which direction they run, leaving compliance officers without an...
- <https://arxiv.org/abs/2606.18021>

6 CONF

ProvenanceGuard: Source-Aware Factuality Verification for MCP-Based LLM Agents

ArXiv CS.AI 50%

arXiv:2606.18037v1 Announce Type: new Abstract: Tool-using LLM agents increasingly use the Model Context Protocol (MCP) to answer from heterogeneous evidence sources, including search, APIs, databases, clinical records, and formulary tools. Standard factuality metrics usually...

<https://arxiv.org/abs/2606.18037>

7 CONF

Agentic AI-based Framework for Mitigating Premature Diagnostic Handoff and Silent Hallucination in Healthcare Applications

ArXiv CS.AI 50%

arXiv:2606.18068v1 Announce Type: new Abstract: Recent advances in Large Language Models (LLMs) and multi-agent systems have driven the rise of Agentic AI, showing promise for medical reasoning. However, open-ended conversational agents remain prone to two critical failure...

<https://arxiv.org/abs/2606.18068>

8 CONF

IsabelleLLM: Automated Theorem Proving Applied to Formally Verifying Consensus

ArXiv CS.AI 50%

arXiv:2606.18098v1 Announce Type: new Abstract: Advances in Artificial Intelligence (AI) have led AI for Theorem Proving to become a promising means of formally verifying computer systems. Whilst formal verification is traditionally reserved for safety-critical systems due to...

<https://arxiv.org/abs/2606.18098>

9 CONF

Your AI Travel Agent Would Book You a Bullfight: An Agentic Benchmark for Implicit Animal Welfare in Frontier AI Models

ArXiv CS.AI 50%

arXiv:2606.18142v1 Announce Type: new Abstract: AI agents are moving from advisors to actors, booking travel, planning menus, and running procurement on behalf of users. Existing benchmarks for AI and animal welfare evaluate model text responses to question-answer prompts,...

<https://arxiv.org/abs/2606.18142>

10 CONF

DRFLOW: A Deep Research Benchmark for Personalized Workflow Prediction

ArXiv CS.AI 50%

arXiv:2606.18191v1 Announce Type: new Abstract: Deep research (DR) systems are increasingly used for complex information-seeking tasks, but existing works mainly focus on generating reports and summaries. In contrast, many enterprise tasks instead require an agent to identify...

<https://arxiv.org/abs/2606.18191>

11 CONF

The Stanford EDGAR Filings Dataset: Reconstructing U.S. Corporate and Financial Disclosures into Layout-Faithful and Token-Efficient Pretraining Data

ArXiv CS.AI 50%

arXiv:2606.18192v1 Announce Type: new Abstract: As high-quality public web corpora become increasingly exhausted, clean long-context documents have become a scarce and expensive source of training data for large language models (LLMs). Existing long-context corpora are often...

<https://arxiv.org/abs/2606.18192>

2 **CONF****Fixed-Point Reasoners: Stable and Adaptive Deep Looped Transformers**

ArXiv CS.AI 50%

arXiv:2606.18206v1 Announce Type: new Abstract: Looped architectures provide an inductive bias toward learning step-by-step procedures for tasks that require compositional reasoning. The number of effective layers reached by looping determines the quality of the solution these...

<https://arxiv.org/abs/2606.18206>

3 **CONF****Correct When Paired, Wrong When Split: Decoupling and Editing Modality-Specific Neurons in MLLMs**

ArXiv CS.AI 50%

arXiv:2606.17057v1 Announce Type: cross Abstract: Although Knowledge Editing provides an efficient mechanism for updating the knowledge of Multimodal Large Language Models (MLLMs), we find that current paradigms still suffer from an important yet remain underexplored issue :...

<https://arxiv.org/abs/2606.17057>

4 **CONF****Towards Distributed Inference of LLMs on a P2P Network**

ArXiv CS.AI 50%

arXiv:2606.17059v1 Announce Type: cross Abstract: Prefix caching can reduce LLM inference latency by reusing KV caches across requests with shared prompts, but cluster-scale reuse is challenging because caches are partitioned across nodes. We propose a decentralized,...

<https://arxiv.org/abs/2606.17059>

5 **CONF****Extracting Semantics: LLM-Guided Automatic Population of Robot Ontology from URDF**

ArXiv CS.AI 50%

arXiv:2606.17073v1 Announce Type: cross Abstract: While commonsense knowledge may suffice for virtual agents, embodied robots interacting with humans require grounded and semantically rich representations of both their environment and their own physical embodiment. In cognitive...

<https://arxiv.org/abs/2606.17073>

6 **CONF****Surveying GenAI-based Automation in Printed Circuit Board Design and Test**

ArXiv CS.AI 50%

arXiv:2606.17074v1 Announce Type: cross Abstract: Generative artificial intelligence (GenAI) is increasingly used for applications in the hardware and software domains. It purports to reduce the manual effort involved in the development and testing of complex systems before...

<https://arxiv.org/abs/2606.17074>

7 **CONF****HRDX: A Large-Scale Vector HD-Map Dataset**

ArXiv CS.AI 50%

arXiv:2606.17080v1 Announce Type: cross Abstract: Reliable autonomous driving requires vectorized HD maps that are geometrically accurate, semantically rich, and scalable to long-horizon driving. However, existing public HD map datasets are limited in scale, provide sparse...

<https://arxiv.org/abs/2606.17080>

8 CONF

The Price of Anarchy in Disaggregated Inference

ArXiv CS.AI 50%

arXiv:2606.17081v1 Announce Type: cross Abstract: Disaggregated inference architectures physically separate prefill and decode phases onto distinct GPU pools, creating competing "agents" that share a fixed hardware budget. We provide, to our knowledge, the first formal...

<https://arxiv.org/abs/2606.17081>

9 CONF

ParkingTransformer: LLM-Enhanced End-to-End Trajectory Planning for Autonomous Parking

ArXiv CS.AI 50%

arXiv:2606.17082v1 Announce Type: cross Abstract: End-to-end autonomous parking has emerged as a critical task within the realm of autonomous driving. However, existing methods suffer from black-box characteristics, lacking high-level semantic understanding and interpretability,...

<https://arxiv.org/abs/2606.17082>

0 CONF

Software Delegation Contracts: Measuring Reviewability in AI Coding-Agent Work

ArXiv CS.AI 50%

arXiv:2606.17099v1 Announce Type: cross Abstract: AI coding agents increasingly accept assigned software tasks, modify repositories under bounded authority, and return work packages for review. Prior work proposed the software delegation contract, covering the task, authority,...

<https://arxiv.org/abs/2606.17099>

1 CONF

Quantum Cinema: An Interactive Cinematic Exploration of Quantum Computing Hardware via Generative World Models

ArXiv CS.AI 50%

arXiv:2606.17102v1 Announce Type: cross Abstract: Quantum computing promises transformative advances across science and industry, yet the physical hardware that enables these computations remains invisible to the public: quantum processors operate inside sealed dilution...

<https://arxiv.org/abs/2606.17102>

2 CONF

Prefill/Decode-Aware Evaluation of LLM Inference on Emerging AI Accelerators

ArXiv CS.AI 50%

arXiv:2606.17104v1 Announce Type: cross Abstract: As large language models (LLMs) are increasingly deployed in latency- and cost-sensitive settings, inference efficiency has become a central systems challenge. While GPUs dominate current deployments, a growing number of AI...

<https://arxiv.org/abs/2606.17104>

3 CONF

Pulling The REINS: Training-Free Safety Alignment of Video Diffusion Models via Representation Steering

ArXiv CS.AI 50%

arXiv:2606.17257v1 Announce Type: cross Abstract: Open-weight video diffusion models can generate photorealistic unsafe content, from violence to misinformation, yet existing defenses either require expensive safety fine-tuning that degrades general capability, or apply external...

<https://arxiv.org/abs/2606.17257>

- 4 **CONF**
- MODE: Modality-Decomposed Expert-Level Mixed-Precision Quantization for MoE Multimodal LLMs**
- ArXiv CS.AI 50%
- arXiv:2606.17118v1 Announce Type: cross Abstract: Mixture-of-Experts Multimodal Large Language Models (MoE-MLLMs) offer remarkable performance but incur prohibitive GPU memory costs, making compression essential. Among PTQ methods, expert-level mixed-precision quantization has...
- <https://arxiv.org/abs/2606.17118>
-
- 5 **CONF**
- Graph neural networks at war: integrating cybersecurity and drone intelligence in the Israeli-Iranian conflict**
- ArXiv CS.AI 50%
- arXiv:2606.17119v1 Announce Type: cross Abstract: Physical cyber systems have brought about new threats and challenges in detection and immediate response. This study examines how Graph Neural Networks (GNNs) can be used to aid cybersecurity and drone management in a physical...
- <https://arxiv.org/abs/2606.17119>
-
- 6 **CONF**
- LineageMark: Multi-user White-box Watermarking for Contribution Tracing in Model Derivation Chains**
- ArXiv CS.AI 50%
- arXiv:2606.17123v1 Announce Type: cross Abstract: In open large language model (LLM) ecosystems, models are frequently adapted across multiple domains and applications, forming multi-stage derivation chains. Consequently, tracking and verifying historical contributions is...
- <https://arxiv.org/abs/2606.17123>
-
- 7 **CONF**
- Statistical Foundations of LLM-based A/B Testing: A Surrogacy Framework for Human Causal Inference**
- ArXiv CS.AI 50%
- arXiv:2606.17165v1 Announce Type: cross Abstract: Organizations and researchers show increasing interest in using large language models (LLMs) in place of human participants in A/B tests, in the hope of experimenting faster and at lower cost. We study when a treatment effect...
- <https://arxiv.org/abs/2606.17165>
-
- 8 **CONF**
- Physics-Informed Attention Mechanism and Generalization Capability of Deep Learning-Based Grain Growth Evolution Prediction**
- ArXiv CS.AI 50%
- arXiv:2606.17235v1 Announce Type: cross Abstract: Machine Learning (ML) models for grain growth prediction are typically trained on idealized synthetic data, yet practical applications require generalization to conditions outside the training distribution. This study evaluated...
- <https://arxiv.org/abs/2606.17235>
-
- 9 **CONF**
- MLLP-VRAIN UPV system for the IWSLT 2026 Simultaneous Speech Translation task**
- ArXiv CS.AI 50%
- arXiv:2606.17255v1 Announce Type: cross Abstract: This work describes the participation of the MLLP-VRAIN research group in the shared task of the IWSLT 2026 Simultaneous Speech Translation track. Our submission utilizes the recently released Parakeet and Qwen 3.5 models to...
- <https://arxiv.org/abs/2606.17255>

0 **CONF****From Democracies to Autocracies: How AI Systems Enable Authoritarianism by Design****ArXiv CS.AI 50%**

arXiv:2606.17286v1 Announce Type: cross Abstract: AI-enabled authoritarianism is not confined to autocracies. In this paper, we provide greater transparency by investigating and mapping the lifecycles of six AI systems deployed in different political regimes, ranging from the US...

<https://arxiv.org/abs/2606.17286>

1 **CONF****Transformer-Based Warm-Starting for Feasible and Optimal Terminal Approach to Tumbling Objects with Space Manipulators****ArXiv CS.AI 50%**

arXiv:2606.17317v1 Announce Type: cross Abstract: Real-time trajectory generation for on-orbit robotic servicing is challenging due to the nonlinear coupling between spacecraft bus motion, manipulator dynamics, visibility cone, and trajectory-level safety constraints. This paper...

<https://arxiv.org/abs/2606.17317>

2 **CONF****Do Large Language Models Always Tell The Same Stories?****ArXiv CS.AI 50%**

arXiv:2606.17350v1 Announce Type: cross Abstract: Recent advances in large language models (LLMs) have enabled the generation of high-quality prose, yet the question of whether these models are capable of generating diverse outputs remains contested. In this work, we investigate...

<https://arxiv.org/abs/2606.17350>

3 **CONF****NarrativeWorldBench: A Frontier-Saturated Benchmark and a Latent World Model for Long-Horizon Co-Creative Audio Drama****ArXiv CS.AI 50%**

arXiv:2606.17391v1 Announce Type: cross Abstract: Long-form serialized audio drama, with arcs that run for 200 to 800 episodes, is a major creative medium and a setting where frontier large language models (LLMs) fail. We benchmark 21 models, spanning classical, fine-tuned,...

<https://arxiv.org/abs/2606.17391>

4 **CONF****SoK: AI-Augmented Binary Reversing****ArXiv CS.AI 50%**

arXiv:2606.17398v1 Announce Type: cross Abstract: Binary reversing is fundamental to software understanding, vulnerability discovery, malware investigation, and firmware auditing. However, it remains inherently challenging due to the irreversible loss of semantic information...

<https://arxiv.org/abs/2606.17398>

5 **CONF****The Discrete-Log Clock: How a Transformer Learns Modular Multiplication****ArXiv CS.AI 50%**

arXiv:2606.17399v1 Announce Type: cross Abstract: When small transformers grok modular multiplication, prior work reports that the learned embedding has a "dense" Fourier spectrum requiring all frequencies. This contrasts with modular addition, where only a sparse set of key...

<https://arxiv.org/abs/2606.17399>

6 CONF

Visored: A Controlled-Natural-Language Prover for LLM-Generated Mathematics

ArXiv CS.AI 50%

arXiv:2606.17581v1 Announce Type: cross Abstract: We present a dependent-type-based prover designed around the way LLMs (and humans) tend to write mathematics, complementing existing systems such as Lean and Rocq. Its core design choices are a surface that imitates mathematical...

<https://arxiv.org/abs/2606.17581>

7 CONF

L-Proto: Language-Aware Episodic Prototypical Training for Multilingual Speaker Verification

ArXiv CS.AI 50%

arXiv:2606.17416v1 Announce Type: cross Abstract: Multilingual speaker verification remains challenging because language-dependent acoustic variability causes speaker identity to become entangled with linguistic characteristics, degrading generalization across languages. In...

<https://arxiv.org/abs/2606.17416>

8 CONF

MODE-RAG: Manifold Outlier Diagnosis and Energy-based Retrieval-Augmented Generation Evaluation

ArXiv CS.AI 50%

arXiv:2606.17449v1 Announce Type: cross Abstract: While Multimodal Retrieval-Augmented Generation (M-RAG) enhances Large Vision-Language Models, it remains highly susceptible to cross-modal hallucinations, causal fabrications, and sycophancy. Furthermore, existing mitigation...

<https://arxiv.org/abs/2606.17449>

9 CONF

AUTOGATE: Automated Clock Gating via Toggling-Aware LLM-based RTL Rewriting

ArXiv CS.AI 50%

arXiv:2606.17461v1 Announce Type: cross Abstract: Fine-grain clock gating (FGCG) is among the most effective techniques for reducing dynamic power, yet current FGCG optimization flows remain largely manual. Recent LLM-based RTL optimization approaches remain limited by two key...

<https://arxiv.org/abs/2606.17461>

0 CONF

AIPatient Arena: EHR-grounded evaluation of large language models in end-to-end clinical consultation workflows

ArXiv CS.AI 50%

arXiv:2606.17474v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly considered for use in clinical consultation tasks, yet most medical evaluations remain static, single-turn, or narrowly outcome-based, limiting their ability to reflect the...

<https://arxiv.org/abs/2606.17474>

1 CONF

Decoding Hidden Deception in Reasoning LLMs: Activation Explainers for Deception Auditing

ArXiv CS.AI 50%

arXiv:2606.17478v1 Announce Type: cross Abstract: As LLMs acquire stronger reasoning capabilities, deceptive behavior becomes an increasingly serious safety concern. Existing deception monitors either score visible transcripts or derive scalar probe scores from representation...

<https://arxiv.org/abs/2606.17478>

2 **CONF****Online LLM Selection via Constrained Bandits with Time-Varying Demand**

ArXiv CS.AI 50%

arXiv:2606.17489v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly deployed in edge-cloud inference systems to handle diverse user tasks with heterogeneous accuracy, latency, and cost profiles. Selecting the appropriate LLM for each incoming task is...

<https://arxiv.org/abs/2606.17489>
3 **CONF****Geometry-Aware Post-Hoc Uncertainty Quantification in Operator Learning**

ArXiv CS.AI 50%

arXiv:2606.17513v1 Announce Type: cross Abstract: Neural operators provide fast surrogates for PDEs but their deterministic predictions limit their use in tasks requiring uncertainty quantification (UQ), especially under geometric variability. Existing approaches primarily model...

<https://arxiv.org/abs/2606.17513>
4 **CONF****Unlocking LLM Code Correction with Iterative Feedback Loops**

ArXiv CS.AI 50%

arXiv:2606.17514v1 Announce Type: cross Abstract: Large Language Models have shown remarkable capabilities in code generation. However, most existing evaluations focus only on single-attempt accuracy and overlook the iterative refinement process that is central to real-world...

<https://arxiv.org/abs/2606.17514>
5 **CONF****OmniDrive: An LLM-Choreographed Multi-Agent World Model with Unified Latent Co-Compression for Multi-View Driving Video Generation**

ArXiv CS.AI 50%

arXiv:2606.17536v1 Announce Type: cross Abstract: Generative world models for autonomous driving face two unresolved tensions: heterogeneous control injection, where free-form language, HD-maps, trajectories, and camera poses reside in incompatible representational spaces, and...

<https://arxiv.org/abs/2606.17536>
6 **CONF****An AI Security Agent for Banking: Multi-Vector Fraud and AML Detection Across Retail and Corporate Accounts**

ArXiv CS.AI 50%

arXiv:2606.17555v1 Announce Type: cross Abstract: Banks simultaneously face signature-based fraud (card-not-present attacks, account takeover, ATM cloning) and behavioural financial crime (structuring, layering, mule networks, business email compromise) -- two threat families...

<https://arxiv.org/abs/2606.17555>
7 **CONF****LLM Features Can Hurt GNNs: Concatenation Interference on Homophilous Graph Benchmarks**

ArXiv CS.AI 50%

arXiv:2606.17579v1 Announce Type: cross Abstract: Adding LLM-generated node features to graph neural networks (GNNs) is widely reported to improve accuracy on standard benchmarks. We document a contrasting observation: when LLM features are introduced through pure input...

<https://arxiv.org/abs/2606.17579>

8 CONF

When LLMs Analyze Scars: From Images to Clinically-Meaningful Features

ArXiv CS.AI 50%

arXiv:2606.18063v1 Announce Type: cross Abstract: Medical image classification faces a fundamental dilemma: while deep learning models achieve remarkable performance at scale, real-world clinical scenarios often suffer from severe data scarcity due to annotation costs, privacy...

<https://arxiv.org/abs/2606.18063>

9 CONF

SkillMoV: Mixture-of-View Routing with Prototype-Conditioned Gating for Unified Multi-View Proficiency Estimation

ArXiv CS.AI 50%

arXiv:2606.17615v1 Announce Type: cross Abstract: Estimating human proficiency from video is a key challenge for automated skill assessment, with applications in sports coaching, music pedagogy, surgical training, and workplace learning. Existing approaches often focus on...

<https://arxiv.org/abs/2606.17615>

0 CONF

Divide, Deliberate, Decide: A Multi-Agent Framework for Fine-Grained Egocentric Action Recognition

ArXiv CS.AI 50%

arXiv:2606.17627v1 Announce Type: cross Abstract: Fine-grained action recognition in egocentric video is challenging for Vision-Language Models (VLMs): actions often differ only in small visual cues, and a single model tends to be biased toward a subset of these cues. We propose...

<https://arxiv.org/abs/2606.17627>

1 CONF

SketchXplain: Intuitive Visual Explanations of Image Classifiers with Sketches

ArXiv CS.AI 50%

arXiv:2606.17646v1 Announce Type: cross Abstract: Saliency map visualizations explain image-based AI predictions by pointing to regions, but these are often unintuitive and semantically unclear, leaving an interpretability gap. We argue that AI explanations should be intuitive...

<https://arxiv.org/abs/2606.17646>

2 CONF

TuneAhead: Predicting Fine-tuning Performance Before Full Training Begins

ArXiv CS.AI 50%

arXiv:2606.17660v1 Announce Type: cross Abstract: Fine-tuning large language models (LLMs) is compute-intensive and error-prone: model performance depends sensitively on data quality and hyperparameter choices, and na\"ive runs can even degrade model performance. This raises a...

<https://arxiv.org/abs/2606.17660>

3 CONF

FacProcessTwin: An LLM-Based System for Process Twin Development

ArXiv CS.AI 50%

arXiv:2606.17666v1 Announce Type: cross Abstract: Process twins provide real-time representations of entire production processes. By capturing how process steps interact, rather than monitoring a single machine in isolation as an asset-based digital twin does, they have the...

<https://arxiv.org/abs/2606.17666>

4 **CONF****ASTEROID: A Spatiotemporal Information Transformer for Forecasting Multi-Step Time Series of Molecular Dynamics****ArXiv CS.AI** **50%**

arXiv:2606.17668v1 Announce Type: cross Abstract: Molecular dynamics (MD) simulation is computationally demanding, particularly for large-scale systems requiring long-term analysis. Accurate forecast of the outcomes of a MD simulation is not only an attractive scientific...

<https://arxiv.org/abs/2606.17668>

5 **CONF****See First, Answer Later: Visual Evidence Pre-Alignment via Sufficiency-Driven RL****ArXiv CS.AI** **50%**

arXiv:2606.17678v1 Announce Type: cross Abstract: Multimodal large language models (MLLMs) integrate strong text reasoning with visual inputs, yet their responses can be inconsistent with the underlying images, indicating ineffective utilization of visual evidence during...

<https://arxiv.org/abs/2606.17678>

6 **CONF****MIVE: A Minimalist Integer Vector Engine for Softmax LayerNorm and RMSNorm Acceleration****ArXiv CS.AI** **50%**

arXiv:2606.17781v1 Announce Type: cross Abstract: The rapid growth of Large Language Models (LLMs) has intensified the need for specialized hardware accelerators that can satisfy stringent inference latency and power constraints. Although matrix multiplications dominate the...

<https://arxiv.org/abs/2606.17781>

7 **CONF****Position: Coding Benchmarks Are Misaligned with Agentic Software Engineering****ArXiv CS.AI** **50%**

arXiv:2606.17799v1 Announce Type: cross Abstract: Coding agents have become a major mode of software engineering, but the benchmarks we use to compare them were designed in a pre-agent era: they collapse model, harness, and environment into a single end-to-end score, typically...

<https://arxiv.org/abs/2606.17799>

8 **CONF****No-Free-Fairness: Fundamental Limits and Trade-offs in Learning Systems****ArXiv CS.AI** **50%**

arXiv:2606.17810v1 Announce Type: cross Abstract: In this paper, we establish a set of theoretical impossibility results, termed the No-Free-Fairness theorems, that identify three fundamental sources of disparity in learning systems. First, we show that when a task exhibits...

<https://arxiv.org/abs/2606.17810>

9 **CONF****High-Fidelity 3D Geometric Reconstruction of Pelvic Organs from MRI: A Hybrid Deep Learning and Iterative Optimization Approach****ArXiv CS.AI** **50%**

arXiv:2606.17836v1 Announce Type: cross Abstract: Patient-specific 3D reconstruction of pelvic organ geometry from MRI is important for pelvic floor modeling and downstream patient-specific analysis. However, while previous studies have focused primarily on either image...

<https://arxiv.org/abs/2606.17836>

00 CONF

A Quantitative Analysis of Multimodal Biomarkers in Alzheimer's Disease

ArXiv CS.AI 50%

arXiv:2606.17867v1 Announce Type: cross Abstract: Despite increasing adoption of multimodal approaches in Alzheimer's Disease (AD) research -- aimed at integrating molecular, structural, clinical, and genetic biomarkers to enhance disease characterization -- the relationships...

<https://arxiv.org/abs/2606.17867>

01 CONF

AI Adoption Across a Multinational Workforce: Sociotechnical Conditions for GenAI Acceptance in Human Resources

ArXiv CS.AI 50%

arXiv:2606.17887v1 Announce Type: cross Abstract: Generative AI (GenAI) deployment in the workplace is accelerating rapidly. Nevertheless, questions of who adopts, who benefits, and who is left behind and why are still understudied. In this paper, we investigate these dynamics...

<https://arxiv.org/abs/2606.17887>

02 CONF

Trustworthy Self-Composable Big-Data-as-a-Service: An LLM-Orchestrated Multi-Agent Framework for Automated Data Engineering, AutoML, MLOps Deployment, and Drift-Aware Lifecycle Optimization

ArXiv CS.AI 50%

arXiv:2606.17915v1 Announce Type: cross Abstract: Big-Data-as-a-Service (BDaaS) platforms require reliable automation across data ingestion, cleaning, feature engineering, model development, deployment, and post-deployment monitoring. However, existing LLM-based data science...

<https://arxiv.org/abs/2606.17915>

03 CONF

Plug-and-Adapt: Multimodal Coreference Resolution at First Sight with a Pretrained Alignment Model

ArXiv CS.AI 50%

arXiv:2606.17950v1 Announce Type: cross Abstract: Visual information helps resolve ambiguity in coreference resolution, leading to notable performance gains. However, existing Multi-modal Coreference Resolution (MCR) methods require training with (partially) annotated data from...

<https://arxiv.org/abs/2606.17950>

04 CONF

SoftMoE: Soft Differentiable Routing for Mixture-of-Experts in LLMs

ArXiv CS.AI 50%

arXiv:2606.17952v1 Announce Type: cross Abstract: Sparse Mixture-of-Experts (MoE) architectures enable scaling LLM parameters under a fixed inference budget by activating only a small subset of experts via top-\$k\$ routing. While this preserves causality and suits autoregressive...

<https://arxiv.org/abs/2606.17952>

05 CONF

A T-API-Compliant ReAct Agentic Loop for Optical Networks: Generic vs. Domain-Specific Tool Abstractions

ArXiv CS.AI 50%

arXiv:2606.18000v1 Announce Type: cross Abstract: Optical networks need intent-driven, closed-loop agentic management, a key enabler for higher autonomy levels. We present the first T-API-compliant reasoning and act (ReAct) loop. We show that domain-specific composite tools...

<https://arxiv.org/abs/2606.18000>

06 CONF

When AI Says "I have been in similar situations": Synthetic Lived Experience in Peer-Like Caregiver Support

ArXiv CS.AI 50%

arXiv:2606.18057v1 Announce Type: cross Abstract: Caregivers often turn to online communities for informational and emotional support. In these spaces, peer supporters frequently draw on personal narratives to respond to emotionally complex caregiving situations. As LLMs are...

<https://arxiv.org/abs/2606.18057>

07 CONF

Security and Privacy Prompts in the Wild: What Users Ask LLMs and How LLMs Respond

ArXiv CS.AI 50%

arXiv:2606.18062v1 Announce Type: cross Abstract: Large language models (LLMs) are widely used to fulfill users' information needs; users ask LLMs about the weather, pose educational questions, and consult them for legal assistance. One particularly understudied area is digital...

<https://arxiv.org/abs/2606.18062>

08 CONF

S4oP: Operator-level Pruning of Structured State Space Models for Resource-Constrained Devices

ArXiv CS.AI 50%

arXiv:2606.18096v1 Announce Type: cross Abstract: Structured State Space Models (SSMs), including the S4 and S4D architectures, have recently emerged as powerful alternatives to attention-based models for capturing long-range dependencies in sequential data. Despite their strong...

<https://arxiv.org/abs/2606.18096>

09 CONF

Ternary Mamba: Grouped Quantization-Aware Training of W1.58A16 State Space Models

ArXiv CS.AI 50%

arXiv:2606.18114v1 Announce Type: cross Abstract: State Space Models (SSMs) such as Mamba-2 offer linear-time inference but their memory footprint limits edge deployment. Prior ternary SSM work (Slender-Mamba) trains from scratch on 150B tokens; we show a pretrained checkpoint...

<https://arxiv.org/abs/2606.18114>

10 CONF

Structural Role Injection in Handlebars-Templated LLM Prompts: Triple-Brace Interpolation, Delimiter Family, and the Limits of HTML Auto-Escaping

ArXiv CS.AI 50%

arXiv:2606.18120v1 Announce Type: cross Abstract: Large language model applications build prompts from templates, and Handlebars is a widely used templating engine and the default prompt-template format in Microsoft Semantic Kernel. Its double-brace `{{x}}` expression HTML-escapes...

<https://arxiv.org/abs/2606.18120>

11 CONF

Embedded Machine Learning for Microcontroller-Class Edge Devices: Data, Feature, Evaluation, and Deployment Pipelines

ArXiv CS.AI 50%

arXiv:2606.18122v1 Announce Type: cross Abstract: Embedded machine learning moves inference from cloud services to resource-constrained devices that must acquire data, preprocess signals, run a model, and act within tight limits on memory, energy, and latency. This paper...

<https://arxiv.org/abs/2606.18122>

12 CONF

IUU+DB: Tracking Illegal, Unreported, and Unregulated Fishing, Seafood Fraud, and Labor Abuse through LLM-driven Information Extraction

ArXiv CS.AI 50%

arXiv:2606.18181v1 Announce Type: cross Abstract: Illegal, unreported, and unregulated fishing (IUU) traditionally refers to fishing activities that violate applicable laws or occur in areas that lack applicable laws. We propose the term IUU+ to capture a broader suite of...

<https://arxiv.org/abs/2606.18181>

13 CONF

Kolmogorov Regression for Robust Diffusion Policies

ArXiv CS.AI 50%

arXiv:2606.18186v1 Announce Type: cross Abstract: Finite-dimensional (FD) diffusion policies exhibit temporal drift owing to discretization artifacts that degrade long-horizon performance (when deployed on physical systems). We introduce a backward Kolmogorov equation that lifts...

<https://arxiv.org/abs/2606.18186>

14 CONF

A Red-Team Study of Anthropic Fable 5 & Opus 4.8 Models

ArXiv CS.AI 50%

arXiv:2606.18193v1 Announce Type: cross Abstract: We evaluate the adversarial robustness of two frontier large language models (LLMs) developed by Anthropic, Fable 5 and Opus 4.8, against four families of automated jailbreak attack across 7 826 harmful intents spanning a...

<https://arxiv.org/abs/2606.18193>

15 CONF

Visual Verification Enables Inference-time Steering and Autonomous Policy Improvement

ArXiv CS.AI 50%

arXiv:2606.18247v1 Announce Type: cross Abstract: Robots deployed in the real world should learn from their experience and improve over time. This requires a mechanism of practicing and learning from feedback. In this paper, we propose VERITAS, a generator-verifier framework for...

<https://arxiv.org/abs/2606.18247>

16 CONF

CausalT5k: Diagnosing Refusal and Failure Modes in Trustworthy Causal Reasoning Across Causal Rungs

ArXiv CS.AI 50%

arXiv:2602.08939v2 Announce Type: replace Abstract: Large language models increasingly produce fluent causal explanations, yet they often fail in ways aggregate accuracy cannot diagnose: confusing association with intervention, abandoning correct judgments under pressure,...

<https://arxiv.org/abs/2602.08939>

17 CONF

OmniSapiens: A Foundation Model for Social Behavior Processing via Heterogeneity-Aware Relative Policy Optimization

ArXiv CS.AI 50%

arXiv:2602.10635v3 Announce Type: replace Abstract: Socially intelligent AI systems must reason across diverse human behavioral tasks and generalize to new social contexts. However, behavioral data is inherently heterogeneous, comprising diverse modalities and prediction targets...

<https://arxiv.org/abs/2602.10635>

18 CONF

Towards Leveraging AutoML for Sustainable Deep Learning: A Multi-Objective HPO Approach on Deep Shift Neural Networks

ArXiv CS.AI 50%

arXiv:2404.01965v3 Announce Type: replace-cross Abstract: Deep Learning (DL) has advanced various fields by extracting complex patterns from large datasets. However, the computational demands of DL models pose environmental and resource challenges. Deep shift neural networks...

<https://arxiv.org/abs/2404.01965>

19 CONF

E2Vec: Feature Embedding with Temporal Information for Analyzing Student Actions in E-Book Systems

ArXiv CS.AI 50%

arXiv:2407.13053v2 Announce Type: replace-cross Abstract: Digital textbook (e-book) systems record student interactions with textbooks as a sequence of events called EventStream data. In the past, researchers extracted meaningful features from EventStream, and utilized them as...

<https://arxiv.org/abs/2407.13053>

20 CONF

LLM-Powered Multi-Agent System for Automated Crypto Portfolio Management

ArXiv CS.AI 50%

arXiv:2501.00826v3 Announce Type: replace-cross Abstract: Cryptocurrency portfolio management requires the fusion of heterogeneous multi-modal signals, including structured price and on-chain time series, unstructured news text, and technical indicators, under high-volatility...

<https://arxiv.org/abs/2501.00826>

21 CONF

Top-Theta Attention: Sparsifying Transformers by Compensated Thresholding

ArXiv CS.AI 50%

arXiv:2502.08363v3 Announce Type: replace-cross Abstract: We present Top-Theta (Top- θ) Attention, a training-free method for sparsifying transformer attention during inference. Our key insight is that static, per-head thresholds can be calibrated to retain the desired...

<https://arxiv.org/abs/2502.08363>

22 CONF

Gaussian DP for Reporting Differential Privacy Guarantees in Machine Learning

ArXiv CS.AI 50%

arXiv:2503.10945v3 Announce Type: replace-cross Abstract: Current practices for reporting differential privacy (DP) guarantees for machine learning (ML) algorithms such as DP-SGD provide an incomplete and potentially misleading picture. For instance, if only a single...

<https://arxiv.org/abs/2503.10945>

23 CONF

AnalogFed: Privacy-Preserving Discovery of Analog Circuits at Scale with Federated Generative AI

ArXiv CS.AI 50%

arXiv:2507.15104v2 Announce Type: replace-cross Abstract: Recent advances in generative AI (GenAI) have shown transformative potential for modern hardware design. However, existing GenAI-driven approaches fall short of enabling large-scale electronic design automation (EDA) due...

<https://arxiv.org/abs/2507.15104>

24 CONF

Blueprint First, Model Second: A Framework for Deterministic LLM Workflow

ArXiv CS.AI 50%

arXiv:2508.02721v2 Announce Type: replace-cross Abstract: While powerful, the inherent non-determinism of large language model (LLM) agents limits their application in structured operational environments where procedural fidelity and predictable execution are strict...

<https://arxiv.org/abs/2508.02721>

25 CONF

Breaking the Code: Security Assessment of AI Code Agents Through Systematic Jailbreaking Attacks

ArXiv CS.AI 50%

arXiv:2510.01359v2 Announce Type: replace-cross Abstract: Code-capable large language model (LLM) agents are embedded in software engineering workflows where they can read, write, and execute code, raising "jailbreak" stakes beyond text-only settings. Prior evaluations emphasize...

<https://arxiv.org/abs/2510.01359>

26 CONF

Adversarial Attacks Leverage Interference Between Features in Superposition

ArXiv CS.AI 50%

arXiv:2510.11709v2 Announce Type: replace-cross Abstract: Why do adversarial examples exist, and why do they transfer between models? Existing explanations appeal to high-dimensional geometry, non-robust patterns in the input, and decision boundary structure, but none provides a...

<https://arxiv.org/abs/2510.11709>

27 CONF

BadScientist: Can a Research Agent Write Convincing but Unsound Papers that Fool LLM Reviewers?

ArXiv CS.AI 50%

arXiv:2510.18003v2 Announce Type: replace-cross Abstract: The convergence of LLM-powered research assistants and AI-based peer review systems creates a critical vulnerability: fully automated publication loops where AI-generated research is evaluated by AI reviewers without...

<https://arxiv.org/abs/2510.18003>

28 CONF

Enhanced Evolutionary Multi-Objective Deep Reinforcement Learning for Reliable and Efficient Wireless Rechargeable Sensor Networks

ArXiv CS.AI 50%

arXiv:2510.21127v2 Announce Type: replace-cross Abstract: Despite rapid advancements in sensor networks, conventional battery-powered sensor networks suffer from limited operational lifespans and frequent maintenance requirements that severely constrain their deployment in...

<https://arxiv.org/abs/2510.21127>

29 CONF

A geometric and deep learning reproducible pipeline for monitoring floating anthropogenic debris in urban rivers using in situ cameras

ArXiv CS.AI 50%

arXiv:2510.23798v2 Announce Type: replace-cross Abstract: The proliferation of floating anthropogenic debris in rivers has emerged as a pressing environmental concern, exerting a detrimental influence on biodiversity, water quality, and human activities such as navigation and...

<https://arxiv.org/abs/2510.23798>

30 CONF

m2sv: A Scalable Benchmark for Map-to-Street-View Spatial Reasoning

ArXiv CS.AI 50%

arXiv:2601.19099v2 Announce Type: replace-cross Abstract: Vision--language models (VLMs) achieve strong performance on many multimodal benchmarks but remain brittle on spatial reasoning tasks that require aligning abstract overhead representations with egocentric views. We...

<https://arxiv.org/abs/2601.19099>

31 CONF

Vulcan: Instance-specialized, Verifiable Systems Heuristics Through LLM-driven Search

ArXiv CS.AI 50%

arXiv:2512.25065v2 Announce Type: replace-cross Abstract: Systems resource management tasks rely primarily on hand-designed heuristics. However, growing hardware heterogeneity and workload diversity require heuristics specialized to particular deployment instances, making manual...

<https://arxiv.org/abs/2512.25065>

32 CONF

Phys4D: Fine-Grained Physics-Consistent 4D Modeling from Video Diffusion

ArXiv CS.AI 50%

arXiv:2603.03485v3 Announce Type: replace-cross Abstract: Recent video diffusion models have achieved impressive capabilities as large-scale generative world models. However, these models often struggle with fine-grained physical consistency, exhibiting physically implausible...

<https://arxiv.org/abs/2603.03485>

33 CONF

CogGen: Cognitive-Load-Inspired Fully Unsupervised Deep Generative Modeling for Compressively Sampled MRI Reconstruction

ArXiv CS.AI 50%

arXiv:2603.04438v3 Announce Type: replace-cross Abstract: Fully unsupervised deep generative modeling (FU-DGM) offers significant potential for compressively sampled magnetic resonance imaging (CS-MRI) reconstruction. Representative FU-DGM formulations, such as deep image prior...

<https://arxiv.org/abs/2603.04438>

34 CONF

Rethinking Multimodal Fusion for Time Series: Text Modalities Need Constrained Fusion

ArXiv CS.AI 50%

arXiv:2603.22372v2 Announce Type: replace-cross Abstract: Recent advances in multimodal learning have motivated the integration of auxiliary modalities such as text or vision into time series (TS) forecasting. However, most existing methods provide limited gains, often improving...

<https://arxiv.org/abs/2603.22372>

35 CONF

From Paper to Program: Knowledge Externalization for AI-Assisted Quantum Many-Body Code Generation

ArXiv CS.AI 50%

arXiv:2604.04089v3 Announce Type: replace-cross Abstract: Large language models can write scientific code, but direct paper-to-program translation remains fragile when correctness depends on tacit conventions in the literature. We identify this bottleneck as \textbf{knowledge}...

<https://arxiv.org/abs/2604.04089>

36 CONF

When Life Gives You BC, Make Q-functions: Extracting Q-values from Behavior Cloning for On-Robot Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.05172v2 Announce Type: replace-cross Abstract: Behavior Cloning (BC) has emerged as a highly effective paradigm for robot learning. However, BC lacks a self-guided mechanism for online improvement after demonstrations have been collected. Existing offline-to-online...

<https://arxiv.org/abs/2605.05172>

37 CONF

Rethinking Cross-Layer Information Routing in Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2605.20708v2 Announce Type: replace-cross Abstract: Diffusion Transformers (DiTs) have become a de facto backbone of modern visual generation, and nearly every major axis of their design -- tokenization, attention, conditioning, objectives, and latent autoencoders -- has...

<https://arxiv.org/abs/2605.20708>

38 CONF

Sustainable Metal-Organic Framework Water Harvesters in the Artificial Intelligence Era

ArXiv CS.AI 50%

arXiv:2605.29179v2 Announce Type: replace-cross Abstract: Metal-organic frameworks (MOFs) are excellent candidates for water harvesting due to their tunable pore environments, which can be precisely engineered to capture and release water in arid conditions. Integrating...

<https://arxiv.org/abs/2605.29179>

39 CONF

Clay-CNN Hybrids: Leveraging Geospatial Foundation Models as Auxiliary Context for Landslide Detection

ArXiv CS.AI 50%

arXiv:2606.14081v3 Announce Type: replace-cross Abstract: Rapid post-event landslide mapping is essential for disaster response but remains difficult to automate due to extreme class imbalance. This study evaluates whether Clay v1.5, a Geospatial Foundation Model (GFM), can...

<https://arxiv.org/abs/2606.14081>

40 CONF

From Agent Traces to Trust: A Survey of Evidence Tracing and Execution Provenance in LLM Agents

ArXiv CS.AI 50%

arXiv:2606.04990v3 Announce Type: replace-cross Abstract: Large language model (LLM)-based agents are evolving from passive text generators into autonomous systems capable of planning, tool use, retrieval, memory access, environmental interaction, and multi-agent collaboration....

<https://arxiv.org/abs/2606.04990>