



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	113	Media
TechCrunch AI	10	Media
The Verge AI	4	Media
HackerNews	4	Community
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 1 story

VERI

Who decides when AI is too dangerous?

The Verge AI

+1 source

80%

On today's episode of Decoder, my guest is Hayden Field, senior AI reporter for The Verge. Often when Hayden comes on the show, it's because something has gone wrong in the world of AI. Last weekend, that something was a pretty intense mix of Anthropic, the Trump administration,...

<https://www.theverge.com/podcast/951542/anthropic-claude-fable-5-mythos-ban-pentagon-ai...>

CONFIRMED SIGNALS — 130 stories

CONF

Show HN: Gerrymandle - Daily puzzle game where you redraw electoral districts

HackerNews

60%

HN 172

No summary — see link for full article.

<https://gerrymandle.cc/>

CONF

Show HN: Fata - Spaced repetition to fight skill rot from AI coding

HackerNews

60%

HN 121

Hi HN, I'm Djoumé. I've been a developer for over 20 years, and like a lot of you I've been coding almost exclusively through an agent in the past few months. It's been amazing to vibe code prototypes in any stack, but when it comes to building something reliable/scalable, I...

<https://fata.dev>

CONF

Show HN: Run Agent Skills with mistral.rs v0.8.10: /v1/skills support and more

HackerNews

50%

HN 16

Hey all! I'm the maintainer of mistral.rs. I just landed support for OpenAI-compatible Agent Skills via a /v1/skills endpoint, and it works with local open models. Until now Skills have basically been locked to closed models, and with the ability to have private, local...

<https://news.ycombinator.com/item?id=48581792>

CONF

Show HN: BeamWeaver - LangChain/DeepAgents-style agents and workflows for Elixir

HackerNews

50%

HN 2

Hi HN, We build agents in Elixir. We kept running into the same issue and found there is no observability for agentic systems. We decided to take the best aspects of LangChain, LangGraph, and DeepAgents and put them into Elixir. BeamWeaver comes with an OTP-native design and: -...

https://github.com/caudena/beam_weaver

CONF

Source: Elastic agrees to buy CRV-backed DeductiveAI for up to \$85M

TechCrunch AI

50%

DeductiveAI, a startup that uses AI to catch and resolve bugs in software, was founded just three years ago.

<https://techcrunch.com/2026/06/18/source-elastic-agrees-to-buy-crv-backed-deductiveai-f...>

CONF

AI inference startup Baseten reportedly raising \$1.5B months after its last mega-round

TechCrunch AI

50%

Startup Baseten is reportedly close to finalizing a \$1.5 billion round at a \$13 billion as the "inference gold rush" marches on.

<https://techcrunch.com/2026/06/18/ai-inference-startup-baseten-reportedly-raising-1-5b-...>

CONF

Snap spins off AI video team into new company, Dotmo, due to costs

TechCrunch AI 50%

The Snapchat maker is spinning off yet another internal unit. Dotmo will be composed of current Snap staff who are leaving the social media company to focus on AI video development.

<https://techcrunch.com/2026/06/18/snap-spins-off-ai-video-team-into-new-company-dotmo-d...>

CONF

OpenAI is bringing on some big guns in the lead-up to its IPO

TechCrunch AI 50%

OpenAI is bulking up before its IPO, landing Transformer co-inventor Noam Shazeer from Google DeepMind and former Trump AI policy official Dean Ball in the same week.

<https://techcrunch.com/2026/06/18/openai-is-bringing-on-some-big-guns-in-the-lead-up-to...>

0

CONF

Almost half of US singles feel negatively about AI in dating, Match says

TechCrunch AI 50%

About 47% of singles look negatively at the use of AI in dating -- but many dating app users are open to AI helping with profile punch-ups and conversation starters.

<https://techcrunch.com/2026/06/18/almost-half-of-u-s-singles-feel-negatively-about-ai-i...>

1

CONF

Amazon hopes to challenge Nvidia more directly by selling its AI chips

TechCrunch AI 50%

AWS is in talks to sell its chips to other data centers. CEO Andy Jassy has said this represents a \$50 billion opportunity for the company.

<https://techcrunch.com/2026/06/18/amazon-hopes-to-challenge-nvidia-more-directly-by-sel...>

2

CONF

AI data centers just got a government-mandated fast lane to the grid

TechCrunch AI 50%

FERC told grid operators to give data centers a fast lane for interconnections, but it failed to address electricity supply shortages.

<https://techcrunch.com/2026/06/18/ai-data-centers-just-got-a-government-mandated-fast-l...>

3

CONF

'Queer Eye' life coach Karamo Brown launches Kē, a wellness app featuring his AI digital clone

TechCrunch AI 50%

After spending a year and a half focusing on his own journey — from fitness and nutrition to meditation, sobriety, relationships, and personal growth — Brown wants to help others do the same.

<https://techcrunch.com/2026/06/18/queer-eyes-life-coach-karamo-brown-launches-ke-a-well...>

4

CONF

General Intuition in talks to raise \$300M at around \$2B valuation

TechCrunch AI 50%

The startup trains embodied AI and world models using Medal's dataset of 2 billion videos per year from 10 million monthly active users.

<https://techcrunch.com/2026/06/18/general-intuition-in-talks-to-raise-300m-at-around-2b...>

5

CONF

The Download: the first brain implant power user and South Korea's AI obsession

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. This man with ALS is the first "power user" of a brain implant that lets him speak Casey Harrell has had a set of electrodes embedded in his...

<https://www.technologyreview.com/2026/06/16/1139010/the-download-brain-implant-power-us...>

6 CONF

Barret Zoph is out at OpenAI again after just five months**The Verge AI** 50%

Five months after returning to OpenAI, Barret Zoph - the company's head of enterprise AI sales - has departed, The Verge has learned. Zoph returned to OpenAI in mid-January after a stint as co-founder and CTO of Thinking Machines Lab, the competing AI company founded by former...

<https://www.theverge.com/ai-artificial-intelligence/952837/barret-zoph-openai-thinking-...>

7 CONF

Photoshop and Premiere now have AI assistants**The Verge AI** 50%

Adobe's plan to stick AI assistants into all of its Creative Cloud suite is now fully underway, with new chatbots now rolling out to its biggest editing and design apps. As part of a public beta launching today, Photoshop, Premiere, Illustrator, InDesign, and Frame.io now each...

[https://www.theverge.com/tech/952099/adobe-ai-assistants-photoshop-premiere-illustrator...](https://www.theverge.com/tech/952099/adobe-ai-assistants-photoshop-premiere-illustrator-...)

8 CONF

Adobe's redesigned AI studio remembers what your creations look like**The Verge AI** 50%

Adobe is introducing some new capabilities for its Firefly AI assistant, alongside a "reimagined" AI studio that lets you edit and generate new designs from a single interface. The new Firefly experience launching today in private beta is designed to give you "persistent..."

<https://www.theverge.com/tech/952104/adobe-firefly-ai-agent-elements-projects-update>

9 CONF

Measuring Curriculum Alignment across Topical Coverage, Competency, and Cognitive Depth: A Longitudinal Framework Applied to CS2013 and CS2023**ArXiv CS.AI** 50%

arXiv:2606.19469v1 Announce Type: new Abstract: Undergraduate computer science is governed by international curricular guidelines revised about once a decade, yet programs lack a reliable, reproducible way to measure how completely they cover the current guidelines and how that...

<https://arxiv.org/abs/2606.19469>

0 CONF

DeFrame: Debiasing Large Language Models Against Framing Effects**ArXiv CS.AI** 50%

arXiv:2602.04306v2 Announce Type: replace-cross Abstract: As large language models (LLMs) are increasingly deployed in real-world applications, ensuring their fair responses across demographics has become crucial. Despite many efforts, an ongoing challenge is hidden bias: LLMs...

<https://arxiv.org/abs/2602.04306>

1 CONF

Hidden Anchors in Multi-Agent LLM Deliberation**ArXiv CS.AI** 50%

arXiv:2606.19494v1 Announce Type: new Abstract: Multi-agent LLM deliberation, where agents exchange and revise answers over several rounds, is increasingly used to improve reasoning and accuracy, yet how and why it works is rarely modelled. Such deliberation mirrors how humans...

<https://arxiv.org/abs/2606.19494>

2 **CONF****LLM Doesn't Know What It Doesn't Know: Detecting Epistemic Blind Spots via Cross-Model Attribution Divergence on Clinical Tabular Data****ArXiv CS.AI** **50%**

arXiv:2606.19509v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly applied to structured clinical data, yet whether they can recognize the limits of their own knowledge on such tasks remains unexplored. We study this question through the lens of...

<https://arxiv.org/abs/2606.19509>

3 **CONF****Emergent Alignment****ArXiv CS.AI** **50%**

arXiv:2606.19527v1 Announce Type: new Abstract: Can Large Language Models (LLMs) discern when their own outputs are misaligned with human ethics? And can they self-correct? We endow an LLM with a conscience step that reviews its own reasoning and outputs, and we extend the...

<https://arxiv.org/abs/2606.19527>

4 **CONF****Multi-View Decompilation for LLM-Based Malware Classification****ArXiv CS.AI** **50%**

arXiv:2606.20436v1 Announce Type: cross Abstract: Malware analysts often inspect compiled binaries through decompiled pseudo-C, when source code is unavailable. Recent work suggests that large language models (LLMs) can assist this process by classifying decompiled code as...

<https://arxiv.org/abs/2606.20436>

5 **CONF****Analyzing the Narration Gap in LLM-Solver Loops****ArXiv CS.AI** **50%**

arXiv:2606.19588v1 Announce Type: new Abstract: Formal tools such as SAT and SMT solvers are increasingly embedded in language model reasoning pipelines when a safety or security critical question can be formulated in logic. Unlike chain of thought whose steps are sampled from...

<https://arxiv.org/abs/2606.19588>

6 **CONF****Configurable Clinical Information Extraction with Agentic RAG: What Works, What Breaks, and Why****ArXiv CS.AI** **50%**

arXiv:2606.19602v1 Announce Type: new Abstract: Patient contexts span hundreds of heterogeneous documents and thousands of structured data points, yet the document-level metadata that AI systems need for retrieval and triage is absent or incomplete. Standard retrieval-augmented...

<https://arxiv.org/abs/2606.19602>

7 **CONF****Which Pairs to Compare for LLM Post-Training?****ArXiv CS.AI** **50%**

arXiv:2606.19607v1 Announce Type: new Abstract: Preference-based post-training has become a central paradigm for aligning language models. A common data-collection strategy is to generate a small set of completions for each prompt and label the resulting comparison pairs....

<https://arxiv.org/abs/2606.19607>

8 CONF

AI4SE and SE4AI Exploration: A Decade Looking Back and Forward

ArXiv CS.AI 50%

arXiv:2606.19630v1 Announce Type: new Abstract: The March 2020 INCOSE INSIGHT special issue on AI and Systems Engineering (SE) became the most downloaded issue in the publication's history and launched a research community that now draws over 250 registrants to its annual...

<https://arxiv.org/abs/2606.19630>

9 CONF

Structuring and Tokenizing Distributed User Interest Context for Generative Recommendation

ArXiv CS.AI 50%

arXiv:2606.20554v1 Announce Type: cross Abstract: Generative recommendation is an emerging paradigm that has shown promise in industrial recommendation systems, aiming to predict users' next interactions from their historical behaviors. At the core of generative recommendation...

<https://arxiv.org/abs/2606.20554>

0 CONF

Beyond Entropy: Learning from Token-Level Distributional Deviations for LLM Reasoning

ArXiv CS.AI 50%

arXiv:2606.19771v1 Announce Type: new Abstract: Reinforcement Learning with Verifiable Rewards (RLVR) has significantly advanced Large Language Model (LLM) reasoning; however, it faces a fundamental optimization instability: uniform token updates precipitate entropy collapse,...

<https://arxiv.org/abs/2606.19771>

1 CONF

A Comparative Study of Pretrained Transformer Models for Quranic ASR: Speech Representations, Label Formats, and Dataset Composition

ArXiv CS.AI 50%

arXiv:2606.19747v1 Announce Type: new Abstract: Quran Automatic Speech Recognition (ASR) aims to convert Quranic recitation into text, enabling applications such as aided memorisation tools and Quranic search engines. However, existing ASR models often exhibit high Word Error...

<https://arxiv.org/abs/2606.19747>

2 CONF

Benchmarking Agentic Review Systems

ArXiv CS.AI 50%

arXiv:2606.19749v1 Announce Type: new Abstract: A new class of agentic review systems are emerging as a remedy to the pressure placed on peer review systems by AI-assisted research, but it is unclear how they should be evaluated. We evaluate two open-source systems (OpenAIReview...

<https://arxiv.org/abs/2606.19749>

3 CONF

Grounded Inference: Principles for Deterministically Encapsulated Generative Models

ArXiv CS.AI 50%

arXiv:2606.19753v1 Announce Type: new Abstract: The incorporation of generative models into traditional computational systems presents both enormous opportunity and tremendous peril. Although many early adopters have realized these perils at great expense, the field still...

<https://arxiv.org/abs/2606.19753>

4 CONF

ORAgentBench: Can LLM Agents Solve Challenging Operations Research Tasks End to End?

ArXiv CS.AI 50%

arXiv:2606.19787v1 Announce Type: new Abstract: Large language models are increasingly deployed as autonomous agents for multi-step tasks in executable environments, yet their ability to perform realistic operations research (OR) work remains unclear. Existing OR evaluations...

<https://arxiv.org/abs/2606.19787>

5 CONF

Beyond Reasoning Gains: Mitigating General-Capability Forgetting in Large Reasoning Models

ArXiv CS.AI 50%

arXiv:2510.21978v2 Announce Type: replace-cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has delivered impressive gains in mathematical and multimodal reasoning and has become a standard post-training paradigm for contemporary language and vision-language...

<https://arxiv.org/abs/2510.21978>

6 CONF

Think Again or Think Longer? Selective Verification for Budget-Aware Reasoning

ArXiv CS.AI 50%

arXiv:2606.19808v1 Announce Type: new Abstract: Test-time reasoning is increasingly used as a serving-time control knob, but extra reasoning is not uniformly valuable: it can repair failed attempts, waste compute on already-correct answers, or introduce harmful answer changes....

<https://arxiv.org/abs/2606.19808>

7 CONF

TelcoAgent: A Scalable 5G Multi-KPM Forecasting With 3GPP-Grounded Explainability

ArXiv CS.AI 50%

arXiv:2606.19821v1 Announce Type: new Abstract: Key Performance Measurement (KPM) forecasting is essential for proactive network management of 5G and next-generation telecom networks. However, existing machine learning (ML) approaches face significant limitations in scalability...

<https://arxiv.org/abs/2606.19821>

8 CONF

StarOR: Synergizing Tree Search and Test-Time Reinforcement Learning for Optimization Modeling

ArXiv CS.AI 50%

arXiv:2606.15197v2 Announce Type: replace-cross Abstract: Optimization modeling is inherently hierarchical, requiring a precise sequence of symbolic commitments. Traditional learning-based automated optimization modeling methods improve modeling policies through large-scale...

<https://arxiv.org/abs/2606.15197>

9 CONF

The Tao of Agency: Autotelic AI, Embedded Agency and Dissolution of the Self

ArXiv CS.AI 50%

arXiv:2606.19924v1 Announce Type: new Abstract: Most artificial intelligence systems are built on the assumption that goals are exogenous and specified by the designer. Exploring what happens when an agent begins generating its own goals opens the field of autotelic AI. Agents...

<https://arxiv.org/abs/2606.19924>

0 CONF

Autonomous Event-Driven Multi-Agent Orchestration for Enterprise AI at Scale

ArXiv CS.AI 50%

arXiv:2606.20058v1 Announce Type: new Abstract: Enterprise AI aims to move toward continuous event monitoring, detection, and action across specialist agents, yet existing multi-agent systems largely assume discrete request-response workflows and remain underexplored at...

<https://arxiv.org/abs/2606.20058>

1 CONF

Augmenting Game AI with Deep Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.20210v1 Announce Type: new Abstract: Immersion in video games depends not only on graphics, audio, and game mechanics, but also on the quality of in-game characters. Producing believable characters, or game AI, remains a significant challenge as behavioral complexity...

<https://arxiv.org/abs/2606.20210>

2 CONF

Residual-Space Evolutionary Optimization via Flow-based Generative Models

ArXiv CS.AI 50%

arXiv:2606.20084v1 Announce Type: new Abstract: Data editing with generative methods typically requires differentiable objectives and gradient-based search. However, these assumptions break down in flow-based settings, where edits are performed through forward and backward...

<https://arxiv.org/abs/2606.20084>

3 CONF

Learning to Prompt: Improving Student Engagement with Adaptive LLM-based High-School Tutoring

ArXiv CS.AI 50%

arXiv:2606.20138v1 Announce Type: new Abstract: LLMs can personalize education, although current static-prompt tutoring systems struggle to adapt to diverse academic disciplines. We develop and test a system with subject-aware prompting, based on 14 pedagogical features (e.g.,...

<https://arxiv.org/abs/2606.20138>

4 CONF

BIM-Edit: Benchmarking Large Language Models for IFC-Based Building Information Modeling

ArXiv CS.AI 50%

arXiv:2606.20146v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly applied to computer-aided design (CAD) to generate design artifacts from textual instructions. In engineering practice, this requires more than creating new geometry, models must also...

<https://arxiv.org/abs/2606.20146>

5 CONF

Apparent Psychological Profiles of Large Language Models are Largely a Measurement Artifact

ArXiv CS.AI 50%

arXiv:2606.20205v1 Announce Type: new Abstract: Psychological instruments designed for humans are increasingly used to assign large language models (LLMs) stable psychological profiles that affect their usability, safety assessment, and use as proxies for human participants in...

<https://arxiv.org/abs/2606.20205>

6 CONF

A Multi-Agent system for Multi-Objective constrained optimization

ArXiv CS.AI 50%

arXiv:2606.20236v1 Announce Type: new Abstract: Many decision-making problems in computing and networking systems can be naturally formulated as cost-minimization problems under performance constraints. In dynamic environments, reinforcement learning (RL) is often used to solve...

<https://arxiv.org/abs/2606.20236>

7 CONF

Navigating Unreliable Parametric and Contextual Knowledge: Explicit Knowledge Conflict Resolution for LLM Inference

ArXiv CS.AI 50%

arXiv:2606.20245v1 Announce Type: new Abstract: Large language models (LLMs) have achieved strong performance across a wide range of language-based tasks by leveraging both extensive parametric knowledge and in-context learning ability, enabling them to incorporate external...

<https://arxiv.org/abs/2606.20245>

8 CONF

Leveraging systems' non-linearity to tackle the scarcity of data in the design of Intelligent Fault Diagnosis Systems

ArXiv CS.AI 50%

arXiv:2606.20323v1 Announce Type: new Abstract: Deep Transfer Learning (DTL) allows for the efficient building of Intelligent Fault Diagnosis Systems (IFDS). On the other hand, DTL methods still heavily rely on large amounts of labelled data. Obtaining such an amount of data can...

<https://arxiv.org/abs/2606.20323>

9 CONF

Rethinking Shrinkage Bias in LLM FP4 Pretraining: Geometric Origin, Systemic Impact, and UFP4 Recipe

ArXiv CS.AI 50%

arXiv:2606.20381v1 Announce Type: new Abstract: FP4 training promises substantial reductions in memory and computation cost for LLM pretraining, yet current FP4 hardware paths and recipes, including NVIDIA Blackwell/Rubin-class systems and AMD MI350-series GPUs, remain centered...

<https://arxiv.org/abs/2606.20381>

10 CONF

Interpretable Sperm Morphology Classification via Attention-Guided Deep Learning

ArXiv CS.AI 50%

arXiv:2606.20438v1 Announce Type: new Abstract: Male infertility is a major cause of couple infertility, often linked to abnormal sperm morphology. While deep learning models offer automated analysis, most lack interpretability, limiting their clinical adoption. This study...

<https://arxiv.org/abs/2606.20438>

11 CONF

What Do Safety-Aligned LLMs Learn From Mixed Compliance Demonstrations?

ArXiv CS.AI 50%

arXiv:2606.20508v1 Announce Type: new Abstract: Prior work has shown that in-context demonstrations can jailbreak language models, but it remains unclear how models interpret different types of compliance demonstrations. We study this by mixing benign compliance demonstrations...

<https://arxiv.org/abs/2606.20508>

2 **CONF****LedgerAgent: Structured State for Policy-Adherent Tool-Calling Agents****ArXiv CS.AI 50%**

arXiv:2606.20529v1 Announce Type: new Abstract: Policy-adherent tool-calling agents in customer-service domains must maintain task states across turns while calling tools and obeying domain policies. Task states consist of relevant facts, identifiers, constraints, and conditions...

<https://arxiv.org/abs/2606.20529>

3 **CONF****Human-AI Agent Interaction in a Business Context****ArXiv CS.AI 50%**

arXiv:2606.18716v1 Announce Type: cross Abstract: As AI agents are increasingly integrated into core business processes, understanding and designing effective interaction patterns between humans and AI agents becomes crucial for value creation. This study identifies and...

<https://arxiv.org/abs/2606.18716>

4 **CONF****Exposing the Unsaid: Visualizing Hidden LLM Bias through Stochastic Path Aggregation****ArXiv CS.AI 50%**

arXiv:2606.19344v1 Announce Type: cross Abstract: Large Language Models (LLMs) exhibit representational and syntactic biases that are difficult to evaluate due to the stochastic nature of text generation. Standard auditing methods rely on a single output inspection or static...

<https://arxiv.org/abs/2606.19344>

5 **CONF****Disentangling Linguistic Relatedness from Task Alignment in Cross-Lingual Transfer****ArXiv CS.AI 50%**

arXiv:2606.19346v1 Announce Type: cross Abstract: We study cross-lingual transfer by fine-tuning seven large language models (4B--671B parameters) on Arabic and evaluating zero-shot reading comprehension on Semitic languages and non-Semitic controls. Across dense and...

<https://arxiv.org/abs/2606.19346>

6 **CONF****How LLMs Fail and Generalize in RTL Coding for Hardware Design?****ArXiv CS.AI 50%**

arXiv:2606.19347v1 Announce Type: cross Abstract: Translating sequential programming priors into the parallel temporal logic of hardware design remains a crucial bottleneck for large language models(LLM). To investigate this, we introduce a new error taxonomy grounded in problem...

<https://arxiv.org/abs/2606.19347>

7 **CONF****Where to Place the Query? Unveiling and Mitigating Positional Bias in In-Context Learning for Diffusion LLMs via Decoding Dynamics****ArXiv CS.AI 50%**

arXiv:2606.19349v1 Announce Type: cross Abstract: While In-Context Learning (ICL) is extensively studied in Autoregressive (AR) LLMs, its mechanism within Diffusion Large Language Models (dLLMs) remains largely unexplored. Unlike AR models restricted by unidirectional causal...

<https://arxiv.org/abs/2606.19349>

8 CONF

Sign-Language Datasets at Scale: A Comprehensive Survey on Resources, Benchmarks, and Annotation Standards

ArXiv CS.AI 50%

arXiv:2606.19352v1 Announce Type: cross Abstract: Sign languages are expressive visual languages used by Deaf and Hard-of-Hearing (DHH) communities. Despite substantial progress in sign-language recognition, translation, and production, advances remain constrained by fragmented...

<https://arxiv.org/abs/2606.19352>

9 CONF

ProMUSE: Progressive Multi-modal Uncertainty-guided Staged Evidential Alzheimer Disease Classification

ArXiv CS.AI 50%

arXiv:2606.19371v1 Announce Type: cross Abstract: Alzheimer's disease (AD) is a fatal disorder that destroys memory and cognitive skills in the elderly population. Most treatments for AD are effective in the early stage, leading to an increasing demand for early AD diagnosis. AD...

<https://arxiv.org/abs/2606.19371>

10 CONF

cAPM: Continual AI-Assisted Pace-Mapping with Active Learning

ArXiv CS.AI 50%

arXiv:2606.19373v1 Announce Type: cross Abstract: Ventricular tachycardia is a life-threatening rhythm disorder and a major cause of sudden cardiac death. Pace-mapping is a clinical procedure for identifying the intervention target during catheter ablation of VT. It requires...

<https://arxiv.org/abs/2606.19373>

11 CONF

Cost-Optimal LLM Routing with Limited User Feedback under User Satisfaction Guarantees

ArXiv CS.AI 50%

arXiv:2606.19376v1 Announce Type: cross Abstract: Inference costs for large language model (LLM) applications are rapidly growing, driven by surging demand and rising infrastructure cost. Users expect high-quality responses, and in commercial settings this is formally codified...

<https://arxiv.org/abs/2606.19376>

12 CONF

How Linear Is a Transformer Feed-Forward Block? Per-Block Linear Recoverability Is Learned, Not Architectural

ArXiv CS.AI 50%

arXiv:2606.19379v1 Announce Type: cross Abstract: Transformer feed-forward networks (FFNs) are often treated as nonlinear stores of computation, yet how nonlinear a trained FFN block actually is has rarely been measured. We treat each FFN as a position-wise input-to-output map...

<https://arxiv.org/abs/2606.19379>

13 CONF

Interpretable and Verifiable Hardware Generation with LLM-Driven Stepwise Refinement

ArXiv CS.AI 50%

arXiv:2606.19387v1 Announce Type: cross Abstract: Large language models (LLMs) have achieved remarkable success in software development. However, they are susceptible to hallucinations, meaning that they can introduce subtle semantic and logical errors. Due to the high stakes in...

<https://arxiv.org/abs/2606.19387>

- 4 **CONF**
- Execution-bound advisory automation for agentic AI: a reproducible AIBOM-driven CSAF-VEX framework**
- ArXiv CS.AI 50%
- arXiv:2606.19390v1 Announce Type: cross Abstract: A protocol driven framework is presented that binds SBOM and AIBOM artefacts to deterministic environment capture and structured runtime telemetry. Exploitability is computed from declared artefacts, observed activation...
- <https://arxiv.org/abs/2606.19390>
-
- 5 **CONF**
- Evaluation of EEG Foundation Models for Event-Based Burst-Suppression Detection in ICU**
- ArXiv CS.AI 50%
- arXiv:2606.20074v1 Announce Type: cross Abstract: Burst suppression (BS) is a clinically relevant electroencephalographic (EEG) pattern used to monitor sedation depth and brain activity in critically ill patients, particularly during induced coma in Intensive Care Units (ICUs)....
- <https://arxiv.org/abs/2606.20074>
-
- 6 **CONF**
- Secure Coding Drift in LLM-Assisted Post-Quantum Cryptography Development: A Gamified Fix**
- ArXiv CS.AI 50%
- arXiv:2606.19474v1 Announce Type: cross Abstract: The transition to Post Quantum Cryptography (PQC) introduces considerable implementation complexity, requiring strict adherence to constant-time execution, side channel resistance, and precise parametrisation. Simultaneously,...
- <https://arxiv.org/abs/2606.19474>
-
- 7 **CONF**
- Techniques for Peak Memory Reduction for LoRA Fine-tuning of LLMs on Edge Devices**
- ArXiv CS.AI 50%
- arXiv:2606.19528v1 Announce Type: cross Abstract: Fine-tuning of Large Language Models (LLMs) using Low-Rank Adaptation (LoRA) on an end-user's data offers personalized experiences while keeping data private, but faces severe memory constraints on consumer hardware. Peak memory...
- <https://arxiv.org/abs/2606.19528>
-
- 8 **CONF**
- VOiLA: Vectorized Online Planning with Learned Diffusion Model for POMDP Agents**
- ArXiv CS.AI 50%
- arXiv:2606.19729v1 Announce Type: cross Abstract: Planning under uncertainty is an essential capability for autonomous robots. The Partially Observable Markov Decision Process (POMDP) provides a powerful framework for such a capability. Although POMDP-based planning has advanced...
- <https://arxiv.org/abs/2606.19729>
-
- 9 **CONF**
- PrefSQA: Pairwise Preference Prediction for Speech Quality Assessment and the Critical Role of High Quality Datasets**
- ArXiv CS.AI 50%
- arXiv:2606.19597v1 Announce Type: cross Abstract: Mean opinion scores (MOS) are widely used for speech quality assessment, yet scalar labels are sensitive to rater variability and listening test differences. This introduces labeling noise, which limits the reliability of MOS...
- <https://arxiv.org/abs/2606.19597>
-

0 CONF

FAPO: Fully Autonomous Prompt Optimization of Multi-Step LLM Pipelines

ArXiv CS.AI 50%

arXiv:2606.19605v1 Announce Type: cross Abstract: Multi-step LLM pipelines fail through interactions among retrieval, reasoning, and formatting steps, so prompt-only optimization can miss bottlenecks in the chain. We present FAPO (Fully Autonomous Prompt Optimization), a...

<https://arxiv.org/abs/2606.19605>

1 CONF

VCG: A Multimodal Retrieval Framework for E-Commerce Video Feeds under Extreme Cold-Start Conditions

ArXiv CS.AI 50%

arXiv:2606.19627v1 Announce Type: cross Abstract: The digital commerce landscape is shifting from static, search-driven catalogs to dynamic, immersive video feeds. This transition introduces an ``extreme cold-start" problem: unlike traditional items, new short-form videos lack...

<https://arxiv.org/abs/2606.19627>

2 CONF

CTS-MoE: Implicit Terrain Adaptation via Mixture-of-Experts for Perceptive Locomotion

ArXiv CS.AI 50%

arXiv:2606.19633v1 Announce Type: cross Abstract: Perceptive legged locomotion over discontinuous terrain (e.g., stairs, gaps, and obstacles) requires adaptive behavior, as a single conservative gait cannot produce the anticipatory maneuvers needed for abrupt topology changes....

<https://arxiv.org/abs/2606.19633>

3 CONF

LOKI: Memory-Free Null-Space Constrained Lifelong Knowledge Editing

ArXiv CS.AI 50%

arXiv:2606.19679v1 Announce Type: cross Abstract: Lifelong knowledge editing aims to efficiently and sequentially update language models over time, as new knowledge becomes available or when the model makes mistakes, while preserving acceptable performance on past knowledge. One...

<https://arxiv.org/abs/2606.19679>

4 CONF

Efficiently Representing Algorithms With Chain-of-Thought Transformers

ArXiv CS.AI 50%

arXiv:2606.19697v1 Announce Type: cross Abstract: The increasing popularity of \emph{reasoning} models -- language models that output a series of reasoning or thought tokens before producing an answer -- is justified, in part, by theoretical results showing that chain-of-thought...

<https://arxiv.org/abs/2606.19697>

5 CONF

FineREX: Fine-Tuned NER-RE for Human Smuggling Knowledge Graphs

ArXiv CS.AI 50%

arXiv:2606.19710v1 Announce Type: cross Abstract: Court proceedings contain valuable evidence about human smuggling networks, but this information is often buried within unstructured, jargon-heavy legal documents. While large language models (LLMs) can support knowledge graph...

<https://arxiv.org/abs/2606.19710>

6 CONF

OnDeFog: Online Decision Transformer under Frame Dropping

ArXiv CS.AI 50%

arXiv:2606.19721v1 Announce Type: cross Abstract: In challenging real-world reinforcement learning applications, communication delays or sensor failures often cause frame dropping, in which the agent cannot receive the dropped states and associated rewards. To address the...

<https://arxiv.org/abs/2606.19721>

7 CONF

Library-Aware Doubles and Iterative Repair for Large Language Model-Generated Unit Tests in OpenSIL Firmware

ArXiv CS.AI 50%

arXiv:2606.19725v1 Announce Type: cross Abstract: Validating changes in low-level C firmware is expensive because unit tests (UTs) are fragile under strict build constraints, where missing headers, unresolved symbols, and dependency mismatches frequently prevent compilation and...

<https://arxiv.org/abs/2606.19725>

8 CONF

QueryGaussian: Scalable and Training-Free Open-Vocabulary 3D Instance Retrieval

ArXiv CS.AI 50%

arXiv:2606.19733v1 Announce Type: cross Abstract: Efficiently retrieving specific 3D instances from large-scale scenes via natural language prompts remains a formidable challenge in multimedia analysis. Existing approaches predominantly follow a "scene-level embedding" paradigm,...

<https://arxiv.org/abs/2606.19733>

9 CONF

SafeSpec: Fast and Safe LLM via Dynamic Reflective Sampling

ArXiv CS.AI 50%

arXiv:2606.19755v1 Announce Type: cross Abstract: Speculative inference accelerates large language model (LLM) decoding but provides no inherent safety guarantees. Existing safety defenses are largely incompatible with speculative inference: they either introduce additional...

<https://arxiv.org/abs/2606.19755>

0 CONF

Data Standards for Humanoid Robotics: The Missing Infrastructure for Physical AI

ArXiv CS.AI 50%

arXiv:2606.19769v1 Announce Type: cross Abstract: The scalability of humanoid robots will depend not only on models and hardware, but also on whether physical experience can accumulate across robots, tasks, organizations, and time. Drawing on the authors' work in developing...

<https://arxiv.org/abs/2606.19769>

1 CONF

Towards Engineering Scaling Laws with Pretraining Data Composition

ArXiv CS.AI 50%

arXiv:2606.19781v1 Announce Type: cross Abstract: Neural scaling laws describe how model performance improves as a power law in compute, model size, and dataset size. While well-established for large language models, these relationships are emerging for large models in particle...

<https://arxiv.org/abs/2606.19781>

- 2

CONF

Improving End-to-End Speech Recognition for Dysarthric Speech through In-Domain Data Augmentation

ArXiv CS.AI

50%

arXiv:2606.19797v1 Announce Type: cross Abstract: Dysarthric speech recognition is crucial for facilitating effective communication among individuals with dysarthria. However, accurately recognizing dysarthric speech poses significant challenges due to varying severity levels...

<https://arxiv.org/abs/2606.19797>
- 3

CONF

Policy-aware Vector Search: A Vision for Fine Grained Access Control in Vector Databases

ArXiv CS.AI

50%

arXiv:2606.19803v1 Announce Type: cross Abstract: Vector databases are increasingly used in security sensitive contexts with Retrieval Augmented Generation and organizational AI pipelines; however, their security capabilities remain limited. Specifically, Fine-grained Access...

<https://arxiv.org/abs/2606.19803>
- 4

CONF

Uncertainty-Aware Reward Modeling for Stable RLHF

ArXiv CS.AI

50%

arXiv:2606.19818v1 Announce Type: cross Abstract: Reinforcement learning from human feedback (RLHF) aligns large language models by training reward models on preference data and optimizing policies to maximize predicted rewards. However, this pipeline faces two fundamental...

<https://arxiv.org/abs/2606.19818>
- 5

CONF

CREDENCE: Claim Reduction for Decomposition & Enhanced Credibility -- Semantic Metrics and Convergence Analysis

ArXiv CS.AI

50%

arXiv:2606.19819v1 Announce Type: cross Abstract: Decomposing compound sentences into atomic, verifiable claims is a prerequisite for reliable automated fact-checking. Prior work has relied on token-overlap (Jaccard) metrics that systematically underestimate decomposition...

<https://arxiv.org/abs/2606.19819>
- 6

CONF

Measuring Biological Capabilities and Risks of AI Agents

ArXiv CS.AI

50%

arXiv:2606.19899v1 Announce Type: cross Abstract: This paper addresses a rapidly emerging policy challenge: how to generate and interpret credible evidence about the biological capabilities and risks of AI scientists, or agentic AI systems capable of autonomously or...

<https://arxiv.org/abs/2606.19899>
- 7

CONF

Spatial-Aware Reduction Framework: Towards Efficient and Faithful Visual State Space Models

ArXiv CS.AI

50%

arXiv:2606.19932v1 Announce Type: cross Abstract: Mamba demonstrates strong efficiency in modeling long visual sequences. However, when token reduction is applied to structurally enhanced Mamba variants, these models exhibit a severe performance collapse. We attribute this...

<https://arxiv.org/abs/2606.19932>

8 CONF

Triangular Consistency as a Universal Constraint for Learning Optical Flow

ArXiv CS.AI 50%

arXiv:2606.19938v1 Announce Type: cross Abstract: We propose triangular consistency as a first-principled constraint for optical flow, which is agnostic to network architecture, supervision type, and dataset, and applies to both image-pair and multi-frame settings. This simple...

<https://arxiv.org/abs/2606.19938>

9 CONF

Confidence Calibration for Multimodal LLMs: An Empirical Study through Medical VQA

ArXiv CS.AI 50%

arXiv:2606.19950v1 Announce Type: cross Abstract: Multimodal Large Language Models (MLLMs) show great potential in medical tasks, but their elicited confidence often misaligns with actual accuracy, potentially leading to misdiagnosis or overlooking correct advice. This study...

<https://arxiv.org/abs/2606.19950>

0 CONF

The Algorithmic-Human Manager: AI, Apps, and Workers in the Indian Gig Economy

ArXiv CS.AI 50%

arXiv:2606.19975v1 Announce Type: cross Abstract: This paper examines the impact of artificial intelligence and digital technologies on the blue-collar gig economy in India, focusing on algorithmic management. This paper examines the impact of artificial intelligence and digital...

<https://arxiv.org/abs/2606.19975>

1 CONF

Tri-Info: Generalizable, Interpretable Failure Prediction for VLA Models via Information Theory

ArXiv CS.AI 50%

arXiv:2606.19998v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models are increasingly deployed across diverse tasks, yet they remain black boxes whose physical interactions can cause irreversible harm, making generalizable and interpretable failure detection...

<https://arxiv.org/abs/2606.19998>

2 CONF

Superhuman Safe and Agile Racing through Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.22748v2 Announce Type: replace-cross Abstract: Autonomous systems have achieved superhuman performance in isolation or simulation, yet they remain brittle in shared, dynamic real-world spaces. This failure stems from the dominant single-agent paradigm for physical...

<https://arxiv.org/abs/2605.22748>

3 CONF

When Lower Privileges Suffice: Investigating Over-Privileged Tool Selection in LLM Agents

ArXiv CS.AI 50%

arXiv:2606.20023v1 Announce Type: cross Abstract: As LLM agents increasingly select tools autonomously, their choices among tools with different privileges become safety-relevant. However, prior tool-selection studies focus on safety-agnostic metadata preferences, leaving...

<https://arxiv.org/abs/2606.20023>

- 4
CONF

Variable-Length Tokenization via Learnable Global Merging for Diffusion Transformers

ArXiv CS.AI
50%

arXiv:2606.20076v1 Announce Type: cross Abstract: Latent Diffusion Models (LDMs) have become dominant in visual synthesis, but their quality-compute trade-off is largely constrained by the tokenizer's fixed compression ratio. Variable-length tokenizers (VLTs) promise adaptive...

<https://arxiv.org/abs/2606.20076>
- 5
CONF

IHUBERT: Vector-Based Semantic Deduplication and Domain-Balanced Pretraining for Persian Resources

ArXiv CS.AI
50%

arXiv:2606.20089v1 Announce Type: cross Abstract: Persian pretrained language models (PLMs) are still limited by the scarcity of large-scale, high-quality pretraining corpora and by insufficient evaluation beyond standard classification and NER tasks. We present IHUBERT, a...

<https://arxiv.org/abs/2606.20089>
- 6
CONF

MakeupMirror: Improving Facial Attribute Preservation in Diffusion Models for Makeup Transfer

ArXiv CS.AI
50%

arXiv:2606.20094v1 Announce Type: cross Abstract: Makeup transfer models enable fun augmented reality (AR) experiences as well as virtual try-on (VTO) for online makeup shopping. While recent state-of-the-art diffusion based solutions such as Stable-Makeup dramatically improve...

<https://arxiv.org/abs/2606.20094>
- 7
CONF

Hybrid Diffusion Transformer for Instruction-Guided Audio Editing via Rectified Flow

ArXiv CS.AI
50%

arXiv:2606.20101v1 Announce Type: cross Abstract: Audio editing aims to modify specific content in an existing audio clip according to a natural language instruction while preserving the remaining acoustic content. Despite the remarkable progress of diffusion models, existing...

<https://arxiv.org/abs/2606.20101>
- 8
CONF

MedRLM: Recursive Multimodal Health Intelligence for Long-Context Clinical Reasoning, Sensor-Guided Screening, Evidence-Grounded Decision Support, and Community-to-Tertiary Referral Optimization

ArXiv CS.AI
50%

arXiv:2606.20164v1 Announce Type: cross Abstract: Real-world clinical decision support requires reasoning over heterogeneous and longitudinal patient information rather than answering isolated medical questions. However, current medical large language models and...

<https://arxiv.org/abs/2606.20164>
- 9
CONF

Evaluating and Enhancing Negation Comprehension in Remote Sensing MLLMs

ArXiv CS.AI
50%

arXiv:2606.20177v1 Announce Type: cross Abstract: Multimodal Large Language Models (MLLMs) have demonstrated remarkable success in various Remote Sensing (RS) tasks. However, their ability to comprehend negation remains underexplored, limiting deployment in real-world...

<https://arxiv.org/abs/2606.20177>

00 CONF

Modeling Day-Long ECG Signals to Predict Heart Failure Risk with Explainable AI

ArXiv CS.AI 50%

arXiv:2601.00014v2 Announce Type: replace-cross Abstract: Heart failure (HF) affects 11.8% of adults aged 65 and older, reducing quality of life and longevity. Preventing HF can reduce morbidity and mortality. We hypothesized that artificial intelligence (AI) applied to 24-hour...

<https://arxiv.org/abs/2601.00014>

01 CONF

ScholarQuest: A Taxonomy-Guided Benchmark for Agentic Academic Paper Search in Open Literature Environments

ArXiv CS.AI 50%

arXiv:2606.20235v1 Announce Type: cross Abstract: Academic paper search is a core step in scientific research, and LLM-based search agents are emerging as a promising paradigm for iterative, intent-driven literature exploration. However, existing benchmarks are insufficient for...

<https://arxiv.org/abs/2606.20235>

02 CONF

Editorial Alignment: A Participatory Approach to Engaging Editorial Expertise in LLM-mediated Knowledge Dissemination

ArXiv CS.AI 50%

arXiv:2606.20258v1 Announce Type: cross Abstract: The emergence of LLM-driven information services is reshaping the conditions under which public knowledge institutions operate, threatening to absorb the editorial function these institutions exist to exercise. While LLMs offer...

<https://arxiv.org/abs/2606.20258>

03 CONF

ELVA: Exploring Ranking-Driven Universal Multimodal Retrieval

ArXiv CS.AI 50%

arXiv:2606.20280v1 Announce Type: cross Abstract: Leveraging Multimodal Large Language Models (MLLMs) via contrastive learning has become a mainstream paradigm for improving the performance of Universal Multimodal Retrieval (UMR). However, previous works have ignored the grain...

<https://arxiv.org/abs/2606.20280>

04 CONF

Boundary Embedding Shaping with Adaptive Contrastive Learning for Graph Structural Disentanglement

ArXiv CS.AI 50%

arXiv:2606.20283v1 Announce Type: cross Abstract: Graph neural networks (GNNs) excel at aggregating neighbor information for classification, yet their performance is hindered by graph structural entanglement, where spurious correlations from semantically irrelevant neighbors...

<https://arxiv.org/abs/2606.20283>

05 CONF

Robust $\mathbb{Q}$ -learning for mean-field control under Wasserstein uncertainty in common noise

ArXiv CS.AI 50%

arXiv:2606.20356v1 Announce Type: cross Abstract: In this article, we present a robust $\mathbb{Q}$ -learning algorithm for discrete-time mean-field control problems under Wasserstein uncertainty in the common noise law. The algorithm combines a quantization-and-projection scheme with a...

<https://arxiv.org/abs/2606.20356>

06 CONF

AutoPass: Evidence-Guided LLM Agents for Compiler Performance Tuning

ArXiv CS.AI 50%

arXiv:2606.20373v1 Announce Type: cross Abstract: Large Language Models (LLMs) show promise for code compilation tasks, but applying them to runtime performance tuning is difficult due to complex microarchitectural effects and noisy runtime measurements. We present AutoPass, a...

<https://arxiv.org/abs/2606.20373>

07 CONF

LLM agent safety, multi-turn red-teaming, jailbreak benchmarks, adversarial robustness, safety-critical systems

ArXiv CS.AI 50%

arXiv:2606.20408v1 Announce Type: cross Abstract: Large language model (LLM) agents are increasingly proposed as supervisory components for safety-critical systems, yet their robustness under sustained, adaptive adversarial pressure remains poorly characterized. We present...

<https://arxiv.org/abs/2606.20408>

08 CONF

Repurposing a Speech Classifier for Guided Diffusion-Based Speech Generation

ArXiv CS.AI 50%

arXiv:2606.20457v1 Announce Type: cross Abstract: Classifier guidance is a way to control diffusion generation by using a noise-conditioned classifier to steer the sampling process toward a target class. One drawback of classifier guidance is that it requires two separately...

<https://arxiv.org/abs/2606.20457>

09 CONF

Analyzing Defensive Misdirection Against Model-Guided Automated Attacks on Agentic AI Systems

ArXiv CS.AI 50%

arXiv:2606.20470v1 Announce Type: cross Abstract: Agentic AI systems increasingly rely on language-model components to interpret instructions, process external data, invoke tools, and coordinate with other agents. These capabilities make prompt-injection and jailbreak attacks...

<https://arxiv.org/abs/2606.20470>

10 CONF

Calibration Without Comprehension: Diagnosing the Limits of Fine-Tuning LLMs for Vulnerability Detection in Systems Software

ArXiv CS.AI 50%

arXiv:2606.20502v1 Announce Type: cross Abstract: Whether LLMs scoring well on vulnerability benchmarks genuinely reason about security or merely pattern-match on contaminated data remains unresolved. We present CWE-Trace, a framework for LLM vulnerability detection built from...

<https://arxiv.org/abs/2606.20502>

11 CONF

How Transparent is DiffusionGemma?

ArXiv CS.AI 50%

arXiv:2606.20560v1 Announce Type: cross Abstract: LLM reasoning transparency is a critical affordance for understanding model decisions, mitigating misuse and misalignment, and debugging surprising model behaviors. However, DiffusionGemma performs a larger fraction of its...

<https://arxiv.org/abs/2606.20560>

12 CONF

Creativity Reconsidered: Generative AI and the Problem of Intentional Agency

ArXiv CS.AI 50%

arXiv:2601.15797v2 Announce Type: replace Abstract: Many theorists maintain that conscious intentional agency is a necessary condition of creativity. We argue that this requirement, which we call the Intentional Agency Condition (IAC), should be abandoned. We motivate this by...

<https://arxiv.org/abs/2601.15797>

13 CONF

PCBSchemaGen: Reward-Guided LLM Code Synthesis for Printed Circuit Boards (PCB) Schematic Design with Structured Verification

ArXiv CS.AI 50%

arXiv:2602.00510v2 Announce Type: replace Abstract: Most LLM code-synthesis benchmarks rely on unit tests as the reward oracle, but PCB schematic design has none: correctness is defined by structured physical constraints over real IC packages and pin-level assignments, per-task...

<https://arxiv.org/abs/2602.00510>

14 CONF

PrototypeNAS: Rapid Design of Deep Neural Networks for Microcontroller Units

ArXiv CS.AI 50%

arXiv:2603.15106v2 Announce Type: replace Abstract: Enabling efficient deep neural network (DNN) inference on edge devices with different hardware constraints is a challenging task that typically requires DNN architectures to be specialized for each device separately. To avoid...

<https://arxiv.org/abs/2603.15106>

15 CONF

CareTransition-Audit: A Benchmark to Audit Discharge Summaries for Efficient Care Transitions

ArXiv CS.AI 50%

arXiv:2604.05435v2 Announce Type: replace Abstract: Incomplete or inconsistent discharge documentation drives care fragmentation and avoidable readmissions. Despite its critical role in patient safety, auditing discharge summaries relies on manual review and does not scale. We...

<https://arxiv.org/abs/2604.05435>

16 CONF

Science Earth: Towards A Planet-Scale Operating System for AI-Native Scientific Discovery

ArXiv CS.AI 50%

arXiv:2606.01316v2 Announce Type: replace Abstract: Scientific discovery demands intelligence, perseverance, and serendipity across vast search spaces. Today, top scientific capabilities remain siloed--one AI system for biological analysis, another for clinical reasoning,...

<https://arxiv.org/abs/2606.01316>

17 CONF

RetailBench: Benchmarking long horizon reasoning and coherent decision making of LLM agents in realistic retail environments

ArXiv CS.AI 50%

arXiv:2606.15862v2 Announce Type: replace Abstract: Large language model (LLM) agents have made rapid progress on short-horizon, well-scoped tasks, yet their ability to sustain coherent decisions in dynamic long-horizon environments remains uncertain. We introduce RetailBench, a...

<https://arxiv.org/abs/2606.15862>

18 **CONF****Searching for Synergy in Shared Workspace Human-AI Collaboration****ArXiv CS.AI** **50%**

arXiv:2606.18413v2 Announce Type: replace Abstract: Automated AI agents are increasingly capable, yet many scientific and professional tasks require human judgment and contextual expertise. We study shared-workspace human-AI teams, where AI agents and human collaborators must...

<https://arxiv.org/abs/2606.18413>

19 **CONF****Wisdom of Committee: Diverse Distillation from Large Foundation Models and Domain Experts****ArXiv CS.AI** **50%**

arXiv:2402.14035v4 Announce Type: replace-cross Abstract: Knowledge distillation from foundation models to compact domain models is challenging due to substantial gaps in capacity, architecture, and modality. For example, in our experiments, distilling from a 76M-parameter...

<https://arxiv.org/abs/2402.14035>

20 **CONF****Global Ease of Living Index: a machine learning framework for longitudinal analysis of major economies****ArXiv CS.AI** **50%**

arXiv:2502.06866v3 Announce Type: replace-cross Abstract: The drastic changes in the global economy, geopolitical conditions, and disruptions such as the COVID-19 pandemic have impacted the cost of living and quality of life. It is essential to comprehend the long-term...

<https://arxiv.org/abs/2502.06866>

21 **CONF****Simulation of Language Evolution under Regulated Social Media Platforms: A Synergistic Approach of Large Language Models and Genetic Algorithms****ArXiv CS.AI** **50%**

arXiv:2502.19193v2 Announce Type: replace-cross Abstract: Social media platforms frequently impose restrictive policies to moderate user content, prompting the emergence of creative evasion language strategies. This paper presents a multi-agent framework based on Large Language...

<https://arxiv.org/abs/2502.19193>

22 **CONF****Bridging Distribution Shift and AI Safety: Conceptual and Methodological Synergies****ArXiv CS.AI** **50%**

arXiv:2505.22829v2 Announce Type: replace-cross Abstract: This paper bridges distribution shift and AI safety through a comprehensive analysis of their conceptual and methodological synergies. While prior discussions often focus on narrow cases or informal analogies, we...

<https://arxiv.org/abs/2505.22829>

23 **CONF****Oranits: Mission Assignment and Task Offloading in Open RAN-based ITS using Metaheuristic and Deep Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2507.19712v3 Announce Type: replace-cross Abstract: In this paper, we explore mission assignment and task offloading in an Open Radio Access Network (Open RAN)-based intelligent transportation system (ITS), where autonomous vehicles leverage mobile edge computing for...

<https://arxiv.org/abs/2507.19712>

24 **CONF****Charting the Future of Scholarly Knowledge with AI: A Community Perspective****ArXiv CS.AI** **50%**

arXiv:2509.02581v2 Announce Type: replace-cross Abstract: Despite the growing availability of tools designed to support scholarly knowledge extraction and organization, many researchers still rely on manual methods, sometimes due to unfamiliarity with existing technologies or...

<https://arxiv.org/abs/2509.02581>

25 **CONF****From Construction to Injection: Edit-Based Fingerprints for Large Language Models****ArXiv CS.AI** **50%**

arXiv:2509.03122v4 Announce Type: replace-cross Abstract: Reliable model fingerprints are essential for protecting large language models (LLMs) against unauthorized redistribution and commercial misuse. In black-box deployment, verification is hindered by defensive filtering of...

<https://arxiv.org/abs/2509.03122>

26 **CONF****Enhancing Generative Auto-bidding with Offline Reward Evaluation and Policy Search****ArXiv CS.AI** **50%**

arXiv:2509.15927v5 Announce Type: replace-cross Abstract: Auto-bidding is a critical tool for advertisers to improve advertising performance. Recent progress has demonstrated that AI-Generated Bidding (AIGB), which learns a conditional generative planner from offline data,...

<https://arxiv.org/abs/2509.15927>

27 **CONF****Bid Farewell to Seesaw: Towards Accurate Long-tail Session-based Recommendation via Dual Constraints of Hybrid Intents****ArXiv CS.AI** **50%**

arXiv:2511.08378v4 Announce Type: replace-cross Abstract: Session-based recommendation (SBR) aims to predict anonymous users' next interaction based on their interaction sessions. In the practical recommendation scenario, low-exposure items constitute the majority of...

<https://arxiv.org/abs/2511.08378>

28 **CONF****Policy-Embedded Graph Expansion: Networked HIV Testing with Diffusion-Driven Network Samples****ArXiv CS.AI** **50%**

arXiv:2601.16233v2 Announce Type: replace-cross Abstract: HIV is a retrovirus that attacks the human immune system and can lead to death without proper treatment. In collaboration with the WHO and the University of Witwatersrand, we study how to improve the efficiency of HIV...

<https://arxiv.org/abs/2601.16233>

29 **CONF****The MAMA-MIA Challenge: Advancing Generalizability and Fairness in Breast MRI Tumor Segmentation and Treatment Response Prediction****ArXiv CS.AI** **50%**

arXiv:2603.01250v2 Announce Type: replace-cross Abstract: Breast cancer is the most frequently diagnosed malignancy among women worldwide and a leading cause of cancer-related mortality. Dynamic contrast-enhanced magnetic resonance imaging plays a central role in tumor...

<https://arxiv.org/abs/2603.01250>

30 CONF

The ACUTE Protocol: Operationalizing Language Model Activations for Better Calibration, Utility, and Trust

ArXiv CS.AI 50%

arXiv:2606.07822v2 Announce Type: replace-cross Abstract: As language models improve and become increasingly deployed to solve a variety of tasks, trustworthiness becomes essential. Calibration is a good proxy for trust: well-calibrated confidence estimates help inform the risk...

<https://arxiv.org/abs/2606.07822>

31 CONF

Gaming-Resistant Insurance Contracts for Autonomous AI Agents: Strategy-Proof Toll Mechanism Design

ArXiv CS.AI 50%

arXiv:2606.16326v2 Announce Type: replace-cross Abstract: Paper A defines a time-consistent actuarial runtime that prices each side-effect-bearing action against a contractually fixed safe default and gates execution against a reserve budget. It treats the operator as passive....

<https://arxiv.org/abs/2606.16326>