



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	119	Media
HackerNews	7	Community
The Verge AI	5	Media
TechCrunch AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Skills for the future software profession: beyond agentic AI!

ArXiv CS.AI +1 source

80%

arXiv:2606.21894v2 Announce Type: replace-cross Abstract: As coding agents are rapidly changing software engineering, a natural question is: what are the core skills needed by future software engineers? To identify where software engineering is headed and thus what skills will...

<https://arxiv.org/abs/2606.21894>

VERI

Hollywood is bending the knee to OpenAI

The Verge AI +1 source

80%

Netflix, A24, Focus Features, and Warner Bros.' Clockwork have all reportedly decided to pass on picking up Artificial - director Luca Guadagnino's new biographical drama about OpenAI cofounder / CEO Sam Altman - for distribution deals. And while Neon and Mubi are still said to...

<https://www.theverge.com/entertainment/954899/luca-guadagnino-artificial-sam-altman-ama...>

CONFIRMED SIGNALS — 130 stories

CONF

Show HN: Shumai - open-source Frame.io alternative for creative work

HackerNews 50% HN 57

Shumai is an open source platform for uploading creative files, managing projects, collecting precise feedback, sharing work, and collaborating with AI agents, all in one simple creative-first workspace. I've always liked the product design of Frame.io, and I wanted to build an...

<https://github.com/shumaiOne/shumai>

CONF

Show HN: Transformer Primitives - A visual explainer you can send to anyone

HackerNews 50% HN 6

I have had a few conversations in the past year with non-technical folks (traditional finance types, consultants) who asked for a simple explainer on how GPTs work. These people generally have the horsepower to grasp new concepts quickly but don't necessarily have the math...

<https://understandgpt.xyz/>

CONF

Show HN: Cachet - A drop-in semantic cache for LLM APIs, 100% local, in Rust

HackerNews 50% HN 5

No summary — see link for full article.

<https://github.com/abhix2112/Cachet>

CONF

Show HN: FuckUI - AI hates your UI. Feed it the web text first

HackerNews 50% HN 4

No summary — see link for full article.

<https://fuckui.com>

CONF

SHOW HN: Huskmail - A disposable email with private inboxes and trial library

HackerNews 50% HN 4

No summary — see link for full article.

<https://huskmail.xyz>

CONF

Show HN: Optimal model routing directly in Claude, Codex and Cursor

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/workweave/router>

CONF

Show HN: Wordit - Change One Letter, Keep the Chain Going

HackerNews 50% HN 3

Hi everyone, I got this idea for a game where, starting from a four letter word you need to go as deep as you can in your vocabulary, changing only one letter per word. bear -> beer -> peer... Each correct word gives you 1 point Each incorrect word takes one life away from you,...

<https://victorribeiro.com/wordit/>

0

CONF

Plants appear to detect the patter of falling rain

MIT Tech Review 50%

MIT engineers have found the first direct evidence that plant seeds can sense sounds in nature: Rice submerged in shallow water germinated 30% to 40% more quickly when exposed to vibrations from water dripping on the surface. They think other types of seeds may respond...

<https://www.technologyreview.com/2026/06/23/1138288/plants-appear-to-detect-the-patter-...>

1

CONF

Why corporate AI super PACs spent \$27 million on a local election

The Verge AI 50%

Hello and welcome to Regulator, the newsletter for Verge subscribers chronicling the misadventures of their favorite tech overlords and Washington swamp creatures. ("Favorite" is, of course, subjective.) Not a subscriber yet? Sign up here, especially if you want the hot scoop on...

<https://www.theverge.com/policy/954970/ai-super-pacs-alex-bores-new-york-12th-district>

2

CONF

Something's off with Midjourney's pivot to body scanners

The Verge AI 50%

Last week, Midjourney, an AI startup best known for its image generator, made an unusual pivot: medical imaging. The company announced a futuristic ultrasound scanner that would dunk users into a vat of water and, hopefully, produce "something as powerful as MRI" yet "as casual..."

<https://www.theverge.com/report/954826/midjourney-medical-ai-ultrasound-body-scanner-la...>

3

CONF

The Fitbit Air takes a smarter approach to the AI health dumpster fire

The Verge AI 50%

Google Health Coach seems to think I'm on the verge of physical collapse. My sleep is not where it needs to be, hence my unimpressive readiness score. My heart rate variability, a measure of how recovered I am, is below baseline. I'm spending too much time in a hot, humid...

<https://www.theverge.com/tech/954768/google-fitbit-air-review-fitness-tracker-wearable-...>

4

CONF

Sony's AI Camera Assistant is exactly as bad as it looks

The Verge AI 50%

When Sony announced the Xperia 1 VIII last month, it promoted the phone by sharing some of the worst photos taken on a Sony camera in years. These weren't just any photos, though: they were taken with Sony's new AI Camera Assistant. After a week with the Xperia 1 VIII, I'm here...

<https://www.theverge.com/tech/954676/sony-xperia-1-viii-ai-camera-assistant>

5 CONF

RIFT-Bench: Dynamic Red-teaming For Agentic AI Systems

ArXiv CS.AI 50%

arXiv:2606.23927v1 Announce Type: new Abstract: Agentic AI systems powered by large language models (LLMs) are rapidly evolving into autonomous decision-making systems, exposing attack vectors beyond those of traditional LLM vulnerabilities. Existing security evaluations are...

<https://arxiv.org/abs/2606.23927>

6 CONF

HOLMES: Evaluating Higher-Order Logical Reasoning in LLMs

ArXiv CS.AI 50%

arXiv:2606.23238v2 Announce Type: replace Abstract: Logical reasoning is essential for reliable AI, yet existing benchmarks are largely first-order-logic-centric, focusing on object-level deduction over fixed predicates. This misses many realistic scenarios where models must...

<https://arxiv.org/abs/2606.23238>

7 CONF

ASALT: Adaptive State Alignment for Lateral Transfer in Multi-agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.24601v1 Announce Type: new Abstract: Multi-agent reinforcement learning (MARL) addresses the problem of training multiple agents that pursue collaborative, competitive, or mixed objectives. Prior work has investigated transfer learning between source and target...

<https://arxiv.org/abs/2606.24601>

8 CONF

Reinforcement Learning Towards Broadly and Persistently Beneficial Models

ArXiv CS.AI 50%

arXiv:2606.24014v1 Announce Type: new Abstract: As AI systems are deployed across increasingly diverse and high-stakes settings, model alignment must generalize beyond the tasks and domains seen during training. This is especially important for reinforcement learning (RL), which...

<https://arxiv.org/abs/2606.24014>

9 CONF

Can Language Model Agents be Helpful Circuit Explainers in Mechanistic Interpretability?

ArXiv CS.AI 50%

arXiv:2606.24026v1 Announce Type: new Abstract: Mechanistic interpretability has made substantial progress in automatically localizing circuits, but explaining what localized components do remains labor-intensive and difficult to standardize. In this work, we study whether...

<https://arxiv.org/abs/2606.24026>

0 CONF

Ensemble Feature Selection and Harris Hawks Optimization for Explainable Mental Health Risk Prediction in Female Sex Workers

ArXiv CS.AI 50%

arXiv:2606.24047v1 Announce Type: new Abstract: One of the significant mental health issues affecting female sex workers (FSWs) is mental disorders, especially depression. Exposure to violence, stigma, and economic hardship further increases their psychological risk. Current...

<https://arxiv.org/abs/2606.24047>

1 CONF

Beyond Trajectory Imitation: Strategy-Guided Policy Optimization for LLM Reasoning

ArXiv CS.AI 50%

arXiv:2606.24064v1 Announce Type: new Abstract: Distilling reasoning capabilities from strong to weak language models typically involves imitating specific solution trajectories, effectively transferring what to answer rather than how to reason. This trajectory-level imitation...

<https://arxiv.org/abs/2606.24064>

2 CONF

ARIA: Adaptive Region-Based Importance Allocation for Conditional Diffusion Distillation

ArXiv CS.AI 50%

arXiv:2606.23898v1 Announce Type: cross Abstract: Distilling conditional diffusion models aims to transfer the behavior of a large teacher to a smaller student while preserving alignment across conditioning inputs. Unlike recognition tasks, knowledge distillation in conditional...

<https://arxiv.org/abs/2606.23898>

3 CONF

T2D-Bench: Evidence-Gated Evaluation of LLM Outputs for Type 2 Diabetes Using a Multi-Layer Clinical-Lifestyle Knowledge Graph

ArXiv CS.AI 50%

arXiv:2606.24145v1 Announce Type: new Abstract: Large language models (LLMs) can produce clinically fluent recommendations for type 2 diabetes while failing to satisfy guideline constraints or explicitly justify lifestyle-related glycemic claims. We present T2D-Bench, a...

<https://arxiv.org/abs/2606.24145>

4 CONF

The Geometry Behind Diffusion and Flow Matching: Gradient Flows and Geodesics in Wasserstein Space

ArXiv CS.AI 50%

arXiv:2606.24157v1 Announce Type: new Abstract: The space $\mathcal{P}_2(\mathbb{R}^d)$ of probability measures with finite second moment carries a natural geometry: the quadratic Wasserstein distance W_2 makes it a complete metric space and, following Otto, a (formal)...

<https://arxiv.org/abs/2606.24157>

5 CONF

An Introduction to Causal Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.24160v1 Announce Type: new Abstract: Causal inference provides a set of principles and tools that allow one to combine data and knowledge about an environment to reason with questions of counterfactual nature, i.e., what would have happened had reality been different,...

<https://arxiv.org/abs/2606.24160>

6 CONF

Navigating User Behavior toward Personalized Multimodal Generation

ArXiv CS.AI 50%

arXiv:2606.24196v1 Announce Type: new Abstract: Modern AIGC pipelines deliver high-fidelity images and videos but presuppose a well-formed creation instruction, while end users rarely articulate visual details, leaving generators misaligned with user demand. We study...

<https://arxiv.org/abs/2606.24196>

7 CONF

Exploring the relationship between human-centric AI and firm idiosyncratic risks

ArXiv CS.AI 50%

arXiv:2606.24224v1 Announce Type: new Abstract: Despite the extensive discussions of human-centric AI (HCAI) in Industry 5.0, its effects on firms' idiosyncratic risks (IR) remains underexplored. This is an imperative issue for firms navigate financial risks during the current...

<https://arxiv.org/abs/2606.24224>

8 CONF

FlowR2A: Learning Reward-to-Action Distribution for Multimodal Driving Planning

ArXiv CS.AI 50%

arXiv:2606.24231v1 Announce Type: new Abstract: Multimodal driving planning faces a long-standing tension between two paradigms: scoring-based methods benefit from dense reward supervision but are confined to a fixed action vocabulary, while anchor-based methods generate...

<https://arxiv.org/abs/2606.24231>

9 CONF

Towards Federated Long-Tailed Graph Learning: An Energy-Guided Dual Decoupling Approach

ArXiv CS.AI 50%

arXiv:2606.24237v1 Announce Type: new Abstract: Federated Graph Learning facilitates collaborative graph modeling across distributed clients while preserving data privacy. However, real-world data categories frequently exhibit long-tailed distributions. Such statistical scarcity...

<https://arxiv.org/abs/2606.24237>

10 CONF

On the Smallness of the Large Language Models Scaling Exponents

ArXiv CS.AI 50%

arXiv:2606.24504v1 Announce Type: new Abstract: We discuss reasons why the scaling exponents of current Large Language Models (LLMs) applications are indicating an unsustainable regime in terms of energy resources. We further show that attributing the smallness of such exponents...

<https://arxiv.org/abs/2606.24504>

11 CONF

Prob-BBDM: a Probabilistic Brownian Bridge Diffusion Model for MRI sequence image-to-image translation

ArXiv CS.AI 50%

arXiv:2606.24313v1 Announce Type: new Abstract: AI-driven image-to-image synthesis is rapidly advancing, with growing applications in medical imaging. Multi-modal image analysis plays a crucial role in optimizing examination quality, yet acquiring multiple imaging modalities in...

<https://arxiv.org/abs/2606.24313>

12 CONF

FLUX3D: High-Fidelity 3D Gaussian Generation with Diffusion-Aligned Sparse Representation

ArXiv CS.AI 50%

arXiv:2606.24874v1 Announce Type: cross Abstract: Sparse voxel representation has emerged as a scalable foundation for image-to-3D Gaussian Splatting (3DGS) generation, yet current methods struggle to preserve high-frequency visual details of input images due to two structural...

<https://arxiv.org/abs/2606.24874>

3 CONF

When Helpfulness Overrides Causal Caution: Context-Dependent Suppression and Recovery in LLMs

ArXiv CS.AI 50%

arXiv:2606.24370v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly integrated into decision-support roles in business and policy contexts. While prior benchmark studies have primarily evaluated LLMs' causal reasoning capabilities, a more fundamental...

<https://arxiv.org/abs/2606.24370>

4 CONF

PHANTOM: A Large-Scale Dataset of Multimodal Adversarial Attacks for Vision-Language Models

ArXiv CS.AI 50%

arXiv:2606.24388v1 Announce Type: new Abstract: We introduce a large-scale, open-source dataset of pre-generated adversarial attacks for vision-language models (VLMs). The dataset is designed to be diverse, representative, and practical, extending existing benchmarks by covering...

<https://arxiv.org/abs/2606.24388>

5 CONF

Age of LLM: A Strategic 1v1 Benchmark for Reasoning, Diplomacy and Reliability of Large Language Models under Fog of War

ArXiv CS.AI 50%

arXiv:2606.24391v1 Announce Type: new Abstract: We introduce Age of LLM, a turn-based 1v1 benchmark in which two LLMs face off on a 13x7 grid to destroy the enemy base. Three stressors are deliberate: fog of war, full diplomacy (messages, ceasefires, ultimatums; uranium kept...

<https://arxiv.org/abs/2606.24391>

6 CONF

Agentic AI for Bilevel Long-Term Optimization of Policy-Driven Physical Layer Systems

ArXiv CS.AI 50%

arXiv:2606.24416v1 Announce Type: new Abstract: Network operators' changing policies, service requirements, and stringent real-time constraints render existing methods designed with fixed objectives and constraints ineffective. This paper presents Agentic long-term performance...

<https://arxiv.org/abs/2606.24416>

7 CONF

ReM-MoA: Reasoning Memory Sustains Mixture-of-Agents Scaling

ArXiv CS.AI 50%

arXiv:2606.24437v1 Announce Type: new Abstract: Mixture-of-Agents (MoA) architectures improve inference-time scaling by organizing multiple LLM agents into layered reasoning pipelines. However, existing MoA variants fail to sustain gains as depth increases, exhibiting...

<https://arxiv.org/abs/2606.24437>

8 CONF

CompressKV: Semantic-Retrieval-Guided KV-Cache Compression for Resource-Efficient Long-Context LLM Inference

ArXiv CS.AI 50%

arXiv:2606.24467v1 Announce Type: new Abstract: Long-context large language model (LLM) inference is increasingly constrained by the memory footprint and decoding cost of key-value (KV) caches, limiting sustainable deployment on resource-constrained hardware. Existing KV cache...

<https://arxiv.org/abs/2606.24467>

9 CONF

A specialized reasoning large language model for accelerating rare disease diagnosis: a randomized AI physician assistance trial

ArXiv CS.AI 50%

arXiv:2606.24510v1 Announce Type: new Abstract: Rare diseases affect millions of individuals worldwide, yet timely diagnosis remains a major public health challenge due to scarcity of specialized clinical expertise. While large language models (LLMs) show promise to support rare...

<https://arxiv.org/abs/2606.24510>

0 CONF

LLMs Prompted for Legal Context Object More: Overrefusal from Small On-Premises LLMs in Criminal Legal Context

ArXiv CS.AI 50%

arXiv:2606.24585v1 Announce Type: new Abstract: While the validity of LLMs' use in the legal context remains subject to ethical and legal debate, legal professionals are already experimenting with personal LLMs, if only for translation and reformulation. However, even such a...

<https://arxiv.org/abs/2606.24585>

1 CONF

AdversaBench: Automated LLM Red-Teaming with Multi-Judge Confirmation and Cross-Model Transferability

ArXiv CS.AI 50%

arXiv:2606.24589v1 Announce Type: new Abstract: Scaling adversarial evaluation of large language models requires both a method for generating hard inputs and a reliable way to confirm that resulting failures are real. We present AdversaBench, an end-to-end red-teaming pipeline...

<https://arxiv.org/abs/2606.24589>

2 CONF

Uncertainty-Aware Longitudinal Forecasting of Alzheimer's Disease Progression Using Deep Learning

ArXiv CS.AI 50%

arXiv:2606.24604v1 Announce Type: new Abstract: Longitudinal modelling of Alzheimer's disease progression is clinically useful only if it can describe not just the most likely next diagnosis, but how a patient may evolve over time and how reliable that forecast is. Most deep...

<https://arxiv.org/abs/2606.24604>

3 CONF

ScaleToT: Generalizing Structured LLM Reasoning for Billion-Scale Low-Activity User Modeling

ArXiv CS.AI 50%

arXiv:2606.24605v1 Announce Type: new Abstract: Accurate user modeling often depends on rich interaction histories, which are unavailable for billions of low-activity users. Large Language Models (LLMs) can infer latent user states from static profiles, but this reasoning...

<https://arxiv.org/abs/2606.24605>

4 CONF

AI Tokenomics: The Economics of Tokens, Computation, and Pricing in Foundation Models

ArXiv CS.AI 50%

arXiv:2606.24616v1 Announce Type: new Abstract: Tokens have become the practical accounting unit for modern foundation model services, linking information processing, computation, memory use, energy expenditure, pricing, and economic value. This paper develops a framework for AI...

<https://arxiv.org/abs/2606.24616>

- 5 **CONF**
- Themis: An explainable AI-enabled framework for Reinforcement Learning with Human Feedback**
- ArXiv CS.AI 50%
- arXiv:2606.24622v1 Announce Type: new Abstract: Training safe Reinforcement Learning (RL) systems is inherently challenging, with no guarantee of avoiding unwanted behaviors. The most effective defenses against this are (i) transparency through explainability and (ii) alignment...
- <https://arxiv.org/abs/2606.24622>
-
- 6 **CONF**
- Decentralised AI Training and Inference with BlockTrain**
- ArXiv CS.AI 50%
- arXiv:2606.24722v1 Announce Type: new Abstract: Frontier AI training is increasingly shaped by access to dense, centrally controlled accelerator clusters. This creates a structural advantage for hyperscalers and large centralized laboratories, and makes open or independent AI...
- <https://arxiv.org/abs/2606.24722>
-
- 7 **CONF**
- Scaling Laws for Task-Specific LLM Distillation**
- ArXiv CS.AI 50%
- arXiv:2606.24747v1 Announce Type: new Abstract: Large Language Models (LLMs) achieve strong performance across a growing range of domains, yet their scale poses deployment challenges in applications where latency and cost constraints are critical. This paper derives empirical...
- <https://arxiv.org/abs/2606.24747>
-
- 8 **CONF**
- Assessing Distribution Shift in Human Activity Recognition for Domain Generalization**
- ArXiv CS.AI 50%
- arXiv:2606.24781v1 Announce Type: new Abstract: While the field of Human Activity Recognition (HAR) continues to draw interest from researchers and advance in important ways, some key challenges remain. One of the most difficult aspects of building HAR models that show good...
- <https://arxiv.org/abs/2606.24781>
-
- 9 **CONF**
- Accuracy and Satisfaction in Multi-Turn LLM Dialogues for NFR Assessment**
- ArXiv CS.AI 50%
- arXiv:2606.24834v1 Announce Type: new Abstract: LLM-based dialogue assistants have become mainstream tools for software developers, yet current evaluation benchmarks focus exclusively on functional correctness. This leaves a critical gap in assessing the quality and accuracy of...
- <https://arxiv.org/abs/2606.24834>
-
- 0 **CONF**
- Matching Tasks to Objectives: Fine-Tuning and Prompt-Tuning Strategies for Encoder-Decoder Pre-trained Language Models**
- ArXiv CS.AI 50%
- arXiv:2606.24841v1 Announce Type: new Abstract: Prompt-based learning has emerged as a dominant paradigm in natural language processing. This study explores the impact of diverse pre-training objectives on the performance of encoder-decoder pre-trained language models across...
- <https://arxiv.org/abs/2606.24841>
-

1 **CONF****Reentrant value fields as delayed coupled reaction-diffusion systems on finite graphs****ArXiv CS.AI 50%**

arXiv:2605.03940v4 Announce Type: cross Abstract: We describe a dynamical system in which a symbolic field is coupled to a geometric field via a bipartite Hilbert-Schmidt kernel. The system is fully described by a retarded functional differential equation (RFDE) on the history...

<https://arxiv.org/abs/2605.03940>

2 **CONF****Beyond the Autoregressive Horizon: A Comprehensive Survey of Diffusion Models, World Modelling, and State Space Models for Code****ArXiv CS.AI 50%**

arXiv:2606.23690v1 Announce Type: cross Abstract: Autoregressive (AR) language models have driven significant progress in automated software engineering, enabling powerful code generation and assistance systems. However, the next-token prediction paradigm introduces structural...

<https://arxiv.org/abs/2606.23690>

3 **CONF****Quantifying Prior Dominance in RAG Systems****ArXiv CS.AI 50%**

arXiv:2606.23695v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) grounds Large Language Models in external knowledge, yet current evaluations rely on discrete heuristics that suffer from "epistemic blindness" - failing to distinguish genuine contextual...

<https://arxiv.org/abs/2606.23695>

4 **CONF****FP8 is All You Need (Part 2): Efficient Ozaki-Bailey Style FFT Through Tensor-core Garner Reformulation and Kulisch Escape Route****ArXiv CS.AI 50%**

arXiv:2606.23698v1 Announce Type: cross Abstract: NVIDIA's Blackwell Ultra (B300) cuts FP64 vector throughput to ~1.3 TFLOPS per GPU, roughly 30x below B200 and well below the level at which bandwidth-limited FP64 workloads stay memory-bound. The Ozaki Scheme II framework...

<https://arxiv.org/abs/2606.23698>

5 **CONF****Self-Recognition Finetuning can Prevent and Reverse Emergent Misalignment****ArXiv CS.AI 50%**

arXiv:2606.23700v1 Announce Type: cross Abstract: Emergent misalignment (EM) has been linked to the activation of misaligned persona vectors and evil character traits, suggesting that EM operates through disruption of the model's aligned character rather than direct learning of...

<https://arxiv.org/abs/2606.23700>

6 **CONF****Evaluating LLM Usage for Efficient and Explainable Numerical and Classified Implicit Sentiment Analysis of Product Desirability****ArXiv CS.AI 50%**

arXiv:2606.23701v1 Announce Type: cross Abstract: Qualitative product feedback can reveal nuanced user experiences, but its implicit sentiment is difficult to measure. This paper presents a scalable and interpretable framework that uses large language models (LLMs) to quantify...

<https://arxiv.org/abs/2606.23701>

7 CONF

Event-Aligned Analysis of Multi-Rater Pain Assessments Using Continuous Wearable Physiology

ArXiv CS.AI 50%

arXiv:2606.23705v1 Announce Type: cross Abstract: Pain is assessed differently by patients, nurses, and clinicians, yet most computational approaches assume a single ground-truth label - effectively ignoring who is doing the rating. We introduce a rater-aware, event-aligned...

<https://arxiv.org/abs/2606.23705>

8 CONF

Legal Reasoning Is Not Lawyering: Rethinking Legal Benchmarks for Pro Se Access to Justice

ArXiv CS.AI 50%

arXiv:2606.23716v1 Announce Type: cross Abstract: Legal AI benchmark research frequently invokes the assumption that large language models can improve access to justice, including for people who cannot access lawyers in order to understand and exercise their legal rights. We...

<https://arxiv.org/abs/2606.23716>

9 CONF

Weight-Space Geometry of Offline Reasoning Training

ArXiv CS.AI 50%

arXiv:2606.23740v1 Announce Type: cross Abstract: Offline reinforcement-learning losses (RFT, RIFT, DFT, Offline GRPO, DPO) are widely used to distill reasoning from large teachers into smaller students, and are typically compared on downstream accuracy alone. We ask whether...

<https://arxiv.org/abs/2606.23740>

0 CONF

Graph Alignment for Benchmarking Graph Neural Networks and Learning Positional Encodings

ArXiv CS.AI 50%

arXiv:2505.13087v2 Announce Type: replace-cross Abstract: We propose a novel benchmarking methodology for graph neural networks (GNNs) based on the graph alignment problem, a combinatorial optimization task that generalizes graph isomorphism by aligning two unlabeled graphs to...

<https://arxiv.org/abs/2505.13087>

1 CONF

Sol Video Inference Engine: Agent-Native Full-Stack Acceleration Framework for Efficient Video Generation

ArXiv CS.AI 50%

arXiv:2606.23743v1 Announce Type: cross Abstract: Modern video diffusion models achieve higher generation quality through scaling, but this also increases inference cost. Although many acceleration methods have been proposed, a central challenge is that the most effective...

<https://arxiv.org/abs/2606.23743>

2 CONF

Synergizing Physically Constrained MCMC and Chemical-Informed Gaussian Processes for Reaction Network Discovery

ArXiv CS.AI 50%

arXiv:2606.23757v1 Announce Type: cross Abstract: Extracting interpretable governing equations from sparse, noisy chemical time-series data remains difficult because discrete reaction topology and continuous kinetic parameters are tightly coupled. We present PC-MCMC-CIGP, a...

<https://arxiv.org/abs/2606.23757>

3 CONF

VeriPilot: An LLM-Powered Verilog Debugging Framework

ArXiv CS.AI 50%

arXiv:2606.23759v1 Announce Type: cross Abstract: Verilog debugging remains one of the most time-consuming stages in digital circuit design. Recent advances in Large Language Models (LLMs) have enabled automated debugging; however, most existing approaches rely solely on test...

<https://arxiv.org/abs/2606.23759>

4 CONF

Neuromorphic Speech Enhancement with Dual-Branch Spiking Neural Networks

ArXiv CS.AI 50%

arXiv:2606.23761v1 Announce Type: cross Abstract: Spiking neural network (SNN)-based neuromorphic speech enhancement has emerged as a promising paradigm due to its energy efficiency, yet it still underperforms classical artificial neural network (ANN)-based approaches owing to...

<https://arxiv.org/abs/2606.23761>

5 CONF

Emergent Relational Order in LLM Agent Societies: From Collective Affect to Authority Stratification

ArXiv CS.AI 50%

arXiv:2606.23764v1 Announce Type: cross Abstract: Fei Xiaotong's Differential Order Pattern characterizes rural society as egocentric and relationally graded, with cooperation attenuating over social distance. Although often treated as culturally specific, its mechanistic basis...

<https://arxiv.org/abs/2606.23764>

6 CONF

Cryptographic certificates of validity for trustworthy AI

ArXiv CS.AI 50%

arXiv:2606.23768v1 Announce Type: cross Abstract: We propose cryptographic certificates of validity for agentic AI systems. The core idea is to formally specify a correctness or policy condition as a logical predicate, compile this predicate to a witness-checking problem over...

<https://arxiv.org/abs/2606.23768>

7 CONF

Ten Digits on a Train: AI-Assisted Verification of Two Eigenvalue Problems

ArXiv CS.AI 50%

arXiv:2606.23821v1 Announce Type: cross Abstract: Accurate numerical eigenvalues are often difficult to certify, especially in singular or non-normal settings. This article reports a human-AI collaboration on two such computations. For a singular self-adjoint Schrödinger...

<https://arxiv.org/abs/2606.23821>

8 CONF

Decentralized Coordination of Autonomous Traffic Through Advanced Air Mobility Corridors

ArXiv CS.AI 50%

arXiv:2606.23832v1 Announce Type: cross Abstract: The use of dedicated corridors for Advanced Air Mobility (AAM) traffic is one of the most commonly proposed pathways to integrating them into existing airspace operations. Most prior research has focused on the design of networks...

<https://arxiv.org/abs/2606.23832>

- 9 **CONF**
- Are Safety Guarantees in Neural Networks Safe? How to Compute Trustworthy Robustness Certifications**
- ArXiv CS.AI 50%
- arXiv:2606.23858v1 Announce Type: cross Abstract: A primary challenge in AI safety is the existence of adversarial examples -- slightly distorted inputs that cause a neural network (NN) to misclassify. To mitigate this problem, recent research focuses on the computation of...
- <https://arxiv.org/abs/2606.23858>
-
- 0 **CONF**
- MGI: Member vs Generated Inference**
- ArXiv CS.AI 50%
- arXiv:2606.23872v1 Announce Type: cross Abstract: As generative models increasingly produce samples that are indistinguishable from human-created content, it becomes difficult to determine whether a given data point was part of a model's natural training set or was generated by...
- <https://arxiv.org/abs/2606.23872>
-
- 1 **CONF**
- Promise and challenges of heart chamber segmentation from non-contrast CT scans using contrastive unpaired image translation: a feasibility study**
- ArXiv CS.AI 50%
- arXiv:2606.23879v1 Announce Type: cross Abstract: Purpose: To evaluate the feasibility and challenges of heart chamber segmentation from non-contrast CT scans using contrastive unpaired image translation and deep learning-based segmentation. Approach: We developed ChameleonNet,...
- <https://arxiv.org/abs/2606.23879>
-
- 2 **CONF**
- Catastrophic Compositional Generation: Why Vanilla Diffusion Models Fail to Extrapolate**
- ArXiv CS.AI 50%
- arXiv:2606.23920v1 Announce Type: cross Abstract: The task of compositional generation involves using a conditional generative model, trained only on a subset of the possible conditions, to produce samples from compositionally-defined target distributions such as a geometric...
- <https://arxiv.org/abs/2606.23920>
-
- 3 **CONF**
- Towards Spec Learning: Inference-Time Alignment from Preference Pairs**
- ArXiv CS.AI 50%
- arXiv:2606.24004v1 Announce Type: cross Abstract: Steering a large language model (LLM) toward a desired behavior typically relies on an iterative process of hand-crafting a prompt based on a careful inspection of the model's responses. This is an involved, brittle, and...
- <https://arxiv.org/abs/2606.24004>
-
- 4 **CONF**
- PixJail: Self-Evolving Paper-to-Pipeline Reproduction for Text-to-Image Jailbreak Evaluation**
- ArXiv CS.AI 50%
- arXiv:2606.24081v1 Announce Type: cross Abstract: As Text-to-Image (T2I) jailbreak techniques evolve rapidly, existing benchmarks and reproduction workflows often struggle to keep pace. More importantly, T2I jailbreak evaluation is not a single prompt-level test, but a...
- <https://arxiv.org/abs/2606.24081>
-

5 CONF

Ensemble Learning for Large Language Models in Text and Code Generation: A Survey

ArXiv CS.AI 50%

arXiv:2503.13505v3 Announce Type: replace-cross Abstract: Generative Pretrained Transformers (GPTs) are foundational Large Language Models (LLMs) for text generation. However, individual LLMs often produce inconsistent outputs and exhibit biases, limiting their representation of...

<https://arxiv.org/abs/2503.13505>

6 CONF

DynaWM: Dynamics-Aware Distillation with World Model and Momentum Targets for Smooth Locomotion over Continuous Stairs

ArXiv CS.AI 50%

arXiv:2606.24089v1 Announce Type: cross Abstract: Recent advances in control have enabled bipedal-wheeled robots to traverse slopes and single-step obstacles, yet long staircase traversal remains challenging as current teacher-student frameworks suffer from weakened...

<https://arxiv.org/abs/2606.24089>

7 CONF

The impact of generative artificial intelligence on academic development of Chinese students in humanities and social sciences

ArXiv CS.AI 50%

arXiv:2606.24104v1 Announce Type: cross Abstract: Generative artificial intelligence(GenAI) is reshaping learning in higher education, with particularly pronounced implications for the humanities and social sciences(HSS), where learning outcomes are commonly expressed through...

<https://arxiv.org/abs/2606.24104>

8 CONF

A Benchmark for Hallucination Detection in VLMs for Gastrointestinal Endoscopy

ArXiv CS.AI 50%

arXiv:2606.24115v1 Announce Type: cross Abstract: Vision-language models (VLMs) are prone to hallucination, which remains a major barrier to their safe deployment in clinical practice. To date, most hallucination detection methods have been evaluated on radiology benchmarks such...

<https://arxiv.org/abs/2606.24115>

9 CONF

DTT-BSR+: A Generative-Regression Cascade for Music Source Restoration

ArXiv CS.AI 50%

arXiv:2606.24127v1 Announce Type: cross Abstract: Music source restoration (MSR) requires jointly addressing source unmixing and the inversion of non-linear production effects. Current methods struggle to achieve accurate target signal reconstruction while maintaining semantic...

<https://arxiv.org/abs/2606.24127>

0 CONF

Breaking Shortcut Learning for Cross-Trial EEG-Guided Target Speech Extraction via Two-Stage Training

ArXiv CS.AI 50%

arXiv:2606.24164v1 Announce Type: cross Abstract: Recent end-to-end models for EEG-guided target speech extraction report impressive results, underscoring potential for neuro-steered hearing technologies. However, our analysis reveals that high within-trial performance can be...

<https://arxiv.org/abs/2606.24164>

1 **CONF****Lightweight Transformer Models for On-Device Fault Detection: A Benchmark Study on Resource-Constrained Deployment****ArXiv CS.AI** **50%**

arXiv:2606.24173v1 Announce Type: cross Abstract: On-device fault detection enables real-time diagnostics without cloud dependency, but deploying machine learning models on resource-constrained hardware demands careful tradeoffs between accuracy, latency, and model size. We...

<https://arxiv.org/abs/2606.24173>

2 **CONF****Deep Learning Approaches for 3D Medical Scene Completion: From Geometric Modeling to Generative Paradigms****ArXiv CS.AI** **50%**

arXiv:2606.24180v1 Announce Type: cross Abstract: Three-dimensional scene completion has evolved as a major problem in computer vision and robotics, and its applications are diverse, including autonomous navigation and augmented reality. In this study, a systematic review has...

<https://arxiv.org/abs/2606.24180>

3 **CONF****Co-occurring associated retained concepts in Diffusion Unlearning****ArXiv CS.AI** **50%**

arXiv:2606.24192v1 Announce Type: cross Abstract: Unlearning has emerged as a key technique to mitigate harmful content generation in diffusion models. However, existing methods often remove not only the target concept, but also benign co-occurring concepts. As illustrated in...

<https://arxiv.org/abs/2606.24192>

4 **CONF****AutoSpec: Safety Rule Evolution for LLM Agents via Inductive Logic Programming****ArXiv CS.AI** **50%**

arXiv:2606.24245v1 Announce Type: cross Abstract: Large language model (LLM) agents increasingly automate complex tasks by integrating language models with external tools and environments. However, their autonomy poses significant safety risks: agents may execute destructive...

<https://arxiv.org/abs/2606.24245>

5 **CONF****SURGELLM: Rethinking Multi-Task Evaluation through Task-Aware Feature Gating with Class-Balanced Normalization****ArXiv CS.AI** **50%**

arXiv:2606.24259v1 Announce Type: cross Abstract: Fine-tuned encoders deployed across heterogeneous NLP tasks face three compounding problems: mismatched inductive biases, class-imbalance corruption of feature statistics, and no mechanism to condition attention on external...

<https://arxiv.org/abs/2606.24259>

6 **CONF****Neural Network-Based Parametric Model Reduction for Predicting Turbulent Flow for Different Vehicle Geometries****ArXiv CS.AI** **50%**

arXiv:2606.24265v1 Announce Type: cross Abstract: Numerical simulations in industrial applications often require performing numerous high-precision computations parameterized by specific experimental conditions. For instance, in vehicle body design, aerodynamic simulations are...

<https://arxiv.org/abs/2606.24265>

- 7 **CONF**
Average Rankings Mask Per-Subject Optimality: A Friedman-Nemenyi Benchmark of EEG Motor-Imagery BCI Decoders
 ArXiv CS.AI 50%
 arXiv:2606.24394v1 Announce Type: cross Abstract: Electroencephalography (EEG) is the dominant non-invasive modality for brain-computer interfaces (BCIs), yet reliable decoding of motor imagery is hampered by inter- and intra-individual variability. A recurring claim is that one...
<https://arxiv.org/abs/2606.24394>
- 8 **CONF**
Detecting AI Coding Agents in Open Source: A Validated Multi-Method Census of 180 Million Repositories
 ArXiv CS.AI 50%
 arXiv:2606.24429v1 Announce Type: cross Abstract: Generative AI coding agents are entering the open-source supply chain, yet their diverse and often invisible traces leave their prevalence poorly understood. We introduce a multi-layered detection framework that integrates...
<https://arxiv.org/abs/2606.24429>
- 9 **CONF**
MedPCFM: Improving Medical Point Cloud Completion by Integrating Point Transformers and Flow Matching
 ArXiv CS.AI 50%
 arXiv:2606.24433v1 Announce Type: cross Abstract: Medical point cloud completion is important for anatomical reconstruction and downstream clinical workflows, yet generative modeling in this setting remains insufficiently studied. We investigate completion through...
<https://arxiv.org/abs/2606.24433>
- 0 **CONF**
The African Language Tax: Quantifying the Cost, Latency, and Context Penalty of Tokenizing African Languages in Frontier LLMs
 ArXiv CS.AI 50%
 arXiv:2606.24460v1 Announce Type: cross Abstract: Commercial large language models bill, scale latency, and budget context per token. Yet tokenizers assign more subword tokens to the same meaning in some languages than in others, so speakers of languages with high...
<https://arxiv.org/abs/2606.24460>
- 1 **CONF**
Adaptive Machine Learning Framework for UAV Trajectory Optimization in O-RAN
 ArXiv CS.AI 50%
 arXiv:2606.24483v1 Announce Type: cross Abstract: The deployment of unmanned aerial vehicles (UAV) as open radio units (O-RUs) in 6G cellular systems presents a promising opportunity to achieve scalable and adaptive network coverage. However, optimizing UAV trajectories in...
<https://arxiv.org/abs/2606.24483>
- 2 **CONF**
A global log for medical AI
 ArXiv CS.AI 50%
 arXiv:2510.04033v2 Announce Type: replace Abstract: Modern computer systems rely on syslog, a universal protocol that records critical events across heterogeneous infrastructure. Medicine's rapidly growing AI stack has no equivalent. As medicine deploys AI tools at scale, there...
<https://arxiv.org/abs/2510.04033>

- 3
CONF
CrossPool: Efficient Multi-LLM Serving for Cold MoE Models through KV-Cache and Weight Disaggregation
ArXiv CS.AI 50%
arXiv:2606.24506v1 Announce Type: cross Abstract: Emerging LLM services increasingly host many sparse MoE models, yet most models receive sparse requests and remain cold. This creates a GPU memory problem: model weights are stable and model-determined, while KV-cache is...
<https://arxiv.org/abs/2606.24506>
- 4
CONF
A Fair Evaluation of Graph Foundation Models for Node Property Prediction
ArXiv CS.AI 50%
arXiv:2606.24509v1 Announce Type: cross Abstract: Due to the wide use of graph-structured data in different fields of industry and science, the development of Graph Foundation Models (GFM) has recently attracted a lot of attention. While many different types of models are...
<https://arxiv.org/abs/2606.24509>
- 5
CONF
Toward Self-Evolution-Ready Workflow Harnesses: A Reversible Migration Path and Convertibility Taxonomy for Expert LLM Pipelines
ArXiv CS.AI 50%
arXiv:2606.24598v1 Announce Type: cross Abstract: While expert-validated "LLM + script" workflows deliver significant value, they remain static: they encode hard-won domain knowledge yet fail to adapt execution based on feedback. Existing agent research predominantly targets...
<https://arxiv.org/abs/2606.24598>
- 6
CONF
Privacy-Preserving RAG via Multi-Agent Semantic Rewriting: Achieving Confidentiality Without Compromising Contextual Fidelity
ArXiv CS.AI 50%
arXiv:2606.24623v1 Announce Type: cross Abstract: Retrieval-Augmented Generation enhances large language models by incorporating external knowledge, but deploying it in sensitive scenarios risks privacy leakage via malicious prompts. To address this, we propose a multi-agent...
<https://arxiv.org/abs/2606.24623>
- 7
CONF
AI-PAVE-Br: Leveraging Large Language Models for Enhanced Product Attribute Value Extraction through a Golden Set Approach
ArXiv CS.AI 50%
arXiv:2606.24655v1 Announce Type: cross Abstract: The explosive growth and complexity of product data within the dynamic Brazilian e-commerce landscape demand robust and specialized methods for structured information extraction. Traditional approaches to Product Attribute Value...
<https://arxiv.org/abs/2606.24655>
- 8
CONF
Context-Aware Prediction of Student Quiz Performance with Multimodal Textbook Features
ArXiv CS.AI 50%
arXiv:2606.24770v1 Announce Type: cross Abstract: Educational platforms often predict student performance from prior interactions, but the assessment content itself also varies in linguistic and visual complexity. This paper studies whether lightweight content features extracted...
<https://arxiv.org/abs/2606.24770>

9 CONF

EG-VQA: Benchmarking Verifiable Video Question Answering with Grounded Temporal Evidence

ArXiv CS.AI 50%

arXiv:2606.24797v1 Announce Type: cross Abstract: Recent advances in Video Large Language Models (Video-LLMs) have yielded promising performance on video question answering (VideoQA). Nevertheless, existing benchmarks are predominantly evaluated through answer correctness, while...

<https://arxiv.org/abs/2606.24797>

00 CONF

IV-CoT: Implicit Visual Chain-of-Thought for Structure-Aware Text-to-Image Generation

ArXiv CS.AI 50%

arXiv:2606.24849v1 Announce Type: cross Abstract: Unified multi-modal large language models (MLLMs) have achieved strong text-to-image generation quality, but still struggle with structure-aware prompt following, where object counts, spatial relations, attribute bindings, and...

<https://arxiv.org/abs/2606.24849>

01 CONF

It's Complicated: On the Design and Evaluation of AI-Powered AAC Interfaces

ArXiv CS.AI 50%

arXiv:2606.24854v1 Announce Type: cross Abstract: Artificial intelligence (AI) can enhance what people who use augmentative and alternative communication (AAC) are able to do with their systems. However, evaluating AI-powered AAC interfaces can be difficult. People are...

<https://arxiv.org/abs/2606.24854>

02 CONF

Random Rule Forest (RRF): Interpretable and Manageable Ensembles of LLM-Generated Questions for Predicting Success from Unstructured Data

ArXiv CS.AI 50%

arXiv:2505.24622v3 Announce Type: replace Abstract: Many high-stakes screening tasks require predicting rare outcomes from unstructured text, where errors are costly and decisions must be auditable. We introduce Random Rule Forest (RRF), an interpretable ensemble that uses a...

<https://arxiv.org/abs/2505.24622>

03 CONF

TIP-Search: Time-Predictable Inference Scheduling for Market Prediction under Uncertain Load

ArXiv CS.AI 50%

arXiv:2506.08026v4 Announce Type: replace Abstract: Real-time market prediction services need correct predictions before a decision deadline; a correct prediction delivered late is not usable. TIP-Search studies time-predictable inference scheduling over fixed market predictors...

<https://arxiv.org/abs/2506.08026>

04 CONF

Representation Interventions Enable Lifelong Knowledge Memory Control in LLMs

ArXiv CS.AI 50%

arXiv:2511.20892v4 Announce Type: replace Abstract: Large language models (LLMs) often produce incorrect or outdated content after being employed. Efficient and accurate knowledge updates without costly retraining are a major challenge. This problem is particularly challenging...

<https://arxiv.org/abs/2511.20892>

05 CONF

Grounded Chess Reasoning in Language Models via Master Distillation

ArXiv CS.AI 50%

arXiv:2603.20510v2 Announce Type: replace Abstract: Language models often lack grounded reasoning capabilities in specialized domains where training data is scarce but bespoke systems excel. We introduce a general framework for distilling expert system reasoning into natural...

<https://arxiv.org/abs/2603.20510>

06 CONF

Subjective-Graph LLM Agents for Simulating Uncertainty in Classroom Social Perception

ArXiv CS.AI 50%

arXiv:2603.20750v2 Announce Type: replace Abstract: Social actors do not observe a common social world: each individual forms judgments from a partial and potentially distorted view of the surrounding network. We study whether graph-local evidence and credibility-weighted...

<https://arxiv.org/abs/2603.20750>

07 CONF

Efficient Test-time Inference for Generative Planning Models with OCL Search

ArXiv CS.AI 50%

arXiv:2606.00618v2 Announce Type: replace Abstract: Generative models have emerged as a powerful paradigm for AI planning, yet their performance remains constrained by the training data distribution. One approach is to improve generated solutions during inference by scaling...

<https://arxiv.org/abs/2606.00618>

08 CONF

When Preferences Fail to Become Incentives: A Utility-Behavior Gap in Large Language Models

ArXiv CS.AI 50%

arXiv:2606.22974v2 Announce Type: replace Abstract: Recent work on preference elicitation in large language models (LLMs) has demonstrated that, when given a series of choices between two outcomes, LLMs reveal a coherent, model-specific utility structure. Notably, this structure...

<https://arxiv.org/abs/2606.22974>

09 CONF

IPO Finance Agent: Evaluation of LLM Financial Analysts beyond Finance Agent v2, with Automated Rubric Generation -- the Case of the SpaceX (SPCX) IPO

ArXiv CS.AI 50%

arXiv:2606.23032v2 Announce Type: replace Abstract: Finance Agent v2 (by Vals AI) has emerged as the reference benchmark for evaluating both Anthropic Claude and OpenAI ChatGPT frontier language models on financial tasks. However, it narrowly deals with periodic reporting from...

<https://arxiv.org/abs/2606.23032>

10 CONF

When AI Meets Finance (StockAgent): Large Language Model-based Stock Trading in Simulated Real-world Environments

ArXiv CS.AI 50%

arXiv:2407.18957v5 Announce Type: replace-cross Abstract: Can AI Agents simulate real-world trading environments to investigate the impact of external factors on stock trading activities (e.g., macroeconomics, policy changes, company fundamentals, and global events)? These...

<https://arxiv.org/abs/2407.18957>

11 CONF

CORE-Bench: Fostering the Credibility of Published Research Through a Computational Reproducibility Agent Benchmark

ArXiv CS.AI 50%

arXiv:2409.11363v2 Announce Type: replace-cross Abstract: AI agents have the potential to aid users on a variety of consequential tasks, including conducting scientific research. To spur the development of useful agents, we need benchmarks that are challenging, but more...

<https://arxiv.org/abs/2409.11363>

12 CONF

Benchmarking LLMs' Mathematical Reasoning with Unseen Random Variables Questions

ArXiv CS.AI 50%

arXiv:2501.11790v5 Announce Type: replace-cross Abstract: Recent studies have raised significant concerns regarding the reliability of current mathematics benchmarks, highlighting issues such as simplistic design and potential data contamination. Consequently, developing a...

<https://arxiv.org/abs/2501.11790>

13 CONF

Societal Alignment Frameworks Can Improve LLM Alignment

ArXiv CS.AI 50%

arXiv:2503.00069v2 Announce Type: replace-cross Abstract: Recent progress in large language models (LLMs) has focused on producing responses that meet human expectations and align with shared values - a process coined alignment. However, aligning LLMs remains challenging due to...

<https://arxiv.org/abs/2503.00069>

14 CONF

Reward-Centered ReST-MCTS: A Robust Decision-Making Framework for Robotic Manipulation in High Uncertainty Environments

ArXiv CS.AI 50%

arXiv:2503.05226v2 Announce Type: replace-cross Abstract: Monte Carlo tree search is attractive for robotic manipulation because it can improve action selection through simulation without requiring a fully differentiable policy. In uncertain domains, however, sparse terminal...

<https://arxiv.org/abs/2503.05226>

15 CONF

Tuning without Peeking: Provable Generalization Bounds and Robust LLM Post-Training

ArXiv CS.AI 50%

arXiv:2507.01752v4 Announce Type: replace-cross Abstract: Gradient-based optimization is the workhorse of deep learning, offering efficient and scalable training via backpropagation. However, exposing gradients during training can leak sensitive information about the underlying...

<https://arxiv.org/abs/2507.01752>

16 CONF

FALCON: Transforming Cyber Threat Intelligence into Deployable IDS Rules with Self-Reflection

ArXiv CS.AI 50%

arXiv:2508.18684v2 Announce Type: replace-cross Abstract: Signature-based Intrusion Detection Systems (IDS) detect malicious activity by matching network or host events against predefined rules. Security analysts manually develop these rules from Cyber Threat Intelligence (CTI)...

<https://arxiv.org/abs/2508.18684>

17 CONF

Breaking the Mirror: Activation-Based Mitigation of Self-Preference in LLM Evaluators

ArXiv CS.AI 50%

arXiv:2509.03647v2 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly serve as automated evaluators, yet they suffer from "self-preference bias": a tendency to favor their own outputs over those of other models. This bias undermines fairness and...

<https://arxiv.org/abs/2509.03647>

18 CONF

VoltanaLLM: Energy-Efficient and SLO-Aware Disaggregated LLM Serving via Adaptive Frequency Control and State-Space Routing

ArXiv CS.AI 50%

arXiv:2509.04827v3 Announce Type: replace-cross Abstract: The energy cost of Large Language Model (LLM) inference is rapidly becoming a barrier to sustainable and scalable deployment. Although modern serving architectures expose distinct prefill and decode behaviors, existing...

<https://arxiv.org/abs/2509.04827>

19 CONF

MOCHA: Multi-modal Objects-aware Cross-architecture Alignment

ArXiv CS.AI 50%

arXiv:2509.14001v5 Announce Type: replace-cross Abstract: Personalized object detection aims to adapt a general-purpose detector to recognize user-specific instances from only a few examples. Lightweight models often struggle in this setting due to their weak semantic priors,...

<https://arxiv.org/abs/2509.14001>

20 CONF

EMFusion: Uncertainty-Aware Conditional Diffusion Model for Multivariate Narrow-band Exposure Forecasting

ArXiv CS.AI 50%

arXiv:2512.15067v4 Announce Type: replace-cross Abstract: The rapid growth in wireless infrastructure has increased the need to accurately estimate and forecast electromagnetic field (EMF) levels to ensure ongoing compliance, assess potential health impacts, and support...

<https://arxiv.org/abs/2512.15067>

21 CONF

Attention in Motion: Secure Platooning via Transformer-based Misbehavior Detection

ArXiv CS.AI 50%

arXiv:2512.15503v3 Announce Type: replace-cross Abstract: Vehicular platooning promises transformative improvements in transportation efficiency and safety through the coordination of multi-vehicle formations enabled by Vehicle-to-Everything (V2X) communication. However, the...

<https://arxiv.org/abs/2512.15503>

22 CONF

Disentangling Aleatoric and Epistemic Uncertainty in Physics-Informed Neural Networks. Application to Insulation Material Degradation Prognostics

ArXiv CS.AI 50%

arXiv:2601.03673v2 Announce Type: replace-cross Abstract: Physics-Informed Neural Networks (PINNs) provide a framework for integrating physical laws with data. However, their application to Prognostics and Health Management (PHM) remains constrained by the limited uncertainty...

<https://arxiv.org/abs/2601.03673>

23 CONF

Are LLM Evaluators Really Narcissists? Sanity Checking Self-Preference Evaluations

ArXiv CS.AI 50%

arXiv:2601.22548v4 Announce Type: replace-cross Abstract: Recent research has shown that large language models (LLMs) favor their own outputs when acting as judges, undermining the integrity of automated post-training and evaluation workflows. However, it is difficult to...

<https://arxiv.org/abs/2601.22548>

24 CONF

No Certificate, No Categorical Speech Act: A Brouwerian Assertibility Constraint for Public Reason

ArXiv CS.AI 50%

arXiv:2603.03971v3 Announce Type: replace-cross Abstract: Generative AI can convert uncertainty into authoritative-seeming verdicts, intensifying the hypersuasive force of automated speech and displacing the justificatory work on which democratic epistemic agency depends. As a...

<https://arxiv.org/abs/2603.03971>

25 CONF

An Approach to Simultaneous Acquisition of Real-Time MRI Video, EEG, and Surface EMG for Articulatory, Brain, and Muscle Activity During Speech Production

ArXiv CS.AI 50%

arXiv:2603.04840v2 Announce Type: replace-cross Abstract: Speech production is a complex process spanning neural planning, motor control, muscle activation, and articulatory kinematics. While the acoustic speech signal is the most accessible product of the speech production act,...

<https://arxiv.org/abs/2603.04840>

26 CONF

HiIDA: Hierarchical Distillation with Diffusion for Advancing Self-Supervised LiDAR Pre-training

ArXiv CS.AI 50%

arXiv:2606.20189v3 Announce Type: replace-cross Abstract: Leveraging Vision Foundation Models (VFMs) for camera-to-LiDAR knowledge distillation offers a promising solution to the scarcity of annotated data needed to represent the immense geometric and kinematic diversity of...

<https://arxiv.org/abs/2606.20189>

27 CONF

Fix Initial Programs and Iteratively Refine Repair Instructions Toward Non-Elimination Multi-Turn Program Correction

ArXiv CS.AI 50%

arXiv:2604.23989v2 Announce Type: replace-cross Abstract: Recent work on large language models (LLMs) has emphasized the importance of scaling inference compute. From this perspective, the state-of-the-art method Scattered Forest Search (SFS) has been proposed, employing Monte...

<https://arxiv.org/abs/2604.23989>

28 CONF

VISTA: An End-to-End Benchmark for Visual Spec-to-Web-App Coding Agents

ArXiv CS.AI 50%

arXiv:2605.26144v3 Announce Type: replace-cross Abstract: We present VISTA (Visual Spec-To-App Benchmark), a benchmark for evaluating the end-to-end web-app generation capabilities of LLM-based agents. Unlike prior code generation benchmarks that focus on algorithmic tasks,...

<https://arxiv.org/abs/2605.26144>

29 CONF

Cross-Dataset, Age, and Gender Generalization: A Comprehensive Analysis of Fine-Tuning Strategies for Low-Resource Children's ASR

ArXiv CS.AI 50%

arXiv:2606.19791v2 Announce Type: replace-cross Abstract: The challenge associated with recognizing dysarthric speech primarily arises from pronounced acoustic variability attributed to impaired articulatory precision. Past research has demonstrated improved recognition through...

<https://arxiv.org/abs/2606.19791>

30 CONF

Does Mixture-of-Experts Actually Help Inference on Consumer and Edge Hardware? An Empirical Study

ArXiv CS.AI 50%

arXiv:2606.21428v2 Announce Type: replace-cross Abstract: Mixture-of-Experts (MoE) language models are often described as ideal for resource-constrained inference. Each token activates only a small subset of experts, so the per-token compute cost, in floating-point operations...

<https://arxiv.org/abs/2606.21428>

31 CONF

Diffusion Integrated Gradients: Controllable Path Generation for Flexible Feature Attribution

ArXiv CS.AI 50%

arXiv:2606.22314v2 Announce Type: replace-cross Abstract: Path-based attribution methods such as Integrated Gradients (IG) are widely adopted for their strong axiomatic properties and effectiveness in attributing model predictions to input features by integrating gradients along...

<https://arxiv.org/abs/2606.22314>

32 CONF

AI Fiction in the Wild

ArXiv CS.AI 50%

arXiv:2606.22748v2 Announce Type: replace-cross Abstract: Some professional authors are beginning to use AI tools to help produce their fiction writing. Are readers using AI to generate fiction, too? Drawing on over 500,000 anonymized, English-language ChatGPT-user conversations...

<https://arxiv.org/abs/2606.22748>