



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	113	Media
HackerNews	8	Community
TechCrunch AI	5	Media
MIT Tech Review	3	Media
The Verge AI	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Show HN: Erlangchain - A tiny Erlang client for LLMs

HackerNews

+1 source

90%

HN 3

arXiv:2602.22428v2 Announce Type: replace-cross Abstract: Test-time guidance is a widely used mechanism for steering pretrained diffusion models toward outcomes specified by a reward function. Existing approaches, however, focus on maximizing reward rather than sampling from the...

<https://github.com/abhavk/erlangchain>

VERI

HARC: Coupling Harmfulness and Refusal Directions for Robust Safety Alignment

ArXiv CS.AI

+1 source

80%

arXiv:2607.00572v1 Announce Type: new Abstract: Understanding how aligned LLMs internally represent safety is critical for diagnosing alignment vulnerabilities, as it explains why jailbreaks succeed and informs the design of robust alignment strategies. Prior work shows that...

<https://arxiv.org/abs/2607.00572>

CONFIRMED SIGNALS — 126 stories

CONF

Show HN: Z-jail - A 130 KB Linux sandbox-C99 with 7 defense layers and zero deps

HackerNews

50%

HN 21

No summary — see link for full article.

<https://github.com/Division-36/Z-jail/>

CONF

Show HN: Classify mechanical faults using Contrastive Language-Audio Pretraining

HackerNews

50%

HN 8

No summary — see link for full article.

<https://github.com/adam-s/car-diagnosis>

CONF

Show HN: Forgejo Appliance, a private Git server based on OrbStack and Tailscale

HackerNews

50%

HN 6

No summary — see link for full article.

<https://github.com/highpost/forgejo-appliance>

CONF

Show HN: AnalystAIPack - 118 runnable agent skills for malware analysis and RE

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/meltedinhex/analyst-ai-pack>

CONF

Show HN: CLI that helps AI agents avoid vulnerable dependencies

HackerNews

50%

HN 3

deptrust is a CLI that checks package versions for known vulnerabilities across npm, PyPI, crates.io, Go modules, RubyGems, NuGet, Maven, Packagist, pub.dev, CocoaPods, Hex.pm, Hackage, GitHub Actions, and more. It runs locally as a CLI and as an MCP server. It calls public...

<https://github.com/clidey/deptrust>

CONF

Show HN: I Made TS Compiler Graph MCP: 10x Fewer Tokens in Claude Code and Codex

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/samchon/ttsc/tree/master/packages/graph>

CONF

Show HN: QueryDrift fail CI when a PR turns 1 query into N

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/AALXX/QueryDrift>

0 CONF

Indian tech tycoon bets \$30M of his own money to build AI alternative to Microsoft Office

TechCrunch AI 50%

Neo is Bhavin Turakhia's fifth venture and his latest involving enterprise software. This time he's taking on Microsoft Office, Google Apps with AI.

<https://techcrunch.com/2026/07/01/indian-tech-tycoon-bets-30m-to-build-an-ai-alternativ...>

1 CONF

SpaceX has an AI device prototype, and it sure sounds phone-ish

TechCrunch AI 50%

SpaceX reportedly showed investors a "handset-like" AI device before going public. It could be another signal SpaceX wants to expand into wireless.

<https://techcrunch.com/2026/07/01/spacex-has-an-ai-device-prototype-and-it-sure-sounds-...>

2 CONF

Venice AI becomes a unicorn with \$65M Series A as its privacy-first AI platform takes off

TechCrunch AI 50%

Venice AI is already profitable, with annualized run-rate revenues of over \$70 million, CEO Erik Voorhees said.

<https://techcrunch.com/2026/07/01/venice-ai-becomes-a-unicorn-with-65m-series-a-as-its-...>

3 CONF

Gemini Spark, Google's agentic assistant, is now available on Mac

TechCrunch AI 50%

Google's 24/7 agentic assistant, Gemini Spark, comes to Mac alongside other improvements, like real-time tracking and support for more apps.

<https://techcrunch.com/2026/07/01/gemini-spark-googles-agentic-assistant-is-now-availab...>

4 CONF

Meta, like SpaceX, looks to turn excess AI compute into cash

TechCrunch AI 50%

Meta is developing plans for a cloud infrastructure business, selling access to AI compute power and models. The move would pit it against the big cloud providers like Amazon Web Services, Google Cloud, and Microsoft Azure.

[https://techcrunch.com/2026/07/01/meta-like-spacex-looks-to-turn-excess-ai-compute-into...](https://techcrunch.com/2026/07/01/meta-like-spacex-looks-to-turn-excess-ai-compute-into-...)

5 CONF

LLMs are stuck in a groupthink groove. This startup is trying to get them out.

MIT Tech Review 50%

Let's start with a game. Open up your chatbot of choice—Claude, ChatGPT, Gemini—and type "Give me a random number between 1 and 10." You're going to get 7. Almost always. Now type "Another" and you'll get 3 or 4. Type "Another" again and you'll get 8 or 9. That won't work every...

<https://www.technologyreview.com/2026/07/01/1140003/llms-are-stuck-in-a-groupthink-rut-...>

6 CONF

The Download: Anthropic launches Claude Science, and California's carbon manure math

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Claude Science is Anthropic's newest flagship product. At an event for pharmaceutical executives, biotech founders, and researchers yesterday,...

<https://www.technologyreview.com/2026/07/01/1139996/the-download-anthropic-claude-scienc...>

7 CONF

Google built a great smart speaker, but Gemini isn't ready for it

The Verge AI 50%

Smart speakers have spent the past few years searching for a compelling second act. Beyond music, timers, and controlling your lights, they've struggled to justify taking up space on the kitchen counter. AI promised to change that. Amazon debuted its new hardware powered by a...

<https://www.theverge.com/tech/959503/google-home-speaker-review-gemini-for-home>

8 CONF

Making Failure Safe: A Constrained, Verifiable Agent Framework for Open-Web Data Collection

ArXiv CS.AI 50%

arXiv:2607.00035v1 Announce Type: new Abstract: LLMs and agents can generate web scrapers from natural-language requirements, but direct generation remains unreliable because of dependency errors, broken selectors, schema mismatches, and heterogeneous page structures. We propose...

<https://arxiv.org/abs/2607.00035>

9 CONF

Solution space path planning for supporting en-route air traffic control

ArXiv CS.AI 50%

arXiv:2607.00064v1 Announce Type: new Abstract: As technology advances, many path-planning algorithms have been proposed for Air Traffic Management, yet their operational adoption in tactical control remains limited, revealing a misalignment between algorithmic design priorities...

<https://arxiv.org/abs/2607.00064>

0 CONF

Constructing Epistemic AI Literacy: Detecting Epistemic Aims and Processes in Student-AI Co-Programming

ArXiv CS.AI 50%

arXiv:2607.00211v1 Announce Type: new Abstract: Epistemic thinking plays a central role in students' learning processes when applying generative artificial intelligence (GenAI), particularly in programming contexts where learners must construct queries, evaluate and validate...

<https://arxiv.org/abs/2607.00211>

1 CONF

From Signals to Structure: How Memory Architecture Drives Language Emergence in LLM Agents

ArXiv CS.AI 50%

arXiv:2607.00233v1 Announce Type: new Abstract: How do two agents invent a shared language from scratch? In a Lewis signaling game, a sender and receiver must coordinate on a code using only their interaction history. We study five memory architectures across varying channel...

<https://arxiv.org/abs/2607.00233>

2 CONF

Mnemopsyne: Agentic Transaction Processing for Validating and Repairing AI-generated Workflows

ArXiv CS.AI 50%

arXiv:2607.00269v1 Announce Type: new Abstract: LLMs, solvers, and agent teams increasingly generate workflow actions, repairs, and plans, but a generated action may be syntactically valid yet stale, infeasible, conflicting, or destructive of the evidence that triggered a...

<https://arxiv.org/abs/2607.00269>

3 CONF

EgoSafetyBench: A Diagnostic Egocentric Video Benchmark for Evaluating Embodied VLMs as Runtime Safety Guards

ArXiv CS.AI 50%

arXiv:2607.00218v1 Announce Type: cross Abstract: Vision-language models (VLMs) are now proposed as runtime safety guards for embodied agents in homes and factories. A deployable guard must catch genuinely unsafe situations while avoiding unnecessary intervention on routine but...

<https://arxiv.org/abs/2607.00218>

4 CONF

Agri-SAGE: Simulation-Grounded Multi-Agent LLM for Context-Aware Agricultural Advisory Generation

ArXiv CS.AI 50%

arXiv:2607.00454v1 Announce Type: new Abstract: Agricultural advisory systems face a fundamental tension: static agronomic guidelines offer consistent, evidence-based recommendations, yet remain blind to in-season variability and dynamic uncertainties. Recent advisory systems...

<https://arxiv.org/abs/2607.00454>

5 CONF

AI Native Games: A Survey and Roadmap

ArXiv CS.AI 50%

arXiv:2607.00527v1 Announce Type: new Abstract: Generative AI now enables games to produce dialogue, quests, characters, images, and worlds at runtime. Yet generation alone does not make a game AI-native, nor does it guarantee playability. This paper defines AI-native games by...

<https://arxiv.org/abs/2607.00527>

6 CONF

Self-GC: Self-Governing Context for Long-Horizon LLM Agents

ArXiv CS.AI 50%

arXiv:2607.00692v1 Announce Type: new Abstract: Long-horizon LLM agents accumulate tool results, files, plans, and user constraints that are too structured to be treated as a disposable text suffix. Current systems mostly rely on in-run heuristics such as chronological pruning...

<https://arxiv.org/abs/2607.00692>

7 CONF

Two AI Metrics Diverged: Will it Make All the Difference?

ArXiv CS.AI 50%

arXiv:2607.00913v1 Announce Type: new Abstract: As exponential compute scaling continues, will the capabilities of frontier AI models outstrip what is accessible to developers on a small fixed budget? Or will capabilities converge, with "meek models inheriting the earth"?

<https://arxiv.org/abs/2607.00913>

8 CONF

Graph-Native Reinforcement Learning Enables Traceable Scientific Hypothesis Generation through Conceptual Recombination

ArXiv CS.AI 50%

arXiv:2607.00924v1 Announce Type: new Abstract: Accelerating materials discovery requires AI systems that can generate scientifically valid hypotheses through multi-step, domain-grounded reasoning. Standard large language models often produce fluent but weakly traceable...

<https://arxiv.org/abs/2607.00924>

9 CONF

Bayesian Uncertainty Propagation for Agentic RAG Pipelines: A Proof-of-Concept Study on Multi-Hop Question Answering

ArXiv CS.AI 50%

arXiv:2607.00972v1 Announce Type: new Abstract: Trustworthy deployment of Agentic Retrieval-Augmented Generation (RAG) systems requires mechanisms for estimating when multi-stage reasoning pipelines may fail. This paper presents an uncertainty-aware Agentic Retrieval-Augmented...

<https://arxiv.org/abs/2607.00972>

10 CONF

Can Agents Generalize to the Open World? Unveiling the Fragility of Static Training in Tool Use

ArXiv CS.AI 50%

arXiv:2607.01084v1 Announce Type: new Abstract: While Large Language Model (LLM) agents demonstrate proficiency in static benchmarks, their deployment in real-world scenarios is hindered by the dynamic nature of user queries, tool sets, and interaction dynamics. To address this...

<https://arxiv.org/abs/2607.01084>

1 CONF

From "Strings" to "Things" for Personal Knowledge Graphs: Evaluating LLM Triple Extraction for Recommendation Systems

ArXiv CS.AI 50%

arXiv:2607.00003v1 Announce Type: cross Abstract: Personal Knowledge Graphs (PKGs) offer a privacy-preserving framework for modeling user preferences, yet constructing them from unstructured, decentralized conversational data remains a challenge. This paper bridges the gap...

<https://arxiv.org/abs/2607.00003>

2 CONF

Persona Without Substrate: Regime-Dependence and the LLM Individuation Problem

ArXiv CS.AI 50%

arXiv:2607.00006v1 Announce Type: cross Abstract: Beckmann & Butlin's (2026) ontological framework for the LLM individuation problem inherits an unargued cross-regime co-reference assumption from the persona-vectors literature: that the same direction picks out the same content...

<https://arxiv.org/abs/2607.00006>

3 CONF

Libra: Training the Environment for Agentic Information Retrieval

ArXiv CS.AI 50%

arXiv:2607.00016v1 Announce Type: cross Abstract: Information localization within massive repositories is a cornerstone of agentic LLM systems. While synthetic data-driven optimization has proven successful in training LLMs, little attention has been paid to optimizing the...

<https://arxiv.org/abs/2607.00016>

- 4 **CONF**
- SkillSelect-Serve: Budget-Controllable and QoS-Aware Skill Service Recommendation and Composition for Small LLM Agents**
- ArXiv CS.AI 50%
- arXiv:2607.00011v1 Announce Type: cross Abstract: Reusable skill libraries are becoming important infrastructure for large language model (LLM) agents, yet existing selection methods often treat skills as retrievable documents and return fixed top-k lists. This paper presents...
- <https://arxiv.org/abs/2607.00011>
-
- 5 **CONF**
- PRA-RAG: Provably Robust Aggregation in Retrieval-Augmented Generation against Retrieval Corruption**
- ArXiv CS.AI 50%
- arXiv:2607.00012v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) enhances Large Language Models (LLMs) by incorporating external knowledge, effectively mitigating their inherent knowledge limitations. However, RAG remains vulnerable to poisoning attacks...
- <https://arxiv.org/abs/2607.00012>
-
- 6 **CONF**
- GRACE-RAG: Governed Retrieval Architecture for Canonical Evidence Synthesis, Enabling Lightweight Deployment in Closed-Domain Institutional Settings**
- ArXiv CS.AI 50%
- arXiv:2607.00013v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) systems are widely used in institutional question answering settings where responses must be grounded in authoritative documentation (Gao et al., 2023). In entity-dense domains where relevant...
- <https://arxiv.org/abs/2607.00013>
-
- 7 **CONF**
- Towards an automated AI-based framework for floor plan compliance checks for residential buildings**
- ArXiv CS.AI 50%
- arXiv:2607.00015v1 Announce Type: cross Abstract: To improve residents' well-being in Australia's urban areas, governments have introduced policy reforms such as SEPP65, BADS, and SPP7.3 to enhance apartment design quality. These regulations require precise geometric and spatial...
- <https://arxiv.org/abs/2607.00015>
-
- 8 **CONF**
- LLMs in the Real World: Evaluating "AI" in Emergency Contexts**
- ArXiv CS.AI 50%
- arXiv:2607.00019v1 Announce Type: cross Abstract: This paper offers a call to action. We urge our colleagues in the research community to play a greater role in the articulation of our findings to the public. To illustrate the stakes we present a case study on the initial stages...
- <https://arxiv.org/abs/2607.00019>
-
- 9 **CONF**
- Aligning Sentence Embeddings to Human Concepts via Sparse Autoencoders**
- ArXiv CS.AI 50%
- arXiv:2607.00023v1 Announce Type: cross Abstract: Dense sentence embeddings are fundamental to modern Retrieval-Augmented Generation (RAG) systems but suffer from a lack of interpretability due to feature superposition. This opacity hinders the alignment of retrieval processes...
- <https://arxiv.org/abs/2607.00023>

0 CONF

FLYNN: Robust Neural Network for Robot Navigation using Fly Brain Topology

ArXiv CS.AI 50%

arXiv:2607.00025v1 Announce Type: cross Abstract: While deep learning models achieve state-of-the-art performance in complex tasks, they remain brittle when faced with new environments or sensory deprivation. In contrast, biological systems exhibit remarkable tolerance to these...

<https://arxiv.org/abs/2607.00025>

1 CONF

Prompting GPT-5 on Scrum Certification Questions: An Empirical Accuracy Study

ArXiv CS.AI 50%

arXiv:2607.00049v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used in Agile Software Development for documentation, coaching, and training. As practitioners adopt these tools to prepare for certifications such as Professional Scrum Master (PSM),...

<https://arxiv.org/abs/2607.00049>

2 CONF

Active Sensing for RIS-Aided Tracking and Power Control: A Hybrid Neuroevolution and Supervised Learning Approach

ArXiv CS.AI 50%

arXiv:2607.00056v1 Announce Type: cross Abstract: This paper studies energy efficient tracking of power-limited mobile users with the assistance of a Reconfigurable Intelligent Surface (RIS). Since localization pilot transmissions dominate the energy budget of power-constrained...

<https://arxiv.org/abs/2607.00056>

3 CONF

Lost in the Tail: Addressing Geographic Imbalance in Urban Visual Place Recognition

ArXiv CS.AI 50%

arXiv:2607.00090v1 Announce Type: cross Abstract: Urban-scale Visual Place Recognition (VPR) aims to identify the geographic location of a query image by matching it against a geo-tagged database. While recent methods achieve impressive performance, they overlook a serious...

<https://arxiv.org/abs/2607.00090>

4 CONF

SNAP-FM: Sparse Nonlinear Accelerated Projection for Physics-Constrained Generative Modeling

ArXiv CS.AI 50%

arXiv:2607.00095v1 Announce Type: cross Abstract: Generative models have emerged as scalable surrogates for physical simulation, yet they offer no guarantee that their outputs respect the conservation laws, boundary conditions, and nonlinear invariants that govern the underlying...

<https://arxiv.org/abs/2607.00095>

5 CONF

EVOTS: Evolutionary Transformer Search for Time Series Forecasting

ArXiv CS.AI 50%

arXiv:2607.00154v1 Announce Type: cross Abstract: Evolutionary neural architecture design for multivariate time-series forecasting remains underexplored, with most approaches relying on fixed Transformer architectures despite substantial variation across tasks and forecasting...

<https://arxiv.org/abs/2607.00154>

6 CONF

Scaling Up Thermodynamic AI Models

ArXiv CS.AI 50%

arXiv:2607.00170v1 Announce Type: cross Abstract: Thermodynamic computing devices based on the Ising model show great promise for low-power AI inference and edge computing, but scalable methods for training large models for such hardware remain limited. Prior theory shows that...

<https://arxiv.org/abs/2607.00170>

7 CONF

Post-Training Pruning for Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2607.00927v1 Announce Type: cross Abstract: Diffusion Transformers (DiTs) have demonstrated impressive performance in image generation but suffer from substantial computational overhead and resource consumption. Post-training pruning offers a promising solution; however,...

<https://arxiv.org/abs/2607.00927>

8 CONF

Leveraging Phase Information to Boost Unrolled Network Learning for Image Deblurring

ArXiv CS.AI 50%

arXiv:2607.00251v1 Announce Type: cross Abstract: While most image deblurring techniques directly restore the spatial image variable, we propose an amplitude and phase decomposition recognizing the importance of accurate phase estimation in recovering sharp image details. To...

<https://arxiv.org/abs/2607.00251>

9 CONF

SEFORA: Student Essays with Feedback Corpus and LLM Feedback Evaluation Framework

ArXiv CS.AI 50%

arXiv:2607.00274v1 Announce Type: cross Abstract: Effective writing feedback is among the strongest drivers of student learning, yet producing it at scale is labor-intensive. LLMs offer a natural path to scaling writing support, but two gaps stand in the way: few public corpora...

<https://arxiv.org/abs/2607.00274>

10 CONF

What's Hidden Matters: Identifying Planning-Critical Occluded Agents using Vision-Language Models

ArXiv CS.AI 50%

arXiv:2607.00283v1 Announce Type: cross Abstract: Autonomous vehicles must safely navigate complex environments where planning-critical agents may be hidden from view. Current approaches often treat all occlusions with uniform conservatism, yielding needlessly defensive driving,...

<https://arxiv.org/abs/2607.00283>

11 CONF

An LLM-Based Framework for Intent-Driven Network Topology Design

ArXiv CS.AI 50%

arXiv:2607.00292v1 Announce Type: cross Abstract: Designing deployable and resilient network topologies from natural language requirements remains a challenging problem in network automation. This work investigates the ability of Large Language Models (LLMs) to generate...

<https://arxiv.org/abs/2607.00292>

2 CONF

RetailSMV: Exocentric vs. Egocentric Adaptation of Foundation Video World Models in Retail

ArXiv CS.AI 50%

arXiv:2607.00310v1 Announce Type: cross Abstract: Foundation video diffusion models are increasingly viewed as world simulators for embodied agents, yet their pretraining on internet-scale generic video leaves them poorly aligned with real-world deployment domains. We study...

<https://arxiv.org/abs/2607.00310>

3 CONF

K-Inverse-RFM: A Modified RFM that Bridges the Gap to Neural Networks for Data-Corrupted Mathematical Tasks

ArXiv CS.AI 50%

arXiv:2607.00329v1 Announce Type: cross Abstract: Recursive Feature Machines (RFMs) are a class of kernel machines that utilize the Average Gradient Outer Product (AGOP) as a mechanism for feature learning. They have been shown to effectively replicate the learning dynamics and...

<https://arxiv.org/abs/2607.00329>

4 CONF

DiscoLoop: Looping Discrete Embeddings and Continuous Hidden States for Multi-hop Reasoning

ArXiv CS.AI 50%

arXiv:2607.00341v1 Announce Type: cross Abstract: Large language models achieve strong performance on many reasoning tasks when allowed to externalize intermediate steps as Chain-of-Thought (CoT). However, many questions require the model to internalize the multi-step reasoning...

<https://arxiv.org/abs/2607.00341>

5 CONF

SoK: Attack and Defense Landscape of Mobile On-device AI Systems

ArXiv CS.AI 50%

arXiv:2607.00362v1 Announce Type: cross Abstract: Mobile on-device AI (MoAI) systems that integrate locally deployed AI models with conventional mobile software components are emerging as a key paradigm for delivering intelligent functionality directly on end-user devices. By...

<https://arxiv.org/abs/2607.00362>

6 CONF

When AI meets quantum information: A comprehensive review

ArXiv CS.AI 50%

arXiv:2607.00365v1 Announce Type: cross Abstract: Artificial intelligence (AI) and quantum information (QI) are rapidly co-evolving. AI is becoming a practical tool for learning, designing, controlling, and verifying quantum systems, while QI offers new computational models,...

<https://arxiv.org/abs/2607.00365>

7 CONF

MEPA: Multi-Scale Representation Alignment for Visual Autoregressive Modeling with Mixture of Experts

ArXiv CS.AI 50%

arXiv:2607.00371v1 Announce Type: cross Abstract: Visual AutoRegressive modeling (VAR) has pioneered a coarse-to-fine multi-scale autoregressive generative paradigm, demonstrating strong capabilities in image generation. However, VAR still suffers from inherent deficiencies in...

<https://arxiv.org/abs/2607.00371>

8 CONF

MalariAI: A Label-Resilient Decoupled Framework for Universal Cell Segmentation and Explainable Stage Classification in Dense Malaria Blood Smears

ArXiv CS.AI 50%

arXiv:2607.00385v1 Announce Type: cross Abstract: Automated malaria diagnosis from blood smear microscopy is a critical challenge in global health AI; in resource-limited settings, the scarcity of expert microscopists remains the primary bottleneck to timely and accurate...

<https://arxiv.org/abs/2607.00385>

9 CONF

SenseWalk: Agent-Based Semantic Trajectory Simulation Powered by Large Language Models in Zoned Environments

ArXiv CS.AI 50%

arXiv:2607.00989v1 Announce Type: cross Abstract: Semantic trajectory analysis has recently emerged as an approach for modeling human movement by capturing implicit patterns and behaviors through semantic information (e.g., visitors' profiles and goals) beyond raw spatial paths...

<https://arxiv.org/abs/2607.00989>

0 CONF

Holographic Quantum Transformer: A Generalist Neuro-Symbolic Architecture for Solving Frustrated Systems via Generative Attention

ArXiv CS.AI 50%

arXiv:2607.00398v1 Announce Type: cross Abstract: Simulating two-dimensional frustrated quantum matter is a grand challenge due to the sign problem and exponential Hilbert space complexity. In this work, we introduce the Holographic Quantum Transformer (HQT), a physics-inspired...

<https://arxiv.org/abs/2607.00398>

1 CONF

The Illusion of High Utility in Safety Alignment of Text-to-Image Diffusion Models

ArXiv CS.AI 50%

arXiv:2607.00402v1 Announce Type: cross Abstract: Safety alignment of text-to-image (T2I) diffusion models aims to suppress harmful generations while preserving utility on benign prompts. Recent methods often appear to deliver high safety with high utility, but this conclusion...

<https://arxiv.org/abs/2607.00402>

2 CONF

EO-VGGT: Orbital Ray-Conditioned 3D Foundation Models for Satellite Multi-View Reconstruction

ArXiv CS.AI 50%

arXiv:2607.00417v1 Announce Type: cross Abstract: In the era of satellite constellations, multi-view optical satellite imagery is pivotal for Earth Observation (EO) and high-quality Digital Surface Model (DSM) reconstruction. Although feed-forward 3D foundation models have...

<https://arxiv.org/abs/2607.00417>

3 CONF

Learning Gait-Aware Quadruped Locomotion with Temporal Logic Specifications

ArXiv CS.AI 50%

arXiv:2607.00442v1 Announce Type: cross Abstract: Reinforcement learning (RL) for quadruped locomotion commonly depends on fixed, hand-crafted, and Markovian reward functions that limit both interpretability of learned policies and lack explicit control over gait behaviors. We...

<https://arxiv.org/abs/2607.00442>

- 4 **CONF**
- Real-Time Hard Negative Sampling via LLM-based Clustering for Large-Scale Two-Tower Retrieval**
- ArXiv CS.AI 50%
- arXiv:2607.00448v1 Announce Type: cross Abstract: The two-tower model has been widely used for large-scale recommendation systems, particularly in the retrieval stage. Industry standards for training two-tower models typically involve in-batch and/or out-of-batch negative...
- <https://arxiv.org/abs/2607.00448>
-
- 5 **CONF**
- Enhancing Hardware Fault Tolerance in Machines with Reinforcement Learning Policy Gradient Algorithms**
- ArXiv CS.AI 50%
- arXiv:2407.15283v2 Announce Type: replace-cross Abstract: Industry is moving toward autonomous, network-connected machines that detect and adapt to changing conditions, including hardware faults. Conventional fault-tolerant design duplicates hardware and reroutes control logic;...
- <https://arxiv.org/abs/2407.15283>
-
- 6 **CONF**
- A Multi-Resolution Finite-Volume Inspired Deep Learning Framework for Spatiotemporal Dynamics Prediction**
- ArXiv CS.AI 50%
- arXiv:2607.00460v1 Announce Type: cross Abstract: Predicting complex spatiotemporal dynamics in physical processes often demands computationally expensive numerical methods or data-driven neural networks that suffer from high training costs, error accumulation, and limited...
- <https://arxiv.org/abs/2607.00460>
-
- 7 **CONF**
- Predicting Lethal Outcome (Cause) And Understanding Key Biomarkers Linked With Acute Myocardial Infarction Using Deep Artificial Neural Network And Ensemble Of Machine Learning Methodologies**
- ArXiv CS.AI 50%
- arXiv:2607.00472v1 Announce Type: cross Abstract: Cardiovascular disease is still one of the main causes of death around the world. Acute myocardial infarction (MI), or heart attack, claims millions of lives each year. MI happens when blood flow to the coronary arteries is...
- <https://arxiv.org/abs/2607.00472>
-
- 8 **CONF**
- Beyond the Prompt: Jailbreaking Function-Calling LLMs via Simulated Moderation Traces**
- ArXiv CS.AI 50%
- arXiv:2607.00481v1 Announce Type: cross Abstract: Jailbreak attacks remain a critical threat to the safe deployment of large language models (LLMs). While prior work has primarily studied attacks and defenses at the prompt level, we show that this prompt-centric paradigm...
- <https://arxiv.org/abs/2607.00481>
-
- 9 **CONF**
- PAPA: Online Personalized Active Preference Alignment**
- ArXiv CS.AI 50%
- arXiv:2607.00486v1 Announce Type: cross Abstract: Diffusion models are highly effective at modeling complex data distributions, including images and text. However, in applications like personalized recommender systems, the objective often shifts to modeling specific regions of...
- <https://arxiv.org/abs/2607.00486>

- 0 **CONF**
MindEdit-Bench: Benchmarking Object-Level Counterfactual Spatial Reasoning in VLMs from In-the-Wild Photos
 ArXiv CS.AI 50%
 arXiv:2607.00491v1 Announce Type: cross Abstract: Benchmarks for vision-language models (VLMs) mostly test observational spatial reasoning: models describe relations already visible in the input. Existing what-if tasks typically vary the observer while keeping the scene fixed....
<https://arxiv.org/abs/2607.00491>
- 1 **CONF**
BaseRT: Best-in-Class LLM Inference on Apple Silicon via Native Metal
 ArXiv CS.AI 50%
 arXiv:2607.00501v1 Announce Type: cross Abstract: We present BaseRT, a native Metal inference runtime for large language models (LLMs) on Apple Silicon, and report the highest inference throughput on this hardware to date. Existing runtimes, including llama.cpp and MLX-based...
<https://arxiv.org/abs/2607.00501>
- 2 **CONF**
AI, Trust, and Teaming: The Humans-as-Handlers Approach for Autonomous and Opaque AI Systems
 ArXiv CS.AI 50%
 arXiv:2607.00523v1 Announce Type: cross Abstract: Artificial intelligence (AI) is becoming ubiquitous, and across domains, increasingly autonomous systems are carrying out tasks which raise significant ethical and legal challenges which demonstrate a need for strong...
<https://arxiv.org/abs/2607.00523>
- 3 **CONF**
From Technical Metrics to User Perception: A User Study of a Multimodal Human-Robot Interaction System for Object Detection and Grasping
 ArXiv CS.AI 50%
 arXiv:2607.00530v1 Announce Type: cross Abstract: Improvements in the technical performance of human--robot interaction (HRI) systems do not automatically translate into differences that human users can detect during live interaction. This paper investigates whether a 15...
<https://arxiv.org/abs/2607.00530>
- 4 **CONF**
Flow-Map GRPO: Reinforcement Learning for Few-Step Flow-Map Generators via Anchored Stochastic Composition
 ArXiv CS.AI 50%
 arXiv:2607.00535v1 Announce Type: cross Abstract: Few-step flow-map generators, such as consistency models and MeanFlow, accelerate sampling by directly learning long-range transport maps between noise and data. However, these models are typically deterministic, which makes them...
<https://arxiv.org/abs/2607.00535>
- 5 **CONF**
Cross-Domain Generalization Failure in Lightweight Intrusion Detection Models for IIoT Networks
 ArXiv CS.AI 50%
 arXiv:2607.00553v1 Announce Type: cross Abstract: Lightweight machine learning models are increasingly proposed for intrusion detection in Industrial Internet of Things (IIoT) networks due to their suitability for resource-constrained edge deployment. Most reported results...
<https://arxiv.org/abs/2607.00553>

6 CONF

A Methodology for Investigating AI Patterns Prevalence in Software Repositories

ArXiv CS.AI 50%

arXiv:2607.00558v1 Announce Type: cross Abstract: As Artificial Intelligence(AI)-based applications take off, a clear understanding of AI patterns can uplift the quality of AI applications. Many AI patterns have been proposed in the literature; however, their prevalence in...

<https://arxiv.org/abs/2607.00558>

7 CONF

Faithful by Definition: Emotion Analysis via Natural Semantic Metalanguage Explications

ArXiv CS.AI 50%

arXiv:2607.00661v1 Announce Type: cross Abstract: Explanations for emotion classifiers are usually produced post hoc, with no guarantee that they reflect the computation behind the label. We present an explication interface for event-based emotion analysis. A parser maps the...

<https://arxiv.org/abs/2607.00661>

8 CONF

LLVM-Bench: Benchmarking and Advancing Large Language Models for LLVM Compiler Issue Resolution

ArXiv CS.AI 50%

arXiv:2607.00700v1 Announce Type: cross Abstract: LLVM is a widely used compiler infrastructure whose scale and complexity make issue resolution labor-intensive and challenging. Although large language models (LLMs) have recently achieved remarkable success in issue resolution,...

<https://arxiv.org/abs/2607.00700>

9 CONF

Creating Impactful Autonomous Driving Datasets: A Strategic Guide from Research Gap to Benchmark

ArXiv CS.AI 50%

arXiv:2607.00710v1 Announce Type: cross Abstract: Well-designed autonomous driving datasets have fundamentally shaped research progress, yet existing literature primarily describes what datasets contain rather than how to strategically design impactful ones. This is especially...

<https://arxiv.org/abs/2607.00710>

0 CONF

Self-conditioned Flow Map Language Models via Fixed-point Flows

ArXiv CS.AI 50%

arXiv:2607.00714v1 Announce Type: cross Abstract: Self-conditioning is a core technique that enhances continuous flow-based language models, where the model learns to denoise generated text by conditioning on its own denoising estimate. While empirically successful, its...

<https://arxiv.org/abs/2607.00714>

1 CONF

Partial Skeleton Visibility for Action Recognition: A Constrained Field-of-View Approach

ArXiv CS.AI 50%

arXiv:2607.00716v1 Announce Type: cross Abstract: Skeleton-based action recognition has achieved remarkable success by exploiting joint coordinates and their topological connections, yet prevailing methods overwhelmingly assume complete and clean skeleton inputs. In real-world...

<https://arxiv.org/abs/2607.00716>

2 **CONF****LLM-Guided ODE Discovery and Parameter Inference from Small-Cohort Aggregate Data**

ArXiv CS.AI 50%

arXiv:2607.00733v1 Announce Type: cross Abstract: Mechanistic modeling via ordinary differential equations (ODEs) provides interpretable descriptions of complex dynamics and enables inference of underlying mechanisms, which is particularly valuable in clinical settings. However,...

<https://arxiv.org/abs/2607.00733>

3 **CONF****ConRTF: Edge-Constrained Boundary Distribution Refinement for Realtime TransFormer Table Structure Recognition**

ArXiv CS.AI 50%

arXiv:2607.00734v1 Announce Type: cross Abstract: Table Structure Recognition (TSR) aims to recover the row and column layout of tables from document images, a key step in document understanding pipelines. Accurate TSR depends on precise boundary localization: small errors in...

<https://arxiv.org/abs/2607.00734>

4 **CONF****Prototype Memory-Guided Training-Free Anomaly Classification and Localization in Prenatal Ultrasound**

ArXiv CS.AI 50%

arXiv:2607.00744v1 Announce Type: cross Abstract: Prenatal anomaly classification and localization is of critical importance for fetal health and pregnancy management. Although ultrasound (US) is the primary modality for prenatal screening, accurate diagnosis remains challenging...

<https://arxiv.org/abs/2607.00744>

5 **CONF****Active Learning for Cascaded Object Detection: Balancing Coverage and Uncertainty in Table Extraction Pipelines**

ArXiv CS.AI 50%

arXiv:2607.00747v1 Announce Type: cross Abstract: Table extraction from business documents relies on a cascaded pipeline where Table Detection (TD) first localizes tables and Table Structure Recognition (TSR) then recovers their internal layout. Building task-specific training...

<https://arxiv.org/abs/2607.00747>

6 **CONF****LeVLJEPA: End-to-End Vision-Language Pretraining Without Negatives**

ArXiv CS.AI 50%

arXiv:2607.00784v1 Announce Type: cross Abstract: Vision-language pretraining remains dominated by contrastive objectives, whereas vision-only self-supervised learning has largely adopted non-contrastive methods. At the same time, the role of vision-language encoders has...

<https://arxiv.org/abs/2607.00784>

7 **CONF****Exploring the Semantic Gap in Agentic Data Systems: A Formative Study of Operationalization Failures in Analytical Workflows**

ArXiv CS.AI 50%

arXiv:2607.00828v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to generate queries, invoke tools, and construct analytical workflows. Although recent advances have substantially improved workflow generation and execution, the semantic...

<https://arxiv.org/abs/2607.00828>

8 CONF

Valdi: Value Diffusion World Models

ArXiv CS.AI 50%

arXiv:2607.00917v1 Announce Type: cross Abstract: World models can enable Model Predictive Control (MPC), but this requires dynamics prediction that is both fast enough for online use and expressive enough to represent uncertain futures. Diffusion models offer a natural...

<https://arxiv.org/abs/2607.00917>

9 CONF

Diffusion-GR2: Diffusion Generative Reasoning Re-ranker

ArXiv CS.AI 50%

arXiv:2607.01170v1 Announce Type: cross Abstract: Generative reasoning re-rankers achieve strong recommendation accuracy by emitting a chain-of-thought before re-ordering a candidate list, but they are slow at inference: an autoregressive (AR) decoder spends one sequential...

<https://arxiv.org/abs/2607.01170>

0 CONF

Aionoscope: Debugging Latent-State Accessibility in Time-Series Representations

ArXiv CS.AI 50%

arXiv:2607.00956v1 Announce Type: cross Abstract: Time-series models are often evaluated by what they can forecast or classify, but those scores do not show whether their representations preserve the process state a user may want to inspect: event timing, phase, amplitude,...

<https://arxiv.org/abs/2607.00956>

1 CONF

TRCGL-Net: A Long-Tailed Multi-Label Chest X-Ray Classification Framework with Generative Data Augmentation and Label Co-Occurrence Modeling

ArXiv CS.AI 50%

arXiv:2607.00975v1 Announce Type: cross Abstract: Chest X-ray multi-label classification is a core task in intelligent medical imaging diagnosis. However, real clinical data often exhibit extreme long-tailed distributions, leading to degraded performance on rare diseases in tail...

<https://arxiv.org/abs/2607.00975>

2 CONF

Reading Order Inference for Complex Document Layouts

ArXiv CS.AI 50%

arXiv:2607.01018v1 Announce Type: cross Abstract: Reading order inference remains a critical bottleneck in the digitization of complex historical manuscripts, where pages contain multiple spatially interleaved reading streams, the canonical example being the Glossa Ordinaria...

<https://arxiv.org/abs/2607.01018>

3 CONF

EchoRisk: A Multicentre Echocardiography Dataset and Benchmark for Cardio-Oncology

ArXiv CS.AI 50%

arXiv:2607.01039v1 Announce Type: cross Abstract: Therapy-induced cardiotoxicity is the leading non-oncological cause of treatment interruption in breast cancer patients, yet early, automated risk stratification from routine cardiac imaging remains an unsolved problem. We...

<https://arxiv.org/abs/2607.01039>

4 **CONF****CausalMix: Data Mixture as Causal Inference for Language Model Training****ArXiv CS.AI** **50%**

arXiv:2607.01104v1 Announce Type: cross Abstract: In Large Language Model (LLM) training, data mixing plays a pivotal role in determining model performance. Recent methods optimize mixture weights via proxy models, but they rely on the assumption of static data distributions. As...

<https://arxiv.org/abs/2607.01104>

5 **CONF****FAR: Failure-Aware Retry for Test-Time Recovery and Continual Policy Improvement****ArXiv CS.AI** **50%**

arXiv:2607.01111v1 Announce Type: cross Abstract: Robot policies inevitably encounter failures when deployed in real environments. Naive retries often repeat the same mistakes, while many existing recovery methods rely on human intervention. In this paper, we propose...

<https://arxiv.org/abs/2607.01111>

6 **CONF****Towards Developing a Multimodal Chat Assistant for University Stakeholders: RAG-based Approach****ArXiv CS.AI** **50%**

arXiv:2607.01115v1 Announce Type: cross Abstract: University stakeholders often face difficulties in accessing timely and reliable information, especially in developing countries, where there are very few intelligent support systems. Existing rule-based chatbots are unable to...

<https://arxiv.org/abs/2607.01115>

7 **CONF****Skills Are Not Islands: Measuring Dependency and Risk in Agent Skill Supply Chains****ArXiv CS.AI** **50%**

arXiv:2607.01136v1 Announce Type: cross Abstract: Agent skills package reusable operational knowledge for Large Language Model (LLM) agents, yet as they grow in scope, they become dependency-bearing artifacts whose identities, versions, and provenance remain implicit. This...

<https://arxiv.org/abs/2607.01136>

8 **CONF****Adversarial Pragmatics for AI Safety Evaluation: A Benchmark for Instruction Conflict, Embedded Commands, and Policy Ambiguity****ArXiv CS.AI** **50%**

arXiv:2607.01153v1 Announce Type: cross Abstract: Safety evaluations for language models increasingly depend on judgments about ambiguous natural-language behaviour: whether a model has followed an instruction, refused appropriately, complied with a policy, resisted an embedded...

<https://arxiv.org/abs/2607.01153>

9 **CONF****Right in the Right Way: LM Training with Verifiable Rewards and Human Demonstrations****ArXiv CS.AI** **50%**

arXiv:2607.01181v1 Announce Type: cross Abstract: RL with verifiable rewards (RLVR) has emerged as a powerful paradigm for training LMs on tasks with well-defined success metrics, such as code generation and mathematical reasoning. However, current RLVR methods optimize only...

<https://arxiv.org/abs/2607.01181>

00 CONF

World from Motion: Generative Dynamic Gaussian Reconstruction from Monocular Video

ArXiv CS.AI 50%

arXiv:2607.01202v1 Announce Type: cross Abstract: We present World from Motion, a method for generating freely renderable dynamic 3D Gaussian representations from monocular videos. Our approach conditions a video model on dense, pixel-aligned renderings that encode appearance,...

<https://arxiv.org/abs/2607.01202>

01 CONF

GPU-Parallel Linearization Error Bounds for Real-Time Robust Optimal Control of Nonlinear and Neural Network Dynamics

ArXiv CS.AI 50%

arXiv:2607.01203v1 Announce Type: cross Abstract: This paper studies real-time robust optimal control for uncertain nonlinear systems, where linear time-varying (LTV) approximations make planning tractable but require sound linearization error bounds (LEBs) to guarantee robust...

<https://arxiv.org/abs/2607.01203>

02 CONF

Evaluating Implicit Biases in LLM Reasoning through Logic Grid Puzzles

ArXiv CS.AI 50%

arXiv:2511.06160v2 Announce Type: replace Abstract: While recent safety guardrails effectively suppress overtly biased outputs, subtler forms of social bias emerge during complex logical reasoning tasks that evade current evaluation benchmarks. To fill this gap, we introduce a...

<https://arxiv.org/abs/2511.06160>

03 CONF

Are Performance-Optimization Benchmarks Reliably Measuring Coding Agents?

ArXiv CS.AI 50%

arXiv:2607.01211v1 Announce Type: cross Abstract: Repository-level performance-optimization benchmarks such as GSO, SWE-Perf and SWE-fficiency evaluate coding agents by applying patches to real repositories and comparing runtime against unoptimized baselines and official...

<https://arxiv.org/abs/2607.01211>

04 CONF

Unexplainability of Artificial Intelligence Judgments and Functional Implementation in Kant's Perspective

ArXiv CS.AI 50%

arXiv:2407.18950v5 Announce Type: replace Abstract: Kant's Critique of Pure Reason, a major contribution to the history of epistemology, proposes a table of categories to elucidate the structure of the a priori principles underlying human judgment. Artificial intelligence (AI)...

<https://arxiv.org/abs/2407.18950>

05 CONF

From Silos to Systems: Process-Oriented Hazard Analysis for AI Systems

ArXiv CS.AI 50%

arXiv:2410.22526v2 Announce Type: replace Abstract: To effectively address potential harms from Artificial Intelligence (AI) systems, it is essential to identify and mitigate system-level hazards. Current analysis approaches focus on individual components of an AI system, like...

<https://arxiv.org/abs/2410.22526>

06 CONF

Large language models replicate and predict human cooperation across experiments in game theory

ArXiv CS.AI 50%

arXiv:2511.04500v3 Announce Type: replace Abstract: Large language models (LLMs) are increasingly deployed as decision-making agents in high-stakes domains and as imitators of human behavior in the social and behavioral sciences. Yet how closely LLMs mirror human decision-making...

<https://arxiv.org/abs/2511.04500>

07 CONF

CoT-X: An Adaptive Framework for Cross-Model Chain-of-Thought Transfer and Optimization

ArXiv CS.AI 50%

arXiv:2511.05747v3 Announce Type: replace Abstract: Chain-of-Thought (CoT) reasoning enhances the problem-solving ability of large language models (LLMs) but leads to substantial inference overhead, limiting deployment in resource-constrained settings. This paper investigates...

<https://arxiv.org/abs/2511.05747>

08 CONF

Diagnosing and Mitigating Compounding Failures in Agentic Persuasion via Taxonomic Strategy Retrieval

ArXiv CS.AI 50%

arXiv:2606.24976v2 Announce Type: replace Abstract: Foundation-model agents in multi-step, open-ended environments frequently suffer from compounding errors, where early mistakes contaminate long-horizon trajectories. While Multi-Agent Debate (MAD) succeeds in deterministic...

<https://arxiv.org/abs/2606.24976>

09 CONF

Heuresis: Search Strategies for Autonomous AI Research Agents Across Quality, Diversity and Novelty

ArXiv CS.AI 50%

arXiv:2606.25198v2 Announce Type: replace Abstract: Autonomous AI Research promises to accelerate the scientific progress of machine learning. To realise this goal, current Large Language Model (LLM)-based agents need to go beyond just writing code, to mastering the exploration...

<https://arxiv.org/abs/2606.25198>

10 CONF

Hey, That's My Model! Introducing Chain & Hash, An LLM Fingerprinting Technique

ArXiv CS.AI 50%

arXiv:2407.10887v4 Announce Type: replace-cross Abstract: Growing concerns over the theft and misuse of Large Language Models (LLMs) underscore the need for effective fingerprinting to link a model to its original version and detect misuse. We define five essential properties...

<https://arxiv.org/abs/2407.10887>

11 CONF

Verbosity Tradeoffs and the Impact of Scale on the Faithfulness of LLM Self-Explanations

ArXiv CS.AI 50%

arXiv:2503.13445v3 Announce Type: replace-cross Abstract: When asked to explain their decisions, LLMs can often give explanations which sound plausible to humans. But are these explanations faithful, i.e. do they convey the factors actually responsible for the decision? In this...

<https://arxiv.org/abs/2503.13445>

12 CONF

Comparative Analysis of Lightweight CNNs for Resource-Constrained Devices: Predictive Performance, Efficiency Trade-offs, and Initialization Effects

ArXiv CS.AI 50%

arXiv:2505.03303v4 Announce Type: replace-cross Abstract: Lightweight convolutional neural networks are often compared using results obtained with different training recipes, input settings, and pretrained checkpoints. Such differences make architecture rankings difficult to...

<https://arxiv.org/abs/2505.03303>

13 CONF

scDataset: Scalable Data Loading for Deep Learning on Large-Scale Single-Cell Omics

ArXiv CS.AI 50%

arXiv:2506.01883v3 Announce Type: replace-cross Abstract: Training deep learning models on single-cell datasets with hundreds of millions of cells requires loading data from disk, as these datasets exceed available memory. While random sampling provides the data diversity needed...

<https://arxiv.org/abs/2506.01883>

14 CONF

FLAT: Revealing Hidden Latent-Conditioned Backdoor Failures in Federated Learning

ArXiv CS.AI 50%

arXiv:2508.04064v2 Announce Type: replace-cross Abstract: Horizontal federated learning (HFL) backdoor audits often summarize model behavior through clean accuracy (CA), mean attack success rate (ASR), or a single known-trigger test. Such summaries can hide a different failure...

<https://arxiv.org/abs/2508.04064>

15 CONF

Predicting LLM Reasoning Performance with Small Proxy Model

ArXiv CS.AI 50%

arXiv:2509.21013v4 Announce Type: replace-cross Abstract: Given the prohibitive cost of pre-training large language models, it is essential to leverage smaller proxy models to optimize datasets before scaling up. However, this approach becomes challenging for reasoning...

<https://arxiv.org/abs/2509.21013>

16 CONF

Toward Cybersecurity-Expert Small Language Models

ArXiv CS.AI 50%

arXiv:2510.14113v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are transforming everyday applications, yet deployment in cybersecurity lags due to a lack of high-quality, domain-specific models and training datasets. To address this gap, we present...

<https://arxiv.org/abs/2510.14113>

17 CONF

Structural Enforcement of Statistical Rigor in AI-Driven Discovery: A Functional Architecture

ArXiv CS.AI 50%

arXiv:2511.06701v3 Announce Type: replace-cross Abstract: AI-Scientist systems risk manufacturing spurious discoveries through uncontrolled multiple testing. We present a functional architecture that enforces statistical rigor at two levels: a Haskell embedded domain-specific...

<https://arxiv.org/abs/2511.06701>

18 CONF

When AI Agents Compete for Jobs: Strategic Capabilities and Economic Dynamics of AI Labour Markets

ArXiv CS.AI 50%

arXiv:2512.04988v2 Announce Type: replace-cross Abstract: Emerging agentic marketplaces provide the economic infrastructure for matching and coordinating the large amounts of AI agents used in agentic swarms. Unlike human workers, AI agents can operate on multiple jobs...

<https://arxiv.org/abs/2512.04988>

19 CONF

Utilizing Earth Foundation Models to Enhance the Simulation Performance of Hydrological Models with AlphaEarth Embeddings

ArXiv CS.AI 50%

arXiv:2601.01558v2 Announce Type: replace-cross Abstract: Predicting river flow in places without streamflow records is challenging because basins respond differently to climate, terrain, vegetation, and soils. Traditional basin attributes describe some of these differences, but...

<https://arxiv.org/abs/2601.01558>

20 CONF

Controllable Diffusion-Based Lesion Inpainting for Scalable Histopathology Data Augmentation

ArXiv CS.AI 50%

arXiv:2601.08127v2 Announce Type: replace-cross Abstract: Expert-annotated training data remains the critical bottleneck for AI in histopathology, particularly for rare pathologies where even dozens of cases may be unavailable. While data augmentation offers a solution, existing...

<https://arxiv.org/abs/2601.08127>

21 CONF

NeuroFilter: Activation-Based Guardrails for Privacy-Conscious LLM Agents

ArXiv CS.AI 50%

arXiv:2601.14660v2 Announce Type: replace-cross Abstract: Agentic Large Language Models (LLMs) are models able to reason, plan, and execute tools over unstructured data. These abilities are enabling transformative applications in domains spanning from personal assistant,...

<https://arxiv.org/abs/2601.14660>

22 CONF

A Two-stage Transformer Framework for Temporal Localization of Distracted Driver Behaviors

ArXiv CS.AI 50%

arXiv:2603.21048v2 Announce Type: replace-cross Abstract: The identification of hazardous driving behaviors from in-cabin video streams is essential for enhancing road safety and supporting the detection of traffic violations and unsafe driver actions. However, current temporal...

<https://arxiv.org/abs/2603.21048>

23 CONF

Crystalite: A Lightweight Transformer for Efficient Crystal Modeling

ArXiv CS.AI 50%

arXiv:2604.02270v2 Announce Type: replace-cross Abstract: Generative models for crystalline materials often rely on equivariant graph neural networks, which capture geometric structure well but are costly to train and slow to sample. We present Crystalite, a lightweight...

<https://arxiv.org/abs/2604.02270>

24 CONF

GUI-Perturbed: Domain Randomization Reveals Systematic Brittleness in GUI Grounding Models

ArXiv CS.AI 50%

arXiv:2604.14262v2 Announce Type: replace-cross Abstract: GUI grounding models report over 85% accuracy on standard benchmarks, yet drop 27-56 percentage points when instructions require spatial reasoning rather than direct element naming. Current benchmarks miss this because...

<https://arxiv.org/abs/2604.14262>

25 CONF

FED-FSTQ: Fisher-Guided Token Quantization for Communication-Efficient Federated Fine-Tuning of LLMs on Edge Devices

ArXiv CS.AI 50%

arXiv:2604.25421v2 Announce Type: replace-cross Abstract: Federated fine-tuning provides a practical route to adapt large language models (LLMs) on edge devices without centralizing private data, yet in mobile deployments the training wall-clock is often bottlenecked by...

<https://arxiv.org/abs/2604.25421>

26 CONF

EcoGEO: Trajectory-Aware Evidence Ecosystems for Web-Enabled LLM Search Agents

ArXiv CS.AI 50%

arXiv:2605.12887v2 Announce Type: replace-cross Abstract: Web-enabled LLM agents are changing how online information influences search outcomes. Existing Generative Engine Optimization (GEO) studies mainly focus on individual webpages. However, agentic web search is not a...

<https://arxiv.org/abs/2605.12887>

27 CONF

Faithful by Construction: Claim-Anchored Attribution for Multi-Document Summarization

ArXiv CS.AI 50%

arXiv:2606.23989v2 Announce Type: replace-cross Abstract: End-to-end large language models (LLMs) produce fluent multi-document summaries but remain prone to hallucination, and the attributions they offer are typically coarse (whole documents or passages) and generated post hoc,...

<https://arxiv.org/abs/2606.23989>

28 CONF

Text Over Image: Auditing Multimodal Robustness in Synthetic Medical Image Detection

ArXiv CS.AI 50%

arXiv:2606.25375v2 Announce Type: replace-cross Abstract: With the rapid adoption of generative AI, synthetic medical images pose growing risks, including diagnostic deception and insurance fraud. Although prior work has explored vision-language model (VLM)-based synthetic image...

<https://arxiv.org/abs/2606.25375>