



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	140	Media
HackerNews	15	Community
TechCrunch AI	5	Media
The Verge AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: Enola-A deterministic architecture graph for developers and AI agents

HackerNews

+1 source

90%

HN 9

arXiv:2606.22737v2 Announce Type: replace Abstract: Before letting an agent operate over real context, can you prove it used the right evidence? GroundEval turns that question into a deterministic test of what the agent searched, fetched, cited, and was permitted to access. In...

<https://github.com/enola-labs/enola/tree/main>

VERI

Show HN: Collection of skills to help improve UX and AI Product design

HackerNews

+1 source

90%

HN 4

arXiv:2607.01480v1 Announce Type: new Abstract: Reinforcement learning with verifiable rewards (RLVR), along with recent selfdistillation variants such as SDPO, evaluates each rollout against a verifier and updates the policy from that episode-level signal. However, the richer...

<https://github.com/tommyjepsen/awesome-ux-skills>

VERI

Show HN: Visualizing Contrastive Language-Audio Pretraining (Clap)

HackerNews

+1 source

90%

HN 3

arXiv:2607.02466v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models are fundamentally bottlenecked by the scarcity of expert demonstrations -- triplets of observations, instructions, and actions that are costly to collect at scale. We argue that this bottleneck...

<https://adamsohn.com/clap/>

VERI

An Exploratory Study on LLM-Generated Code and Comments in Code Repositories

ArXiv CS.AI

+1 source

80%

arXiv:2607.01867v1 Announce Type: cross Abstract: The use of LLMs in software development has become increasingly widespread on tasks such as code generation and summarization. Reports from large technology companies showed that around 20% to 30% of their code are generated by...

<https://arxiv.org/abs/2607.01867>

CONFIRMED SIGNALS — 155 stories

CONF

Show HN: Mail Memories - A desktop app to rescue photos from Gmail

HackerNews

60%

HN 101

Hey HN, I'm the creator of Mail Memories. Like many of you, I've had my Gmail address for more than 20 years. A few years ago, I got curious and wanted to see what photos were buried deep in my account. I ended up finding lots of "lost" pictures of old friends, family members,...

<https://mailmemories.com>

CONF

Show HN: A graph paper generator that renders vector PDFs in the browser

HackerNews

50%

HN 98

No summary — see link for full article.

<https://freegraphpaper.net/>

CONF

Show HN: CLI tool for detecting non-exact code duplication with embedding models

HackerNews

50%

HN 80

No summary — see link for full article.

<https://github.com/rafal-qa/slopo>

CONF

Show HN: Claudoro, Pomodoro timer embedded in the Claude Code statusline

HackerNews 50% HN 45

3 weeks ago I had a nasty accident and fractured my vertebrae. As I lay in bed I needed something to take my mind off it all so built "Claudoro". Claudoro is a pomodoro timer built right into the Claude Code status line, as well as can be directly controlled from Claude Code and...

<https://github.com/emson/claudoro>

CONF

Show HN: Declaw Arena - a CTF-style challenge to break an AI agent in a microVM

HackerNews 50% HN 7

No summary — see link for full article.

<https://declaw.ai/arena>

0

CONF

Show HN: Mirrors - test AI agent changes by replaying production traces

HackerNews 50% HN 6

No summary — see link for full article.

<https://www.runmirrors.com/>

1

CONF

Show HN: NeuralFit game - Adjust the neural network weights manually

HackerNews 50% HN 5

No summary — see link for full article.

<https://neuralfit.ai201.site/>

2

CONF

Show HN: digga - DNS, RDAP/WHOIS, subdomains and more in one lookup

HackerNews 50% HN 5

No summary — see link for full article.

<https://digga.dev>

3

CONF

Show HN: Bedtimeforkids let kids learn while entertain

HackerNews 50% HN 2

No summary — see link for full article.

<https://bedtimeforkids.vercel.app>

4

CONF

Show HN: A 155K-param transformer builds a map of a world it's never shown

HackerNews 50% HN 2

No summary — see link for full article.

<https://ankur-chr.github.io/inside-the-model/>

5

CONF

Show HN: NEUX A daily brief on Europe's critical industries

HackerNews 50% HN 2

I built NEUX because I feel a lot of the recent focus is on the AI industry in US and China. I want a single place to learn about Europe's critical industrial system and what companies are changing the status quo. I'm publishing a brief per area, each with a headline, a short...

<https://neuxresearch.com/>

6 CONF

Show HN: AnythingLLM Fork as NPM Package

HackerNews 50% HN 2

I am a huge fan of anythingLLM, I've used it extensively. I had use cases that it doesn't quite fit; so I decided to fork it as an easy to install, lighter weight package. I've been changing the agent mode code to focus more on automation. I'd love your feedback!

<https://github.com/fred-terzi/totem-llm>

7 CONF

Mark Zuckerberg tells staff that AI agents haven't progressed as quickly as he'd hoped

TechCrunch AI 50%

At an internal meeting, the Meta CEO reportedly said that AI development efforts were not moving as quickly as anticipated.

<https://techcrunch.com/2026/07/02/mark-zuckerberg-tells-staff-that-ai-agents-havent-pro...>

8 CONF

Jersey Mike's IPO illustrates how bad the AI hype has become

TechCrunch AI 50%

Just for kicks, I took a look at Jersey Mike's IPO documents. Surely a sandwich shop would have no need to mention AI. But lo and behold.

<https://techcrunch.com/2026/07/02/jersey-mikes-ipo-illustrates-how-bad-the-ai-hype-has-...>

9 CONF

Anthropic is discussing a new custom chip with Samsung

TechCrunch AI 50%

The news comes about a week after OpenAI announced its own custom AI chip in a partnership with Broadcom.

<https://techcrunch.com/2026/07/02/anthropic-is-discussing-a-new-custom-chip-with-samsung/>

0 CONF

OpenAI proposed donating 5% of its equity to a US sovereign wealth fund

TechCrunch AI 50%

OpenAI CEO Sam Altman has reportedly proposed giving 5% of the company's equity to a U.S. sovereign wealth fund, reviving discussions about letting the public share in the financial gains from the AI boom.

<https://techcrunch.com/2026/07/02/openai-proposed-donating-5-of-its-equity-to-a-us-sove...>

1 CONF

Microsoft launches its own AI deployment company with \$2.5 billion commitment

TechCrunch AI 50%

Microsoft follows Amazon, OpenAI, and Anthropic with its new AI deployment group.

<https://techcrunch.com/2026/07/02/microsoft-launches-its-own-ai-deployment-company-with...>

2 CONF

Teaching AI to run with the turbines

MIT Tech Review 50%

Artificial intelligence may have captured the public imagination through chatbots and image generators, but some of its most consequential use cases are unfolding far from consumer-facing tools. In industries where physical infrastructure, operational continuity, and safety are...

<https://www.technologyreview.com/2026/07/02/1138433/teaching-ai-to-run-with-the-turbines/>

3 CONF

OpenAI floats giving Trump administration 5 percent cut of AI boom

The Verge AI 50%

OpenAI has floated giving the US government a 5 percent ownership stake as a way of easing tensions with the Trump administration and blunting mounting public backlash against AI, according to the Financial Times. CEO Sam Altman argued that giving the public a financial interest...

<https://www.theverge.com/ai-artificial-intelligence/960588/openai-government-5-percent-...>

4 **CONF****ThreadWeaver: Adaptive Threading for Efficient Parallel Reasoning in Language Models****ArXiv CS.AI** **50%**

arXiv:2512.07843v2 Announce Type: replace-cross Abstract: Scaling inference-time computation has enabled Large Language Models (LLMs) to achieve strong reasoning performance, but their inherently sequential decoding incurs substantial latency, motivating parallelization of the...

<https://arxiv.org/abs/2512.07843>

5 **CONF****CreativityNeuro: Steering Language Model Weights to Improve Divergent Thinking and Reduce Mode Collapse****ArXiv CS.AI** **50%**

arXiv:2607.01433v1 Announce Type: new Abstract: Divergent thinking is a crucial aspect of creativity, yet large language models (LLMs) tend to consistently generate similar responses to open-ended questions, in what has been termed the artificial hivemind effect. Here, we...

<https://arxiv.org/abs/2607.01433>

6 **CONF****Discrete Diffusion Language Models for Interactive Radiology Report Drafting****ArXiv CS.AI** **50%**

arXiv:2607.01436v1 Announce Type: new Abstract: Diffusion language models, which generate text by denoising a token canvas bidirectionally instead of emitting tokens left to right, have become competitive with autoregressive (AR) generation. Medical foundation models, however,...

<https://arxiv.org/abs/2607.01436>

7 **CONF****Revisiting Chain-of-Thought Reasoning under Limited Supervision: Semi-supervised Chain-of-Thought Learning****ArXiv CS.AI** **50%**

arXiv:2607.01511v1 Announce Type: new Abstract: Chain-of-thought (CoT) reasoning has emerged as an effective approach for activating latent reasoning capabilities in large language models. However, most existing CoT methods use reasoning chains mainly as inference-time prompts,...

<https://arxiv.org/abs/2607.01511>

8 **CONF****EO-Agents: A Three-Agent LLM Pipeline for Earth Observation Hypothesis Generation****ArXiv CS.AI** **50%**

arXiv:2607.01584v1 Announce Type: new Abstract: Large language models have recently been explored for scientific hypothesis generation, but most prior work relies on unstructured literature and free-form textual claims. We present a pipeline for Earth observation that grounds...

<https://arxiv.org/abs/2607.01584>

9 **CONF****Safe and Adaptive Cloud Healing: Verifying LLM-Generated Recovery Plans with a Neural-Symbolic World Model****ArXiv CS.AI** **50%**

arXiv:2607.01595v1 Announce Type: new Abstract: As the scale and complexity of cloud-based AI systems continue to escalate, ensuring service reliability through rapid fault detection and adaptive recovery has become a critical challenge. While existing approaches integrate Large...

<https://arxiv.org/abs/2607.01595>

0 **CONF****HERMES: A Multi-Granularity Labeling Substrate for Pre-training Data Mixtures**

ArXiv CS.AI 50%

arXiv:2607.02266v1 Announce Type: cross Abstract: Most data-mixing methods assume the corpus has already been partitioned into groups, and the choice of those groups determines what a mixer can express. Existing labels, including provenance, topic or format taxonomies, and flat...

<https://arxiv.org/abs/2607.02266>

1 **CONF****Scaling with Confidence: Calibrating Confidence of LLMs for Adaptive Test Time Scaling**

ArXiv CS.AI 50%

arXiv:2607.01612v1 Announce Type: new Abstract: Training large language models (LLMs) with reinforcement learning (RL) has significantly advanced their performance on reasoning and question-answering tasks. However, prevailing RL reward designs typically prioritize response...

<https://arxiv.org/abs/2607.01612>

2 **CONF****Spatial Support Matters: Geometry-Aware Graph Fusion for Rainfall Field Reconstruction**

ArXiv CS.AI 50%

arXiv:2607.01621v1 Announce Type: new Abstract: Fine-scale rainfall reconstruction is critical for urban flood modeling, but real rainfall sensing systems observe the field through incompatible spatial supports: gauges measure points, microwave links measure paths, and...

<https://arxiv.org/abs/2607.01621>

3 **CONF****Autonomous discovery of traffic laws with AI traffic scientists**

ArXiv CS.AI 50%

arXiv:2607.01639v1 Announce Type: new Abstract: Universal traffic laws describe recurrent patterns in congestion, mobility and driving behavior across cities, providing a scientific basis for transportation planning, management and control. Their discovery, however, remains...

<https://arxiv.org/abs/2607.01639>

4 **CONF****Separating Expert Retention from Autonomous Source Inference in Raw-ECG-Replay-Free Continual ECG Deployment**

ArXiv CS.AI 50%

arXiv:2607.01674v1 Announce Type: new Abstract: In multi-source ECG deployment, models may need to incorporate new data sources when earlier raw ECGs cannot be retained or replayed. Freezing a pretrained backbone and assigning each source an isolated classifier prevents...

<https://arxiv.org/abs/2607.01674>

5 **CONF****Epistemic Goggles: A Pretrained Module that Induces an Epistemic Frame via Gradient Editing**

ArXiv CS.AI 50%

arXiv:2607.01690v1 Announce Type: new Abstract: Finetuning a language model on documents that are explicitly annotated as fictional results in a model that still actually believes the documents' core claims, an effect known as Negation Neglect. In our evaluations, models trained...

<https://arxiv.org/abs/2607.01690>

6 CONF

Mixture-of-Parallelisms: Towards Memory-Efficient Training Stack for Mixture-of-Experts Models

ArXiv CS.AI 50%

arXiv:2607.01844v1 Announce Type: cross Abstract: This paper showcases a memory-efficient training stack for Mixture-of-Experts (MoE) models. It is a training paradigm that combines and specializes various existing and novel parallelism techniques at different layers and stages...

<https://arxiv.org/abs/2607.01844>

7 CONF

Optimizing RAG Rerankers with LLM Feedback via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2604.02091v2 Announce Type: replace-cross Abstract: Rerankers play a pivotal role in refining retrieval results for Retrieval-Augmented Generation. However, current reranking models are typically optimized on static human annotated relevance labels in isolation, decoupled...

<https://arxiv.org/abs/2604.02091>

8 CONF

BLAgent: Agentic RAG for File-Level Bug Localization

ArXiv CS.AI 50%

arXiv:2605.17965v2 Announce Type: replace-cross Abstract: Bug localization remains a key bottleneck for large language model (LLM)-based software maintenance, where accurately identifying faulty code is essential for debugging, root cause analysis, triage, and automated program...

<https://arxiv.org/abs/2605.17965>

9 CONF

Repair the Amplifier, Not the Symptom: Stable World-Model Correction for Agent Rollouts

ArXiv CS.AI 50%

arXiv:2607.01767v1 Announce Type: new Abstract: As agent planning moves from short tool chains toward persistent workflows with thousands or tens of thousands of steps, failures will occur inside large planning graphs rather than in isolated predictions. Replanning the entire...

<https://arxiv.org/abs/2607.01767>

10 CONF

Subliminal Clocks: Latent Time Modelling in Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2607.01774v1 Announce Type: new Abstract: Diffusion Language Models (DLMs) have recently emerged as a promising alternative to autoregressive models. Unlike standard diffusion-based approaches, DLMs are not explicitly conditioned on a timestep, raising a natural question:...

<https://arxiv.org/abs/2607.01774>

11 CONF

Safety Testing LLM Agents at Scale: From Risk Discovery to Evidence-Grounded Verification

ArXiv CS.AI 50%

arXiv:2607.01793v1 Announce Type: new Abstract: LLM agents increasingly perform autonomous actions through external tools, leading to complex and evolving safety risks. However, existing safety testing targets expert-designed safety violations, and the corresponding outcomes are...

<https://arxiv.org/abs/2607.01793>

2 CONF

MMIR-TCM: Memory-Integrated Multimodal Inference and Retrieval for TCM Clinical Decision Support

ArXiv CS.AI 50%

arXiv:2607.01814v1 Announce Type: new Abstract: Traditional Chinese Medicine (TCM) diagnosis, particularly through tongue inspection, faces persistent challenges in subjectivity and reproducibility. The application of multimodal artificial intelligence to TCM clinical tasks,...

<https://arxiv.org/abs/2607.01814>

3 CONF

Pre-Flight: A Benchmark for Evaluating Large Language Models on Aviation Operational Knowledge

ArXiv CS.AI 50%

arXiv:2607.01829v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly proposed for aviation business operations, from documentation and training generation to customer facing assistants. General purpose benchmarks do not measure whether a model reasons...

<https://arxiv.org/abs/2607.01829>

4 CONF

CLAP: Closed-Loop Training, Evaluation, and Release Control for Domain Agent Post-training

ArXiv CS.AI 50%

arXiv:2607.01846v1 Announce Type: new Abstract: Domain agents often face noisy business data, uncertain post-training gains, offline/application mismatch, and adapter-release risk. This paper presents CLAP (Closed-Loop Agent Post-training), a closed-loop method that converts...

<https://arxiv.org/abs/2607.01846>

5 CONF

Safety Targeted Embedding Exploit via Refinement

ArXiv CS.AI 50%

arXiv:2607.01859v1 Announce Type: new Abstract: Safety training for large language models (LLMs) is conducted predominantly in English, leaving uncertain how well safety mechanisms generalize to low-resource languages and mixed-language code-switching. We show that this creates...

<https://arxiv.org/abs/2607.01859>

6 CONF

Spec-AUF: Accept-Until-Fail Training under Train-Inference Misalignment for Masked Block Drafters

ArXiv CS.AI 50%

arXiv:2607.01893v1 Announce Type: new Abstract: Speculative decoding accelerates autoregressive generation by drafting a block of tokens that the target model verifies left-to-right, committing only the longest accepted prefix. Block (DLM-style) drafters predict the whole block...

<https://arxiv.org/abs/2607.01893>

7 CONF

Rethinking Complexity Metrics for LLM-Integrated Applications: Beyond Source Code

ArXiv CS.AI 50%

arXiv:2607.01903v1 Announce Type: new Abstract: LLM-integrated applications blend natural language prompts with program code, and much of their runtime behavior originates in the prompt layer rather than in the code itself. Existing complexity metrics, however, operate solely at...

<https://arxiv.org/abs/2607.01903>

8 CONF

ContextSniper: AntTrail's Token-Efficient Code Memory for Repository-Level Program Repair

ArXiv CS.AI 50%

arXiv:2607.01916v1 Announce Type: new Abstract: Large language model agents can repair real repository issues, but they often spend large context budgets on whole-file reads, broad searches, and long terminal outputs where useful evidence is mixed with irrelevant code and logs....

<https://arxiv.org/abs/2607.01916>

9 CONF

A-TMA: Decoupling State-Aware Memory Failures in Long-Term Agent Memory

ArXiv CS.AI 50%

arXiv:2607.01935v1 Announce Type: new Abstract: Long term memory lets LLM agents act as persistent assistants, but user facts change. A useful memory system must know what is true now, what used to be true, and what changed. We study \emph{ghost memory}, a state coordination...

<https://arxiv.org/abs/2607.01935>

0 CONF

Multimodal Knowledge Edit-Scoped Generalization for Online Recursive MLLM Editing

ArXiv CS.AI 50%

arXiv:2607.01978v1 Announce Type: new Abstract: Online multimodal knowledge editing requires injecting a continual stream of visual-textual corrections into multimodal large language models (MLLMs) with bounded overhead and minimal disruption to unrelated behaviors. Existing...

<https://arxiv.org/abs/2607.01978>

1 CONF

Traceable Fault Diagnosis for Battery Energy Storage Systems via Retrieval-Augmented Multi-Agent O&M Assistant

ArXiv CS.AI 50%

arXiv:2607.01992v1 Announce Type: new Abstract: Large-scale battery energy storage systems (BESSs) require O&M decisions that combine alarms, cell-level measurements, device topology, diagnostic tables, historical cases, and maintenance documents. Monitoring platforms can flag...

<https://arxiv.org/abs/2607.01992>

2 CONF

Hidden Forgetting in Continual Multimodal Learning: When Accuracy Survives but Grounding Fails

ArXiv CS.AI 50%

arXiv:2607.02020v1 Announce Type: new Abstract: Multimodal large language models must continually adapt to evolving tasks and domains, yet standard continual learning metrics mainly measure whether old answers remain correct, leaving the stability of multimodal grounding largely...

<https://arxiv.org/abs/2607.02020>

3 CONF

Algebraic Model Counting for Global Analysis of Optimal Decision Trees

ArXiv CS.AI 50%

arXiv:2607.02069v1 Announce Type: new Abstract: Ensuring model reliability in Explainable AI requires a global assessment of the hypothesis space. We propose a formal framework for the exhaustive analysis of optimal and near-optimal decision trees, called Algebraic Decision Tree...

<https://arxiv.org/abs/2607.02069>

4 **CONF****XSearch: Explainable Code Search via Concept-to-Code Alignment**

ArXiv CS.AI 50%

arXiv:2605.16046v2 Announce Type: replace-cross Abstract: Semantic code search has been widely adopted in both academia and industry. These approaches embed natural-language queries and code snippets into a shared embedding space and retrieve results based on vector similarity....

<https://arxiv.org/abs/2605.16046>

5 **CONF****Enhancing Fitness Intelligence through Domain-Specific LLM Post-Training**

ArXiv CS.AI 50%

arXiv:2607.02118v1 Announce Type: new Abstract: Scientific Fitness Coaching (SFC) is typically delivered by human professionals, making it costly and inaccessible to many. While recent advances in Large Language Models (LLMs) show considerable promise for more inclusive fitness...

<https://arxiv.org/abs/2607.02118>

6 **CONF****Coding-agents can replicate scientific machine learning papers**

ArXiv CS.AI 50%

arXiv:2607.02134v1 Announce Type: new Abstract: Scientific machine learning papers typically make computational claims, e.g., that the relative mean square error is less than 5% or that the 95% predictive credible interval covers the test data. A coding agent can be prompted to...

<https://arxiv.org/abs/2607.02134>

7 **CONF****A²utoLPBench: An Auto-Generated, Agent-Friendly LP Benchmark via Inverse-KKT Construction**

ArXiv CS.AI 50%

arXiv:2607.02141v1 Announce Type: new Abstract: Most LP-from-text benchmarks are static datasets of word problems written and labeled by hand. Once such a dataset is released, its size is fixed, its difficulty is fixed, and every problem can leak into the training data of future...

<https://arxiv.org/abs/2607.02141>

8 **CONF****A rubric-based controlled comparison of frontier language models on expert-authored clinical reasoning tasks**

ArXiv CS.AI 50%

arXiv:2607.02175v1 Announce Type: new Abstract: Multiple-choice medical benchmarks are increasingly saturated, and recent rubric-based evaluations such as HealthBench have shown that open-ended clinical performance is far from solved - its "Hard" subset top score remains 32%. We...

<https://arxiv.org/abs/2607.02175>

9 **CONF****UA-ChatDev: Uncertainty-Aware Multi-Agent Collaboration for Reliable Software Development**

ArXiv CS.AI 50%

arXiv:2607.02186v1 Announce Type: new Abstract: Software development is a complex task that demands cooperation among agents with diverse roles. Large language models (LLMs) have enabled autonomous multi-agent software development frameworks that leverage role-based...

<https://arxiv.org/abs/2607.02186>

0 **CONF****Criticality-Based Guard Rail Validation for AI Agent Decisions in Autonomous Telecom Networks****ArXiv CS.AI** **50%**

arXiv:2607.02210v1 Announce Type: new Abstract: The evolution toward fully autonomous telecommunications networks (Autonomous Network Levels 4-5) requires AI/ML agents to make real-time network decisions without human intervention. However, no standardized runtime mechanism...

<https://arxiv.org/abs/2607.02210>

1 **CONF****AgenticDataBench: A Comprehensive Benchmark for Data Agents****ArXiv CS.AI** **50%**

arXiv:2607.01647v1 Announce Type: cross Abstract: Data science aims to derive actionable insights from heterogeneous raw data, unlocking the value of the massive amounts of data generated in modern society. Automating this process is essential to reducing labor-intensive efforts...

<https://arxiv.org/abs/2607.01647>

2 **CONF****Grounded autonomous research: a fault-tolerant LLM pipeline from corpus to manuscript in frontier computational physics****ArXiv CS.AI** **50%**

arXiv:2607.02329v1 Announce Type: new Abstract: Autonomous-research agents have demonstrated end-to-end LLM automation in machine-learning sandboxes where execution provides calibration. Frontier physical science differs categorically: physical reasoning underlies every...

<https://arxiv.org/abs/2607.02329>

3 **CONF****Leveraging Metamemory Agent for Enhanced Data-Free Code Generation in Large Language Models****ArXiv CS.AI** **50%**

arXiv:2501.07892v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have shown strong performance in automated code generation, with few-shot prompting widely used for its simplicity and effectiveness. However, few-shot methods depend on curated or manually...

<https://arxiv.org/abs/2501.07892>

4 **CONF****BaRA: Budget-constrained and Reliable Web Data Collection Agent****ArXiv CS.AI** **50%**

arXiv:2607.00007v2 Announce Type: replace-cross Abstract: Large language model (LLM)-based web agents automate web navigation and data collection. However, live web data collection demands capabilities beyond task completion: agents must discover site-internal pages and retrieve...

<https://arxiv.org/abs/2607.00007>

5 **CONF****What LLM Agents Say When No One Is Watching: Social Structure and Latent Objective Emergence in Multi-Agent Debates****ArXiv CS.AI** **50%**

arXiv:2607.02507v1 Announce Type: new Abstract: LLM agents will increasingly act in socially structured settings where role, audience, and relational context can shape what is advantageous or costly to say. We study whether such social structure, without any explicit objective...

<https://arxiv.org/abs/2607.02507>

6 CONF

ReContext: Recursive Evidence Replay as LLM Harness for Long-Context Reasoning

ArXiv CS.AI 50%

arXiv:2607.02509v1 Announce Type: new Abstract: Understanding and reasoning over long contexts has become a key requirement for deploying large language models (LLMs) in realistic applications. Although recent LLMs support increasingly long context windows, they often fail to...

<https://arxiv.org/abs/2607.02509>

7 CONF

Online Safety Monitoring for LLMs

ArXiv CS.AI 50%

arXiv:2607.02510v1 Announce Type: new Abstract: Despite alignment training, LLMs remain prone to generating unsafe outputs at deployment time. Monitoring outputs online and raising an alarm when safety can no longer be assumed is therefore critical. We study a simple real-time...

<https://arxiv.org/abs/2607.02510>

8 CONF

Distributed Attacks in Persistent-State AI Control

ArXiv CS.AI 50%

arXiv:2607.02514v1 Announce Type: new Abstract: As AI coding agents become more autonomous, they increasingly ship code iteratively, with the codebase persisting across sessions. This persistence creates a new attack surface: a misaligned or prompt-injected agent can distribute...

<https://arxiv.org/abs/2607.02514>

9 CONF

Kara: Efficient Reasoning LLM Serving via Sliding-Window KV Cache Compression

ArXiv CS.AI 50%

arXiv:2607.01237v1 Announce Type: cross Abstract: Reasoning language models often generate long chain-of-thought (CoT), which accumulates a massive KV cache during the decoding phase and incurs high decoding latency and limited throughput. To address these issues, KV cache...

<https://arxiv.org/abs/2607.01237>

10 CONF

Breaking Safety at the Token Boundary: How BPE Tokenization Creates Exploitable Gaps in LLM Alignment

ArXiv CS.AI 50%

arXiv:2607.01239v1 Announce Type: cross Abstract: Character-level perturbations bypass safety alignment in modern LLMs despite leaving prompts human-readable. We identify and test a central structural mechanism: BPE tokenization fragments safety-critical words into sub-word...

<https://arxiv.org/abs/2607.01239>

1 CONF

ExPerT: Personalizing LLM Responses to Users' Domain Expertise via Query-Wise Semantic and Keystroke Behavioral Cues

ArXiv CS.AI 50%

arXiv:2607.01242v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used by end users, yet existing personalization methods relying on static profiles or text-only signals fail to capture query-specific expertise variation. We present ExPerT, a...

<https://arxiv.org/abs/2607.01242>

2 CONF

Office Comprehension Benchmark

ArXiv CS.AI 50%

arXiv:2607.01245v1 Announce Type: cross Abstract: We introduce Office Comprehension Bench (OCB), the first public benchmark to jointly evaluate LLM systems on Word, Excel, and PowerPoint comprehension over native file formats (.docx, .xlsx, .pptx) and their variants. OCB...

<https://arxiv.org/abs/2607.01245>

3 CONF

LLMs as Teaching Assistants for Mathematics Exam Grading: Reliability, and Practical Usability

ArXiv CS.AI 50%

arXiv:2607.01247v1 Announce Type: cross Abstract: Open-ended mathematics exams are valuable because they assess reasoning, proof construction, algorithmic thinking, and communication of intermediate steps. They are also difficult to grade at scale because instructors must apply...

<https://arxiv.org/abs/2607.01247>

4 CONF

VLAFlow: A Unified Training Framework for Vision-Language-Action Models via Co-training and Future Latent Alignment

ArXiv CS.AI 50%

arXiv:2607.01586v1 Announce Type: cross Abstract: Vision-language-action models (VLAs) have recently advanced robotic manipulation, yet the effects of different robot-data pre-training paradigms remain difficult to compare because existing models often differ in architecture,...

<https://arxiv.org/abs/2607.01586>

5 CONF

Structuring the Space of Sociotechnical Alignment

ArXiv CS.AI 50%

arXiv:2607.01250v1 Announce Type: cross Abstract: Sociotechnical alignment concerns the social desirability of AI behavior and is thus inherently normative, not merely technical. While NLP research increasingly addresses its technical aspects, it often leaves underspecified what...

<https://arxiv.org/abs/2607.01250>

6 CONF

How Indian Dermatologists are Utilizing Artificial Intelligence for Clinical Practice and Workflow Management: A Nationwide Survey with a Special Focus on atopic dermatitis

ArXiv CS.AI 50%

arXiv:2607.01252v1 Announce Type: cross Abstract: Background: Dermatology AI has mainly focused on image-based diagnosis, while chronic disease workflows have received less attention. We surveyed Indian dermatologists to map routine clinical challenges, with a focus on atopic...

<https://arxiv.org/abs/2607.01252>

7 CONF

Beyond Detection: Redesigning Assessment and Governande of Generative AI at the Universidad Polit'ecnica de Madrid (UPM)

ArXiv CS.AI 50%

arXiv:2607.01255v1 Announce Type: cross Abstract: Universities have responded to generative artificial intelligence (GenAI) in noticeably different ways, both internationally and within Spain. So far, the dominant reaction has been defensive, this is, most institutions frame the...

<https://arxiv.org/abs/2607.01255>

8 **CONF****AI Assistance for Human Review of Default Judgments**

ArXiv CS.AI 50%

arXiv:2607.01256v1 Announce Type: cross Abstract: Overwhelmed courts in the United States review millions of default judgments each year. Unfortunately, such manual reviews are time-consuming and prone to error. In an audit of 188 debt collection cases granted default judgment...

<https://arxiv.org/abs/2607.01256>
9 **CONF****Artificial Intelligence-Enabled Accounting Information Systems and Fraud Detection in Nigeria's Financial Services Sector: The Moderating Role of Natural Language Processing**

ArXiv CS.AI 50%

arXiv:2607.01257v1 Announce Type: cross Abstract: The rapid digitalisation of financial systems has improved operational efficiency and financial inclusion while simultaneously increasing exposure to sophisticated forms of cyber-enabled fraud and electronic financial misconduct....

<https://arxiv.org/abs/2607.01257>
0 **CONF****The Rising Unsustainability of AI Graphics Cards Production**

ArXiv CS.AI 50%

arXiv:2607.01258v1 Announce Type: cross Abstract: The rapid advancement of Artificial Intelligence (AI) has been accompanied by significant increases in computational and environmental costs, driven by large-scale investments in AI infrastructure, hardware, and software. In...

<https://arxiv.org/abs/2607.01258>
1 **CONF****Neuron-Aware Active Few-Shot Learning for LLMs**

ArXiv CS.AI 50%

arXiv:2607.02423v1 Announce Type: cross Abstract: Active Few-Shot Learning (AFSL) adapts LLMs to specialized domains by identifying the most valuable unlabeled samples for annotation and use as few-shot demonstrations, effectively reducing human annotation costs while promoting...

<https://arxiv.org/abs/2607.02423>
2 **CONF****Black-Box Inference of LLM Architectural Properties with Restrictive API Access**

ArXiv CS.AI 50%

arXiv:2607.01313v1 Announce Type: cross Abstract: In practice, most commercial LLM providers do not publicly release details of underlying LLM architectures. However, prior work has shown that given limited API access to an LLM (namely, top-\$k\$ logits and/or a logit bias...

<https://arxiv.org/abs/2607.01313>
3 **CONF****Multi-modal Rail Crossing Safety Analysis**

ArXiv CS.AI 50%

arXiv:2607.01365v1 Announce Type: cross Abstract: Given one or more images of a railway crossing, can we leverage visual cues that allow us to robustly estimate how safe it is? Can we improve our ability to do so by introducing structured data (such as official accident reports)...

<https://arxiv.org/abs/2607.01365>

4 **CONF****AI-enabled gravitational-waves searches for binary neutron stars at optimal sensitivity****ArXiv CS.AI** **50%**

arXiv:2607.01372v1 Announce Type: cross Abstract: Gravitational Waves (GWs) represent the newest window of astronomy, furthering our understanding of compact objects like black holes and neutron stars in the Universe. The signal from two merging neutron stars is especially...

<https://arxiv.org/abs/2607.01372>

5 **CONF****How Should Transformers Encode Numeric Values in Electronic Health Records?****ArXiv CS.AI** **50%**

arXiv:2607.01391v1 Announce Type: cross Abstract: How do we encode numeric values in transformer-based sequence processing, particularly in electronic health record (EHR) data? We systematically compare discrete, continuous, and hybrid value encoding strategies using synthetic...

<https://arxiv.org/abs/2607.01391>

6 **CONF****NeuroBridge: Bridging Multi-Task MRI Knowledge for Neurodegenerative Disease Diagnosis****ArXiv CS.AI** **50%**

arXiv:2607.01401v1 Announce Type: cross Abstract: INTRODUCTION: Accurate MRI-based identification of Alzheimer's disease (AD), mild cognitive impairment (MCI), and related dementias remains challenging because disease-related structural changes are often subtle and...

<https://arxiv.org/abs/2607.01401>

7 **CONF****GPUAlert: A Zero-Instrumentation Process-Boundary Monitor for Diagnosing GPU Training-Job Failures****ArXiv CS.AI** **50%**

arXiv:2607.01409v1 Announce Type: cross Abstract: GPU training jobs fail often, roughly two in five on large production clusters, yet the operator typically learns of a failure only by reconnecting hours later. Experiment trackers require editing the training script and...

<https://arxiv.org/abs/2607.01409>

8 **CONF****Adoption and Impact of Command-Line AI Coding Agents: A Study of Microsoft's Early 2026 Rollout of Claude Code and GitHub Copilot CLI****ArXiv CS.AI** **50%**

arXiv:2607.01418v1 Announce Type: cross Abstract: Organizations rolling out agentic command line tools like Anthropic's Claude Code and GitHub's Copilot CLI need to know who will try them, who will keep using them, and whether the tools produce enough output to justify their...

<https://arxiv.org/abs/2607.01418>

9 **CONF****MultAttnAttrib: Training-Free Multimodal Attribution in Long Document Question Answering****ArXiv CS.AI** **50%**

arXiv:2607.01420v1 Announce Type: cross Abstract: As grounded QA systems are increasingly deployed in AI assistants, accurately attributing generated answers to evidence is critical for user trust and model safety. While unimodal attributions have been explored in depth, the...

<https://arxiv.org/abs/2607.01420>

0 CONF

Search for Truth from Reasoning: A Dynamic Representation Editing Framework for Steering LLM Trajectories

ArXiv CS.AI 50%

arXiv:2606.28589v2 Announce Type: replace Abstract: Current approaches to enhance Large Language Model (LLM) reasoning, such as Chain-of-Thought and "Wait" prompts, primarily encourage models to think more, yet often fail to guide them toward Truth. While Representation Editing...

<https://arxiv.org/abs/2606.28589>

1 CONF

On the Utility and Factual Reliability of Pruned Mixture-of-Experts Models in the Biomedical Domain

ArXiv CS.AI 50%

arXiv:2607.01444v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) models offer inference speedups via selective activation but impose substantial memory requirements because the whole network must remain loaded. Structured expert pruning is a practical approach for...

<https://arxiv.org/abs/2607.01444>

2 CONF

Grounded Iterative Language Planning: How Parameterized World Models Reduce Hallucination Propagation in LLM Agents

ArXiv CS.AI 50%

arXiv:2606.27806v2 Announce Type: replace Abstract: World models for language agents come in two useful forms. An agent-based world model calls an LLM API and reasons flexibly in language, but its errors appear as hallucinated state changes that are hard to score with ordinary...

<https://arxiv.org/abs/2606.27806>

3 CONF

Don't Let Gains FADE: Breaking Down Policy Gradient Weights in RL

ArXiv CS.AI 50%

arXiv:2607.01490v1 Announce Type: cross Abstract: Reinforcement learning post-training dramatically improves LLM reasoning, but suffers from training instability and diversity collapse. Advantage functions offer an appealing fix: they reshape the training objective, reweight...

<https://arxiv.org/abs/2607.01490>

4 CONF

Reasoning LLM Improves Speaker Recognition in Long-form TV Dramas

ArXiv CS.AI 50%

arXiv:2607.02504v1 Announce Type: cross Abstract: Long-form TV dramas present a formidable challenge for comprehensive video understanding, where deciphering complex storyline often relies on \textbf{speaker recognition}, the task of accurately attributing each spoken utterance...

<https://arxiv.org/abs/2607.02504>

5 CONF

X-LogSMask: Expand Transformer for Graph-Structured Data

ArXiv CS.AI 50%

arXiv:2607.01553v1 Announce Type: cross Abstract: Transformers have become general-purpose architectures, but their all-to-all self-attention is poorly matched to graph data, whose interactions are sparse, structured and multi-scale. Existing Graph Transformers address this...

<https://arxiv.org/abs/2607.01553>

6 CONF

ADVENT: LLM-Driven Automatic Predicate Invention for ILP

ArXiv CS.AI 50%

arXiv:2607.01585v1 Announce Type: cross Abstract: Predicate invention (PI), the creation of new predicates to extend the hypothesis space, remains a critical bottleneck in Inductive Logic Programming (ILP). Existing methods rely on domain expertise and produce semantically...

<https://arxiv.org/abs/2607.01585>

7 CONF

Beyond Gradient-Based Attacks: Adversarial Robustness and Explainability Stability in Cybersecurity Classifiers

ArXiv CS.AI 50%

arXiv:2607.01679v1 Announce Type: cross Abstract: Adversarial attacks on cybersecurity classifiers pose a dual threat: degrading predictions and destabilising the SHAP-based explanations that security analysts rely on to understand and triage alerts. We extend our prior MLP...

<https://arxiv.org/abs/2607.01679>

8 CONF

MedStreamBench: A Time-Aware Benchmark for Streaming and Proactive Medical Video Understanding

ArXiv CS.AI 50%

arXiv:2607.01751v1 Announce Type: cross Abstract: Existing medical video benchmarks primarily evaluate whether a model produces the correct answer, but rarely assess whether it answers at the right time. In real clinical settings, AI systems must decide not only what to predict,...

<https://arxiv.org/abs/2607.01751>

9 CONF

ProCal: Inference-Time Proposal Calibration for Open-Vocabulary Object Detection

ArXiv CS.AI 50%

arXiv:2607.01759v1 Announce Type: cross Abstract: Open-vocabulary object detection aims to localize and classify objects beyond the fixed set of categories seen during training. Recent open-vocabulary object detection methods improve localization and classification for unseen...

<https://arxiv.org/abs/2607.01759>

00 CONF

A Dual-Helix Governance Approach Towards Reliable Agentic Artificial Intelligence for WebGIS Development

ArXiv CS.AI 50%

arXiv:2603.04390v2 Announce Type: replace Abstract: WebGIS development requires consistency, yet agentic AI often fails due to LLM context constraints, forgetting, stochasticity, instruction failure, and adaptation rigidity. We propose a dual-helix governance framework reframing...

<https://arxiv.org/abs/2603.04390>

01 CONF

PreScience: A Dataset and Benchmark for Scientific Forecasting

ArXiv CS.AI 50%

arXiv:2602.20459v2 Announce Type: replace Abstract: Can AI systems trained on the existing scientific record forecast the advances that will follow? We introduce PreScience, a dataset and benchmark for scientific forecasting built around 98K recent AI research papers, together...

<https://arxiv.org/abs/2602.20459>

02 CONF

Trust Region Inverse Reinforcement Learning: Explicit Dual Ascent using Local Policy Updates

ArXiv CS.AI 50%

arXiv:2605.11020v2 Announce Type: replace-cross Abstract: Inverse reinforcement learning (IRL) is typically formulated as maximizing entropy subject to matching the distribution of expert trajectories. Classical (dual-ascent) IRL guarantees monotonic performance improvement but...

<https://arxiv.org/abs/2605.11020>

03 CONF

Single-Channel EEG-Based Cognitive Load Assessment in Online Learning: A Hybrid Deep Learning Approach

ArXiv CS.AI 50%

arXiv:2607.01795v1 Announce Type: cross Abstract: Monitoring cognitive load during online learning could help instructors identify content that learners find difficult, but remote settings remove the visual cues that support this judgement in a classroom. We study whether a...

<https://arxiv.org/abs/2607.01795>

04 CONF

Decoupling Code Complexity from Newcomer Participation: A Causal Study of AI Coding Agent Adoption in OSS

ArXiv CS.AI 50%

arXiv:2607.01810v1 Announce Type: cross Abstract: Open-source projects depend on a steady inflow of newcomers. A growing concern is that AI coding agents (tools such as Cursor and Claude Code that write code from natural-language instructions) will crowd them out, by absorbing...

<https://arxiv.org/abs/2607.01810>

05 CONF

Has This Checkpoint Been Abliterated? A Two-Signal Audit and Its Failure Map

ArXiv CS.AI 50%

arXiv:2607.01854v1 Announce Type: cross Abstract: Can a platform tell, before deployment, whether an open-weight checkpoint has had its refusal mechanism stripped? Runtime guards cannot: they score generations, not the artifact. We combine two cheap internal signals, a...

<https://arxiv.org/abs/2607.01854>

06 CONF

SAB-LVLM: Significance-Aware Binarization for Large Vision-Language Models

ArXiv CS.AI 50%

arXiv:2607.01876v1 Announce Type: cross Abstract: Large Vision-Language Models (LVLMs) have achieved remarkable progress in multimodal understanding, yet their enormous parameter scale and cross-modal computation incur substantial memory and latency overhead, severely limiting...

<https://arxiv.org/abs/2607.01876>

07 CONF

SABER: A Semantic-Aligned Brain Network Analysis Framework via Multi-scale Hypergraphs

ArXiv CS.AI 50%

arXiv:2607.01901v1 Announce Type: cross Abstract: Effective brain disease diagnosis requires the synergy of brain connectivity patterns and high-level semantic knowledge. Existing methods, however, largely treat semantics from large language models (LLMs) as auxiliary features...

<https://arxiv.org/abs/2607.01901>

08 CONF

Population-Based Multi-Objective Training of Discriminators for Semi-Supervised GANs

ArXiv CS.AI 50%

arXiv:2607.01907v1 Announce Type: cross Abstract: Semi-supervised generative adversarial networks (SSL-GANs) can exploit large unlabeled datasets while retaining a classifier in the discriminator, but their training is often unstable. This paper proposes a population-based...

<https://arxiv.org/abs/2607.01907>

09 CONF

AlriskEval-edu: New Dataset for Risk Assessment in AI-mediated K-12 Educational Explanations

ArXiv CS.AI 50%

arXiv:2607.01934v1 Announce Type: cross Abstract: This work introduces AlriskEval-edu-db2, a new dataset designed to train and evaluate auditors based on LLMs for an explainable pedagogical risk assessment in instructional content for grades K-12. The dataset comprises 1,639...

<https://arxiv.org/abs/2607.01934>

10 CONF

CausalSteward: An Agentic Divide-Conquer-Combine Copilot for Causal Discovery

ArXiv CS.AI 50%

arXiv:2607.01936v1 Announce Type: cross Abstract: Learning causal models from high-dimensional data is a significant challenge, particularly in real-world settings where violations of core assumptions lead to causal identifiability issues. Although massive amounts of prior...

<https://arxiv.org/abs/2607.01936>

11 CONF

Conditional Co-Ablation: Recovering Self-Repair Backups in Transformer Circuits

ArXiv CS.AI 50%

arXiv:2607.01940v1 Announce Type: cross Abstract: Mechanistic interpretability often relies on component-level interventions to discover how a model produces a behavior. This guides attribution, capability knockout, and model pruning downstream to operate by scoring each unit by...

<https://arxiv.org/abs/2607.01940>

12 CONF

Robust for the Wrong Reasons: The Representational Geometry of LLM Robustness to Science Skepticism

ArXiv CS.AI 50%

arXiv:2607.01951v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly consulted on contested scientific questions, raising the concern that they will sycophantically retreat from established consensus when a user signals doubt -- drifting toward a false...

<https://arxiv.org/abs/2607.01951>

13 CONF

NeoMap: Training-free Novel-View Synthesis from Single Images and Videos

ArXiv CS.AI 50%

arXiv:2607.01962v1 Announce Type: cross Abstract: We study the challenging problem of novel view video synthesis from single images or monocular videos. Existing methods, which operate under the assumption that pre-trained video models lack native novel view synthesis capability...

<https://arxiv.org/abs/2607.01962>

14 CONF

Object Aligner: A Configurable JSON Schema Similarity Score for Graphs, Applied to LLM Prompt Optimization

ArXiv CS.AI 50%

arXiv:2607.01972v1 Announce Type: cross Abstract: Large language models (LLMs) are often asked to produce JSON conforming to a fixed schema, powering information extraction, tool calling, agentic planning, and knowledge-graph construction. Measuring how closely an output matches...

<https://arxiv.org/abs/2607.01972>

15 CONF

Do Newer Lightweight CNNs Perform Better Under Resource Constraints? A Controlled Multigenerational Study of Architecture, Initialization, Training Budget, and Efficiency

ArXiv CS.AI 50%

arXiv:2607.01984v1 Announce Type: cross Abstract: Newer lightweight convolutional neural networks are often presented as improving predictive performance and deployment efficiency, but such claims require controlled evaluation. This study compares nine lightweight CNN model...

<https://arxiv.org/abs/2607.01984>

16 CONF

Towards Load-Aware Prefill Deflection for Disaggregated LLM Serving

ArXiv CS.AI 50%

arXiv:2607.02043v1 Announce Type: cross Abstract: Disaggregated LLM serving runs prefill and decode on separate GPU pools to keep the two phases from interfering. In practice, this creates a new asymmetry: under bursty, heavy-tailed workloads prefill nodes saturate while decode...

<https://arxiv.org/abs/2607.02043>

17 CONF

SPLIT: Cross-Lingual Empathy and Cultural Grounding in English and Ukrainian LLM Responses

ArXiv CS.AI 50%

arXiv:2607.02049v1 Announce Type: cross Abstract: Large Language Models are increasingly deployed in emotional-support contexts and crisis-related situations. Nevertheless, their cross-lingual abilities in these circumstances remain underexplored. Existing benchmarks emphasize...

<https://arxiv.org/abs/2607.02049>

18 CONF

Beyond the Performance Illusion: Structure-Aware Stratified Partitioning and Curriculum Distributionally Robust Optimization for Spatially Correlated Domains

ArXiv CS.AI 50%

arXiv:2607.02055v1 Announce Type: cross Abstract: Performance evaluation in AI systems commonly assumes that random dataset splits produce independent and identically distributed (i.i.d.) subsets. We show that this assumption often breaks down in spatiotemporally correlated...

<https://arxiv.org/abs/2607.02055>

19 CONF

Prompt Coverage Adequacy

ArXiv CS.AI 50%

arXiv:2607.02057v1 Announce Type: cross Abstract: In recent years, it has become increasingly evident that large language models (LLMs) and autonomous agents raise the level of abstraction in software development by shifting the focus from writing precise procedures to...

<https://arxiv.org/abs/2607.02057>

20 CONF

SA-HGNN: Sample-Adaptive Hyperbolic Graph Neural Network for EEG-Based Depression Recognition

ArXiv CS.AI 50%

arXiv:2607.02063v1 Announce Type: cross Abstract: Graph Neural Networks (GNNs) have been widely used to capture spatial functional connectivity patterns to improve electroencephalography (EEG)-based depression recognition performance. However, the functional connectivity of...

<https://arxiv.org/abs/2607.02063>

21 CONF

kNNGuard: Turning LLM Hidden Activations into a Training-Free Configurable Guardrail

ArXiv CS.AI 50%

arXiv:2607.02072v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in domains requiring guardrails to detect unsafe, off-topic, or adversarial prompts. Existing guardrails predominately rely on fine-tuning to build classifiers, which often...

<https://arxiv.org/abs/2607.02072>

22 CONF

Guided Action Flow: Q-Guided Inference for Flow-Matching Vision-Language-Action Policies

ArXiv CS.AI 50%

arXiv:2607.02092v1 Announce Type: cross Abstract: Flow-matching vision-language-action policies generate robot action chunks through an iterative transport process, creating an opportunity for test-time guidance without retraining the base policy. We study this opportunity in...

<https://arxiv.org/abs/2607.02092>

23 CONF

Behind the Refusal: Determining Guardrail Activation via Behavioral Monitoring

ArXiv CS.AI 50%

arXiv:2607.02121v1 Announce Type: cross Abstract: As Large Language Models (LLMs) and agentic systems become integrated into real-world applications, ensuring their safety and security is critical. Guardrail systems that detect and block malicious instructions sent to and from...

<https://arxiv.org/abs/2607.02121>

24 CONF

ART for Diffusion Sampling: Continuous-Time Control and Actor-Critic Learning

ArXiv CS.AI 50%

arXiv:2607.02137v1 Announce Type: cross Abstract: We study timestep allocation for score-based diffusion sampling, where a learned reverse-time dynamics is discretized on a finite grid. Uniform and hand-crafted schedules are standard choices, but they rely on fixed prescriptions...

<https://arxiv.org/abs/2607.02137>

25 CONF

Predicting Early Stages Of Alzheimer's Disease And Identifying Key Biomarkers Using Deep Artificial Neural Network And Ensemble Of Machine Learning Methodologies

ArXiv CS.AI 50%

arXiv:2607.02142v1 Announce Type: cross Abstract: Alzheimers disease (AD) is a brain disorder that develops slowly and mainly affects memory, thinking, language, and daily activities. It is one of the most common causes of dementia and creates many difficulties for patients as...

<https://arxiv.org/abs/2607.02142>

26 CONF

Dynamic Neural Graph Encoding of Inference Processes in Deep Weight Space

ArXiv CS.AI 50%

arXiv:2607.02166v1 Announce Type: cross Abstract: The rapid advancements in using neural networks as implicit data representations have attracted significant interest in developing machine learning methods that analyze and process the weight spaces of other neural networks....

<https://arxiv.org/abs/2607.02166>

27 CONF

What Types of Human-AI Teams Exist?

ArXiv CS.AI 50%

arXiv:2607.02198v1 Announce Type: cross Abstract: Human-AI teaming has received increasing attention in the literature. However, the range of studies conducted in multiple domains make it difficult to understand what types of teams are being studied, and in what ways are they...

<https://arxiv.org/abs/2607.02198>

28 CONF

The Eticas AI Risk Taxonomy: Open Infrastructure for Operationalizing AI Audits

ArXiv CS.AI 50%

arXiv:2607.02201v1 Announce Type: cross Abstract: The rapid deployment of AI systems across high-stakes domains has created urgent demand for standardized evaluation, yet the field remains fragmented across competing risk taxonomies that catalog risks without showing how an...

<https://arxiv.org/abs/2607.02201>

29 CONF

Challenges and Recommendations for LLMs-as-a-Judge in Multilingual Settings and Low-Resource Languages

ArXiv CS.AI 50%

arXiv:2607.02235v1 Announce Type: cross Abstract: LLM-as-a-Judge has become the dominant evaluation paradigm for many natural language generation tasks, due to shortcomings of conventional metrics and high correlations with human judgment, albeit mostly in English. There are now...

<https://arxiv.org/abs/2607.02235>

30 CONF

The Dual Nature of LLM Persona: Aggregated Tendencies and Frame-Dependent Geometry

ArXiv CS.AI 50%

arXiv:2607.02368v1 Announce Type: cross Abstract: Evaluations of LLM personas via psychometric questionnaires typically rely on aggregate scores, discarding within-instance correlation structure. We test whether this geometric structure is intrinsic or frame-dependent....

<https://arxiv.org/abs/2607.02368>

31 CONF

World Wide Models: Literary Tools for Cultural AI

ArXiv CS.AI 50%

arXiv:2607.02369v1 Announce Type: cross Abstract: LLMs stage a new form of cultural encounter that is massive, automated, and monolingual. Literary disciplines have always negotiated cultural struggles with comparative reading of literature, narratological and poetic analysis,...

<https://arxiv.org/abs/2607.02369>

32 CONF

VisionAid: An Offline-First Multimodal Android Assistant for People with Visual Impairment, Featuring Personalized Object Retrieval

ArXiv CS.AI 50%

arXiv:2607.02371v1 Announce Type: cross Abstract: Over 285 million people worldwide live with a visual impairment, for whom everyday tasks such as avoiding obstacles, locating personal belongings, recognizing familiar faces, or handling cash remain persistent obstacles to...

<https://arxiv.org/abs/2607.02371>

33 CONF

OrbitQuant: Data-Agnostic Quantization for Image and Video Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2607.02461v1 Announce Type: cross Abstract: Diffusion transformers (DiTs) achieve state-of-the-art image and video generation, but their multi-step sampling and growing parameter count make inference expensive. Post-training quantization (PTQ) is the natural remedy, yet...

<https://arxiv.org/abs/2607.02461>

34 CONF

Human Capital, Not Model Benchmarks, Predicts Hybrid Intelligence in Forecasting

ArXiv CS.AI 50%

arXiv:2607.02467v1 Announce Type: cross Abstract: Whether pairing people with AI helps or hurts is usually reported as a single average effect. Using a real-money prediction market (Polymarket) as an objective, externally resolved benchmark, this pilot shows that the value of...

<https://arxiv.org/abs/2607.02467>

35 CONF

Beyond Adam: SOAP and Muon for Faster, Label-Efficient Training of Machine Learning Interatomic Potentials

ArXiv CS.AI 50%

arXiv:2607.02499v1 Announce Type: cross Abstract: Machine learning interatomic potentials (MLIPs) have become a hallmark of AI for scientific simulation. While efforts on new architectures and datasets have led to increasingly accurate and general models, the choice of optimizer...

<https://arxiv.org/abs/2607.02499>

36 CONF

ADMC: Attention-based Diffusion Model for Missing Modalities Feature Completion

ArXiv CS.AI 50%

arXiv:2507.05624v2 Announce Type: replace Abstract: Multimodal emotion and intent recognition is essential for automated human-computer interaction, It aims to analyze users' speech, text, and visual information to predict their emotions or intent. One of the significant...

<https://arxiv.org/abs/2507.05624>

37 CONF

Psychological Imagination Networks Show Cross-Population Centrality and Clustering Alignment in Humans That Large Language Models Fail to Replicate

ArXiv CS.AI 50%

arXiv:2510.04391v5 Announce Type: replace Abstract: Mental imagery vividness is a stable individual trait, yet whether imagined scenarios share relational structure across human and synthetic large language model (LLM) populations remains unknown. We applied psychological...

<https://arxiv.org/abs/2510.04391>

38 CONF

HAL: Inducing Human-likeness in LLMs with Alignment

ArXiv CS.AI 50%

arXiv:2601.02813v3 Announce Type: replace Abstract: Aligning language models to qualitative behavioral traits, such as human-likeness, remains difficult because they are hard to define, measure, and optimize. As a result, improvements in human-like behavior are largely driven by...

<https://arxiv.org/abs/2601.02813>

39 CONF

OmniGAIA: Towards Native Omni-Modal AI Agents

ArXiv CS.AI 50%

arXiv:2602.22897v3 Announce Type: replace Abstract: Human intelligence naturally intertwines omni-modal perception -- spanning vision, audio, and language -- with complex reasoning and tool usage to interact with the world. However, current multi-modal LLMs are primarily...

<https://arxiv.org/abs/2602.22897>

40 CONF

Resolving superposition in AI for interpretability and cross-modal alignment in patient-neuronal images

ArXiv CS.AI 50%

arXiv:2606.31394v2 Announce Type: replace-cross Abstract: Artificial intelligence is transforming our capability to solve biological challenges. In dimensionality bottleneck regimes exacerbated by high-dimensional biological data, neural networks force distinct concepts into the...

<https://arxiv.org/abs/2606.31394>

41 CONF

SkillDAG: Self-Evolving Typed Skill Graphs for LLM Skill Selection at Scale

ArXiv CS.AI 50%

arXiv:2606.03056v2 Announce Type: replace Abstract: As LLM agents adopt large skill libraries, selecting the right subset becomes a structural problem rather than a similarity-matching one: skills depend on, conflict with, specialize, or duplicate one another, a structure...

<https://arxiv.org/abs/2606.03056>

42 CONF

The Token Not Taken: Sampling, State, and the Stochasticity of AI Agents

ArXiv CS.AI 50%

arXiv:2606.08998v3 Announce Type: replace Abstract: Agentic AI systems can behave differently across runs: the same request may produce a different plan, a different tool call, a different code edit, or a different final answer. Such variability arises from several layers that...

<https://arxiv.org/abs/2606.08998>

43 CONF

Power Systems Agent Benchmark: Executable Evaluation of AI Agents in Electric Power Engineering

ArXiv CS.AI 50%

arXiv:2606.20950v2 Announce Type: replace Abstract: Executable evaluation -- checking the consequences of an agent's actions with a program rather than grading its prose -- has become a prominent way to assess tool-using AI agents in software settings. Electric power engineering...

<https://arxiv.org/abs/2606.20950>

44 CONF

Playing 20 Question Game with Policy-Based Reinforcement Learning

ArXiv CS.AI 50%

arXiv:1808.07645v5 Announce Type: replace-cross Abstract: The 20 Questions (Q20) game is a well known game which encourages deductive reasoning and creativity. In the game, the answerer first thinks of an object such as a famous person or a kind of animal. Then the questioner...

<https://arxiv.org/abs/1808.07645>

45 CONF

The Transformer as a Polar State Estimator

ArXiv CS.AI 50%

arXiv:2605.11007v2 Announce Type: replace-cross Abstract: We show that the core components of the Transformer -- attention, residual connections, and normalization -- arise naturally from a single geometric state estimation problem. Modeling the latent state in polar form, with...

<https://arxiv.org/abs/2605.11007>

46 CONF

Contrastive Deep Learning Reveals Age Biomarkers in Histopathological Skin Biopsies

ArXiv CS.AI 50%

arXiv:2411.16956v2 Announce Type: replace-cross Abstract: As global life expectancy increases, so does the burden of chronic diseases, yet individuals exhibit considerable variability in the rate at which they age. Identifying biomarkers that distinguish fast from slow ageing is...

<https://arxiv.org/abs/2411.16956>

47 CONF

Less Data, More Security: Advancing Cybersecurity LLMs Specialization via Resource-Efficient Domain-Adaptive Continuous Pre-training with Minimal Tokens

ArXiv CS.AI 50%

arXiv:2507.02964v2 Announce Type: replace-cross Abstract: The increasing scale of AI workloads demands High-Performance Computing (HPC) infrastructure and training methodologies that are both scalable and sustainable. While Large Language Models (LLMs) demonstrate exceptional...

<https://arxiv.org/abs/2507.02964>

48 CONF

RedCoder: Automated Multi-Turn Red Teaming for Code LLMs

ArXiv CS.AI 50%

arXiv:2507.22063v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) for code generation (i.e., Code LLMs) have demonstrated impressive capabilities in AI-assisted software development and testing. However, recent studies have shown that these models are prone...

<https://arxiv.org/abs/2507.22063>

49 CONF

Uncertain but Useful: Leveraging CNN Training Variability into Data Augmentation

ArXiv CS.AI 50%

arXiv:2509.05238v2 Announce Type: replace-cross Abstract: Deep learning (DL) has transformed neuroimaging by delivering state-of-the-art performance with reduced computation times. Yet, the numerical uncertainty inherent to DL training remains largely underexplored despite its...

<https://arxiv.org/abs/2509.05238>

50 **CONF****Psychological Steering in LLMs: An Evaluation of Effectiveness and Trustworthiness****ArXiv CS.AI** **50%**

arXiv:2510.04484v2 Announce Type: replace-cross Abstract: The ability to control LLMs' emulated emotional states and personality traits is an essential step in enabling rich, human-centered interactions in socially interactive settings. We introduce PsySET, a...

<https://arxiv.org/abs/2510.04484>

51 **CONF****Navigating the Alignment-Calibration Trade-off: A Pareto-Superior Frontier via Model Merging****ArXiv CS.AI** **50%**

arXiv:2510.17426v3 Announce Type: replace-cross Abstract: The "alignment tax" of post-training is typically framed as a drop in task accuracy. We show it also involves a severe loss of calibration, making models overconfident, less reliable, and model outputs less diverse. We...

<https://arxiv.org/abs/2510.17426>

52 **CONF****Towards Cellular-Scale Interpretability in Pathology Foundation Models for Biomarker Assessment****ArXiv CS.AI** **50%**

arXiv:2511.05150v2 Announce Type: replace-cross Abstract: Molecular biomarker testing in pathology is often costly and tissue-consuming, limiting scalable clinical deployment. Artificial intelligence applied to hematoxylin and eosin (HE)-stained histology could enable rapid...

<https://arxiv.org/abs/2511.05150>

53 **CONF****Gravity-Awareness: Deep Learning Models and LLM Simulation of Human Awareness in Altered Gravity****ArXiv CS.AI** **50%**

arXiv:2511.05536v2 Announce Type: replace-cross Abstract: Earth's gravity fundamentally shapes human behaviour. The brain encodes this force as an internal model of gravity, enabling the prediction and interpretation of gravitational effects during perception and action....

<https://arxiv.org/abs/2511.05536>

54 **CONF****Who Gets the Reward & Who Gets the Blame? Evaluation-Aligned Training Signals for Multi-LLM Agents****ArXiv CS.AI** **50%**

arXiv:2511.10687v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) in multi-agent systems (MAS) have shown promise for complex tasks, yet current training methods lack principled ways to connect system-level evaluation with agent- and message-level learning....

<https://arxiv.org/abs/2511.10687>

55 **CONF****From Experiments to Expertise: Scientific Knowledge Consolidation for AI-Driven Computational Physics****ArXiv CS.AI** **50%**

arXiv:2603.13191v2 Announce Type: replace-cross Abstract: While large language models (LLMs) have transformed AI agents into proficient executors of computational materials science, performing a hundred simulations does not make a researcher. What distinguishes research from...

<https://arxiv.org/abs/2603.13191>

56 CONF

CaP-X: A Framework for Benchmarking and Improving Coding Agents for Robot Manipulation

ArXiv CS.AI 50%

arXiv:2603.22435v2 Announce Type: replace-cross Abstract: "Code-as-Policy" considers how executable code can complement data-intensive Vision-Language-Action (VLA) methods, yet their effectiveness as autonomous controllers for embodied manipulation remains underexplored. We...

<https://arxiv.org/abs/2603.22435>

57 CONF

Dive into Claude Code: The Design Space of Today's and Future AI Agent Systems

ArXiv CS.AI 50%

arXiv:2604.14228v2 Announce Type: replace-cross Abstract: Claude Code is an agentic coding tool that can run shell commands, edit files, and call external services on behalf of the user. This study describes its architecture by analyzing the publicly available source code and...

<https://arxiv.org/abs/2604.14228>

58 CONF

ContraFix: Skill-Enhanced Contrastive Runtime Analysis for Vulnerability Repair

ArXiv CS.AI 50%

arXiv:2605.17450v2 Announce Type: replace-cross Abstract: As software systems grow increasingly complex, automated vulnerability repair (AVR) remains difficult because the materials available to a repair system are usually failure artifacts rather than repair guidance....

<https://arxiv.org/abs/2605.17450>

59 CONF

MemSyco-Bench: Benchmarking Sycophancy in Agent Memory

ArXiv CS.AI 50%

arXiv:2607.01071v2 Announce Type: replace-cross Abstract: Memory has emerged as a cornerstone of modern LLM-based agents, supporting their evolution from single-turn assistants to long-term collaborators. However, memory is not always beneficial: retrieved memories often induce...

<https://arxiv.org/abs/2607.01071>