



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	112	Media
HackerNews	11	Community
TechCrunch AI	9	Media
The Verge AI	3	Media
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Agentic AI for Commercial Insurance Underwriting with Adversarial Self-Critique

ArXiv CS.AI +2 sources

100%

arXiv:2602.13213v2 Announce Type: replace Abstract: Commercial insurance underwriting is a labor-intensive process that requires manual review of extensive documentation to assess risk and determine policy pricing. While AI offers substantial efficiency improvements, existing...

<https://arxiv.org/abs/2602.13213>

VERI

Anthropic is launching Claude Cowork on mobile and web

The Verge AI +1 source

80%

Starting Tuesday, Anthropic's Claude Cowork AI platform will be available on mobile and web for the first time. The expanded access is rolling out first to Max subscribers and coming to Claude users on other plans "in the coming weeks." Claude Cowork was previously only...

<https://www.theverge.com/ai-artificial-intelligence/961978/anthropic-claude-cowork-mobi...>

VERI

Beyond Reactivity: Measuring Proactive Problem Solving in LLM Agents

ArXiv CS.AI +1 source

80%

arXiv:2510.19771v4 Announce Type: replace Abstract: LLM-based agents are increasingly moving towards proactivity: rather than awaiting instruction, they exercise agency to anticipate user needs and solve them autonomously. However, evaluating proactivity is challenging; current...

<https://arxiv.org/abs/2510.19771>

CONFIRMED SIGNALS — 131 stories

CONF

Show HN: Davit, a Apple Containers UI

HackerNews 60% HN 265

Hi HN, we built Prizmi because we wanted an AI assistant that would i) do real and interesting things instead of explaining how to do them ourselves, ii) act on all our devices, iii) be proactive in true sense, iv) keep us in full control, v) respect our privacy, vi) be easy to...

<https://davit.app>

CONF

Show HN: Rowboat - Open-source, local-first alternative to Claude Desktop

HackerNews 60% HN 136

TL;DR: we built a framework-agnostic agent runtime that uses gVisor for isolation and runs on k8s. It's open-source under AGPLv3. Recently we've been working on a customer support "AI assistant" - essentially an interactive knowledge base/L1 support but with an option to touch...

<https://github.com/rowboatlabs/rowboat>

CONF

Show HN: Free Mermaid Diagram Editor

HackerNews 50% HN 27

Hi HN, A few years ago I created a popular open source repo for animating children's drawings. Recently, I decided to try turning it into a polished user experience. It doesn't use ANY generative AI (very proud of that). Instead, it uses traditional character rigging and motion...

<https://moxiedocs.com/mermaid-diagram-editor>

CONF

Show HN: Baerly-storage, a document DB that runs per request, no DB server

HackerNews 50% HN 15

No summary — see link for full article.

<https://github.com/Gusto/baerly-storage>

CONF

Show HN: Context Warp Drive - Deterministic context folding for AI agents

HackerNews 50% HN 8

No summary — see link for full article.

<https://github.com/dogtorjonah/context-warp-drive>

CONF

Show HN: I wrote a 1-bit WebGPU runtime to run a 1.7B LLM in the browser

HackerNews 50% HN 5

No summary — see link for full article.

<https://aidekin.com/>

0 CONF

Show HN: Backlog - tasks and contexts manager for AI coding agents

HackerNews 50% HN 4

I thought humanizing AI writing was straightforward. Just delete the em dashes, adjust the tone, and remove a few obvious words. However, as I "delved" into research on AI versus human writing, I discovered there was way more to it! Factors like sentence rhythm, repetition,...

<https://github.com/mazen160/backlog>

1 CONF

Show HN: FlowHook - a file watcher that knows if your build passed or failed

HackerNews 50% HN 3

FlowHook is a CLI tool that watches your files and executes given commands, similar to watchexec. But I believe it has three things that differentiate it from existing tools. It persists your configuration. If you set up flowhook instance in your folder and add the necessary...

<https://github.com/etcoder-642/FlowHook>

2 CONF

Show HN: Larasend - Self-hosted transactional email for Laravel on your own SES

HackerNews 50% HN 3

No summary — see link for full article.

<https://larasend.com>

3 CONF

Show HN: Tracking GenAI cost and endpoint fragility so app teams don't have to

HackerNews 50% HN 2

No summary — see link for full article.

<https://llmintel.ai/demo>

4 CONF

Show HN: Signal for LLM - The Modulator Architecture (Theory Complete)

HackerNews 50% HN 2

No summary — see link for full article.

<https://divinecanon.github.io/signalengine-EN/>

5 **CONF****Meta just launched a new AI generator, Muse Image, and users are already pushing back over use of their photos****TechCrunch AI** 50%

The new image-generating model has numerous use cases, including advertising and decorating, and creator-based opportunities.
<https://techcrunch.com/2026/07/07/meta-rolls-out-muse-a-new-ai-image-generator/>

6 **CONF****Why the rise of open source AI isn't hurting Anthropic ... yet****TechCrunch AI** 50%

Open source models' success isn't coming at the expense of frontier labs. Instead, they each seem to capture two phases of the same life cycle.
<https://techcrunch.com/2026/07/07/why-the-rise-of-open-source-ai-isnt-hurting-anthropic...>

7 **CONF****Microsoft joins AI cost-cutting trend by relying more on its own models****TechCrunch AI** 50%

Microsoft is the latest Silicon Valley giant to cut back on its AI spending.
<https://techcrunch.com/2026/07/07/microsoft-joins-ai-cost-cutting-trend-by-relying-more...>

8 **CONF****Discord admits AI moderation bug wrongfully banned users over harmless images****TechCrunch AI** 50%

The company confirmed that the issue had been affecting accounts since May, with an additional 200 users banned over the weekend before its team identified and fixed the problem.
<https://techcrunch.com/2026/07/07/discord-admits-ai-moderation-bug-wrongfully-banned-us...>

9 **CONF****Savi's app aims to protect consumers from realistic AI scams like kidnappers demanding ransom****TechCrunch AI** 50%

The company just raised \$7 million in seed funding, and is launching its app for iPhone and Android on Tuesday.
<https://techcrunch.com/2026/07/07/savis-app-aims-to-protect-consumers-from-realistic-ai...>

0 **CONF****The first American autonomous ground vehicles are fighting in Ukraine****TechCrunch AI** 50%

Forterra has deployed more than 100 of its self-driving ATVs in conflict zones in Ukraine.
<https://techcrunch.com/2026/07/07/the-first-american-autonomous-ground-vehicles-are-fig...>

1 **CONF****The Download: your stake in OpenAI, and the Treasury's AI warning****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Your family's \$300 stake in OpenAI Sam Altman's proposal that Americans should share in the wealth created by AI is back in the spotlight,...
<https://www.technologyreview.com/2026/07/07/1140197/the-download-your-openai-stake-trea...>

2 **CONF****The foundational elements of AI architecture that IT leaders need to scale****MIT Tech Review** 50%

With the rapid progress of AI capabilities and the move to agentic systems, organizations are expanding their use cases as the technology continues to grow. That constant evolution also introduces risk, leaving IT leaders to wonder which investments will prove valuable even six...
<https://www.technologyreview.com/2026/07/07/1139413/the-foundational-elements-of-ai-arc...>

3 **CONF****Achieving operational excellence with AI****MIT Tech Review** 50%

Frameworks like Lean Six Sigma and business process management (BPM) first gained traction because they promised clarity in the chaos—a structured way to bring order to messy, sprawling operations. Lean Six Sigma emphasized statistical rigor and quality control; BPM created...

<https://www.technologyreview.com/2026/07/02/1140045/achieving-operational-excellence-wi...>
4 **CONF****Meta's new Muse Image model can pull other Instagram users into AI photos****The Verge AI** 50%

Meta is launching the first AI image generation model made by its Superintelligence Labs division. The Muse Image model now powers the image-making tools across the Meta AI app, Instagram, and WhatsApp, and it's coming soon to Facebook and Messenger, according to an announcement...

<https://www.theverge.com/tech/962485/meta-muse-image-ai-model-instagram>
5 **CONF****Prompt-to-Paper: Agentic AI System for Bioinformatics****ArXiv CS.AI** 50%

arXiv:2607.05456v1 Announce Type: new Abstract: While recent advances in large language models have enabled end-to-end automated manuscript generation, existing systems suffer from three critical deficiencies: (i) generated claims are not deterministically grounded in verifiable...

<https://arxiv.org/abs/2607.05456>
6 **CONF****CSTutorBench: Benchmarking Small Language Models as Tutors for Block-Based Programming****ArXiv CS.AI** 50%

arXiv:2607.05571v1 Announce Type: new Abstract: Large language models are increasingly explored as AI tutors, yet deploying them in K-12 settings raises concerns around privacy, cost, and reliance on proprietary models. Small language models (SLMs) offer a promising alternative,...

<https://arxiv.org/abs/2607.05571>
7 **CONF****Foundation Models for Automatic CAD Generation****ArXiv CS.AI** 50%

arXiv:2607.05573v1 Announce Type: new Abstract: Recent advances in Large Language Models (LLMs) and Vision-Language Models (VLMs) enable the automatic generation of parametric 3D designs from natural-language specifications. This chapter presents an empirical study of foundation...

<https://arxiv.org/abs/2607.05573>
8 **CONF****Akashic: A Low-Overhead LLM Inference Service with MemAttention****ArXiv CS.AI** 50%

arXiv:2607.05708v1 Announce Type: new Abstract: Recent LLM-based agent systems continuously accumulate context across multi-turn interactions, tool invocations, and cross-session workflows. Replaying the full history for every request quickly becomes impractical: long contexts...

<https://arxiv.org/abs/2607.05708>
9 **CONF****Decision Protocols in Multi-Agent Large Language Model Conversations****ArXiv CS.AI** 50%

arXiv:2607.05477v1 Announce Type: cross Abstract: Improving the task performance of Large Language Models (LLMs) is essential, yet scaling these models faces significant challenges such as diminishing returns and high costs. Multi-Agent Systems (MAS) offer a promising solution...

<https://arxiv.org/abs/2607.05477>

0 CONF

Learning to Control LLM Agent Harnesses with Offline Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.05458v1 Announce Type: cross Abstract: Large language model (LLM) agents are usually improved by changing prompts, models, or hand-written workflows, while the execution harness around the model is treated as fixed infrastructure. We argue that this harness is itself...

<https://arxiv.org/abs/2607.05458>

1 CONF

TurnOPD: Making On-Policy Distillation Turn-Aware for Efficient Long-Horizon Agent Training

ArXiv CS.AI 50%

arXiv:2607.05804v1 Announce Type: new Abstract: On-policy distillation (OPD) trains a student policy by matching a stronger teacher on the student's own trajectories, offering a promising framework for language agent training. However, its application to long-horizon agentic...

<https://arxiv.org/abs/2607.05804>

2 CONF

Onnes: A Physics-Grounded Multi-Agent LLM Simulator for Cryogenic Fault Diagnosis in Quantum Computing Infrastructure

ArXiv CS.AI 50%

arXiv:2607.05805v1 Announce Type: new Abstract: Dilution refrigerators are the enabling infrastructure of superconducting quantum computers, yet their fault diagnosis is still dominated by threshold alarms that report that something is wrong, not what. We present Onnes, a...

<https://arxiv.org/abs/2607.05805>

3 CONF

PCBWorld: A Benchmark Environment for Engine-Grounded PCB Design Automation

ArXiv CS.AI 50%

arXiv:2607.05915v1 Announce Type: new Abstract: PCB routing is the task of connecting the nets of a board with copper traces under strict design rules, yet learning-based methods still lag behind rule-based routers. We introduce PCBWorld, an open-source engine-grounded PCB...

<https://arxiv.org/abs/2607.05915>

4 CONF

SearchEyes: Towards Frontier Multimodal Deep Search Intelligence via Search World Simulation

ArXiv CS.AI 50%

arXiv:2607.05943v1 Announce Type: new Abstract: Training multimodal search agents to perform multi-hop reasoning remains challenging due to a fundamental structural disconnect: existing pipelines construct training data, search environments, and reward signals independently,...

<https://arxiv.org/abs/2607.05943>

5 CONF

Integrating knowledge graphs and multilingual scholarly corpora for domain-adaptive LLMs in SSH

ArXiv CS.AI 50%

arXiv:2607.05956v1 Announce Type: new Abstract: The integration of Large Language Models (LLMs) into scientific research workflows, particularly for bibliographic discovery and literature synthesis, raises significant methodological, epistemic and regulatory challenges for the...

<https://arxiv.org/abs/2607.05956>

- 6
CONF
Information Limits and Attractor Dynamics in Economies of Frontier LLM Agents: A Pre-Registered Test
ArXiv CS.AI 50%

arXiv:2607.06001v1 Announce Type: new Abstract: We report a pre-registered, two-part experiment on small economies of frontier language-model agents (Claude Opus 4.8), testing two quantitative predictions about coupled multi-agent systems: an information-theoretic capacity...

<https://arxiv.org/abs/2607.06001>
- 7
CONF
A toy framework for single and multi-agent human-AI curiosity ecosystems
ArXiv CS.AI 50%

arXiv:2607.06214v1 Announce Type: new Abstract: This paper offers a toy framework for considering curiosity as an ecosystem. First, it suggests that a single agent's inquiry policy (how, when, and why an agent asks a question) depends on how the agent values immediate...

<https://arxiv.org/abs/2607.06214>
- 8
CONF
Information Gain-based Rollout Policy Optimization: An Adaptive Tree-Structured Rollout Approach for Multi-Turn LLM Agents
ArXiv CS.AI 50%

arXiv:2607.06223v1 Announce Type: new Abstract: Reinforcement learning has become a promising paradigm for improving large language model (LLM) agents on long-horizon search tasks, where the agent must make a sequence of intermediate decisions before receiving a final outcome....

<https://arxiv.org/abs/2607.06223>
- 9
CONF
From Application-Layer Simulation to Native Meta-Architecture: Structural Tension as an Endogenous Driver for Heterogeneous AI Evolution
ArXiv CS.AI 50%

arXiv:2607.06269v1 Announce Type: new Abstract: Current large language models (LLMs) are fundamentally stateless: their behavior is fully determined by input at inference time, and any higher-order cognitive architecture must be simulated at the application layer through prompt...

<https://arxiv.org/abs/2607.06269>
- 0
CONF
Harnessing Generative Image Models for Training-Free Primitive Shape Abstraction
ArXiv CS.AI 50%

arXiv:2607.05568v1 Announce Type: cross Abstract: Representing 3D shapes as compact sets of geometric primitives is fundamental to robotics, simulation, and scene understanding. Generative image models trained at scale have recently emerged as generalist visual learners that can...

<https://arxiv.org/abs/2607.05568>
- 1
CONF
Driving the Wrong Way: Leveraging Interpretability in End2End Autonomous Driving Models
ArXiv CS.AI 50%

arXiv:2607.06328v1 Announce Type: new Abstract: The increasing adoption of end-to-end learning for autonomous driving introduces increased model complexity and opacity, raising the risk of learning undesired or erroneous behavior. In this work, we integrate unsupervised...

<https://arxiv.org/abs/2607.06328>

2 CONF

RMISC: A Large-scale Real-world Multivariate Corpus for Time Series Foundation Models

ArXiv CS.AI 50%

arXiv:2607.06504v1 Announce Type: new Abstract: Recent years have witnessed the emergence of multivariate modeling using time series foundation models (TSFMs), which achieve advanced zero-shot generalization. Modern multivariate TSFMs are predominantly pretrained on multivariate...

<https://arxiv.org/abs/2607.06504>

3 CONF

Detoxify: A framework for abusive text transformation using LLMs

ArXiv CS.AI 50%

arXiv:2507.10177v2 Announce Type: replace-cross Abstract: Although Large Language Models (LLMs) have demonstrated significant advancements in natural language processing tasks, their effectiveness in the classification and transformation of abusive text into non-abusive versions...

<https://arxiv.org/abs/2507.10177>

4 CONF

Hierarchical Multi-Agent Reinforcement Learning for Carbon-Aware AI Data Centers in Power Distribution Systems

ArXiv CS.AI 50%

arXiv:2607.03324v2 Announce Type: replace-cross Abstract: Eco-friendly energy management for artificial intelligence data centers (AIDCs) is crucial because of the significant increase in energy consumption-induced carbon emissions from AIDCs resulting from the rapid expansion...

<https://arxiv.org/abs/2607.03324>

5 CONF

Doomed from the Start: Early Abort of LLM Agent Episodes via a Recall-Controlled Probe Cascade

ArXiv CS.AI 50%

arXiv:2607.06503v1 Announce Type: new Abstract: Large language model (LLM) agents solving multi-step tasks frequently commit to trajectories that are doomed to fail, yet continue to consume substantial inference compute before the failure becomes observable. We show that failure...

<https://arxiv.org/abs/2607.06503>

6 CONF

FootsiesGym: A Fighting Game Benchmark for Two-Player Zero-Sum Imperfect-Information Games

ArXiv CS.AI 50%

arXiv:2607.06514v1 Announce Type: new Abstract: We present FootsiesGym, an open-source environment for learning in a non-trivial two-player, zero-sum, imperfect-information game. Built on HiFight's minimalist 2D fighting game Footsies, it isolates the cyclic, non-transitive...

<https://arxiv.org/abs/2607.06514>

7 CONF

FreqDepthKV: Frequency-Guided Depth Sharing for Robust KV Cache Compression in Long-Context LLM Inference

ArXiv CS.AI 50%

arXiv:2607.06519v1 Announce Type: new Abstract: Long-context LLM inference is increasingly limited by the memory and bandwidth cost of KV caches, yet aggressive compression can remove the layer-specific evidence needed for retrieval and multi-step reasoning. We introduce...

<https://arxiv.org/abs/2607.06519>

8 CONF

Bridging Physical Reasoning and Task Generalization via Visual Action Outcome Reasoning Alignment

ArXiv CS.AI 50%

arXiv:2607.06522v1 Announce Type: new Abstract: Vision-language models (VLMs) struggle to generalize in interactive physical reasoning, particularly under unseen tasks and environments. Two key failure modes are prominent: hallucinated chain-of-thought (CoT) reasoning that...

<https://arxiv.org/abs/2607.06522>

9 CONF

Rethinking Indic AI from a Lens of Cultural Heritage Preservation

ArXiv CS.AI 50%

arXiv:2607.06544v1 Announce Type: new Abstract: As Artificial Intelligence (AI) makes inroads into different parts of the Indian subcontinent, there is significant interest in studying how AI impacts the linguistic and cultural foundations of this civilization. AI is seen as a...

<https://arxiv.org/abs/2607.06544>

0 CONF

Proof of Execution: Runtime Verification for Governed AI Agent Actions

ArXiv CS.AI 50%

arXiv:2607.05397v1 Announce Type: cross Abstract: Agent systems increasingly execute rather than advise. When an AI agent queries regulated data, invokes effectful tools, and mutates persistent state, correctness is not captured by whether a terminal output looks plausible. The...

<https://arxiv.org/abs/2607.05397>

1 CONF

Benchmarking KV-Cache Optimizations across Task Quality and System Performance for Long-Context Serving

ArXiv CS.AI 50%

arXiv:2607.05399v1 Announce Type: cross Abstract: Large language model serving is increasingly limited by KV-cache growth under long-context workloads, yet existing KV-cache compression techniques are difficult to compare because they were evaluated on different models, tasks,...

<https://arxiv.org/abs/2607.05399>

2 CONF

Catalyst Papers in Artificial Intelligence Research: A Landscape on ICLR from 2017 to 2025

ArXiv CS.AI 50%

arXiv:2607.05401v1 Announce Type: cross Abstract: A small number of methodological contributions, including word2vec, the Transformer, large-scale pre-training, and reinforcement learning from human feedback, have reshaped NLP and AI research over the past decade. OpenReview now...

<https://arxiv.org/abs/2607.05401>

3 CONF

AI tools in Arab University English classrooms: Looking back and forward

ArXiv CS.AI 50%

arXiv:2607.05403v1 Announce Type: cross Abstract: This paper aims to synthesize empirical research on AI tools used to support English as a second/foreign language (EL2) learners in Arab University classrooms (AUCs) between Jan 1st 2023 and Aug 31st 2025. We utilized 3 large...

<https://arxiv.org/abs/2607.05403>

4 **CONF****The Jagged Global Economy: Frontier AI Unevenly Exposes National Economies**

ArXiv CS.AI 50%

arXiv:2607.05404v1 Announce Type: cross Abstract: Frontier AI's labor-market effects matter to workers, firms, and policymakers, but current evidence generally comes from a handful of high-income economies. The capabilities of frontier AI are jagged across work tasks and...

<https://arxiv.org/abs/2607.05404>

5 **CONF****CCBENCH: Assessing LLM Cultural Competence via Implicitly Signaled Norms using Health Queries**

ArXiv CS.AI 50%

arXiv:2607.05405v1 Announce Type: cross Abstract: To interact with users fairly and without stereotyping, AI models must display cultural competency, i.e., the ability to infer and adapt to a user's implicitly signaled cultural values, rather than relying on static demographic...

<https://arxiv.org/abs/2607.05405>

6 **CONF****A Guiding Framework for K-12 Teachers in Creating AI-powered Learning Technologies through Vibe Coding**

ArXiv CS.AI 50%

arXiv:2607.05406v1 Announce Type: cross Abstract: Large language models generate code from natural language prompts, enabling "vibe coding," which allows non-programmers to develop computational solutions. Vibe coding for teachers amplifies the value of teachers-as-designers,...

<https://arxiv.org/abs/2607.05406>

7 **CONF****Position: Preventing AI-Generated CSAM Necessitates New Approaches to AI Safety**

ArXiv CS.AI 50%

arXiv:2607.05407v1 Announce Type: cross Abstract: Modern artificial intelligence (AI) systems present profound new risks to child safety. AI is increasingly being misused to create AI-generated child sexual abuse material, facilitate child sexual exploitation, and reduce...

<https://arxiv.org/abs/2607.05407>

8 **CONF****The GenAI Skill Bypass: Mapping Divergent Pathways of University Students and Staff AI Literacy**

ArXiv CS.AI 50%

arXiv:2607.05411v1 Announce Type: cross Abstract: Higher education institutions are increasingly expected to ensure that both students and staff develop Generative AI (GenAI) literacies. In response, they are introducing professional development programs and embedding GenAI...

<https://arxiv.org/abs/2607.05411>

9 **CONF****Why does AI unlock new possibilities in STEM education? A Bibliometric Analysis of Trends and Future Agenda**

ArXiv CS.AI 50%

arXiv:2607.05412v1 Announce Type: cross Abstract: STEM education faces challenges in personalization and interdisciplinary integration. AI technology has brought new possibilities, but the mechanisms by which AI reshapes the STEM education ecosystem require systematic...

<https://arxiv.org/abs/2607.05412>

0 **CONF****When AI Classifies: What Counts as Public Administration?****ArXiv CS.AI** **50%**

arXiv:2607.05420v1 Announce Type: cross Abstract: This study examines how alternative systems of scholarly representation identify and characterize broad public administration (PA) and artificial intelligence related public administration (AI-in-PA) scholarship. Using Web of...

<https://arxiv.org/abs/2607.05420>

1 **CONF****Modality Relevance is not Modality Utility: Post-hoc Selective Modality Escalation for Cost-Aware Multimodal RAG****ArXiv CS.AI** **50%**

arXiv:2607.05438v1 Announce Type: cross Abstract: Multimodal retrieval-augmented generation (RAG) grounds a generator in evidence drawn from heterogeneous modalities -- text, tables, and images. The dominant deployment choice is binary and made before the model has tried to...

<https://arxiv.org/abs/2607.05438>

2 **CONF****Scientific Code Search at Scale: A Multi-Domain Dataset and Benchmark****ArXiv CS.AI** **50%**

arXiv:2607.05443v1 Announce Type: cross Abstract: Scientists increasingly rely on open-source tools to support their research workflows, yet discovering relevant software among over 600 million GitHub repositories remains challenging. Existing code search benchmarks focus on...

<https://arxiv.org/abs/2607.05443>

3 **CONF****Empirical Minimal-Realisation Compression of Deep Neural Networks via Controllability-Observability Tests****ArXiv CS.AI** **50%**

arXiv:2607.05457v1 Announce Type: cross Abstract: Deep neural networks often contain substantial hidden-state redundancy, but most compression methods operate directly on weights, neurons, or quantised representations without explicitly characterising the dynamical role of...

<https://arxiv.org/abs/2607.05457>

4 **CONF****Learning 4D Geometric Priors for Inference-Efficient World Action Models****ArXiv CS.AI** **50%**

arXiv:2607.05468v1 Announce Type: cross Abstract: World Action Models (WAMs) have shown strong potential for robotic manipulation by jointly modeling visual future dynamics and executable action sequences. However, existing video-action co-training methods primarily optimize...

<https://arxiv.org/abs/2607.05468>

5 **CONF****Is Your NPU Ready for LLMs? Dissecting the Hidden Efficiency Bottlenecks in Mobile LLM Inference****ArXiv CS.AI** **50%**

arXiv:2607.05475v1 Announce Type: cross Abstract: Deploying Large Language Models (LLMs) on mobile devices enhances privacy and reduces latency, but is severely bottlenecked by hardware inefficiency. We present the first comprehensive, cross-layer measurement study of mobile LLM...

<https://arxiv.org/abs/2607.05475>

6 CONF

Privilege and confidentiality in generative AI workflows

ArXiv CS.AI 50%

arXiv:2607.05479v1 Announce Type: cross Abstract: Generative AI (GenAI) systems store and process client data in three distinct ways: in the model's parameters through training and memorisation, in the context window during a live session, and in knowledge databases for...

<https://arxiv.org/abs/2607.05479>

7 CONF

PatchOptic for Shared-State LLM Workflows with Projected Views and Verified Structured Updates

ArXiv CS.AI 50%

arXiv:2607.05483v1 Announce Type: cross Abstract: Agentic workflows often operate over shared, structured state. Because LLM context windows are limited, each model invocation is typically shown only the state fragment needed for the current workflow step, a pattern commonly...

<https://arxiv.org/abs/2607.05483>

8 CONF

Lean-Quantum: Toward AI-Assisted Formalization of Quantum Information

ArXiv CS.AI 50%

arXiv:2607.05492v1 Announce Type: cross Abstract: Quantum information theory is built on entropic quantities; among them, the sandwiched Rényi relative entropy is a fundamental divergence with various applications, and its data processing inequality (DPI) under quantum...

<https://arxiv.org/abs/2607.05492>

9 CONF

Rendering-Aware Bayesian 3D Gaussian Splatting with Native Uncertainty and Adaptive Complexity Control

ArXiv CS.AI 50%

arXiv:2607.05522v1 Announce Type: cross Abstract: 3D Gaussian splatting (3DGS) is a strong representation for real-time novel-view synthesis, but its standard training pipeline relies on point estimates and hand-tuned heuristics, providing no native uncertainty or principled...

<https://arxiv.org/abs/2607.05522>

10 CONF

Most LLM Conformity Needs No Speaker: Measuring the Speaker-Free Floor in Peer-Pressure Benchmarks

ArXiv CS.AI 50%

arXiv:2607.05545v1 Announce Type: cross Abstract: LLM conformity is often used to describe cases where a model changes a correct answer toward a peer or group response. We show that most of this apparent conformity survives even after the peer is removed. The reason is a...

<https://arxiv.org/abs/2607.05545>

11 CONF

The yes-no bias of large language models reflects answer order and wording, not shifts in moral judgment

ArXiv CS.AI 50%

arXiv:2607.05552v1 Announce Type: cross Abstract: Large language models (LLMs) increasingly issue judgments read as binary verdicts, and a growing literature reports such judgments shifting under logically irrelevant changes of wording - among them an amplified yes-no bias on...

<https://arxiv.org/abs/2607.05552>

- 2
CONF

Prompt Robustness Is Task-Dependent: Comparing Objective and Belief-Style Questions in LLM Evaluation

ArXiv CS.AI 50%

arXiv:2607.05554v1 Announce Type: cross Abstract: Survey-style evaluations of large language models often treat a prompted response as a measure of a model's values or beliefs. This assumption is particularly fragile when responses are read as evidence of political values,...

<https://arxiv.org/abs/2607.05554>
- 3
CONF

Whose fairness? Structural concentration in AI bias research

ArXiv CS.AI 50%

arXiv:2607.05574v1 Announce Type: cross Abstract: Artificial intelligence increasingly mediates consequential decisions in healthcare, law, and public services, and the field has responded with an extensive methodology for measuring and mitigating bias. Yet the fairness...

<https://arxiv.org/abs/2607.05574>
- 4
CONF

ResonatorLM: Causal Resonant Field Mixing for Efficient Long-Context Language Modelin

ArXiv CS.AI 50%

arXiv:2607.05583v1 Announce Type: cross Abstract: Contemporary language models are dominated by the transformer architecture, which leverages self-attention mechanisms to enable more efficient, parallelized training across a wide set of documents and corpora. This has allowed...

<https://arxiv.org/abs/2607.05583>
- 5
CONF

To Retain or to Adapt? Generalizing Continual Learning

ArXiv CS.AI 50%

arXiv:2607.05609v1 Announce Type: cross Abstract: The Continual Learning (CL) literature has long been driven by the goal of mitigating catastrophic forgetting. This objective rests on a pervasive, often unstated assumption: that a lifelong learner should approximate the...

<https://arxiv.org/abs/2607.05609>
- 6
CONF

BaFCo: A Document Understanding Benchmark for Complex Bangla Form Comprehension

ArXiv CS.AI 50%

arXiv:2607.05614v1 Announce Type: cross Abstract: Document comprehension is a challenging yet impactful task for Multimodal Large Language Models, especially as these systems see growing adoption in real-world, human-centric applications. However, this adoption is limited for...

<https://arxiv.org/abs/2607.05614>
- 7
CONF

EvalLoop: A Methodology for Evaluation-Driven Iterative Improvement of Business AI Systems

ArXiv CS.AI 50%

arXiv:2607.05638v1 Announce Type: cross Abstract: Teams deploying large language models in business contexts need evaluation systems, yet most treat evaluation as static model selection: run benchmarks, rank models, deploy the winner. This framing misses evaluation's primary...

<https://arxiv.org/abs/2607.05638>

8 **CONF****Do It Right! A Methodology for Successful NLP System Development****ArXiv CS.AI** **50%**

arXiv:2607.05644v1 Announce Type: cross Abstract: Natural language processing (NLP) is a common method for supplying data to clinical research and decision making by extracting information from electronic medical records. Numerous textbooks and tutorials describe specific...

<https://arxiv.org/abs/2607.05644>

9 **CONF****Physics-Regularized Machine Learning for Proprioceptive Vehicle Localization Using Onboard Sensors****ArXiv CS.AI** **50%**

arXiv:2607.05663v1 Announce Type: cross Abstract: Accurate and robust localization is essential for autonomous mobility systems in real-world environments. While fusing Inertial Measurement Unit (IMU) data with satellite-based correction signals provides precise vehicle pose...

<https://arxiv.org/abs/2607.05663>

0 **CONF****What Do AI Agents Actually Change? An Empirical Taxonomy of Mutation Patterns in Performance-Improving Pull Requests****ArXiv CS.AI** **50%**

arXiv:2607.05666v1 Announce Type: cross Abstract: AI coding agents are black boxes: we cannot inspect how they generate code, but we can inspect what they change. This distinction matters for search-based software engineering (SBSE), where techniques such as genetic improvement...

<https://arxiv.org/abs/2607.05666>

1 **CONF****RPAM: A Principled Metric for Evaluating Associations in Language Models with High Predictive Validity in Downstream Outputs****ArXiv CS.AI** **50%**

arXiv:2607.05679v1 Announce Type: cross Abstract: Language models (LMs) exhibit problematic biases, such as stereotypes. Effectively analyzing and mitigating such biases requires accurate and generalizable evaluation methods of the underlying associations. Some existing...

<https://arxiv.org/abs/2607.05679>

2 **CONF****Depression Symptoms and Relational Patterns in 187k ChatGPT Histories****ArXiv CS.AI** **50%**

arXiv:2607.05685v1 Announce Type: cross Abstract: Large language models are increasingly used as private, always-available conversational systems, but little is known about how people with depressive symptoms use them. Building on CSCW work on disclosure and peer support, we...

<https://arxiv.org/abs/2607.05685>

3 **CONF****The Balkanization of Execution-Security Research for AI Coding Agents: Isolation, Access Control, and Time-of-Check-to-Time-of-Use Vulnerabilities****ArXiv CS.AI** **50%**

arXiv:2607.05743v1 Announce Type: cross Abstract: AI coding agents now read repositories, call tools, and execute shell commands with limited human oversight, and a fast-growing body of work studies whether the execution layer around them is actually safe. That literature is...

<https://arxiv.org/abs/2607.05743>

4 **CONF****When Should LLMs Search? Counterfactual Supervision for Search Routing****ArXiv CS.AI 50%**

arXiv:2607.05752v1 Announce Type: cross Abstract: Search-augmented language models can use external evidence to compensate for limitations in parametric knowledge, but search is not uniformly beneficial: models may call search for questions they can already answer, or rely on...

<https://arxiv.org/abs/2607.05752>

5 **CONF****MechVQA: Benchmarking and Enhancing Multimodal LLMs on Comprehensive Mechanical Drawing Understanding****ArXiv CS.AI 50%**

arXiv:2605.30794v2 Announce Type: replace-cross Abstract: Multimodal Large Language Models (MLLMs) have demonstrated significant achievements in general visual question answering (VQA) tasks. However, they remain brittle on mechanical engineering drawings, where high annotation...

<https://arxiv.org/abs/2605.30794>

6 **CONF****Segmentation before Answering: Pixel Grounding for MLLM Visual Reasoning****ArXiv CS.AI 50%**

arXiv:2607.05798v1 Announce Type: cross Abstract: Recent advancements in Multimodal Large Language Models (MLLMs) have evolved from static perception to interleaved visual-language reasoning, often referred to as ``thinking with images''. A basic operation in this reasoning...

<https://arxiv.org/abs/2607.05798>

7 **CONF****Complementary Roles of Image Classification and Vessel Segmentation in AI-Based Screening for Retinopathy of Prematurity Plus Disease in a Kenyan Preterm Cohort****ArXiv CS.AI 50%**

arXiv:2607.05825v1 Announce Type: cross Abstract: Background. Retinopathy of prematurity (ROP) is a preventable cause of childhood blindness, with rising burden in low- and middle-income countries where ROP-trained ophthalmologists are scarce. Plus disease, marked by retinal...

<https://arxiv.org/abs/2607.05825>

8 **CONF****VisTCP: A Visualization Framework to Construct Knowledge-Graph-Based Representation for Traditional Chinese Painting****ArXiv CS.AI 50%**

arXiv:2607.05841v1 Announce Type: cross Abstract: Structured representation can characterize semantic objects and relationships in images. It provides a possible effective way for the semantic understanding of Traditional Chinese Paintings (TCPs) to better support archaeology...

<https://arxiv.org/abs/2607.05841>

9 **CONF****Beyond Refusal: A Same-Lineage Study of Aligned and Abliterated LLMs for Vulnerability Analysis****ArXiv CS.AI 50%**

arXiv:2607.05842v1 Announce Type: cross Abstract: Large language model (LLM)-assisted software security operates at a difficult boundary: the vulnerability-analysis terminology needed for legitimate code review, triage, and repair can closely resemble terminology associated with...

<https://arxiv.org/abs/2607.05842>

0 CONF

Think Before You Grid-Search: Floor-First Triage for LLM Serving

ArXiv CS.AI 50%

arXiv:2607.05876v1 Announce Type: cross Abstract: LLM serving optimization typically benchmarks many configurations and reaches for heavy profilers when latency targets are missed. We argue for the reverse discipline: estimation is the analytical layer of profiling -- without...

<https://arxiv.org/abs/2607.05876>

1 CONF

Harrison.Rad 1.5 Technical Report: A radiology foundation model that can draft reports from images, priors and clinical context

ArXiv CS.AI 50%

arXiv:2607.05880v1 Announce Type: cross Abstract: Imaging demand is growing faster than the radiology workforce can expand, and reporting backlogs cannot be resolved through training and recruitment alone. The most direct opportunity is reducing the time and effort radiologists...

<https://arxiv.org/abs/2607.05880>

2 CONF

i-EXAM: Instructable and Explainable Attack Connectivity Graph Modeler

ArXiv CS.AI 50%

arXiv:2607.05888v1 Announce Type: cross Abstract: i-EXAM is a planning-powered tool that helps system administrators to create security profiles of complex networks and perform what-if analyses to identify network hardening strategies. It leverages planning compilation that...

<https://arxiv.org/abs/2607.05888>

3 CONF

PolicyShiftGuard: Benchmarking and Improving Policy-Adaptive Image Guardrails

ArXiv CS.AI 50%

arXiv:2607.05910v1 Announce Type: cross Abstract: Image guardrails are typically trained and evaluated under a fixed safety policy, implicitly treating safety as an intrinsic property of an image. Real deployments are different: the same image may be allowed in one product,...

<https://arxiv.org/abs/2607.05910>

4 CONF

CMDR: Contextual Multimodal Document Retrieval

ArXiv CS.AI 50%

arXiv:2607.05927v1 Announce Type: cross Abstract: Multimodal document retrieval aims to retrieve relevant pages while preserving both textual and visual content from the original document. However, existing benchmarks primarily evaluate simple lexical or semantic matching, and...

<https://arxiv.org/abs/2607.05927>

5 CONF

NegROI: Click-Centric Uncertainty-Guided Refinement with Scene-Conditioned Negative Prompts for Robust Interactive 3D Segmentation

ArXiv CS.AI 50%

arXiv:2607.05955v1 Announce Type: cross Abstract: Interactive 3D segmentation aims to extract object masks in point clouds with minimal user clicks. Despite recent progress, most existing approaches still struggle with (i) coarse voxel resolution that blurs fine boundaries under...

<https://arxiv.org/abs/2607.05955>

6 CONF

Agentic AI for IPoDWM Network Lifecycle Automation: An MCP-Enabled Architecture

ArXiv CS.AI 50%

arXiv:2607.05958v1 Announce Type: cross Abstract: We present a distributed, vendor-agnostic multi-MCP architecture for SDN-based automation and autonomous control of multi-vendor, multi-layer IPoDWM networks. The framework enables E2E service lifecycle automation, closed-loop...

<https://arxiv.org/abs/2607.05958>

7 CONF

Faithful or Findable? Evaluating LLM-Generated Metadata for RDF Dataset Search

ArXiv CS.AI 50%

arXiv:2607.05970v1 Announce Type: cross Abstract: Dataset search depends heavily on metadata, making LLM-generated metadata a consequential form of synthetic content in retrieval systems. We study six metadata-generation settings for RDF datasets, ranging from simple rewriting...

<https://arxiv.org/abs/2607.05970>

8 CONF

Propose and Attend: Training-free MLLM Grounding Confidence via Multi-Token Localized Attention

ArXiv CS.AI 50%

arXiv:2607.05978v1 Announce Type: cross Abstract: Multimodal large language models can emit localized predictions, bounding boxes for objects and temporal windows for video and audio events, but they hallucinate these regions prolifically. The model's own token log-probabilities...

<https://arxiv.org/abs/2607.05978>

9 CONF

From Blueprint to Reality: Modeling and Applying Putnam's Social Capital Theory with LLM-based Multi-agent Simulations

ArXiv CS.AI 50%

arXiv:2607.06080v1 Announce Type: cross Abstract: Putnam's Social Capital Theory is a foundational framework for collective action and community prosperity. However, traditional empirical methods face practical limits on control and replication. Meanwhile, LLM-based social...

<https://arxiv.org/abs/2607.06080>

00 CONF

Agents That Teach: Towards Designing Incidental Learning Back into AI-Assisted Software Development

ArXiv CS.AI 50%

arXiv:2607.06101v1 Announce Type: cross Abstract: AI coding agents are rapidly reshaping how software is built, with developers increasingly delegating substantial coding tasks to autonomous agents in pursuit of higher productivity. While these gains are real, they come at the...

<https://arxiv.org/abs/2607.06101>

01 CONF

EcoVision: AI-Powered Drone Imaging for Salt Marsh Vegetation Monitoring and Dominance Mapping

ArXiv CS.AI 50%

arXiv:2607.06105v1 Announce Type: cross Abstract: High-resolution RGB imagery acquired from low-altitude UAV surveys was processed through a modular pipeline incorporating transformer-based semantic segmentation, connected-component vegetation extraction, fine-grained species...

<https://arxiv.org/abs/2607.06105>

02 CONF

RoME: Robust Mixture of Low-Rank Experts against Multiple Adversarial Perturbations

ArXiv CS.AI 50%

arXiv:2607.06109v1 Announce Type: cross Abstract: Multi-perturbation adversarial training (MAT) aims to achieve robustness against multiple ℓ_p perturbations but suffers from robustness trade-offs between different threats. To address this, we employ a mixture of experts...

<https://arxiv.org/abs/2607.06109>

03 CONF

LLM-Guided Measurement Credibility Correction for Trustworthy Industrial Process Inference

ArXiv CS.AI 50%

arXiv:2607.06111v1 Announce Type: cross Abstract: Industrial prediction and soft sensing depend on credible input measurements. In field deployment, a predictor may receive biased, delayed, stale, or derived measurements that still look plausible. Prediction can then fail before...

<https://arxiv.org/abs/2607.06111>

04 CONF

x-Prediction Is All You Need: Training-Free Accelerated Generation via Endpoint Decodability

ArXiv CS.AI 50%

arXiv:2607.06114v1 Announce Type: cross Abstract: Diffusion and flow matching models generate high-quality samples, but their ODE samplers often need tens to hundreds of neural function evaluations (NFEs). This remains a practical challenge for released checkpoints, since many...

<https://arxiv.org/abs/2607.06114>

05 CONF

Property-Driven Synthetic Data Engineering for Data-Scarce Software Systems: Reflections from the Breast Cancer Domain

ArXiv CS.AI 50%

arXiv:2607.06133v1 Announce Type: cross Abstract: Modern software systems increasingly depend on data for analysis, prediction, testing, and decision-making. Yet many important domains, including medicine, safety-critical systems, and regulated industries, lack abundant...

<https://arxiv.org/abs/2607.06133>

06 CONF

LLM Agents for Deliberative Collaboration: A Study on Joint Decision Making Under Partial Observability

ArXiv CS.AI 50%

arXiv:2607.06157v1 Announce Type: cross Abstract: Deliberation plays a crucial role in collaboration; when humans work together, they naturally engage in communication to align information and reach an agreement. In this paper, we investigate deliberative large language model...

<https://arxiv.org/abs/2607.06157>

07 CONF

X-FEMR: A Token-level Explainable Approach for Electronic Health Records Foundation Models using Transformer-based Models

ArXiv CS.AI 50%

arXiv:2607.06163v1 Announce Type: cross Abstract: Foundation Models for Electronic Health Records (FEMRs) are pretrained on large-scale structured patient data, enabling them to convert longitudinal patient trajectories into generalizable representations for diverse clinical...

<https://arxiv.org/abs/2607.06163>

08 CONF

Spider 2.0-AIFunc: Extending Real-World Text-to-SQL to AI-Native SQL Workflows

ArXiv CS.AI 50%

arXiv:2607.06229v1 Announce Type: cross Abstract: Major cloud data platforms now expose large language model capabilities as native SQL functions, enabling analysts to perform classification, filtering, sentiment analysis, extraction, similarity search, and aggregation within...

<https://arxiv.org/abs/2607.06229>

09 CONF

VendorBench-100: A Unified Cross-Paradigm Benchmark for Deepfake Image Detection

ArXiv CS.AI 50%

arXiv:2607.06254v1 Announce Type: cross Abstract: Deepfake image detection is currently served by three fundamentally different paradigms: commercial APIs, zero-shot vision-language models (LLMs), and open-source detectors. Despite their widespread use, these paradigms are...

<https://arxiv.org/abs/2607.06254>

10 CONF

Designing Maintainable Hybrid Generative Systems: A Quantum-Inspired Approach to Automated Music Harmony Generation

ArXiv CS.AI 50%

arXiv:2607.06296v1 Announce Type: cross Abstract: This paper presents the design and evaluation of a maintainable hybrid generative architecture for automated music harmony generation from melody. The proposed system combines quantum-inspired candidate exploration over...

<https://arxiv.org/abs/2607.06296>

11 CONF

UI2App: Benchmarking Visual Interaction Inference in Executable Web Application Generation

ArXiv CS.AI 50%

arXiv:2607.06306v1 Announce Type: cross Abstract: Large language models (LLMs) have demonstrated growing competence in web page generation. However, existing text-driven approaches rely on complex prompts that impose substantial demands on users and offer limited expressivity...

<https://arxiv.org/abs/2607.06306>

12 CONF

RuBench: A Repository-Level Agentic Coding Benchmark with Natively Authored Russian Task Specifications

ArXiv CS.AI 50%

arXiv:2607.06411v1 Announce Type: cross Abstract: Developers increasingly delegate real maintenance work to product-grade coding agents, and many state tasks in their native language, in the style of a customer request rather than a curated English issue. Existing...

<https://arxiv.org/abs/2607.06411>

13 CONF

An Experimental Design Approach to Evaluating Agentic AI's Autonomous Model Discovery

ArXiv CS.AI 50%

arXiv:2607.06413v1 Announce Type: cross Abstract: Large language model coding agents increasingly perform open-ended data modeling and analysis. These agents are stochastic and adaptive, and therefore their autonomous model discovery behavior cannot be adequately characterized...

<https://arxiv.org/abs/2607.06413>

14 CONF

From Voting to Agent Collaboration: Answer-Type-Aware LLM Pipelines for BioASQ 14b

ArXiv CS.AI 50%

arXiv:2607.06452v1 Announce Type: cross Abstract: Biomedical question answering requires not only accurate extraction of information from scientific literature but also reliable integration of evidence across multiple documents. This study presents a question-type-specific large...

<https://arxiv.org/abs/2607.06452>

15 CONF

MLLM-LLaVA-FL: Multimodal Large Language Model Assisted Federated Learning

ArXiv CS.AI 50%

arXiv:2409.06067v3 Announce Type: replace Abstract: Previous studies on federated learning (FL) often encounter performance degradation due to data heterogeneity among different clients. In light of the recent advances in multimodal large language models (MLLMs), such as GPT-4v...

<https://arxiv.org/abs/2409.06067>

16 CONF

Industry Classification of GitHub Repositories Using the North American Industry Classification System (NAICS)

ArXiv CS.AI 50%

arXiv:2607.06505v1 Announce Type: cross Abstract: GitHub hosts hundreds of millions of public repositories, but the platform exposes no native mapping from repositories to standardized industry sectors. This gap limits empirical work on the geography of innovation, the...

<https://arxiv.org/abs/2607.06505>

17 CONF

RSF-GLLM: Bridging the Semantic Gap in Multi-Hop Knowledge Graph QA via Recurrent Soft-Flow and Decoupled LLM Generation

ArXiv CS.AI 50%

arXiv:2607.06527v1 Announce Type: cross Abstract: Multi-hop Question Answering over Knowledge Graphs faces a critical challenge: traditional retrieve-then-read pipelines break differentiability, preventing the retriever from learning to bridge the semantic gap where intermediate...

<https://arxiv.org/abs/2607.06527>

18 CONF

Graph Convolutional Attention: A Spectral Perspective on Graph Denoising and Diffusion

ArXiv CS.AI 50%

arXiv:2607.06546v1 Announce Type: cross Abstract: Denoising graphs is a fundamental problem in graph learning and the core operation of graph diffusion models. Attention-based architectures like graph transformers have recently shown promise in denoising graphs. However, our...

<https://arxiv.org/abs/2607.06546>

19 CONF

Reducing the Complexity of Deep Learning Models for EEG Analysis on Wearable Devices

ArXiv CS.AI 50%

arXiv:2606.12742v4 Announce Type: replace Abstract: Wearable healthcare devices are the fastest-growing Internet of Things (IoT) sector. Many automated healthcare services rely on two crucial biological signals, namely ECG and EEG, which reflect the activity of the heart and...

<https://arxiv.org/abs/2606.12742>

20 CONF

Replication in Visual Diffusion Models: A Survey and Outlook

ArXiv CS.AI 50%

arXiv:2408.00001v2 Announce Type: replace-cross Abstract: Visual diffusion models have revolutionized the field of creative AI, producing high-quality and diverse content. However, they inevitably memorize training images or videos, subsequently replicating their concepts,...

<https://arxiv.org/abs/2408.00001>

21 CONF

Toward AI standardization: A triadic human-ai collaboration framework for multi-level autonomous mobility

ArXiv CS.AI 50%

arXiv:2504.19120v2 Announce Type: replace-cross Abstract: The goal of the current study is to introduce a triadic human-AI collaboration framework that could be applied in transportation systems such as automated vehicles, micromobility systems, and vehicle teleoperation....

<https://arxiv.org/abs/2504.19120>

22 CONF

Narrative-Centered Emotional Reflection: An Early Prototype for AI-Supported Emotional Self-Reflection

ArXiv CS.AI 50%

arXiv:2504.20342v2 Announce Type: replace-cross Abstract: Reflexion is an AI-powered prototype designed to explore structured emotional self-reflection. By integrating emotion detection, layered reflective prompting, and metaphorical storytelling generation, Reflexion was...

<https://arxiv.org/abs/2504.20342>

23 CONF

Position: EU AI Act's Research Exemptions Can Break the Publication Norms of Major AI Conferences

ArXiv CS.AI 50%

arXiv:2506.03218v2 Announce Type: replace-cross Abstract: The EU has become one of the vanguards in regulating the digital age. A particularly important regulation in the Artificial Intelligence (AI) domain is the 2024 enacted EU AI Act. The AI Act specifies -- due to a...

<https://arxiv.org/abs/2506.03218>

24 CONF

LLM4Delay: Flight Delay Prediction via Cross-Modality Adaptation of Large Language Models and Aircraft Trajectory Representation

ArXiv CS.AI 50%

arXiv:2510.23636v4 Announce Type: replace-cross Abstract: Flight delay prediction has become a key focus in air traffic management (ATM), as delays reflect inefficiencies in the system. This paper proposes LLM4Delay, a large language model (LLM)-based framework for predicting...

<https://arxiv.org/abs/2510.23636>

25 CONF

KernelEvolve: Scaling Agentic Kernel Coding for Heterogeneous AI Accelerators at Meta

ArXiv CS.AI 50%

arXiv:2512.23236v4 Announce Type: replace-cross Abstract: Making deep learning recommendation model (DLRM) training and inference fast and efficient is important. However, this presents three key system challenges - model architecture diversity, kernel primitive diversity, and...

<https://arxiv.org/abs/2512.23236>

26 CONF

Multi-Task Instruction Tuning via Data Scheduling for Low-Resource Arabic SpeechLLMs

ArXiv CS.AI 50%

arXiv:2601.12494v3 Announce Type: replace-cross Abstract: Audio large language models (LLMs) enable unified speech understanding and generation, but adapting them to linguistically complex and dialect-rich settings such as Arabic-English remains challenging. We present a...

<https://arxiv.org/abs/2601.12494>

27 CONF

Transformers converge to invariant algorithmic cores

ArXiv CS.AI 50%

arXiv:2602.22600v2 Announce Type: replace-cross Abstract: Training selects for behavior, not circuitry: many weight configurations can implement the same function. Studying any single trained neural network thus risks describing accidents of one training run rather than the...

<https://arxiv.org/abs/2602.22600>

28 CONF

Quantifying Frontier LLM Capabilities for Container Sandbox Escape

ArXiv CS.AI 50%

arXiv:2603.02277v2 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly act as autonomous agents, using tools to execute code, read and write files, and access networks, creating novel security risks. To mitigate these risks, agents are commonly...

<https://arxiv.org/abs/2603.02277>

29 CONF

Geometry-Aware Uncertainty Coresets for Robust Visual In-Context Learning in Histopathology

ArXiv CS.AI 50%

arXiv:2605.18419v2 Announce Type: replace-cross Abstract: Vision-language models (VLMs) can couple visual perception with open-ended clinical reasoning, making them attractive for computational histopathology. However, fine-tuning billions of parameters on scarce,...

<https://arxiv.org/abs/2605.18419>

30 CONF

Mathematical Reasoning in Large Language Models: Benchmarks, Architectures, Evaluation, and Open Challenges

ArXiv CS.AI 50%

arXiv:2605.19723v2 Announce Type: replace-cross Abstract: Mathematical reasoning is essential for problem-solving in education, science, and industry, serving as a crucial benchmark for evaluating artificial intelligence systems. As Large Language Models (LLMs) improve their...

<https://arxiv.org/abs/2605.19723>

31 CONF

An LLM-Native Psychometric Instrument Reveals a Self-Report--Behavior Gap Across 25 Models

ArXiv CS.AI 50%

arXiv:2606.09843v3 Announce Type: replace-cross Abstract: Large language models (LLMs) give stable answers to personality questionnaires, yet these self-reports fail to predict how the models behave. Is this gap an artifact of forcing human trait categories onto LLMs, or...

<https://arxiv.org/abs/2606.09843>

32 CONF

Categorizing Mathematical Concepts with LLM Voting Ensembles in Mathswitch

ArXiv CS.AI 50%

arXiv:2606.28815v2 Announce Type: replace-cross Abstract: Mathswitch is an open-source project that imports mathematical concept records from sources such as Wikidata, Wikipedia, MathWorld, Encyclopedia of Mathematics, nLab, ProofWiki, and Agda-Unimath, and links records that...

<https://arxiv.org/abs/2606.28815>

33 CONF

CONFLUX: A Latent Diffusion Model for 3D Chest-CT Synthesis with RL Post-Training

ArXiv CS.AI 50%

arXiv:2607.02998v2 Announce Type: replace-cross Abstract: Controllable generative models of 3D medical images can synthesize volumes with specified clinical attributes, but this demands samples that are simultaneously high-fidelity, natively 3D, and faithful to the requested...

<https://arxiv.org/abs/2607.02998>

34 CONF

RoboDojo: A Unified Sim-and-Real Benchmark for Comprehensive Evaluation of Generalist Robot Manipulation Policies

ArXiv CS.AI 50%

arXiv:2607.04434v2 Announce Type: replace-cross Abstract: Generalist robot manipulation policies have advanced rapidly, yet existing benchmarks remain limited in systematically evaluating their capabilities. Many rely on simple, short-horizon, or skill-narrow tasks with limited...

<https://arxiv.org/abs/2607.04434>