



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	94	Media
TechCrunch AI	11	Media
HackerNews	8	Community
The Verge AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Show HN: Microsoft releases Flint, a visualization language for AI agents

HackerNews

+1 source

100%

HN 254

Data visualizations are the bridge between user and data. But building AI agents that can generate visualizations reliably can be very tricky: - simple chart specs can be reliable, but generated charts are often of low quality due to reliance on system defaults; - complex chart...

<https://microsoft.github.io/flint-chart/#/>

VERI

L-GTA: Latent Generative Modeling for Time Series Augmentation

ArXiv CS.AI

+1 source

80%

arXiv:2507.23615v2 Announce Type: replace-cross Abstract: Data augmentation is becoming increasingly important across various areas of time series analysis, including forecasting, classification, and anomaly detection. We introduce the Latent Generative Temporal Augmentation...

<https://arxiv.org/abs/2507.23615>

VERI

Future Confidence Distillation in Large Language Models

ArXiv CS.AI

+1 source

80%

arXiv:2607.07626v1 Announce Type: cross Abstract: Reliable confidence estimation is essential for deploying large language models (LLMs) in confidence-aware systems, where downstream decisions such as retrieval, tool use, and adaptive computation depend on accurately estimating...

<https://arxiv.org/abs/2607.07626>

CONFIRMED SIGNALS — 110 stories

CONF

Show HN: Freeport - a P2P ride-hailing marketplace I built on Nostr

HackerNews

50%

HN 34

For the past 3 weeks, I've been hacking on a personal AI that you can text. It's curious about you & proactively figures out ways to help you, while still feeling like a friend. Its personality & proactivity is probably what I've spent the most time working on for this project,...

<https://ptrinh.github.io/freeport/>

CONF

Show HN: Onboard-CLI, a LLM powered and AST-based tool to visualize codebase

HackerNews

50%

HN 22

No summary — see link for full article.

<https://github.com/animesh-94/Onboard-CLI>

CONF

Show HN: Frugon - Find which LLM calls a cheaper model could handle (local, MIT)

HackerNews

50%

HN 18

I started leaning in on AI heavily this year, as I wanted to get more done autonomously, but then my token usage climbed dramatically to the point where my weekly quota would run out before the end of the week, sometimes a couple of days into the week. I realised I had to do...

<https://github.com/Rodiun/frugon>

CONF

Show HN: Monogram, AI app with on the fly generation

HackerNews 50% HN 9

As someone who doesn't use any of the more "advanced" features on sites like ChatGPT or Claude (agentic mode, memory, image generation, deep research, etc), I found the bloat of these services, both in UX and in performance, to be pretty tedious. There are a million AI chat...

<https://www.monogram.ai/>

CONF

Show HN: Nicasa - A native macOS image viewer with AI tools, OCR, and annotation

HackerNews 50% HN 3

No summary — see link for full article.

<https://nicasa.w3cub.com>

CONF

Show HN: REST - Living Without Burnout. A manifesto about sustainable discipline

HackerNews 50% HN 2

Lately, I've been thinking a lot about what gives me energy and what slowly takes it away. Those thoughts eventually turned into a small manifesto I called REST.

<https://themanifesto.rest/>

0 CONF

Show HN: InterviewWatch - Detect AI Assistance During Live Technical Interviews

HackerNews 50% HN 2

No summary — see link for full article.

<https://www.interviewwatch.com/>

1 CONF

SpaceXAI releases Grok 4.5, which Elon describes as an 'Opus-class model'

TechCrunch AI 50%

Elon Musk's tech company released the newest version of Grok on Wednesday, promising a cheaper, more efficient alternative to other powerful AI models.

<https://techcrunch.com/2026/07/08/spacexai-releases-grok-4-5-which-elon-describes-as-an-...>

2 CONF

This startup thinks robotics is about to have its ChatGPT moment

TechCrunch AI 50%

General Intuition is betting millions of hours of video game data can train the foundation models for physical AI, making it easier to build smarter robots with minimal real-world data.

<https://techcrunch.com/2026/07/08/this-startup-thinks-robotics-is-about-to-have-its-cha...>

3 CONF

Google Photos adds a new AI 'Video Remix' tool

TechCrunch AI 50%

The feature can do things like apply cinematic relighting to brighten up a dark clip, swap out a plain background for something fun, or add artistic styles to videos.

<https://techcrunch.com/2026/07/08/google-photos-adds-a-new-ai-video-remix-tool/>

4 CONF

Why this CEO thinks video games make better training data than the internet

TechCrunch AI 50%

When it comes to achieving artificial general intelligence (AGI), large language models just don't have what it takes. Models like ChatGPT and Claude are great at text, but they're less skilled at understanding how things actually move through space and time — an essential skill...

<https://techcrunch.com/video/why-this-ceo-thinks-video-games-make-better-training-data-...>

5 CONF

Meta wants its AI glasses to seem less creepy. Its AI strategy says otherwise.

TechCrunch AI 50%

Meta is adding a new safeguard to stop people from secretly recording others with its AI glasses. But the update comes as the company continues to expand how much personal data its AI products collect and use.

<https://techcrunch.com/2026/07/08/meta-wants-its-ai-glasses-to-seem-less-creepy-its-ai-...>

6 CONF

Prime Intellect raises \$130M Series A to help enterprises build their own AI agents

TechCrunch AI 50%

Founded in 2024, Prime Intellect's goal is to give organizations capabilities to train their own agentic systems without relying on frontier AI labs.

<https://techcrunch.com/2026/07/08/prime-intellect-raises-130m-series-a-to-help-enterpri...>

7 CONF

These AI startups are growing revenue at faster and faster rates

TechCrunch AI 50%

There are a lot of fast-growing AI startups, but some are growing even faster, they say.

<https://techcrunch.com/2026/07/08/these-ai-startups-are-growing-revenue-at-faster-and-f...>

8 CONF

Former OpenAI exec Kevin Weil is now on the board of Stoke Space

TechCrunch AI 50%

Kevin Weil's new role at Stoke Space suggests reusable rockets are the next hot thing in Silicon Valley.

<https://techcrunch.com/2026/07/08/former-openai-exec-kevin-weil-is-now-on-the-board-of-...>

9 CONF

Hot French startup ZML releases free product to speed inference across lots of AI chips

TechCrunch AI 50%

ZML, a hot French AI startup endorsed by Turing Award winner Yann LeCun, has now released ZML/LLMD, software that could make running AI less costly.

<https://techcrunch.com/2026/07/08/hot-french-startup-zml-releases-free-product-to-speed...>

0 CONF

AI chip maker SambaNova raises \$1B at \$11B valuation, 5 months after last mega round

TechCrunch AI 50%

AI chip maker SambaNova has raised at an \$11 billion valuation months after Intel was rumored to be trying to buy it for about \$1.6 billion.

<https://techcrunch.com/2026/07/08/sambanova-draws-1b-at-11b-valuation-in-series-f-first...>

1 CONF

EmTech AI 2026: The Rise of the AI Platform

MIT Tech Review 50%

No summary — see link for full article.

<https://www.technologyreview.com/2026/07/08/1140223/emtech-ai-2026-the-rise-of-the-ai-p...>

2 CONF

ChatGPT's upgraded voice mode is better at shutting up

The Verge AI 50%

OpenAI is overhauling ChatGPT's voice mode with a new model that it says is more like "talking to another person." The new GPT-Live-1 is designed to interrupt you less and will also wait for you to continue speaking if you pause mid-conversation. During a press briefing, OpenAI...

<https://www.theverge.com/ai-artificial-intelligence/962856/chatgpt-upgraded-voice-mode-...>

3 CONF

LLM-powered reasoning in agent-based modeling

ArXiv CS.AI 50%

arXiv:2607.06757v1 Announce Type: new Abstract: Agent-based modeling (ABM) has the capability to model millions of individuals and their interactions, which is useful for policy making. However, ABMs have traditionally relied on static prior, which prevents the models from...

<https://arxiv.org/abs/2607.06757>

4 CONF

Evaluating SageMath-Augmented LLM Agents for Computational and Experimental Mathematics

ArXiv CS.AI 50%

arXiv:2607.06820v1 Announce Type: new Abstract: Recent advances in AI for Mathematics have focused largely on autoformalization and theorem proving, leaving the role of Computer Algebra Systems (CAS) in agentic LLM workflows underexplored. We propose a ReAct-style agentic setup...

<https://arxiv.org/abs/2607.06820>

5 CONF

The Harness Effect: How Orchestration Design Sets the Token Economics of Enterprise Agentic AI

ArXiv CS.AI 50%

arXiv:2607.06906v1 Announce Type: new Abstract: Agentic AI development today runs on token maging: buying capability with tokens -- longer reasoning traces, more turns, wider tool payloads, bigger replayed contexts -- so tokens per task grow faster than task value. Falling...

<https://arxiv.org/abs/2607.06906>

6 CONF

Large Behavior Model: A Promptable Digital Twin of the Retail Customer

ArXiv CS.AI 50%

arXiv:2607.06993v1 Announce Type: new Abstract: Customer behavior modeling underpins recommendation, marketing, and decision support, yet existing approaches either optimize predictive accuracy without explaining decisions or simulate users without grounding them in real...

<https://arxiv.org/abs/2607.06993>

7 CONF

Learning social norms enhances compatibility in dynamic human-AI coordination

ArXiv CS.AI 50%

arXiv:2607.07021v1 Announce Type: new Abstract: Humans continuously coordinate with others in dynamic interactions, often through implicit, hard-to-quantify social norms that act as shared tacit expectations among interacting agents. As AI agents, including large language models...

<https://arxiv.org/abs/2607.07021>

8 CONF

Operational Reframing and Approval-Framed Delegation in Multi-Agent LLM Safety

ArXiv CS.AI 50%

arXiv:2607.07097v1 Announce Type: new Abstract: Safety evaluations of multi-agent LLM systems often compare a direct prompt with a planner-executor pipeline and report the difference as a single "pipeline effect." We argue that this aggregate is difficult to interpret because it...

<https://arxiv.org/abs/2607.07097>

9 CONF

Reasoning Consistency Scanning: A Framework for Auditing Chain-of-Thought Validity in AI Safety Evaluations

ArXiv CS.AI 50%

arXiv:2607.07229v1 Announce Type: new Abstract: Prior work has shown that chain-of-thought (CoT) reasoning is often unfaithful: a model's stated reasoning does not reliably reflect the process that produced its output. Detecting unfaithfulness, though, requires controlled...

<https://arxiv.org/abs/2607.07229>

0 CONF

From Beats to Breaches:How Offensive AI Infers Sensitive User Information from Playlists

ArXiv CS.AI 50%

arXiv:2605.04724v2 Announce Type: replace-cross Abstract: The pervasive integration of AI has enabled Offensive AI: the exploitation of AI for malicious ends across the cyber-kill chain. A critical manifestation is the user attribute inference attack, where AI infers sensitive...

<https://arxiv.org/abs/2605.04724>

1 CONF

Physics-Audited Agentic Discovery in Scientific Machine Learning

ArXiv CS.AI 50%

arXiv:2607.07379v1 Announce Type: new Abstract: In agentic scientific machine learning (SciML), large language model (LLM) agents can discover surrogate models and select one by an automated score, typically an error metric. A low error, however, does not establish that the...

<https://arxiv.org/abs/2607.07379>

2 CONF

MIRA-Math: A Benchmark for Minimal Information Requesting and Mathematical Reasoning

ArXiv CS.AI 50%

arXiv:2607.07391v1 Announce Type: new Abstract: Mathematical reasoning benchmarks typically provide all facts needed to solve each problem, while interactive benchmarks often mix reasoning with tools, retrieval, and long-horizon dialogue. We introduce MIRA-Math, a benchmark for...

<https://arxiv.org/abs/2607.07391>

3 CONF

Reason Less, Verify More: Deterministic Gates Recover a Silent Policy-Violation Failure Mode in Tool-Using LLM Agents

ArXiv CS.AI 50%

arXiv:2607.07405v1 Announce Type: new Abstract: Tool-using LLM agents can violate the very policies they are deployed to enforce while appearing to complete the task successfully. In policy-permissive environments, a tool may execute any well-formed call even when the...

<https://arxiv.org/abs/2607.07405>

4 CONF

SpaCellAgent: A Self-Evolving LLM-Based Multi-Agent Framework for Trajectory Analysis

ArXiv CS.AI 50%

arXiv:2607.07467v1 Announce Type: new Abstract: Spatial and Single-cell transcriptomics are transformative in deciphering cellular dynamics. As the fundamental paradigm for reconstructing cell developmental paths, trajectory inference (TI) is critical. However, existing methods...

<https://arxiv.org/abs/2607.07467>

5 CONF

Explain Before You Answer: A Survey on Compositional Visual Reasoning

ArXiv CS.AI 50%

arXiv:2508.17298v3 Announce Type: replace-cross Abstract: Compositional visual reasoning has emerged as a key research frontier in multimodal AI, aiming to endow machines with the human-like ability to decompose visual scenes, ground intermediate concepts, and perform multi-step...

<https://arxiv.org/abs/2508.17298>

6 CONF

Do LLM-Generated Skills Make Better AI Data Scientists? A Component Ablation Across Data-Science Workflows

ArXiv CS.AI 50%

arXiv:2607.07504v1 Announce Type: new Abstract: Product data scientists often ask LLM-based agents to help with recurring execution tasks such as cleaning data, writing SQL, choosing statistical tests, and formatting results. Reusable skill files are meant to avoid prompting...

<https://arxiv.org/abs/2607.07504>

7 CONF

Recursive Self-Improvement in AI: From Bounded Self-Refinement to Autonomous Research Loops

ArXiv CS.AI 50%

arXiv:2607.07663v1 Announce Type: new Abstract: AI systems increasingly participate in their own improvement: revising their outputs, adapting their own harnesses during deployment, training on data they generate, and, increasingly, conducting AI research itself. This literature...

<https://arxiv.org/abs/2607.07663>

8 CONF

SkillCenter: A Large-Scale Source-Grounded Skill Library for Autonomous AI Agents

ArXiv CS.AI 50%

arXiv:2607.07676v1 Announce Type: new Abstract: Autonomous AI agents can execute complex tasks with limited human review, yet they often lack the grounded operational knowledge to make their outputs not just executable but correct, secure, and maintainable. We introduce...

<https://arxiv.org/abs/2607.07676>

9 CONF

Deep Reinforcement Learning for Reliability Based Bi-Objective Portfolio Optimization

ArXiv CS.AI 50%

arXiv:2607.06610v1 Announce Type: cross Abstract: Portfolio optimization under uncertainty is inherently a multi-objective decision problem involving complex interactions among return, risk, market dynamics, and practical investment constraints. Existing reliability based...

<https://arxiv.org/abs/2607.06610>

0 CONF

LLM-Guided Task-Semantic Field Factorization for Industrial Process Forecasting

ArXiv CS.AI 50%

arXiv:2607.06623v1 Announce Type: cross Abstract: Process industries rely on time-series forecasting and soft sensing to estimate quality variables that are hard to measure online. Labeled data are scarce, operating regimes change frequently, and retraining models or rebuilding...

<https://arxiv.org/abs/2607.06623>

1 CONF

Non-contact, Real-time, Heart-rate Measurement using Image Processing with Commodity Cameras and AI Agents

ArXiv CS.AI 50%

arXiv:2607.06598v1 Announce Type: cross Abstract: Heart rate measurement is one of the key requirements for real-time health monitoring, in particular for health caring of elderly people. Traditional heart rate measurement relies on contact sensing mechanisms such as some heart...

<https://arxiv.org/abs/2607.06598>

2 CONF

MiLSD: A Micro Line-Segment Detector for Resource-Constrained Devices

ArXiv CS.AI 50%

arXiv:2607.06600v1 Announce Type: cross Abstract: Line segment detection is a key building block in visual SLAM, 3D reconstruction, and industrial inspection. Recent deep learning methods have greatly improved accuracy, yet even the smallest models require several megabytes of...

<https://arxiv.org/abs/2607.06600>

3 CONF

Do Counterfactually Fair Image Classifiers Satisfy Group Fairness? -- A Theoretical and Empirical Study

ArXiv CS.AI 50%

arXiv:2607.06603v1 Announce Type: cross Abstract: The notion of algorithmic fairness has been actively explored from various aspects of fairness, such as counterfactual fairness (CF) and group fairness (GF). However, the exact relationship between CF and GF remains to be...

<https://arxiv.org/abs/2607.06603>

4 CONF

Security and Privacy in Agentic AI: Grand Challenges and Future Directions

ArXiv CS.AI 50%

arXiv:2607.06608v1 Announce Type: cross Abstract: We present key challenges and future research directions in the security and privacy of agentic AI, based on a horizon-scanning exercise that brought together thirty leading international experts from academia, industry, and...

<https://arxiv.org/abs/2607.06608>

5 CONF

D2PO: Optimizing Diffusion Samplers via Dynamic Preference

ArXiv CS.AI 50%

arXiv:2607.06609v1 Announce Type: cross Abstract: We propose D2PO (Dynamic Direct Preference Optimization), a principled framework for optimizing diffusion sampling policies with respect to timestep schedules and classifier-free guidance (CFG) weights. Our work is motivated by a...

<https://arxiv.org/abs/2607.06609>

6 CONF

STAGformer: A Spatio-temporal Agent Graph Transformer for Micro Mobility Demand Forecasting

ArXiv CS.AI 50%

arXiv:2607.06614v1 Announce Type: cross Abstract: Accurate station-level demand forecasting is essential for the efficient operation of bike-sharing systems, yet it remains challenging due to complex spatio-temporal dependencies and the large scale of urban networks. This paper...

<https://arxiv.org/abs/2607.06614>

7 **CONF****LoCA: Spatially-Aware Low-Rank Convolutional Adaptation of Vision Foundation Models****ArXiv CS.AI 50%**

arXiv:2607.06918v1 Announce Type: cross Abstract: Pre-trained Vision Foundation Models (VFMs) provide strong visual representations for diverse downstream tasks. The key challenge of VFM adaptation stems from the prohibitive costs of full fine-tuning and catastrophic forgetting....

<https://arxiv.org/abs/2607.06918>

8 **CONF****Overview of the NLPCC 2026 Shared Task 1: Difficulty-Aware Multilingual and Multimodal Medical Instructional Video Understanding Evaluation****ArXiv CS.AI 50%**

arXiv:2607.06618v1 Announce Type: cross Abstract: Following the CMIVQA, MMI-VQA, and M4IVQA challenges in NLPCC 2023--2025, we introduce the Difficulty-Aware Medical Instructional Video Question Answering (DA-MIVQA) shared task for NLPCC 2026. DA-MIVQA extends previous...

<https://arxiv.org/abs/2607.06618>

9 **CONF****ProMoE-FL: Prototype-conditioned Mixture of Experts for Multimodal Federated Learning with Missing Modalities****ArXiv CS.AI 50%**

arXiv:2607.06633v1 Announce Type: cross Abstract: In this paper, we address the problem of multimodal federated learning with missing modality. Existing methods utilize an additional public dataset or perform naive feature synthesis that is based solely on the available...

<https://arxiv.org/abs/2607.06633>

0 **CONF****Specification Grounding Drives Test Effectiveness for LLM Code****ArXiv CS.AI 50%**

arXiv:2607.06636v1 Announce Type: cross Abstract: Large language models frequently generate code that appears correct on typical inputs yet fails on edge cases, invalid inputs, and other specification-defined corner conditions. A popular fix has the model write its own tests and...

<https://arxiv.org/abs/2607.06636>

1 **CONF****Healthier LLMs: Retrieval-Augmented Generation for Public Health Question Answering****ArXiv CS.AI 50%**

arXiv:2607.06641v1 Announce Type: cross Abstract: Large language models (LLMs) achieve promising results on medical question answering benchmarks, yet their use in public health is constrained by hallucinations and the rapid evolution of official guidance. Retrieval-Augmented...

<https://arxiv.org/abs/2607.06641>

2 **CONF****DiPhon: Diffusion on Graphons for Scalable Graph Generation****ArXiv CS.AI 50%**

arXiv:2607.07232v1 Announce Type: cross Abstract: Diffusion models represent a leading paradigm for graph generation, with notable impact in domains such as molecular design. Yet, scaling these models to large graphs remains an open problem. We approach this question in the...

<https://arxiv.org/abs/2607.07232>

3 CONF

Digital Fragmentation and Generative AI Use Across 103 Million Application Events

ArXiv CS.AI 50%

arXiv:2607.06681v1 Announce Type: cross Abstract: Knowledge workers switch between applications thousands of times per day, spending nearly a tenth of the work year transitioning between digital applications in a process called digital fragmentation. Whether this fragmentation...

<https://arxiv.org/abs/2607.06681>

4 CONF

tsbootstrap: Distribution-Free Uncertainty Quantification and Conformal Prediction for Time Series

ArXiv CS.AI 50%

arXiv:2607.06690v1 Announce Type: cross Abstract: Finance, sensing, and demand streams violate the exchangeability that IID conformal prediction and the IID bootstrap assume, and existing libraries implement either a general resampling engine or conformal calibration without the...

<https://arxiv.org/abs/2607.06690>

5 CONF

SPEAR: A Simulator for Photorealistic Embodied AI Research

ArXiv CS.AI 50%

arXiv:2607.06701v1 Announce Type: cross Abstract: Interactive simulators have become powerful tools for training embodied agents and generating synthetic visual data, but existing photorealistic simulators suffer from limited generality, programmability, and rendering speed. We...

<https://arxiv.org/abs/2607.06701>

6 CONF

SmartHomeSecure: Automated Detection and Repair of Smart Home Configuration Errors Using Large Language Models

ArXiv CS.AI 50%

arXiv:2607.06748v1 Announce Type: cross Abstract: Smart home automation platforms increasingly rely on user-authored YAML configuration files to define device behaviors, but these files are prone to syntax, formatting, and semantic logic errors that can cause automation failures...

<https://arxiv.org/abs/2607.06748>

7 CONF

AirPASS: Over-the-Air Federated Learning via Pinching Antenna Systems

ArXiv CS.AI 50%

arXiv:2607.06768v1 Announce Type: cross Abstract: This paper investigates over-the-air federated learning (AirFL) in wireless systems where the access point is equipped with a multi-waveguide pinching antenna system (PASS). We adopt the widely studied learning-oriented AirFL...

<https://arxiv.org/abs/2607.06768>

8 CONF

From Agentic to Autogenic Network Management for AI-Native 6G and Beyond: A Standards Perspective

ArXiv CS.AI 50%

arXiv:2607.06786v1 Announce Type: cross Abstract: Standards bodies, including TM Forum, 3GPP, and ETSI, are converging on Agentic AI as the foundation for next-generation network management, where Large AI Model (LAM)-based agents autonomously interpret intent, coordinate...

<https://arxiv.org/abs/2607.06786>

9 CONF

Enhancing deep learning models for time series classification via knowledge distillation

ArXiv CS.AI 50%

arXiv:2607.06796v1 Announce Type: cross Abstract: Deep learning has achieved remarkable success in various domains including time series analysis, computer vision and natural language processing. However, high computational and memory demands of state-of-the-art architectures...

<https://arxiv.org/abs/2607.06796>

0 CONF

A Multi-Analyst LLM Pipeline for Auditable Rule Discovery Across 68 Public Physiological Corpora

ArXiv CS.AI 50%

arXiv:2607.06802v1 Announce Type: cross Abstract: Open physiological corpora are heterogeneous: they use different sensors, labels, sampling rates, recording settings, and clinical endpoints. They can support detector design, but they do not directly specify which detector rules...

<https://arxiv.org/abs/2607.06802>

1 CONF

Ad Headline Generation using Self-Critical Masked Language Model

ArXiv CS.AI 50%

arXiv:2607.06818v1 Announce Type: cross Abstract: For any E-commerce website it is a nontrivial problem to build enduring advertisements that attract shoppers. It is hard to pass the creative quality bar of the website, especially at a large scale. We thus propose a programmatic...

<https://arxiv.org/abs/2607.06818>

2 CONF

Gradient-Based Speech-to-Text Alignment for Any ASR Model: From CTC to Speech LLMs

ArXiv CS.AI 50%

arXiv:2607.06831v1 Announce Type: cross Abstract: Speech-to-text alignment means finding the temporal boundaries of each word in the audio. Some models provide such an alignment directly and others do not. Connectionist temporal classification (CTC) and transducer models have an...

<https://arxiv.org/abs/2607.06831>

3 CONF

ReMoDEx: A Local-to-Global Relevance-Based Model Decision Explainability Framework for large-Scale Image Datasets

ArXiv CS.AI 50%

arXiv:2607.06889v1 Announce Type: cross Abstract: Deep learning image classifiers achieve strong predictive performance yet remain opaque in how decisions are formed. A model may predict correctly while relying on irrelevant cues, shortcut associations, peripheral structures, or...

<https://arxiv.org/abs/2607.06889>

4 CONF

LiveOIBench: Can Large Language Models Outperform Human Contestants in Informatics Olympiads?

ArXiv CS.AI 50%

arXiv:2510.09595v3 Announce Type: replace Abstract: Competitive programming problems are increasingly used to evaluate the coding capabilities of large language models (LLMs) due to their complexity and ease of verification. Yet, current coding benchmarks face limitations such...

<https://arxiv.org/abs/2510.09595>

5 CONF

Self-Supervised Pretraining Improves Cross-Site and Cross-Scale Robustness of Point Cloud Leaf-Wood Segmentation

ArXiv CS.AI 50%

arXiv:2607.06948v1 Announce Type: cross Abstract: The accuracy of existing leaf-wood segmentation methods for tree point clouds varies across forest types and sites. Self-supervised learning (SSL) on point clouds has improved the generalization of deep learning models for...

<https://arxiv.org/abs/2607.06948>

6 CONF

Large Language Models (LLMs) and Generative AI in Cybersecurity and Privacy: A Survey of Dual-Use Risks, AI-Generated Malware, Explainability, and Defensive Strategies

ArXiv CS.AI 50%

arXiv:2607.06963v1 Announce Type: cross Abstract: Large Language Models (LLMs) and generative AI (GenAI) systems, such as ChatGPT, Claude, Gemini, LLaMA, Copilot, Stable Diffusion by OpenAI, Anthropic, Google, Meta, Microsoft, Stability AI, respectively, are revolutionizing...

<https://arxiv.org/abs/2607.06963>

7 CONF

End-to-End LLM Flight Planning with RAG-based Memory and Multi-modal Coach Agent

ArXiv CS.AI 50%

arXiv:2607.06964v1 Announce Type: cross Abstract: Bridging the gap between human pilot intent and autonomous flight operation is critical for real-world electric vertical takeoff and landing (eVTOL) aircraft deployment. Flight planning traditionally relies on classic algorithms...

<https://arxiv.org/abs/2607.06964>

8 CONF

Multimodal Spatiotemporal-Frequency Fusion with Peak Enhancement for Cellular Traffic Forecasting

ArXiv CS.AI 50%

arXiv:2607.07016v1 Announce Type: cross Abstract: Accurate forecasting of cellular network traffic is essential for network planning, resource allocation, and quality-of-service assurance in modern mobile communication systems. Real-world traffic often exhibits bursty endogenous...

<https://arxiv.org/abs/2607.07016>

9 CONF

Latent graph encoding of multimodal neuroimaging features with generative AI architectures

ArXiv CS.AI 50%

arXiv:2607.07027v1 Announce Type: cross Abstract: While generative models enable encoding of complex neuroimaging data for feature generation and reconstruction, developing optimal architectural frameworks with appropriate encoding and latent space processes is crucial for...

<https://arxiv.org/abs/2607.07027>

10 CONF

Riemannian Geometry for Pre-trained Language Model Embeddings

ArXiv CS.AI 50%

arXiv:2607.07047v1 Announce Type: cross Abstract: Understanding the geometric structure of pre-trained language model embeddings matters for interpretability and safety. We ask whether sentence-level classification signal lives in the Riemannian geometry of contextual token...

<https://arxiv.org/abs/2607.07047>

1 **CONF****Multiplication Beyond Groups: Stratified Fourier Mechanisms in Transformer Circuits**

ArXiv CS.AI 50%

arXiv:2607.07066v1 Announce Type: cross Abstract: Transformers have demonstrated a remarkable ability to learn algorithmic reasoning, yet mechanistic analyses have mostly focused on globally invertible operations such as cyclic addition and group composition. In this work, we...

<https://arxiv.org/abs/2607.07066>

2 **CONF****Navigating Hierarchy: Hyperbolic Learning on Brain Graphs for Disorder Diagnosis**

ArXiv CS.AI 50%

arXiv:2607.07077v1 Announce Type: cross Abstract: Functional brain networks exhibit a hierarchical organization across ROI, community, and whole-brain levels, supporting local processing, inter-community coordination, and global integration. Recent studies have demonstrated that...

<https://arxiv.org/abs/2607.07077>

3 **CONF****AT-Attn: Temporal-Aware Cross-Attention for Longitudinal Multimodal Alzheimer's Disease Diagnosis**

ArXiv CS.AI 50%

arXiv:2607.07091v1 Announce Type: cross Abstract: In longitudinal Alzheimer's disease (AD) diagnosis support, clinical and imaging information is often collected at irregular visits. Integrating these multimodal observations may improve diagnostic assessment, but naive fusion...

<https://arxiv.org/abs/2607.07091>

4 **CONF****Single-Rollout Asynchronous Optimization for Agentic Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2607.07508v1 Announce Type: cross Abstract: Reinforcement learning (RL) is becoming increasingly important for post-training large language models (LLMs). Previous RL pipelines for LLMs were mostly synchronous and batch-interleaved, which is inefficient for long-horizon...

<https://arxiv.org/abs/2607.07508>

5 **CONF****Predicting LLM Safety Before Release by Simulating Deployment**

ArXiv CS.AI 50%

arXiv:2607.07184v1 Announce Type: cross Abstract: Pre-deployment safety evaluations aim to inform the downstream risks of releasing a new AI model. Yet most evaluations provide limited evidence about how often undesired model behavior will occur in deployment: they generally...

<https://arxiv.org/abs/2607.07184>

6 **CONF****Memory Scarcity, Open Models, and the Restructuring of the AI Industry, 2026-2030 -- A quantitative scenario analysis of inference economics, training-cost divergence, and infrastructure solvency**

ArXiv CS.AI 50%

arXiv:2607.07207v1 Announce Type: cross Abstract: We analyze how four forces restructure the AI industry over 2026-2030: the DRAM/HBM price surge, frontier-capable open-weight models (GLM-5.2), rapid inference-efficiency gains (near-Shannon-limit KV-cache compression,...

<https://arxiv.org/abs/2607.07207>

7 **CONF****Vision Foundation Models in Radiology: A Scoping Review of Data, Methodology, Evaluation and Clinical Translation****ArXiv CS.AI** **50%**

arXiv:2607.07219v1 Announce Type: cross Abstract: Vision foundation models (VFM) are increasingly being developed for radiological imaging, yet their definition, development and evaluation remain heterogeneous. We conducted a PRISMA ScR scoping review of peer-reviewed studies...

<https://arxiv.org/abs/2607.07219>

8 **CONF****ORCAID: Oblique Rule-Based Continuous-Action Interpretation for Deep RL Policies****ArXiv CS.AI** **50%**

arXiv:2607.07235v1 Announce Type: cross Abstract: Explainability remains a key issue in reinforcement learning (RL). Distilling an interpretable policy from an agent trained in a complex environment is particularly challenging when the action space is continuous. We introduce...

<https://arxiv.org/abs/2607.07235>

9 **CONF****CarbonCLIP: Enhance Carbon Prediction from Satellite Imagery via Integrated Street-View Semantics and Temporal Context Training****ArXiv CS.AI** **50%**

arXiv:2607.07292v1 Announce Type: cross Abstract: Accurately estimating urban carbon emissions is critical for sustainable urban planning, yet many existing approaches remain difficult to apply consistently across cities due to data-source heterogeneity and the lack of...

<https://arxiv.org/abs/2607.07292>

10 **CONF****Multimodal Voice Activity Projection for Turn-Taking in Social Robots with Voice-Activity-Related Pretrained Encoders****ArXiv CS.AI** **50%**

arXiv:2607.07294v1 Announce Type: cross Abstract: Turn-taking prediction is a key requirement for social robots involved in human-human interaction, particularly in mediator settings, where the robot must anticipate conversational dynamics rather than merely react to pauses....

<https://arxiv.org/abs/2607.07294>

11 **CONF****HAIJv2-CrowdCount: Zero-Shot Benchmark for Dense Crowd Counting****ArXiv CS.AI** **50%**

arXiv:2607.07322v1 Announce Type: cross Abstract: Automated crowd counting in Hajj video is difficult not because current models lack capacity, but because the footage violates the assumptions those models were built on: cameras observe the crowd from steep, near-vertical...

<https://arxiv.org/abs/2607.07322>

12 **CONF****Heterogeneity-Adaptive Diffusion Schrodinger Bridge for PET-Guided Whole-Body MRI Translation****ArXiv CS.AI** **50%**

arXiv:2607.07401v1 Announce Type: cross Abstract: While whole-body multimodal medical imaging scanners have been increasingly recognized for more effective medical applications, the excessive long acquisition time in PET-MR scanning is a major obstacle in more efficient clinical...

<https://arxiv.org/abs/2607.07401>

3 CONF

Latency-Aware Bid Acceptance under Operational Feasibility: A Public Benchmark with Hindsight Ceilings

ArXiv CS.AI 50%

arXiv:2607.07343v1 Announce Type: cross Abstract: Online truckload bid acceptance is a closed-loop stochastic decision problem in which a carrier or broker must, in real time, accept or reject a tendered load subject to operational feasibility, fleet repositioning costs, and...

<https://arxiv.org/abs/2607.07343>

4 CONF

HumAIN: Human-Aware Implicit Social Robot Navigation

ArXiv CS.AI 50%

arXiv:2607.07357v1 Announce Type: cross Abstract: Effective social robot navigation requires sensitivity to human behavior, often revealed through subtle skeletal cues like gait and orientation. We present Human-Aware Implicit Social Robot Navigation (HumAIN), a novel framework...

<https://arxiv.org/abs/2607.07357>

5 CONF

Multi-Agent AI Control: Distributed Attacks Hamper Per-Instance Monitors

ArXiv CS.AI 50%

arXiv:2607.07368v1 Announce Type: cross Abstract: AI control is a family of techniques to prevent an AI with malicious goals from subverting its operator's intent. AI Control usually studies a single agent in one trajectory, but real deployments run many agents over shared...

<https://arxiv.org/abs/2607.07368>

6 CONF

On Adversarial Vulnerability of Vision-Language Models through the Lens of Intermediate Spectral Subspaces

ArXiv CS.AI 50%

arXiv:2607.07375v1 Announce Type: cross Abstract: Adversarial vulnerability in deep neural networks (DNNs) has been studied from the perspectives of decision-boundary geometry, feature robustness, input-output Jacobians, and the instability of inverse problems. Here, we focus on...

<https://arxiv.org/abs/2607.07375>

7 CONF

SynthAVE: Scalable Synthetic Labeling for E-Commerce with LLM-Arena Validation

ArXiv CS.AI 50%

arXiv:2607.07469v1 Announce Type: cross Abstract: Fine-tuning large language models (LLMs) for e-commerce attribute extraction requires labeled data representative across thousands of product types, attributes, and multiple languages. This combinatorial scale translates to...

<https://arxiv.org/abs/2607.07469>

8 CONF

Where to Intervene? Benchmarking Fairness-Aware Learning on Differentially Private Synthetic Tabular Data

ArXiv CS.AI 50%

arXiv:2607.07471v1 Announce Type: cross Abstract: Machine learning models are increasingly deployed in high-stakes domains, raising concerns about both privacy and fairness. Differential Privacy (DP) has become a gold standard for privacy-preserving data analysis, while...

<https://arxiv.org/abs/2607.07471>

9 CONF

Beyond Attack-Success Rate: Action-Graded Severity Scale for Tool-Using AI Agents

ArXiv CS.AI 50%

arXiv:2607.07474v1 Announce Type: cross Abstract: Agentic red-teaming benchmarks report whether an injected agent was compromised as a single bit: the attack succeeded, or it did not. We argue that this binary attack-success rate discards the information a defender most needs,...

<https://arxiv.org/abs/2607.07474>

0 CONF

Creativity from Friction: Human-AI Interaction for Exploratory Structural Design

ArXiv CS.AI 50%

arXiv:2607.07521v1 Announce Type: cross Abstract: AI agents that generate final answers based on user input often do not meet the needs of creative fields. Fields such as structural design and architecture need interactive systems that help users externalise and develop ideas,...

<https://arxiv.org/abs/2607.07521>

1 CONF

ALER-TI: Aligned Latent Embedding Retrieval for Time Series Imputation

ArXiv CS.AI 50%

arXiv:2607.07640v1 Announce Type: cross Abstract: Deep learning has significantly advanced time series imputation, yet most existing architectures primarily rely on localized temporal context within the corrupted input sequence. This reliance can be limiting in real-world...

<https://arxiv.org/abs/2607.07640>

2 CONF

DiaLLM: An Investigation into the Robustness-Generation Gap in English Dialect Adaptation

ArXiv CS.AI 50%

arXiv:2607.07669v1 Announce Type: cross Abstract: Large language models increasingly \emph{understand} dialectal English, yet still \emph{produce} only standard, US-leaning English, leaving dialectal generation, the harder half of the problem, largely unaddressed. We introduce...

<https://arxiv.org/abs/2607.07669>

3 CONF

Selective Timestep Weighting and Advantage-Based Replay for Sample-Efficient Diffusion RLHF

ArXiv CS.AI 50%

arXiv:2607.07693v1 Announce Type: cross Abstract: Reinforcement learning from human feedback (RLHF) has emerged as a powerful paradigm for aligning generative models with human preferences. However, applying RLHF to diffusion models remains highly feedback inefficient, as...

<https://arxiv.org/abs/2607.07693>

4 CONF

Breaking Database Lock-in: Agentic Regeneration of High Performance Storage Readers for Database Bypass

ArXiv CS.AI 50%

arXiv:2607.07696v1 Announce Type: cross Abstract: Analytical workloads operating on data stored in external database systems face a fundamental bottleneck: data access is guarded entirely by the database driver, like JDBC or ODBC, forcing all reads through query execution and...

<https://arxiv.org/abs/2607.07696>

5 **CONF****Co-LMLM: Continuous-Query Limited Memory Language Models**

ArXiv CS.AI 50%

arXiv:2607.07707v1 Announce Type: cross Abstract: Limited memory language models (LMLMs) externalize factual knowledge during pretraining to a knowledge base (KB), rather than memorizing it in their weights. During generation, the model then fetches knowledge from the KB as...

<https://arxiv.org/abs/2607.07707>

6 **CONF****Successor-Generator Planning with LLM-generated Heuristics**

ArXiv CS.AI 50%

arXiv:2501.18784v5 Announce Type: replace Abstract: Heuristics are a central component of deterministic planning, particularly in domain-independent settings where general applicability is prioritized over task-specific tuning. This work revisits that paradigm in light of recent...

<https://arxiv.org/abs/2501.18784>

7 **CONF****AGAPI-Agents: An Open-Access Agentic AI Platform for Accelerated Materials Design on AtomGPT.org**

ArXiv CS.AI 50%

arXiv:2512.11935v2 Announce Type: replace Abstract: Agentic AI systems increasingly connect large language models (LLMs) to external scientific tools, yet whether and when tool access improves prediction accuracy remains uncharacterized. We present AGAPI (AtomGPT.org API), an...

<https://arxiv.org/abs/2512.11935>

8 **CONF****Neutral Substrates: A Design Constraint for Shared Records Under Persistent Interpretive Disagreement**

ArXiv CS.AI 50%

arXiv:2601.14271v2 Announce Type: replace Abstract: Shared accountability records are often used by parties who may never agree about causation, responsibility, or normative interpretation. For such records, neutrality cannot be achieved by omitting contested information,...

<https://arxiv.org/abs/2601.14271>

9 **CONF****Framing Instability in LLM Ethical Stance: Auditing Negation Sensitivity in Moral Dilemmas**

ArXiv CS.AI 50%

arXiv:2601.21433v2 Announce Type: replace Abstract: Language models are increasingly consulted on ethically consequential questions, yet the stance a model expresses may not survive a change in framing. We audit 16 models across 14 ethically fraught dilemmas using...

<https://arxiv.org/abs/2601.21433>

00 **CONF****AI Chatbot Suicide Risk Detection and Response: Human Validation Study of the Open-Source VERA-MH Safety Evaluation**

ArXiv CS.AI 50%

arXiv:2602.05088v4 Announce Type: replace Abstract: Millions of people now use generative AI chatbots for psychological support. Despite their promise, the most pressing question in AI for mental health is whether these tools are safe. The field currently lacks a validated,...

<https://arxiv.org/abs/2602.05088>

01 CONF

SOMtime the World Ain'\$t Fair: Violating Fairness Using Self-Organizing Maps

ArXiv CS.AI 50%

arXiv:2602.18201v2 Announce Type: replace Abstract: Unsupervised representations are widely assumed to be neutral with respect to sensitive attributes when those attributes are withheld from training. We show that this assumption is false. Using SOMtime, a topology-preserving...

<https://arxiv.org/abs/2602.18201>

02 CONF

Power and Limitations of Aggregation in Compound AI Systems

ArXiv CS.AI 50%

arXiv:2602.21556v2 Announce Type: replace Abstract: When designing compound AI systems, a common approach is to query multiple copies of the same model and aggregate the responses to produce a synthesized output. Given the homogeneity of these models, this raises the question of...

<https://arxiv.org/abs/2602.21556>

03 CONF

Participatory provenance as representational auditing for AI-mediated public consultation

ArXiv CS.AI 50%

arXiv:2604.20711v2 Announce Type: replace Abstract: Artificial intelligence is increasingly deployed to synthesize large-scale public input in policy consultations and participatory processes. Yet no formal framework exists for auditing whether these summaries faithfully...

<https://arxiv.org/abs/2604.20711>

04 CONF

Refine Thought: A Test-Time Inference Method for Embedding Model Reasoning

ArXiv CS.AI 50%

arXiv:2511.13726v2 Announce Type: replace-cross Abstract: We propose RT (Refine Thought), a method that can enhance the semantic reasoning ability of text embedding models. The method obtains the final semantic representation by running multiple forward passes of the text...

<https://arxiv.org/abs/2511.13726>

05 CONF

Thinking Ahead: Foresight Intelligence in MLLMs and World Model

ArXiv CS.AI 50%

arXiv:2511.18735v3 Announce Type: replace-cross Abstract: In this work, we define Foresight Intelligence as the capability to anticipate and interpret future events-an ability essential for applications such as autonomous driving, yet largely overlooked by existing research. To...

<https://arxiv.org/abs/2511.18735>

06 CONF

FDRMFL: Multimodal Federated Feature Extraction Model Based on Information Maximization and Contrastive Learning

ArXiv CS.AI 50%

arXiv:2512.02076v2 Announce Type: replace-cross Abstract: We propose FDRMFL, a task-driven multimodal feature extraction framework for federated regression under non-IID data distributions. Extracting predictive features from high-dimensional multimodal inputs is particularly...

<https://arxiv.org/abs/2512.02076>

07 CONF

Named-Entity Recognition in the Crime Domain (CrimeNER): Case Study and Dataset

ArXiv CS.AI 50%

arXiv:2603.02150v2 Announce Type: replace-cross Abstract: The extraction of critical information from crime-related documents is a crucial task for law enforcement agencies. The extraction of this information can be interpreted as a Named-Entity Recognition (NER) task. However,...

<https://arxiv.org/abs/2603.02150>

08 CONF

Exploration of Fast-Slow Latent Recurrence for Train-Short, Test-Long Generalization

ArXiv CS.AI 50%

arXiv:2604.01577v3 Announce Type: replace-cross Abstract: We study out of distribution generalization in streaming tasks where models are trained on short sequences but must operate over much longer, unknown horizons under bounded memory. Our focus is on a persistent fast slow...

<https://arxiv.org/abs/2604.01577>

09 CONF

Diversity Without Fidelity: A Solver-Sampler Mismatch in Multi-Agent LLM Negotiation Simulation

ArXiv CS.AI 50%

arXiv:2604.11840v3 Announce Type: replace-cross Abstract: Language models are increasingly used to simulate people: survey respondents, negotiators, stakeholders in policy exercises. In that role a model should reproduce how people plausibly behave, hesitating, conceding late,...

<https://arxiv.org/abs/2604.11840>

10 CONF

CogAdapt: Adapting Clinical ECG Foundation Models for Wearable Cognitive Load Assessment

ArXiv CS.AI 50%

arXiv:2605.22774v4 Announce Type: replace-cross Abstract: Assessing cognitive load continuously and at low latency would help adaptive human-computer interaction, but it remains hard because labeled data are scarce and models generalize poorly across subjects. Recent ECG...

<https://arxiv.org/abs/2605.22774>

11 CONF

Trust, but Don't Verify: Epistemic Blind Spots in LLM Source Evaluation

ArXiv CS.AI 50%

arXiv:2606.05403v2 Announce Type: replace-cross Abstract: Language models increasingly act as epistemic proxies, synthesizing evidence from multiple sources to inform decisions. Whether they evaluate the quality of that evidence, or merely aggregate it based on surface...

<https://arxiv.org/abs/2606.05403>

12 CONF

The Signs Were Always There: Training-Free Concept Detection and Steering in Raw Transformer Dimensions

ArXiv CS.AI 50%

arXiv:2606.12629v3 Announce Type: replace-cross Abstract: The standard basis of transformer hidden states is a training-free, architecture-general feature basis for detecting concepts and, in language models, steering them; with no learned dictionary. Individual dimensions act...

<https://arxiv.org/abs/2606.12629>

13 **CONF****Can Trustless Agents Be Trusted? An Empirical Study of the ERC-8004 Decentralized AI Agent Ecosystem****ArXiv CS.AI 50%**

arXiv:2606.26028v2 Announce Type: replace-cross Abstract: As autonomous AI agents increasingly transact across organizational boundaries, a fundamental trust challenge emerges: how can an agent assess whether an unknown counterpart is trustworthy? The ERC-8004 protocol addresses...

<https://arxiv.org/abs/2606.26028>