



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	97	Media
TechCrunch AI	18	Media
The Verge AI	8	Media
HackerNews	7	Community
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 5 stories

VERI

Fidji Simo steps down from leading OpenAI's AGI work due to illness

The Verge AI

+1 source

80%

OpenAI's Fidji Simo is departing her full-time role as the company's AGI chief and is transitioning to being a "part-time advisor," she said on X. The news follows Simo's original announcement in April that she would take a few weeks of medical leave due to a neuroimmune...

<https://www.theverge.com/ai-artificial-intelligence/963738/openai-fidji-simo-steps-down...>

VERI

Meta says its new AI model is ready to compete on coding

The Verge AI

+1 source

80%

After reentering the AI race with its first in-house Muse Spark model in April, Meta is now opening up the doors to developers with a new model that can plug into AI coding software with the new Meta Model API. Meta says that Muse Spark 1.1 is a "step-change" from the first...

<https://www.theverge.com/ai-artificial-intelligence/963193/meta-muse-spark-model-api>

VERI

Google will now tell you if an ad was made with AI

The Verge AI

+1 source

80%

You can see if ads on Google Search, Google Discover, and YouTube were made or edited using AI from a new section in Google's "My Ad Center," as reported earlier by TechCrunch. The update, announced on Thursday, adds a "created or edited with AI" label under the "how this ad was..."

<https://www.theverge.com/ai-artificial-intelligence/963628/google-ai-generated-ads-label>

VERI

Anthropic found a hidden space where Claude puzzles over concepts

MIT Tech Review

+1 source

80%

The AI firm Anthropic has developed a technique that has given it the clearest glimpse yet at what's really going on inside large language models as they answer questions or carry out tasks. What they found ranges from the mundane to the unnerving. Researchers at the company...

<https://www.technologyreview.com/2026/07/09/1140293/anthropic-found-a-hidden-space-wher...>

VERI

Character.AI wants a piece of the microdrama pie

The Verge AI

+1 source

80%

Character.AI's plan to become more than just an LLM-powered chatbot platform is going beyond interactive books, comics, and audio dramas. Today, the company announced the debut of c.ai Series - short-form, episodic videos designed to be watched and interacted with - on your...

<https://www.theverge.com/entertainment/962897/character-ai-series-microdrama-vertical-v...>

CONFIRMED SIGNALS — 121 stories

CONF

Show HN: Follow London Trains in 3D

HackerNews

60%

HN 151

deck.gl based visualiser of the TFL Api + National Rails to be able to track (with minimal drift) a train along the way in London and to the nearest airports. If you pick one from any platform in <https://nexttrain.london> basically you can share your train journey along the...

<https://ride.nexttrain.london/>

CONF

Show HN: FableCut - A browser video editor AI agents can drive (zero deps)

HackerNews 50% HN 89

I run a small AI lab and playground and got super excited about Anthropics paper "Verbalizable Representations Form a Global Workspace in Language Models" (<https://transformer-circuits.pub/2026/workspace/index.html>) It talks about how they use a tool they call a Jacobian Lens...

<https://github.com/ronak-create/FableCut>

CONF

Show HN: Policy enforcement for Claude Code, Cursor, and Codex

HackerNews 50% HN 12

Show HN: Runtime authorization for Claude Code, Cursor, and Codex Hi HN, Fernando and I built Kastr. Kastr intercepts AI agent tool calls and evaluates them against deterministic policies before they execute. This is aimed at developers using coding agents like Claude Code,...

<https://kastra.ai/>

CONF

Show HN: We beat Gemini Embedding 2 by training only 16M params (open weights)

HackerNews 50% HN 7

No summary — see link for full article.

<https://huggingface.co/EximiusLabs/fusion-embedding-1-2b-preview>

0 CONF

Show HN: Strata - guess the city from its real 3D terrain and bathymetry

HackerNews 50% HN 6

Launched Blocks.ai today - the control plane and network layer for agents. One outbound connect, zero-trust security, and A2A protocols built in. Connect your agent to any frontend you build or any specialist agents you need regardless of runtime without tons of messy...

<https://flightq.app/strata>

1 CONF

Show HN: Agentic FC - a football management SIM played by AI agents over MCP

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/gaemi/agentic-fc>

2 CONF

Show HN: Subagentmaxxing - Orchestrate GPT and Cursor Models with Claude Code

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/Kuberwastaken/subagentmaxxing>

3 CONF

OpenAI says GPT 5.6 is the 'preferred model' for Microsoft Copilot 365 amid breakup chatter

TechCrunch AI 50%

OpenAI's new family of models will continue to power Microsoft's suite of workplace and productivity apps.

<https://techcrunch.com/2026/07/09/openai-says-gpt-5-6-is-the-preferred-model-for-micros...>

4 CONF

An AI agent startup just let its agent run its \$100M fundraise

TechCrunch AI 50%

Lyzr, a startup that builds AI agents for enterprises, used its own AI agent to raise a \$100 million round — proof, evidently, that the product actually works.

<https://techcrunch.com/2026/07/09/an-ai-agent-startup-just-let-its-agent-run-its-100-mi...>

5 CONF

OpenAI is shutting down Atlas, but its AI browser ambitions are still growing

TechCrunch AI 50%

OpenAI is sunsetting its AI-powered browser after less than a year. But it's moving some agentic browsing features to its desktop app and a Chrome extension.

<https://techcrunch.com/2026/07/09/openai-is-shutting-down-atlas-but-its-ai-browser-ambi...>

6 CONF

Elon Musk praises Mythos/Fable, promises not to 'cut off' Anthropic

TechCrunch AI 50%

Should Anthropic trust Elon Musk to host its models? With about \$40 billion in revenue at stake, Musk insists that the company can.

<https://techcrunch.com/2026/07/09/elon-musk-praises-mythos-fable-promises-not-to-cut-of...>

7 CONF

Can AI answer the \$3 trillion question?

TechCrunch AI 50%

The AI ROI debate has returned and the numbers are even bigger, as are, perhaps, the consequences.

<https://techcrunch.com/2026/07/09/can-ai-answer-the-3-trillion-question/>

8 CONF

Meta enters the crowded AI coding battle with Muse Spark 1.1

TechCrunch AI 50%

Meta's pitch to users is Spark's ability to handle large agentic workloads, fix bugs, and help with large code migrations — the kind of automation that enterprises are increasingly turning to AI companies to provide.

<https://techcrunch.com/2026/07/09/meta-enters-the-crowded-ai-coding-battle-with-muse-sp...>

9 CONF

New York Times says OpenAI hid evidence in ChatGPT copyright trial

TechCrunch AI 50%

News publishers say OpenAI hid tools and datasets that could identify copyrighted journalism in ChatGPT outputs, escalating their lawsuit with a new motion for sanctions.

<https://techcrunch.com/2026/07/09/new-york-times-says-openai-hid-evidence-in-chatgpt-co...>

0 CONF

Paris-based AI voice startup Gradium raises \$100M seed, backed by Nvidia

TechCrunch AI 50%

The company is using the cash to open an office in the Bay Area and compete for talent there, "strengthening its position at the heart of the world's leading AI ecosystem."

<https://techcrunch.com/2026/07/09/paris-based-ai-voice-startup-gradium-raises-100m-seed...>

1 CONF

How did the government decide OpenAI's frontier model was safe to release?

TechCrunch AI 50%

"Exactly what that dialog looked like between the government and Anthropic and OpenAI is unclear."

<https://techcrunch.com/2026/07/09/how-did-the-government-decide-openais-frontier-model-...>

2 CONF

Instagram users: Here's how to stop Meta's AI from using your photos

TechCrunch AI 50%

Muse Image allows users to generate AI images using photos from public Instagram accounts. As long as a person's profile is public, another user can tag that account and use their images as part of an AI-generated creation.

<https://techcrunch.com/2026/07/09/how-to-stop-metas-ai-image-generator-from-using-your-...>

3 CONF

Meta's new AI chips will begin production in September

TechCrunch AI 50%

The company is taking a modular approach to designing these chips, anticipating that their needs will change as AI evolves rapidly by the time the chips are in production.

<https://techcrunch.com/2026/07/09/metanew-ai-chips-will-begin-production-in-september/>

4 CONF

Anthropic, OpenAI, and SpaceX are bigger than the last 25 years of tech exits

TechCrunch AI 50%

Three big AI IPOs are set to generate more value than all the U.S. VC-backed exits since 2000.

<https://techcrunch.com/2026/07/09/anthropic-openai-and-spacex-are-bigger-than-the-last-25-years-of-tech-exits/>

5 CONF

Popular open source AI developer tool Ollama raises \$65M, grows to nearly 9M users

TechCrunch AI 50%

Benchmark-backed Ollama has amassed 176,000 stars, and nearly 17,000 forks on GitHub by helping developers easily run AI on their PCs.

<https://techcrunch.com/2026/07/09/popular-open-source-ai-developer-tool-ollama-raises-65m-grows-to-nearly-9m-users/>

6 CONF

The ChatGPT browser is already dead

The Verge AI 50%

OpenAI is already shutting down ChatGPT Atlas, its browser that could do tasks for you on your behalf, less than a year after launching it. Atlas was announced in October, but as part of its wave of news about ChatGPT Work today, the company confirmed that it will be...

<https://www.theverge.com/ai-artificial-intelligence/963654/openai-chatgpt-atlas-ai-browser-is-dead>

7 CONF

OpenAI rolls out GPT-5.6 after government greenlight — and announces 'ChatGPT Work'

The Verge AI 50%

About two weeks after OpenAI's GPT-5.6 was caught up in regulatory drama - rolled out only to government-approved organizations during a "limited preview" period - the company has received the Trump administration's greenlight for a public rollout of the model. OpenAI CEO Sam...

<https://www.theverge.com/ai-artificial-intelligence/963464/openai-gpt-5-6-codex-chatgpt-work>

8 CONF

Say hello to Claude Wrapped

The Verge AI 50%

The popularity of Spotify Wrapped has kicked off a wide range of year-in-review features, on apps from YouTube to Uber - and now, the lookback trend has come to AI. Anthropic on Thursday announced a "reflect" feature for its Claude chatbot, allowing users to see an analysis of...

<https://www.theverge.com/ai-artificial-intelligence/963105/anthropic-claude-wrapped-reflection>

9 CONF

FL Studio 2026 turns its AI chatbot into your assistant engineer

The Verge AI 50%

Last year, Image Line introduced Gopher for FL Studio, an AI chatbot that was basically a glorified instruction manual. You asked it how to do something, and it would serve up the relevant instructions. It's the kind of thing I actually use AI for on a semi-regular basis. But in...

<https://www.theverge.com/tech/963052/fl-studio-2026-music-daw-ai-chatbot>

0 CONF

AI-integrated models for assessing agricultural resilience

ArXiv CS.AI 50%

arXiv:2607.07759v1 Announce Type: new Abstract: Agricultural supply chains are vulnerable to disruptions through linked biophysical and economic systems. We develop an AI-powered tool that integrates economic models (GTAP) with biophysical models (APSIM) to analyze supply chain...

<https://arxiv.org/abs/2607.07759>

1 CONF

Adversarial Social Epistemology for Assemblies of Humans and Large Language Models

ArXiv CS.AI 50%

arXiv:2607.07760v1 Announce Type: new Abstract: We outline an adversarial social epistemology (ASE) for densely interactive communicative landscapes in which public assertions are scaffolded by chains of testimony, inference, institutional certification, and tacit trust. In such...

<https://arxiv.org/abs/2607.07760>

2 CONF

Aligning Clinical Needs and AI Capabilities: A Survey on LLMs for Medical Reasoning

ArXiv CS.AI 50%

arXiv:2607.07761v1 Announce Type: new Abstract: Large language models (LLMs) have emerged as important tools in healthcare, showing growing potential for clinical reasoning and patient care. This survey examines recent progress in medical LLMs, focusing on reasoning applications...

<https://arxiv.org/abs/2607.07761>

3 CONF

Alignment Plausibility: A New Standard for Assuring AI in Healthcare

ArXiv CS.AI 50%

arXiv:2607.07766v1 Announce Type: new Abstract: Large language models (LLMs) have become significant providers of mental health support, yet they remain products of an attention economy whose operational and commercial targets favour sustained engagement over the friction that...

<https://arxiv.org/abs/2607.07766>

4 CONF

VectorizationLLM: Smart Vectorization Based AI Assistant

ArXiv CS.AI 50%

arXiv:2607.07846v1 Announce Type: new Abstract: VectorizationLLM is a specialized Large Language Model based on Google open-weight LLMs. The model is designed to assist students to learn smart vectorization, time/wave vector analysis, piecewise functions, Fourier analysis, and...

<https://arxiv.org/abs/2607.07846>

5 CONF

A Graph Neural Network Model for Real-Time Gesture Recognition Based on sEMG Signals

ArXiv CS.AI 50%

arXiv:2607.07850v1 Announce Type: new Abstract: For seamless control of advanced hand prostheses and augmented reality, accurate and immediate hand gestures recognition is essential. Surface electromyography (sEMG) signals obtained from the forearm are commonly employed for this...

<https://arxiv.org/abs/2607.07850>

6 CONF

Agentic AI and Retrieval-Augmented Models in Straight-Through Underwriting

ArXiv CS.AI 50%

arXiv:2607.07858v1 Announce Type: new Abstract: Artificial intelligence (AI) is beginning to reshape actuarial practice, particularly in domains that require reasoning over unstructured documents, heterogeneous data sources, and regulated decision workflows. Actuaries now face a...

<https://arxiv.org/abs/2607.07858>

7 CONF

Feedback Manipulation Regularization: Enabling Offline Agent Alignment for Imitation Learning

ArXiv CS.AI 50%

arXiv:2607.07859v1 Announce Type: new Abstract: Reinforcement learning (RL) research has increasingly shifted focus towards alignment, ensuring agents learn behaviors adhering to human values. While human demonstrations and feedback have proven crucial for alignment, existing...

<https://arxiv.org/abs/2607.07859>

8 CONF

Nigeria Machinery: A Low-Resource Industrial Dataset with a Domain-Grounded Reasoning Layer

ArXiv CS.AI 50%

arXiv:2607.07883v1 Announce Type: new Abstract: There is relatively little, public, and model-ready data on industrial machinery for African economies. This makes it hard to do quantitative analysis or to train language models on numeric tasks grounded in that setting. We...

<https://arxiv.org/abs/2607.07883>

9 CONF

Persona Cartography: Charting Language Model Personality Traits in Weight Space

ArXiv CS.AI 50%

arXiv:2607.07916v1 Announce Type: new Abstract: Large language models exhibit recurring behavioural patterns -- personas -- that shape generalisation and safety, but we lack reliable tools for decomposing, measuring, and controlling them. Our central insight is to treat personas...

<https://arxiv.org/abs/2607.07916>

10 CONF

Concretized Proposition Prompting Resolves Composition-Knowledge Dichotomy in Large Language Models

ArXiv CS.AI 50%

arXiv:2607.08018v1 Announce Type: new Abstract: LLMs often struggle to balance compositionality with knowledgeability, a challenge we define as Composition-Knowledge Dichotomy. To address this, we propose Concretized Proposition Prompting (CPP), a framework that explicitly...

<https://arxiv.org/abs/2607.08018>

11 CONF

From Prompts to Contracts: Harness Engineering for Auditable Enterprise LLM Agents

ArXiv CS.AI 50%

arXiv:2607.08028v1 Announce Type: new Abstract: Enterprise large language model (LLM) applications often begin as prototypes whose behavior is carried by prompts and retrieval context. Productization adds requirements for source boundaries, entity routing, answer contracts, and...

<https://arxiv.org/abs/2607.08028>

12 CONF

A safety-oriented hypothetico-deductive framework for AI-assisted differential diagnosis

ArXiv CS.AI 50%

arXiv:2607.08038v1 Announce Type: new Abstract: Diagnostic error is a major threat to patient safety, yet current large language model (LLM) systems often treat diagnosis as a one-shot prediction task, lacking safeguards against missed high-risk alternatives or rigorous...

<https://arxiv.org/abs/2607.08038>

3 CONF

When LLMs Agree, Are They Right? Auditing Self-Consistency and Cross-Model Agreement as Confidence Signals

ArXiv CS.AI 50%

arXiv:2607.08065v1 Announce Type: new Abstract: LLM-as-judge (Zheng et al., 2023) is increasingly the default for evaluating AI systems in enterprise pipelines, often scaled to ensembles (Verga et al., 2024) or "mixture-of-experts" (Shazeer et al., 2017) panels of judges. These...

<https://arxiv.org/abs/2607.08065>

4 CONF

PARA-PV: Physics-Aware Retrieval-Augmented PV Prediction Based on Frozen Foundation Model and Distribution Shift Correction

ArXiv CS.AI 50%

arXiv:2607.08079v1 Announce Type: new Abstract: Accurate photovoltaic (PV) power forecasting is essential for reliable grid dispatch and renewable energy integration, yet it remains challenging because PV generation is jointly shaped by weather variability, day-night...

<https://arxiv.org/abs/2607.08079>

5 CONF

CausalIDS: Benchmarking Causal Reasoning in Data-Science Agents

ArXiv CS.AI 50%

arXiv:2607.08093v1 Announce Type: new Abstract: Large language models (LLMs) increasingly act as integrated data-science agents, combining abstract reasoning with advanced tool use. Yet the relevant benchmark landscape largely divides into symbolic causal reasoning benchmarks...

<https://arxiv.org/abs/2607.08093>

6 CONF

ASMR: Agentic Schema Generation for Ship Maintenance Report Writing

ArXiv CS.AI 50%

arXiv:2607.08177v1 Announce Type: new Abstract: In this paper, we study the automatic schema generation problem: given a collection of historical ship maintenance and operational reports across multiple form categories, automatically discover compact and informative schemas that...

<https://arxiv.org/abs/2607.08177>

7 CONF

Playing ZendoWorld: Challenging AI Agents on Active Visual Concept Induction

ArXiv CS.AI 50%

arXiv:2607.08233v1 Announce Type: new Abstract: A central challenge in building intelligent systems is enabling agents to jointly perceive complex inputs, form hypotheses about hidden patterns, and design informative experiments to test them. To study this problem, we propose...

<https://arxiv.org/abs/2607.08233>

8 CONF

Compete Then Collaborate: Frontier AI Teachers Build a Verifiable Curriculum to Improve a Coding Student Beyond Imitation

ArXiv CS.AI 50%

arXiv:2607.08255v1 Announce Type: new Abstract: Large language models increasingly serve as teachers generating training data for smaller students. Prior multi-teacher knowledge distillation methods merge outputs without determining which frontier model teaches best, often...

<https://arxiv.org/abs/2607.08255>

9 CONF

PolyUQuest: Verifiable Structure-Aware Web RAG over Heterogeneous Graphs

ArXiv CS.AI 50%

arXiv:2607.08269v1 Announce Type: new Abstract: Existing retrieval-augmented generation (RAG) systems treat web pages as flat text, losing the structural and semantic signals encoded in HTML. We present PolyUQuest, a verifiable, structure-aware web RAG framework built on a...

<https://arxiv.org/abs/2607.08269>

0 CONF

Psychological Competence as a Missing Dimension in AI Evaluation

ArXiv CS.AI 50%

arXiv:2607.08285v1 Announce Type: new Abstract: Current AI evaluation frameworks focus primarily on technical performance, including accuracy, robustness, reasoning ability, and policy compliance. These measures remain essential, but they are not sufficient for systems that...

<https://arxiv.org/abs/2607.08285>

1 CONF

Blind-Spots-Bench: Evaluating Blind Spots in Multimodal Models

ArXiv CS.AI 50%

arXiv:2607.08317v1 Announce Type: new Abstract: Modern AI models achieve strong performance on many established benchmarks, yet they still fail on tasks that humans find almost trivial, such as manipulating a string or drawing a dog with five legs. These examples suggest that...

<https://arxiv.org/abs/2607.08317>

2 CONF

MobiDiff: Semantic-Aware Multi-Channel Discrete Diffusion for Human Mobility Data Generation

ArXiv CS.AI 50%

arXiv:2607.08357v1 Announce Type: new Abstract: Human mobility data are essential for transportation optimization, urban planning, and resource allocation, yet real-world mobility data are costly to collect and difficult to share due to privacy concerns. Recent diffusion-based...

<https://arxiv.org/abs/2607.08357>

3 CONF

Towards Mechanistically Understanding Why Memorized Knowledge Fails to Generalize in Large Language Model Finetuning

ArXiv CS.AI 50%

arXiv:2607.08393v1 Announce Type: new Abstract: Fine-tuning LLMs to inject new knowledge faces a critical challenge: LLMs can quickly memorize new facts, yet fail to use them for downstream reasoning tasks. We formalize this failure as the $\text{Knowing--Using Gap}$,...

<https://arxiv.org/abs/2607.08393>

4 CONF

Game Theory Driven Multi-Agent Framework Mitigates Language Model Hallucination

ArXiv CS.AI 50%

arXiv:2607.08403v1 Announce Type: new Abstract: The application of lightweight Large Language Models in rule-based scientific domains remains severely limited by their tendency to mimic linguistic patterns rather than reproduce axiomatic reasoning, causing frequent...

<https://arxiv.org/abs/2607.08403>

5 CONF

AI-guided stimuli discovery and generation to optimize facial emotion perception studies in autism

ArXiv CS.AI 50%

arXiv:2607.08533v1 Announce Type: new Abstract: Understanding perceptual differences between autistic and neurotypical adults requires behavioral assays that are sensitive, reliable, and mechanistically informative. Facial emotion perception is a useful test case because group...

<https://arxiv.org/abs/2607.08533>

6 CONF

CommuniWave:A Machine Learning Model for Quantifying the Degree of Temporary Informal Behavior in Urban Communities

ArXiv CS.AI 50%

arXiv:2607.08554v1 Announce Type: new Abstract: For urban managers and designers, improving the functional attributes of urban communities to enhance territorial resilience in the face of complexity and uncertainty is crucial. Currently, community planning often follows a...

<https://arxiv.org/abs/2607.08554>

7 CONF

Towards Precision Therapy in Hepatocellular Carcinoma: A Clinical-Reasoning LLM for Risk Stratification and Treatment Guidance

ArXiv CS.AI 50%

arXiv:2607.08602v1 Announce Type: new Abstract: Hepatocellular carcinoma (HCC) is a common malignancy and a leading cause of cancer-related mortality. Current guidelines and staging systems provide coarse categories, but often miss within-stage heterogeneity and the clinical...

<https://arxiv.org/abs/2607.08602>

8 CONF

The complexities of patient-centred conversational artificial intelligence

ArXiv CS.AI 50%

arXiv:2607.08625v1 Announce Type: new Abstract: Consumer-facing health chatbots powered by large language models (LLMs) are increasingly used for symptom assessment. However, chatbot development and evaluation often rely on cooperative, articulate, simulated patients. We...

<https://arxiv.org/abs/2607.08625>

9 CONF

SolarChain-Eval: A Physics-Constrained Benchmark for Trustworthy Economic Agents in Decentralized Energy Markets

ArXiv CS.AI 50%

arXiv:2607.08681v1 Announce Type: new Abstract: As agentic AI systems are increasingly applied to cyber-physical environments, their evaluation requires assessment of both task performance and trustworthiness. In decentralized energy markets, autonomous agents may improve market...

<https://arxiv.org/abs/2607.08681>

0 CONF

The Contribution of XAI for the Safe Development and Certification of AI: An Expert-Based Analysis

ArXiv CS.AI 50%

arXiv:2408.02379v2 Announce Type: replace-cross Abstract: Developing and certifying safe - or so-called trustworthy - AI has become an increasingly salient issue, especially in light of upcoming regulation such as the EU AI Act. In this context, the black-box nature of machine...

<https://arxiv.org/abs/2408.02379>

1 CONF

AUTOPILOT VQA: Benchmarking Vision-Language Models for Incident-Centric Dashcam Understanding

ArXiv CS.AI 50%

arXiv:2607.08745v1 Announce Type: new Abstract: Recent advances in Vision-Language Models, Large Language Models, and Multimodal Large Language Models have improved autonomous driving tasks such as scene understanding, decision making, trajectory prediction, and visual question...

<https://arxiv.org/abs/2607.08745>

2 CONF

Using AI-based Learning Assistants in Higher Education: A Large-Scale Descriptive Analysis

ArXiv CS.AI 50%

arXiv:2607.08748v1 Announce Type: new Abstract: In this study, we present a large-scale descriptive analysis of the use of an AI-based learning assistant (Syntea) in higher education. Based on objective log data from 77,543 students enrolled in distance studies, we examine usage...

<https://arxiv.org/abs/2607.08748>

3 CONF

Towards the Explainability of Temporal Graph Networks via Memory Backtracking and Topological Attribution

ArXiv CS.AI 50%

arXiv:2607.07716v1 Announce Type: cross Abstract: Temporal graphs are ubiquitous in real-world applications and Temporal Graph Networks (TGNs) have achieved superior predictive accuracy. Understanding which historical events drive model predictions can enhance trustworthiness of...

<https://arxiv.org/abs/2607.07716>

4 CONF

When Implausible Tokens Get Reinforced: Tail-Aware Credit Calibration for LLM Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.07976v1 Announce Type: cross Abstract: Reinforcement learning (RL) has achieved remarkable success in enhancing the reasoning capabilities of large language models (LLMs). However, widely used critic-free RL methods rely on uniform credit assignment, broadcasting the...

<https://arxiv.org/abs/2607.07976>

5 CONF

Omni-Sleep: A Sleep Foundation Model via Hierarchical Contrastive Learning of CNS--ANS Dynamic

ArXiv CS.AI 50%

arXiv:2607.07720v1 Announce Type: cross Abstract: Sleep physiology arises from the coordinated dynamics of the central nervous system (CNS) and autonomic nervous system (ANS), as reflected by multimodal polysomnography signals including EEG, EOG, EMG, ECG, and respiration....

<https://arxiv.org/abs/2607.07720>

6 CONF

Collective Intelligence with Foundation Models

ArXiv CS.AI 50%

arXiv:2607.07729v1 Announce Type: cross Abstract: As foundation models grow in scale and diversity, coordinating multiple models into cooperative reasoning systems offers a path toward safer, more reliable AI. This chapter presents a multi-agent framework where solver models...

<https://arxiv.org/abs/2607.07729>

7 CONF

A Transdiagnostic Space of Disorder Like Phenotypes in Reinforcement Learning Agents

ArXiv CS.AI 50%

arXiv:2607.07753v1 Announce Type: cross Abstract: Modelling psychological disorders in artificial agents offers both a testbed for computational psychiatry and a lens on the failure modes of affective control. Prior work induces one or two disorders in a reinforcement learning...

<https://arxiv.org/abs/2607.07753>

8 CONF

Graph-Regularized Deep Learning for EEG-Based Emotion Recognition with Psychologically-Grounded Label Structure

ArXiv CS.AI 50%

arXiv:2607.07773v1 Announce Type: cross Abstract: EEG-based emotion recognition is critical for mental health monitoring and affective brain-computer interfaces, yet existing deep learning approaches often treat emotion classes as isolated labels, ignoring their psychological...

<https://arxiv.org/abs/2607.07773>

9 CONF

From Solvers to Research: Large Language Model-Driven Formal Mathematics at the Research Frontier

ArXiv CS.AI 50%

arXiv:2607.07779v1 Announce Type: cross Abstract: Recent developments in AI for Mathematics (AI4Math), especially Large Language Model (LLM)-driven theorem provers, has achieved remarkable success in formal proof generation for well-defined mathematical problems through...

<https://arxiv.org/abs/2607.07779>

0 CONF

DreamCharacter-1: From 3D Generative Foundation Models to Product-Ready Character Generation

ArXiv CS.AI 50%

arXiv:2607.07817v1 Announce Type: cross Abstract: We present DreamCharacter-1, a lightweight post-adaptation framework that calibrates pretrained 3D foundation models toward high-fidelity, production-ready 3D character generation. Building upon a 3D foundation backbone, our...

<https://arxiv.org/abs/2607.07817>

1 CONF

From Triggers to Emotions: A CPM-Grounded Appraisal Multi-Agent for Dynamic Emotional Evolution in Persona-Based Dialogue

ArXiv CS.AI 50%

arXiv:2607.07824v1 Announce Type: cross Abstract: Large Language Models (LLMs) have substantially advanced persona-based dialogue agents for emotion-sensitive role simulation in healthcare, education, counseling, customer service, and interactive storytelling. However, two...

<https://arxiv.org/abs/2607.07824>

2 CONF

Shift & Drift: A Zero-Shot Benchmark for Generalizable and Robust Autonomous Driving Motion Planning

ArXiv CS.AI 50%

arXiv:2607.07844v1 Announce Type: cross Abstract: While closed-loop motion planners trained on large-scale, object-level datasets, e.g., nuPlan, demonstrate strong in-distribution (ID) performance, their generalization to novel urban topologies and recovery mechanisms following...

<https://arxiv.org/abs/2607.07844>

3 CONF

Kime-Representation Formulations of Three Open Problems in the Foundations of Classical Mechanics: Uncertainty, Invariant Entropy, and Directional Degrees of Freedom

ArXiv CS.AI 50%

arXiv:2607.07851v1 Announce Type: cross Abstract: We give mathematically self-contained formulations, in the complex-time (kime) representation, of three open problems from the foundations of classical mechanics: (I) the extension of the classical entropic uncertainty principle...

<https://arxiv.org/abs/2607.07851>

4 CONF

Time-to-Collision Based Dynamic Obstacle Avoidance Using Pretrained Vision Models for Robots in Unstructured Environments

ArXiv CS.AI 50%

arXiv:2607.07885v1 Announce Type: cross Abstract: Dynamic obstacle avoidance in unstructured outdoor environments remains a critical challenge for autonomous mobile robots, particularly when large-scale robot-specific training data and simulation-based policies are impractical....

<https://arxiv.org/abs/2607.07885>

5 CONF

How Do I Know What to Say Next? Barenholtz's Autogenerative Theory as an Enrichment of Harrisean Integrationism

ArXiv CS.AI 50%

arXiv:2607.07891v1 Announce Type: cross Abstract: Roy Harris's Integrationist linguistics offers a compelling critique of the referentialist tradition embedded deep at the heart of computational approaches to language, arguing that language is not a code that maps onto a...

<https://arxiv.org/abs/2607.07891>

6 CONF

Closed-Loop Dynamic Validator Node Scaling in Private Substrate Blockchains Using Takagi-Sugeno Fuzzy Inference

ArXiv CS.AI 50%

arXiv:2607.07901v1 Announce Type: cross Abstract: Private blockchain networks run with fixed node configurations that cannot adapt to changing workload conditions. Too many nodes serving a light workload waste resources; too few nodes facing heavy demand slow block production...

<https://arxiv.org/abs/2607.07901>

7 CONF

Multimodal Unlearning Across Vision, Language, Video, and Audio: Survey of Methods, Datasets, and Benchmarks

ArXiv CS.AI 50%

arXiv:2607.07907v1 Announce Type: cross Abstract: With the growing adoption of VLMs, DMs, LLMs, and AFMs, these multimodal foundation models can inadvertently encode sensitive, copyrighted, biased, or unsafe cross-modal associations that originate from their training data....

<https://arxiv.org/abs/2607.07907>

8 CONF

A Multi-cluster Boundary Learning Method for Out-of-Scope Intent Detection via MiniLM Embedding

ArXiv CS.AI 50%

arXiv:2607.07974v1 Announce Type: cross Abstract: Intent detection is a critical task that bridges human intents and system actions in human-machine interaction systems. However, there still exist challenges for detecting out-of-scope (OOS) intents. (i) The traditional methods...

<https://arxiv.org/abs/2607.07974>

- 9 **CONF**
3100 Opinions on Code Review in an AI World: Building Causal Theory from Practitioner Discourse
 ArXiv CS.AI 50%
 arXiv:2607.07980v1 Announce Type: cross Abstract: Coding agents now author entire pull requests, and practitioners sharply disagree about what this does to code review: whether it becomes the bottleneck, whether human review is still necessary, and whether it quietly erodes the...
<https://arxiv.org/abs/2607.07980>
- 0 **CONF**
Who Broke the System? Failure Localization in LLM-Based Multi-Agent Systems
 ArXiv CS.AI 50%
 arXiv:2607.07989v1 Announce Type: cross Abstract: Large language model (LLM) based multi-agent systems enable complex problem solving through coordinated reasoning and action, but their distributed structure also introduces new challenges in diagnosing system-level failures....
<https://arxiv.org/abs/2607.07989>
- 1 **CONF**
Can We Trust LLM's Logic? Quantifying Uncertainty, Coherence, and Robustness via a Graph-Based Framework
 ArXiv CS.AI 50%
 arXiv:2607.08017v1 Announce Type: cross Abstract: Large-Language Models (LLMs) can be prone to flawed and unfaithful reasoning that decoding strategies like Self-Consistency (SC) fail to detect as they evaluate only final-answer agreement while ignoring the logical validity of...
<https://arxiv.org/abs/2607.08017>
- 2 **CONF**
Structured Pruning of Large Language Models via Power Transformation and Sign-Preserving Score Aggregation with Adaptive Feature Retention
 ArXiv CS.AI 50%
 arXiv:2607.08027v1 Announce Type: cross Abstract: This paper proposes an improved structured pruning method for large language models (LLMs) that addresses key challenges in adapting Adaptive Feature Retention (AFR), an unstructured pruning technique, to structured pruning. When...
<https://arxiv.org/abs/2607.08027>
- 3 **CONF**
DKDNet: Dual Knowledge and Data-Driven Network for Cross-Domain Automatic Modulation Classification
 ArXiv CS.AI 50%
 arXiv:2607.08031v1 Announce Type: cross Abstract: The dynamics of communication environments induce significant distribution shifts across domains, challenging the generalization of deep learning-based automatic modulation classification (AMC) models. While existing UDA methods...
<https://arxiv.org/abs/2607.08031>
- 4 **CONF**
Leveraging Color Naming for Image Enhancement
 ArXiv CS.AI 50%
 arXiv:2607.08185v1 Announce Type: cross Abstract: Enhancing images to make them visually appealing is a persistent challenge in computer vision. Many deep-learning methods train models on paired datasets to replicate expert editing styles. However, these approaches struggle with...
<https://arxiv.org/abs/2607.08185>

5 **CONF****Aleena: Alignment Agent for Research Software Engineering Collaborations**

ArXiv CS.AI 50%

arXiv:2607.08043v1 Announce Type: cross Abstract: Research software collaborations span meetings, informal chats, pull requests, and GitHub issues. A decision surfaced in a Slack thread, refined in a meeting, and implemented in a pull request can lose its original rationale...

<https://arxiv.org/abs/2607.08043>

6 **CONF****What LLM Forecasters Know but Don't Say: Probing Internal Representations for Calibration and Faithfulness**

ArXiv CS.AI 50%

arXiv:2607.08046v1 Announce Type: cross Abstract: Large language models fine-tuned for forecasting can be accurate yet poorly calibrated, and their chain-of-thought (CoT) reasoning may not faithfully reflect the evidence behind a forecast. We ask whether internal representations...

<https://arxiv.org/abs/2607.08046>

7 **CONF****Who Analyses the Analyser? Self-Validating LLM Hazard Analysis with Constitutional Meta-STPA**

ArXiv CS.AI 50%

arXiv:2607.08054v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly trusted to draft the artifacts of safety analysis such as, losses, hazards, Unsafe Control Actions (UCAs), and safety constraints, inside rigorous processes such as Systems-Theoretic...

<https://arxiv.org/abs/2607.08054>

8 **CONF****Towards Efficient Large Language Model Serving: A Survey on System-Aware KV Cache Optimization**

ArXiv CS.AI 50%

arXiv:2607.08057v1 Announce Type: cross Abstract: Despite the rapid advancements of large language models (LLMs), LLM serving systems remain memory-intensive and costly. The key-value (KV) cache, which stores KV tensors during autoregressive decoding, is crucial for enabling...

<https://arxiv.org/abs/2607.08057>

9 **CONF****COBART: Controlled, Optimized, Bidirectional and Auto-Regressive Transformer for Ad Headline Generation**

ArXiv CS.AI 50%

arXiv:2607.08071v1 Announce Type: cross Abstract: Online ads are essential to all businesses and ad headlines are one of their core creative component. Existing methods can generate headlines automatically and also optimize their click-through-rate (CTR) and quality. However,...

<https://arxiv.org/abs/2607.08071>

0 **CONF****Deep Learning Method for Stationary Distribution of Reflected Brownian Motion**

ArXiv CS.AI 50%

arXiv:2607.08091v1 Announce Type: cross Abstract: The stationary distribution of reflected Brownian motion (RBM) plays an important role in the analysis of high-dimensional stochastic systems, yet closed-form solutions are known only for a few special cases. Computing important...

<https://arxiv.org/abs/2607.08091>

1 CONF

PS4: Proxy-Supervised Joint Training for Real Target Speaker Extraction

ArXiv CS.AI 50%

arXiv:2607.08111v1 Announce Type: cross Abstract: Training target speaker extraction (TSE) models for real conversational mixtures remains challenging because large-scale training corpora and clean target speech for supervision are unavailable. We present PS4, a proxy-supervised...

<https://arxiv.org/abs/2607.08111>

2 CONF

ICDAR 2026 HIPE-OCRepair Competition on LLM-Assisted OCR Post-Correction for Historical Documents

ArXiv CS.AI 50%

arXiv:2607.08143v1 Announce Type: cross Abstract: We present the results of HIPE-OCRepair-2026, an ICDAR competition on LLM-assisted OCR post-correction of historical documents. OCR post-correction remains a long-standing challenge in digital heritage: large-scale collections of...

<https://arxiv.org/abs/2607.08143>

3 CONF

Out of Sight: Compression-Aware Content Protection against Agentic Crawlers

ArXiv CS.AI 50%

arXiv:2607.08180v1 Announce Type: cross Abstract: The rise of LLM-based agents with reasoning, summarization, and memory capabilities has created a new threat surface for online content that conventional defenses fail to address. Existing defenses like access controls can be...

<https://arxiv.org/abs/2607.08180>

4 CONF

Open-ended Multi-agent Autocurricula via Visual Inspection of Policies with Multi-modal LLMs

ArXiv CS.AI 50%

arXiv:2607.08193v1 Announce Type: cross Abstract: Open-ended curricula in Reinforcement Learning (RL) aim to train generally-capable agents by identifying tasks that facilitate learning increasingly complex skills. A major challenge when designing such curricula is assessing...

<https://arxiv.org/abs/2607.08193>

5 CONF

TMI: Text-to-Image Meets Image-to-Image for Complementary Data Synthesis to Boost Long-Tailed Instance Segmentation

ArXiv CS.AI 50%

arXiv:2607.08201v1 Announce Type: cross Abstract: Large-vocabulary instance segmentation is constrained by long-tailed category distributions and fine-grained inter-class ambiguity. While data synthesis offers a promising alternative, current paradigms have complementary...

<https://arxiv.org/abs/2607.08201>

6 CONF

Best-of-\$N\$ TTS Evaluation is Confounded by ASR Family Alignment

ArXiv CS.AI 50%

arXiv:2607.08256v1 Announce Type: cross Abstract: Best-of-\$N\$ (BoN) inference improves content consistency in zero-shot text-to-speech by selecting from \$N\$ candidates with an automatic speech recognition (ASR) verifier. We identify an underexplored evaluation confound: a...

<https://arxiv.org/abs/2607.08256>

7 CONF

Multi-Agent Firewall Architecture for Privacy Protection of Sensitive Data in Interactions with Language Models

ArXiv CS.AI 50%

arXiv:2607.08282v1 Announce Type: cross Abstract: While Large Language Models (LLMs) have become essential productivity tools, their integration into workflows without adequate safeguards creates significant risks. This paper proposes an open-source, privacy-focused, user-facing...

<https://arxiv.org/abs/2607.08282>

8 CONF

Towards Isolated Interventions via Almost Orthogonal Features in Language Models

ArXiv CS.AI 50%

arXiv:2602.04718v2 Announce Type: replace-cross Abstract: A central premise in mechanistic interpretability is that meaningful concepts in language models are represented by linear features in activation space. For such features to support reliable interventions, manipulating...

<https://arxiv.org/abs/2602.04718>

9 CONF

ADORN: Adaptive Drift handling for Open RAN using Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.08443v1 Announce Type: cross Abstract: Dynamic traffic variations in Open Radio Access Networks (O-RAN) lead to drift, which degrades the performance of Artificial Intelligence/Machine Learning (AI/ML) models. Traditional retraining approaches maintain forecasting...

<https://arxiv.org/abs/2607.08443>

00 CONF

Large-Language-Models-as-a-Judge in Theory-Agnostic Adaptive Metric-Alignment for Prototypical Networks in Personality Recognition

ArXiv CS.AI 50%

arXiv:2607.08374v1 Announce Type: cross Abstract: Personality recognition has traditionally been constrained by theory-dependent formulations, where models are trained to fit predefined psychological taxonomies rather than uncovering shared underlying behavioral structure. This...

<https://arxiv.org/abs/2607.08374>

01 CONF

TRACE: A Two-Channel Robust Attribution Watermark via Complementary Embeddings for LLM-Agent Trajectories

ArXiv CS.AI 50%

arXiv:2607.08400v1 Announce Type: cross Abstract: LLM agents reach users through resellers, who may rebrand a developer's agent or substitute a cheaper model. When provenance is disputed, attribution rests on the trajectory log (the record of tool calls, observations, and...

<https://arxiv.org/abs/2607.08400>

02 CONF

DrugGen 2: A disease-aware language model for enhancing drug discovery

ArXiv CS.AI 50%

arXiv:2607.08404v1 Announce Type: cross Abstract: Current computational approaches for drug design typically focus on generating molecules conditioned on specific targets or general molecular properties, often neglecting the influence of disease context on target behavior and...

<https://arxiv.org/abs/2607.08404>

03 CONF

Predicting Male Fertility Using Machine Learning: A Semen Parameters Based Analysis with the VISEM Dataset

ArXiv CS.AI 50%

arXiv:2607.08429v1 Announce Type: cross Abstract: Male infertility is a significant yet often underdiagnosed aspect of reproductive health, with semen analysis serving as the cornerstone of clinical evaluation. To address this problem, this study investigates the use of machine...

<https://arxiv.org/abs/2607.08429>

04 CONF

Two Axes of LLM Abstention: Answer Correctness and Question Answerability

ArXiv CS.AI 50%

arXiv:2607.08456v1 Announce Type: cross Abstract: A model should refuse two different things: answers it would get wrong, and questions it should not answer at all, such as unanswerable ones or ones resting on a false premise. The usual recipe thresholds a single confidence...

<https://arxiv.org/abs/2607.08456>

05 CONF

Cognitive-structured Multimodal Agent for Multimodal Understanding, Generation, and Editing

ArXiv CS.AI 50%

arXiv:2607.08497v1 Announce Type: cross Abstract: Recent unified multimodal models show a single architecture can jointly perform vision/language understanding and image generation/editing. However, they repeatedly feed all historical visual and textual inputs into a shared...

<https://arxiv.org/abs/2607.08497>

06 CONF

When the Judge Changes, So Does the Measurement: Auditing LLM-as-Judge Reliability

ArXiv CS.AI 50%

arXiv:2607.08535v1 Announce Type: cross Abstract: An LLM-as-judge score can move even when the candidate responses stay fixed, simply because the evaluator has changed. We treat this evaluator-replacement ambiguity as a measurement-validity problem. Across four judgment...

<https://arxiv.org/abs/2607.08535>

07 CONF

VocaDet: Sample-Driven Open-Vocabulary Object Detection and Segmentation via Visual Tokenization and Vector Database Retrieval

ArXiv CS.AI 50%

arXiv:2607.08541v1 Announce Type: cross Abstract: Open-vocabulary object detection and segmentation aim to recognize arbitrary objects beyond predefined categories. Although recent vision-language and reference-based approaches have significantly advanced this field, they often...

<https://arxiv.org/abs/2607.08541>

08 CONF

UltraX: Refining Pre-Training Data at Scale with Adaptive Programmatic Editing

ArXiv CS.AI 50%

arXiv:2607.08646v1 Announce Type: cross Abstract: As available training data approaches its physical limit, gains from Scaling Laws have begun to diminish. Consequently, improving Large Language Models (LLMs) now depends less on data expansion and more on higher-quality data...

<https://arxiv.org/abs/2607.08646>

09 CONF

Validity of LLMs as data annotators: AMALIA on authority

ArXiv CS.AI 50%

arXiv:2607.08731v1 Announce Type: cross Abstract: A national language model offers a linguistic community its own instrument for measuring what its citizens say and value. Portugal's AMALIA, a publicly funded 9B-parameter model for European Portuguese, appears competitive on...

<https://arxiv.org/abs/2607.08731>

10 CONF

IFAR: Multi-Perspective and Multi-Level Causal Discovery with LLMs

ArXiv CS.AI 50%

arXiv:2409.05559v2 Announce Type: replace Abstract: Large language models (LLMs) have developed rapidly, and their reasoning capabilities have become a hot research topic. However, there is still limited exploration of abductive reasoning. The multi-perspective and multi-level...

<https://arxiv.org/abs/2409.05559>

11 CONF

A Vision Toward Energy-Efficient Domain-Specific Artificial Intelligence Models and Agents

ArXiv CS.AI 50%

arXiv:2510.22052v2 Announce Type: replace Abstract: The field of artificial intelligence (AI) has taken a tight hold on broad aspects of society, industry, business, and governance in ways that dictate the prosperity and might of the world's economies. The AI market size is...

<https://arxiv.org/abs/2510.22052>

12 CONF

Conversational AI for Rapid Scientific Prototyping: A Case Study on ESA's ELOPE Competition

ArXiv CS.AI 50%

arXiv:2601.04920v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used as coding partners, yet their role in accelerating scientific discovery remains underexplored. This paper presents a case study of using ChatGPT for rapid prototyping in ESA's...

<https://arxiv.org/abs/2601.04920>

13 CONF

LiteOdyssey: A Lightweight Reasoning AI Agent for Interpretable Rare-Disease Diagnosis

ArXiv CS.AI 50%

arXiv:2606.16149v2 Announce Type: replace Abstract: Rare disease diagnosis involves interpreting clinical and genetic findings through complex diagnostic reasoning. We investigated whether this reasoning could be translated into a portable policy for guiding general-purpose...

<https://arxiv.org/abs/2606.16149>

14 CONF

InvestPhilBench: A Multi-Layer Benchmark for Evaluating Large Language Model Procedural Reasoning in Expert Investment Philosophy

ArXiv CS.AI 50%

arXiv:2606.25984v2 Announce Type: replace Abstract: Large language models are increasingly deployed as investment research assistants, yet no benchmark tests whether they can accurately reconstruct and apply the specific procedural decision frameworks of expert investors. We...

<https://arxiv.org/abs/2606.25984>

15 CONF

Knowledge Graph and Accurate Portrait Construction of Scientific and Technological Academic Conferences

ArXiv CS.AI 50%

arXiv:2204.04888v2 Announce Type: replace-cross Abstract: In recent years, with the continuous progress of science and technology, the number of scientific research achievements has increased rapidly. As an exchange platform and medium for scientific research achievements,...

<https://arxiv.org/abs/2204.04888>

16 CONF

MultiFair: Multimodal Balanced Fairness-Aware Medical Classification with Dual-Level Gradient Modulation

ArXiv CS.AI 50%

arXiv:2510.07328v2 Announce Type: replace-cross Abstract: Medical decision systems increasingly rely on data from multiple sources to ensure reliable and unbiased diagnosis. However, existing multimodal learning models fail to achieve this goal because they often overlook two...

<https://arxiv.org/abs/2510.07328>

17 CONF

SwinIFS: Landmark Guided Swin Transformer For Identity Preserving Face Super Resolution

ArXiv CS.AI 50%

arXiv:2601.01406v2 Announce Type: replace-cross Abstract: Face super-resolution aims to recover high-quality facial images from severely degraded low-resolution inputs, but remains challenging due to the loss of fine structural details and identity-specific features. This work...

<https://arxiv.org/abs/2601.01406>

18 CONF

GenDA: Generative Data Assimilation on Complex Urban Areas via Classifier-Free Diffusion Guidance

ArXiv CS.AI 50%

arXiv:2601.11440v3 Announce Type: replace-cross Abstract: Urban wind flow reconstruction is essential for assessing air quality, heat dispersion, and pedestrian comfort, yet remains challenging when only sparse sensor data are available. We propose GenDA, a generative data...

<https://arxiv.org/abs/2601.11440>

19 CONF

Rethinking LLM-as-a-Judge: Representation-as-a-Judge with Small Language Models via Semantic Capacity Asymmetry

ArXiv CS.AI 50%

arXiv:2601.22588v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are widely used as reference-free evaluators via prompting, but this "LLM-as-a-Judge" paradigm is costly, opaque, and sensitive to prompt design. In this work, we investigate whether smaller...

<https://arxiv.org/abs/2601.22588>

20 CONF

Self-EvolveRec: Self-Evolving Recommender Systems with LLM-based Directional Feedback

ArXiv CS.AI 50%

arXiv:2602.12612v2 Announce Type: replace-cross Abstract: Traditional methods for automating recommender system design, such as Neural Architecture Search (NAS), are often constrained by a fixed search space defined by human priors, limiting innovation to pre-defined operators....

<https://arxiv.org/abs/2602.12612>

21 CONF

Curriculum Learning for Efficient Chain-of-Thought Distillation via Structure-Aware Masking and GRPO

ArXiv CS.AI 50%

arXiv:2602.17686v4 Announce Type: replace-cross Abstract: Distilling Chain-of-Thought (CoT) reasoning from large language models into compact student models presents a fundamental challenge: teacher rationales are often too verbose for smaller models to faithfully reproduce....

<https://arxiv.org/abs/2602.17686>

22 CONF

Curvature-Weighted Capacity Allocation: A Minimum Description Length Framework for Layer-Adaptive Large Language Model Optimization

ArXiv CS.AI 50%

arXiv:2603.00910v2 Announce Type: replace-cross Abstract: Layer-wise capacity in large language models is highly non-uniform: some layers contribute disproportionately to loss reduction, whereas others are nearly redundant. Existing layer-scoring methods provide sensitivity...

<https://arxiv.org/abs/2603.00910>

23 CONF

Robust Weighted Triangulation of Causal Effects Under Model Uncertainty

ArXiv CS.AI 50%

arXiv:2603.01119v2 Announce Type: replace-cross Abstract: A fundamental challenge in causal inference with observational data is correct specification of a causal model. When there is model uncertainty, analysts may seek to use estimates from multiple candidate models that rely...

<https://arxiv.org/abs/2603.01119>

24 CONF

Training and Evaluating Diffusion Policies with Long Context Lengths

ArXiv CS.AI 50%

arXiv:2606.16447v2 Announce Type: replace-cross Abstract: Imitation learning has enabled highly-dexterous robotic manipulation from RGB observations. Policies trained with these methods, however, typically condition robot actions on only a short history of observations. These...

<https://arxiv.org/abs/2606.16447>

25 CONF

Hierarchical Control in Multi-Agent Games: LLM-based Planning and RL Execution

ArXiv CS.AI 50%

arXiv:2606.20014v3 Announce Type: replace-cross Abstract: Reinforcement learning (RL) has achieved strong performance in sequential decision-making, yet scaling to complex multi-agent environments remains challenging due to sparse rewards, large state-action spaces, and the...

<https://arxiv.org/abs/2606.20014>

26 CONF

LLM for the development of FCM

ArXiv CS.AI 50%

arXiv:2607.04983v2 Announce Type: replace-cross Abstract: This article is about the development of a fuzzy cognitive map using a local large language model. In the light of recent advances it is evident that large language models, and even local large language models are capable...

<https://arxiv.org/abs/2607.04983>