



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	72	Media
HackerNews	5	Community
The Verge AI	4	Media
TechCrunch AI	2	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

A Sovereign, Open-Source Foundation Model for German and English

ArXiv CS.AI +1 source

80%

arXiv:2607.09424v1 Announce Type: cross Abstract: We present Soofi S 30B-A3B, a sovereign, open-source Mixture-of-Experts (MoE) hybrid Mamba Transformer foundation model for German and English. Its hybrid design activates only 3B of 30B parameters per token and keeps the...

<https://arxiv.org/abs/2607.09424>

VERI

Knowledge Graphs and Explainable AI as Complementary Resources for Urban Mining

ArXiv CS.AI +1 source

80%

arXiv:2607.09578v1 Announce Type: new Abstract: Pre-demolition assessment, the regulated audit process at the heart of urban mining, is an information process in which AI support must serve qualified auditors who remain accountable for the decisions taken. The relevant unit of...

<https://arxiv.org/abs/2607.09578>

VERI

The Patchwork Problem in LLM-Generated Code

ArXiv CS.AI +1 source

80%

arXiv:2607.08981v1 Announce Type: cross Abstract: LLM-generated code often compiles, passes tests, and appears correct, yet breaks once deployed. The root cause is frequently structural rather than logical. A generated endpoint references configuration keys never declared in the...

<https://arxiv.org/abs/2607.08981>

CONFIRMED SIGNALS — 77 stories

CONF

Show HN: SayItDev LLM and Speech capabilities with 0 dependencies and no models

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/innovatorved/sayitdev>

CONF

Show HN: Kote - Capture and reuse engineering context from AI chats and Git

HackerNews 50% HN 3

I kept running into the same problem: I'd solve something with the help of an AI assistant, spend time debugging an issue, or make an architectural decision... and a few weeks later I couldn't remember where that information was. Most of the time I simply didn't bother saving it...

<https://github.com/pedroaugusto04/Kote>

CONF

Show HN: Tenclicker - a WebLLM clicker game about whatever you want

HackerNews 50% HN 2

No summary — see link for full article.

<https://tenclicker.com>

CONF

Show HN: A subjective AI eval. Arcade games built by AI

HackerNews 50% HN 2

No summary — see link for full article.

<https://ai-arcade.app>

CONF

Show HN: Isitsecure - 1-command SAST and DAST and LLM security scanner for web apps

HackerNews 50% HN 2

As vibe coding has started to become more mainstream, so has the realization that typical AI-generated code is insecure by default. What started as a free security scan tool for vibe coded apps turned into an open source project for a much deeper Static Application Security Test...

<https://github.com/jaurakunal/isitsecure>

CONF

Google will now disclose which ads are made with AI

TechCrunch AI 50%

While Google prohibits misleading and deceptive ads, an ad can still leverage AI to create some type of synthetic or digitally altered content. Until now, that's something Google only required election ads to disclose.

<https://techcrunch.com/2026/07/09/google-will-now-disclose-which-ads-are-made-with-ai/>

0 CONF

Lorde says Ray-Ban Meta AI glasses are 'not sexy'

The Verge AI 50%

Lorde was performing at the Real Cool Festival in Madrid on Thursday and took some time during her set to speak out against AI glasses. While she didn't specify any brands in particular, it's likely she was taking a shot at festival sponsor Ray-Ban, which has collaborated with...

<https://www.theverge.com/ai-artificial-intelligence/964539/lorde-says-ray-ban-meta-ai-g...>

1 CONF

Apple's failed self-driving car program left a legacy of powerful AI chips

The Verge AI 50%

Apple's self-driving car program never really got off the ground, but it may have been what made the company's chips the powerful AI performers they are. Early in the development of the self-driving platform, Apple realized that it would need powerful on-device AI processing....

<https://www.theverge.com/tech/964519/apple-silicon-self-driving-car-ai-m7-ultra>

2 CONF

A Unified Approach to Interpreting Knowledge Distillation for Large Language Models via Interactions

ArXiv CS.AI 50%

arXiv:2607.08776v1 Announce Type: cross Abstract: Despite the success of knowledge distillation (KD) in Large Language Models (LLMs), the underlying mechanism behind its efficacy remains unclear. In this paper, we propose a unified approach to explore the common mechanism of...

<https://arxiv.org/abs/2607.08776>

3 CONF

A Formalization of the Mean-Field Derivation of the Vlasov Equation: AI-Assisted Lean Formalization as a Strategy Game

ArXiv CS.AI 50%

arXiv:2607.08986v1 Announce Type: new Abstract: We formalize a research result in the Lean 4 proof assistant by having a mathematician direct an AI system, and frame the activity as a formalization game. The objective is to turn a LaTeX document into Lean. The game is won when...

<https://arxiv.org/abs/2607.08986>

- 4 **CONF**
- Neuro-Agentic Control: A Deep Learning-based LLM-Powered Agentic AI Framework for Controlling Security Controls**
- ArXiv CS.AI 50%
- arXiv:2607.09076v1 Announce Type: new Abstract: Cyberattacks on operational technology are increasingly causing costly downtime and physical damage, exposing the limitations of traditional rule-based monitoring in industrial IoT environments. While Large Language Models (LLMs)...
- <https://arxiv.org/abs/2607.09076>
- 5 **CONF**
- MedRealMM: A Real-World Multimodal Benchmark for Chinese Online Medical Consultation**
- ArXiv CS.AI 50%
- arXiv:2607.09142v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly deployed in online medical consultation, yet existing benchmarks remain poorly aligned with real clinical practice. Many rely on synthetic conversations or patient simulators, omit...
- <https://arxiv.org/abs/2607.09142>
- 6 **CONF**
- Toward Auditable AI Scientists: A Hypothesis Evolution Protocol for LLM Agents**
- ArXiv CS.AI 50%
- arXiv:2607.09195v1 Announce Type: new Abstract: Large language model (LLM) agents are increasingly expected to play a central role in AI-driven scientific discovery. Equipped with broad knowledge, flexible reasoning, and tool use, they have the potential to autonomously explore...
- <https://arxiv.org/abs/2607.09195>
- 7 **CONF**
- WILDTRACE: Benchmarking Natural Evidence Trails in Long-Context Reasoning**
- ArXiv CS.AI 50%
- arXiv:2607.09328v1 Announce Type: cross Abstract: Answering complex questions over long documents frequently requires integrating evidence that the source itself disperses naturally across distant passages. In an incident report, the operating condition, design flaw, and missed...
- <https://arxiv.org/abs/2607.09328>
- 8 **CONF**
- Communication-Efficient Digital-Twin Coordination for Heterogeneous LLM Embodied Agents over Computing Power Networks**
- ArXiv CS.AI 50%
- arXiv:2607.09330v1 Announce Type: new Abstract: Embodied agent teams powered by heterogeneous large language models (LLMs) are being widely deployed in physical artificial intelligence such as smart factories, warehouses, and service robotics. To enable collaboration among such...
- <https://arxiv.org/abs/2607.09330>
- 9 **CONF**
- Fictional Worldbuilding: Multi-Agent LLM Collaboration with Hierarchical Context Compression and Iterative Review**
- ArXiv CS.AI 50%
- arXiv:2607.09403v1 Announce Type: new Abstract: Worldbuilding, the construction of coherent fictional worlds, is a foundational task in game design and literary creation. Large Language Models (LLMs) offer new possibilities for automated content generation, but their application...
- <https://arxiv.org/abs/2607.09403>

- 0 **CONF**
- How Does Bayesian Causal Discovery Fail? Characterising Structural Consequences in Linear Gaussian Networks under Latent Confounding**
- ArXiv CS.AI 50%
- arXiv:2607.09449v1 Announce Type: new Abstract: Bayesian causal discovery is widely used for its ability to quantify epistemic uncertainty over directed acyclic graphs (DAGs) through posterior inference. However, its behaviour under latent confounding remains poorly understood,...
- <https://arxiv.org/abs/2607.09449>
-
- 1 **CONF**
- QAgent: An LLM-based Multi-Agent System for Autonomous OpenQASM programming**
- ArXiv CS.AI 50%
- arXiv:2508.20134v2 Announce Type: replace Abstract: Programming quantum circuits at the OpenQASM level is essential for achieving hardware-aware optimization and reliable execution on noisy intermediate-scale quantum (NISQ) devices, yet it remains challenging due to the need for...
- <https://arxiv.org/abs/2508.20134>
-
- 2 **CONF**
- Ceci n'est pas une pipe: AI systems as semantic abstractions**
- ArXiv CS.AI 50%
- arXiv:2607.09489v1 Announce Type: new Abstract: An AI system's output is not the fact or world state it appears to describe, but rather an engineered representation. We propose a semantic framework to describe AI systems, to be able to examine the correctness of such...
- <https://arxiv.org/abs/2607.09489>
-
- 3 **CONF**
- Shortcut Trajectory Planning for Efficient Offline Reinforcement Learning**
- ArXiv CS.AI 50%
- arXiv:2607.09336v1 Announce Type: cross Abstract: Diffusion-based trajectory planners have shown strong performance in offline reinforcement learning, but their iterative denoising process often incurs high inference cost. Consistency-based planners reduce the number of sampling...
- <https://arxiv.org/abs/2607.09336>
-
- 4 **CONF**
- Self-Guided Test-Time Training for Long-Context LLMs**
- ArXiv CS.AI 50%
- arXiv:2607.09415v1 Announce Type: cross Abstract: Long-context processing has become increasingly important for large language models (LLMs), but simply extending the context window does not guarantee effective utilization of long inputs. As input length grows, accuracy often...
- <https://arxiv.org/abs/2607.09415>
-
- 5 **CONF**
- SAGEAgent: A Self-Evolving Agent for Cost-Aware Modality Acquisition in Multimodal Survival Prediction**
- ArXiv CS.AI 50%
- arXiv:2607.09521v1 Announce Type: new Abstract: Does every cancer patient truly need a complete diagnostic workup for accurate survival prediction? In multimodal clinical oncology, diagnostic modalities follow a clinically mandated order of escalating burden -- from demographics...
- <https://arxiv.org/abs/2607.09521>
-

6 CONF

NL-PAC: Specification Ambiguity and Certified Minimax Risk Floors in LLM-Mediated Supervision

ArXiv CS.AI 50%

arXiv:2607.08961v1 Announce Type: cross Abstract: Large language models increasingly provide labels, evaluations, and feedback for tasks specified in natural language. When a specification admits multiple readings but the supervision channel does not reveal which is operative,...

<https://arxiv.org/abs/2607.08961>

7 CONF

A Personalized Computational Framework for Assessing the Sufficiency of Partially Observed Data in Healthcare AI models

ArXiv CS.AI 50%

arXiv:2607.09165v1 Announce Type: cross Abstract: Achieving early and timely diagnosis and treatment for disease is a major challenge. Recent applications of machine learning (ML) algorithms trained on patient data have shown promise in many different settings for predicting the...

<https://arxiv.org/abs/2607.09165>

8 CONF

Agora: Enhancing LLM Agent Reasoning Via Auction-Based Task Allocation

ArXiv CS.AI 50%

arXiv:2607.09600v1 Announce Type: new Abstract: Enhancing the reasoning capabilities of large language model (LLM) agents requires effective orchestration of diverse expert models and tools. However, existing frameworks typically call APIs based on coarse-grained matching...

<https://arxiv.org/abs/2607.09600>

9 CONF

ConceptSMILE: Auditing the Trustworthiness of Concept-Based Explainable AI

ArXiv CS.AI 50%

arXiv:2607.09649v1 Announce Type: new Abstract: Concept-based explainable artificial intelligence (AI) can make model reasoning more human-understandable, but concept-level outputs are not automatically trustworthy. We introduce ConceptSMILE, a model-agnostic perturbation-based...

<https://arxiv.org/abs/2607.09649>

0 CONF

REFORGE: A Method for Benchmarking LLMs' Reverse Engineering Capabilities in Decompiled Binary Function Naming

ArXiv CS.AI 50%

arXiv:2607.07738v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly applied to reverse-engineering tasks, and recent threat-intelligence reporting shows them operating inside live offensive-security workflows. Claims about their capability, however,...

<https://arxiv.org/abs/2607.07738>

1 CONF

iLENS: Interpretable LLM-Guided Mixture-of-Experts for Neuroimaging Survival Analysis

ArXiv CS.AI 50%

arXiv:2607.08778v1 Announce Type: cross Abstract: Alzheimer's Disease (AD) is a complex neurodegenerative disorder that continues to impact millions of people worldwide. Predicting AD conversion during the prodromal stage remains critical for disease understanding and patient...

<https://arxiv.org/abs/2607.08778>

2 CONF

Sticky Routing: Training MoE Models for Memory-Efficient Inference

ArXiv CS.AI 50%

arXiv:2607.08780v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) models activate only a sparse subset of experts per token, yet consecutive tokens frequently activate different experts -- causing constant weight swapping between slow storage and fast memory on edge...

<https://arxiv.org/abs/2607.08780>

3 CONF

Reward Transport: Property Control in Flow Matching via Noise-Space Alignment

ArXiv CS.AI 50%

arXiv:2607.08781v1 Announce Type: cross Abstract: The coupling in flow matching -- the rule pairing noise vectors with data points -- is typically treated as a computational choice. We show that this coupling can instead serve as an alignment interface: by matching noise and...

<https://arxiv.org/abs/2607.08781>

4 CONF

HERO: A Heterogeneity-Aware Benchmark Library for Federated Continual Learning

ArXiv CS.AI 50%

arXiv:2607.08784v1 Announce Type: cross Abstract: Federated continual learning (FCL) evaluates how distributed clients learn from changing data streams while retaining previously learned knowledge. Existing evaluations are difficult to compare because they often change datasets,...

<https://arxiv.org/abs/2607.08784>

5 CONF

Augmenting Fundamental Analysis with Large Language Models: A RAG-Based System for Generating Investor Briefs

ArXiv CS.AI 50%

arXiv:2607.09121v1 Announce Type: cross Abstract: In this study, we examine the opportunities brought by Large Language Models (LLMs) to various aspects of fundamental analysis of companies based on their reports as well as data and documents describing macroeconomic situation...

<https://arxiv.org/abs/2607.09121>

6 CONF

LLM-Driven Evolutionary Generation of Multi-Objective Bayesian Optimization Algorithms

ArXiv CS.AI 50%

arXiv:2607.08791v1 Announce Type: cross Abstract: Designing effective multi-objective Bayesian optimization (MOBO) algorithms requires balancing many interdependent design choices whose optimal configuration is problem-dependent and typically demands deep expertise. We extend...

<https://arxiv.org/abs/2607.08791>

7 CONF

EHR-MPC: Inference-Time Control for Sepsis Treatment with Generative Patient Digital Twins

ArXiv CS.AI 50%

arXiv:2607.08793v1 Announce Type: cross Abstract: Sepsis is a leading cause of mortality, yet optimal treatment policies remain contested. Existing reinforcement learning (RL) approaches learn fixed strategies for sepsis treatment, limiting adaptability to changing clinical...

<https://arxiv.org/abs/2607.08793>

8 CONF

Multi-Attribute Steering of Language Models via Targeted Intervention

ArXiv CS.AI 50%

arXiv:2502.12446v3 Announce Type: replace-cross Abstract: Inference-time intervention (ITI) has emerged as a promising method for steering large language model (LLM) behavior in a particular direction (e.g., improving helpfulness) by intervening on token representations without...

<https://arxiv.org/abs/2502.12446>

9 CONF

TheBioCollection: Unified Pre-Training Scale LLM Corpus for Biology

ArXiv CS.AI 50%

arXiv:2607.08803v1 Announce Type: cross Abstract: The push toward large language models for biology (BioLM) has created a need for training corpora that can endow models with a genuine understanding of biology. However, existing biological resources, such as molecular databases,...

<https://arxiv.org/abs/2607.08803>

0 CONF

Eluna: An Agentic LLM System for Automating Warehouse Operations with Reasoning and Task Execution

ArXiv CS.AI 50%

arXiv:2607.08960v1 Announce Type: cross Abstract: Warehouse operations are governed by Standard Operating Procedures (SOPs) that encode complex, multi-system decision logic, which must be executed reliably under strict time constraints, yet LLM agents lack mechanisms to enforce...

<https://arxiv.org/abs/2607.08960>

1 CONF

MultiView-Bench: A Diagnostic Benchmark for World-Centric Multi-View Integration in VLMs

ArXiv CS.AI 50%

arXiv:2607.08970v1 Announce Type: cross Abstract: Recent benchmarks for VLMs largely assess single- or limited-view perception, leaving untested the core cognitive ability to integrate observations across viewpoints into a coherent, world-centric (allocentric) 3D mental model....

<https://arxiv.org/abs/2607.08970>

2 CONF

Correlation-Aware Contextual Bandits with Surrogate Rewards for LLM Routing

ArXiv CS.AI 50%

arXiv:2607.09015v1 Announce Type: cross Abstract: We study contextual bandit problems with correlated arms and access to surrogate reward signals produced by a machine learning model, motivated by applications such as large language model (LLM) routing. Unlike classical...

<https://arxiv.org/abs/2607.09015>

3 CONF

PRecG: Legal Precedent Retrieval with Graph Neural Networks and Rhetorical Role Segmentation

ArXiv CS.AI 50%

arXiv:2607.09094v1 Announce Type: cross Abstract: Legal precedent retrieval is a fundamental task in legal case preparation, planning, litigation strategy, and legal research. Current approaches for automatic precedent retrieval map legal documents to a low-dimensional semantic...

<https://arxiv.org/abs/2607.09094>

4 **CONF****Integrating Large Language Models and Graph Convolutional Networks for Semi-Supervised Image Classification****ArXiv CS.AI** **50%**

arXiv:2607.09104v1 Announce Type: cross Abstract: While the growing availability of image data has driven significant advances, labeling datasets remains costly and time-consuming. Therefore, semi-supervised approaches such as Graph Convolutional Networks (GCNs), which learn...

<https://arxiv.org/abs/2607.09104>

5 **CONF****Event Stream based Multi-Modal Video Anomaly Detection: A Benchmark Dataset and Algorithms****ArXiv CS.AI** **50%**

arXiv:2607.09114v1 Announce Type: cross Abstract: Video anomaly detection (VAD) is critical for automated surveillance but remains fragile under challenging conditions such as illumination variations, fast motion, and complex backgrounds when relying solely on visible light...

<https://arxiv.org/abs/2607.09114>

6 **CONF****STEEL: Sparsity-Aware Fused Attention for Energy-Efficient Long-Sequence Inference on AMD's XDNA NPU****ArXiv CS.AI** **50%**

arXiv:2607.09385v1 Announce Type: cross Abstract: The growing adoption of large language model-based agents within operating system workflows has increased the importance of energy-efficient inference on laptop-class systems-on-chip (SoCs). While cloud offloading remains common,...

<https://arxiv.org/abs/2607.09385>

7 **CONF****Attention to Detail: Evaluating Energy, Performance, and Accuracy Trade-offs Across vLLM Configurations****ArXiv CS.AI** **50%**

arXiv:2607.09172v1 Announce Type: cross Abstract: Large Language Models are reshaping how software is developed and maintained. They are typically deployed in production using inference engines such as vLLM, which can efficiently serve pre-trained, highly configurable models....

<https://arxiv.org/abs/2607.09172>

8 **CONF****Generative Communications: Overview, Technologies, and Trends****ArXiv CS.AI** **50%**

arXiv:2607.09183v1 Announce Type: cross Abstract: The groundbreaking development of generative artificial intelligence (AI) is rapidly boosting the ability to generate content such as images and videos, reshaping communication paradigms. This article introduces generative...

<https://arxiv.org/abs/2607.09183>

9 **CONF****LLMs for health: Perceived benefits, risks, intention to use AI chatbots, and willingness to self-disclose across sensitive health topics****ArXiv CS.AI** **50%**

arXiv:2607.09253v1 Announce Type: cross Abstract: AI chatbots are increasingly used for answering health-related questions. This study examines the role of topic type discussed with an AI chatbot and individual characteristics on perceived benefits and risks, intention to use an...

<https://arxiv.org/abs/2607.09253>

0 CONF

Blockchain-Linked Auditable Decision Management for Telecom/IoT Fraud-Control Requests

ArXiv CS.AI 50%

arXiv:2607.09259v1 Announce Type: cross Abstract: Telecom fraud-control studies often stop at detector-level classification, but deployment use requires request-level policy resolution, lifecycle traceability, and auditability. This paper reframes fraud control as...

<https://arxiv.org/abs/2607.09259>

1 CONF

Geopolitical alignment: Endorsement effects in large language models

ArXiv CS.AI 50%

arXiv:2607.09262v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to summarize and evaluate policy-relevant information, but it remains unclear whether their judgments are implicitly shaped by geopolitical cues. I study this question with an...

<https://arxiv.org/abs/2607.09262>

2 CONF

Automatic Thematic Indexing of Large Literary Corpora: A Machine Learning Approach to Voltaire's Complete Works

ArXiv CS.AI 50%

arXiv:2607.09316v1 Announce Type: cross Abstract: Thematic indexing -- the practice of assigning structured conceptual labels to sections of text -- is essential to scholarly access in large-scale literary and historical editions, yet it remains a largely manual,...

<https://arxiv.org/abs/2607.09316>

3 CONF

Letting the Data Speak: Extracting Keywords from Crowdsourced Collections with AI

ArXiv CS.AI 50%

arXiv:2607.09324v1 Announce Type: cross Abstract: Identifying and assigning keywords at scale is a technical, practical, and ethical challenge for crowdsourced collections. This article reports the findings of the "Extracting Keywords from Crowdsourced Collections" project,...

<https://arxiv.org/abs/2607.09324>

4 CONF

Deceptive Grounding: Entity Attribution Failure in Clinical Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2607.09349v1 Announce Type: cross Abstract: Retrieval-augmented generation evaluation checks whether model claims are factually grounded in retrieved documents. It does not check whether retrieved evidence is attributed to the correct entity. A clinical RAG response can...

<https://arxiv.org/abs/2607.09349>

5 CONF

When Routes Run Out: Adversarial Co-Learning and Explainable Robustness in Quantum Repeater Networks

ArXiv CS.AI 50%

arXiv:2607.09378v1 Announce Type: cross Abstract: We study an adversarial bandit problem for entanglement-based quantum-network routing over a modest graph corpus. Alice selects an end-to-end repeater route for an Ekert-91 protocol (E91) representing her move, while Eve selects...

<https://arxiv.org/abs/2607.09378>

- 6
CONF
SVF-CR: Synchronized Visual-Facial Cross-Refinement for Multimodal Ambivalence and Hesitancy Recognition
ArXiv CS.AI 50%

arXiv:2607.09417v1 Announce Type: cross Abstract: Ambivalence and hesitancy are subtle behavioral states that are expressed through a combination of verbal content, facial behavior, visual context, and acoustic cues. Effective recognition therefore requires not only extracting...

<https://arxiv.org/abs/2607.09417>
- 7
CONF
Practical Source Code Recovery from Binary Functions Using Anchor-Based Retrieval and LLM Reasoning
ArXiv CS.AI 50%

arXiv:2607.09452v1 Announce Type: cross Abstract: We present a practical pipeline for recovering source code from stripped binary functions by combining reverse engineering, anchor-based source code retrieval, and large language model reasoning. Our binary-to-source-code...

<https://arxiv.org/abs/2607.09452>
- 8
CONF
All Explanations are Wrong, But Many Are Useful: Exploring the Rashomon Explanation Set with Large Language Models
ArXiv CS.AI 50%

arXiv:2607.09502v1 Announce Type: cross Abstract: Explaining machine-learning models is increasingly important for decision-making and consumer trust, yet it is widely believed to come at a cost: existing Explainable AI (XAI) methods suffer from a persistent...

<https://arxiv.org/abs/2607.09502>
- 9
CONF
What VGGT Knows About Overlap: Probing Geometric Foundation Models for Co-Visibility
ArXiv CS.AI 50%

arXiv:2607.09503v1 Announce Type: cross Abstract: A fundamental challenge in 3D reconstruction and robotic localization is co-visibility: determining which image pairs share overlapping visible surfaces, particularly in scenarios with minimal overlap. We demonstrate that VGGT...

<https://arxiv.org/abs/2607.09503>
- 0
CONF
Failure as a Process: An Anatomy of CLI Coding Agent Trajectories
ArXiv CS.AI 50%

arXiv:2607.09510v1 Announce Type: cross Abstract: Large language model (LLM) coding agents are increasingly deployed to autonomously perform software engineering tasks in terminal-based environments, making their reliability a growing concern. Existing empirical studies...

<https://arxiv.org/abs/2607.09510>
- 1
CONF
Seeing is Free, Speaking is Not: Uncovering the True Energy Bottleneck in Edge VLM Inference
ArXiv CS.AI 50%

arXiv:2607.09520v1 Announce Type: cross Abstract: Vision-Language Models (VLMs) are the perceptual backbone of embodied AI, but their energy footprint on edge hardware remains poorly understood. Existing efficiency efforts focus predominantly on reducing visual tokens,...

<https://arxiv.org/abs/2607.09520>

2 **CONF****ALICE: Learning a General-Purpose Pathology Foundation Model from Vision, Vision-Language, and Slide-Level Experts****ArXiv CS.AI** **50%**

arXiv:2607.09526v1 Announce Type: cross Abstract: Foundation models are reshaping computational pathology, yet their capabilities remain shaped by pretraining objectives, data sources, and spatial scales, fragmenting complementary expertise across separate backbones. Here we...

<https://arxiv.org/abs/2607.09526>

3 **CONF****PAC-ACT: Post-training Actor-Critic for Action Chunking Transformers****ArXiv CS.AI** **50%**

arXiv:2607.09590v1 Announce Type: cross Abstract: Precision industrial contact manipulation requires reliable robot policies under pose perturbations and contact-force constraints. Vision-language-action models offer broad generalization but often introduce high inference...

<https://arxiv.org/abs/2607.09590>

4 **CONF****The LLMbda Calculus: AI Agents, Conversations, and Information Flow****ArXiv CS.AI** **50%**

arXiv:2602.20064v2 Announce Type: replace-cross Abstract: Large language models are increasingly deployed as agents: they plan, call tools, read untrusted data, and act on the results. This exposes them to prompt injection: data meant only to be read is obeyed as an instruction....

<https://arxiv.org/abs/2602.20064>

5 **CONF****VEXAIoT: Autonomous IoT Vulnerability Exploitation using AI Agents****ArXiv CS.AI** **50%**

arXiv:2607.09653v1 Announce Type: cross Abstract: Internet of Things (IoT) systems are inherently vulnerable due to constrained hardware, outdated firmware, and insecure default configurations, creating a need for scalable and adaptive security testing approaches. While recent...

<https://arxiv.org/abs/2607.09653>

6 **CONF****Evolution of Accuracy and Visual-Cognitive Errors in a Decade of Vision-Language AI Models****ArXiv CS.AI** **50%**

arXiv:2607.09654v1 Announce Type: cross Abstract: Vision language models (VLMs) have made remarkable progress in visual reasoning during the last decade. Most evaluations have used simple scenes (MS-COCO) that do not showcase complex human interactions or behaviors, only a...

<https://arxiv.org/abs/2607.09654>

7 **CONF****Leveraging Multi-Agent System (MAS) and Fine-Tuned Small Language Models (SLMs) for Automated Telecom Network Troubleshooting****ArXiv CS.AI** **50%**

arXiv:2511.00651v2 Announce Type: replace Abstract: Telecom networks are rapidly growing in scale and complexity, making effective management, operation, and optimization increasingly challenging. Although Artificial Intelligence (AI) has been applied to many telecom tasks,...

<https://arxiv.org/abs/2511.00651>

8 CONF

Programming over Thinking: Efficient and Robust Multi-Constraint Planning

ArXiv CS.AI 50%

arXiv:2601.09097v4 Announce Type: replace Abstract: Multi-constraint planning involves identifying, evaluating, and refining candidate plans while satisfying multiple, potentially conflicting constraints. Existing large language model (LLM) approaches face fundamental...

<https://arxiv.org/abs/2601.09097>

9 CONF

PACE: A Personalized Adaptive Curriculum Engine for 9-1-1 Call-taker Training

ArXiv CS.AI 50%

arXiv:2603.05361v2 Announce Type: replace Abstract: 9-1-1 call-taking training requires mastery of over a thousand interdependent skills, covering diverse incident types and protocol-specific nuances. A nationwide labor shortage is already straining training capacity, but...

<https://arxiv.org/abs/2603.05361>

0 CONF

Rectification Difficulty and Optimal Sample Allocation in LLM-Augmented Surveys

ArXiv CS.AI 50%

arXiv:2604.17267v2 Announce Type: replace Abstract: Large Language Models can generate synthetic survey responses at low cost, but their accuracy varies unpredictably across questions. We study the design problem of allocating a fixed budget of human respondents across...

<https://arxiv.org/abs/2604.17267>

1 CONF

Heterogeneous Information-Bottleneck Coordination Graphs for Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2605.17393v2 Announce Type: replace Abstract: Coordination graphs are a central abstraction in cooperative multi-agent reinforcement learning (MARL), yet existing sparse-graph learners lack a theoretically grounded mechanism to decide which edges should exist and how much...

<https://arxiv.org/abs/2605.17393>

2 CONF

Latent Reward Steering: An Adaptive Inference-Time Framework that Implicitly Promotes Cognitive Behaviors in Reasoning LLMs

ArXiv CS.AI 50%

arXiv:2606.00726v2 Announce Type: replace Abstract: Strong reasoning depends not only on model knowledge but also on how effectively cognitive behaviors are deployed during generation. Existing methods often rely on explicit behavior-level control, making them insufficiently...

<https://arxiv.org/abs/2606.00726>

3 CONF

GrAInS: Gradient-based Attribution for Inference-Time Steering of LLMs and VLMs

ArXiv CS.AI 50%

arXiv:2507.18043v2 Announce Type: replace-cross Abstract: Inference-time steering methods offer a lightweight alternative to fine-tuning large language models (LLMs) and vision-language models (VLMs) by modifying internal activations at test time without updating model weights...

<https://arxiv.org/abs/2507.18043>

4 **CONF****Explaining Human Choice Probabilities with Simple Vector Representations****ArXiv CS.AI 50%**

arXiv:2511.03643v3 Announce Type: replace-cross Abstract: We formalize human choice behavior in a probabilistic hide-and-seek task. In our geometric construction, vectors represent participant choice frequencies as well as probability matching and maximizing strategies. We...

<https://arxiv.org/abs/2511.03643>5 **CONF****H3Former: Hypergraph-based Semantic-Aware Aggregation via Hyperbolic Hierarchical Contrastive Loss for Fine-Grained Visual Classification****ArXiv CS.AI 50%**

arXiv:2511.10260v2 Announce Type: replace-cross Abstract: Fine-Grained Visual Classification (FGVC) remains a challenging task due to subtle inter-class differences and large intra-class variations. Existing approaches typically rely on feature-selection mechanisms or...

<https://arxiv.org/abs/2511.10260>6 **CONF****Point of Order: Action-Aware LLM Persona Modeling for Data-Grounded Civic Deliberation****ArXiv CS.AI 50%**

arXiv:2511.17813v3 Announce Type: replace-cross Abstract: LLM-based simulations can enable controlled studies of civic deliberation, but current systems lack speaker-attributed data and methods for evaluating long-form institutional behavior. ASR transcripts typically use...

<https://arxiv.org/abs/2511.17813>7 **CONF****Principles of Lipschitz continuity in neural networks****ArXiv CS.AI 50%**

arXiv:2602.04078v2 Announce Type: replace-cross Abstract: Deep learning has achieved remarkable success across a wide range of domains, significantly expanding the frontiers of what is achievable in artificial intelligence. Yet, despite these advances, critical challenges remain...

<https://arxiv.org/abs/2602.04078>8 **CONF****Preference Conditioned Multi-Objective Reinforcement Learning: Decomposed, Diversity-Driven Policy Optimization****ArXiv CS.AI 50%**

arXiv:2602.07764v2 Announce Type: replace-cross Abstract: Multi-objective reinforcement learning (MORL) seeks to train agents capable of balancing conflicting objectives. While single preference-conditioned policies offer a highly scalable solution, existing approaches remain...

<https://arxiv.org/abs/2602.07764>9 **CONF****SWE-Milestone: Evaluating AI Agents on Continuous Software Evolution****ArXiv CS.AI 50%**

arXiv:2603.13428v3 Announce Type: replace-cross Abstract: Real-world software must continuously evolve to meet ever-changing and open-ended requirements. AI agents, increasingly deployed as long-running systems, are now entrusted to drive this evolution. Yet, existing benchmarks...

<https://arxiv.org/abs/2603.13428>

0

CONF

Estimating Uncertainty from Reasoning: A Large-Scale Study of Multi- and Crosslingual MCQA Performance in LLMs

ArXiv CS.AI 50%

arXiv:2607.06327v2 Announce Type: replace-cross Abstract: Uncertainty estimation (UE) enables LLM-powered systems to recognize when to abstain, yet existing research has predominantly focused on English. We present the first large-scale evaluation of UE methods across 22...

<https://arxiv.org/abs/2607.06327>