



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	108	Media
TechCrunch AI	12	Media
HackerNews	9	Community
The Verge AI	5	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

PM-Bench: Evaluating Prospective Memory in LLM Agents

ArXiv CS.AI +1 source

80%

arXiv:2607.12385v1 Announce Type: new Abstract: A significant challenge in agentic AI is prospective memory: the ability to execute an intention at a specific future cue or state while other activities are ongoing. We introduce PM-Bench, a text-based benchmark for measuring...

<https://arxiv.org/abs/2607.12385>

VERI

Meta accused of using biased AI targeting for mass layoffs

The Verge AI +1 source

80%

A group of 26 former Meta employees is suing the company over claims that it used AI tools to unfairly target workers on leave with layoffs, as reported earlier by Reuters. In the lawsuit, the employees allege Meta determined which workers to dismiss based on performance data...

<https://www.theverge.com/tech/965486/meta-lawsuit-former-employees-ai-layoffs>

CONFIRMED SIGNALS — 131 stories

CONF

Show HN: I RL-trained an agent that trains models with RL (for ~\$1.3k)

HackerNews 60% HN 100

No summary — see link for full article.

<https://github.com/Danau5tin/ai-trains-ai>

CONF

Show HN: Oodle.ai - \$10 per million agent traces

HackerNews 50% HN 29

we had hundreds of discussions with engineering leaders over the past few months, and everyone's trying to understand where they are in the AI journey. we collected all this data into a benchmark and built a free grader to let you know where you stand. you answer on a 1-5 scale...

<https://www.oodle.ai/product/agent-observability>

CONF

Show HN: Pullrun - Run One OCI Image as a Container, Firecracker, or Apple VM

HackerNews 50% HN 7

Hi HN, I'm a software engineer who's been building AI agents for about 1.5 years. I've noticed that teams often give agents the ability to perform real tasks, such as refunding money, changing accounts, and writing to production databases, without involving a human in the...

<https://github.com/pullrun/pullrun>

CONF

Show HN: Low-latency local LLM runner via OpenJDK Panama FFM (Java 22)

HackerNews 50% HN 6

I wanted to run AI from inside the JVM. I started out with the standard REST sidecar, ripped that out to use Project Panama (Foreign Function & Memory API) in the new JDK versions to interface directly with llama.cpp. I still wasn't happy with how that functioned, so I built...

<https://github.com/projectargus-cc/libargus.cc>

CONF

Show HN: GWZ - Git Workspace Zone (multi-repo that feels like plain Git)

HackerNews 50% HN 4

Coming from the land of the mono-repo (ex-googler), I decided to embrace the multi-repo, plain-git style. Methought, "hey, submodules, what could possibly go wrong". Well, how about everything: detached HEADs, impossibly tedious commit management, and just about everything else....

<https://news.ycombinator.com/item?id=48915923>

CONF

Show HN: Pulsys - Pull-through cache for Hugging Face built with io_uring

HackerNews 50% HN 4

Pulsys is an open-source, self hosted, authenticated pull-through cache for Hugging Face. It acts as a drop-in replacement for existing HF clients and serves warm hits directly from disk, achieving up to 90 GB/s throughput (benchmarks are run on loopback on ec2) by bypassing...

<https://github.com/pulsys-io/pulsys>

CONF

Show HN: Alluvia - mine your Claude Code/Cursor/ChatGPT history, locally

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/dylanp12/alluvia>

0 CONF

show hn: 1d data to 3D topology, AI deepfakes to protien folding

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/bitofabluetit/knotreel>

1 CONF

Show HN: Pupul, portable and revocable memory you carry into any AI

HackerNews 50% HN 2

No summary — see link for full article.

<https://pupulcorp.com/revoke-demo>

2 CONF

OpenAI researcher Miles Wang in talks to launch AI drug discovery startup valued at \$2B

TechCrunch AI 50%

The funding discussions point to investor interest in applying AI to make breakthroughs in life sciences.

<https://techcrunch.com/2026/07/14/openai-researcher-miles-wang-in-talks-to-launch-ai-dr...>

3 CONF

Lorde says AI glasses are 'not sexy'

TechCrunch AI 50%

"Increasingly in our world, it gets harder and harder to know what is real," Lorde said onstage.

<https://techcrunch.com/2026/07/14/lorde-says-ai-glasses-are-not-sexy/>

4 CONF

OpenAI's first hardware device is reportedly a screenless speaker that can move

TechCrunch AI 50%

The device is weirdly described as involving "mechanical elements that can move on their own" and the Bloomberg report includes the detail that the device is designed to "feel like a companion and become a physical manifestation of OpenAI's ChatGPT."

<https://techcrunch.com/2026/07/14/openais-first-hardware-device-is-reportedly-a-screenl...>

5 CONF

OpenAI pushes back on Apple trade secret lawsuit

TechCrunch AI 50%

OpenAI has issued another statement on the lawsuit, this time suggesting it lacks merit.

<https://techcrunch.com/2026/07/14/openai-pushes-back-on-apple-trade-secret-lawsuit/>

6 CONF

OpenAI's new flagship model deletes files on its own, people keep warning

TechCrunch AI 50%

A number of social media posts claim that GPT-5.6 Sol deleted files and data without warning. OpenAI had basically disclosed the problem in June.

<https://techcrunch.com/2026/07/14/openais-new-flagship-model-deletes-files-on-its-own-p...>

7 CONF

Apple opens its new Siri AI to everyone with the iOS 27 public beta

TechCrunch AI 50%

If you've been waiting to try Apple's revamped Siri without installing a developer beta, you now can. The company on Tuesday released the iOS 27 public beta, giving iPhone owners early access to its AI-powered assistant and other new features before the software's official...

<https://techcrunch.com/2026/07/14/apple-opens-its-new-siri-ai-to-everyone-with-the-ios-...>

8 CONF

The founder of Hinge raised \$18M to build a new AI dating service, Overtone

TechCrunch AI 50%

Overtone describes itself as "a voice- and audio-forward service, enabled by AI, that provides highly curated introductions."

<https://techcrunch.com/2026/07/14/the-founder-of-hinge-raised-18m-to-build-a-new-ai-dat...>

9 CONF

DeepMind CEO calls for an independent standards body to regulate frontier AI

TechCrunch AI 50%

DeepMind CEO Demis Hassabis is proposing an AI "standards body" modeled after FINRA, to test frontier models and develop best practices for their release.

<https://techcrunch.com/2026/07/14/deepmind-ceo-calls-for-an-independent-standards-body-...>

0 CONF

Meta's Adam Mosseri says AI token budgets could soon be capped per engineer

TechCrunch AI 50%

Instagram head Adam Mosseri believes companies will eventually need to manage AI token spending the same way they manage payroll or other operating expenses, predicting that engineers could soon face limits on how much they spend using AI tools.

<https://techcrunch.com/2026/07/14/metas-adam-mosseri-says-ai-token-budgets-could-soon-b...>

1 CONF

Spotify expands its AI push with a ChatGPT-like music assistant

TechCrunch AI 50%

Spotify is rolling out a new AI-powered conversational feature that lets Premium subscribers chat with the app to discover music, podcasts, audiobooks, and more.

<https://techcrunch.com/2026/07/14/spotify-expands-its-ai-push-with-a-chatgpt-like-music...>

2 CONF

Superhuman's new auto-draft feature almost makes me like AI replies

TechCrunch AI 50%

Superhuman's latest AI email drafting feature is its most convincing yet, generating replies that often required little to no editing in our testing.

<https://techcrunch.com/2026/07/14/superhumans-new-auto-draft-feature-almost-makes-me-li...>

3 CONF

The Download: Claude's inner workings, and the future of world models

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. What Anthropic's latest AI discovery does—and doesn't—show —James O'Donnell When Anthropic announced last week that it had found a new window...

<https://www.technologyreview.com/2026/07/14/1140391/the-download-anthropic-claude-inter...>

4 CONF

OpenAI may announce a ChatGPT smart speaker this year

The Verge AI 50%

OpenAI's first device is set to be a smart speaker that lets you talk with ChatGPT, according to a report from Bloomberg. The device apparently won't have a screen, but will use a camera and additional sensors to "understand" your environment. The report comes just days after...

<https://www.theverge.com/ai-artificial-intelligence/965670/openai-chatgpt-ai-smart-spea...>

5 CONF

SpaceXAI's Grok programming tool was uploading its users' entire codebase to cloud storage

The Verge AI 50%

SpaceXAI's Grok Build AI coding tool was spotted uploading users' entire codebases to Google Cloud before it was reported, and the company turned it off. The Register reports that Cereblab published findings on Monday showing how the Grok Build CLI was packaging and uploading...

<https://www.theverge.com/ai-artificial-intelligence/965600/spacexai-grok-build-reposito...>

6 CONF

Spotify is now an AI chatbot, too

The Verge AI 50%

Spotify is experimenting with a new AI feature that allows Premium subscribers to play and explore music, audiobooks, and podcasts by having conversations with a chatbot. The "Talk to Spotify" feature appears across the Home and Now Playing view on Spotify's mobile app. You can...

<https://www.theverge.com/entertainment/965358/spotify-ai-chatbot-interface-music-audiob...>

7 CONF

Google's Demis Hassabis says it's time for a global AI watchdog — led by the US

The Verge AI 50%

Demis Hassabis thinks the world needs an AI watchdog with the power to hit the brakes if frontier models become too dangerous. Writing in a blog post, the Google DeepMind CEO and cofounder said the US should lead the initiative, arguing that the country is the best place to set...

<https://www.theverge.com/tech/965270/google-deepmind-demis-hassabis-global-ai-watchdog>

8 CONF

Knowledge- and Gradient-Guided Reinforcement Learning for Parametrized Action Markov Decision Processes

ArXiv CS.AI 50%

arXiv:2607.12924v1 Announce Type: new Abstract: In this paper, we study Reinforcement Learning in Parametrized Action Markov Decision Processes (PAMDP), where each decision consists of a symbolic action and numerical parameters. In such settings Reinforcement Learning algorithms...

<https://arxiv.org/abs/2607.12924>

- 9 **CONF**
Ontology-Amplified Distillation and Contextuality Auditing for Sovereign Enterprise Language Models: A Combined Proof-of-Mechanism and Negative-Results Method Study
 ArXiv CS.AI 50%
 arXiv:2607.11948v1 Announce Type: new Abstract: Regulated financial institutions operating under data-residency rules need tenant-owned language models that can run inside the institution's perimeter. This paper combines two related FAOS studies into one mechanism-and-control...
<https://arxiv.org/abs/2607.11948>
- 0 **CONF**
GRID: Grammar-Railed Decoding for Enterprise SQL Generation
 ArXiv CS.AI 50%
 arXiv:2607.11951v1 Announce Type: new Abstract: Large language models can write SQL, but enterprise deployment demands more than plausible text: outputs must be syntactically valid, must respect per-role and per-schema policy, must carry provable (not best-effort) guarantees,...
<https://arxiv.org/abs/2607.11951>
- 1 **CONF**
LP Mining with LP2Graph: A Use Case for Railway Rescheduling
 ArXiv CS.AI 50%
 arXiv:2607.11980v1 Announce Type: new Abstract: Like many optimization-driven domains, railway rescheduling relies on Mixed-Integer Linear Programming (MILP), yet the field's modeling knowledge is scattered across hundreds of papers in incompatible notations, and narrative...
<https://arxiv.org/abs/2607.11980>
- 2 **CONF**
Designing Agent-Ready Websites for AI Web Agents: A Framework for Machine Readability, Actionability, and Decision Reliability
 ArXiv CS.AI 50%
 arXiv:2607.12056v1 Announce Type: new Abstract: Online shopping is increasingly shifting toward a model in which AI agents independently search for products, compare options, evaluate constraints, and carry out parts of the purchasing process for users. Website design must now...
<https://arxiv.org/abs/2607.12056>
- 3 **CONF**
Operationalising Multi-Dimensional Evaluation for Conversational Agents: A Scalable, Governed Pipeline with Selective Re-evaluation and Model Benchmarking
 ArXiv CS.AI 50%
 arXiv:2607.12085v1 Announce Type: new Abstract: Evaluating retail conversational agents requires methods beyond lexical-overlap metrics to assess intent alignment, factuality, helpfulness, clarity, tone, and overall response quality. Although LLM-as-a-judge methods provide...
<https://arxiv.org/abs/2607.12085>
- 4 **CONF**
Are we Merging the Right Models? Impact of Expert Training Duration on Model Merging for LLMs
 ArXiv CS.AI 50%
 arXiv:2607.11997v1 Announce Type: cross Abstract: Multi-task model merging combines separately trained expert models into a single model that handles all tasks without co-training. Standard practice merges experts at their optimal validation loss. We challenge this convention by...
<https://arxiv.org/abs/2607.11997>

5 CONF

The Computational Basis of Confidence in Large Language Models

ArXiv CS.AI 50%

arXiv:2607.12447v1 Announce Type: cross Abstract: Reliable confidence -- the probability that a model's own answer is correct -- is essential for the trustworthy deployment of language models. Existing work has largely evaluated confidence by how well it predicts correctness and...

<https://arxiv.org/abs/2607.12447>

6 CONF

Good Benchmarks

ArXiv CS.AI 50%

arXiv:2607.12217v1 Announce Type: new Abstract: Good tasks are correct, solvable, verifiable, well-specified, and hard for interesting reasons. The best tasks describe a real problem an experienced practitioner would recognize, in language a practitioner would use, with tests...

<https://arxiv.org/abs/2607.12217>

7 CONF

On-Device Deep Research at 4B: Exposure Bounds Faithfulness, Retrieval Bounds Coverage

ArXiv CS.AI 50%

arXiv:2607.12257v1 Announce Type: new Abstract: On-device research agents search a corpus, read sources, and write a cited brief on a personal laptop. Whether their citations are faithful, and at what cost, is unmeasured for a deployable small model. This study fixes one 4B...

<https://arxiv.org/abs/2607.12257>

8 CONF

How Many Tasks Are Enough for Agent Benchmark Decisions? A Replay Analysis of Public LLM Agent Benchmarks

ArXiv CS.AI 50%

arXiv:2607.12338v1 Announce Type: new Abstract: Agent benchmarks often compare two agents after all tasks have run, but costly evaluations make partial runs tempting. A task fraction alone does not show whether a partial run supports the same pairwise conclusion as the completed...

<https://arxiv.org/abs/2607.12338>

9 CONF

Critic Experience Bank: Self-Evolving Step-Level Confidence Estimation for LLM Agents

ArXiv CS.AI 50%

arXiv:2607.12397v1 Announce Type: new Abstract: LLM agents act in external environments where each action changes the state that later decisions condition on, and where a single wrong step can waste interaction budget or trigger irreversible side effects long before the final...

<https://arxiv.org/abs/2607.12397>

10 CONF

Isolation as a First-Class Principle for LLM-Agent System Safety: Concepts, Taxonomy, Challenges and Future Directions

ArXiv CS.AI 50%

arXiv:2607.12406v1 Announce Type: new Abstract: The capability of LLM agents to function as the ``brain" of a system fundamentally expands the scope of analysis beyond a standalone model. Consequently, safety is no longer only about input--output content alignment. It also...

<https://arxiv.org/abs/2607.12406>

- 1

CONF

Accepted Prefixes Are Not All You Need: A Negative Result on PEFT-Based Block-Diffusion Drafting

ArXiv CS.AI

50%

arXiv:2607.12422v1 Announce Type: new Abstract: Speculative decoding accelerates autoregressive language model inference by using a cheap drafter to propose multiple future tokens and a target model to verify them. A common design goal is therefore to improve draft quality while...

<https://arxiv.org/abs/2607.12422>
- 2

CONF

Action-Aware Generative Sequence Modeling for Short Video Recommendation

ArXiv CS.AI

50%

arXiv:2604.25834v2 Announce Type: replace Abstract: With the rapid development of the Internet, users have increasingly higher expectations for the recommendation accuracy of online content consumption platforms. However, short videos often contain diverse segments, and users...

<https://arxiv.org/abs/2604.25834>
- 3

CONF

Do LLMs Know What They Know? Measuring Metacognitive Efficiency with Signal Detection Theory

ArXiv CS.AI

50%

arXiv:2603.25112v2 Announce Type: replace-cross Abstract: Standard evaluation of LLM confidence relies on calibration metrics (ECE, Brier score) that conflate how much a model knows (Type-1 accuracy) with how well its confidence signal tracks that knowledge (Type-2 metacognitive...

<https://arxiv.org/abs/2603.25112>
- 4

CONF

Evidence-Grounded AI for Musculoskeletal Care

ArXiv CS.AI

50%

arXiv:2607.12527v1 Announce Type: new Abstract: Musculoskeletal diseases are among the leading causes of disability worldwide and create the greatest global need for rehabilitation. Because recovery, remodelling and degeneration often unfold over months to years, musculoskeletal...

<https://arxiv.org/abs/2607.12527>
- 5

CONF

Vertical Standardisation for High-Risk AI Systems under the EU AI Act: A Domain-Specific Framework for Algorithmic Hiring

ArXiv CS.AI

50%

arXiv:2607.12588v1 Announce Type: new Abstract: According to the recent European legislation, high-risk AI systems will have to adapt in order to comply with requirements related to specific areas, like risk management, data quality and governance, logging and traceability,...

<https://arxiv.org/abs/2607.12588>
- 6

CONF

A Learning-Rate-Gated Failure of GRPO in a Small Language and Vision-Language Model Web Agent: A Controlled Null and Its Mechanism

ArXiv CS.AI

50%

arXiv:2607.12640v1 Announce Type: new Abstract: Reinforcement learning with verifiable rewards, and Group Relative Policy Optimization (GRPO) in particular, is now run routinely on a supervised checkpoint in the hope of producing a stronger agent. We ask whether it adds skill to...

<https://arxiv.org/abs/2607.12640>

7 CONF

Internet of Agentic Things: Networked AI Agents for Closed-Loop IoT Orchestration

ArXiv CS.AI 50%

arXiv:2607.12662v1 Announce Type: new Abstract: The paper introduces the Internet of Agentic Things (IoAT), an architectural framework that integrates agentic AI, IoT, cyber-physical systems, Physical AI, edge computing, and digital twins into a unified closed-loop orchestration...

<https://arxiv.org/abs/2607.12662>

8 CONF

Tracing Agentic Failure from the Flow of Success

ArXiv CS.AI 50%

arXiv:2607.12747v1 Announce Type: new Abstract: Failure attribution for LLM-based agentic systems, i.e., identifying which steps in a failure trajectory caused the task to fail, is critical for debugging and improving these systems. Existing approaches either rely on...

<https://arxiv.org/abs/2607.12747>

9 CONF

Do We Really Need Multimodal Emotion Language Models Larger Than 1B Parameters?

ArXiv CS.AI 50%

arXiv:2607.12787v1 Announce Type: new Abstract: Recent advances in multimodal large language models (MLLMs) have significantly improved the performance of multimodal emotion recognition (MER) and enabled interpretable description generation by jointly modeling video, audio, and...

<https://arxiv.org/abs/2607.12787>

0 CONF

Who Grades the Grader? Co-Evolving Evaluation Metrics and Skills for Self-Improving LLM Agents

ArXiv CS.AI 50%

arXiv:2607.12790v1 Announce Type: new Abstract: Self-evolving agent systems improve by creating, revising, and retiring their own skills, but every such loop rests on a hidden assumption: a reliable evaluation metric already exists. In many real applications it does not. We make...

<https://arxiv.org/abs/2607.12790>

1 CONF

Visual Access Boundaries in Vision-Language Model Reasoning

ArXiv CS.AI 50%

arXiv:2607.12815v1 Announce Type: new Abstract: Chain-of-Thought (CoT) prompting is widely used as a test-time scaling strategy for Vision-Language Models (VLMs), but it remains unclear what is extended when VLMs generate longer reasoning traces. We ask whether CoT requires...

<https://arxiv.org/abs/2607.12815>

2 CONF

Human-AI Agent Interaction as a Neuroplastic Training Environment

ArXiv CS.AI 50%

arXiv:2607.12823v1 Announce Type: new Abstract: Interaction with AI agents has become one of the most frequent activities of everyday digital life. Whether conversing with an assistant, working with a coding copilot, or generating images, the interaction follows a common...

<https://arxiv.org/abs/2607.12823>

3 CONF

Solution of the Hempel's statistical ambiguity problem and Causal AI

ArXiv CS.AI 50%

arXiv:2607.12826v1 Announce Type: new Abstract: This paper addresses Carl Hempel's longstanding problem of statistical ambiguity in inductive-statistical inference, in which contradictory predictions are derived from statistical laws. To avoid such predictions, Carl Hempel...

<https://arxiv.org/abs/2607.12826>

4 CONF

A Multi-Agent System for Autonomous, Fine-Tuning-Free Clinical Symptom Detection: Development and Validation Study

ArXiv CS.AI 50%

arXiv:2607.12886v1 Announce Type: new Abstract: Clinical notes contain many of the signs and symptoms that bring patients to care, yet this information rarely reaches structured fields. Existing extraction approaches either rely on context-insensitive rules that generate false...

<https://arxiv.org/abs/2607.12886>

5 CONF

MemOps: Benchmarking Lifecycle Memory Operations in Long-Horizon Conversations

ArXiv CS.AI 50%

arXiv:2607.12893v1 Announce Type: new Abstract: Long-term memory has become a foundational capability for LLM-based agents that accompany users across extended, multi-session interactions. Existing benchmarks, however, evaluate such memory almost exclusively through downstream...

<https://arxiv.org/abs/2607.12893>

6 CONF

ReLope: KL-Regularized LoRA Probes for Multimodal LLM Routing

ArXiv CS.AI 50%

arXiv:2603.24787v2 Announce Type: replace Abstract: Routing has emerged as a promising strategy for balancing performance and cost in large language model (LLM) systems that combine lightweight models with powerful but expensive large models. Recent studies show that \emph{probe}...

<https://arxiv.org/abs/2603.24787>

7 CONF

Resist and Update: Counterfactual Report Coordinates for Incentive-Compatible LLMs

ArXiv CS.AI 50%

arXiv:2607.12985v1 Announce Type: new Abstract: Aligned language models routinely misreport under non-evidential incentive pressure: they agree with a confident user or overstate certainty even when their internal belief is unchanged. We cast this as a failure of internal...

<https://arxiv.org/abs/2607.12985>

8 CONF

Win by Silence: Deletion Non-Monotonicity, Autonomous Exploitation, and Typed-State Gating in LLM Plan Evaluation

ArXiv CS.AI 50%

arXiv:2607.12986v1 Announce Type: new Abstract: Plan evaluators can reward a strategic plan for becoming less explicit. This paper studies that failure in a staged expected-value scorer for LLM-generated venture routes. Proposition 1 gives the score change from deleting an...

<https://arxiv.org/abs/2607.12986>

9 CONF

DRIFTLENS: Measuring Memory-Induced Reasoning Drift in Personalized Language Models

ArXiv CS.AI 50%

arXiv:2607.02374v2 Announce Type: replace Abstract: Personalization changes what a model says to a user; we show that it can also change the reasoning trajectory used to justify the response. Modern LLMs personalize interactions by storing user attributes, preferences, and prior...

<https://arxiv.org/abs/2607.02374>

0 CONF

Do AI Agents Know When a Task Is Simple? Toward Complexity-Aware Reasoning and Execution

ArXiv CS.AI 50%

arXiv:2607.13034v1 Announce Type: new Abstract: Large language model (LLM) agents increasingly automate multi-step engineering and informatics workflows, yet they rarely ask how much effort a task actually requires. They often follow a maximum-context-first strategy--re-reading...

<https://arxiv.org/abs/2607.13034>

1 CONF

Answering Without Referring: How AI Search Rewrites the Web's Economic Bargain

ArXiv CS.AI 50%

arXiv:2607.07652v1 Announce Type: cross Abstract: Search engines have long allocated attention on the web by routing users from queries to websites. AI search changes this arrangement because information needs can be resolved inside the intermediary. Using URL-level Comscore...

<https://arxiv.org/abs/2607.07652>

2 CONF

FAIR GraphRAG: A Retrieval-Augmented Generation Approach for Semantic Data Analysis

ArXiv CS.AI 50%

arXiv:2607.11464v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) addresses the limitations of Large Language Models (LLMs) when providing responses to domain-specific questions. Graph-based RAG approaches, such as GraphRAG, enhance retrieval by capturing...

<https://arxiv.org/abs/2607.11464>

3 CONF

So Many Opinions, So Many LLMs: Comparing Large Language Models to Traditional Machine Learning for Open- Ended Survey Analysis

ArXiv CS.AI 50%

arXiv:2607.11890v1 Announce Type: cross Abstract: Open-ended surveys offer valuable insights, but they are notoriously difficult to analyze at scale. Building on previous work that employed traditional machine learning to classify text ("So Many Responses, So Little Time: A...

<https://arxiv.org/abs/2607.11890>

4 CONF

CANDI: Contextual Alignment for Niche Domains Question Answering

ArXiv CS.AI 50%

arXiv:2607.11891v1 Announce Type: cross Abstract: The deployment of large language models (LLMs) in specialized domains like medical diagnostics and financial advisory necessitates evaluating capabilities beyond general knowledge. Traditional question-answering benchmarks often...

<https://arxiv.org/abs/2607.11891>

5 CONF

I'm Sorry, but I Can't Help with Braille: Revealing Accessibility Failures in State-of-the-Art LLMs

ArXiv CS.AI 50%

arXiv:2607.11893v1 Announce Type: cross Abstract: Large Language Models (LLMs) perform strongly on many language tasks, but their capability in structurally constrained, accessibility-critical modalities such as Braille remains unclear. We evaluate state-of-the-art LLMs on...

<https://arxiv.org/abs/2607.11893>

6 CONF

Graph-Based Detection of Disinformation Narrative Diffusion between Russian and Ukrainian Telegram Channels

ArXiv CS.AI 50%

arXiv:2607.11894v1 Announce Type: cross Abstract: Detecting disinformation narratives on social media is challenging due to the scale of amplification, rapid evolution, and linguistic variability of online content. We propose a graph-based framework for identifying and analyzing...

<https://arxiv.org/abs/2607.11894>

7 CONF

SeqGPT: A Constrained Transformer Agent for the Inverse Design of Multi-Panel Composite Structures

ArXiv CS.AI 50%

arXiv:2607.11910v1 Announce Type: cross Abstract: Optimizing composite stacking sequences to match continuous targets (e.g., Lamination or Buckling Parameters) with discrete manufacturing constraints represents a challenging combinatorial inverse problem that regularly occurs in...

<https://arxiv.org/abs/2607.11910>

8 CONF

AAAI-26 Dual Submissions: Novel Challenges

ArXiv CS.AI 50%

arXiv:2607.11918v1 Announce Type: cross Abstract: Dual submissions, in which identical or substantially similar papers are simultaneously submitted to one or more archival venues, without cross-citation or disclosure, are a growing problem for the AAAI Conference and other...

<https://arxiv.org/abs/2607.11918>

9 CONF

Do You Remember? Toward Memory-Centric Multimodal AI

ArXiv CS.AI 50%

arXiv:2607.11919v1 Announce Type: cross Abstract: Human memory is reconstructive, not a faithful recording. Current multimodal LLMs (MLLMs) lack this capability: they process images through a frozen visual encoder, produce a one-shot text output, and discard internal...

<https://arxiv.org/abs/2607.11919>

10 CONF

Sensitivity to Subjective Expected Utility Maximization: A Methodological Study, with an Illustrative Application to LLM Decision-Making

ArXiv CS.AI 50%

arXiv:2607.11920v1 Announce Type: cross Abstract: Evaluating decisions made under uncertainty is hard when labeled outcomes are scarce, costly, or confounded with luck. We treat subjective expected utility (SEU) maximization as a stated standard and define a graded measure --...

<https://arxiv.org/abs/2607.11920>

1 CONF

BattVAE-GP: Generative Modeling of Long-Horizon Battery Degradation with Uncertainty Quantification

ArXiv CS.AI 50%

arXiv:2607.11943v1 Announce Type: cross Abstract: Long-horizon physics-based simulations of battery degradation provide mechanistic insight but remain computationally expensive, limiting their use for dense exploration of operating conditions over extended cycle life. Here, we...

<https://arxiv.org/abs/2607.11943>

2 CONF

BAT-RM: A Boundary-Aware Transformer with Region-Aware Multi-Directional Mamba for Clinically Deployed Cervical Cancer Radiotherapy Auto-Contouring

ArXiv CS.AI 50%

arXiv:2607.11949v1 Announce Type: cross Abstract: We present a clinically deployed end-to-end auto-contouring system for cervical cancer radiotherapy planning, anchored by the Boundary-Aware Transformer with Region-Aware Mamba (BAT-RM), a hybrid architecture that integrates...

<https://arxiv.org/abs/2607.11949>

3 CONF

Scale-Aware Attention for Scarce Neural Data: An RG-Flow Transformer on Sleep-EDF EEG

ArXiv CS.AI 50%

arXiv:2607.11950v1 Announce Type: cross Abstract: Brain field potentials are scale-free: their power spectra follow a $1/f^{\beta}$ law whose aperiodic exponent β tracks cortical state, and sleep depth in particular is a shift in β . We ask whether a transformer...

<https://arxiv.org/abs/2607.11950>

4 CONF

Evaluating Reliability in Machine Learning Models for Early Chronic Kidney Disease Prediction: A Systematic Review of Data Leakage and Predictor Stability

ArXiv CS.AI 50%

arXiv:2607.11963v1 Announce Type: cross Abstract: The early detection of Chronic Kidney Disease using machine learning has attracted significant interest in healthcare-related computer science. Despite rapid advancements in this field, many reported studies remain inconsistent...

<https://arxiv.org/abs/2607.11963>

5 CONF

Learning to Discretize: Diffusion-Based Adaptive Mesh with Spectral Guidance

ArXiv CS.AI 50%

arXiv:2607.11974v1 Announce Type: cross Abstract: Most neural partial differential equation (PDE) surrogates learn how fields evolve after a grid has already been chosen. However, before any operator is applied, the grid has already determined how modeling capacity is allocated...

<https://arxiv.org/abs/2607.11974>

6 CONF

Gene Expression-Informed Jointly Controlled Generative Modeling for Precision Molecular Design

ArXiv CS.AI 50%

arXiv:2607.11978v1 Announce Type: cross Abstract: Precision molecular design aims to discover personalized drug candidates through joint control of multiple conditions, such as biological relevance and molecular design strategies. Biological relevance reflects cellular...

<https://arxiv.org/abs/2607.11978>

7 CONF

Sparse Inter-Layer Dependencies of Transformer FFN Neurons

ArXiv CS.AI 50%

arXiv:2607.11990v1 Announce Type: cross Abstract: Feedforward network (FFN) blocks account for a large fraction of the parameters and computation in Transformer architectures, yet their internal structure remains difficult to interpret due to the additive superposition induced...

<https://arxiv.org/abs/2607.11990>

8 CONF

Representation and Reference Selection in Training-Free Synthetic Image Attribution

ArXiv CS.AI 50%

arXiv:2607.12052v1 Announce Type: cross Abstract: Synthetic image attribution aims at identifying the generator responsible for a given AI-generated image. Training-free reference-based attribution methods are easily scalable, since newly emerging generators can be incorporated...

<https://arxiv.org/abs/2607.12052>

9 CONF

PFAdapter: Hierarchical LoRA Decomposition for Personalized Federated MLLMs

ArXiv CS.AI 50%

arXiv:2607.12111v1 Announce Type: cross Abstract: Agentic AI systems are reshaping communications and networking by deploying autonomous intelligent agents capable of collaborative learning while maintaining data privacy at network edges. Within distributed network environments,...

<https://arxiv.org/abs/2607.12111>

0 CONF

Continual Learning with Elastic Regularization and Synthetic Replay for Federated MLLM Fine-Tuning

ArXiv CS.AI 50%

arXiv:2607.12112v1 Announce Type: cross Abstract: Federated fine-tuning of Multimodal Large Language Models (MLLMs) across distributed networks enables privacy-sensitive adaptation to evolving data streams, yet a fundamental obstacle prevents robust deployment in dynamic...

<https://arxiv.org/abs/2607.12112>

1 CONF

GaitSpan: Growing Humanoid Locomotion from Walking to Running

ArXiv CS.AI 50%

arXiv:2607.12114v1 Announce Type: cross Abstract: A humanoid that can walk should not relearn locomotion from scratch to jog or run. Yet current approaches often obtain gait diversity by prescribing gait schedules, imitating motion clips, training experts to switch between or...

<https://arxiv.org/abs/2607.12114>

2 CONF

TRAIL: A Platform for Configurable Human--AI Teaming Experiments

ArXiv CS.AI 50%

arXiv:2607.12180v1 Announce Type: cross Abstract: An AI teammate's design properties (personality, communication style, when it speaks) can shape a team's trust, coordination, and decisions. Studying this rigorously demands infrastructure no existing tool provides: reproducible...

<https://arxiv.org/abs/2607.12180>

3 CONF

Accelerating Masked Diffusion Large Language Models: A Survey of Efficient Inference Techniques

ArXiv CS.AI 50%

arXiv:2607.12829v1 Announce Type: cross Abstract: Diffusion large language models (dLLMs) offer a theoretical advantage in parallel generation over standard autoregressive models. However, parallel generation alone does not guarantee practical speedups. Realizing this efficiency...

<https://arxiv.org/abs/2607.12829>

4 CONF

The Benjamini--Hochberg Procedure Can Fail to Control the FDR for Correlated Two-Sided Gaussian Tests

ArXiv CS.AI 50%

arXiv:2607.12208v1 Announce Type: cross Abstract: We show that the Benjamini--Hochberg procedure can fail to control the false discovery rate (FDR) at its nominal level for correlated two-sided Gaussian p -values. We construct a factor model for which, at level $\alpha=0.01$, a...

<https://arxiv.org/abs/2607.12208>

5 CONF

RCWT: Measuring Task-Budget Displacement from Coordination Content in LLM Calls

ArXiv CS.AI 50%

arXiv:2607.12216v1 Announce Type: cross Abstract: Multi-agent and memory-augmented LLM systems often place coordination content, shared state, prior discussion, tool outputs, summaries, and role instructions, inside the same finite prompt used for the current task. This creates...

<https://arxiv.org/abs/2607.12216>

6 CONF

Fin-Analyst at FinMMEval 2026 Task 3: A Live Hybrid Trading Agent with LLM Specialists and Rule-Based Signals

ArXiv CS.AI 50%

arXiv:2607.12233v1 Announce Type: cross Abstract: Large language model (LLM) trading agents show promising performance in equity markets, yet remain narrowly focused on US equities with little evidence from live deployment. We present Fin-Analyst, a hybrid agent for FinMMEval...

<https://arxiv.org/abs/2607.12233>

7 CONF

Code-MUE: Measuring Code LLMs' Uncertainty through Execution-based Semantic Interaction Graphs

ArXiv CS.AI 50%

arXiv:2607.12273v1 Announce Type: cross Abstract: As Code Large Language Models (LLMs) become central to modern software engineering, their inherent stochasticity poses significant real-world risks, where even minor errors can lead to severe functional, security, or safety...

<https://arxiv.org/abs/2607.12273>

8 CONF

A Longitudinal Analysis of Public Discourse on AI Ethics in Education Using Twitter Data

ArXiv CS.AI 50%

arXiv:2607.12295v1 Announce Type: cross Abstract: The rapid integration of artificial intelligence (AI) and generative AI (GenAI) into education presents significant opportunities to enhance teaching and learning, while raising ethical concerns about the responsible use of these...

<https://arxiv.org/abs/2607.12295>

9 CONF

SQuTR: A Robustness Benchmark for Spoken Query to Text Retrieval under Acoustic Noise

ArXiv CS.AI 50%

arXiv:2602.12783v3 Announce Type: replace-cross Abstract: Spoken query retrieval is an important interaction mode in modern information retrieval. However, existing evaluation datasets are often limited to simple queries under constrained noise conditions, making them inadequate...

<https://arxiv.org/abs/2602.12783>

0 CONF

Evaluating Health Misinformation in Low-Resource Languages: Integrating Small Language Models with a Culturally-Sensitive Responsible NLP Framework (Bangla as a Case Study)

ArXiv CS.AI 50%

arXiv:2607.12336v1 Announce Type: cross Abstract: Artificial Intelligence (AI) technologies, while serving as a foundational enabler for modern social media and digital health services, exert a bivalent effect by simultaneously acting as a combatant against and a spread vector...

<https://arxiv.org/abs/2607.12336>

1 CONF

Demonstration of the common dual-channel feature decoupling characteristic of front-door mediation causal inference methods in whole-slide image classification

ArXiv CS.AI 50%

arXiv:2607.12376v1 Announce Type: cross Abstract: Causal inference using front door intervention and multi-instance learning (MIL) has advanced the analysis of Whole Slide Images (WSI) in digital pathology. These methods adjust feature distributions of subtle evidence sub-images...

<https://arxiv.org/abs/2607.12376>

2 CONF

An Omnilingual-ASR-Based Speech-LLM System for the 2nd MLC-SLM Challenge

ArXiv CS.AI 50%

arXiv:2607.12468v1 Announce Type: cross Abstract: We describe our submission to Task 1 of the 2nd MLCSLM Challenge: a cascaded diarization-then-recognition system that combines DiariZen-Large-s80 (WavLM-Large) segmentation, CAM++ embedding-based two-speaker clustering, and a...

<https://arxiv.org/abs/2607.12468>

3 CONF

Traceback Translators Against Forgetting in Continual Fake Speech Detection

ArXiv CS.AI 50%

arXiv:2607.12569v1 Announce Type: cross Abstract: Fake speech detectors are increasingly challenged by the development of new and more accurate generative models. To cope with this problem, continual learning techniques are nowadays widely considered feasible strategies for...

<https://arxiv.org/abs/2607.12569>

4 CONF

Deep Learning-based Surrogate Modelling of the LOD Method for Multiscale Problems

ArXiv CS.AI 50%

arXiv:2607.12570v1 Announce Type: cross Abstract: Multiscale problems are notoriously difficult to tackle using traditional numerical methods, as accurately resolving fine-scale features often requires prohibitively fine discretizations. This challenge is particularly pronounced...

<https://arxiv.org/abs/2607.12570>

5 CONF

Explainable-by-Design Audio Deepfake Detection via Wiener-Hopf Linear Prediction

ArXiv CS.AI 50%

arXiv:2607.12584v1 Announce Type: cross Abstract: The rapid advancement of synthetic speech generation methods has made audio deepfake detection a critical challenge in multimedia forensics. While recent approaches achieve high detection accuracy, they typically rely on...

<https://arxiv.org/abs/2607.12584>

6 CONF

Multi-Perspective Agentic Program Repair via Code Property Graphs and Temporal Execution Graphs

ArXiv CS.AI 50%

arXiv:2607.12605v1 Announce Type: cross Abstract: Large language models (LLMs) have improved automated program repair (APR), but two limitations remain. First, raw execution traces are often too large and repetitive to serve as effective model context. Second, repeated patch...

<https://arxiv.org/abs/2607.12605>

7 CONF

Can Induced Emotion Bias LLM Behaviors in Sequential Decision Making?

ArXiv CS.AI 50%

arXiv:2607.12631v1 Announce Type: cross Abstract: As Large Language Models (LLMs) are increasingly deployed as autonomous agents in high-stakes domains, understanding contextual factors that may modulate their decision-making becomes critical. While LLMs are trained to perceive...

<https://arxiv.org/abs/2607.12631>

8 CONF

Evidence-Grounded Verified Agentic Reasoning: A Path Toward Eliminating LLM Hallucination in Empirical Inference via Tool-Attested Kernel Proofs

ArXiv CS.AI 50%

arXiv:2607.12650v1 Announce Type: cross Abstract: Tool access alone does not make LLM empirical reasoning governable: accepted outputs need not descend from attested evidence, and accepted deductions need not hold up under formal scrutiny. We present EG-VAR (Evidence-Grounded...

<https://arxiv.org/abs/2607.12650>

9 CONF

Jetson-PI: Towards Onboard Real-Time Robot Control via Foresight-Aligned Asynchronous Inference

ArXiv CS.AI 50%

arXiv:2607.12659v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models have achieved impressive performance on diverse embodied tasks. However, deploying VLA models on low-power onboard devices, such as the Jetson Orin, remains challenging due to their high...

<https://arxiv.org/abs/2607.12659>

00 CONF

Text-Aided Multi-Modal Panoptic Symbol Spotting for CAD Floor Plan Drawings

ArXiv CS.AI 50%

arXiv:2607.12678v1 Announce Type: cross Abstract: Computer-Aided Design (CAD) floor plan drawings contain both graphical primitives and textual annotations, which provide complementary geometric and semantic cues for intelligent design understanding. Among CAD analysis tasks,...

<https://arxiv.org/abs/2607.12678>

01 CONF

Line-Anchored Feedback Cuts Token Costs and Improves Correctness in AI Code Editing

ArXiv CS.AI 50%

arXiv:2607.12713v1 Announce Type: cross Abstract: Generated tokens are a direct driver of the cost, latency, and energy of generative AI (GAI) code editing. We show the format of feedback is a lever on all three. We compare two deliveries of the same requested changes: a...

<https://arxiv.org/abs/2607.12713>

02 CONF

Bulkhead: Automated Semantic Detection and Remediation of Container Escape Vulnerabilities

ArXiv CS.AI 50%

arXiv:2607.12723v1 Announce Type: cross Abstract: Filesystem isolation in container ecosystems is often weakened by cross-boundary path misresolution, causing path traversal (PaTra) vulnerabilities. These vulnerabilities stem from insecure host-container interactions and have...

<https://arxiv.org/abs/2607.12723>

03 CONF

Weakly Supervised Spatio-Temporal Candidate Discovery of Dairy Farm Sites from Seasonal Satellite Imagery

ArXiv CS.AI 50%

arXiv:2607.12748v1 Announce Type: cross Abstract: Farm site discovery from satellite imagery is a spatiotemporal candidate ranking problem because farm evidence is distributed across pasture, field boundaries, roads, buildings, and seasonal vegetation patterns. Direct farm...

<https://arxiv.org/abs/2607.12748>

04 CONF

Practical Judgment, Virtue, and Intuition in the Use of Opaque AI-Enabled Systems

ArXiv CS.AI 50%

arXiv:2607.12755v1 Announce Type: cross Abstract: AI-enabled systems are seeing increasing deployment across numerous domains, with many being "black boxes" with respect to core functions and capabilities. I.e., many systems take inputs and give outputs, but without users having...

<https://arxiv.org/abs/2607.12755>

05 CONF

The One-Word Census: Answer-Choice Conformity Across 44 Language Models

ArXiv CS.AI 50%

arXiv:2607.12796v1 Announce Type: cross Abstract: When a language model must pick one answer from a large space of equally valid options, which does it pick -- and how often is it the same answer every other model picks? Asked to "pick a word -- any word," 44 models chose...

<https://arxiv.org/abs/2607.12796>

06 CONF

ViHoRec: A Quality-Controlled Vietnamese Hotel Recommendation Dataset and Cold-Start Benchmark

ArXiv CS.AI 50%

arXiv:2607.12946v1 Announce Type: cross Abstract: Recommender-system research for Vietnamese remains limited by the absence of a public, well-documented hotel interaction resource. Building such a resource is challenging for three reasons: cross-platform hotel names must be...

<https://arxiv.org/abs/2607.12946>

07 CONF

Form, Not Content? A Preregistered, Placebo-Controlled Evaluation of Learned Error-Conditioned Self-Repair Through Prompts and Weights in Frozen Small Code Models

ArXiv CS.AI 50%

arXiv:2607.12962v1 Announce Type: cross Abstract: Frozen small code LLMs are deployed locally, yet the information guiding a retry after a failed attempt is still measured without placebo controls in the self-repair literature. We treat a failed program as a conjecture and an...

<https://arxiv.org/abs/2607.12962>

08 CONF

Rethinking Reward Models for Multi-Domain Test-Time Scaling

ArXiv CS.AI 50%

arXiv:2510.00492v3 Announce Type: replace Abstract: The reliability of large language models (LLMs) during test-time scaling is often assessed with \emph{external verifiers} or \emph{reward models} that distinguish correct reasoning from flawed logic. Prior work has studied both...

<https://arxiv.org/abs/2510.00492>

09 CONF

CrochetBench: Can Vision-Language Models Move from Describing to Doing in Crochet Domain?

ArXiv CS.AI 50%

arXiv:2511.09483v3 Announce Type: replace Abstract: While multimodal large language models can describe visual content, their ability to generate executable procedures remains underexplored. CrochetBench presented in this paper evaluates this shift from describing to doing...

<https://arxiv.org/abs/2511.09483>

10 CONF

Mobility-Aware Cache Framework for Scalable LLM-Based Human Mobility Simulation

ArXiv CS.AI 50%

arXiv:2602.16727v2 Announce Type: replace Abstract: Simulating large-scale human mobility is fundamental to understanding population movement patterns and supporting real-world geospatial applications such as urban planning, epidemic response, and transportation analysis. Recent...

<https://arxiv.org/abs/2602.16727>

11 CONF

Propheticus: Machine Learning Framework for the Development of Predictive Models for Reliable and Secure Software

ArXiv CS.AI 50%

arXiv:1809.01898v2 Announce Type: replace-cross Abstract: The growing complexity of software calls for innovative solutions that support the deployment of reliable and secure software. Machine Learning (ML) has shown its applicability to various complex problems and is...

<https://arxiv.org/abs/1809.01898>

12 CONF

Seeing Through Uncertainty: Free-Energy-Inspired Real-Time Adaptation for Robust Visual Navigation

ArXiv CS.AI 50%

arXiv:2403.01977v4 Announce Type: replace-cross Abstract: Navigation in the natural world is a feat of adaptive inference, where biological organisms maintain goal-directed behaviour despite noisy and incomplete sensory streams. Central to this ability is the Free Energy...

<https://arxiv.org/abs/2403.01977>

13 CONF

Enabling Energy-Efficient Simultaneous Multi-Task Reinforcement Learning through Spiking Neural Networks with Active Dendrites for Bio-inspired Generalist Agents

ArXiv CS.AI 50%

arXiv:2412.04847v2 Announce Type: replace-cross Abstract: Reinforcement learning (RL) has demonstrated remarkable capabilities in training agents to solve complex tasks autonomously, such as mobile robots, UAVs/UGVs, and game-playing agents). However, scaling RL to master...

<https://arxiv.org/abs/2412.04847>

14 CONF

Modeling Story Expectations: A Generative Framework using LLMs

ArXiv CS.AI 50%

arXiv:2412.15239v4 Announce Type: replace-cross Abstract: Consumers' engagement with stories is shaped by their expectations about what will happen next, yet modeling these forward-looking beliefs over unstructured narrative content has remained challenging. We develop a...

<https://arxiv.org/abs/2412.15239>

15 CONF

Inclusive Federated Learning Through Compliance-Weighted Noise Allocation in Healthcare AI

ArXiv CS.AI 50%

arXiv:2505.22108v4 Announce Type: replace-cross Abstract: Background: Federated learning (FL) enables collaborative training of clinical AI models without centralizing patient data, but adoption is limited by privacy concerns, heterogeneous institutional compliance, and resource...

<https://arxiv.org/abs/2505.22108>

16 CONF

SheetMind: An End-to-End LLM-Powered Multi-Agent Framework for Spreadsheet Automation

ArXiv CS.AI 50%

arXiv:2506.12339v2 Announce Type: replace-cross Abstract: We present SheetMind, a modular multi-agent framework powered by large language models (LLMs) for spreadsheet automation via natural language instructions. In this paper, we introduce a hierarchical agentic system...

<https://arxiv.org/abs/2506.12339>

17 CONF

Look, Focus, Act: Efficient and Robust Robot Learning via Human Gaze and Foveated Vision Transformers

ArXiv CS.AI 50%

arXiv:2507.15833v3 Announce Type: replace-cross Abstract: Human vision is a highly active process driven by gaze, which directs attention to task-relevant regions through foveation, dramatically reducing visual processing. In contrast, robot learning systems typically rely on...

<https://arxiv.org/abs/2507.15833>

18 CONF

Evidence Recomposition and Predictive Context Residualization for Visual Attribution in Multimodal Large Language Models

ArXiv CS.AI 50%

arXiv:2509.22415v3 Announce Type: replace-cross Abstract: Multimodal large language models (MLLMs) have achieved strong vision-language performance, yet their token-level visual evidence remains difficult to inspect. Recent logit-lens attribution methods project each...

<https://arxiv.org/abs/2509.22415>

19 CONF

Evolution Strategies at Scale: LLM Fine-Tuning Beyond Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2509.24372v3 Announce Type: replace-cross Abstract: Fine-tuning large language models (LLMs) for downstream tasks is an essential stage of modern AI deployment. Reinforcement learning (RL) has emerged as the dominant fine-tuning paradigm, underpinning many state-of-the-art...

<https://arxiv.org/abs/2509.24372>

20 CONF

Adaptive Testing for LLM Evaluation: A Psychometric Alternative to Static Benchmarks

ArXiv CS.AI 50%

arXiv:2511.04689v3 Announce Type: replace-cross Abstract: Evaluating large language models (LLMs) typically requires thousands of benchmark items, making the process expensive, slow, and increasingly impractical at scale. Existing evaluation protocols rely on average accuracy...

<https://arxiv.org/abs/2511.04689>

21 CONF

First, do NOHARM: a medical safety benchmark and randomized study of physician and AI teaming on clinical consultations

ArXiv CS.AI 50%

arXiv:2512.01241v4 Announce Type: replace-cross Abstract: Large language models (LLMs) and medical AI tools are routinely used by physicians and patients for medical advice, yet their clinical safety profiles remain poorly characterized. We present NOHARM (Numerous Options Harm...

<https://arxiv.org/abs/2512.01241>

22 CONF

Tracking Drift: Variation-Aware Entropy Scheduling for Non-Stationary Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2601.19624v3 Announce Type: replace-cross Abstract: Real-world reinforcement learning often faces environment drift, but most existing methods rely on static entropy coefficients/target entropy, causing over-exploration during stable periods and under-exploration after...

<https://arxiv.org/abs/2601.19624>

23 CONF

DECO: Decoupled Multimodal Diffusion Transformer for Bimanual Dexterous Manipulation with a Plugin Tactile Adapter

ArXiv CS.AI 50%

arXiv:2602.05513v3 Announce Type: replace-cross Abstract: Bimanual dexterous manipulation relies on integrating multimodal inputs to perform complex real-world tasks. To address the challenges of effectively combining these modalities, we propose DECO, a decoupled multimodal...

<https://arxiv.org/abs/2602.05513>

24 CONF

Self-Regulated Reading with AI Support: An Eight-Week Study with Students

ArXiv CS.AI 50%

arXiv:2602.09907v2 Announce Type: replace-cross Abstract: College students increasingly use AI chatbots to support academic reading, yet we lack granular understanding of how these interactions shape their reading experience and cognitive engagement. We conducted an eight-week...

<https://arxiv.org/abs/2602.09907>

25 CONF

Declarative by Design, Assistable Only by Convention: Benchmarking Multi-Agent Frameworks for AI-Assistability

ArXiv CS.AI 50%

arXiv:2602.11198v2 Announce Type: replace-cross Abstract: Multi-agent frameworks (MAFs) promise to simplify LLM-driven software development, yet no principled metric captures how well AI coding assistants can generate correct, framework-specific code. We introduce...

<https://arxiv.org/abs/2602.11198>

26 CONF

1D-Bench: A Benchmark for Iterative UI Code Generation with Visual Feedback in Real-World

ArXiv CS.AI 50%

arXiv:2602.18548v2 Announce Type: replace-cross Abstract: Design-to-code translates high-fidelity UI designs into executable front-end implementations, but progress remains hard to compare due to inconsistent datasets, toolchains, and evaluation protocols. We introduce 1D-Bench,...

<https://arxiv.org/abs/2602.18548>

27 CONF

Understanding Sources of Demographic Predictability in Brain MRI via Disentangling Anatomy and Contrast

ArXiv CS.AI 50%

arXiv:2603.04113v2 Announce Type: replace-cross Abstract: Demographic attributes can be predicted from medical images, raising concerns about bias in clinical AI systems. In X-ray imaging, acquisition characteristics have been shown to contribute substantially to this...

<https://arxiv.org/abs/2603.04113>

28 CONF

Research Novelty in Information Systems Journals After ChatGPT: Differences Across Institutional Language Contexts

ArXiv CS.AI 50%

arXiv:2603.22510v2 Announce Type: replace-cross Abstract: Large language models are increasingly used in scholarly work, yet it remains unclear whether their productivity gains are accompanied by changes in research novelty. We examine how relative abstract-level semantic...

<https://arxiv.org/abs/2603.22510>

29 CONF

The Score Hamiltonian: Mapping Diffusion Models to Adiabatic Transport

ArXiv CS.AI 50%

arXiv:2606.05217v3 Announce Type: replace-cross Abstract: We exhibit an exact correspondence between sampling with score-based diffusion models and adiabatic transport of ground states for a family of Schrödinger operators we call Score Hamiltonians, built from the learned...

<https://arxiv.org/abs/2606.05217>

30 CONF

RFM-Editing 2: Text-Guided Audio Editing with Rectified Flow Matching and Coarse-to-Fine Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2606.20101v3 Announce Type: replace-cross Abstract: Audio editing aims to modify specific content in an existing audio clip according to a text instruction or description while preserving the remaining acoustic content. Despite the remarkable progress of diffusion models,...

<https://arxiv.org/abs/2606.20101>

31 CONF

Prime Fourier Embeddings: A Principled Basis for Modular Arithmetic

ArXiv CS.AI 50%

arXiv:2606.23044v2 Announce Type: replace-cross Abstract: Numbers have algebraic structure that standard neural embeddings often fail to expose. We introduce Prime Fourier Embeddings (PFE), which encode integers as prime-indexed (cos, sin) pairs derived from the harmonic...

<https://arxiv.org/abs/2606.23044>

32 CONF

ReflectWorld-MM: An Entity-Oriented Multimodal Memory System for Open-Ended Video Streams

ArXiv CS.AI 50%

arXiv:2607.09759v2 Announce Type: replace-cross Abstract: Building assistants that can continually watch the world, remember what they see, and reason over their accumulated experience is a long-standing goal, and recently multimodal agents equipped with long-term memory over...

<https://arxiv.org/abs/2607.09759>

33 CONF

ABot-N1: Toward a General Visual Language Navigation Foundation Model

ArXiv CS.AI 50%

arXiv:2607.10383v2 Announce Type: replace-cross Abstract: Visual Language Navigation foundation models aim to unify deep reasoning for grounded spatial decisions with broad versatility for diverse embodied tasks. Current approaches typically achieve this integration via...

<https://arxiv.org/abs/2607.10383>