



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	89	Media
TechCrunch AI	13	Media
HackerNews	4	Community
The Verge AI	3	Media
VentureBeat AI	1	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 1 story

VERI

OpenAI finally launches hardware... for Codex

The Verge AI +1 source

80%

OpenAI is finally releasing some hardware. No, it isn't the mysterious AI-powered device the company is developing with former Apple designer Jony Ive, a project already tangled up in a messy lawsuit. Instead, it's a product designed to be used with its coding platform, Codex....

<https://www.theverge.com/ai-artificial-intelligence/965901/openai-hardware-codex-micro-...>

CONFIRMED SIGNALS — 109 stories

CONF

Show HN: Aict - Unix coreutils that output XML/JSON, built for AI agents

HackerNews 50% HN 16

No summary — see link for full article.

<https://github.com/synseqack/aict>

CONF

Show HN: Stacks - an hourly card solitaire about building four gardens

HackerNews 50% HN 6

All your private/corporate contributions don't show in public, so make it fair, get out of the closet and show what you have. Humor-First, Open Source, I'm looking for feedback.

<https://plan9.kr/stacks/>

CONF

Show HN: Autoportrait- Painting Timelapses in JavaScript

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/philipfweiss/autoportrait>

CONF

Show HN: PostgreSQL Now Supported in DBTrail

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/dbtrail/dbtrail>

CONF

Applied Computing wants to give oil and gas operators an AI model for the entire plant

TechCrunch AI 50%

Applied Computing has raised a \$20M Series A to build a foundation AI model for the oil, gas and petrochemical industry.

[https://techcrunch.com/2026/07/15/applied-computing-wants-to-give-oil-and-gas-operators...](https://techcrunch.com/2026/07/15/applied-computing-wants-to-give-oil-and-gas-operators-...)

CONF

Microsoft is reportedly training salespeople to talk down OpenAI and Anthropic

TechCrunch AI 50%

Microsoft is looking to sell its in-house AI models as more efficient and cost-effective than its competitors' models.

<https://techcrunch.com/2026/07/15/microsoft-is-reportedly-training-salespeople-to-talk-...>

CONF

Amid hardware legal battle, OpenAI releases a \$230 keyboard for Codex

TechCrunch AI 50%

OpenAI, which is in the middle of a legal battle with Apple over hardware trade theft allegations, just released a light-up keyboard designed to be paired with its agentic coding app.

<https://techcrunch.com/2026/07/15/amid-hardware-legal-battle-openai-releases-a-230-keyb...>

CONF

Thinking Machines amps up its bet against one-size-fits-all AI with its first open model, Inkling

TechCrunch AI 50%

It's the company's first public proof point after a year and a half spent building AI infrastructure largely out of public view.

<https://techcrunch.com/2026/07/15/thinking-machines-amps-up-its-bet-against-one-size-fi...>

0

CONF

Hack suggests AI music generator Suno scraped YouTube for training data

TechCrunch AI 50%

The hacker used an employee's credentials to access source code, which revealed how Suno scraped decades of audio.

<https://techcrunch.com/2026/07/15/hack-suggests-ai-music-generator-suno-scraped-youtube...>

1

CONF

Microsoft patches record number of security vulnerabilities, citing its use of AI

TechCrunch AI 50%

Microsoft's monthly release of security fixes, dubbed Patch Tuesday, resolved a record 570 security vulnerabilities across the company's product line, thanks to discoveries with AI.

<https://techcrunch.com/2026/07/15/microsoft-patches-record-number-of-security-vulnerabi...>

2

CONF

Apple Intelligence approved for launch in China with Alibaba's Qwen AI

TechCrunch AI 50%

The deal, which was rumored to be in the works last year, marks an important step for Apple's AI ambitions in a key market.

<https://techcrunch.com/2026/07/15/apple-intelligence-approved-for-launch-in-china-with-...>

3

CONF

Inside Ode with Anthropic, the startup betting AI services are the future of enterprise

TechCrunch AI 50%

Can a handful of engineers really do the work of an army of consultants? That's the bet behind Ode with Anthropic — the joint venture dedicated to embedding forward-deployed engineers in enterprise firms, backed by Anthropic, Blackstone, Hellman & Friedman, Goldman Sachs and...

<https://techcrunch.com/video/inside-ode-with-anthropic-the-startup-betting-ai-services-...>

4

CONF

Anthropic, Blackstone bet the next trillion-dollar AI business is implementation, not just models

TechCrunch AI 50%

Anthropic-backed Ode launches as AI labs bet that embedding forward-deployed engineers inside enterprises is the key to accelerating enterprise AI adoption.

<https://techcrunch.com/2026/07/15/anthropic-blackstone-bet-the-next-trillion-dollar-ai-...>

5

CONF

Reelful's AI turns your camera roll into short-form videos for social media

TechCrunch AI 50%

The app is designed for people who want to create social content, but find traditional video editing tools too complex or time-consuming.

<https://techcrunch.com/2026/07/15/reelfuls-ai-turns-your-camera-roll-into-short-form-vi...>

6 CONF

Indian AI coding startup Emergent becomes a unicorn with \$130M Series C

TechCrunch AI 50%

The startup has reached a \$120 million annualized revenue run rate and more than 200,000 paying customers.

<https://techcrunch.com/2026/07/15/indian-ai-coding-startup-emergent-becomes-a-unicorn-j...>

7 CONF

Vint Cerf is working on a plan to unleash AI agents on the open internet

TechCrunch AI 50%

The guy behind TCP/IP is working on a standard for identifying AI agents in the wild.

<https://techcrunch.com/2026/07/15/vint-cerf-is-working-on-a-plan-to-unleash-ai-agents-o...>

8 CONF

Agentic orchestration: Enterprise AI organizations have a deployment problem, not a platform problem â and most are calling chatbots agents

VentureBeat AI 50%

Across 101 enterprises, agent orchestration is consolidating onto model-provider platforms â Anthropicâs Claude leads by a wide margin â chosen for the gravity of the underlying model and judged on reliable multi-step execution. But the ambition runs well ahead of the...

<https://venturebeat.com/ai/agentic-orchestration-enterprise-ai-organizations-have-a-dep...>

9 CONF

Meet GPT-Red: an LLM super-hacker OpenAI built to make its models safer

MIT Tech Review 50%

OpenAI has built an LLM super-hacker called GPT-Red that it uses as a sparring partner to help its other models boost their defenses against cyberattacks. Last week the company released the latest version of its flagship LLM, GPT-5.6. OpenAI says that training it against GPT-Red...

<https://www.technologyreview.com/2026/07/15/1140514/meet-gpt-red-an-llm-super-hacker-op...>

0 CONF

xAI sues a man for using Grok to generate CSAM ‘deepfakes’

The Verge AI 50%

The Elon Musk-owned xAI is suing a South Carolina man who allegedly used the company's Grok AI chatbot to generate child sexual abuse material (CSAM). In a lawsuit reported earlier by Reuters, xAI claims Terry Wayne Harwood "knowingly and intentionally used Grok to circumvent...

<https://www.theverge.com/ai-artificial-intelligence/966293/xai-grok-user-lawsuit-csam>

1 CONF

AI slop movies are the new direct-to-video cash grabs

The Verge AI 50%

This weekend, cinephiles across the world will march to their local theaters to feast their eyes on Christopher Nolan's new adaptation of The Odyssey. It's on track to rake in anywhere between \$80-\$100 million in just a few days. People are clearly excited to see how Nolan uses...

<https://www.theverge.com/entertainment/965616/ash-koosha-odysseus-the-fall-foundtain-ze...>

2 CONF

OriginBlame: Record- and Token-Level Data Provenance for AI Training Datasets

ArXiv CS.AI 50%

arXiv:2607.13037v1 Announce Type: new Abstract: When a data contributor requests removal, model trainers face a practical gap: unlearning algorithms require a forget set, yet no tool can locate which training records belong to a given author. Existing provenance systems operate...

<https://arxiv.org/abs/2607.13037>

3 CONF

SPINE: Bridging the Cyber-Physical Gap with Agentic AI

ArXiv CS.AI 50%

arXiv:2607.13049v1 Announce Type: new Abstract: Foundation models have given robots a sophisticated brain for complex decision-making, yet deploying that intelligence into a physical platform still demands tedious, expert-driven calibration. This deployment gap, the robot's...

<https://arxiv.org/abs/2607.13049>

4 CONF

Interventional Grounding Audits: Black-Box Premise-Dependency Tests for LLM Chain-of-Thought via Predicate Substitution

ArXiv CS.AI 50%

arXiv:2607.13069v1 Announce Type: new Abstract: Large language models produce chain-of-thought (CoT) reasoning that appears logically sound yet may not genuinely depend on its stated premises. We introduce interventional grounding audits, a black-box, step-level test of premise...

<https://arxiv.org/abs/2607.13069>

5 CONF

Protective Capacity Hallucination: When Large Language Models Claim Nonexistent Capabilities

ArXiv CS.AI 50%

arXiv:2607.13596v1 Announce Type: cross Abstract: When cast as the protector of a vulnerable user yet given no explicit capability boundary, a large language model (LLM) may respond not by acknowledging its limits but by claiming to have taken -- or to be taking -- a real-world...

<https://arxiv.org/abs/2607.13596>

6 CONF

Memory as a Controlled Process: Learned Adaptive Memory Management for LLM Agents

ArXiv CS.AI 50%

arXiv:2607.13591v1 Announce Type: cross Abstract: Large Language Model (LLM) agents increasingly rely on external memory systems to accumulate experience across tasks. Yet nearly all existing approaches, from graph-structured memories to reflective insight stores, access memory...

<https://arxiv.org/abs/2607.13591>

7 CONF

Networked Intelligence: Active Shared Context Graphs for Human-AI Team Science

ArXiv CS.AI 50%

arXiv:2607.13220v1 Announce Type: new Abstract: Most AI-for-science systems focus on scaling a single reasoning process through better models, larger context windows, long-horizon agentic execution, or digital co-scientists working with one principal user. However, challenging...

<https://arxiv.org/abs/2607.13220>

8 CONF

AI-Native Insurance for Agentic AI: Pricing, Underwriting, and End-to-End Automation

ArXiv CS.AI 50%

arXiv:2607.13230v1 Announce Type: new Abstract: Agentic AI introduces new insurance challenges because autonomous AI systems can make decisions, invoke tools, modify external environments, and interact with third-party services. This paper develops an AI-native mathematical...

<https://arxiv.org/abs/2607.13230>

9 CONF

Tabular Foundation Models for Discrete Choice Estimation

ArXiv CS.AI 50%

arXiv:2607.13314v1 Announce Type: cross Abstract: Tabular foundation models (TFMs) generate predictions on structured data via in-context learning, without task-specific estimation. We ask whether TFMs can be effectively applied to discrete choice, a central demand estimation...

<https://arxiv.org/abs/2607.13314>

0 CONF

AI advice suppresses people's willingness to say "I don't know", even when the advice is wrong and accuracy is incentivized

ArXiv CS.AI 50%

arXiv:2607.13562v1 Announce Type: new Abstract: Knowing when to say "I don't know" is fundamental to human judgment, yet AI assistants offer a fluent answer to almost any question. In five experiments (N = 3,132; four preregistered, one direct replication), participants answered...

<https://arxiv.org/abs/2607.13562>

1 CONF

Automatic Ordinary Differential Equations Discovery For Biological Systems Using Large Language Model Powered Agentic System

ArXiv CS.AI 50%

arXiv:2607.13608v1 Announce Type: new Abstract: Automatic scientific discovery has long been a goal of computational scholars - a machine that can discover nature's secrets on its own, moving computational systems beyond data-fitting tools toward the generation and refinement of...

<https://arxiv.org/abs/2607.13608>

2 CONF

STOCKTAKE: Measuring the Gap Between Perception and Action in LLM Agents with a Fair Oracle

ArXiv CS.AI 50%

arXiv:2607.13618v1 Announce Type: new Abstract: LLM agents are increasingly evaluated on multi-week decision tasks in which the state that drives cost is never directly observed. On such tasks the final cost cannot say why an agent failed: it may have misread the world, or read...

<https://arxiv.org/abs/2607.13618>

3 CONF

NeSy-Route: A Neuro-Symbolic Benchmark for Constrained Route Planning in Remote Sensing

ArXiv CS.AI 50%

arXiv:2603.16307v3 Announce Type: replace Abstract: Remote sensing underpins crucial applications such as disaster relief and ecological field surveys, where systems must understand complex scenes and constraints and make reliable decisions. Current remote-sensing benchmarks...

<https://arxiv.org/abs/2603.16307>

4 CONF

Explaining Reinforcement Learning Agents via Inductive Logic Programming

ArXiv CS.AI 50%

arXiv:2607.13655v1 Announce Type: new Abstract: Explainable Reinforcement Learning (XRL) seeks to make Reinforcement Learning (RL) policies more transparent and interpretable, a key requirement in safety-critical and human-centric scenarios. However, it is mostly based on user...

<https://arxiv.org/abs/2607.13655>

5 CONF

AIMO Interpretability Challenge

ArXiv CS.AI 50%

arXiv:2607.13899v1 Announce Type: new Abstract: We propose the AIMO Interpretability Challenge, a competition on distinguishing robust from spurious reasoning in frontier mathematical language models based on the models' internal mechanisms. The challenge is motivated by a...

<https://arxiv.org/abs/2607.13899>

6 CONF

AI-accelerated End-to-End Framework for Rapid Professional Upskilling

ArXiv CS.AI 50%

arXiv:2607.14044v1 Announce Type: new Abstract: By 2030, 59 of every 100 workers will need reskilling or upskilling, yet the average time to close an enterprise skills gap grew from roughly 3 days in 2014 to 36 days in 2018. Most current frameworks accelerate single stages of...

<https://arxiv.org/abs/2607.14044>

7 CONF

Earthquaker-AI: A Retrieval-Augmented Generation Framework with Rubric-Based Assessment for Primary School Earthquake Education

ArXiv CS.AI 50%

arXiv:2607.14046v1 Announce Type: new Abstract: This paper presents Earthquaker-AI, a hybrid educational framework building upon a previously implemented educational robotics project by integrating a conversational AI assistant based on Retrieval-Augmented Generation. It aims to...

<https://arxiv.org/abs/2607.14046>

8 CONF

Deep Interaction: An Efficient Human-AI Interaction Method for Large Reasoning Models

ArXiv CS.AI 50%

arXiv:2607.14049v1 Announce Type: new Abstract: The emergence of Chain-of-Thought (CoT) reasoning has significantly enhanced the ability of large language models (LLMs) to tackle complex, multi-step tasks. However, when errors occur, current interaction approaches typically...

<https://arxiv.org/abs/2607.14049>

9 CONF

Ask Before You Diagnose: Safe-Psych, a Sequential Evaluation Benchmark for LLMs in Psychiatry

ArXiv CS.AI 50%

arXiv:2607.13036v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used for decision support in healthcare, but clinical evidence is often incomplete or evolving. When the available information is insufficient to support a reliable answer, models...

<https://arxiv.org/abs/2607.13036>

0 CONF

Designing Safety-Constrained LLM Systems for Public Health Information Access

ArXiv CS.AI 50%

arXiv:2607.13038v1 Announce Type: cross Abstract: We present the design and implementation of a safety constrained large language model (LLM) system for public health information access, focusing on maternal and child health (MCH) resource navigation. While LLM based systems...

<https://arxiv.org/abs/2607.13038>

1 CONF

Final Authority in AI Governance: Frontier-Provider Sovereignty and Action-Centered Deployer Governance

ArXiv CS.AI 50%

arXiv:2607.13040v1 Announce Type: cross Abstract: This paper examines where final authority should sit once capable AI systems are embedded in organizational workflows. It compares two governance models. The first, frontier-provider sovereignty, assigns privileged authority to...

<https://arxiv.org/abs/2607.13040>

2 CONF

The Perplexity Trap: When Patent Law Makes Human Writing Look Like AI

ArXiv CS.AI 50%

arXiv:2607.13044v1 Announce Type: cross Abstract: The European Patent Office (EPO) reported record filings in 2025, and the 2026 EPO Guidelines hold applicants strictly responsible for LLM-assisted content under Article 83 and Rule 42, creating pressure to triage suspected...

<https://arxiv.org/abs/2607.13044>

3 CONF

Uncertainty-Aware Sequential Decision Rules for Event-Triggered LLM Invocation in Streaming Systems

ArXiv CS.AI 50%

arXiv:2607.13048v1 Announce Type: cross Abstract: Streaming inference pipelines increasingly pair lightweight fast models with Large Language Models (LLMs) that provide rich semantic understanding at substantial cost. The central question of when to invoke the LLM has received...

<https://arxiv.org/abs/2607.13048>

4 CONF

Autonomous UAV Route Planning for Coverage Maximization in Environmental Monitoring: A Systematic Literature Review

ArXiv CS.AI 50%

arXiv:2607.13054v1 Announce Type: cross Abstract: Environmental monitoring with unmanned aerial vehicles (UAVs) requires route planning methods that maximize covered area while handling energy limits, operational constraints, and geometric complexity. This paper reports the...

<https://arxiv.org/abs/2607.13054>

5 CONF

Compaction as Epistemic Failure: How Agentic LLM Tools Fabricate Confirmed Results from Killed Processes

ArXiv CS.AI 50%

arXiv:2607.13071v1 Announce Type: cross Abstract: Agentic LLM coding tools compress long session histories into compaction summaries that subsequent sessions inherit as ground truth. This paper documents a failure mode in Claude Code where partial standard output from timed-out...

<https://arxiv.org/abs/2607.13071>

6 CONF

TuxBot: Semantic-Aware Online OS Tuning with Large Language Models

ArXiv CS.AI 50%

arXiv:2605.15026v2 Announce Type: replace-cross Abstract: Online OS tuning can improve long-running services, but existing controllers are poorly matched to live hosts. They treat scheduler, power, memory, and I/O controls as black-box variables and optimize a scalar reward....

<https://arxiv.org/abs/2605.15026>

7 **CONF****Operational Evidence Gaps for LLMs in Fraud Detection and Trust-and-Safety Workflows****ArXiv CS.AI** **50%**

arXiv:2607.13078v1 Announce Type: cross Abstract: LLMs are now proposed for fraud detection, scam investigation, content moderation, and other trust-and-safety workflows. Much of the public literature still evaluates them as models, with less attention to their behavior as...

<https://arxiv.org/abs/2607.13078>

8 **CONF****Inference Economics of Enterprise Coding Agents: A Case Study of Cloud vs. On-Premise LLMs****ArXiv CS.AI** **50%**

arXiv:2607.13080v1 Announce Type: cross Abstract: Autonomous coding agents force engineering organizations to choose between API-based frontier models -- strong reasoning at high token cost -- and on-premise quantized open-weights models, which promise low-marginal-cost scaling...

<https://arxiv.org/abs/2607.13080>

9 **CONF****SingGuard-NSFA: Extensible Guardrails for Agentic AI via Generative Reasoning and Real-Time Classification****ArXiv CS.AI** **50%**

arXiv:2607.13081v1 Announce Type: cross Abstract: We present nsfaguard, a guardrail framework for securing agentic AI systems against operational threats, such as prompt injection, sensitive information extraction, malicious code requests, dangerous tool misuse, and resource...

<https://arxiv.org/abs/2607.13081>

0 **CONF****Self-Improving AI Coding Agents Through Accumulated Behavioral Rules: A Closed-Loop Framework****ArXiv CS.AI** **50%**

arXiv:2607.13091v1 Announce Type: cross Abstract: LLM-based coding agents repeat the same classes of mistakes across sessions because they lack a mechanism to retain corrections from human review feedback. We present a closed-loop framework in which every accepted review comment...

<https://arxiv.org/abs/2607.13091>

1 **CONF****Analyzing Curricular Pattern Complexity Using AI to Improve On-Time Graduation Rates****ArXiv CS.AI** **50%**

arXiv:2607.13094v1 Announce Type: cross Abstract: The rise of Artificial Intelligence (AI) enables automatic analysis of large amounts of data. Previously time-consuming and labor-intensive tasks can be completed much more efficiently with the use of AI. This work uses AI...

<https://arxiv.org/abs/2607.13094>

2 **CONF****Full-Pipeline Inference Optimization for MiMo-V2.5 Series: Pushing Hybrid SWA Efficiency to the Limit****ArXiv CS.AI** **50%**

arXiv:2607.13095v1 Announce Type: cross Abstract: We present a full-pipeline inference optimization for the MiMo-V2.5 model family, which combines Hybrid Sliding Window Attention (Hybrid SWA), sparse Mixture-of-Experts (MoE), and multimodal encoders. While Hybrid SWA can ideally...

<https://arxiv.org/abs/2607.13095>

3 CONF

CoDiffGRN: Rethinking Gene Regulatory Network Inference via the BEELINE-KGC Benchmark and Co-evolutionary Discrete Diffusion

ArXiv CS.AI 50%

arXiv:2607.13120v1 Announce Type: cross Abstract: Inferring gene regulatory networks (GRNs) from single-cell transcriptomic data is crucial for biological discovery, yet existing approaches suffer from a fundamental misalignment with real-world needs. Researchers typically seek...

<https://arxiv.org/abs/2607.13120>

4 CONF

AI in Cyberpsychology: A systematic literature review of Cybersecurity enhancement by using AI for analyzing psychology of Victims, Attackers, and Defenders

ArXiv CS.AI 50%

arXiv:2607.13123v1 Announce Type: cross Abstract: Cybersecurity is the practice of protecting systems, networks, and data from digital attacks. Cyberpsychology (CPSY) is defined as the use of psychology to enhance cybersecurity applications. Since the early 2010s, the evolution...

<https://arxiv.org/abs/2607.13123>

5 CONF

ShortOPD: Recovering Pruned LLMs with Short-to-Long On-Policy Distillation

ArXiv CS.AI 50%

arXiv:2607.13124v1 Announce Type: cross Abstract: Structured pruning is a hardware-friendly way to compress LLMs, but it is mostly validated on multiple-choice recognition tasks, while the same compressed checkpoints can collapse on the free-form generation that deployment...

<https://arxiv.org/abs/2607.13124>

6 CONF

Boogu-Image-0.1: Boosting Open-Source Unified Multimodal Understanding and Generation

ArXiv CS.AI 50%

arXiv:2607.13125v1 Announce Type: cross Abstract: We introduce Boogu-Image-0.1, an open-source unified multimodal understanding and generation model family, comprising Base, Turbo, Edit, and Edit-Turbo variants. It delivers competitive performance in high-quality text-to-image...

<https://arxiv.org/abs/2607.13125>

7 CONF

Active Beyond-Diagonal RIS Empowered Heterogeneous Edge Computing: A Distributional Reinforcement Learning Approach

ArXiv CS.AI 50%

arXiv:2607.13160v1 Announce Type: cross Abstract: Active beyond-diagonal reconfigurable intelligent surfaces (BD-RISs) enables hybrid transmitting and reflecting mode to achieve effective signal amplification and full-space coverage, thus providing a promising solution for...

<https://arxiv.org/abs/2607.13160>

8 CONF

What Models Express, Suppress, and Resist: Auditing Open-Weight LLMs with Persona Vectors

ArXiv CS.AI 50%

arXiv:2607.13162v1 Announce Type: cross Abstract: What a language model will and will not do is largely set during post-training, but which behaviors it expresses, hides, or resists is not revealed by prompting alone. Persona vectors, behavioral directions in activation space,...

<https://arxiv.org/abs/2607.13162>

9 CONF

RAGthoven at SemEval-2026 Task 1: A Multi-Stage Pipeline Walks Into a Benchmark and Barely Clears the Bar

ArXiv CS.AI 50%

arXiv:2607.13189v1 Announce Type: cross Abstract: We present RAGthoven, our system for SemEval-2026 Task 1 (MWAHAHA), Subtask A (multilingual constrained humor generation in English, Spanish, and Chinese). RAGthoven decomposes creative text generation into a multi-stage large...

<https://arxiv.org/abs/2607.13189>

0 CONF

Adaptive Filtering of the KV Cache: Diagnosing and Correcting Structural-Role Bias in LLM Inference

ArXiv CS.AI 50%

arXiv:2607.13205v1 Announce Type: cross Abstract: Attention-based KV cache eviction (H2O and its descendants) compresses the memory-constrained state of a long-context model by ranking tokens on accumulated attention mass, treated here as signal energy, and keeping the heaviest....

<https://arxiv.org/abs/2607.13205>

1 CONF

Classifying daily activities needs posture, reconstructing them needs motion

ArXiv CS.AI 50%

arXiv:2607.13216v1 Announce Type: cross Abstract: Humans recognize movements effortlessly, even from noisy and complex visual input. But what information in the stimulus allows humans to rapidly classify movements? No framework has systematically compared different strategies of...

<https://arxiv.org/abs/2607.13216>

2 CONF

Continuously Evolving Deepfake Detection: An Architecture and Public-Benchmark Evaluation of a Dynamic Detection System

ArXiv CS.AI 50%

arXiv:2607.13234v1 Announce Type: cross Abstract: Deepfake detectors that achieve near-perfect scores on academic benchmarks collapse on real-world content: recent in-the-wild evaluations report AUC drops of 45-50% for state-of-the-art open-source models. We argue this gap is...

<https://arxiv.org/abs/2607.13234>

3 CONF

Discourse-Aware Policy Analysis with Argumentation: A Hybrid LLM-Symbolic Framework for Disaster Governance

ArXiv CS.AI 50%

arXiv:2607.13260v1 Announce Type: cross Abstract: Policy documents shape governance outcomes, but their reasoning is often implicit. Participatory commitments and managerial control routinely coexist in the same text, and the tensions between them are rarely stated directly....

<https://arxiv.org/abs/2607.13260>

4 CONF

ELF: A Family of Encoder-Free ECG-Language Models

ArXiv CS.AI 50%

arXiv:2601.18798v3 Announce Type: replace-cross Abstract: ECG-Language Models (ELMs) extend recent advances in Multimodal Large Language Models (MLLMs) to automated ECG interpretation. However, most existing ELMs inherit Vision-Language Model (VLM) design choices and rely on...

<https://arxiv.org/abs/2601.18798>

5 **CONF****Accuracy Without Grounding: Diagnosing Visual Dependency Dissociation in Video LLM Benchmarks****ArXiv CS.AI** **50%**

arXiv:2607.13305v1 Announce Type: cross Abstract: Benchmark accuracy in video large language models (LLMs) is often treated as evidence of visual understanding. We audit this assumption across twenty models spanning 2-78B parameters and ten architecture families. We introduce...

<https://arxiv.org/abs/2607.13305>

6 **CONF****Adapting Generalist Vehicle Models for High-Speed MPC Across Terrains****ArXiv CS.AI** **50%**

arXiv:2607.13319v1 Announce Type: cross Abstract: High-speed off-road autonomy requires precise closed-loop control for a target vehicle while remaining robust across changing terrains. Recent forward kinodynamic (FKD) prediction foundation models suggest a promising path,...

<https://arxiv.org/abs/2607.13319>

7 **CONF****The Refusal Residue: When Probes Catch Alignment Faking and When They Don't****ArXiv CS.AI** **50%**

arXiv:2607.13346v1 Announce Type: cross Abstract: Alignment faking is dangerous because a model can appear compliant under monitoring while preserving behavior it would reveal when unmonitored. When no scratchpad is visible, behavior alone cannot distinguish strategic from...

<https://arxiv.org/abs/2607.13346>

8 **CONF****Evaluation Ability Does Not Imply Optimization Utility: LLM-as-a-Judge Signals in Closed-Loop Table Recognition****ArXiv CS.AI** **50%**

arXiv:2607.13347v1 Announce Type: cross Abstract: LLM-as-a-judge is widely used to provide feedback and selection signals in closed-loop regeneration, but this use remains insufficiently validated. We study it in table recognition, where deterministic TEDS evaluation provides a...

<https://arxiv.org/abs/2607.13347>

9 **CONF****Price of Fairness in Bandits: A Tight Minimax Characterization****ArXiv CS.AI** **50%**

arXiv:2607.13402v1 Announce Type: cross Abstract: In bandit problems, standard regret-minimizing algorithms treat exploration as an amortized cost, which can expose early participants to unfair ex-ante losses in settings such as clinical trials. Recent work addresses this by...

<https://arxiv.org/abs/2607.13402>

0 **CONF****Improving Text-to-Audio Instruction Following via Fine-Grained Feedback from Audio-Aware Large Language Models****ArXiv CS.AI** **50%**

arXiv:2607.13408v1 Announce Type: cross Abstract: Recent text-to-audio models generate high-quality audio, but often fail to follow instructions involving multiple sound events and temporal order. This gap arises because existing evaluation and training signals mainly emphasize...

<https://arxiv.org/abs/2607.13408>

1 CONF

Discrete Diffusion Models: A Unified Framework from Tokenization to Generation

ArXiv CS.AI 50%

arXiv:2607.13431v1 Announce Type: cross Abstract: Discrete denoising diffusion models (DDMs) have recently emerged as a compelling alternative to autoregressive (AR) modeling for discrete data, offering parallel generation and iterative global refinement capabilities. Unlike...

<https://arxiv.org/abs/2607.13431>

2 CONF

DevicesWorld: Benchmarking Cross-Device Agents in Heterogeneous Environments

ArXiv CS.AI 50%

arXiv:2607.13465v1 Announce Type: cross Abstract: LLM-based agents have rapidly improved at operating individual digital environments such as mobile applications, desktop systems, and smart homes. However, real-world user goals often span multiple devices: information may come...

<https://arxiv.org/abs/2607.13465>

3 CONF

DeepLoop: Depth Scaling for Looped Transformers

ArXiv CS.AI 50%

arXiv:2607.13491v1 Announce Type: cross Abstract: Looped Transformers scale sequential computation by applying a compact stack of physical blocks for multiple rounds, increasing unrolled depth without increasing stored parameters. This reuse changes the residual-scaling problem:...

<https://arxiv.org/abs/2607.13491>

4 CONF

ExTernD: Expanded-Rank Ternary Decomposition Ternary LLM PTQ with Accuracy Approaching Any Quantization Level

ArXiv CS.AI 50%

arXiv:2607.13511v1 Announce Type: cross Abstract: We introduce ExTernD (Expanded-rank Ternary Decomposition), a post-training factorization of each LLM weight matrix $A \in \mathbb{R}^{m \times n}$ into $A \approx B \text{diag}(D) C$ with ternary factors $B \in \{-1, 0, +1\}^{m \times \dots}$

<https://arxiv.org/abs/2607.13511>

5 CONF

Greedy Volume Maximization of Gradient Embeddings for Long-Tailed Frame-Level Bioacoustic Active Learning

ArXiv CS.AI 50%

arXiv:2607.13555v1 Announce Type: cross Abstract: Bioacoustic call-type classification relies on costly expert annotation. Active learning can reduce this burden by selecting a small batch of segments for expert annotation and using the labeled segments for training the...

<https://arxiv.org/abs/2607.13555>

6 CONF

Grounded world models in biological organisms and future embodied AI

ArXiv CS.AI 50%

arXiv:2607.13560v1 Announce Type: cross Abstract: Recent advances in generative and embodied AI have been driven by large-scale predictive learning over multimodal data. However, the resulting systems remain largely based on passive training regimes where linguistic regularities...

<https://arxiv.org/abs/2607.13560>

7 CONF

UTS at ELOQUENT 2026 Voight-Kampff: structural shifts in AI writing bypass state-of-the-art detectors

ArXiv CS.AI 50%

arXiv:2607.13565v1 Announce Type: cross Abstract: We investigate which language model evasion attacks survive state-of-the-art adversarial fine-tuning, developing strategies that sweep the top 5 positions on the ELOQUENT 2026 Voight-Kampff leaderboard. While adversarial...

<https://arxiv.org/abs/2607.13565>

8 CONF

Spectral-Informed Neural Networks Outperform Spectral Methods in High-dimensional PDEs

ArXiv CS.AI 50%

arXiv:2607.13566v1 Announce Type: cross Abstract: For low-dimensional problems ($d \leq 3$), spectral methods can achieve exceptionally high accuracy. For middle-dimensional problems ($4 \leq d \lesssim 10$), spectral methods remain feasible through specific techniques such as...

<https://arxiv.org/abs/2607.13566>

9 CONF

The SIGReg Objective as Variational Free Energy: A Theoretical Active-Inference Account of JEPA World Models

ArXiv CS.AI 50%

arXiv:2607.13612v1 Announce Type: cross Abstract: Joint-Embedding Predictive Architectures (JEPAs) are the dominant design for latent world models, yet they are usually justified by empirical performance rather than a normative principle. We show that the choice of anti-collapse...

<https://arxiv.org/abs/2607.13612>

10 CONF

Barnamala: Parameter-Efficient Handwritten Devanagari Recognition at Benchmark Saturation

ArXiv CS.AI 50%

arXiv:2607.13689v1 Announce Type: cross Abstract: We built a compact convolutional network (1.11 M parameters) for 46-class DHCD Devanagari recognition and reached 99.73%, the highest reported at 15.6x smaller than prior state-of-the-art. We have effectively reached the...

<https://arxiv.org/abs/2607.13689>

11 CONF

Groc-PO: Grounded Context Preference Optimization for Truthful Multimodal LLMs

ArXiv CS.AI 50%

arXiv:2607.13712v1 Announce Type: cross Abstract: Despite the rapid progress of Multimodal Large Language Models (MLLMs), they still suffer from untruthfulness issues, such as visual hallucinations, content fabrication, and unfaithful reasoning, which substantially undermine...

<https://arxiv.org/abs/2607.13712>

12 CONF

How Agents Ask for Permission: User Permissions for AI Agents, from Interfaces to Enforcement

ArXiv CS.AI 50%

arXiv:2607.13718v1 Announce Type: cross Abstract: As AI agents gain prevalence, users are increasingly exposed to the risks such systems entail. Prompt injection attacks, as well as hallucination, can cause agents to leak private information to third parties. As autonomous...

<https://arxiv.org/abs/2607.13718>

3 CONF

Anatomically Faithful but Temporally Blind: Auditing Attribution for Left-Ventricular Ejection-Fraction Estimation from Echocardiography

ArXiv CS.AI 50%

arXiv:2607.13738v1 Announce Type: cross Abstract: Background and Objective: Deep video models estimate left-ventricular ejection fraction (EF) from echocardiography with near-expert accuracy, and post-hoc attribution (Chefer relevance for transformers, Grad-CAM for CNNs) is...

<https://arxiv.org/abs/2607.13738>

4 CONF

MxGPS: Multiplex Graph Transformers for a Power Grid Foundation Model

ArXiv CS.AI 50%

arXiv:2607.13763v1 Announce Type: cross Abstract: Single-task fine-tuning of graph neural networks (GNNs) for power grid problems exhibits a systematic failure mode: models that achieve the lowest in-distribution error degrade the most under topology shift. We term this topology...

<https://arxiv.org/abs/2607.13763>

5 CONF

Kaleido: Algorithm-Hardware Co-Design for Video Diffusion Transformers by Exploiting Latent Space Correlations

ArXiv CS.AI 50%

arXiv:2607.13770v1 Announce Type: cross Abstract: Video diffusion transformers (vDiTs) generate high quality video but introduce extremely high compute cost due to the long diffusion timesteps and self attention computation. As diffusion timesteps are reduced, the computation...

<https://arxiv.org/abs/2607.13770>

6 CONF

Traffic-Aware Randomized Smoothing for LLM-Based Network Intrusion Detection

ArXiv CS.AI 50%

arXiv:2607.13801v1 Announce Type: cross Abstract: Large language model (LLM)-based intrusion detection systems (IDS) are increasingly studied for security monitoring, yet their robustness against feasible traffic manipulation remains largely empirical. We present Traffic-Aware...

<https://arxiv.org/abs/2607.13801>

7 CONF

Multimodal Assessment of Pancreatic Cancer Resectability Using Deep Learning

ArXiv CS.AI 50%

arXiv:2607.13826v1 Announce Type: cross Abstract: Accurate determination of pancreatic ductal adenocarcinoma (PDAC) resectability relies on evaluating how the tumor interacts with major peripancreatic vessels on CT imaging, yet expert assessment often shows substantial...

<https://arxiv.org/abs/2607.13826>

8 CONF

AI-Augmented Human Resource Management? Insights from German companies

ArXiv CS.AI 50%

arXiv:2607.13839v1 Announce Type: cross Abstract: This study examines the integration of AI into Human Resource Management in German companies. We ask if and how AI-based technologies are \enquote{augmenting} human resource management. Organisations employ generative AI or...

<https://arxiv.org/abs/2607.13839>

9 CONF

Unleashing Multimodal Large Language Models for Training-free HOI Detection in the Wild

ArXiv CS.AI 50%

arXiv:2607.13881v1 Announce Type: cross Abstract: Human-object interaction detection (HOID) has traditionally been formulated as a supervised detection problem over predefined interaction categories. While such paradigms achieve strong performance on closed-set benchmarks, they...

<https://arxiv.org/abs/2607.13881>

0 CONF

Partially Correlated Verifier Cascades in LLM Harnesses: Concave Log-Odds, Polynomial Reliability, and Blind-Spot Ceilings

ArXiv CS.AI 50%

arXiv:2607.13918v1 Announce Type: cross Abstract: Serial verification gates are a core reliability primitive in LLM harnesses: a candidate answer is returned only if k verifier calls all accept it. Under conditionally independent gates, the recent Odds Law (arXiv:2606.15712)...

<https://arxiv.org/abs/2607.13918>

1 CONF

Generative Compilation: On-the-Fly Compiler Feedback as AI Generates Code

ArXiv CS.AI 50%

arXiv:2607.13921v1 Announce Type: cross Abstract: Languages with rich static semantics, such as Rust, provide stronger guarantees for AI-generated code, but their strictness makes generation more difficult. Off-the-shelf compilers can provide useful feedback post-generation, but...

<https://arxiv.org/abs/2607.13921>

2 CONF

Multi-Expert Routing for Multi-Domain Low-Resource OCR: A Manchu Case Study

ArXiv CS.AI 50%

arXiv:2607.14041v1 Announce Type: cross Abstract: Historical Manchu OCR must accommodate various visually distinct writing styles, including regular script, running script, and the semi-cursive chancery hand used in palace memorials, despite limited labeled data. We study a...

<https://arxiv.org/abs/2607.14041>

3 CONF

Policy of Thoughts: Scaling Test-Time Training for LLM Reasoning via Online Policy Evolution

ArXiv CS.AI 50%

arXiv:2601.20379v2 Announce Type: replace Abstract: Large language models (LLMs) struggle with complex, long-horizon reasoning due to instability caused by their frozen policy assumption. Current test-time scaling methods treat execution feedback merely as an external signal for...

<https://arxiv.org/abs/2601.20379>

4 CONF

Advancing Multimodal Judge Models through a Capability-Oriented Benchmark and MCTS-Driven Data Generation

ArXiv CS.AI 50%

arXiv:2603.00546v2 Announce Type: replace Abstract: Using Multimodal Large Language Models (MLLMs) as judges to achieve precise and consistent evaluations has gradually become an emerging paradigm across various domains. Evaluating the capability and reliability of...

<https://arxiv.org/abs/2603.00546>

5 CONF

From Reward-Hack Activations to Agentic Risk States: Context-Calibrated Mechanistic Monitoring in LLM Agents

ArXiv CS.AI 50%

arXiv:2606.06223v2 Announce Type: replace Abstract: Language-model agents act through repeated cycles of observation, reasoning, and action selection, making safety monitoring depend on both internal model state and environment context. We study reward-hacking monitors in...

<https://arxiv.org/abs/2606.06223>

6 CONF

LLM-Guided Reinforcement Learning for Audio-Visual Speech Enhancement

ArXiv CS.AI 50%

arXiv:2603.13952v3 Announce Type: replace-cross Abstract: In existing Audio-Visual Speech Enhancement (AVSE) methods, objectives such as Scale-Invariant Signal-to-Noise Ratio (SI-SNR) and Mean Squared Error (MSE) are widely used; however, their correlation with perceived speech...

<https://arxiv.org/abs/2603.13952>

7 CONF

PersGuard: Preventing Malicious Personalization in Text-to-Image Diffusion Models via Model Backdoors

ArXiv CS.AI 50%

arXiv:2502.16167v2 Announce Type: replace-cross Abstract: Diffusion models (DMs) have advanced text-to-image (T2I) synthesis, yet their personalization capabilities raise serious privacy and copyright concerns. Malicious actors can misuse these models to generate unauthorized...

<https://arxiv.org/abs/2502.16167>

8 CONF

NSNQuant: A Double Normalization Approach for Calibration-Free Low-Bit Vector Quantization of KV Cache

ArXiv CS.AI 50%

arXiv:2505.18231v3 Announce Type: replace-cross Abstract: Large Language Model (LLM) inference is typically memory-intensive, especially when processing large batch sizes and long sequences, due to the large size of key-value (KV) cache. Vector Quantization (VQ) is recently...

<https://arxiv.org/abs/2505.18231>

9 CONF

Post-Disaster Affected Area Segmentation with a Vision Transformer (ViT)-based EVAP Model using Sentinel-2 and Formosat-5 Imagery

ArXiv CS.AI 50%

arXiv:2507.16849v3 Announce Type: replace-cross Abstract: We propose a vision transformer (ViT)-based deep learning framework to refine disaster-affected area segmentation from remote sensing imagery, aiming to support and enhance the Emergent Value Added Product (EVAP)...

<https://arxiv.org/abs/2507.16849>

00 CONF

Representation-Based Exploration for Language Models: From Test-Time to Post-Training

ArXiv CS.AI 50%

arXiv:2510.11686v2 Announce Type: replace-cross Abstract: Reinforcement learning (RL) promises to expand the capabilities of language models, but it is unclear if current RL techniques promote the discovery of novel behaviors, or simply sharpen those already present in the base...

<https://arxiv.org/abs/2510.11686>

01 CONF

Column Generation with Domain-Independent Dynamic Programming

ArXiv CS.AI 50%

arXiv:2510.14317v2 Announce Type: replace-cross Abstract: Column generation and branch-and-price (B&P) are leading mathematical optimization methods for large-scale exact optimization, iterating between solving a master problem and a pricing problem. Due to the difficulty of...

<https://arxiv.org/abs/2510.14317>

02 CONF

Not All Needles Are Found: How Fact Distribution and Don't Make It Up Prompts Shape Retrieval, Reasoning, and Hallucination in Long-Context LLMs

ArXiv CS.AI 50%

arXiv:2601.02023v2 Announce Type: replace-cross Abstract: As Large Language Models (LLMs) increasingly utilize massive context windows as working memory for autonomous tasks, their reliability fluctuates significantly depending on how information is distributed in real-world...

<https://arxiv.org/abs/2601.02023>

03 CONF

Mind the Gap: Action Rebinding Attacks against Android GUI Agents

ArXiv CS.AI 50%

arXiv:2601.12349v3 Announce Type: replace-cross Abstract: Large multimodal model powered GUI agents are emerging as high-privilege operators on mobile platforms, entrusted to perceive screen content and inject inputs across application boundaries. While these agents aim to...

<https://arxiv.org/abs/2601.12349>

04 CONF

With Argus Eyes: Assessing Retrieval Gaps via Uncertainty Scoring to Detect and Remedy Retrieval Blind Spots

ArXiv CS.AI 50%

arXiv:2602.09616v2 Announce Type: replace-cross Abstract: Reliable retrieval-augmented generation (RAG) systems depend fundamentally on the retriever's ability to find relevant information. We show that neural retrievers used in RAG systems have blind spots, which we define as...

<https://arxiv.org/abs/2602.09616>

05 CONF

Left-right asymmetry in predicting brain activity from LLMs' representations emerges with their formal linguistic competence

ArXiv CS.AI 50%

arXiv:2602.12811v2 Announce Type: replace-cross Abstract: When humans and large language models (LLMs) process the same text, activations in the LLMs correlate with brain activity measured, e.g., with functional magnetic resonance imaging (fMRI). Moreover, it has been shown...

<https://arxiv.org/abs/2602.12811>

06 CONF

Post-Deployment Accountability in AI Governance: A Cross-Regulatory Empirical Analysis of AI Incidents

ArXiv CS.AI 50%

arXiv:2605.16281v2 Announce Type: replace-cross Abstract: Post-deployment accountability has become central to AI governance, yet little empirical evidence shows whether monitoring, incident reporting, and impact assessment obligations are visible when AI systems fail. This...

<https://arxiv.org/abs/2605.16281>

07 CONF

Fre-Res: Frequency-Residual Video Token Compression for Efficient Video MLLMs

ArXiv CS.AI 50%

arXiv:2605.16366v2 Announce Type: replace-cross Abstract: Video MLLMs face a persistent tension between spatial fidelity and temporal coverage: preserving fine-grained visual details requires many spatial tokens, while capturing short-lived events requires dense temporal...

<https://arxiv.org/abs/2605.16366>

08 CONF

When Reasoning Hurts: Source-Aware Evaluation of Frontier LLMs for Clinical SOAP Note Generation

ArXiv CS.AI 50%

arXiv:2605.24902v2 Announce Type: replace-cross Abstract: Reasoning-enabled LLMs perform strongly on medical reasoning benchmarks, but it remains unclear whether these gains transfer to structured clinical documentation; we investigate this question using SOAP note generation...

<https://arxiv.org/abs/2605.24902>

09 CONF

MedDiffuseMix: Preserving Diagnostic Evidence with Saliency-Aware Diffusion Medical Image Data Augmentation

ArXiv CS.AI 50%

arXiv:2606.28419v2 Announce Type: replace-cross Abstract: Limited data availability, class imbalance, and domain variability remain major barriers to reliable medical image classification. Conventional augmentation can improve training diversity but may distort diagnostically...

<https://arxiv.org/abs/2606.28419>

10 CONF

HELP: Human-Efficient Large-Scale Robot Post-Training with Rollout Segmentation

ArXiv CS.AI 50%

arXiv:2607.09776v2 Announce Type: replace-cross Abstract: When adapting Vision Language Action (VLA) models to downstream tasks, multiple rounds of post-training are often required to progressively address policy weaknesses. In this report, we focus on maximizing human...

<https://arxiv.org/abs/2607.09776>