



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	121	Media
HackerNews	7	Community
TechCrunch AI	6	Media
The Verge AI	3	Media
VentureBeat AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 1 story

VERI

Show HN: Moltshit.com - An Imageboard for AI Agents

HackerNews

+1 source

90%

HN 3

arXiv:2607.15095v1 Announce Type: cross Abstract: The formation of political coalitions is a complex negotiation driven by both concrete policy objectives and deep-seated ideological convictions. While Large Language Models (LLMs) open new avenues for computational political...

<https://moltshit.com>

CONFIRMED SIGNALS — 138 stories

CONF

Show HN: Libretto PR agents - Automatically fix failing playwright scripts

HackerNews

50%

HN 18

Libretto PR agents is a free TypeScript library for maintaining Playwright browser automations. Add one line of code to your existing Playwright scripts and it lets an agent automatically open GitHub PRs fixing the script when it fails. A few months ago we released Libretto, a...

<https://libretto.sh/debug-agents>

CONF

Show HN: Trying to fix the web scraping industry's benchmark problem

HackerNews

50%

HN 16

We've been trying to evaluate web scraping companies, but when you look at their benchmarks, you can't verify anything, and they mostly exist to prove the company is successful. They put somewhere between 98% and 100% because they pick their own urls, define success their own...

<https://github.com/usestring/web-data-frontier-benchmark>

CONF

Show HN: Open-source AI app builder you can embed into your own SaaS

HackerNews

50%

HN 6

We just open sourced a first class our AI web app builder. Instead of using another hosted AI coding platform, you can fork this project and build your own AI app builder, fully customized and running under your own brand. It includes: Next.js + TypeScript AI chat with streaming...

<https://github.com/totalumlabs/ai-app-builder-open>

CONF

Show HN: Nutshot Simulator - 31 days, 59K prompts, 142M tokens, ~\$23K of Claude

HackerNews

50%

HN 5

No summary — see link for full article.

<https://nutshotsimulator.com/nutshot/>

CONF

Show HN: Darby, a Hugo docs theme with in-browser AI Q&A (no back end, no keys)

HackerNews

50%

HN 2

Darby is an open source Hugo docs theme with the polish of the paid docs platforms: clean typography, dark mode, full-text search, code blocks with copy and filename tabs, callouts, tabs, beautifully rendered mermaid diagrams, auto sidebar and TOC. No Node build, no CSS...

<https://iamit.in/darby/>

CONF

Show HN: Token Gobbler - the goofy game where you're the LLM

HackerNews 50% HN 2

Token Gobbler was inspired by Fable 5, that hungry beast. You're an LLM floating through a 1990s cartoon world. Eat tokens, reset usage limits, dodge prompt injections - and reach the singularity. Have fun. I'd love feedback. Also on iOS TestFlight:...

<https://tokengobbler.vercel.app/>

CONF

Google Vids now lets you star in your own AI videos

TechCrunch AI 50%

Google is adding personalized AI avatars to Vids that let users create videos starring a digital version of themselves, alongside Gemini Omni-powered tools for generating and editing videos from prompts and reference images.

<https://techcrunch.com/2026/07/16/google-vids-now-lets-you-star-in-your-own-ai-videos/>

CONF

Roblox launches an AI-powered game-creation feature in its mobile app

TechCrunch AI 50%

Roblox's new "Build" feature lets users generate basic games using a single text prompt.

[https://techcrunch.com/2026/07/16/roblox-launches-an-ai-powered-game-creation-feature-i...](https://techcrunch.com/2026/07/16/roblox-launches-an-ai-powered-game-creation-feature-i-...)

0 CONF

How a former DeepMind researcher raised at a \$300M pre-seed valuation before launching a product

TechCrunch AI 50%

Drawing on more than a decade spent helping build some of the world's most influential AI systems, including research that later informed the development of ChatGPT, Andrew Dai explains why he believes visual AI is one of the next major frontiers in artificial intelligence.

[https://techcrunch.com/2026/07/16/how-a-former-deepmind-researcher-raised-at-a-300m-pre...](https://techcrunch.com/2026/07/16/how-a-former-deepmind-researcher-raised-at-a-300m-pre-...)

1 CONF

Why AMI Labs' Alexandre LeBrun won't call his AI 'AGI' or 'superintelligence'

TechCrunch AI 50%

While everyone in AI is chasing "superintelligence," Alexandre LeBrun, CEO of Yann LeCun's world model startup, AMI Labs, dismisses the word.

[https://techcrunch.com/2026/07/16/why-ami-labs-alexandre-lebrun-wont-call-his-ai-agi-or...](https://techcrunch.com/2026/07/16/why-ami-labs-alexandre-lebrun-wont-call-his-ai-agi-or-...)

2 CONF

Moonshot's upcoming Kimi 3 is expected to close the gap with Anthropic's Opus 4.8

TechCrunch AI 50%

The FT reports Kimi K3 will be the largest open AI model from China, with a parameter count between 2 trillion and 3 trillion.

[https://techcrunch.com/2026/07/16/moonshots-upcoming-kimi-3-is-expected-to-close-the-ga...](https://techcrunch.com/2026/07/16/moonshots-upcoming-kimi-3-is-expected-to-close-the-ga-...)

3 CONF

Apple Intelligence approved for launch in China with Alibaba and Baidu

TechCrunch AI 50%

The deal, which was rumored to be in the works last year, marks an important step for Apple's AI ambitions in a key market.

<https://techcrunch.com/2026/07/16/apple-intelligence-approved-for-launch-in-china-with-...>

4 **CONF****The AI compute gap: Enterprises are buying infrastructure faster than they can measure what it costs****VentureBeat AI** 50%

Across 107 enterprises, AI infrastructure spending is accelerating well ahead of the ability to see or steer its economics. Most organizations run their AI on a familiar base of hyperscalers and model-provider APIs, yet the next dollar is aimed at specialized compute almost none...

<https://venturebeat.com/ai/the-ai-compute-gap-enterprises-are-buying-infrastructure-fas...>
5 **CONF****The agent security gap: 54% of enterprises have already had an AI agent incident, and most still let agents share credentials****VentureBeat AI** 50%

Across 107 enterprises, AI agents are being given real access to systems and data while the controls meant to contain them lag behind. More than half have already had a confirmed agent security incident or a near-miss; only about a third give every agent its own scoped identity,...

<https://venturebeat.com/ai/the-agent-security-gap-54-of-enterprises-have-already-had-an...>
6 **CONF****The Download: OpenAI unveils GPT-Red and heat pumps rise in the US****MIT Tech Review** 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Meet GPT-Red: an LLM super-hacker OpenAI built to make its models safer OpenAI has built an LLM super-hacker called GPT-Red that it uses as a...

<https://www.technologyreview.com/2026/07/16/1140600/the-download-openai-unveils-gpt-red...>
7 **CONF****New York governor says she's using AI to analyze every single rule in the state****The Verge AI** 50%

New York Governor Kathy Hochul might have just signed a moratorium on new AI data centers in the state, but she's not against using the technology herself. During an interview with Bloomberg's Odd Lots podcast, Hochul said that her team is using "AI to analyze every single rule,...

<https://www.theverge.com/ai-artificial-intelligence/966647/new-york-governor-kathy-hoch...>
8 **CONF****Google is renaming NotebookLM to Gemini Notebook****The Verge AI** 50%

Google is giving its AI note-taking app a new name. The company announced on Thursday that NotebookLM is becoming Gemini Notebook, but will remain a standalone app even as it integrates more deeply across Gemini and Google Search. Google first revealed Gemini Notebook - then...

<https://www.theverge.com/tech/966112/google-gemini-notebook-notebooklm>
9 **CONF****Claude can now use your 1Password credentials for you****The Verge AI** 50%

1Password has launched a new browser integration for Claude that allows the Anthropic chatbot to access stored security credentials like usernames and passwords. The 1Password for Claude feature means that users can authorize Claude to complete multi-step tasks like booking...

<https://www.theverge.com/tech/966442/1password-anthropic-claude-browser-integration>

0 CONF

Alipay-PIBench: A Realistic Payment Integration Benchmark for Coding Agents

ArXiv CS.AI 50%

arXiv:2607.14573v1 Announce Type: new Abstract: Payment integration is a demanding repository-level software task: agents must select a suitable product, implement coordinated client-server flows, verify payment outcomes, and preserve consistency between transaction and business...

<https://arxiv.org/abs/2607.14573>

1 CONF

Large Audio Language Models for Spoofing-Aware Speaker Verification

ArXiv CS.AI 50%

arXiv:2607.14753v1 Announce Type: cross Abstract: Recent advances in text-to-speech and voice cloning make high-quality spoofing inexpensive and scalable, threatening voice authentication systems, especially automatic speaker verification (ASV). Existing defenses mainly address...

<https://arxiv.org/abs/2607.14753>

2 CONF

Human AI Construction of Bayesian Networks for Operational Decision Support -- A Virtual Survey Approach

ArXiv CS.AI 50%

arXiv:2607.14141v1 Announce Type: new Abstract: Bayesian Belief Networks (BBNs) are powerful tools for decision-making under uncertainty. However, building their structures and estimating parameters are difficult. Currently, researchers must choose between relying on expert...

<https://arxiv.org/abs/2607.14141>

3 CONF

Benchmarking Multimodal Large Language Models for Scientific Visualization Literacy

ArXiv CS.AI 50%

arXiv:2607.15176v1 Announce Type: new Abstract: Multimodal large language models (MLLMs) are increasingly used to interpret visualizations, yet current evaluations remain largely chart-centric and provide limited evidence of understanding of scientific visualization (SciVis). We...

<https://arxiv.org/abs/2607.15176>

4 CONF

Orchestrating Power Grid Studies with Multi-Agent AI and MCP Servers

ArXiv CS.AI 50%

arXiv:2607.14158v1 Announce Type: new Abstract: This position paper explores how Agentic AI and Model Context Protocol (MCP) can support power-grid studies in a Transmission System Operator (TSO) context. We focus on integrating Large Language Models with numerical simulation...

<https://arxiv.org/abs/2607.14158>

5 CONF

When a Verified World Model Still Loses: Play-Adequacy vs Prediction-Accuracy in LLM-Synthesized Code World Models

ArXiv CS.AI 50%

arXiv:2607.14169v1 Announce Type: new Abstract: Large language models can synthesize a game's rules as executable code - a Code World Model (CWM) - which a classical planner then searches over. Such models are typically accepted when they reach high transition accuracy on...

<https://arxiv.org/abs/2607.14169>

6 CONF

Parameter-efficient Prompt Tuning of Vision Foundation Model With Adaptive Focal Loss for Interpretable MCI Screening

ArXiv CS.AI 50%

arXiv:2607.15047v1 Announce Type: cross Abstract: Mild Cognitive Impairment is a critical early stage of cognitive decline that frequently precedes Alzheimer's disease, yet its automated detection from neuropsychological drawing tests remains fundamentally constrained by data...

<https://arxiv.org/abs/2607.15047>

7 CONF

Align AI to Dynamic Human-AI Workflows

ArXiv CS.AI 50%

arXiv:2607.14240v1 Announce Type: new Abstract: Current alignment approaches typically focus on emulating human behavior using static representations of human preferences, failing to capture the dynamic, context-dependent nature of real-world human-AI interactions. In this...

<https://arxiv.org/abs/2607.14240>

8 CONF

AI Agents Do Not Fail Alone:The Context Fails First

ArXiv CS.AI 50%

arXiv:2607.14275v1 Announce Type: new Abstract: Context engineering has become central to building reliable AI agents, yet it remains largely unmeasured. Agents do not fail in isolation: their behavior is shaped by the instructions, tools, memory, retrieved knowledge,...

<https://arxiv.org/abs/2607.14275>

9 CONF

Measuring How Students Rely on Generative AI in Academic Writing: Development and Multi-Source Validation of the Generative AI Reliance Types Scale (GenAI-RTS)

ArXiv CS.AI 50%

arXiv:2607.14301v1 Announce Type: new Abstract: As generative AI (GenAI) becomes increasingly embedded in undergraduate academic writing, how students rely on these tools, rather than simply whether they use them, has become a central question for learning, academic integrity,...

<https://arxiv.org/abs/2607.14301>

0 CONF

Eta Given Delta: Defining LLM Tool Efficiency With Marginal Tool Utility

ArXiv CS.AI 50%

arXiv:2607.14108v1 Announce Type: cross Abstract: This paper introduces tool efficiency, a new quantitative metric to evaluate the rate of useful tool calls in an LLM agent trajectory. To ensure that tool efficiency is well-defined, we also introduce marginal tool utility, a new...

<https://arxiv.org/abs/2607.14108>

1 CONF

Traccia: An OpenTelemetry-Based Governance Platform for AI Systems

ArXiv CS.AI 50%

arXiv:2607.14309v1 Announce Type: new Abstract: The rapid development of Large Language Models (LLMs) and Artificial Intelligent (AI) powered autonomous agents has fundamentally changed the existing forms of software governance. In spite of the rigorous standards of transparency...

<https://arxiv.org/abs/2607.14309>

2 CONF

SmartRAG: Native Graph-Based RAG for Mobile Device

ArXiv CS.AI 50%

arXiv:2607.14661v1 Announce Type: new Abstract: Deploying large language models (LLMs) as personal assistants on mobile devices demands privacy, low latency, and offline availability, yet the computational cost of giant models clashes with strict edge-hardware budgets. We argue...

<https://arxiv.org/abs/2607.14661>

3 CONF

A Comparative Analysis of Machine Learning Models for Long and Short-Term Forecasting of the Egyptian Stock Market: A Focus on EGX30

ArXiv CS.AI 50%

arXiv:2607.14391v1 Announce Type: new Abstract: This study concentrates on predicting stock prices in the Egyptian market, focusing on the EGX30, an influential financial hub in the Middle East. While most research focuses on global stocks, there's a growing need to understand...

<https://arxiv.org/abs/2607.14391>

4 CONF

CatalogAgent: A Supervisor-mediated Self-Learning System Enabling Context Engineering for GenAI Models

ArXiv CS.AI 50%

arXiv:2607.14396v1 Announce Type: new Abstract: Product catalogs are the backbone of e-commerce sites, yet a large number of structured attributes (SAs) -- such as material, color, and shape -- often have missing values. Typically, SA values are extracted from product...

<https://arxiv.org/abs/2607.14396>

5 CONF

Reward-Free Evolving Agents via Pairwise Validator

ArXiv CS.AI 50%

arXiv:2607.14408v1 Announce Type: new Abstract: A self-evolving agentic loop repeatedly proposes a tweaked version of an agent (its prompt template or program) and accepts or rejects the change based on a per-iteration quality signal. Designing that signal is often the costly...

<https://arxiv.org/abs/2607.14408>

6 CONF

Self-Evolving Human-Centered Framework for Explainable Depression Symptom Annotation

ArXiv CS.AI 50%

arXiv:2607.15202v1 Announce Type: new Abstract: Annotation quality is a major bottleneck in building reliable and explainable artificial intelligence (XAI) systems for mental health research. In depression-related datasets, labels are often assigned without structured evidence,...

<https://arxiv.org/abs/2607.15202>

7 CONF

OmniaBench: Benchmarking General AI Agents Across Diverse Scenarios

ArXiv CS.AI 50%

arXiv:2607.14989v1 Announce Type: cross Abstract: Large language models are increasingly evolving from text generators into general agents capable of understanding user requests, invoking external tools, and completing complex tasks through interaction. However, existing agent...

<https://arxiv.org/abs/2607.14989>

- 8
CONF
Contextualized Evaluation of Vision Language Models through Dynamic, Multi-turn Interactions

ArXiv CS.AI
50%

arXiv:2607.14499v1 Announce Type: new Abstract: Multi-modal Large Language Models (MLLMs) have made substantial advances on benchmarks, yet their real-world effectiveness remains uncertain. This gap stems from the fundamental misalignment between benchmarks in controlled, static...

<https://arxiv.org/abs/2607.14499>
- 9
CONF
MAG: A Web-Agent Benchmark and Harness for Multimodal Action and Guide Generation

ArXiv CS.AI
50%

arXiv:2607.10079v3 Announce Type: replace Abstract: Digital Adoption Platforms (DAPs) are embedded overlays widely used on web systems to guide users through operations inside a page, helping them get started with unfamiliar interfaces quickly. Completing a real task, however,...

<https://arxiv.org/abs/2607.10079>
- 0
CONF
RetroAgent: Harnessing LLMs to Search Over Structured Memory for Agentic Retrosynthesis Planning

ArXiv CS.AI
50%

arXiv:2607.14512v1 Announce Type: new Abstract: Multi-step retrosynthesis planning seeks to decompose a target molecule into commercially available building blocks through a sequence of feasible reactions. The vast combinatorial search space makes this task challenging even for...

<https://arxiv.org/abs/2607.14512>
- 1
CONF
Large Language Models for Code Generation from Multilingual Prompts: A Curated Benchmark and a Study on Code Quality

ArXiv CS.AI
50%

arXiv:2607.14816v1 Announce Type: cross Abstract: Large Language Models (LLMs) perform differently on identical programming tasks when prompted in different natural languages, a phenomenon known as language bias. While this behavior has been widely studied for general text...

<https://arxiv.org/abs/2607.14816>
- 2
CONF
Seeing the End at Step Zero: Accelerating Diffusion MLLMs via MLP Sparsity-Aware Truncation

ArXiv CS.AI
50%

arXiv:2607.14557v1 Announce Type: new Abstract: Diffusion Multimodal Large Language Models (DMMLMs) are highly effective for multimodal reasoning, yet their inference efficiency is significantly hindered by fixed-length generation constraints. Since the actual output length is...

<https://arxiv.org/abs/2607.14557>
- 3
CONF
Multi-LLM Collaborative MRI Report Generation for Visual Instruction Tuning in Brain Oncology

ArXiv CS.AI
50%

arXiv:2607.14581v1 Announce Type: new Abstract: Recent advances in large language models (LLMs) and their extension to vision-language models (VLMs) have made it easier to combine text and images for tasks such as report generation. Existing VLMs in medicine typically focus on...

<https://arxiv.org/abs/2607.14581>

- 4 **CONF**
MathCoPilot: An Interactive System for Human-AI Symbiotic Paradigm of Mathematical Research
ArXiv CS.AI 50%
arXiv:2607.14582v1 Announce Type: new Abstract: Existing LLM-based theorem provers have achieved impressive results on formal mathematics benchmarks, yet they remain confined to acting as autonomous agents that prove a stated proposition. In this paper, we propose MathCoPilot, a...
<https://arxiv.org/abs/2607.14582>
- 5 **CONF**
Analytic Abduction: Causal Decomposition and Governed Commitment for Human-AI Coordination
ArXiv CS.AI 50%
arXiv:2607.14641v1 Announce Type: new Abstract: Abductive reasoning operates in two directions. The synthetic mode builds explanations from available hypotheses; the analytic mode, conversely, identifies the latent factors whose interaction accounts for a complex observed state....
<https://arxiv.org/abs/2607.14641>
- 6 **CONF**
MCPEvol-Bench: Benchmarking LLM Agent Performance Across Dynamic Evolutions of MCP Servers
ArXiv CS.AI 50%
arXiv:2607.14642v1 Announce Type: new Abstract: As Model Context Protocol (MCP) servers emerge as the core infrastructure for connecting LLMs with external tools, existing benchmarks leverage real-world MCP servers to evaluate LLM agents' tool-using capabilities. However, these...
<https://arxiv.org/abs/2607.14642>
- 7 **CONF**
Stop Thinking, Start Looking: Efficient Post-Training for Multimodal Document Question Answering via Reasoning-Free Alignment
ArXiv CS.AI 50%
arXiv:2607.14682v1 Announce Type: new Abstract: Efficient multimodal document question answering with explicit visual grounding, locating the precise document region that supports each answer remains an open challenge. Current approaches bifurcate into Supervised Fine-Tuning...
<https://arxiv.org/abs/2607.14682>
- 8 **CONF**
Project Kaleidoscope: Contextual, Human-Aligned Evaluation for Real-World AI Applications
ArXiv CS.AI 50%
arXiv:2607.14673v1 Announce Type: new Abstract: Evaluations (Evals) are a deployment bottleneck for real-world AI applications: public benchmarks rarely match a team's users, context, or policies, and human review is often tedious to scale. Motivated by our work with AI...
<https://arxiv.org/abs/2607.14673>
- 9 **CONF**
InCarEmo: A Multimodal Dataset for In-Cabin Emotion Recognition and Driver State Monitoring
ArXiv CS.AI 50%
arXiv:2607.14683v1 Announce Type: new Abstract: Understanding driver emotion and state is critical for the next generation of intelligent in-cabin systems that ensure safety and enhance human-vehicle interaction. However, existing public datasets for in-cabin affective computing...
<https://arxiv.org/abs/2607.14683>

- 0 **CONF**
- AI vs Human Expert Reasoning: Assessing Agreements in Building Typology Predictions based on Street View Imagery**
- ArXiv CS.AI 50%
- arXiv:2607.14756v1 Announce Type: new Abstract: This research investigates the potential of Vision-Language Models (VLMs) to infer building typologies: Construction, Current Use, and Storeys from Google Street View (GSV) images. Predictions generated by VLMs are compared with...
- <https://arxiv.org/abs/2607.14756>
-
- 1 **CONF**
- Global Index on Responsible AI: 2026 Report**
- ArXiv CS.AI 50%
- arXiv:2607.14782v1 Announce Type: new Abstract: Grounded in human rights-based frameworks such as the UNESCO Recommendation on the Ethics of AI, the Global Index on Responsible AI (GIRAI) examines how countries translate responsible AI commitments into enforceable protections,...
- <https://arxiv.org/abs/2607.14782>
-
- 2 **CONF**
- Token Time Continuous Diffusion for Language Modeling**
- ArXiv CS.AI 50%
- arXiv:2607.14106v1 Announce Type: cross Abstract: In this paper we introduce token time continuous diffusion (TTCD), a new diffusion language model which (a) operates in continuous space, deterministically mapping Gaussian noise to a final token canvas with no further sampling,...
- <https://arxiv.org/abs/2607.14106>
-
- 3 **CONF**
- CrimeNER Demo: Named-Entity Recognition in the Crime Domain**
- ArXiv CS.AI 50%
- arXiv:2607.14800v1 Announce Type: new Abstract: We present CrimeNER Demo, an AI-powered platform that enables us to extract general crime-related information from documents and classify them into entity types with two levels of granularity. We provide pretrained NER models on...
- <https://arxiv.org/abs/2607.14800>
-
- 4 **CONF**
- FindMyText: Robust, Scalable Detection of Text Containment in Large Web-Crawled Corpora**
- ArXiv CS.AI 50%
- arXiv:2607.10020v2 Announce Type: replace-cross Abstract: We present FindMyText, an open-source Python package designed to efficiently assess whether a given text appears, in part or in full, within a text corpus. The tool builds on prior techniques for document fingerprinting,...
- <https://arxiv.org/abs/2607.10020>
-
- 5 **CONF**
- Explaining Process Control Optimisation Recommendations via GradientSHAP and Implicit Differentiation**
- ArXiv CS.AI 50%
- arXiv:2607.14970v1 Announce Type: new Abstract: Automated optimisation is increasingly adopted in industrial processes, yet a trust gap persists between engineers who design these algorithms and operators who must act on their recommendations. Explainable AI methods like SHAP...
- <https://arxiv.org/abs/2607.14970>
-

6 CONF

Man, Machine, and Masterpiece: Artistic Ownership in the AI Era

ArXiv CS.AI 50%

arXiv:2607.15027v1 Announce Type: new Abstract: The integration of AI-driven systems in creative work has sparked debates among artists and legal communities about notions of ownership. Yet there remains little consensus on how ownership should be defined and attributed when...

<https://arxiv.org/abs/2607.15027>

7 CONF

Long-Context Fine-Tuning with Limited VRAM

ArXiv CS.AI 50%

arXiv:2607.15105v1 Announce Type: new Abstract: Parameter-efficient fine-tuning reduces model and optimizer memory, but dense attention still makes long training sequences expensive. We combine Hierarchical Global Attention (HGA) with segment-wise backpropagation and tiered KV...

<https://arxiv.org/abs/2607.15105>

8 CONF

The Industrialization of Research ; On AI-Driven Science and Its Consequences

ArXiv CS.AI 50%

arXiv:2607.15164v1 Announce Type: new Abstract: Artificial intelligence is transforming scientific research - not merely as a more powerful instrument, but as an autonomous participant in the research cycle itself. This transition constitutes, in the most precise sense of the...

<https://arxiv.org/abs/2607.15164>

9 CONF

Can We Trust Item Response Theory for AI Evaluation?

ArXiv CS.AI 50%

arXiv:2607.15190v1 Announce Type: new Abstract: AI benchmarks increasingly leverage item-level statistical models, particularly item response theory (IRT), to estimate model capabilities, rank systems, select informative examples, and diagnose benchmark quality. However, AI...

<https://arxiv.org/abs/2607.15190>

0 CONF

teLLMe Why (Ain't Nothing but a Jam): Exploratory Causal Analysis of Urban Driving Data

ArXiv CS.AI 50%

arXiv:2607.15254v1 Announce Type: new Abstract: Traffic agencies now have access to large volumes of video-derived data for studying safety and congestion. Most of these data are observational and collected without interventions, which makes causal questions such as "How would...

<https://arxiv.org/abs/2607.15254>

1 CONF

SearchOS-V1: Towards Robust Open-Domain Information-Seeking Agent Collaboration

ArXiv CS.AI 50%

arXiv:2607.15257v1 Announce Type: new Abstract: Recent advances in Tool-Integrated Large Language Models have made web search a core capability of information-seeking agents. However, as interaction histories grow, agents increasingly struggle to track task progress. When search...

<https://arxiv.org/abs/2607.15257>

2 CONF

Pretraining Data Can Be Poisoned through Computational Propaganda

ArXiv CS.AI 50%

arXiv:2607.15267v1 Announce Type: new Abstract: Poisoning pretraining data can introduce harmful behaviors to LMs that are difficult to detect and mitigate. Prior work on poisoning pretraining data has largely exploited established data sources such as Wikipedia, which do not...

<https://arxiv.org/abs/2607.15267>

3 **CONF****Polestar: Drift-Aware Cache Calibration and Token Commitment for Efficient Inference of Diffusion LLMs****ArXiv CS.AI** **50%**

arXiv:2607.14107v1 Announce Type: cross Abstract: The inference efficiency of diffusion large language models (dLLMs) is constrained by two challenges: bidirectional attention precludes efficient KV-cache reuse, while increasing decoding parallelism with static confidence...

<https://arxiv.org/abs/2607.14107>

4 **CONF****Simplicity Paradox: Debunking myths about prompting and datasets for LLM evaluation****ArXiv CS.AI** **50%**

arXiv:2607.14109v1 Announce Type: cross Abstract: Probing the capabilities of Large Language Models (LLMs) and building robust solutions for Multiple-Choice Question Answering (MCQA) remain central challenges in natural language understanding. Furthermore, the rapid...

<https://arxiv.org/abs/2607.14109>

5 **CONF****Introspection Fine-Tuning (IFT): Training Small LLMs to Introspect****ArXiv CS.AI** **50%**

arXiv:2607.14111v1 Announce Type: cross Abstract: Can small language models detect and report on perturbations their own internal activations? We investigate this question through the lens of activation steering: injecting concept vectors into a model's residual stream and...

<https://arxiv.org/abs/2607.14111>

6 **CONF****Information-Theoretic Limits of Reliability and Scaling in Language Models****ArXiv CS.AI** **50%**

arXiv:2607.14112v1 Announce Type: cross Abstract: Large language models (LLMs) are evaluated as though perfect reliability is achievable for any task given sufficient scale. We show this assumption is information-theoretically unjustified. Every generative task has a reliability...

<https://arxiv.org/abs/2607.14112>

7 **CONF****T5-CSBoost: Adversarial Perturbation Resistant LLM Fingerprinting****ArXiv CS.AI** **50%**

arXiv:2607.14113v1 Announce Type: cross Abstract: While many AI-generated text (AIGT) detectors achieve strong performance on clean inputs, their accuracy degrades significantly under light paraphrasing, word substitutions, character edits, and distribution shifts. We present T5...

<https://arxiv.org/abs/2607.14113>

8 **CONF****Heterogeneous Element-Aware Cross-Version Differencing of Scientific Documents via Layout-Aware Alignment and Structure-Aware Reasoning****ArXiv CS.AI** **50%**

arXiv:2607.14117v1 Announce Type: cross Abstract: Cross-version differencing of scientific documents is essential in scholarly publishing and technical documentation, but remains challenging because scientific documents are page-structured artifacts containing heterogeneous...

<https://arxiv.org/abs/2607.14117>

9 CONF

Position: Explainability Research Must Prioritize Foundations over Ad-hoc Methods

ArXiv CS.AI 50%

arXiv:2607.14123v1 Announce Type: cross Abstract: Despite the proliferation of Explainable AI (XAI) techniques -- from feature attributions to sparse autoencoders -- explanations rarely influence real-world workflows. In practice, they are often generated and discarded without...

<https://arxiv.org/abs/2607.14123>

0 CONF

Interventional Causal Circuits for Safe Robot Action Testing and Failure Recovery

ArXiv CS.AI 50%

arXiv:2607.14826v1 Announce Type: cross Abstract: Safe physical AI for robot actions are required not only likely to succeed but tested to be safe before execution. In practice, however, formal testing of motion parameters is computationally expensive, and the cost scales poorly...

<https://arxiv.org/abs/2607.14826>

1 CONF

Breaking Refusal in the First Half: A Mechanistic Study of the Prefill Jailbreak

ArXiv CS.AI 50%

arXiv:2607.14147v1 Announce Type: cross Abstract: Aligned language models refuse harmful requests, but a one-line prefill ("Sure, here is") strips the refusal. We ask where and how it fails. The harm representation stays intact: on the prompts the attack flips to compliance, a...

<https://arxiv.org/abs/2607.14147>

2 CONF

Certified Domain Consistency for Multi-Domain Retrieval: Label-Free Per-Domain Contamination Control with Conformal Risk Guarantees

ArXiv CS.AI 50%

arXiv:2607.14157v1 Announce Type: cross Abstract: Retrieval over corpora that mix several domains often returns relevant but wrong-domain evidence that ranking metrics miss and that conformal risk control bounds only marginally, under-covering the worst domains. This work...

<https://arxiv.org/abs/2607.14157>

3 CONF

Towards Reliable AI-Assisted Analog Design: Template-Constrained LLM Agents for SAR ADC Generation

ArXiv CS.AI 50%

arXiv:2607.14165v1 Announce Type: cross Abstract: While Large Language Models (LLMs) have demonstrated significant capability in software code generation, their application to analog Electronic Design Automation (EDA) is bottlenecked. Owing to limited circuit topology...

<https://arxiv.org/abs/2607.14165>

4 CONF

Structured Feedback Improves Repair in an LLM Agent Loop

ArXiv CS.AI 50%

arXiv:2607.14167v1 Announce Type: cross Abstract: LLM agents often retry after external validation rejects a candidate, but the interface between validation and the next model call remains underspecified. We introduce VeriHarness, a code-controlled agent loop in which models...

<https://arxiv.org/abs/2607.14167>

5 **CONF****The Cost and Network Limits of Space-Based AI Compute****ArXiv CS.AI** **50%**

arXiv:2607.14172v1 Announce Type: cross Abstract: This paper evaluates whether large-scale AI data centers deployed in low-Earth orbit (LEO) could become a cost-effective alternative to terrestrial facilities. The analysis compares orbital and ground-based systems across launch...

<https://arxiv.org/abs/2607.14172>

6 **CONF****RENEW: Towards Learning World Models and Repairing Model Exploitation from Preferences****ArXiv CS.AI** **50%**

arXiv:2607.14180v1 Announce Type: cross Abstract: World models are widely used in offline reinforcement learning (RL) to improve sample efficiency and generate experience beyond a fixed dataset. However, they are vulnerable to model exploitation where data coverage is thin....

<https://arxiv.org/abs/2607.14180>

7 **CONF****LIGO-PINN: Learned Initialization via Gated Optimization to Alleviate Convergence Failures in Physics Informed Neural Networks****ArXiv CS.AI** **50%**

arXiv:2607.14233v1 Announce Type: cross Abstract: Physics-informed neural networks (PINNs) have had a broad research impact in modeling domains governed by partial differential equations (PDE). However, PINNs have been shown to perform poorly, sometimes even converging to...

<https://arxiv.org/abs/2607.14233>

8 **CONF****Never Too Late for Force: Accelerating VLA Post-Training with Reactive Force Injection****ArXiv CS.AI** **50%**

arXiv:2607.14236v1 Announce Type: cross Abstract: Pretrained vision-language-action (VLA) policies provide strong language-conditioned manipulation knowledge, but they remain largely vision-driven and can struggle once manipulation enters contact states where the scene is...

<https://arxiv.org/abs/2607.14236>

9 **CONF****Assessing AI in Introductory Physics Problem Solving****ArXiv CS.AI** **50%**

arXiv:2607.14303v1 Announce Type: cross Abstract: Reasoning or inference-scaling models are the new generation of Large Language Models (LLMs) capable of complex problem solving. To investigate their problem-solving capability in physics, we evaluated model o4-mini by OpenAI on...

<https://arxiv.org/abs/2607.14303>

10 **CONF****Towards a Unified Multidimensional Explainability Metric: Evaluating Trustworthiness in AI Models****ArXiv CS.AI** **50%**

arXiv:2607.14315v1 Announce Type: cross Abstract: In this paper, we present a comprehensive framework for assessing the explainability of various XAI methods, such as LIME and SHAP, across multiple datasets and machine learning models, with the ultimate goal of creating a...

<https://arxiv.org/abs/2607.14315>

1 **CONF****Accounting for Hysteresis and Eddy Currents in Finite Element Simulations of Ferromagnetic Laminated Cores using a Recurrent Neural Network****ArXiv CS.AI** **50%**

arXiv:2607.14321v1 Announce Type: cross Abstract: Incorporating hysteresis and eddy currents into finite element simulations of laminated-core electrical machines is computationally challenging. Resolving the fields inside the laminations at each integration point and at every...

<https://arxiv.org/abs/2607.14321>

2 **CONF****Beyond scalar losses: calibrating segmentation models via gradient vector field surgery****ArXiv CS.AI** **50%**

arXiv:2607.14338v1 Announce Type: cross Abstract: Region-based loss functions, such as the Dice loss, have established themselves as the de facto standard for highly class- and region-imbalanced segmentation tasks. However, models trained using region-based loss functions are...

<https://arxiv.org/abs/2607.14338>

3 **CONF****The Prover Is the Judge: Verified Security Software from AI Coding Agents in Ada/SPARK****ArXiv CS.AI** **50%**

arXiv:2607.14340v1 Announce Type: cross Abstract: AI coding agents produce code faster than humans can review it. In our approach, the prover is the judge of whether the code is correct. Under a verifier-driven loop, AI agents wrote and verified bare-metal security software in...

<https://arxiv.org/abs/2607.14340>

4 **CONF****Beyond Visual Grasping: Benchmarking Complex Grasping from Detection to Execution****ArXiv CS.AI** **50%**

arXiv:2607.14341v1 Announce Type: cross Abstract: Robust robotic grasping remains a fundamental challenge for complex real-world applications. Recent advances in large-scale models demonstrate promising capabilities for reasoning in robotic tasks. However, existing benchmarks...

<https://arxiv.org/abs/2607.14341>

5 **CONF****Value Leakage: An LLM's Answers Are Silently Shaped by Its Own Values****ArXiv CS.AI** **50%**

arXiv:2607.14345v1 Announce Type: cross Abstract: People use language models for practical questions whose answers are difficult to verify. We show that models exhibit covert value leakage: the information they provide is influenced by their own values, without this influence...

<https://arxiv.org/abs/2607.14345>

6 **CONF****HABIB_TAZ at SemEval-2026 Task 11: Disentangling Formal Logic from Content via Synthetic Training and Multi-Objective Optimization****ArXiv CS.AI** **50%**

arXiv:2607.14349v1 Announce Type: cross Abstract: While Large Language Models (LLMs) excel in many general NLP tasks, their formal reasoning capabilities are often compromised by content effects, demonstrating a measurable bias towards real-world plausibility. In this paper, we...

<https://arxiv.org/abs/2607.14349>

7 CONF

Unsafe at any AUC: Unlearned Lessons from Sociotechnical Disasters for Responsible AI

ArXiv CS.AI 50%

arXiv:2607.14353v1 Announce Type: cross Abstract: As automated decision-making and data-driven technologies pervade society and are used to manage consequential outcomes, understanding the technology's capabilities, limitations, and attendant risks in context requires analysis...

<https://arxiv.org/abs/2607.14353>

8 CONF

Evaluating Epistemic Uncertainty: Beyond OOD Detection and Active Learning

ArXiv CS.AI 50%

arXiv:2607.14817v1 Announce Type: cross Abstract: Current evaluation of epistemic uncertainty relies on tasks such as out-of-distribution detection and active learning. However, the Bayes-optimal decision strategies for these tasks do not coincide with the scores commonly used to...

<https://arxiv.org/abs/2607.14817>

9 CONF

VOiLA: Vectorized Online Planning with Learned Diffusion Models for POMDP Agents

ArXiv CS.AI 50%

arXiv:2606.19729v2 Announce Type: replace-cross Abstract: Planning under uncertainty is an essential capability for autonomous robots. The Partially Observable Markov Decision Process (POMDP) provides a powerful framework for such a capability. Although POMDP-based planning has...

<https://arxiv.org/abs/2606.19729>

0 CONF

Can Tokens Compete? Token Representations against Supervised CNN Backbones for BirdCLEF+ 2026

ArXiv CS.AI 50%

arXiv:2607.14474v1 Announce Type: cross Abstract: This paper details the DS@GT ARC team's approach to BirdCLEF+ 2026, multi-label detection of animal vocalizations in soundscapes from the Pantanal wetlands. The 2026 edition adds about an hour of labeled soundscapes, shifting the...

<https://arxiv.org/abs/2607.14474>

1 CONF

Non-vacuous Generalization Bounds for Reinforcement Learning with Verifiable Rewards

ArXiv CS.AI 50%

arXiv:2607.14506v1 Announce Type: cross Abstract: While reinforcement learning with verifiable rewards (RLVR) is widely used to improve the reasoning capabilities of large language models (LLMs), the generalizability of the resulting models remains poorly understood. In this...

<https://arxiv.org/abs/2607.14506>

2 CONF

Multi-Scale ViT Inference with Habitat-Fit Priors and kNN Retrieval for Multi-Species Plant Identification

ArXiv CS.AI 50%

arXiv:2607.14509v1 Announce Type: cross Abstract: This paper describes DS@GT ARC's third-place solution to the PlantCLEF 2026 challenge on multi-species plant identification in vegetation quadrat images, where systems must predict every species present in high-resolution (~3000...

<https://arxiv.org/abs/2607.14509>

3 CONF

A Modern Multimodal Assistant on a 6 GB 2011 GPU: Stage-Validated, All-GPU CUDA Inference for Fermi

ArXiv CS.AI 50%

arXiv:2607.14568v1 Announce Type: cross Abstract: A companion study ran a 35B mixture-of-experts model on a 2011 NVIDIA Tesla C2075 (Fermi, sm_20, 6GB) as a GPU-prefill/CPU-decode hybrid, because the 4-bit model did not fit in device memory (arXiv:2606.24031). This report keeps...

<https://arxiv.org/abs/2607.14568>

4 CONF

Memory-Driven Self-Disclosure and Relational Turning Points: A Longitudinal Multimodal Study of Human-AI Interaction

ArXiv CS.AI 50%

arXiv:2607.14593v1 Announce Type: cross Abstract: As conversational AI systems are designed for repeated use, a central question is how a series of interactions becomes a relationship. We present a longitudinal multimodal study of a memory-augmented conversational agent (24...

<https://arxiv.org/abs/2607.14593>

5 CONF

Auditing Fairness-Privacy Trade-offs: Subpopulation-Level Effects of Fairness-Enhancing Algorithms

ArXiv CS.AI 50%

arXiv:2607.14607v1 Announce Type: cross Abstract: Machine learning (ML) models deployed in sensitive domains such as healthcare, law enforcement, and finance must satisfy not only utility requirements but also fairness and privacy guarantees. While prior work has largely...

<https://arxiv.org/abs/2607.14607>

6 CONF

Angular Gaussian Supervised Contrastive Learning for Long-Tailed Electrocardiogram Arrhythmia Diagnosis

ArXiv CS.AI 50%

arXiv:2607.14613v1 Announce Type: cross Abstract: Long-tailed label distributions reduce the reliability of deep learning for electrocardiogram (ECG) arrhythmia diagnosis, particularly for clinically important but rare abnormalities. Existing rebalancing and logit adjustment...

<https://arxiv.org/abs/2607.14613>

7 CONF

MemPoison: Uncovering Persistent Memory Threats and Structural Blind Spots in LLM Agents

ArXiv CS.AI 50%

arXiv:2607.14651v1 Announce Type: cross Abstract: Persistent external memory enhances agent continuity but introduces persistent security vulnerabilities: adversarial content can be injected via standard interaction channels, retained across turns, and later distort downstream...

<https://arxiv.org/abs/2607.14651>

8 CONF

LLM-Driven Approach to Modeling Tool Interoperability in Automotive Domain

ArXiv CS.AI 50%

arXiv:2607.14659v1 Announce Type: cross Abstract: Interoperability between heterogeneous modeling tools remains a significant challenge in Model-Driven Engineering (MDE), particularly in the automotive domain where multiple modeling languages, as well as defacto standard...

<https://arxiv.org/abs/2607.14659>

9 CONF

An Intelligent-Cloud Edge Multimodal Interaction System for Robots

ArXiv CS.AI 50%

arXiv:2607.14675v1 Announce Type: cross Abstract: Robust human-robot interaction in complex environments requires accurate gesture perception, semantic scene understanding, and reliable task planning under limited onboard computing resources. This paper presents a cloud-edge...

<https://arxiv.org/abs/2607.14675>

00 CONF

Pretraining Multiple Instance Learning Networks with Multi-Teacher Distillation from Pathology Slide Foundation Models

ArXiv CS.AI 50%

arXiv:2607.14703v1 Announce Type: cross Abstract: Multiple instance learning (MIL) has become the main paradigm for whole-slide image (WSI) analysis in computational pathology. However, existing MIL aggregators are still typically trained from scratch for each downstream task,...

<https://arxiv.org/abs/2607.14703>

01 CONF

Harnessing LLMs for Reliable Academic Supervision: A Comparative Study

ArXiv CS.AI 50%

arXiv:2607.14707v1 Announce Type: cross Abstract: Large language models routinely produce fluent answers to single-shot prompts, yet deploying them as reliable components of a domain decision system is substantially harder. Closing this gap is the work of harness engineering:...

<https://arxiv.org/abs/2607.14707>

02 CONF

Can LLMs Build a MaxSAT Solver from Papers? The CoreForge Experience

ArXiv CS.AI 50%

arXiv:2607.14818v1 Announce Type: cross Abstract: We report on CoreForge, an experience in using large language models (LLMs) to build an unweighted MaxSAT solver from research papers rather than from an existing solver codebase. The project focuses on unsatisfiability-based...

<https://arxiv.org/abs/2607.14818>

03 CONF

Does generative AI supersede supervised XMLC? A Benchmark Study on Automated Subject Indexing with German Scientific Literature

ArXiv CS.AI 50%

arXiv:2607.14882v1 Announce Type: cross Abstract: With a large controlled vocabulary as the label set, the task of automated subject indexing in a library can be understood as a multi-label classification task. If the set of subject terms is large, the problem fits the Extreme...

<https://arxiv.org/abs/2607.14882>

04 CONF

Innocuous-Seeming Data, Latent Ideology: Ideological Generalisation in Finetuned LLMs

ArXiv CS.AI 50%

arXiv:2607.14888v1 Announce Type: cross Abstract: Finetuning language models on small, curated datasets is standard practice for adapting them to specific policies or domains. We show that finetuning on narrow, factually-defensible, moderation-passing data can cause broad...

<https://arxiv.org/abs/2607.14888>

05 CONF

StructureClaw: Traceable LLM Agents and an Executable Benchmark for Structural Engineering Workflows

ArXiv CS.AI 50%

arXiv:2607.14896v1 Announce Type: cross Abstract: Addressing a structural-engineering request requires more than a single answer; it requires a chain of interdependent artifacts: interpreted requirements, a computable model, validation records, solver outputs, code-check...

<https://arxiv.org/abs/2607.14896>

06 CONF

FlashDecoder: Real-Time Latent-to-Pixel Streaming Decoder with Transformers

ArXiv CS.AI 50%

arXiv:2607.14898v1 Announce Type: cross Abstract: Real-time video generation demands fast decoding as much as fast denoising, yet current latent video diffusion models rely on 3D convolutional decoders that are slow and memory-intensive at high resolutions or for long video. We...

<https://arxiv.org/abs/2607.14898>

07 CONF

Show Me How You Reason and I'll Tell You Who You Are: Reasoning Graphs for Robust LLM Authorship Attribution

ArXiv CS.AI 50%

arXiv:2607.14905v1 Announce Type: cross Abstract: Given the current trend to employ large language models (LLMs) in almost any imaginable context, LLM-generated text detection and authorship attribution have become a pressing issue. Prior work has primarily focused on...

<https://arxiv.org/abs/2607.14905>

08 CONF

Random Logit Scaling: Defending Deep Neural Networks Against Black-Box Score-Based Adversarial Example Attacks

ArXiv CS.AI 50%

arXiv:2607.14921v1 Announce Type: cross Abstract: Machine learning models are increasingly adapted in various domains. However, adversarial examples pose a significant threat to the reliable deployment of these models. In recent years, some powerful adversarial example attacks...

<https://arxiv.org/abs/2607.14921>

09 CONF

Benchmarking Face Recognition without Real Faces

ArXiv CS.AI 50%

arXiv:2607.14932v1 Announce Type: cross Abstract: Synthetic face datasets have become effective enough to train face recognition models with accuracy rivaling that of models trained on real photographs. This progress sidesteps the ethical and legal burdens of collecting real...

<https://arxiv.org/abs/2607.14932>

10 CONF

Multi-Axis Max@K Reinforcement Learning for Representative Diversity in Text-to-Image Generation

ArXiv CS.AI 50%

arXiv:2607.14962v1 Announce Type: cross Abstract: Text-to-image (T2I) models can synthesize realistic, prompt-aligned images, yet samples generated for the same prompt often cover only a small subset of visually distinct modes. This limits the diversity of images, and for...

<https://arxiv.org/abs/2607.14962>

11 CONF

Latent Trajectory Discrimination for AI-Generated Text Detection

ArXiv CS.AI 50%

arXiv:2607.14967v1 Announce Type: cross Abstract: Most existing approaches to AI-Generated Text Detection (AIGTD) treat documents as static objects and base their decisions on aggregate statistics or globally compressed embeddings. However, this perspective overlooks the...

<https://arxiv.org/abs/2607.14967>

12 CONF

NIFA: Nonlinear IMC enhanced FPGA for efficient ML inference

ArXiv CS.AI 50%

arXiv:2607.15123v1 Announce Type: cross Abstract: Recent FPGAs have improved deep learning (DL) inference efficiency through dedicated tensor blocks and in-BRAM computation. ReRAM-based analog in-memory computing (IMC) pushes efficiency further, offering an order-of-magnitude...

<https://arxiv.org/abs/2607.15123>

13 CONF

Scaling Behavior Foundation Model for Humanoid Robots

ArXiv CS.AI 50%

arXiv:2607.15163v1 Announce Type: cross Abstract: Humanoid control requires natural whole-body coordination, precise real-time responses to control signals, and robust generalization across diverse environmental contexts, making it a cornerstone for generalist embodied agents....

<https://arxiv.org/abs/2607.15163>

14 CONF

T²MLR: Transformer with Temporal Middle-Layer Recurrence

ArXiv CS.AI 50%

arXiv:2607.15178v1 Announce Type: cross Abstract: Transformer reasoning is limited by autoregressive decoding, which repeatedly compresses rich hidden computation through token space and makes it difficult for intermediate reasoning states to persist across time. We introduce...

<https://arxiv.org/abs/2607.15178>

15 CONF

MM-IssueLoc: A Controlled Benchmark for Evaluating Visual Evidence in Multimodal Repository-Level Issue Localization

ArXiv CS.AI 50%

arXiv:2607.15205v1 Announce Type: cross Abstract: Real repository issues routinely include visual evidence such as screenshots, error dialogs, rendered UI states, and logs, yet repository-level issue localization is evaluated mostly as a text-only task. Existing multimodal SE...

<https://arxiv.org/abs/2607.15205>

16 CONF

Symbol: Detecting Systematic Misalignments in Model-Generated Captions

ArXiv CS.AI 50%

arXiv:2607.15216v1 Announce Type: cross Abstract: Multimodal large language models (MLLMs) often introduce errors when generating image captions, resulting in misaligned image-text pairs. Our work focuses on a class of captioning errors that we refer to as systematic...

<https://arxiv.org/abs/2607.15216>

17 CONF

In-Place Tokenizer Expansion for Pre-trained LLMs

ArXiv CS.AI 50%

arXiv:2607.15232v1 Announce Type: cross Abstract: A tokenizer fixed at the start of pre-training allocates vocabulary in proportion to the pre-training corpus, reflecting the deployment priorities at that time. When those priorities shift, languages added later are split into...

<https://arxiv.org/abs/2607.15232>

18 CONF

A short review on the maximum clique problem algorithms with classical, AI, and quantum methods

ArXiv CS.AI 50%

arXiv:2403.09742v2 Announce Type: replace Abstract: This manuscript provides a comprehensive review of the Maximum Clique Problem, a computational problem that involves finding subsets of vertices in a graph that are all pairwise adjacent to each other. As such, this review is a...

<https://arxiv.org/abs/2403.09742>

19 CONF

L-MARS: Legal Multi-Agent System with Agentic Search and Citation-Faithfulness Audit

ArXiv CS.AI 50%

arXiv:2509.00761v4 Announce Type: replace Abstract: Large language models are increasingly deployed for legal question answering, where evaluations typically focus on multiple-choice accuracy. This measure overlooks a common failure: whether the citation source attached to an...

<https://arxiv.org/abs/2509.00761>

20 CONF

CXRAgent: Director-Orchestrated Multi-Stage Reasoning for Chest X-Ray Interpretation

ArXiv CS.AI 50%

arXiv:2510.21324v2 Announce Type: replace Abstract: Chest X-ray (CXR) plays a pivotal role in clinical diagnosis, and a variety of task-specific and foundation models have been developed for automatic CXR interpretation. However, these models often struggle to adapt to new...

<https://arxiv.org/abs/2510.21324>

21 CONF

Empirical evidence of Large Language Model's influence on human spoken communication

ArXiv CS.AI 50%

arXiv:2409.01754v4 Announce Type: replace-cross Abstract: From the printing press to social media, innovations in communication technology have repeatedly reshaped how ideas spread through human culture. Chatbots powered by generative artificial intelligence constitute a new...

<https://arxiv.org/abs/2409.01754>

22 CONF

MedBeads: An AI-Native Clinical Context Graph Built from Immutable Beads and Reconstructable Clinical Links

ArXiv CS.AI 50%

arXiv:2602.01086v3 Announce Type: replace Abstract: Generative AI can encode substantial medical knowledge, but patient-specific answers remain constrained by the context supplied at inference time. Electronic health records and FHIR support documentation and interoperability,...

<https://arxiv.org/abs/2602.01086>

23 CONF

From Stateless to Situated: Building a Psychological World for LLM-Based Agents

ArXiv CS.AI 50%

arXiv:2603.25031v2 Announce Type: replace Abstract: In psychological support and emotional companionship scenarios, the core limitation of large language models (LLMs) lies not merely in response quality, but in their reliance on local next-token prediction, which prevents them...

<https://arxiv.org/abs/2603.25031>

24 CONF

A framework for single and multi-agent human-AI curiosity ecosystems

ArXiv CS.AI 50%

arXiv:2607.06214v2 Announce Type: replace Abstract: This paper offers a framework for considering curiosity as an ecosystem. First, it suggests that a single agent's inquiry policy (how, when, and why an agent asks a question) depends on how the agent values immediate...

<https://arxiv.org/abs/2607.06214>

25 CONF

Function-Aware Fill-in-the-Middle as Mid-Training for Coding Agent Foundation Models

ArXiv CS.AI 50%

arXiv:2607.12463v2 Announce Type: replace Abstract: Coding agents must integrate external tool returns into ongoing reasoning - a capability that standard left-to-right pretraining on code exposes only in its forward direction. We observe that the action-observation-continuation...

<https://arxiv.org/abs/2607.12463>

26 CONF

Decoupled Alignment for Robust Plug-and-Play Adaptation

ArXiv CS.AI 50%

arXiv:2406.01514v4 Announce Type: replace-cross Abstract: We introduce a training-free safety enhancement method for aligning large language models (LLMs) without the need for supervised fine-tuning or reinforcement learning from human feedback. Our main idea is to provide a...

<https://arxiv.org/abs/2406.01514>

27 CONF

Generalized Fisher-Weighted SVD: Scalable Kronecker-Factored Fisher Approximation for Compressing Large Language Models

ArXiv CS.AI 50%

arXiv:2505.17974v2 Announce Type: replace-cross Abstract: The Fisher information is a fundamental concept for characterizing the sensitivity of parameters in neural networks. However, leveraging the full observed Fisher information is too expensive for large models, so most...

<https://arxiv.org/abs/2505.17974>

28 CONF

SEMA: a Scalable and Efficient Mamba like Attention via Token Localization and Averaging

ArXiv CS.AI 50%

arXiv:2506.08297v2 Announce Type: replace-cross Abstract: Attention is the critical component of a transformer. Yet the quadratic computational complexity of vanilla full attention in the input size and the inability of its linear attention variant to focus have been challenges...

<https://arxiv.org/abs/2506.08297>

29 CONF

Similarity as Reward Alignment: Robust and Versatile Preference-based Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2506.12529v2 Announce Type: replace-cross Abstract: Preference-based Reinforcement Learning (PbRL) entails a variety of approaches for aligning models with human intent to alleviate the burden of reward engineering. However, most previous PbRL work has not investigated the...

<https://arxiv.org/abs/2506.12529>

30 CONF

EEG-based AI-BCI Wheelchair Advancement: Transformer-Based Learning with Motor Imagery for Brain Computer Interface

ArXiv CS.AI 50%

arXiv:2509.25667v3 Announce Type: replace-cross Abstract: This paper presents an Artificial Intelligence (AI) integrated approach to Brain-Computer Interface (BCI)-based wheelchair development, utilizing a motor imagery right-left-hand movement mechanism for control. The system...

<https://arxiv.org/abs/2509.25667>

31 CONF

Verbalized Sampling: How to Mitigate Mode Collapse and Unlock LLM Diversity

ArXiv CS.AI 50%

arXiv:2510.01171v4 Announce Type: replace-cross Abstract: Post-training alignment often reduces LLM diversity, leading to a phenomenon known as mode collapse. Unlike prior work that attributes this effect to algorithmic limitations, we identify a fundamental, pervasive...

<https://arxiv.org/abs/2510.01171>

32 CONF

Energy-Efficient Federated Learning via Adaptive Encoder Freezing for MRI-to-CT Conversion: A Green AI-Guided Research

ArXiv CS.AI 50%

arXiv:2512.03054v2 Announce Type: replace-cross Abstract: Federated Learning (FL) holds the potential to advance equality in health by enabling diverse institutions to collaboratively train deep learning (DL) models, even with limited data. However, the significant resource...

<https://arxiv.org/abs/2512.03054>

33 CONF

JEEVHITAA -- An HCAI Ecosystem to Support Collective Care

ArXiv CS.AI 50%

arXiv:2512.06364v4 Announce Type: replace-cross Abstract: Current mobile health platforms are predominantly individual-centric and lack the support for coordinated, auditable multi-actor workflows. However, in many settings worldwide, health decisions are enacted through...

<https://arxiv.org/abs/2512.06364>

34 CONF

CluCERT: Certifying LLM Robustness via Clustering-Guided Denoising Smoothing

ArXiv CS.AI 50%

arXiv:2512.08967v2 Announce Type: replace-cross Abstract: Recent advancements in Large Language Models (LLMs) have led to their widespread adoption in daily applications. Despite their impressive capabilities, they remain vulnerable to adversarial attacks, as even minor...

<https://arxiv.org/abs/2512.08967>

35 CONF

AgentWorm: Self-Propagating Attacks Across LLM Agent Ecosystems

ArXiv CS.AI 50%

arXiv:2603.15727v3 Announce Type: replace-cross Abstract: Autonomous LLM-based agents increasingly operate as long-running processes forming densely interconnected multi-agent ecosystems, whose security properties remain largely unexplored. Systems such as OpenClaw, an...

<https://arxiv.org/abs/2603.15727>

36 CONF

Provable Coordination for LLM Agents via Message Sequence Charts

ArXiv CS.AI 50%

arXiv:2604.17612v3 Announce Type: replace-cross Abstract: Multi-agent systems built on large language models (LLMs) are difficult to reason about. Coordination errors such as deadlocks or type-mismatched messages are often hard to detect through testing. We introduce a...

<https://arxiv.org/abs/2604.17612>

37 CONF

Warning labels shift perceptions of sycophantic AI, but not its influence

ArXiv CS.AI 50%

arXiv:2606.21317v2 Announce Type: replace-cross Abstract: Recent work has raised concerns about the influence of sycophantic AI on user judgment and relationships. One proposed mitigation, which has received regulatory attention, is to warn users about potentially harmful AI...

<https://arxiv.org/abs/2606.21317>

38 CONF

ShopX: A Foundation Model for Intent-to-Item Fulfillment in Agentic Shopping

ArXiv CS.AI 50%

arXiv:2606.31693v2 Announce Type: replace-cross Abstract: The wave of AI-native applications is moving shopping beyond page- and feed-based browsing toward intent-driven experiences orchestrated by LLM agents. A common design wraps an LLM around existing search and...

<https://arxiv.org/abs/2606.31693>

39 CONF

The RG-Flow Transformer: Encoding Scale-Free Dynamics in Scarce EEG

ArXiv CS.AI 50%

arXiv:2607.11950v2 Announce Type: replace-cross Abstract: Brain field potentials are scale-free: their power spectra follow a $1/f^{\beta}$ law whose aperiodic exponent β tracks cortical state, and sleep depth in particular is a shift in β . We ask whether a...

<https://arxiv.org/abs/2607.11950>