



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	73	Media
HackerNews	9	Community
TechCrunch AI	4	Media
The Verge AI	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Apple's lawsuit couldn't come at a worse time for OpenAI

TechCrunch AI +1 source

80%

Apple filed a trade secrets lawsuit against OpenAI last Friday, and it's not messing around. The complaint alleges a pattern of misconduct reaching all the way up to OpenAI's chief hardware officer and claims more than 400 former Apple employees now work at the company. OpenAI's...

<https://techcrunch.com/podcast/apples-lawsuit-couldnt-come-at-a-worse-time-for-openai/>

VERI

Toward a mechanistic understanding of inference in visual cortex and diffusion models

ArXiv CS.AI +1 source

80%

arXiv:2607.15693v1 Announce Type: cross Abstract: We describe a model of perceptual inference in primary visual cortex (V1) equivalent to a minimal diffusion model whose function can be readily understood from its parameters. The model is based on sparse coding with a...

<https://arxiv.org/abs/2607.15693>

CONFIRMED SIGNALS — 83 stories

CONF

Show HN: LoopGain - Stop agent loops with control theory, not max_iterations

HackerNews 50% HN 29

Hello all, I'm a software developer. Over the last few months more and more of my work has turned into using coding agents instead of typing the whole code myself. Usually a few claude sessions at once, sometimes codex, one per feature or per revealed bug. I ran them in a split...

<https://github.com/loopgain-ai/loopgain>

CONF

Show HN: PilotCite - Get your brand cited by ChatGPT, Gemini, and more

HackerNews 50% HN 6

No summary — see link for full article.

<https://www.pilotcite.com>

CONF

Show HN: Chalie - AI peer not employee

HackerNews 50% HN 5

Hey everyone, just added Bedrock and Vertex support to FlexInference. It now can do provider fallback on 429 and 5xx. Plus there is upstream retries you can setup via SDK. This way you can do routing for free through OpenAI, Anthropic, Google, Vertex, and Bedrock. Working on...

<https://github.com/chalie-ai/chalie>

CONF

Show HN: A self-hosted AI that turns Hacker News into a daily briefing

HackerNews 50% HN 4

Hi everyone, I read HN every day, but there was always more great content than I had time to read. I know there are already several HN summarizers on GitHub, and I tried some of them. They just didn't fit the workflow I wanted, so I decided to build my own. My project is a...

<https://github.com/RecNes/hn-ai-summarizer>

CONF

Show HN: Zlvox - 25 Free Developer Tools (AI, JSON, PDF) with Zero Tracking

HackerNews 50% HN 4

No summary — see link for full article.

<https://zlvox.com/>

CONF

Show HN: Knowledge graph skill for Claude/Kimi Code

HackerNews 50% HN 3

Working with multiple Claude Code instances I found issues with requirements traceability, test coverage, and inter-instance communication. I was working with knowledge graphs so built a knowledge graph skill for Claude Code - every user story, use-case, feature, test-case is a...

<https://org-edge.com/sysedge.html>

CONF

Show HN: Regularly, hand-picked daily logic and number puzzles, no ads

HackerNews 50% HN 3

Hi HN. I wanted to present to you Regularly (<https://regularly.co>), a daily logic/math puzzle website I made to scratch my own itch: can we have a fun and engaging puzzle app that doesn't eventually devolve into an iap blockaded mess? For background, a few years ago I released...

<https://regularly.co>

0 CONF

Show HN: AI Buddy - Free STT and Screenshot Tool for Mac (Requires Gemini)

HackerNews 50% HN 2

I created this tool because I honestly didn't want to pay a monthly subscription fee for other apps, and I also needed it to be a bit more reliable and wanted to be able to take screenshots in a particular way that's useful for working with Claude Code.

<https://claudete.co/ai-buddy>

1 CONF

Show HN - Online developer profiles with embedded AI hypeman

HackerNews 50% HN 2

No summary — see link for full article.

<https://engineerprofiles.com>

2 CONF

'Odyssey' director Christopher Nolan calls AI an obvious 'Trojan horse'

TechCrunch AI 50%

"Everybody knows the Greeks are inside."

<https://techcrunch.com/2026/07/19/odyssey-director-christopher-nolan-calls-ai-an-obviou...>

3 CONF

Nonprofit Current AI is racing to build the World Wide Web of AI, free for all

TechCrunch AI 50%

Current AI, a non-profit building AI that leaves no one culture behind, has made remarkable progress across devices, AI chat and more.

<https://techcrunch.com/2026/07/19/nonprofit-current-ai-is-racing-to-build-the-world-wid...>

4 CONF

Causal-Audit: Explicit and Auditable Graph-based Reasoning via Target-Aware Causal Chain Construction

ArXiv CS.AI 50%

arXiv:2607.15281v1 Announce Type: new Abstract: Causal and intervention-based question answering is fundamental to advancing large language models (LLMs) toward reasoning beyond surface-level correlations and understanding underlying causal mechanisms. However, existing...

<https://arxiv.org/abs/2607.15281>

- 5
CONF
AnovaX: A Local, Multi-Agent Voice Assistant with LLM Planning, Typed Executors, and Adaptive Recovery
ArXiv CS.AI 50%

arXiv:2607.15367v1 Announce Type: new Abstract: Desktop voice assistants are still dominated by cloud pipelines that ship raw audio off the machine and expose a fixed set of skills. We describe AnovaX, a small local-first assistant that runs entirely on the user's computer and...

<https://arxiv.org/abs/2607.15367>
- 6
CONF
DrawingVQA: A Real-World Benchmark for Multi-Depth Visual-Textual Reasoning on Construction Drawings
ArXiv CS.AI 50%

arXiv:2607.15418v1 Announce Type: new Abstract: We introduce DrawingVQA, the first benchmark designed to evaluate multimodal large language models (MLLMs) on real-world construction drawings -- a core media in architecture, civil, and many other engineering practices. Unlike...

<https://arxiv.org/abs/2607.15418>
- 7
CONF
S1-Omni: A Unified Multimodal Reasoning Model for Scientific Understanding, Prediction, and Generation
ArXiv CS.AI 50%

arXiv:2607.15686v1 Announce Type: new Abstract: We present S1-Omni, a unified multimodal reasoning model for scientific understanding, prediction, and generation. AI for Science (AI4S) has advanced significantly through domain-specific models, tool-augmented LLMs, and scientific...

<https://arxiv.org/abs/2607.15686>
- 8
CONF
From Black Box to Executable Logic: Explainable Reinforcement Learning through Prolog Expert Systems
ArXiv CS.AI 50%

arXiv:2607.15459v1 Announce Type: new Abstract: A trained deep reinforcement learning policy is a black box, and we ask whether it can be made explainable by rewriting it as an executable logic program that reproduces its behaviour and that a person can read, a logic engine can...

<https://arxiv.org/abs/2607.15459>
- 9
CONF
A Critical Analysis of Trustworthy AI Tools, Mark Frameworks, and the Implementation Chasms
ArXiv CS.AI 50%

arXiv:2607.15480v1 Announce Type: new Abstract: As artificial intelligence (AI) systems increasingly impact society, ensuring their ethical and trustworthy deployment has become a global priority. While a myriad of high-level ethical guidelines have emerged, criticism persists...

<https://arxiv.org/abs/2607.15480>
- 0
CONF
Logic, Optimization, and Artificial Intelligence
ArXiv CS.AI 50%

arXiv:2607.15532v1 Announce Type: new Abstract: Logic and optimization can, in combination, make valuable contributions to rule-based AI. Logic is the obvious medium for encoding a rule base and drawing inferences from it, while optimization provides a powerful technology for...

<https://arxiv.org/abs/2607.15532>

1 **CONF****MGDT: MLLM-Guided Diffusion Transformer with Relation-Adaptive Mixture-of-Experts for Multimodal Knowledge Graph Completion****ArXiv CS.AI** **50%**

arXiv:2607.15592v1 Announce Type: new Abstract: Multimodal Knowledge Graph Completion (MKGC) requires inferring missing entities from structural, textual, and visual cues. Existing diffusion-based MKGC methods usually denoise directly on raw multimodal features. Such a design...

<https://arxiv.org/abs/2607.15592>

2 **CONF****Neuro-Symbolic AI for LEED compliance: Document-Centric Benchmarking, Deterministic Numeric Checking, and When Multimodal Hurts****ArXiv CS.AI** **50%**

arXiv:2607.15647v1 Announce Type: new Abstract: LEED v4.1 BD+C certification remains a document-intensive process that requires reviewers to read hundreds of pages of project evidence and apply credit-specific threshold logic by hand. This paper investigates whether small,...

<https://arxiv.org/abs/2607.15647>

3 **CONF****CTC: The Composite Task Challenge for Cooperative Multi-Agent Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2502.00345v2 Announce Type: replace-cross Abstract: The critical role of division of labor (DOL) in enhancing cooperation is well-recognized in real-world applications. Consequently, many cooperative multi-agent reinforcement learning (MARL) methods have incorporated DOL...

<https://arxiv.org/abs/2502.00345>

4 **CONF****NeurOWL: An LLM-Based Neural-symbolic Framework for Incomplete OWL Ontology Reasoning****ArXiv CS.AI** **50%**

arXiv:2607.15776v1 Announce Type: new Abstract: OWL ontologies provide a formal knowledge representation framework that enables semantic reasoning, and have been widely adopted across domains such as healthcare and bioinformatics. In practice, however, real-world ontologies are...

<https://arxiv.org/abs/2607.15776>

5 **CONF****Efficient and Privacy Aware Edge Cloud Collaborative Inference for Large Language Models****ArXiv CS.AI** **50%**

arXiv:2607.13093v2 Announce Type: replace-cross Abstract: On-device LLM inference faces a trilemma of response latency, limited hardware resources and user privacy. Full cloud inference delivers strong computing power but exposes user prompts and dialogue data, while standalone...

<https://arxiv.org/abs/2607.13093>

6 **CONF****AI Trading: Evaluating Large Language Models for Technical Market Analysis****ArXiv CS.AI** **50%**

arXiv:2607.15414v1 Announce Type: cross Abstract: Large Language Models (LLMs) have emerged as powerful tools for processing the heterogeneous information environments of modern financial markets. This paper presents a systematic, comparative evaluation of five prominent LLMs:...

<https://arxiv.org/abs/2607.15414>

7 **CONF****Closing the AI Trust Gap: The Case for Independent Certification for Trustworthy AI****ArXiv CS.AI** **50%**

arXiv:2607.15992v1 Announce Type: new Abstract: Over the past decade, responsible AI (RAI) has produced a substantial body of practice for identifying and mitigating the risks AI poses in high-stakes settings. Yet this work has not produced a market that rewards trustworthiness....

<https://arxiv.org/abs/2607.15992>

8 **CONF****SciForge: An AI-Native, Multimodal Workbench for Scientific Discovery****ArXiv CS.AI** **50%**

arXiv:2607.16038v1 Announce Type: new Abstract: Scientific work increasingly spans heterogeneous artifacts -- papers, code, datasets, scientific file formats, model outputs, figures, manuscripts, and team decisions -- yet general-purpose AI assistants rarely preserve these...

<https://arxiv.org/abs/2607.16038>

9 **CONF****Harmonizing AI Safety Thresholds****ArXiv CS.AI** **50%**

arXiv:2607.16112v1 Announce Type: new Abstract: Frontier AI companies have published capability thresholds that differ substantially, making it difficult for third parties to verify whether a threshold has been crossed or to compare requirements across companies. Moreover,...

<https://arxiv.org/abs/2607.16112>

0 **CONF****CRAFT: Clustering Rubrics to Diagnose Weak LLM Capabilities and Generate Targeted Fine-Tuning Data****ArXiv CS.AI** **50%**

arXiv:2607.16122v1 Announce Type: new Abstract: Evaluations should do more than measure a models current performance. They should tell us what to fix for the next model iteration and provide a way to generate targeted post training data. Most evaluation pipelines identify weak...

<https://arxiv.org/abs/2607.16122>

1 **CONF****RL-Struct: A Lightweight Reinforcement Learning Framework for Reliable Structured Output in LLMs****ArXiv CS.AI** **50%**

arXiv:2512.00319v3 Announce Type: replace Abstract: The Structure Gap between probabilistic LLM generation and deterministic schema requirements hinders automated workflows. We propose RL-Struct, a lightweight framework using Gradient Regularized Policy Optimization (GRPO) with...

<https://arxiv.org/abs/2512.00319>

2 **CONF****Partial Information Decomposition as a Multi-Contrast 3D MRI Selection Strategy for Resource-Constrained Deep Neural Network Training in Brain Tumor Segmentation****ArXiv CS.AI** **50%**

arXiv:2607.15396v1 Announce Type: cross Abstract: Multi-contrast 3D MRI segmentation can be computationally demanding when all available sequences are used. We evaluate a pre-training Partial Information Decomposition framework that ranks input pairs according to their...

<https://arxiv.org/abs/2607.15396>

- 3
CONF
Coercion and Deception in AI-to-AI Management: An Agentic Benchmark of Unprompted Escalation
ArXiv CS.AI 50%
arXiv:2607.15434v1 Announce Type: cross Abstract: Multi-agent systems routinely place one AI agent in authority over another. When a subordinate refuses a task, the manager chooses the outcome: it can renegotiate, report the failure honestly, coerce the subordinate, or lie about...
<https://arxiv.org/abs/2607.15434>
- 4
CONF
Verbalizable Representations Form a Global Workspace in Language Models
ArXiv CS.AI 50%
arXiv:2607.15495v1 Announce Type: cross Abstract: Out of everything the human brain processes, only a small fraction is consciously accessible, in the sense of being available for verbal report, deliberate control, and flexible reasoning. In this paper, we present evidence that...
<https://arxiv.org/abs/2607.15495>
- 5
CONF
LLM-Driven AutoML for Cross-Lingual Handwritten OCR: Closed-Loop Neural Architecture Search with GPT-5, GPT-4o, and Claude Sonnet 4
ArXiv CS.AI 50%
arXiv:2607.15509v1 Announce Type: cross Abstract: We present a fully automated closed-loop AutoML framework that uses GPT-5, GPT-4o, and Claude Sonnet 4 as autonomous neural architecture designers for cross-lingual handwritten optical character recognition. Each large language...
<https://arxiv.org/abs/2607.15509>
- 6
CONF
Cache-Aware Prompt Compression:A Two-Tier Cost Model for LLM API Caching
ArXiv CS.AI 50%
arXiv:2607.15516v1 Announce Type: cross Abstract: Production LLM deployments combine two cost-reduction primitives: prompt caching (a discounted rate for re-used token prefixes) and prompt compression (fewer tokens sent). The compression literature has standardized on...
<https://arxiv.org/abs/2607.15516>
- 7
CONF
SLAPBench: Benchmarking Multimodal Large Language Models for Four-Finger SLAP Fingerprint Verification
ArXiv CS.AI 50%
arXiv:2607.15517v1 Announce Type: cross Abstract: Four-finger SLAP fingerprints are flat live-scan impressions of the index, middle, ring, and little fingers of one hand, used for identity verification in border control and law enforcement. No benchmark has evaluated whether...
<https://arxiv.org/abs/2607.15517>
- 8
CONF
Evolutionary Algorithm-Guided LLMs for Physics-Informed Neural Network Design
ArXiv CS.AI 50%
arXiv:2607.15560v1 Announce Type: cross Abstract: Physics-informed neural networks (PINNs) are unusually sensitive to interacting choices of architecture, activation, loss weighting, collocation, optimization, and constraint enforcement. Large language models (LLMs) can propose...
<https://arxiv.org/abs/2607.15560>

9 CONF

Ask Twice, Look Twice: Prompt Echoing Resolves the Question-First Paradox in Vision-Language Models

ArXiv CS.AI 50%

arXiv:2607.15565v1 Announce Type: cross Abstract: Where should the question go in a vision-language model (VLM) prompt: before the image or after it? Intuition says before: knowing what is asked should tell the model where to look. Yet across visual question answering...

<https://arxiv.org/abs/2607.15565>

0 CONF

MemoGuard: An Adaptive Runtime for Guarding Against Memory Traps in Communication-Limited Robot Navigation

ArXiv CS.AI 50%

arXiv:2607.15589v1 Announce Type: cross Abstract: Communication-limited robots in mission-critical scenarios such as disaster inspection and search-and-rescue must make reliable onboard decisions without access to remote operators or high-capacity reasoning services. Episodic...

<https://arxiv.org/abs/2607.15589>

1 CONF

Scalable LLM Agent Tool Access in the Cloud

ArXiv CS.AI 50%

arXiv:2607.15593v1 Announce Type: cross Abstract: LLM agents increasingly rely on tool calling to act on external systems, and the Model Context Protocol (MCP) has quickly become its de facto interface. Operating MCP at cloud scale, however, becomes difficult. On the tool...

<https://arxiv.org/abs/2607.15593>

2 CONF

Think at 5 Hz, Act at 20 Hz: Asynchronous Fast-Slow Vision-Language-Action Inference for Closed-Loop Driving

ArXiv CS.AI 50%

arXiv:2607.15621v1 Announce Type: cross Abstract: Large language models bring instruction following and scene reasoning to end-to-end driving, but their inference latency collides with the control rate a vehicle requires. Existing closed-loop agents hide this gap by invoking the...

<https://arxiv.org/abs/2607.15621>

3 CONF

Efficient Difficulty-Aware Dynamic Routing for Diffusion-Based Real-World Image Super-Resolution

ArXiv CS.AI 50%

arXiv:2607.15711v1 Announce Type: cross Abstract: Diffusion-based methods have achieved impressive performance in real-world image super-resolution (Real-ISR) by leveraging large pre-trained stable diffusion (SD) models as powerful generative priors. However, these methods still...

<https://arxiv.org/abs/2607.15711>

4 CONF

Debiasing Text-to-Image Evaluation via Implicit Cultural Alignment Reward Modeling

ArXiv CS.AI 50%

arXiv:2607.15740v1 Announce Type: cross Abstract: As Text-to-Image (T2I) systems rapidly advance, evaluating the cultural authenticity of synthesized content has become increasingly important for fair and trustworthy generative AI. Existing T2I evaluation metrics and multimodal...

<https://arxiv.org/abs/2607.15740>

- 5
CONF
GeoChrono: Benchmarking and Rethinking Long-Term Temporal Understanding in Remote Sensing
ArXiv CS.AI 50%
arXiv:2607.15768v1 Announce Type: cross Abstract: Remote sensing offers an unparalleled vantage point for observing the Earth's long-term surface evolution, yet it demands that a model not only perceive land cover at isolated moments, but also track changes, memorize evolution...
<https://arxiv.org/abs/2607.15768>
- 6
CONF
Modularized Dynamic-Granularity Video LLM for Multi-Event Long Video Understanding
ArXiv CS.AI 50%
arXiv:2607.15778v1 Announce Type: cross Abstract: Video Large Language Models (Video LLMs) have made significant advancements in various video understanding tasks. However, long-video scenarios remain challenging due to the tension between limited visual token budgets and the...
<https://arxiv.org/abs/2607.15778>
- 7
CONF
In-context learning of closed form solution to simple linear regression task using transformer with linear self-attention
ArXiv CS.AI 50%
arXiv:2607.15819v1 Announce Type: cross Abstract: In-context learning is a remarkable property of transformers and has recently received a lot of interest. In many studies of in-context learning, it has been shown that transformers are capable of implementing solver for linear...
<https://arxiv.org/abs/2607.15819>
- 8
CONF
Agentic Synthesis against Counterexample-Supplemented Sketches
ArXiv CS.AI 50%
arXiv:2607.15854v1 Announce Type: cross Abstract: Coding agents can fix a failing example without preserving the domain rule that made it fail, so later generations can repeat the same plausible mistake. We present agentic synthesis against counterexample-supplemented sketches,...
<https://arxiv.org/abs/2607.15854>
- 9
CONF
Induction in Both Directions: A Mechanistic Analysis of In-Context Learning in Masked Diffusion Language Models
ArXiv CS.AI 50%
arXiv:2607.15893v1 Announce Type: cross Abstract: While the internal mechanisms of autoregressive (AR) transformers have been studied extensively, much less is known about diffusion language models (DLMs), an emerging alternative that generates text by iterative denoising. In...
<https://arxiv.org/abs/2607.15893>
- 0
CONF
On the Failure of Boundary-Seeking Distillation in Bottlenecked Generative Architectures
ArXiv CS.AI 50%
arXiv:2607.15919v1 Announce Type: cross Abstract: Data-free knowledge distillation transfers the knowledge encoded in a teacher model to a student model without access to the original training data. Prior work such as Contrastive Abductive Knowledge Extraction (CAKE) achieves...
<https://arxiv.org/abs/2607.15919>

1 CONF

When Not to Automate: A Formal Protocol for Human Preservation in AI-Optimized Organizations

ArXiv CS.AI 50%

arXiv:2607.15944v1 Announce Type: cross Abstract: Standard automation ROI misses four categories of systemic risk -- tacit knowledge erosion, resilience reduction, regulatory exposure, and socio-institutional capital degradation -- that affect long-term organizational...

<https://arxiv.org/abs/2607.15944>

2 CONF

Robustness of Reinforcement Learning-Based Congestion Management in Low-Voltage Grids

ArXiv CS.AI 50%

arXiv:2607.16004v1 Announce Type: cross Abstract: Increases in photovoltaic generation, charging of electric vehicles and heat-pump demand challenge operating limits in low-voltage distribution grids. This requires curative curtailment methods that can operate under sparse...

<https://arxiv.org/abs/2607.16004>

3 CONF

Revisiting data-driven dynamic security assessment with a tabular foundation model

ArXiv CS.AI 50%

arXiv:2607.16031v1 Announce Type: cross Abstract: Data-driven pre-fault dynamic security assessment (DSA) rapidly evaluates the dynamic risk of credible contingencies on a power system using machine learning. Existing approaches face two limitations. First, they require a large...

<https://arxiv.org/abs/2607.16031>

4 CONF

Frontier AI performance across the business disciplines: a case-grounded benchmark of knowledge work and analytical reasoning

ArXiv CS.AI 50%

arXiv:2607.16057v1 Announce Type: cross Abstract: Large language models (LLMs) are improving rapidly as reflected in benchmark scores, yet these AI benchmarks largely test capabilities such as factual recall, narrow question answering, mathematical problem-solving, and coding...

<https://arxiv.org/abs/2607.16057>

5 CONF

When Model Merging Rivals Joint Multi-Task Reinforcement Learning: A Task-Vector Geometry Analysis

ArXiv CS.AI 50%

arXiv:2607.16062v1 Announce Type: cross Abstract: Model merging is promoted as a substitute for joint multi-task training, yet in the reinforcement-learning setting this substitution is essentially never tested against the baseline it claims to replace: methods merge...

<https://arxiv.org/abs/2607.16062>

6 CONF

Spatial Normalization for Cross-Domain Retinal Layer Segmentation in Optical Coherence Tomography

ArXiv CS.AI 50%

arXiv:2607.16065v1 Announce Type: cross Abstract: Retinal layer segmentation in Optical Coherence Tomography (OCT) is a fundamental step for extracting quantitative biomarkers of retinal structure. Indeed, there is a growing interest in the analysis of OCTs in the context of...

<https://arxiv.org/abs/2607.16065>

- 7
CONF
LLM-Powered Agentic AI for 5G/6G Networks: A Tutorial and Survey on Architectures, Protocols, and Standardization
ArXiv CS.AI 50%

arXiv:2607.16066v1 Announce Type: cross Abstract: Agentic Artificial Intelligence (AI), enabled by Large Language Models, marks a shift from rule-based automation toward autonomous, goal-driven control of Next-Generation Networks (NGNs). Existing surveys treat the two domains in...

<https://arxiv.org/abs/2607.16066>
- 8
CONF
JoyNexus: Service-Oriented Multi-Tenant Post-Training for VLA Models
ArXiv CS.AI 50%

arXiv:2607.16074v1 Announce Type: cross Abstract: The post-training of Vision-Language-Action (VLA) models is essential due to the diversity of simulators, robot embodiments, and task objectives. Existing compute services, whether offered as direct accelerator rental or...

<https://arxiv.org/abs/2607.16074>
- 9
CONF
A Methodology for Auditable Trustworthiness Levels in AI Lifecycle Governance
ArXiv CS.AI 50%

arXiv:2607.16130v1 Announce Type: cross Abstract: AI governance increasingly requires judgments about whether an AI system remains adequately trustworthy over time, whether observed changes are tolerable, and how such judgments should be documented in a transparent and...

<https://arxiv.org/abs/2607.16130>
- 0
CONF
Evaluating Open-Weight LLMs for Generating Structured Threat Information for Autonomous Vehicle Vulnerabilities
ArXiv CS.AI 50%

arXiv:2607.16175v1 Announce Type: cross Abstract: Connected and Autonomous Vehicles (CAVs) rely on interconnected software and hardware components, including sensors, Electronic Control Units, in-vehicle infotainment systems, and telematics units, where vulnerabilities can...

<https://arxiv.org/abs/2607.16175>
- 1
CONF
Human-Aligned Procedural Level Generation Reinforcement Learning via Text-Level-Sketch Shared Representation
ArXiv CS.AI 50%

arXiv:2508.09860v2 Announce Type: replace Abstract: Human-aligned AI is a critical component of co-creativity, as it enables models to accurately interpret human intent and generate controllable outputs that align with design goals in collaborative content creation. This...

<https://arxiv.org/abs/2508.09860>
- 2
CONF
Mechanistic Interpretability of Cognitive Complexity in LLMs via Linear Probing using Bloom's Taxonomy
ArXiv CS.AI 50%

arXiv:2602.17229v2 Announce Type: replace Abstract: The black-box nature of Large Language Models necessitates novel evaluation frameworks that transcend surface-level performance metrics. This study investigates the internal neural representations of cognitive complexity using...

<https://arxiv.org/abs/2602.17229>

3 CONF

The AI Fiction Paradox

ArXiv CS.AI 50%

arXiv:2603.13545v2 Announce Type: replace Abstract: AI development has a fiction dependency problem. Developers have treated large corpora of modern books, including fiction, as valuable enough to accept substantial cost and legal risk, yet current models still struggle to...

<https://arxiv.org/abs/2603.13545>

4 CONF

Internal Pluralism and the Limits of Pairwise Comparisons

ArXiv CS.AI 50%

arXiv:2607.02672v2 Announce Type: replace Abstract: Local pairwise comparisons are a standard tool for learning how people want decision rules to work, e.g., in participatory design or alignment. However, their use builds in two strong assumptions: that local comparisons are...

<https://arxiv.org/abs/2607.02672>

5 CONF

MaxSAT-Based Feedback for Guiding Vision-Language Models in Sudoku

ArXiv CS.AI 50%

arXiv:2607.12711v2 Announce Type: replace Abstract: Vision--Language Models (VLMs) have recently demonstrated promising performance on structured visual reasoning tasks, including grid-based puzzles. However, despite strong perceptual capabilities, these models lack explicit...

<https://arxiv.org/abs/2607.12711>

6 CONF

BrainPilot: Automating Brain Discovery with Agentic Research

ArXiv CS.AI 50%

arXiv:2607.15079v2 Announce Type: replace Abstract: Understanding the brain increasingly depends on integrating evidence across scales, modalities, and disciplines. Addressing a single research question therefore requires a coordinated sequence of operations, from surveying...

<https://arxiv.org/abs/2607.15079>

7 CONF

Perception-Aligned AI Outputs: End-to-End Visual Prediction for Uncertainty Communication in Clinical Decision-Making

ArXiv CS.AI 50%

arXiv:2205.04599v2 Announce Type: replace-cross Abstract: Explainable Artificial Intelligence (XAI) is essential for trustworthy AI in healthcare, yet many existing methods rely on technical explanations that are difficult for clinicians and patients to interpret. We introduce...

<https://arxiv.org/abs/2205.04599>

8 CONF

A Scaffolded GenAI Lab in Early Undergraduate CS: A Mixed-Methods, Multi-Course Evaluation

ArXiv CS.AI 50%

arXiv:2505.00100v2 Announce Type: replace-cross Abstract: Background and Context. Generative AI (GenAI) tools are increasingly used in programming courses, but we have limited evidence about how brief instruction can foster responsible, learning-oriented use. Objectives. We...

<https://arxiv.org/abs/2505.00100>

9 CONF

A Bit of Freedom Goes a Long Way: Classical and Quantum Algorithms for Reinforcement Learning under a Generative Model

ArXiv CS.AI 50%

arXiv:2507.22854v3 Announce Type: replace-cross Abstract: We propose novel classical and quantum online algorithms for learning finite- and infinite-horizon Markov Decision Processes (MDPs). Our algorithms are based on a hybrid online-offline reinforcement learning model wherein...

<https://arxiv.org/abs/2507.22854>

0 CONF

Unsupervised Deep Learning for Inverse Problems in Computed Tomography

ArXiv CS.AI 50%

arXiv:2508.05321v4 Announce Type: replace-cross Abstract: Assume you encounter an inverse problem that shall be solved for a large number of data, but no ground-truth data is available. To emulate this, in this study we assume it is unknown how to solve the imaging problem of...

<https://arxiv.org/abs/2508.05321>

1 CONF

Latent Fusion Jailbreak: Blending Harmful and Harmless Representations to Elicit Unsafe LLM Outputs

ArXiv CS.AI 50%

arXiv:2508.10029v3 Announce Type: replace-cross Abstract: Safety-aligned large language models can still be manipulated through white-box interventions that modify their internal representations. We introduce Latent Fusion Jailbreak (LFJ), which works by pairing a harmful query...

<https://arxiv.org/abs/2508.10029>

2 CONF

A Systematic Study of Large Language Models for Task and Motion Planning With PDDLStream

ArXiv CS.AI 50%

arXiv:2510.00182v2 Announce Type: replace-cross Abstract: While we know that large language models (LLMs) can solve some planning problems, we do not understand the extent of these capabilities for robotics. One promising direction is to integrate the semantic knowledge of LLMs...

<https://arxiv.org/abs/2510.00182>

3 CONF

Are Heterogeneous Graph Neural Networks Truly Effective for Node Classification? A Causal Perspective

ArXiv CS.AI 50%

arXiv:2510.05750v2 Announce Type: replace-cross Abstract: Graph neural networks (GNNs) have achieved remarkable success in node classification. Building on this progress, heterogeneous graph neural networks (HGNNs) integrate relation types and node and edge semantics to leverage...

<https://arxiv.org/abs/2510.05750>

4 CONF

Analysing Moral Bias in Finetuned LLMs through Mechanistic Interpretability

ArXiv CS.AI 50%

arXiv:2510.12229v3 Announce Type: replace-cross Abstract: Large language models (LLMs) have been shown to internalize human-like biases during finetuning, yet the mechanisms by which these biases manifest remain unclear. In this work, we investigated whether the well-known Knobe...

<https://arxiv.org/abs/2510.12229>

5 CONF

DiffuMamba: High-Throughput Diffusion LMs with Mamba Backbone

ArXiv CS.AI 50%

arXiv:2511.15927v4 Announce Type: replace-cross Abstract: Diffusion language models (DLMs) have emerged as a promising alternative to autoregressive (AR) generation, yet their reliance on Transformer backbones limits inference efficiency due to quadratic attention or KV-cache...

<https://arxiv.org/abs/2511.15927>

6 CONF

Latency-Response Theory Model: Evaluating Large Language Models via Response Accuracy and Chain-of-Thought Length

ArXiv CS.AI 50%

arXiv:2512.07019v3 Announce Type: replace-cross Abstract: The proliferation of Large Language Models (LLMs) necessitates valid evaluation methods to provide guidance for both downstream applications and actionable future improvements. The Item Response Theory (IRT) model with...

<https://arxiv.org/abs/2512.07019>

7 CONF

Hide and Seek in Embedding Space: Geometry-based Steganography and Detection in Large Language Models

ArXiv CS.AI 50%

arXiv:2601.22818v2 Announce Type: replace-cross Abstract: Fine-tuned LLMs can covertly encode prompt secrets into outputs via steganographic channels. Prior work demonstrated this threat but relied on trivially recoverable encodings. We formalize payload recoverability via...

<https://arxiv.org/abs/2601.22818>

8 CONF

Jailbreak Foundry: From Papers to Runnable Attacks for Reproducible Benchmarking

ArXiv CS.AI 50%

arXiv:2602.24009v4 Announce Type: replace-cross Abstract: Jailbreak techniques for large language models (LLMs) evolve faster than benchmarks, making robustness estimates stale and difficult to compare across papers due to drift in datasets, harnesses, and judging protocols. We...

<https://arxiv.org/abs/2602.24009>

9 CONF

Evaluating LLM-Based 0-to-1 Software Generation in End-to-End CLI Tool Scenarios

ArXiv CS.AI 50%

arXiv:2604.06742v2 Announce Type: replace-cross Abstract: The evolution of Large Language Models (LLMs) has catalyzed a paradigm shift towards intent-driven software development, where autonomous agents are expected to design and deliver complete, runnable software systems from...

<https://arxiv.org/abs/2604.06742>

0 CONF

What Is the Minimum Architecture for Prolepsis? Early Irrevocable Commitment Across Tasks in Small Transformers

ArXiv CS.AI 50%

arXiv:2604.15010v2 Announce Type: replace-cross Abstract: When do transformers commit to a decision, and what prevents them from correcting it? We introduce prolepsis: a transformer commits early, task-specific attention heads sustain the commitment, and no layer corrects it....

<https://arxiv.org/abs/2604.15010>

1 CONF

Energy-based Transport for Amortized Bayesian Inference

ArXiv CS.AI 50%

arXiv:2605.15407v3 Announce Type: replace-cross Abstract: We consider amortized Bayesian inference for nonlinear inverse problems using only samples from the joint distribution of parameters and observations, including problems with unknown functions in a Banach space. Classical...

<https://arxiv.org/abs/2605.15407>

2 CONF

Learning the Brain's Dynamics as a Port-Hamiltonian System: A GNN-Surrogate Metriplectic Twin for Non-Equilibrium Cortical Dynamics and Closed-Loop Neuromodulation

ArXiv CS.AI 50%

arXiv:2607.10439v2 Announce Type: replace-cross Abstract: We model human motor cortex, recorded during rest and motor-imagery BCI conditions, as a port-Hamiltonian system: a conservative interconnection (skew-symmetric coupling between band-limited neural phasors) together with...

<https://arxiv.org/abs/2607.10439>

3 CONF

Faithful Autoformalization of Natural Language Assertions

ArXiv CS.AI 50%

arXiv:2607.13303v2 Announce Type: replace-cross Abstract: Formal contracts are essential for software testing and verification, yet writing them remains labor-intensive and error-prone. LLMs offer a promising path toward autoformalization: synthesizing executable assertions from...

<https://arxiv.org/abs/2607.13303>

4 CONF

LLM-as-a-Judge Scores Are Unreliable Optimization Signals in Closed-Loop Table Recognition

ArXiv CS.AI 50%

arXiv:2607.13347v2 Announce Type: replace-cross Abstract: LLM-as-a-judge is widely used to provide feedback and selection signals in closedloop regeneration, but this use remains insufficiently validated. We study it in table recognition, where deterministic TEDS evaluation...

<https://arxiv.org/abs/2607.13347>

5 CONF

NexForge: Scaling Agent Capabilities through Requirement-Driven Task Synthesis for LLMs

ArXiv CS.AI 50%

arXiv:2607.14186v2 Announce Type: replace-cross Abstract: Synthesizing training data to scale agent capabilities in LLM post-training is bottlenecked by substrate-bound task synthesis: tasks are generated from fixed tools, repositories, or skill graphs, so expanding coverage...

<https://arxiv.org/abs/2607.14186>