



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	123	Media
TechCrunch AI	11	Media
HackerNews	9	Community
The Verge AI	6	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 7 stories

VERI

Show HN: Language Model Builder (an app to learn about and build models)

HackerNews

+1 source

90%

HN 14

arXiv:2607.16235v2 Announce Type: replace-cross Abstract: Mobile and wearable devices offer an unprecedented opportunity for continuous, passive health monitoring and active health coaching. However, the largest wearable datasets are not publicly available for research, and...

<https://languagemodelbuilder.com/>

VERI

Show HN: Observability for Coding Agents and LLM Applications

HackerNews

+1 source

90%

HN 3

arXiv:2607.18268v1 Announce Type: new Abstract: Real-world applications that use closed-source large language models (LLMs) need advanced safety measures that go beyond the basic content filters. Content moderation filters such as toxicity and bias have relatively standard...

<https://telemetry.dev/>

VERI

Substack adds an AI detector to help spot blogs written by no one

The Verge AI

+1 source

80%

Substack will now help users determine whether what they're reading may have been written by AI. A new tool coming to the platform can scan posts, notes, replies, and comments to provide an estimate of how much text could be AI-generated or written with AI assistance, according...

<https://www.theverge.com/ai-artificial-intelligence/968855/substack-pangram-ai-detectin...>

VERI

OpenAI says it accidentally hacked Hugging Face with a new AI system

The Verge AI

+1 source

80%

OpenAI says its AI models mistakenly breached open-source AI platform Hugging Face during internal testing. In a blog post on Tuesday, OpenAI writes that GPT-5.6 Sol and "an even more capable pre-release model" discovered vulnerabilities within their sandboxed testing...

<https://www.theverge.com/ai-artificial-intelligence/968988/openai-hugging-face-hack-ai>

VERI

Google launches a cheaper alternative to large AI security models like Mythos

The Verge AI

+1 source

80%

Google is launching Gemini 3.6 Flash alongside a new security model dedicated to quickly finding and patching security vulnerabilities. In a blog post on Tuesday, Google describes Gemini 3.5 Flash Cyber as a "cost-efficient and highly capable alternative" to larger, more...

<https://www.theverge.com/tech/968572/google-gemini-flash-cyber-ai-security-model>

VERI

Anthropic's \$1.5 billion book piracy settlement approved by judge

The Verge AI

+1 source

80%

A federal judge has signed off on Anthropic's \$1.5 billion class action settlement with authors who accused the company of training its AI models on copyrighted books, as reported earlier by Reuters. In an order on Monday, Judge Araceli Martínez-Olguín writes that the settlement...

<https://www.theverge.com/ai-artificial-intelligence/968724/anthropic-authors-settlement...>

VERI

Vector-Bench: Can Models Surgically Edit SVG Code?

ArXiv CS.AI

+1 source

80%

arXiv:2607.19056v1 Announce Type: new Abstract: Instruction-based vector editing requires two capabilities: making a requested change and leaving everything else alone. The second is easy to miss when an output is judged only as a raster image. We introduce Vector-Bench, a...

<https://arxiv.org/abs/2607.19056>

CONFIRMED SIGNALS — 136 stories

CONF

Show HN: OSS Cross-Harness self hosted registry and analytics for AI Agents

HackerNews

50%

HN 21

Hey HN, I built Superserve, a compute layer that lets AI agents live inside isolated Firecracker microVMs with no session time limits. The problem I kept running into: most sandbox providers kill your agent after 24 hours. If you're running something autonomous that needs to...

<https://github.com/Observal/Observal>

CONF

Show HN: Anypost - transactional and broadcast email API, 8¢ per 1k

HackerNews

50%

HN 19

We think it's too damn expensive to send email these days. We built Anypost on modern (non-cloudy) infrastructure to ensure we can offer a top-shelf product without having to hit you with the cloud tax. We're not looking to be the next unicorn. We just want a sustainable product...

<https://anypost.com>

CONF

Show HN: TokenPath - token-level citations for LLM output, read from attention

HackerNews

50%

HN 5

I'm an architect at my company and everyone on my team is now using Claude/Codex agents for their work. Managing context in claude.md/agents.md files is fine but so much of the knowledge required doesn't fit cleanly in one repo or the other. Plus, much of the team (product,...

<https://tokenpath.ai>

CONF

Show HN: Jaybase: Append-only fact store for AI doing business workflows

HackerNews

50%

HN 2

Hey HN, I build custom AI agents for businesses, and one of the things I've noticed is that everyone wants to use AI to do real, critical work, but it's hard to trust it because sometimes agents can go a little off the rails. No one wants their agent deleting all their customer...

<https://github.com/kyle-visner/jaybase>

CONF

Show HN: Findborg - ad-free find engine with community, AI, and real web results

HackerNews

50%

HN 2

No summary — see link for full article.

<https://www.findborg.com/>

CONF

Show HN: Hoop - A sandboxed P2P live collaboration harness for Claude Code

HackerNews

50%

HN 2

No summary — see link for full article.

<https://github.com/bruno-de-queiroz/hoop>

4 **CONF****Show HN: Sorted Receipts - clients dump receipts in one link, LLM sorts them****HackerNews** 50% HN 2

Solo founder. Xero retired Hubdoc in May; the displaced users are mostly solo bookkeepers with 15-30 clients whose replacement options are an enterprise suite (Dext) or per-document credits (AutoEntry). Sorted Receipts is the boring version of what they keep asking for: each...

<https://upload.sortedreceipts.com/demo>
5 **CONF****The Anthropic-Physical Intelligence rumor roiling AI Twitter****TechCrunch AI** 50%

Anthropic and OpenAI's aggressive 2026 acquisition sprees set the stage for a weekend rumor.

<https://techcrunch.com/2026/07/21/the-anthropic-physical-intelligence-rumor-roiling-ai-...>
6 **CONF****Jack Dorsey is taking on Slack with Buzz, a group chat platform for teams and their AI agents****TechCrunch AI** 50%

Buzz is a group chat platform for the workplace that puts humans and their AI agents in the same conversation.

<https://techcrunch.com/2026/07/21/jack-dorsey-is-taking-on-slack-with-buzz-a-group-chat-...>
7 **CONF****AI and the rise of the universal entertainment app****TechCrunch AI** 50%

Over the past decade, streaming platforms competed by dominating individual formats like music, video, podcasts, or audiobooks. Now, as AI makes it easier to create, organize, and recommend content, those distinctions are fading, pushing companies like Spotify, Netflix, YouTube,...

<https://techcrunch.com/2026/07/21/ai-and-the-rise-of-the-universal-entertainment-app/>
8 **CONF****US threatens sanctions against Chinese AI models over IP theft****TechCrunch AI** 50%

Treasury Secretary Scott Bessent said the U.S. could sanction Chinese open AI models over alleged IP theft, expanding the Trump administration's campaign to slow China's AI advances.

<https://techcrunch.com/2026/07/21/us-threatens-sanctions-against-chinese-ai-models-over-...>
9 **CONF****Music streamer Deezer says more than 50% of daily uploads are AI-generated****TechCrunch AI** 50%

Deezer said more than 90,000 AI-generated tracks were uploaded daily on the platform in June.

<https://techcrunch.com/2026/07/21/music-streamer-deezer-says-more-than-50-of-daily-uplo-...>
0 **CONF****Can an Apple lawsuit derail OpenAI's hardware plans?****TechCrunch AI** 50%

On the latest episode of Equity, we debate whether Apple's lawsuit will cast a shadow over OpenAI's much-discussed plans to get into hardware and go public.

<https://techcrunch.com/2026/07/19/can-an-apple-lawsuit-derail-openais-hardware-plans/>
1 **CONF****Advancing next-gen AI with materials science innovation****MIT Tech Review** 50%

The conversation about AI often centers on algorithms, computing power, or huge investments in new semiconductor fabrication plants and hyperscale data centers. But beneath each of these advances is another layer of innovation that makes them possible: advanced materials. Every...

<https://www.technologyreview.com/2026/07/21/1140602/advancing-next-gen-ai-with-material-...>

2 **CONF****Neill Blomkamp's new zombie AI 'film' is just slop warmed over****The Verge AI** 50%

On Monday, District 9 and Gran Turismo director Neill Blomkamp unveiled his latest project: a 13-minute sci-fi short titled Nightborne that's loosely based on Peter Watts' 2014 novel Echopraxia. The short comes from Blomkamp's new AI startup / production company, Barley Studios,...

<https://www.theverge.com/entertainment/968703/neill-blomkamps-nightborne-barley-studios...>

3 **CONF****America needs to stop getting shocked by Chinese AI****The Verge AI** 50%

Last week, two Chinese AI companies unveiled models they say can credibly compete with the best systems from OpenAI and Anthropic. The response was swift and predictable. Markets wobbled, commentators declared Silicon Valley shooketh, and policymakers reached for the familiar...

<https://www.theverge.com/ai-artificial-intelligence/968136/chinese-ai-models-another-sp...>

4 **CONF****SysAdmin: Measuring Instrumental Power-Seeking in Frontier AI****ArXiv CS.AI** 50%

arXiv:2607.18239v1 Announce Type: new Abstract: Power-seeking defined as behaviors where AI systems acquire resources, evade oversight, or resist termination beyond task requirements is identified as a key driver of Loss of Control (LoC) risk. In this work, we introduce...

<https://arxiv.org/abs/2607.18239>

5 **CONF****Calibrated Selective Fact-Checking via Evidence Chain Evaluation****ArXiv CS.AI** 50%

arXiv:2607.18240v1 Announce Type: new Abstract: Large language models (LLMs) can achieve strong fact-checking accuracy, yet forced binary decisions conceal a critical reliability problem: systems may issue confident verdicts even when supporting evidence is weak, sparse, or...

<https://arxiv.org/abs/2607.18240>

6 **CONF****BatchDAG: LLM-Planned Execution Graphs for Scalable Ad-Hoc Analysis Over Enterprise Data****ArXiv CS.AI** 50%

arXiv:2607.18241v1 Announce Type: new Abstract: Large language models (LLMs) excel at analyzing individual documents but break down on exhaustive, cross-entity analytical questions over enterprise-scale datasets due to context overflow, loss of per-entity attribution, and linear...

<https://arxiv.org/abs/2607.18241>

7 **CONF****AI Tool Discovery at Scale: All You Need is DNS****ArXiv CS.AI** 50%

arXiv:2607.18242v1 Announce Type: new Abstract: The coming era of autonomous AI agents demands a discovery mechanism capable of navigating millions of tools, yet existing solutions buckle under $O(N)$ complexity and centralized governance. Instead of building another fragile...

<https://arxiv.org/abs/2607.18242>

8 CONF

From Trajectories to Instructions: Language-Conditioned Meta-Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.18830v1 Announce Type: cross Abstract: Model-Agnostic Meta-Learning (MAML) is a widely used framework for reinforcement learning (RL) that enables efficient transfer by learning global policy parameters that can be rapidly adapted to new tasks. MAML training proceeds...

<https://arxiv.org/abs/2607.18830>

9 CONF

Phionyx: A Deterministic AI Runtime Architecture with Structured State Management and Pre-Response Governance

ArXiv CS.AI 50%

arXiv:2607.18246v1 Announce Type: new Abstract: We present Phionyx, a deterministic AI runtime architecture derived from the broader Echoism interaction framework that introduces a governance-first approach to AI engineering: treating large language model (LLM) outputs as noisy...

<https://arxiv.org/abs/2607.18246>

0 CONF

Beyond Accuracy and Cost: Latency-Aware LLM Query Routing for Dynamic Workloads

ArXiv CS.AI 50%

arXiv:2607.18253v1 Announce Type: new Abstract: Modern language query routers improve inference efficiency by assigning each query to a model that balances response quality and monetary cost. However, current query routers are largely latency-agnostic and do not consider the...

<https://arxiv.org/abs/2607.18253>

1 CONF

Cross-Dialect Generalization Without Retraining: Benchmarks and Evaluation of Schema-Derived Constrained Decoding for MLIR

ArXiv CS.AI 50%

arXiv:2607.18254v1 Announce Type: new Abstract: Multi-Level Intermediate Representation (MLIR) underlies modern ML compiler infrastructure (TensorFlow, JAX/StableHLO, PyTorch Inductor, IREE), yet appears only in trace amounts in code-LM pretraining corpora. MLIR is also...

<https://arxiv.org/abs/2607.18254>

2 CONF

AgentDebugX: An Open-Source Toolkit for Failure Observability, Attribution, and Recovery in LLM Agents

ArXiv CS.AI 50%

arXiv:2607.18754v1 Announce Type: new Abstract: LLM agent failures are difficult to debug because the step where an error surfaces is often not the one that caused it. Existing observability tools replay execution traces but provide little support for identifying the root cause...

<https://arxiv.org/abs/2607.18754>

3 CONF

Probabilistic Concept-Aware Steering for Trustworthy LLM Inference

ArXiv CS.AI 50%

arXiv:2607.18259v1 Announce Type: new Abstract: Steering vectors (SVs), an inference-time intervention technique for large language models (LLMs), guide the generation process by adding a concept-specific direction vector to intermediate activations during inference. However,...

<https://arxiv.org/abs/2607.18259>

4 CONF

Relay-Bench: Evaluating LLMs on Multi-Domain Reasoning Chains

ArXiv CS.AI 50%

arXiv:2607.18438v1 Announce Type: cross Abstract: Introducing Relay-Bench, an unsaturated, holistic, text-only benchmark that measures LLMs' ability to complete an assortment of tasks from distinct domains in a single prompt. The leading model, GPT-5.5 (xHigh), scores 43.3%. The...

<https://arxiv.org/abs/2607.18438>

5 CONF

When JSON Is Not Enough: Semantic Reliability of Schema-Constrained LLM Ordering Agents

ArXiv CS.AI 50%

arXiv:2607.18261v1 Announce Type: new Abstract: LLM agents are increasingly used as transaction compilers: a user states an intent in natural language, and the model emits a structured object that an API can execute. JSON Schema and provider-level structured-output modes are...

<https://arxiv.org/abs/2607.18261>

6 CONF

ProbSPARQL: Querying Knowledge Graphs with Multi-dimensional, Uncertain Numeric Data

ArXiv CS.AI 50%

arXiv:2607.18262v1 Announce Type: new Abstract: The SFB 1574 Circular Factory is building a shared knowledge graph infrastructure for integrating data about returned products. A central challenge is that circular-factory data include numeric measurements that (i) originate from...

<https://arxiv.org/abs/2607.18262>

7 CONF

Position: AI/ML Deepfake Research is Misaligned with AI-Generated Non-Consensual Intimate Imagery (AIG-NCII)

ArXiv CS.AI 50%

arXiv:2607.18263v1 Announce Type: new Abstract: AI-generated non-consensual intimate imagery (AIG-NCII) is not adequately addressed in AI/ML literature regarding AI-generated media, commonly referred to as "deepfakes". While research on deepfakes currently focuses on its...

<https://arxiv.org/abs/2607.18263>

8 CONF

State Compression in Two-Agent LLM Relays: A Closed-World Study of Constraint Preservation

ArXiv CS.AI 50%

arXiv:2607.18265v1 Announce Type: new Abstract: Long-running Large Language Model (LLM)-based agents often accumulate large intermediate traces containing audits, eliminations, and numeric calculations. In practice, this state is compressed before handing it to a downstream...

<https://arxiv.org/abs/2607.18265>

9 CONF

PlotTwist: A Creative Plot Generation Framework with Small Language Models

ArXiv CS.AI 50%

arXiv:2603.16410v2 Announce Type: replace-cross Abstract: Creative plot generation presents a fundamental challenge for language models: transforming a concise premise into a coherent narrative that sustains global coherence, character development, pacing, tone consistency, and...

<https://arxiv.org/abs/2603.16410>

0 CONF

Wisdom of LLM Crowds: Aggregation and Contamination in Language Model Ensembles

ArXiv CS.AI 50%

arXiv:2607.18269v1 Announce Type: new Abstract: The wisdom of crowds -- the finding that aggregating judgments across individuals often outperforms the best individual -- has been extensively studied with human forecasters. Whether the same phenomenon emerges when the ``crowd"...

<https://arxiv.org/abs/2607.18269>

1 CONF

Using LLMs for Explainable, Data-Driven Insight Generation from Time Series

ArXiv CS.AI 50%

arXiv:2607.18271v1 Announce Type: new Abstract: Time series forecasts are widely used in decision-critical domains, where they are rarely consumed without accompanying explanations. Producing such explanations is usually a manual and costly process, and attempts to automate it...

<https://arxiv.org/abs/2607.18271>

2 CONF

Preference-Conditioned Multi-Objective Reinforcement Learning for Runtime-Tunable Transit Signal Priority

ArXiv CS.AI 50%

arXiv:2607.18286v1 Announce Type: cross Abstract: Transit signal priority (TSP) requires balancing competing objectives: reducing bus delay while limiting adverse impacts on non-bus traffic and avoiding extreme waits for a subset of vehicles. Existing reinforcement-learning (RL)...

<https://arxiv.org/abs/2607.18286>

3 CONF

Operational Hallucination and Safety Drift in AI Agents

ArXiv CS.AI 50%

arXiv:2607.18366v1 Announce Type: new Abstract: Large language models (LLMs) serving as planners in tool-using autonomous agents introduce dynamic reliability risks in multi-turn execution. While single-turn safety mechanisms are relatively mature, extended interactions reveal...

<https://arxiv.org/abs/2607.18366>

4 CONF

MAGE: Human-Like Macro Placement via Agentic Multimodal Reasoning

ArXiv CS.AI 50%

arXiv:2607.18536v1 Announce Type: new Abstract: Macro placement still requires substantial manual refinement in industrial physical design flows. We present MAGE (Macro Placement Agentic Engine), a multimodal multi-agent framework for macro placement refinement. MAGE decomposes...

<https://arxiv.org/abs/2607.18536>

5 CONF

Engineering Trustworthy Agentic AI for Critical Systems

ArXiv CS.AI 50%

arXiv:2607.18548v1 Announce Type: new Abstract: Agentic artificial intelligence systems, capable of autonomous perception, planning, tool use, and multi-step action, are increasingly proposed for critical engineering domains where decisions carry physical, operational, or...

<https://arxiv.org/abs/2607.18548>

- 6 **CONF**
- When Does Machine Learning Beat Value Sorting? A Three-Dataset Diagnostic of Exposure-Weighted Shipment Prioritization**
- ArXiv CS.AI 50%
- arXiv:2607.18573v1 Announce Type: new Abstract: Delay-risk models are usually judged by predictive accuracy. What matters in practice is narrower: with capacity to review only a few shipments, which ones should a manager check first? We evaluate whether machine learning clears a...
- <https://arxiv.org/abs/2607.18573>
-
- 7 **CONF**
- QuArch: A Benchmark for Evaluating LLM Reasoning in Computer Architecture**
- ArXiv CS.AI 50%
- arXiv:2510.22087v2 Announce Type: replace-cross Abstract: The field of computer architecture, which bridges high-level software abstractions and low-level hardware implementations, remains absent from current large language model (LLM) evaluations. To this end, we present QuArch...
- <https://arxiv.org/abs/2510.22087>
-
- 8 **CONF**
- Do AI-Native Biotechs Need Departments? Benchmarking Company World Models for AI-Driven Drug Development**
- ArXiv CS.AI 50%
- arXiv:2607.18696v1 Announce Type: new Abstract: AI-native biotechnology companies are often designed by copying human biotech org charts into agent roles. We argue for a different abstraction: a Company World Model, defined as a persistent asset-to-value state representation...
- <https://arxiv.org/abs/2607.18696>
-
- 9 **CONF**
- One Rewrite to Fix Them All? Type-Aware Repair Allocation for Text-to-Image Prompt Optimization**
- ArXiv CS.AI 50%
- arXiv:2607.18724v1 Announce Type: new Abstract: Text-to-image (T2I) generators often fail to follow their prompts faithfully, producing wrong counts, swapped attributes, ambiguous relations, and illegible text. Prompt optimization repairs such failures by rewriting the user...
- <https://arxiv.org/abs/2607.18724>
-
- 0 **CONF**
- AI Tour Meeting: Group Travel Planning by LLM Agents**
- ArXiv CS.AI 50%
- arXiv:2607.18806v1 Announce Type: new Abstract: This paper proposes AI Tour Meeting, a group travel planning framework powered by multiple Large Language Model (LLM)-based agents. The agents are instantiated with distinct personas and collaboratively seek an itinerary that...
- <https://arxiv.org/abs/2607.18806>
-
- 1 **CONF**
- Evaluating medical AI under missing information: same-provider judges and human raters change apparent safety**
- ArXiv CS.AI 50%
- arXiv:2607.18828v1 Announce Type: new Abstract: Readiness stress-testing of medical AI has focused on closed-ended and multimodal benchmarks. We extend it to open-ended clinical conversation under missing information, where safe behavior means recognizing absent information and...
- <https://arxiv.org/abs/2607.18828>
-

2 **CONF****PhoenixRepair: Rethinking Repair Strategy Exploration in Software Agents**

ArXiv CS.AI 50%

arXiv:2607.18859v1 Announce Type: new Abstract: While Large Language Models have greatly advanced automated issue resolution, existing agent-based methods exhibit a fundamental limitation in their insufficient exploration of repair strategies. This insufficiency manifests in two...

<https://arxiv.org/abs/2607.18859>

3 **CONF****NaviAIS: A Scenario-Level Vessel Trajectory Prediction Dataset with Vectorized Lane Priors and the NaviLane Forecasting Framework**

ArXiv CS.AI 50%

arXiv:2607.18887v1 Announce Type: new Abstract: Vessel trajectory prediction in complex maritime environments is essential for traffic management, collision warning, route planning, and autonomous navigation. Although AIS-based learning methods have progressed rapidly, existing...

<https://arxiv.org/abs/2607.18887>

4 **CONF****Enhancing Transformer-based Routing by Encoding Distance via Relative Positional Encoding**

ArXiv CS.AI 50%

arXiv:2607.18909v1 Announce Type: new Abstract: This paper explores Relative Positional Encoding (RPE) as an additive bias in Transformer architectures to solve the Team Orienteering Problem. By embedding in the attention mechanism pairwise spatial relationships among nodes of...

<https://arxiv.org/abs/2607.18909>

5 **CONF****OntoBook: Ontology-Grounded Synthetic Textbooks for Medical Encoder Pretraining**

ArXiv CS.AI 50%

arXiv:2607.18927v1 Announce Type: new Abstract: We present OntoBook, a method that converts medical ontology structure into pretraining signal for encoder language models. Our approach has three stages: random walks through ontology graphs capture hierarchical and causal...

<https://arxiv.org/abs/2607.18927>

6 **CONF****What General Intelligence Requires: Non-Reducible Constraints Across Levels of Description**

ArXiv CS.AI 50%

arXiv:2607.18943v1 Announce Type: new Abstract: General intelligence, of the kind that underwrites the full range of human cognitive achievement, is not a property of computational architecture alone. This paper advances a single thesis: the structural constraints on general...

<https://arxiv.org/abs/2607.18943>

7 **CONF****From Dependency to Compositionality: A Neurosymbolic Lifting of LLM Outputs via Combinatory Categorical Grammar**

ArXiv CS.AI 50%

arXiv:2607.18961v1 Announce Type: new Abstract: Large language models (LLMs) generate fluent text by incrementally predicting the next token from a prefix. Critics in the generative tradition argue that such systems lack genuine grammar; influential replies from the...

<https://arxiv.org/abs/2607.18961>

8 CONF

Mi-Memory: A Lifecycle Memory Framework for Personal AI

ArXiv CS.AI 50%

arXiv:2607.18975v1 Announce Type: new Abstract: Personal AI is moving beyond chat-only interaction toward continuous services that span phones, cars, homes, wearables, cameras, and tools. In this setting, memory cannot remain a cache of prior conversations. It should serve as a...

<https://arxiv.org/abs/2607.18975>

9 CONF

Fishing Out Free Riders: Shapley-Based Reward Attribution for Parallel Reasoning via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.18979v1 Announce Type: new Abstract: Large Language Models (LLMs) excel at multi-step reasoning, yet current parallel reasoning approaches often fail to distinguish the contributions of individual reasoning paths. Many paths may be redundant, misleading, or even...

<https://arxiv.org/abs/2607.18979>

0 CONF

Athena-Brain Technical Report: An Efficient Robot Brain for General Intelligence and Embodied Interaction

ArXiv CS.AI 50%

arXiv:2607.18985v1 Announce Type: new Abstract: Large language models (LLMs) have demonstrated remarkable capabilities in language understanding, reasoning, and world knowledge. As embodied agents become increasingly capable, there is a growing demand for compact models that can...

<https://arxiv.org/abs/2607.18985>

1 CONF

Appearance Pointers -- Multimodal Region Control of Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2607.19344v1 Announce Type: cross Abstract: Controllable image generation remains challenging for creative professionals, who often require precise regional control over materials, object identities, and spatial arrangements that cannot be reliably achieved through text...

<https://arxiv.org/abs/2607.19344>

2 CONF

On the Effectiveness of Pretraining for Graph Combinatorial Optimization

ArXiv CS.AI 50%

arXiv:2607.19072v1 Announce Type: new Abstract: This paper introduces a self-supervised pretraining framework for graph combinatorial optimization specifically designed to address the nature of routing problems like the Traveling Salesman Problem. By utilizing graph contrastive...

<https://arxiv.org/abs/2607.19072>

3 CONF

OpenRTAG: A Comprehensive Benchmark for Robust Text-Attributed Graph Learning under Data Quality Degradation

ArXiv CS.AI 50%

arXiv:2607.19108v1 Announce Type: new Abstract: Text-attributed graphs (TAGs) are an important graph data form that combine relational structure with rich node text. However, real-world TAGs are often imperfect, with quality issues arising from text, structure, and labels, and...

<https://arxiv.org/abs/2607.19108>

- 4 **CONF**
- BioSecBench-Surveillance: A Verifiable Benchmark for AI Agents in Pathogen Genomic Surveillance**
- ArXiv CS.AI 50%
- arXiv:2607.19262v1 Announce Type: new Abstract: As pathogen genomic surveillance scales, the bottleneck is shifting from data generation to analysis. We present BioSecBench-Surveillance, a verifiable benchmark of 100 evaluations testing whether AI agents can infer the right...
- <https://arxiv.org/abs/2607.19262>
-
- 5 **CONF**
- Graph-Based Agentic AI with LangGraph: Workflow Pathways for Long-Running Stateful Business Processes**
- ArXiv CS.AI 50%
- arXiv:2607.19297v1 Announce Type: new Abstract: This paper is a practitioner guide to graph-based workflow pathways for long-running, stateful, multi-step generative AI systems in business processes. Rather than treating LangGraph, a low-level orchestration framework for...
- <https://arxiv.org/abs/2607.19297>
-
- 6 **CONF**
- LLM Detection as an Intervention: Downstream Impact under Strategic User Behavior**
- ArXiv CS.AI 50%
- arXiv:2607.19300v1 Announce Type: new Abstract: As LLM adoption becomes more widespread, there is a growing interest in detecting LLM-generated content, for example through LLM detection tools and through heuristics based on language patterns. Detectors operate as an...
- <https://arxiv.org/abs/2607.19300>
-
- 7 **CONF**
- Broken Gates: Re-evaluating Web Bot Defenses in the Age of LLM Agents**
- ArXiv CS.AI 50%
- arXiv:2607.18659v1 Announce Type: cross Abstract: LLM-based browser agents are rapidly changing the threat landscape for web security. Unlike traditional automation frameworks that execute predefined scripts, these agents can autonomously navigate websites, reason about page...
- <https://arxiv.org/abs/2607.18659>
-
- 8 **CONF**
- Active Electrosensing and Communication in MARL-trained Weakly Electric Fish Collectives**
- ArXiv CS.AI 50%
- arXiv:2511.08436v2 Announce Type: replace-cross Abstract: How complex collective behavior emerges from individual interactions is a fundamental scientific question, but experimental cost and difficulty of simultaneous multi-brain recordings limit direct study in animals. Here we...
- <https://arxiv.org/abs/2511.08436>
-
- 9 **CONF**
- An LLM-powered Agentic Recommendation System for Connected TV Content Discovery**
- ArXiv CS.AI 50%
- arXiv:2607.09988v2 Announce Type: replace-cross Abstract: Recommendation systems, from traditional multi-stage to recent unified generative architectures, face challenges in incorporating diverse contextual signals, such as trending topics, breaking news, cultural events, and...
- <https://arxiv.org/abs/2607.09988>
-

- 0
CONF

Assistant or Actor? Student Trust, Control, and Delegation Regret When Using a General-Purpose AI Agent

ArXiv CS.AI 50%

arXiv:2607.18257v1 Announce Type: cross Abstract: When AI agents shift from answering questions to taking actions, users face a new problem: deciding what to delegate, to a system whose action space they cannot fully anticipate. We call the resulting dissatisfaction delegation...

<https://arxiv.org/abs/2607.18257>
- 1
CONF

The Economics of Autonomy: Real-Time Risk Indexing for Insurable AI-Driven 6G Systems

ArXiv CS.AI 50%

arXiv:2607.18267v1 Announce Type: cross Abstract: The transition to sixth-generation (6G) networks transforms wireless infrastructure into a cognitive substrate supporting Vehicle-to-Everything (V2X), Industrial IoT (IIoT), and Integrated Sensing and Communication (ISAC). In...

<https://arxiv.org/abs/2607.18267>
- 2
CONF

Domain Design for the Cops and Robbers Problem

ArXiv CS.AI 50%

arXiv:2607.18274v1 Announce Type: cross Abstract: Cops and Robbers is a well-studied problem in graph theory. The setting consists of a robber and one or more cops placed on an undirected graph. Taking turns moving throughout the graph, the cops try to capture the robber. The...

<https://arxiv.org/abs/2607.18274>
- 3
CONF

FALCON-Discover: Discovering Concentrated False-Confidence Regions for Calibration

ArXiv CS.AI 50%

arXiv:2607.18278v1 Announce Type: cross Abstract: Calibration is usually evaluated in aggregate, but the most dangerous failures are often local: predictions that remain highly confident despite being wrong. We study this failure mode as false-confidence concentration, the...

<https://arxiv.org/abs/2607.18278>
- 4
CONF

Beyond Single-Dimensional Compression: The Compound Sparsity Frontier of Large Language Models

ArXiv CS.AI 50%

arXiv:2607.18280v1 Announce Type: cross Abstract: Large language models (LLMs) are often compressed through static parameter pruning or dynamic token-level computation, yet aggressive sparsification can trigger rapid performance degradation beyond an essential sparsity boundary....

<https://arxiv.org/abs/2607.18280>
- 5
CONF

FedCC: A Low-Resource Federated Adaptation of Foundation Models for Robust Corpus Callosum localization in Fetal Ultrasound Images

ArXiv CS.AI 50%

arXiv:2607.18283v1 Announce Type: cross Abstract: Accurate localization of the corpus callosum (CC) in fetal ultrasound (US) images is crucial for the early identification of neurodevelopmental abnormalities. However, this task remains highly challenging due to the intrinsic...

<https://arxiv.org/abs/2607.18283>

- 6
CONF

Compressing What Matters: Neuron Importance Meets Data-Aware Low Rank Approximation for Language Model Compression

ArXiv CS.AI 50%

arXiv:2607.18284v1 Announce Type: cross Abstract: To excel at their domain large language models are comprised of billions of parameters. Yet this comes at the cost of huge memory requirements restricting their applicability in resource-constrained environments. To address the...

<https://arxiv.org/abs/2607.18284>
- 7
CONF

Edge-Efficient Transformer for End-to-End RF Spectrum Monitoring

ArXiv CS.AI 50%

arXiv:2607.18285v1 Announce Type: cross Abstract: We present E-SpecFormer (Edge Spectrum monitoring Transformer) for end-to-end automatic modulation and covert channel (CC) recognition. We introduce LiTAN (Linear Tanh Attention Network), a Softmax- and LayerNorm-free attention...

<https://arxiv.org/abs/2607.18285>
- 8
CONF

BearingNAS: Obtaining In-Sensor Intelligent Fault Diagnosis Systems for Bearings Using a Laptop

ArXiv CS.AI 50%

arXiv:2607.18287v1 Announce Type: cross Abstract: This paper introduces BearingNAS, a Hardware-Aware Neural Architecture Search (HW-NAS) framework designed to shift the intelligence directly onto the sensor die via in-sensor processing. BearingNAS frames the search as a...

<https://arxiv.org/abs/2607.18287>
- 9
CONF

Towards Principled Continual Anomaly Detection: A Systematic Framework and Benchmark Scenarios

ArXiv CS.AI 50%

arXiv:2607.18289v1 Announce Type: cross Abstract: Continual anomaly detection (CAD) studies how models can adapt to evolving data distributions while retaining performance on previously observed regimes. CAD benchmarks, however, depend critically on how tasks are defined,...

<https://arxiv.org/abs/2607.18289>
- 0
CONF

Dual-domain fused LSTM modeling for efficient time-dependent reliability analysis

ArXiv CS.AI 50%

arXiv:2607.18291v1 Announce Type: cross Abstract: Time-dependent reliability analysis is crucial for ensuring the long-term safety and performance of engineering systems under uncertainties. However, traditional surrogate model methods often struggle to incorporate...

<https://arxiv.org/abs/2607.18291>
- 1
CONF

Agentic Calibration of Grey-Box Simulation Models: An LLM-Driven Alternative

ArXiv CS.AI 50%

arXiv:2607.18308v1 Announce Type: cross Abstract: Calibration of grey-box simulation models is a constrained optimization problem in which model evaluations are expensive, the parameter space can be high-dimensional, and the search must respect plausibility constraints. Although...

<https://arxiv.org/abs/2607.18308>

- 2

CONF

Distribution-First Population Simulation: Collapse, Calibration, and Recall in Non-WEIRD LLM Persona Modeling

ArXiv CS.AI

50%

arXiv:2607.18310v1 Announce Type: cross Abstract: Synthetic-population tools increasingly run every individual as an independent large language model (LLM) agent. Using real survey microdata, we show that this paradigm has a basic failure mode, and we set a distribution-first...

<https://arxiv.org/abs/2607.18310>
- 3

CONF

Approximating SPR Distance Between Phylogenetic Trees with Graph Neural Networks

ArXiv CS.AI

50%

arXiv:2607.18311v1 Announce Type: cross Abstract: Comparing phylogenetic tree topologies is essential for understanding epidemic dynamics, yet biologically meaningful distances such as the Subtree Prune and Regraft (SPR) distance are NP-hard to compute and intractable on large...

<https://arxiv.org/abs/2607.18311>
- 4

CONF

Dynamic Loss Balancing for Joint SOH and RUL Prediction of Lithium-Ion Batteries via a Rotary SOH-Injected Prior Battery Transformer

ArXiv CS.AI

50%

arXiv:2607.18329v1 Announce Type: cross Abstract: The deployment of reliable lithium-ion battery management systems is crucial for accelerating electrification, yet the joint prognosis of State of Health (SOH) and Remaining Useful Life (RUL) remains severely hindered by task...

<https://arxiv.org/abs/2607.18329>
- 5

CONF

Physics-Guided Masked Multi-Task Network for Edge-Friendly Battery Health Diagnostics from Stochastically Fragmented Charging Profiles

ArXiv CS.AI

50%

arXiv:2607.18330v1 Announce Type: cross Abstract: The deployment of reliable lithium-ion battery management systems is crucial for accelerating electrification, yet the joint prognosis of State of Health (SOH) and Remaining Useful Life (RUL) remains severely hindered by task...

<https://arxiv.org/abs/2607.18330>
- 6

CONF

Federated Lightweight Fine-Tuning

ArXiv CS.AI

50%

arXiv:2607.18343v1 Announce Type: cross Abstract: Federated fine-tuning is bottlenecked by communication: FedAvg and pseudo-gradient schemes transmit a payload that scales with the model, and gradient compression shrinks it by only a constant factor. We take a different lever....

<https://arxiv.org/abs/2607.18343>
- 7

CONF

FSDBN: Foreground-Aware EEG--Visual Alignment via Dynamic Brain Networks

ArXiv CS.AI

50%

arXiv:2607.18344v1 Announce Type: cross Abstract: EEG-based visual decoding provides a non-invasive pathway for interpreting visual semantics. However, existing methods often overlook the perceptual asymmetry between foreground and background in complex scenes, leading to...

<https://arxiv.org/abs/2607.18344>

8 CONF

Addressing Limited Data in Auditory Attention Decoding with Diffusion Generative Models

ArXiv CS.AI 50%

arXiv:2607.18345v1 Announce Type: cross Abstract: Limited training data constrains deep learning models for Auditory Attention Decoding (AAD) in hearing aids (HAs). AAD uses electroencephalogram (EEG) data to decode listener's attention, enabling real-time tracking of specific...

<https://arxiv.org/abs/2607.18345>

9 CONF

An Analysis of Residual-Stream Geometry Across Transformer Depth

ArXiv CS.AI 50%

arXiv:2607.18348v1 Announce Type: cross Abstract: We propose a transition-centred geometric analysis of transformer residual streams. Relative displacement measures how \emph{far} representations move between consecutive layers, and orthogonal Procrustes analysis separates each...

<https://arxiv.org/abs/2607.18348>

0 CONF

Multi-layer MIMO Relay as Deep Physical Neural Networks: Power Amplifiers as Activation Functions

ArXiv CS.AI 50%

arXiv:2607.18354v1 Announce Type: cross Abstract: Wireless physical neural networks (WPNNs) embed neural computation directly into analog hardware, offering lower energy consumption and latency than conventional digital implementations. In this paper, we propose a deep WPNN in...

<https://arxiv.org/abs/2607.18354>

1 CONF

Decode-Time Grammars: Constrained LLM Generation over a Refinement Order of Grammar Fragments

ArXiv CS.AI 50%

arXiv:2607.18357v1 Announce Type: cross Abstract: Large language models now write a growing share of the world's code, increasingly inside agents and serving systems that compile, execute, or dispatch generated code without line-by-line review. This works well for mainstream...

<https://arxiv.org/abs/2607.18357>

2 CONF

HALLMARK: Diagnosing Three Failure Modes in LLM Citation Verifiers

ArXiv CS.AI 50%

arXiv:2607.18360v1 Announce Type: cross Abstract: Large language models (LLMs) now routinely draft literature reviews and assist with academic writing, which means a higher risk of fabricated references: GPTZero found 53 papers with hallucinated citations among NeurIPS 2025's...

<https://arxiv.org/abs/2607.18360>

3 CONF

A Controlled Study of Attention-Only Transformers

ArXiv CS.AI 50%

arXiv:2607.18363v1 Announce Type: cross Abstract: Feed-forward networks hold two thirds of a transformer's non-embedding parameters, yet the architecture has not received a necessity test that controls parameters, compute, and depth at once. We pretrain attention-only decoder...

<https://arxiv.org/abs/2607.18363>

- 4

CONF

Adversarial Robustness of Phishing Email Detection: A Comparative Study of TF-IDF + Logistic Regression and Fine-Tuned DistilBERT

ArXiv CS.AI

50%

arXiv:2607.18429v1 Announce Type: cross Abstract: Phishing emails remain one of the most persistent cybersecurity threats, and machine-learning classifiers are widely used to detect them. Most reported detection accuracies, however, are measured on clean, in-distribution test...

<https://arxiv.org/abs/2607.18429>
- 5

CONF

Trusted Credentials, Untrusted Behavior: Benchmarking LLM-Agent Security in High-Performance Computing

ArXiv CS.AI

50%

arXiv:2607.18485v1 Announce Type: cross Abstract: Large language model (LLM) agents are starting to take on routine work in high-performance computing (HPC), including monitoring Slurm jobs, diagnosing failed builds, inspecting simulation output, and coordinating scientific...

<https://arxiv.org/abs/2607.18485>
- 6

CONF

Towards an Automated Test of LLM Security Knowledge

ArXiv CS.AI

50%

arXiv:2607.18496v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used for a range of software, hardware and human-centered security tasks. Consequently, LLM performance on security tasks is an active area of measurement and research, often with a...

<https://arxiv.org/abs/2607.18496>
- 7

CONF

Querying Multimodal Scientific Papers with AI: Practices and Preferences Across Blind, Low-Vision, and Sighted Scientists

ArXiv CS.AI

50%

arXiv:2607.18514v1 Announce Type: cross Abstract: Visual diagrams, figures, and tables are central to scientific papers, and convey information beyond what is captured in text. While blind or low-vision (BLV) scientists have traditionally relied on static alternative text to...

<https://arxiv.org/abs/2607.18514>
- 8

CONF

EduPanel: A Three-Agent LLM Judge for Teaching Videos -- Reliability, Complementarity, and Human Trust Calibration

ArXiv CS.AI

50%

arXiv:2607.18529v1 Announce Type: cross Abstract: Teaching videos are becoming a major medium for education, creating a growing need for scalable evaluation of their pedagogical quality. Existing automatic judges do not fully address this setting because teaching quality depends...

<https://arxiv.org/abs/2607.18529>
- 9

CONF

Censoring-Aware In-Context Learning for Generalized Supplier Lead Time Estimation in Supply Chain Planning

ArXiv CS.AI

50%

arXiv:2607.18530v1 Announce Type: cross Abstract: Supplier lead time forecasting is a central input to material requirements planning, inventory optimization, and supply chain risk management. However, many industrial lead time datasets are naturally right-censored: at the time...

<https://arxiv.org/abs/2607.18530>

00 CONF

Operational Proto-Introspection in Looped Language Models: Process-Quality Taps, Executable Branching, and the Readout-Control Boundary

ArXiv CS.AI 50%

arXiv:2607.18553v1 Announce Type: cross Abstract: Can a language model read the quality of ongoing computation, and can an external intervention turn that readout into better outcomes? We test both questions in a frozen 2.6B looped transformer, Ouro-RLTT. On GSM8K, a strict...

<https://arxiv.org/abs/2607.18553>

01 CONF

The Story Shapes the Agent: Narrative Priors in LLM Behavior

ArXiv CS.AI 50%

arXiv:2607.18566v1 Announce Type: cross Abstract: Persona prompting is widely used to steer LLM agent behavior, yet the narrative framing of a task can matter more than the assigned persona. We isolate this effect through structural isomorphism, constructing three text-based...

<https://arxiv.org/abs/2607.18566>

02 CONF

Intelligent Multi-UAV Navigation in ITNTNs: A Hierarchical LLM Approach

ArXiv CS.AI 50%

arXiv:2607.18604v1 Announce Type: cross Abstract: The deployment of high-speed Uncrewed Aerial Vehicles (UAVs) in 3D aerial highways necessitates robust coordination of physical flight kinematics and multi-tier network handovers. While Deep Reinforcement Learning (DRL) offers...

<https://arxiv.org/abs/2607.18604>

03 CONF

Deep Learning Estimation of Sex, Age, Height, and Weight from CT-derived Digitally Reconstructed Radiographs

ArXiv CS.AI 50%

arXiv:2607.18638v1 Announce Type: cross Abstract: Purpose: To develop and validate a deep learning ensemble for estimating adult sex, age, height, and weight from coronal digitally reconstructed radiographs (DRRs) generated from diagnostic CT. Materials and Methods: This...

<https://arxiv.org/abs/2607.18638>

04 CONF

Attributes Should Come from Images, Not Class Names: Distribution-Conditioned Attribute Selection for Vision-Language Models

ArXiv CS.AI 50%

arXiv:2607.18695v1 Announce Type: cross Abstract: A popular route to interpretable zero-shot classification asks a large language model (LLM) to describe each class name and prompts CLIP with the resulting descriptors. We show that these descriptors carry little visual evidence...

<https://arxiv.org/abs/2607.18695>

05 CONF

Find Before You Fine-Tune: A Diagnostic Study of Small LLMs for Cybersecurity QA

ArXiv CS.AI 50%

arXiv:2607.18725v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly fine-tuned for critical-domain Question-Answering (QA), yet choosing which small model to adapt, before paying the cost of adaptation, remains difficult. Fine-tuning can improve...

<https://arxiv.org/abs/2607.18725>

06 CONF

Bounding Boxes to Improve Small Language Model Performance on Vision-Based Grading Tasks

ArXiv CS.AI 50%

arXiv:2607.18767v1 Announce Type: cross Abstract: The deployment of Small Language Models (SLMs) in educational settings offers significant advantages in terms of privacy, cost, and scalability. However, SLMs often struggle with complex vision-based tasks, such as grading...

<https://arxiv.org/abs/2607.18767>

07 CONF

LakeQuest: A Three-Domain Benchmark for Grounded Question Answering across Data Lakes

ArXiv CS.AI 50%

arXiv:2607.12310v2 Announce Type: replace-cross Abstract: While modern question answering (QA) systems excel on clean, schema-aligned corpora, real-world knowledge is rarely so neatly packaged. Answering questions over enterprise and scientific data lakes requires systems to...

<https://arxiv.org/abs/2607.12310>

08 CONF

Cross-Agent Campaign Attribution: Linking Asynchronous Attacks Across LLM Agents

ArXiv CS.AI 50%

arXiv:2607.18826v1 Announce Type: cross Abstract: LLM-agent defenses are typically evaluated one session at a time. In deployment, however, attacks can be distributed across independent agents, teams, and runtimes, leaving each local guardrail with only a sparse fragment. We...

<https://arxiv.org/abs/2607.18826>

09 CONF

Public perceptions of AI-driven decision-making in healthcare: A structural equation modeling approach

ArXiv CS.AI 50%

arXiv:2607.18884v1 Announce Type: cross Abstract: Artificial intelligence (AI) is increasingly integrated into healthcare to support diagnostics, decision-making, and administrative processes. However, the successful implementation of AI depends not only on technical performance...

<https://arxiv.org/abs/2607.18884>

10 CONF

Circuit Claims Depend on What Is Extracted and How It Is Compared

ArXiv CS.AI 50%

arXiv:2607.18921v1 Announce Type: cross Abstract: Circuit extraction identifies a small set of model components whose presence preserves a target behavior under ablation, and the resulting circuit is often read as the mechanism behind that behavior. We argue that this reading is...

<https://arxiv.org/abs/2607.18921>

11 CONF

Functional Equivalence and Geometric Diversity in Neural Network Approximations: An Empirical Characterization

ArXiv CS.AI 50%

arXiv:2607.18930v1 Announce Type: cross Abstract: The Universal Approximation Theorem states that a neural network with a single hidden layer is sufficient to approximate any continuous univariate function on a compact domain to arbitrary error. However, the uniqueness of such...

<https://arxiv.org/abs/2607.18930>

12 CONF

AutoJourn: Multi-Perspective Summarisation, Bias Detection and Bias Neutralisation for LLM-Generated News in Automated Journalism

ArXiv CS.AI 50%

arXiv:2607.18983v1 Announce Type: cross Abstract: We present AutoJourn, a demonstration system for multi-perspective news generation and bias-aware evaluation using large language models (LLMs). The system tackles three core challenges in responsible automated journalism:...

<https://arxiv.org/abs/2607.18983>

13 CONF

Computational Humor with Multimodal LLMs: Methods, Datasets, Evaluation, and Challenges

ArXiv CS.AI 50%

arXiv:2607.19011v1 Announce Type: cross Abstract: Multimodal humor in memes, cartoons, and comics remains difficult for AI systems because intended meaning depends on non-literal mechanisms, shared cultural knowledge, and communicative intent rather than literal scene...

<https://arxiv.org/abs/2607.19011>

14 CONF

Spectral Higher-Order Neural Networks Have Sharp Expressivity Bounds

ArXiv CS.AI 50%

arXiv:2607.19042v1 Announce Type: cross Abstract: Neural hypergraphs are a natural generalization of neural networks, the reference models in modern machine learning. Yet, their deployment has proven demanding: the number of weighted hyperedges required leads to an intractable...

<https://arxiv.org/abs/2607.19042>

15 CONF

Where Should Optimizer State Live? Tiered State Allocation for Memory-Efficient Mixture-of-Experts Training

ArXiv CS.AI 50%

arXiv:2607.19058v1 Announce Type: cross Abstract: Optimizer state is the largest single line item in the memory budget of mixture-of-experts (MoE) training: on a 6.78B-parameter MoE language model, AdamW keeps 50.6 GB of first and second moments to update 12.6 GB of bfloat16...

<https://arxiv.org/abs/2607.19058>

16 CONF

Deep learning-based prediction of time-resolved adhesive forces in viscoelastic Hertzian contacts

ArXiv CS.AI 50%

arXiv:2607.19060v1 Announce Type: cross Abstract: Fast prediction of the response of adhesive soft viscoelastic contacts represents a current challenge in soft robotics and for gripping and manipulation tasks. Determining the complete time-resolved force trajectory requires full...

<https://arxiv.org/abs/2607.19060>

17 CONF

DAIS: Dependency-Aware Intermediate QA Supervision for Complex Reasoning

ArXiv CS.AI 50%

arXiv:2607.19088v1 Announce Type: cross Abstract: Chain-of-thought (CoT) supervision exposes intermediate rationales, but flat rationale targets usually optimize a single reasoning sequence and provide limited supervision on how local conclusions should support later decisions....

<https://arxiv.org/abs/2607.19088>

18 CONF

MEDIC: Comprehensive Evaluation of Leading Indicators for LLM Safety and Utility in Clinical Applications

ArXiv CS.AI 50%

arXiv:2409.07314v3 Announce Type: replace-cross Abstract: While Large Language Models (LLMs) achieve superhuman performance on standardized medical licensing exams, these static benchmarks have become saturated and increasingly disconnected from the functional requirements of...

<https://arxiv.org/abs/2409.07314>

19 CONF

MIRAGE: Multi-scale Lesion-Informed Representation with Auxiliary Guidance for MRI Contrast Enhancement

ArXiv CS.AI 50%

arXiv:2607.19137v1 Announce Type: cross Abstract: Inferring contrast enhancement from one pre-contrast breast MRI slice is underdetermined: post-contrast appearance contains physiological information that is not uniquely encoded in baseline anatomy. Optimizing only paired pixel...

<https://arxiv.org/abs/2607.19137>

20 CONF

Breaking the Homogeneity Assumption: Specialized Multi-Generator Adversarial Learning for Rare Failure Detection in Predictive Maintenance

ArXiv CS.AI 50%

arXiv:2607.19153v1 Announce Type: cross Abstract: Supervised learning models in the predictive maintenance field are regularly trained on highly imbalanced industrial datasets: machine failures occur rarely but have a disproportionate effect on operations. In addition to the...

<https://arxiv.org/abs/2607.19153>

21 CONF

MIRA-Ev:A Benchmark for Granular Evidence Detection and Relational Reasoning in Clinical Exams

ArXiv CS.AI 50%

arXiv:2607.19201v1 Announce Type: cross Abstract: Clinical NLP evaluation remains dominated by multiple-choice question answering (MCQA), which scores only final-answer accuracy and cannot detect when a model reaches the correct diagnosis while grounding it in irrelevant...

<https://arxiv.org/abs/2607.19201>

22 CONF

The Price of Reasoning: Cost-Quality Tradeoffs in Reinforcement Learning for Neural Machine Translation

ArXiv CS.AI 50%

arXiv:2607.19226v1 Announce Type: cross Abstract: Reinforcement learning with verifiable rewards (RLVR) has been established as a viable paradigm for the post-training of Large Language Models (LLMs), including downstream tasks, such as Neural Machine Translation (NMT). With the...

<https://arxiv.org/abs/2607.19226>

23 CONF

Inference-Time Steering for Cross-Lingual Factual Consistency in LLMs

ArXiv CS.AI 50%

arXiv:2607.19243v1 Announce Type: cross Abstract: Although Large Language Models (LLMs) demonstrate remarkable multilingual fluency, their internal knowledge representations remain disproportionately biased toward high-resource languages. This leads to cross-lingual factual...

<https://arxiv.org/abs/2607.19243>

24 CONF

Prompt Design at Scale: How Format, Instruction Count, and Context Length Shape Instruction Adherence and Hallucination in Large Language Models

ArXiv CS.AI 50%

arXiv:2607.19257v1 Announce Type: cross Abstract: Practitioners make three prompt-design decisions with almost no controlled evidence behind them: how to format instructions and context (markdown, plain text, prose, or tabular), how many simultaneous instructions a system prompt...

<https://arxiv.org/abs/2607.19257>

25 CONF

Benchmarking Generalization in Financial Statement Fraud Detection: robust evaluation and novel tasks

ArXiv CS.AI 50%

arXiv:2607.19259v1 Announce Type: cross Abstract: Financial statement fraud detection (FSFD) is crucial for market integrity but faces challenges from increasingly sophisticated schemes and under-utilized textual data in financial reports. Existing methods often rely on random...

<https://arxiv.org/abs/2607.19259>

26 CONF

The safety failures we are not instrumenting: a perspective on hidden safety-critical challenges in modern AI systems

ArXiv CS.AI 50%

arXiv:2607.19292v1 Announce Type: cross Abstract: Current AI safety discourse still focuses disproportionately on visible failures, including obvious harms, dramatic misuse, and hypothetical catastrophic scenarios. That focus is incomplete. In deployed systems, many of the most...

<https://arxiv.org/abs/2607.19292>

27 CONF

Riemannian Deep Learning: Modules, Networks, and Geometries

ArXiv CS.AI 50%

arXiv:2607.19305v1 Announce Type: cross Abstract: Deep neural networks on manifold-valued representations have attracted growing interest, but many basic components remain tied to specific manifolds, rely on Euclidean approximations, or require costly and numerically fragile...

<https://arxiv.org/abs/2607.19305>

28 CONF

Provable diffusion-based posterior sampling for linear inverse problems via DDIM

ArXiv CS.AI 50%

arXiv:2607.19333v1 Announce Type: cross Abstract: Diffusion-based methods have achieved remarkable empirical success in solving inverse problems. However, many existing posterior samplers either lack rigorous theoretical guarantees or incur substantial computational overhead. We...

<https://arxiv.org/abs/2607.19333>

29 CONF

Dr. Zero: Self-Evolving Search Agents without Training Data

ArXiv CS.AI 50%

arXiv:2601.07055v2 Announce Type: replace Abstract: As high-quality data becomes increasingly difficult to obtain, self-evolution without curated training data has emerged as a promising paradigm. This approach allows large language models (LLMs) to autonomously generate and...

<https://arxiv.org/abs/2601.07055>

30 CONF

Context-Masked Truncated Reasoning Audits for Answer-Key Dependence in LLM Tutors

ArXiv CS.AI 50%

arXiv:2607.04572v2 Announce Type: replace Abstract: Large language model (LLM) tutors may have access to teacher notes, answer keys, rubrics, or retrieved solutions while producing student-facing explanations. We study whether truncated reasoning probes can distinguish direct...

<https://arxiv.org/abs/2607.04572>

31 CONF

FormalAnalyticGeo: A Neural-Symbolic Based Framework for Multimodal Analytic Geometry Problem Generation

ArXiv CS.AI 50%

arXiv:2607.12982v2 Announce Type: replace Abstract: Math reasoning has achieved significant progress with the rapid advancement of Multimodal Large Language Models (MLLMs), however analytic geometry remains largely underexplored, primarily due to the scarcity of annotated...

<https://arxiv.org/abs/2607.12982>

32 CONF

Chi-Square Wavelet Graph Neural Networks for Heterogeneous Graph Anomaly Detection

ArXiv CS.AI 50%

arXiv:2505.18934v2 Announce Type: replace-cross Abstract: Graph Anomaly Detection (GAD) in heterogeneous networks presents unique challenges due to node and edge heterogeneity. Existing Graph Neural Network (GNN) methods primarily focus on homogeneous GAD and thus fail to...

<https://arxiv.org/abs/2505.18934>

33 CONF

RoboInspector: Unveiling the Unreliability of Policy Code for LLM-enabled Robotic Manipulation

ArXiv CS.AI 50%

arXiv:2508.21378v2 Announce Type: replace-cross Abstract: Large language models (LLMs) demonstrate remarkable capabilities in reasoning and code generation, enabling robotic manipulation to be initiated with just a single instruction. The LLM carries out various tasks by...

<https://arxiv.org/abs/2508.21378>

34 CONF

Hyperdimensional Probe: Decoding LLM Representations via Vector Symbolic Architectures

ArXiv CS.AI 50%

arXiv:2509.25045v3 Announce Type: replace-cross Abstract: Despite their capabilities, Large Language Models (LLMs) remain opaque with limited understanding of their internal representations. Current interpretability methods either focus on input-oriented feature extraction, such...

<https://arxiv.org/abs/2509.25045>

35 CONF

Breaking the MoE LLM Trilemma: Dynamic Expert Clustering with Structured Compression

ArXiv CS.AI 50%

arXiv:2510.02345v4 Announce Type: replace-cross Abstract: Mixture-of-Experts (MoE) Large Language Models (LLMs) face a trilemma of load imbalance, parameter redundancy, and communication overhead. We introduce a unified framework based on dynamic expert clustering and structured...

<https://arxiv.org/abs/2510.02345>

36 CONF

ImplicitRDP: An End-to-End Visual-Force Diffusion Policy with Structural Slow-Fast Learning

ArXiv CS.AI 50%

arXiv:2512.10946v2 Announce Type: replace-cross Abstract: Human-level contact-rich manipulation relies on the distinct roles of two key modalities: vision provides spatially rich but temporally slow global context, while force sensing captures rapid local contact dynamics....

<https://arxiv.org/abs/2512.10946>

37 CONF

Training and Simulation of Quadrupedal Robot in Adaptive Stair Climbing and Descending for Indoor Firefighting: An End-to-End Reinforcement Learning Approach

ArXiv CS.AI 50%

arXiv:2602.03087v2 Announce Type: replace-cross Abstract: Quadruped robots are used for primary searches during the early stages of indoor fires. A typical primary search involves quickly and thoroughly looking for victims under hazardous conditions and monitoring flammable...

<https://arxiv.org/abs/2602.03087>

38 CONF

M-RAG: Semantic Key-Value Indexing for Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2603.26667v2 Announce Type: replace-cross Abstract: Retrieval-augmented generation (RAG) turns external documents into evidence for large language models. In practice, this is also a data access problem: a system must decide what to index, what to retrieve, and what...

<https://arxiv.org/abs/2603.26667>

39 CONF

Don't Blame the Large Language Model: How Agent Harness Evolution Shapes Coding Agent Quality

ArXiv CS.AI 50%

arXiv:2607.03691v2 Announce Type: replace-cross Abstract: Coding agents, autonomous systems that use large language models (LLMs) to resolve software engineering tasks, rely on agent harness: a middleware layer in between a developer and a large language model that orchestrates...

<https://arxiv.org/abs/2607.03691>

40 CONF

"I understand your perspective": LLM Persuasion through the Lens of Communicative Action Theory

ArXiv CS.AI 50%

arXiv:2606.08076v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) can generate high-quality arguments, yet their ability to engage in nuanced and persuasive communicative actions remains largely unexplored. This work explores the persuasive potential of LLMs...

<https://arxiv.org/abs/2606.08076>

41 CONF

Reframing AI Loss of Control: What Control Is, How to Have It, How to Lose It

ArXiv CS.AI 50%

arXiv:2606.12442v2 Announce Type: replace-cross Abstract: At present, loss of control risks have gained much prominence in public discussion, particularly in relation to AI, with extensive discourse present among academics, frontier labs, and even governments. However, in the...

<https://arxiv.org/abs/2606.12442>

42 CONF

Team RAS in 11th ABAW Competition: Multimodal Ambivalence Recognition Approach

ArXiv CS.AI 50%

arXiv:2607.14702v2 Announce Type: replace-cross Abstract: Automatic recognition of ambivalence and hesitancy is challenging because these states may be expressed through inconsistent linguistic, acoustic, facial, and contextual patterns, while top-performing systems often rely...
<https://arxiv.org/abs/2607.14702>

43 CONF

Real-World Evaluation of an AI Agent Drafting Translational Impact Summaries

ArXiv CS.AI 50%

arXiv:2607.16989v2 Announce Type: replace-cross Abstract: Introduction. Clinical and Translational Science Award (CTSA) programs must document their scholars' research impact, but assembling each scholar's record by hand takes staff an estimated 15 hours and does not scale to a...
<https://arxiv.org/abs/2607.16989>