



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	92	Media
HackerNews	13	Community
TechCrunch AI	2	Media
The Verge AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: TamedTable, AI ETL in Natural Language

HackerNews

+1 source

90%

HN 8

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. A fundamental flaw leaves LLMs strikingly vulnerable to attack It is impossible to make large language models fully secure against hacks...

<https://www.tamedtable.com/>

VERI

Show HN: CI33-opLM, a 236M language model with a reversible operator bottleneck

HackerNews

+1 source

90%

HN 2

arXiv:2509.17192v3 Announce Type: replace Abstract: LLM-based social simulations can make a generated transcript look like a single behavioral signal, but the model behind that transcript may be doing several different jobs: choosing what an actor says or does, deciding what...

<https://ci33.t3atlas.dev/>

VERI

Show HN: Growth-Ratio Energy Function as Leading Indicator of Agent Task Failure

HackerNews

+1 source

90%

HN 2

arXiv:2607.29062v1 Announce Type: new Abstract: Model capabilities have improved in large part due to scaling chain of thought. This has been a promising development for AI safety--where models verbalize their reasoning, it is possible to monitor it. However, in some cases,...

<https://vishalvermalabs.com/papers/empirical-lyapunov-stability-agent-failure/>

VERI

MOSAIC: Masked Outsourcing of Secure AI Computations

ArXiv CS.AI

+1 source

80%

arXiv:2607.29221v1 Announce Type: cross Abstract: We address the challenge of securely and efficiently outsourcing AI computations from a trusted but computationally weak client to an untrusted but powerful server, in the setting where the client holds both the input and the...

<https://arxiv.org/abs/2607.29221>

CONFIRMED SIGNALS — 102 stories

CONF

Show HN: MicroCodex Coding Agent - OpenAI/codex reimplemented in C++ <1MB binary

HackerNews

50%

HN 19

That is a tool that I have developed based on my own needs since it was a shame not to use AI for requirement generation. There are ancient tools like IBM DOORS but they require buying seat licences and they don't have any support for LLMs. I see that requirement management is a...

<https://github.com/paoloanzn/microcodex>

CONF

Show HN: Alfa. Killing AI hallucinations with resonance

HackerNews

50%

HN 12

This code predicts the exact time and price of a future price point, and can also construct a curve of future prices. It can be used to decipher any process that has a graph—for example, a graph of mutual understanding between artificial intelligence and a human.

<https://github.com/ConsciousGroupMind/SKYNET-800---Collective-Intelligence-Forecasting-...>

CONF

Show HN: Btfy - a blockchain that uses weather observations

HackerNews 50% HN 10

Author here. ~660 notes published from a private Obsidian vault of 8,360 (3M words). Started in OneNote around 2011, moved to Obsidian in 2021 [1], imported everything with a Python script [2]. Notes are compounding over time, and that is where my writing comes from. The...

<https://github.com/kotagit75/btfy>

CONF

Show HN: Do Codex skills save tokens? A six-run task-size benchmark

HackerNews 50% HN 2

No summary — see link for full article.

<https://codex-howto-benchmark.nguyenvantamdk2.chatgpt.site>

CONF

Show HN: AI Visibility: Check whether ChatGPT/Claude/Gemini recommend your brand

HackerNews 50% HN 2

No summary — see link for full article.

<https://aivisibility.pro/>

0 CONF

Show HN: Chinese are offering Claude/Codex offers 90% off

HackerNews 50% HN 2

Chinese resellers are offering *Claude and Codex API access at up to 90% discounts* compared with Anthropic and OpenAI's official API pricing. I've heard that some of them use open-source projects like *Sub2API* to achieve this. Their business model appears to involve purchasing...

<https://news.ycombinator.com/item?id=49151751>

1 CONF

Show HN: Kota - Bring AI agent CLIs into the same room

HackerNews 50% HN 2

A couple months ago, I was tired of keep copy-paste across different Agent Chatbot tabs, feeling like a slave for AI sessions, so I created Kota for my own work. Kota was designed around a concept of Knights of the Round Table. It treats human as part of a project team, and -...

<https://www.kota.place/>

2 CONF

Show HN: Atomadic- Zero-LLM Sub-200us MCP Action Interlock (Storefront Restored)

HackerNews 50% HN 2

No summary — see link for full article.

<https://atomadic.tech/orchard.html>

3 CONF

Show HN: Mapping Jensen Huang's five-layer AI stack to engineering degrees

HackerNews 50% HN 2

Hi HN, I mapped Jensen Huang's five layers of AI, energy, chips, infrastructure, models, and applications, to engineering disciplines. It is aimed at students asking two questions: Where does my degree fit into this landscape? Should I specialise now or stay broad? I would...

<https://ai-five-layer-map.pages.dev/>

4 CONF

Show HN: Mousecrack - Bypass captchas with deep learning

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/puffinsoft/mousecrack>

5 **CONF****Sam Altman and AI's decel debate****TechCrunch AI** 50%

On the latest episode of Equity, we discuss why Sam Altman has calling on the industry to "pace the rate of AI development."
<https://techcrunch.com/2026/08/02/sam-altman-and-ais-decel-debate/>

6 **CONF****Fender's CEO seems to think your bandmates are just analog AI****The Verge AI** 50%

Fender CEO Edward "Bud" Cole gave an interview to T3 in May celebrating the 75th anniversary of the Telecaster with comments on AI and music that initially flew under the radar. But it has started making the rounds recently, pouring more fuel on an already raging fire of bad PR...
<https://www.theverge.com/ai-artificial-intelligence/974265/fender-ceo-bud-cole-ai-music>

7 **CONF****Is paying artists enough to convince them to embrace AI?****The Verge AI** 50%

Illustrators have spent years sounding the alarm about generative artificial intelligence startups training their models on artists' work without permission. They've pointed out how the practice is tantamount to theft, and in response, many gen AI boosters have argued that it's...
<https://www.theverge.com/ai-artificial-intelligence/974018/pippa-seedance-artist-royalties>

8 **CONF****OpenClaw and Ollama in Agentic AI: Toward Fully Autonomous and Scalable AI Agent Systems****ArXiv CS.AI** 50%

arXiv:2607.28629v1 Announce Type: new Abstract: The rapid transition from reactive large language models (LLMs) to persistent, action-capable systems has exposed critical gaps in the architectural understanding of Agentic AI, particularly in separating inference, orchestration,...
<https://arxiv.org/abs/2607.28629>

9 **CONF****Can AI Evaluate AI Scientists? A Benchmarking Study of Autonomous Research Generation Systems Using Automated Multi-Model Review****ArXiv CS.AI** 50%

arXiv:2607.28631v1 Announce Type: new Abstract: AI Scientist systems capable of autonomous research have the potential to significantly accelerate scientific discovery. However, evaluating and comparing the quality of AI-generated papers remains an open challenge. We propose and...
<https://arxiv.org/abs/2607.28631>

0 **CONF****LLM Framework for Discovering Major Mathematical Conjectures: AI's Quest for the Next Riemann Hypothesis****ArXiv CS.AI** 50%

arXiv:2607.28632v1 Announce Type: new Abstract: Major mathematical conjectures still depend heavily on expert intuition, so a unified method for the systematic generation and validation of conjectures with substantial mathematical potential remains unavailable. We present a...
<https://arxiv.org/abs/2607.28632>

1 **CONF****TAPR: Enhancing LLM Performance with a Task-Aware Prompt Rewriter****ArXiv CS.AI** 50%

arXiv:2607.28657v1 Announce Type: new Abstract: Large Language Models (LLMs) often require carefully crafted prompts to unlock their full potential, which can be a barrier for non-expert users. This work addresses the challenge by introducing a Task-Aware Prompt Rewriter (TAPR),...
<https://arxiv.org/abs/2607.28657>

2 **CONF****Empowering Cross-Domain Sequential Recommendation with Hybrid Tokenization and Serial-Parallel Decoding****ArXiv CS.AI** **50%**

arXiv:2607.28659v1 Announce Type: new Abstract: Cross-domain sequential recommendation (CDSR) aims to model users' dynamic interest transitions and sequential patterns across multiple domains. Recently, generative recommendation (GR) has emerged. It first learns semantic...

<https://arxiv.org/abs/2607.28659>

3 **CONF****How Hard Does It Think? Analyzing Step-Aware Reasoning Energy in LLM Chain-of-Thought Trajectories****ArXiv CS.AI** **50%**

arXiv:2607.28674v1 Announce Type: new Abstract: Understanding how computational effort is allocated across individual chain-of-thought (CoT) reasoning steps remains an open challenge: existing interpretability methods rely on output-level signals or collapse processing depth...

<https://arxiv.org/abs/2607.28674>

4 **CONF****Reasoning in Real World Clinical Care: Why Large Language Models Are Not Yet Safe for Autonomous Clinical Decision Support****ArXiv CS.AI** **50%**

arXiv:2607.28677v1 Announce Type: new Abstract: LLM now pass medical licensing examinations and, in curated cases, can rival physicians at diagnostic reasoning. These developments have accelerated the use of LLMs for symptom assessment and clinical decision support in diagnostic...

<https://arxiv.org/abs/2607.28677>

5 **CONF****Multi-Agent Planning with Spatio-Temporal and Topological Constraints using STL-GO****ArXiv CS.AI** **50%**

arXiv:2607.28679v1 Announce Type: new Abstract: Multi-agent planning problems arise in a variety of engineering applications, such as multi-robot wildfire fighting and unmanned aerial inspection in factories. A particular challenge is the existence of spatio-temporal (i.e., when...

<https://arxiv.org/abs/2607.28679>

6 **CONF****Safety, or Just Capability? A Validity Audit of Agent-Safety Benchmarks****ArXiv CS.AI** **50%**

arXiv:2607.28685v1 Announce Type: new Abstract: Agent-safety benchmarks measure different behaviors, and their scores get quoted interchangeably as an agent's safety. We treat four of them (R-Judge, InjecAgent, AgentHarm, AgentDojo) as measurements to be validated, running each...

<https://arxiv.org/abs/2607.28685>

7 **CONF****EarlyDx: An Admission-Anchored Benchmark for Open-Ended Generation of Evidence-Supported ED-Encounter Diagnoses****ArXiv CS.AI** **50%**

arXiv:2607.28788v1 Announce Type: new Abstract: Clinical diagnosis at hospital admission must be made rapidly from limited, incomplete evidence. Existing diagnosis-prediction benchmarks are poorly suited to this setting: they restrict prediction to closed code sets, exclude...

<https://arxiv.org/abs/2607.28788>

8 CONF

HERO: History-Enriched Rollout Training for Long-Horizon Autoregressive Neural Operators

ArXiv CS.AI 50%

arXiv:2607.29135v1 Announce Type: cross Abstract: Neural operators provide fast surrogates for time-dependent partial differential equations (PDEs) by applying a learned evolution operator recursively to its own predictions, but this autoregressive rollout feeds every prediction...

<https://arxiv.org/abs/2607.29135>

9 CONF

Fragility of Value under Imperfect Alignment

ArXiv CS.AI 50%

arXiv:2607.28881v1 Announce Type: new Abstract: As more responsibility is placed upon AI systems, it becomes increasingly important to guarantee that these systems are aligned with humanity. A common fear in AI safety is that human value is fragile -- that is, optimizing too...

<https://arxiv.org/abs/2607.28881>

0 CONF

NeSyFS: A Neuro-symbolic Fast-Slow Thinking Framework for LLM Agent under Partial Observability

ArXiv CS.AI 50%

arXiv:2607.28942v1 Announce Type: new Abstract: Recently Large Language Models (LLMs) have been increasingly deployed as autonomous agents in applications such as self-reflection, retrieval-augmented generation, and scientific discovery. In these settings, agents must act based...

<https://arxiv.org/abs/2607.28942>

1 CONF

MerchantBench: Benchmarking LLM Agents for Long-Term Coherence in E-Commerce Operations

ArXiv CS.AI 50%

arXiv:2607.28956v1 Announce Type: new Abstract: Large language model agents are increasingly evaluated as autonomous tool users, yet most benchmarks focus on bounded tasks with immediate success criteria. Real-world deployments often require Long-Term Coherence, the capacity to...

<https://arxiv.org/abs/2607.28956>

2 CONF

DungeonBench: A Benchmark for Rules-Rich Tactical Reasoning in Dungeons & Dragons Combat

ArXiv CS.AI 50%

arXiv:2607.29577v1 Announce Type: new Abstract: Games and simulators make valuable benchmarks by turning decisions into measurable outcomes, but many current suites under-test rules-rich tactical reasoning: the ability to choose well when geometry, timing, resources, objectives,...

<https://arxiv.org/abs/2607.29577>

3 CONF

Evidence-Grounded Constraint Checking in Construction Documents

ArXiv CS.AI 50%

arXiv:2607.29058v1 Announce Type: new Abstract: Professional-document review is a constraint-checking problem in which decisions depend on relations among text, geometry, pages, and document revisions. We present an evidence-grounded pipeline that normalizes extracted facts,...

<https://arxiv.org/abs/2607.29058>

4 **CONF****CAGE: Certified Authorization under Typed-Return Uncertainty for Tool-Using Agents****ArXiv CS.AI 50%**

arXiv:2607.29190v1 Announce Type: new Abstract: Tool-using LLM agents act on typed tool returns, records pairing provenance and categorical fields with numerical values. Runtime permission gates generally authorize the observed return and action, leaving the decision unprotected...

<https://arxiv.org/abs/2607.29190>
5 **CONF****MirrorCraft: Paired Evaluation under Hidden Rule Changes in Minecraft****ArXiv CS.AI 50%**

arXiv:2607.29218v1 Announce Type: new Abstract: With the prosperity of the large language models (LLMs), it has become an interesting topic: how do LLM-based agents work in Minecraft? Unfortunately, most existing benchmarks evaluate them under fixed game mechanics. High...

<https://arxiv.org/abs/2607.29218>
6 **CONF****To Add Is Machine, To Delete Is Human: Measuring and Mitigating Deletion Avoidance in LLM Code Editing****ArXiv CS.AI 50%**

arXiv:2607.28887v1 Announce Type: cross Abstract: Large language models increasingly write and repair production code, yet evidence is mounting that their test-passing patches leave codebases harder to maintain. We identify one concrete source: deletion avoidance, the systematic...

<https://arxiv.org/abs/2607.28887>
7 **CONF****Beyond Component Testing: Validating Agentic AI Systems****ArXiv CS.AI 50%**

arXiv:2607.29405v1 Announce Type: new Abstract: Agentic AI systems act through multi-step trajectories that combine planning, tool use, memory, interaction, and adaptation. This behavior stretches validation practice beyond component testing and one-shot input--output...

<https://arxiv.org/abs/2607.29405>
8 **CONF****Compiled AI: Deterministic Code Generation for LLM-Based Workflow Automation****ArXiv CS.AI 50%**

arXiv:2604.05150v2 Announce Type: replace-cross Abstract: We study compiled AI, a paradigm in which large language models generate executable code artifacts during a compilation phase, after which workflows execute deterministically without further model invocation. This...

<https://arxiv.org/abs/2604.05150>
9 **CONF****A Unified Benchmark of Deep Learning Models for Multi-task 3D Brain Tumor Segmentation from Magnetic Resonance Imaging****ArXiv CS.AI 50%**

arXiv:2607.28858v1 Announce Type: cross Abstract: Automatic brain tumor segmentation from magnetic resonance imaging (MRI) has become a fundamental task in computer-assisted diagnosis, treatment planning, and disease monitoring. Although numerous deep learning architectures have...

<https://arxiv.org/abs/2607.28858>

- 0
CONF

Scaffolding Critical Engagement with GenAI: Transforming Ethnic Minority Preparatory Students' Collaborative Discourse in Prompt Engineering Tasks

ArXiv CS.AI
50%

arXiv:2607.28630v1 Announce Type: cross Abstract: Generative AI (GenAI) holds significant promise for advancing educational equity among ethnic minority students by broadening access to learning resources and mitigating linguistic barriers. However, these benefits are...

<https://arxiv.org/abs/2607.28630>
- 1
CONF

Topology-Aware Data Movement for Disaggregated GPU Inference

ArXiv CS.AI
50%

arXiv:2607.28633v1 Announce Type: cross Abstract: Disaggregated LLM inference creates a datacenter networking problem that no existing system solves correctly. When prefill and decode run on separate GPU pools, the KV cache must be transferred between them. For a 70B model this...

<https://arxiv.org/abs/2607.28633>
- 2
CONF

The Asymmetric Effects of Knowledge Distillation on Bias in Small Language Models

ArXiv CS.AI
50%

arXiv:2607.28639v1 Announce Type: cross Abstract: We show that knowledge distillation in small instruction-tuned language models has asymmetric effects on bias. On unambiguous tasks (BBQ-disambig), response-based distillation from a Gemma-2-9B teacher improves context-following:...

<https://arxiv.org/abs/2607.28639>
- 3
CONF

The Formalism Trap: Are LLM-as-a-Judge Evaluators Blinded by Consensus Mimicry under Social Load?

ArXiv CS.AI
50%

arXiv:2607.28641v1 Announce Type: cross Abstract: We introduce the \textit{Agentic Formalism Trap} and the Evaluative Dissonance Index ($\$D_E\$$), quantifying how LLM-as-a-Judge systems conflate structural proceduralism with semantic truth under adversarial load. Analyzing 22,500...

<https://arxiv.org/abs/2607.28641>
- 4
CONF

ConnectED: A Curriculum-Aligned AI System for Vietnamese Instructional Lesson Planning and Student Learning

ArXiv CS.AI
50%

arXiv:2607.28647v1 Announce Type: cross Abstract: This paper presents ConnectED, a human-centered AI system that supports the full instructional lifecycle in Vietnamese education by linking curriculum-aligned lesson design, interactive student learning, and feedback-driven...

<https://arxiv.org/abs/2607.28647>
- 5
CONF

COSI-Lab: Conference Living Lab for Modeling Multi-Perspective Multimodal Social Intention

ArXiv CS.AI
50%

arXiv:2607.28649v1 Announce Type: cross Abstract: COSI-Lab presents a multimodal, multi-sensor dataset of an interdisciplinary scientific workshop containing 32 academics at an international conference. It captures ecologically valid social interactions in a weakly scripted...

<https://arxiv.org/abs/2607.28649>

6 CONF

Unanticipated Effects of Generative AI on Expertise Pathways and Performance Perception in System Administration

ArXiv CS.AI 50%

arXiv:2607.28650v1 Announce Type: cross Abstract: While industry discourse often emphasizes immediate productivity gains and frames GenAI primarily as a tool for automation, the integration of GenAI into system administration may involve deeper shifts in professional practice...

<https://arxiv.org/abs/2607.28650>

7 CONF

HenTwin: A Multimodal Digital Twin Framework for Longitudinal Biological State Monitoring in Laying Hens

ArXiv CS.AI 50%

arXiv:2607.28652v1 Announce Type: cross Abstract: Early-life monitoring in laying hens remains constrained by fragmented single-modality sensing and the absence of formal system-level state representations. HenTwin, a multimodal digital twin framework implemented as a five-layer...

<https://arxiv.org/abs/2607.28652>

8 CONF

Evaluating Federated Pre-Training: On the Reliability of Downstream Fine-Tuning and Intrinsic Evaluation

ArXiv CS.AI 50%

arXiv:2607.28658v1 Announce Type: cross Abstract: Federated pre-training offers a way to train foundation models on private or distributed data without centralizing the underlying datasets. However, evaluating federated pre-training remains challenging because differences in...

<https://arxiv.org/abs/2607.28658>

9 CONF

Sensitivity Analysis of GRU, LSTM and Transformer Encoder in Classification of Automated Driving Systems

ArXiv CS.AI 50%

arXiv:2607.28665v1 Announce Type: cross Abstract: Automated driving systems (ADSs) are becoming ubiquitous. Future Software Defined Vehicles (SDVs) may be able to run multiple ADSs, both native and aftermarket such as Comma.ai's Openpilot. Monitoring systems to independently...

<https://arxiv.org/abs/2607.28665>

0 CONF

Guarantees on Dynamical System Distinguishability for LLM Token Generation

ArXiv CS.AI 50%

arXiv:2607.28667v1 Announce Type: cross Abstract: Recent work has shown that classifying large language models (LLMs)' responses can be distinguished by modeling token embeddings as trajectories of a black-box dynamical system (DS) and comparing prediction residuals of two DSs....

<https://arxiv.org/abs/2607.28667>

1 CONF

LAWFUL: Law-Aligned Witness for Faithful Use of Latents

ArXiv CS.AI 50%

arXiv:2607.28672v1 Announce Type: cross Abstract: When a neural network predicts a physical system accurately, has it learned the governing law as formal, structured knowledge, and if so, does the network's internal computation actually use that representation throughout the...

<https://arxiv.org/abs/2607.28672>

- 2

CONF

Metaphor-Induced Algorithmic Steering: Cross-Domain Procedural Transfer in LLM Code Generation

ArXiv CS.AI

50%

arXiv:2607.28683v1 Announce Type: cross Abstract: Large language models benefit from elements in natural language, such as metaphors and analogies in training data and inference input to achieve generalisability across different domains. However, these language elements may also...

<https://arxiv.org/abs/2607.28683>
- 3

CONF

Technological Advances in Detecting and Managing Cognitive Impairment in Older Adults: Trends, Challenges, and Future Directions

ArXiv CS.AI

50%

arXiv:2607.28687v1 Announce Type: cross Abstract: As populations age, cognitive decline from mild cognitive impairment (MCI) to dementia is a defining health challenge of the coming decades, yet routine assessment often misses its earliest signs. This article critically...

<https://arxiv.org/abs/2607.28687>
- 4

CONF

Predicting Steel Fatigue Life from Micrographs Using Physics-Informed Deep Learning

ArXiv CS.AI

50%

arXiv:2607.28695v1 Announce Type: cross Abstract: Here is the plain text version optimized for arXiv's submission form. Custom macros (like \CV and \SI) have been converted to standard text/math so they render correctly on the webpage: Evaluating the fatigue life of structural...

<https://arxiv.org/abs/2607.28695>
- 5

CONF

Wait for the Signal: Simple Frequency-Aware Flow-Matching

ArXiv CS.AI

50%

arXiv:2607.28760v1 Announce Type: cross Abstract: As image generation models scale to ever higher resolutions, global coherence, local detail, and texture fidelity become critical axes for generation quality. However, standard flow matching treats all spatial frequencies...

<https://arxiv.org/abs/2607.28760>
- 6

CONF

Benchmarks Are Not Monolithic: Sample-Level Auditing and Orchestration for LLM Evaluation

ArXiv CS.AI

50%

arXiv:2607.28801v1 Announce Type: cross Abstract: Benchmark datasets are central to evaluating Large Language Models (LLMs), yet they are typically conceived as monolithic tasks, obscuring substantial variation in the demands of individual samples. We introduce a dataset-centric...

<https://arxiv.org/abs/2607.28801>
- 7

CONF

Rolling With Resistance: Preference-Optimized LLM Counselors Can Trade Goal Persistence for Relational Attunement in Motivational Interviewing

ArXiv CS.AI

50%

arXiv:2607.28814v1 Announce Type: cross Abstract: In Motivational Interviewing (MI), a client's sustain talk (arguments for the status quo) calls for the counselor to roll with resistance, a move that can fail in two opposite ways: capitulation (abandoning the change agenda to...

<https://arxiv.org/abs/2607.28814>

8 CONF

Hypergradient-based Bilevel Reinforcement Learning with Improved Sample Complexity

ArXiv CS.AI 50%

arXiv:2607.28849v1 Announce Type: cross Abstract: Bilevel reinforcement learning (RL) is an important framework within the literature of RL that can be used to formalize various categories of problems, such as meta-learning, hierarchical task decomposition, and reinforcement...

<https://arxiv.org/abs/2607.28849>

9 CONF

TextCloak: Thwarting Unauthorized LLM Exploitation via RL-Driven Unlearnable Text

ArXiv CS.AI 50%

arXiv:2607.28862v1 Announce Type: cross Abstract: The rapid development of Large Language Models (LLMs) has led to significant advances across a wide range of language tasks, while simultaneously raising growing concerns about unauthorized data exploitation and privacy leakage....

<https://arxiv.org/abs/2607.28862>

0 CONF

Validation Evidence in LLM Repair Agents: How Much of What Passes Actually Tests the Bug?

ArXiv CS.AI 50%

arXiv:2607.28871v1 Announce Type: cross Abstract: When a repair agent runs a test and sees it pass, the result is treated as evidence about the reported defect. We measure how often that treatment is warranted. BSG-VA (buggy-state/candidate-state/gold-fix validation analysis)...

<https://arxiv.org/abs/2607.28871>

1 CONF

Human-LLM Collaborative Inductive Coding for Conceptualizing K-12 Educator AI Use

ArXiv CS.AI 50%

arXiv:2607.28889v1 Announce Type: cross Abstract: Qualitative researchers increasingly encounter interaction corpora whose scale exceeds what manual coding alone can address, and large language models (LLMs) are frequently proposed as analytic assistants. The open questions are...

<https://arxiv.org/abs/2607.28889>

2 CONF

Agreement Is Not Quality: Blind Expert Verification of Human and LLM Qualitative Coding When Human Consensus Is Not Ground Truth

ArXiv CS.AI 50%

arXiv:2607.28890v1 Announce Type: cross Abstract: Evaluations of LLM-assisted qualitative coding almost universally measure model performance as agreement with human coders, a practice that presumes human coding is the standard to approximate. This study provides empirical...

<https://arxiv.org/abs/2607.28890>

3 CONF

FairFund-Bench: Evaluating Distributive Bias in LLM Resource Allocation

ArXiv CS.AI 50%

arXiv:2607.28934v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly involved in the distribution of scarce resources, raising concerns about biased allocations based on characteristics like race and gender. Recent LLM audits have produced inconsistent...

<https://arxiv.org/abs/2607.28934>

4 **CONF****DiffAttack: Evasion Attacks Against Face Recognition via Latent Diffusion Models****ArXiv CS.AI** **50%**

arXiv:2607.28936v1 Announce Type: cross Abstract: Facial biometric identification relies on the distinctiveness of user attributes within a high-dimensional embedding space. However, the decision boundaries of deep face recognition (FR) systems are often sufficiently narrow that...

<https://arxiv.org/abs/2607.28936>

5 **CONF****Efficient LLM Adversarial Training via Low-Rank Defense and Circuit-Guided Surrogates****ArXiv CS.AI** **50%**

arXiv:2607.28959v1 Announce Type: cross Abstract: Adversarial training is one of the most effective defenses against adversarial attacks, yet the computational cost remains prohibitive at modern scales, especially for large language models (LLMs). While existing mitigation...

<https://arxiv.org/abs/2607.28959>

6 **CONF****A robust association between LLM use and scientific productivity: Assessing stopping-time selection****ArXiv CS.AI** **50%**

arXiv:2607.28968v1 Announce Type: cross Abstract: Renault, Bergeaud, and Bosquet (hereafter RBB) argue that dating LLM adoption as the first month in which an author's abstract is flagged induces a stopping-time selection that can produce a positive event-study path even when...

<https://arxiv.org/abs/2607.28968>

7 **CONF****RAID: Towards Robust AI-Generated Image Detection with Bit-Reversed Images****ArXiv CS.AI** **50%**

arXiv:2607.28974v1 Announce Type: cross Abstract: The rapid advancement of image generation models has made it increasingly difficult for people to distinguish AI-generated images from real ones. To prevent the potential risks associated with the misuse of fake images,...

<https://arxiv.org/abs/2607.28974>

8 **CONF****Adjudicated Captioning: Multi-Agent Alignment Scoring and Consensus-Distilled Beam Arbitration for Strict Zero-Shot Image Captioning****ArXiv CS.AI** **50%**

arXiv:2607.28986v1 Announce Type: cross Abstract: Zero-shot image captioning (ZIC) describes images without paired image-caption supervision during captioner training, relying on text-only corpora and frozen pretrained image-text scorers. Existing retrieval-augmented methods...

<https://arxiv.org/abs/2607.28986>

9 **CONF****TerraNova: A Foundation Model for the Anthropocene****ArXiv CS.AI** **50%**

arXiv:2607.29527v1 Announce Type: cross Abstract: A defining problem of the Anthropocene is to model the physical Earth and human societies as one coupled system, yet no learned representation spans their observational breadth. We argue the obstacle is geometric: the physical...

<https://arxiv.org/abs/2607.29527>

0 CONF

Improving scDiffusion with Sparsity-Biased Classifier-Free Guidance

ArXiv CS.AI 50%

arXiv:2607.29043v1 Announce Type: cross Abstract: Single-cell RNA sequencing (scRNA-seq) has become an essential tool in modern cellular biology, and generating accurate synthetic scRNA-seq data is becoming increasingly important. Although diffusion models have achieved...

<https://arxiv.org/abs/2607.29043>

1 CONF

Learning Lookahead Lemmas for Neural Network Verification

ArXiv CS.AI 50%

arXiv:2607.29051v1 Announce Type: cross Abstract: State-of-the-art neural network verifiers use the branch-and-bound procedure as their core solving mechanism. We introduce an inprocessing framework for neural network verification driven by the lookahead procedure. Under this...

<https://arxiv.org/abs/2607.29051>

2 CONF

SeekBrain: An Autonomous Multi-Agent System for Accelerating Neuroscience Discovery

ArXiv CS.AI 50%

arXiv:2607.29347v1 Announce Type: cross Abstract: Modern neuroscience relies on integrating multi-scale, multimodal datasets to uncover the neural principles underlying intelligence. However, analytical challenges posed by highly heterogeneous data and fragmented workflows...

<https://arxiv.org/abs/2607.29347>

3 CONF

Semantics of Subterfuge: Benchmarking Legal Deception Detection Against General-domain State-of-the-Art

ArXiv CS.AI 50%

arXiv:2607.29066v1 Announce Type: cross Abstract: Deception detection has critical implications for legal proceedings, law enforcement, and online security. Although human judgment is limited in accuracy and scalability, Natural Language Processing (NLP) offers a data-driven...

<https://arxiv.org/abs/2607.29066>

4 CONF

A Human-Centered Validation of the Explainability-Performance Coefficient

ArXiv CS.AI 50%

arXiv:2607.29614v1 Announce Type: cross Abstract: The rapid adoption of deep learning models in high-risk domains has intensified the need for trustworthy Explainable Artificial Intelligence (XAI). However, objectively evaluating explanation fidelity and aligning XAI metrics...

<https://arxiv.org/abs/2607.29614>

5 CONF

DoubleHelix: Structured Cross-Modal Fusion for Audio-Visual Speech Recognition with LLMs

ArXiv CS.AI 50%

arXiv:2607.29112v1 Announce Type: cross Abstract: Audio-visual speech recognition (AVSR) relies on effective fusion of audio and visual modalities, yet existing approaches treat cross-modal interaction as a single-step operation without structured iterative refinement. We...

<https://arxiv.org/abs/2607.29112>

6 CONF

Multi-Granularity Position Embedding of Graphs via Granular-Ball for Link Prediction

ArXiv CS.AI 50%

arXiv:2607.29115v1 Announce Type: cross Abstract: Link prediction aims to identify potential or future connections within a given graph structure. Position information is essential for link prediction, as it distinguishes homogeneous nodes through their relative relationships,...

<https://arxiv.org/abs/2607.29115>

7 CONF

Have I Seen You? Embedding Behavior Signals Synthetic Face Dataset Membership

ArXiv CS.AI 50%

arXiv:2607.29144v1 Announce Type: cross Abstract: Synthetic face datasets are increasingly used to reduce privacy exposure and data access constraints in biometric recognition. Yet the generators that produce these datasets are trained on real faces, so synthetic data may still...

<https://arxiv.org/abs/2607.29144>

8 CONF

Memory Provenance Laundering in LLM Agents: A Non-Amplification Firewall for Persistent Memory

ArXiv CS.AI 50%

arXiv:2607.29167v1 Announce Type: cross Abstract: Long-term memory lets large language model(LLM) agents reuse prior preferences and work flows, but it also turns untrusted observations into persistent action context. We identify memory provenance laundering: during LLM-based...

<https://arxiv.org/abs/2607.29167>

9 CONF

CLIFT: Turning Gemini Robotics On-Device into Humanoid Specialists via Non-Invasive Closed-Loop Iterative Fine-Tuning

ArXiv CS.AI 50%

arXiv:2607.29172v1 Announce Type: cross Abstract: While robot foundation models are growing increasingly capable, the strongest models are typically trained on proprietary data and remain closed-source, limiting downstream users' ability to adapt them to new tasks, embodiments,...

<https://arxiv.org/abs/2607.29172>

0 CONF

FBFM: A Training-Free Asynchronous Feedback Mechanism for Flow-Matching in World-Action Models Execution

ArXiv CS.AI 50%

arXiv:2607.29235v1 Announce Type: cross Abstract: Although world-action models (WAMs) enhance long-horizon robot control by predicting visual evolution before acting, long-horizon reliability demands repeated re-grounding in real observations--not recursive rollout. Existing...

<https://arxiv.org/abs/2607.29235>

1 CONF

Small Is Enough: Per-User Style Rewriting of AI-Edited Text via LoRA Adapters

ArXiv CS.AI 50%

arXiv:2607.29238v1 Announce Type: cross Abstract: InMyStyle is a privacy first, single user system that adapts small language models to rewrite AI-edited text towards an individual user's writing style without an instruction prompt at inference. Given a user's documents, it uses...

<https://arxiv.org/abs/2607.29238>

2 **CONF****TAVI-TEC: An AI-Based Tool for Procedural Planning of Transcatheter Aortic Valve Implantation**

ArXiv CS.AI 50%

arXiv:2607.29243v1 Announce Type: cross Abstract: Computed tomography angiography (CTA) is crucial for preprocedural TAVI planning, providing the anatomical information required for prosthesis sizing and vascular access assessment. As the volume of TAVI procedure increases,...

<https://arxiv.org/abs/2607.29243>

3 **CONF****The persuasive power of large language models does not depend on their perceived national origin**

ArXiv CS.AI 50%

arXiv:2607.29334v1 Announce Type: cross Abstract: Conversational AI developed by geopolitical rivals reaches citizens worldwide, raising concerns that it could sway public opinion or be rejected as foreign propaganda, with consequences for democratic discourse and information...

<https://arxiv.org/abs/2607.29334>

4 **CONF****DualDiT: A Conditional Dual-Output Diffusion Transformer for Joint OCT Image and Segmentation Mask Generation**

ArXiv CS.AI 50%

arXiv:2607.29337v1 Announce Type: cross Abstract: Background and Objective: Generating realistic medical images with anatomically accurate segmentation masks helps address the shortage of annotated data in medical imaging, particularly in optical coherence tomography (OCT) of...

<https://arxiv.org/abs/2607.29337>

5 **CONF****AgenticRepair: Multi-Faceted Program Context Engineering for Agentic Vulnerability Repair**

ArXiv CS.AI 50%

arXiv:2607.29422v1 Announce Type: cross Abstract: Automated vulnerability repair aims to reduce the time and effort required to patch security flaws from a vulnerability triage report. Recent agentic AI approaches have shown promising results in automated program repair...

<https://arxiv.org/abs/2607.29422>

6 **CONF****QR-Structured Thermal Triggers for Targeted Semantic Attacks on Infrared Vision-Language Models**

ArXiv CS.AI 50%

arXiv:2607.29445v1 Announce Type: cross Abstract: Infrared vision-language models (IR-VLMs) extend thermal perception to open-vocabulary classification, image captioning, and visual question answering. However, their robustness to structured thermal perturbations and the...

<https://arxiv.org/abs/2607.29445>

7 **CONF****From Code Review to Code Critique: Intent, Drift, and Spotlight for AI-Generated Diffs at Scale**

ArXiv CS.AI 50%

arXiv:2607.29516v1 Announce Type: cross Abstract: AI coding agents are generating code at volumes that exceed the capacity of traditional peer review. At the same time, existing AI code review tools over-index on low-value suggestions such as style and best practices while...

<https://arxiv.org/abs/2607.29516>

8 CONF

Towards the Holographic Characteristic of LLMs for Efficient Short-text Generation

ArXiv CS.AI 50%

arXiv:2601.22546v2 Announce Type: replace-cross Abstract: The recent advancements in Large Language Models (LLMs) have attracted interest in exploring their in-context learning abilities and chain-of-thought capabilities. However, there are few studies investigating the specific...

<https://arxiv.org/abs/2601.22546>

9 CONF

The Theoretical Foundation of Socratic Tests: Dynamic, Multimodal, Conversational Examinations

ArXiv CS.AI 50%

arXiv:2607.29624v1 Announce Type: cross Abstract: Traditional static assessments rely on a subtractive, deficit-based grading model that often penalizes ambition and obscures diagnostic feedback. Conversely, traditional face-to-face oral examinations introduce severe...

<https://arxiv.org/abs/2607.29624>

0 CONF

Combining Large Language Models and Symbolic Reasoning for Multi-Robot Temporal Planning through Explainable Knowledge Bases

ArXiv CS.AI 50%

arXiv:2502.19135v2 Announce Type: replace Abstract: We present PLANTOR, a framework for generating and executing multi-robot task plans from natural-language task descriptions through LLM-assisted knowledge-base construction. The approach uses large language models to synthesize...

<https://arxiv.org/abs/2502.19135>

1 CONF

Deconstructing Off-Policy Ratios: Entropy-Scaled Trust Regions for Asynchronous Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.22186v2 Announce Type: replace Abstract: Asynchronous reinforcement learning (RL) accelerates large language model (LLM) post-training by overlapping rollout generation with policy optimization, but the resulting stale, off-policy data can destabilize optimization and...

<https://arxiv.org/abs/2607.22186>

2 CONF

What Makes a Sale? Simulating End-to-End Seller--Buyer Retail Dynamics with LLM Agents

ArXiv CS.AI 50%

arXiv:2604.04468v2 Announce Type: replace Abstract: Evaluating retail strategies before deployment is difficult, as outcomes are determined across multiple stages, from seller-side persuasion through buyer-seller interaction to purchase decisions. However, existing retail...

<https://arxiv.org/abs/2604.04468>

3 CONF

Dual-Dimensional Consistency: Balancing Budget and Quality in Adaptive Inference-Time Scaling

ArXiv CS.AI 50%

arXiv:2605.15100v2 Announce Type: replace Abstract: Large Language Models (LLMs) have demonstrated remarkable abilities in reasoning. However, maximizing their potential through inference-time scaling faces challenges in trade-off between sampling budget and reasoning quality....

<https://arxiv.org/abs/2605.15100>

4 **CONF****PAIR: Prefix-Aware Internal Reward Model for Multi-Turn Agent Optimization****ArXiv CS.AI** **50%**

arXiv:2605.17877v2 Announce Type: replace Abstract: A significant hurdle for current LLMs is the execution of complex, multi-stage tasks. Group Relative Policy Optimization (GRPO) has been emerging as a leading choice, but its reliance on sparse outcome rewards severely limits...

<https://arxiv.org/abs/2605.17877>

5 **CONF****The Self-Correction Illusion: Role Relabeling Gates Explicit Error Flagging in Large Language Models****ArXiv CS.AI** **50%**

arXiv:2606.05976v2 Announce Type: replace Abstract: Recent works show that LLM agents struggle to correct errors in their own reasoning traces, despite their ability to correct errors from external sources. We ask whether this reflects a capability deficit or an artifact of the...

<https://arxiv.org/abs/2606.05976>

6 **CONF****A Multi-Agent System for Motor Design Optimization via an FEA-AI Hybrid Approach****ArXiv CS.AI** **50%**

arXiv:2606.09037v2 Announce Type: replace Abstract: This study presents a large language model (LLM)-based multi-agent framework for interior permanent magnet synchronous motor (IPMSM) design optimization that mitigates limitations of conventional workflows: expertise-dependent...

<https://arxiv.org/abs/2606.09037>

7 **CONF****Solution Space Path Planning: A Real-Time Human-Centered Path Planning Algorithm for En-Route Air Traffic Control****ArXiv CS.AI** **50%**

arXiv:2607.00064v2 Announce Type: replace Abstract: As technology advances, various algorithms have been proposed for air traffic management, yet their operational adoption in tactical control remains limited. This gap motivates a human-centered design emphasizing algorithmic...

<https://arxiv.org/abs/2607.00064>

8 **CONF****Dual-Force: Enhanced Offline Diversity Maximization under Imitation Constraints****ArXiv CS.AI** **50%**

arXiv:2501.04426v2 Announce Type: replace-cross Abstract: Offline diversity maximization under imitation constraints can transform demonstration data into a set of distinct behavioral policies, improving robustness to distribution shift without additional environment...

<https://arxiv.org/abs/2501.04426>

9 **CONF****Pay for The Second-Best Service: A Game-Theoretic Approach Against Dishonest LLM Providers****ArXiv CS.AI** **50%**

arXiv:2511.00847v5 Announce Type: replace-cross Abstract: The widespread adoption of Large Language Models (LLMs) through Application Programming Interfaces (APIs) induces a critical vulnerability: the potential for dishonest manipulation by service providers. This manipulation...

<https://arxiv.org/abs/2511.00847>

00 CONF

Knowledge Restoration-driven Prompt Optimization: Unlocking LLM Potential for Open-Domain Relational Triplet Extraction

ArXiv CS.AI 50%

arXiv:2601.15037v2 Announce Type: replace-cross Abstract: Open-domain Relational Triplet Extraction (ORTE) aims to mine structured knowledge without predefined relation schemas. Large Language Models (LLMs) have advanced ORTE toward a prompt-driven paradigm through powerful...

<https://arxiv.org/abs/2601.15037>

01 CONF

Alvilization v0: Toward Large-Scale Artificial Social Simulation with a Unified Agent Architecture and Adaptive Agent Profiles

ArXiv CS.AI 50%

arXiv:2602.10429v2 Announce Type: replace-cross Abstract: Alvilization v0 is a publicly deployed large-scale artificial society that couples a resource-constrained sandbox with a unified LLM-agent architecture, aiming to sustain long-horizon autonomy while remaining executable...

<https://arxiv.org/abs/2602.10429>

02 CONF

ActionParty: Multi-Subject Action Binding in Generative Video Games

ArXiv CS.AI 50%

arXiv:2604.02330v2 Announce Type: replace-cross Abstract: Recent advances in video diffusion have enabled the development of "world models" capable of simulating interactive environments. However, these models are largely restricted to single-agent settings, failing to control...

<https://arxiv.org/abs/2604.02330>

03 CONF

Leveraging Image Generators to Address Data Scarcity: The Gen4Regen Dataset for Forest Regeneration Mapping

ArXiv CS.AI 50%

arXiv:2605.05627v2 Announce Type: replace-cross Abstract: Sustainable forest management relies on precise species composition mapping, yet traditional ground surveys are labour-intensive and geographically constrained. While Uncrewed Aerial Vehicles (UAVs) offer scalable data...

<https://arxiv.org/abs/2605.05627>

04 CONF

The Metanym Game: A Self-Contained, Self-Consistent LLM Peer-Community Benchmark for Structural Intelligence

ArXiv CS.AI 50%

arXiv:2606.21008v2 Announce Type: replace-cross Abstract: The metanym game is a competitive word game for LLMs that measures structural intelligence against established cognitive-science constructs. No content is given in advance; the contestants create all of it -- a new kind...

<https://arxiv.org/abs/2606.21008>

05 CONF

Copy Less, Ground More: Overcoming Repetitive Copying in Long-Context Reasoning via Evidence-Aware Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2607.19345v2 Announce Type: replace-cross Abstract: Large language models that generate step-by-step reasoning traces have achieved strong performance on complex tasks, and extending them to long-context settings has emerged as an important frontier. However, we identify a...

<https://arxiv.org/abs/2607.19345>

06 CONF

LabEvolver: Training-Free Experience Evolution for Safe and Grounded Wet-Lab Agents

ArXiv CS.AI 50%

arXiv:2607.27690v2 Announce Type: replace-cross Abstract: We introduce LabEvolver, a training-free framework that equips safe and grounded wet-lab agents with episodic memory from execution experience. LabEvolver couples a state-grounded inner trial loop for adaptive perception,...

<https://arxiv.org/abs/2607.27690>