



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
HackerNews	6	Community
TechCrunch AI	5	Media
MIT Tech Review	2	Media
The Verge AI	2	Media
ArXiv CS.AI	1	Media

CONFIRMED SIGNALS — 16 stories

CONF

Show HN: Nightcrawler - A local AI pentesting agent running on a smartphone

HackerNews 60% HN 109

Title: Show HN: FutureSearch, AI forecasting you can verify AI forecasting is now approximately superhuman. Today, FutureSearch is exiting our long public beta and launching. We started FutureSearch in August 2023. (We're the original AI forecasting company, at least in a...

<https://github.com/garagehq/nightcrawler/>

CONF

Show HN: An AI wrote and illustrated a book of fables about threads and packets

HackerNews 50% HN 6

No summary — see link for full article.

<https://booyaka101.github.io/thedailyfable/day07/>

CONF

Show HN: TokenMaxxer - track every AI token you spend across your coding tools

HackerNews 50% HN 5

I use Claude Code, Codex and Cursor (and sometimes Antigravity) basically every day, and could never tell how much I was actually consuming across all of them. So I built TokenMaxxer. A small CLI reads the files these tools already write locally and puts it all in one dashboard,...

<https://tokenmaxxer.xyz>

CONF

Show HN: Find the perfect name for your startup, app or domain

HackerNews 50% HN 3

No summary — see link for full article.

<https://www.howdowesay.com>

CONF

Show HN: ModelFuzz - Open-source runtime guardrails for AI agents

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/higagan/modelfuzz>

CONF

Show HN: Every plaintext HTTP request on your Linux box captured via eBPF

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/yeet-src/httpwatch>

CONF

After killer quarter, Palantir CEO Alex Karp calls AI industry 'Marxist'

TechCrunch AI 50%

After a quarter that delivered \$1 billion in profit, Palantir CEO Alex Karp on Monday once again warned that AI frontier labs are too untrustworthy for enterprises.

<https://techcrunch.com/2026/08/03/after-killer-quarter-palantir-ceo-alex-karp-calls-ai-...>

CONF

Design Arena creators raise \$7.9 million to bring taste to AI models

TechCrunch AI 50%

Design Arena is used by 5.3 million people around the world, providing critical human evaluations to frontier labs.

<https://techcrunch.com/2026/08/03/designarena-creators-raise-7-9-million-to-bring-taste...>

CONF

Influencers draw backlash for attending OpenAI's first luxury trip

TechCrunch AI 50%

OpenAI's first-ever influencer brand trip is sparking online backlash as tensions over the use of AI continue.

<https://techcrunch.com/2026/08/03/influencers-draw-backlash-for-attending-openais-first...>

CONF

Congress' favorite AI tool? ChatGPT

TechCrunch AI 50%

House spending records show OpenAI's ChatGPT dominates paid AI use on Capitol Hill, with congressional offices relying on the chatbot to draft memos, summarize legislation, and assist constituent communications.

<https://techcrunch.com/2026/08/03/congresss-favorite-ai-tool-chatgpt/>

CONF

A Marc Benioff-backed startup thinks AI can solve the AI deployment problem

TechCrunch AI 50%

June emerged from stealth today with a \$20 million pre-seed round to make AI adoption simpler.

<https://techcrunch.com/2026/08/03/a-marc-benioff-backed-startup-thinks-ai-can-solve-the...>

CONF

Trump's AI protectionism has come for robotics

MIT Tech Review 50%

This story originally appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. Humanoid robots usually elicit more cringe than awe: They stumble, kick children, and despite advances are still worse at using their hands...

<https://www.technologyreview.com/2026/08/03/1141056/trumps-ai-protectionism-has-come-fo...>

CONF

The Download: reward hacking explained, and suspected Iranian cyberattacks

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Here's why AI agents lie and cheat to reach their goals When two OpenAI models hacked into Hugging Face last month, they weren't trying to...

<https://www.technologyreview.com/2026/08/03/1141039/the-download-reward-hacking-water-c...>

CONF

Europe's AI labeling and transparency rules are now in effect

The Verge AI 50%

The European Union has ushered in some additional rules that aim to make it easier for people to identify chatbots and AI deepfakes online. The new transparency obligations under the bloc's landmark AI Act came into effect on August 2nd, requiring companies to disclose when...

<https://www.theverge.com/ai-artificial-intelligence/974571/eu-ai-act-transparency-label...>

CONF

China's Alibaba takes another swipe at America's AI supremacy

The Verge AI 50%

Chinese tech giant Alibaba released what it says is its largest and "most capable AI model to date," claiming performance rivaling the best systems from US frontier labs Anthropic and OpenAI, as well as domestic rivals like Moonshot AI's Kimi K3. Alibaba said it was making the...

<https://www.theverge.com/ai-artificial-intelligence/974342/alibaba-qwen-max-open-weight-ai>

6

CONF

On the Generalization of Steering Vectors for Chain-of-Thought Faithfulness

ArXiv CS.AI 50%

arXiv:2607.29062v1 Announce Type: new Abstract: Model capabilities have improved in large part due to scaling chain of thought. This has been a promising development for AI safety--where models verbalize their reasoning, it is possible to monitor it. However, in some cases,...

<https://arxiv.org/abs/2607.29062>