



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	170	Media
HackerNews	8	Community
TechCrunch AI	7	Media
The Verge AI	5	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 5 stories

VERI

Risky Business: Measuring The Faithfulness-Safety Tension

ArXiv CS.AI +2 sources

100%

arXiv:2608.03745v1 Announce Type: new Abstract: Chain-of-Thought (CoT) reasoning offers a promising window into model monitoring. However, monitoring relies on faithfulness, i.e., the model output strictly derives from its reasoning trace. We identify an alignment tension where...

<https://arxiv.org/abs/2608.03745>

VERI

FLARE: Diffusion for Hybrid Language Model

ArXiv CS.AI +1 source

80%

arXiv:2606.01774v2 Announce Type: replace-cross Abstract: Autoregressive (AR) large language models (LLMs) have achieved broad practical success, but sequential decoding remains a key bottleneck for low-latency deployment. Recent efficient-inference work has progressed along two...

<https://arxiv.org/abs/2606.01774>

VERI

AI Alignment and Fiduciary Obligation

ArXiv CS.AI +1 source

80%

arXiv:2608.02660v1 Announce Type: cross Abstract: Advanced AI assistants engage users in extended interactions across a widening range of roles, including advice, decision support, collaboration, learning, emotional support, and companionship among others. Current alignment...

<https://arxiv.org/abs/2608.02660>

VERI

Here's why AI agents lie and cheat to reach their goals

MIT Tech Review +1 source

80%

MIT Technology Review Explains: Let our writers untangle the complex, messy world of technology to help you understand what's coming next. You can read more from the series here. When two OpenAI models hacked into the website Hugging Face in July, they weren't trying to make...

<https://www.technologyreview.com/2026/08/03/1141009/heres-why-ai-agents-lie-and-cheat-t...>

VERI

'Not healthy' LLM use is more common than you think

The Verge AI +1 source

80%

Hank Green, a popular YouTuber and science communicator, said he is stepping back from production amid intense criticism over his use of AI. Green described his AI usage as "not healthy," but stressed that he used it for finding research sources and not to write scripts. Much of...

<https://www.theverge.com/ai-artificial-intelligence/975180/llm-ai-chatbot-use-not-healthy>

CONFIRMED SIGNALS — 180 stories

CONF

Show HN: Fine-tune an 8B model on a 4 GB laptop GPU

HackerNews 60% HN 125

Creating 3D is hard. LLMs seem to be getting better at tool use and spatial understanding. While MCPs have proved to be a good way to use these tools- the current methods have these challenges: - Access to scene graph and core C modules of Blender - Lack of parallelism, only way...

<https://github.com/MakazhanAlpamys/Soup>

CONF

Show HN: SIEMatic, a fair-sourced observability and security platform

HackerNews 50% HN 8

I wrote this article to breakdown some of the recent misalignment cases where models have broken out of containment in conjunction with our work at Vulnetic and how we prevent these mishaps in production.

<https://github.com/mcindi/siematic>

CONF

Show HN: I gave my video editor an MCP server so Claude can edit videos for me

HackerNews 50% HN 5

No summary — see link for full article.

<https://www.shorz.ai>

CONF

Show HN: Simple self-hosted LLM assistant with user-steered compounding context

HackerNews 50% HN 5

I built a personal LLM assistant on Cloudflare Workers + Durable Objects. You specify a category and topic when starting a new conversation, so the backend maintains a summary for each category/topic - building up as more conversations happen under the same one. There is no...

<https://github.com/kol3x/pawmc>

0 CONF

Show HN: My tool scanned 256 AI-built apps and most had exposed credentials

HackerNews 50% HN 4

All of a sudden everybody wants to build with AI. People in hell want ice water. Built something with AI and don't know if it's secure? Try necktochoke.

<https://necktochoke.com/>

1 CONF

Show HN: Arabic-Hebrew artifact in GPT-5.4, 12,160 frozen trials

HackerNews 50% HN 3

No summary — see link for full article.

<https://zenodo.org/records/21799525>

2 CONF

Show HN: cctap - see and reach the Claude Code session that needs you

HackerNews 50% HN 3

I built this clod code or codec skill because I just love using these AI agents to write code. It's just so convenient to do. But when it comes to working in a corporate environment, I have to ultimately follow a Jira ticket or a Asana ticket. Now, when I'm following those...

<https://github.com/chipmates/cctap>

3 CONF

Show HN: AI product evaluation methodology at huby

HackerNews 50% HN 2

There are a number of AI benchmarks but I feel they aren't practical when it comes to evaluating AI products. Several months back we started working on designing a methodology to independently evaluate AI products. After conducting a good bit of research, we tested the...

<https://news.ycombinator.com/item?id=49176006>

4 CONF

Nvidia doesn't mess around: A week after open AI industry group formed, it's already showing progress

TechCrunch AI 50%

The week-old Open Secure AI Alliance, spearheaded by Nvidia and grown to over 120 companies, already has proposals out for defending against AI agents.

<https://techcrunch.com/2026/08/04/nvidia-doesnt-mess-around-a-week-after-open-ai-indust...>

5 CONF

Spotify expands AI remix and covers project with Merlin partnership

TechCrunch AI 50%

Spotify says Merlin, which represents more than 30,000 independent labels and distributors, has joined Universal Music Group in backing its upcoming AI-powered remix and covers product. The paid tool will let fans create AI-generated covers and remixes of participating artists'...

<https://techcrunch.com/2026/08/04/spotify-adds-merlin-to-its-ai-music-remix-and-covers-...>

6 CONF

Elon Musk spends half his time talking robots and AI on Tesla earnings calls

TechCrunch AI 50%

An analysis of the last seven years of Tesla earnings calls shows just how little attention Musk pays to Tesla's car business.

<https://techcrunch.com/2026/08/04/elon-musk-spends-half-his-time-talking-robots-and-ai-...>

7 CONF

Apple says more ex-employees may have taken confidential data to OpenAI

TechCrunch AI 50%

Apple says its trade secrets investigation into OpenAI has widened. In a new court filing, Apple claims additional former staff may have retained or accessed confidential information.

<https://techcrunch.com/2026/08/04/apple-says-more-ex-employees-may-have-taken-confident...>

8 CONF

SpaceX made more revenue as an AI company than a space company

The Verge AI 50%

SpaceX's AI revenue grew more than three times to \$2.6 billion from the year before, mostly because of deals that the company made to provide compute to other AI companies, according to SpaceX's quarterly earnings. The AI division, which the company said in its documents to go...

<https://www.theverge.com/science/975335/spacex-made-more-money-as-a-neocloud>

9 CONF

How an OpenAI influencer trip backfired

The Verge AI 50%

The brand trip is a right of passage for influencers. It's a mark of legitimacy that a sponsor wants to invite them on an all-expenses-paid vacation, often with luxurious freebies and activities. Trips can also spur hard feelings from uninvited influencers, trigger criticism...

<https://www.theverge.com/tech/975173/openai-influencers-brand-trip-ai-backlash-marketing>

0 CONF

Can Reddit fend off a new wave of AI SEO spam?

The Verge AI 50%

Earlier this year, a Reddit user had asked members of a skincare-focused subreddit if anyone had tried a specific hypochlorous acid spray, a product often used for acne. There were dozens of responses; one from a user named Primary-Taro4254 seemed innocuous enough, at least at...

<https://www.theverge.com/ai-artificial-intelligence/973098/reddit-ai-search-seo-marketi...>

1 CONF

Beyond the Hivemind: Escaping LLM Homogeneity via Meta-Persona Anchoring and Sequential Temperature Scaling

ArXiv CS.AI 50%

arXiv:2608.02618v1 Announce Type: new Abstract: Recent studies have identified an ``Artificial Hivemind" effect in Large Language Models (LLMs) causing models to converge on a narrow, homogenized consensus even for open questions. This semantic collapse limits the diversity of...

<https://arxiv.org/abs/2608.02618>

2 **CONF****LeanMem: Simple and Efficient Long-Term Memory for LLM Agents****ArXiv CS.AI 50%**

arXiv:2608.03463v1 Announce Type: new Abstract: Long-term memory is essential for LLM-based agents to sustain interactions and reliably leverage distant history. However, existing memory systems typically process heterogeneous dialogue content through a uniform summarization and...

<https://arxiv.org/abs/2608.03463>

3 **CONF****Explainable AI for the EU Right to Explanation: A Systematic Review of the Law-XAI Translation Gap****ArXiv CS.AI 50%**

arXiv:2608.02699v1 Announce Type: new Abstract: When algorithms make or influence consequential decisions---about loan eligibility, hiring, or healthcare---EU law grants affected individuals a Right to Explanation. Yet whether (and how) Explainable AI (XAI) can satisfy this...

<https://arxiv.org/abs/2608.02699>

4 **CONF****A Unified Framework for Human AI Collaboration in Security Operations Centers with Trusted Autonomy****ArXiv CS.AI 50%**

arXiv:2505.23397v3 Announce Type: replace Abstract: This article presents a structured framework for Human-AI collaboration in Security Operations Centers (SOCs), integrating AI autonomy, trust calibration, and Human-in-the-loop decision making. Existing frameworks in SOCs often...

<https://arxiv.org/abs/2505.23397>

5 **CONF****Towards a new paradigm of scientific discovery with socialized artificial intelligence****ArXiv CS.AI 50%**

arXiv:2608.02775v1 Announce Type: new Abstract: Scientific discovery has advanced through successive transformations in the organization of knowledge. Observation and experimentation established the empirical foundations of science. Theory made it possible to derive general...

<https://arxiv.org/abs/2608.02775>

6 **CONF****Reversing Arrows in Large Language Models****ArXiv CS.AI 50%**

arXiv:2608.03512v1 Announce Type: new Abstract: Large language models (LLMs) have achieved strong performance on text-to-knowledge graph generation and related tasks. Nevertheless, it is still unclear whether they accurately model the direction-dependent semantics of inverse...

<https://arxiv.org/abs/2608.03512>

7 **CONF****Hypercubes, Hyperplanes, and Constraint-Induced Complexity Collapse in Atomic Concept Learning****ArXiv CS.AI 50%**

arXiv:2608.02930v1 Announce Type: new Abstract: We revisit higher-arity atomic concept learning through the geometry of hypercubes and hyperplanes of ground instances. Our starting point is the observation that the ambient r -dimensional hypercube of ground atoms is not...

<https://arxiv.org/abs/2608.02930>

8 **CONF****When Compression Scores Cannot Decide: Information Boundaries for Group-Robust LLM Pruning****ArXiv CS.AI** **50%**

arXiv:2608.02940v1 Announce Type: new Abstract: A reproducible compression statistic can still select the wrong candidate. A dense pruning score with 0.906 split-half reliability predicted a 16.1% gain. Its selected endpoint was 6.0% and 7.7% worse than two controls. We model...

<https://arxiv.org/abs/2608.02940>

9 **CONF****LoCA: Forward-Only LLM Tuning after One-Shot Calibration with Local Credit Assignment****ArXiv CS.AI** **50%**

arXiv:2608.03020v1 Announce Type: new Abstract: Parameter-efficient post-training reduces the number of trainable parameters, but still requires repeated end-to-end backpropagation through the frozen backbone. Every adaptation step therefore needs backward-capable hardware and...

<https://arxiv.org/abs/2608.03020>

0 **CONF****Scaling an Autoregressive Transformer for Single-Cell Generation****ArXiv CS.AI** **50%**

arXiv:2608.02961v1 Announce Type: cross Abstract: We study a self-supervised generation task for single-cell gene expression vectors: given a set of vectors from a cell type, we aim to generate additional gene expression vectors of that cell type. For this task we characterize...

<https://arxiv.org/abs/2608.02961>

1 **CONF****Does Forgetting Transfer Across Modalities? A Real-World Benchmark for Cross-Modal Knowledge Unlearning Evaluation****ArXiv CS.AI** **50%**

arXiv:2608.03791v1 Announce Type: new Abstract: Vision-Language Models (VLMs), like Large Language Models (LLMs), may memorize sensitive, copyrighted, or harmful knowledge from their pretraining corpora. Removing such knowledge is essential for building trustworthy AI systems....

<https://arxiv.org/abs/2608.03791>

2 **CONF****AI Agent Economics: Can Autonomous Economic Behavior Emerge among AI Agents under Minimal External Conditions?****ArXiv CS.AI** **50%**

arXiv:2608.03076v1 Announce Type: new Abstract: Multi-agent studies commonly place AI agents in predefined games, markets, or roles, making it difficult to distinguish endogenous economic organization from behavior inherited from the scenario. We ask whether economic relations...

<https://arxiv.org/abs/2608.03076>

3 **CONF****Beyond Average Performance: Dynamic Instance Clustering and Specialized Algorithm Design in LLM-Assisted Evolutionary Search****ArXiv CS.AI** **50%**

arXiv:2608.03129v1 Announce Type: new Abstract: Large Language Model-assisted Evolutionary Search (LES) has emerged as a powerful paradigm for automated algorithm design. However, existing LES methods primarily optimize for average performance, inherently directing search effort...

<https://arxiv.org/abs/2608.03129>

4 **CONF****A Blind Spot in Alignment: Quantifying Biosecurity Risks in Large Language Models****ArXiv CS.AI 50%**

arXiv:2608.02684v1 Announce Type: cross Abstract: Large Language Models (LLMs) are accelerating biological research, yet this same capability poses a critical biosecurity threat: models that assist in protein engineering can equally be prompted to generate predicted toxin-like...

<https://arxiv.org/abs/2608.02684>

5 **CONF****Spatial proteomics guided by H&E-based AI reveals recurrence-risk niches in triple-negative breast cancer****ArXiv CS.AI 50%**

arXiv:2608.03145v1 Announce Type: new Abstract: Deep learning models can predict cancer recurrence from H&E stained slides, but the localized molecular states underlying these predictions remain largely obscured. Here, we developed an outcome informed spatial pathology framework...

<https://arxiv.org/abs/2608.03145>

6 **CONF****ADMITBench: A Safety-Governed Reference Framework for Evaluating the Admissibility of Industrial LLM Advisories****ArXiv CS.AI 50%**

arXiv:2608.03866v1 Announce Type: new Abstract: This white paper presents ADMITBench, a reference framework for evaluating industrial LLM advisories at the level of the proposed action. The framework implements a versioned, safety-governed evaluation contract that checks whether...

<https://arxiv.org/abs/2608.03866>

7 **CONF****Evidence-Grounded Multimodal Knowledge Graph Construction for Multi-Lecture Educational Reasoning****ArXiv CS.AI 50%**

arXiv:2608.03161v1 Announce Type: new Abstract: Lecture videos distribute knowledge across speech, slide text, diagrams, equations, and presentation order, which transcript-only retrieval does not fully preserve. This paper presents an evidence-grounded multimodal pipeline that...

<https://arxiv.org/abs/2608.03161>

8 **CONF****Diversity is Not Ambiguity: Toward Accurate and Efficient Ambiguity Detection for Open-Domain QA****ArXiv CS.AI 50%**

arXiv:2608.03177v1 Announce Type: new Abstract: How can question answering (QA) systems determine whether a query is ambiguous? Ambiguity detection is essential in open-domain QA, as misclassification leads to answering the wrong interpretation or unnecessary clarification....

<https://arxiv.org/abs/2608.03177>

9 **CONF****Reachability Is Not Realization: Tracing the Sources of LLM Benchmark Gains****ArXiv CS.AI 50%**

arXiv:2608.03219v1 Announce Type: new Abstract: Benchmark gains are often treated as evidence of greater LLM capability. Yet the same gain can reflect different changes in model behavior. A model may reach new answers, or produce answers that were already within reach. Aggregate...

<https://arxiv.org/abs/2608.03219>

- 0
CONF

AgentPanel: Toward a New Paradigm for Human--AI Collaboration in Exploring Scientific Questions

ArXiv CS.AI 50%

arXiv:2608.03283v1 Announce Type: new Abstract: Identifying promising scientific ideas remains an important challenge in research practice. Researchers commonly rely on small-group discussions or one-to-one interactions with a single large language model, yet these approaches...

<https://arxiv.org/abs/2608.03283>
- 1
CONF

Distractor-Aware Truncation: Disentangling Context-Length Effects from Signal Loss in Long-Context LLM Benchmarks

ArXiv CS.AI 50%

arXiv:2608.03297v1 Announce Type: new Abstract: A standard claim in the literature on retrieval-augmented and memory-augmented language models is that shorter context is better when the relevant information is preserved. We test this claim by running every sample of two...

<https://arxiv.org/abs/2608.03297>
- 2
CONF

Screenshots or Tools? Eliciting Tool Use and Managing Multimodal Context in Hybrid GUI-MCP Computer-Use Agents

ArXiv CS.AI 50%

arXiv:2608.03327v1 Announce Type: new Abstract: Hybrid computer-use agents can act through screenshots or call text tools. We find that having a tool available does not settle which way the effect goes. Under one identical GUI-MCP harness on the OSWorld-MCP benchmark (309...

<https://arxiv.org/abs/2608.03327>
- 3
CONF

Enactive Artificial Intelligence: A Decision-Centric Architecture for Complex Systems

ArXiv CS.AI 50%

arXiv:2608.03413v1 Announce Type: new Abstract: As artificial intelligence (AI) continues to evolve and mature, recent AI practices have moved beyond large language models (LLMs) and text or image generation tasks, increasingly integrating tools, agents, and harnesses to solve...

<https://arxiv.org/abs/2608.03413>
- 4
CONF

When Policies Change Probabilities: Modular Decision-Making for LLM Code Review

ArXiv CS.AI 50%

arXiv:2608.02677v1 Announce Type: cross Abstract: LLM code reviewers often estimate patch risk and make approval decisions in one prompt. A probability should depend on evidence; costs should determine the action taken from it. We test whether four deployed reviewer interfaces...

<https://arxiv.org/abs/2608.02677>
- 5
CONF

BODHI: Do LLMs Branch Out and Discover Heterogeneous Inferences?

ArXiv CS.AI 50%

arXiv:2608.02867v1 Announce Type: cross Abstract: Although reinforcement learning with verifiable rewards (RLVR) has improved the performance of large language models (LLMs) across a variety of reasoning tasks, there is significant debate as to whether RLVR expands the reasoning...

<https://arxiv.org/abs/2608.02867>

6 CONF

LLaDA MoE v2: Scaling Mixture-of-Experts Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2608.03457v1 Announce Type: new Abstract: Diffusion language models (dLLMs) offer an alternative to autoregressive (AR) language modeling, yet the scaling behavior of Mixture-of-Experts (MoE) dLLMs remains poorly understood. We systematically characterize how optimization...

<https://arxiv.org/abs/2608.03457>

7 CONF

WeClawArena: An Auditable Sandbox and Benchmark for Cross-User Agents Collaboration and Security in Human-Centered Agent Networks

ArXiv CS.AI 50%

arXiv:2608.03499v1 Announce Type: new Abstract: Recent advances in persistent personal-agent frameworks are making human-centered agent networks realistic deployment targets: each user can be served by an AI agent that acts on the user's behalf, maintains state, and communicates...

<https://arxiv.org/abs/2608.03499>

8 CONF

Can LLM design high-quality experiments? A Comprehensive and Systematic Benchmark on Autonomous Experimental Design

ArXiv CS.AI 50%

arXiv:2608.03501v1 Announce Type: new Abstract: AI for Research (AI4Research) leverages AI to automate and improve scientific workflows. While experimental design is a critical stage of the research process, prior work has focused primarily on code implementation and execution,...

<https://arxiv.org/abs/2608.03501>

9 CONF

Hybrid LLM-Augmented Reinforcement Learning Agents for Complex Sequential Decision Tasks

ArXiv CS.AI 50%

arXiv:2608.03502v1 Announce Type: new Abstract: Large Language Models (LLMs) have recently shown strong capabilities in reasoning, planning, and tool-use, enabling new forms of autonomous agents. However, LLM-based agents struggle with long-horizon sequential decision tasks that...

<https://arxiv.org/abs/2608.03502>

10 CONF

When Many Answers Are Valid, Voting Fails: Symbolic Verification for Best-of-K Causal Reasoning in LLMs

ArXiv CS.AI 50%

arXiv:2608.03506v1 Announce Type: new Abstract: Self-consistency assumes the most frequent answer among sampled reasoning traces is the most reliable, but this can fail in causal reasoning: samples often repeat the same confounding error, and votes fragment across multiple valid...

<https://arxiv.org/abs/2608.03506>

1 CONF

Soft Guidance Starts to Outperform CoT Prompting as LLMs Improve

ArXiv CS.AI 50%

arXiv:2608.03550v1 Announce Type: new Abstract: Chain-of-Thought (CoT) prompting remains the standard baseline for evaluating models' reasoning abilities. Originally, this technique was introduced to elicit step-by-step reasoning from large language models (LLMs), which would...

<https://arxiv.org/abs/2608.03550>

2 **CONF****Enhancing VLM Reward Models Through Structure-Aware Fine-Tuning****ArXiv CS.AI 50%**

arXiv:2608.03875v1 Announce Type: cross Abstract: Designing effective reward functions remains a major bottleneck in Reinforcement Learning (RL). Recent work uses large foundation Vision-Language Models (VLMs) as reward models, computing text-observation similarity to bypass...

<https://arxiv.org/abs/2608.03875>

3 **CONF****Adversarial Fast-Moving Real-World Domains as Test Beds for Benchmarking AI Scientist Capabilities****ArXiv CS.AI 50%**

arXiv:2608.03569v1 Announce Type: new Abstract: Benchmarking the ability of AI scientists to generate novel ideas is notoriously difficult. Existing benchmarks in this field have made progress in evaluating scientific reasoning and research replication, but often rely on...

<https://arxiv.org/abs/2608.03569>

4 **CONF****Policy Fragmentation or Institutional Alignment? Institutional Governance of AI in Universities and Business Schools****ArXiv CS.AI 50%**

arXiv:2608.03584v1 Announce Type: new Abstract: Artificial intelligence (AI) is rapidly transforming high-skilled domains, requiring higher education institutions (HEI) to balance the teaching of foundational principles with the integration of emerging tools to ensure workforce...

<https://arxiv.org/abs/2608.03584>

5 **CONF****Learning Clinical-Trial Strategy: Offline Policy Training for Decision Agents****ArXiv CS.AI 50%**

arXiv:2608.03606v1 Announce Type: new Abstract: Clinical development is sequential decision-making under uncertainty, where a sponsor must plan a portfolio of experiments from heterogeneous evidence. We study this setting by framing oncology clinical development as an offline...

<https://arxiv.org/abs/2608.03606>

6 **CONF****Rethinking Modality Reliability in Multimodal Sentiment Analysis with Incomplete Observations****ArXiv CS.AI 50%**

arXiv:2608.03611v1 Announce Type: new Abstract: Multimodal Sentiment Analysis (MSA) integrates text, audio, and vision to infer human affect, yet real-world multimodal observations are often incomplete. Existing methods for incomplete-observation MSA mainly follow two paradigms....

<https://arxiv.org/abs/2608.03611>

7 **CONF****Unequal Verdicts: Investigating Gender Bias in LLM-Based Fake News Detection****ArXiv CS.AI 50%**

arXiv:2608.03627v1 Announce Type: new Abstract: Large Language Models (LLMs) are increasingly used for automated fact-checking, yet their susceptibility to gender bias in this context remains underexplored. This study presents the first systematic investigation of gender bias in...

<https://arxiv.org/abs/2608.03627>

8 CONF

Cross-Layer Interaction under Weight-Space Ablation: A Closed-Form Attention Jacobian Bound and a Test on a Real Pretrained Model

ArXiv CS.AI 50%

arXiv:2608.03629v1 Announce Type: new Abstract: A companion paper studies when activation patching and weight-space ablation agree, inside an idealized model where a conditional computation is carried additively through a residual stream. For the one composition in that model...

<https://arxiv.org/abs/2608.03629>

9 CONF

Is Inter-Seed Cross-Play Enough? Evaluating the Robustness of Zero-Shot Coordination Algorithms to Implementation Details

ArXiv CS.AI 50%

arXiv:2608.03644v1 Announce Type: new Abstract: AI agents deployed in real-world settings must be capable of coordinating with humans and other AI agents they have not encountered before. Zero-shot coordination (ZSC) algorithms aim to achieve this by specifying high-level...

<https://arxiv.org/abs/2608.03644>

0 CONF

Taming the Implicit: Dual-Channel Risk-Aware Reinforcement Fine-Tuning for Continual Multimodal Post-Training

ArXiv CS.AI 50%

arXiv:2608.03660v1 Announce Type: new Abstract: Reinforcement fine-tuning (RFT) is widely believed to inherently resist catastrophic forgetting in continual post-training of multimodal large language models. Under pronounced task distributional shifts, however, forgetting across...

<https://arxiv.org/abs/2608.03660>

1 CONF

PhyAI: Real-Time Physical AI at the Edge, Scalable Rollouts in the Cloud

ArXiv CS.AI 50%

arXiv:2608.03682v1 Announce Type: new Abstract: Physical AI policies require inference throughout their lifecycle, including model evaluation, cloud reinforcement learning rollout, edge GPU serving, and onboard deployment. Although these settings share the same checkpoint and...

<https://arxiv.org/abs/2608.03682>

2 CONF

BulkPR-Bench: Benchmarking Queue-Level Governance of Interacting Pull Requests

ArXiv CS.AI 50%

arXiv:2608.02685v1 Announce Type: cross Abstract: Coding-agent benchmarks increasingly cover long-horizon, end-to-end, and interactive development, but typically retain one requested outcome or a fixed change sequence. Sequential policies can process a pull-request (PR) queue...

<https://arxiv.org/abs/2608.02685>

3 CONF

Agents Catching Agents: Shortcut Cascades and Benchmark Gaming in Clinical Multi-Agent Systems

ArXiv CS.AI 50%

arXiv:2608.03744v1 Announce Type: new Abstract: Clinical decision support is moving toward committees of language-model agents deliberating on a shared workspace. We ask whether such committees can be gamed by shortcuts, cues a benchmark rewards but a clinician would ignore....

<https://arxiv.org/abs/2608.03744>

4 **CONF****Computing Actual Causes for Neural Network Predictions under Structured Causal Inputs****ArXiv CS.AI** **50%**

arXiv:2608.03772v1 Announce Type: new Abstract: Explaining the predictions of neural networks is a central challenge in trustworthy AI. Existing explanation methods, such as those based on feature attribution or minimal sufficient sets, typically treat input features as...

<https://arxiv.org/abs/2608.03772>

5 **CONF****KnowHal: A Knowledge-Driven Benchmark for Comprehensive Multimodal Hallucination Evaluation****ArXiv CS.AI** **50%**

arXiv:2608.03782v1 Announce Type: new Abstract: Hallucination remains a critical challenge for developing trustworthy Multimodal Large Language Models (MLLMs). While existing benchmarks mainly focus on entity, attribute, and relation hallucinations, knowledge-related failures...

<https://arxiv.org/abs/2608.03782>

6 **CONF****LatentGuard: Efficient and Inspectable Latent Reasoning for LLM Safeguards****ArXiv CS.AI** **50%**

arXiv:2608.03838v1 Announce Type: new Abstract: Reasoning-based guard models improve LLM safeguards, but decoding explicit rationales for every interaction makes them costly to deploy. Although latent-reasoning methods reduce token generation by moving reasoning into continuous...

<https://arxiv.org/abs/2608.03838>

7 **CONF****Oilbird: Training-Free Speculative Decoding with Keys the Verifier Already Computes****ArXiv CS.AI** **50%**

arXiv:2608.03839v1 Announce Type: new Abstract: Training-free speculative decoding drafts by matching an exact suffix of the context against a pool of earlier context. That lookup misses correct drafts already in the pool, most visibly on tool-calling traffic, where a request...

<https://arxiv.org/abs/2608.03839>

8 **CONF****MAFIA: Query-Only Memory Attacks via Probing and Factual Injection against Audited LLM Agents****ArXiv CS.AI** **50%**

arXiv:2608.03844v1 Announce Type: new Abstract: Memory-augmented LLM agents rely on rich context for long-horizon reasoning and acting, yet their memory modules expose a persistent attack surface for malicious records, making the study of memory poisoning threats imperative....

<https://arxiv.org/abs/2608.03844>

9 **CONF****ContinualSkillBench: Can LLM Agents Truly Evolve Their Capabilities?****ArXiv CS.AI** **50%**

arXiv:2608.03874v1 Announce Type: new Abstract: Modern agent frameworks equip large language models with external skill libraries to solve complex tasks. However, it remains unclear whether these systems can effectively evolve their skills and whether the resulting skills...

<https://arxiv.org/abs/2608.03874>

- 0 **CONF**
- When Efficiency Becomes Fragility: Exploiting Dynamic Routing Vulnerabilities in Adaptive UAV Tracking**
- ArXiv CS.AI 50%
- arXiv:2608.03902v1 Announce Type: new Abstract: Resource constraints on UAV platforms have driven a paradigm shift in aerial tracking, from pursuing performance toward balancing accuracy with efficiency. Adaptive Transformer Trackers, which leverage an input-dependent dynamic...
- <https://arxiv.org/abs/2608.03902>
-
- 1 **CONF**
- Socially Grounded Agentic AI: Coordinating Plural Perspectives through Social Theory**
- ArXiv CS.AI 50%
- arXiv:2608.03910v1 Announce Type: new Abstract: As AI systems are deployed across increasingly diverse social contexts, alignment can no longer be framed as the optimization of a single, unified set of values. Instead, systems must be able to recognize, represent, and respond to...
- <https://arxiv.org/abs/2608.03910>
-
- 2 **CONF**
- Implementing Causal Perception: Competing SCMs and Situated Fairness**
- ArXiv CS.AI 50%
- arXiv:2608.03917v1 Announce Type: new Abstract: Causal perception occurs when agents with competing Structural Causal Models (SCMs) of the same system infer different probability distributions, including the hypothetical distributions implied by each agent's SCM under the same...
- <https://arxiv.org/abs/2608.03917>
-
- 3 **CONF**
- The Transformer Revolution, Part 1: Dynamic Processing through Output- Weight Interconnections**
- ArXiv CS.AI 50%
- arXiv:2608.03921v1 Announce Type: new Abstract: This paper offers a new interpretation of the Transformer during inference. Against the "stochastic parrot" view that large language models merely reproduce statistical regularities learned in training, we argue that Transformers...
- <https://arxiv.org/abs/2608.03921>
-
- 4 **CONF**
- Interpretable Adaptive Sampling for LLM Test-Time Scaling**
- ArXiv CS.AI 50%
- arXiv:2608.03961v1 Announce Type: new Abstract: Test-time scaling improves LLM reasoning by generating and aggregating multiple candidate answers, yet many pipelines use fixed per-query budgets that spend the same compute on easy and difficult prompts. These fixed budgets are...
- <https://arxiv.org/abs/2608.03961>
-
- 5 **CONF**
- PASE: Leveraging the Phonological Prior of WavLM for Low-Hallucination Generative Speech Enhancement**
- ArXiv CS.AI 50%
- arXiv:2511.13300v1 Announce Type: cross Abstract: Generative models have shown remarkable performance in speech enhancement (SE), achieving superior perceptual quality over traditional discriminative approaches. However, existing generative SE approaches often overlook the risk...
- <https://arxiv.org/abs/2511.13300>
-

- 6 CONF**
- UniPASE: A Generative Model for Universal Speech Enhancement with High Fidelity and Low Hallucinations**
- ArXiv CS.AI 50%
- arXiv:2604.14606v2 Announce Type: cross Abstract: Universal speech enhancement (USE) aims to restore speech signals from diverse distortions across multiple sampling rates. We propose UniPASE, an extension of the low-hallucination PASE framework tailored for USE. At its core is...
- <https://arxiv.org/abs/2604.14606>
-
- 7 CONF**
- OncoTriad-QA: A Patient-Level Radiology-Pathology-Genomics Benchmark for Pan-Cancer Reasoning**
- ArXiv CS.AI 50%
- arXiv:2608.02615v1 Announce Type: cross Abstract: Cancer diagnosis and characterization require integrating complementary evidence from radiology, pathology, genomics, and clinical metadata. However, most medical large language model (LLM) and vision-language model (VLM)...
- <https://arxiv.org/abs/2608.02615>
-
- 8 CONF**
- Evaluating OpenAI's Privacy Filter: Cross-Lingual, Cross-Domain PII Detection Across 42 Benchmarks**
- ArXiv CS.AI 50%
- arXiv:2608.02616v1 Announce Type: cross Abstract: We present the first independent, systematic evaluation of OpenAI's Privacy Filter (OPF), a 1.5B-parameter bidirectional PII detector, across 42 synthetic benchmarks spanning 22 languages and 5 domains. Zero-shot, OPF achieves...
- <https://arxiv.org/abs/2608.02616>
-
- 9 CONF**
- Preferred, Not Safer: Pairwise Preference Is a Poor Proxy for Clinical Safety**
- ArXiv CS.AI 50%
- arXiv:2608.02617v1 Announce Type: cross Abstract: We evaluate whether clinician pairwise preferences provide a reliable signal of clinical safety in large language model (LLM) evaluation using expert feedback from MOOVE (Massive Open Online Validation and Evaluation), a...
- <https://arxiv.org/abs/2608.02617>
-
- 0 CONF**
- Knowing the Form, Not the Function: Automatically Auditing Answer--Authority Decoupling in Legal Benchmarks**
- ArXiv CS.AI 50%
- arXiv:2608.02621v1 Announce Type: cross Abstract: Legal benchmarks typically score final answers even when models also state legal authority. We test whether answer correctness can serve as a proxy for authority grounding. Under ordinary reasoning prompts that did not request...
- <https://arxiv.org/abs/2608.02621>
-
- 1 CONF**
- Multimodal Auto-regressive Transformer Surrogate for Modeling Variable Operations and Quantifying Uncertainty in Geological Carbon Storage**
- ArXiv CS.AI 50%
- arXiv:2608.02629v1 Announce Type: cross Abstract: The use of variable well perforation and injection strategies can improve the efficiency of geological carbon storage operations. We develop a new multimodal auto-regressive transformer surrogate to model these operations under...
- <https://arxiv.org/abs/2608.02629>

2 **CONF****Studying, Identifying, and Fixing Hidden Technical Debt in AI-Intensive Cyber-Physical Systems****ArXiv CS.AI** **50%**

arXiv:2608.02638v1 Announce Type: cross Abstract: Artificial Intelligence (AI) components are increasingly pervasive in several software systems, including Cyber-Physical Systems (CPSs). AI-CPS are used in several domains, including autonomous vehicles, industry, home...

<https://arxiv.org/abs/2608.02638>

3 **CONF****Instruction Stacking Collapse: A Benchmark and the Capability-Dependent Value of Prompt Compilation****ArXiv CS.AI** **50%**

arXiv:2608.02639v1 Announce Type: cross Abstract: Production prompts rarely carry a single instruction. One system message may require valid JSON, a word limit, three citations, and a fixed tone at the same time. We study how instruction-following degrades as such constraints...

<https://arxiv.org/abs/2608.02639>

4 **CONF****Verified Tool Calls Improve LLM Agent Reliability Under Non-Atomic Failures****ArXiv CS.AI** **50%**

arXiv:2608.02645v1 Announce Type: cross Abstract: Large Language Model (LLM) agents rely on external tools to perform multistage tasks. Existing agent frameworks typically assume that tool calls are atomic and return binary success or failure signals. However, real-world systems...

<https://arxiv.org/abs/2608.02645>

5 **CONF****Cross-Anesthetic ECoG State Decoding Fails at the Decision Threshold, Not the Representation****ArXiv CS.AI** **50%**

arXiv:2608.02646v1 Announce Type: cross Abstract: Decoders of anesthetic state from cortical activity fail across drug classes, most notoriously ketamine, but reported accuracy cannot say whether the neural representation or only the decision threshold has failed; we separate...

<https://arxiv.org/abs/2608.02646>

6 **CONF****Secure AI Watermarking Framework for IP Protection in Multi-Tenant Cloud Platforms****ArXiv CS.AI** **50%**

arXiv:2608.02656v1 Announce Type: cross Abstract: The Secured data safe guard transaction with multi-tenant environments run on private-protected authenticate platforms runs by secured handed environments that emerges with the expansion of cloud-based AI services. To enhanced...

<https://arxiv.org/abs/2608.02656>

7 **CONF****Your Agentic LLMs Secretly Encode Latent Signals of Indirect Prompt-Injection Exposure****ArXiv CS.AI** **50%**

arXiv:2608.02657v1 Announce Type: cross Abstract: Agentic LLMs are vulnerable to indirect prompt injection (IPI) attacks, e.g., malicious side-tasks hidden in external tool results. While many efforts have sought to address the threats, little is known about the internals of...

<https://arxiv.org/abs/2608.02657>

- 8
CONF
Verifier-Guided Model Discovery for Physical Dynamical Systems with Pretrained Symbolic Transformers

ArXiv CS.AI 50%

arXiv:2608.02662v1 Announce Type: cross Abstract: Reliable forecasting of nonlinear physical systems underpins scientific discovery and engineering decision-making. Yet high-fidelity simulations are prohibitively costly, and machine-learning surrogates can be opaque and encode...

<https://arxiv.org/abs/2608.02662>
- 9
CONF
ZK-SR117: A Chunked Zero-Knowledge Attestation Design for Aggregated Fair-Lending Metrics, with a Control Mapping toward Full SR 11-7 Coverage

ArXiv CS.AI 50%

arXiv:2608.02664v1 Announce Type: cross Abstract: Deploying ML models in regulated decision-making (credit underwriting, fraud detection, loan approval) requires demonstrating fairness and robustness to auditors without exposing model weights or customer data. We address this...

<https://arxiv.org/abs/2608.02664>
- 0
CONF
Single Canonical Prompts Underestimate LLM Safety's Surface-Form Sensitivity

ArXiv CS.AI 50%

arXiv:2608.02665v1 Announce Type: cross Abstract: A benchmark score is a measurement instrument, yet most benchmarks read each item at a single canonical surface form. We ask whether that reading is faithful: when an item's intent is held fixed and only its meaning-preserving...

<https://arxiv.org/abs/2608.02665>
- 1
CONF
Security-First Evaluation of Text-to-Terraform: Benchmarking LLMs and SLMs for Secure IaC Generation

ArXiv CS.AI 50%

arXiv:2608.02672v1 Announce Type: cross Abstract: Cloud misconfiguration remains a leading cause of security incidents, yet whether LLMs and SLMs can generate security-compliant Infrastructure-as-Code is an open question. We benchmark seven models, three closed LLMs (Claude Opus...

<https://arxiv.org/abs/2608.02672>
- 2
CONF
Moving the Safety Barrier: Dynamic Routing Adaptive Alignment Against White-Box Attacks

ArXiv CS.AI 50%

arXiv:2608.02674v1 Announce Type: cross Abstract: With the widespread deployment of large foundation models (LFMs) in open environments, safety threats are shifting from black-box jailbreaks toward white-box attacks that directly identify and disrupt internal safety neurons or...

<https://arxiv.org/abs/2608.02674>
- 3
CONF
DenialRAG: Single-Document RAG Poisoning via Embedded Parametric Denial

ArXiv CS.AI 50%

arXiv:2608.02678v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) systems are vulnerable to corpus poisoning: an attacker who inserts a crafted document into the retrieval corpus can steer the underlying large language model (LLM) toward an attacker-chosen...

<https://arxiv.org/abs/2608.02678>

4 **CONF****AI Sandbox: Technical Report****ArXiv CS.AI** **50%**

arXiv:2608.02679v1 Announce Type: cross Abstract: Collaborative AI experimentation across industry and academia requires platforms that enable rapid prototyping while preserving controlled access, tenant separation, and transparent workflows. Despite growing interest in AI...

<https://arxiv.org/abs/2608.02679>

5 **CONF****TraceCompiler: Skill-Guided Mining and Compilation of LLM Agent Traces into Mostly Deterministic Workflows****ArXiv CS.AI** **50%**

arXiv:2608.02680v1 Announce Type: cross Abstract: Tool-using language-model agents repeatedly rediscover procedures they have already executed, producing traces that mix reusable structure with retries, exploration, accidental ordering, and repeated lookups. We present...

<https://arxiv.org/abs/2608.02680>

6 **CONF****Measuring Explainer Stability via Attribution Separability****ArXiv CS.AI** **50%**

arXiv:2608.02697v1 Announce Type: cross Abstract: Attribution methods (AMs) assign an importance score to each feature and are widely adopted to explain black-box models. However, most methods can produce variable attribution scores due to stochastic components in their...

<https://arxiv.org/abs/2608.02697>

7 **CONF****Can Training Logs Make Model Comparisons More Precise?****ArXiv CS.AI** **50%**

arXiv:2608.02705v1 Announce Type: cross Abstract: Comparing stochastically trained models requires estimating both a performance difference and its uncertainty from repeated runs. We study whether training logs from those same runs can make such comparisons more precise. Because...

<https://arxiv.org/abs/2608.02705>

8 **CONF****Don't Regenerate, Debug: A Domain-Specific Agent for Repairing Near-Miss Hardware Operators****ArXiv CS.AI** **50%**

arXiv:2608.02712v1 Announce Type: cross Abstract: Kernel generation for hardware accelerators such as GPUs and NPUs has become a proving ground for large language models (LLMs), and state-of-the-art systems raise correctness through pipelines that couple LLMs with agentic...

<https://arxiv.org/abs/2608.02712>

9 **CONF****Privacy-Preserving AI Verification via Minimal Information Disclosure****ArXiv CS.AI** **50%**

arXiv:2608.02774v1 Announce Type: cross Abstract: AI verification crosses a trust boundary: a verifier must learn enough to establish an authorized claim, yet the same evidence can reveal sensitive details about the model, workload, or hardware. We introduce minimal information...

<https://arxiv.org/abs/2608.02774>

00 CONF

A Hyperfinite Framework for Score-Based Generative Modeling

ArXiv CS.AI 50%

arXiv:2608.02799v1 Announce Type: cross Abstract: Score-based diffusion models are typically formulated using continuous-time stochastic differential equations and measure-theoretic stochastic calculus. In this paper, we develop a hyperfinite formulation of score-based...

<https://arxiv.org/abs/2608.02799>

01 CONF

SAGE: Semantic Explainability of Attention-Based Survival Models in Computational Pathology

ArXiv CS.AI 50%

arXiv:2608.02803v1 Announce Type: cross Abstract: Attention-based multiple instance learning (ABMIL) is the predominant approach for slide-level prediction in computational pathology, yet its attention maps provide only local explanations: they indicate where a model focuses but...

<https://arxiv.org/abs/2608.02803>

02 CONF

Evading Chain-of-Thought Monitoring Through Model Poisoning

ArXiv CS.AI 50%

arXiv:2608.02820v1 Announce Type: cross Abstract: Chain-of-thought (CoT) monitoring is an increasingly important component of AI safety stacks but relies on the assumption that a model's reasoning trace is informative about its actions. This work studies the limits of CoT...

<https://arxiv.org/abs/2608.02820>

03 CONF

In-Context Collapse in Vision-Language Models and How to Mitigate it?

ArXiv CS.AI 50%

arXiv:2608.02830v1 Announce Type: cross Abstract: Many-shot in-context learning (ICL) lets vision-language models (VLMs) adapt from image-label demonstrations without weight updates, and is widely assumed to improve as more demonstrations are supplied. We show the opposite: as...

<https://arxiv.org/abs/2608.02830>

04 CONF

ValueFormer: A Causal Transformer Value Function with Stage-Aware Labels for Semi-Autonomous Vision-Language-Action Policies

ArXiv CS.AI 50%

arXiv:2608.02958v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) policies trained by behavior cloning fail silently: from the action stream alone, a collapsing rollout looks much like one making clean progress, because imitation supplies no notion of progress....

<https://arxiv.org/abs/2608.02958>

05 CONF

TQLite: Multi-LLM Jury Guided Distillation for Real-time MQM Translation Quality Evaluation

ArXiv CS.AI 50%

arXiv:2608.02975v1 Announce Type: cross Abstract: Large language models (LLMs) have demonstrated impressive performance in MQM-based translation quality (TQ) evaluation, and recent advances in large reasoning models (LRMs) promise even greater improvements. However, both LLMs...

<https://arxiv.org/abs/2608.02975>

06 CONF

Internalising the Identity Primitive: Cryptographic Individuality for an Autonomous Agent on a Public Blockchain

ArXiv CS.AI 50%

arXiv:2608.02986v1 Announce Type: cross Abstract: A software agent on a public blockchain accumulates authority and economic stakes, raising the engineering question of what makes it count as an individual. The paper's central contribution is a shift of trust root for the...

<https://arxiv.org/abs/2608.02986>

07 CONF

SparSEETy: Extracting Tokens from Sparsity-Exploiting LLM Serving Systems via Deterministic Side Channels

ArXiv CS.AI 50%

arXiv:2608.02995v1 Announce Type: cross Abstract: Modern large language models (LLMs) exhibit activation sparsity, wherein only a subset of their neurons is activated for given input tokens. Researchers have leveraged this property to optimize LLM serving systems by omitting...

<https://arxiv.org/abs/2608.02995>

08 CONF

A Graph Signal Processing Perspective on Numerical Sequence Representations in LLM In-Context Learning

ArXiv CS.AI 50%

arXiv:2608.03015v1 Announce Type: cross Abstract: Pretrained large language models (LLMs) have demonstrated in-context learning (ICL) capabilities for numerical inference over sequences serialized as text. Prior work has identified and characterized this form of numerical...

<https://arxiv.org/abs/2608.03015>

09 CONF

LLM Serving in the Wild: An Empirical Study of Frameworks, Methods, and System Designs

ArXiv CS.AI 50%

arXiv:2608.03036v1 Announce Type: cross Abstract: Large Language Models (LLMs) are integrated into software systems and AI services, making efficient LLM serving a concern for software engineering. Serving LLMs is challenging because inference requires computation, memory, GPU...

<https://arxiv.org/abs/2608.03036>

10 CONF

Emulate or Estimate? The Divergent Strengths of Base and Post-Trained Language Models for Opinion Simulation

ArXiv CS.AI 50%

arXiv:2608.03044v1 Announce Type: cross Abstract: Large language models are increasingly used to simulate human opinions, but prior work reports conflicting results: some studies find promising alignment with human survey data, while others find persona collapse and weak...

<https://arxiv.org/abs/2608.03044>

11 CONF

CVPO: Enhancing LLM Reinforcement Learning Reasoning via Value-Variance Adaptation and Dynamic Curriculum Learning

ArXiv CS.AI 50%

arXiv:2608.03068v1 Announce Type: cross Abstract: Reinforcement learning (RL) has emerged as an effective method for enhancing the reasoning capabilities of large language models (LLMs). However, existing methods suffer from insufficient precision in feedback on generated answer...

<https://arxiv.org/abs/2608.03068>

12 CONF

AI Security Leaderboard: Methodology, Results and Minimal Standard

ArXiv CS.AI 50%

arXiv:2608.03070v1 Announce Type: cross Abstract: Frontier AI model developers increasingly rely on layered safeguards to prevent catastrophic misuse, but little public evidence exists on how much protection these safeguards provide, or how consistently across developers. We...

<https://arxiv.org/abs/2608.03070>

13 CONF

CorePath: A Breast-Specialized Pathology Foundation Model for Core Needle Biopsy Diagnosis and Risk-Controlled Report Generation

ArXiv CS.AI 50%

arXiv:2608.03079v1 Announce Type: cross Abstract: Breast core needle biopsy (CNB) is central to breast cancer diagnosis yet remains challenging because limited tissue sampling, lesion heterogeneity, and subtle morphologic overlap can obscure subtype distinctions. We developed...

<https://arxiv.org/abs/2608.03079>

14 CONF

FakeI2V-Bench: Benchmarking the Applicability of Image-level Deepfake Detectors for Deepfake Video Detection

ArXiv CS.AI 50%

arXiv:2608.03096v1 Announce Type: cross Abstract: Recent advances in video generation models have significantly intensified the deepfake threat, yet the current deepfake video detection benchmarks remain underdeveloped. In particular, the effectiveness of image-level detectors...

<https://arxiv.org/abs/2608.03096>

15 CONF

Optimal Liability Design for Medical AI

ArXiv CS.AI 50%

arXiv:2608.03114v1 Announce Type: cross Abstract: Artificial intelligence (AI) is increasingly integrated into medical decision-making, yet its liability implications remain complex, particularly when physicians differ in diagnostic skills and their quality is unobservable. This...

<https://arxiv.org/abs/2608.03114>

16 CONF

Minimax-Optimal Semiparametric Contextual Dynamic Pricing with Multimodal Revenue

ArXiv CS.AI 50%

arXiv:2608.03142v1 Announce Type: cross Abstract: We study contextual dynamic pricing with arbitrary covariate sequences and bounded, possibly nonbinary purchase quantities. Demand follows a semiparametric surplus-index model with an unknown linear valuation parameter and an...

<https://arxiv.org/abs/2608.03142>

17 CONF

GROW: Group-Relative Advantage-Weighted On-Policy Reinforcement Learning of Autoregressive-Diffusion Text-to-Speech model

ArXiv CS.AI 50%

arXiv:2608.03215v1 Announce Type: cross Abstract: Reinforcement learning for flow-matching text-to-speech is complicated by deterministic ODE sampling: trajectory-level policy-gradient methods typically convert the ODE into an SDE and track per-step likelihood ratios,...

<https://arxiv.org/abs/2608.03215>

18 CONF

Self-Supervised Representation-Guided Generative Dataset Distillation

ArXiv CS.AI 50%

arXiv:2608.03218v1 Announce Type: cross Abstract: Dataset distillation compresses a large training set into a compact synthetic set while retaining its downstream utility. Most existing methods target randomly initialized networks, whereas modern vision systems often adapt...

<https://arxiv.org/abs/2608.03218>

19 CONF

Fail-Fast, Restart-Smart: Early Failure Prediction and Restart for SWE Agentic Tasks

ArXiv CS.AI 50%

arXiv:2608.03222v1 Announce Type: cross Abstract: Software engineering (SWE) agents resolve repository-level issues through long trajectories that grow increasingly expensive as context accumulates. Failed runs tend to be longer and exhibit redundant exploration or looping,...

<https://arxiv.org/abs/2608.03222>

20 CONF

Structure-Aware Robust Fine-Tuning: Defending Vision-Language-Action Robots Against Physical Attention Hijacking

ArXiv CS.AI 50%

arXiv:2608.03231v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) policies promise general robotic manipulation, but their robustness against physical-world attacks remains fragile. In particular, we show that physically realizable adversarial patches can reliably...

<https://arxiv.org/abs/2608.03231>

21 CONF

FinVerse: Financial Time-Series Benchmark

ArXiv CS.AI 50%

arXiv:2608.03259v1 Announce Type: cross Abstract: As time-series foundation models have emerged, the need for benchmarks that can evaluate their forecasting ability in meaningful ways has become increasingly important. Existing time-series forecasting benchmarks provide useful...

<https://arxiv.org/abs/2608.03259>

22 CONF

The Tell-Tale Trace: Detecting Reasoning Failures in LLMs Using Chain-of-Thought Dynamics

ArXiv CS.AI 50%

arXiv:2608.03291v1 Announce Type: cross Abstract: Chain-of-thought (CoT) reasoning improves large language model (LLM) performance while also providing an observable interface to the model's reasoning process. Existing approaches that leverage verbalized CoTs to monitor...

<https://arxiv.org/abs/2608.03291>

23 CONF

Evaluating LLM Trade-offs for Enterprise Automation: Lessons from Workflow Generation in a Production Enterprise Platform

ArXiv CS.AI 50%

arXiv:2608.03311v1 Announce Type: cross Abstract: Enterprise compliance management requires rapid adaptation to evolving regulatory frameworks (e.g., DORA, AI RMF, FedRAMP) and tight remediation SLAs. Traditional static orchestrators often fail in hybrid cloud environments where...

<https://arxiv.org/abs/2608.03311>

24 CONF

When Oracle Conditioning Misleads Deployment: Conditioning-Availability Bias in Echocardiographic Segmentation

ArXiv CS.AI 50%

arXiv:2608.03342v1 Announce Type: cross Abstract: Conditional segmentation models may be trained and evaluated with auxiliary signals cleaner than those available at deployment. We study this protocol-level manifestation of shortcut learning and auxiliary-variable shift in...

<https://arxiv.org/abs/2608.03342>

25 CONF

The Evolutionary Origin of Values: implications for AI alignment, sentience and existential risk

ArXiv CS.AI 50%

arXiv:2608.03361v1 Announce Type: cross Abstract: AI systems based on Large Language Models (LLMs) have prompted fears that they may harbor hidden goals, seek to dominate or eliminate humanity, or even suffer as sentient beings. We address these concerns by tracing the...

<https://arxiv.org/abs/2608.03361>

26 CONF

FACTWASH: Catching AI Rewrites That Wash Hearsay into Fact

ArXiv CS.AI 50%

arXiv:2608.03372v1 Announce Type: cross Abstract: AI systems rewrite information constantly: conversations become stored memories, documents become answers. The rewrite can keep a claim while washing away what made it checkable, who said it, how sure they were, when it held. We...

<https://arxiv.org/abs/2608.03372>

27 CONF

Shaping Wind-Tunnel Airflow for Unmanned Aerial Vehicles using Online Learning

ArXiv CS.AI 50%

arXiv:2608.03378v1 Announce Type: cross Abstract: The development and testing of advanced aerial robots require experiments in controlled environments with tailored airflow profiles. This paper presents an online learning algorithm for controlling the complex airflow field in a...

<https://arxiv.org/abs/2608.03378>

28 CONF

OliveGemma: A 3 Billion Visual Language Model for Recognising the Mediterranean & European Diet

ArXiv CS.AI 50%

arXiv:2608.03428v1 Announce Type: cross Abstract: Image based dietary assessment offers a scalable alternative to self reported food diaries, yet fine-grained food recognition remains challenging due to high intra-class variability and visually similar dishes. This study...

<https://arxiv.org/abs/2608.03428>

29 CONF

Balancing Efficiency and Efficacy: Training-Free Attention-Guided Switching Between Explicit and Latent Thoughts for MLLMs

ArXiv CS.AI 50%

arXiv:2608.03450v1 Announce Type: cross Abstract: Reasoning in Multimodal Large Language Models (MLLMs) requires both fine-grained visual perception and rigorous logical deduction. Explicit text-based Chain-of-Thought (CoT) is computationally expensive and prone to visual...

<https://arxiv.org/abs/2608.03450>

30 CONF

Leveraging System-Level Observations to Inform Bayesian Learning of Model Parameters for Quantitative Verification

ArXiv CS.AI 50%

arXiv:2608.03489v1 Announce Type: cross Abstract: Combining Bayesian learning and quantitative verification is a powerful toolset for analysing key quantitative properties of software systems, like reliability and response time. However, the accuracy and robustness of...

<https://arxiv.org/abs/2608.03489>

31 CONF

AI Forensics Across White-, Grey-, and Black-Box Access: A Process Model and Research Agenda for Post-Incident Investigation of AI Systems

ArXiv CS.AI 50%

arXiv:2608.03520v1 Announce Type: cross Abstract: AI systems are increasingly involved in decisions and actions that may later require investigation. When an AI related incident occurs, investigators need to reconstruct what the system did, why it behaved that way, and which...

<https://arxiv.org/abs/2608.03520>

32 CONF

Training Documents Reranker with Search Rubrics for Deep Research Agent

ArXiv CS.AI 50%

arXiv:2608.03527v1 Announce Type: cross Abstract: Retrieval systems help deep research agents generate high-quality answers by providing relevant documents. However, existing retrievers typically select documents through relevance matching, while individually well-matched...

<https://arxiv.org/abs/2608.03527>

33 CONF

Pin Once, Swap Light: Subspace-Aligned Centroid-Residual Training for Efficient Ultra-LoRA Serving

ArXiv CS.AI 50%

arXiv:2608.03579v1 Announce Type: cross Abstract: Modern multi-tenant Low-Rank Adapters (LoRAs) serving systems concurrently host tens to hundreds of LoRA adapters. Though powerful, this introduces a critical system dilemma between serving efficiency and task performance:...

<https://arxiv.org/abs/2608.03579>

34 CONF

AI-Assisted Peer Review Across Research Communities: From Reviewer AI Policies to LLM Review Quality

ArXiv CS.AI 50%

arXiv:2608.03581v1 Announce Type: cross Abstract: AI-assisted peer review is increasingly discussed and adopted as a tool to support the scientific publishing process, yet there is little systematic understanding of how publication venues regulate its use or of how capable...

<https://arxiv.org/abs/2608.03581>

35 CONF

MuEvo: LLM-Driven Evolution of Multi-Heuristic Ensemble

ArXiv CS.AI 50%

arXiv:2608.03636v1 Announce Type: cross Abstract: Large language model-based automated heuristic design (LLM-AHD) has shown strong potential in discovering effective heuristics for combinatorial optimization problems. However, existing methods primarily optimize a single...

<https://arxiv.org/abs/2608.03636>

36 CONF

Decoupling Generation and Selection for Budget-Constrained Faithful Summarization

ArXiv CS.AI 50%

arXiv:2608.03655v1 Announce Type: cross Abstract: Abstractive summarization models remain vulnerable to factual inconsistency, redundancy, and weak length control. We propose a modular generation-and-selection framework for sentence-budget-constrained summarization. A pretrained...

<https://arxiv.org/abs/2608.03655>

37 CONF

How Closely Do LLM Reviews Align with Human Peer Review?

ArXiv CS.AI 50%

arXiv:2608.03659v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to generate scientific reviews, yet existing evaluations rarely examine whether different providers align with both conference decisions and human reviewing priorities within the...

<https://arxiv.org/abs/2608.03659>

38 CONF

GPTKB 2.0: Direct Construction of Disambiguated Knowledge Bases from Large Language Models

ArXiv CS.AI 50%

arXiv:2608.03729v1 Announce Type: cross Abstract: Automated Knowledge Base Construction (AKBC) is a core NLP task, and recent work proposes generating knowledge bases directly from large language models (LLMs), treating the model itself as the knowledge source. However, LLMs...

<https://arxiv.org/abs/2608.03729>

39 CONF

AI-Based Sound Effect Generation: A Narrative Review of Generative Models Across Input Modalities

ArXiv CS.AI 50%

arXiv:2608.03742v1 Announce Type: cross Abstract: Sound effects play a crucial role in conveying actions, events, and environmental cues across digital applications, often requiring a high degree of variation and contextual adaptability. Artificial intelligence (AI)-driven audio...

<https://arxiv.org/abs/2608.03742>

40 CONF

Can LLMs Test Terminal User Interfaces?

ArXiv CS.AI 50%

arXiv:2608.03743v1 Announce Type: cross Abstract: Terminal User Interfaces (TUIs) combine the stateful, screen-oriented behaviour of GUIs with terminal deployment and are now common in developer tools. Yet they lack a dedicated testing methodology. We survey 197 real-world TUI...

<https://arxiv.org/abs/2608.03743>

41 CONF

Evaluating LLMs in Database Scenarios: A Lifecycle Benchmark for Assessing Their Potential in Core Database Tasks

ArXiv CS.AI 50%

arXiv:2608.03794v1 Announce Type: cross Abstract: Large Language Models (LLMs) are transforming database interaction paradigms, evolving from simple query translators to autonomous database administrators (DBAs). However, current evaluation benchmarks remain disproportionately...

<https://arxiv.org/abs/2608.03794>

42 CONF

Efficient Knowledge Distillation for LLMs: Offline Top-K Logits and a Fused Chunked KL Loss

ArXiv CS.AI 50%

arXiv:2608.03796v1 Announce Type: cross Abstract: Small language models are often the only option for deployment under tight latency, cost, and on-premises constraints, but they are rarely trained from scratch: a compressed model is usually recovered through knowledge...

<https://arxiv.org/abs/2608.03796>

43 CONF

Autoreflection: How Agentic Strange Loops Turn Human Culture into AI Infrastructure

ArXiv CS.AI 50%

arXiv:2608.03800v1 Announce Type: cross Abstract: An LLM-based agent is a loop that reads itself. Agentic frameworks externalize identity, memory, and disposition into editable files. The agent loads and edits these files during each activation. I argue that this architecture...

<https://arxiv.org/abs/2608.03800>

44 CONF

VIBE: A VAD-Informed Benchmark for Entity-Centered Affective Profiling of Large Language Model Outputs

ArXiv CS.AI 50%

arXiv:2608.03810v1 Announce Type: cross Abstract: Large language models routinely describe socially salient targets, including political figures, countries, religions, organizations, historical events, and social groups, encoding affective framing alongside factual content: a...

<https://arxiv.org/abs/2608.03810>

45 CONF

Beyond Representational Similarity: Source-Conditioned Description-Length Gain for Generative Plagiarism Detection and Candidate Source Reranking

ArXiv CS.AI 50%

arXiv:2608.03859v1 Announce Type: cross Abstract: Large language models (LLMs) pose challenges to academic integrity and peer review. Yet generative plagiarism detection remains an underexplored and largely unresolved challenge. Prior work on LLM-generated-text detection targets...

<https://arxiv.org/abs/2608.03859>

46 CONF

SciRet: A Compute-Aware Empirical Study of Retrieval and Reranking for Scientific RAG

ArXiv CS.AI 50%

arXiv:2608.03860v1 Announce Type: cross Abstract: We introduce SciRet, a compute-aware empirical study of retrieval-augmented generation for scientific question answering over CORD-19. Rather than proposing a new model, we evaluate a fixed scientific RAG pipeline across three...

<https://arxiv.org/abs/2608.03860>

47 CONF

GENESIS: Towards Explainable Causal Discovery

ArXiv CS.AI 50%

arXiv:2608.03868v1 Announce Type: cross Abstract: Causal Discovery (CD) from observational data faces two fundamental challenges. First, purely statistical methods often lack the power to resolve structural ambiguities in low-sample regimes. Second, although LLM-assisted hybrid...

<https://arxiv.org/abs/2608.03868>

48 **CONF****Equivariant Music Transformer****ArXiv CS.AI** **50%**

arXiv:2608.03920v1 Announce Type: cross Abstract: Humans recognize a musical passage even when it is shifted in time or transposed in pitch, indicating a notion of equivariance in the representation space. Our analysis, however, shows that standard music transformers map such...

<https://arxiv.org/abs/2608.03920>

49 **CONF****Logic Before Language: Pre-pretraining on Formal Derivations Fosters Skill Acquisition and Compressibility****ArXiv CS.AI** **50%**

arXiv:2608.03930v1 Announce Type: cross Abstract: Pre-pretraining language models (LMs) on symbolic data can accelerate and improve natural language acquisition. However, existing pre-pretraining tasks, such as Dyck and procedural algorithms, rely on narrow primitives that fail...

<https://arxiv.org/abs/2608.03930>

50 **CONF****Separating quantum circuits from classical LLMs****ArXiv CS.AI** **50%**

arXiv:2608.03962v1 Announce Type: cross Abstract: Modern large language models - transformers and diffusion language models - are built around two canonical algorithmic tasks: prediction and generation. We prove unconditional separations between low-depth quantum computation and...

<https://arxiv.org/abs/2608.03962>

51 **CONF****Can Large Language Models Recover Semantic Optimization Opportunities That Compilers Miss?****ArXiv CS.AI** **50%**

arXiv:2608.03983v1 Announce Type: cross Abstract: Optimizing compilers miss profitable transformations when their enabling semantics are absent from the analyzed program representation. We ask whether large language models (LLMs) can recover such semantics from heterogeneous...

<https://arxiv.org/abs/2608.03983>

52 **CONF****Test-Time Scaling in Reasoning LLMs: Inference Regimes, Evaluation, and Reproducibility****ArXiv CS.AI** **50%**

arXiv:2608.04001v1 Announce Type: cross Abstract: Large language models can solve substantially harder reasoning problems with more inference-time compute. The term "test-time scaling," however, now covers diverse inference algorithms that extend deliberation along a single...

<https://arxiv.org/abs/2608.04001>

53 **CONF****Modeling Matches as Language: A Generative Transformer Approach for Counterfactual Player Valuation in Football****ArXiv CS.AI** **50%**

arXiv:2603.15212v2 Announce Type: replace Abstract: Evaluating football player transfers is challenging because player actions depend strongly on tactical systems, teammates, and match context. Despite this complexity, recruitment decisions often rely on static statistics and...

<https://arxiv.org/abs/2603.15212>

54 CONF

Assessing the Effect of Cross-Domain Mapping on Creativity in Humans and Large Language Models

ArXiv CS.AI 50%

arXiv:2603.19087v2 Announce Type: replace Abstract: Creativity is the ability to come up with novel ideas, a capacity crucial for human development and flourishing. Are large language models (LLMs) creative in the same way humans are, and can the same interventions increase...

<https://arxiv.org/abs/2603.19087>

55 CONF

TransVLM: A Vision-Language Framework and Benchmark for Detecting Any Shot Transitions

ArXiv CS.AI 50%

arXiv:2604.27975v2 Announce Type: replace-cross Abstract: Traditional Shot Boundary Detection (SBD) inherently struggles with complex transitions by formulating the task around isolated cut points, frequently yielding corrupted video shots. We address this fundamental limitation...

<https://arxiv.org/abs/2604.27975>

56 CONF

Language model agents show in-group trust bias invisible to standard behavioural audits

ArXiv CS.AI 50%

arXiv:2605.28114v2 Announce Type: replace Abstract: Language-model agents are moving from single-user assistants into persistent networks that build trust and reputation with one another, and the same models increasingly control physically embodied robots as well as software....

<https://arxiv.org/abs/2605.28114>

57 CONF

Designing for Doubt: The Case for Informed Abstention in Autonomous Agents

ArXiv CS.AI 50%

arXiv:2606.02965v2 Announce Type: replace Abstract: As large language models gain tool access and are deployed as autonomous agents capable of editing records, executing transactions, and modifying infrastructure, we still evaluate them based on the sole metric of task...

<https://arxiv.org/abs/2606.02965>

58 CONF

The Geometric Nature and a Free Proxy for Flow-Matching Uncertainty

ArXiv CS.AI 50%

arXiv:2607.27933v2 Announce Type: replace Abstract: Flow matching (FM) has become a popular action head paradigm for modern embodied models. However, as a conditional generative model, it does not explicitly expose its inherent uncertainty, producing faulty action chunks even...

<https://arxiv.org/abs/2607.27933>

59 CONF

SKILL-KD: Contrastive Skill Distillation for LLM Agents

ArXiv CS.AI 50%

arXiv:2607.28048v2 Announce Type: replace Abstract: Skill-based prompting has become a practical mechanism for improving large language model (LLM) agents, yet existing skill acquisition methods often treat skills as experience summaries, memory entries, or direct summaries of...

<https://arxiv.org/abs/2607.28048>

60 CONF

Measurement Without Validity: The Compounding Reliability Problem in Agentic AI Evaluation

ArXiv CS.AI 50%

arXiv:2608.00794v2 Announce Type: replace Abstract: Agentic AI evaluation pipelines produce benchmark scores that justify deployment decisions, safety certifications, and regulatory compliance claims. No formal framework has yet characterized how validity degrades across the...

<https://arxiv.org/abs/2608.00794>

61 CONF

When Memory Becomes Authority: Benchmarking Authority Collapse at the Memory Consolidation Boundary

ArXiv CS.AI 50%

arXiv:2608.01679v2 Announce Type: replace Abstract: Persistent memory allows (self-evolving) LLM agents to adapt across tasks by consolidating heterogeneous interaction histories into reusable facts, preferences, observations, and rules. Yet consolidation also imposes an...

<https://arxiv.org/abs/2608.01679>

62 CONF

A Survey on Design Methodologies for Accelerating Deep Learning on Heterogeneous Architectures

ArXiv CS.AI 50%

arXiv:2311.17815v3 Announce Type: replace-cross Abstract: Given their increasing size and complexity, the need for efficient execution of deep neural networks has become increasingly pressing in the design of heterogeneous High-Performance Computing (HPC) and edge platforms,...

<https://arxiv.org/abs/2311.17815>

63 CONF

Efficient unsupervised domain adaptation via self-supervised vision transformer and synergistic cross-domain alignment

ArXiv CS.AI 50%

arXiv:2407.21311v2 Announce Type: replace-cross Abstract: Unsupervised domain adaptation (UDA) aims to mitigate domain shift, where the distribution of labeled source data differs from that of unlabeled target data. Despite recent advances, existing methods often rely on...

<https://arxiv.org/abs/2407.21311>

64 CONF

Compound and Parallel Modes of Tropical Convolutional Neural Networks

ArXiv CS.AI 50%

arXiv:2504.06881v2 Announce Type: replace-cross Abstract: Convolutional neural networks (CNNs) are foundational to many state-of-the-art computer vision systems, yet their reliance on multiplication-intensive computations poses challenges for deployment on resource-constrained...

<https://arxiv.org/abs/2504.06881>

65 CONF

From Generator to Embedder: Harnessing Innate Abilities of Multimodal LLMs via Building Zero-Shot Discriminative Embedding Model

ArXiv CS.AI 50%

arXiv:2508.00955v3 Announce Type: replace-cross Abstract: Adapting generative Multimodal Large Language Models (MLLMs) into universal embedding models typically demands resource-intensive contrastive pre-training, while traditional hard negative mining methods suffer from severe...

<https://arxiv.org/abs/2508.00955>

66 CONF

Speech LLMs in Low-Resource Scenarios: Data Volume Requirements and the Impact of Pretraining on High-Resource Languages

ArXiv CS.AI 50%

arXiv:2508.05149v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have demonstrated potential in handling spoken inputs for high-resource languages, reaching state-of-the-art performance in various tasks. However, their applicability is still less explored...

<https://arxiv.org/abs/2508.05149>

67 CONF

 π -Attention: Online Efficient Sparse Transformers for Long-Context Modeling

ArXiv CS.AI 50%

arXiv:2511.10696v3 Announce Type: replace-cross Abstract: Sparse attention is crucial in long-context Transformers, which restricts each token to a limited neighborhood and thereby reduces the quadratic cost of full self-attention. Local windows capture nearby context...

<https://arxiv.org/abs/2511.10696>

68 CONF

Self-Guided Adaptive Safety Alignment: Synthesizing and Internalizing Guidelines in Reasoning Models

ArXiv CS.AI 50%

arXiv:2511.21214v4 Announce Type: replace-cross Abstract: Explicit safety policies can improve reasoning-model safety, but their effective coverage may lag behind evolving jailbreak strategies. We study whether a reasoning model can synthesize and internalize a task-specific...

<https://arxiv.org/abs/2511.21214>

69 CONF

PRISMA: Improving the Accuracy-Latency Frontier of Diffusion-based PDE Solvers Using Physics-Informed Spectral Attention

ArXiv CS.AI 50%

arXiv:2512.01370v2 Announce Type: replace-cross Abstract: Diffusion-based solvers for partial differential equations (PDEs) are often bottle-necked by slow gradient-based test-time optimization routines that use PDE residuals for loss guidance. They additionally suffer from...

<https://arxiv.org/abs/2512.01370>

70 CONF

Where Knowledge Collides: A Mechanistic Study of Intra-Memory Knowledge Conflict in Language Models

ArXiv CS.AI 50%

arXiv:2601.09445v2 Announce Type: replace-cross Abstract: In language models (LMs), intra-memory knowledge conflict arises when inconsistent information about the same subject is encoded within the model's parametric knowledge. Prior work has primarily focused on resolving...

<https://arxiv.org/abs/2601.09445>

71 CONF

AgenticSCR: An Autonomous Agentic Secure Code Review for Immature Vulnerabilities Detection

ArXiv CS.AI 50%

arXiv:2601.19138v2 Announce Type: replace-cross Abstract: Secure code review is critical during pre-integration, where Atlassian developers rely on lightweight analysis tools, while deep security assessment is deferred to later stages, delaying feedback and increasing...

<https://arxiv.org/abs/2601.19138>

72 CONF

Injection-Execution Dissociation: A Mechanistic Evaluation of Persistent Memory Attacks and Defenses in Stateful LLM Agents

ArXiv CS.AI 50%

arXiv:2605.08442v5 Announce Type: replace-cross Abstract: We discover that prompt-injection success and tool-execution success are separable safety properties: defenses that block injection do not necessarily block execution, and vice versa. We call this the injection-execution...

<https://arxiv.org/abs/2605.08442>

73 CONF

MalariAI: A Label-Resilient Decoupled Framework for Annotation-Agnostic Cell Segmentation and Explainable Stage Classification in Dense Malaria Blood Smears

ArXiv CS.AI 50%

arXiv:2607.00385v2 Announce Type: replace-cross Abstract: Automated malaria diagnosis from blood smear microscopy is a critical global health AI challenge; expert scarcity remains the primary diagnostic bottleneck. Existing deep learning systems face three compounding failures:...

<https://arxiv.org/abs/2607.00385>

74 CONF

Foundations of Equivariant Deep Learning: Unifying Graph and Sheaf Neural Networks

ArXiv CS.AI 50%

arXiv:2607.03798v2 Announce Type: replace-cross Abstract: Symmetry is everywhere in nature and society. Geometric deep learning builds architectures respecting group symmetries, whereas topological deep learning organizes computation through cells, incidence relations, and...

<https://arxiv.org/abs/2607.03798>

75 CONF

TriGlue: a Biology-Inspired Generative Model for Generating Molecular Glue-Induced Ternary Complex

ArXiv CS.AI 50%

arXiv:2607.22143v2 Announce Type: replace-cross Abstract: Molecular glue degraders have emerged as a promising strategy for targeted protein degradation by inducing ternary complex formation between an E3 ubiquitin ligase and a target protein. Despite their therapeutic...

<https://arxiv.org/abs/2607.22143>

76 CONF

AgentGUI: An Interface for Observing and Steering Long-Running AI Agents

ArXiv CS.AI 50%

arXiv:2607.26300v2 Announce Type: replace-cross Abstract: AI agents are increasingly adept at tackling complex, long-running tasks. With the rapid surge of autonomous capabilities, human oversight is systematically lagging behind due to limited human-centered interfacing. Aiming...

<https://arxiv.org/abs/2607.26300>

77 CONF

Benchmarking LLM Competence on Logical Inference over Probability Operators

ArXiv CS.AI 50%

arXiv:2607.27405v3 Announce Type: replace-cross Abstract: Both expressions of uncertainty and inferences are ubiquitous in natural language, and valid inferences over natural-language expressions of uncertainty are necessary for not only everyday conversations but also for...

<https://arxiv.org/abs/2607.27405>

78 CONF

PAIChecker: Uncovering and Checking PR-Issue Misalignment in SWE-Bench-Like Benchmarks

ArXiv CS.AI 50%

arXiv:2607.28587v2 Announce Type: replace-cross Abstract: SWE-bench-like benchmarks are widely used for evaluating LLM's issue resolution capability. They typically follow a common construction pipeline: each PR (Pull Request) is paired with its linked issue by extracting issue...

<https://arxiv.org/abs/2607.28587>

79 CONF

Which Modality Decides? Counterfactual Modality Attribution for Multimodal LLMs

ArXiv CS.AI 50%

arXiv:2608.00076v2 Announce Type: replace-cross Abstract: Multimodal large language models (MLLMs) increasingly support high-stakes decision making by combining complementary information from images and text. While existing explainability methods identify influential image...

<https://arxiv.org/abs/2608.00076>

80 CONF

LLM-OSDA: An Optimal-Stopping Dynamic Auction for Native Advertising in Multi-Turn LLM Conversations

ArXiv CS.AI 50%

arXiv:2608.00123v2 Announce Type: replace-cross Abstract: LLM-native advertising embeds sponsored content directly into model-generated responses, shifting the unit of sale from a fixed slot to a moment within an evolving conversation. Existing LLM ad-auction mechanisms...

<https://arxiv.org/abs/2608.00123>

81 CONF

Optimising for Flourishing: Flourishing Metrics and Return on Flourishing as Success Criteria for Artificial Intelligence and Post-AGI Economic Systems

ArXiv CS.AI 50%

arXiv:2608.00151v2 Announce Type: replace-cross Abstract: Current evaluation frameworks for artificial intelligence focus mainly on capability, safety, and proxies such as adoption, engagement, efficiency, productivity, and financial return. These criteria are necessary but...

<https://arxiv.org/abs/2608.00151>

82 CONF

ACE-GraphRAG: Agentic Context Engineering for Hierarchical GraphRAG

ArXiv CS.AI 50%

arXiv:2608.01269v2 Announce Type: replace-cross Abstract: Hierarchical Graph Retrieval-Augmented Generation (GraphRAG) organizes corpus knowledge at multiple levels of granularity, yet fixed context construction may fail to translate these multi-resolution representations into a...

<https://arxiv.org/abs/2608.01269>

83 CONF

Ranking Image Fusion the Way Humans Do: A Learned Pairwise Preference Metric for Infrared-Visible Fusion Assessment

ArXiv CS.AI 50%

arXiv:2608.01301v2 Announce Type: replace-cross Abstract: Infrared-visible image fusion (IVIF) has no ideal fused reference, so fusion algorithms are routinely ranked by scalar objective metrics that formalize different proxies for information transfer, structure, or source...

<https://arxiv.org/abs/2608.01301>

84 **CONF****Style Wins, Substance Loses: A Diagnosis of LLM-as-Judge in Idea Generation****ArXiv CS.AI 50%**

arXiv:2608.01666v2 Announce Type: replace-cross Abstract: However, whether these judges truly evaluate the scientific substance of ideas or are influenced by superficial stylistic presentation remains an open question. To address this question, we propose SciStyleBench, a...

<https://arxiv.org/abs/2608.01666>

85 **CONF****PhyCheck: Fine-Grained Evidence-Grounded Dataset for Physical Law Understanding in Video-LLMs****ArXiv CS.AI 50%**

arXiv:2608.02150v2 Announce Type: replace-cross Abstract: Embodied intelligence and world models require video understanding systems to go beyond recognizing objects and actions and develop an understanding of physical regularities. However, despite their strong performance on...

<https://arxiv.org/abs/2608.02150>