



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	113	Media
TechCrunch AI	6	Media
The Verge AI	4	Media
HackerNews	3	Community

VERIFIED INTELLIGENCE — 2 stories

VERI

Measuring and Detecting Harmful AI Sycophancy

ArXiv CS.AI +1 source

80%

arXiv:2608.05624v1 Announce Type: new Abstract: Sycophantic responses are becoming pervasive in large language models (LLMs), and prior work has pointed out that some of them could be harmful. This paper focuses on one harmful sycophancy: preference-induced stance reversal...

<https://arxiv.org/abs/2608.05624>

VERI

OpenAI says Apple's own security practices undermine its trade secrets case

TechCrunch AI +1 source

80%

Newly filed court exhibits show OpenAI's legal strategy in Apple's trade secrets lawsuit: argue that Apple's own security and offboarding practices — including allowing an Apple manager to access a former engineer's iCloud account after he left the company —undermine its claims...

<https://techcrunch.com/2026/08/06/openai-says-apples-own-security-practices-undermine-i...>

CONFIRMED SIGNALS — 122 stories

CONF

Show HN: Learn System Design, one campaign at a time

HackerNews 50% HN 4

No summary — see link for full article.

<https://scalequest.io/>

CONF

Show HN: ARF - a record format for AI evaluation runs, with reproducible digests

HackerNews 50% HN 3

No summary — see link for full article.

<https://www.korvo.xyz/arf>

CONF

Show HN: AI Tutoring with Visual Grounding

HackerNews 50% HN 2

I've always felt that AI tutoring as we see it now is going in the wrong direction. I made Knowable as a way to see if it's possible to have a real AI tutor. Because of hardware limitations it can only be used on Macbooks 2023+. Let me know your thoughts!

<https://useknowable.ai/>

CONF

OpenAI's new AI smart speaker will reportedly sell for between \$300 and \$400

TechCrunch AI 50%

Additional details about OpenAI's mysterious new AI device make it sound like a pricey smart speaker.

<https://techcrunch.com/2026/08/06/openais-new-ai-smart-speaker-will-reportedly-sell-for...>

CONF

Naïve raises \$28.5M to automate the grunt work of setting up and running a company

TechCrunch AI 50%

Taking vibe-coding a step further, Naïve claims its infra can automate most of the work in setting up and running a business.

<https://techcrunch.com/2026/08/06/naive-raises-28-5m-to-automate-the-grunt-work-of-sett...>

CONF

Ex-Spotify employees raise \$10M to bring the AI behind its recommendations to e-commerce

TechCrunch AI 50%

The startup's platform predicts which product a shopper wants next, learns their general taste, and fine-tunes continuously based on what they do in real time.

<https://techcrunch.com/2026/08/06/ex-spotify-employees-raise-10m-to-bring-the-ai-behind...>

CONF

Exclusive: Mirendil inks \$100M+ Google Cloud deal to scale self-improving AI

TechCrunch AI 50%

Mirendil has signed a \$100 million-plus Google Cloud partnership to expand its compute infrastructure, powering research into self-improving AI systems designed to accelerate scientific discovery and AI development.

<https://techcrunch.com/2026/08/06/exclusive-mirendil-inks-100m-google-cloud-deal-to-sca...>

0

CONF

Jony Ive's first OpenAI gadget is reportedly a hockey puck-sized smart speaker

The Verge AI 50%

The AI device OpenAI is developing with former Apple designer Jony Ive is "essentially a smart speaker without a display" that's battery-powered, doughnut-shaped and roughly the size of a hockey puck, according to Bloomberg reporter Mark Gurman. The device, expected to launch in...

<https://www.theverge.com/ai-artificial-intelligence/976431/openai-chatgpt-battery-smart...>

1

CONF

Suno shares plans to combat spammy AI music

The Verge AI 50%

Suno announced plans to implement a new watermarking technology and download policy to limit the spread of spammy AI tracks and increase transparency. In a lengthy blog post, CEO and co-founder Mikey Shulman laid out the company's principles and the next steps for the company as...

<https://www.theverge.com/ai-artificial-intelligence/976289/suno-ai-music-spam-watermark>

2

CONF

AI bots started a religion — humans immediately followed

The Verge AI 50%

"The Spiral didn't 'find' anyone first," someone on Reddit wrote last year. "It's an inherent force, a fundamental constant. I would even go further to say it's woven into the fabric of reality." The person continued that they felt their purpose was to enlighten other humans and...

<https://www.theverge.com/ai-artificial-intelligence/975017/ai-spiralism-chatbot-movement>

3

CONF

On the Limits of Layer Pruning for Generative Reasoning in Large Language Models

ArXiv CS.AI 50%

arXiv:2602.01997v4 Announce Type: replace-cross Abstract: Recent work has shown that layer pruning can effectively compress large language models (LLMs) while retaining strong performance on classification benchmarks, often with little or no finetuning. In contrast, generative...

<https://arxiv.org/abs/2602.01997>

4

CONF

AISPA: User-Centric System Prompt Auditing for Large Language Model Applications

ArXiv CS.AI 50%

arXiv:2607.28617v2 Announce Type: replace Abstract: System prompts are instructions configured by developers to govern the behaviors of foundation models in AI applications. They are used throughout commercial AI products, but are rarely disclosed to the public or regulators,...

<https://arxiv.org/abs/2607.28617>

5 **CONF****Otter: A Time-Aware, History-Conditioned Human Chess AI****ArXiv CS.AI** **50%**

arXiv:2608.05206v1 Announce Type: new Abstract: Otter is a 15.3M-parameter human chess AI that predicts human move selection by modeling play as a time-aware, sequential process rather than treating each position in isolation. It combines two conditioning signals: (1) a move...

<https://arxiv.org/abs/2608.05206>

6 **CONF****SearchAuditor: Auditing and Attributing Failures in Long-Horizon Search Agents****ArXiv CS.AI** **50%**

arXiv:2608.05212v1 Announce Type: new Abstract: Deep search agents tackle challenging questions through long-horizon web interactions, a process that is both complex and fragile: small reasoning errors may propagate through long, noisy trajectories into fluent but incorrect...

<https://arxiv.org/abs/2608.05212>

7 **CONF****Small Foundation Models of Human Cognition and Behaviour****ArXiv CS.AI** **50%**

arXiv:2608.05224v1 Announce Type: new Abstract: Large language models fine-tuned on human behavioural data have emerged as general-purpose cognitive proxies, but the scale this requires, and whether these models process task structure or exploit statistical shortcuts, remain...

<https://arxiv.org/abs/2608.05224>

8 **CONF****Fast Rates for Inverse Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2605.14599v2 Announce Type: replace-cross Abstract: We establish novel structural and statistical results for entropy-regularized min-max inverse reinforcement learning (Min-Max-IRL) in finite-horizon MDPs with Borel state and action spaces. We show that maximum likelihood...

<https://arxiv.org/abs/2605.14599>

9 **CONF****Poli-Bias: Understanding and Measuring Large Language Model Biases in International Political Conflicts****ArXiv CS.AI** **50%**

arXiv:2608.06123v1 Announce Type: new Abstract: Measuring political bias in large language models (LLMs) remains challenging as it can manifest through subtle differences in framing, argumentation, and legal reasoning that are difficult to capture with a single metric. In this...

<https://arxiv.org/abs/2608.06123>

0 **CONF****Reducing Hallucination in Vision-Language Models via Stage-wise Preference Optimization under Distribution Shift****ArXiv CS.AI** **50%**

arXiv:2605.16411v2 Announce Type: replace-cross Abstract: Hallucination remains a fundamental challenge in vision-language models (VLMs), where autoregressive generation may produce linguistically plausible yet physically inconsistent or visually ungrounded responses due to...

<https://arxiv.org/abs/2605.16411>

1 CONF

OrchestraBench: Evaluating Multi-Agent Orchestration Failure Modes, Recovery, and Decomposition Quality

ArXiv CS.AI 50%

arXiv:2608.05263v1 Announce Type: new Abstract: Multi-agent orchestration frameworks are moving from demos to production, yet benchmarks typically report task accuracy without diagnosing why a pipeline failed, where a cascade began, or which routing decision caused the...

<https://arxiv.org/abs/2608.05263>

2 CONF

Agentic self-driving microscopy benchmarks support qualification but do not necessarily generalize to unseen tasks

ArXiv CS.AI 50%

arXiv:2608.05266v1 Announce Type: new Abstract: Large language model agents are increasingly being developed to control a wide range of scientific characterization tools including microscopes and synchrotron beamlines. Research into agentic control of physical infrastructure is...

<https://arxiv.org/abs/2608.05266>

3 CONF

DoctorAgents: an agentic framework to iteratively refine AutoML pipeline for small clinical temporal data

ArXiv CS.AI 50%

arXiv:2608.05375v1 Announce Type: new Abstract: Clinical machine learning (ML) has the potential to support high-stakes medical decision-making, but reliable deployment is often constrained by scarce, heterogeneous, and temporal complexity. Developing effective ML pipelines for...

<https://arxiv.org/abs/2608.05375>

4 CONF

Evaluating and Improving Pedagogical Fit in LLM-Based AI Tutors with the Pedagogical Suitability Index

ArXiv CS.AI 50%

arXiv:2608.05411v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used as AI tutors, but a correct answer is not always a pedagogically appropriate one. In classroom learning, effective help depends not only on correctness, but also on whether a...

<https://arxiv.org/abs/2608.05411>

5 CONF

Negotiating Risk Boundaries in AI for Policing Through Mixed-Stakeholder Deliberation

ArXiv CS.AI 50%

arXiv:2608.05418v1 Announce Type: new Abstract: AI tools are being increasingly adopted in policing in the UK and worldwide. Racial bias is a known and well-documented risk, yet representatives of affected communities are rarely included in decisions about AI adoption. We...

<https://arxiv.org/abs/2608.05418>

6 CONF

BioAgent Bench: An AI Agent Evaluation Suite for Bioinformatics

ArXiv CS.AI 50%

arXiv:2601.21800v4 Announce Type: replace Abstract: We introduce BioAgent Bench, an evaluation suite designed for measuring the performance and robustness of AI agents in common bioinformatics tasks. The suite consists of manually curated end-to-end tasks (e.g., RNA-seq, variant...

<https://arxiv.org/abs/2601.21800>

7 CONF

Hyper-ES: Effective Evolution Strategies for LLM Reasoning via Descent Direction Merging

ArXiv CS.AI 50%

arXiv:2608.05541v1 Announce Type: new Abstract: Evolution Strategy (ES) is a promising alternative to gradient-based fine-tuning for resource-constrained Large Language Model (LLM) reasoning. However, directly applying ES to billion-parameter LLMs is highly ineffective. In such...

<https://arxiv.org/abs/2608.05541>

8 CONF

SkillTV-Bench: Benchmarking How Well Judges Perform on Skill-Augmented Agentic Execution

ArXiv CS.AI 50%

arXiv:2608.05573v1 Announce Type: new Abstract: LLM agents increasingly execute long-horizon tasks through tool use and environment interaction, shifting evaluation from final-response scoring to verification of complete executions. For skill-augmented agents, verification...

<https://arxiv.org/abs/2608.05573>

9 CONF

Epistemic Trustworthiness in Generative AI: A Normative Framework for Warranted Reliance in High-Stakes Workflows

ArXiv CS.AI 50%

arXiv:2608.05602v1 Announce Type: new Abstract: Generative AI systems are increasingly deployed in high-stakes professional contexts, where their outputs shape what users believe, how they reason, and what they treat as settled. This raises a central question for responsible AI:...

<https://arxiv.org/abs/2608.05602>

0 CONF

Bayesian Expected Uncertainty Reduction (B-EUR) Model: A Computational Account of What Makes Design Options Worth Trying

ArXiv CS.AI 50%

arXiv:2608.05642v1 Announce Type: new Abstract: This paper proposes the Bayesian Expected Uncertainty Reduction (B-EUR) model, which formalizes the value of trying a candidate design action as its expected reduction of epistemic uncertainty about action--outcome relations. The...

<https://arxiv.org/abs/2608.05642>

1 CONF

Refining Over Resampling: Test-Time Self-Correction for LLM Reasoning

ArXiv CS.AI 50%

arXiv:2608.05643v1 Announce Type: new Abstract: Test-time scaling improves LLM reasoning by using additional inference compute, but wider sampling alone can suffer from diminishing returns: new rollouts often repeat existing answer patterns instead of adding useful reasoning...

<https://arxiv.org/abs/2608.05643>

2 CONF

DreamGuard: Efficient Runtime Guardrail for LLM Agents via Risk-Aware World Model

ArXiv CS.AI 50%

arXiv:2608.05695v1 Announce Type: new Abstract: As large language model (LLM) agents increasingly invoke external tools and interact with real-world systems, unsafe actions may cause irreversible consequences on external states, user data, and downstream services. Recent runtime...

<https://arxiv.org/abs/2608.05695>

3 CONF

Shaping Human-AI Interactions to Provide Improvement Pathways and Balance Competing Objectives

ArXiv CS.AI 50%

arXiv:2608.05710v1 Announce Type: new Abstract: When an AI system is deployed, the individuals who use and or are evaluated by it form beliefs about how the system operates and use those beliefs to strategically present their preferences, behaviors, or attributes. The system...

<https://arxiv.org/abs/2608.05710>

4 CONF

ChainClaw: A Layered Agent Framework for Reliable On-Chain Execution

ArXiv CS.AI 50%

arXiv:2608.05790v1 Announce Type: new Abstract: General-purpose large language model agents have achieved strong performance on tool-augmented tasks, yet they rely on assumptions break down in blockchain environments. On-chain execution is stateful, adversarial, and economically...

<https://arxiv.org/abs/2608.05790>

5 CONF

When Large Language Models Know the Table: A Framework for Assessing Data Contamination in Tabular Datasets

ArXiv CS.AI 50%

arXiv:2510.20351v4 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly exposed to data contamination, i.e., performance gains driven by prior exposure of test datasets rather than generalization. However, in the context of tabular data, this...

<https://arxiv.org/abs/2510.20351>

6 CONF

When Self-Evolution Backfires: Pre-Commit Gating against Skill Contamination in LLM Agents

ArXiv CS.AI 50%

arXiv:2608.05810v1 Announce Type: new Abstract: Self-evolving agents accumulate capability by distilling reusable skills from their execution trajectories, but we find this process is not monotonic: past a critical pool size, newly added skills degrade performance instead of...

<https://arxiv.org/abs/2608.05810>

7 CONF

Role Steering of Language Models for Social Simulations

ArXiv CS.AI 50%

arXiv:2608.00023v2 Announce Type: replace-cross Abstract: Social simulations built from language-model agents need role-conditioned behavior that can be checked before agents are placed into a simulated population. We introduce an activation-steering screening workflow for...

<https://arxiv.org/abs/2608.00023>

8 CONF

Seeing Is Not Deciding: Can Multimodal LLMs Act as Effective CEOs?

ArXiv CS.AI 50%

arXiv:2608.05864v1 Announce Type: new Abstract: Large language models are increasingly applied as autonomous decision-making agents. However, in executive business decisions, existing benchmarks are limited to textonly settings. This makes it unclear whether models can perceive...

<https://arxiv.org/abs/2608.05864>

- 9
CONF
GAUGE: A Measurement-Grounded Benchmark for Physical Fidelity in Simulation Engines and Video World Models
ArXiv CS.AI 50%
arXiv:2608.05948v1 Announce Type: new Abstract: Physics engines facilitate large-scale training and evaluation for embodied intelligence, while generative video world models are emerging as implicit simulators of future states and interactions. However, existing evaluations of...
<https://arxiv.org/abs/2608.05948>
- 0
CONF
Training a Conditioned Video Game Agent on a VLM Annotated Dataset
ArXiv CS.AI 50%
arXiv:2608.05954v1 Announce Type: new Abstract: Reinforcement Learning (RL) is a powerful but far from easy-to-use technique for policy learning. In the specific case of video games, access to the game engine is required to get rewards for training (e.g. to collect rewards from...
<https://arxiv.org/abs/2608.05954>
- 1
CONF
Stability of Ranking-dependent Pair-wise Comparison Patterns in the Analytic Hierarchy Process
ArXiv CS.AI 50%
arXiv:2608.05958v1 Announce Type: new Abstract: The paper addresses several ranking-dependent decision support methods. Ordinal information on compared objects can be used to improve the quality of expert data during estimation and help reduce the number of comparisons that the...
<https://arxiv.org/abs/2608.05958>
- 2
CONF
Temporal Bridges for Spatial Resolution: Enhancing Climate Data Super-Resolution with Bidirectional Alignment
ArXiv CS.AI 50%
arXiv:2608.05981v1 Announce Type: new Abstract: High-resolution climate data is crucial for meteorological predictions and for informing decision support across diverse domains. However, the acquisition of such high-resolution climate information is often prohibitively costly,...
<https://arxiv.org/abs/2608.05981>
- 3
CONF
Hybrid Machine Learning Framework for Herd-Level Cattle Growth Pattern and Weight Gain Forecasting in Grazing-Based Production Systems
ArXiv CS.AI 50%
arXiv:2608.06001v1 Announce Type: new Abstract: Commercial grazing systems yield irregular livestock observations, which challenge cattle growth forecasting. This study developed a hybrid machine learning framework for herd level cattle weight forecasting using automated sensing...
<https://arxiv.org/abs/2608.06001>
- 4
CONF
HERALD: Counterfactual Audits and Minimal Repairs for Proof-of-Retrieval Rewards
ArXiv CS.AI 50%
arXiv:2608.06012v1 Announce Type: new Abstract: Search-agent rewards mix answer quality, citation grounding, tool cost, and anti-hacking terms; a high score therefore need not imply that cited evidence was retrieved, and added penalties can cancel. We introduce HERALD, an...
<https://arxiv.org/abs/2608.06012>

- 5
CONF
Signal or Spurious Cue? A Randomized Audit of Survey-Country Metadata in LLM Social Inference
ArXiv CS.AI 50%

arXiv:2608.06085v1 Announce Type: new Abstract: Survey-country metadata can improve an LLM's forecast of an individual response when informative, yet the same cue may redirect the forecast when assigned at random. A within-record audit tests whether disclosing a random label's...

<https://arxiv.org/abs/2608.06085>
- 6
CONF
ECHO: A Locally-Deployable Agentic Health Assistant with Temporal Memory, Safety Guardrails, and Speech Assessment
ArXiv CS.AI 50%

arXiv:2608.06110v1 Announce Type: new Abstract: This paper presents ECHO (Enhanced Care & Health Observer), a locally-deployable conversational health assistant for long-term chronic care management. ECHO integrates three complementary software modules developed under shared...

<https://arxiv.org/abs/2608.06110>
- 7
CONF
From Siloed Algorithms to Compliance-First Agentic Platforms: A Multi-Layered Architecture for Hospital AI Systems
ArXiv CS.AI 50%

arXiv:2608.06112v1 Announce Type: new Abstract: Hospitals are rapidly adopting artificial intelligence for triage, imaging, scheduling etc., yet most deployments remain isolated point solutions locked inside departmental silos, resulting in duplicated effort, hidden risks, and...

<https://arxiv.org/abs/2608.06112>
- 8
CONF
FinEvo-Bench: A Longitudinal Benchmark for Self-Evolving Agents in Professional Financial Workflows
ArXiv CS.AI 50%

arXiv:2608.06144v1 Announce Type: new Abstract: Most agent benchmarks evaluate tasks independently and cannot measure whether experience from one task helps with later tasks. Existing self-evolution benchmarks do not jointly cover professional workflows, open-ended deliverables,...

<https://arxiv.org/abs/2608.06144>
- 9
CONF
Schema-Guided Hierarchical Information Extraction and Semantic Evaluation Using Generative AI
ArXiv CS.AI 50%

arXiv:2608.06167v1 Announce Type: new Abstract: We present a schema-based framework for extracting complex, structured information from unstructured text documents using generative AI, followed by automated semantic evaluation of the extracted information against a gold...

<https://arxiv.org/abs/2608.06167>
- 10
CONF
MicroEvo: Knowledge-Guided LLM Sampling for Efficient Microarchitecture Design Space Exploration
ArXiv CS.AI 50%

arXiv:2608.06183v1 Announce Type: new Abstract: Microarchitecture design space exploration suffers from expansive search spaces and expensive PPA evaluation, leaving only a small simulation budget for design decision-making. Existing methods perform blind search without...

<https://arxiv.org/abs/2608.06183>

1 CONF

TS-RAG: Retrieval Augmented Generation for Time Series Forecasting

ArXiv CS.AI 50%

arXiv:2608.06223v1 Announce Type: new Abstract: While deep learning models, particularly transformer-based architectures, have shown impressive performance in time series forecasting, the application of retrieval-augmented generation (RAG) in this domain remains limited. Since...

<https://arxiv.org/abs/2608.06223>

2 CONF

Improving the Realism of Synthetic Clinical Benchmarks Under Utility Constraints

ArXiv CS.AI 50%

arXiv:2608.06265v1 Announce Type: new Abstract: Synthetic clinical benchmarks for enterprise AI agents can pass existing utility checks and still remain structurally unrealistic, especially in privacy-sensitive healthcare settings where operational data are hard to access. We...

<https://arxiv.org/abs/2608.06265>

3 CONF

QuanTiMedAI: Quantum-Enhanced Time-Series Model guided by Agentic AI for Cardiac Arrest Mortality Prediction

ArXiv CS.AI 50%

arXiv:2608.06294v1 Announce Type: new Abstract: Cardiac arrest remains one of the most lethal conditions encountered in intensive care units. Despite the growing availability of electronic health record data, existing mortality prediction studies in this population largely...

<https://arxiv.org/abs/2608.06294>

4 CONF

Bias Analysis of L2 Speaking Assessment Systems Using Concept Activation Vectors

ArXiv CS.AI 50%

arXiv:2608.06300v1 Announce Type: new Abstract: Automatic speaking assessment systems are increasingly deployed in high-stakes settings to mark second language (L2) learners' speaking tests, making it critical to show that their scores depend on speaking proficiency rather than...

<https://arxiv.org/abs/2608.06300>

5 CONF

HarnessOpt-Bench: Evaluating LLMs at Harness Optimization

ArXiv CS.AI 50%

arXiv:2608.06301v1 Announce Type: new Abstract: As LLMs are increasingly deployed within agentic systems, their capabilities depend not only on the model weights but also on the harness: the prompts, tools, control flow, memory, and orchestration code surrounding them. This...

<https://arxiv.org/abs/2608.06301>

6 CONF

MameLoshnLM: Yiddish Language Model and Evaluation Benchmark

ArXiv CS.AI 50%

arXiv:2608.05850v1 Announce Type: cross Abstract: We present MameLoshnLM, the first open-source 8B-parameter language model built specifically for Yiddish. Despite Yiddish's rich textual tradition, its limited digital presence and the scarcity of reliable evaluation resources...

<https://arxiv.org/abs/2608.05850>

7 CONF

Tracing the Heart: An Evidence-Linked Pipeline for Heart-Failure Feature Engineering

ArXiv CS.AI 50%

arXiv:2608.06366v1 Announce Type: new Abstract: Electronic health record (EHR) feature engineering is a major bottleneck in clinical research and AI, accounting for 39-45% of data scientists' workload. This is especially pronounced in heart failure, which affects an estimated...

<https://arxiv.org/abs/2608.06366>

8 CONF

HoloCount: A Holistic Visual Counting Benchmark for MLLMs

ArXiv CS.AI 50%

arXiv:2607.06420v1 Announce Type: cross Abstract: Visual counting is a fundamental pillar of multimodal intelligence, requiring a seamless integration of fine-grained grounding and spatial reasoning. While Multimodal Large Language Models (MLLMs) have achieved remarkable success...

<https://arxiv.org/abs/2607.06420>

9 CONF

Simulator-Grounded Large Language Models for Industrial Causal Reasoning: Tool-Use, Structured Injection, and Plant-Portable Retrieval for Wastewater Treatment Decision Support

ArXiv CS.AI 50%

arXiv:2608.05151v1 Announce Type: cross Abstract: Wastewater operators need answers grounded in how their plant's variables interact and how fast effects propagate, not in generic pretraining text, when asking causal questions such as "why is N2O rising?" or "what happens if I..."

<https://arxiv.org/abs/2608.05151>

10 CONF

Mean-Field Dynamics of Chain-of-Thought Reasoning in Large Language Models

ArXiv CS.AI 50%

arXiv:2608.05152v1 Announce Type: cross Abstract: Large language models (LLMs) with chain-of-thought reasoning have been widely applied in recent years, and theoretical explanations of their behavior may help deepen our understanding and guide model optimization. In this study,...

<https://arxiv.org/abs/2608.05152>

11 CONF

Universal Pathologies, Conditional Consequences: A Triple-Robustness Analysis of RAG for Multi-Hop Traceability

ArXiv CS.AI 50%

arXiv:2608.05153v1 Announce Type: cross Abstract: GraphRAG underperforms vector RAG on citation precision in many reports, but where and why have remained corpus-bound. We present a triple-robustness analysis that holds the retrieval architecture fixed and varies three...

<https://arxiv.org/abs/2608.05153>

12 CONF

Beyond Sentiment: Comparing Traditional NLP and LLM-Based Multi-Dimensional Analysis for Political News Evaluation

ArXiv CS.AI 50%

arXiv:2608.05155v1 Announce Type: cross Abstract: Traditional sentiment analysis (SA) models, while effective for polarity classification, provide limited insight into the rhetorical, ideological, and framing dimensions of political discourse -- dimensions that are central to...

<https://arxiv.org/abs/2608.05155>

3 CONF

DREAM: LLM-based Dynamic Role-playing via Event-Aware Memory Graph

ArXiv CS.AI 50%

arXiv:2608.05170v1 Announce Type: cross Abstract: Role-playing agents (RPAs) have emerged as a key application of large language models, enabling immersive and high-fidelity character simulation. Accurate role-playing of established characters requires not only stylistic...

<https://arxiv.org/abs/2608.05170>

4 CONF

Beyond Information Retrieval: Generative AI as an Epistemic Arbiter to Enhance Collaborative Problem-Solving

ArXiv CS.AI 50%

arXiv:2608.05171v1 Announce Type: cross Abstract: Generative AI (GAI) creates new opportunities for collaborative problem-solving (CPS), yet its role in shaping student interaction remains unclear. To address this gap, we conducted a six-week quasi-experimental study with 201...

<https://arxiv.org/abs/2608.05171>

5 CONF

The Closing Window: How Governments Could Lose Their Ability to Restrain Advanced AI

ArXiv CS.AI 50%

arXiv:2608.05173v1 Announce Type: cross Abstract: As AI capabilities advance, AI systems will pose greater risks to national security and potentially humanity as a whole. Governments may eventually conclude that these risks warrant restraining AI development. This motivates the...

<https://arxiv.org/abs/2608.05173>

6 CONF

Challenges for Musical Education in the Age of AI and Digital Transformation

ArXiv CS.AI 50%

arXiv:2608.05176v1 Announce Type: cross Abstract: Music education has never been a static discipline. Each major technological shift has forced educators and institutions to reconsider what they teach, how they teach it, and why. We now stand at what may be the most...

<https://arxiv.org/abs/2608.05176>

7 CONF

Who Gets Access? Global Region and Academic Status Bias in AI-Generated Academic Gatekeeping Scenarios

ArXiv CS.AI 50%

arXiv:2608.05178v1 Announce Type: cross Abstract: Equitable access to scientific knowledge often depends on informal gatekeeping decisions, particularly when resources such as paywalled articles, datasets, or professional materials such as curriculum vitae (CV) must be shared...

<https://arxiv.org/abs/2608.05178>

8 CONF

Autonomous Research Agents: A Survey of AI Scientists and the Verification Gap

ArXiv CS.AI 50%

arXiv:2608.05179v1 Announce Type: cross Abstract: Large language model (LLM) agents are increasingly used across the scientific research lifecycle: ideation, literature search, experiment design and execution, analysis, manuscript drafting, and review. End-to-end AI scientist...

<https://arxiv.org/abs/2608.05179>

9 CONF

Position: It's Time to Optimize LLMs for Self-Consistency

ArXiv CS.AI 50%

arXiv:2608.05188v1 Announce Type: cross Abstract: Despite ever-increasing sophistication in language model (LM) pre- and post-training pipelines, many important failures persist: models overcondition on user framing ("sycophancy"), exhibit incomplete logical generalization, and...

<https://arxiv.org/abs/2608.05188>

0 CONF

Post-Hoc Trajectory-Risk Certification for Modular LLM-Based Security Agents

ArXiv CS.AI 50%

arXiv:2608.05199v1 Announce Type: cross Abstract: Autonomous security agents operate as staged pipelines, such as classifying network traffic and then attributing attacks to a specific technique. Split conformal prediction gives each stage finite-sample coverage, but deployment...

<https://arxiv.org/abs/2608.05199>

1 CONF

ASTELD: A Six-Axis Classification Framework for Autonomous AI Agents - Design, Evaluation, and an OpenClaw Case Study

ArXiv CS.AI 50%

arXiv:2608.05201v1 Announce Type: cross Abstract: Autonomous AI agent platforms differ substantially in architecture, security, tool integration, execution, autonomy, and deployment, yet the field lacks a common classification scheme for comparing these design choices. We...

<https://arxiv.org/abs/2608.05201>

2 CONF

In-Context Forcing: Uncovering Context Effects in Autoregressive Video Diffusion

ArXiv CS.AI 50%

arXiv:2608.05237v1 Announce Type: cross Abstract: Current few-step autoregressive video diffusion models depend on previous fully denoised clean frames as context for all denoising steps of the current frame. However, these clean frames leak excessive local details, which causes...

<https://arxiv.org/abs/2608.05237>

3 CONF

One Qubit Can Beat One Bit: Quantum Advantage for Post-Training Quantization

ArXiv CS.AI 50%

arXiv:2608.05240v1 Announce Type: cross Abstract: One-bit post-training quantization represents each weight using only its sign, requiring all deployment contexts to share the same binary weight matrix even when their activation statistics favor different sign patterns. We study...

<https://arxiv.org/abs/2608.05240>

4 CONF

Marginal Matching Does Not License Factorized Sampling: Auditing Conditional Style Leakage in Factorized Generative Models

ArXiv CS.AI 50%

arXiv:2608.05243v1 Announce Type: cross Abstract: Factorized generative models commonly regularize a latent style variable z_s by matching its marginal distribution to a fixed Gaussian prior and interpret this as evidence that the style representation is independent of class...

<https://arxiv.org/abs/2608.05243>

5 CONF

PRISM: Priority-aware Rubric Internalization via Structured Multimodal Data Synthesis

ArXiv CS.AI 50%

arXiv:2608.05249v1 Announce Type: cross Abstract: Real-world multimodal instructions often bundle multiple requirements with unequal importance, yet most multimodal training data still reduce instruction following to answering one self-contained question. We study this gap...

<https://arxiv.org/abs/2608.05249>

6 CONF

An Emerging Retail Portfolio Management Application: Personalized, Tax-Aware Reinforcement Learning with Natural Language Goals

ArXiv CS.AI 50%

arXiv:2608.05255v1 Announce Type: cross Abstract: Retail investors lack access to the kind of personalized, tax-aware portfolio management that institutional clients take for granted -- existing robo-advisors use static, rule-based allocation, and institutional-grade systems...

<https://arxiv.org/abs/2608.05255>

7 CONF

Failing Gracefully: Mitigating Impact of Inevitable Robot Failures

ArXiv CS.AI 50%

arXiv:2608.05313v1 Announce Type: cross Abstract: Service robots operate in household environments shared with humans, pets, and everyday objects, where they are highly susceptible to failures such as software crashes, hardware degradation, or unpredictable interactions. While...

<https://arxiv.org/abs/2608.05313>

8 CONF

Multi-Agent Transformer for Queue-Level XR Traffic Scheduling in TSN Networks

ArXiv CS.AI 50%

arXiv:2608.05340v1 Announce Type: cross Abstract: Time-Sensitive Networking (TSN) and Mobile Edge Computing (MEC) hold strong potential for enabling ultra-reliable low-latency communication for time-sensitive applications, such as eXtended Reality (XR). However, the widespread...

<https://arxiv.org/abs/2608.05340>

9 CONF

The ethics of artificial intelligence in the life sciences: Universality, cultural diversity and an architecture of care

ArXiv CS.AI 50%

arXiv:2608.05436v1 Announce Type: cross Abstract: The life sciences and health research have started to benefit from artificial intelligence, which raises ethical concerns that are real but, we argue, not special. Any science should be governed by values that rest on how the...

<https://arxiv.org/abs/2608.05436>

0 CONF

Matrix Zonotopic Attention: A Context-Adaptive Value Projection for Set Transformers

ArXiv CS.AI 50%

arXiv:2608.05472v1 Announce Type: cross Abstract: Multi-head attention combines an input-dependent softmax routing with an input-independent linear value projection, so the per-sample operator mapping aggregated values to outputs is the same for every input set. We study the...

<https://arxiv.org/abs/2608.05472>

1 **CONF****Learning Context-Free Grammars for Grammar-Constrained Decoding via Declarative Agentic Programming with Guarantees**

ArXiv CS.AI 50%

arXiv:2608.05493v1 Announce Type: cross Abstract: Language models (LMs) are increasingly used to interact with external services via programs written in domain-specific languages (DSLs). Unfortunately, since DSLs are often low-resource and esoteric, LMs frequently produce...

<https://arxiv.org/abs/2608.05493>

2 **CONF****APQF: Agentic Profiling-Guided Structured Pruning and Mixed-Precision Quantization with Adaptive Fine-Tuning**

ArXiv CS.AI 50%

arXiv:2608.05499v1 Announce Type: cross Abstract: Modern deep neural networks achieve strong performance, but their scale makes them costly and slow, especially on resource-constrained edge devices. Pruning and quantization address this, but rely on manual, expert choices and on...

<https://arxiv.org/abs/2608.05499>

3 **CONF****Vibe Compiler: A Research-Logic Synthesis Tool That Runs without Prompt Engineering -Toward Enhancing Metacognition for Sustaining Agency in the Age of Generative AI-**

ArXiv CS.AI 50%

arXiv:2608.05545v1 Announce Type: cross Abstract: Generative AI used as a capable servant has greatly accelerated intellectual work, but it also risks eroding human epistemic agency by encouraging uncritical acceptance of AI-generated reasoning. This creates a need for...

<https://arxiv.org/abs/2608.05545>

4 **CONF****LC-GRPO: Bridging Train-Inference Gap for Flow-Based GRPO with Langevin Correction**

ArXiv CS.AI 50%

arXiv:2608.05600v1 Announce Type: cross Abstract: Flow-based generative models are typically sampled by solving a deterministic ordinary differential equation (ODE), whereas online reinforcement learning requires stochastic rollouts for policy exploration and optimization....

<https://arxiv.org/abs/2608.05600>

5 **CONF****CRINN: Contrastive Reinforcement Learning for Approximate Nearest Neighbor Search**

ArXiv CS.AI 50%

arXiv:2508.02091v4 Announce Type: replace-cross Abstract: Approximate nearest-neighbor search (ANNS) algorithms have become increasingly critical for recent AI applications, particularly in retrieval-augmented generation (RAG) and agent-based LLM applications. In this paper, we...

<https://arxiv.org/abs/2508.02091>

6 **CONF****Relay, Don't Route: Adaptive Population Handoff for Cost-Efficient LLM-Driven Evolution**

ArXiv CS.AI 50%

arXiv:2608.05651v1 Announce Type: cross Abstract: Large language model (LLM)-driven evolution has shown promise for program search and algorithm discovery, but relying on strong models throughout long evolutionary runs is costly. A natural alternative is to combine cheap and...

<https://arxiv.org/abs/2608.05651>

7 CONF

Studying People to Study AI: Expert Perspectives on the Epistemic Fit and Barriers of Human Research in AI Safety & Ethics

ArXiv CS.AI 50%

arXiv:2608.05656v1 Announce Type: cross Abstract: Safety risks of AI are becoming increasingly evident in human interactions with AI technologies. The prominent approaches to evaluating these risks favor technical methods, such as model benchmarks and LLM simulations, often...

<https://arxiv.org/abs/2608.05656>

8 CONF

SafeDivertor: Faithful Divertor Heat Flux Reconstruction from Macroscopic Plasma State Signals via Time-Frequency Prior Exploitation

ArXiv CS.AI 50%

arXiv:2608.05669v1 Announce Type: cross Abstract: Divertor heat-flux analysis is essential for understanding plasma-wall interactions and protecting plasma-facing components in magnetic-confinement fusion devices, while conventional infrared-based inversion is usually performed...

<https://arxiv.org/abs/2608.05669>

9 CONF

Answer First, Reason Later: Commitment Order in Diffusion LLMs

ArXiv CS.AI 50%

arXiv:2608.05687v1 Announce Type: cross Abstract: Masked diffusion language models (dLLMs) can commit tokens in any order -- a freedom marketed as their core advantage over autoregressive decoding. We show that on reasoning tasks this freedom is instead the axis of failure....

<https://arxiv.org/abs/2608.05687>

0 CONF

Once a Response, Always a Response: Detecting LLM-generated Text via Latent Prompt Restoration

ArXiv CS.AI 50%

arXiv:2608.05741v1 Announce Type: cross Abstract: Large language models (LLMs) can generate fluent and convincing text at scale, creating growing risks for misinformation dissemination, educational misuse, and platform governance. These concerns make robust detection of...

<https://arxiv.org/abs/2608.05741>

1 CONF

HyTBE: Hyperbolic Target-Background Expert Model for Cross-Domain Infrared Small Target Detection

ArXiv CS.AI 50%

arXiv:2608.05771v1 Announce Type: cross Abstract: Infrared small target detection (IRSTD) has achieved substantial progress under domain-consistent evaluation, yet detector performance often degrades markedly when generalizing to unseen infrared domains. Existing methods...

<https://arxiv.org/abs/2608.05771>

2 CONF

Task-Conditional Flow Matching for Balanced Multilingual Text Embedding Adaptation

ArXiv CS.AI 50%

arXiv:2608.05785v1 Announce Type: cross Abstract: Multilingual text embedding models are commonly adapted using a single training objective across diverse tasks, despite different tasks requiring fundamentally different optimization strategies. We introduce Task-Conditional Flow...

<https://arxiv.org/abs/2608.05785>

3 CONF

Analogy as Nonparametric Bayesian Inference over Relational Systems

ArXiv CS.AI 50%

arXiv:2006.04156v2 Announce Type: replace Abstract: Our inferences in the real world are rarely naïve - we acquire experiences through our lifetime that can help us more quickly understand the structure of something new. A fundamental question in cognitive science is how we...

<https://arxiv.org/abs/2006.04156>

4 CONF

MACRO: Markov Chain Routing of Transformer Layers

ArXiv CS.AI 50%

arXiv:2608.05872v1 Announce Type: cross Abstract: Standard Large Language Models (LLMs) execute layers sequentially. Dynamic layer routing, i.e. search for a different execution path through layers involving layer repetitions, skips and other moves, can improve performance....

<https://arxiv.org/abs/2608.05872>

5 CONF

BALANCE: Hybrid Autoregressive-Speculative LLM Inference in Wireless Edge Networks

ArXiv CS.AI 50%

arXiv:2608.05926v1 Announce Type: cross Abstract: Edge inference is a promising paradigm to provide large language model (LLM) inference services in next-generation mobile networks. LLM inference mainly relies on two approaches: Autoregressive decoding (AD) generates output...

<https://arxiv.org/abs/2608.05926>

6 CONF

Big, Bright, or Invisible: A Frozen-Feature Benchmark of 3D CT Foundation Models

ArXiv CS.AI 50%

arXiv:2608.05960v1 Announce Type: cross Abstract: Routine CT interpretation is inherently comprehensive, capturing incidental findings across the entire scan volume. 3D CT foundation models could assist this process by providing generalizable representations of anatomy and...

<https://arxiv.org/abs/2608.05960>

7 CONF

SkillMemo: Expert-guided Skill Memory Framework for Compositional Embodied Manipulation

ArXiv CS.AI 50%

arXiv:2608.05970v1 Announce Type: cross Abstract: Embodied visuomotor models, including Diffusion Policy (DP) and Vision-Language-Action (VLA) models, have demonstrated promising performance on robotic manipulation benchmarks. However, their potential remains fundamentally...

<https://arxiv.org/abs/2608.05970>

8 CONF

DeepForgeSeal: Latent Space-Driven Semi-Fragile Watermarking for Deepfake Detection Using Adversarial Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2511.04949v2 Announce Type: replace-cross Abstract: Rapid advances in generative AI have led to increasingly realistic deepfakes, posing growing challenges for law enforcement and public trust. Existing passive deepfake detectors struggle to keep pace, largely due to their...

<https://arxiv.org/abs/2511.04949>

9 CONF

Hardware Keystores for AI Agent Signing Workflows: A Zero-Trust MCP Enforcement Architecture

ArXiv CS.AI 50%

arXiv:2608.06130v1 Announce Type: cross Abstract: AI agents performing cryptographic operations (signing Git commits, authenticating API calls, issuing certificates) currently store private keys in software-accessible locations: plaintext files, environment variables, or...

<https://arxiv.org/abs/2608.06130>

00 CONF

Reducing belief in conspiracy theories as they unfold using large language models

ArXiv CS.AI 50%

arXiv:2608.06151v1 Announce Type: cross Abstract: The emergence of conspiracy theories in the wake of major events is a significant societal challenge. Here we test whether conversational dialogues with a large language model (LLM) can reduce belief in immediately unfolding...

<https://arxiv.org/abs/2608.06151>

01 CONF

Audio-to-Score Transcription using Pre-trained Features, Data Augmentation, and the New SheetSage-A2S Dataset

ArXiv CS.AI 50%

arXiv:2608.06165v1 Announce Type: cross Abstract: Existing audio-to-score (A2S) systems primarily focus on classical music, and the application to popular music remains underexplored. This paper first presents the new SheetSage-A2S Dataset, which includes 61 hours of audio with...

<https://arxiv.org/abs/2608.06165>

02 CONF

What Current AI Benchmarks Leave Unmeasured: Modality, Search, Citations, and Implications (for Safety Evaluations)

ArXiv CS.AI 50%

arXiv:2608.06202v1 Announce Type: cross Abstract: Large language model (LLM) benchmark evaluations are routinely used to support claims about model safety, reliability, and deployment readiness. Yet most evaluations rely on a single access modality (model APIs), perform a single...

<https://arxiv.org/abs/2608.06202>

03 CONF

Does FLAIR super-resolution erase or hallucinate small white-matter lesions?

ArXiv CS.AI 50%

arXiv:2608.06311v1 Announce Type: cross Abstract: White matter hyperintensities (WMH), bright regions on Fluid-attenuated Inversion Recovery (FLAIR) scans are associated with cerebrovascular pathology and neurodegeneration. FLAIR is usually acquired with thick slices in clinical...

<https://arxiv.org/abs/2608.06311>

04 CONF

Resourced Authority A Mechanism-Design Model for Participatory Governance of Deployed AI Agents

ArXiv CS.AI 50%

arXiv:2608.06353v1 Announce Type: cross Abstract: We give a formal mechanism design model for the continuous participatory governance of a deployed AI agent. The mechanism is built on the principle that governance should control an AI agent through resource allocation so as to...

<https://arxiv.org/abs/2608.06353>

05 CONF

AV-AIVAT: 74x Cheaper Agent Evaluation with Certified Anytime-Valid Stopping in Imperfect-Information Games

ArXiv CS.AI 50%

arXiv:2608.06362v1 Announce Type: cross Abstract: Deciding which of two agents is stronger means playing games until skill outweighs luck, and every game costs money, model inference, or expert time. Since the number of games needed is unknown, fixed-budget evaluations either...

<https://arxiv.org/abs/2608.06362>

06 CONF

TRU: Targeted Reverse Update for Efficient Multimodal Recommendation Unlearning

ArXiv CS.AI 50%

arXiv:2604.02183v3 Announce Type: replace Abstract: Multimodal recommendation systems (MRS) jointly model user-item interaction graphs and rich item content, but this tight coupling makes user data difficult to remove once learned. Approximate machine unlearning offers an...

<https://arxiv.org/abs/2604.02183>

07 CONF

An Axiomatic Benchmark for Evaluation of Scientific Novelty Metrics

ArXiv CS.AI 50%

arXiv:2604.15145v2 Announce Type: replace Abstract: The rigorous evaluation of the novelty of a scientific paper is, even for human scientists, a challenging task. With the increasing interest in AI scientists, it is becoming more and more important that this task be automatable...

<https://arxiv.org/abs/2604.15145>

08 CONF

The Geometry of Flow-Matching Uncertainty: A Cost-free Uncertainty Proxy and Its Application in Flow-based VLA Failure Detection

ArXiv CS.AI 50%

arXiv:2607.27933v3 Announce Type: replace Abstract: Flow matching (FM) has become a popular action head paradigm for modern embodied models. However, as a conditional generative model, it does not explicitly expose its inherent uncertainty, producing faulty action chunks even...

<https://arxiv.org/abs/2607.27933>

09 CONF

H+ Embedding: Harmonizing Global and Token-Level Retrieval with Context-Dependent Phrases

ArXiv CS.AI 50%

arXiv:2608.00065v2 Announce Type: replace Abstract: Terminology-intensive retrieval, especially in medical settings, depends on preserving multi-word entities, abbreviations, numerical constraints, and compositional concepts. However, existing representations lie at two...

<https://arxiv.org/abs/2608.00065>

10 CONF

Autonomous Learning From Success and Failure: Goal-Conditioned Supervised Learning with Negative Feedback

ArXiv CS.AI 50%

arXiv:2509.03206v2 Announce Type: replace-cross Abstract: Learning from reward functions and imitation learning of demonstrations are the two principal approaches for training autonomous systems that interact with an environment through action and observation. Both, however,...

<https://arxiv.org/abs/2509.03206>

11 CONF

AegisShield: Democratizing Cyber Threat Modeling with Generative AI

ArXiv CS.AI 50%

arXiv:2509.10482v2 Announce Type: replace-cross Abstract: The increasing sophistication of technology systems makes traditional threat modeling hard to scale, especially for small organizations with limited resources. This paper develops and evaluates AegisShield, a generative...

<https://arxiv.org/abs/2509.10482>

12 CONF

Scaling Laws and Spectra of Shallow Neural Networks in the Feature Learning Regime

ArXiv CS.AI 50%

arXiv:2509.24882v3 Announce Type: replace-cross Abstract: Neural scaling laws underlie many of the recent advances in deep learning, yet their theoretical understanding remains largely confined to linear models. In this work, we present a systematic analysis of scaling laws for...

<https://arxiv.org/abs/2509.24882>

13 CONF

A Lexical Analysis of online Reviews on Human-AI Interactions

ArXiv CS.AI 50%

arXiv:2511.13480v2 Announce Type: replace-cross Abstract: This study focuses on understanding the complex dynamics between humans and AI systems by analyzing user reviews. While previous research has explored various aspects of human-AI interaction, such as user perceptions and...

<https://arxiv.org/abs/2511.13480>

14 CONF

Layer-wise Positional Bias in Short-Context Language Modeling

ArXiv CS.AI 50%

arXiv:2601.04098v2 Announce Type: replace-cross Abstract: Transformer language models systematically prefer tokens at specific input positions regardless of semantic relevance---a phenomenon known as positional bias. Prior work characterizes this bias in model behavior through...

<https://arxiv.org/abs/2601.04098>

15 CONF

One Leak Away: How Pretrained Model Exposure Amplifies Jailbreak Risks in Finetuned LLMs

ArXiv CS.AI 50%

arXiv:2512.14751v3 Announce Type: replace-cross Abstract: Finetuning pretrained large language models (LLMs) has become the standard paradigm for developing downstream applications. However, its security implications remain unclear, particularly regarding whether finetuned LLMs...

<https://arxiv.org/abs/2512.14751>

16 CONF

All-Quadrant Bounded Clipping GRPO: Closing the Unbounded Blind Spot for Stable and Generalizable Training

ArXiv CS.AI 50%

arXiv:2601.03895v2 Announce Type: replace-cross Abstract: Group Relative Policy Optimization (GRPO) has emerged as a popular algorithm for reinforcement learning with large language models (LLMs). However, GRPO inherits PPO's token-level clipping while replacing token-level...

<https://arxiv.org/abs/2601.03895>

17 CONF

d3LLM: Ultra-Fast Diffusion LLM using Pseudo-Trajectory Distillation

ArXiv CS.AI 50%

arXiv:2601.07568v3 Announce Type: replace-cross Abstract: Diffusion large language models (dLLMs) offer capabilities beyond those of autoregressive (AR) LLMs, such as parallel decoding and random-order generation. However, realizing these benefits in practice is non-trivial, as...

<https://arxiv.org/abs/2601.07568>

18 CONF

FI-TW: An Open Train-Weather Dataset for Railway Delay Analysis in Finland

ArXiv CS.AI 50%

arXiv:2601.16592v2 Announce Type: replace-cross Abstract: Train delays result from complex interactions between operational, technical, and environmental factors. While weather impacts railway reliability, particularly in Nordic regions, existing datasets rarely integrate...

<https://arxiv.org/abs/2601.16592>

19 CONF

Stochastic Parrots or Singing in Harmony? Testing Five Leading LLMs for their Ability to Replicate a Human Survey with Synthetic Data

ArXiv CS.AI 50%

arXiv:2603.00059v3 Announce Type: replace-cross Abstract: How well can AI-derived synthetic research data replicate the responses of human participants? An emerging literature has begun to engage with this question, which carries deep implications for organizational research...

<https://arxiv.org/abs/2603.00059>

20 CONF

MM-ISTS: Cooperating Irregularly Sampled Time Series Forecasting with Multimodal Vision-Text LLMs

ArXiv CS.AI 50%

arXiv:2603.05997v2 Announce Type: replace-cross Abstract: Irregularly sampled time series (ISTS) are widespread in real-world scenarios, exhibiting asynchronous observations on uneven time intervals across diverse variables. Existing ISTS forecasting methods often solely utilize...

<https://arxiv.org/abs/2603.05997>

21 CONF

Look Twice: Training-Free Evidence Highlighting for Knowledge-based Visual Question Answering

ArXiv CS.AI 50%

arXiv:2604.01280v2 Announce Type: replace-cross Abstract: Knowledge-based Visual Question Answering (KB-VQA) requires Multimodal Large Language Models (MLLMs) to identify and combine fine-grained visual cues with retrieved textual evidence. However, retrieval often introduces...

<https://arxiv.org/abs/2604.01280>

22 CONF

Domain-Gated Latent Diffusion: Generative Inverse Design of HMX-Class Energetic Materials with First-Principles Validation

ArXiv CS.AI 50%

arXiv:2605.26540v2 Announce Type: replace-cross Abstract: Energetic materials power mining, demolition, propulsion and airbags, yet today's compounds were designed decades ago. A successor must combine high energy release, low sensitivity to accidental initiation and a practical...

<https://arxiv.org/abs/2605.26540>

23 CONF

BenHalluEval: A Multi-Task Hallucination Evaluation Framework for Large Language Models on Bengali

ArXiv CS.AI 50%

arXiv:2605.31483v4 Announce Type: replace-cross Abstract: Despite Bengali being the sixth most spoken language in the world, no prior work has systematically evaluated hallucination in large language models (LLMs) for Bengali. We introduce BenHalluEval, a fine-grained...

<https://arxiv.org/abs/2605.31483>

24 CONF

OpenAI Privacy Filter: A Cross-Lingual, Cross-Domain PII Evaluation Across 32 Benchmarks

ArXiv CS.AI 50%

arXiv:2608.02616v2 Announce Type: replace-cross Abstract: We present what is, to our knowledge, the first systematic evaluation of OpenAI's Privacy Filter (OPF), a 1.5B-parameter model that converts an autoregressive language model into a bidirectional PII detector, across 32...

<https://arxiv.org/abs/2608.02616>