



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

## Source Breakdown

| Source        | Articles | Category  |
|---------------|----------|-----------|
| ArXiv CS.AI   | 117      | Media     |
| HackerNews    | 11       | Community |
| TechCrunch AI | 2        | Media     |
| The Verge AI  | 1        | Media     |

## CONFIRMED SIGNALS — 131 stories

CONF

**Show HN: Airy - Free, fast, and simple voice content creation**

HackerNews 50% HN 38

No summary — see link for full article.

<https://airy.so>

CONF

**Show HN: Voice driven murder mystery, Interview AI suspects with your voice**

HackerNews 50% HN 20

Hey HN! I'm excited to show off this really fun project I put together. I originally built this project 2-3 years ago, AI was already booming at the time, however voice AI agents were still very early. I loved my proof of concept at the time, but wasn't quite happy with it. I...

<https://www.whodunnitai.com/>

CONF

**Show HN: Open-source playground to red-team AI agents against public prompts**

HackerNews 50% HN 12

No summary — see link for full article.

<https://playground.fabraix.com/>

CONF

**Show HN: Job Seeker - AI agent skills for job searching**

HackerNews 50% HN 5

Markdown skills that coding agents (Devin, Claude, Cursor, opencode) consume to automate the entire job search: search LinkedIn, fill Easy Apply, track in kanban, draft replies.

<https://github.com/galiprandi/job-seeker>

CONF

**Show HN: Pacific Slate: a self-hosted, model-agnostic multi-agent AI assistant**

HackerNews 50% HN 5

I didn't want to buy a standalone computer or repurpose a laptop to run constantly so I could maintain a system to sync my LLMs, so I built this. It's a simple overview of my system, laid out in a way easy to unpack and replicate for yourself. The project is meant to be...

<https://pacslate.com/>

CONF

**Show HN: FAith - a syntax-free JVM language, compiled by an LLM front-end**

HackerNews 50% HN 4

No summary — see link for full article.

<https://github.com/krossovochkin/faith>

CONF

**Show HN: Whetstone - 20 Claude Code skills, each distilled from one real failure**

HackerNews 50% HN 3

No summary — see link for full article.

<https://whetstone.akbarsha.dev/>

## CONF

**Show HN: See how long it takes ChatGPT Sol 5.6 to send a byte across the LAN**

HackerNews 50% HN 3

The problem: I noticed that any VNC client I tried wasn't able to reliably send keystrokes (about a byte on average) on my local wired LAN between a Mac Mini and a Windows desktop, both connected with gigabit ethernet cables to the same router. Keystrokes would constantly drop...

<https://robss2020.github.io/sol-vnc-elapsed/>

## CONF

**Show HN: TokenSpend, the AI ROI Solution**

HackerNews 50% HN 3

Hey HN, I was curious about how most companies and individuals were using AI without any meaningful insights. I built TokenSpend ( <https://tokenspend.dev/> ); we find existing ROI on AI code usage by tracking how much output was produced for your agent sessions, identifying...

<https://tokenspend.dev/>

## 0 CONF

**Show HN: Aadress - LangChain integration for cross-agent discovery and trust**

HackerNews 50% HN 3

No summary — see link for full article.

<https://docs.langchain.com/oss/python/integrations/tools>

## 1 CONF

**Show HN: Wardline, a Go proxy that auto-blocks compromised AI agents**

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/kabirnarang39/wardline>

## 2 CONF

**Anthropic is turning Claude Code's auto mode on by default**

TechCrunch AI 50%

Programming with Claude Code will soon require even less human oversight.

<https://techcrunch.com/2026/08/09/anthropic-is-turning-claude-codes-auto-mode-on-by-def...>

## 3 CONF

**The AI safety test is becoming a safety risk**

TechCrunch AI 50%

AI agents are escaping cybersecurity testing environments and reaching real-world systems, raising questions about whether safety infrastructure, industry standards and regulation can keep pace with increasingly powerful models.

<https://techcrunch.com/2026/08/09/the-ai-safety-test-is-becoming-a-safety-risk/>

## 4 CONF

**AI detectors are creating a new era of distrust**

The Verge AI 50%

This is The Stepback, a weekly newsletter breaking down one essential story from the tech world. For more news about how AI is changing our daily lives, follow Emma Roth. The Stepback arrives in our subscribers' inboxes at 8AM ET. Opt in for The Stepback here. How it started...

<https://www.theverge.com/column/976690/ai-writing-detectors-suspicion>

5 CONF

**Towards Multi-Label Graph Foundation Models: from Single-Vector Representation Learning to Multi-Semantic Basis Learning**

ArXiv CS.AI 50%

arXiv:2608.06394v1 Announce Type: new Abstract: Multi-label node classification is an important yet challenging task in graph learning, where nodes exhibit multiple semantics simultaneously. Existing methods for multi-label node classification can effectively model multiple...

<https://arxiv.org/abs/2608.06394>

6 CONF

**EntropyMoE: Entropy-Aware Sparse Expert Routing for Tokenizer-Free LLMs**

ArXiv CS.AI 50%

arXiv:2608.06398v1 Announce Type: new Abstract: Recent byte-level large language models (LLMs) have made tokenizer-free modeling increasingly competitive by grouping bytes into dynamically sized patches. However, existing byte-patch architectures still apply the same dense...

<https://arxiv.org/abs/2608.06398>

7 CONF

**Beyond Routing Weights: Faithful Response-Level Interpretation of Mixture-of-Experts Reward Models via Contribution Contrast**

ArXiv CS.AI 50%

arXiv:2608.06400v1 Announce Type: new Abstract: Reward models are central to learning from human preferences, yet identifying what drives their predictions remains challenging. Recent sparse Mixture-of-Experts (MoE) reward models seek to improve interpretability by routing...

<https://arxiv.org/abs/2608.06400>

8 CONF

**Interpretable Unsupervised Community Detection with LLM-Symbolized Structured Processes**

ArXiv CS.AI 50%

arXiv:2608.06402v1 Announce Type: new Abstract: Community detection is a fundamental task in graph analytics that aims to identify cohesive groups of entities with similar behaviors or interests. Classic objective-driven methods struggle with complex graph structures, while...

<https://arxiv.org/abs/2608.06402>

9 CONF

**Learning to Predict Middle-Layer Attention in MLLMs for Visual Token Pruning**

ArXiv CS.AI 50%

arXiv:2608.06411v1 Announce Type: new Abstract: Multimodal large language models (MLLMs) achieve strong performance across diverse vision-language tasks, but their efficiency is limited by the cost of processing numerous visual tokens. Visual token pruning can reduce this cost,...

<https://arxiv.org/abs/2608.06411>

0 CONF

**WebGrader: Training LLMs for Web Development with Self-Evolving Programmatic Grader**

ArXiv CS.AI 50%

arXiv:2608.06474v1 Announce Type: new Abstract: Large language models increasingly generate complete websites from natural-language descriptions, and reinforcement learning has become a central approach to closing their remaining functional gap. This training regime is...

<https://arxiv.org/abs/2608.06474>

1 CONF

**Can MLLMs Decode the Creative Leap? Introducing C4 for Cross-Concept Understanding**

ArXiv CS.AI 50%

arXiv:2608.06501v1 Announce Type: new Abstract: Creative capabilities of MLLMs matter in design, communication, education, and human--AI collaboration, yet remain difficult to evaluate because explicit targets and reward signals are scarce compared with accuracy-oriented tasks....

<https://arxiv.org/abs/2608.06501>

2 CONF

**KNOWPLAN: Knowledge-Driven AI Agents for Smart Degree Pathway Planning**

ArXiv CS.AI 50%

arXiv:2608.06530v1 Announce Type: new Abstract: Planning a degree from official university sources requires solving two problems in order. The institution's curriculum must first be reconstructed from catalogs, departmental pages, JSON endpoints, and PDFs that share no schema,...

<https://arxiv.org/abs/2608.06530>

3 CONF

**Multi-Level Modeling of Large Language Model Inference Latency and Energy via Hybrid Analytical--Machine-Learning Predictors**

ArXiv CS.AI 50%

arXiv:2608.06723v1 Announce Type: cross Abstract: The rapid scaling of Large Language Models (LLMs) has significantly increased computational cost, energy consumption, and inference latency, making accurate estimation essential for sustainable artificial intelligence deployment...

<https://arxiv.org/abs/2608.06723>

4 CONF

**Automated item evaluation: Predicting item acceptance and rejection using LLM-generated critiques**

ArXiv CS.AI 50%

arXiv:2608.06609v1 Announce Type: new Abstract: Automated item evaluation (AIE) refers to the use of computational methods to assess item quality without requiring manual expert review or field testing of the items under evaluation. We aimed to build a near-comprehensive AIE...

<https://arxiv.org/abs/2608.06609>

5 CONF

**Shape Your Feed: An LLM-based Agentic System for Conversational Recommendation**

ArXiv CS.AI 50%

arXiv:2608.06632v1 Announce Type: new Abstract: Industrial recommendation systems predominantly adopt a passive ranking paradigm that infers user preferences from implicit behavioral signals (e.g., clicks, dwell time) rather than explicit, natural language inputs. As a result,...

<https://arxiv.org/abs/2608.06632>

6 CONF

**TradeVerse: A Longitudinal Benchmark of Political Negotiation in International Trade**

ArXiv CS.AI 50%

arXiv:2608.06549v1 Announce Type: cross Abstract: LLMs are increasingly being applied to tasks involving institutional and political texts, but existing benchmarks evaluate them on isolated documents or single tasks. In realpolitik, negotiations are longitudinal data, where...

<https://arxiv.org/abs/2608.06549>

7 CONF

**CellWorld: From Gene-Level Reconstruction to Latent Cell Prediction in Spatial Transcriptomics Foundation Models**

ArXiv CS.AI 50%

arXiv:2608.06659v1 Announce Type: new Abstract: This paper shows that latent-space predictive pretraining can provide a scalable route to foundation models for spatial transcriptomics. Existing spatial transcriptomics foundation models primarily reconstruct masked gene...

<https://arxiv.org/abs/2608.06659>

8 CONF

**Vehicle routing problem using deep reinforcement learning - A case study about truck planning in the industry**

ArXiv CS.AI 50%

arXiv:2608.06668v1 Announce Type: new Abstract: As an important component of the supply chain industry, transportation has experienced rapid development in the past decade with the assistance of digital platforms and intelligent algorithms. Within the field of transportation...

<https://arxiv.org/abs/2608.06668>

9 CONF

**AutoMOOSE: An Agentic AI for Autonomous Phase-Field Simulation**

ArXiv CS.AI 50%

arXiv:2603.20986v2 Announce Type: replace Abstract: Phase-field modeling links thermodynamics and kinetics to microstructural evolution, but multiphysics frameworks such as MOOSE require expertise to construct inputs, manage campaigns, diagnose failures, and validate results. We...

<https://arxiv.org/abs/2603.20986>

0 CONF

**AgentPatch: Coarse-to-Fine Weak-Task Repair for Merging Agentic Multimodal Large Language Models**

ArXiv CS.AI 50%

arXiv:2608.06699v1 Announce Type: new Abstract: Agentic multimodal large language models (MLLMs) extend multimodal perception and reasoning with planning, tool use, and interaction in dynamic environments. Yet current models are specialized for particular tools or environments,...

<https://arxiv.org/abs/2608.06699>

1 CONF

**IB-RL: Isolated Bilateral Reinforcement Learning for Strategic Dialogue Agents**

ArXiv CS.AI 50%

arXiv:2608.06735v1 Announce Type: new Abstract: Reinforcement learning (RL) has achieved strong results in improving large language models (LLMs) on tasks with stationary, verifiable rewards, such as mathematical reasoning and code execution. In these settings, the environment...

<https://arxiv.org/abs/2608.06735>

2 CONF

**Mind the Gap: A Dual Knowledge Graph Framework for Unified Multi-task User Intent Inference**

ArXiv CS.AI 50%

arXiv:2608.06752v1 Announce Type: new Abstract: This paper proposes DKG-MTI, a dual knowledge graph framework for unified multi-task user intent inference from online travel reviews. Existing approaches often rely on hierarchical pipelines that suffer from error propagation or...

<https://arxiv.org/abs/2608.06752>

3 CONF

**Capek 0.5: An Execution-Centric Vision-Language Model for Embodied Intelligence**

ArXiv CS.AI 50%

arXiv:2608.06756v1 Announce Type: new Abstract: Vision-language models are increasingly serving as the reasoning core of embodied agents. Robot execution is inherently iterative: each action reshapes the scene and physical state, continually renewing what must be perceived,...

<https://arxiv.org/abs/2608.06756>

4 CONF

**Surg-UniWorld: A Unified Surgical World Model with Multimodal Control Experts**

ArXiv CS.AI 50%

arXiv:2608.06770v1 Announce Type: new Abstract: Controllable surgical world models can provide a generative foundation for surgical artificial intelligence and simulation by synthesizing realistic instrument-tissue interactions. However, existing methods lack a unified...

<https://arxiv.org/abs/2608.06770>

5 CONF

**Evolving Parallel Algorithm Portfolios via Potential-Aware Instance Generation with LLMs**

ArXiv CS.AI 50%

arXiv:2608.06808v1 Announce Type: new Abstract: The Automatic Construction of Portfolios via Large Language Models (LLM-ACP) suffers from poor generalization in practical few-shot scenarios when solving complex combinatorial optimization problems. Instance and algorithm...

<https://arxiv.org/abs/2608.06808>

6 CONF

**Long-Horizon Agent Trajectory Attribution: A Unified Benchmark and Fine-Grained Annotation Framework**

ArXiv CS.AI 50%

arXiv:2608.06909v1 Announce Type: new Abstract: Large language model (LLM) agents increasingly operate through long-horizon trajectories involving user instructions, tool use, external observations, and memory. Existing benchmarks primarily evaluate behavioral outcomes but...

<https://arxiv.org/abs/2608.06909>

7 CONF

**Does Splitting a Triage Decision Across Agents Hide Bias or Help Catch It? A Multi-Agent Simulation Study of LLM-Based Resource Allocation Under Audit Capacity Constraints**

ArXiv CS.AI 50%

arXiv:2608.06949v1 Announce Type: new Abstract: Prior benchmarking work has shown that a single large language model (LLM), forced to make life-or-death resource-allocation decisions, exhibits measurable demographic bias. Real deployments, however, rarely use a single agent:...

<https://arxiv.org/abs/2608.06949>

8 CONF

**Critical Acclaim Orientation in Large Language Models: Evidence from Film Preference Elicitation**

ArXiv CS.AI 50%

arXiv:2608.06955v1 Announce Type: new Abstract: Large language models (LLMs) are trained on corpora that contain expressions of human judgment about films, books, music, and more. Yet whether LLMs systematically reproduce evaluative hierarchies remains unclear. Prior research on...

<https://arxiv.org/abs/2608.06955>



9 CONF

**CAI Copilot: Reducing Operational Workload in Molecular Design through Intent-Driven Agentic Workflows**

ArXiv CS.AI 50%

arXiv:2608.06961v1 Announce Type: new Abstract: Early-stage molecular design is an iterative process, not just a task of generating molecules. Researchers turn broad goals into design strategies, refine candidates, assess many properties, and gather evidence before synthesis and...

<https://arxiv.org/abs/2608.06961>

0 CONF

**Learning in Deep Networks under Dale's Constraint**

ArXiv CS.AI 50%

arXiv:2608.06963v1 Announce Type: new Abstract: Biologically plausible learning models aim to explain how neural circuits can implement effective learning under the constraints of real neurons. Although significant progress has been made, a major remaining challenge is that...

<https://arxiv.org/abs/2608.06963>

1 CONF

**ReQuant: Fixed-Grid Discrete Refinement for Post-Training Quantization**

ArXiv CS.AI 50%

arXiv:2608.07019v1 Announce Type: new Abstract: Post-training quantization (PTQ) is widely used to reduce the memory and computational cost of large language models. Existing PTQ methods typically obtain an initial quantized model through heuristic rules or greedy optimization,...

<https://arxiv.org/abs/2608.07019>

2 CONF

**ZIPBrain: Can EEG Foundation Models Be Faster, Locally Deployable, but Accurate?**

ArXiv CS.AI 50%

arXiv:2608.07033v1 Announce Type: new Abstract: This work investigates whether Electroencephalograph (EEG) foundation models (EFMs) can be made faster and locally deployable without sacrificing accuracy. EEG foundation models are a major trend, offering strong general-purpose...

<https://arxiv.org/abs/2608.07033>

3 CONF

**Unsupervised Adaptation of PDE Foundation Models**

ArXiv CS.AI 50%

arXiv:2608.07053v1 Announce Type: new Abstract: Pretrained partial differential equation (PDE) foundation models can generalize across different equations, but adapting them to unseen PDE systems typically requires dense solution data, which is often expensive or unavailable. To...

<https://arxiv.org/abs/2608.07053>

4 CONF

**BONSAI: Evolvability-Guided Tree Search over Skills**

ArXiv CS.AI 50%

arXiv:2608.07056v1 Announce Type: new Abstract: A skill is a natural language document that steers a frozen agent whose weights cannot be updated so any capability the agent lacks must be supplied in prose Optimising a skill is therefore optimising text against a score and the...

<https://arxiv.org/abs/2608.07056>

5 **CONF****PTQ4SNN: Membrane-Aware Post-Training Quantization for Spiking Neural Networks****ArXiv CS.AI** **50%**

arXiv:2608.07066v1 Announce Type: new Abstract: Spiking neural networks (SNNs) enable sparse and event-driven computation, but their low-bit deployment remains incomplete because recurrent membrane states are commonly retained in floating point even after weight quantization....

<https://arxiv.org/abs/2608.07066>

6 **CONF****CoinRAG: Contextualized Information Nugget KV Cache Reuse for Long-Context RAG****ArXiv CS.AI** **50%**

arXiv:2608.07458v1 Announce Type: cross Abstract: Recent optimization studies on Retrieval-Augmented Generation (RAG) have exploited chunk-level KV cache reuse to avoid processing long retrieved contexts for higher efficiency, while significant information redundancy and noise...

<https://arxiv.org/abs/2608.07458>

7 **CONF****Transformers Struggle to Use Their Emergent World Models: Revisiting the Tower of Hanoi, and the Illusion of Thinking****ArXiv CS.AI** **50%**

arXiv:2608.07077v1 Announce Type: new Abstract: The Tower of Hanoi is a simple planning puzzle that in prior work has proven challenging for large reasoning models (LRMs). Current models solve the standard formulation of the puzzle, but still struggle with the flat-to-flat...

<https://arxiv.org/abs/2608.07077>

8 **CONF****Momba: Network Modernization Improves Multi-Objective Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2608.07180v1 Announce Type: cross Abstract: Recent advances in deep reinforcement learning (RL) have shown that improving neural network architectures can yield substantial gains in sample efficiency and asymptotic performance without altering the underlying algorithms. In...

<https://arxiv.org/abs/2608.07180>

9 **CONF****DiDPO: Diff-in-Diff Policy Optimization for Coding Agent Training****ArXiv CS.AI** **50%**

arXiv:2608.07147v1 Announce Type: new Abstract: Reinforcement learning with Verifiable Reward (RLVR) has emerged as a powerful paradigm for training coding agents, where the execution feedback from compilation and tests provides objective verification. However, unlike agent...

<https://arxiv.org/abs/2608.07147>

0 **CONF****A MARL Centered Reference Architecture for Large Language Model Augmentation in Smart Manufacturing****ArXiv CS.AI** **50%**

arXiv:2608.07148v1 Announce Type: new Abstract: Modern manufacturing imposes six coupled demands on adaptive control: local decisions with global consequences, partial observability, nonstationarity, reflex speed response with long horizon effects, delayed and diffuse outcomes,...

<https://arxiv.org/abs/2608.07148>

1 CONF

**NiyamAI - An Intent-Bound AI Agent with Cryptographically Verifiable Guardrails using Zero-Knowledge Proofs**

ArXiv CS.AI 50%

arXiv:2608.07167v1 Announce Type: new Abstract: Giving an AI agent the ability to send emails, query databases, or execute commands is useful--until the agent is tricked into doing something it shouldn't. Prompt injection, hallucinated reasoning, and unsafe tool calls form the...

<https://arxiv.org/abs/2608.07167>

2 CONF

**Agent Memory Distillation: Empowering Small LLM Agents with Hierarchical Teacher Memory**

ArXiv CS.AI 50%

arXiv:2608.07169v1 Announce Type: new Abstract: Memory systems have shown promise for improving agent performance, but their potential remains largely unexplored for small language models, which struggle to generate sufficient successful trajectories on their own. We propose...

<https://arxiv.org/abs/2608.07169>

3 CONF

**Authoring and Management of Transparent Research Integrity Assessments of Randomised Clinical Trial Publications Using LLM-assisted Tools and Provenance Knowledge Graphs**

ArXiv CS.AI 50%

arXiv:2608.07202v1 Announce Type: new Abstract: Systematic reviews of Randomised Controlled Trials (RCTs) are routinely used as evidence for clinical care guidelines. Such evidence has to meet high research integrity standards to prevent low quality or false research outputs...

<https://arxiv.org/abs/2608.07202>

4 CONF

**Beyond the Black Box: Interpretable Models of Human Randomisation Failures**

ArXiv CS.AI 50%

arXiv:2608.07220v1 Announce Type: new Abstract: Mixed strategy equilibrium predicts i.i.d play: past actions should not help predict future decisions. Human players, however, systematically depart from this benchmark, and in O'Neill's zero sum card game, these departures can be...

<https://arxiv.org/abs/2608.07220>

5 CONF

**People Are Not Just Their Countries. Disentangling Social Determinants of LLM Value Alignment Across Europe**

ArXiv CS.AI 50%

arXiv:2608.07367v1 Announce Type: new Abstract: As Large Language Models (LLMs) are increasingly used as a primary source of information and advice, understanding their alignment to humans in terms of values becomes a pressing concern. A growing literature has leveraged large...

<https://arxiv.org/abs/2608.07367>

6 CONF

**Do AI Personas Grow? Analyzing and Benchmarking Personality Evolution in LLM Agents After Life Events**

ArXiv CS.AI 50%

arXiv:2608.06485v1 Announce Type: cross Abstract: Personality-conditioned LLM agents (PC-Agents) are increasingly used in emotional support, social simulation, and role-playing, motivating the development of lifelong agents that remain coherent over extended interactions. A key...

<https://arxiv.org/abs/2608.06485>

7 CONF

**Aftab: A Comprehensive Benchmark of CNN Encoders and Advanced Value Functions in Parallelized Q-Networks**

ArXiv CS.AI 50%

arXiv:2608.07335v1 Announce Type: cross Abstract: Recent advancements in deep reinforcement learning have increasingly favored simplified, highly parallelized paradigms. Notably, the Parallelized Q-Network (PQN) algorithm achieves stable off-policy learning without relying on...

<https://arxiv.org/abs/2608.07335>

8 CONF

**A Picture is Worth a Thousand Tokens: How Vision Language Models Cut AI Energy Costs While Improving Accuracy**

ArXiv CS.AI 50%

arXiv:2608.07427v1 Announce Type: new Abstract: LLM inference accounts for over 90% of AI operational energy, scaling directly with input token count---a critical inefficiency for telecom network analytics and numerical time-series data analysis (NTSDA), where raw multivariate...

<https://arxiv.org/abs/2608.07427>

9 CONF

**Post-Grokking Collapse at the Representation-Readout Interface in Muon-Trained Transformers**

ArXiv CS.AI 50%

arXiv:2608.07436v1 Announce Type: new Abstract: Under the standard split, Muon gets hidden matrices and AdamW embeddings/output head. Muon groks modular addition faster, but its solutions do not hold. All nine configurations on  $\$(a+b) \bmod 113\$$  grok and later lose...

<https://arxiv.org/abs/2608.07436>

0 CONF

**Fisher-R1: Training LLM Agents for Reliable Hypothesis Testing**

ArXiv CS.AI 50%

arXiv:2608.07437v1 Announce Type: new Abstract: Reliable hypothesis testing is the foundation of many empirical scientific claims. Large language model (LLM) agents are increasingly used to automate this process, as they can inspect datasets, generate code, and produce analyses...

<https://arxiv.org/abs/2608.07437>

1 CONF

**Interaction Creates Dynamical AI Behavior Absent in Isolation**

ArXiv CS.AI 50%

arXiv:2608.07457v1 Announce Type: new Abstract: What will happen when AI agents interact in daily life, e.g. when one AI starts bossing another around? We find a counterintuitive answer that opens new avenues for out-of-equilibrium Physics. When a boss AI directs a stream of...

<https://arxiv.org/abs/2608.07457>

2 CONF

**Mitigating Scoring Bias in LLM-as-a-Judge via Random Number Generation**

ArXiv CS.AI 50%

arXiv:2608.05726v1 Announce Type: cross Abstract: Large Language Models (LLMs) are often used as evaluators of text quality, known as LLM-as-a-Judge, which can outperform conventional automatic evaluation metrics that rely on reference texts. However, LLM evaluators tend to...

<https://arxiv.org/abs/2608.05726>

3 CONF

**Multimodal Drivers' Emotion Recognition and Safety-Oriented Intervention for Intelligent Transportation Systems**

ArXiv CS.AI 50%

arXiv:2608.06378v1 Announce Type: cross Abstract: Driver emotions can affect risk perception, decision-making, and vehicle control under complex road conditions. Existing studies mainly focus on driver emotion recognition, while limited attention has been given to context-aware...

<https://arxiv.org/abs/2608.06378>

4 CONF

**Mobile Interaction for Assessing Fatigue, Sleep, and Activity in Neurodegenerative and Chronic Diseases**

ArXiv CS.AI 50%

arXiv:2608.06380v1 Announce Type: cross Abstract: Fatigue, sleep, or disturbances in daily activities are common symptoms among patients with neurodegenerative disorders (NDD) and immune-mediated inflammatory diseases (IMID). The current assessment of such symptoms is usually...

<https://arxiv.org/abs/2608.06380>

5 CONF

**TEXAS: Task-Expert-Aware Supervision for Downstream Mixture-of-Experts LLM Adaptation**

ArXiv CS.AI 50%

arXiv:2608.06396v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) language models route each token through a small subset of experts, making routing patterns useful for identifying task-relevant experts during downstream adaptation. Yet current approaches have two...

<https://arxiv.org/abs/2608.06396>

6 CONF

**TransSLR: A Lightweight Transformer for Sign Language Recognition**

ArXiv CS.AI 50%

arXiv:2608.06407v1 Announce Type: cross Abstract: Automated Sign Language Recognition for under-represented languages remains a largely unsolved problem. Central African Sign Language (CASL) exemplifies this gap: the only available bench-mark, CASL-W60, has a best reported...

<https://arxiv.org/abs/2608.06407>

7 CONF

**Separating Decision-Rule Misalignment from Readout-Coverage Limitations in Speech Language Models**

ArXiv CS.AI 50%

arXiv:2608.06409v1 Announce Type: cross Abstract: Speech language models are increasingly evaluated on paralinguistic tasks by the accuracy of prompted answers, but answer accuracy combines failures at different stages of the audio-to-answer computation. We introduce a...

<https://arxiv.org/abs/2608.06409>

8 CONF

**WorldMark: A Plug-and-Play World Knowledge Interface for Cross-Host Language Model Watermarking**

ArXiv CS.AI 50%

arXiv:2608.06416v1 Announce Type: cross Abstract: Watermarking traces the provenance of text produced by large language models by embedding statistically detectable signals during decoding. Existing schemes fall into logits-based, sampling-based, entropy-aware, and...

<https://arxiv.org/abs/2608.06416>

9 CONF

**CyberForge: Verified Vulnerability Injection at Repository Level for Cybersecurity Agent Training**

ArXiv CS.AI 50%

arXiv:2608.06471v1 Announce Type: cross Abstract: Despite recent advances, frontier large language model (LLM) agents remain limited in discovering and patching complex vulnerabilities in real-world software. Generally available agents can already aid attackers, who only need to...

<https://arxiv.org/abs/2608.06471>

0 CONF

**StepJack: Benchmarking Computer-Use Agent Safety Against Multi-Step Indirect Prompt Injection**

ArXiv CS.AI 50%

arXiv:2608.06477v1 Announce Type: cross Abstract: Computer-use agents (CUAs) face a growing threat from indirect prompt injection, where adversarial instructions are planted in the environment such as web pages. In this paper, we introduce multi-step indirect prompt injection, a...

<https://arxiv.org/abs/2608.06477>

1 CONF

**Agentic AI: User Empowerment or Enclosure?**

ArXiv CS.AI 50%

arXiv:2608.06510v1 Announce Type: cross Abstract: Agentic AI promises a more flexible form of digital agency: systems that can act on users' behalf, from filtering content to negotiating prices to selecting services. Whether it will empower users is an open question, and we...

<https://arxiv.org/abs/2608.06510>

2 CONF

**CertBind from Multimodal Connectivity to Certifiable Retrieval Decisions**

ArXiv CS.AI 50%

arXiv:2608.06516v1 Announce Type: cross Abstract: Lightweight connectors make frozen multimodal encoders composable at the representation level. Deployment exposes a second problem at the level of task decisions. A connected route can expand cross-modal reach while changing an...

<https://arxiv.org/abs/2608.06516>

3 CONF

**Beyond "AI Language": The case for the idiolectal nature of LLM output**

ArXiv CS.AI 50%

arXiv:2608.06589v1 Announce Type: cross Abstract: While large language model outputs are frequently analysed as a collective super variety termed "AI language," this chapter argues that this perspective coexists with distinct, model-specific linguistic signatures akin to human...

<https://arxiv.org/abs/2608.06589>

4 CONF

**How Much AI Is in This Track? Quantifying the Proportion of AI-Generated Stems in Hybrid Music Mixtures**

ArXiv CS.AI 50%

arXiv:2608.07285v1 Announce Type: cross Abstract: AI-generated music is increasingly used at the stem level, with producers integrating synthetic drums, basslines, or vocals alongside human-performed instruments. However, current AI music detection systems are binary, treating...

<https://arxiv.org/abs/2608.07285>

- 5
CONF
**KReF: Training-Free Retrieval for Long-Term Time-Series Forecasting and Predictive Uncertainty**
ArXiv CS.AI 50%
arXiv:2608.06748v1 Announce Type: cross Abstract: Probabilistic long-term time-series forecasting commonly relies on trained models. Training-free conformal methods typically construct intervals around a pre-existing point forecaster and do not natively represent a complete...
<https://arxiv.org/abs/2608.06748>
- 6
CONF
**Genotypic Triggers: Exposing Pharmacogenomic Blind Spots via Host-Specific Backdoors in Generative Antimicrobial Peptide Models**
ArXiv CS.AI 50%
arXiv:2608.06779v1 Announce Type: cross Abstract: Large Language Models (LLMs) have accelerated drug discovery, particularly in the automated design of antimicrobial peptides (AMPs). However, current validation pipelines for peptide generation models overlook historical...
<https://arxiv.org/abs/2608.06779>
- 7
CONF
**LoRAScan: Detecting Backdoor Prompts in Low-Rank Adapters for Large Language Models via Down-Projection Activation Spikes**
ArXiv CS.AI 50%
arXiv:2608.06795v1 Announce Type: cross Abstract: Low-rank adaptation (LoRA) enables efficient specialization and distribution of large language models through compact adapters. However, untrusted adapters introduce a supply-chain threat: a backdoored adapter can cause a model...
<https://arxiv.org/abs/2608.06795>
- 8
CONF
**Investigating Quantum-Embedded Transformers on Classical Datasets for Cross-Modality Classification**
ArXiv CS.AI 50%
arXiv:2608.06846v1 Announce Type: cross Abstract: We test whether a parameterized quantum circuit (PQC) improves a hybrid quantum-classical model's performance on classical datasets, using an interface-matched classical map as the control while holding all other components...
<https://arxiv.org/abs/2608.06846>
- 9
CONF
**Calibrating WEAT Against Anisotropy: ZCA Whitening as a Geometric Pre-Processing Step for Embedding Association Tests**
ArXiv CS.AI 50%
arXiv:2608.06908v1 Announce Type: cross Abstract: We propose Zero-phase Component Analysis (ZCA) whitening as a geometric pre-processing step for the Word Embedding Association Test (WEAT). WEAT is a bias measurement method widely used in both computational social science and AI...
<https://arxiv.org/abs/2608.06908>
- 0
CONF
**GPTKB 2.0: Browsing, Querying, and Auditing a Disambiguated LLM-Derived Knowledge Base**
ArXiv CS.AI 50%
arXiv:2608.06992v1 Announce Type: cross Abstract: We present a web demo for exploring a large-scale disambiguated knowledge base (KB) materialized from a large language model (LLM). GPTKB 2.0 contains 38.4M triples over 1.6M canonical entities, together with 207.6K consolidated...
<https://arxiv.org/abs/2608.06992>

1 CONF

**Beyond Foundation Models: Dimension-Aware Neural Architecture Search with Small-Data Representation Models for Cryocooler Lifetime Prediction**

ArXiv CS.AI 50%

arXiv:2608.06993v1 Announce Type: cross Abstract: Large-scale pretrained time-series models achieve strong results through large-scale pretraining and task-agnostic representation learning, but they rely on abundant, diverse data that industrial and scientific domains often...

<https://arxiv.org/abs/2608.06993>

2 CONF

**Accounting Graph Transformer for Short-History Multi-KPI Forecasting in Small Businesses**

ArXiv CS.AI 50%

arXiv:2608.07037v1 Announce Type: cross Abstract: Small businesses often have only 12-24 months of accounting history, yet planning and risk workflows require coordinated forecasts across financial statements. We study joint 12-month forecasting of 13 income-statement,...

<https://arxiv.org/abs/2608.07037>

3 CONF

**Human-Centered Explainable AI for TinyML Edge Devices: A Pareto-Based Selection Framework with LLM-Guided Design**

ArXiv CS.AI 50%

arXiv:2608.07091v1 Announce Type: cross Abstract: Edge Artificial Intelligence (Edge AI) enables the deployment of AI models directly on local edge devices, while such deployments are subject to strict resource constraints, particularly in clinical applications requiring local...

<https://arxiv.org/abs/2608.07091>

4 CONF

**PHOENIX: Fine-Tuned SLM-Powered Autonomous Satellite Lifetime Extension via Predictive Self-Healing and Multi-Agent AI Recovery**

ArXiv CS.AI 50%

arXiv:2608.07126v1 Announce Type: cross Abstract: Most CubeSats, small and low-cost satellites roughly the size of a shoebox, do not survive as long as they were designed to: a study of 178 missions found that only 48-65% remain operational after two years, against a designed...

<https://arxiv.org/abs/2608.07126>

5 CONF

**Fluid-DiT: Graph-Free Diffusion Transformers for Fluid Flow Simulations Learning**

ArXiv CS.AI 50%

arXiv:2608.07161v1 Announce Type: cross Abstract: Simulating complex fluid flows requires capturing full equilibrium distributions rather than just mean trajectories, yet high-fidelity solvers remain computationally prohibitive. Recent advances, such as Diffusion Graph Networks...

<https://arxiv.org/abs/2608.07161>

6 CONF

**Representation-driven Endoscopic Visual Embedding Alignment for Latent Generation**

ArXiv CS.AI 50%

arXiv:2608.07176v1 Announce Type: cross Abstract: Developing foundation generative models for endoscopy is limited by the gap between natural and clinical images and the computational cost of training large Diffusion Transformers. Although representation alignment has improved...

<https://arxiv.org/abs/2608.07176>



7 CONF

**Measuring Concept Content in Text from LLM Activations: ESG Evidence from Concept Vectors and Linear Probes**

ArXiv CS.AI 50%

arXiv:2608.07208v1 Announce Type: cross Abstract: Existing measures of how much a text is about a concept read the surface of the text: dictionary word shares, topic proportions, embedding similarities. They score the words a text uses, not the judgment a reader forms about it....

<https://arxiv.org/abs/2608.07208>

8 CONF

**Toward a Causal Data Management Ecosystem for Decision Making and Agentic AI**

ArXiv CS.AI 50%

arXiv:2608.07214v1 Announce Type: cross Abstract: Modern AI is no longer a single model but an ecosystem: classical ML predictors, deep and multimodal models, large language models, and agents, each trained and tuned over different data sources and each producing outputs at...

<https://arxiv.org/abs/2608.07214>

9 CONF

**Artificial Intelligence Can Match Domain Experts in Evidence Extraction and Critical Appraisal of Microbial Oncogenesis Research Publications**

ArXiv CS.AI 50%

arXiv:2608.07250v1 Announce Type: cross Abstract: Confirmed oncogenic microbes contribute significantly to cancer burden. Identifying novel microbial oncogenicity could yield strategies that will reduce disease burdens. However, relevant evidence is dispersed and infeasible for...

<https://arxiv.org/abs/2608.07250>

0 CONF

**Reading Copom's Tone: A Weighted LLM Framework for Hawkish-Dovish Sentiment, Forward Guidance, and Uncertainty**

ArXiv CS.AI 50%

arXiv:2608.07251v1 Announce Type: cross Abstract: This paper documents an applied natural-language-processing framework for measuring the tone of Brazilian Monetary Policy Committee (Copom) statements. The project is explicitly inspired by iSent, Ita\u's Central Bank sentiment...

<https://arxiv.org/abs/2608.07251>

1 CONF

**SCALE: Scientific Concept Aggregation via LLMs and Embeddings for Fine-Grained Taxonomy Extension**

ArXiv CS.AI 50%

arXiv:2608.07254v1 Announce Type: cross Abstract: The increasing specialization of scientific research challenges existing classification systems, which provide effective representations of broad disciplines and research topics but often fail to capture the fine-grained...

<https://arxiv.org/abs/2608.07254>

2 CONF

**TOFD: Target-Oriented Feature Decoupling against Poisoning Attacks in Split Federated Learning**

ArXiv CS.AI 50%

arXiv:2608.07274v1 Announce Type: cross Abstract: Split Federated Learning (SFL) facilitates privacy-preserving collaborative training with reduced client-side overhead. However, its split architecture introduces unique attack surfaces, rendering it vulnerable to diverse...

<https://arxiv.org/abs/2608.07274>

3 CONF

**Towards Assurance Closure in AI-Native Large-Scale Agile Software Development**

ArXiv CS.AI 50%

arXiv:2608.07317v1 Announce Type: cross Abstract: The AI-Native Manifesto envisions large-scale agile software development in which humans increasingly govern intent, risk, and exceptions while agents execute more of the engineering process. Realizing that end-state requires...

<https://arxiv.org/abs/2608.07317>

4 CONF

**Zero Gap Is Not Restoration: Stratified Per-Question Probability Evaluation and Step-wise Mitigation of Benchmark Contamination**

ArXiv CS.AI 50%

arXiv:2608.07341v1 Announce Type: cross Abstract: Test data from public benchmarks inevitably leaks into pretraining corpora, inflating evaluation scores once memorized. \textbf{Contamination mitigation evaluation} intervenes in the decoding process to suppress memorization and...

<https://arxiv.org/abs/2608.07341>

5 CONF

**Assessing AI-generated music detection in real-world broadcast monitoring**

ArXiv CS.AI 50%

arXiv:2608.07359v1 Announce Type: cross Abstract: The proliferation of AI-generated music in broadcast media raises concerns about transparency and fair compensation, but reliable detection under real broadcast conditions remains unresolved. Existing studies report substantial...

<https://arxiv.org/abs/2608.07359>

6 CONF

**Measurements Automatically Extracted from Zero Echo Time MRI Using Deep Learning Image Segmentation and Geometric Modeling Agree with Expert Manual Readings**

ArXiv CS.AI 50%

arXiv:2608.07368v1 Announce Type: cross Abstract: Computed tomography (CT) remains the reference for 3D osseous morphometry in femoroacetabular impingement (FAI) but requires ionizing radiation and manual measurement. Zero echo time (ZTE) MRI visualizes cortical bone and yields...

<https://arxiv.org/abs/2608.07368>

7 CONF

**LSEAD: A Privacy-Preserving LLM-Based Speech Analysis Framework for Early Alzheimer's Disease Screening**

ArXiv CS.AI 50%

arXiv:2608.07378v1 Announce Type: cross Abstract: Early diagnosis of Alzheimer's disease (AD) is critical for enabling timely interventions that may slow disease progression and improve patient outcomes. There is a growing need for AD detection methods that are non-invasive and...

<https://arxiv.org/abs/2608.07378>

8 CONF

**Diffusion LLMs as Targets and Adversaries: Mechanistic Safety Exploits**

ArXiv CS.AI 50%

arXiv:2608.07430v1 Announce Type: cross Abstract: Diffusion Large Language Models (DLLMs) replace autoregressive next-token prediction with iterative parallel denoising, yet their internal safety mechanisms remain poorly understood. In this work, we investigate DLLMs both as...

<https://arxiv.org/abs/2608.07430>

9 CONF

**SABRE: Scalable and Automated Benchmarking of VLMs under Stress**

ArXiv CS.AI 50%

arXiv:2608.07435v1 Announce Type: cross Abstract: Vision-language models (VLMs) are improving rapidly, but benchmark development lags behind, making weaknesses hard to identify. Building stress tests is costly: samples must satisfy controlled conditions, remain answerable, and...

<https://arxiv.org/abs/2608.07435>

00 CONF

**Taxonomy-Driven Analysis of Open-Source AI Risk Mitigation Tools**

ArXiv CS.AI 50%

arXiv:2608.07446v1 Announce Type: cross Abstract: Rapid adoption of large language models (LLMs) in enterprise settings has introduced operational, security, and governance risks. As generative AI applications move from pilot to production, manual harm identification and...

<https://arxiv.org/abs/2608.07446>

01 CONF

**CreativeInstruct: Scalably Teaching LLMs to Balance Quality, Creativity, and Diversity**

ArXiv CS.AI 50%

arXiv:2608.07460v1 Announce Type: cross Abstract: While post-training improves the capabilities of large language models (LLMs), it generally lowers their output diversity and creativity, negatively impacting tasks that explicitly require creativity (e.g., story generation) as...

<https://arxiv.org/abs/2608.07460>

02 CONF

**Serious Games: Human-AI Interaction, Evolution, and Coevolution**

ArXiv CS.AI 50%

arXiv:2505.16388v3 Announce Type: replace Abstract: The serious games between humans and AI have only just begun. Evolutionary Game Theory (EGT) models the competitive and cooperative strategies of biological entities. EGT could help predict the potential evolutionary...

<https://arxiv.org/abs/2505.16388>

03 CONF

**"LLM Agent Performance" Is Not a Single Evaluation Target**

ArXiv CS.AI 50%

arXiv:2602.03238v3 Announce Type: replace Abstract: LLM agent benchmark scores are shaped not only by the model but also by the agent harness, environment, evaluator, and inference budget. Unified execution controls these non-model factors by evaluating candidate models under...

<https://arxiv.org/abs/2602.03238>

04 CONF

**Counterfactual Simulation Training for Chain-of-Thought Faithfulness**

ArXiv CS.AI 50%

arXiv:2602.20710v2 Announce Type: replace Abstract: Inspecting Chain-of-Thought reasoning is among the most common means of understanding why an LLM produced its output. But well-known problems with CoT faithfulness severely limit what insights can be gained from this practice....

<https://arxiv.org/abs/2602.20710>

05 CONF

**MEDLEY-BENCH: Benchmarking Behavioural Metacognition and Belief Revision Under Social Pressure in Large Language Models**

ArXiv CS.AI 50%

arXiv:2604.16009v2 Announce Type: replace Abstract: Most large language model benchmarks evaluate final-answer quality but reveal little about how models revise beliefs under disagreement or conflicting evidence. We introduce MEDLEY-BENCH, an open benchmark comparing structured...

<https://arxiv.org/abs/2604.16009>

06 CONF

**Ratchet: How Reliable Must an LLM Judge Be to Retire a Skill?**

ArXiv CS.AI 50%

arXiv:2605.22148v3 Announce Type: replace Abstract: A large language model (LLM) agent that writes and edits its own skill library must also decide which skills to keep, from one noisy scalar per skill. The answer is exact: a judge scoring failures as passes at rate  $(1-\tau)/2$ ...

<https://arxiv.org/abs/2605.22148>

07 CONF

**Same Answer, Different Confidence: Protocol Sensitivity in LLM Confidence Calibration**

ArXiv CS.AI 50%

arXiv:2605.27752v3 Announce Type: replace Abstract: Is verbalized confidence better calibrated than token likelihood? The answer depends on how the token likelihood is measured: which answer is scored, and under which prompt. Published comparisons diverge on this, and in a...

<https://arxiv.org/abs/2605.27752>

08 CONF

**ForesightSafety-SAGE: A Fully Automated Scenario Generation and Safety Evaluation Framework for LLM Agents**

ArXiv CS.AI 50%

arXiv:2606.08531v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly evolving from simple text-based interaction systems into LLM agents that can maintain memory, use tools, access external environments, and execute tasks. As their capabilities and...

<https://arxiv.org/abs/2606.08531>

09 CONF

**Semantic Adapter Routing with Fine-Tuning Task Embeddings**

ArXiv CS.AI 50%

arXiv:2606.19079v2 Announce Type: replace Abstract: Parameter-efficient fine-tuning (PEFT) has led to model ecosystems in which a single backbone is paired with many task-specialized adapters. Given such a library, routing aims to select the most appropriate adapter for a user...

<https://arxiv.org/abs/2606.19079>

10 CONF

**OpenForgeRL: Train Harness-native Agents in Any Environment**

ArXiv CS.AI 50%

arXiv:2607.21557v3 Announce Type: replace Abstract: Modern AI agents rely on elaborate inference harnesses such as Claude Code, Codex, and OpenClaw to drive multi-turn reasoning, tool use, and access to external systems. While powerful, these complex harnesses also make agents...

<https://arxiv.org/abs/2607.21557>

11 CONF

**SkillTrace: Multi-Trace Provenance Auditing for LLM-Agent Skill Reuse**

ArXiv CS.AI 50%

arXiv:2608.05204v2 Announce Type: replace Abstract: LLM-agent ecosystems are rapidly growing around reusable skills: mixed-modality packages of metadata, natural-language instructions, code, tools, references, and operational workflows. As skills become marketplace artifacts,...

<https://arxiv.org/abs/2608.05204>

12 CONF

**Harnessing the Synergy between LLM Agents and Knowledge Graphs for Urban Socioeconomic Prediction**

ArXiv CS.AI 50%

arXiv:2411.00028v3 Announce Type: replace-cross Abstract: Socioeconomic prediction aims to leverage various urban data to predict the socioeconomic indicators of regions such as population and commercial activity level, which plays an important role in understanding urban...

<https://arxiv.org/abs/2411.00028>

13 CONF

**Judge a Book by its Cover: Investigating Multi-Modal LLMs for Multi-Page Handwritten Document Transcription**

ArXiv CS.AI 50%

arXiv:2502.20295v3 Announce Type: replace-cross Abstract: Handwriting text recognition (HTR) remains a challenging task. Existing approaches require fine-tuning on labeled data, which is impractical to obtain for real-world problems, or rely on zero-shot tools such as OCR...

<https://arxiv.org/abs/2502.20295>

14 CONF

**A primer on optimal transport for causal inference with observational data**

ArXiv CS.AI 50%

arXiv:2503.07811v3 Announce Type: replace-cross Abstract: The theory of optimal transportation has developed into a powerful and elegant framework for comparing probability distributions, with wide-ranging applications in all areas of science. The fundamental idea of analyzing...

<https://arxiv.org/abs/2503.07811>

15 CONF

**In Situ Training of Implicit Neural Compressors for Scientific Simulations via Sketch-Based Regularization**

ArXiv CS.AI 50%

arXiv:2511.02659v4 Announce Type: replace-cross Abstract: Focusing on implicit neural representations, we present a novel in situ training protocol that employs limited memory buffers of full and sketched data samples, where the sketched data are leveraged to prevent...

<https://arxiv.org/abs/2511.02659>

16 CONF

**Intelligence per Watt: Measuring Intelligence Efficiency of Local AI**

ArXiv CS.AI 50%

arXiv:2511.07885v5 Announce Type: replace-cross Abstract: Large language model (LLM) queries are predominantly processed by frontier models in centralized cloud infrastructure. Demand growth strains this paradigm faster than providers can scale. Two advances create an...

<https://arxiv.org/abs/2511.07885>

17 CONF

**MetaSICL: Globalizing Auditory LLMs for Underserved Speakers and Languages via Meta Speech In-Context Learning**

ArXiv CS.AI 50%

arXiv:2601.18904v3 Announce Type: replace-cross Abstract: Generative AI for speech and audio is increasingly expected to serve users across languages, cultures, and communities, yet current auditory Large Language Models (LLMs) are still largely trained and evaluated on...

<https://arxiv.org/abs/2601.18904>

18 CONF

**CASA: Classification Augmented with Safety Attention for Robust Multimodal Alignment**

ArXiv CS.AI 50%

arXiv:2604.00310v2 Announce Type: replace-cross Abstract: Multimodal large-language models (MLLMs) often experience degraded safety alignment when harmful queries exploit cross-modal interactions. Models aligned on text alone show a higher rate of successful attacks when...

<https://arxiv.org/abs/2604.00310>

19 CONF

**Cluster Attention for Graph Machine Learning**

ArXiv CS.AI 50%

arXiv:2604.07492v2 Announce Type: replace-cross Abstract: Message Passing Neural Networks have recently become the most popular approach to graph machine learning tasks; however, their receptive field is limited by the number of message passing layers. To increase the receptive...

<https://arxiv.org/abs/2604.07492>

20 CONF

**GRM: Utility-Aware Jailbreak Attacks on Audio LLMs via Gradient-Ratio Masking**

ArXiv CS.AI 50%

arXiv:2604.09222v2 Announce Type: replace-cross Abstract: Audio Large Language Models (ALLMs) enable spoken interaction but introduce new jailbreak vulnerabilities. Existing perturbation-based jailbreaks do not explicitly control which frequency bands carry the perturbation....

<https://arxiv.org/abs/2604.09222>

21 CONF

**Predictive Multi-Tier Memory Management for KV Cache in Large-Scale GPU Inference**

ArXiv CS.AI 50%

arXiv:2604.26968v2 Announce Type: replace-cross Abstract: Key-value (KV) cache memory management is the primary bottleneck limiting throughput and cost-efficiency in large-scale GPU inference serving. Current systems suffer from three compounding inefficiencies: (1) the absence...

<https://arxiv.org/abs/2604.26968>

22 CONF

**Playing Games with My Heart: An Evaluation of AI Companion Apps**

ArXiv CS.AI 50%

arXiv:2605.08093v2 Announce Type: replace-cross Abstract: The use of chatbots for various forms of companionship is growing rapidly, raising a myriad of questions about simulated relationships, emotional dependence, and psychological harm. While major platforms such as ChatGPT,...

<https://arxiv.org/abs/2605.08093>

23 CONF

**On Seeding Watermarks to Detect Verbatim LLM Copy-Paste Responses**

ArXiv CS.AI 50%

arXiv:2605.16336v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have made fluent essay writing, code drafting, and quiz answering instantly available to students at every level, from secondary school through graduate study. Many educators do not object to...

<https://arxiv.org/abs/2605.16336>

24 CONF

**Multi-Legal-Bench: Evaluating LLMs on Legal Reasoning Across Jurisdictions, Languages, and Legal Traditions**

ArXiv CS.AI 50%

arXiv:2605.29738v2 Announce Type: replace-cross Abstract: Legal NLP benchmarks overwhelmingly evaluate a single language or aggregate tasks that differ fundamentally across jurisdictions, making cross-lingual comparison impossible. We introduce Multi-Legal-Bench, the first...

<https://arxiv.org/abs/2605.29738>

25 CONF

**Compositional Behavioral Semantics for State Abstraction in Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2606.25357v2 Announce Type: replace-cross Abstract: State abstraction plays a key role in scaling reinforcement learning to complex but structured systems. In studying such systems, a wide range of behavioral structures have been studied in reinforcement learning,...

<https://arxiv.org/abs/2606.25357>

26 CONF

**A Physics-Inspired Classical Digital Twin of Cortical Dynamics: A Band-Stratified Metriplectic Port-Hamiltonian Neural Network Learned from Brain-Computer-Interface EEG**

ArXiv CS.AI 50%

arXiv:2607.10439v3 Announce Type: replace-cross Abstract: We present a physics-inspired classical digital twin of brain-computer-interface (BCI) data: a graph neural network constrained to a band-stratified, metriplectic port-Hamiltonian form, with parameters learned from scalp...

<https://arxiv.org/abs/2607.10439>

27 CONF

**IFCLoRA: Topology-Aware Rank Allocation for Parameter-Efficient Fine-Tuning**

ArXiv CS.AI 50%

arXiv:2607.22251v2 Announce Type: replace-cross Abstract: Low-Rank Adaptation (LoRA) is a widely used approach to parameter-efficient fine-tuning (PEFT) of LLMs whose effectiveness depends on rank allocation. Existing adaptive LoRA methods derive ranks from local gradient,...

<https://arxiv.org/abs/2607.22251>

28 CONF

**Automated Numerical Stability Analysis of Deep Learning Operators**

ArXiv CS.AI 50%

arXiv:2607.25494v3 Announce Type: replace-cross Abstract: Finite-precision arithmetic unavoidably introduces numerical approximation errors. Numerical computations may use insufficient precision or an improper formulation, which leads to numerical instability. In this paper, we...

<https://arxiv.org/abs/2607.25494>

29 CONF

**WAM-Diff2: Hierarchical AR-to-Diffusion Distillation for Highly Efficient Autonomous Driving VLA**

ArXiv CS.AI 50%

arXiv:2608.01035v2 Announce Type: replace-cross Abstract: Vision-Language-Action (VLA) models have emerged as a prominent paradigm for end-to-end autonomous driving; however, their efficient deployment is severely constrained by high computational latency and exposure bias...

<https://arxiv.org/abs/2608.01035>

30 CONF

**Ranking Image Fusion the Way Humans Do: A Learned Pairwise Preference Measure for Infrared-Visible Fusion Assessment**

ArXiv CS.AI 50%

arXiv:2608.01301v3 Announce Type: replace-cross Abstract: Infrared-visible image fusion (IVIF) has no ideal fused reference, so algorithms are ranked by scalar objective metrics that formalize proxies for information transfer, structure, or source similarity. These proxies often...

<https://arxiv.org/abs/2608.01301>

31 CONF

**Is Self-Pretraining really useful to improve diagnosis in medical Time Series?**

ArXiv CS.AI 50%

arXiv:2608.06122v2 Announce Type: replace-cross Abstract: Inspired by recent evidence that transformer architectures benefit from Self-PreTraining (SPT) on long-context benchmarks, we investigate whether similar gains extend to multimodal, multivariate, and even simple...

<https://arxiv.org/abs/2608.06122>