



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	136	Media
HackerNews	11	Community
TechCrunch AI	10	Media
The Verge AI	5	Media
MIT Tech Review	1	Media

CONFIRMED SIGNALS — 163 stories

CONF

Show HN: Line9 - A Mermaid rendering engine with its own layout

HackerNews 50% HN 84

Hi HN, Today I'm showcasing Trunchbull, a benchmarking platform designed for authoring benchmarks and running them against different models. We have direct support for benchmarks that use the harbor authoring system, custom tool authoring via the vercel ai sdk and configuration...

<https://line9.ai/diagram>

CONF

Show HN: Meteor shower, planet alignment, eclipse kit

HackerNews 50% HN 8

I made this little walkthrough for finding the visible planets tonight and a meteor shower shooting start counter. Plus a little model of the solar system to try to wrap my head around what's going on. It's all based on your location if you provide it. lifesprites.com/solar.html...

<https://lifesprites.com/share/0ACA25>

CONF

Show HN: I asked Claude to write DDIA as a live comic

HackerNews 50% HN 6

No summary — see link for full article.

<https://systemscomic.com/ddia/read>

CONF

Show HN: Posts grew 6x since ChatGPT, but success rate remained relatively flat

HackerNews 50% HN 5

No summary — see link for full article.

<https://www.orangecrumbs.com/stories/show-hn>

CONF

Show HN: Agent Memory Leaderboard - first public results for AI memory systems

HackerNews 50% HN 3

We just released the first public results from the Agent Memory Leaderboard (AML). The first evaluation focuses on Text Memory across two tracks: - Open-source Methods - Commercial Products 136 teams registered, and 69 representative memory frameworks completed the first...

<https://agentmemoryleaderboard.ai/leaderboard/academic/textual>

CONF

Show HN: CrewScore - coverage check for AI agent prompts (0.6.11)

HackerNews 50% HN 3

No summary — see link for full article.

<https://crewscore.ai/>

CONF

Show HN: WindStrap - Bootstrap 5.3.8 to Tailwind CSS Using DeepSeek V4 Flash

HackerNews 50% HN 2

bootstrap to tailwinds using DeepSeek V4 Flash

<https://github.com/rip747/windstrap>

CONF

Show HN: Harbofly, reclaim the 20-100GB of dev caches on your Mac

HackerNews 50% HN 2

No summary — see link for full article.

<https://harbofly.app>

CONF

Show HN: Tanchi - Open-source AI prospecting agent that only automates email

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/tanchihq/tanchi>

0 CONF

Show HN: FEDERaiDE, a TUI harness with P2P multi-agent routing and built in IDE

HackerNews 50% HN 2

Hi! Federaide is a general purpose multi-agent harness that runs in your terminal. It is meant for recreational programming and automating your scripts. The agents are just named instances of language models which have their own memories, backstories et cetera, and can...

<https://federaide.rocklab.in>

1 CONF

Show HN: Tss - Simple Tailscale Host Selector command-line tool

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/berggren/tss>

2 CONF

Some Claude users are mad that Anthropic's new watermarks will catch them using it at their jobs, classes

TechCrunch AI 50%

Is Anthropic's new watermarking system a travesty? Some have taken to social media to complain that it is.

<https://techcrunch.com/2026/08/12/some-claude-users-are-mad-that-anthropics-new-waterma...>

3 CONF

Amazon will train on Twitch streamers' content by default, unless they opt out

TechCrunch AI 50%

"If this was opt-in, nobody would opt in," Twitch CPO Mike Minton said on a livestream responding to user feedback. "That's honestly the answer."

<https://techcrunch.com/2026/08/12/amazon-will-train-on-twitch-streamers-content-by-defa...>

4 CONF

AI coding startup Cognition reportedly already in talks to raise at \$40B valuation

TechCrunch AI 50%

Cognition may be looking to raise another mega round just a few months after raising \$1 billion at a \$26 billion valuation.

<https://techcrunch.com/2026/08/12/ai-coding-startup-cognition-reportedly-already-in-tal...>

5 CONF

As AI safety concerns mount, three pioneers make the case for staying open

TechCrunch AI 50%

At Ai4, three of the world's most respected AI experts — Geoffrey Hinton, Fei-Fei Li, and Andrew Ng — debated regulation, open source access, and how America can compete as China advances in Asia.

<https://techcrunch.com/2026/08/12/as-ai-safety-concerns-mount-three-pioneers-make-the-c...>

6 CONF

Why Stream ring-maker Sandbar says the future of AI wearables is voice

TechCrunch AI 50%

AI notetaking hardware has taken off over the past couple of years, with credit-card-sized devices, pendants, pins, and even transcribing earbuds all promising to capture your meetings and turn them into summaries and action items. Now, a whole wave of wearables — rings...

<https://techcrunch.com/video/why-stream-ring-maker-sandbar-says-the-future-of-ai-wearab...>

7 CONF

Lovable confirms new \$13.3B valuation, raises another \$400M

TechCrunch AI 50%

This new funding comes after Lovable hit \$500 million in annualized run rate revenue in June, the startup told TechCrunch.

<https://techcrunch.com/2026/08/12/lovable-confirms-new-13-3b-valuation-raises-another-4...>

8 CONF

Why Sandbar thinks it's voice-enabled ring can avoid the AI hardware graveyard

TechCrunch AI 50%

AI notetaking hardware has taken off over the past couple of years, with credit-card-sized devices, pendants, pins, and even transcribing earbuds all promising to capture your meetings and turn them into summaries and action items. Now, a whole wave of wearables — rings...

<https://techcrunch.com/podcast/why-sandbar-thinks-its-voice-enabled-ring-can-avoid-the-...>

9 CONF

Everything announced at Made by Google '26: Pixel 11, Pixel Watch 5, Pixel Tag, and tons of Gemini features

TechCrunch AI 50%

From the Pixel 11 series and a brand new competitor to Apple's AirTag, here are all the announcements from the Made by Google 2026 event.

<https://techcrunch.com/2026/08/12/google-unveils-pixel-11-lineup-new-airtag-rival-and-g...>

0 CONF

AI code-testing startup Blacksmith's valuation jumps almost 10x in less than a year

TechCrunch AI 50%

Blacksmith says revenue has grown more than tenfold over the past year.

<https://techcrunch.com/2026/08/12/blacksmiths-valuation-jumps-10x-to-550m-as-ai-coding-...>

1 CONF

Anthropic says it will watermark text generated by its AI models

TechCrunch AI 50%

Anthropic will extend support for watermarking AI generations for older models as well.

<https://techcrunch.com/2026/08/11/anthropic-says-it-will-watermark-text-generated-by-it...>

2 CONF

Scaling AI agents with trustworthy data

MIT Tech Review 50%

Business and technology leaders need no convincing that the time of agentic AI is here. Organizations are rapidly adopting agents, and few executives doubt the technology's potential to transform work. But many organizations find that realizing the desired return on investment...

<https://www.technologyreview.com/2026/08/12/1141032/scaling-ai-agents-with-trustworthy-...>

3 CONF

Twitch streamers can now opt out from training Amazon's AI

The Verge AI 50%

Twitch users can now opt out of allowing their content to be used to train Amazon's generative AI models. Opting out means that "your streams, VODs, clips, stream chats, and pictures and text on your channel" won't be used in "future training" of an Amazon AI model "whose...

<https://www.theverge.com/tech/979112/twitch-streamers-can-now-opt-out-from-training-ama...>

4 **CONF****Guitar company D'Addario admits that AI music was used in a promotional video****The Verge AI** 50%

After weeks of controversy and speculation, music company D'Addario has admitted that AI, specifically Suno, was used as part of a recent promotional video. For nearly two weeks, the company has denied the allegations, even as evidence piled up against it. It offered various...

<https://www.theverge.com/ai-artificial-intelligence/978982/daddario-guitar-ai-music-suno>
5 **CONF****Google's Pixel Watch 5 dives deeper into AI and health****The Verge AI** 50%

The \$399 Google Pixel Watch 5 isn't about the hardware. Sure, there's a new satin pyrite case finish, a few new strap colors, and a Steph Curry Special Edition. Under the hood, there's a slightly faster Qualcomm processor and an itty-bitty battery bump. There's a \$50 price hike...

<https://www.theverge.com/tech/978094/pixel-watch-5-hands-on-made-by-google-gemini-weara...>
6 **CONF****Of course the ChatGPT dog cancer vaccine spawned a startup****The Verge AI** 50%

Remember that much-hyped story about an Australian tech entrepreneur using ChatGPT, Grok, and other AI tools to craft a personalized cancer vaccine for his dog? Well, surprise: He's launched a startup. That entrepreneur is Paul Conyngham, who says he is launching Gamgee to offer...

<https://www.theverge.com/ai-artificial-intelligence/978671/ai-cured-dog-cancer-mrna-vac...>
7 **CONF****Grok is now an AI 'teammate' you can assign work****The Verge AI** 50%

SpaceXAI has introduced Grok Bot, an always-on AI agent service designed to behave like independent "AI teammates" that can do your work for you. The bots share their own cloud-based computer environment, and can sign into apps, tools, and websites you already use to complete...

<https://www.theverge.com/ai-artificial-intelligence/978666/spacexai-grok-bot-ai-agent-b...>
8 **CONF****VICBench: A Multi-Language Benchmark for Code Vulnerability Detection****ArXiv CS.AI** 50%

arXiv:2608.12246v1 Announce Type: cross Abstract: Evaluating security vulnerability detection tools requires benchmark datasets with vulnerability-inducing commits (VICs) - the commits that first introduce vulnerabilities into codebases. VICs are essential for determining the...

<https://arxiv.org/abs/2608.12246>
9 **CONF****Poor Man's Agentic Modeling: Simulating Large LLM-Agent Societies on a Laptop****ArXiv CS.AI** 50%

arXiv:2608.11215v1 Announce Type: new Abstract: Simulating societies of many large language model (LLM) agents is expensive, yet the questions asked of such simulations are usually macroscopic: phase behaviour, stylised facts, and scaling with the number of agents \$N\$, not the...

<https://arxiv.org/abs/2608.11215>

0 CONF

RealisticTritonBench: A Benchmark for Triton-Kernel Generation in Real-World AI Frameworks

ArXiv CS.AI 50%

arXiv:2608.12004v1 Announce Type: cross Abstract: In modern AI frameworks, GPU kernels are key to overall system performance. Combining usability, portability, and near-handwritten CUDA performance, Triton is widely adopted for implementing GPU kernels. Recent advances show the...

<https://arxiv.org/abs/2608.12004>

1 CONF

LLMs in Process Diagram Engineering: From Optimal PFDs to Validated P&IDs

ArXiv CS.AI 50%

arXiv:2608.11220v1 Announce Type: new Abstract: Nowadays, the creation of a process flow diagram (PFD) and its subsequent transformation into a piping and instrumentation diagram (P&ID) is predominantly performed manually. Applying artificial intelligence in the task could...

<https://arxiv.org/abs/2608.11220>

2 CONF

Harnessing agent memory to build lifelong AI partners for materials scientists

ArXiv CS.AI 50%

arXiv:2608.11224v1 Announce Type: new Abstract: Materials research advances through accumulated experience - scripts that work, protocols that are trusted, warnings attached to failed calculations or experiments, and judgement that links a new question to an old result. This...

<https://arxiv.org/abs/2608.11224>

3 CONF

AgenticTwin: An Agentic LLM Framework Integrated with Digital Twin for Anomaly Detection

ArXiv CS.AI 50%

arXiv:2608.11679v1 Announce Type: new Abstract: Digital twins are increasingly used to monitor and simulate the behavior of cyber-physical systems. Even with skilled operators, interpreting anomalies detected within digital twin pipelines is challenging, as the sheer complexity...

<https://arxiv.org/abs/2608.11679>

4 CONF

Cutting AI Datacenter Energy with Reinforcement Learning: Measured Power Control of LLM Training from One GPU to the Fleet

ArXiv CS.AI 50%

arXiv:2608.11226v1 Announce Type: new Abstract: Reinforcement-learning post-training dominates modern language-model development, yet its power behavior on GPU hardware has not been characterized, and datacenters manage GPU power with workload-blind mechanisms, static caps and...

<https://arxiv.org/abs/2608.11226>

5 CONF

LinearKV: One Cached State Suffices for Position-Independent Caching in Hybrid LLMs

ArXiv CS.AI 50%

arXiv:2608.11231v1 Announce Type: new Abstract: LLM serving is increasingly accelerated by position-independent caching (PIC). Existing PIC methods, however, are built for full-attention models, where a token-indexed KV cache underlies its core operations: matching reusable...

<https://arxiv.org/abs/2608.11231>

6 CONF

APEX: Adaptive Expert Prefetching for Memory-Efficient Edge MoE Inference

ArXiv CS.AI 50%

arXiv:2608.11688v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) models are attractive for edge deployment because they provide high model capacity while activating only a small subset of parameters per token, improving compute efficiency. However, MoE inference at the...

<https://arxiv.org/abs/2608.11688>

7 CONF

Towards Query-Agnostic RAG Evaluation via Query Coverage and Claim Verifiability

ArXiv CS.AI 50%

arXiv:2608.11238v1 Announce Type: new Abstract: Retrieval-augmented generation improves the factuality of large language models by grounding responses in retrieved evidence, yet existing evaluation frameworks struggle to provide consistent, fine-grained diagnostics across the...

<https://arxiv.org/abs/2608.11238>

8 CONF

VQ-bench: A Composable Vector Quantization Framework

ArXiv CS.AI 50%

arXiv:2608.11240v1 Announce Type: new Abstract: Vector quantization is an old problem but has recently become central to AI infrastructure. It is therefore experiencing a surge of renewed engineering and research activity. This paper provides a unified framework for developing...

<https://arxiv.org/abs/2608.11240>

9 CONF

RecSys Factory: Bounding LLM Agent Autonomy to Decision Points in the Industrial Recommender Lifecycle

ArXiv CS.AI 50%

arXiv:2608.11241v1 Announce Type: new Abstract: Deploying LLM agents into industrial recommender operations exposes a three-way tension we frame as the autonomy-determinism-efficiency trilemma: general autonomy (interpreting operator intent, generating glue code zero-shot),...

<https://arxiv.org/abs/2608.11241>

0 CONF

The Off-Support Barrier: Why Semantic Safety Constraints Are Not Learning-Problem Invariants, and What Follows for Prior Design, Containment, and Verification

ArXiv CS.AI 50%

arXiv:2608.11243v1 Announce Type: new Abstract: We argue that a single structural fact organizes a wide range of phenomena in contemporary AI safety: a semantic safety constraint (e.g., the agent does not escape its sandbox) is an off-support object. Formally, if q is the data...

<https://arxiv.org/abs/2608.11243>

1 CONF

Constructing Dynamic Master Logic Models as Knowledge Graphs for Complex System Diagnostics Using Retrieval-Augmented Large Language Models

ArXiv CS.AI 50%

arXiv:2608.12304v1 Announce Type: new Abstract: Dynamic Master Logic (DML) provides a hierarchical framework for representing system behavior by linking functional objectives to underlying structural elements. However, DML construction typically relies on expert interpretation...

<https://arxiv.org/abs/2608.12304>

2 **CONF****Reproducing, Analyzing, and Detecting Reward Hacking in Rubric-Based Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2606.04923v2 Announce Type: replace-cross Abstract: Rubric-based reinforcement learning (RL) uses an LLM-as-a-Judge (LaaJ) to score model outputs according to rubrics as rewards. However, policy models may exploit latent biases in the judge, leading to reward hacking and...

<https://arxiv.org/abs/2606.04923>
3 **CONF****Social Meaning in Large Language Models: Structure, Magnitude, and Pragmatic Prompting**

ArXiv CS.AI 50%

arXiv:2604.02512v2 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly exhibit human-like patterns of pragmatic and social reasoning. This paper addresses two related questions: do LLMs approximate human social meaning not only qualitatively but also...

<https://arxiv.org/abs/2604.02512>
4 **CONF****Program Semantic Inequivalence Game with Large Language Models**

ArXiv CS.AI 50%

arXiv:2505.03818v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) can achieve strong performance on everyday coding tasks, but they can fail on complex tasks that require non-trivial reasoning about program semantics. Finding training examples to teach LLMs...

<https://arxiv.org/abs/2505.03818>
5 **CONF****Symbolic Machine Learning for Vapor-Liquid Equilibrium Prediction in Cx-N2 Binary Mixtures**

ArXiv CS.AI 50%

arXiv:2608.11255v1 Announce Type: new Abstract: Accurate prediction of vapor-liquid equilibrium (VLE) for hydrocarbon-nitrogen mixtures remains challenging for cubic equations of state, particularly across broad ranges of composition and hydrocarbon chain length. While deep...

<https://arxiv.org/abs/2608.11255>
6 **CONF****Glance, Scrutinize, and Think: Advancing Video Anomaly Detection from Training-Free to Agentic Reasoning**

ArXiv CS.AI 50%

arXiv:2608.11260v1 Announce Type: new Abstract: Video Anomaly Detection (VAD) aims to identify anomalous events and localize their temporal intervals. Existing approaches exhibit a "when-what" dissociation: traditional DNN-based methods localize when anomalies occur but lack...

<https://arxiv.org/abs/2608.11260>
7 **CONF****MBA: Multimodal Benchmark and Agents for Real-World Business Ideation**

ArXiv CS.AI 50%

arXiv:2608.11616v1 Announce Type: new Abstract: Agentic systems powered by large language models (LLMs) have opened new opportunities for business ideation. Yet existing approaches remain confined to a text-only paradigm, despite the inherently multimodal nature of real-world...

<https://arxiv.org/abs/2608.11616>

8 CONF

Can Frontier LLMs Match Natively Multimodal Embeddings? A Comparison on Hard-Negative Text-to-Image Retrieval

ArXiv CS.AI 50%

arXiv:2608.11343v1 Announce Type: new Abstract: Multimodal retrieval and classification across different types of media, spanning text, images, video and audio, has traditionally relied on dual-encoder models that align visual and textual representations through contrastive...

<https://arxiv.org/abs/2608.11343>

9 CONF

M-Net: Integrating Spectral Features and Physical Field Operators into Deep Learning for Medical Image Segmentation

ArXiv CS.AI 50%

arXiv:2608.12196v1 Announce Type: cross Abstract: Purpose: Deep learning-based medical image segmentation has achieved remarkable success, yet purely data-driven approaches often fail to exploit the rich mathematical structure inherent in medical images. We investigate whether...

<https://arxiv.org/abs/2608.12196>

0 CONF

When Self-Consistency Backfires: Majority Vote Hurts the Majority of Hard Science Problems for Small LLMs

ArXiv CS.AI 50%

arXiv:2608.11403v1 Announce Type: new Abstract: Self-consistency (SC) via majority vote is a widely used way to spend inference-time compute: sample N chains of thought, return the plurality answer. On the full GPQA Diamond benchmark (198 graduate-level science questions),...

<https://arxiv.org/abs/2608.11403>

1 CONF

Social Chain of Thought: A Multi-Agent Architecture Grounded in Medical Differential Diagnosis Methodology

ArXiv CS.AI 50%

arXiv:2608.11420v1 Announce Type: new Abstract: Medical diagnostic reasoning is a high-impact use case for LLMs that carries significant implications for the health and wellbeing of users. When OpenAI (2026) reports that more than 5% of ChatGPT messages globally are...

<https://arxiv.org/abs/2608.11420>

2 CONF

On Benchmarking Human-Like Intelligence in Machines

ArXiv CS.AI 50%

arXiv:2502.20502v2 Announce Type: replace Abstract: Recent advances in Artificial Intelligence (AI) have yielded powerful computational models that, by learning from vast amounts of human-generated data, are increasingly posited as approximate models of human cognition. However,...

<https://arxiv.org/abs/2502.20502>

3 CONF

Socioduality: A Relational Process Framework for Human-AI Interaction

ArXiv CS.AI 50%

arXiv:2608.11322v1 Announce Type: cross Abstract: Human-AI research often evaluates individual capabilities, combined performance, or final outputs, but these approaches do not preserve how one party's response becomes part of the conditions under which the other party's next...

<https://arxiv.org/abs/2608.11322>

4 **CONF****Localizing Safety Alignment: MLP Layers and Mid-Network Blocks Encode Refusal Behavior in Large Language Models**

ArXiv CS.AI 50%

arXiv:2608.11583v1 Announce Type: new Abstract: Safety alignment in large language models is often treated as a distributed property of the entire network, yet its practical brittleness suggests that refusal behavior may be concentrated in a smaller set of parameters. This work...

<https://arxiv.org/abs/2608.11583>

5 **CONF****EnterpriseRAG: Benchmarking LLM Instruction Adherence and Robustness under Non-Ideal Enterprise Retrieval**

ArXiv CS.AI 50%

arXiv:2608.11584v1 Announce Type: new Abstract: Enterprise RAG deployments face a critical reliability gap: while LLMs satisfy 80% of individual constraints, only 26.8% of responses meet all requirements simultaneously, revealing a 57-point orchestration gap. Existing benchmarks...

<https://arxiv.org/abs/2608.11584>

6 **CONF****XBridge: Entity-Grounded Latent Bridge for Heterogeneous LLM Communication**

ArXiv CS.AI 50%

arXiv:2608.11676v1 Announce Type: new Abstract: Heterogeneous multi-agent LLM systems, where agents are powered by different model families, can outperform homogeneous configurations by reducing redundant reasoning patterns. Yet existing communication protocols either operate...

<https://arxiv.org/abs/2608.11676>

7 **CONF****Making Your LLMs More Objective: Stabilizing LLM Safety Behavior Across Traits with Trait-Invariant Safety Tuning**

ArXiv CS.AI 50%

arXiv:2608.11705v1 Announce Type: new Abstract: Aligned large language models (LLMs) are expected to exhibit safety behavior based on the content of the user request: they should refuse unsafe requests and comply with safe ones. However, we show that the same request can elicit...

<https://arxiv.org/abs/2608.11705>

8 **CONF****Agent Skills Can Be Harmful: An Empirical Study of Skill-Induced Failures in LLM Agents**

ArXiv CS.AI 50%

arXiv:2608.11888v1 Announce Type: new Abstract: Agent skills are the de facto mechanism for extending LLM agents with reusable guidance. A skill can shape the agent's task execution, including planning, tool use, problem-solving, and validation. Prior work reported mixed results...

<https://arxiv.org/abs/2608.11888>

9 **CONF****OEIS Open: How many conjectures can language models turn into theorems?**

ArXiv CS.AI 50%

arXiv:2608.11941v1 Announce Type: new Abstract: We construct OEIS Open, a benchmark based on 492 open mathematical conjectures from the OEIS, formalized in Lean by Tsoukalas et al. Whereas these conjectures had previously been attempted only with a bespoke agent, our open-source...

<https://arxiv.org/abs/2608.11941>

- 0 **CONF**
- Retry, Switch, or Abstain? Learning Strategy-Aware Tool-Use Policies via Controlled Error Injection**
- ArXiv CS.AI 50%
- arXiv:2608.11977v1 Announce Type: new Abstract: Tool-using LLM agents are commonly trained and evaluated in environments where tool calls succeed reliably, yet deployed tools can fail transiently, persistently, or silently. Robust recovery therefore requires more than repeated...
- <https://arxiv.org/abs/2608.11977>
-
- 1 **CONF**
- Claim-Level Reliability Assessment for Efficient Test-Time Reasoning**
- ArXiv CS.AI 50%
- arXiv:2608.11994v1 Announce Type: new Abstract: We propose claim-level falsification as a principle for test-time scaling and instantiate it through Claim-Level Reliability Assessment (CLR), a training-free framework that reallocates test-time compute from additional solution...
- <https://arxiv.org/abs/2608.11994>
-
- 2 **CONF**
- CTBench: Evaluating Troubleshooting Capabilities of AI Agents in Realistic Telecom Network Operations**
- ArXiv CS.AI 50%
- arXiv:2608.12002v1 Announce Type: new Abstract: Agents are increasingly considered for automating network operations and maintenance, where engineers must diagnose network faults, optimize configurations to enhance services, and reduce operational costs while acting under strict...
- <https://arxiv.org/abs/2608.12002>
-
- 3 **CONF**
- Mechanist: AI as a Scientific Instrument for Discovering the Mechanisms of Intelligence**
- ArXiv CS.AI 50%
- arXiv:2608.12036v1 Announce Type: new Abstract: AI models have achieved remarkable success across diverse domains, yet the mechanisms underlying their capabilities and the risks they may pose remain poorly understood. As AI development becomes faster and increasingly automated,...
- <https://arxiv.org/abs/2608.12036>
-
- 4 **CONF**
- Graph-Structured Rubrics: Compiling Rubrics into Typed Evaluation Graphs for LLM Judges**
- ArXiv CS.AI 50%
- arXiv:2608.12097v1 Announce Type: new Abstract: Rubric-based evaluators commonly treat rubrics as prompt context or flat criteria: they specify what to judge but leave criterion composition implicit, even when natural-language rules state it. We introduce Graph-Structured...
- <https://arxiv.org/abs/2608.12097>
-
- 5 **CONF**
- Who Thinks Best Depends on How Long You Let Them: Budget-Dependent Rankings in LLM Evaluation**
- ArXiv CS.AI 50%
- arXiv:2608.12150v1 Announce Type: new Abstract: Standard evaluation of large language models assumes stable model rankings across inference conditions. We challenge this assumption by varying the token generation budget, i.e., the maximum tokens a model may produce, across seven...
- <https://arxiv.org/abs/2608.12150>

6 **CONF****Evaluating LLM Generated Detection Rules in Cybersecurity****ArXiv CS.AI** **50%**

arXiv:2509.16749v1 Announce Type: cross Abstract: LLMs are increasingly pervasive in the security environment, with limited measures of their effectiveness, which limits trust and usefulness to security practitioners. Here, we present an open-source evaluation framework and...

<https://arxiv.org/abs/2509.16749>

7 **CONF****GRPO for Financial Advice Generation: Outperforming Commercial LLMs under CATE Evaluation****ArXiv CS.AI** **50%**

arXiv:2608.11787v1 Announce Type: cross Abstract: Generating actionable financial advice from business records demands that models integrate numerical reasoning, domain knowledge, and sound judgment, while avoiding recommendations that could harm the business. Direct supervision...

<https://arxiv.org/abs/2608.11787>

8 **CONF****Variable Selection in the Context of AI Fairness****ArXiv CS.AI** **50%**

arXiv:2608.11251v1 Announce Type: cross Abstract: Fairness in AI systems has become more important with recent regulatory demands, such as the EU AI Act. Traditional approaches often do not take into account philosophical ethics and social awareness. Variable selection...

<https://arxiv.org/abs/2608.11251>

9 **CONF****Methodologies for Improving the Quality of AI Tutoring in K-12 Education****ArXiv CS.AI** **50%**

arXiv:2608.11259v1 Announce Type: cross Abstract: Many AI tutors leverage large language models (LLMs) today. Given that LLMs are opaque black boxes, robust evaluation and live experimentation to measure the impact of every change are essential. We pioneered AI-powered tutoring...

<https://arxiv.org/abs/2608.11259>

0 **CONF****Uncertainty-Aware and Explainable Ensemble Deep Learning Framework for Multi-Class Skin Lesion Classification****ArXiv CS.AI** **50%**

arXiv:2608.11280v1 Announce Type: cross Abstract: Skin cancer diagnosis from dermoscopic images remains challenging due to high intra-class variability, inter-class similarity, class imbalance, and the limited interpretability of deep learning models. This paper proposes an...

<https://arxiv.org/abs/2608.11280>

1 **CONF****Chemically Meaningful Textualization Enables Explainable Validation of Metal-Organic Frameworks by Large Language Models****ArXiv CS.AI** **50%**

arXiv:2608.11283v1 Announce Type: cross Abstract: Computation-ready metal-organic framework (MOF) databases are essential for high-throughput screening, yet many reported crystal structures remain chemically unreasonable or disordered, compromising simulation fidelity. Existing...

<https://arxiv.org/abs/2608.11283>

2 **CONF****CLEAR: Class-wise Expert Aggregation with Structured Sampling for Long-Tailed Classification****ArXiv CS.AI** **50%**

arXiv:2608.11287v1 Announce Type: cross Abstract: Long-tailed classification poses a reliability challenge because models trained on imbalanced data are unevenly reliable across frequent and underrepresented classes. While existing methods address imbalance through re-balancing,...

<https://arxiv.org/abs/2608.11287>

3 **CONF****Backdoor Decontamination Dynamics in LLM Agents****ArXiv CS.AI** **50%**

arXiv:2608.11295v1 Announce Type: cross Abstract: Open-weight LLM agents are vulnerable to backdoors installed during fine-tuning, which may be undetectable if the trigger conditions are never met during testing. Assuming defenders do not know the existing trigger, they cannot...

<https://arxiv.org/abs/2608.11295>

4 **CONF****Clinical Feasibility of Low-Magnification Fluorescence Imaging for Breast Cancer Margin Detection Using Texture Analysis and Deep Learning****ArXiv CS.AI** **50%**

arXiv:2608.11317v1 Announce Type: cross Abstract: High-resolution images of unprocessed surgical breast tissue can be obtained using microscopy with ultraviolet surface excitation (MUSE). This technique is considered a promising method for checking surgical margins during breast...

<https://arxiv.org/abs/2608.11317>

5 **CONF****Contextual Quality-Diversity Evolutionary Reinforcement Learning for HVAC Control in Tropical Commercial Buildings****ArXiv CS.AI** **50%**

arXiv:2608.11324v1 Announce Type: cross Abstract: This paper proposes a contextual quality-diversity evolutionary reinforcement-learning controller, CQD-ERL, for the supervisory control of a tropical, water-cooled chiller plant and its associated air side. Rather than converging...

<https://arxiv.org/abs/2608.11324>

6 **CONF****Dual-Domain Cross-Modal Decoding for Clinical Text-Guided Medical Image Segmentation****ArXiv CS.AI** **50%**

arXiv:2608.11335v1 Announce Type: cross Abstract: Clinical text can narrow down what to segment, but recent text-guided designs emphasize spatial alignment while overlooking frequency content that governs texture and boundaries. We propose Dual-Domain Cross-Modal Decoding...

<https://arxiv.org/abs/2608.11335>

7 **CONF****Governing Agentic AI in FinTech****ArXiv CS.AI** **50%**

arXiv:2608.11344v1 Announce Type: cross Abstract: Financial institutions are delegating consequential decisions to agentic AI systems that decompose goals, coordinate models and tools, and act with little oversight. Yet agentic AI governance in FinTech is under-investigated. We...

<https://arxiv.org/abs/2608.11344>

8 CONF

AI Guardrail Survival under Single-Cycle Agentic Self-Summarization

ArXiv CS.AI 50%

arXiv:2608.11392v1 Announce Type: cross Abstract: Long-running agents periodically compact their context, replacing the transcript with a model-generated summary. Recent work shows that dropping a standing safety constraint during compaction drives behavioral violations...

<https://arxiv.org/abs/2608.11392>

9 CONF

TangPoetryBench: A Multi-Dimensional Benchmark and Rubric-Conditioned Evaluator for Poetry-to-Image Generation

ArXiv CS.AI 50%

arXiv:2608.11452v1 Announce Type: cross Abstract: Text-to-image (T2I) models are increasingly asked to illustrate literary and cultural content, yet we cannot measure how well an image renders the meaning of a poem. The task is many-sided: a good illustration must be visually...

<https://arxiv.org/abs/2608.11452>

0 CONF

The Next Challenge for Agentic Cybersecurity: A Realistic, Contamination-Free Reverse Engineering Benchmark

ArXiv CS.AI 50%

arXiv:2608.11469v1 Announce Type: cross Abstract: AI agents are rapidly improving in cybersecurity capabilities when the source code is available for analysis, yet much of the software most consequential to cybersecurity, including malware, firmware, and proprietary...

<https://arxiv.org/abs/2608.11469>

1 CONF

HyperFix: Combinatorial Nonlinear Correction for Task Vector Merging

ArXiv CS.AI 50%

arXiv:2608.11499v1 Announce Type: cross Abstract: Task vectors enable model merging without joint retraining. In practice, the subset of task vectors to be merged may vary, but many existing methods use scalar tuning for a particular subset, requiring repeated tuning across...

<https://arxiv.org/abs/2608.11499>

2 CONF

Conflict and Congruency Effects in Large Language Models: In-Weight and In-Context Competition in a Verbal Conflict Task

ArXiv CS.AI 50%

arXiv:2608.11510v1 Announce Type: cross Abstract: Congruency effects, observed in conflict tasks such as Stroop and flanker tasks, have been investigated for nearly a century in psychology and neuroscience, but their mechanistic basis is not fully understood. We introduce a...

<https://arxiv.org/abs/2608.11510>

3 CONF

Let it Cook: Learning to Wait in Sequential Decision Making

ArXiv CS.AI 50%

arXiv:2608.11511v1 Announce Type: cross Abstract: In sequential decision making, an agent typically observes its environment and acts at every timestep. However, such active participation may not always be necessary; tasks such as brewing coffee include periods that are served...

<https://arxiv.org/abs/2608.11511>

4 **CONF****Do Influence Tactics Matter? Investigating Prompt Framing Effects in LLM Code Generation****ArXiv CS.AI** **50%**

arXiv:2608.11513v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly integrated into software engineering workflows, helping developers write, debug, test, and maintain code. While prompt wording and structure are known to influence model performance,...

<https://arxiv.org/abs/2608.11513>5 **CONF****Generative Semantic Segmentation via an Observable Semantic-Image Interface and Hierarchical Generator Evidence Alignment****ArXiv CS.AI** **50%**

arXiv:2608.11537v1 Announce Type: cross Abstract: Generative semantic segmentation exposes structured predictions as images, but direct color decoding is susceptible to color drift and boundary mixing, whereas latent-feature decoders that predict a separate output distribution...

<https://arxiv.org/abs/2608.11537>6 **CONF****A Conceptual Framework for Enhancing Workforce Readiness for Smart Manufacturing in the AI Era****ArXiv CS.AI** **50%**

arXiv:2608.11540v1 Announce Type: cross Abstract: The convergence of artificial intelligence (AI), Industrial Internet of Things, cyber-physical systems, and advanced robotics is reshaping manufacturing faster than engineering curricula can adapt, widening the gap between the...

<https://arxiv.org/abs/2608.11540>7 **CONF****Beyond Single-Turn Confidence: Trajectory-Adapted Uncertainty Quantification for LLM Agents****ArXiv CS.AI** **50%**

arXiv:2608.11552v1 Announce Type: cross Abstract: Uncertainty quantification (UQ) methods for language models are typically evaluated on single-turn outputs, where uncertainty is attached to one generated answer. For LLM agents, however, the unit of observation is an interactive...

<https://arxiv.org/abs/2608.11552>8 **CONF****From Synthesis to Removal: Physics-Grounded Reflection Simulation and Diffusion-Based Video Dereflection****ArXiv CS.AI** **50%**

arXiv:2608.11562v1 Announce Type: cross Abstract: Videos captured through glass often contain reflections that degrade visual quality and interfere with downstream vision tasks. Although single-image reflection removal has been extensively studied, video reflection removal...

<https://arxiv.org/abs/2608.11562>9 **CONF****Reinforcing Step-level Reasoning for Effective Self-Correction in LLMs****ArXiv CS.AI** **50%**

arXiv:2608.11573v1 Announce Type: cross Abstract: Achieving effective self-correction, where models verify and correct their own mistakes, remains a fundamental challenge for large language models (LLMs). In this work, we propose Self-Fix Step-DPO (SFS-DPO), a reinforcement...

<https://arxiv.org/abs/2608.11573>

0 CONF

A Hybrid Framework of Vision Transformer and Gated Recurrent Unit for Detection of Mosquito Diseases

ArXiv CS.AI 50%

arXiv:2608.11582v1 Announce Type: cross Abstract: Identifying dengue virus-infected mosquitoes from control mosquitoes is a major challenge in analyzing mosquito locomotion behavior due to the small size and complexity of the video background. Conventional AI methods are often...

<https://arxiv.org/abs/2608.11582>

1 CONF

FM-LLM: A frequency-enhanced mixture-of-experts framework for adapting LLMs to time series forecasting

ArXiv CS.AI 50%

arXiv:2608.11623v1 Announce Type: cross Abstract: Recent advances in Large Language Models (LLMs) have spurred cross-modal solutions for time-series forecasting. However, existing methods rely heavily on textual prompts for modality alignment-introducing nontrivial computational...

<https://arxiv.org/abs/2608.11623>

2 CONF

Learning to Persuade Exposes How Easily LLMs Abandon Correct Beliefs

ArXiv CS.AI 50%

arXiv:2608.11624v1 Announce Type: cross Abstract: Persuasion is a core dynamic of natural language communication, shaping how large language models (LLMs) update beliefs, resolve disagreements, and reach decisions. As LLMs increasingly debate, advise, and think collaboratively...

<https://arxiv.org/abs/2608.11624>

3 CONF

Deep Learning Based Relative Transfer Matrix Estimation for Multiple Sources and Multiple Microphones

ArXiv CS.AI 50%

arXiv:2608.11627v1 Announce Type: cross Abstract: The Relative Transfer Matrix (ReTM), recently introduced as a generalization of the relative transfer function for multiple receivers and sources, shows promising performance when applied to speech enhancement in noisy...

<https://arxiv.org/abs/2608.11627>

4 CONF

Beyond Memory: A Transactional Continuity Kernel for Long-Lived AI Agents

ArXiv CS.AI 50%

arXiv:2608.11632v1 Announce Type: cross Abstract: Persistent AI agents accumulate versioned state across long horizons, but storage retention alone does not identify authoritative state. Without an explicit control plane, unmediated updates by models, tools, and background...

<https://arxiv.org/abs/2608.11632>

5 CONF

Motion-as-Prompt: Enhancing Motion Reasoning in Multimodal Large Language Models via Motion-Guided Cross-Frame Visual Prompting

ArXiv CS.AI 50%

arXiv:2608.11655v1 Announce Type: cross Abstract: Motion-centric video reasoning is fundamental to interactive applications such as robotic manipulation and autonomous navigation. However, multimodal large language models (MLLMs) typically process videos through sparse uniform...

<https://arxiv.org/abs/2608.11655>

6 CONF

Causal Agent based on Large Language Model

ArXiv CS.AI 50%

arXiv:2408.06849v3 Announce Type: replace Abstract: The large language model (LLM) has achieved significant success across various domains. However, the inherent complexity of causal problems and causal theory poses challenges in accurately describing them in natural language,...

<https://arxiv.org/abs/2408.06849>

7 CONF

Is Per-Agent Policy Composition Safe? Rethinking Successor-Feature Transfer in Cooperative Multi-Agent Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2608.11658v1 Announce Type: cross Abstract: Many reinforcement learning systems, from fleet management to traffic signal control, must serve an objective that changes dynamically after deployment, and retraining a policy for each new objective is prohibitively expensive....

<https://arxiv.org/abs/2608.11658>

8 CONF

GCPO: Diagnosing and Constraining Subspace Geometry in Rollout RL for LLMs

ArXiv CS.AI 50%

arXiv:2608.11674v1 Announce Type: cross Abstract: On-policy rollout methods such as GRPO are central to post-training of large language models, yet they frequently suffer from training instabilities, cross-task capability degradation, and response-length inflation. Although...

<https://arxiv.org/abs/2608.11674>

9 CONF

The Wording Effect: Quantifying Two-Way Drift in LLM Benchmark Performance

ArXiv CS.AI 50%

arXiv:2608.11694v1 Announce Type: cross Abstract: A benchmark score comes from a single phrasing of each problem. That single phrasing is treated as if it stood for the whole space of ways the same problem could be asked, but it does not. We show that rephrasing a problem while...

<https://arxiv.org/abs/2608.11694>

00 CONF

When the API Speaks the Wrong Language: Revisiting Post-Training for Multilingual Tool Use

ArXiv CS.AI 50%

arXiv:2608.11715v1 Announce Type: cross Abstract: The reliability of Large Language Models (LLMs) for API calling degrades in multilingual settings. A common failure occurs when a model selects the correct tool but generates argument values in an inconsistent language, which we...

<https://arxiv.org/abs/2608.11715>

01 CONF

Advancing MLLM-based UAV Image Understanding and Reasoning: A Benchmark and a Training-Free Multi-Agent System

ArXiv CS.AI 50%

arXiv:2608.11738v1 Announce Type: cross Abstract: Multimodal Large Language Model (MLLM)-based UAV aerial image understanding and reasoning is essential for aerial intelligence yet poses distinct challenges arising from extreme scale variation, arbitrary camera orientations, and...

<https://arxiv.org/abs/2608.11738>

02 CONF

JieZi: A Large-Scale Expert-Audited Dataset and Benchmark for Ancient Chinese Character Exegesis

ArXiv CS.AI 50%

arXiv:2608.11741v1 Announce Type: cross Abstract: The scholarly exegesis of ancient Chinese characters demands integrating visual observation, linguistic analysis, and historical context. However, existing computational approaches focus narrowly on subtasks such as character...

<https://arxiv.org/abs/2608.11741>

03 CONF

Instruction Alignment for Binary Code Representation Learning

ArXiv CS.AI 50%

arXiv:2608.11766v1 Announce Type: cross Abstract: Binary code representation learning is a fundamental problem in software security and reverse engineering. Existing methods mainly learn function-level embeddings that capture coarse-grained semantic relationships between binary...

<https://arxiv.org/abs/2608.11766>

04 CONF

TELLME: Test-Enhanced Learning for Language Model Enrichment

ArXiv CS.AI 50%

arXiv:2608.11788v1 Announce Type: cross Abstract: Continual pre-training (CPT) has been widely adopted as a method for domain adaptation in large language models. However, CPT has consistently been accompanied by challenges, such as the difficulty of acquiring large-scale...

<https://arxiv.org/abs/2608.11788>

05 CONF

Toward Meaningful Transparency for AI Chatbots: Disclosing Persuasive Intent Reduces Persuasion

ArXiv CS.AI 50%

arXiv:2608.11794v1 Announce Type: cross Abstract: The growing role of AI-generated content and AI-enabled systems in public communication has led regulators to demand clear disclosure of content provenance and AI involvement. But the effects of such disclosures remain uncertain....

<https://arxiv.org/abs/2608.11794>

06 CONF

How China-Origin Vision-Language Models Move from Refusal to Reframing in State Alignment

ArXiv CS.AI 50%

arXiv:2608.11816v1 Announce Type: cross Abstract: State-aligned distortion has been documented in China-origin text-based large language models (LLMs), but whether, and in what form, it arises in multimodal systems has not been systematically examined. We construct a balanced...

<https://arxiv.org/abs/2608.11816>

07 CONF

Redistribution-based Cost Inference Improves Sparse Safe Offline RL

ArXiv CS.AI 50%

arXiv:2608.12306v1 Announce Type: cross Abstract: Safe offline RL typically assumes access to dense per-step cost annotations, but in practice supervisors provide only trajectory-level stop-feedback: a binary signal at the first unsafe transition, with no per-step attribution....

<https://arxiv.org/abs/2608.12306>

08 CONF

Benchmark-Based Comparative Assessment of Publicly Benchmarked Indian Foundation Models: A Capability and Evaluation-Maturity Framework

ArXiv CS.AI 50%

arXiv:2608.11891v1 Announce Type: cross Abstract: Governments increasingly fund indigenous foundation models to strengthen national AI capability, digital sovereignty, and multilingual computing. Assessing the progress of such national ecosystems is complicated by inconsistent...

<https://arxiv.org/abs/2608.11891>

09 CONF

Do You See What You Draw? A Semantic Closed-Loop Framework for Holistic Evaluation of Unified Multimodal Models

ArXiv CS.AI 50%

arXiv:2608.11907v1 Announce Type: cross Abstract: As Large Vision-Language Models increasingly aim to integrate visual generation and understanding within a single parameter space, evaluating such structural unification in a cohesive manner remains a critical challenge. Current...

<https://arxiv.org/abs/2608.11907>

10 CONF

Hamilton-Zero: A Neural Tensor-Network Foundation Model for Ground States of Arbitrary Quadratic Qubit Hamiltonians

ArXiv CS.AI 50%

arXiv:2608.11911v1 Announce Type: cross Abstract: A central promise of useful quantum advantage is the ability to compute ground states of Hamiltonian systems beyond the reach of classical simulation methods. Here we demonstrate that this problem can be effectively amortized...

<https://arxiv.org/abs/2608.11911>

11 CONF

TailBooster: A Dual-Layer Generative Framework for Extreme Value Augmentation with Operational Validity Enforcement

ArXiv CS.AI 50%

arXiv:2608.11951v1 Announce Type: cross Abstract: Extreme events in air transport, such as severe arrival delays and abnormal air times, cause cascading network disruptions with substantial operational, economic, and safety costs. Such events are rare in historical records,...

<https://arxiv.org/abs/2608.11951>

12 CONF

Causal inference for group-contaminated structured outcomes: observable quotients, lossless reduction and exact randomization inference

ArXiv CS.AI 50%

arXiv:2608.11954v1 Announce Type: cross Abstract: Structured potential outcomes such as microscopy images may be recorded after an unknown, unit-specific transformation. If that transformation can depend on treatment, covariates or the intrinsic outcome, raw-coordinate analyses...

<https://arxiv.org/abs/2608.11954>

13 CONF

HCGRec: Hint-Conditioned Generative Recommendation with Semantic IDs

ArXiv CS.AI 50%

arXiv:2608.11980v1 Announce Type: cross Abstract: Semantic-ID generative recommenders represent each item as a short sequence of discrete semantic tokens and predict the next item by autoregressively generating this token sequence. This paradigm enables a unified generation...

<https://arxiv.org/abs/2608.11980>

14 CONF

Remote Sensing and Machine Learning-Based Analysis of Land Use and Vegetation Change in Dhaka District, Bangladesh

ArXiv CS.AI 50%

arXiv:2608.12001v1 Announce Type: cross Abstract: Rapid urbanization in Dhaka District, Bangladesh has triggered substantial alterations in land use and environmental conditions, necessitating systematic monitoring for informed urban planning and ecological sustainability. This...

<https://arxiv.org/abs/2608.12001>

15 CONF

Dual-Model Sentiment Analysis of Consumer Reviews in the Retail Coffee Sector Using Machine Learning and Deep Learning Approaches

ArXiv CS.AI 50%

arXiv:2608.12007v1 Announce Type: cross Abstract: Consumer reviews play an important role in shaping brand perception and business strategies, particularly in service-driven industries such as retail coffee. This study presents a comparative sentiment analysis framework for...

<https://arxiv.org/abs/2608.12007>

16 CONF

From Safety Documentation to Safety Knowledge Support: An Evidence-Grounded LLM Framework for Medical Devices

ArXiv CS.AI 50%

arXiv:2608.12025v1 Announce Type: cross Abstract: Medical devices are becoming more software-intensive, connected, and AI-enabled. Their development requires risk-management evidence aligned with ISO 14971 and, for software, IEC 62304. This evidence must be kept consistent...

<https://arxiv.org/abs/2608.12025>

17 CONF

LoSA: Near-Lossless Sparse Attention for Training-Free Video Diffusion Acceleration

ArXiv CS.AI 50%

arXiv:2608.12032v1 Announce Type: cross Abstract: Video diffusion transformers are costly to sample: every denoising step applies self-attention over a long 3D token sequence, a quadratic cost that dominates as resolution and duration grow. Sparse attention reduces this cost...

<https://arxiv.org/abs/2608.12032>

18 CONF

How Far from Clinical Deployment? Evaluating the Complete Unsupervised Domain Adaptation Pipeline in Medical Imaging

ArXiv CS.AI 50%

arXiv:2608.12035v1 Announce Type: cross Abstract: Deploying unsupervised domain adaptation (UDA) in clinical practice requires choosing which algorithm to use and which of its trained models to ship. However, the deployment (target) domain is unlabeled, so models cannot be...

<https://arxiv.org/abs/2608.12035>

19 CONF

Faithful, Sufficient and Understandable: Rethinking Graph Counterfactual Explanations via Discrete Diffusion Inversion

ArXiv CS.AI 50%

arXiv:2608.12083v1 Announce Type: cross Abstract: Graph Neural Networks (GNNs) achieve strong predictive performance on graph-structured data across domains such as chemistry, biology, and network analysis, yet they provide no intrinsic explanation of their predictions. This...

<https://arxiv.org/abs/2608.12083>

20 **CONF****Confidence Calibration of Deep Learning Systems****ArXiv CS.AI** **50%**

arXiv:2608.12100v1 Announce Type: cross Abstract: In high-stakes applications, reliable confidence estimates are as important as the predictions themselves. Confidence calibration ensures that predicted probabilities reflect the likelihood of correctness, making it essential for...

<https://arxiv.org/abs/2608.12100>

21 **CONF****No One to Blame: A Framework of Constitutive AI Unaccountability****ArXiv CS.AI** **50%**

arXiv:2608.12104v1 Announce Type: cross Abstract: The increasing deployment of autonomous, agentic AI systems challenges traditional accountability mechanisms. Existing research predominantly frames AI accountability gaps as barriers that can be overcome through better...

<https://arxiv.org/abs/2608.12104>

22 **CONF****Ready Cohorts: Bounding GPU Opportunity and Avoiding Host Round Trips in LLM-Agent Control****ArXiv CS.AI** **50%**

arXiv:2608.12123v1 Announce Type: cross Abstract: LLM-agent services repeatedly execute small deterministic transitions between model and tool calls: route an outcome, update state, and emit the next effect. We ask when this control path exposes enough concurrent work for GPU...

<https://arxiv.org/abs/2608.12123>

23 **CONF****Do LLMs Take Care of Their Own? Similarity Signals Can Induce Cooperation****ArXiv CS.AI** **50%**

arXiv:2608.12125v1 Announce Type: cross Abstract: As LLM-based agents with user-instructed goals are becoming widely deployed, they increasingly encounter each other in strategic interactions, and face challenges of finding mutually beneficial outcomes. Prior literature has...

<https://arxiv.org/abs/2608.12125>

24 **CONF****A corpus-specific clinical RAG system matches or outperforms newer frontier LLMs on HealthBench****ArXiv CS.AI** **50%**

arXiv:2608.12138v1 Announce Type: cross Abstract: General-purpose large language models (LLMs) have recently been reported to match or exceed specialized clinical AI tools on medical benchmarks, but such comparisons draw on a narrow set of systems and on benchmarks developed...

<https://arxiv.org/abs/2608.12138>

25 **CONF****Co-constructing sociotechnical AI governance: participatory system mapping using algorithm registers****ArXiv CS.AI** **50%**

arXiv:2608.12166v1 Announce Type: cross Abstract: Algorithm registers have been championed as a means of providing transparency on the use of algorithms in public services. Yet potential publics differ in their expectations of what should be made transparent and how, as well as...

<https://arxiv.org/abs/2608.12166>

26 CONF

HSTGFormer: Hyper Spatial-Temporal Graph Transformer for 3D Human Pose Estimation

ArXiv CS.AI 50%

arXiv:2608.12187v1 Announce Type: cross Abstract: Transformer-based methods have achieved strong performance in monocular 3D human pose estimation, but most existing approaches organise spatial and temporal reasoning as separate stages, which may weaken unified spatial-temporal...

<https://arxiv.org/abs/2608.12187>

27 CONF

Machine Learning-Based Cyber Defense for Cloud Infrastructure: An Adaptive Deep Q-Network Architecture for Intelligent Intrusion Detection and Automated Threat Mitigation

ArXiv CS.AI 50%

arXiv:2608.12190v1 Announce Type: cross Abstract: With the increasing complexity of cyber assaults in cloud environments, adaptable security solutions are needed that can support real-time detection and autonomous response. In this paper, we propose a reinforcement...

<https://arxiv.org/abs/2608.12190>

28 CONF

Information Abundance Paradox: Long-Context Training Undermines Parametric Knowledge

ArXiv CS.AI 50%

arXiv:2608.12218v1 Announce Type: cross Abstract: Large language models are increasingly trained and deployed with long contexts that span documents, code repositories, and interaction histories. This scaling reflects the implicit assumption that training on longer contexts will...

<https://arxiv.org/abs/2608.12218>

29 CONF

SCOUT: Unlocking Enhanced Spatial Reasoning via Structured Chain-of-Thought and Multi-Objective Process Reward

ArXiv CS.AI 50%

arXiv:2608.12220v1 Announce Type: cross Abstract: Existing Vision-Language Models (VLMs) exhibits a critical bottleneck in robust spatial reasoning. Recent reinforcement learning (RL) methods aim to close this gap with verifiable outcomes, yet they suffer from poor credit...

<https://arxiv.org/abs/2608.12220>

30 CONF

Domain-Aware Lightweight Spectral-Grouped Convolutions for Hyperspectral Fish Freshness Classification

ArXiv CS.AI 50%

arXiv:2608.12227v1 Announce Type: cross Abstract: Hyperspectral imaging (HSI) offers nondestructive assessment of fish freshness by detecting biochemical alterations across spectral bands. However, conventional deep learning approaches do not fully address the particular...

<https://arxiv.org/abs/2608.12227>

31 CONF

How Organizations Use AI: Evidence from ChatGPT

ArXiv CS.AI 50%

arXiv:2608.12236v1 Announce Type: cross Abstract: We study how organizations use frontier generative AI by linking ChatGPT Enterprise account records to usage, worker roles, task classifications, and public-company financial data through March 2026. These linked data enable a...

<https://arxiv.org/abs/2608.12236>

32 CONF

HAMP-LIC: Hessian-Aware Mixed-Precision Post-Training Quantization for Learned Image Compression

ArXiv CS.AI 50%

arXiv:2608.12239v1 Announce Type: cross Abstract: Use this plain-text version for the arXiv abstract field: Learned image compression (LIC) models achieve strong rate-distortion performance but are hindered by high computational complexity and encoding-decoding mismatches across...

<https://arxiv.org/abs/2608.12239>

33 CONF

Diagram-MMU: A Multi-Modal Benchmark for Scientific Diagrams

ArXiv CS.AI 50%

arXiv:2608.12262v1 Announce Type: cross Abstract: Multimodal Large Language Models (MLLMs) have been growing the capability for scientific writing and collaboration. For example, OpenAI Prism is a free workspace for scientific writing and collaboration. One important feature in...

<https://arxiv.org/abs/2608.12262>

34 CONF

A Neighborhood Attention Transformer Network for Enhanced 3D Segmentation of the Left Anterior Descending Artery

ArXiv CS.AI 50%

arXiv:2608.12274v1 Announce Type: cross Abstract: Background: Accurate segmentation of the Left Anterior Descending (LAD) artery in 3D free-breathing, non-contrast CT is critical for cardiac dose sparing in thoracic radiotherapy. The LAD is extremely small, has poor soft-tissue...

<https://arxiv.org/abs/2608.12274>

35 CONF

Structural Silence: When AI Infrastructure Fails Speakers of Underrepresented Languages

ArXiv CS.AI 50%

arXiv:2608.12278v1 Announce Type: cross Abstract: Artificial intelligence tools for education and language support are increasingly framed as scalable responses to access gaps in under-resourced communities. Yet the infrastructure underlying these tools, including training...

<https://arxiv.org/abs/2608.12278>

36 CONF

Class Activation Mapping in Explainable Computer Vision: A Method-Centered Review of CNN, Transformer, and Foundation-Model-Era Visual Explanations

ArXiv CS.AI 50%

arXiv:2608.12299v1 Announce Type: cross Abstract: Class activation mapping (CAM) is one of the most widely used visual explanation families in explainable artificial intelligence. Its purpose is intuitive: it converts internal model evidence into a heatmap that highlights the...

<https://arxiv.org/abs/2608.12299>

37 CONF

AI4AI at Test-Time: Strong-to-Weak Capability Transfer via Harnesses

ArXiv CS.AI 50%

arXiv:2608.12307v1 Announce Type: cross Abstract: Recent work on distillation transfers the capabilities of large models to smaller ones often by updating the latter's parameters, through teacher forcing, on-policy distillation, and related training-time methods. In this paper,...

<https://arxiv.org/abs/2608.12307>

38 CONF

DreamFly: Causal Memory and Receding-Horizon Diffusion Planning for Aerial Vision-Language Navigation

ArXiv CS.AI 50%

arXiv:2608.12308v1 Announce Type: cross Abstract: Aerial vision-language navigation (VLN) requires an embodied agent to integrate visual evidence over time, plan future actions, and determine when it has reached a navigation goal under partial observability. Although recent VLA...

<https://arxiv.org/abs/2608.12308>

39 CONF

SteeringSafety: Benchmarking Representation Steering in LLMs Across Safety Perspectives

ArXiv CS.AI 50%

arXiv:2509.13450v3 Announce Type: replace Abstract: We introduce SteeringSafety, a benchmark for evaluating representation steering methods across nine safety perspectives spanning 18 datasets. While prior work highlights the general capabilities of representation steering, we...

<https://arxiv.org/abs/2509.13450>

40 CONF

Behavior and Representation in Open-Weight Large Language Models for Combinatorial Optimization: From Feature Extraction to Algorithm Selection

ArXiv CS.AI 50%

arXiv:2512.13374v2 Announce Type: replace Abstract: Recent advances in Large Language Models (LLMs) open new perspectives for automation in optimization, yet little is known about whether their internal representations capture problem structure or algorithmic behavior. We...

<https://arxiv.org/abs/2512.13374>

41 CONF

Credo: Declarative Control of LLM Pipelines via Beliefs and Policies

ArXiv CS.AI 50%

arXiv:2604.14401v2 Announce Type: replace Abstract: Agentic AI systems are becoming commonplace in domains that require long-lived, stateful decision-making in continuously evolving conditions. As such, correctness depends not only on the output of individual model calls, but...

<https://arxiv.org/abs/2604.14401>

42 CONF

Hallucination Mitigation with Agentic AI, Nested Learning, and AI Sustainability via Semantic Caching

ArXiv CS.AI 50%

arXiv:2605.29055v2 Announce Type: replace Abstract: This paper describes an approach to hallucination detection and mitigation using a HOPE-inspired Nested Learning architecture with Continuum Memory Systems (CMS) and semantic similarity caching, tested on a hybrid benchmark of...

<https://arxiv.org/abs/2605.29055>

43 CONF

Moxia: A Trust-First Neuro-Symbolic Execution Architecture for Self-Explaining Mathematical Reasoning

ArXiv CS.AI 50%

arXiv:2606.00671v3 Announce Type: replace Abstract: We present Moxia (formerly AXIOM), a trust-first neuro-symbolic architecture for self-explaining mathematical reasoning over natural-language input. Its language model is strictly a canonicalizer: it rewrites informal problem...

<https://arxiv.org/abs/2606.00671>

44 CONF

A foundation model of numerical intelligence with cross-disciplinary generalization

ArXiv CS.AI 50%

arXiv:2607.28432v3 Announce Type: replace Abstract: Intelligence is commonly understood as the ability to acquire and apply knowledge, adapt to unfamiliar situations and solve new problems. Large language models exhibit this capacity by inferring task-relevant knowledge from...

<https://arxiv.org/abs/2607.28432>

45 CONF

ComBodied Agents: a New Paradigm of Human-Centric Agentic AI

ArXiv CS.AI 50%

arXiv:2608.10915v2 Announce Type: replace Abstract: After an older adult misses a medication dose, a software agent can send another reminder and an embodied agent can bring the medication. Yet neither explains whether the person forgot, is confused, has side effects, or...

<https://arxiv.org/abs/2608.10915>

46 CONF

ReXrank: A Public Leaderboard for AI-Powered Radiology Report Generation

ArXiv CS.AI 50%

arXiv:2411.15122v2 Announce Type: replace-cross Abstract: AI-driven models have demonstrated significant potential in automating radiology report generation for chest X-rays. However, there is no standardized benchmark for objectively evaluating their performance. To address...

<https://arxiv.org/abs/2411.15122>

47 CONF

Explainability in Practice: A Survey of Explainable NLP Across Various Domains

ArXiv CS.AI 50%

arXiv:2502.00837v3 Announce Type: replace-cross Abstract: Natural Language Processing (NLP) is now embedded in critical sectors including healthcare, finance, and customer relationship management, where models such as GPT-4o, Gemini, and BERT increasingly inform decisions. The...

<https://arxiv.org/abs/2502.00837>

48 CONF

COLORA: Efficient Fine-Tuning for Convolutional Models with a Study Case on Optical Coherence Tomography Image Classification

ArXiv CS.AI 50%

arXiv:2505.18315v3 Announce Type: replace-cross Abstract: We introduce CoLoRA (Convolutional Low-Rank Adaptation), a parameter-efficient fine-tuning method for convolutional neural networks (CNNs). CoLoRA extends LoRA to convolutional layers by decomposing kernel...

<https://arxiv.org/abs/2505.18315>

49 CONF

P2MFDS: A Privacy-Preserving Multimodal Fall Detection System for Elderly People in Bathroom Environments

ArXiv CS.AI 50%

arXiv:2506.17332v2 Announce Type: replace-cross Abstract: By 2050, people aged 65 and over are projected to make up 16% of the global population. As aging is closely associated with increased fall risk, particularly in wet and confined environments such as bathrooms where over...

<https://arxiv.org/abs/2506.17332>

50 **CONF****Small Data Explainer -- The impact of small data methods in everyday life**

ArXiv CS.AI 50%

arXiv:2507.11773v2 Announce Type: replace-cross Abstract: The emergence of breakthrough artificial intelligence (AI) techniques has led to a renewed focus on how small data settings, i.e., settings with limited information, can benefit from such developments. This includes...

<https://arxiv.org/abs/2507.11773>

51 **CONF****Commonsense on Demand: Generating and Selectively Integrating Commonsense Knowledge for Natural Language Inference**

ArXiv CS.AI 50%

arXiv:2507.15100v3 Announce Type: replace-cross Abstract: Natural Language Inference (NLI) determines whether a premise entails, contradicts, or is neutral with respect to a hypothesis. The task is often framed as emulating human inference, in which commonsense knowledge plays a...

<https://arxiv.org/abs/2507.15100>

52 **CONF****Empowering Children to Create AI-Enabled Augmented Reality Experiences**

ArXiv CS.AI 50%

arXiv:2508.08467v2 Announce Type: replace-cross Abstract: Despite their potential to enhance children's learning experiences, AI-enabled AR technologies are predominantly used in ways that position children as consumers rather than creators. We introduce Capybara, an AR-based...

<https://arxiv.org/abs/2508.08467>

53 **CONF****CORE-3D: Context-aware Open-vocabulary Retrieval by Embeddings in 3D**

ArXiv CS.AI 50%

arXiv:2509.24528v4 Announce Type: replace-cross Abstract: Object retrieval from a scene has become a new trend of research due to its numerous applications. Recent approaches achieve zero-shot, open-vocabulary 3D semantic mapping by assigning embedding vectors to 2D...

<https://arxiv.org/abs/2509.24528>

54 **CONF****A-3PO: Accelerating Asynchronous LLM Training with Staleness-aware Proximal Policy Approximation**

ArXiv CS.AI 50%

arXiv:2512.06547v4 Announce Type: replace-cross Abstract: Decoupled PPO has been a successful reinforcement learning (RL) algorithm to deal with the high data staleness under the asynchronous RL setting. Decoupled loss used in decoupled PPO improves coupled-loss style of...

<https://arxiv.org/abs/2512.06547>

55 **CONF****Probably Approximately Correct Maximum A Posteriori Inference**

ArXiv CS.AI 50%

arXiv:2601.16083v2 Announce Type: replace-cross Abstract: Computing the conditional mode of a distribution, better known as the maximum a posteriori (MAP) assignment, is a fundamental task in probabilistic inference. However, MAP is generally intractable, and remains hard even...

<https://arxiv.org/abs/2601.16083>

56 CONF

Superposition Without Interference? Towards Isolated Interventions via Almost Orthogonal Features in Language Models

ArXiv CS.AI 50%

arXiv:2602.04718v5 Announce Type: replace-cross Abstract: A central premise in mechanistic interpretability is that meaningful concepts in language models are represented by linear features in activation space. For such features to support reliable interventions, manipulating...

<https://arxiv.org/abs/2602.04718>

57 CONF

LLM-Powered Automatic Translation and Urgency in Crisis Scenarios

ArXiv CS.AI 50%

arXiv:2602.13452v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly proposed for crisis preparedness and response, particularly for multilingual communication. However, their suitability for high-stakes crisis contexts remains insufficiently...

<https://arxiv.org/abs/2602.13452>

58 CONF

How effective are VLMs in assisting humans in inferring the quality of mental models from Multimodal short answers?

ArXiv CS.AI 50%

arXiv:2603.00056v2 Announce Type: replace-cross Abstract: STEM Mental models can play a critical role in assessing students' conceptual understanding of a topic. They not only offer insights into what students know but also into how effectively they can apply, relate to, and...

<https://arxiv.org/abs/2603.00056>

59 CONF

VLM2Rec: Resolving Modality Collapse in Vision-Language Model Embedders for Multimodal Sequential Recommendation

ArXiv CS.AI 50%

arXiv:2603.17450v2 Announce Type: replace-cross Abstract: Sequential Recommendation (SR) in multimodal settings typically relies on small frozen pretrained encoders, which limits semantic capacity and prevents Collaborative Filtering (CF) signals from being fully integrated into...

<https://arxiv.org/abs/2603.17450>

60 CONF

TMRL: Diffusion Timestep-Modulated Pretraining Enables Exploration for Efficient Policy Finetuning

ArXiv CS.AI 50%

arXiv:2605.12236v2 Announce Type: replace-cross Abstract: Fine-tuning pre-trained robot policies with reinforcement learning (RL) often inherits the bottlenecks introduced by pre-training with behavioral cloning (BC), which produces narrow action distributions that lack the...

<https://arxiv.org/abs/2605.12236>

61 CONF

Cryptographically verifiable authorization for autonomous AI agents: A falsifiable hypothesis and proof-of-concept

ArXiv CS.AI 50%

arXiv:2607.21325v2 Announce Type: replace-cross Abstract: Autonomous AI agents increasingly execute actions, invoke tools, and operate on protected resources with limited human oversight. Existing authentication and authorization mechanisms establish identity and delegate...

<https://arxiv.org/abs/2607.21325>

62 CONF

Epistemic Transfer in AI-Assisted Verification: A Framework and Evaluation Protocol

ArXiv CS.AI 50%

arXiv:2608.08882v2 Announce Type: replace-cross Abstract: AI tools that help people judge online claims are usually evaluated while the tool is present. This paper asks a different question: after using such a tool, what can the user still do on their own? I call this epistemic...

<https://arxiv.org/abs/2608.08882>

63 CONF

ENTLORE: A Graph-Grounded Benchmark for Latent Organizational Reasoning in Enterprise Question Answering

ArXiv CS.AI 50%

arXiv:2608.10679v2 Announce Type: replace-cross Abstract: Enterprise question answering is framed as retrieving internal documents and generating grounded answers. Routine enterprise records, however, are work by-products in which required organizational relations remain...

<https://arxiv.org/abs/2608.10679>