



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	113	Media
HackerNews	5	Community
TechCrunch AI	3	Media
The Verge AI	3	Media

CONFIRMED SIGNALS — 124 stories

CONF

Show HN: A public AI whose memory is shared across all users

HackerNews 50% HN 69

No summary — see link for full article.

<https://wildstatic.com/>

CONF

Show HN: Deltix - AI Driven Testing

HackerNews 50% HN 53

Hey HN, I'm Charles, the solo founder of Continuum (<https://continuum-app.xyz/>). Continuum is a financial planning web app for people with 1-15M. It helps you create simulations of job offers, house purchases, and life events. Its goal is to replace human advisors that cost...

<https://app.deltix.ai>

CONF

Show HN: Widen, a native Postgres GUI using Apple's on-device LLM

HackerNews 50% HN 9

No summary — see link for full article.

<https://github.com/betocmn/widen>

CONF

Show HN: Sib - Unixy LLM Client using Git to store conversations, instead SQLite

HackerNews 50% HN 4

Hi all :) The reason I like git is that the source tree snapshot and the commit structure itself are always immutable, and destructive operations are just renaming a ref file. Once you understand that, no matter how hard the CLI is to make sense of, you never hesitate to run a...

<https://github.com/sib-project/sib>

CONF

Show HN: Who is using AI? A daily tracker, 1,600+ city heatmap, DOI dataset

HackerNews 50% HN 2

No summary — see link for full article.

<https://whoisusingai.com/tracker>

CONF

Stripe will reportedly acquire AI gateway startup OpenRouter for \$7B+

TechCrunch AI 50%

OpenRouter's CEO recently described the startup as Stripe for AI.

<https://techcrunch.com/2026/08/16/stripe-will-reportedly-acquire-ai-gateway-startup-ope...>

CONF

Anthropic CEO says AI backlash is 'fundamentally a crisis of trust'

TechCrunch AI 50%

Dario Amodei is pushing back against the idea that he's been painting an overly pessimistic picture of AI.

<https://techcrunch.com/2026/08/16/anthropic-ceo-says-ai-backlash-is-fundamentally-a-crisi...>

CONF

Microsoft kills off unsuccessful AI features while merging its separate Copilot apps

TechCrunch AI 50%

Microsoft is simplifying Copilot by combining its consumer and business apps, and dropping AI-generated podcasts, Group Chats, Deep Research, and its Mico character.

<https://techcrunch.com/2026/08/13/microsoft-kills-off-unsuccessful-ai-features-while-me...>

CONF

OpenAI reportedly disbanded its preparedness team

The Verge AI 50%

According to the Financial Times, OpenAI disbanded its preparedness team at the end of last month. The job of the preparedness team was to assess if models posed serious risks and develop ways to mitigate those risks. (You know, like the possibility that it could go rogue and...

<https://www.theverge.com/ai-artificial-intelligence/980817/openai-disbands-preparedness...>

0 CONF

ChatGPT's Computer History tracks your clicks and keystrokes

The Verge AI 50%

ChatGPT's desktop app on macOS has a new feature called Computer History that turns your actions into training data, learning how you work, suggesting automations, and even picking up tasks you left half done. It uses your activity to build a timeline that ChatGPT and Codex can...

<https://www.theverge.com/ai-artificial-intelligence/980742/chatgpts-computer-history-tr...>

1 CONF

Rogue AI aren't science fiction anymore

The Verge AI 50%

This is The Stepback, a weekly newsletter breaking down one essential story from the tech world. For more on AI safety, follow Robert Hart. The Stepback arrives in our subscribers' inboxes at 8AM ET. Opt in for The Stepback here. How it started It all started in July, when one...

<https://www.theverge.com/column/980337/rogue-ai-science-fiction-openai>

2 CONF

A Year in LLM Serving: Workload Evolution, Caching and Load-Balancing

ArXiv CS.AI 50%

arXiv:2608.13573v1 Announce Type: new Abstract: Large Language Model (LLM) serving has become a critical cloud workload, and realistic traces are essential for motivating and benchmarking serving systems. However, existing LLM serving workload studies remain limited in scale and...

<https://arxiv.org/abs/2608.13573>

3 CONF

ArGENt: Arbitrary Geometry-encoded Transformer for Operator Learning

ArXiv CS.AI 50%

arXiv:2602.11626v3 Announce Type: replace-cross Abstract: Learning solution operators on arbitrary geometries remains a central challenge in scientific machine learning, especially for many-query simulation, physics-informed learning, and evolving geometries requiring accurate,...

<https://arxiv.org/abs/2602.11626>

4 CONF

AI Evaluation Should Work With Humans

ArXiv CS.AI 50%

arXiv:2608.13577v1 Announce Type: new Abstract: This position paper argues that the dominant paradigm of AI evaluation (which focuses on superhuman autonomous performance and so implicitly targets the goal of replacing humans) is guiding AI development in the wrong direction....

<https://arxiv.org/abs/2608.13577>

5 CONF

Stable Miscalibration in Large Language Models: A Practical View of High-Confidence Errors

ArXiv CS.AI 50%

arXiv:2608.13591v1 Announce Type: new Abstract: High-confidence errors in large language models are often treated as evidence of fragile internal inference. We study a different possibility: stable miscalibration, where a confident wrong answer remains locally stable under small...

<https://arxiv.org/abs/2608.13591>

6 CONF

BAT: Learning to Reason about Spatial Sounds with Large Language Models

ArXiv CS.AI 50%

arXiv:2402.01591v4 Announce Type: replace-cross Abstract: Spatial sound reasoning is a fundamental human skill, enabling us to navigate and interpret our surroundings based on sound. In this paper we present BAT, which combines the spatial sound perception ability of a binaural...

<https://arxiv.org/abs/2402.01591>

7 CONF

Mind the Long Tail: Understanding the Difficulty of Delay Detection in Business Processes

ArXiv CS.AI 50%

arXiv:2608.14367v1 Announce Type: cross Abstract: The early detection of delayed cases in business processes is a critical capability for organizations. Predictive process monitoring (PPM) supports this task by using historical event logs to predict the remaining time of ongoing...

<https://arxiv.org/abs/2608.14367>

8 CONF

No Universal Signal Predicts Sample-Level LLM Regression under Version Updates

ArXiv CS.AI 50%

arXiv:2608.13607v1 Announce Type: new Abstract: Frontier LLMs are updated frequently and typically outperform their predecessors in aggregate. But aggregate gains say little about individual samples: an update can still cause sample-level regression, where a response correct...

<https://arxiv.org/abs/2608.13607>

9 CONF

SemPlan: Benchmarking Structured Semantic Planning for LLM-Based Queries over Enterprise Data

ArXiv CS.AI 50%

arXiv:2608.13612v1 Announce Type: new Abstract: Natural-language interfaces to enterprise data must translate underspecified requests into governed, executable behavior while controlling invalid queries, policy failures, cost, and nondeterminism. SemPlan Benchmark evaluates this...

<https://arxiv.org/abs/2608.13612>

10 CONF

Your Probabilistic JEPa Is Secretly a Hidden Markov Model: A State-Space Interpretation of Joint-Embedding Predictive Learning

ArXiv CS.AI 50%

arXiv:2608.13621v1 Announce Type: new Abstract: A hidden Markov model (HMM) combines three roles: inference of a hidden-state belief from observations, propagation through a Markov transition, and emission back to observation space. We show that full, time-indexed Predictive...

<https://arxiv.org/abs/2608.13621>

1 CONF

ARC: Fair Relative Advantage Comparison in Open-Ended Real-World Interaction

ArXiv CS.AI 50%

arXiv:2608.13622v1 Announce Type: new Abstract: Open-ended real-world interaction admits multiple valid behaviors: an agent may answer directly, ask for clarification, provide progress updates, or confirm before acting. This flexibility breaks a core assumption behind...

<https://arxiv.org/abs/2608.13622>

2 CONF

Second Thought: Reasoning in Parallel as LLM Agents Act and Observe

ArXiv CS.AI 50%

arXiv:2608.13667v1 Announce Type: new Abstract: LLM agents in the ReAct paradigm alternate between reasoning, acting, and observing, but deliberate reasoning is confined to the Thought phase: while the agent serializes an action and waits for the environment, its reasoning is...

<https://arxiv.org/abs/2608.13667>

3 CONF

Learning to Assemble Novel Structures with Unfamiliar Parts under Semantic Constraints

ArXiv CS.AI 50%

arXiv:2608.13684v1 Announce Type: new Abstract: This paper describes a neurosymbolic architecture for learning to assemble novel structures using evidence from embodied conversations and task demonstrations. We focus on scenarios where an agent encounters, after deployment,...

<https://arxiv.org/abs/2608.13684>

4 CONF

Coverage Aware Active Evaluation for Failure Discovery with Paired Systems

ArXiv CS.AI 50%

arXiv:2608.13719v1 Announce Type: new Abstract: Autonomous systems can fail in rare and heterogeneous ways, making real-world failure discovery difficult under limited testing budgets. Although cheaper proxies such as simulators, lower-fidelity systems, or related policies can...

<https://arxiv.org/abs/2608.13719>

5 CONF

Simulation-Aware In-Context Policy Improvement for LLM-Aided Analog Layout Refinement

ArXiv CS.AI 50%

arXiv:2608.13767v1 Announce Type: new Abstract: Analog IC layout design remains a labor-intensive iterative process dominated by simulation-driven refinement. Although end-to-end layout generators accelerate initial placement and routing, they still require experts to manually...

<https://arxiv.org/abs/2608.13767>

6 CONF

SKILLER: Language-Level Reinforcement Learning for Reusable Skill Extraction in Small Language Models

ArXiv CS.AI 50%

arXiv:2608.10538v2 Announce Type: replace Abstract: Agent skills represent a standardized format for packaging procedural knowledge and domain expertise, serving within agent harness systems as an essential mechanism to continually constrain a language model's behavior space for...

<https://arxiv.org/abs/2608.10538>

7 CONF

Joint Optimization of Memory and Computing Frequency for Energy-Efficient DNN Inference

ArXiv CS.AI 50%

arXiv:2608.13863v1 Announce Type: new Abstract: Deep neural network (DNN) inference on mobile devices often incurs high latency and energy consumption due to limited computing and memory resources. To enable energy-efficient DNN inference, most existing studies focus on dynamic...

<https://arxiv.org/abs/2608.13863>

8 CONF

When Personal Memory Has No Single Answer: Evaluating LLM Agents under Irreducible Conflict

ArXiv CS.AI 50%

arXiv:2608.13921v1 Announce Type: new Abstract: LLM agents increasingly maintain personal memory across sessions, but it can conflict. Preferences depend on context, behavior evolves, and sources can conflict. When a query lacks context, time, or source authority to interpret...

<https://arxiv.org/abs/2608.13921>

9 CONF

Never the Number: Structural Abstention for AI Systems Whose Answers Are Consumed as Fact

ArXiv CS.AI 50%

arXiv:2608.13926v1 Announce Type: new Abstract: Large language models have made natural language interfaces to databases (NLIDB) newly credible, but LLM text-to-SQL systems fail in a way that matters for deployment: a hallucinated column or a mis-aggregated total yields a fluent...

<https://arxiv.org/abs/2608.13926>

0 CONF

Training Fair Tabular Foundation Models

ArXiv CS.AI 50%

arXiv:2608.14211v1 Announce Type: cross Abstract: Tabular Foundation Models (TFMs) have emerged as leading methods for tabular predictive tasks, leveraging in-context learning to predict on new data without task-specific training. Despite the increased use of TFMs in high-stakes...

<https://arxiv.org/abs/2608.14211>

1 CONF

Implementing Computational Law in Wolfram Language for the Governance of Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2608.13958v1 Announce Type: new Abstract: How do we govern AI systems whose reasoning we cannot fully inspect? Governance does not require understanding a system's reasoning. It requires stating what the system is obliged, permitted, and forbidden to do, and checking...

<https://arxiv.org/abs/2608.13958>

2 CONF

Nanbeige4.2-3B on Apple Silicon: Fixing Deployment Bugs and Decreasing Looped Transformer Memory Overhead

ArXiv CS.AI 50%

arXiv:2608.13987v1 Announce Type: new Abstract: Nanbeige4.2-3B is a 3B-parameter agentic model built around a Looped Transformer (LT) that reuses one stack of layers for a second forward pass, adding effective depth without additional parameters. Evaluated on Apple Silicon...

<https://arxiv.org/abs/2608.13987>

3 CONF

Simulation-Driven Vehicular Traffic Data Augmentation: Extending Sensor Coverage Through Virtual Sensing

ArXiv CS.AI 50%

arXiv:2608.13993v1 Announce Type: new Abstract: Urban traffic management relies on sensor networks whose spatial coverage is limited by deployment costs and privacy regulations. Machine learning models trained on such sparse data cannot generalize to unmonitored locations and...

<https://arxiv.org/abs/2608.13993>

4 CONF

Benchmarking data-driven material models on the classic Treloar dataset

ArXiv CS.AI 50%

arXiv:2608.14063v1 Announce Type: new Abstract: Machine learning is rapidly reshaping constitutive modeling, offers new ways to learn material behavior directly from experimental data, and challenges long-established modeling paradigms. But with a growing number of...

<https://arxiv.org/abs/2608.14063>

5 CONF

Scaling Domain Data Repetition in LLM Pretraining

ArXiv CS.AI 50%

arXiv:2608.14071v1 Announce Type: new Abstract: As large language models scale, their training-token budgets must also increase to maintain an appropriate tokens-per-parameter ratio (TPP). However, high-quality domain data is much harder to scale than general web...

<https://arxiv.org/abs/2608.14071>

6 CONF

Mandato: Protocol-Level Enforcement of Digitally Signed Mandates on AI Agent Actions with Cryptographically Chained Audit Trails

ArXiv CS.AI 50%

arXiv:2608.14074v1 Announce Type: new Abstract: AI agents increasingly act on external systems through standardized tool-calling protocols such as the Model Context Protocol (MCP), yet no infrastructure layer constrains their actions to what a principal has verifiably...

<https://arxiv.org/abs/2608.14074>

7 CONF

A Pathway to General-Purpose Scientific AI: Multimodal Comprehension of Scientific Images

ArXiv CS.AI 50%

arXiv:2608.14075v1 Announce Type: new Abstract: Scientific figures and tables encode essential experimental evidence, yet remain difficult for digital libraries and multimodal AI systems to retrieve and interpret. The ALD/E-ImageMiner benchmark and ICDAR 2026 Competition on...

<https://arxiv.org/abs/2608.14075>

8 CONF

RL-Index: Reinforcement Learning for Retrieval Index Reasoning

ArXiv CS.AI 50%

arXiv:2606.16316v2 Announce Type: replace-cross Abstract: Retrieving external knowledge is crucial for real-world tasks but remains difficult when queries and relevant knowledge are linked by implicit reasoning (e.g., shared theorems or coding logic). Existing methods rely...

<https://arxiv.org/abs/2606.16316>

- 9
CONF
Reinforcement Learning-Based Production Scheduling in an Industry-Based Coating Scenario Using the Digital Model Playground
ArXiv CS.AI 50%

arXiv:2608.14122v1 Announce Type: new Abstract: Production scheduling in complex manufacturing environments is challenging when sequence-dependent setup times, stochastic disturbances, and due-date constraints must be addressed simultaneously. While reinforcement learning (RL)...

<https://arxiv.org/abs/2608.14122>
- 0
CONF
Towards Efficient Multimodal and Multilingual Opinion Extraction for STI: A QLoRA-Based Fine-Tuning Approach
ArXiv CS.AI 50%

arXiv:2608.14152v1 Announce Type: new Abstract: Recent advances in large language models (LLMs) have reshaped semantic analysis. Opinion Extraction (OE) for Science and Technology Intelligence (STI) requires concise core opinions from large information streams. Off-the-shelf...

<https://arxiv.org/abs/2608.14152>
- 1
CONF
Removing Temporal Note Redundancy Improves Multimodal Reinforcement Learning for Medicine
ArXiv CS.AI 50%

arXiv:2608.14157v1 Announce Type: new Abstract: Mechanical ventilation is a critical life-support intervention, requiring dynamic adjustments to ventilator settings as a patient's condition evolves. While reinforcement learning (RL) offers a promising framework for optimizing...

<https://arxiv.org/abs/2608.14157>
- 2
CONF
BiasTrace: Linking Reasoning Behaviours to Biased Outputs in LLMs
ArXiv CS.AI 50%

arXiv:2608.14161v1 Announce Type: new Abstract: LLMs exhibit social biases that can produce inaccurate and discriminatory inferences, posing risks in high-stakes applications. While prior work has made progress in measuring and mitigating bias, it largely focuses on final...

<https://arxiv.org/abs/2608.14161>
- 3
CONF
LUNAR: Benchmarking Personalized Large Language Models on UNiversal User BehAvioR Logs
ArXiv CS.AI 50%

arXiv:2608.05246v2 Announce Type: replace Abstract: Existing personalized LLM benchmarks primarily rely on textual personas or isolated behavioral signals, providing limited evaluation of cross-domain behavioral personalization, where responses must be grounded in heterogeneous...

<https://arxiv.org/abs/2608.05246>
- 4
CONF
APTER: Adaptive Post-Training with Expert-Grounded Rubrics
ArXiv CS.AI 50%

arXiv:2608.14212v1 Announce Type: new Abstract: As large language models enter professional domains, they must satisfy domain constraints, include critical evidence, and provide complete reasoning rather than merely produce fluent responses. Existing post-training methods often...

<https://arxiv.org/abs/2608.14212>

5 CONF

Forecast Collapse in Time-Series Foundation Models

ArXiv CS.AI 50%

arXiv:2608.14106v1 Announce Type: cross Abstract: When forecasting hourly returns for 1,000 US equities, we observe an unexpected phenomenon: predictions become nearly flat and show poor stock ranking, as measured by cross-sectional correlation. We call this forecast collapse....

<https://arxiv.org/abs/2608.14106>

6 CONF

Grounding Without Corrective Control: Truth-Tracking Profiles for Large Language Models

ArXiv CS.AI 50%

arXiv:2608.14252v1 Announce Type: new Abstract: Recent work suggests that some large language model representations have content or reference. Grounding can secure either without supplying live routes for correction. This paper asks what follows from that gap. An output is...

<https://arxiv.org/abs/2608.14252>

7 CONF

AnchorBench: A Multi-Pathway Benchmark for the Anchoring Effect in LLMs

ArXiv CS.AI 50%

arXiv:2608.14320v1 Announce Type: new Abstract: The anchoring effect is a cognitive bias in which an initial reference value shifts a later judgment toward itself. This effect is well established in human judgment and decision-making, and recent work suggests that large language...

<https://arxiv.org/abs/2608.14320>

8 CONF

Intern-S2-Mobius: Foundation Model with Decoupled Knowledge and Reasoning

ArXiv CS.AI 50%

arXiv:2608.14290v1 Announce Type: new Abstract: We introduce Mobius-v0, an architecture that comprises a globally shared Memory (FFN) that stores knowledge vectors and multiple Reasoners (Self-Attn) that iteratively achieve compositional reasoning. Using hidden states as cache...

<https://arxiv.org/abs/2608.14290>

9 CONF

Program-space Diffusion for Morphology-to-Transcriptomics Prediction

ArXiv CS.AI 50%

arXiv:2608.14330v1 Announce Type: new Abstract: Spatial transcriptomics (ST) enables genome-wide gene expression profiling while preserving tissue architecture, but its cost and limited scalability remain major bottlenecks. This has motivated models that predict spatial...

<https://arxiv.org/abs/2608.14330>

0 CONF

Clearing the Fog: Towards Installing and Refining Proactive Exploration Capabilities in LLM Agents

ArXiv CS.AI 50%

arXiv:2608.14339v1 Announce Type: new Abstract: We study proactive exploration in LLM agents, i.e., the ability to explore an environment to acquire information that improves future decision-making. In this regard, we first identify two fundamental bottlenecks that hinder this...

<https://arxiv.org/abs/2608.14339>

1 CONF

Designing Sustainable Federated Learning as a Service using Neural Architecture Search

ArXiv CS.AI 50%

arXiv:2608.14359v1 Announce Type: new Abstract: The sustainability constraints of FLaaS consumers pose significant challenges to maintaining carbon-feasible federated training in FLaaS environments. These constraints often lead to infeasible consumer participation and unstable...

<https://arxiv.org/abs/2608.14359>

2 CONF

LLMs Don't Pay for the Jump

ArXiv CS.AI 50%

arXiv:2608.14397v1 Announce Type: new Abstract: Zahavy [2026] argues that Large Language Models, despite their capabilities in induction and deduction, cannot perform the abductive "jump" that produced Einstein's equivalence principle, and attributes this limitation to the...

<https://arxiv.org/abs/2608.14397>

3 CONF

Why we need an AI-resilient society

ArXiv CS.AI 50%

arXiv:1912.08786v3 Announce Type: replace-cross Abstract: Three generations of software have transformed the role of artificial intelligence in society. In the first, programmers wrote explicit logic. In the second, neural networks learned programs from data. In the third, large...

<https://arxiv.org/abs/1912.08786>

4 CONF

Knowing When to Stop: Bayesian Optimal Stopping for LLM Evaluations

ArXiv CS.AI 50%

arXiv:2608.14425v1 Announce Type: new Abstract: LLM evaluations often use fixed sampling budgets, testing every item the same number of times even after estimates are precise. We introduce optstop, a precision-based adaptive stopping framework that treats evaluation as a...

<https://arxiv.org/abs/2608.14425>

5 CONF

Participatory Moral AI Is Not Neutral: The Invisible Hand of Developers

ArXiv CS.AI 50%

arXiv:2608.14522v1 Announce Type: new Abstract: As AI systems make more morally loaded decisions across society, one response has been moral preference elicitation. In this approach, researchers poll participants on hypothetical dilemmas and use the aggregated votes to train a...

<https://arxiv.org/abs/2608.14522>

6 CONF

Proxy-Validated LLM UX Micro-Simulations: An Artifact-First Protocol for Early-Stage Decision Support

ArXiv CS.AI 50%

arXiv:2608.13563v1 Announce Type: cross Abstract: Early-stage teams often lack users, time, and budget to run repeated UX studies, yet still need decision-oriented signals to iterate safely. We study an LLM-driven UX micro-simulation pipeline that generates structured...

<https://arxiv.org/abs/2608.13563>

7 CONF

Don't Claim Benchmark-Oriented Optimization Improves General Coding Capability -- Diverse Evaluation Is Required

ArXiv CS.AI 50%

arXiv:2608.13566v1 Announce Type: cross Abstract: Post-training papers, model cards, and blog posts often treat scores on a small set of coding benchmarks (e.g., SWE-bench and LiveCodeBench) as evidence of broad coding capability, both for research artifacts and user-facing...

<https://arxiv.org/abs/2608.13566>

8 CONF

Think in Latent, Explain in Language: Self-Explainable Latent Reasoning

ArXiv CS.AI 50%

arXiv:2608.13570v1 Announce Type: cross Abstract: Latent reasoning has emerged as a powerful alternative to text-based Chain-of-Thought (CoT), offering significant gains in computational efficiency by compressing verbose reasoning into compact embeddings. However, compressing...

<https://arxiv.org/abs/2608.13570>

9 CONF

Not All Tokens Are Equal: Inflation-Aware Routing for Agentic LLM Systems

ArXiv CS.AI 50%

arXiv:2608.13571v1 Announce Type: cross Abstract: When a language model fails to answer a query on the first attempt, an agentic system retries, consuming additional tokens each time. This retry overhead creates a gap between what a model's per-token price implies and what a...

<https://arxiv.org/abs/2608.13571>

0 CONF

The Architect: Interactive Visualization of Deep Learning Mathematics Directly in Microsoft Excel

ArXiv CS.AI 50%

arXiv:2608.13572v1 Announce Type: cross Abstract: We present The Architect, a system that turns Microsoft Excel into an interactive view of deep learning mathematics. A user describes a neural network in a compact table. The system then generates a workbook that shows the full...

<https://arxiv.org/abs/2608.13572>

1 CONF

BCMT: Blockwise Causal Memory Transformer

ArXiv CS.AI 50%

arXiv:2608.13578v1 Announce Type: cross Abstract: Transformer architectures rely on dense self-attention to model long-range dependencies, but this mechanism exhibits quadratic complexity with respect to sequence length. We introduce BCMT (Blockwise Causal Memory Transformer),...

<https://arxiv.org/abs/2608.13578>

2 CONF

From Prediction to Intervention: Personalized Meal-Level Glucose Regulation via an LLM Agent

ArXiv CS.AI 50%

arXiv:2608.13581v1 Announce Type: cross Abstract: Personalized glucose regulation remains a central yet unresolved challenge in precision nutrition, as postprandial glucose response varies substantially across individuals. Existing approaches based on glycemic indices fail to...

<https://arxiv.org/abs/2608.13581>

- 3
CONF
Context Aware AI Assistant and AR Interface for Lunar Extravehicular Activity (EVA) Procedural Guidance
ArXiv CS.AI 50%
arXiv:2608.13589v1 Announce Type: cross Abstract: As human space exploration returns to the Moon, astronauts need rapid access to procedural information during extravehicular activities (EVAs), where attention is divided across navigation, repair tasks, tool handling, and...
<https://arxiv.org/abs/2608.13589>
- 4
CONF
Overcoming Shortcut Learning in Graph Neural Networks through Active Explanation Guidance
ArXiv CS.AI 50%
arXiv:2608.14121v1 Announce Type: cross Abstract: Graph Neural Networks (GNNs) can solve prediction tasks by unintentionally exploiting shortcuts---that is, edges, nodes, and features that correlate with but are not causal for the prediction---which compromise their reliability...
<https://arxiv.org/abs/2608.14121>
- 5
CONF
Secret-Stego Dissimilarity as a Design Axis: Invertible Coverless Image Steganography with Diffusion Models
ArXiv CS.AI 50%
arXiv:2608.13597v1 Announce Type: cross Abstract: Coverless image steganography (CIS) synthesizes a stego image rather than modifying an existing cover image, enabling authorized recipients to reconstruct the original secret image from the stego. Existing diffusion-based CIS...
<https://arxiv.org/abs/2608.13597>
- 6
CONF
Fine-Tuning Qwen3-27B for C-to-Rust Code Translation: A Three-Stage Curriculum of Pretraining, Debugging-Aware SFT, and Task-Specific SFT
ArXiv CS.AI 50%
arXiv:2608.13681v1 Announce Type: cross Abstract: Translating C code into safe, idiomatic Rust is a longstanding software-engineering goal because it can eliminate entire classes of memory-safety vulnerabilities while preserving the functional behavior of legacy systems. Large...
<https://arxiv.org/abs/2608.13681>
- 7
CONF
SAGE: Surrogate-gradient Adaptation via Attention-Guided Entropy for Spiking Transformers
ArXiv CS.AI 50%
arXiv:2608.13702v1 Announce Type: cross Abstract: Spiking neural networks (SNNs) offer an energy-efficient alternative to conventional deep neural networks by exploiting sparse event-driven computation, but their training remains challenging because the non-differentiable spike...
<https://arxiv.org/abs/2608.13702>
- 8
CONF
TeachMateGPT: A Multi-Agent Knowledge-Grounded Framework for Pedagogical Assessment Generation from Science Curriculum Materials
ArXiv CS.AI 50%
arXiv:2608.13708v1 Announce Type: cross Abstract: Automatically generating textbook-grounded assessment items can reduce science teachers' workload, but existing retrieval-augmented generation (RAG) systems rely on flat retrieval, support only single-question generation, lack...
<https://arxiv.org/abs/2608.13708>

- 9 **CONF**
- Building AI-Intensive Software with AI: Early Results and a Cautionary Tale on Measuring Development Cost**
- ArXiv CS.AI 50%
- arXiv:2608.13730v1 Announce Type: cross Abstract: Empirical reports on the true cost of AI-intensive software development remain scarce, and the few that exist are easy to get wrong in ways that never surface in the final number. We report early results from an ongoing case...
- <https://arxiv.org/abs/2608.13730>
-
- 0 **CONF**
- Does ISO-Grounded NFR Specification Improve LLM Code Generation? A Comparison of Rich and Structured Interventions against a Natural-Language Baseline**
- ArXiv CS.AI 50%
- arXiv:2608.13742v1 Announce Type: cross Abstract: In LLM-based code generation, Non-Functional Requirements (NFRs) are often specified as terse one-line phrases. We ask whether grounding those specifications in ISO/IEC 25010 Quality Model, either as rich natural-language prose...
- <https://arxiv.org/abs/2608.13742>
-
- 1 **CONF**
- CutClean: Neural Network Pruning for Privacy-Preserving Inference**
- ArXiv CS.AI 50%
- arXiv:2608.13773v1 Announce Type: cross Abstract: Neural networks are increasingly deployed in high-stakes applications with growing privacy leakage concerns. We show that this privacy leakage can occur even in the absence of representation imbalances that lead to traditional...
- <https://arxiv.org/abs/2608.13773>
-
- 2 **CONF**
- Do AI chatbots find what experts would? Effects of model, user role, and sample size on study retrieval for medical questions**
- ArXiv CS.AI 50%
- arXiv:2608.13786v1 Announce Type: cross Abstract: Large language model (LLM) chatbots are increasingly used to answer clinical questions with citations to relevant clinical studies. Prior research has largely focused on citation fabrication, leaving a gap in evaluating the...
- <https://arxiv.org/abs/2608.13786>
-
- 3 **CONF**
- Towards Practical Algorithm Selection for Unsupervised Domain Adaptation in Medical Imaging**
- ArXiv CS.AI 50%
- arXiv:2607.28125v2 Announce Type: replace-cross Abstract: Numerous unsupervised domain adaptation (UDA) algorithms exist, but for clinical practice, selecting the best-suited one along with proper hyperparameters often remains unclear, as the unlabeled deployment (target) domain...
- <https://arxiv.org/abs/2607.28125>
-
- 4 **CONF**
- ASSERT: A Measurement Pipeline for GenAI Audits**
- ArXiv CS.AI 50%
- arXiv:2608.13840v1 Announce Type: cross Abstract: Audits of generative AI (GenAI) systems often summarize behavior as a reported rate: how often the audited system complies with policy. Researchers and stakeholders use that rate to compare systems, track regressions, and gate...
- <https://arxiv.org/abs/2608.13840>
-

- 5 **CONF**
- Engineering Signals of Human-AI Collaboration in the Agentic Coding Era: A Longitudinal Analysis of 33,228 Pull Requests from vLLM and SGLang with Implications for Biomedical AI Agents and Bioinformatics Pipeline Developmen**
- ArXiv CS.AI 50%
- arXiv:2608.13884v1 Announce Type: cross Abstract: The rapid adoption of AI coding assistants and autonomous agentic development systems has coincided with major changes in the pace and structure of open-source software engineering. Yet empirical longitudinal evidence of these...
- <https://arxiv.org/abs/2608.13884>
-
- 6 **CONF**
- CForce: Boosting Parallel Decoding for dLLMs via Consistency Forcing**
- ArXiv CS.AI 50%
- arXiv:2608.13925v1 Announce Type: cross Abstract: Diffusion large language models (dLLMs) accelerate language generation by predicting multiple masks in a single forward pass. However, existing dLLMs can suffer from unreliable predictions in early denoising stages under...
- <https://arxiv.org/abs/2608.13925>
-
- 7 **CONF**
- CMCNet: Aligning Ultrasound Image Embeddings with Textual TI-RADS Representations for Fine-Grained Thyroid Classification**
- ArXiv CS.AI 50%
- arXiv:2608.13939v1 Announce Type: cross Abstract: Ultrasound is the primary imaging modality for assessing thyroid nodules, and the ACR TI-RADS framework standardizes diagnosis through five ultrasound feature categories that are aggregated into five risk levels (TR1-TR5)....
- <https://arxiv.org/abs/2608.13939>
-
- 8 **CONF**
- Musical Mirrors: The LLM as Sounding Board in Songwriting**
- ArXiv CS.AI 50%
- arXiv:2608.13944v1 Announce Type: cross Abstract: This paper examines a use of AI in creative practice as an interpretive sounding board for human-generated material, rather than the more familiar pattern of AI generation followed by human curation. Through the lens of resonance...
- <https://arxiv.org/abs/2608.13944>
-
- 9 **CONF**
- EchoRec: Multi-Item Prediction-Empowered Generative Recommendation via Cycle-Consistent Preference Alignment**
- ArXiv CS.AI 50%
- arXiv:2608.14011v1 Announce Type: cross Abstract: Generative recommendation autoregressively generates the semantic IDs of the target item, unifying preference modeling and index retrieval within the shared token space. Recent attempts have introduced Multi-Token Prediction...
- <https://arxiv.org/abs/2608.14011>
-
- 0 **CONF**
- Content Based Video Narration of Gameplay with Vision Language Models**
- ArXiv CS.AI 50%
- arXiv:2608.14016v1 Announce Type: cross Abstract: Live game commentary is scarce: it exists for professional esports broadcasts and almost nowhere else. We present a content-based video narration system that produces spoken, esports-style commentary for arbitrary gameplay...
- <https://arxiv.org/abs/2608.14016>
-

1 CONF

ForgeWM: Progressive Causal Training for Few-Step Action-Conditioned Video World Models

ArXiv CS.AI 50%

arXiv:2608.14022v1 Announce Type: cross Abstract: Action-conditioned video world models require low-latency causal generation and reliable responses to game-native controls. Although causal distillation enables one- or few-step video synthesis, extending it to interactive world...

<https://arxiv.org/abs/2608.14022>

2 CONF

HAM-RAG: Hierarchy-Aware Multimodal RAG for Structure-Faithful Interleaved Generation

ArXiv CS.AI 50%

arXiv:2608.14032v1 Announce Type: cross Abstract: Existing multimodal RAG methods often flatten structured documents into isolated text and image units, weakening the source organization and local text-image logic needed for faithful evidence selection and placement. We propose...

<https://arxiv.org/abs/2608.14032>

3 CONF

Rethinking Automated Program Repair: The Impact of Bug Complexity, Fault Localization, and LLM Cost-efficiency

ArXiv CS.AI 50%

arXiv:2608.14065v1 Announce Type: cross Abstract: Background: Software bugs remain a critical challenge in development, necessitating effective Automated Program Repair (APR) techniques. While Large Language Model (LLM)-based APR systems have shown promise, prior studies...

<https://arxiv.org/abs/2608.14065>

4 CONF

P2Skill: Privacy Preserving Skill Distillation for Cloud-Local LLM Inference Systems

ArXiv CS.AI 50%

arXiv:2608.14094v1 Announce Type: cross Abstract: Cloud-local LLM inference systems have the potential to use the reasoning capability of large cloud models while protecting sensitive user data on personal devices. Cloud-bound requests must exclude personally identifiable...

<https://arxiv.org/abs/2608.14094>

5 CONF

Rewrite Once, Validate Anywhere: Producing OWL-Aware SHACL Constraints (Extended Version)

ArXiv CS.AI 50%

arXiv:2608.14104v1 Announce Type: cross Abstract: The Shapes Constraint Language (SHACL) is a W3C recommendation to express syntactic constraints, called shapes, on RDF graphs. SHACL validators are used to test whether a given graph adheres to such a shape. However, RDF graphs...

<https://arxiv.org/abs/2608.14104>

6 CONF

Act2Intention: A Benchmark For Developing Active Mobile Agents Through Inferring User Intention from GUI Actions

ArXiv CS.AI 50%

arXiv:2608.14132v1 Announce Type: cross Abstract: Mobile GUI Agents powered by multimodal large language models (MLLMs) show promise in human-computer intelligence. However, current research primarily focuses on reactive task execution while lacking a comprehensive...

<https://arxiv.org/abs/2608.14132>

7 **CONF****Correcting What You Cannot See: Credit Assignment for Perception Distillation in Multimodal Reasoners****ArXiv CS.AI** **50%**

arXiv:2607.28336v3 Announce Type: replace Abstract: On-policy distillation provides dense supervision for multimodal reasoners, but its trajectory-level reward cannot determine whether a failed answer arose from perception or subsequent reasoning. Perception Success Rate (PSR),...

<https://arxiv.org/abs/2607.28336>

8 **CONF****Concept Guidance: Precise, Training-Free Latent Control for Text-to-Image Generation****ArXiv CS.AI** **50%**

arXiv:2608.14172v1 Announce Type: cross Abstract: Text-to-image diffusion models have two major drawbacks that severely limit their practical utility: (1) standard models lack an intrinsic mechanism for continuous, concept-specific guidance (e.g., for precisely controlling how...

<https://arxiv.org/abs/2608.14172>

9 **CONF****From Field Data to Global Food Systems Intelligence: A Semantic Graph Framework for Sustainable Wheat Production****ArXiv CS.AI** **50%**

arXiv:2502.19507v2 Announce Type: replace Abstract: In response to the growing need for structured, interoperable agricultural data, this paper presents the Sustainable Wheat Production Datahub, a modular, graph-based framework that brings diverse wheat production datasets...

<https://arxiv.org/abs/2502.19507>

0 **CONF****How Much Do Legal RAG Systems Still Hallucinate?****ArXiv CS.AI** **50%**

arXiv:2608.14210v1 Announce Type: cross Abstract: Hallucination is a major challenge for retrieval-augmented generation (RAG) systems in the legal domain, where ungrounded answers can lead to serious consequences. To better understand this problem, we conduct a fine-grained...

<https://arxiv.org/abs/2608.14210>

1 **CONF****Acoustic UAV Detection in Battlefield Scenarios: Handling Noise, Domain Shift, and Weak Labels****ArXiv CS.AI** **50%**

arXiv:2608.14287v1 Announce Type: cross Abstract: Passive acoustic sensing offers a critical, cost-efficient, and, crucially, passive alternative for detecting small unmanned aerial vehicles. However, the practical deployment of acoustic systems is discouraged by extreme...

<https://arxiv.org/abs/2608.14287>

2 **CONF****A Four-Axis Trustworthiness Benchmark for LLM-as-Judge in Principle-Based Regulation****ArXiv CS.AI** **50%**

arXiv:2608.14329v1 Announce Type: cross Abstract: Principle-based regulation, with evaluative standards such as "fair, clear, and not misleading" or "deliver good outcomes", cannot be reduced to binary predicates, and LLM-as-judge is increasingly used as the substitute. Our...

<https://arxiv.org/abs/2608.14329>

- 3

CONF

A Hybrid LLM-Based Framework for Automated Security Annotation Generation in Business Process Models

ArXiv CS.AI

50%

arXiv:2608.14370v1 Announce Type: cross Abstract: The modelling and analysis of secure business processes require the incorporation of security annotations into process models. Although BPMN extensions, including SecBPMN2, exist for this purpose, the derivation of accurate and...

<https://arxiv.org/abs/2608.14370>
- 4

CONF

CFM-Bench: A Unified Multi-Domain, Multi-Task Benchmark for Channel Foundation Models

ArXiv CS.AI

50%

arXiv:2607.14975v2 Announce Type: replace Abstract: Channel foundation models (CFMs) are commonly evaluated in model-specific pipelines that differ in data, radio configurations, partitions, adaptation procedures, task definitions, and metrics, preventing reproducible comparison...

<https://arxiv.org/abs/2607.14975>
- 5

CONF

Can We Defend Against AI-Generated Video Attacks on Real-World Crisis Events? A Systematic Evaluation of Detectors, Generators and Social Dissemination

ArXiv CS.AI

50%

arXiv:2608.14391v1 Announce Type: cross Abstract: Recent video generators can fabricate realistic depictions of wars, disasters, public emergencies, and other real-world crises, creating substantial risks of misinformation. Existing benchmarks, however, provide limited evidence...

<https://arxiv.org/abs/2608.14391>
- 6

CONF

AI-Assisted Discovery and Construction of a Counterexample to the Convergence of Three-Block ADMM with the Identity Matrix as its Third Constraint Block

ArXiv CS.AI

50%

arXiv:2608.14396v1 Announce Type: cross Abstract: The alternating direction method of multipliers (ADMM), as a landmark algorithm, has attracted tremendous research attention and extensive practical applications over the past two decades. It is well known that, although the...

<https://arxiv.org/abs/2608.14396>
- 7

CONF

Whose doctor does the AI recommend? An algorithm audit of reputation and demographic signals in large language model-assisted physician choice

ArXiv CS.AI

50%

arXiv:2608.14399v1 Announce Type: cross Abstract: Patients increasingly ask large language model (LLM) assistants which doctor to see, making these systems AI infomediaries: algorithms that intermediate one person's choice among other people and thereby decide, silently and at...

<https://arxiv.org/abs/2608.14399>
- 8

CONF

Ensuring Safe Physical AI in Urban Mobility via Hazard-Informed Synthesized Envelopes

ArXiv CS.AI

50%

arXiv:2608.14481v1 Announce Type: cross Abstract: As heterogeneous robotic systems deploy across diverse urban zones, maintaining safety amid complex human-robot interactions remains a critical challenge. We present a unified framework that bridges systematic hazard analysis and...

<https://arxiv.org/abs/2608.14481>

9 CONF

Optimal Scheduling of Road Maintenance Jobs Considering Impact on Traffic Flows

ArXiv CS.AI 50%

arXiv:2608.14491v1 Announce Type: cross Abstract: Network-level maintenance planning requires repeated evaluations of equilibrium traffic flows under road capacity reductions. While equilibrium traffic assignment models are well established, their repeated solution quickly...

<https://arxiv.org/abs/2608.14491>

00 CONF

Generating Benchmark Health Data Using a Tabular Diffusion Transformer

ArXiv CS.AI 50%

arXiv:2608.14496v1 Announce Type: cross Abstract: Cross-Tabular Data Generation (CTDG) seeks to learn a generative model from multiple heterogeneous tables and produce new synthetic tabular datasets. However, existing synthetic tabular data generation methods are largely...

<https://arxiv.org/abs/2608.14496>

01 CONF

RecipeNet: A Hierarchical Transformer for Recipe Data

ArXiv CS.AI 50%

arXiv:2608.14505v1 Announce Type: cross Abstract: Recipe data arises in domains such as materials synthesis, pharmaceutical formulation, and industrial manufacturing, where procedures are represented as ordered sequences of steps containing heterogeneous structured fields....

<https://arxiv.org/abs/2608.14505>

02 CONF

Marionette: Predicting World States, Rendering Geometry, Painting Appearance

ArXiv CS.AI 50%

arXiv:2608.14530v1 Announce Type: cross Abstract: Interactive game world models typically autoregress visual observations directly in pixel or latent space, forcing structured properties such as pose, geometry, and occlusion to be implicitly maintained by the same generative...

<https://arxiv.org/abs/2608.14530>

03 CONF

Decoding the Past: An Uncertainty-Aware Deep Learning Framework for Sex Attribution in Prehistoric Hand Stencils

ArXiv CS.AI 50%

arXiv:2608.14539v1 Announce Type: cross Abstract: Determining the biological sex of the individuals who created Upper Paleolithic hand stencils remains a challenging problem due to the absence of ground truth, population differences between contemporary and prehistoric groups,...

<https://arxiv.org/abs/2608.14539>

04 CONF

The Metacognitive Bottleneck: Japanese Riddles Reveal Fundamental Limits of Machine Insight and Self-Evaluation in Reasoning AI

ArXiv CS.AI 50%

arXiv:2509.14704v3 Announce Type: replace Abstract: Benchmark saturation and training-data contamination increasingly obscure whether reported gains in large language models (LLMs) reflect genuine advances in reasoning or familiarity with recurring patterns in benchmark...

<https://arxiv.org/abs/2509.14704>

05 CONF

BUZZY: Contrastive Scoring to Mitigate Text-Induced Bias in Multimodal Multiple-Choice QA

ArXiv CS.AI 50%

arXiv:2603.28026v2 Announce Type: replace Abstract: Multimodal multiple-choice question answering (MCQA) provides a standardized and objectively measurable setting for evaluating vision-language models (VLMs). However, because the MCQA format incorporates the candidate choices...

<https://arxiv.org/abs/2603.28026>

06 CONF

GRASP: Gated Regression-Aware Skill Proposer for Self-Improving LLM Agents

ArXiv CS.AI 50%

arXiv:2605.29668v2 Announce Type: replace Abstract: LLM agents acting in structured environments fail in operational rather than conversational ways, and reliability depends on procedural knowledge of the environment. Prior self-improvement methods accumulate natural-language...

<https://arxiv.org/abs/2605.29668>

07 CONF

EchoChange: A Diffusion Language Model with Dual Pass Remasking for Factual Remote Sensing Disaster Change Captioning

ArXiv CS.AI 50%

arXiv:2608.01856v2 Announce Type: replace Abstract: Bi-temporal remote-sensing disaster change captioning often needs to identify sparse and spatially localized changes across large pre- and post-event scenes and then translate them into coherent, factual descriptions. However,...

<https://arxiv.org/abs/2608.01856>

08 CONF

Making AI-Generated Feedback Matter: A Large-Scale Study of Feedback Workflows and Student Enactment

ArXiv CS.AI 50%

arXiv:2608.11625v2 Announce Type: replace Abstract: Feedback processes strongly influence student learning, yet their educational value depends on addressing two distinct challenges: providing high-quality, timely, and individualised feedback at scale, and supporting students to...

<https://arxiv.org/abs/2608.11625>

09 CONF

LLM-Guided Graph Generation for Structure-Based Local Improvement Methods

ArXiv CS.AI 50%

arXiv:2608.13333v2 Announce Type: replace Abstract: Large neighborhood search normally selects a random subset of decision variables for iterative optimization. To efficiently solve various problems, researchers tend to design variable selection strategies that take into account...

<https://arxiv.org/abs/2608.13333>

10 CONF

Leveraging Few-Shot Learning and Large Language Models for Analyzing Blood Pressure Variations Across Biological Sex from Scientific Literature

ArXiv CS.AI 50%

arXiv:2402.01826v2 Announce Type: replace-cross Abstract: Current blood pressure (BP) technologies and standards were established decades ago, and these standards are still used worldwide today, often without adjusting BP readings for individual demographic factors such as sex...

<https://arxiv.org/abs/2402.01826>

11 CONF

LLM-Advisor: An LLM Advisor for Cost-efficient Path Planning across Multiple Terrains

ArXiv CS.AI 50%

arXiv:2503.01236v3 Announce Type: replace-cross Abstract: This paper addresses fixed-graph terrain-aware path refinement, in which a global planner is restricted to a predefined route space and may remain optimal within that space while missing lower-cost terrain corridors...

<https://arxiv.org/abs/2503.01236>

12 CONF

Redefining Generalization in Visual Domains: A Two-Axis Framework for Fake Image Detection with FusionDetect

ArXiv CS.AI 50%

arXiv:2510.05740v2 Announce Type: replace-cross Abstract: The rapid development of generative models has made it increasingly crucial to develop detectors that can reliably detect synthetic images. Although most of the work has now focused on cross-generator generalization, we...

<https://arxiv.org/abs/2510.05740>

13 CONF

Lost in Phonation: Voice Quality Variation as an Evaluation Dimension for Speech Foundation Models

ArXiv CS.AI 50%

arXiv:2510.25577v2 Announce Type: replace-cross Abstract: Recent advances in Speech Foundation Models (SFMs) enable direct processing of raw audio, allowing models to respond to subtle paralinguistic variation. However, how these models interpret non-lexical cues remains largely...

<https://arxiv.org/abs/2510.25577>

14 CONF

PEFT-MuTS: A Multivariate Parameter-Efficient Fine-Tuning Framework for Remaining Useful Life Prediction based on Cross-domain Time Series Representation Model

ArXiv CS.AI 50%

arXiv:2601.22631v2 Announce Type: replace-cross Abstract: The application of data-driven remaining useful life (RUL) prediction has long been constrained by the availability of large amount of degradation data. Mainstream solutions such as domain adaptation and meta-learning...

<https://arxiv.org/abs/2601.22631>

15 CONF

AtomBridge: Agentic VLA Inference Plugin for Long-Horizon Tasks in Scientific Experiments

ArXiv CS.AI 50%

arXiv:2602.09430v2 Announce Type: replace-cross Abstract: Robotic laboratories play a critical role in autonomous scientific discovery by enabling scalable, continuous experimental execution. Recent vision-language-action (VLA) models offer a promising foundation for robotic...

<https://arxiv.org/abs/2602.09430>

16 CONF

A Unified Assessment of the Poverty of the Stimulus Argument for Neural Language Models

ArXiv CS.AI 50%

arXiv:2602.09992v2 Announce Type: replace-cross Abstract: Several recent contributions have evaluated the Poverty of the Stimulus Hypothesis (PoSH) using Artificial Neural Networks (ANNs). The results suggest that ANN-based language models can acquire certain structure-dependent...

<https://arxiv.org/abs/2602.09992>

17 CONF

A Systematic Comparison of Training Objectives for Out-of-Distribution Detection in Image Classification

ArXiv CS.AI 50%

arXiv:2603.07571v3 Announce Type: replace-cross Abstract: Out-of-distribution (OOD) detection is critical in safety-sensitive applications. While this challenge has been addressed from various perspectives, the influence of training objectives on OOD behavior remains...

<https://arxiv.org/abs/2603.07571>

18 CONF

Beyond aggregate scores: Deployment-aware and non-compensatory benchmarking of vision-based eye-state recognition models for driver monitoring

ArXiv CS.AI 50%

arXiv:2606.08123v2 Announce Type: replace-cross Abstract: Model selection for safety-relevant visual recognition is often based on clean aggregate performance, although robustness, transfer, embedded latency, and explanation faithfulness may produce different preferences. This...

<https://arxiv.org/abs/2606.08123>

19 CONF

Breaking Chains with Trees: Model-Parallel Deep Learning with $\mathcal{O}(\log N)$ Time Complexity

ArXiv CS.AI 50%

arXiv:2606.21497v2 Announce Type: replace-cross Abstract: Modern deep neural networks are trained using error backpropagation, which requires sequential forward and backward computations across network layers. As these networks become deeper, this introduces limitations, since...

<https://arxiv.org/abs/2606.21497>

20 CONF

TypeProbe: Recovering Type Representations from Hidden States of Pre-trained Code Models

ArXiv CS.AI 50%

arXiv:2607.08339v2 Announce Type: replace-cross Abstract: State-of-the-art code models achieve impressive performance, yet the extent to which they internally encode type information remains poorly understood. We probe the residual streams of pretrained code models for internal...

<https://arxiv.org/abs/2607.08339>

21 CONF

A Negative-Control Protocol for Clinical EEG Foundation-Model Benchmarks: Dataset Identity and External-Cohort Stress Testing

ArXiv CS.AI 50%

arXiv:2607.24519v3 Announce Type: replace-cross Abstract: EEG foundation-model gains may depend on cohort, montage, or probe design. We evaluated five models on five tasks across four benchmark datasets plus Korean CAUEEG, using subject-disjoint validation where identifiers...

<https://arxiv.org/abs/2607.24519>

22 CONF

Cloud-ScPO: Hidden-State Geometry for Semi-Supervised Preference Optimization in LLM Reasoning

ArXiv CS.AI 50%

arXiv:2608.01014v2 Announce Type: replace-cross Abstract: Preference optimization improves mathematical reasoning in large language models (LLMs), but reliable chosen-rejected pairs usually require verified answers, human annotations, or external reward models. We investigate...

<https://arxiv.org/abs/2608.01014>

23 CONF

Can AI Agents Simulate A/B Test Outcomes? A Validation Framework for Agentic Experimentation

ArXiv CS.AI 50%

arXiv:2608.02345v2 Announce Type: replace-cross Abstract: A/B testing remains the standard for rolling out new features in the technology industry. Each experiment, however, consumes real traffic, engineering effort, and weeks of wall-clock time. Can AI agents---conditioned on...

<https://arxiv.org/abs/2608.02345>

24 CONF

Discovering Efficient and Explainable Communication Topologies for LLM-based Multi-Agent Systems via Causal Inference

ArXiv CS.AI 50%

arXiv:2608.12921v2 Announce Type: replace-cross Abstract: The performance of large language model (LLM)-based multi-agent systems (MAS) largely depends on effective communication topologies. Existing topology generation methods, however, typically learn communication topologies...

<https://arxiv.org/abs/2608.12921>