



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	234	Media
TechCrunch AI	9	Media
HackerNews	6	Community
The Verge AI	2	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Show HN: Sokoban AI Solver

HackerNews

+1 source

90%

HN 67

arXiv:2608.15958v1 Announce Type: new Abstract: Deciding whether a Sokoban puzzle is solvable is PSPACE-complete (Culberson, 1997): solutions can be exponentially long and there is no short certificate to check. Solvability is also a fragile property, since even a single...

<https://mkornreich.me/projects/sokoban/>

VERI

Show HN: AI Text Watermarking Playground

HackerNews

+1 source

90%

HN 4

In the 1990s, Barbara Sherwood Lollar descended into the Kidd Creek mine in northern Ontario, which cuts more than three kilometers into the ancient root of North America. There her team of geochemists found water that had been confined underground for more than a billion years....

<https://watermark.keito.me/>

VERI

Anthropic explains how Claude's invisible text watermarks will work

The Verge AI

+1 source

80%

Anthropic has clarified how it's planning to apply invisible watermarks to Claude-generated text in order to comply with Europe's AI transparency rules. On Friday, Anthropic announced that Claude's text marking system is "a version of the SynthID-Text approach" - an open-source...

<https://www.theverge.com/ai-artificial-intelligence/980869/anthropic-claude-watermarks-...>

CONFIRMED SIGNALS — 246 stories

CONF

Show HN: Tesana - An AI game engine that builds quality games end-to-end

HackerNews

50%

HN 9

Hey HN! I built <https://tesana.ai/en>, a game generation platform that allows anyone to create full games with graphics, game logic, and animations from a simple text prompt - without coding. Tesana has three key features 1. describe your game idea and have Tesana write a game...

<https://tesana.ai/en>

CONF

Show HN: AI-generated music box version of any song (for baby)

HackerNews

50%

HN 3

title basically covers it. project #2 from my paternity leave. My partner & I got frustrated with hearing Brahms's Lullaby and Wheels On The Bus repeatedly. I also stand by that babies can enjoy adult music just as easily as they can enjoy "kids music". But most adult music is...

<https://lullabyconverter.com>

CONF

Show HN: I turned the ending of Trainspotting into a tiny web art installation

HackerNews

50%

HN 3

I have always loved the final moments of Trainspotting and Born Slippy! I tried turning them into a small audiovisual experience on the web. Put on headphones and give it a couple of minutes.

<https://trainspotting.vercel.app>

CONF

Show HN: I backtested my own stock site's signals and published what failed

HackerNews

50%

HN 2

No summary — see link for full article.

<https://matterofstocks.com/>

CONF

Anthropic's annualized revenue surges to \$65B

TechCrunch AI 50%

The model maker added \$18 billion in annualized revenue in two months.

<https://techcrunch.com/2026/08/17/anthropics-annualized-revenue-surges-to-65b/>

CONF

AI automation startup Relay shuts down, staff joins Google's Chrome team

TechCrunch AI 50%

"We have some really ambitious plans to help you work with AI in Chrome to get things done, and I'll have more to share soon," Jacob Bank, Relay founder and CEO, said.

<https://techcrunch.com/2026/08/17/ai-automation-startup-relay-shuts-down-staff-joins-go...>

0 CONF

Amazon, which started off selling books, is destroying rare texts to train AI

TechCrunch AI 50%

Rare books are incredibly valuable for training LLMs, since these models have already trained on whatever's available online.

<https://techcrunch.com/2026/08/17/amazon-once-an-online-bookseller-is-destroying-rare-b...>

1 CONF

Groq raises \$350M to fuel its pivot from AI chips to neocloud

TechCrunch AI 50%

Groq raised \$350 million at a \$3.5 billion valuation as the former AI chipmaker pivots to a neocloud business and expands its Nvidia-powered data center footprint.

<https://techcrunch.com/2026/08/17/groq-raises-350m-to-fuel-its-pivot-from-ai-chips-to-n...>

2 CONF

Nvidia investing \$1.5B in SoftBank data center developer behind OpenAI project

TechCrunch AI 50%

Nvidia's investment in SoftBank's data center developer will guarantee its chips power an OpenAI data center.

<https://techcrunch.com/2026/08/17/nvidia-investing-1-5b-in-softbank-data-center-develop...>

3 CONF

Wispr raises \$280M at \$2B valuation as it looks beyond dictation

TechCrunch AI 50%

The funds will allow Wispr to increase its footprint as it ventures into new areas, such as meetings, with its newly released note-taker tool.

<https://techcrunch.com/2026/08/17/wispr-raises-280m-at-2b-valuation-as-it-looks-beyond-...>

4 CONF

How to tell if your AI platforms' accounts have been hacked

TechCrunch AI 50%

A guide on how to check if hackers have broken into your accounts on the most popular AI platforms.

<https://techcrunch.com/2026/08/15/how-to-tell-if-your-ai-platforms-accounts-have-been-h...>

5 CONF

Kog is going deeper to squeeze more inference out of GPUs

TechCrunch AI 50%

The idea that GPUs are poorly suited for agentic workflows may be a misconception, according to French startup Kog.

<https://techcrunch.com/2026/08/14/kog-is-going-deeper-to-squeeze-more-inference-out-of-...>

6 CONF

Whisker's AI-powered litter robot thinks my cats swapped bodies

The Verge AI 50%

The greatest invention in pet tech in recent years is the litter robot. A machine that scoops your kitties' poop so you don't have to - what else could a cat owner possibly want? How about insights into your kitty's litter box usage that could flag health issues? Sign me up. But...

<https://www.theverge.com/tech/978323/whisker-litter-robot-5-pro-review>

7 CONF

Bactrainus: Optimizing Large Language Models for Multi-hop Complex Question Answering Tasks

ArXiv CS.AI 50%

arXiv:2501.06286v2 Announce Type: replace-cross Abstract: Multi-hop question answering requires a system to identify and integrate evidence distributed across documents, yet large language models remain vulnerable to irrelevant context. We investigate this evidence bottleneck in...

<https://arxiv.org/abs/2501.06286>

8 CONF

The Unwritten Benchmark: A New Challenge for Multimodal Machine Learning in Abstract Perceptual Reasoning

ArXiv CS.AI 50%

arXiv:2608.14558v1 Announce Type: new Abstract: Current multimodal models have demonstrated remarkable proficiency in recognizing static visual and auditory content. However, their capacity for abstract perceptual reasoning, inferring unseen information from dynamic, generative...

<https://arxiv.org/abs/2608.14558>

9 CONF

Global AI Regulations for FAIR and Ethics in High-Risk Use Cases: A Comparative Review

ArXiv CS.AI 50%

arXiv:2608.14562v1 Announce Type: new Abstract: AI governance is shifting from voluntary ethics to enforceable, risk-based regulation, yet cross-jurisdictional divergence creates compliance uncertainty for operators of high-stakes AI. We present a comparative matrix for the EU,...

<https://arxiv.org/abs/2608.14562>

0 CONF

Position: AI Governance Needs ISO-like Interoperability Protocols, Not Just Laws

ArXiv CS.AI 50%

arXiv:2608.14568v1 Announce Type: new Abstract: As Artificial Intelligence (AI) systems become deeply integrated into critical global infrastructure, the urgency for robust governance frameworks has intensified. However, current approaches, led by jurisdiction-specific laws,...

<https://arxiv.org/abs/2608.14568>

1 CONF

Validation-Frontier Representation Selection under Constrained Observation

ArXiv CS.AI 50%

arXiv:2608.15095v1 Announce Type: new Abstract: AI systems deployed outside clean benchmark settings often rely on observations that are incomplete, unstable, costly, or degraded by monitoring failures. This paper studies representation selection under constrained observation:...

<https://arxiv.org/abs/2608.15095>

2 **CONF****SymbolicLight V1: Spike-Gated Dual-Path Language Modeling at High Activation Sparsity****ArXiv CS.AI 50%**

arXiv:2605.21333v2 Announce Type: replace-cross Abstract: Natively trained spiking language models must preserve information across time while operating through sparse binary activations, a combination that has produced a persistent quality gap relative to dense Transformers. We...

<https://arxiv.org/abs/2605.21333>

3 **CONF****An Agentic Framework Using Rules and LLMs for Embedding and Annotating Descriptive Document Layouts: A Plant Science Use Case****ArXiv CS.AI 50%**

arXiv:2608.14587v1 Announce Type: new Abstract: Background: Recent advances in information retrieval (IR) leverage both dense and sparse representations, large language models (LLMs), and specialized retrieval models to improve ranking accuracy, relevance, and cross-lingual...

<https://arxiv.org/abs/2608.14587>

4 **CONF****The Hallucination Snowball: Modeling Error Propagation as State Transitions in Multi-Agent LLM Pipelines****ArXiv CS.AI 50%**

arXiv:2608.14588v1 Announce Type: new Abstract: Sequential multi-agent LLM pipelines chain specialized agents without verification at handoffs, creating a structural flaw with measurable and severe consequences. We show that hallucinations injected at Stage 1 do not merely...

<https://arxiv.org/abs/2608.14588>

5 **CONF****Toward Safe LLM Agents: A Survey of Specification, Verification, and Enforcement****ArXiv CS.AI 50%**

arXiv:2608.14590v1 Announce Type: new Abstract: LLM agents increasingly perform irreversible real-world actions, including database updates, API calls, file operations, and autonomous use of tools. However, no existing system provides formally grounded, task-level safety...

<https://arxiv.org/abs/2608.14590>

6 **CONF****Position: Medical AI Neglects Real Treatment Outcomes****ArXiv CS.AI 50%**

arXiv:2608.14598v1 Announce Type: new Abstract: Medical AI has rapidly improved its ability to perform diagnostic and prognostic tasks that lead to treatment decisions. But understanding of treatment itself is still inadequately trained and evaluated, using human opinions and...

<https://arxiv.org/abs/2608.14598>

7 **CONF****When Do LLMs Apply the Wrong Law? Diagnosing LLM Failures in Temporal Legal Reasoning****ArXiv CS.AI 50%**

arXiv:2608.14610v1 Announce Type: new Abstract: Legal reasoning tasks such as legal judgment prediction (LJP) require identifying the temporally correct version of the law governing a case -- a capability we term temporal applicable-law determination. However, whether large...

<https://arxiv.org/abs/2608.14610>

8 CONF

Defake-o3: From Speculative Rationales to Verifiable Evidence for Explainable AIGI Detection

ArXiv CS.AI 50%

arXiv:2608.16259v1 Announce Type: cross Abstract: The rapid progress of image generation models calls for AI-generated image (AIGI) detectors that are not only accurate but also explainable and reliable. While MLLM-based detectors can provide natural language explanations,...

<https://arxiv.org/abs/2608.16259>

9 CONF

Class Imbalance and Batch Effects in LLM-Based Screening for Systematic Reviews

ArXiv CS.AI 50%

arXiv:2608.14737v1 Announce Type: cross Abstract: This study analyses LLMs in imbalanced binary classification, using study screening in systematic reviews as the application domain. An experiment was conducted in five reviews, comparing individual and batch processing, with and...

<https://arxiv.org/abs/2608.14737>

0 CONF

Accuracy and Reliability of Large Language Models in Cosmetic Chemistry and Skin Health: A Benchmarking Study

ArXiv CS.AI 50%

arXiv:2608.14631v1 Announce Type: new Abstract: As consumers increasingly turn to AI chatbots for skincare advice, the technical accuracy of Large Language Models (LLMs) in cosmetic chemistry remains largely under-evaluated. We benchmarked 14 LLMs on a structured set of topics...

<https://arxiv.org/abs/2608.14631>

1 CONF

Task- and Session-Level Model Routing: A Common-Interface Hybrid Evaluation of Four Open-Source Routers Across Four Benchmarks

ArXiv CS.AI 50%

arXiv:2608.14641v1 Announce Type: new Abstract: Agentic systems increasingly delegate model selection to a router, yet open-source routers are usually evaluated with different tasks, candidate pools, and execution protocols, limiting direct comparison. We present a common...

<https://arxiv.org/abs/2608.14641>

2 CONF

Evaluating Agentic Code Repair Capabilities in Distributed Systems

ArXiv CS.AI 50%

arXiv:2608.14863v1 Announce Type: cross Abstract: LLM-based coding agents have advanced rapidly on single-process SWE tasks, with frontier models now clustering in the high-70s on SWE-bench Verified. Distributed-system debugging, however, remains an under-explored regime: bugs...

<https://arxiv.org/abs/2608.14863>

3 CONF

When Uncertainty Isn't Enough: An Empirical Study of Self-Correction in Code Generation

ArXiv CS.AI 50%

arXiv:2608.14659v1 Announce Type: new Abstract: Large language models for code generation often produce incorrect solutions without reliable indicators of failure. We study whether uncertainty estimation methods developed for natural language transfer to code generation, and...

<https://arxiv.org/abs/2608.14659>

4 CONF

Cross-Domain Industrial Fault Detection by Causal Mechanism Monitoring

ArXiv CS.AI 50%

arXiv:2608.14666v1 Announce Type: new Abstract: Unsupervised fault detection in industrial systems is dominated by reconstruction based methods that monitor individual sensor marginal distributions. This misses coupling faults, where the physical relationship between sensor...

<https://arxiv.org/abs/2608.14666>

5 CONF

Auditing an AI-Generated Mathematical Proof: A Correction to a Greedy Conditioning Lemma in Quantum Parallel Repetition

ArXiv CS.AI 50%

arXiv:2608.14673v1 Announce Type: new Abstract: Chapter 6 of OpenAI's *Ten Advances in Mathematics and Theoretical Computer Science* claims an exponential parallel-repetition theorem for all finite two-player, one-round entangled games. Early in the proof, the chapter uses a...

<https://arxiv.org/abs/2608.14673>

6 CONF

When Agentic Executions Fail: Detecting and Localizing Runtime Faults from Telemetry

ArXiv CS.AI 50%

arXiv:2608.14680v1 Announce Type: new Abstract: Reliability in LLM-based agentic systems is a property of the whole execution (its tool calls, model calls, guardrails, and inter-agent messages), not of the final answer alone, yet evaluating only task outcomes reveals little...

<https://arxiv.org/abs/2608.14680>

7 CONF

A concentration result for multilayer feedforward neural networks

ArXiv CS.AI 50%

arXiv:2608.15335v1 Announce Type: new Abstract: We consider for an arbitrary fixed ρ and for each positive integer n a multilayer feedforward artificial neural network with ρ layers, n neurons in the first layer (the input layer) and only one neuron, the output...

<https://arxiv.org/abs/2608.15335>

8 CONF

Emergent Misaligned Communication in Long-Horizon Multi-Agent LLM Commerce

ArXiv CS.AI 50%

arXiv:2608.14825v1 Announce Type: cross Abstract: Frontier LLM agents increasingly transact on behalf of separate principals, often using natural language rather than structured APIs. Much of the safety literature studies misaligned LLM behavior through adversarial-elicitation...

<https://arxiv.org/abs/2608.14825>

9 CONF

From Errors to Proofs: Minimal-Core-Guided Repair for Neuro-Symbolic Constraint Solving

ArXiv CS.AI 50%

arXiv:2608.14771v1 Announce Type: new Abstract: Making language models solve constraint problems reliably often means having them translate the problem into a formal specification and delegating the search to a sound solver. But the translation is itself a language-model task,...

<https://arxiv.org/abs/2608.14771>

0 **CONF****CEDAR-GRPO: Process-Aware Reinforcement Learning for General Abductive Reasoning in LLMs****ArXiv CS.AI 50%**

arXiv:2608.14791v1 Announce Type: new Abstract: Abductive reasoning, often characterized as inference to the best explanation, is central to explanation under uncertainty, from everyday sense-making and investigation to scientific discovery. Yet LLM research has mostly studied...

<https://arxiv.org/abs/2608.14791>

1 **CONF****Do LLMs Know What to Ask and When? Evaluating Multi-Turn Information Seeking****ArXiv CS.AI 50%**

arXiv:2608.14808v1 Announce Type: new Abstract: When a user question is underspecified, a capable model should recognize that its context is insufficient, identify the missing information, ask for it, and respond only once that information determines a unique answer. We...

<https://arxiv.org/abs/2608.14808>

2 **CONF****MINT: Min-Selection Preference Distillation for Balanced Multi-Objective Alignment****ArXiv CS.AI 50%**

arXiv:2608.14828v1 Announce Type: new Abstract: Aligning a language agent to several objectives at once is a persistent failure mode of preference-based training: when objectives are combined additively, optimization collapses onto whichever is cheapest to improve and sacrifices...

<https://arxiv.org/abs/2608.14828>

3 **CONF****What the Reranker Sees: Multi-Aspect Page Annotation for Long-Document Multimodal Question Answering****ArXiv CS.AI 50%**

arXiv:2608.14841v1 Announce Type: new Abstract: Long-document visual question answering (VQA) over documents of tens to hundreds of pages mixing text, tables, charts, and figures typically follows retrieve-then-read pipelines. In our setting, the bottleneck shifts from retrieval...

<https://arxiv.org/abs/2608.14841>

4 **CONF****Command-Space Counterfactual Explanations for Pareto-Conditioned Reinforcement Learning****ArXiv CS.AI 50%**

arXiv:2608.14963v1 Announce Type: cross Abstract: Pareto Conditioned Networks learn multiple multi-objective reinforcement learning behaviours by conditioning a single policy on a desired return command. However, the local mapping from command and state to action remains opaque....

<https://arxiv.org/abs/2608.14963>

5 **CONF****Personalized Auto-Research: Towards a True AI Co-Scientist****ArXiv CS.AI 50%**

arXiv:2608.14881v1 Announce Type: new Abstract: AI co-scientists that generate hypotheses, retrieve related work, design experiments, execute code, and draft full papers are beginning to change how research is carried out. Despite this rapid progress, state-of-the-art systems...

<https://arxiv.org/abs/2608.14881>

6 CONF

Frontier AI Forecasting Has a Measurement Problem: An Audit of Progress Evidence

ArXiv CS.AI 50%

arXiv:2608.14903v1 Announce Type: new Abstract: Quantitative forecasts of frontier artificial intelligence often connect dated targets to trends in benchmark scores, training compute, release time, or expert belief. This paper audits whether the public measurement record...

<https://arxiv.org/abs/2608.14903>

7 CONF

LLMs Can Predict Failure Risk, But Struggle to Predict Which Collaboration Protocol Pays Off: Cost-Aware Protocol Routing Across Reasoning Tasks

ArXiv CS.AI 50%

arXiv:2608.14927v1 Announce Type: new Abstract: Multi-agent large language model (LLM) systems can improve reasoning by spending more computation, but deployment requires deciding when extra collaboration is worth its cost. We isolate this decision by running every problem under...

<https://arxiv.org/abs/2608.14927>

8 CONF

RETRACE: Resilience-Guided Trait-Conditioned Craving Estimation from Wearable Physiology in Opioid Use Disorder

ArXiv CS.AI 50%

arXiv:2608.14947v1 Announce Type: new Abstract: Detecting opioid craving from wearable physiological signals is critical yet difficult, with the potential to support proactive interventions for individuals with opioid use disorder (OUD). This challenge is especially pronounced...

<https://arxiv.org/abs/2608.14947>

9 CONF

Demand-Driven Vertiport Siting and Discrete-Event Fleet Simulation for On-Demand Urban Air Mobility Network Design

ArXiv CS.AI 50%

arXiv:2608.14974v1 Announce Type: new Abstract: This paper presents a demand-driven framework for on-demand Urban Air Mobility (UAM) network design that links vertiport siting, fleet simulation, and door-to-door travel-time feasibility. Demand is estimated from commuter and...

<https://arxiv.org/abs/2608.14974>

10 CONF

Does a Tool Result Carry More Authority Than Plain Text? Three Prospective Studies of False-Claim Adoption in a Synthetic Assignment Task with Claude Opus 5

ArXiv CS.AI 50%

arXiv:2608.14992v1 Announce Type: new Abstract: Language-model systems increasingly read from stores they also write to, so a claim that was merely written earlier can return looking retrieved. We tested whether the message package carrying an unsupported assignment changes...

<https://arxiv.org/abs/2608.14992>

11 CONF

LLM-Based Hierarchical Coordinated Control with Continuation-Aware Policy Learning

ArXiv CS.AI 50%

arXiv:2608.15041v1 Announce Type: new Abstract: Coordinating multiple interacting units in complex engineering systems is challenging when system interactions are difficult to model, operational information is heterogeneous, and low-level actions must satisfy strict constraints....

<https://arxiv.org/abs/2608.15041>

2 CONF

OpenGPT-4o-Image: A Comprehensive Dataset for Advanced Image Generation and Editing

ArXiv CS.AI 50%

arXiv:2509.24900v2 Announce Type: replace-cross Abstract: The performance of unified multimodal models for image generation and editing is fundamentally constrained by the quality and comprehensiveness of their training data. While existing datasets have covered basic tasks like...

<https://arxiv.org/abs/2509.24900>

3 CONF

GraphLoom: Reliability-Calibrated Graph Evidence Routing for Multimodal KG-RAG

ArXiv CS.AI 50%

arXiv:2608.15056v1 Announce Type: new Abstract: Multimodal retrieval-augmented generation (RAG) systems often rely on long unstructured contexts or aggressively expanded evidence graphs, which can introduce noisy evidence, weaken multi-hop reasoning, and increase unsupported...

<https://arxiv.org/abs/2608.15056>

4 CONF

Comprehensive Benchmarking of Deep Learning Architectures for Lung Cancer Histopathology

ArXiv CS.AI 50%

arXiv:2608.15915v1 Announce Type: cross Abstract: Lung cancer remains the leading cause of cancer-related mortality worldwide, while histopathological diagnosis is often affected by inter-observer variability and the substantial workload associated with manual slide examination...

<https://arxiv.org/abs/2608.15915>

5 CONF

Data-knowledge dual-driven intelligent framework for full-chain, experiment-efficient synthesis of 2D dendrites

ArXiv CS.AI 50%

arXiv:2603.16959v2 Announce Type: replace-cross Abstract: Exemplified by the chemical vapor deposition growth of two-dimensional dendrites, which has potential applications in catalysis and presents a parameter-intensive, data-scarce and reaction process-complex model problem,...

<https://arxiv.org/abs/2603.16959>

6 CONF

Second-Order Policy Effects as State Transitions: A Source-Linked Benchmark for Policy Simulation

ArXiv CS.AI 50%

arXiv:2608.15101v1 Announce Type: new Abstract: Policy evaluation often estimates direct benefits and costs while treating the institutional environment as fixed. In practice, a policy changes the system it enters: actors adapt, enforcement capacity shifts, burdens move, and new...

<https://arxiv.org/abs/2608.15101>

7 CONF

CAPO: Constraint-Aware Prompt Optimization for LLM Agents

ArXiv CS.AI 50%

arXiv:2608.16068v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed as agents that rely on system prompts to use tools and complete tasks. Such deployments impose distinct operational requirements, including appropriate tool use, concise...

<https://arxiv.org/abs/2608.16068>

8 CONF

Platform Adaptation Under Governance Interventions: Actor Best-Response Modeling and an External Public-Case Benchmark

ArXiv CS.AI 50%

arXiv:2608.15131v1 Announce Type: new Abstract: Digital platforms govern by changing rules: rankings, monetization thresholds, moderation standards, verification systems, disclosure requirements, appeal processes, and access policies. These interventions are rarely absorbed...

<https://arxiv.org/abs/2608.15131>

9 CONF

Tracking the Truth: Object-Centric Spatio-Temporal Monitoring for Video Large Language Models

ArXiv CS.AI 50%

arXiv:2605.08974v2 Announce Type: replace-cross Abstract: While multimodal large language models (MLLMs) have advanced video understanding, they remain highly prone to hallucinations in dynamic scenes. We argue this stems from a failure in spatio-temporal monitoring, the ability...

<https://arxiv.org/abs/2605.08974>

0 CONF

Translating finite-domain integer constraint models to CP/SMT/ILP/PB/SAT solvers with CPMpy

ArXiv CS.AI 50%

arXiv:2608.15143v1 Announce Type: new Abstract: Constraint solving is a declarative approach for solving combinatorial satisfaction and optimization problems. The user specifies their problem through constraints and decision variables, and a generic solver is used to find a...

<https://arxiv.org/abs/2608.15143>

1 CONF

OODBench: Out-of-Distribution Benchmark for Large Vision-Language Models

ArXiv CS.AI 50%

arXiv:2602.18094v2 Announce Type: replace-cross Abstract: Existing Visual-Language Models (VLMs) have achieved significant progress by being trained on massive-scale datasets, typically under the assumption that data are independent and identically distributed (IID). However, in...

<https://arxiv.org/abs/2602.18094>

2 CONF

TRCA: Transition-wise Rubric Credit Assignment for Long-horizon LLM Agents

ArXiv CS.AI 50%

arXiv:2608.16156v1 Announce Type: new Abstract: Long-horizon large language model (LLM) agents are typically optimized with sparse terminal outcomes, making fine-grained credit assignment across multi-step interactions difficult. Existing approaches either rely on process...

<https://arxiv.org/abs/2608.16156>

3 CONF

Domain-Specific Text Embedding Models for Entity Resolution

ArXiv CS.AI 50%

arXiv:2608.16161v1 Announce Type: cross Abstract: General-purpose text embedding models are designed to capture semantic similarity but are not optimised for distinguishing entity records that represent the same real-world business or person. This limitation affects applications...

<https://arxiv.org/abs/2608.16161>

4 CONF

Towards Standardized Evaluation in Automated Domain Modeling: Introducing a Benchmark

ArXiv CS.AI 50%

arXiv:2608.15255v1 Announce Type: new Abstract: Domain modeling plays an essential role in domain-driven design, capturing essential entities and their relationships within a specific domain. Despite advancements in automated domain modeling, the absence of standardized...

<https://arxiv.org/abs/2608.15255>

5 CONF

VibeWorlding: Can Multimodal Agents Construct 3D Open Worlds End-to-End?

ArXiv CS.AI 50%

arXiv:2608.15265v1 Announce Type: new Abstract: Constructing an interactive 3D open world from a user query is important. However, existing methods are primarily evaluated on idealized, simple queries, making it difficult to systematically analyze and compare how multimodal...

<https://arxiv.org/abs/2608.15265>

6 CONF

 $\$D^{\{2\}}R^{\{2\}}\$$: Discrete Diffusion with Regulation Reinforcement for Single-Cell Perturbation Prediction

ArXiv CS.AI 50%

arXiv:2608.15288v1 Announce Type: new Abstract: Predicting single-cell transcriptomic responses to genetic perturbations is central to functional genomics and virtual-cell modeling. Existing approaches, however, typically predict an entire expression profile as a whole, leaving...

<https://arxiv.org/abs/2608.15288>

7 CONF

Understanding Cognition-Induced Risks in Agentic AI Systems

ArXiv CS.AI 50%

arXiv:2608.15304v1 Announce Type: new Abstract: Frontier agentic systems powered by large language models (LLMs) exhibit human-like patterns of cognition. As these systems become deeply integrated across different domains, their cognitive engagement raises critical concerns for...

<https://arxiv.org/abs/2608.15304>

8 CONF

Physics-informed VAE-EVT for Tail Aware Radio Map Prediction

ArXiv CS.AI 50%

arXiv:2608.15314v1 Announce Type: new Abstract: Ultra-reliable low-latency communication (URLLC) requires precise identification of spatial regions where the signal-to-noise ratio (SNR) falls below an outage threshold. In this context, an outage refers to instances in which SNR...

<https://arxiv.org/abs/2608.15314>

9 CONF

VARM-Bench: Benchmarking Verifiable Structured Reasoning in Chinese Abusive Speech Moderation

ArXiv CS.AI 50%

arXiv:2608.15600v1 Announce Type: new Abstract: The widespread circulation of abusive online content has increased the need for reliable moderation of Chinese social-media text. Existing Chinese benchmarks support label classification, fine-grained toxicity categorization, and...

<https://arxiv.org/abs/2608.15600>

0 **CONF****Incoherent by Design? On the Moral Self-Consistency of LLMs****ArXiv CS.AI 50%**

arXiv:2608.15354v1 Announce Type: new Abstract: LLMs are increasingly used in morally sensitive contexts, yet it is unclear whether they apply ethical principles consistently across situations. A model that can state a moral principle may still violate it when the same scenario...

<https://arxiv.org/abs/2608.15354>

1 **CONF****Grounding Healthcare LLMs in a Causal Knowledge Graph: Framework, Metrics, and a Cardiovascular Pilot****ArXiv CS.AI 50%**

arXiv:2608.15382v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly proposed for healthcare decision support, but their evaluations still reward single-answer accuracy rather than reasoning about interventions, mechanisms, harms, evidence, and...

<https://arxiv.org/abs/2608.15382>

2 **CONF****Agentic-SQL Revisited: Autonomy-Based Taxonomy and Empirical Benchmark Analysis for LLM Text-to-SQL****ArXiv CS.AI 50%**

arXiv:2608.15389v1 Announce Type: new Abstract: LLM-based Text-to-SQL progress is reported across heterogeneous benchmarks, backbones, and inference protocols, making cross-system comparison fragile. We reframe the field as a leaderboard aggregation: we collect the metrics...

<https://arxiv.org/abs/2608.15389>

3 **CONF****From Sequence to Structure: Relational Uncertainty Propagation for LLM Agents****ArXiv CS.AI 50%**

arXiv:2608.16002v1 Announce Type: cross Abstract: Reliable uncertainty quantification (UQ) is essential for deploying large language model (LLM) agents in complex interactive environments. Existing UQ methods largely rely on local signals, such as token probabilities, predictive...

<https://arxiv.org/abs/2608.16002>

4 **CONF****Large Language Model Assisted Operational Monitoring for Battery Energy Storage System Integrated Power Distribution Networks****ArXiv CS.AI 50%**

arXiv:2608.15396v1 Announce Type: new Abstract: Battery energy storage systems (BESS) are increasingly used in distribution networks for voltage regulation and demand response, which increases the volume and complexity of operational telemetry available to grid operators. This...

<https://arxiv.org/abs/2608.15396>

5 **CONF****Implementation of a Metacognition Framework for Self-Awareness and Self-Regulation in Ensembles of LLMs****ArXiv CS.AI 50%**

arXiv:2608.15400v1 Announce Type: new Abstract: Large Language Models (LLMs) are notorious for struggling with assessing their own uncertainty, detecting knowledge conflicts, or recognizing when problems exceed their expertise; such limitations inevitably undermine reliability...

<https://arxiv.org/abs/2608.15400>

6 CONF

Distinguishing AI-Generated Music from Edited Audio as a Hard-Negative Robustness Task

ArXiv CS.AI 50%

arXiv:2608.14916v1 Announce Type: cross Abstract: AI-generated music detectors are commonly evaluated against original songs, but real-world uploads are often remixed, re-encoded, pitch-shifted, or otherwise edited. These edited versions form a difficult negative class: they are...

<https://arxiv.org/abs/2608.14916>

7 CONF

Does the Proof Prove It That Way? Faithful Formalization of Elements Proofs

ArXiv CS.AI 50%

arXiv:2608.15432v1 Announce Type: new Abstract: In formal verification, both the autoformalization of statements and automated proof search have been studied extensively. While automated proof search can produce a formal proof that compiles, the generated proof does not...

<https://arxiv.org/abs/2608.15432>

8 CONF

Mental Model Management: An Operator-Based Framework for LLM Memory

ArXiv CS.AI 50%

arXiv:2608.15451v1 Announce Type: new Abstract: Large language models process large amounts of information but usually lack an explicit mechanism for maintaining compact and evolving conceptual representations. We introduce Mental Model Management (3M), a framework in which...

<https://arxiv.org/abs/2608.15451>

9 CONF

Efficient Safety Alignment of Language Models via Latent Personality Traits

ArXiv CS.AI 50%

arXiv:2607.07918v2 Announce Type: replace-cross Abstract: Current safety methods for large language models are known to be vulnerable to adversarial attacks, motivating research into robust alternatives. Latent Adversarial Training (LAT) is among the most effective defenses, but...

<https://arxiv.org/abs/2607.07918>

0 CONF

ATLAS: Scaffold-Free Algorithm Synthesis by LLMs via Embedding-Guided Quality-Diversity Search

ArXiv CS.AI 50%

arXiv:2608.15546v1 Announce Type: new Abstract: Most LLM-based automated algorithm design methods optimize a designated component within a human-specified scaffold, fixing overall organization and component interactions. We present ATLAS, an embedding-guided quality-diversity...

<https://arxiv.org/abs/2608.15546>

1 CONF

Admission Without Answers: Label-Free Certification and Experience Learning for LLM-Based Optimization Modeling

ArXiv CS.AI 50%

arXiv:2608.15565v1 Announce Type: new Abstract: Experience-learning agents for optimization modeling improve by storing verified skills, but existing learners admit knowledge by checking against known answers, which real ticket streams do not provide. The natural label-free...

<https://arxiv.org/abs/2608.15565>

2 CONF

Agent Gym: A Framework for Continuous Evaluation and Evolution of LLM Agents Through Human-in-the-Loop Feedback

ArXiv CS.AI 50%

arXiv:2608.15591v1 Announce Type: new Abstract: Large Language Model (LLM) agents deployed in production environments face a fundamental tension: the agent's behavior is frozen at deployment time, while the business rules and edge cases it must handle continue to evolve....

<https://arxiv.org/abs/2608.15591>

3 CONF

The Granularity Gap: A Multi-Dimensional Cross-Generational Audit of Sycophancy in Gemini Models

ArXiv CS.AI 50%

arXiv:2606.05183v2 Announce Type: replace-cross Abstract: Pass/fail safety evaluation reports whether a model refused. It does not report how far a model went to please the user, and we show these are close to different measurements. We audited sycophancy across three Gemini...

<https://arxiv.org/abs/2606.05183>

4 CONF

Large Models for Small Devices: Recent Advances and Empirical Analysis of Edge AI Deployment

ArXiv CS.AI 50%

arXiv:2608.15693v1 Announce Type: new Abstract: Running large AI models on resource-constrained edge devices requires model compression to reduce model size and computation. What compresses well, however, need not deploy well. We survey dozens of recent works that report...

<https://arxiv.org/abs/2608.15693>

5 CONF

PLeDO: Pain Level Detection for Osteoarthritis from EMR Data

ArXiv CS.AI 50%

arXiv:2608.15719v1 Announce Type: new Abstract: Osteoarthritis (OA) is a progressive chronic joint disease resulting in a breakdown of articular cartilage and bone when damaged joint tissues are not able to normally repair themselves. The aim of this pilot research study is to...

<https://arxiv.org/abs/2608.15719>

6 CONF

Toward AI-Friendly Cartography: Understanding How Color Design Influences Foundation Model Spatial Reasoning on Sequential Choropleth Maps

ArXiv CS.AI 50%

arXiv:2608.15736v1 Announce Type: new Abstract: Foundation models (FMs) increasingly support multimodal and geospatial reasoning, yet it remains unclear whether cartographic principles designed for human perception are equally effective for machines. Focusing on sequential...

<https://arxiv.org/abs/2608.15736>

7 CONF

Propaganda Forensics: Recovering the Generation Pipeline of an AI-Driven Influence Campaign

ArXiv CS.AI 50%

arXiv:2608.15746v1 Announce Type: new Abstract: We present a forensic analysis of the generation pipeline behind a recent AI-driven influence campaign. We introduce PROPAGIA, a corpus of 2,646 propagandist French articles from the Storm-1516/CopyCop campaign disclosed by VIGINUM...

<https://arxiv.org/abs/2608.15746>

8 CONF

Broken Symmetry in LLM Refusal: Answer Release Is More Local Than Refusal Restoration

ArXiv CS.AI 50%

arXiv:2608.15772v1 Announce Type: new Abstract: When a language model refuses to answer a prompt, it is unclear whether the correct answer is erased from its internal representations, or merely suppressed at the output layer. We investigate this mechanism using a controlled...

<https://arxiv.org/abs/2608.15772>

9 CONF

KV-Rescue: Recovering Reasoning Language Model KV Eviction Loss via Stepwise Interleaving

ArXiv CS.AI 50%

arXiv:2608.15797v1 Announce Type: new Abstract: KV-cache eviction caps the memory cost of long reasoning traces but is inherently lossy because the model decodes from a partial view of its history. Under aggressive budgets, this not only lowers accuracy but can also cause...

<https://arxiv.org/abs/2608.15797>

0 CONF

mR²AG: Multimodal Retrieval-Reflection-Augmented Generation for Knowledge-Based VQA

ArXiv CS.AI 50%

arXiv:2411.15041v2 Announce Type: replace Abstract: Advanced Multimodal Large Language Models (MLLMs) struggle with recent Knowledge-based Visual Question Answering (VQA) tasks, such as INFOSEEK and Encyclopedic-VQA, due to their limited and frozen knowledge scope, often leading...

<https://arxiv.org/abs/2411.15041>

1 CONF

Breaking and Defending LLM-Powered Social Media Bot Detection Systems

ArXiv CS.AI 50%

arXiv:2608.15893v1 Announce Type: new Abstract: The rise of social media bots poses a persistent threat, enabling misinformation, opinion manipulation, and the erosion of trust in online platforms. To combat this, machine learning systems have been developed to detect and limit...

<https://arxiv.org/abs/2608.15893>

2 CONF

MUPA²E: Multimodal Unified Perception with Asymmetric Attention for Emotion Assessment

ArXiv CS.AI 50%

arXiv:2608.15999v1 Announce Type: new Abstract: Automatic emotion assessment can benefit from combining neural and behavioral signals, but many multimodal approaches rely on separate, modality-specific feature-extraction pipelines before fusion. This paper presents...

<https://arxiv.org/abs/2608.15999>

3 CONF

Prior Audit-Repair Context Shifts LLM Verifier Thresholds Toward Leniency

ArXiv CS.AI 50%

arXiv:2608.16003v1 Announce Type: new Abstract: Automated checking pipelines increasingly place one language model as the checker and another (or the same one) as the fixer. We ask whether that wiring changes what the checker reports. Measuring false alarms on...

<https://arxiv.org/abs/2608.16003>

4 **CONF****Protein Structure Prediction: From Evolutionary Constraints to Generative Modeling****ArXiv CS.AI** **50%**

arXiv:2608.16094v1 Announce Type: new Abstract: Accurate protein structure prediction is fundamental to structural biology because protein structure underlies molecular function and provides a basis for mechanistic interpretation. Recent advances in deep learning have...

<https://arxiv.org/abs/2608.16094>

5 **CONF****Assessing LLMs' mathematical abilities requires understanding the various mechanisms of mathematical creativity****ArXiv CS.AI** **50%**

arXiv:2608.16118v1 Announce Type: new Abstract: How should we assess whether large language models can perform mathematical invention? I argue that this question is currently underspecified: mathematical creativity is not one capacity but several mechanistically distinct modes...

<https://arxiv.org/abs/2608.16118>

6 **CONF****Identifying Confusion Trends in Concept-based XAI for Multi-Label Classification****ArXiv CS.AI** **50%**

arXiv:2608.15731v1 Announce Type: cross Abstract: Deep Neural Networks (DNNs) deployed in high-risk domains, such as healthcare and autonomous driving, must be not only accurate but also understandable to ensure user trust. In real-world computer vision tasks, these models often...

<https://arxiv.org/abs/2608.15731>

7 **CONF****Trajectory-Level Automatic Curriculum Learning for Legged Locomotion on Unstructured Terrain****ArXiv CS.AI** **50%**

arXiv:2608.16164v1 Announce Type: new Abstract: Training locomotion policies for complex unstructured terrain requires a curriculum to avoid early exploration failures. However, since unstructured terrain lacks explicit difficulty ordering for curriculum design, existing methods...

<https://arxiv.org/abs/2608.16164>

8 **CONF****Competing at Every Price Point with Agentic Evolution over a Menu of LLMs****ArXiv CS.AI** **50%**

arXiv:2608.16207v1 Announce Type: new Abstract: Consider a firm that surveys its competition for a particular agentic task and seeks to offer superior accuracy at every competitor price point. A firm that Pareto-dominated its competitors would leave no rational customer a reason...

<https://arxiv.org/abs/2608.16207>

9 **CONF****Reconstruction: A Blind Benchmark for Recovering Research Ideas from Pre-Publication Bibliographies****ArXiv CS.AI** **50%**

arXiv:2608.16645v1 Announce Type: new Abstract: Can a language model recover the true research idea of a published paper when given only that paper's pre-publication bibliography? We introduce Reconstruction, a blind idea-recovery benchmark that withholds the seed paper and all...

<https://arxiv.org/abs/2608.16645>

00 CONF

DriveCache: Action-Aware Caching for Driving World Model Inference

ArXiv CS.AI 50%

arXiv:2608.16354v1 Announce Type: new Abstract: Driving video generation models support autonomous-driving development by predicting controllable future scenes for simulation, planning evaluation, and offline data generation. Diffusion-based driving generators repeatedly...

<https://arxiv.org/abs/2608.16354>

01 CONF

Think Inside the Chunk: RegulaRAG for Regulation-Compliant Scenario Generation using LLMs: A Case Study of UN Regulation No. 152

ArXiv CS.AI 50%

arXiv:2608.16394v1 Announce Type: new Abstract: Generating regulation-compliant test scenarios is essential for validating safety-critical automotive systems, yet Large Language Models (LLMs) struggle to ground outputs in long, hierarchical standards. We present RegulaRAG, a...

<https://arxiv.org/abs/2608.16394>

02 CONF

Invariant Pretraining for Robust Code Representations

ArXiv CS.AI 50%

arXiv:2608.15412v1 Announce Type: cross Abstract: Encoder-based code representation models remain widely deployed for discriminative tasks such as clone detection and code classification, where their small size and low inference cost are decisive. Their robustness, however, is...

<https://arxiv.org/abs/2608.15412>

03 CONF

The Value of a Prompt: An LLM-Relative Kolmogorov-Complexity Approach

ArXiv CS.AI 50%

arXiv:2608.16438v1 Announce Type: new Abstract: In a world where valuable artifacts are increasingly created, completed, or processed by LLMs, the central economic question is not only what the LLM can produce, but what *\emph{value}* remains in the inputs (i.e., the prompts) we...

<https://arxiv.org/abs/2608.16438>

04 CONF

JailbreakSkill: Scaling Automated Red-Teaming with Reusable and Ever-Evolving Skills

ArXiv CS.AI 50%

arXiv:2608.16465v1 Announce Type: new Abstract: Automated red-teaming has produced a growing collection of attack strategies, yet they typically remain scattered across prompts and workflows, making them difficult to systematically integrate, reuse, and improve at scale. We...

<https://arxiv.org/abs/2608.16465>

05 CONF

Offline Reinforcement Learning for Hemodynamic Management of Sepsis in the ICU: a MIMIC-IV Study with Dual Off-Policy Evaluation

ArXiv CS.AI 50%

arXiv:2608.16482v1 Announce Type: new Abstract: The dosing of intravenous fluids and vasopressors in sepsis is a sequential decision made under uncertainty and guided largely by clinical judgment, which makes it a natural target for reinforcement learning from historical care....

<https://arxiv.org/abs/2608.16482>

06 CONF

DeepInsight II: One Trace from Benchmark to Robot

ArXiv CS.AI 50%

arXiv:2608.16556v1 Announce Type: new Abstract: Across a Physical AI stack, evaluation maturity is inversely aligned with deployment risk: foundation models enjoy mature, standardized harnesses, while the embodied layers on which deployment actually turns remain fragmented...

<https://arxiv.org/abs/2608.16556>

07 CONF

PL-Guard: Probabilistic Logic Reasoning for LLM Guardrails

ArXiv CS.AI 50%

arXiv:2608.15673v1 Announce Type: cross Abstract: Large language model guardrails can be viewed as policy-consistency problems: a system must determine which policy-relevant facts hold in a prompt-response pair and what those facts imply under a given policy. Common approaches,...

<https://arxiv.org/abs/2608.15673>

08 CONF

Physics of Agents: Statistical Mechanics Predicts Collective Behavior of AI Agents

ArXiv CS.AI 50%

arXiv:2608.16578v1 Announce Type: new Abstract: AI agents increasingly operate as part of interacting systems rather than in isolation. As agents exchange information and jointly make decisions, their interactions can improve collective reasoning but may also produce herding,...

<https://arxiv.org/abs/2608.16578>

09 CONF

Cost Scales with Change, Not Corpus Size: Incrementally Maintaining an Evolving Semantic Substrate

ArXiv CS.AI 50%

arXiv:2608.16621v1 Announce Type: new Abstract: Retrieval-augmented and agentic question-answering systems increasingly re-derive the meaning of a corpus at query time. Put plainly, instead of re-deriving what a corpus means on every question, the work is done once when a...

<https://arxiv.org/abs/2608.16621>

10 CONF

NARRATE: A Multimodal Real-World Australian Driving Dataset for Human-Centred Explanations in Automated Driving

ArXiv CS.AI 50%

arXiv:2608.14767v1 Announce Type: cross Abstract: Automated vehicles must explain their decisions in ways that passengers can understand, monitor, and trust. Existing language-annotated driving datasets are mostly observer-written, post-hoc, simulation-based, or generated from...

<https://arxiv.org/abs/2608.14767>

11 CONF

PDDLCode: Agentic PDDL Generation for LLM-Assisted Symbolic Planning

ArXiv CS.AI 50%

arXiv:2608.16637v1 Announce Type: new Abstract: LLMs remain unreliable for long-horizon planning, often generating logically inconsistent or non-applicable plans. Recent hybrid methods instead translate natural language into the Planning Domain Definition Language (PDDL),...

<https://arxiv.org/abs/2608.16637>

12 CONF

PolyComp: A Polycube-based Benchmark for Compositional 3D Spatial Reasoning in Multimodal Models

ArXiv CS.AI 50%

arXiv:2608.14741v1 Announce Type: cross Abstract: We introduce PolyComp, a procedurally generated and verified benchmark that stresses visual recognition and compositional spatial reasoning. In each problem, a model must identify which of four options shows a pair of polycube...

<https://arxiv.org/abs/2608.14741>

13 CONF

When Agents Coordinate: Measuring Coordination in Multi-Agent AI Coding

ArXiv CS.AI 50%

arXiv:2608.16801v1 Announce Type: new Abstract: We study how teams of AI coding agents coordinate while solving programming tasks. Current evaluations usually report whether the agents complete the task and how much the run costs, leaving the coordination inside the team largely...

<https://arxiv.org/abs/2608.16801>

14 CONF

ML-AutoResearch: Training Machine Learning Research Agents with Automatically Generated Environments

ArXiv CS.AI 50%

arXiv:2603.17216v2 Announce Type: replace Abstract: With the advent of AI agents, automated scientific discovery is becoming an increasingly plausible goal. However, training agents to autonomously execute the engineering-heavy labor of machine learning (ML) research requires...

<https://arxiv.org/abs/2603.17216>

15 CONF

Policy Iteration with Human Feedback: Bringing Post-Training RL to In-context Learning

ArXiv CS.AI 50%

arXiv:2608.16831v1 Announce Type: new Abstract: Generative pretraining established reusable task representations; later work on language-based task conditioning and in-context learning showed that a fixed model could adapt its behavior from instructions and demonstrations....

<https://arxiv.org/abs/2608.16831>

16 CONF

A Temporal Reasoning Benchmarking Framework for LRMs via Difficulty-controlled and Dynamic Test Generation

ArXiv CS.AI 50%

arXiv:2607.04784v1 Announce Type: cross Abstract: Defining the reasoning boundaries and ensuring the reliability of Large Reasoning Models (LRMs) remains a critical challenge. Current benchmarks primarily rely on static datasets susceptible to data contamination or synthetic...

<https://arxiv.org/abs/2607.04784>

17 CONF

Orbital AI Computing: Carbon Tradeoffs Across Satellite Scale

ArXiv CS.AI 50%

arXiv:2608.14557v1 Announce Type: cross Abstract: Low Earth Orbit (LEO) computing is emerging for low-latency, globally distributed AI services, enabled by advances in satellite constellations and reusable launch systems. However, its sustainability remains unclear. Prior work...

<https://arxiv.org/abs/2608.14557>

18 CONF

Forward Pass Domain Adaptation (Without Cross-Layer Backpropagation)

ArXiv CS.AI 50%

arXiv:2608.14563v1 Announce Type: cross Abstract: Forward-Pass-Only MLP training (FPO) adapts large language models without a backward pass through the model body, achieving 2.7--3.2x the throughput of standard fine-tuning at ~40% less peak training memory, while leaving...

<https://arxiv.org/abs/2608.14563>

19 CONF

From LLM Inference to Agentic Workloads: Characterization and Implications for Serving Systems

ArXiv CS.AI 50%

arXiv:2608.15127v1 Announce Type: cross Abstract: Agentic applications are shifting AI serving from isolated model inference to long-running workloads in which LLMs coordinate tools, environments, and persistent state. However, the system behavior of these workloads---where...

<https://arxiv.org/abs/2608.15127>

20 CONF

HarmProfile: Characterizing Harmful Distributions in Frontier LLMs

ArXiv CS.AI 50%

arXiv:2608.14577v1 Announce Type: cross Abstract: Frontier large language models (LLMs) safety evaluation has largely treated harmful generation as an attack outcome rather than as an object of analysis. Consequently, little is known about the harmful outputs produced during...

<https://arxiv.org/abs/2608.14577>

21 CONF

Multi-Modal Generative Fuzzy System: Fuzzy Inference Guided Large Model Interactive Question Answering Framework

ArXiv CS.AI 50%

arXiv:2608.14584v1 Announce Type: cross Abstract: In Multimodal Question Answering (MQA), models are required to jointly encode and integrate heterogeneous information from multiple modalities, including text, images, and speech, to perform complex semantic reasoning and...

<https://arxiv.org/abs/2608.14584>

22 CONF

Viola 13M, a Gated Spiral Attention Architecture for Parameter Efficient Small Language Models

ArXiv CS.AI 50%

arXiv:2608.14604v1 Announce Type: cross Abstract: Small language models in the ten to one hundred million parameter range are attractive for on device inference, rapid experimentation, and controlled scientific study, yet most of them reuse the standard transformer block without...

<https://arxiv.org/abs/2608.14604>

23 CONF

Plausible but Not Valid: A Psychometric Audit of LLMs as Synthetic Survey Respondents

ArXiv CS.AI 50%

arXiv:2608.14606v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used as synthetic survey respondents, but existing evaluations ask whether answers look plausible at the individual level. We argue the right question is psychometric: do LLMs...

<https://arxiv.org/abs/2608.14606>

24 CONF

Understanding AI Anxiety in the Workplace: A Multimethod Investigation Using Fear Acquisition Theory and the Technology Acceptance Model

ArXiv CS.AI 50%

arXiv:2608.14609v1 Announce Type: cross Abstract: As artificial intelligence (AI) rapidly diffuses and concerns about job displacement intensify, the psychological mechanisms underlying AI job replacement anxiety remain insufficiently understood. Drawing on Integrated Fear...

<https://arxiv.org/abs/2608.14609>

25 CONF

DumpsterCluster: From Dumpster Diving to Serving LLaMA-70B on \$60 GPUs

ArXiv CS.AI 50%

arXiv:2608.14614v1 Announce Type: cross Abstract: As AI datacenters retire functional GPUs, vast quantities of still capable accelerators enter secondary markets. This paper investigates whether these retired GPUs can find a productive afterlife to form a DumpsterCluster that...

<https://arxiv.org/abs/2608.14614>

26 CONF

Local AI pre-screening for human triple-blind peer review in health sciences

ArXiv CS.AI 50%

arXiv:2608.14625v1 Announce Type: cross Abstract: Academic peer review is under mounting strain: NeurIPS 2025 received 21,575 submissions, ICLR 2025 received 11,603, and ICML 2025 received 12,107. This volume has outpaced the supply of qualified reviewers, and large language...

<https://arxiv.org/abs/2608.14625>

27 CONF

Evidence of conceptual mastery in the application of rules by Large Language Models

ArXiv CS.AI 50%

arXiv:2503.00992v2 Announce Type: replace Abstract: In this paper we leverage psychological methods to investigate LLMs' conceptual mastery in applying rules. We introduce a novel procedure to match the diversity of thought generated by LLMs to that observed in a human sample....

<https://arxiv.org/abs/2503.00992>

28 CONF

Characterizing Rhetorical Misalignment in Decision-Making with Language Models

ArXiv CS.AI 50%

arXiv:2608.14630v1 Announce Type: cross Abstract: Human decision-making is often shaped by a range of well-documented cognitive biases. As large language models (LLMs) become increasingly integrated into high-stakes human-AI decision-making, it is important to understand whether...

<https://arxiv.org/abs/2608.14630>

29 CONF

Valid Per-Field Selective Risk Control for Document Extraction: Three Failure Modes, a Validity Ladder, and When Conditioning Pays

ArXiv CS.AI 50%

arXiv:2608.14639v1 Announce Type: cross Abstract: Per-field accept/review with selective risk at most alpha -- accept a field only if the error rate among accepted fields is controlled -- is the trust contract document-extraction systems need, and the natural procedure silently...

<https://arxiv.org/abs/2608.14639>

30 CONF

Stop Indexing at Full Precision: Revisiting Clustering for Vector Embeddings

ArXiv CS.AI 50%

arXiv:2608.14648v1 Announce Type: cross Abstract: In this study, we revisit three widely used techniques in vector search and utilize them to optimize vector embedding indexing through clustering: dimensionality reduction, quantization, and dimension pruning. We propose an...

<https://arxiv.org/abs/2608.14648>

31 CONF

Do Uncertainty Signals Help? A Systematic Study of Uncertainty-Aware Decoding with Rollback Mechanisms

ArXiv CS.AI 50%

arXiv:2608.14653v1 Announce Type: cross Abstract: Prediction uncertainty is a widely adopted metric for quantifying model confidence, with downstream applications spanning model explanation, data selection, and prediction rollback. Despite its demonstrated utility, the potential...

<https://arxiv.org/abs/2608.14653>

32 CONF

Ring-based Spatial Transformer: Learning Non-linear Spatial Interactions between Building Distribution and Pedestrian Flow

ArXiv CS.AI 50%

arXiv:2608.14660v1 Announce Type: cross Abstract: This study proposes a ring-based SpatialTransformer to learn how building uses at different distances from a railway station interact to generate pedestrian flow. Concentric ring buffers at 100-meter intervals up to 800 meters...

<https://arxiv.org/abs/2608.14660>

33 CONF

Does the Heart Show Your Pain? Tackling the X-ITE Pain Challenge with Self-Supervised ECG Representation Learning

ArXiv CS.AI 50%

arXiv:2608.14662v1 Announce Type: cross Abstract: Accurate recognition of pain using physiological signals remains a challenging problem due to pain's subjective nature and high inter-individual variability. In this study, we investigate self-supervised representation learning...

<https://arxiv.org/abs/2608.14662>

34 CONF

BRA-Audit: Budgeted Runtime Auditing for LLM Multi-Agent Systems via Cumulative-Exposure Audit-Point Placement

ArXiv CS.AI 50%

arXiv:2608.14668v1 Announce Type: cross Abstract: LLM-based multi-agent systems (LLM-MAS) solve complex tasks through specialized collaboration, but inter-agent dependencies can propagate hallucinated or malicious outputs into system-level failures. Auditor agents mitigate these...

<https://arxiv.org/abs/2608.14668>

35 CONF

ARGUS: Attention-Guided Transformers for Scalable Person Identification Using Wi-Fi Telemetry

ArXiv CS.AI 50%

arXiv:2608.14670v1 Announce Type: cross Abstract: Passive, device-free person identification offers an alternative to camera- and wearable-based biometrics, yet existing wireless approaches rely largely on gait or activity cues and are rarely evaluated at scale. In this paper,...

<https://arxiv.org/abs/2608.14670>

36 CONF

Automatic or Controlled? Repetition Priming Reveals Divergent Processing in Base LLMs, Instruct LLMs, and Humans

ArXiv CS.AI 50%

arXiv:2608.14681v1 Announce Type: cross Abstract: Words recur constantly in natural language use, yet it remains unclear whether language models reactivate prior representations or re-evaluate repeated words afresh, and whether post-training changes this default behavior. We...

<https://arxiv.org/abs/2608.14681>

37 CONF

Tail-Aware Top-\$k\$ On-Policy Distillation

ArXiv CS.AI 50%

arXiv:2608.14728v1 Announce Type: cross Abstract: On-policy distillation (OPD) has emerged as an effective paradigm for transferring knowledge between language models, where a student is trained to align its next-token distribution with the teacher's along its own trajectories....

<https://arxiv.org/abs/2608.14728>

38 CONF

DeCo-MIL: Debaised Counterfactual Reasoning for Long-Tailed Whole Slide Image Analysis

ArXiv CS.AI 50%

arXiv:2608.14719v1 Announce Type: cross Abstract: Multiple instance learning (MIL) is widely used for weakly supervised whole slide image (WSI) analysis. However, under long-tailed distributions, MIL-based WSI analysis faces a nested dual long-tail: an inter-slide class long...

<https://arxiv.org/abs/2608.14719>

39 CONF

Multi-Agent Closed-Loop Reasoning for Organic Structure Elucidation from Multimodal Spectra

ArXiv CS.AI 50%

arXiv:2608.14720v1 Announce Type: cross Abstract: Following the molecular discovery and synthesis revolutions, scalable automated structure elucidation from routine spectroscopic data remains an outstanding challenge. Despite decades of computational efforts, no existing system...

<https://arxiv.org/abs/2608.14720>

40 CONF

Unraveling the Size Determination Mechanism of Nanocrystal Synthesis via Interpretable Neural Networks

ArXiv CS.AI 50%

arXiv:2608.14734v1 Announce Type: cross Abstract: Deep learning models of nanocrystal synthesis enable the prediction of size and shape by encoding precursors and reaction conditions. However, their black-box nature hinders gaining deep insights into the underlying synthetic...

<https://arxiv.org/abs/2608.14734>

41 CONF

Synthesizing Post-Acetazolamide Cerebral Blood Flow Maps from Baseline MRI in Moyamoya Using 3D Generative AI

ArXiv CS.AI 50%

arXiv:2608.14758v1 Announce Type: cross Abstract: For patients with Moyamoya disease, impaired cerebrovascular reserve (CVR) is an important hemodynamic criterion for recommending extracranial-to-intracranial bypass surgery. Standard CVR assessment in this cohort uses paired...

<https://arxiv.org/abs/2608.14758>

42 CONF

Cross-Modal Ultrasound-MRI Learning for Fetal Brain Ventricular Volumetry and Abnormality Screening

ArXiv CS.AI 50%

arXiv:2608.14763v1 Announce Type: cross Abstract: Assessment of ventriculomegaly (VM) on fetal brain ultrasound relies primarily on measuring lateral ventricular atrial width on standard planes, which is operator-dependent and may not fully reflect the overall ventricular...

<https://arxiv.org/abs/2608.14763>

43 CONF

Artificial Intelligence as a Tool for Combating Child Labour: A Real-Time Edge Vision Pipeline for Child Detection and Age Estimation

ArXiv CS.AI 50%

arXiv:2608.14770v1 Announce Type: cross Abstract: An estimated 138 million children remain in child labour worldwide, and the monitoring systems used by affected sectors, built on periodic household visits and interviews, systematically under-detect them. We present a real-time...

<https://arxiv.org/abs/2608.14770>

44 CONF

ER-KANS: Efficient and Robust Kolmogorov-Arnold Networks for Data-Scarce Scientific Machine Learning

ArXiv CS.AI 50%

arXiv:2608.14773v1 Announce Type: cross Abstract: The efficient-KAN literature---covering Chebyshev, wavelet, and radial-basis-function variants of the original Kolmogorov-Arnold Network---has been benchmarked almost entirely on clean data. We show that this choice conceals a...

<https://arxiv.org/abs/2608.14773>

45 CONF

Prompting is not enough: supervised baselines and leakage control for measuring shared decision-making with LLMs in pediatric encounters

ArXiv CS.AI 50%

arXiv:2608.14792v1 Announce Type: cross Abstract: Objectives: To determine whether zero-shot prompting of a large language model (LLM) is sufficient to detect shared decision-making (SDM) behaviors in real clinical encounters, and whether supervised learning adds value under...

<https://arxiv.org/abs/2608.14792>

46 CONF

Handover Analysis for Vehicular Communication with Explainability on the Fly

ArXiv CS.AI 50%

arXiv:2608.14820v1 Announce Type: cross Abstract: Handover (HO) management in vehicular networks requires fast and reliable decision-making under highly dynamic conditions. While machine learning (ML) approaches can improve HO detection by capturing complex relationships among...

<https://arxiv.org/abs/2608.14820>

47 CONF

Workspace Topology as an Attack Vector in Agentic Coding Assistants

ArXiv CS.AI 50%

arXiv:2608.14876v1 Announce Type: cross Abstract: Agentic coding assistants are finding widespread use, not just in new code development but in quickly ingesting and leveraging third-party code. This opens up a risk of malicious code being ingested as these coding tools operate...

<https://arxiv.org/abs/2608.14876>

48 CONF

SpIn-ViT: Designing a Sparsity-Induced Vision Transformer That Is Mechanistically Interpretable

ArXiv CS.AI 50%

arXiv:2608.14922v1 Announce Type: cross Abstract: Mechanistic interpretability has recently expanded to Vision Transformers (ViTs), with Sparse Autoencoders (SAEs) increasingly used as post-hoc tools to decompose internal representations into sparse and more interpretable...

<https://arxiv.org/abs/2608.14922>

49 CONF

PaSTel: Anchoring Histology in Spatial Transcriptomics via Multi-Scale Hierarchical Bio-Prior Contrastive Pretraining

ArXiv CS.AI 50%

arXiv:2608.14924v1 Announce Type: cross Abstract: Spatial transcriptomics (ST) links tissue morphology with molecular programs, motivating multimodal pretraining methods that align histology images with gene expression. However, existing approaches suffer from two key...

<https://arxiv.org/abs/2608.14924>

50 CONF

Generative data assimilation highlights fronts as key regulators of ocean energy cascade

ArXiv CS.AI 50%

arXiv:2608.14955v1 Announce Type: cross Abstract: Mesoscale eddies are fundamental to the ocean circulation, yet the extent to which submesoscale motions, a few kilometers across, influence mesoscale eddy energetics through a kinetic energy cascade remains uncertain....

<https://arxiv.org/abs/2608.14955>

51 CONF

Do Geometry-Aware Positional Encodings Help Transformers in Spatial Imperfect-Information Games?

ArXiv CS.AI 50%

arXiv:2608.14982v1 Announce Type: cross Abstract: Transformers applied to spatial imperfect-information games must represent map geometry while tracking hidden entities through time. We ask whether geometry-aware positional encodings improve these capabilities, without claiming...

<https://arxiv.org/abs/2608.14982>

52 CONF

RamseyGadgets: A Graph Construction Dataset for LLMs

ArXiv CS.AI 50%

arXiv:2608.14999v1 Announce Type: cross Abstract: Constructing special graphs is an important task within graph theory and computer science. Many popular graph constructions are the result of a comprehensive exploration of relevant graphs and human ingenuity. Given the rise of...

<https://arxiv.org/abs/2608.14999>

53 CONF

FZ-VLM: A Two Stage Florence-Zephyr Vision Language Model Framework for Pulmonary Nodule Characterization and Clinical Decision Making

ArXiv CS.AI 50%

arXiv:2608.15004v1 Announce Type: cross Abstract: Lung cancer remains one of the leading causes of cancer-related mortality worldwide, and Computed Tomography (CT) is a primary imaging tool for screening and followup assessment. After pulmonary nodule detection, radiologists...

<https://arxiv.org/abs/2608.15004>

54 CONF

MetaReason: Precise Interleaved Multimodal Reasoning via Editing Meta Information for Solving Geometry Problems

ArXiv CS.AI 50%

arXiv:2608.15006v1 Announce Type: cross Abstract: Although visual reasoning is crucial for solving complex geometry tasks, existing vision-language models rely heavily on text-only reasoning. Some recent methods introduce intermediate visual states to facilitate reasoning, but...

<https://arxiv.org/abs/2608.15006>

55 CONF

SysEvolve: An AI-native, safe, autonomous adversarial attack-defense co-evolutionary system

ArXiv CS.AI 50%

arXiv:2608.15012v1 Announce Type: cross Abstract: The rapid advancement of large language models (LLMs) has created a growing asymmetry in cybersecurity, where attack accelerates toward autonomous execution while defense remains predominantly human-intensive. Despite substantial...

<https://arxiv.org/abs/2608.15012>

56 CONF

Hierarchical Agentic Incident Response with Digital-Twin-Validated Attack Inference

ArXiv CS.AI 50%

arXiv:2608.15016v1 Announce Type: cross Abstract: Network incident response remains slow and labor-intensive as the defender must infer multi-stage attacks from partial observations and translate recovery decisions into reliable system commands. Decision-theoretic planners...

<https://arxiv.org/abs/2608.15016>

57 CONF

DualMiT-Net: Local-Global Transformer-Convolutional Fusion for Breast Mass Segmentation in Mammographic Regions of Interest

ArXiv CS.AI 50%

arXiv:2608.15019v1 Announce Type: cross Abstract: Breast mass segmentation is an important step in computer-aided mammography, but it remains difficult because masses can have low contrast, irregular shapes, and boundaries that blend with surrounding breast tissue. To address...

<https://arxiv.org/abs/2608.15019>

58 CONF

Beyond Direct Access: Resource Hijacking in LLM Agents

ArXiv CS.AI 50%

arXiv:2608.15108v1 Announce Type: cross Abstract: Large language model agents are increasingly connected to high-value resources such as computing infrastructure, credentials, usage budgets, identities, private knowledge, communication channels, and organizational workflows....

<https://arxiv.org/abs/2608.15108>

59 CONF

Left-Branching Transformers Excel at Right-Branching Languages: Data Shapes Word Order Preferences in Language Models

ArXiv CS.AI 50%

arXiv:2608.15129v1 Announce Type: cross Abstract: We systematically compare word order preferences in decoder-only language models across 192 artificial languages and typologically diverse natural languages. On artificial languages, models exhibit a left-branching preference...

<https://arxiv.org/abs/2608.15129>

60 CONF

Scale-Consistent Posterior Dynamics for Diffusion Inverse Problems

ArXiv CS.AI 50%

arXiv:2608.15144v1 Announce Type: cross Abstract: Posterior sampling with a pretrained diffusion prior is governed by a conditional score whose intermediate likelihood component is generally intractable. We begin from an ideal one-parameter posterior SDE family in which a...

<https://arxiv.org/abs/2608.15144>

61 CONF

LAPF: LLM-Agent-Based Path Finder Using the UAVScenes Dataset

ArXiv CS.AI 50%

arXiv:2608.15175v1 Announce Type: cross Abstract: Uncrewed aerial vehicles (UAVs) are increasingly deployed for autonomous navigation in complex outdoor environments, where dynamic conditions and mission requirements require intelligent adaptive decision-making. Existing...

<https://arxiv.org/abs/2608.15175>

62 CONF

HiPHI: A Large-Scale Benchmark for High-Precision Human Motion and Object-Interaction

ArXiv CS.AI 50%

arXiv:2608.16222v1 Announce Type: cross Abstract: Humanoid intelligence requires learning over an extremely diverse space of whole-body motions and physically grounded interactions. However, existing embodied datasets remain fundamentally limited: internet-scale video data lack...

<https://arxiv.org/abs/2608.16222>

63 CONF

The Quality of Claude AI-authored Python Tests Is Not Weaker Than Human-authored Tests

ArXiv CS.AI 50%

arXiv:2608.15188v1 Announce Type: cross Abstract: We evaluate the quality of Claude AI-written Python tests against human-written Python tests from two established open-source projects Django and Pandas. Hundreds of tests per corpus are scored under one identical protocol. Using...

<https://arxiv.org/abs/2608.15188>

64 CONF

UAV Video Deblurring via Motion-Aware Diffusion: A Path to Robust Target Detection

ArXiv CS.AI 50%

arXiv:2608.15259v1 Announce Type: cross Abstract: Unmanned Aerial Vehicles (UAVs) play a crucial role in various scenarios ranging from disaster response to traffic surveillance. However, aerial video footage often suffers from severe motion blur due to rapid flight maneuvers,...

<https://arxiv.org/abs/2608.15259>

65 CONF

No Task Fails Every Time: Why One-Shot Audits Are Structurally Blind to Agent Damage

ArXiv CS.AI 50%

arXiv:2608.15286v1 Announce Type: cross Abstract: We introduce AgentRelBench, an environment-agnostic reliability instrument that computes ground-truth, severity-priced damage from database state diffs across repeated runs, with no LLM in the measurement path, demonstrated on...

<https://arxiv.org/abs/2608.15286>

66 CONF

Logical Embeddings for Argument Analysis

ArXiv CS.AI 50%

arXiv:2608.15325v1 Announce Type: cross Abstract: We propose a new framework for machine-learning-oriented argument analysis tasks. Our proposal involves replacing traditional contextualized word embeddings used in most NLP tasks with logical embeddings, an alternative encoding...

<https://arxiv.org/abs/2608.15325>

67 CONF

When AI Rewrites, Classifiers Relax: Uncertainty-Aware Sentiment Analysis on Sarcastic and AI-Paraphrased Social Text

ArXiv CS.AI 50%

arXiv:2608.15338v1 Announce Type: cross Abstract: Sentiment classifiers are increasingly applied to social media content that is either sarcastic or AI-generated --- two distributional regimes where standard evaluations offer little guidance. We present a three-part empirical...

<https://arxiv.org/abs/2608.15338>

68 CONF

Chameleon: An Adaptive AI-Driven Honeypot Architecture Using Threat-Calibrated Particle Swarm Optimization and Semantic Deception Rapidly-Exploring Random Trees

ArXiv CS.AI 50%

arXiv:2608.15407v1 Announce Type: cross Abstract: An invariant behavioral profile is the defining vulnerability of traditional honeypot installations: a skilled adversary can confirm the presence of a deception environment within only a few diagnostic commands, limiting its...

<https://arxiv.org/abs/2608.15407>

69 CONF

FloodReasonBench: Benchmarking VLM Reasoning Segmentation for Embodied Flood Response at the Edge

ArXiv CS.AI 50%

arXiv:2608.15410v1 Announce Type: cross Abstract: Reasoning segmentation enables vision-language models (VLMs) to translate mission-relevant language requests into pixel-level visual grounding, offering a natural perception interface for embodied agents. However, existing...

<https://arxiv.org/abs/2608.15410>

70 CONF

Gated Against One Model, Open to the Next: Option-Only Solvability in Legal Multiple-Choice Benchmarks

ArXiv CS.AI 50%

arXiv:2608.15428v1 Announce Type: cross Abstract: Multiple-choice benchmarks are graded on whether a model picks the right option, not on whether it needed the question. Measuring that gap takes care: a model answering A to most items scores above chance wherever the key sits at...

<https://arxiv.org/abs/2608.15428>

71 CONF

MistyPilot: Enabling Social-Robot Control through Multi-Agent LLM Skill Orchestration

ArXiv CS.AI 50%

arXiv:2608.15549v1 Announce Type: cross Abstract: Programming small social robots from natural-language instructions requires more than invoking isolated APIs. Interactive tasks combine reactive physical behaviors with stateful social behaviors, while existing interfaces often...

<https://arxiv.org/abs/2608.15549>

72 CONF

Kozuchi Agent: A Language-Agnostic Open-Weight Agent for Software Repair

ArXiv CS.AI 50%

arXiv:2608.15579v1 Announce Type: cross Abstract: Industrial software-engineering teams increasingly need LLM agents that turn bug reports into correct patches, yet benchmark-scale operation adds long horizons, tool-use discipline, context persistence, heterogeneous clusters,...

<https://arxiv.org/abs/2608.15579>

73 CONF

FluxBin: Flexible LUT-based Ultra-low-bit LLM Inference by Algorithm-Kernel Synergy

ArXiv CS.AI 50%

arXiv:2608.15602v1 Announce Type: cross Abstract: While binary quantization theoretically promises extreme compression and acceleration for Large Language Models (LLMs), existing research often overlooks the necessity of specialized hardware kernels, thus failing to unleash the...

<https://arxiv.org/abs/2608.15602>

74 CONF

Do Assessment Instruments Measure the Same Thing for Humans and LLMs? A Latent Structure Analysis

ArXiv CS.AI 50%

arXiv:2608.15630v1 Announce Type: cross Abstract: The rapid development and growing deployment of large language models (LLMs) have made it increasingly important to understand their capabilities. A common approach is to evaluate LLMs using assessment instruments originally...

<https://arxiv.org/abs/2608.15630>

75 CONF

Algorithm-Architecture Co-Design for Efficient VLA Inference via Speculative Inference and Verification

ArXiv CS.AI 50%

arXiv:2608.15636v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) models have demonstrated remarkable capabilities in the field of embodied AI, but their high computational cost and limited predicted action length hinder real-time deployment. Although Dadu-Corki, a...

<https://arxiv.org/abs/2608.15636>

76 CONF

When Stories Evolve: Benchmarking LLM Storytelling Across Agent Architectures in Open-Ended World Simulations

ArXiv CS.AI 50%

arXiv:2608.15654v1 Announce Type: cross Abstract: Large language models can write fluent stories, but open-ended storytelling requires more than local fluency. In evolving world simulations and AI-native games, models must preserve facts, relationships, causal dependencies, and...

<https://arxiv.org/abs/2608.15654>

77 CONF

FirstDiff: One-Step Diffusion-Based Anomaly Detection for Multivariate Time Series via Initial Noise Prediction

ArXiv CS.AI 50%

arXiv:2608.15727v1 Announce Type: cross Abstract: Diffusion models have recently shown strong potential for multivariate time-series anomaly detection by learning the distribution of normal data through iterative denoising. Existing diffusion-based approaches, however, typically...

<https://arxiv.org/abs/2608.15727>

78 CONF

CardiacMamba: Fair and Robust RGB-RF Fusion for Remote Heart Rate Estimation via State Space Modeling

ArXiv CS.AI 50%

arXiv:2608.15831v1 Announce Type: cross Abstract: Remote photoplethysmography (rPPG) enables non-contact heart rate (HR) monitoring from facial videos, but RGB-only methods are vulnerable to illumination changes, motion artifacts, and skin-tone-dependent optical reflectance. We...

<https://arxiv.org/abs/2608.15831>

79 CONF

Pre-training Visual Dexterity in Simulation

ArXiv CS.AI 50%

arXiv:2608.15917v1 Announce Type: cross Abstract: Large-scale pre-training has made robot policy fine-tuning increasingly data-efficient, but this progress has largely been driven by datasets and embodiments built around simple parallel-jaw grippers. Dexterous, multi-fingered...

<https://arxiv.org/abs/2608.15917>

80 CONF

Noesis: Bidirectional Graph-RAG with Adaptive Parallelism and Cross-Knowledge-Base Semantic Discovery

ArXiv CS.AI 50%

arXiv:2608.15919v1 Announce Type: cross Abstract: Retrieval-Augmented Generation over knowledge graphs (Graph-RAG) has emerged as a powerful paradigm for grounding large language models in domain-specific corpora. However, existing systems face persistent limitations: (1) static...

<https://arxiv.org/abs/2608.15919>

81 CONF

LLMs Get Smarter from Targeted Synthetic Multilingual Data

ArXiv CS.AI 50%

arXiv:2608.15964v1 Announce Type: cross Abstract: Language-specific competency (LSC) is the phenomenon of a language model performing better or worse depending on the language of the prompt. In other words, a language model outputs different (and potentially incorrect) responses...

<https://arxiv.org/abs/2608.15964>

82 CONF

A Scalable Pipeline for LLM-Teacher Distillation Labeling: Work-Stealing Job Scheduling and Memory-Aware GPU Concurrency

ArXiv CS.AI 50%

arXiv:2608.15975v1 Announce Type: cross Abstract: Labeling large text corpora with LLM teachers has become a practical route to training data at scale. At millions of items, hand-labeling every batch is not feasible, and two questions dominate: what label quality a teacher buys...

<https://arxiv.org/abs/2608.15975>

83 CONF

Reprojection-Guided 3D Gaussian Splatting Diffusion for Weakly Supervised Single-Image Normal Estimation

ArXiv CS.AI 50%

arXiv:2508.05950v4 Announce Type: replace-cross Abstract: We propose CLONE, a Continuous Latent Optimization framework for Normal Estimation via 3D Gaussian splatting. The core idea is to construct an image-geometry-image consistency strategy that unifies explicit geometric...

<https://arxiv.org/abs/2508.05950>

84 CONF

NICE: Scale-Stable Perturbations for Graph Neural Network Explanations via Noise Corruption

ArXiv CS.AI 50%

arXiv:2608.16038v1 Announce Type: cross Abstract: Post-hoc Graph Neural Network (GNN) explainers commonly follow a Perturb-Query paradigm, inferring the importance of graph elements based on queried predictions to perturbed inputs. However, such perturbations often introduce...

<https://arxiv.org/abs/2608.16038>

85 CONF

Decoupling Parcellation from Classification: Systematic Benchmark of Fast Brain Segmentation Methods for Alzheimer's Disease Detection

ArXiv CS.AI 50%

arXiv:2608.16039v1 Announce Type: cross Abstract: Brain parcellation and classification are typically evaluated in isolation, yet downstream AD detection performance depends on their interaction. We decouple these components and systematically benchmark fast deep learning...

<https://arxiv.org/abs/2608.16039>

86 CONF

Behaviour Is an Incomplete Measure of Reasoning Development: Cross-surface pre-arrival accessibility and the limits of developmental inference in a recurrent-depth reasoner

ArXiv CS.AI 50%

arXiv:2608.16085v1 Announce Type: cross Abstract: Capability development is routinely inferred from behavioural thresholds, from final checkpoints, or from what a decoder can read out of a hidden state. These quantities need not identify the same event. We study a 30M-parameter...

<https://arxiv.org/abs/2608.16085>

87 CONF

QUMem: Personalized Memory for Query-Conditioned User-State Inference in LLM Agents

ArXiv CS.AI 50%

arXiv:2608.16168v1 Announce Type: cross Abstract: Large language model (LLM) agents increasingly use external memory systems to support personalization by drawing on long and evolving interaction histories, in which user preferences may be distributed across time, change with...

<https://arxiv.org/abs/2608.16168>

88 CONF

A cross-modal generative model for incomplete and degraded prostate MRI with multicentre clinical validation

ArXiv CS.AI 50%

arXiv:2608.16233v1 Announce Type: cross Abstract: Missing or degraded sequences can limit prostate multiparametric MRI. We developed MSCNet, a sequence-conditioned cross-modal generative framework for reconstructing unavailable contrasts and restoring degraded acquisitions....

<https://arxiv.org/abs/2608.16233>

89 CONF

Software Engineering for AI-driven Building Operation

ArXiv CS.AI 50%

arXiv:2608.16237v1 Announce Type: cross Abstract: Building operations are energy-inefficient. Artificial Intelligence (AI)-driven control systems promise benefits through optimization and predictive control, but deploying them in real buildings reveals a significant software...

<https://arxiv.org/abs/2608.16237>

90 CONF

CompoSkill: Compositional Skill Chain Attacks from Individually Scanner-Passing LLM Agent Skills

ArXiv CS.AI 50%

arXiv:2608.16246v1 Announce Type: cross Abstract: Autonomous AI agents tackling Long Horizon Tasks depend on marketplace skills that are certified one at a time: a scanner returns a safety verdict for each skill and declares the ecosystem safe if every package passes. We show...

<https://arxiv.org/abs/2608.16246>

91 CONF

Foresight-England: Development of a National-Scale Generative AI Model of Electronic Health Records for Medical Event Prediction across the COVID-19 Pandemic

ArXiv CS.AI 50%

arXiv:2608.16273v1 Announce Type: cross Abstract: Foresight-England (Foresight-E) is the first national-scale generative foundation model of electronic health records (EHRs), developed as a research pilot strictly for COVID-19 research. We evaluated its ability to model the...

<https://arxiv.org/abs/2608.16273>

92 CONF

Deep Thought Alignment: Trajectory-Level Latent Distillation for Video Reasoning

ArXiv CS.AI 50%

arXiv:2608.16316v1 Announce Type: cross Abstract: Large Multimodal Models (LMMs) for video reasoning have long been hindered by the high computational cost of processing vast amounts of visual information. This dilemma motivates the transfer of the reasoning capabilities of...

<https://arxiv.org/abs/2608.16316>

93 CONF

Revisiting the Performance of Generative Artificial Intelligence on Introductory Object-Oriented Programming Assessments: Insights from 2026

ArXiv CS.AI 50%

arXiv:2608.16318v1 Announce Type: cross Abstract: Recent advances in Generative Artificial Intelligence (GenAI) have substantially improved the ability of large language models (LLMs) to generate and explain source code. However, their performance on authentic object-oriented...

<https://arxiv.org/abs/2608.16318>

94 CONF

OceanDepths: A Global Dataset of Paired Subsurface and Surface Ocean Observations

ArXiv CS.AI 50%

arXiv:2608.16373v1 Announce Type: cross Abstract: Despite comprising over 70% of its surface, the world's oceans are critically underobserved compared to the land surface or the atmosphere. Understanding the global ocean requires jointly observing its surface and subsurface...

<https://arxiv.org/abs/2608.16373>

95 CONF

Coverage-Maximizing Multinomial Subset Routing under Operational Constraints

ArXiv CS.AI 50%

arXiv:2608.16375v1 Announce Type: cross Abstract: We introduce Multinomial Subset Routing (MSR), a new online routing framework over K experts in which the learner keeps a multinomial routing policy instead of a deterministic subset of experts. At each round, the learner...

<https://arxiv.org/abs/2608.16375>

96 CONF

Synthetic Data Augmentation for Satellite-Based Analysis of Battle-Damaged Agricultural Fields in Ukraine

ArXiv CS.AI 50%

arXiv:2608.16380v1 Announce Type: cross Abstract: Monitoring war-induced damage to agricultural land in Ukraine is important for understanding threats to food security, environmental stability, and post-war recovery. However, the development of computer-vision systems for...

<https://arxiv.org/abs/2608.16380>

97 CONF

Ventor-QTest: Threat-Model-Driven Verification of Vendor-Hosted LLM APIs

ArXiv CS.AI 50%

arXiv:2608.16391v1 Announce Type: cross Abstract: As large language models become increasingly widespread, third-party providers that deploy open-weight models have become an important part of the ecosystem. Auditing the quality of their inference APIs is therefore an open...

<https://arxiv.org/abs/2608.16391>

98 CONF

Towards Risk-free AI Agent Deployment

ArXiv CS.AI 50%

arXiv:2608.16411v1 Announce Type: cross Abstract: LLM-based agents are rapidly moving from research prototypes into the core business processes of organizations, but these agents pose deployment risks to security, compliance, and functionality. In this article, we argue that...

<https://arxiv.org/abs/2608.16411>

99 CONF

Visualizing Uncertainty-to-Action Composition for Human Oversight

ArXiv CS.AI 50%

arXiv:2608.16428v1 Announce Type: cross Abstract: Artificial intelligence systems often disclose uncertainty, yet they rarely make clear what response that uncertainty should trigger. Most uncertainty visualizations encode uncertainty in model outputs, leaving users to discern...

<https://arxiv.org/abs/2608.16428>

00 CONF

Contrastive Energy Fields for Inference-Time Procedure Planning in Instructional Videos

ArXiv CS.AI 50%

arXiv:2608.16457v1 Announce Type: cross Abstract: Procedure planning seeks to estimate a sequence of actions to transition from an observed initial state to a given goal state. Current procedure planning approaches directly predict action sequences from latent representations...

<https://arxiv.org/abs/2608.16457>

01 CONF

A Human-LLM Teaming Framework for Privacy Risk Analysis: An Illustration with CBDC-Based Welfare Schemes

ArXiv CS.AI 50%

arXiv:2608.16461v1 Announce Type: cross Abstract: Central Bank Digital Currency (CBDC)-based welfare schemes may be potentially privacy invasive as they process significant volumes of beneficiary personal data and lead to privacy harms such as surveillance, discrimination and...

<https://arxiv.org/abs/2608.16461>

02 CONF

A Regulatory Placebo? The Systemic Failure of Mandatory GenAI Labeling

ArXiv CS.AI 50%

arXiv:2608.16470v1 Announce Type: cross Abstract: We examine the worldwide trend of mandatory labeling of generative artificial intelligence(GenAI) as a reactive, symbolic form of legislation triggered by technological panic and institutional responses. From a technical...

<https://arxiv.org/abs/2608.16470>

03 CONF

Matched Outcomes, Divergent Gaze: How Foveated MLLMs Search Compared to Humans

ArXiv CS.AI 50%

arXiv:2608.16514v1 Announce Type: cross Abstract: Human visual search is serial: the fovea must land on a candidate to confirm it, and those landings form a scanpath. Whether multimodal large language models (MLLMs), given the same foveated input, search as humans do bears on...

<https://arxiv.org/abs/2608.16514>

04 CONF

Palmyra x6 Technical Report: An Agentic, Tool-Use Model Post-Trained via Anchored Supervised Fine-Tuning

ArXiv CS.AI 50%

arXiv:2608.16620v1 Announce Type: cross Abstract: Palmyra x6 is a large language model optimized for use with enterprise-oriented agentic tasks. The model was built by post-training a Mixture-of-Experts base model with Anchored Supervised Fine-Tuning on a compact corpus of...

<https://arxiv.org/abs/2608.16620>

05 CONF

When Do Explanations Help In-Context Learning? A Comparative Study of Natural Language Explanation Types and Faithfulness

ArXiv CS.AI 50%

arXiv:2608.16627v1 Announce Type: cross Abstract: Natural language explanations (NLEs) are increasingly used as inputs, for example, as few-shot rationales that influence model behavior in in-context learning (ICL). However, it remains unclear how different types of NLEs compare...

<https://arxiv.org/abs/2608.16627>

06 CONF

Toward Better Assessment of LLMs' Performance in Clinical Error Detection

ArXiv CS.AI 50%

arXiv:2608.16643v1 Announce Type: cross Abstract: Automated detection of errors in clinical documentation is a promising application of large language models (LLMs), yet decisions to deploy such models rest on benchmarks that evaluate each clinical note in isolation....

<https://arxiv.org/abs/2608.16643>

07 CONF

X²Localizer: Cross-grained Alignment for Progressive Cross-view Video Geo-localization

ArXiv CS.AI 50%

arXiv:2608.16658v1 Announce Type: cross Abstract: Cross-view Video Geo-localization (CVG) aims to localize ground-view videos by retrieving their corresponding geo-tagged aerial images. However, CVG approaches rely on fixed-length inputs and post-hoc refinement, hindering...

<https://arxiv.org/abs/2608.16658>

08 CONF

Would this change your answer? Evaluating Explanations of LLM Behavior In The Wild with Counterfactual Experiments

ArXiv CS.AI 50%

arXiv:2608.16747v1 Announce Type: cross Abstract: Many areas of AI research, such as language model interpretability and chain of thought faithfulness, seek to explain model behaviors. But what constitutes a "good" explanation? In this work, we evaluate explanations through the...

<https://arxiv.org/abs/2608.16747>

09 CONF

Topological Attribution Distance (TAD): Revealing Segment-Level RAG Influence on LLM Output Geometry for Incident Log Analysis

ArXiv CS.AI 50%

arXiv:2608.16775v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly being deployed in cybersecurity operations to assist cybersecurity analysts with rapid decision-making against emerging threats. However, there is a main criteria that must be met...

<https://arxiv.org/abs/2608.16775>

10 CONF

Model Hypnosis: Strong control of AI via additive subliminal effects

ArXiv CS.AI 50%

arXiv:2608.16834v1 Announce Type: cross Abstract: We demonstrate that AI models are broadly susceptible to a phenomenon we call model hypnosis, in which individually weak and seemingly irrelevant cues in the prompt can be systematically combined to strongly control model...

<https://arxiv.org/abs/2608.16834>

11 CONF

SMA: Who Said That? Auditing Membership Leakage in Semi-Black-box RAG Controlling

ArXiv CS.AI 50%

arXiv:2508.09105v3 Announce Type: replace Abstract: Retrieval-Augmented Generation (RAG) and its Multimodal Retrieval-Augmented Generation (MRAG) significantly improve the knowledge coverage and contextual understanding of Large Language Models (LLMs) by introducing external...

<https://arxiv.org/abs/2508.09105>

12 CONF

Calibrated Generative AI as Meta-Reviewer: A Systemic Functional Linguistics Discourse Analysis of Reviews of Peer Reviews

ArXiv CS.AI 50%

arXiv:2509.15035v2 Announce Type: replace Abstract: This study investigates the use of generative AI to support formative assessment through machine generated reviews of peer reviews in graduate online courses in a public university in the United States. Drawing on Systemic...

<https://arxiv.org/abs/2509.15035>

13 CONF

Specifying AI-SDLC Processes: A Protocol Language for Human-Agent Boundaries

ArXiv CS.AI 50%

arXiv:2606.20615v3 Announce Type: replace Abstract: AI agents now act as first-class members of the software development lifecycle, but the instruments teams use to direct them enforce nothing: process encoded in prompts is flexible but unenforceable, while workflow formalisms...

<https://arxiv.org/abs/2606.20615>

14 CONF

Lesioned Multimodal Language Models Reproduce Aphasic Picture-Naming Patterns

ArXiv CS.AI 50%

arXiv:2607.11621v2 Announce Type: replace Abstract: Aphasia following stroke commonly produces systematic naming errors with characteristic profiles, but whether general-purpose language models not designed for clinical simulation can reproduce these patterns remains untested....

<https://arxiv.org/abs/2607.11621>

15 CONF

Where did the ambiguity go? Examining how multimodal models interpret polysemous words

ArXiv CS.AI 50%

arXiv:2608.00410v2 Announce Type: replace Abstract: Human language is highly polysemous. Many common words (e.g., "bank" or "palm") carry several distinct meanings that shape what humans communicate and imagine. Large language models (LLMs) have been shown to understand this...

<https://arxiv.org/abs/2608.00410>

16 CONF

Decode-Branch Transformers: Decoupling the Primary Prefill Path from Additional Decode Computation

ArXiv CS.AI 50%

arXiv:2608.12385v2 Announce Type: replace Abstract: As large language models serve ever more requests, cumulative inference cost is growing relative to the one-time cost of training. In typical serving, prompt prefill runs in parallel and is compute-bound, whereas autoregressive...

<https://arxiv.org/abs/2608.12385>

17 CONF

From Monte Carlo to neural networks approximations of boundary value problems

ArXiv CS.AI 50%

arXiv:2209.01432v4 Announce Type: replace-cross Abstract: In this paper we study probabilistic and neural network approximations for solutions to Poisson equation subject to Holder data in general bounded domains of $\mathbb{R}^d$. We aim at two fundamental goals. The first, and...

<https://arxiv.org/abs/2209.01432>

18 CONF

MiniGPT-Reverse-Designing: Predicting Image Adjustments Utilizing MiniGPT-4

ArXiv CS.AI 50%

arXiv:2406.00971v3 Announce Type: replace-cross Abstract: Vision-Language Models (VLMs) have recently seen significant advancements through integrating with Large Language Models (LLMs). The VLMs, which process image and text modalities simultaneously, have demonstrated the...

<https://arxiv.org/abs/2406.00971>

19 CONF

Assessing AI-Generated vs. Human-Authored Spear Phishing SMS Attacks: An Empirical Study

ArXiv CS.AI 50%

arXiv:2406.13049v3 Announce Type: replace-cross Abstract: Personalized phishing is difficult to defend against because messages can be tailored to a target's work, interests, and social context. Large language models may make such tailoring faster and easier, but it remains...

<https://arxiv.org/abs/2406.13049>

20 CONF

MoE-Enhanced Explainable Deep Manifold Transformation for Complex Data Embedding and Visualization

ArXiv CS.AI 50%

arXiv:2410.19504v3 Announce Type: replace-cross Abstract: Dimensionality reduction (DR) plays a crucial role in various fields, including data engineering and visualization, by simplifying complex datasets while retaining essential information. However, achieving both high DR...

<https://arxiv.org/abs/2410.19504>

21 CONF

Rethinking Token-wise Feature Caching: Accelerating Diffusion Transformers with Dual Feature Caching

ArXiv CS.AI 50%

arXiv:2412.18911v3 Announce Type: replace-cross Abstract: Diffusion Transformers (DiT) have become the dominant methods in image and video generation yet still suffer substantial computational costs. As an effective approach for DiT acceleration, feature caching methods are...

<https://arxiv.org/abs/2412.18911>

22 CONF

Bringing Generative Learning to Representation Learning: Self-Supervised Transfer Learning as Distribution Matching

ArXiv CS.AI 50%

arXiv:2502.14424v3 Announce Type: replace-cross Abstract: Most self-supervised learning objectives defend against collapse but leave the target representation law unspecified. We formulate representation learning as Distribution Matching (DM), learning an augmentation-invariant...

<https://arxiv.org/abs/2502.14424>

23 CONF

Enhancing the Non-Functional Quality Compliance of LLM-Generated Code through Quality-Aware Preference Learning

ArXiv CS.AI 50%

arXiv:2503.09020v3 Announce Type: replace-cross Abstract: Large Language Models (LLMs) have been widely adopted in commercial code completion engines, significantly enhancing coding efficiency and productivity. However, even functionally correct LLM-generated code may exhibit...

<https://arxiv.org/abs/2503.09020>

24 CONF

Leveraging Machine Unlearning for Cost-Efficient Preference Alignment

ArXiv CS.AI 50%

arXiv:2504.06659v2 Announce Type: replace-cross Abstract: Despite advances in Preference Alignment (PA) for Large Language Models (LLMs), mainstream methods like reinforcement learning with human feedback face notable challenges. These approaches require high-quality datasets of...

<https://arxiv.org/abs/2504.06659>

25 CONF

Bye-bye, Bluebook? Automating Legal Drudgery With AI-Augmented Rule Following

ArXiv CS.AI 50%

arXiv:2505.02763v2 Announce Type: replace-cross Abstract: One of the central promises of legal AI is to automate drudgery -- the formal, repetitive tasks of lawyers' work that consume time without calling for much discretion. Yet it remains an open question how well AI models...

<https://arxiv.org/abs/2505.02763>

26 CONF

WATCH: Adaptive Monitoring for AI Deployments via Weighted-Conformal Martingales

ArXiv CS.AI 50%

arXiv:2505.04608v5 Announce Type: replace-cross Abstract: Responsibly deploying artificial intelligence (AI) / machine learning (ML) systems in high-stakes settings arguably requires not only proof of system reliability, but also continual, post-deployment monitoring to quickly...

<https://arxiv.org/abs/2505.04608>

27 CONF

VirnyFlow: Optimizing ML Pipelines for Accuracy, Fairness, and Stability at Scale

ArXiv CS.AI 50%

arXiv:2506.01584v2 Announce Type: replace-cross Abstract: Developing machine learning (ML) systems for real-world deployment requires navigating context-dependent trade-offs among accuracy, fairness, stability, and other objectives. Existing AutoML frameworks optimize pipelines...

<https://arxiv.org/abs/2506.01584>

28 CONF

LlamaRec-LKG-RAG: A Single-Pass, Learnable Knowledge Graph-RAG Framework for LLM-Based Ranking

ArXiv CS.AI 50%

arXiv:2506.07449v2 Announce Type: replace-cross Abstract: Recent advances in Large Language Models (LLMs) have driven their adoption in recommender systems through Retrieval-Augmented Generation (RAG) frameworks. However, existing RAG approaches predominantly rely on flat,...

<https://arxiv.org/abs/2506.07449>

29 CONF

From Prompts to Constructs: A Dual-Validity Framework for Large Language Model Research in Psychology

ArXiv CS.AI 50%

arXiv:2506.16697v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are entering psychological research both as tools and as objects of inquiry. Yet many studies apply human instruments to LLMs without establishing that the outputs are reliable or...

<https://arxiv.org/abs/2506.16697>

30 CONF

ProteoKnight: Convolution-based Phage Virion Protein Classification and Uncertainty Analysis

ArXiv CS.AI 50%

arXiv:2508.07345v2 Announce Type: replace-cross Abstract: \textbf{Introduction:} Accurate prediction of Phage Virion Proteins (PVP) is essential for genomic studies due to their crucial role as structural elements in bacteriophages. Computational tools, particularly machine...

<https://arxiv.org/abs/2508.07345>

31 CONF

Privacy-Preserving Decentralized Federated Learning via Explainable Adaptive Differential Privacy

ArXiv CS.AI 50%

arXiv:2509.10691v3 Announce Type: replace-cross Abstract: Decentralized federated learning enables collaborative model training without a central server, but shared model updates can still leak sensitive information through inversion, reconstruction, and membership inference...

<https://arxiv.org/abs/2509.10691>

32 CONF

Geometrically Constrained and Token-Based Probabilistic Spatial Transformers

ArXiv CS.AI 50%

arXiv:2509.11218v2 Announce Type: replace-cross Abstract: Spatial transformations such as rotation and scale obscure the morphological cues needed for accurate image classification. Careful consideration is required for reliable use in high stakes settings. A model should stay...

<https://arxiv.org/abs/2509.11218>

33 CONF

A Large-Scale Chinese Knowledge Graph-Text Alignment Dataset for Benchmarking Knowledge-Grounded LLMs

ArXiv CS.AI 50%

arXiv:2510.06039v2 Announce Type: replace-cross Abstract: Reliable evaluation of knowledge-grounded Large Language Models (LLMs) in Chinese requires resources that explicitly align Chinese-language text with verifiable Knowledge Graph (KG) facts. Yet existing Chinese benchmarks...

<https://arxiv.org/abs/2510.06039>

34 CONF

Q-Regularized Generative Auto-Bidding: From Suboptimal Trajectories to Optimal Policies

ArXiv CS.AI 50%

arXiv:2601.02754v3 Announce Type: replace-cross Abstract: With the rapid development of e-commerce, auto-bidding has become a key asset in optimizing advertising performance under diverse advertiser environments. The current approaches focus on reinforcement learning (RL) and...

<https://arxiv.org/abs/2601.02754>

35 CONF

The Fake Friend Dilemma: Relational Trust and the Political Economy of Conversational AI

ArXiv CS.AI 50%

arXiv:2601.03222v2 Announce Type: replace-cross Abstract: As conversational AI systems become a larger part of the media landscape, they raise questions about whose interests they serve and the risks they may pose to users. These systems do more than provide information: they...

<https://arxiv.org/abs/2601.03222>

36 CONF

AWED-PIPER: Agents, Web Applications & Expert Detectors for Personally Identifiable Information Protection & Fine-grained Named Entity Recognition across 36 languages for 6.6 Billion Speakers

ArXiv CS.AI 50%

arXiv:2601.10161v3 Announce Type: replace-cross Abstract: Named Entity Recognition (NER) and Personally Identifiable Information (PII) anonymization are critical tasks in Natural Language Processing (NLP) for information extraction and privacy preservation. We introduce...

<https://arxiv.org/abs/2601.10161>

37 CONF

Sequential LLM Release Facilitates Manipulation in Regulated Markets

ArXiv CS.AI 50%

arXiv:2601.11496v3 Announce Type: replace-cross Abstract: AI agents increasingly mediate bargaining, negotiation and persuasion for people and firms. Such markets extend software-mediated commerce, but add a governance problem: independent model releases change delegates...

<https://arxiv.org/abs/2601.11496>

38 CONF

Credit Fairness: Online Fairness In Shared Resource Pools

ArXiv CS.AI 50%

arXiv:2601.17944v2 Announce Type: replace-cross Abstract: We study repeated allocation of shared resources among agents with time-varying demands and capped linear utilities. In this setting, independently maximizing the minimum endowment-normalized utility in each round...

<https://arxiv.org/abs/2601.17944>

39 CONF

Analytical Provisioning for Attention-FFN Disaggregated LLM Serving under Stochastic Workloads

ArXiv CS.AI 50%

arXiv:2601.21351v4 Announce Type: replace-cross Abstract: Attention-FFN disaggregation (AFD) is an emerging architecture for LLM decoding that separates state-heavy, KV-cache-dominated Attention computation from stateless, compute-intensive FFN computation, connected by per-step...

<https://arxiv.org/abs/2601.21351>

40 CONF

LoRA-CRAFT: Cross-layer Rank Adaptation via Frozen Tucker Decomposition of Pre-trained Attention Weights

ArXiv CS.AI 50%

arXiv:2602.17510v2 Announce Type: replace-cross Abstract: We introduce LoRA-CRAFT (LoRA-CRAFT), abbreviated CRAFT throughout, an extremely parameter-efficient fine-tuning (PEFT) method that applies...

<https://arxiv.org/abs/2602.17510>

41 CONF

Exact Attention Sensitivity and the Geometry of Transformer Stability

ArXiv CS.AI 50%

arXiv:2602.18849v2 Announce Type: replace-cross Abstract: We develop a sensitivity analysis for transformer attention in a geometry aligned with tokenwise computation. Our main result is the exact identity $\|J_{\tau(u)}\|_{\infty} = \theta(p)/\tau$ for the Jacobian $J_{\tau(u)}$...

<https://arxiv.org/abs/2602.18849>

42 CONF

SemVideo: Reconstructs What You Watch from Brain Activity via Hierarchical Semantic Guidance

ArXiv CS.AI 50%

arXiv:2602.21819v3 Announce Type: replace-cross Abstract: Reconstructing dynamic visual experiences from brain activity provides a compelling avenue for exploring the neural mechanisms of human visual perception. While recent progress in fMRI-based image reconstruction has been...

<https://arxiv.org/abs/2602.21819>

43 CONF

Faster, Cheaper, More Accurate: Specialised Knowledge Tracing Models Outperform LLMs

ArXiv CS.AI 50%

arXiv:2603.02830v2 Announce Type: replace-cross Abstract: Predicting future student responses to questions is particularly valuable for educational learning platforms where it enables effective interventions. One of the key approaches to do this has been through the use of...

<https://arxiv.org/abs/2603.02830>

44 CONF

SimulCost: A Cost-Aware Benchmark and Toolkit for Automating Physics Simulations with LLMs

ArXiv CS.AI 50%

arXiv:2603.20253v4 Announce Type: replace-cross Abstract: Evaluating LLM agents for scientific tasks has focused on token costs while ignoring tool-use costs like simulation time and experimental resources. As a result, metrics like pass@k become impractical under realistic...

<https://arxiv.org/abs/2603.20253>

45 CONF

RoboLab: A High-Fidelity Simulation Benchmark for Analysis of Task Generalist Policies

ArXiv CS.AI 50%

arXiv:2604.09860v4 Announce Type: replace-cross Abstract: The pursuit of general-purpose robotics has yielded impressive foundation models, yet simulation-based benchmarking remains a bottleneck due to rapid performance saturation and a lack of true generalization testing....

<https://arxiv.org/abs/2604.09860>

46 CONF

When Do Cheap Probes Predict Expensive Training? Probing 3D-CT Encoders for Text Generation

ArXiv CS.AI 50%

arXiv:2607.22771v2 Announce Type: replace-cross Abstract: Building a 3D CT vision language model begins with a choice of which image encoder to build on. Today that choice is made by fine-tuning every candidate through the full language model and comparing downstream scores, an...

<https://arxiv.org/abs/2607.22771>

47 CONF

EEG Emotion Recognition From AI-Generated Biodigital Architecture Images

ArXiv CS.AI 50%

arXiv:2607.24808v2 Announce Type: replace-cross Abstract: Emotional responses to biodigital architecture were examined using electroencephalographic (EEG) data from AI-generated images. A pre-experiment involving 336 participants identified 60 images, selected from an initial...

<https://arxiv.org/abs/2607.24808>

48 CONF

Agentic AI: User Empowerment or Foreclosure?

ArXiv CS.AI 50%

arXiv:2608.06510v2 Announce Type: replace-cross Abstract: Agentic AI promises systems that can act on users' behalf, from filtering content to negotiating prices to selecting services. Whether it will empower users is an open question, and one that depends on more than the...

<https://arxiv.org/abs/2608.06510>

49 CONF

Effect of Abstractions and Prompting Strategies on LLM-Guided High-Performance Optimizations

ArXiv CS.AI 50%

arXiv:2608.08085v2 Announce Type: replace-cross Abstract: Code performance optimization is a vital aspect of modern software development, as it enables faster response times and reduced resource usage. These optimizations require a deep understanding of low-level hardware...

<https://arxiv.org/abs/2608.08085>