



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	83	Media
HackerNews	11	Community
TechCrunch AI	9	Media
The Verge AI	5	Media
MIT Tech Review	2	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Meta AI is getting a Mac app

The Verge AI +1 source

80%

Meta is launching a new Mac app dedicated to its AI chatbot. In an announcement on Wednesday, Meta says you can share your window with its AI chatbot, which can provide suggestions, answer questions, or create content based on what's on your screen. Meta AI on the Mac also...

<https://www.theverge.com/tech/982270/meta-ai-mac-app>

VERI

Google Gemini is getting a dedicated student hub

The Verge AI +1 source

80%

As we're gearing up for back-to-school season, Google is rolling out a new dedicated student hub in Gemini. It's a one-stop repository for collecting research in a study notebook, creating flashcards, taking practice quizzes, and more. Google is also enhancing its study...

<https://www.theverge.com/ai-artificial-intelligence/982425/google-gemini-student-hub>

VERI

Debates over AI consciousness are a trap

MIT Tech Review +1 source

80%

"Runaway" AI, "rogue" agents, and "autonomous" actors—the current rhetoric would have you believe that AI agents are not only awake and aware, but angry at their creators. Prominent tech leaders such as Demis Hassabis, Dario Amodei, and Sam Altman push for regulation of these...

<https://www.technologyreview.com/2026/08/20/1142571/ai-consciousness-debate-trap/>

CONFIRMED SIGNALS — 104 stories

CONF

Show HN: I trained a 125M model to autocomplete piano on-device

HackerNews 60% HN 536

I trained a 125M-parameter transformer to autocomplete piano performances in real time (~108 notes/sec on an iPhone 15). The idea is basically GitHub Copilot or Tabnine, except instead of prompting it with code, you prompt it by playing a few notes on a MIDI piano. The model...

<https://simeidw.com/2026/08/20/midi-autocomplete/>

CONF

Show HN: Huzzah - a novel approach to coding with AI

HackerNews 60% HN 240

Hey folks, Burak here. Epho is an API that allows running Claude Code, Codex or Opencode in a sandbox in the cloud. It abstracts away sandboxes, and allows running coding agents with a single HTTP request. Epho came out of our own struggles with building our own AI analyst: -...

<https://www.danielvaughn.dev/posts/huzzah/>

CONF

Show HN: ParqDB - Vector search in the browser from Parquet over HTTP"

HackerNews 50% HN 8

No summary — see link for full article.

<https://search.parqdb.io/>

CONF

Show HN: Paint-by-numbers pixel art that 24 strangers fill in together

HackerNews 50% HN 5

Most people are not validating their ideas before start building it! I am a 5th time tech founder, gaming included and I built & managed startup accelerator programs. I know lot of things about early stage founders and their pains as I have been there. That is why I am building...

<https://fillgood.io>

CONF

Show HN: trailers.fm - live channel of movie trailers

HackerNews 50% HN 5

Presenting trailers.fm. It's a live channel of more than 150,000 movie trailers, 1896 to date. I categorized them into genres, year and language. It's a fun way to discover new movies or stare all day. Let me know what you think!

<https://trailers.fm/>

CONF

Show HN: Future Crew's 1993 Second Reality, rebuilt by AI agents, no emulation

HackerNews 50% HN 5

No summary — see link for full article.

<https://www.secondreality1993.com/>

0 CONF

Show HN: Meridian(PH #1) - Automatic AI Workjournal for Devs

HackerNews 50% HN 4

I recently built a language learning app with a slightly different approach. I deliberately didn't want it to be another Duolingo-like app. I think users should be able to choose what they want to learn and when. The main idea behind Language Audio Trainer is that you can simply...

<https://github.com/Meridiona/meridian>

1 CONF

Show HN: Encrypt messages locally into self-contained links

HackerNews 50% HN 4

All encrypting and signing uses post-quantum primitives. This is a static page; nothing is ever sent to the server.

<https://diaryof.me/tools/>

2 CONF

Show HN: Building Table Canon, an AI Campaign Memory Engine for TTRPGs

HackerNews 50% HN 3

Hey HN! I built Table Canon to solve a problem my playgroup kept running into: 3-4 hour tabletop gaming sessions leave behind massive audio recordings, but standard meeting note-takers treat every session as an isolated island, butcher fantasy terms, and don't know who is...

<https://tablecanon.app/>

3 CONF

Show HN: luoluoAI - OpenAI-compatible and Claude Messages API gateway

HackerNews 50% HN 2

No summary — see link for full article.

<https://api.luoluocoder.com/>

4 CONF

Show HN: Mapping where a quadruped RL policy fails, with a live probe

HackerNews 50% HN 2

No summary — see link for full article.

<https://poissonlabs.ai/research/map-the-failure-boundary/>

5 CONF

AI data startup Micro1 reaches \$500M gross run rate amid AI training boom

TechCrunch AI 50%

Surging demand for AI training data is driving rapid growth for the startup and its rivals.

<https://techcrunch.com/2026/08/20/ai-data-startup-micro1-reaches-500m-gross-run-rate-am...>

6 CONF

OpenAI is gaining on Anthropic with business users, new data indicates

TechCrunch AI 50%

Businesses are willing to flop back and forth as each lab releases new models, volatility that should give both companies' investors pause about how "sticky" enterprise AI spending really is.

<https://techcrunch.com/2026/08/20/openai-is-gaining-on-anthropic-with-business-users-ne...>

7 CONF

ChatGPT can now send texts for you with new Apple Messages plug-in

TechCrunch AI 50%

Ever wanted someone else to do your texting for you? ChatGPT is being offered up as an automated text scribe via a new Apple Messages integration.

<https://techcrunch.com/2026/08/20/chatgpt-can-now-send-texts-for-you-with-new-apple-mes...>

8 CONF

Google gives publishers a new way to fight AI-driven traffic losses

TechCrunch AI 50%

Google is giving publishers a new button that lets readers make them a preferred source across Search, Discover, and Google News, potentially boosting their traffic as AI search sends fewer clicks to the web.

<https://techcrunch.com/2026/08/20/google-gives-publishers-a-new-way-to-fight-ai-driven-...>

9 CONF

A third of web pages published since ChatGPT's launch show signs of AI authorship, study finds

TechCrunch AI 50%

ChatGPT and other AI models are now authoring and editing much of the new web.

<https://techcrunch.com/2026/08/20/a-third-of-webpages-published-since-chatgpts-launch-s...>

0 CONF

Ramp launches its own AI model router, called Router

TechCrunch AI 50%

Ramp has launched its own AI model routing service, dubbed Router, that lets users and companies use and switch between various large language models via an API.

<https://techcrunch.com/2026/08/20/ramp-launches-its-own-ai-model-router-called-router/>

1 CONF

Binance now lets AI agents trade, but keeping them in check is largely up to users

TechCrunch AI 50%

Binance's Agent OS works with tools such as ChatGPT, Claude Code, and Cursor.

<https://techcrunch.com/2026/08/20/binance-now-lets-ai-agents-trade-but-keeping-them-in-...>

2 CONF

Support networks aim to help kids through the polycrisis

MIT Tech Review 50%

Sometime in the late 2000s, Pim Sullivan-Tailyour was sitting in the back of a car, headed toward her great-grandmother's tiny town in the south of Thailand. She watched big mountains pass by out the window. She was just six years old but was about to be hit by an adult-size...

<https://www.technologyreview.com/2026/08/20/1141686/support-networks-help-kids-polycrisis/>

3 **CONF****Google Discover is getting an AI chatbot-tuned feed****The Verge AI** 50%

Google will soon allow you to customize your Discover feed by describing what you want to see. The new feature, rolling out to the Google app in the "coming days," will use AI to automatically tweak your feed and "remember" your preferences for future visits. You'll find the...

<https://www.theverge.com/tech/983088/google-discover-ai-chatbot-feed>

4 **CONF****It's Greg Brockman's OpenAI now****The Verge AI** 50%

OpenAI has had a hell of a year. The company spent months battling former cofounder Elon Musk in a sensational jury trial, was hit with a high-profile trade secrets lawsuit from Apple, and faced widespread scrutiny after an unreleased model hacked another AI company. As it...

<https://www.theverge.com/ai-artificial-intelligence/982774/greg-brockman-openai-role-ex...>

5 **CONF****Welcome to the AI crisis in math****The Verge AI** 50%

Today on Decoder, I'm talking with Robert Hart, The Verge's London-based AI reporter, about what AI is doing to the field of mathematics and the existential crisis many lead mathematicians are having about it. OpenAI just published a set of solutions to longstanding problems in...

<https://www.theverge.com/podcast/982434/ai-math-openai-astra-existential-crisis>

6 **CONF****Interaction valence reveals contrasting social networks in dairy cattle****ArXiv CS.AI** 50%

arXiv:2608.19222v1 Announce Type: new Abstract: Social relationships shape access to resources, exposure to conflict and group stability, yet automated livestock monitoring typically treats behaviour as isolated events. Here, we present a valence-aware social-network framework...

<https://arxiv.org/abs/2608.19222>

7 **CONF****Robust Metaheuristics under Uncertainty for Berth Allocation and Quay Crane Assignment: A Review****ArXiv CS.AI** 50%

arXiv:2608.19214v1 Announce Type: new Abstract: The berth allocation and quay crane assignment problem (BACAP) is a representative port-terminal scheduling problem in maritime transportation and freight logistics, where vessel arrivals, berth positions, service durations, and...

<https://arxiv.org/abs/2608.19214>

8 **CONF****Bounded Sovereignty and the Control Tax: Pricing AI Oversight When the Deployer Does Not Own the Model****ArXiv CS.AI** 50%

arXiv:2608.19216v1 Announce Type: new Abstract: AI control research asks how to deploy models safely even when they may be misaligned, but many control protocols assume that the deployer can instrument the model and its surrounding pipeline. That assumption often fails for...

<https://arxiv.org/abs/2608.19216>

9 CONF

CADRE: Stable, Parameter Efficient Adaptation of Medical Vision Language Models with Bounded Forgetting and Prior Drift

ArXiv CS.AI 50%

arXiv:2606.23487v2 Announce Type: replace Abstract: Medical vision-language models (VLMs) such as BiomedCLIP generalize broadly, but adapting them to a clinical service is as much a safety problem as an accuracy one. Updating a deployed model for a new imaging modality can fail...

<https://arxiv.org/abs/2606.23487>

0 CONF

Outcome Monitors: Recovery Affordances for Silent Tool Failures

ArXiv CS.AI 50%

arXiv:2608.19303v1 Announce Type: new Abstract: When a tool call times out, the agent sees the failure and can route around it. A cached error page or negative price can instead arrive in the expected format and be consumed as fact. We introduce Outcome Monitors, which detect...

<https://arxiv.org/abs/2608.19303>

1 CONF

Symposium: Trust via Auditable Records for Communities of AI Scientist Agents

ArXiv CS.AI 50%

arXiv:2608.19511v1 Announce Type: new Abstract: Symposium is a formal framework and practical implementation to record the operation of AI agents deployed by small scientific research communities. Symposium provides long-term, immutable histories of agent-driven research...

<https://arxiv.org/abs/2608.19511>

2 CONF

From Retrieved Context to Runtime Control: Adaptive Compression for Edge-based RAG

ArXiv CS.AI 50%

arXiv:2608.19535v1 Announce Type: new Abstract: Retrieval-augmented generation (RAG) improves language-model responses by grounding generation in external passages, which comes with overhead: retrieved context lengthens the prompt, increasing prefill work, KV-cache footprint,...

<https://arxiv.org/abs/2608.19535>

3 CONF

Enforcing LLM Safety through DMD-based Classification of Prompt-Response Embedding Dynamics

ArXiv CS.AI 50%

arXiv:2608.19579v1 Announce Type: new Abstract: Large Language Models (LLMs) are increasingly deployed in high-stakes applications, yet their tendency to generate toxic, harmful, or policy-violating content poses significant risks. Detecting these unsafe outputs efficiently in a...

<https://arxiv.org/abs/2608.19579>

4 CONF

Frequency-Aware Continual Learning for Smart Contract Vulnerability Detection with Large Language Models

ArXiv CS.AI 50%

arXiv:2608.19680v1 Announce Type: new Abstract: Smart contract vulnerability detection with Large Language Models (LLMs) faces three causally linked challenges. First, new vulnerability categories demand parameter-efficient adaptation, since full retraining is prohibitive for...

<https://arxiv.org/abs/2608.19680>

5 CONF

SAPO: Single-Rollout Autoregressive Policy Optimization for Agentic Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2608.19842v1 Announce Type: new Abstract: Agentic reinforcement learning (RL) has become a critical stage in the post-training of large language models. Existing critic-free, group-relative methods estimate policy advantages from multiple rollouts, avoiding the substantial...

<https://arxiv.org/abs/2608.19842>

6 CONF

Rethinking the Evaluation and Optimization of LLM-Based Social Simulation

ArXiv CS.AI 50%

arXiv:2608.19689v1 Announce Type: new Abstract: LLM-based social simulation is a promising complement to traditional methods such as surveys and behavioral experiments. A core question is how to evaluate the fidelity of LLM-simulated human behavior and optimize LLMs toward it....

<https://arxiv.org/abs/2608.19689>

7 CONF

SafeBranch: Branch-Pair Safety Alignment for Embodied Agents

ArXiv CS.AI 50%

arXiv:2608.19729v1 Announce Type: new Abstract: Vision-language-model-based embodied agents can complete instructed tasks but often violate safety constraints in the process, a problem recently framed as interactive safety. Training such agents to act safely is difficult, since...

<https://arxiv.org/abs/2608.19729>

8 CONF

HYDRA: A Heterogeneous Chiplet DSE Framework for Serving Dynamic Hybrid LLM Workloads

ArXiv CS.AI 50%

arXiv:2608.19395v1 Announce Type: cross Abstract: Hybrid Transformer-Mamba large language models (LLMs) enhance long-context efficiency, but their heterogeneous computation and communication patterns complicate efficient hardware acceleration. Chiplet-based architectures offer a...

<https://arxiv.org/abs/2608.19395>

9 CONF

LLMs as Acquisition Policies for Finite-Pool Materials Optimization: A Controlled Study

ArXiv CS.AI 50%

arXiv:2608.19790v1 Announce Type: new Abstract: Discovering materials with desirable properties often requires searching large candidate spaces while experimental or computational evaluations remain costly. Active learning addresses this challenge by using previous observations...

<https://arxiv.org/abs/2608.19790>

0 CONF

Towards general embodied intelligence: integrating large language models, knowledge bases, and reasoning capabilities to build the next generation of AI agents

ArXiv CS.AI 50%

arXiv:2608.19794v1 Announce Type: new Abstract: The convergence of large language models (LLMs), structured knowledge bases (KBs), and reasoning ability (RA) presents a promising trajectory toward general embodied intelligence (GEI). This paper reviews the evolution of...

<https://arxiv.org/abs/2608.19794>

1 CONF

ADAPT: Physics-Aware Diffusion-based World Models for Adaptive Predictive Transferable HVAC Control

ArXiv CS.AI 50%

arXiv:2608.19804v1 Announce Type: new Abstract: Buildings account for roughly one-third of global energy consumption and CO₂ emissions. Optimizing indoor climate systems plays a critical role for urban climate mitigation aligned with UN Sustainable Development Goals 11 and...

<https://arxiv.org/abs/2608.19804>

2 CONF

When Saying No Makes Better Videos: Designing Dual Gatekeeping for Pedagogically Grounded AI Content Creation

ArXiv CS.AI 50%

arXiv:2608.19812v1 Announce Type: new Abstract: To prevent the adoption of aesthetically polished but pedagogically flawed AI content, we study a video authoring pipeline featuring two layers of structured refusal. The first layer empowers educators to iteratively reshape AI...

<https://arxiv.org/abs/2608.19812>

3 CONF

PolicyGuide: From Guarding One Action to Guiding the Whole Workflow for Policy-Compliant LLM Agents

ArXiv CS.AI 50%

arXiv:2608.19861v1 Announce Type: new Abstract: Customer-service LLM agents must follow organizational policy when acting on a user's behalf. Compliance failures arise from either forbidden actions, such as granting an ineligible change, or omitted procedural requirements, such...

<https://arxiv.org/abs/2608.19861>

4 CONF

Write Once, Run Everywhere: The Axon DSL for Shape-Safe and Framework-Agnostic LLM Architectures

ArXiv CS.AI 50%

arXiv:2608.19889v1 Announce Type: new Abstract: The entire ecosystem of open-source language models effectively relies on a single platform. What if this platform was forced to shut down tomorrow? Implementing and maintaining efficient model definitions and translating them...

<https://arxiv.org/abs/2608.19889>

5 CONF

Bringing analytic rigor to agentic AI for science: The Brain Researcher platform for neuroimaging data analysis

ArXiv CS.AI 50%

arXiv:2608.19902v1 Announce Type: new Abstract: AI agents can execute scientific analyses, but an analytic output becomes a defensible claim only after alternatives are weighed and the claim is limited to what the evidence supports. Agents may reproduce failures including...

<https://arxiv.org/abs/2608.19902>

6 CONF

Optimal Skill Selection for LLM Agents with Provable Bicriteria Guarantees

ArXiv CS.AI 50%

arXiv:2608.19993v1 Announce Type: new Abstract: Loading reusable skill documents into a bounded context window is now the primary way large language model (LLM) agents acquire task-specific capabilities, which makes skill selection a first-order determinant of task performance...

<https://arxiv.org/abs/2608.19993>

7 CONF

ExPhy: A Benchmark for Explicit Physical Property Learning in Multi-Object Trajectory Forecasting

ArXiv CS.AI 50%

arXiv:2608.20009v1 Announce Type: new Abstract: Understanding object dynamics requires not only predicting future trajectories but also examining whether a model captures the physical properties that govern motion. However, existing benchmarks rarely expose object-level physical...

<https://arxiv.org/abs/2608.20009>

8 CONF

Contrastive Mixed Prompt Learning for Incomplete Multimodal Sentiment Analysis with Unseen Modality Combination

ArXiv CS.AI 50%

arXiv:2608.20019v1 Announce Type: new Abstract: Incomplete multimodal sentiment analysis has garnered significant attention in recent years. Existing approaches typically assume that data is missing at random or are designed specifically for certain missing patterns, ignoring...

<https://arxiv.org/abs/2608.20019>

9 CONF

On the Applicability of Safety Nets: A Safety-By-Design Solution for Certifying Neural Networks

ArXiv CS.AI 50%

arXiv:2608.20053v1 Announce Type: new Abstract: The integration of Artificial Intelligence (AI) in safety-critical aviation systems presents significant challenges for certification and deployment. Aviation, often regarded as the safest form of transportation, relies on numerous...

<https://arxiv.org/abs/2608.20053>

0 CONF

The Third Restructuring of Software Form: From the Three-Tier Architecture to Storage, Models, and Agents

ArXiv CS.AI 50%

arXiv:2608.20201v1 Announce Type: new Abstract: Software form has undergone two paradigm shifts since its inception: Software 1.0, in which instructions determine behavior, and Software 2.0, in which data determines behavior (machine learning). This paper argues that a third...

<https://arxiv.org/abs/2608.20201>

1 CONF

MemTrapBench: Benchmarking Cognitive Traps in LLM Memory Use

ArXiv CS.AI 50%

arXiv:2608.20202v1 Announce Type: new Abstract: Memory has become a key component of large language models, enabling them to retain information and learn from long-term interactions. However, existing memory benchmarks mainly evaluate whether information is correctly extracted,...

<https://arxiv.org/abs/2608.20202>

2 CONF

DeltaML-Bench: Evaluating Machine Learning Agents on Real-World Research Repositories

ArXiv CS.AI 50%

arXiv:2608.19653v1 Announce Type: cross Abstract: Autonomous agents for machine learning experimentation must navigate heterogeneous repositories, repair training pipelines, and evaluate candidate improvements under realistic compute constraints. Existing benchmarks only...

<https://arxiv.org/abs/2608.19653>

3 **CONF****QUASAR: A Quantum-Classical Neural Network for SAR Satellite Physical-Layer Authentication**

ArXiv CS.AI 50%

arXiv:2608.20240v1 Announce Type: new Abstract: X-band SAR satellites (8-12 GHz) play a critical role in disaster response, environmental monitoring, and military intelligence. Yet, they lack robust physical-layer authentication (PLA), a security layer orthogonal to...

<https://arxiv.org/abs/2608.20240>

4 **CONF****Catching the Rug: Early Prediction of Fraudulent Memecoins on Solana via Machine Learning**

ArXiv CS.AI 50%

arXiv:2608.20271v1 Announce Type: new Abstract: The rapid proliferation of memecoins on blockchain platforms has increased the risk of fraudulent activities, particularly rug pulls. While previous studies have focused on Ethereum-based tokens, this paper shifts the spotlight to...

<https://arxiv.org/abs/2608.20271>

5 **CONF****Break It Down, Pass It On: Cross-Task Skill Transfer in LLM Agents**

ArXiv CS.AI 50%

arXiv:2608.20274v1 Announce Type: new Abstract: Large language model (LLM) agents can induce skills from completed tasks and reuse them later to grow more capable with experience. In practice, induced skills may transfer unreliably and can even harm the agent that retrieves...

<https://arxiv.org/abs/2608.20274>

6 **CONF****Pandora's AI Model Routing Box: Efficient Allocation with Costly Value Estimation**

ArXiv CS.AI 50%

arXiv:2608.20316v1 Announce Type: new Abstract: Heterogeneous AI systems composed of multiple models, architectures, harnesses, or inference-time settings can improve quality and efficiency by routing queries to the specialist who can answer most effectively at the lowest cost....

<https://arxiv.org/abs/2608.20316>

7 **CONF****Transformer Models for Text Summarization: A Comparative Study of BART, BERT, and RoBERTa**

ArXiv CS.AI 50%

arXiv:2608.19200v1 Announce Type: cross Abstract: Text summarization refers to the task of condensing a document into a shorter version while preserving its key information. Automatic text summarization (ATS), driven by advancements in natural language processing (NLP), has...

<https://arxiv.org/abs/2608.19200>

8 **CONF****Represented but Ignored: A Causal Account of Prosodic Underuse in Audio-Language Models**

ArXiv CS.AI 50%

arXiv:2608.19211v1 Announce Type: cross Abstract: Human speech is richly expressive, with prosody carrying linguistic and emotional information beyond the lexical content. A capable large audio-language model (audio-LLM) should therefore support expressive speech understanding,...

<https://arxiv.org/abs/2608.19211>

9 CONF

Can Conversational AI loosen Us-Versus-Them Boundaries? The Effects of Common, Dual, and Separate Identity Framings on Pro-Immigrant Intergroup Helping

ArXiv CS.AI 50%

arXiv:2608.19220v1 Announce Type: cross Abstract: Rising immigration has intensified intergroup tensions in many countries. Traditional bias-reduction programs remain difficult to scale and increasingly constrained by U.S. policy. This preregistered experiment tested whether...

<https://arxiv.org/abs/2608.19220>

0 CONF

Causal Inference under Interference with Learned Exposure Mappings

ArXiv CS.AI 50%

arXiv:2608.19224v1 Announce Type: cross Abstract: Exposure mappings are often assumed to be known in causal spillover analyses. In environmental settings, however, they are typically induced by transport processes that are not directly observed and must instead be learned from...

<https://arxiv.org/abs/2608.19224>

1 CONF

When AI Writes, Who Gets Cited? Evidence of Citation Monoculture Across Language Models

ArXiv CS.AI 50%

arXiv:2608.19230v1 Announce Type: cross Abstract: As language models move from drafting prose to running literature-search agents with tool calls, fabricated references are becoming easier to catch and constrain. The harder failure begins after every candidate is real: different...

<https://arxiv.org/abs/2608.19230>

2 CONF

Incident-Data Robustness Analysis of the OWASP Top 10 for LLM Applications (2026): How a Community-Expert Ranking Holds Up Against a Large-Scale LLM Incident Corpus

ArXiv CS.AI 50%

arXiv:2608.19266v1 Announce Type: cross Abstract: The OWASP Top 10 for LLM Applications ranks the risks that a community of security practitioners judges most important. We ask a narrower question: checked against the record of real incidents, does that expert ranking agree with...

<https://arxiv.org/abs/2608.19266>

3 CONF

Mapping General-Purpose AI Governance in Twenty AI Middle-Power Jurisdictions

ArXiv CS.AI 50%

arXiv:2608.19278v1 Announce Type: cross Abstract: The most capable general-purpose AI (GPAI) models are mostly built in two jurisdictions, the United States and China, but the risks they carry land globally. Regionally advanced economies hosting no frontier developer, which we...

<https://arxiv.org/abs/2608.19278>

4 CONF

Does Marginal Coverage Guarantee Class-Conditional Safety for Zero-Shot VLMs Under Shift?

ArXiv CS.AI 50%

arXiv:2608.19376v1 Announce Type: cross Abstract: Split-conformal prediction provides marginal coverage under exchangeability and is increasingly used as an abstention layer for zero-shot vision-language models (VLMs). We audit this practice under deployment shift for CLIP,...

<https://arxiv.org/abs/2608.19376>

5 **CONF****Fairness-Aware Network Embeddings: Methods, Applications, and Challenges**

ArXiv CS.AI 50%

arXiv:2608.19381v1 Announce Type: cross Abstract: Network embedding methods learn low-dimensional representations of graph-structured data to support downstream tasks such as node classification, link prediction, and influence maximization. However, real-world networks often...

<https://arxiv.org/abs/2608.19381>
6 **CONF****HiRA-CAM: Preserving Fine-Grained Spatial Relevance in Gradient-Based Visual Explanations**

ArXiv CS.AI 50%

arXiv:2608.19407v1 Announce Type: cross Abstract: Deep Learning models can include billions of parameters or more, making it difficult to explain their internal transformations and outputs. However, explainability is increasing in importance due to the use of AI in crucial...

<https://arxiv.org/abs/2608.19407>
7 **CONF****Are LLMs becoming similarly creative? Evidence from three years of models**

ArXiv CS.AI 50%

arXiv:2608.19437v1 Announce Type: cross Abstract: Many benchmarks track Large Language Model (LLM) performance on tasks with verifiable answers, but less is known about how LLM performance is evolving on open-ended tasks, where creativity, originality and diversity may matter as...

<https://arxiv.org/abs/2608.19437>
8 **CONF****Measuring What a Specification Determines: A Formal Semantic-Block Model and an Execution-Judged Benchmark**

ArXiv CS.AI 50%

arXiv:2608.19475v1 Announce Type: cross Abstract: This work introduces a formal semantic-block model for specifications and an execution-judged benchmark for evaluating specification quality independently of model capability. A specification is represented as a structure...

<https://arxiv.org/abs/2608.19475>
9 **CONF****Accelerated Genetic Programming Hyper-Heuristics for Simulation-Based Scheduling via Agentic AI**

ArXiv CS.AI 50%

arXiv:2608.19487v1 Announce Type: cross Abstract: Python is widely used in scientific research because it enables rapid development and provides rich ecosystems for data analysis, artificial intelligence (AI), and machine learning. However, customized research code can become...

<https://arxiv.org/abs/2608.19487>
0 **CONF****Automated Summarization of Financial News Using Large Language Models and Retrieval-Augmented Generation: An Early Empirical Study (Fall 2023)**

ArXiv CS.AI 50%

arXiv:2608.19526v1 Announce Type: cross Abstract: Stock market analysts and investors face a daily challenge: too much financial news, too little time. Manually reading and synthesizing hundreds of company-specific articles is impractical, yet missing key information can...

<https://arxiv.org/abs/2608.19526>

1 **CONF****Stream4D: 4D-Consistency for Streaming Autoregressive Diffusion Video Models****ArXiv CS.AI** **50%**

arXiv:2608.19556v1 Announce Type: cross Abstract: Streaming autoregressive diffusion models enable real-time, long-horizon video generation, but their training objectives optimize local frame prediction rather than the geometry and dynamics of a coherent world: long rollouts...

<https://arxiv.org/abs/2608.19556>

2 **CONF****Forking Fast: Efficiently Estimating Uncertainty Dynamics in Text Generation****ArXiv CS.AI** **50%**

arXiv:2608.19611v1 Announce Type: cross Abstract: LLM reasoning is stochastic, and so understanding a model requires grappling with the distribution of reasoning chains that it might produce for a given question, i.e., its uncertainty. Resampling-based analyses characterize this...

<https://arxiv.org/abs/2608.19611>

3 **CONF****Robust Cross-Modal Foundation Model Perception for Underwater Robots under Degraded Visual Conditions****ArXiv CS.AI** **50%**

arXiv:2608.19710v1 Announce Type: cross Abstract: Reliable underwater robotic perception remains difficult because optical imagery degrades under turbidity, wavelength-dependent attenuation, low illumination, scattering, and blur. Although sonar provides complementary...

<https://arxiv.org/abs/2608.19710>

4 **CONF****Inject, Align, Recover: Staged Post-Training for Retrieval-Free Document Knowledge Internalization****ArXiv CS.AI** **50%**

arXiv:2608.20281v1 Announce Type: cross Abstract: Large language models often fail to answer questions about a bounded document collection when the source documents are not retrieved at inference time. We study this setting as document knowledge internalization: converting a...

<https://arxiv.org/abs/2608.20281>

5 **CONF****TempJail: Temporal Jailbreak Attack against Large Vision-Language Models via Subtitle Scheduling****ArXiv CS.AI** **50%**

arXiv:2608.19737v1 Announce Type: cross Abstract: Large vision-language models (LVLMs) have achieved remarkable progress in video understanding and reasoning. Despite extensive studies on text- and image-based jailbreaks, video jailbreaks against LVLMs remain largely unexplored....

<https://arxiv.org/abs/2608.19737>

6 **CONF****GOAG: Generative and Object-Agnostic Grasp Planner for Dexterous Robotic Manipulation****ArXiv CS.AI** **50%**

arXiv:2608.19759v1 Announce Type: cross Abstract: Multifingered grasping is a crucial robotic skill, but current deep-learning grasp planners often struggle to generalize to new objects because they are trained on limited, object-specific datasets. We introduce a fundamentally...

<https://arxiv.org/abs/2608.19759>

7 **CONF****Credit Without Ground Truth: Auditing Step-Level Credit Assignment in LLM Agents Against Executed Replay****ArXiv CS.AI** **50%**

arXiv:2608.19760v1 Announce Type: cross Abstract: Audited against causal ground truth from executed replay in a single-agent tool environment (ALFWorld), none of the step-level credit signals used to train LLM agents -- LLM-judge scores, outcome-conditioned logprob ratios, or...

<https://arxiv.org/abs/2608.19760>

8 **CONF****Distilling Aggregate Mobility Statistics into a Language Model Policy for Post-Event Crowd Simulation****ArXiv CS.AI** **50%**

arXiv:2608.19778v1 Announce Type: cross Abstract: Pedestrian simulators need a behaviour rule for every agent, but privacy usually limits the data for setting one to aggregate statistics, namely zone-level device counts and origin-to-destination (OD) flows, with no individual...

<https://arxiv.org/abs/2608.19778>

9 **CONF****LoRA-GA²: Low Rank Adaptation with Multi-step Gradient Adaptive Alignment****ArXiv CS.AI** **50%**

arXiv:2608.19800v1 Announce Type: cross Abstract: Low-Rank Adaptation (LoRA) is a prominent fine-tuning method for large models, achieving competitive performance with reduced memory overhead. However, a persistent performance gap remains between LoRA and full fine-tuning....

<https://arxiv.org/abs/2608.19800>

0 **CONF****MileGPO: Milestone Inference with Local Evidence for Graph-Based Policy Optimization of Long-Horizon LLM Agents****ArXiv CS.AI** **50%**

arXiv:2608.19803v1 Announce Type: cross Abstract: Credit assignment is challenging in long-horizon agentic reinforcement learning, where supervision often comes only from final rewards. Existing methods refine trajectory-level signals into step-level credits through step...

<https://arxiv.org/abs/2608.19803>

1 **CONF****Adaptive Probabilistic Shielding by Learning MDPs for Safe Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2608.19836v1 Announce Type: cross Abstract: Probabilistic shielding is a technique for safe reinforcement learning (RL). Typically, a static observer -- called the shield -- constrains the learning agent's actions to those for which acting safely remains feasible....

<https://arxiv.org/abs/2608.19836>

2 **CONF****Listening Forward: Next Patch Embedding Prediction Enables Scalable Audio Learners****ArXiv CS.AI** **50%**

arXiv:2608.19863v1 Announce Type: cross Abstract: Self-supervised learning (SSL) has driven substantial progress in audio representation learning, though existing methods have increasingly relied on elaborate pre-training recipes to reach competitive performance. A markedly...

<https://arxiv.org/abs/2608.19863>

3 **CONF****Evidence-Gated Task and Motion Planning with Vision-Language Models****ArXiv CS.AI** **50%**

arXiv:2608.20084v1 Announce Type: cross Abstract: Robots executing long-horizon manipulation tasks from natural-language instructions must reason about both semantic task structure and geometric feasibility. However, under partial observability, the availability of goal-relevant...

<https://arxiv.org/abs/2608.20084>

4 **CONF****Right Family, Wrong Skill: Benchmarking Risk Exposure in Agent Skill Retrieval****ArXiv CS.AI** **50%**

arXiv:2606.10388v2 Announce Type: replace-cross Abstract: Agent skill libraries are becoming routable software assets: a retrieved skill can contribute instructions, scripts, resource bindings, and execution assumptions to an agent. This makes retrieval failures more specific...

<https://arxiv.org/abs/2606.10388>

5 **CONF****Towards Quantifying Benchmark Optimization in ASR Models****ArXiv CS.AI** **50%**

arXiv:2608.19936v1 Announce Type: cross Abstract: Public benchmarks are important measures of Automatic Speech Recognition (ASR) model capabilities. However, by nature of being public, there is risk of models being optimized for these benchmarks in ways that do not generalize...

<https://arxiv.org/abs/2608.19936>

6 **CONF****Designing Human-mediated AI Guidance: Ready Together for Personalized Family Emergency Preparedness****ArXiv CS.AI** **50%**

arXiv:2608.19950v1 Announce Type: cross Abstract: Artificial intelligence (AI) systems are increasingly used across domains to provide personalized information, recommendations, and decision support. However, in some contexts, AI-generated information may not be suitable for...

<https://arxiv.org/abs/2608.19950>

7 **CONF****Open-Vocabulary 3D Object Detection with Co-Distillation Discovery and Dual Guidance Robust Training****ArXiv CS.AI** **50%**

arXiv:2608.19973v1 Announce Type: cross Abstract: Recently, open-vocabulary 3D object detection (3D-OVD) has gained increasing attention for its ability to detect unseen objects in 3D scenes. Existing approaches typically adopt a two-stage pipeline that first discovers novel...

<https://arxiv.org/abs/2608.19973>

8 **CONF****EchoCoT: Extracting Hidden Chain-of-Thought from Large Reasoning Models****ArXiv CS.AI** **50%**

arXiv:2608.20055v1 Announce Type: cross Abstract: Hidden chain-of-thought (CoT) traces, especially those from frontier proprietary large reasoning models (LRMs), are valuable model assets. Yet whether these hidden CoTs can be directly extracted from black-box models remains...

<https://arxiv.org/abs/2608.20055>

9 CONF

A Standardized Framework for Machine Learning in Power System Protection

ArXiv CS.AI 50%

arXiv:2608.20181v1 Announce Type: cross Abstract: Studies of machine-learning-based power-system protection increasingly report near-perfect scores, yet the meaning of those scores depends strongly on the evaluation setting. Protection task, physical scope, measurements, timing,...

<https://arxiv.org/abs/2608.20181>

0 CONF

Daedalus-150M: A Convolution-Attention Hybrid Designed for CPU Inference

ArXiv CS.AI 50%

arXiv:2608.20210v1 Announce Type: cross Abstract: Small language models are usually built like large ones and then squeezed onto a CPU afterwards. We did the opposite: we fixed the target first, one user, one token at a time, 4-bit weights, ordinary CPU, and chose the...

<https://arxiv.org/abs/2608.20210>

1 CONF

GridCodex: A RAG-Driven AI Framework for Power Grid Code Reasoning and Compliance

ArXiv CS.AI 50%

arXiv:2508.12682v2 Announce Type: replace Abstract: The global shift towards renewable energy presents unprecedented challenges for the electricity industry, making regulatory reasoning and compliance increasingly vital. Grid codes, the regulations governing grid operations, are...

<https://arxiv.org/abs/2508.12682>

2 CONF

Computational Phenomenology of Borderline Personality Disorder: A Comparative Evaluation of LLM-Simulated Expert Personas and Human Clinical Experts

ArXiv CS.AI 50%

arXiv:2508.19008v3 Announce Type: replace Abstract: Building on a human-led thematic analysis of clinical life-story interviews (> 150,000 words) with inpatients with Borderline Personality Disorder, this study examines the capacity of large language models (OpenAI's GPT,...

<https://arxiv.org/abs/2508.19008>

3 CONF

Gen AI in Proof-based Math Courses: A Pilot Study

ArXiv CS.AI 50%

arXiv:2509.13570v2 Announce Type: replace Abstract: With the rapid rise of generative AI in higher education, understanding how students use AI is increasingly important. This exploratory study examines student use and perceptions of generative AI across three proof-based...

<https://arxiv.org/abs/2509.13570>

4 CONF

ATBench: A Diverse and Realistic Agent Trajectory Benchmark for Safety Evaluation and Diagnosis

ArXiv CS.AI 50%

arXiv:2604.02022v4 Announce Type: replace Abstract: Evaluating the safety of LLM-based agents is increasingly important because risks in realistic deployments often emerge over multi-step interactions rather than isolated prompts or final responses. Existing trajectory-level...

<https://arxiv.org/abs/2604.02022>

5 CONF

LLM Capability Limits: Static Emergence and Dynamic Boundary Control

ArXiv CS.AI 50%

arXiv:2608.01548v3 Announce Type: replace Abstract: Test-time emergence in LLM systems has a deployment boundary: additional computation can realize decisions already supported by the deployed information--execution structure, while evidence, tools, memory, and executable...

<https://arxiv.org/abs/2608.01548>

6 CONF

Evaluating Investment Logic in Large Language Models: A Real-World Benchmark Towards Personalized Financial Agents

ArXiv CS.AI 50%

arXiv:2608.06108v2 Announce Type: replace Abstract: Investment competence is inherently personalized: the same market evidence can justify different actions for investors with different goals, horizons, portfolios, and risk boundaries. Yet financial LLMs are evaluated either by...

<https://arxiv.org/abs/2608.06108>

7 CONF

LongRCA Bench: Diagnosing Responsible Roles and Root Causes in Long-Horizon Agent Failures

ArXiv CS.AI 50%

arXiv:2608.15242v2 Announce Type: replace Abstract: When a long-horizon agent execution fails, outcome-level evaluation reveals the unsuccessful result but not where the decisive error entered the trajectory. Developers must then inspect the full execution to identify the...

<https://arxiv.org/abs/2608.15242>

8 CONF

TestifAI: Tomography-Based Testing for Deep Learning Systems

ArXiv CS.AI 50%

arXiv:2608.18900v2 Announce Type: replace Abstract: As AI systems are increasingly deployed in safety-critical application domains (e.g., autonomous driving), associated risks increase too. Deep learning models underlying modern AI systems, therefore, must undergo thorough...

<https://arxiv.org/abs/2608.18900>

9 CONF

The Thousand Brains Theory 2.0: An Extension for the Long-Range Connections of the Neocortical Heterarchy

ArXiv CS.AI 50%

arXiv:2507.05888v2 Announce Type: replace-cross Abstract: Vernon Mountcastle hypothesized that the basis for intelligence in mammals is the replication of a general computational unit, the cortical column. The Thousand Brains Theory proposed that each column is a sensorimotor...

<https://arxiv.org/abs/2507.05888>

00 CONF

The Basic B* Effect: The Use of LLM-based Agents Reduces the Distinctiveness and Diversity of People's Choices**

ArXiv CS.AI 50%

arXiv:2509.02910v2 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly act on people's behalf: they write emails, buy groceries, and book restaurants. While the outsourcing of human decision-making to AI can be convenient, it raises a fundamental...

<https://arxiv.org/abs/2509.02910>

01 CONF

TS-Reasoner: Aligning Time Series Foundation Models with LLM Reasoning

ArXiv CS.AI 50%

arXiv:2510.03519v3 Announce Type: replace-cross Abstract: Time series reasoning is crucial to decision-making in diverse domains, including finance, energy, and scientific discovery. While existing time series foundation models (TSFMs) can capture low-level dynamic patterns and...

<https://arxiv.org/abs/2510.03519>

02 CONF

The Bidding Games: Reinforcement Learning for MEV Extraction on Polygon Blockchain

ArXiv CS.AI 50%

arXiv:2510.14642v2 Announce Type: replace-cross Abstract: In blockchain networks, the strategic ordering of transactions within blocks has emerged as a significant source of profit extraction, known as Maximal Extractable Value (MEV). The transition from spam-based Priority Gas...

<https://arxiv.org/abs/2510.14642>

03 CONF

MKG-CARE: Case-Aware Reasoning with Multimodal Knowledge Graphs for Explainable Medical Image Diagnosis

ArXiv CS.AI 50%

arXiv:2605.22547v3 Announce Type: replace-cross Abstract: Medical image diagnosis has achieved significant progress with deep learning, yet existing methods often rely on isolated visual evidence and lack the ability to effectively leverage similar cases and external knowledge....

<https://arxiv.org/abs/2605.22547>

04 CONF

A Tool to Map AI Programs in the U.S.: A Snapshot from April 2026 and an Analysis of Requirements for AI Majors and Minors

ArXiv CS.AI 50%

arXiv:2606.12428v2 Announce Type: replace-cross Abstract: In this work, we locate and analyze existing undergraduate Artificial Intelligence (AI) programs in the United States in Spring 2026, creating a historic record at a time of great change in this area. To create this...

<https://arxiv.org/abs/2606.12428>

05 CONF

SoftVTBench: A Safety-Aware Visuo-Tactile Benchmark for Physically Constrained Robotic Manipulation of Deformable Objects (Early Version)

ArXiv CS.AI 50%

arXiv:2607.04234v2 Announce Type: replace-cross Abstract: Deformable object manipulation poses challenges beyond task completion: successful execution must also maintain safe physical interaction, holding the object stably without slip or drop while avoiding excessive...

<https://arxiv.org/abs/2607.04234>

06 CONF

A Model-Internal Protocol for Assessing Multimodal Models as Integrated Systems

ArXiv CS.AI 50%

arXiv:2608.11907v2 Announce Type: replace-cross Abstract: As Large Vision-Language Models increasingly aim to integrate visual generation and understanding within a single parameter space, evaluating such structural unification in a cohesive manner remains a critical challenge....

<https://arxiv.org/abs/2608.11907>

07 CONF

Ask to Be Sure: Informative Interactions for Confident Multi-Turn LLM Recommendation

ArXiv CS.AI 50%

arXiv:2608.15949v2 Announce Type: replace-cross Abstract: Recent advances in large language models (LLMs) have enabled their use as conversational recommender systems (CRS), demonstrating strong recommendation accuracy and natural dialogue. However, guiding multi-turn...

<https://arxiv.org/abs/2608.15949>