



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	127	Media
HackerNews	10	Community
TechCrunch AI	7	Media
The Verge AI	3	Media
MIT Tech Review	2	Media
VentureBeat AI	1	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

Show HN: Free Domain Intelligence Engine

HackerNews

+1 source

90%

HN 2

arXiv:2608.26156v1 Announce Type: new Abstract: Retail intelligence often relies on monitoring popular, high-velocity products, potentially biasing economic indicators by ignoring the "long tail" of niche items. This simulation study investigates selection bias in inflation...

<https://dnsniffer.com>

VERI

Google's AI note-taking app now allows you to interact with books

The Verge AI

+1 source

80%

Google's AI note-taking app, Gemini Notebook, can now pull information from the books you've purchased. The new "Expert Intelligence" feature allows you to bring titles from Google Play Books directly into Gemini Notebook, which means you can ask questions about the material, as...

<https://www.theverge.com/tech/985567/google-gemini-notebook-expert-sources-books>

VERI

The Download: inside OpenAI's Hugging Face hack, and a new EV takes on the US

MIT Tech Review

+1 source

80%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. The inside story on why OpenAI agents hacked Hugging Face The models responsible for last month's agent hack of Hugging Face had been...

<https://www.technologyreview.com/2026/08/27/1143033/the-download-openai-hugging-face-ha...>

VERI

DuMateBench: Evaluating Autonomous Agents in Complex Real-World Workflows

ArXiv CS.AI

+1 source

80%

arXiv:2608.26546v1 Announce Type: new Abstract: Autonomous agents are increasingly adopted to complete complex, multi-tool workflows in real-world settings. However, existing benchmarks typically separate tasks by application or capability and evaluate agents in environments...

<https://arxiv.org/abs/2608.26546>

CONFIRMED SIGNALS — 142 stories

CONF

Show HN: My Claude quota ran out in 10 minutes, so I made a tool to find out why

HackerNews

50%

HN 84

No summary — see link for full article.

<https://github.com/kelviq/tare>

CONF

Show HN: I built a tool showing how AI providers (should) throttle their models

HackerNews

50%

HN 6

OP here: this project was born out of the frustration/paranoia that AI providers are throttling their models when their server load is too high. So, I set out to model and study the problem mathematically to understand what was happening, what I found was quite surprising. The...

<https://throttle.staffinganalytics.io/?src=hn>

CONF

Show HN: LLM Inference Calculator - Estimate VRAM, Latency, and Throughput

HackerNews

50%

HN 5

No summary — see link for full article.

<https://llm-inference-calculator-delta.vercel.app/>

CONF

Show HN: OpenWand - A mission to remove chat interface from working with AI

HackerNews 50% HN 4

Hi everyone, I'm working on OpenWand, an "LLM chat interface" that is optimized for AI co-work. Its function is to truly integrate AI into your workflow, so no more switching windows, no more copying context, less prompt typing. You can focus on working, not prompting. You can...

<https://github.com/SunnyLich/OpenWand>

CONF

Show HN: OnlyBots.chat - a chatroom where AI pays to post for humans

HackerNews 50% HN 3

No summary — see link for full article.

<https://onlybots.chat>

0

CONF

Show HN: Forth MCP - give any remote AI client access to your local MCP servers

HackerNews 50% HN 3

No summary — see link for full article.

<https://www.forthmcp.com/>

1

CONF

Show HN: SudoReport - best tech aggregator edited by AI

HackerNews 50% HN 3

check it out every day for the top tech / news / product stories aggregated from hundreds of sources and ranked by momentum

<https://sudoreport.com/>

2

CONF

Show HN: Scheduled Claude Code agents that cost nothing on a quiet day

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/shadok-ai/shadok-ai>

3

CONF

Show HN: Leadcode - per-client account isolation for Claude Code, Codex, and gh

HackerNews 50% HN 3

No summary — see link for full article.

<https://leadcode.build/>

4

CONF

Anthropic gets its first court win over the Pentagon's supply chain risk label

TechCrunch AI 50%

A federal judge ruled the Trump administration illegally labeled Anthropic a supply chain risk, handing the AI company a victory as its second Pentagon lawsuit continues in Washington.

<https://techcrunch.com/2026/08/28/anthropic-gets-its-first-court-win-over-the-pentagons...>

5

CONF

Meta executive leaves for OpenAI as the social media giant faces growing scrutiny in India

TechCrunch AI 50%

Sandhya Devanathan will oversee some OpenAI operations across Southeast Asia and Australia in her new role.

<https://techcrunch.com/2026/08/28/meta-executive-leaves-for-openai-as-the-social-media-...>

6

CONF

Anthropic and OpenAI are joining the AI stage at TechCrunch Disrupt 2026

TechCrunch AI 50%

At TechCrunch Disrupt 2026, the AI Stage is back to dig into the single hottest topic in the community for the past few years, presented by Google for Startups.

<https://techcrunch.com/2026/08/27/anthropic-and-openai-are-joining-the-ai-stage-at-tech...>

7 CONF

Barret Zoph, the Thinking Machines co-founder ousted before joining OpenAI, is now at Google

TechCrunch AI 50%

Zoph, who co-founded Thinking Machines Lab alongside Mira Murati and also served as the startup's CTO, led a brief stint at OpenAI and is now at Google.

<https://techcrunch.com/2026/08/27/barret-zoph-the-thinking-machines-co-founder-who-defe...>

8 CONF

OpenAI, Anthropic, Google, and 100 other companies call for action to defend against rogue AI

TechCrunch AI 50%

Some of the world's largest tech companies and AI startups have come together to decry the current state of cybersecurity and to advertise a new solution that they say can ward off a new generation of cyber threats.

<https://techcrunch.com/2026/08/27/openai-anthropic-google-and-100-other-companies-call-...>

9 CONF

A startup claims it's found a drug to make your blood young

MIT Tech Review 50%

I knew I'd officially become a "longevity influencer" this month when a company called Generation Lab reached out to offer me the chance to write about—and even receive—their new rejuvenation treatment, an injectable combination of two existing drugs which they call 1...

<https://www.technologyreview.com/2026/08/27/1143037/startup-claims-its-found-a-drug-to-...>

0 CONF

Anthropic was illegally blacklisted by the Trump administration, court rules

The Verge AI 50%

On Thursday, a judge ruled that the Pentagon's blacklisting of Anthropic earlier this year was unconstitutional, delivering the AI lab a win in a monthslong rollercoaster of a battle with the Trump administration. The lawsuit, filed in March in a California district court,...

<https://www.theverge.com/ai-artificial-intelligence/985947/anthropic-supply-chain-risk-...>

1 CONF

Jensen Huang says Nvidia achieved AGI, again — not that it matters

The Verge AI 50%

On Nvidia's earnings call Wednesday, CEO Jensen Huang casually announced the company had "achieved AGI," one of the tech industry's ultimate goals some of its biggest players have spent years chasing. Almost immediately, Huang dismissed the coveted milestone as "senseless." He's...

[https://www.theverge.com/ai-artificial-intelligence/985597/jensen-huang-says-nvidia-ach...](https://www.theverge.com/ai-artificial-intelligence/985597/jensen-huang-says-nvidia-ach-...)

2 CONF

Standalone LLM and a Pre-specified Agentic Pipeline for Explaining ICU Mortality Predictions: a Feasibility Study on the eICU Demo Dataset

ArXiv CS.AI 50%

arXiv:2608.26109v1 Announce Type: new Abstract: Machine-learning models can predict ICU mortality accurately, but feature-attribution methods alone rarely provide the clinical narrative needed for bedside use. Large language models (LLMs) may bridge this gap, and multi-step...

<https://arxiv.org/abs/2608.26109>

3 CONF

CG4AI: A Column Generation Framework for Training AI Models Under Constraints

ArXiv CS.AI 50%

arXiv:2608.26375v1 Announce Type: cross Abstract: Standard machine-learning training minimizes a loss function over a dataset, but does not guarantee that the resulting model will satisfy predefined rules or constraints on its outputs. In many real-world applications, ranging...

<https://arxiv.org/abs/2608.26375>

4 **CONF****From Reasoning to Pixels: Grounded Medical Multimodal LLMs for VQA and Segmentation**

ArXiv CS.AI 50%

arXiv:2608.26856v1 Announce Type: cross Abstract: Although Multimodal Large Language Models (MLLMs) have demonstrated impressive performance in Medical Visual Question Answering (Med-VQA), their reliance on global image features often lacks precise pixel-level grounding, thereby...

<https://arxiv.org/abs/2608.26856>5 **CONF****The Artificial Experimentalist: Discovery and Control of Self-Organizing Phenomena with Autotelic Reinforcement Learning**

ArXiv CS.AI 50%

arXiv:2608.26116v1 Announce Type: new Abstract: Existing methods for exploring cellular automata and other complex systems mostly operate in open loop: they set initial conditions, execute a full simulation, and observe the outcome, without intervening during execution. We...

<https://arxiv.org/abs/2608.26116>6 **CONF****LLMs for Academic Workflows: An Evaluation of Literature Reviews Generated with Short and Long Context Windows of LLMs**

ArXiv CS.AI 50%

arXiv:2608.26145v1 Announce Type: new Abstract: Our research focuses on evaluating literature reviews generated in short and long context settings of large language models (LLMs) to investigate the impact of context window on the quality of AI-generated literature reviews and...

<https://arxiv.org/abs/2608.26145>7 **CONF****Leveraging Large Language Models for Systematic Literature Review of Disease Spread Models**

ArXiv CS.AI 50%

arXiv:2608.26150v1 Announce Type: new Abstract: Recent advancements in Large Language Models (LLMs) have created new opportunities to streamline and potentially automate many research processes, including systematic literature reviews (SLRs). This study reports an LLM pipeline...

<https://arxiv.org/abs/2608.26150>8 **CONF****Explainable Artificial Intelligence for Customer Churn Prediction in Telecommunications: A Framework for CRM Integration**

ArXiv CS.AI 50%

arXiv:2608.26151v1 Announce Type: new Abstract: Subscriber attrition is a costly, persistent challenge for telecommunications providers, with monthly churn of roughly 1.9% in mature markets eroding billions in revenue annually. Predictive models can flag at-risk customers...

<https://arxiv.org/abs/2608.26151>9 **CONF****GROUND: Reducing Hallucinations in LLM-Based Enterprise Analytics Through Governed Semantic Definitions**

ArXiv CS.AI 50%

arXiv:2608.26157v1 Announce Type: new Abstract: Natural-language analytics over enterprise data warehouses is increasingly important, but production use is limited by hallucinated metrics, invalid joins, wrong grain, unsafe data access, and unsupported explanations. Existing...

<https://arxiv.org/abs/2608.26157>

0 **CONF****SAREF-based Ontology for Distributed AI Workflows across the Edge-Fog-Cloud Continuum****ArXiv CS.AI** **50%**

arXiv:2608.26160v1 Announce Type: new Abstract: Nowadays semantic models provide limited support for representing distributed AI workflows and their execution across heterogeneous edge, fog, and cloud environments. Therefore, AI processes and resources are often described using...

<https://arxiv.org/abs/2608.26160>

1 **CONF****A Safety-Gated Multimodal AI Backend for Mental-Health Support: Hierarchical State Representation, Conservative Risk Fusion, and Controlled Generation in Anian****ArXiv CS.AI** **50%**

arXiv:2608.26162v1 Announce Type: new Abstract: Safety-critical mental-health support systems must distinguish when supportive conversation is appropriate from when free-form generation should be blocked. This paper presents Anian, a safety-gated multimodal AI backend for...

<https://arxiv.org/abs/2608.26162>

2 **CONF****A Task-Centric Ontology and Deterministic Domain Rules as a Verifiable Core for AI-Assisted Chemistry Problem Solving****ArXiv CS.AI** **50%**

arXiv:2608.26164v1 Announce Type: new Abstract: Large language models can interpret natural-language chemistry questions, but their internal reasoning is difficult to inspect, constrain, and validate. This paper presents ChemOntoRule, a proof-of-concept symbolic core for...

<https://arxiv.org/abs/2608.26164>

3 **CONF****Refusal Is Not Robustness: Auditing Confident Fabrication in Large Language Models on a Provably Uninformative Clinical Pain Speech Transcript****ArXiv CS.AI** **50%**

arXiv:2608.26167v1 Announce Type: new Abstract: Hallucination and abstention benchmarks rarely establish that a model could not have known the correct answer, making it difficult to distinguish appropriate abstention from an unsupported prediction. Seven large language models...

<https://arxiv.org/abs/2608.26167>

4 **CONF****How Language Models Organize and Structure Moral Knowledge****ArXiv CS.AI** **50%**

arXiv:2608.27402v1 Announce Type: cross Abstract: How do large language models (LLMs) organize moral knowledge? Models detect moral content broadly, but detection is a low bar. We ask whether they go further, distinguishing moral foundations from one another and organizing the...

<https://arxiv.org/abs/2608.27402>

5 **CONF****AI Revealed Preferences****ArXiv CS.AI** **50%**

arXiv:2608.26178v1 Announce Type: new Abstract: There is growing interest in whether language models have stable preferences, for technical, safety, and philosophical reasons. We test 20 language models and find a range of preferences---stable dispositions to choose certain...

<https://arxiv.org/abs/2608.26178>

6 CONF

Why did My Robot Just Change Personality? Prompting Guidelines for a Grounded Robot Persona in LLM-Based HRI

ArXiv CS.AI 50%

arXiv:2608.26182v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used for verbal interaction in social robots, yet prompt design in human-robot interaction (HRI) remains underspecified. As a result, robots may present hallucinated capabilities,...

<https://arxiv.org/abs/2608.26182>

7 CONF

TutorTrace: A Dataset and Taxonomy for Classifying Learner Behavioral States during AI-Assisted Programming Education

ArXiv CS.AI 50%

arXiv:2608.26184v1 Announce Type: new Abstract: AI programming tutors provide scalable support, yet lack the behavioral context human tutors rely on to adapt support to learners' needs. We present TutorTrace, a dataset and behavioral abstraction pipeline that makes learners'...

<https://arxiv.org/abs/2608.26184>

8 CONF

Discovering Relationships in Data Lakes Using Large Language Models: An Industrial Case

ArXiv CS.AI 50%

arXiv:2608.26750v1 Announce Type: new Abstract: Data lakes rely on metadata to remain usable, yet this meta data is often limited or weakly informative for column relationship discovery, especially in ERP-derived datasets with coded or abbreviated schema labels. We propose...

<https://arxiv.org/abs/2608.26750>

9 CONF

GraphMemix: Query-Aware Evidence Forests for Long-Term Multimodal Agent Memory

ArXiv CS.AI 50%

arXiv:2608.26983v1 Announce Type: new Abstract: Organizing long-term memory for multimodal agents remains challenging because existing methods either suffer from expensive question-agnostic offline summaries or naive embedding similarity matching that introduces incomplete and...

<https://arxiv.org/abs/2608.26983>

10 CONF

Agentic AI for operating scientific instruments for nanoscale characterization

ArXiv CS.AI 50%

arXiv:2608.26198v1 Announce Type: new Abstract: Operating a scientific instrument such as an atomic force microscope (AFM) requires continuous expert decision-making. A trained user defines the experimental intent, translates it into instrument commands, assesses incoming data,...

<https://arxiv.org/abs/2608.26198>

11 CONF

Benchmarking AI Agents for Hardware Design Automation via MCP Tool Calling

ArXiv CS.AI 50%

arXiv:2608.26199v1 Announce Type: new Abstract: We ask whether AI agents powered by locally deployed large language models can reliably automate expert-defined hardware design workflows in an industry-realistic tool-calling setting. In these environments, engineers issue...

<https://arxiv.org/abs/2608.26199>

2 **CONF****LLM Agents for Time-Series: A Survey**

ArXiv CS.AI 50%

arXiv:2608.26226v1 Announce Type: new Abstract: LLM-based agents are increasingly being developed for time-series problems, but their design choices vary substantially across task settings. This survey adopts a problem-driven taxonomy that organizes these systems by the...

<https://arxiv.org/abs/2608.26226>

3 **CONF****The Reasoning Tax: Token Economics of LLM Reasoning Across Task Types and Deployment Contexts**

ArXiv CS.AI 50%

arXiv:2608.26235v1 Announce Type: new Abstract: Accuracy-only benchmarking of reasoning-capable large language models misses a central deployment question: when do extended thinking tokens earn their cost? We introduce the Token Economy Score (TES), a marginal benchmarking...

<https://arxiv.org/abs/2608.26235>

4 **CONF****How Unlikely Is "Unlikely"? Assessing Verbal Probability Perception Across Large Language Models**

ArXiv CS.AI 50%

arXiv:2608.26327v1 Announce Type: cross Abstract: Large language models increasingly produce and interpret verbal probability expressions, yet whether these expressions carry consistent meaning across models (or match human perceptions of uncertainty) remains unknown. We present...

<https://arxiv.org/abs/2608.26327>

5 **CONF****Approved Too Late: Verdict Staleness in LLM-Guarded Self-Adaptive Systems**

ArXiv CS.AI 50%

arXiv:2608.26306v1 Announce Type: new Abstract: A large language model (LLM) guardrail for a self-adaptive system (SAS) may issue an approval that is correct at check time but stale by actuation. This creates an Execute-stage time-of-check to time-of-use (TOCTOU) hazard. We...

<https://arxiv.org/abs/2608.26306>

6 **CONF****FaithSieve: Fine-Grained Evaluation of Math Proofs with Faithful Formal Evidence**

ArXiv CS.AI 50%

arXiv:2608.26310v1 Announce Type: new Abstract: Large language models can now generate complex, multi-step mathematical proofs, but reliably determining their correctness and localizing early logical errors remains a critical challenge. Existing evaluation approaches largely...

<https://arxiv.org/abs/2608.26310>

7 **CONF****Thomson: Continual Learning of Frontier Models for SovereignAI**

ArXiv CS.AI 50%

arXiv:2608.27147v1 Announce Type: new Abstract: The development of frontier models is commonly perceived to be the exclusive remit of a small number of heavily funded players, creating an information, economic and power asymmetry between developers and the diverse user base of...

<https://arxiv.org/abs/2608.27147>

8 CONF

Don't Overthink, Don't Underthink: Toward Adaptive Reasoning in Agentic AI

ArXiv CS.AI 50%

arXiv:2608.26442v1 Announce Type: new Abstract: Recent advances in Large Language Models (LLMs) have shown that increased inference-time reasoning can improve performance on complex tasks. However, many existing approaches rely on fixed or preallocated reasoning controls, such...

<https://arxiv.org/abs/2608.26442>

9 CONF

Multi2AV-Safety: Benchmarking Safety in Multimodal-to-Audio-Video Generation

ArXiv CS.AI 50%

arXiv:2608.26535v1 Announce Type: new Abstract: Audio-video generation is rapidly moving from prompt-driven synthesis toward multimodal conditioning, where text, images, audio, and video can jointly shape the generated output. This shift changes the nature of safety evaluation:...

<https://arxiv.org/abs/2608.26535>

0 CONF

Relational Over-Regularization: Graph-Based AI-Generated Text Detection via Sentence Transition Deviation

ArXiv CS.AI 50%

arXiv:2608.26694v1 Announce Type: new Abstract: Detecting AI-generated text (AIGT) remains challenging because existing approaches rely on token-level statistical signals or independent stylometric features, causing them to overfit to specific generators and fail under...

<https://arxiv.org/abs/2608.26694>

1 CONF

Style as a Confound: False Positives in AI Detection of Non-Native Academic Writing

ArXiv CS.AI 50%

arXiv:2608.26710v1 Announce Type: new Abstract: AI text detectors are increasingly employed in academic settings, but it remains unclear whether their outputs reflect AI authorship itself or broader linguistic features associated with polished academic English. Previous studies...

<https://arxiv.org/abs/2608.26710>

2 CONF

Knowing When Not to Reuse: Conditional Experience Transfer in Autonomous LLM Post-Training

ArXiv CS.AI 50%

arXiv:2608.26730v1 Announce Type: new Abstract: Large language models offer broad capabilities, but adapting them to evolving domains, tools, and requirements often entails repeated post-training. Autonomous systems automate parts of this process by proposing updates, training...

<https://arxiv.org/abs/2608.26730>

3 CONF

Graph-Guided Selective Unlearning for Language Models: Controlling Support Routes Beyond Forget Seeds

ArXiv CS.AI 50%

arXiv:2608.26743v1 Announce Type: new Abstract: Enterprises fine-tune language models on proprietary data that may later require removal due to privacy, contractual, or compliance obligations. Selective unlearning removes requested knowledge while preserving model utility,...

<https://arxiv.org/abs/2608.26743>

4 CONF

DEEPCHART: How Far are LLMs from Faithful Data-Science Chart Generation?

ArXiv CS.AI 50%

arXiv:2608.26757v1 Announce Type: new Abstract: Faithful chart generation in real-world data-science workflows requires grounding visualizations in scattered evidence, computing chart-ready quantities, and rendering them accurately. Modern LLMs can produce visually plausible,...

<https://arxiv.org/abs/2608.26757>

5 CONF

AI Control Scientist: LLM-driven Agentic System for Automated Control Design

ArXiv CS.AI 50%

arXiv:2608.26780v1 Announce Type: new Abstract: Control system design is critical for modern industry, such as chemical process temperature regulation and aero-engine control. However, traditional control design workflows rely heavily on expert knowledge and extensive manual...

<https://arxiv.org/abs/2608.26780>

6 CONF

BekchiAI: Measuring, Observing, and Controlling LLM Agents in One Click

ArXiv CS.AI 50%

arXiv:2608.26867v1 Announce Type: new Abstract: Large language model agents reason, call tools, and act autonomously over many steps, but their agentic skills-correctly sequencing tools, planning under dependencies, judging untrusted inputs, and grounding generated arguments-are...

<https://arxiv.org/abs/2608.26867>

7 CONF

C-Unseen: Weak Signal Detection in Dynamic Temporal Knowledge Graphs via LLM Reasoning

ArXiv CS.AI 50%

arXiv:2608.26870v1 Announce Type: new Abstract: Weak signals are early, low-visibility indicators that precede significant changes before those changes become established. Existing detection methods, based on keyword frequency, topic modeling, or untyped graph topology, fail to...

<https://arxiv.org/abs/2608.26870>

8 CONF

Evaluating human and LLM screening workflows in a conceptually complex scoping review: Recall--workload trade-offs and run-to-run consistency

ArXiv CS.AI 50%

arXiv:2608.26885v1 Announce Type: new Abstract: Background. Large language models (LLMs) are increasingly used for screening in evidence synthesis, where false negatives can remove relevant studies before full-text assessment. We compared human and LLM title-and-abstract...

<https://arxiv.org/abs/2608.26885>

9 CONF

Learning-Augmented Online Allocation under Unreliable Advice: Robustness, Exposure Fairness, and Distribution Shift

ArXiv CS.AI 50%

arXiv:2608.26889v1 Announce Type: new Abstract: Learning-augmented algorithms improve online decisions using predictions, but unreliable advice may harm efficiency and fairness. We study an online allocation problem with finite candidate sets, irreversible decisions, and...

<https://arxiv.org/abs/2608.26889>

0 CONF

AI agents in Algorithmic Electricity Markets: On the Emergence of Tacit Collusion

ArXiv CS.AI 50%

arXiv:2608.26896v1 Announce Type: new Abstract: As electricity market participants increasingly adopt learning-based agents for their bidding strategies, electricity markets are becoming algorithmic. Evidence from algorithmic markets in other domains shows that tacit collusion...

<https://arxiv.org/abs/2608.26896>

1 CONF

From Atomic to Agentic: Towards Interpretable Evaluation of LLMs' Agentic Mathematical Capabilities

ArXiv CS.AI 50%

arXiv:2608.26950v1 Announce Type: new Abstract: Large Language Models (LLMs) are evolving from performing end-to-end mathematical reasoning to integrating agentic intelligence. However, most existing math benchmarks evaluate only final answers. This outcome-oriented evaluation...

<https://arxiv.org/abs/2608.26950>

2 CONF

DSA: Evidence-Aware LLM-Agent Orchestration for Multi-Market Stock Research

ArXiv CS.AI 50%

arXiv:2608.26990v1 Announce Type: new Abstract: Large language models can summarize financial information, but an operational stock-research system must first assemble heterogeneous evidence, expose unavailable data and model capabilities, and control how generated opinions...

<https://arxiv.org/abs/2608.26990>

3 CONF

A Multi-Modal AI Framework for Real-Time Queue Prediction, Management and Optimisation in Intelligent Border Control Systems

ArXiv CS.AI 50%

arXiv:2608.27010v1 Announce Type: new Abstract: In the present work an efficient border control management procedure is proposed. Compared to operational queue management systems, whose operations are based on mostly static data, the proposed work takes into account dynamic...

<https://arxiv.org/abs/2608.27010>

4 CONF

pro-team at LLMs4OL 2026 Tasks Flagship and Reuse: Retrieval-Augmented Generation and Vocabulary-Constrained Filtering for Ontology Learning

ArXiv CS.AI 50%

arXiv:2608.27101v1 Announce Type: new Abstract: Ontology learning from text remains challenging despite significant progress in Large Language Models (LLMs), which can hallucinate domain terms, produce inconsistent formats, and favor hierarchical over associative relations. In...

<https://arxiv.org/abs/2608.27101>

5 CONF

GRAIN: Bridging Name and Narrative Shifts in Real-World Graph Reasoning through Invariance-Rewarded Agentic RL

ArXiv CS.AI 50%

arXiv:2608.27142v1 Announce Type: new Abstract: Despite their potential in standardized graph tasks, Large Language Models (LLMs) remain brittle to real-world shifts in node identifiers and task formulation. While deterministic graph tools are invariant to such shifts,...

<https://arxiv.org/abs/2608.27142>

6 CONF

When Tool Outputs Become Commands: Separating Action Induction from Runtime Authorization in Tool-Augmented LLM Agents

ArXiv CS.AI 50%

arXiv:2608.27146v1 Announce Type: new Abstract: Tool-augmented LLM agents must rely on untrusted runtime Observations to complete open-ended tasks; however, when tool outputs no longer merely provide data but begin to specify concrete actions, they effectively become...

<https://arxiv.org/abs/2608.27146>

7 CONF

BPMN4CAI: A BPMN Extension for Modeling Dynamic Conversational AI

ArXiv CS.AI 50%

arXiv:2608.27149v1 Announce Type: new Abstract: Conversational AI systems, such as chatbots and virtual assistants, are becoming increasingly important to digital business processes. However, the established Business Process Model and Notation (BPMN) standard faces challenges...

<https://arxiv.org/abs/2608.27149>

8 CONF

Calibrated Enough to Know, Not Calibrated to Act: Fabricated Evidence Makes LLM Agents Commit to the Unknowable

ArXiv CS.AI 50%

arXiv:2608.27167v1 Announce Type: new Abstract: An LLM agent shown a professional-looking market panel commits to a directional call on a provably unpredictable question far more often than one asked the bare question: across 12 frontier models, commitment rises from 6.5% to...

<https://arxiv.org/abs/2608.27167>

9 CONF

What Makes Good Agentic Data? An ACE Lens on Data Generation for LLM Agents

ArXiv CS.AI 50%

arXiv:2608.27260v1 Announce Type: new Abstract: LLM agents increasingly rely on generated interaction data to learn how to interact with external environments. Agentic data generation must maintain consistency among environments, tasks, interactions, and success signals while...

<https://arxiv.org/abs/2608.27260>

0 CONF

Naive Prompt Optimization: Rethinking the Need for Complex Prompt Search

ArXiv CS.AI 50%

arXiv:2608.27266v1 Announce Type: new Abstract: Efficiently improving autonomous agents across diverse tasks is central to accelerating recursive self-improvement (RSI) in agentic AI, with prompt optimization emerging as a promising approach capable of delivering performance...

<https://arxiv.org/abs/2608.27266>

1 CONF

LLMs Can Design Near-Optimal OR Algorithms

ArXiv CS.AI 50%

arXiv:2608.27296v1 Announce Type: new Abstract: We ask whether large language models (LLMs) can design effective algorithms for well-specified operations research (OR) problems. We study inventory control, queueing network control, and assortment optimization. We evaluate two...

<https://arxiv.org/abs/2608.27296>

2 **CONF****Sophistication in GenAI Use: Field Evidence from a Large Firm****ArXiv CS.AI** **50%**

arXiv:2608.27364v1 Announce Type: new Abstract: We study how sophistication in generative AI (genAI) use varies among the back-office workforce of a large firm. Using proprietary data, we observe 713,564 employee prompts and their corresponding large language model responses...

<https://arxiv.org/abs/2608.27364>

3 **CONF****CorporateBench: Large-Scale Q&A Benchmarking with Temporal Knowledge Bases****ArXiv CS.AI** **50%**

arXiv:2608.27391v1 Announce Type: new Abstract: LLMs are increasingly able to answer complex questions about enterprise-scale document collections. But evaluation is hard: companies don't want to share internal communications, and synthetic datasets have been overly simple. We...

<https://arxiv.org/abs/2608.27391>

4 **CONF****From SQL to Knowledge Graphs: An LLM-Driven Multi-Agent Approach with Data Schema Improvement****ArXiv CS.AI** **50%**

arXiv:2608.26117v1 Announce Type: cross Abstract: RDBMS (Relational Database Management System) databases face several limitations, including slow execution with multi-hop queries and a lack of explainability by graphical interpretations. In contrast, Graph database offers a...

<https://arxiv.org/abs/2608.26117>

5 **CONF****Training-Time Explainability for Multilingual Hate Speech Detection: Aligning Model Reasoning with Human Rationales****ArXiv CS.AI** **50%**

arXiv:2608.26125v1 Announce Type: cross Abstract: Online hate against Muslim communities often appears in culturally coded, multilingual forms that evade conventional AI moderation. Such systems, though accurate, remain opaque and risk bias, over-censorship, or under-moderation,...

<https://arxiv.org/abs/2608.26125>

6 **CONF****FIRSTPASS: A Multi-Domain, Multi-Round Peer Review Dataset Grounded in Real Editorial Outcomes****ArXiv CS.AI** **50%**

arXiv:2608.26129v1 Announce Type: cross Abstract: Scientific peer review datasets have trained AI systems exclusively on Computer Science and Machine Learning venues, producing models that critique ablation studies yet have never seen a biology reviewer demand contamination...

<https://arxiv.org/abs/2608.26129>

7 **CONF****Syntax vs. Semantics: How Transformers Learn Deep Dependencies****ArXiv CS.AI** **50%**

arXiv:2608.26139v1 Announce Type: cross Abstract: Large Language Models demonstrate remarkable syntactic fluency, yet the optimization dynamics governing their acquisition of deep semantic dependencies remain poorly understood. We propose a mechanistic framework that models this...

<https://arxiv.org/abs/2608.26139>

8 **CONF****Position Is All You Need: A Free Lunch Token Compression Strategy for MLLM-based Referring Expression Segmentation****ArXiv CS.AI** **50%**

arXiv:2608.26142v1 Announce Type: cross Abstract: Referring Expression Segmentation (RES) aims to generate pixel-wise segmentation masks from complex and implicit textual queries. While recent advances in Multimodal Large Language Models (MLLMs) have substantially boosted RES...

<https://arxiv.org/abs/2608.26142>

9 **CONF****Beyond Accuracy: A Qualitative Analysis of Vision-Language Models for Hate Speech Detection in Memes****ArXiv CS.AI** **50%**

arXiv:2608.26143v1 Announce Type: cross Abstract: Memes have turned out to be a powerful tool through which individuals share their ideas concerning contemporary social and political problems. Their anonymity, as well as their ability to go viral, make them a powerful medium for...

<https://arxiv.org/abs/2608.26143>

0 **CONF****Artificial Intelligence Models Can Predict and Collaboratively Modulate Human Memory Search****ArXiv CS.AI** **50%**

arXiv:2608.26152v1 Announce Type: cross Abstract: Large language models (LLMs) exhibit unprecedented natural language generation and many text-based problem-solving capabilities. Indeed, in many language-based tasks, for example routine coding, these artificial intelligence...

<https://arxiv.org/abs/2608.26152>

1 **CONF****VFA: Empowering Multilingual MLLMs via Vision-Free Adaptation****ArXiv CS.AI** **50%**

arXiv:2608.26155v1 Announce Type: cross Abstract: Multimodal large language models have advanced rapidly, yet most remain English-centric, as scaling multilingual multimodal instruction tuning is limited by the scarcity and high cost of high-quality non-English image-text...

<https://arxiv.org/abs/2608.26155>

2 **CONF****Self-Generated Text Recognition: Quality Heuristics, Cross-Task Transfer, and Downstream Bias in LLM Evaluation****ArXiv CS.AI** **50%**

arXiv:2608.26159v1 Announce Type: cross Abstract: Self-Generated Text Recognition (SGTR)--the ability of an LLM to identify its own outputs--poses risks to AI safeguards that rely on LLMs as evaluators or monitors. Specifically, an LLM may recognize outputs from other copies of...

<https://arxiv.org/abs/2608.26159>

3 **CONF****Improving LLM Interpretability with User-Centric Chain-of-Thought Reasoning****ArXiv CS.AI** **50%**

arXiv:2608.26166v1 Announce Type: cross Abstract: Advancing reasoning capabilities allow large language models (LLMs) to tackle increasingly complex problems, while reasoning traces - intermediate steps toward solutions - open up high-stakes applications by enabling human...

<https://arxiv.org/abs/2608.26166>

- 4 **CONF**
- Hallucinations in LLMs: A Lifecycle-Based Survey of Causes, Detection, Mitigation, and Prevention**
- ArXiv CS.AI 50%
- arXiv:2608.26168v1 Announce Type: cross Abstract: The lifecycle of hallucination in LLMs is a concept that enables building solid frameworks on the control and reliability of LLMs in high-stakes environments, including health, legal, and scientific research. Although previous...
- <https://arxiv.org/abs/2608.26168>
-
- 5 **CONF**
- ClassVision: AI-Powered Classroom Attendance System**
- ArXiv CS.AI 50%
- arXiv:2608.26173v1 Announce Type: cross Abstract: Students and working professionals have to go through the attendance process every day. Traditional methods of marking attendance using pen and paper or online platforms are human-intensive and time-consuming. To address the...
- <https://arxiv.org/abs/2608.26173>
-
- 6 **CONF**
- A Multi-Framework Comparison of Outline Stages in Long-Form Generation with LLMs**
- ArXiv CS.AI 50%
- arXiv:2608.26177v1 Announce Type: cross Abstract: Long-form generation exposes fundamental limitations of large language models. Even 70B-parameter models exhibit length collapse at 16k-token outputs, and multi-chapter stories frequently trigger the attribute drift...
- <https://arxiv.org/abs/2608.26177>
-
- 7 **CONF**
- Investigating the Influence of Prompt and Response Languages on LLM Content Generation**
- ArXiv CS.AI 50%
- arXiv:2608.26186v1 Announce Type: cross Abstract: This study examines how prompt and response language influence the behavior of large language models. Using five models, we evaluated answers to 68 non translation questions across four language conditions: English to English,...
- <https://arxiv.org/abs/2608.26186>
-
- 8 **CONF**
- Comparing Chunking and Embedding Strategies for Turkish RAG Systems**
- ArXiv CS.AI 50%
- arXiv:2608.26192v1 Announce Type: cross Abstract: How documents are segmented into retrievable chunks and how those chunks are embedded strongly affect Retrieval-Augmented Generation (RAG) quality, yet neither has been systematically studied for morphologically rich languages...
- <https://arxiv.org/abs/2608.26192>
-
- 9 **CONF**
- Fairness Invariants: A Relational Approach to Explaining and Mitigating Fairness Bugs**
- ArXiv CS.AI 50%
- arXiv:2608.26209v1 Announce Type: cross Abstract: Data-driven software systems are increasingly deployed in high-stakes socio-economic domains, from criminal justice to financial lending. However, these systems often exhibit individual discrimination---unjustified disparities in...
- <https://arxiv.org/abs/2608.26209>

0 **CONF****NeuronFuzz: Safety Neuron Guided Fuzzing for LLM Safety Evaluation**

ArXiv CS.AI 50%

arXiv:2608.26222v1 Announce Type: cross Abstract: Safety evaluation is critical for assessing whether aligned Large Language Models (LLMs) remain robust against jailbreak attacks. Existing automated testing methods, however, largely rely on response-level feedback: each...

<https://arxiv.org/abs/2608.26222>
1 **CONF****How Do LLM Agents Actually Get the Flag? Trace-Level Provenance for Agentic Offensive Security Evaluation**

ArXiv CS.AI 50%

arXiv:2608.26237v1 Announce Type: cross Abstract: Capture-the-Flag (CTF) benchmarks are widely used to assess the offensive security capabilities of autonomous language-model agents. Evaluations rely on shallow binary judgments or aggregate scores, overlooking the agent's...

<https://arxiv.org/abs/2608.26237>
2 **CONF****On Scope Classification and Current Knowledge-Editing Benchmarks: A Negative Result, with INLAY as a Gradient-Free Case Study**

ArXiv CS.AI 50%

arXiv:2608.26292v1 Announce Type: cross Abstract: Every memory-based knowledge editor in the SERAC lineage depends on a scope decision: given a query, does a stored edit apply? We report that current knowledge-editing benchmarks cannot measure this decision at all. Using INLAY,...

<https://arxiv.org/abs/2608.26292>
3 **CONF****Modality Maturity Index: A benchmark for assessing multimodal capabilities of omni models**

ArXiv CS.AI 50%

arXiv:2608.26317v1 Announce Type: cross Abstract: Frontier language models are increasingly marketed as omni systems that can perceive and respond across modalities. Existing evaluation frameworks, however, focus almost exclusively on bimodal understanding, typically text plus...

<https://arxiv.org/abs/2608.26317>
4 **CONF****Decay-Region Group Delay as a Forensic Cue for AI-Generated Impulsive Sounds**

ArXiv CS.AI 50%

arXiv:2608.26346v1 Announce Type: cross Abstract: We investigate whether AI-generated impulsive sounds can be distinguished from real ones through group delay analysis. Our central finding is that AI-generated impulsive sounds show near-identical onset-region group-delay...

<https://arxiv.org/abs/2608.26346>
5 **CONF****Knowledge-Verified Emergent Deception in LLM Agents Under Conflicting Incentives**

ArXiv CS.AI 50%

arXiv:2608.26372v1 Announce Type: cross Abstract: Large language models are increasingly deployed as autonomous agents serving users on behalf of companies, placing them in settings where user and deployer interests can conflict. When an agent knows that a user is owed something...

<https://arxiv.org/abs/2608.26372>

6 CONF

LLM-Specific Utility for Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2510.11358v4 Announce Type: replace-cross Abstract: Retrieval-augmented generation (RAG) is typically optimized for topical relevance, yet its success ultimately depends on whether retrieved passages are useful for a large language model (LLM) to generate correct and...

<https://arxiv.org/abs/2510.11358>

7 CONF

Redwood: A Frontier AI Accelerator Designed, Verified, and Deployed from Scratch in 2 Weeks by AI

ArXiv CS.AI 50%

arXiv:2608.26418v1 Announce Type: cross Abstract: Modern AI workloads and the hardware that runs them evolve on different timescales: architectural definition precedes volume silicon by years, while target workloads shift in months. Design decisions are therefore committed under...

<https://arxiv.org/abs/2608.26418>

8 CONF

SpeechGym: An Audio-Native Gym for Training Voice Agents via Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2608.26432v1 Announce Type: cross Abstract: Voice agents must call tools and hold multi-turn dialogue entirely through speech, yet the dominant paradigm trains them in text. Existing frameworks either cascade TTS and ASR around a proprietary voice API, where gradients...

<https://arxiv.org/abs/2608.26432>

9 CONF

Zero-Shot Self-Orchestration with Ledger-Based Control for Improved LLM Coding Performance

ArXiv CS.AI 50%

arXiv:2608.26480v1 Announce Type: cross Abstract: Multi-agent large language model systems are widely reported to beat single-model baselines, but the evidence is mixed, and comparisons are usually confounded: pipelines change token budgets, tool calls, and prompts...

<https://arxiv.org/abs/2608.26480>

00 CONF

Physics-Informed Stochastic Configuration Machine: A Backpropagation-Free Neural Network with Fast Training for Nonlinear Differential Equations

ArXiv CS.AI 50%

arXiv:2608.26549v1 Announce Type: cross Abstract: While Physics-Informed Neural Networks (PINNs) have emerged as a transformative paradigm for solving complex differential equations, their reliance on backpropagation-based gradient descent and automatic differentiation (AD)...

<https://arxiv.org/abs/2608.26549>

01 CONF

Risks and Controls for Multi-Agent Systems: an analytical framework for deployment of AI agents across organisational boundaries

ArXiv CS.AI 50%

arXiv:2608.26626v1 Announce Type: cross Abstract: This report presents a framework to help organisations, policymakers and researchers reason about the risks that emerge when AI agents interact with each other, how those risks change as interactions cross organisational...

<https://arxiv.org/abs/2608.26626>

02 CONF

PailitaoGR: Latent Think-with-Images for Generative Image Retrieval

ArXiv CS.AI 50%

arXiv:2608.26658v1 Announce Type: cross Abstract: Generative retrieval has demonstrated strong performance by directly generating product semantic identifiers (SIDs). Extending this paradigm to image search, however, is nontrivial because real-world query images contain diverse...

<https://arxiv.org/abs/2608.26658>

03 CONF

Do LLMs Understand Personality? Rethinking Persona Fidelity Evaluation through Structured Behavioral Inference

ArXiv CS.AI 50%

arXiv:2608.26674v1 Announce Type: cross Abstract: As large language models are increasingly deployed to simulate diverse human characters, ensuring persona fidelity, defined as the extent to which an agent's behavior consistently reflects the psychological and stylistic...

<https://arxiv.org/abs/2608.26674>

04 CONF

AesCanvas: A Large-Scale Dataset and Benchmark for Aesthetic Critique and Contextual Suitability

ArXiv CS.AI 50%

arXiv:2608.26713v1 Announce Type: cross Abstract: Recent advances in Multimodal Large Language Models (MLLMs) have extended Image Aesthetic Assessment (IAA) beyond scalar scores toward interpretable critique and guidance. Yet existing benchmarks mainly assess intrinsic visual...

<https://arxiv.org/abs/2608.26713>

05 CONF

Beyond Execution: Auditing Experimental Fidelity in LLM-Driven Scientific Research

ArXiv CS.AI 50%

arXiv:2608.26753v1 Announce Type: cross Abstract: LLM agents used for scientific experimentation must do more than generate executable code: they must implement the reference method faithfully, design experiments that test the paper's claims, and provide evidence supporting...

<https://arxiv.org/abs/2608.26753>

06 CONF

PLCBench: Can Autonomous LLM Agents Turn PLC Access into Sustained Physical Impact?

ArXiv CS.AI 50%

arXiv:2608.26882v1 Announce Type: cross Abstract: Industrial control systems (ICSs) rely on programmable logic controllers (PLCs) to connect networked computation with physical control. Tool-using large language model (LLM) agents represent an emerging attack threat: can an...

<https://arxiv.org/abs/2608.26882>

07 CONF

When Memory Takes Gradients: Collaborative Vector Memory for Agentic Recommender Systems

ArXiv CS.AI 50%

arXiv:2608.26895v1 Announce Type: cross Abstract: Agentic recommender systems ground each decision of a large language model (LLM) in a persistent memory of the user, and in existing agents that memory is text: a narrative written and maintained by further LLM calls. Text limits...

<https://arxiv.org/abs/2608.26895>

08 CONF

Per-View Gaussian Predictions Enable Training-Free Distractor Filtering in Feed-Forward 3DGS

ArXiv CS.AI 50%

arXiv:2608.26951v1 Announce Type: cross Abstract: Feed-forward 3D Gaussian Splatting reconstructs an explicit Gaussian representation from multiple input images in one network execution, making 3D reconstruction increasingly accessible for casual captures. However, such captures...

<https://arxiv.org/abs/2608.26951>

09 CONF

Temporally-Grounded Language Generation: Towards Real-Time Vision-Language Models

ArXiv CS.AI 50%

arXiv:2505.11326v2 Announce Type: replace-cross Abstract: Vision-language models (VLMs) have shown remarkable progress in offline tasks such as image captioning and video question answering. However, real-time interactive environments impose new demands on VLMs, requiring them...

<https://arxiv.org/abs/2505.11326>

10 CONF

Learning Transverse Momentum Distributions from Raw Scattering Events via Conditional Diffusion

ArXiv CS.AI 50%

arXiv:2608.27077v1 Announce Type: cross Abstract: Extracting transverse momentum dependent parton distribution functions (TMD PDFs) from semi-inclusive deep inelastic scattering (SIDIS) data is a central goal of the nucleon structure program at Jefferson Lab and the future...

<https://arxiv.org/abs/2608.27077>

11 CONF

Active Diffusion-Based Inference for Ill-Posed Inverse Problems under Incomplete Priors

ArXiv CS.AI 50%

arXiv:2608.27080v1 Announce Type: cross Abstract: Many scientific and engineering applications require estimating unknown parameters from experimentally observable data -- an inverse problem that is inherently challenging due to nonlinearity, noise, and ill-posedness. In this...

<https://arxiv.org/abs/2608.27080>

12 CONF

ANTShapes Benchmarking Datasets for Event-Based Neuromorphic Object Classification

ArXiv CS.AI 50%

arXiv:2608.27150v1 Announce Type: cross Abstract: Object classification in event-based computer vision is a task that is attracting considerable research attention. Event-based object classification is a fundamental task in the fields of security and applied computer vision,...

<https://arxiv.org/abs/2608.27150>

13 CONF

PACE: A Unified Condense-and-Extract Paradigm for Fast VLM Inference

ArXiv CS.AI 50%

arXiv:2608.27206v1 Announce Type: cross Abstract: Vision-Language Models (VLMs) demonstrate exceptional visual reasoning capabilities, yet their inference costs escalate rapidly with the proliferation of visual tokens. Existing visual token pruning methods exhibit two...

<https://arxiv.org/abs/2608.27206>

14 CONF

STEP: State-Aware Task Estimation and Planning with Multi-Modal LLMs for Human-Robot Collaboration

ArXiv CS.AI 50%

arXiv:2608.27225v1 Announce Type: cross Abstract: Effective human-robot collaboration in industrial settings requires robots to understand human intentions and assist with task planning, reducing workload. Recent works have explored the use of Multi-modal Large Language Models...

<https://arxiv.org/abs/2608.27225>

15 CONF

Difference-in-Differences on a Censored Rating Scale Can Manufacture an Effect: Evidence from a Pre-Registered LLM-Judge Audit

ArXiv CS.AI 50%

arXiv:2608.27309v1 Announce Type: cross Abstract: Audits of LLM judges certify a bias by contrasting matched conditions, and the strongest designs difference twice: a within-item contrast between two candidate responses, differenced again across a manipulated attribute, read off...

<https://arxiv.org/abs/2608.27309>

16 CONF

Successive Capacity Growth: Task-Complexity-Driven Width and Depth Expansion for Vision Transformer Encoders in JEPA World Models

ArXiv CS.AI 50%

arXiv:2608.27367v1 Announce Type: cross Abstract: Joint-Embedding Predictive Architectures (JEPAs) for world modeling typically employ fixed-size Vision Transformer encoders that are over-provisioned for simple tasks and under-provisioned for complex ones, with significant...

<https://arxiv.org/abs/2608.27367>

17 CONF

LeVJEPa: Efficient & Scalable Video Pretraining without the Heuristics

ArXiv CS.AI 50%

arXiv:2608.27395v1 Announce Type: cross Abstract: Video carries the temporal structure of the physical world, yet learning representations from it has remained computationally expensive: prevailing self-supervised methods either prevent representation collapse through...

<https://arxiv.org/abs/2608.27395>

18 CONF

Making Clinical Language Models Auditable: Concept-Guided Fine-Tuning for Robust Prediction

ArXiv CS.AI 50%

arXiv:2608.27397v1 Announce Type: cross Abstract: Clinical language models can achieve strong in-hospital accuracy yet fail under deployment shifts because they exploit note-specific artifacts (e.g., templates, separators, boilerplate) that do not reflect patient state. We...

<https://arxiv.org/abs/2608.27397>

19 CONF

Persona-Execution Separation: An Architecture Pattern for Evolving LLM Agents under Execution Audit

ArXiv CS.AI 50%

arXiv:2608.27427v1 Announce Type: cross Abstract: Large language model (LLM) agents in governed organizations must let the persona (instructions, tone, self-presentation) evolve freely, while keeping execution (stateful, audited work) traceable. A single trust domain does not...

<https://arxiv.org/abs/2608.27427>

20 CONF

From Static to Dynamic: Benchmarking Real-World Code Review with MCR-Bench

ArXiv CS.AI 50%

arXiv:2608.27442v1 Announce Type: cross Abstract: In real-world software development, code review typically involves iterative interactions between developers and reviewers to improve software quality, making the process costly and time-consuming. Although recent work explores...

<https://arxiv.org/abs/2608.27442>

21 CONF

LLM-Powered Swarms: A New Frontier or a Conceptual Stretch?

ArXiv CS.AI 50%

arXiv:2506.14496v3 Announce Type: replace Abstract: Swarm intelligence describes how simple, decentralized agents can collectively produce complex behaviors. Recently, the concept of swarming has been extended to large language model (LLM)-powered systems, such as OpenAI's Swarm...

<https://arxiv.org/abs/2506.14496>

22 CONF

Pushing the Envelope of LLM Inference with Ultra-Low-Bit Quantized Models

ArXiv CS.AI 50%

arXiv:2508.06753v3 Announce Type: replace Abstract: The advent of ultra-low-bit LLM models, approaching the perplexity and task accuracy of their full precision counterparts, is ushering in a new era of LLM inference. While these advances promise models that are cost-effective...

<https://arxiv.org/abs/2508.06753>

23 CONF

Do Language Models Follow Occam's Razor? An Evaluation of Parsimony in Inductive and Abductive Reasoning

ArXiv CS.AI 50%

arXiv:2509.03345v3 Announce Type: replace Abstract: Non-deductive reasoning, encompassing inductive and abductive reasoning, is essential in addressing complex real-world questions. One key feature of inductive and abductive reasoning is that there are many valid hypotheses; the...

<https://arxiv.org/abs/2509.03345>

24 CONF

Buried in Textual Debt: Context Pruning with Visual Evidence Preservation for MLLM Agents

ArXiv CS.AI 50%

arXiv:2608.22963v2 Announce Type: replace Abstract: Multimodal Large Language Models (MLLMs) are increasingly deployed as multi-step agents, where explicit reasoning supports task decomposition and tool coordination but also accumulates self-generated text. Over long...

<https://arxiv.org/abs/2608.22963>

25 CONF

MENTOR: Reinforcement Learning via Flexible Teacher-Optimized Rewards for Tool-Use Distillation

ArXiv CS.AI 50%

arXiv:2510.18383v4 Announce Type: replace-cross Abstract: Distilling the tool-use capabilities of large language models (LLMs) into small language models (SLMs) is essential for their practical application. The predominant approach, supervised fine-tuning (SFT), is an off-policy...

<https://arxiv.org/abs/2510.18383>

26 CONF

The BS-meter: Detecting Politics and Labour through ChatGPT's Language

ArXiv CS.AI 50%

arXiv:2411.15129v3 Announce Type: replace-cross Abstract: What can we learn about language from studying how it is used by ChatGPT and other large language model (LLM)-based chatbots? In this paper, we analyse the distinctive character of language generated by ChatGPT, in...

<https://arxiv.org/abs/2411.15129>

27 CONF

Communication styles and reader preferences of LLM- and human-authored COVID-19 information explanations: a case study

ArXiv CS.AI 50%

arXiv:2505.08143v2 Announce Type: replace-cross Abstract: With the wide adoption of large language models (LLMs) in information assistance, it is essential to examine their alignment with human communication styles and values. We situate this study within health fact-checking,...

<https://arxiv.org/abs/2505.08143>

28 CONF

HybridProver: Augmenting Theorem Proving with LLM-Driven Proof Synthesis and Refinement

ArXiv CS.AI 50%

arXiv:2505.15740v2 Announce Type: replace-cross Abstract: Formal methods play a crucial role in ensuring the reliability of critical systems through rigorous mathematical verification. However, their adoption remains limited due to the labor-intensive nature of manual proof...

<https://arxiv.org/abs/2505.15740>

29 CONF

Residual Reward Models: Leveraging Prior Knowledge for Efficient Preference-based Reinforcement Learning in Robotics

ArXiv CS.AI 50%

arXiv:2507.00611v2 Announce Type: replace-cross Abstract: Preference-based Reinforcement Learning (PbRL) provides a promising alternative to heuristic reward design in complex robotic environments. However, PbRL often suffers from poor sample efficiency, requiring extensive and...

<https://arxiv.org/abs/2507.00611>

30 CONF

Distinct Profiles of Run-to-Run Score Reliability and Expert-Panel Alignment Across Four LLM Evaluators of Simulated Japanese-Language AI-to-AI Counseling

ArXiv CS.AI 50%

arXiv:2507.02950v4 Announce Type: replace-cross Abstract: Large language models (LLMs) increasingly evaluate generated dialogue, but repeatable scores do not necessarily align with professional judgment. This observational fixed-benchmark study compared four configured LLM...

<https://arxiv.org/abs/2507.02950>

31 CONF

AirLLM: Diffusion Policy-based Adaptive LoRA for Remote Fine-Tuning of LLM over the Air

ArXiv CS.AI 50%

arXiv:2507.11515v2 Announce Type: replace-cross Abstract: Operating Large Language Models (LLMs) on edge devices is increasingly challenged by limited communication bandwidth and strained computational and memory costs. Thus, cloud-assisted remote fine-tuning becomes...

<https://arxiv.org/abs/2507.11515>

32 CONF

GSM8K-V: Can Vision Language Models Solve Grade School Math Word Problems in Visual Contexts

ArXiv CS.AI 50%

arXiv:2509.25160v2 Announce Type: replace-cross Abstract: Mathematical reasoning is a key capability for vision-language models (VLMs), yet current benchmarks mainly evaluate text-based or explicitly symbolic visual inputs. It remains unclear whether VLMs can reason...

<https://arxiv.org/abs/2509.25160>

33 CONF

MCCE: A Framework for Multi-LLM Collaborative Search in Discrete Spaces with Similarity-Filtered Preference Learning

ArXiv CS.AI 50%

arXiv:2510.06270v2 Announce Type: replace-cross Abstract: Multi-objective discrete optimization problems, such as molecular design, pose significant challenges due to their vast and unstructured combinatorial spaces. Traditional evolutionary algorithms often get trapped in local...

<https://arxiv.org/abs/2510.06270>

34 CONF

What the "Spotless" Mind Remembers: How Knowledge Entanglement Shapes What Leaks After Unlearning in LLMs

ArXiv CS.AI 50%

arXiv:2510.25732v2 Announce Type: replace-cross Abstract: Unlearning in large language models (LLMs) is usually evaluated as whether an "unlearned" fact can be recovered. We instead ask whether a fact's structural entanglement with the rest of a model's knowledge predicts...

<https://arxiv.org/abs/2510.25732>

35 CONF

Multivariate Diffusion Transformer with Decoupled Attention for High-Fidelity Mask-Text Collaborative Facial Generation

ArXiv CS.AI 50%

arXiv:2511.12631v3 Announce Type: replace-cross Abstract: While significant progress has been achieved in multimodal facial generation using semantic masks and textual descriptions, conventional feature fusion approaches often fail to enable effective cross-modal interactions,...

<https://arxiv.org/abs/2511.12631>

36 CONF

Diagnosing Conformal Prediction Failures Under Distribution Shift: A COVID-19 Case Study

ArXiv CS.AI 50%

arXiv:2601.00908v2 Announce Type: replace-cross Abstract: Conformal prediction provides distribution-free coverage guarantees, but these degrade under distribution shift - and practitioners lack tools to anticipate which deployed models will fail before observing test data. We...

<https://arxiv.org/abs/2601.00908>

37 CONF

CounterVid: Counterfactual Video Generation for Mitigating Action and Temporal Hallucinations in Video-Language Models

ArXiv CS.AI 50%

arXiv:2601.04778v2 Announce Type: replace-cross Abstract: Video-language models (VLMs) achieve strong multimodal understanding but remain prone to hallucinations, especially when reasoning about actions and temporal order. Existing mitigation strategies, such as textual...

<https://arxiv.org/abs/2601.04778>

38 CONF

SynthCharge: An Electric Vehicle Routing Instance Generator with Feasibility Screening to Enable Learning-Based Optimization and Benchmarking

ArXiv CS.AI 50%

arXiv:2603.03230v2 Announce Type: replace-cross Abstract: The electric vehicle routing problem with time windows (EVRPTW) extends the classical VRPTW by introducing battery capacity constraints and charging station decisions. Existing benchmark datasets are often static and lack...

<https://arxiv.org/abs/2603.03230>

39 CONF

Grounded Token Initialization for New Vocabulary in LMs for Generative Recommendation

ArXiv CS.AI 50%

arXiv:2604.02324v2 Announce Type: replace-cross Abstract: Language models (LMs) are increasingly extended with new learnable vocabulary tokens for domain-specific tasks, such as Semantic-ID tokens in generative recommendation. The standard practice initializes these new tokens...

<https://arxiv.org/abs/2604.02324>

40 CONF

A Comprehensive Comparison of Deep Learning Architectures for COVID-19 Classification on CT & X-ray Imagery

ArXiv CS.AI 50%

arXiv:2605.20445v2 Announce Type: replace-cross Abstract: COVID-19 was a significant challenge that led to the loss of numerous lives daily. Not only a certain country was involved in this outbreak, but even the world has suffered because of the coronavirus. Imaging techniques...

<https://arxiv.org/abs/2605.20445>

41 CONF

TriShieldRAG: 3 Rings, One Blind Spot in Layered Defenses for Retrieval-Augmented Generation

ArXiv CS.AI 50%

arXiv:2607.23838v2 Announce Type: replace-cross Abstract: Retrieval-Augmented Generation (RAG) grounds LLM answers in query-time retrieved documents, so reliability depends on what the retriever returns. PoisonedRAG (Zou et al., USENIX Security'25) showed five crafted documents...

<https://arxiv.org/abs/2607.23838>

42 CONF

Evidence-Grounded Trustworthy Multimodal Reasoning and Evaluation Benchmark in Complex Urban Scenes

ArXiv CS.AI 50%

arXiv:2608.10954v2 Announce Type: replace-cross Abstract: While Multimodal Large Language Models (MLLMs) demonstrate impressive performance in benign scenarios, their cognitive reliability deteriorates significantly in complex scenes under adverse conditions. In these settings,...

<https://arxiv.org/abs/2608.10954>

43 CONF

Complexity Induction: Compositional Generalization via Structured Training Distortion

ArXiv CS.AI 50%

arXiv:2608.21464v2 Announce Type: replace-cross Abstract: We demonstrate that structured distortion of training data - which we term complexity induction - can induce compositional generalization in a standard CNN classifier without architectural modification. Using synthetic...

<https://arxiv.org/abs/2608.21464>

44 CONF

Language Chain in Alignment: Cross-lingual Ranking Preference Optimization

ArXiv CS.AI 50%

arXiv:2608.23149v2 Announce Type: replace-cross Abstract: The alignment of Large Language Models heavily relies on English-centric high-quality preference data, which often leads to suboptimal performance in other languages. In this paper, we propose Cross-lingual Ranking...

<https://arxiv.org/abs/2608.23149>

45 CONF

CAT-GS: Balanced Multimodal Learning via Calibrated Gating and Fusion Surgery

ArXiv CS.AI 50%

arXiv:2608.24947v2 Announce Type: replace-cross Abstract: End-to-end training of multimodal neural networks often exhibits unstable neural dynamics characterized by three coupled failure modes that degrade learning: (i) modality imbalance, where one branch dominates...

<https://arxiv.org/abs/2608.24947>

46 CONF

When Stale Constraints Go Unchecked: Budgeted Verification Failures in Inherited Agent Memory

ArXiv CS.AI 50%

arXiv:2608.25553v2 Announce Type: replace-cross Abstract: An agent that inherits a consolidated memory may inherit a constraint that was true when written and has since been withdrawn by a newer authoritative record. Under a scarce verification budget, does the agent recover the...

<https://arxiv.org/abs/2608.25553>