



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	243	Media
TechCrunch AI	8	Media
HackerNews	6	Community
The Verge AI	5	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Evaluating the Hidden Costs of Personalization in Large Language Models

ArXiv CS.AI

+1 source

80%

arXiv:2608.28833v1 Announce Type: new Abstract: While Large language models (LLMs) incorporate user personalization signals to improve usability and helpfulness, they increasingly shift from providing balanced, informative responses toward optimizing for user satisfaction when...

<https://arxiv.org/abs/2608.28833>

VERI

Look It Up: Analysing Internal Web Search Capabilities of Modern LLMs

ArXiv CS.AI

+1 source

80%

arXiv:2511.18931v2 Announce Type: replace-cross Abstract: Modern large language models increasingly integrate internal web-based retrieval to provide real-time answers, yet it remains unclear how effectively these systems identify information need, trigger retrieval, and use...

<https://arxiv.org/abs/2511.18931>

VERI

HumanStudy-Bench: Towards AI Agent Design for Participant Simulation

ArXiv CS.AI

+1 source

80%

arXiv:2602.00685v2 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used as simulated participants in social science experiments, but their behavior is often unstable and highly sensitive to design choices. Prior evaluations frequently conflate base...

<https://arxiv.org/abs/2602.00685>

CONFIRMED SIGNALS — 257 stories

CONF

Show HN: Saccade - Live semantic browser truth for AI agents

HackerNews

50%

HN 5

I build one this application is to resolve the most of the agent cannot handle the brother they extremely slow. So I come up with this idea, we install one of the extntion on chrome or edge to give continuously compile the tabs authorized by the user into semantically meaningful...

<https://github.com/nanlogic/saccade>

CONF

Show HN: Guess the AI model from its web design

HackerNews

50%

HN 5

No summary — see link for full article.

<https://raq.com/blog/can-you-spot-ai-web-design/quiz>

CONF

Show HN: Xyzzy - AI teamwork in one Python process, with a tamper-evident log

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/Project-Nexus-YR/XYZZY>

CONF

Show HN: Claude Code Arcade

HackerNews 50% HN 2

This week-end, I wondered if we could add a little fun animation in Claude Code CLI's status line on macOS, based on the LLM current work. First I created a fish tank, then a safari scene, only using emojis. Being an 80's kid, I challenged Fable 5 to add a Ms Pac-Man, with good...

<https://github.com/jystervinou/claude-code-arcade>

CONF

Show HN: VajraClaw - Deterministic <1μs execution guardrail for AI agents

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/Top-Celestial-Company-Ltd/DROS-VajraClaw-Hacker>

CONF

Show HN: Sonde, a local code graph for AI agents that refuses to guess

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/anishmoncivarghese/sonde>

0 CONF

Apple shares 'shocking evidence' against former employee accused of stealing company data for OpenAI

TechCrunch AI 50%

Apple says it has evidence that a former employee destroyed evidence of data theft after learning he was under investigation.

<https://techcrunch.com/2026/08/31/apple-shares-shocking-evidence-against-former-employee...>

1 CONF

Instagram puts new limits on undisclosed AI profiles

TechCrunch AI 50%

As frustration over AI influencers has been growing, Instagram is limiting the reach of undisclosed AI profiles.

<https://techcrunch.com/2026/08/31/instagram-puts-new-limits-on-undisclosed-ai-profiles/>

2 CONF

Harvard Law dropout raises \$6M for Blue Voice to build a 'Harvey for police officers'

TechCrunch AI 50%

Blue Voice is trained on department-specific laws, local ordinances, protocols, and guidelines that general-purpose AI tools can't access on the public internet.

<https://techcrunch.com/2026/08/31/harvard-law-dropout-raises-6m-for-blue-voice-to-build...>

3 CONF

Nvidia's \$3.5B MediaTek bet reveals its plan for tackling Big Tech's AI chip buildout

TechCrunch AI 50%

Nvidia invests \$3.5 billion into Taiwanese chipmaker MediaTek. The deal shows how Nvidia plans to stay essential to AI infrastructure as Big Tech begins to build its own AI chips.

<https://techcrunch.com/2026/08/31/nvidias-3-5b-mediatek-bet-reveals-its-plan-for-tackli...>

4 CONF

Sony Music, Warner sue Anthropic, alleging a 'brazen campaign' of intellectual property theft

TechCrunch AI 50%

This latest lawsuit is particularly broad and homes in on accusations of illegal piracy.

<https://techcrunch.com/2026/08/29/sony-music-warner-sue-anthropic-alleging-a-brazen-cam...>

5 **CONF****At TechBBQ, Europe's AI conversations kept coming back to: Who's actually in control?****TechCrunch AI** 50%

Investors, founders, and operators from across Europe arrived for the annual Nordic TechBBQ conference to talk about how humans can have agency over AI.

<https://techcrunch.com/2026/08/29/at-techbbq-europes-ai-conversations-kept-coming-back-...>

6 **CONF****The Hugging Face hack could indicate cultural issues at OpenAI****MIT Tech Review** 50%

This story originally appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. By now you've probably heard about last month's major AI security incident, in which OpenAI agents escaped their sandbox and hacked into the...

<https://www.technologyreview.com/2026/08/31/1143180/hugging-face-hack-could-indicate-cu...>

7 **CONF****Debian won't ban AI code from its Linux distribution****The Verge AI** 50%

Debian voted to allow developers to use AI tools in their contributions to the Linux distribution's "development, maintenance, [and] documentation." The new policy on AI acknowledges that "responsible" use of AI can improve developers' productivity, and goes on to say,...

<https://www.theverge.com/tech/986789/linux-debian-generative-ai-policy>

8 **CONF****New York Governor Kathy Hochul thinks AI should be 'less evil'****The Verge AI** 50%

Today, I'm talking with New York Governor Kathy Hochul, and I'll just warn you — this episode moves really fast. It's an election year, after all, with a shocking amount of tech policy at stake, and Governor Hochul has taken strong positions on almost every major tech issue...

<https://www.theverge.com/podcast/986661/ny-gov-kathy-hochul-ai-data-centers-ban-3d-prin...>

9 **CONF****ChatGPT to face tougher regulation in the EU****The Verge AI** 50%

OpenAI will soon be held accountable for mitigating risks related to ChatGPT's impact on minors, user mental health, and the spread of illegal content in the European Union. That's because ChatGPT is now considered a Very Large Online Search Engine under the EU's Digital...

<https://www.theverge.com/ai-artificial-intelligence/986682/openai-chatgpt-eu-dsa>

0 **CONF****Instagram cracks down on AI accounts pretending to be human****The Verge AI** 50%

Instagram is finally taking steps to address the rise of fake AI-influencer accounts that have gotten harder to spot. It's also renaming the "AI creator" label to "AI-generated profile" to make it clear when a profile features an AI-generated person that's not a real human...

<https://www.theverge.com/tech/986593/instagram-addresses-fake-ai-profile-slop>

1 **CONF****EviAnchor: Mitigating Hallucinations in Large Vision-Language Models via Regional Visual Evidence Compensation****ArXiv CS.AI** 50%

arXiv:2608.29092v1 Announce Type: new Abstract: Large vision-language models (LVLMs) frequently generate content unsupported by visual inputs. Preliminary experiments show that visual evidence is primarily incorporated into answer-side representations in early-to-middle decoder...

<https://arxiv.org/abs/2608.29092>

- 2

CONF

From Question-First to Analyst-First: Domain-Expert Skills and Verified Knowledge Compilation for Proactive Enterprise Analytics

ArXiv CS.AI

50%

arXiv:2608.28594v1 Announce Type: new Abstract: Conversational analytics systems assume the user already has a well-formed question, leaving a non-expert facing a blank query box on an unfamiliar enterprise schema. Commercial 'proactive' tools narrow this gap only by detecting...

<https://arxiv.org/abs/2608.28594>
- 3

CONF

The Race between Agentic AI Capabilities and Data Quality Control in Online Surveys

ArXiv CS.AI

50%

arXiv:2608.28597v1 Announce Type: new Abstract: Online surveys are a foundational data collection instrument in a variety of fields, with attention checks serving as critical guardians of response quality. However, the rapid emergence of agentic AI (goal directed systems powered...

<https://arxiv.org/abs/2608.28597>
- 4

CONF

SHAPE of Chain-of-Thought in Math Reasoning

ArXiv CS.AI

50%

arXiv:2608.28600v1 Announce Type: new Abstract: Large language models (LLMs) achieve strong performance on mathematical reasoning benchmarks, yet the mathematically meaningful skills underlying their reasoning remain underexplored. We introduce \texttt{SHAPE}, a framework that...

<https://arxiv.org/abs/2608.28600>
- 5

CONF

Leveraging Generative AI to Design Accessible Interactive Visualizations for Undergraduate Mathematics: A Six-Phase Workflow

ArXiv CS.AI

50%

arXiv:2608.28601v1 Announce Type: new Abstract: Interactive visualizations support conceptual understanding in undergraduate mathematics, but building them has required programming expertise most instructors lack. Using a design-based research approach, we develop, deploy, and...

<https://arxiv.org/abs/2608.28601>
- 6

CONF

C3-UniMM: Causal Cycle-Consistent Unified Multimodal Modeling via Super Alignment and Shared Decoding Space

ArXiv CS.AI

50%

arXiv:2608.28603v1 Announce Type: new Abstract: Unified Multimodal Models aim to achieve any-to-any understanding and generation across arbitrary modalities. However, existing methods primarily rely on modeling implicit statistical correlations and lack cross-modal structural...

<https://arxiv.org/abs/2608.28603>
- 7

CONF

RegDivergence-101: An LLM Benchmark for Cross-Jurisdiction Regulatory Contradiction Detection in Life Sciences

ArXiv CS.AI

50%

arXiv:2608.28607v1 Announce Type: new Abstract: Pharmaceutical sponsors developing a drug for both the United States and the European Union must reconcile guidance issued independently by the FDA and the EMA. Where the two agencies require substantively the same thing, a sponsor...

<https://arxiv.org/abs/2608.28607>

8 CONF

InternReviewer & InternAdvocate: Objective Reward and Evaluation for Agentic Reinforcement Learning in Peer Review and Rebuttal

ArXiv CS.AI 50%

arXiv:2608.28612v1 Announce Type: new Abstract: Generating professional scholarly content, such as peer reviews and rebuttals, requires an intricate synergy between domain reasoning and factual grounding. This work presents a comprehensive framework for the development and...

<https://arxiv.org/abs/2608.28612>

9 CONF

Machine Learning-Enhanced Tabu Search for Tactical Wireless Network Design

ArXiv CS.AI 50%

arXiv:2608.28627v1 Announce Type: new Abstract: Designing high-performance tactical wireless networks under realistic operational constraints gives rise to challenging combinatorial optimization problems, where the evaluation of candidate solutions relies on detailed physical...

<https://arxiv.org/abs/2608.28627>

10 CONF

AI Scientist Mission Control (AIMC): Visual Analytics for Human Oversight of Autonomous Scientific Discovery

ArXiv CS.AI 50%

arXiv:2608.28637v1 Announce Type: new Abstract: Autonomous scientific discovery systems can generate large numbers of research ideas, experiments, and manuscripts with minimal human intervention. As these systems become increasingly capable, scientists require effective...

<https://arxiv.org/abs/2608.28637>

1 CONF

From Extraction to Governed Memory: Multi-Agent Knowledge Graph Construction with Domain-Expert Review

ArXiv CS.AI 50%

arXiv:2608.28642v1 Announce Type: new Abstract: Knowledge graphs used by agentic systems are often treated as flat stores of extracted triples, with little record of who owns a fact, why it was admitted, or how it should be used downstream. We argue that reliable agentic...

<https://arxiv.org/abs/2608.28642>

2 CONF

BiasMix-Finance: Post-Generation KYC Guardrails for LLM Portfolio Advice

ArXiv CS.AI 50%

arXiv:2608.28646v1 Announce Type: new Abstract: Large language models (LLMs) can generate plausible-sounding ETF portfolios while silently violating basic KYC-style constraints on risk, fees, and diversification. This is especially problematic in agentic multi-turn advisory...

<https://arxiv.org/abs/2608.28646>

3 CONF

Self-Specialized Teachers for Domain Post-Training

ArXiv CS.AI 50%

arXiv:2608.28647v1 Announce Type: new Abstract: Target-only post-training can improve performance in a specialized domain while degrading behaviors that a general-purpose base model acquired before adaptation. We study this problem when target-domain data are available but a...

<https://arxiv.org/abs/2608.28647>

- 4
CONF

A Unifying Perspective on Language Model Representations: From Filler-Role Structure to Mechanistic Interpretability

ArXiv CS.AI 50%

arXiv:2608.29034v1 Announce Type: cross Abstract: A wide range of methods have been proposed for interpreting language models, delivering important insights into their inner workings. However, different methods and their resulting insights stand in relative isolation: what could...

<https://arxiv.org/abs/2608.29034>
- 5
CONF

A Generalized Optimization Engine (GOE) for Edge AI Inference Acceleration

ArXiv CS.AI 50%

arXiv:2608.28652v1 Announce Type: new Abstract: Artificial intelligence (AI) models have demonstrated remarkable capabilities across various domains, yet their widespread deployment is impeded by significant computational costs, particularly on resource-constrained devices. This...

<https://arxiv.org/abs/2608.28652>
- 6
CONF

FRAMEWORKERS: A Dynamic Multi-Agent Framework for AI-Generated Video Production

ArXiv CS.AI 50%

arXiv:2608.29814v1 Announce Type: new Abstract: Modern video generators excel at synthesizing individual clips, but complete video production requires coordinating a long sequence of interdependent creative steps, including scripting, storyboarding, generation, and editing. It...

<https://arxiv.org/abs/2608.29814>
- 7
CONF

ProgRouter: Online Progress-Guided Orchestration for Multi-Agent LLM Workflows under Quality-Cost Tradeoffs

ArXiv CS.AI 50%

arXiv:2608.25992v2 Announce Type: replace Abstract: Multi-agent large language model (LLM) workflows have emerged as a powerful paradigm for solving complex, open-ended tasks through collaborative reasoning among specialized LLM agents, but they incur substantial operating costs...

<https://arxiv.org/abs/2608.25992>
- 8
CONF

PermitGPT: A Unified Generative-AI Pipeline for Construction Hazard Forecasting, Permit Prediction, and Community Impact

ArXiv CS.AI 50%

arXiv:2608.28728v1 Announce Type: new Abstract: Urban construction governance requires early decisions that connect workplace safety, permitting requirements, and community impact, yet the relevant evidence is often scattered across separate municipal and regulatory data...

<https://arxiv.org/abs/2608.28728>
- 9
CONF

Frequency Selective Neural Networks as a Foundation Architecture for Time Series Learning

ArXiv CS.AI 50%

arXiv:2608.29012v1 Announce Type: new Abstract: Time-series data across physical and biological domains are fundamentally driven by complex, non-stationary oscillatory modes. While deep learning models, such as Convolutional Neural Networks (CNNs), Recurrent Neural Networks, and...

<https://arxiv.org/abs/2608.29012>

0 **CONF****Gurukul AI: An Interactive AI-Driven Educational Platform for Indian Education System****ArXiv CS.AI** **50%**

arXiv:2608.28611v1 Announce Type: cross Abstract: Recent advances in large language models (LLMs) like ChatGPT and LLaMA have transformed AI-driven education, but these systems are predominantly trained on Western-centric data, making them ill-suited for regional curricula like...

<https://arxiv.org/abs/2608.28611>

1 **CONF****MineCEraft: Evaluating Language Models as Construction Engineers in the World of Minecraft****ArXiv CS.AI** **50%**

arXiv:2608.28884v1 Announce Type: new Abstract: We introduce MineCEraft (Minecraft Construction Engineering Benchmark, pronounced mine-see-ee-raft), an easy-to-use, open-source benchmark designed to systematically evaluate the reliability and limitations of LLMs for construction...

<https://arxiv.org/abs/2608.28884>

2 **CONF****Automated Researchers Can Reliably Mitigate Alignment Failures****ArXiv CS.AI** **50%**

arXiv:2608.28945v1 Announce Type: new Abstract: Automating alignment research may accelerate progress toward aligned AI, but whether it does is hard to measure. Luckily, many alignment failures, such as deception, sycophancy, and jailbreaks, are already measurable by public...

<https://arxiv.org/abs/2608.28945>

3 **CONF****The Role of Network Topology and Opponent Information in Shaping Cooperation in Multi-Agent Reinforcement Learning Systems****ArXiv CS.AI** **50%**

arXiv:2608.28977v1 Announce Type: new Abstract: Several works have investigated the influence of graph topology on cooperation among artificial agents, while the majority of the literature has focused on modelling agents' adaptation through strategy imitation, which relies...

<https://arxiv.org/abs/2608.28977>

4 **CONF****Agentic AI uncovers conserved cross-tissue protein co-abundance programs inaccessible to single-dataset analysis****ArXiv CS.AI** **50%**

arXiv:2608.28990v1 Announce Type: new Abstract: Protein co-abundance clusters preserved across tissues can reveal shared disease mechanisms and candidate therapeutic targets, particularly when proteins implicated in organ-confined diseases converge in peripheral or accessible...

<https://arxiv.org/abs/2608.28990>

5 **CONF****Facts Without Rules: Boundary Metadata Collapse in Multi-Agent LLM Handoffs****ArXiv CS.AI** **50%**

arXiv:2608.29028v1 Announce Type: new Abstract: Multi-agent LLM systems often coordinate by compressing an upstream interaction into a handoff artifact that downstream agents treat as shared state. We show that this handoff step is a structural source of privacy leakage:...

<https://arxiv.org/abs/2608.29028>

6 CONF

EmoLASP: Emotion Recognition with Language Models and Answer Set Programming

ArXiv CS.AI 50%

arXiv:2608.29035v1 Announce Type: new Abstract: Emotion recognition in conversations is increasingly tackled with language models, but these models can be unstable and expensive to fine-tune or to prompt with long dialogue histories. We propose EmoLASP, a framework that combines...

<https://arxiv.org/abs/2608.29035>

7 CONF

PQMass: Probabilistic Assessment of the Quality of Generative Models using Probability Mass Estimation

ArXiv CS.AI 50%

arXiv:2402.04355v4 Announce Type: replace-cross Abstract: We propose a likelihood-free method for comparing two distributions given samples from each, with the goal of assessing the quality of generative models. The proposed approach, PQMass, provides a statistically rigorous...

<https://arxiv.org/abs/2402.04355>

8 CONF

Revolutionizing Turn-by-Turn Navigation with Cloud-Edge Deep Learning

ArXiv CS.AI 50%

arXiv:2608.29073v1 Announce Type: new Abstract: Turn-by-turn (TBT) navigation systems are integral to modern driving experiences, providing real-time audio instructions to guide drivers safely to destinations. However, existing audio instruction policy often relies on rule-based...

<https://arxiv.org/abs/2608.29073>

9 CONF

SGM: Safety Glasses for Multimodal Large Language Models via Neuron-Level Detoxification

ArXiv CS.AI 50%

arXiv:2512.15052v4 Announce Type: replace-cross Abstract: Disclaimer: Samples in this paper may be harmful and cause discomfort. Multimodal large language models (MLLMs) enable multimodal understanding but inherit toxic signals from weakly curated pretraining corpora, leading to...

<https://arxiv.org/abs/2512.15052>

0 CONF

Denoising as Projection: Constrained Optimization with Gradient-Guided Diffusion

ArXiv CS.AI 50%

arXiv:2608.29507v1 Announce Type: cross Abstract: Diffusion models are increasingly used not only for sampling from learned data distributions, but also for generating samples that optimize task-specific objectives. A common approach is to guide the reverse diffusion process...

<https://arxiv.org/abs/2608.29507>

1 CONF

Emergent Misalignment Is Not Magical

ArXiv CS.AI 50%

arXiv:2608.29118v1 Announce Type: new Abstract: Fine-tuning large language models (LLMs) on narrowly harmful datasets can lead to misalignment broadly, a phenomenon known as emergent misalignment (EM). EM poses a challenge for AI safety and our understanding of LLMs. Prior work...

<https://arxiv.org/abs/2608.29118>

2 **CONF****BLOOM-WILT: Logit Tilting for Behaviour Elicitation in Automated LLM Auditing****ArXiv CS.AI** **50%**

arXiv:2608.31105v1 Announce Type: new Abstract: Users of a deployed language model routinely encounter behaviours that testing almost never surfaces, since deployment puts the model through orders of magnitude more interactions than any evaluation can simulate. Automated...

<https://arxiv.org/abs/2608.31105>

3 **CONF****Concepts Whisper: Spectral Anti-Concentration and the Dual Geometry of Transformer Representations****ArXiv CS.AI** **50%**

arXiv:2605.01609v2 Announce Type: replace-cross Abstract: We find that transformer concept representations systematically anti-concentrate in the spectral tail of the unembedding covariance, encoding word-level concepts in low-variance directions across a 17-model core suite and...

<https://arxiv.org/abs/2605.01609>

4 **CONF****Cross-Relational Preference Learning for Better LLM Instruction Following****ArXiv CS.AI** **50%**

arXiv:2608.29352v1 Announce Type: new Abstract: Large Language Models (LLMs) still exhibit limited capability in following complex instructions. While existing approaches often rely on preference learning to enhance this ability, they typically overlook the relationships between...

<https://arxiv.org/abs/2608.29352>

5 **CONF****How Identity and Opinion Shape Political Sycophancy in LLMs****ArXiv CS.AI** **50%**

arXiv:2608.29198v1 Announce Type: new Abstract: As Large Language Models (LLMs) increasingly encourage users to disclose personal profiles for tailored assistance, measuring their political alignment becomes increasingly important. However, many existing benchmarks for assessing...

<https://arxiv.org/abs/2608.29198>

6 **CONF****Localizing Emergent Failures in Agentic AI: Recovering Minimal Repair Families via Counterfactual Replay****ArXiv CS.AI** **50%**

arXiv:2608.29228v1 Announce Type: new Abstract: Failures in agentic AI systems can arise from interactions among messages exchanged by multiple large language model (LLM) agents. Pointwise attribution cannot distinguish a jointly necessary repair from alternative singleton...

<https://arxiv.org/abs/2608.29228>

7 **CONF****Validating FKG.in: Soundness Assessment in LLM-Augmented Indian Food Knowledge****ArXiv CS.AI** **50%**

arXiv:2608.29249v1 Announce Type: new Abstract: The online culinary ecosystem is increasingly populated by recipe content generated, modified, or summarized by Large Language Models (LLMs). While often plausible, such outputs may contain hallucinated ingredients, misrepresented...

<https://arxiv.org/abs/2608.29249>

8 CONF

RACER: Reinforced Agent Collaboration for Explainable Reasoning on Knowledge Graphs

ArXiv CS.AI 50%

arXiv:2608.29263v1 Announce Type: new Abstract: Large Language Models (LLMs) often suffer from hallucination and struggle with complex reasoning tasks requiring multi-hop domain knowledge. While integrating Knowledge Graphs (KGs) provides a structured and verifiable information...

<https://arxiv.org/abs/2608.29263>

9 CONF

EpaCache: Error-Propagation-Aware Caching for Accelerating Diffusion-Based Visual Generation

ArXiv CS.AI 50%

arXiv:2608.29264v1 Announce Type: new Abstract: Diffusion-based visual generative models deliver strong image and video synthesis quality but incur high inference costs because sequential samplers repeatedly evaluate large networks. Caching-based methods reduce inference latency...

<https://arxiv.org/abs/2608.29264>

0 CONF

Measuring the Depth of LLM Unlearning via Activation Patching

ArXiv CS.AI 50%

arXiv:2605.24614v2 Announce Type: replace-cross Abstract: Large language model (LLM) unlearning has emerged as a crucial post-hoc mechanism for privacy protection and AI safety, yet auditing whether target knowledge is truly erased remains challenging. Existing output-level...

<https://arxiv.org/abs/2605.24614>

1 CONF

BIRD-History: A Benchmark for History-Driven Text-to-SQL with Fine-Grained Knowledge Annotations

ArXiv CS.AI 50%

arXiv:2608.29345v1 Announce Type: new Abstract: While recent Large Language Model (LLM)-based text-to-SQL systems achieve impressive performance on standard benchmarks, they struggle when user queries implicitly rely on domain-specific knowledge, such as business logic, data...

<https://arxiv.org/abs/2608.29345>

2 CONF

Plant-Inspired AI: Plants as Inspiration for Novel Problem Formulations, and Two Case Studies

ArXiv CS.AI 50%

arXiv:2608.29356v1 Announce Type: new Abstract: Artificial Intelligence (AI) has long been inspired by studies of biological intelligence. Reinforcement learning, for instance, drew inspiration from studies involving animal learning and is now a powerful paradigm for solving...

<https://arxiv.org/abs/2608.29356>

3 CONF

LiteSearch-VL: Small Multimodal Search Agents via Trajectory Distillation and Synthetic Step-DPO

ArXiv CS.AI 50%

arXiv:2608.29357v1 Announce Type: new Abstract: Multimodal search agents answer visual questions by interleaving image understanding, web retrieval, tool use, and evidence synthesis. Strong systems exist, but in two expensive regimes: proprietary frontier models such as GPT-5...

<https://arxiv.org/abs/2608.29357>

4 **CONF****Evaluating Tiny Recursive Models Across Training for Code Generation****ArXiv CS.AI 50%**

arXiv:2608.29376v1 Announce Type: new Abstract: Code generation increasingly relies on large transformer models, whose capability advances with scale. Yet such a scale is costly, creating demand for small models, especially where data is limited. Recursive models address this by...

<https://arxiv.org/abs/2608.29376>
5 **CONF****LLMs Interpret, Embeddings Organize, Graphs Emerge: Agent-Driven Compilation of Scientific Knowledge****ArXiv CS.AI 50%**

arXiv:2608.29612v1 Announce Type: new Abstract: Sustained scientific work requires a knowledge substrate that carries interpretation across tasks and preserves paths to source evidence. We call this process \emph{scientific knowledge compilation} and implement it in ASKS, the...

<https://arxiv.org/abs/2608.29612>
6 **CONF****Beyond Ranking Accuracy: Evaluating LLM-Cited Feature Rationales for Next Basket Repurchase Recommendation****ArXiv CS.AI 50%**

arXiv:2608.30333v1 Announce Type: cross Abstract: Next-basket repurchase recommendation is commonly formulated as a ranking task: given a customer's purchase history, the system ranks previously purchased items that may be needed again. In production settings, however, ranking...

<https://arxiv.org/abs/2608.30333>
7 **CONF****Cross-Regional Grapevine Cold Hardiness Prediction via Learned Multimodal Latent Representations****ArXiv CS.AI 50%**

arXiv:2608.31097v1 Announce Type: new Abstract: Accurate daily predictions of cold hardiness in woody plants are critical in regions where freezing temperatures can damage dormant buds and reduce seasonal yield. Existing biophysical, hybrid, and deep learning models have shown...

<https://arxiv.org/abs/2608.31097>
8 **CONF****Prefix-Adaptive Block Diffusion for Efficient Document Recognition****ArXiv CS.AI 50%**

arXiv:2605.16861v2 Announce Type: replace-cross Abstract: Block Diffusion Models (BDMs) support parallel generation, flexible-length output, and KV caching, making them promising for efficient document parsing. However, existing BDMs bind denoising and cache commitment to fixed...

<https://arxiv.org/abs/2605.16861>
9 **CONF****Review Before Trust: Source-Grounded Integrity Gates for AI-Assisted Personal Health Records****ArXiv CS.AI 50%**

arXiv:2608.29965v1 Announce Type: new Abstract: Large language models can convert medical documents into structured data, but plausible output may still be unsupported by the source. Persisting such output in a longitudinal health record, a record that accumulates patient...

<https://arxiv.org/abs/2608.29965>

- 0
CONF

An Open-Source, Event-Driven Pipeline for Cryptocurrency Market Data: Ingestion, Forecasting, and On-Chain Fraud Detection

ArXiv CS.AI 50%

arXiv:2608.29973v1 Announce Type: new Abstract: Cryptocurrency markets generate high-frequency, multi-source data that is expensive to work with unless a team already has commercial-grade streaming and warehousing infrastructure in place. This paper describes a fully open-source...

<https://arxiv.org/abs/2608.29973>
- 1
CONF

AutoCRAT: Within-trajectory Joint Control of Stochasticity and Compute for LLM Reasoning

ArXiv CS.AI 50%

arXiv:2608.29988v1 Announce Type: new Abstract: Large language models (LLMs) achieve strong reasoning performance, which depends critically on inference-time decisions. Yet these decisions are commonly handled by static, one-size-fits-all policies, limiting adaptation to diverse...

<https://arxiv.org/abs/2608.29988>
- 2
CONF

NeuroBreak: Unveil Internal Jailbreak Mechanisms in Large Language Models

ArXiv CS.AI 50%

arXiv:2509.03985v3 Announce Type: replace-cross Abstract: Jailbreak attacks bypass the safety alignment of large language models (LLMs) to elicit harmful outputs, yet the vast parameter space makes diagnosing the underlying failure mechanisms extremely challenging. We present...

<https://arxiv.org/abs/2509.03985>
- 3
CONF

Beyond Uncertainty: Multi-Solver Disagreement Rewards for Self-Evolving Reasoning Curricula

ArXiv CS.AI 50%

arXiv:2608.30035v1 Announce Type: new Abstract: Self-evolving reasoning frameworks train a Challenger to generate questions exposing a Solver's weaknesses, creating adaptive curricula without human data. However, existing approaches use a single solver's sampling uncertainty as...

<https://arxiv.org/abs/2608.30035>
- 4
CONF

Balance of Benchmarks: Semantic Density Reweighting for Benchmark Multiplicity and Task-Conditioned Evaluation

ArXiv CS.AI 50%

arXiv:2608.30044v1 Announce Type: new Abstract: Language models are commonly compared by averaging scores across a benchmark list with equal weight. Such lists grow through publication outside an explicit measurement design, so equal weighting turns the density of published...

<https://arxiv.org/abs/2608.30044>
- 5
CONF

Can LLM Agents Discover? Evaluating Creativity on ML Engineering Tasks

ArXiv CS.AI 50%

arXiv:2608.30047v1 Announce Type: new Abstract: Recent AI systems promise autonomous scientific discovery, claiming to discover algorithms and produce research papers, yet understanding whether they exhibit creativity, the capacity to produce solutions that are both novel and...

<https://arxiv.org/abs/2608.30047>

6 **CONF****LaMoC: Loss-Aware Modular Compression for LLMs****ArXiv CS.AI 50%**

arXiv:2608.30226v1 Announce Type: new Abstract: Modular compression has enabled considerable parameter reduction in LLMs while preserving strong language understanding and downstream task accuracy. However, existing joint modular compression methods primarily rely on activation...

<https://arxiv.org/abs/2608.30226>

7 **CONF****CoLa-ICD: A Knowledge-Enhanced Framework for Long-Tail Automated Medical Coding****ArXiv CS.AI 50%**

arXiv:2608.30234v1 Announce Type: new Abstract: Automatic medical coding assigns ICD codes to clinical notes, but it remains challenging due to long documents, imbalanced label distributions, and diverse terms. These challenges are especially severe for rare codes, which have...

<https://arxiv.org/abs/2608.30234>

8 **CONF****LLM-Based Knowledge Graph Completion Combining Discrete Structural Coding with Similar Entity Information****ArXiv CS.AI 50%**

arXiv:2608.30235v1 Announce Type: new Abstract: Knowledge graph completion requires models to use both textual descriptions and relational structure. Existing LLM-based methods either encode KG structure as discrete tokens or refine a restricted set of candidate entities, and...

<https://arxiv.org/abs/2608.30235>

9 **CONF****Generating Workflow DAGs from Natural Language with Non-Reasoning LLMs****ArXiv CS.AI 50%**

arXiv:2608.30250v1 Announce Type: new Abstract: This paper addresses the problem of translating natural-language routing rules written by business administrators into executable workflow graphs for enterprise contact centers. Each target is a directed acyclic graph (DAG) of...

<https://arxiv.org/abs/2608.30250>

0 **CONF****SimCRAFT: Distilling Remote Sensing Agents via Synthetic Trajectories and Contextual Retrieval-Augmented Fine-Tuning****ArXiv CS.AI 50%**

arXiv:2608.30277v1 Announce Type: new Abstract: The unprecedented surge in Earth observation data volume and diversity has exposed a critical bottleneck for traditional manual workflows, catalyzing the emergence of Remote Sensing (RS) Agents. However, the practical deployment of...

<https://arxiv.org/abs/2608.30277>

1 **CONF****Ignorance or Incompetence? Constructing Knowledge-Gated, Verifiable Tasks for LLM Agents****ArXiv CS.AI 50%**

arXiv:2608.30322v1 Announce Type: new Abstract: Professional agent tasks often depend on conventions that are absent from public corpora, yet benchmarks rarely control whether an agent has access to those conventions. We introduce a knowledge-gated task-construction protocol...

<https://arxiv.org/abs/2608.30322>

2 **CONF****Answer Probing-Guided Search for Diverse Solution Exploration of LLMs****ArXiv CS.AI 50%**

arXiv:2608.30345v1 Announce Type: new Abstract: Generating multiple diverse and high-quality solutions is valuable for many applications, such as code-test generation and drug discovery. However, Large Language Models (LLMs) tend to converge on a single high-confidence solution...

<https://arxiv.org/abs/2608.30345>

3 **CONF****Will the User Ever Know? Covert Indirect Prompt Injection on Tool-Using LLM Agents****ArXiv CS.AI 50%**

arXiv:2608.30362v1 Announce Type: new Abstract: As LLM agents take real-world actions through tools, indirect prompt injection (IPI) has emerged as a serious threat. The standard metric, Attack Success Rate (ASR), counts whether an injection succeeds but ignores what the user...

<https://arxiv.org/abs/2608.30362>

4 **CONF****TuringLLM: Efficiently Scaling Foundation Models Toward Physical AI****ArXiv CS.AI 50%**

arXiv:2608.30567v1 Announce Type: new Abstract: We present Turing-20B-A2B, a 20B-parameter Mixture-of-Experts language model that activates approximately 2B parameters per token, designed for long-context and latency-sensitive physical AI applications. The model adopts Quantile...

<https://arxiv.org/abs/2608.30567>

5 **CONF****DERELAB: Probing Defeasible Reasoning and Confirmation Bias in LLMs with a Generative Benchmark****ArXiv CS.AI 50%**

arXiv:2608.30413v1 Announce Type: new Abstract: Defeasible reasoning is a type of reasoning where inferences are drawn from plausible current evidence, but can be retracted upon the introduction of newer evidence. Although recent studies have examined language-model behaviors in...

<https://arxiv.org/abs/2608.30413>

6 **CONF****CM2: Multimodal Cultural Reasoning via an Integrated Multi-Agent Framework****ArXiv CS.AI 50%**

arXiv:2608.30498v1 Announce Type: new Abstract: Multimodal Large Language Models (MLLMs) have shown remarkable success in STEM domains, where progress is often driven by vertical, step-by-step deduction under relatively stable symbol systems. Their horizontal, interdisciplinary...

<https://arxiv.org/abs/2608.30498>

7 **CONF****Designing an Auditable LLM-Supported Workflow for Qualitative Thematic Analysis****ArXiv CS.AI 50%**

arXiv:2608.30543v1 Announce Type: new Abstract: Large Language Models (LLMs) offer new possibilities for scaling qualitative analysis, but existing applications often provide limited methodological transparency regarding how qualitative methods are translated into computational...

<https://arxiv.org/abs/2608.30543>

8 CONF

MM-BrowseComp: A Comprehensive Benchmark for Multimodal Browsing Agents

ArXiv CS.AI 50%

arXiv:2508.13186v2 Announce Type: replace-cross Abstract: AI agents with advanced reasoning and tool-use capabilities have demonstrated impressive performance in web browsing for deep search. However, existing benchmarks such as BrowseComp primarily focus on textual content,...

<https://arxiv.org/abs/2508.13186>

9 CONF

Automated Testing of LLM-Based Post Hoc Explainers Using Model Checking as an Oracle

ArXiv CS.AI 50%

arXiv:2608.30581v1 Announce Type: new Abstract: Large language models (LLMs) are used as post hoc explainers of sequential decision-making policies, producing natural-language explanations of why an action was chosen. However, LLMs often generate plausible but incorrect...

<https://arxiv.org/abs/2608.30581>

0 CONF

AgenticRag-R1: Agentic Reinforcement Learning with Stack Memory for Multi-Step Reasoning, Retrieval and Memorizing

ArXiv CS.AI 50%

arXiv:2608.29622v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) improves the factuality of large language models (LLMs), yet existing RAG systems often struggle with complex, multi-step reasoning that requires adaptive retrieval and continuous revision of...

<https://arxiv.org/abs/2608.29622>

1 CONF

The Halt Vector: Internalizing a Causal Steering Intervention for Efficient Reasoning

ArXiv CS.AI 50%

arXiv:2608.28859v1 Announce Type: cross Abstract: Reasoning models do not stop when they know the answer. On DeepSeek-R1-Distill-Qwen-7B the chain of thought runs about twice as long as the model's own answer probability takes to settle, and how much of that excess is removable...

<https://arxiv.org/abs/2608.28859>

2 CONF

VFR-Audit: Verdict-Level Reliability for Fairness Audits in Hospital Length-of-Stay Prediction

ArXiv CS.AI 50%

arXiv:2608.30846v1 Announce Type: new Abstract: Fairness audits in clinical Artificial Intelligence convert continuous fairness metrics into binary pass-or-fail verdicts against operational thresholds, where hospital governance boards, payers, and regulators act on the resulting...

<https://arxiv.org/abs/2608.30846>

3 CONF

Predicting Residential Rents in Dakar Using Machine Learning

ArXiv CS.AI 50%

arXiv:2608.30865v1 Announce Type: new Abstract: Dakar's residential rental market remains poorly documented despite its economic and social importance: 54.4% of households are renters, compared to 23.3% nationally. This study develops a complete machine learning pipeline to...

<https://arxiv.org/abs/2608.30865>

- 4 **CONF**
- Responsible Integration of AI in Cancer Genomics: Barriers, Risks, and Pathways to Trustworthy Clinical Translation**
- ArXiv CS.AI 50%
- arXiv:2608.30912v1 Announce Type: new Abstract: Artificial intelligence (AI) and natural language processing (NLP) are increasingly used to extract, integrate, and interpret biomedical knowledge relevant to cancer genomics, yet their translation into routine clinical oncology...
- <https://arxiv.org/abs/2608.30912>
-
- 5 **CONF**
- Learning Action Models with Conditional and Quantified Effects via Uncertainty-Guided Exploration**
- ArXiv CS.AI 50%
- arXiv:2608.30955v1 Announce Type: new Abstract: Accurate action models are critical for effective planning. Existing action-model learning methods largely assume simple action representations or become computationally intractable when learning conditional and quantified effects....
- <https://arxiv.org/abs/2608.30955>
-
- 6 **CONF**
- MNIST-PRO: MNIST is Back as a Partially Observable World for AI Agents**
- ArXiv CS.AI 50%
- arXiv:2608.31022v1 Announce Type: new Abstract: AI agents in partially observable environments need to coordinate active sensing with working memory to maintain an evolving perceptual state. However, existing benchmarks struggle to isolate this perceptual-state construction and...
- <https://arxiv.org/abs/2608.31022>
-
- 7 **CONF**
- Wrong Prediction, Right Answer: Recovering Evidence from Collapsed LLM Sequence Scores**
- ArXiv CS.AI 50%
- arXiv:2608.31068v1 Announce Type: new Abstract: When a large language model fails a reasoning task, it is often assumed to lack the underlying capability. However, this conflates a genuine absence of reasoning with a late-stage output bottleneck. We observe a consistent readout...
- <https://arxiv.org/abs/2608.31068>
-
- 8 **CONF**
- When Do Larger Batches Help Scale LLM Reinforcement Learning?**
- ArXiv CS.AI 50%
- arXiv:2608.29296v1 Announce Type: cross Abstract: Larger batches reduce the variance of stochastic gradients per update and are therefore often expected to accelerate training. Yet whether this statistical benefit translates into lower wall-clock time-to-target remains unclear,...
- <https://arxiv.org/abs/2608.29296>
-
- 9 **CONF**
- OntoAligner-Ensemble: Voting-Based Fusion across Heterogeneous Ontology Alignment Techniques**
- ArXiv CS.AI 50%
- arXiv:2608.31137v1 Announce Type: new Abstract: Ontology alignment (OA) has evolved through several methodological paradigms, ranging from lexical and structural aligners to knowledge graph embedding (KGE) models and, more recently, Large Language Model (LLM)-based approaches....
- <https://arxiv.org/abs/2608.31137>
-

00 CONF

Agent-Based Model Framework for the North Carolina Modeling Infectious Diseases Program (NC MinD ABM) Overview, Design Concepts, and Details Protocol

ArXiv CS.AI 50%

arXiv:2202.06853v2 Announce Type: cross Abstract: To help facilitate a variety of simulations related to healthcare facilities in North Carolina, we have developed an agent-based model (ABM) to accurately simulate patient (i.e., agent) movement to and from these facilities. This...

<https://arxiv.org/abs/2202.06853>

01 CONF

PowerSlider: Exploiting Phase Asymmetry for LLM Serving under Demand Response

ArXiv CS.AI 50%

arXiv:2608.21719v1 Announce Type: cross Abstract: AI inference clusters are increasingly constrained by instantaneous power, not just energy: grid operators condition new capacity on demand response, imposing time-varying power caps. Existing LLM serving systems optimize a...

<https://arxiv.org/abs/2608.21719>

02 CONF

NLP-Driven Knowledge Extraction and Thematic Classification of Translated Ancient Indian Medical Texts

ArXiv CS.AI 50%

arXiv:2608.28608v1 Announce Type: cross Abstract: Ancient Indian medical texts like Sushruta Samhita have extensive information on diseases, treatments, and surgical techniques. Yet, their ancient format and use of intricate vocabulary pose difficulties in accessibility and...

<https://arxiv.org/abs/2608.28608>

03 CONF

Parametric Multimodal User Memory: Storing What Captions Cannot Carry

ArXiv CS.AI 50%

arXiv:2608.28609v1 Announce Type: cross Abstract: A personalized agent needs a user memory: a persistent model of who its user is. Today it is almost always text -- transcripts and captions retrieved by similarity. This serves the captionable half of a person ("my cat is named...

<https://arxiv.org/abs/2608.28609>

04 CONF

From GenAI Virtual Patient Dialogue Logs to Teacher-Interpretable Process Evidence: A Learning Analytics Study in Higher Education

ArXiv CS.AI 50%

arXiv:2608.28619v1 Announce Type: cross Abstract: Medical history taking is a dialogue-based clinical reasoning task in which learners must gather, organise, and integrate patient information while the consultation unfolds. Generative AI-powered virtual patients (GenAI VPs) make...

<https://arxiv.org/abs/2608.28619>

05 CONF

Looking Again: Measuring Sycophancy in the Reasoning Chains of Multimodal Models Under Pressure

ArXiv CS.AI 50%

arXiv:2608.28623v1 Announce Type: cross Abstract: Large multimodal reasoning models (LMRMs) are getting increasingly capable, primarily through generating explicit chain-of-thought reasoning before answering. In language models it has been observed that this performance often...

<https://arxiv.org/abs/2608.28623>

06 CONF

MA-RAG: Multi-Agent Retrieval-Augmented Generation for Query-Driven Summarization of Longitudinal Parkinson's Disease Assessments

ArXiv CS.AI 50%

arXiv:2608.28624v1 Announce Type: cross Abstract: Accurate interpretation of single-visit and longitudinal clinical assessments for Parkinson's disease is time-consuming and often depends on specialist expertise. Although large language models (LLMs) can generate natural...

<https://arxiv.org/abs/2608.28624>

07 CONF

Do large language models scrutinise what they review? A multimodal audit of scoring calibration, error detection, and author-identity effects

ArXiv CS.AI 50%

arXiv:2608.28626v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to generate peer reviews, prompting examination of their capacity for critical evaluation. This study evaluates two multimodal LLMs, Qwen2.5-VL-72B and Pixtral-Large-124B, as...

<https://arxiv.org/abs/2608.28626>

08 CONF

Error Detection for PET/CT Radiology Reports: Domain-Specific vs Large Language Models

ArXiv CS.AI 50%

arXiv:2608.30021v1 Announce Type: cross Abstract: Errors in radiology reports can adversely affect patient treatment, yet automated report quality assurance remains challenging because errors are often subtle and require domain expertise to detect. Although large language models...

<https://arxiv.org/abs/2608.30021>

09 CONF

Do MLLMs Really Understand Low-Resource Khmer Documents? A Pilot Study on Khmer Document VQA

ArXiv CS.AI 50%

arXiv:2608.28635v1 Announce Type: cross Abstract: Recent multimodal large language models (MLLMs) have advanced document understanding, visual question answering, and text extraction. However, their reliability in low-resource, non-Latin settings remains uncertain. Khmer form...

<https://arxiv.org/abs/2608.28635>

10 CONF

Terminal-Bench-LILT: Multilingual Agentic Coding Benchmark Grounded in Language, Region, and Culture

ArXiv CS.AI 50%

arXiv:2608.28641v1 Announce Type: cross Abstract: Most evaluations for coding agents are conducted exclusively in English, which does not reflect real-world multilingual deployment. We present Terminal-Bench-LILT, a suite of 300 authentic coding tasks in ten languages: Arabic,...

<https://arxiv.org/abs/2608.28641>

11 CONF

Cross Lingual Transfer in Tulu Legal Comprehension: Script-Dependent Improvement and RAG-Induced Knowledge Conflict

ArXiv CS.AI 50%

arXiv:2608.28645v1 Announce Type: cross Abstract: Low-resource languages without an adequate training corpus often use a related, higher-resource language as a scaffold for comprehension. Still, there is a need to develop rigorous evaluation methods to identify when models fail...

<https://arxiv.org/abs/2608.28645>

12 CONF

Can Large Language Models Identify Meaningful Touchpoints in Conversion Attribution?

ArXiv CS.AI 50%

arXiv:2608.28649v1 Announce Type: cross Abstract: Touchpoint selection in conversion attribution, namely identifying meaningful touchpoints contributing to conversions, is essential for e-commerce recommendation and online advertising. Current selection methods rely heavily on...

<https://arxiv.org/abs/2608.28649>

13 CONF

GreenBench: Benchmarking Energy Efficiency and Carbon Footprint of Open-Source LLM Inference on Apple Silicon

ArXiv CS.AI 50%

arXiv:2608.28667v1 Announce Type: cross Abstract: The rapid proliferation of Large Language Models (LLMs) has raised concerns about their environmental impact during inference. While Green AI research has focused on datacenter GPUs and embedded platforms, the energy profile of...

<https://arxiv.org/abs/2608.28667>

14 CONF

Measuring Similarity between Artistic and AI Generated Images using Siamese Neural Networks

ArXiv CS.AI 50%

arXiv:2608.28671v1 Announce Type: cross Abstract: AI-generated art has sparked debates around potential plagiarism, as these images may closely resemble existing artworks. This research quantifies the similarity between original pieces and AI-generated counterparts, particularly...

<https://arxiv.org/abs/2608.28671>

15 CONF

Defending Wearable VLMs Against Private Attribute Inference

ArXiv CS.AI 50%

arXiv:2608.28691v1 Announce Type: cross Abstract: Wearable VLM pipelines promise continuous multimodal assistance from egocentric visual capture: a user asks a task-driven question about the surrounding scene, and the system uses compact visual tokens to support language...

<https://arxiv.org/abs/2608.28691>

16 CONF

The reach of a verification tool decides its value: A controlled study of verification surface, artifact quality, and cost in AI coding agents

ArXiv CS.AI 50%

arXiv:2608.28795v1 Announce Type: cross Abstract: Modern artificial-intelligence coding agents can be equipped with tools for checking their own work e.g. a linter, a boot probe, a shell, a screenshot tool. We call this set the agent's verification surface. This study asks...

<https://arxiv.org/abs/2608.28795>

17 CONF

Text-Driven Artistic Staging: Pose, Lighting, and Camera References from Paintings

ArXiv CS.AI 50%

arXiv:2608.28823v1 Announce Type: cross Abstract: Artists coordinate human pose, illumination, and camera placement to convey narrative and emotion, but existing generative methods typically model these elements independently. We introduce text-to-editable 3D staging, a task...

<https://arxiv.org/abs/2608.28823>

18 CONF

Federated Learning for MRI-based BrainAGE: a multicenter study on post-stroke functional outcome prediction

ArXiv CS.AI 50%

arXiv:2506.15626v3 Announce Type: replace-cross Abstract: $\textbf{Objective:}$ Brain-predicted age difference (BrainAGE) is a neuroimaging biomarker reflecting brain health. However, training robust BrainAGE models requires large datasets, often restricted by privacy concerns....

<https://arxiv.org/abs/2506.15626>

19 CONF

Curvature Cryptanalysis of Smooth Transformer Feed-Forward Networks

ArXiv CS.AI 50%

arXiv:2608.28843v1 Announce Type: cross Abstract: We show that smooth two-layer feed-forward networks (FFNs) expose an additional structural model extraction channel under a chosen-input raw-output oracle at the FFN branch; consider transformer FFN branches with GELU or SiLU...

<https://arxiv.org/abs/2608.28843>

20 CONF

A rigor-matched audit of periodic-step layer skipping for efficient llm inference: conflagrators versus swift, with a supplemental analysis of trained routing alternatives

ArXiv CS.AI 50%

arXiv:2608.28846v1 Announce Type: cross Abstract: Layer-skipping methods for efficient LLM inference decide, at some granularity, which transformer layers to execute for a given input. We present a rigor-matched, three-seed audit of two periodic-step, search-based methods that...

<https://arxiv.org/abs/2608.28846>

21 CONF

Equivariant Sheaf Neural Networks: Learning Geometric Transport on Graphs

ArXiv CS.AI 50%

arXiv:2608.28853v1 Announce Type: cross Abstract: Equivariant graph neural networks provide a principled way to model geometric systems, but efficient first-order architectures remain limited in how vector information can be transformed as it moves across a graph. We introduce...

<https://arxiv.org/abs/2608.28853>

22 CONF

Hybrid Offline-Online Multi-Agent Decision Transformers for Wireless Resource Management

ArXiv CS.AI 50%

arXiv:2608.28878v1 Announce Type: cross Abstract: This paper develops a hybrid offline-online multi-agent reinforcement learning framework based on decision transformers. The policy is first pretrained offline via supervised sequence modeling of trajectories generated by...

<https://arxiv.org/abs/2608.28878>

23 CONF

Moving the Mean Toward the Known Good, Not Beyond It: What Inference-Time Interventions and Weight Consolidation Buy in Open-Ended Generation

ArXiv CS.AI 50%

arXiv:2608.28886v1 Announce Type: cross Abstract: What does a generation loop gain from learning on its own verified successes? In cycles of generate, verify, select and LoRA-consolidate on online bin packing, training on value-filtered candidates shifts what the model writes on...

<https://arxiv.org/abs/2608.28886>

24 CONF

From the Loss Landscape to Diverse Feature Learning in Neural Networks

ArXiv CS.AI 50%

arXiv:2608.28948v1 Announce Type: cross Abstract: Over the course of the last decade, neural networks have grown from an academic curiosity to moving the markets of nations. Despite this explosion in both research and deployment, relatively little is understood about how they...

<https://arxiv.org/abs/2608.28948>

25 CONF

The information geometry of product-reference discrete diffusion: Interaction growth complexity and optimal scheduling

ArXiv CS.AI 50%

arXiv:2608.28949v1 Announce Type: cross Abstract: We study a class of product-reference diffusion algorithms for sampling from a discrete distribution. We show that their sampling performance can be characterized using a path-based measure of data geometry that we call the...

<https://arxiv.org/abs/2608.28949>

26 CONF

Diagnose, Then Refine: A Closed-Loop TTS System with AudioLLM-Guided Correction

ArXiv CS.AI 50%

arXiv:2608.28970v1 Announce Type: cross Abstract: Current TTS systems typically rely on open-loop, single-pass generation and can produce sporadic local prosodic defects, such as misplaced stress, unnatural pauses, or flattened intonation, that utterance-level metrics often fail...

<https://arxiv.org/abs/2608.28970>

27 CONF

Free Speech and Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2608.28973v1 Announce Type: cross Abstract: Philosophers and legal scholars are engaged in debates about the implications of artificial intelligence for freedom of expression. This paper analyzes the free speech issues raised by two distinct AI technologies: social media...

<https://arxiv.org/abs/2608.28973>

28 CONF

The Illusion of Replacement: Rethinking Specialized Machine Learning Models in the Foundation Model Era

ArXiv CS.AI 50%

arXiv:2608.28980v1 Announce Type: cross Abstract: Can the specialized architectures that machine learning has traditionally built for structured data be replaced by language-based models? This question is examined through a review of 159 papers (2016--2026) across nine...

<https://arxiv.org/abs/2608.28980>

29 CONF

Development of an Autonomous AI Coding Agent using Monte Carlo Tree Search (MCTS) and Gemini LLM Frameworks

ArXiv CS.AI 50%

arXiv:2608.29096v1 Announce Type: cross Abstract: The ongoing changes in software engineering requirements have created a substantial need for automated tools which can create secure source code from natural language input. The performance of traditional Large Language Models...

<https://arxiv.org/abs/2608.29096>

30 CONF

CGFM-Nav: Cognitive Graph-Field Memory for Semantic-Guided Lifelong Multimodal Embodied Navigation

ArXiv CS.AI 50%

arXiv:2608.29114v1 Announce Type: cross Abstract: Vision-and-Language Navigation (VLN) requires agents to reason over accumulated observations while continuously exploring unseen regions. However, existing environment representations often struggle to jointly support explicit...

<https://arxiv.org/abs/2608.29114>

31 CONF

HEAR Who Said What: Unlocking Speaker-Attributed Reasoning via Counterfactual Voice Grounding

ArXiv CS.AI 50%

arXiv:2608.29120v1 Announce Type: cross Abstract: Speech Language Models (SLMs) are increasingly deployed in multi-speaker environments, yet their ability to attribute speech to the correct speaker and reason over speaker identities remains unclear. Hence, we introduce HEAR, a...

<https://arxiv.org/abs/2608.29120>

32 CONF

Triggering Chain-of-Thought via Latent Feature Interventions in Large Language Models

ArXiv CS.AI 50%

arXiv:2601.08058v2 Announce Type: replace-cross Abstract: Chain-of-Thought (CoT) prompting often improves the reasoning performance of large language models (LLMs), but the internal signal that triggers this behavior remains poorly understood. Leveraging the sparse features...

<https://arxiv.org/abs/2601.08058>

33 CONF

STARLINC: Satellite Trail Artifact Removal using Inter-Frame Correlation

ArXiv CS.AI 50%

arXiv:2608.29145v1 Announce Type: cross Abstract: The rapid expansion of low Earth orbit satellites such as Starlink is increasingly contaminating astronomical surveys. In practice, contaminated images are often identified through inspection. However, modern surveys generate...

<https://arxiv.org/abs/2608.29145>

34 CONF

Training and Agentic Inference Strategies for LLM-based Manim Animation Generation

ArXiv CS.AI 50%

arXiv:2604.18364v2 Announce Type: replace Abstract: Generating programmatic animation using libraries such as Manim presents unique challenges for Large Language Models (LLMs), requiring spatial reasoning, temporal sequencing, and familiarity with domain-specific APIs that are...

<https://arxiv.org/abs/2604.18364>

35 CONF

TAAL: Mitigating Early Beam Pruning in Generative Recommendation via Temporal Autoregressive Alignment

ArXiv CS.AI 50%

arXiv:2608.29179v1 Announce Type: cross Abstract: Generative recommendation encodes items as hierarchical semantic identifiers (SIDs) and retrieves the next item through autoregressive decoding. Standard next-token prediction, however, does not explicitly cover the multimodal...

<https://arxiv.org/abs/2608.29179>

36 CONF

Rate-Coding Bundle Memory: A Unified Model of Memory and Control for Symbolic Computation in the Brain

ArXiv CS.AI 50%

arXiv:2608.29189v1 Announce Type: cross Abstract: We propose a neurobiologically plausible model of cognition that combines the advantages of connectionist and symbolic systems, and that can explain a wide range of cognitive phenomena. This model, called Rate-Coding Bundle...

<https://arxiv.org/abs/2608.29189>

37 CONF

PokaiTrainer: Scaling Belief-State Search to Competitive Pok\emon VGC

ArXiv CS.AI 50%

arXiv:2608.29197v1 Announce Type: cross Abstract: Decision-time equilibrium search carried poker to superhuman play, but it has so far relied on tractable subgames: a handful of actions per decision, chance confined to card deals, one player moving at a time. Competitive...

<https://arxiv.org/abs/2608.29197>

38 CONF

Measurement Validity in LLM Cultural Alignment

ArXiv CS.AI 50%

arXiv:2608.29266v1 Announce Type: cross Abstract: Researchers increasingly treat LLM survey responses as a proxy for human cultural values. This includes projecting model outputs onto instruments like the Inglehart-Welzel Cultural Map and drawing conclusions about which cultures...

<https://arxiv.org/abs/2608.29266>

39 CONF

ST-EVO: Towards Generative Spatio-Temporal Evolution of Multi-Agent Communication Topologies

ArXiv CS.AI 50%

arXiv:2602.14681v5 Announce Type: replace-cross Abstract: LLM-powered Multi-Agent Systems (MAS) have emerged as an effective approach towards collaborative intelligence, and have attracted wide research interests. Among them, ``self-evolving" MAS, treated as a more flexible and...

<https://arxiv.org/abs/2602.14681>

40 CONF

RAGDiffusion++: From Macro-Retrieval to Micro-Fidelity Alignment for Garment Generation

ArXiv CS.AI 50%

arXiv:2608.29280v1 Announce Type: cross Abstract: Standard clothing asset generation---restoring forward-facing flat-lay garment images from diverse real-world contexts---holds immense commercial value yet demands both macroscopic topological accuracy and microscopic physical...

<https://arxiv.org/abs/2608.29280>

41 CONF

Evaluating LLM-based AI agents integrated with materials synthesis tools: the case of atomic layer deposition

ArXiv CS.AI 50%

arXiv:2608.29309v1 Announce Type: cross Abstract: This work provides an overview of the different strategies that can be used to evaluate the performance of AI models and agents based on large language models (LLMs) for materials synthesis. After providing a brief overview of...

<https://arxiv.org/abs/2608.29309>

42 CONF

Arabic Safety Alignment as Selective Refusal: An Empirical Study of SFT, DPO, and Guard Calibration

ArXiv CS.AI 50%

arXiv:2608.29378v1 Announce Type: cross Abstract: Arabic large language models must refuse harmful prompts without over-refusing benign or sensitive prompts, yet a single refusal rate hides this trade-off. We evaluate it using benign refusal B and harmful-prompt refusal H, where...

<https://arxiv.org/abs/2608.29378>

43 CONF

AI Can Be Easily Persuaded in Clinical Decision Making

ArXiv CS.AI 50%

arXiv:2608.29453v1 Announce Type: cross Abstract: As AI becomes increasingly integrated into clinical practice, it is playing a growing role in medical decision making. Medicine, however, is a high stakes and evidence based field, where decisions can directly affect patients'...

<https://arxiv.org/abs/2608.29453>

44 CONF

Reference-Grafting Matches Fine-Tuning at Eliciting Sandbagged Capabilities

ArXiv CS.AI 50%

arXiv:2608.29458v1 Announce Type: cross Abstract: Sandbagging, in which a model deliberately underperforms on an evaluation despite retaining the underlying capability, threatens the safety evaluations that frontier-model governance depends on. The Elicitation Game found that...

<https://arxiv.org/abs/2608.29458>

45 CONF

Benchmark Contamination: A Taxonomy Organized by Defeated Mitigation

ArXiv CS.AI 50%

arXiv:2608.29463v1 Announce Type: cross Abstract: A benchmark score is a joint property of the model, the evaluation harness, the elicitation budget, the sampled population, and contamination status. Leaderboards publish the model and the score, so capability and leakage stay...

<https://arxiv.org/abs/2608.29463>

46 CONF

Chain-of-Thought Faithfulness of Reasoning Models Varies with Where and How Preference Cues Are Delivered

ArXiv CS.AI 50%

arXiv:2608.29464v1 Announce Type: cross Abstract: Chain-of-thought (CoT) monitoring assumes that reasoning traces faithfully record the information that shapes a model's answer. Existing faithfulness tests often place explicit bias cues in the user message, while agents may...

<https://arxiv.org/abs/2608.29464>

47 CONF

Applications of Risk Science to AI Fairness Evaluation: Principles, Challenges, and Best Practices

ArXiv CS.AI 50%

arXiv:2608.29478v1 Announce Type: cross Abstract: Scholarly work which aims to describe potential societal impacts (e.g., risks) of proliferating technology (especially related to artificial intelligence or other algorithmic systems) is likely to have an impact beyond the...

<https://arxiv.org/abs/2608.29478>

48 **CONF****Generalizable Multi-Agent Planning from Signal Temporal Logic Specifications via Diffusion****ArXiv CS.AI** **50%**

arXiv:2608.29490v1 Announce Type: cross Abstract: Multi-agent systems in the real-world (e.g., drone swarms, autonomous cars, warehouse robots) must satisfy rich, temporal tasks while avoiding collisions. Signal Temporal Logic (STL) elegantly encodes such objectives, but current...

<https://arxiv.org/abs/2608.29490>

49 **CONF****Argument-Aware Semantic Alignment of Normative Texts: A Toulmin-Based Neuro-Symbolic Approach****ArXiv CS.AI** **50%**

arXiv:2608.29529v1 Announce Type: cross Abstract: Semantic alignment between specialized normative texts is challenging when equivalent requirements use different terms, syntax, and levels of abstraction. Lexical overlap, distributional embeddings, and semantic similarity...

<https://arxiv.org/abs/2608.29529>

50 **CONF****SemPOI-RL: Aligning LLM Semantic Reasoning for Interpretable Out-of-Town POI Sequential Generation****ArXiv CS.AI** **50%**

arXiv:2608.30399v1 Announce Type: cross Abstract: Large language models (LLMs) exhibit strong semantic reasoning and open-ended generation abilities, but aligning these abilities with structured sequential generation remains challenging. This challenge is particularly evident in...

<https://arxiv.org/abs/2608.30399>

51 **CONF****SUP-MIMIC: A Multi-Task Clinical Diagnosis Benchmark for Evaluating LLMs' Robustness to Contradictory Evidence****ArXiv CS.AI** **50%**

arXiv:2608.29582v1 Announce Type: cross Abstract: Current evaluations of large language models (LLMs) primarily focus on factual knowledge retrieval, overlooking the fundamental challenge of navigating the complex, non-bijective mappings between clinical indicators and...

<https://arxiv.org/abs/2608.29582>

52 **CONF****MI-Distillation: Selecting from Model-Interpolated Instruct-Reasoning Data Spectrum for Chain-of-Thought Distillation****ArXiv CS.AI** **50%**

arXiv:2608.29623v1 Announce Type: cross Abstract: Recent advances in large reasoning models (LRMs) have shown strong performance on complex problems through long chain-of-thought (Long CoT) reasoning. However, distilling such trajectories into smaller student models remains...

<https://arxiv.org/abs/2608.29623>

53 **CONF****Conducting Stylistic Analysis of Paintings through an Art-History Agent****ArXiv CS.AI** **50%**

arXiv:2608.29644v1 Announce Type: cross Abstract: Attributing an artwork to an artist has traditionally relied on detailed visual observations and descriptions, known as stylistic analysis in art history. By contrast, current artificial intelligence (AI) models used in the field...

<https://arxiv.org/abs/2608.29644>

54 CONF

Higher-Dimensional Rotary Position Embedding

ArXiv CS.AI 50%

arXiv:2608.29715v1 Announce Type: cross Abstract: Transformers rely on position embedding mechanisms in long context modeling in most cases. Rotary Position Embedding (RoPE) embeds positional information with independent 2D rotations, forming relative position terms in...

<https://arxiv.org/abs/2608.29715>

55 CONF

R^{\$}^2\$A: Learning Persona Policies Through Persona Representation Learning and Runtime Alignment

ArXiv CS.AI 50%

arXiv:2608.29798v1 Announce Type: cross Abstract: The same Persona behavior can be beneficial in one context but harmful in another, causing static Persona elicitation to perform inconsistently across tasks. We introduce the Persona Selection--Realization Framework, which models...

<https://arxiv.org/abs/2608.29798>

56 CONF

REIGN: Refurbished Embeddings with Integrated Guidance Networks for Efficient Context-Length Scaling

ArXiv CS.AI 50%

arXiv:2608.29899v1 Announce Type: cross Abstract: Dense retrieval over long documents is expensive. Token-level encoders scale quadratically in sequence length, and most long-context embedding models reach 32K tokens only through architectural workarounds or by stretching...

<https://arxiv.org/abs/2608.29899>

57 CONF

INTERVenE: Temporal-Abstraction-Interval Based Transformers for Short-Horizon Medical Event Prediction

ArXiv CS.AI 50%

arXiv:2608.29901v1 Announce Type: cross Abstract: Electronic Health Record (EHR) prediction models in the intensive care unit must learn from sparse and irregular measurements while preserving the clinical meaning of time and supporting transparent decision-making. We present...

<https://arxiv.org/abs/2608.29901>

58 CONF

When Less is More: Understanding When Token Filtering Helps and Fails in AI-generated Text Detection

ArXiv CS.AI 50%

arXiv:2608.29903v1 Announce Type: cross Abstract: The rapid advancement of large language models (LLMs) has made AI-generated text detection increasingly critical. Existing zero-shot detectors assume that more token-level evidence leads to more reliable detection. However, our...

<https://arxiv.org/abs/2608.29903>

59 CONF

IndicDetect: Evaluating Cross-Lingual LLM-Generated Text Detection for Hindi, Telugu, and Tamil

ArXiv CS.AI 50%

arXiv:2608.29919v1 Announce Type: cross Abstract: The rapid proliferation of LLMs has further heightened the need to develop dependable AI-generated text detection, especially beyond English. Nevertheless, current benchmarks pay little attention to Indic languages and test...

<https://arxiv.org/abs/2608.29919>

60 CONF

Sleight of Word Benchmark: Can Language Models Notice If Their Own Output Was Tampered With?

ArXiv CS.AI 50%

arXiv:2608.29921v1 Announce Type: cross Abstract: The output of a Language Model can be tampered with \emph{while} the model is writing it. A simple test can thus be constructed by evaluating the model's perception of this external perturbation. In this spirit, a simple...

<https://arxiv.org/abs/2608.29921>

61 CONF

Influence Is Not Authority: When Causal Guardrail Signals Make Legitimate Tool Use Look Like an Attack in Tool-Using LLM Agents

ArXiv CS.AI 50%

arXiv:2608.29942v1 Announce Type: cross Abstract: The key limitation of current state-of-the-art influence-based guardrails is that they do not reliably distinguish a legitimate, user-authorized action from a malicious, unauthorized action when both rely on external tool...

<https://arxiv.org/abs/2608.29942>

62 CONF

The Policy Deficit in AI x Social-Emotional Learning Research

ArXiv CS.AI 50%

arXiv:2608.29950v1 Announce Type: cross Abstract: As artificial intelligence (AI) is increasingly integrated into social-emotional learning (SEL) initiatives, the need for evidence-based policy has become paramount. We systematically reviewed 65 peer-reviewed papers that examine...

<https://arxiv.org/abs/2608.29950>

63 CONF

Beyond Fluency: A Rubric-Based Benchmark for Evaluating Saudi Dialect and Cultural Competence in Large Language Models

ArXiv CS.AI 50%

arXiv:2608.29990v1 Announce Type: cross Abstract: Large language models are increasingly deployed in Arabic-speaking markets, yet standard benchmarks overwhelmingly reward Modern Standard Arabic (MSA) fluency while leaving dialectal and culturally grounded competence unmeasured....

<https://arxiv.org/abs/2608.29990>

64 CONF

"Act Like a 5th Grader" is Not Enough: Bounding Knowledge in LLM-Based User Simulators

ArXiv CS.AI 50%

arXiv:2608.30033v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used to simulate human behavior but frequently fail to exhibit realistic cognitive constraints, suffering from a "superhuman bias." Using a dataset of over 71,000 reading...

<https://arxiv.org/abs/2608.30033>

65 CONF

Reachability-Based Capability Confinement for LLM Agents under Indirect Prompt Injection

ArXiv CS.AI 50%

arXiv:2608.30041v1 Announce Type: cross Abstract: Large language model agents place outputs from external skills into their execution context, allowing attacker-controlled data to influence later privileged actions. Existing defenses mainly classify untrusted content or...

<https://arxiv.org/abs/2608.30041>

66 CONF

Forget or Fine-tune? A Comparative Study of Machine Unlearning Strategies for Noisy Label Correction

ArXiv CS.AI 50%

arXiv:2608.30046v1 Announce Type: cross Abstract: Noisy labels remain a critical challenge for training deep neural networks, since memorizing incorrect labels degrades generalization. Once noisy samples are identified after training, the standard solution is to retrain the...

<https://arxiv.org/abs/2608.30046>

67 CONF

Pak3H: Evaluating the Cost of Cultural Mismatch in LLM Alignment with a Human-Contextualized Urdu Benchmark

ArXiv CS.AI 50%

arXiv:2608.30065v1 Announce Type: cross Abstract: Large language models (LLMs) demonstrate strong Helpfulness, Harmlessness, and Honesty (3H) alignment in English-centric settings, but these gains transfer poorly to low-resource languages due to cultural mismatches. Existing...

<https://arxiv.org/abs/2608.30065>

68 CONF

Arkios: An Open Bilingual English-Nepali Language Model Trained From Scratch, with a Devanagari-Aware Tokenizer

ArXiv CS.AI 50%

arXiv:2608.30092v1 Announce Type: cross Abstract: We present Arkios, a 1.04B-parameter dense transformer pretrained from scratch on 150B tokens of bilingual English-Nepali text, using a custom single-file C/CUDA training stack and a Devanagari-aware byte-level BPE tokenizer...

<https://arxiv.org/abs/2608.30092>

69 CONF

Can LLMs Take the Pulse of the Economy? A Real-Time Evaluation of LLM Nowcasts on Macroeconomic Indicators

ArXiv CS.AI 50%

arXiv:2608.30110v1 Announce Type: cross Abstract: Nowcasting headline macroeconomic indicators, i.e., estimating an indicator's value for the current reference period before its official release, is critical for monetary policy and financial markets, and central banks devote...

<https://arxiv.org/abs/2608.30110>

70 CONF

E-SENS: Exclusion-Sensitive Penalization for Negative-Constraint Retrieval

ArXiv CS.AI 50%

arXiv:2608.30130v1 Announce Type: cross Abstract: Retrieval-augmented language models can fail to respect negative constraints when the retriever supplies evidence about concepts the user explicitly excluded. Beyond explicit negation, queries may ask for answers that include one...

<https://arxiv.org/abs/2608.30130>

71 CONF

CPR for LLMs: Critical-Point Routing against Catastrophic Forgetting in Domain Adaptation

ArXiv CS.AI 50%

arXiv:2608.30158v1 Announce Type: cross Abstract: Supervised fine-tuning (SFT) is the de facto standard for adapting large language models (LLMs) to target domains, but it often degrades the model's general capabilities, a phenomenon known as catastrophic forgetting. Existing...

<https://arxiv.org/abs/2608.30158>

72 CONF

The Differential Reasoning Router: Operationalizing Cost-Aware LLM Annotation in E-commerce

ArXiv CS.AI 50%

arXiv:2608.30224v1 Announce Type: cross Abstract: Large Language Models (LLMs) are increasingly used to annotate structured product data in e-commerce, but early deployment often begins as a cold-start problem: only limited pre-launch labels are available, the value of expensive...

<https://arxiv.org/abs/2608.30224>

73 CONF

ScenePilot: Grow-and-Repair Policy for Text-Driven 3D Indoor Scene Generation

ArXiv CS.AI 50%

arXiv:2608.30307v1 Announce Type: cross Abstract: Text-driven 3D indoor scene generation has advanced from dataset-bound layout modeling to open-vocabulary synthesis with large language and vision-language models. Yet existing methods remain limited: one-pass generators often...

<https://arxiv.org/abs/2608.30307>

74 CONF

Tail-Replay: Escaping the Curse of Linear Attention in Prefix Caching for Hybrid LLMs

ArXiv CS.AI 50%

arXiv:2608.30310v1 Announce Type: cross Abstract: Hybrid large language models interleave full-attention layers with linear-attention layers to reduce the cost of long-context inference. This structure complicates prefix caching: full-attention key-value caches are...

<https://arxiv.org/abs/2608.30310>

75 CONF

One AI Signal, Many Human Judgments: A Bayesian Cascade Analysis of AI-based Credibility Indicators in Online Information Spread

ArXiv CS.AI 50%

arXiv:2608.30311v1 Announce Type: cross Abstract: Social media platforms increasingly use AI-based credibility indicators to help users judge misinformation. Unlike individual human-AI decision-making, these indicators are embedded in information spread: users see both an AI...

<https://arxiv.org/abs/2608.30311>

76 CONF

Online Estimation of Dynamic Origin-Destination Matrices Using Reinforcement Learning with Link-Flow Propagation Guidance

ArXiv CS.AI 50%

arXiv:2608.30317v1 Announce Type: cross Abstract: Online dynamic origin-destination (OD) matrix estimation (DODE) calibrates time-dependent OD demand to reproduce observed link-flow trajectories. In online, OD demand should be estimated from current observations and propagated...

<https://arxiv.org/abs/2608.30317>

77 CONF

Beyond Token-Level Guidance: Inference-Time Alignment of Specialized LLMs via Cross-Family Representation Steering

ArXiv CS.AI 50%

arXiv:2608.30319v1 Announce Type: cross Abstract: Large language models (LLMs) finetuned for specialized domains represent crucial high-impact applications. Inference-time alignment improves safety degraded from specialization finetuning without requiring substantial...

<https://arxiv.org/abs/2608.30319>

78 **CONF****TopGQ: Fast GNN Post-Training Quantization Leveraging Topology Information****ArXiv CS.AI** **50%**

arXiv:2608.30394v1 Announce Type: cross Abstract: Existing GNN quantization methods suffer from considerable quantization overhead, which severely limits their practical usage in real-world scenarios. To this end, we present TopGQ, an accurate post-training GNN quantization...

<https://arxiv.org/abs/2608.30394>

79 **CONF****ImageEval 2026: Culturally Grounded Arabic Multimodal Evaluation****ArXiv CS.AI** **50%**

arXiv:2608.30475v1 Announce Type: cross Abstract: We present an overview of the ImageEval 2026 shared task on culturally grounded Arabic multimodal evaluation. It includes two tasks: (i) AynVQA, covering spoken visual question answering and image-grounded hallucination detection...

<https://arxiv.org/abs/2608.30475>

80 **CONF****Preference Shapes Relevance: Cross-component Hierarchical Semantic Alignment for Personalized Generative Retrieval****ArXiv CS.AI** **50%**

arXiv:2608.30553v1 Announce Type: cross Abstract: Generative Retrieval (GR) has emerged as a promising paradigm by mapping queries directly to Semantic IDs (SIDs) with powerful representation capabilities for candidate items. However, existing SIDs derived solely from item...

<https://arxiv.org/abs/2608.30553>

81 **CONF****DiffSAC: Diffusion-guided Sampling for Consensus-based Robust Estimation****ArXiv CS.AI** **50%**

arXiv:2608.30603v1 Announce Type: cross Abstract: Robust estimation is a core computer vision task frequently tackled using sample consensus. However, traditional methods suffer from inefficient sampling as they struggle to identify effective minimum sets before hypothesis...

<https://arxiv.org/abs/2608.30603>

82 **CONF****Generative Retrieval for E-commerce: Jointly Learning Embedding and Codebook with Same Product Cluster****ArXiv CS.AI** **50%**

arXiv:2608.30606v1 Announce Type: cross Abstract: With the development of large language models (LLMs), generative retrieval is becoming increasingly important in e-commerce scenarios. Current mainstream approaches typically use a two-stage training strategy: first train a...

<https://arxiv.org/abs/2608.30606>

83 **CONF****BiG-SURE - Bipartite Graph for Semantic Uncertainty and Reliability Estimation of LLMs****ArXiv CS.AI** **50%**

arXiv:2608.30646v1 Announce Type: cross Abstract: Reliable uncertainty estimation is a crucial requirement for deploying large language models (LLMs) and vision-language models (VLMs) in safety-critical settings, especially when the model parameters are not accessible...

<https://arxiv.org/abs/2608.30646>

84 CONF

Learning Dynamics of Logits Debiasing for Long-Tailed Semi-Supervised Learning

ArXiv CS.AI 50%

arXiv:2608.30699v1 Announce Type: cross Abstract: Long-tailed distributions are prevalent in real-world semi-supervised learning (SSL), where pseudo-labels tend to favor majority classes, leading to degraded generalization. While many long-tailed semi-supervised learning (LTSSL)...

<https://arxiv.org/abs/2608.30699>

85 CONF

RailSyn: Diagnosis-Guided Image Generation for Traceable Data Completion in Railway Foreign Object Detection

ArXiv CS.AI 50%

arXiv:2608.30709v1 Announce Type: cross Abstract: Railway foreign object detection (RFOD) is critical to safe railway operation, yet scarce real positive samples incompletely represent task-relevant variations in object scale, intrusion relation, railway scene, illumination, and...

<https://arxiv.org/abs/2608.30709>

86 CONF

RailGen: Improving Railway Intrusion Detection via Agent-Guided Small-Scale Foreign Object Generation

ArXiv CS.AI 50%

arXiv:2608.30727v1 Announce Type: cross Abstract: Small-object detection under long-tailed data distributions is a fundamental yet challenging problem in multimedia. Railway Foreign Object Detection (RFOD) epitomizes this challenge with easily confused small intrusions and...

<https://arxiv.org/abs/2608.30727>

87 CONF

Calibrating Small Language Models for Claim Check-Worthiness Detection

ArXiv CS.AI 50%

arXiv:2608.30731v1 Announce Type: cross Abstract: Assessing claim check-worthiness is an essential first step in automated fact-checking pipelines. This work is motivated by a real deployment challenge at an early-stage startup: running large language models (LLMs) over every...

<https://arxiv.org/abs/2608.30731>

88 CONF

Stick to What You Know: A Study of Knowledge-Aligned Supervised Fine-Tuning

ArXiv CS.AI 50%

arXiv:2608.30987v1 Announce Type: cross Abstract: Supervised fine-tuning (SFT) trains a base language model to imitate target responses, and these targets may require knowledge the base model has not robustly internalized. We study this as a source of hallucinations and frame a...

<https://arxiv.org/abs/2608.30987>

89 CONF

Pretrained, Curriculum-Tuned, and Ensembled: A Tracer-Aware Interactive Segmentation Pipeline for AutoPET V

ArXiv CS.AI 50%

arXiv:2608.30844v1 Announce Type: cross Abstract: Interactive lesion segmentation in whole-body PET/CT requires a model to provide a strong initial prediction while also responding efficiently to sparse corrective scribbles during inference. This setting is particularly...

<https://arxiv.org/abs/2608.30844>

90 CONF

TAMI: Temporally Aligned, Missingness-Aware, and Interpretable Multimodal Fusion for Mental Health Assessment in Older Adults with Mild Cognitive Impairment

ArXiv CS.AI 50%

arXiv:2608.30857v1 Announce Type: cross Abstract: Depression and anxiety in older adults with Mild Cognitive Impairment (MCI) are frequently underdiagnosed due to limited access to care. Multimodal analysis of remote clinical interviews is a scalable screening approach, but...

<https://arxiv.org/abs/2608.30857>

91 CONF

Personas Differ from Native-Language Generation: Language Pathways Shape LLM Interpersonal Advice

ArXiv CS.AI 50%

arXiv:2608.30873v1 Announce Type: cross Abstract: LLMs are increasingly used for interpersonal advice and as tools for studying social behavior across languages and cultures. A common shortcut for eliciting language- or culture-related variation is to ask a model to answer as a...

<https://arxiv.org/abs/2608.30873>

92 CONF

Evaluating and Improving LLM Self-Modeling

ArXiv CS.AI 50%

arXiv:2608.30980v1 Announce Type: cross Abstract: We study self-modeling: an LLM's ability to answer questions about its own behavior. We focus on verifiable behavioral questions, such as whether a prompt edit would change the model's final answer. To measure this capability, we...

<https://arxiv.org/abs/2608.30980>

93 CONF

Towards Stream Learning on Embedded Systems: Benchmarking the Memory Consumption of Stream Learning Methods

ArXiv CS.AI 50%

arXiv:2608.30923v1 Announce Type: cross Abstract: Stream learning is commonly evaluated through predictive performance and adaptation to concept drift. However, sustained operation of a stream learner also requires predictable and bounded resource usage even on long streams....

<https://arxiv.org/abs/2608.30923>

94 CONF

Stride-k Subsampling: Train-Free Audio Token Reduction for Whisper

ArXiv CS.AI 50%

arXiv:2608.30927v1 Announce Type: cross Abstract: Whisper exposes speech through a fixed 1500-token encoder interface, now a default representation for ASR decoders and Whisper-based speech language models (SpeechLMs), yet its redundancy remains largely unexamined. We propose...

<https://arxiv.org/abs/2608.30927>

95 CONF

MusGU+: Toward a Musician-Centered Evaluation Framework and Discovery Tool for Generative Music AI

ArXiv CS.AI 50%

arXiv:2608.30940v1 Announce Type: cross Abstract: Generative music systems are increasingly presented as tools that democratize music creation, yet their practical suitability for musicians remains underexplored. Prior work includes openness-focused evaluation frameworks, such...

<https://arxiv.org/abs/2608.30940>

96 CONF

AgentTRIM: Tool Risk Mitigation for Agentic AI

ArXiv CS.AI 50%

arXiv:2601.12449v2 Announce Type: replace-cross Abstract: AI agents are autonomous systems that combine LLMs with external tools to solve complex tasks. While such tools extend capability, improper tool permissions introduce security risks such as indirect prompt injection and...

<https://arxiv.org/abs/2601.12449>

97 CONF

LLM Judges Verify Presence, Not Absence: Omission Blindness in AI Clinical Notes and What Recovers It

ArXiv CS.AI 50%

arXiv:2608.31016v1 Announce Type: cross Abstract: Ambient AI scribes draft clinical notes, and published audits find their dominant error is omission: information the encounter established that the note fails to record. The standard check is an LLM judge: a second model reads...

<https://arxiv.org/abs/2608.31016>

98 CONF

One note in three: a verified census of three deployed AI scribes, and the instrument that counted it

ArXiv CS.AI 50%

arXiv:2608.31017v1 Announce Type: cross Abstract: Ambient AI scribes draft clinical notes under the reassurance that a clinician signs every note. We audited three commercial AI scribes on the same 142 consultations: 565 notes from recorded UK primary-care and US ambulatory...

<https://arxiv.org/abs/2608.31017>

99 CONF

LLM Post-Training as Brownfield Maintenance: An Industrial Perspective on Dataware Engineering

ArXiv CS.AI 50%

arXiv:2608.31102v1 Announce Type: cross Abstract: Industrial post-training is a brownfield regime. Teams inherit a deployed checkpoint and must land targeted improvements under fixed compute and mixture budgets without regressing the rest. The maintained artifact is increasingly...

<https://arxiv.org/abs/2608.31102>

00 CONF

AssetOpsBench: Benchmarking AI Agents for Task Automation in Industrial Asset Operations and Maintenance

ArXiv CS.AI 50%

arXiv:2506.03828v4 Announce Type: replace Abstract: AI for Industrial Asset Lifecycle Management aims to automate complex operational workflows, such as condition monitoring and maintenance scheduling, to minimize system downtime. While traditional AI/ML approaches solve narrow...

<https://arxiv.org/abs/2506.03828>

01 CONF

Beyond Pixels: Exploring DOM Downsampling for LLM-Based Web Agents

ArXiv CS.AI 50%

arXiv:2508.04412v4 Announce Type: replace Abstract: The advent of large language models (LLMs) has sparked an evolution of autonomous web browsing agents: given a web browsing task and serialised user interface (UI) state, an LLM is expected to suggest input actions that...

<https://arxiv.org/abs/2508.04412>

02 CONF

Stop Before You Fail: Operational Capability Boundaries for Mitigating Unproductive Reasoning in Large Reasoning Models

ArXiv CS.AI 50%

arXiv:2509.24711v4 Announce Type: replace Abstract: Current answering paradigms for Large Reasoning Models (LRMs) often fail to account for the fact that some questions may lie beyond the model's operational capability boundary, leading to long but unproductive reasoning. In...

<https://arxiv.org/abs/2509.24711>

03 CONF

MobileDreamer: Generative Sketch World Model for GUI Agent

ArXiv CS.AI 50%

arXiv:2601.04035v2 Announce Type: replace Abstract: Mobile GUI agents have shown strong potential in real-world automation and practical applications. However, most existing agents remain reactive, making decisions mainly from current screen, which limits their performance on...

<https://arxiv.org/abs/2601.04035>

04 CONF

SPRInG: Continual LLM Personalization via Selective Parametric Adaptation and Retrieval-Interpolated Generation

ArXiv CS.AI 50%

arXiv:2601.09974v3 Announce Type: replace Abstract: Personalizing Large Language Models typically relies on static retrieval or one-time adaptation, assuming user preferences remain invariant over time. However, real-world interactions are dynamic, where user interests...

<https://arxiv.org/abs/2601.09974>

05 CONF

Real-Time Deadlines Reveal Fragile Temporal Adaptation in LLM Strategic Dialogues

ArXiv CS.AI 50%

arXiv:2601.13206v2 Announce Type: replace Abstract: Large Language Models (LLMs) generate text token-by-token in discrete time, yet real-world communication, from therapy sessions to business negotiations, critically depends on continuous time constraints. We use simulated...

<https://arxiv.org/abs/2601.13206>

06 CONF

Defining Operational Conditions for Safety-Critical AI-Based Systems from Data

ArXiv CS.AI 50%

arXiv:2601.22118v3 Announce Type: replace Abstract: Artificial Intelligence (AI) has been on the rise in many domains, including numerous safety-critical applications. However, for complex systems in the real world, defining the underlying environmental conditions in which the...

<https://arxiv.org/abs/2601.22118>

07 CONF

Beyond Dense States: Sparse Transcoders as Causally Testable Operators for LLM Latent Reasoning

ArXiv CS.AI 50%

arXiv:2602.01695v2 Announce Type: replace Abstract: Latent reasoning reduces the token-generation cost of chain-of-thought reasoning by replacing explicit intermediate tokens with continuous latent transitions. However, existing latent reasoning methods usually rely on dense and...

<https://arxiv.org/abs/2602.01695>

08 CONF

AgentRx: Diagnosing AI Agent Failures from Execution Trajectories

ArXiv CS.AI 50%

arXiv:2602.02475v2 Announce Type: replace Abstract: AI agents often fail in ways that are difficult to localize because executions are probabilistic, long-horizon, multi-agent, and mediated by noisy tool outputs. We address this gap by manually annotating failed agent runs and...

<https://arxiv.org/abs/2602.02475>

09 CONF

Towards Natural Personalization: Evaluating Long-Horizon Preference Following in Personalized User-LLM Interactions

ArXiv CS.AI 50%

arXiv:2603.04191v2 Announce Type: replace Abstract: Large Language Models (LLMs) are increasingly serving as personal assistants, where users may share individual preferences over extended interactions. However, assessing how well LLMs can follow these preferences in natural,...

<https://arxiv.org/abs/2603.04191>

10 CONF

A Multi-Agent Human-LLM Collaborative Framework for Closed-Loop Scientific Literature Summarization

ArXiv CS.AI 50%

arXiv:2604.01452v2 Announce Type: replace Abstract: Scientific discovery is slowed by fragmented literature that requires excessive human effort to gather, analyze, and understand. AI tools, including autonomous summarization and question answering, have been developed to aid in...

<https://arxiv.org/abs/2604.01452>

11 CONF

SciAtlas: A Computable Atlas of Science for Knowledge-Grounded AI Research

ArXiv CS.AI 50%

arXiv:2605.22878v2 Announce Type: replace Abstract: Artificial intelligence is rapidly entering the core workflows of scientific research. Yet reliable scientific reasoning requires access to accumulated scientific knowledge with sufficient breadth, depth, and standardization....

<https://arxiv.org/abs/2605.22878>

12 CONF

PortBench: A Correlation-Aware, Full-Pipeline Benchmark for LLM-Driven Portfolio Management

ArXiv CS.AI 50%

arXiv:2605.27887v4 Announce Type: replace Abstract: Large language models (LLMs) have shown strong performance across diverse financial tasks, yet portfolio management (PM) remains poorly benchmarked. Existing benchmarks exhibit two gaps: they are often equity-only and ignore...

<https://arxiv.org/abs/2605.27887>

13 CONF

Hidden Thoughts Are Not Secret: Reasoning Trace Exposure in LLMs

ArXiv CS.AI 50%

arXiv:2606.00642v2 Announce Type: replace Abstract: Reasoning traces have become a valuable form of learning signals for improving and transferring the capabilities of large language models. In particular, detailed traces can help distill reasoning behavior from stronger teacher...

<https://arxiv.org/abs/2606.00642>

14 CONF

When Should Models Change Their Minds? Contextual Belief Management in Large Language Models

ArXiv CS.AI 50%

arXiv:2605.30219v2 Announce Type: replace Abstract: Long-horizon interactions require language models to manage accumulating information: when to update their state, when to preserve their state, and what to ignore. We study this challenge as Contextual Belief Management (CBM):...

<https://arxiv.org/abs/2605.30219>

15 CONF

Self-Correction Can Amplify Hallucinations: Fact-Level Repair with Graph-Based Evidence Routing in Multimodal Generation

ArXiv CS.AI 50%

arXiv:2606.00232v2 Announce Type: replace Abstract: We study fact-level repair for multimodal generation, where a fluent output may contain specific facts that are not supported by the input. Existing inference-time repair methods often generate feedback by jointly conditioning...

<https://arxiv.org/abs/2606.00232>

16 CONF

Beyond Helpfulness: A Teaching-over-Solving Diagnostic for Measuring Educational Impact in LLM Tutors

ArXiv CS.AI 50%

arXiv:2606.16206v2 Announce Type: replace Abstract: Large language models are increasingly proposed as educational tutors, yet stronger task-solving ability does not necessarily imply stronger learning support. Motivated by recent calls to measure the social impact of NLP...

<https://arxiv.org/abs/2606.16206>

17 CONF

DART: Draft-Agreement Routing for Training-Free Adaptive Thinking Budgets in Hybrid Reasoning Models

ArXiv CS.AI 50%

arXiv:2606.23181v2 Announce Type: replace Abstract: Hybrid reasoning models can answer directly or spend extra tokens on extended thinking. A practical router should choose between these modes for each query, so easy problems avoid unnecessary reasoning and hard problems receive...

<https://arxiv.org/abs/2606.23181>

18 CONF

PsychoAgent: An Affect-Sensitive Cognitive Architecture for Conflict-Aware Memory in LLM Agents

ArXiv CS.AI 50%

arXiv:2608.07438v2 Announce Type: replace Abstract: Human-like cognition does not select past experience by topical similarity alone: affective significance and unresolved conflict also shape what becomes accessible. We present PsychoAgent, a cognitive architecture for LLM...

<https://arxiv.org/abs/2608.07438>

19 CONF

Robustness Analysis of Agentic AI to Inconsistent and Incomplete Tool Responses

ArXiv CS.AI 50%

arXiv:2608.22676v2 Announce Type: replace Abstract: Tool-using agents increasingly rely on external tools to complete multi-step tasks, but tool returns can fail in different ways and require different recovery actions. Existing robustness studies often use uncertainty-based...

<https://arxiv.org/abs/2608.22676>

20 **CONF****When Seeing Is Not Enough: Benchmarking Interactive Visual Grounding in LVLMs****ArXiv CS.AI 50%**

arXiv:2608.23978v2 Announce Type: replace Abstract: Visual grounding is typically evaluated as a one-shot mapping from an informative referring expression to a visual target. This formulation misses a central property of real-world reference: initial referring expressions are...

<https://arxiv.org/abs/2608.23978>

21 **CONF****post-graph-rag: A PostgreSQL-Native Bi-Temporal Graph RAG Engine with Temporal Grounding at Synthesis****ArXiv CS.AI 50%**

arXiv:2608.24921v2 Announce Type: replace Abstract: Graph RAG connects facts no single passage states, but implementations pay three times: in infrastructure, keeping vector store, graph database and document store in sync; in quality, because a pipeline that never refuses...

<https://arxiv.org/abs/2608.24921>

22 **CONF****Candidate supply and answer selection shape the value of LLM judging in multi-agent systems****ArXiv CS.AI 50%**

arXiv:2608.25937v2 Announce Type: replace Abstract: Multi-agent systems (MAS) sometimes already have the potential to answer correctly, but still report a wrong answer. Explaining this outcome is difficult because generation, communication and final answer-selection rules...

<https://arxiv.org/abs/2608.25937>

23 **CONF****LongGuard: Mechanistic Analysis and Training-Free Mitigation of Long-Context Failure in Safety Guardrails****ArXiv CS.AI 50%**

arXiv:2608.27580v2 Announce Type: replace Abstract: Safety guardrails serve as the last line of defense against harmful inputs and outputs of large language models (LLMs), yet they are trained and evaluated almost exclusively on short text. We present LongGuard, a framework that...

<https://arxiv.org/abs/2608.27580>

24 **CONF****Multimodal Large Language Models Predict Urban Safety Perception but Encode Non-Neutral Demographic Priors****ArXiv CS.AI 50%**

arXiv:2503.00610v2 Announce Type: replace-cross Abstract: Understanding how people perceive urban environments is essential for inclusive planning, yet conventional surveys are costly and difficult to scale. We investigate whether Multimodal Large Language Models (MLLMs) can...

<https://arxiv.org/abs/2503.00610>

25 **CONF****HeTGB: A Comprehensive Benchmark for Heterophilic Text-Attributed Graphs****ArXiv CS.AI 50%**

arXiv:2503.04822v2 Announce Type: replace-cross Abstract: Graph neural networks (GNNs) have demonstrated success in modeling relational data primarily under the assumption of homophily. However, many real-world graphs exhibit heterophily, where linked nodes belong to different...

<https://arxiv.org/abs/2503.04822>

26 CONF

Watch your steps: Dormant Adversarial Behaviors that Activate upon LLM Finetuning

ArXiv CS.AI 50%

arXiv:2505.16567v4 Announce Type: replace-cross Abstract: Finetuning open-weight Large Language Models (LLMs) is standard practice for achieving task-specific performance improvements. Until now, finetuning has been regarded as a controlled and secure process in which training...

<https://arxiv.org/abs/2505.16567>

27 CONF

EquiReg: Equivariance Regularized Diffusion for Inverse Problems

ArXiv CS.AI 50%

arXiv:2505.22973v3 Announce Type: replace-cross Abstract: Diffusion models represent the state-of-the-art for solving inverse problems such as image restoration tasks. Diffusion-based inverse solvers incorporate a likelihood term to guide prior sampling, generating data...

<https://arxiv.org/abs/2505.22973>

28 CONF

Training-free LLM Verification via Recycling Few-shot Examples

ArXiv CS.AI 50%

arXiv:2506.17251v3 Announce Type: replace-cross Abstract: Although large language models (LLMs) have achieved remarkable performance, the inherent stochasticity of their reasoning processes and varying conclusions present significant challenges. Majority voting or Best-of-N with...

<https://arxiv.org/abs/2506.17251>

29 CONF

Proof2Hybrid: Automatic Mathematical Benchmark Synthesis for Proof-Centric Problems

ArXiv CS.AI 50%

arXiv:2508.02208v3 Announce Type: replace-cross Abstract: Evaluating the mathematical capability of Large Language Models (LLMs) is a critical yet challenging frontier. Existing benchmarks fall short, particularly for proof-centric problems, as manual creation is unscalable and...

<https://arxiv.org/abs/2508.02208>

30 CONF

StructSynth: Dependency Graphs as Generation Plans for Low-Data Tabular Synthesis with Language Models

ArXiv CS.AI 50%

arXiv:2508.02601v2 Announce Type: replace-cross Abstract: Tabular data derives its value from inter-feature dependencies, yet preserving them during synthesis is fragile when samples are scarce. Existing approaches either learn dependencies implicitly through distribution...

<https://arxiv.org/abs/2508.02601>

31 CONF

Turning the Spell Around: Lightweight Alignment Amplification via Rank-One Safety Injection

ArXiv CS.AI 50%

arXiv:2508.20766v2 Announce Type: replace-cross Abstract: Safety alignment in Large Language Models (LLMs) often involves mediating internal representations to refuse harmful requests. Recent research has demonstrated that these safety mechanisms can be bypassed by ablating or...

<https://arxiv.org/abs/2508.20766>

32 CONF

Operationalising AI Regulatory Sandboxes: Activities, Requirements, and Technical Assessment under the EU AI Act

ArXiv CS.AI 50%

arXiv:2509.25256v4 Announce Type: replace-cross Abstract: The systematic assessment of AI systems is increasingly vital as these technologies enter high-stakes domains. To address this, the EU's Artificial Intelligence Act introduces AI Regulatory Sandboxes (AIRS): supervised...

<https://arxiv.org/abs/2509.25256>

33 CONF

Correctness Forensics for Batch Speculative Decoding: Diagnosing the Ragged Tensor Problem

ArXiv CS.AI 50%

arXiv:2510.22876v4 Announce Type: replace-cross Abstract: Inference optimizations are routinely evaluated by throughput alone, without verifying output correctness. We conduct a forensic analysis of batch speculative decoding and find that several widely-used implementations...

<https://arxiv.org/abs/2510.22876>

34 CONF

ScalePRM: Training Process Reward Models by Scaling Verification Compute Without Ground Truth

ArXiv CS.AI 50%

arXiv:2512.03244v2 Announce Type: replace-cross Abstract: Training process reward models (PRMs) requires step-level correctness labels, obtained either through expensive human annotation or by relying on ground-truth answers, limiting the ability to scale process-level...

<https://arxiv.org/abs/2512.03244>

35 CONF

Kinship Data Benchmark for Multi-hop Reasoning

ArXiv CS.AI 50%

arXiv:2601.07794v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly evaluated on their ability to perform multi-hop reasoning, i.e., to combine multiple pieces of information into a coherent inference. We introduce KinshipQA, a benchmark...

<https://arxiv.org/abs/2601.07794>

36 CONF

Attention-Discounted Adaptive Sampler for Masked Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2606.10829v3 Announce Type: replace-cross Abstract: Masked diffusion language models can reduce inference steps by revealing multiple tokens per denoising iteration, but this parallelism is fragile: positions that are individually confident may be unsafe to commit together...

<https://arxiv.org/abs/2606.10829>

37 CONF

Securing Time Integrity in Energy IoT Against Clock Drift and Y2K38 Failures

ArXiv CS.AI 50%

arXiv:2601.23147v3 Announce Type: replace-cross Abstract: Time integrity across distributed Internet of Things (IoT) devices is fundamental to reliable sensing, control, and security in energy cyber-physical systems. However, operational energy IoT systems remain vulnerable to...

<https://arxiv.org/abs/2601.23147>

38 CONF

Robustness of Vision Language Models Against Split-Image Harmful Input Attacks

ArXiv CS.AI 50%

arXiv:2602.08136v2 Announce Type: replace-cross Abstract: Vision-Language Models (VLMs) are now a core part of modern AI. Recent work proposed several visual jailbreak attacks using single/ holistic images. However, contemporary VLMs demonstrate strong robustness against such...

<https://arxiv.org/abs/2602.08136>

39 CONF

Reverse N-Wise Output-Oriented Testing for AI/ML and Quantum Computing Systems

ArXiv CS.AI 50%

arXiv:2602.14275v2 Announce Type: replace-cross Abstract: Artificial intelligence/machine learning (AI/ML) systems and emerging quantum computing software present unprecedented testing challenges characterized by high-dimensional/continuous input spaces,...

<https://arxiv.org/abs/2602.14275>

40 CONF

Toward Generalizable Deep Learning Based Peatland Fire Detection via Walsh Hadamard Transform and Domain Adaptation

ArXiv CS.AI 50%

arXiv:2603.02465v2 Announce Type: replace-cross Abstract: Machine learning-based wildfire detection has advanced significantly using deep learning models trained on large wildfire image and video datasets. However, peatland fires exhibit distinct characteristics, including...

<https://arxiv.org/abs/2603.02465>

41 CONF

Risk-Adjusted Harm Scoring for Automated Red Teaming for LLMs in Financial Services

ArXiv CS.AI 50%

arXiv:2603.10807v2 Announce Type: replace-cross Abstract: Existing LLM safety evaluations rely on binary attack-success rates and domain-agnostic taxonomies, leaving regulated Banking, Financial Services, and Insurance (BFSI) deployments exposed to failures elicited through...

<https://arxiv.org/abs/2603.10807>

42 CONF

Mixture-Greedy for Online Generative Model Selection: Is UCB Necessary in Diversity-Aware Multi-Armed Bandits?

ArXiv CS.AI 50%

arXiv:2603.21716v2 Announce Type: replace-cross Abstract: Efficient selection among multiple generative models is increasingly important in modern generative AI, where sampling from suboptimal models is costly. This problem can be viewed as a multi-armed bandit (MAB) task. Under...

<https://arxiv.org/abs/2603.21716>

43 CONF

Do We Still Need Humans in the Loop? Human vs. LLM Annotation in Active Learning for TikTok Hate Speech Detection

ArXiv CS.AI 50%

arXiv:2604.13899v5 Announce Type: replace-cross Abstract: Annotating data remains a costly bottleneck for supervised NLP. Active learning (AL) reduces the number of human labels needed by selecting only the most informative instances, while instruction-tuned LLMs attack the same...

<https://arxiv.org/abs/2604.13899>

44 CONF

Perturbation Sensitivity of Maximum-Likelihood Pairwise Ranking in Computational Decision Systems

ArXiv CS.AI 50%

arXiv:2604.17805v3 Announce Type: replace-cross Abstract: Maximum-likelihood pairwise ranking is a common computational mechanism for prioritization, reputation estimation, and comparison-driven decision support. Despite its broad use, the perturbation sensitivity of this...

<https://arxiv.org/abs/2604.17805>

45 CONF

Making the Discrete Continuous: Synthetic RAW Augmentations for Fine-Grained Evaluation of Person Detection Performance in Low Light

ArXiv CS.AI 50%

arXiv:2605.22455v2 Announce Type: replace-cross Abstract: Real-world deployment of AI vision models is both fueled and limited by the data available for training and testing. Real datasets are sparse and uneven: long-tailed or unbalanced distributions hinder generalization, and...

<https://arxiv.org/abs/2605.22455>

46 CONF

Semantic Flow Regularization: Teaching LLMs to Generate Diverse Yet Coherent Responses

ArXiv CS.AI 50%

arXiv:2605.27971v2 Announce Type: replace-cross Abstract: When large language models are fine-tuned to generate persona- or tone-conditioned responses, their output diversity is severely limited--a failure we term Cross-Style Collapse. We trace this collapse to the cross-entropy...

<https://arxiv.org/abs/2605.27971>

47 CONF

Reverse Probing: Supervised Token-level Uncertainty Quantification for Large Language Models in Clinical Text

ArXiv CS.AI 50%

arXiv:2605.28740v2 Announce Type: replace-cross Abstract: As large language models are increasingly deployed for clinical text, ensuring they can reliably signal their own uncertainty becomes critical. Most existing uncertainty quantification (UQ) methods are designed for...

<https://arxiv.org/abs/2605.28740>

48 CONF

DASH: Dual-Branch Score Distillation for Guidance-Calibrated Compact Diffusion Models

ArXiv CS.AI 50%

arXiv:2606.00798v2 Announce Type: replace-cross Abstract: Parameter compression of class-conditional diffusion models exposes a structural limitation in output-level distillation: supervising only the guided output leaves the two score branches non-identifiable, so the...

<https://arxiv.org/abs/2606.00798>

49 CONF

Divide-and-Conquer Modeling for the CTF-4-Science Lorenz Benchmark

ArXiv CS.AI 50%

arXiv:2606.10084v2 Announce Type: replace-cross Abstract: This submission documents the divide-and-conquer modeling strategy developed for the CTF-4-Science Lorenz Chaotic Systems Challenge at AI-DEEDS 2026. The challenge uses the CTF-4-Science Lorenz benchmark to evaluate...

<https://arxiv.org/abs/2606.10084>

50 CONF

Cultural Bias Without a Cultural Self: A Disassociation Study of LLM's Persona and Bias

ArXiv CS.AI 50%

arXiv:2607.02368v3 Announce Type: replace-cross Abstract: Language models prompted with cultural personas increasingly stand in for human respondents in cross-cultural research. Their responses separate personas cleanly, and that separation is read as evidence of a cultural...

<https://arxiv.org/abs/2607.02368>

51 CONF

Criterion-Conditional In-Context Learning: Evaluating Criterion-Shift Adaptation in Vision-Language Models

ArXiv CS.AI 50%

arXiv:2607.02575v2 Announce Type: replace-cross Abstract: Vision-language models can perform new tasks without parameter updates through in-context learning (ICL), whose core mechanism is utilizing the support set for task induction. In the standard ICL setting, once the task is...

<https://arxiv.org/abs/2607.02575>

52 CONF

A Few Teacher Steps Go a Long Way: Cost-Efficient On-Policy Data Augmentation for Agent Post-Training

ArXiv CS.AI 50%

arXiv:2607.04574v2 Announce Type: replace-cross Abstract: For LLM agents, supervised fine-tuning is not only about teacher labels' quality, but also about which interaction contexts those labels condition on. Pure behavioral cloning uses full teacher demonstrations, creating a...

<https://arxiv.org/abs/2607.04574>

53 CONF

Decoupled Structure-Feature Alignment via Alternating Optimization for Graph Learning

ArXiv CS.AI 50%

arXiv:2607.11577v2 Announce Type: replace-cross Abstract: Conventional Graph Neural Networks (GNNs) couple feature transformation and neighborhood aggregation, which often renders them vulnerable to topological noise and heterophilous connections. To decouple this dependency, we...

<https://arxiv.org/abs/2607.11577>

54 CONF

TCS-BENCH: Benchmarking State-of-the-Art Generative AI Theoretical Computer Science Research Ability

ArXiv CS.AI 50%

arXiv:2608.09538v3 Announce Type: replace-cross Abstract: We introduce TCS-Bench, a benchmark for evaluating Large Language Models (LLMs) on research-level Theoretical Computer Science (TCS) proof generation. TCS-Bench consists of theorem-proving tasks from papers published at...

<https://arxiv.org/abs/2608.09538>

55 CONF

Modeling the Structure of Human Behavior with AI Prompt Vectors

ArXiv CS.AI 50%

arXiv:2608.18265v2 Announce Type: replace-cross Abstract: We introduce a general, easy-to-implement AI-based method for modeling and analyzing the structure and complexity of human behavior. We assign a large language model a "type vector" and then prompt it to choose actions...

<https://arxiv.org/abs/2608.18265>

56 CONF

Decomposing Wrong-Consensus Agreement in LLM Self-Consistency

ArXiv CS.AI 50%

arXiv:2608.18795v2 Announce Type: replace-cross Abstract: Agreement among repeated samples of a language model is routinely read as evidence about answer reliability, yet wrong answers can agree just as strongly as right ones. This paper asks what information wrong-consensus...

<https://arxiv.org/abs/2608.18795>

57 CONF

What Does an Evaluation License? A Commit-Bound Census of Claim Replay in Inspect Evals

ArXiv CS.AI 50%

arXiv:2608.19269v4 Announce Type: replace-cross Abstract: Benchmarks can run without determining what their results license. We freeze a large evaluation collection and attempt to replay its historical claims. Most units stop because the evidence required for replay is not...

<https://arxiv.org/abs/2608.19269>

58 CONF

A Statistical Audit of Physical AI Benchmark Redundancy

ArXiv CS.AI 50%

arXiv:2608.25940v2 Announce Type: replace-cross Abstract: Physical AI models are evaluated on suites of benchmarks that differ across model reports, leaving the model-by-benchmark matrix sparse and the relationship between benchmarks unmeasured. We construct a matrix of 51...

<https://arxiv.org/abs/2608.25940>

59 CONF

ICON Decomposition: Auditing Deep Neural Networks with Multivariate Variance-based Concept-level Explanations

ArXiv CS.AI 50%

arXiv:2608.26083v2 Announce Type: replace-cross Abstract: Deep neural networks often exploit spurious associations, a failure known as shortcut learning. Auditing for shortcuts requires testing many candidate concepts, such as acquisition settings or demographics. Current...

<https://arxiv.org/abs/2608.26083>

60 CONF

Deriving Scaling Laws for OpenEuroLLM Models: Learning Rate, Batch Size and Loss

ArXiv CS.AI 50%

arXiv:2608.28308v2 Announce Type: replace-cross Abstract: We study the scaling behavior of learning rate and batch size in pretraining dense large language models on English-prevalent corpora. Beyond scaling jointly optimal learning rates and batch sizes, we investigate their...

<https://arxiv.org/abs/2608.28308>