



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	152	Media
The Verge AI	8	Media
TechCrunch AI	8	Media
HackerNews	7	Community
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Anthropic launches Claude Fable 5.1 and says it's up to 45 percent cheaper for agentic work

The Verge AI

+1 source

80%

Anthropic says its newest AI models, Fable 5.1 and Mythos 5.1, address criticisms from customers about price, data retention, and overzealous safeguards. The company claims Claude Fable 5.1 offers stronger performance than Fable 5, but costs around 25 percent less typically and...

<https://www.theverge.com/ai-artificial-intelligence/987830/anthropic-claude-fable-mytho...>

VERI

John Deere launched an AI chatbot for farmers

The Verge AI

+1 source

80%

John Deere is testing a new "JD" AI assistant that it says can help farmers make more money, with answers about best practices and historical trends that are based on their own data. It uses their "field, machine and operational data" to answer questions on topics like equipment...

<https://www.theverge.com/ai-artificial-intelligence/987486/john-deere-jd-ai-chatbot>

VERI

SARTM: Segment Any RGB Thermal Model with Language aided Distillation

ArXiv CS.AI

+1 source

80%

arXiv:2505.01950v2 Announce Type: replace-cross Abstract: The recent Segment Anything Model (SAM) demonstrates strong instance segmentation performance across various downstream tasks. However, SAM is trained solely on RGB data, limiting its direct applicability to RGB-thermal...

<https://arxiv.org/abs/2505.01950>

CONFIRMED SIGNALS — 172 stories

CONF

Show HN: Weedout - Safari extension that hides YouTube AI-labeled videos

HackerNews

60%

HN 125

I'm the developer. I built this because AI-generated conspiracy videos were getting out of control in my own YouTube feed. Weedout is a \$1.99 Safari extension for macOS that removes videos YouTube labels "Made with AI" from feed, search, related videos, playlists and Shorts. It...

<https://masteranza.github.io/weedout/>

CONF

Show HN: What Happens When You Give Your AI Agents a Voice and an Attitude

HackerNews

50%

HN 10

If you are using React, you've probably used or heard of React Email by Resend. Ironically once we started building Unitpost we hit the limit: React Email works when an engineer owns a .tsx file. Add a designer or an agent and you get a second copy that drifts. So we built...

<https://fellowgeek.github.io/mcp-speak/>

CONF

Show HN: Claude and ChatGPT need a datacenter. This runs on my phone

HackerNews

50%

HN 9

This week-end, I wondered if we could add a little fun animation in Claude Code CLI's status line on macOS, based on the LLM current work. First I created a fish tank, then a safari scene, only using emojis. Being an 80's kid, I challenged Fable 5 to add a Ms Pac-Man, with good...

<https://llmobi.pages.dev>

CONF

Show HN: I built a live stream site where every spot is instantly AI-generated

HackerNews 50% HN 9

No summary — see link for full article.

<https://videoai.me/live>

CONF

Show HN: Rails Baseline - a Rails SaaS starter with agent-readable architecture

HackerNews 50% HN 4

I have built a lot of SaaS Rails apps over the last decade+...and agents are making their way into my workflow more then ever. So I decided I wanted a starter that picks a LOT of vanilla Rails flavors, a little bit of battle tested additions and a lot of agent discoverable...

<https://railsbaseline.com/>

CONF

Show HN: Three Pins a daily 3-hole golf game in the browser

HackerNews 50% HN 4

No summary — see link for full article.

<https://threepins.golf/>

0 CONF

Show HN: OSS, K8s-native AI platform for distributed multi-model inference

HackerNews 50% HN 4

Hi everyone, I'm one of the co-founders of axem. We recently open sourced Shaide, a project we've been working on to make running multiple LLMs on your own infrastructure less painful. It started pretty simply. Running one model wasn't the hard part. The hard part came when we...

<https://github.com/axem-solutions/shaide>

1 CONF

OpenAI's Astra model is on the way — and very good at breaking into computer systems

TechCrunch AI 50%

OpenAI previewed the precautions it is taking as it prepares to release Astra, its newest, cyber-critical LLM.

<https://techcrunch.com/2026/09/01/open-ais-astra-model-is-on-the-way-and-very-good-at-b...>

2 CONF

Google's answer to Canva is an AI tool where you prompt instead of design

TechCrunch AI 50%

With Google Pics, Google is pushing deeper into the creative software market dominated by Canva and Adobe, but with a distinctly AI-first approach.

<https://techcrunch.com/2026/09/01/googles-answer-to-canva-is-an-ai-tool-where-you-promp...>

3 CONF

ChatGPT Health adds Epic integration for clinicians to import patient data

TechCrunch AI 50%

OpenAI said that the integration provides read-only access to health records for clinicians.

<https://techcrunch.com/2026/09/01/chatgpt-health-adds-epic-integration-for-clinicians-t...>

4 CONF

AIR raises \$50M to help companies vet the skills and add-ons AI agents use

TechCrunch AI 50%

AIR's platform can discover agents running at a company, continuously vets any skills and add-ons they use, and blocks any unwanted behavior.

<https://techcrunch.com/2026/09/01/air-raises-50m-to-help-companies-vet-the-skills-and-a...>

5 CONF

The Pentagon now has its own version of ChatGPT and Grok

TechCrunch AI 50%

Versions of OpenAI's ChatGPT and SpaceXAI's Grok will join Google's Gemini on the Pentagon's central portal for AI tools.
<https://techcrunch.com/2026/08/31/the-pentagon-now-has-its-own-version-of-chatgpt-and-g...>

6 CONF

Clipto uses AI to search terabytes of video and is now valued at \$250M

TechCrunch AI 50%

The three-year-old startup says it reached \$15 million in ARR and profitability before raising its latest \$15 million round.
<https://techcrunch.com/2026/08/31/three-year-old-ai-media-search-startup-clipto-hits-a-...>

7 CONF

How AI plotted an interstellar journey to Alpha Centauri

MIT Tech Review 50%

A nonprofit organization called the Fermi Explorer Mission announced today that it intends to launch a spacecraft to our nearest star system by the end of 2029. It's a hugely ambitious mission—if all goes well, the spacecraft could take up to 80,000 years to arrive at Alpha...
<https://www.technologyreview.com/2026/09/01/1143247/ai-interstellar-journey-alpha-centa...>

8 CONF

Making the AI-powered case for legacy modernization

MIT Tech Review 50%

For years, legacy technology has been a problem companies knew they needed to solve, but one they often struggled to tackle. The cost, complexity, and risk of replacing business-critical systems could make modernization feel like a disruption to manage instead of an opportunity...
<https://www.technologyreview.com/2026/09/01/1142180/making-the-ai-powered-case-for-lega...>

9 CONF

The Download: engineered microbes for crops, and OpenAI's culture problem

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. How engineered microbes could help feed the world's crops Fertilizer is crucial for the global food supply, but making it uses a lot of...
<https://www.technologyreview.com/2026/09/01/1143199/the-download-engineered-microbes-op...>

0 CONF

Google needs Hollywood more than the studios need AI

The Verge AI 50%

Google has reportedly been reaching out to a number of Hollywood's biggest studios, hoping to strike licensing agreements that would allow it to train its AI models on copyrighted material in exchange for massive piles of cash. In theory, these deals would be a win-win: a huge...
<https://www.theverge.com/tech/987429/google-needs-hollywood-more-than-the-studios-need-ai>

1 CONF

OpenAI delayed its new model's development after the Hugging Face hack

The Verge AI 50%

After an unreleased OpenAI model wreaked enough havoc to make international headlines, OpenAI delayed the development of a different unreleased model suite, Astra, in order to shore up its safety work, the company wrote Tuesday in a blog post. In July, an unreleased OpenAI model...
<https://www.theverge.com/ai-artificial-intelligence/987695/openai-astra-unreleased-mode...>

2 **CONF****The rise of AI ‘civilizations’ and the fall of corporate responsibility****The Verge AI** 50%

Depending on who you ask, developer platform Hugging Face was recently attacked by OpenAI - after it lost control of its own AI tools - or by a succession of AI "civilizations." Welcome to the linguistic battlefield of AI safety, where word choices can shift responsibility for a...

<https://www.theverge.com/ai-artificial-intelligence/987566/ai-civilizations-opeai-huggi...>

3 **CONF****Apple accuses OpenAI of destroying evidence****The Verge AI** 50%

Apple is pushing for "expedited discovery" in its legal battle against OpenAI over concerns the company is actively destroying evidence, as reported earlier by Bloomberg. In a filing on Monday, Apple alleges OpenAI only just handed over a MacBook used by a former employee at the...

<https://www.theverge.com/tech/987575/apple-openai-destroying-evidence-trade-secrets-law...>

4 **CONF****Google Pics is like Canva, but with even more AI****The Verge AI** 50%

Google has a new suite of creative design tools for Workspace users called Google Pics, which aims to make editing and generating "professional-grade" AI images less cumbersome for businesses. Built around Gemini and the Nano Banana generative AI model, Google Pics is designed...

<https://www.theverge.com/tech/987423/google-pics-ai-image-editor-generator>

5 **CONF****CARE: Contrastive Anchor-based Rubric Evolution for Large Language Model Post-Training****ArXiv CS.AI** 50%

arXiv:2609.00892v1 Announce Type: new Abstract: Rubric-based reinforcement learning decomposes open-ended instructions into prompt-specific, flexible rubrics, making it better suited than reinforcement learning with verifiable rewards for post-training LLMs on open-ended tasks....

<https://arxiv.org/abs/2609.00892>

6 **CONF****Long-Horizon State Tracking in LLMs: Executing MD5 through a Deep Sequence of Dependent Tool Calls****ArXiv CS.AI** 50%

arXiv:2609.00012v1 Announce Type: new Abstract: Long-horizon tasks remain uncommon in large language model (LLM) evaluation, and for a reason: when each step depends on the last, per-step accuracy that looks excellent in isolation decays catastrophically, as errors cascade and...

<https://arxiv.org/abs/2609.00012>

7 **CONF****OpenAgentFlow: Enabling System-Wide Safety Boundaries for Heterogeneous AI Agent Fleets****ArXiv CS.AI** 50%

arXiv:2609.00015v1 Announce Type: new Abstract: AI agents powered by large language models are evolving from isolated assistants into heterogeneous systems in which multiple agents, planners, controllers, and execution backends operate over the same user or enterprise...

<https://arxiv.org/abs/2609.00015>

8 CONF

SCAFFOLD: A Large-Scale Structured Dataset of Computer Science Research Figures with Diagram QA and Chain-of-Thought Reasoning Traces

ArXiv CS.AI 50%

arXiv:2609.00018v1 Announce Type: new Abstract: Computer science papers rely heavily on diagrams: architecture drawings, system flowcharts, and pipeline schematics that often carry more information than the text around them. There is currently no public dataset that pairs this...

<https://arxiv.org/abs/2609.00018>

9 CONF

MiNER: Fine-Tuned Biomedical Natural Language Processing for Malaria Disease Entity Recognition in Clinical Texts

ArXiv CS.AI 50%

arXiv:2609.00073v1 Announce Type: new Abstract: Malaria remains a significant global health burden, necessitating continuous research efforts to understand its complex molecular mechanisms, epidemiology, and potential therapeutic interventions. Extracting essential biomedical...

<https://arxiv.org/abs/2609.00073>

0 CONF

AI Morbidity and Mortality: A Framework for Clinical AI Failure Review

ArXiv CS.AI 50%

arXiv:2609.00076v1 Announce Type: new Abstract: Clinical artificial intelligence is increasingly embedded in real-world care, yet existing safety mechanisms are poorly suited to reconstructing and learning from individual AI-related errors and near-misses. Aggregate model...

<https://arxiv.org/abs/2609.00076>

1 CONF

Recursive Criticality of AI Self-Improvement

ArXiv CS.AI 50%

arXiv:2609.00137v1 Announce Type: new Abstract: AI is increasingly used in the R\&D process that produces future AI systems. We study the conditions under which this feedback becomes self-amplifying. Our model describes how the rate of AI capability growth depends on baseline...

<https://arxiv.org/abs/2609.00137>

2 CONF

IMPACT: Attention Is the Interaction Map for Scalable Interaction-Aware World Model Training

ArXiv CS.AI 50%

arXiv:2609.00161v1 Announce Type: new Abstract: World models have made remarkable progress in action-conditioned future prediction for embodied agents, yet still struggle to model physically plausible interactions. Existing approaches address this limitation by constraining the...

<https://arxiv.org/abs/2609.00161>

3 CONF

LLM-Driven Autonomous Vehicles Inherit Human Driver Biases in Pedestrian Yielding: Results and Implications From A New Benchmark

ArXiv CS.AI 50%

arXiv:2609.00192v1 Announce Type: new Abstract: Public trust in Autonomous Vehicles (AVs) may depend not only on technical success but also on the fairness of their decision making. While a recent trend in AV research involves using general purpose "common sense" models to guide...

<https://arxiv.org/abs/2609.00192>

- 4 **CONF**
- AI Should Not Only Be Helpful. It Should Be Contingent. Artificial Intimacy, Sycophancy, and the Future of Social Learning**
- ArXiv CS.AI 50%
- arXiv:2609.00211v1 Announce Type: new Abstract: Conversational artificial intelligence is increasingly embedded in everyday social environments, where it functions as both an informational tool and a source of interpersonal feedback. This perspective introduces contingency,...
- <https://arxiv.org/abs/2609.00211>
-
- 5 **CONF**
- Learning What to Retain: Gated-Memory Routing for Efficient Collaboration in Multi-Agent LLM Systems**
- ArXiv CS.AI 50%
- arXiv:2609.00237v1 Announce Type: new Abstract: Large language model (LLM)-based multi-agent systems tackle complex reasoning by orchestrating how multiple agents are configured and how they collaborate. A central challenge is to adapt orchestration to the evolving collaboration...
- <https://arxiv.org/abs/2609.00237>
-
- 6 **CONF**
- Autoresearch for Marketplace Catalogs: From Legacy Forms to AI-Native Matching**
- ArXiv CS.AI 50%
- arXiv:2609.00274v1 Announce Type: new Abstract: Two-sided service marketplaces are moving from deterministic request-form intake to AI-native probabilistic matching, enabled by large language models (LLMs) that infer intent, preferences, and latent constraints from natural...
- <https://arxiv.org/abs/2609.00274>
-
- 7 **CONF**
- Human-AI Co-Interpretation for Responsible AI: A Hermeneutic Perspective**
- ArXiv CS.AI 50%
- arXiv:2609.00334v1 Announce Type: new Abstract: Across law, education, policy analysis, and public moral argumentation, LLM outputs are being used often for work that requires interpretations to be justified with textual evidence and explicit normative standards. Yet a recurrent...
- <https://arxiv.org/abs/2609.00334>
-
- 8 **CONF**
- Good Memory Has ECC: Evaluating the Memory of Vision-Language Models Beyond Accuracy**
- ArXiv CS.AI 50%
- arXiv:2609.00103v1 Announce Type: cross Abstract: Memory is widely viewed as an important unsolved problem for LLMs and VLMs, and current benchmarks typically evaluate it by testing accuracy over long text or video. However, accuracy alone misses properties that matter for real...
- <https://arxiv.org/abs/2609.00103>
-
- 9 **CONF**
- Dr. Claw: An AI Scientist Workspace for Vibe Research**
- ArXiv CS.AI 50%
- arXiv:2609.00365v1 Announce Type: new Abstract: Command-line coding agents (e.g., Claude Code, Gemini CLI) can already read and write files and sustain long sessions, yet end-to-end research still fragments across chat tools, IDEs, terminals, and writing environments, and the...
- <https://arxiv.org/abs/2609.00365>
-

0 **CONF****Learning to Remember: End-to-End Training of Memory Agents for Long-Context Reasoning****ArXiv CS.AI 50%**

arXiv:2602.18493v2 Announce Type: replace-cross Abstract: Long-context LLMs and Retrieval-Augmented Generation defer state tracking and evidence consolidation to query time, which is brittle when facts evolve and answers depend on latent states. We introduce Unified Memory Agent...

<https://arxiv.org/abs/2602.18493>

1 **CONF****Conversation Coach: A Voice-enabled AI System that Helps Practice Difficult Workplace Conversations****ArXiv CS.AI 50%**

arXiv:2609.00441v1 Announce Type: new Abstract: Effective manager-employee communication is critical for retaining high performers and developing underperformers, yet training managers in these skills remains costly. Text-based chatbots offer a scalable approach but cannot...

<https://arxiv.org/abs/2609.00441>

2 **CONF****UniACE: A Unified Framework for Evaluating LLM Agentic Capabilities****ArXiv CS.AI 50%**

arXiv:2605.27898v3 Announce Type: replace Abstract: Agent benchmarks are increasingly used to compare large language models (LLMs) across domains, yet a reported score reflects a complete model--harness--environment configuration rather than the model alone. Benchmark packages...

<https://arxiv.org/abs/2605.27898>

3 **CONF****SinkPruner: Sink-Free Visual Token Pruning for Multimodal Large Language Models****ArXiv CS.AI 50%**

arXiv:2609.01004v1 Announce Type: cross Abstract: Despite their strong multimodal understanding ability, multimodal large language models (MLLMs) incur substantial computational overhead when processing long visual token sequences. To reduce inference costs, recent studies have...

<https://arxiv.org/abs/2609.01004>

4 **CONF****The Privacy-Hallucination Tradeoff in Differentially Private Language Models****ArXiv CS.AI 50%**

arXiv:2609.00492v1 Announce Type: new Abstract: Both privacy and factual accuracy are paramount in high-stakes domains like healthcare. Concerningly, we uncover and investigate a privacy-hallucination tradeoff in differentially private (DP) language models. First, we empirically...

<https://arxiv.org/abs/2609.00492>

5 **CONF****CoVer: Conflict-Aware Claim Verification****ArXiv CS.AI 50%**

arXiv:2609.00508v1 Announce Type: new Abstract: Social media fact-checking has long been challenged by evidence-level and aggregation-level conflicts, where erroneous evidence mimics authoritative news sources. To capture this challenge and support conflict verification tasks,...

<https://arxiv.org/abs/2609.00508>

6 **CONF****ISO-RAG: Isoperimetric Noise Control for Retrieval-Augmented Generation****ArXiv CS.AI 50%**

arXiv:2609.00513v1 Announce Type: new Abstract: Retrieval-Augmented Generation (RAG) mitigates large language models (LLMs) hallucinations, yet conventional dense retrieval struggles with the complex reasoning paths of multi-hop question answering (QA). Graph-based RAG captures...

<https://arxiv.org/abs/2609.00513>

7 **CONF****Residual Sparsification via Output Importance for Compressing Mixture-of-Experts LLMs****ArXiv CS.AI 50%**

arXiv:2609.00575v1 Announce Type: new Abstract: Mixture-of-experts (MoE) architectures scale large language models efficiently, but they demand massive GPU memory. To cope with such demand, models are commonly compressed to reduce their memory footprint. Residual sparsification...

<https://arxiv.org/abs/2609.00575>

8 **CONF****Consistency Without Alignment: Item-Sensitive Language Models Indistinguishable From Random****ArXiv CS.AI 50%**

arXiv:2609.00576v1 Announce Type: new Abstract: Item-sensitivity, defined as whether a model's choice depends on the specific input rather than on its own output prior, is widely reported as evidence of task competence. We show this evidence is necessary but not sufficient using...

<https://arxiv.org/abs/2609.00576>

9 **CONF****Socrates went Nuclear: Comparing Interaction Strategies for AI systems in a Learning Context using Brain Sensing****ArXiv CS.AI 50%**

arXiv:2609.00584v1 Announce Type: new Abstract: Does unrestricted AI access bypass the cognitive effort required for learning, or does it streamline knowledge acquisition? This paper reports on a study where we compare three designs for user-AI interaction in a learning context:...

<https://arxiv.org/abs/2609.00584>

0 **CONF****Control-Data Flow Separation: Stable Prompt Optimization in Multi-Agent LLMs****ArXiv CS.AI 50%**

arXiv:2609.00621v1 Announce Type: new Abstract: Prompt optimization can improve multi-agent LLM systems, but the prompts being optimized often serve two entangled roles: generating task-relevant content and specifying execution-critical protocols, such as message routing, output...

<https://arxiv.org/abs/2609.00621>

1 **CONF****WorldBench: Culturally Grounded Benchmark for Multilingual Agents****ArXiv CS.AI 50%**

arXiv:2609.01056v1 Announce Type: new Abstract: Despite the growing use of LLM-powered agents to solve multi-step tasks in complex environments, existing benchmarks rarely test state preservation, performance across languages, and application to realistic, grounded scenarios. To...

<https://arxiv.org/abs/2609.01056>

2 **CONF****Self-Reports Are Not Verification: Environment-Grounded Auditing of LLM Operators in Evolutionary Search****ArXiv CS.AI** **50%**

arXiv:2609.00652v1 Announce Type: new Abstract: Language model agents increasingly propose actions, observe external feedback, and explain their own behavior. Their confidence and rationales are convenient monitoring signals, but convenience is not verification. We introduce an...

<https://arxiv.org/abs/2609.00652>

3 **CONF****SciTrue: Reliable Scientific Claim Validation with Frontier and Open Language Models at the NTCIR SciClaimEval Task****ArXiv CS.AI** **50%**

arXiv:2609.00654v1 Announce Type: new Abstract: We describe the SciTrue team's participation in both subtasks of the NTCIR-19 SciClaimEval task~\cite{sciclaimeval}, which asks systems to verify scientific claims against the tables and figures of a paper. Rather than tuning a...

<https://arxiv.org/abs/2609.00654>

4 **CONF****QTEA: Ternary LLMs with Sparse Residual Salient Weight and By-Column Optimization****ArXiv CS.AI** **50%**

arXiv:2609.00224v1 Announce Type: cross Abstract: Weight-only post-training quantization (PTQ) can alleviate the computational burden of serving large language models (LLMs) at scale. However, existing PTQ methods often fail to generalize across models and suffer severe accuracy...

<https://arxiv.org/abs/2609.00224>

5 **CONF****Triple-Bottom-Line Sustainability of Language Models for Edge AI: A Comparison Between SLMs and Quantized LLMs****ArXiv CS.AI** **50%**

arXiv:2609.00665v1 Announce Type: new Abstract: Edge-AI model selection is commonly driven by one isolated metric - accuracy, latency, memory, energy, or safety, even though a deployable language model must balance all five. Our work focuses on answering the question whether na...

<https://arxiv.org/abs/2609.00665>

6 **CONF****Prompt-Robust Language Models: Which Training Strategies Work?****ArXiv CS.AI** **50%**

arXiv:2609.01217v1 Announce Type: new Abstract: Despite their strong performance, large language models remain highly sensitive to prompt formulation. Prior work addresses this through refined data construction or through dedicated robustness objectives. We reproduce and compare...

<https://arxiv.org/abs/2609.01217>

7 **CONF****HBQ: Hierarchical Scaling Block Quantization with Hardware-Efficiency-Aware Design for Accurate LLM Inference****ArXiv CS.AI** **50%**

arXiv:2609.00450v1 Announce Type: cross Abstract: Block Quantization (BQ) is a promising approach for efficient deployment of large language models (LLMs), enabling low-precision computation with controlled accuracy degradation. Compared to scalar weight-only quantization (WoQ),...

<https://arxiv.org/abs/2609.00450>

8 **CONF****FLaG: Frequency-Domain Latent-attention Gated Pooling for Token Aggregation****ArXiv CS.AI** **50%**

arXiv:2609.00831v1 Announce Type: new Abstract: Token aggregation converts token-level representations into fixed-dimensional sample representations, but most pooling methods operate only in the original token space. We introduce Frequency-Domain Latent-attention Gated Pooling...

<https://arxiv.org/abs/2609.00831>

9 **CONF****Reinforcement Learning Enhanced LLM Agents for Complex Vehicle Routing Problems****ArXiv CS.AI** **50%**

arXiv:2609.00859v1 Announce Type: new Abstract: Vehicle Routing Problems (VRPs) are fundamental combinatorial optimization problems with widespread applications in various scenarios. The advanced optimization solvers can effectively solve such problems. However, modeling complex...

<https://arxiv.org/abs/2609.00859>

0 **CONF****Towards reliable multimodal disaster severity assessment through preference optimization and explainable vision-language reasoning****ArXiv CS.AI** **50%**

arXiv:2609.00879v1 Announce Type: new Abstract: Reliable disaster damage assessment requires models that provide both accurate predictions and transparent explanations. However, existing multimodal approaches are limited by scarce annotated data and insufficient evaluation of...

<https://arxiv.org/abs/2609.00879>

1 **CONF****Denoising Diffusion Generative Models Secretly Calculate Attentions****ArXiv CS.AI** **50%**

arXiv:2609.00885v1 Announce Type: new Abstract: Denoising diffusion models are the dominant architecture for image generation, whereas most natural language generation and modeling are primarily handled by well-known transformer architectures employing attention mechanism. Here,...

<https://arxiv.org/abs/2609.00885>

2 **CONF****In-Context Neurofeedback: Can LLMs Control Their Internal Representations through Privileged Access?****ArXiv CS.AI** **50%**

arXiv:2609.00904v1 Announce Type: new Abstract: Whether large language models (LLMs) can control their own internal representations matters for both machine metacognition and AI safety. A recent study applied neurofeedback to LLMs and claimed that they can control their internal...

<https://arxiv.org/abs/2609.00904>

3 **CONF****Few-Shot Out of Domain Intent Detection with Covariance Corrected Mahalanobis Distance****ArXiv CS.AI** **50%**

arXiv:2609.00961v1 Announce Type: new Abstract: Conversational agents like chatbots and voice assistants are trained to understand and respond to user intents. On encountering an utterance with an intent different from the ones they have been trained on, these agents are...

<https://arxiv.org/abs/2609.00961>

4 **CONF****Spawn Freely, Act Sparingly: Progressive Risk Vesting for Recursive LLM-Agent Trees****ArXiv CS.AI 50%**

arXiv:2609.01035v1 Announce Type: new Abstract: Recursive LLM agents can broaden their search by spawning specialists. Some branches later request tools that send data or deploy code. When should a branch receive authority to act? We distinguish sandbox spawning, in which...

<https://arxiv.org/abs/2609.01035>

5 **CONF****User Representation via Cross Multi-source Behavior Pre-training for Mobile Games****ArXiv CS.AI 50%**

arXiv:2609.01057v1 Announce Type: new Abstract: User representation pre-training has become a fundamental paradigm for alleviating data sparsity in downstream personalization tasks. However, existing studies predominantly focus on single-app or app-level behaviors, overlooking...

<https://arxiv.org/abs/2609.01057>

6 **CONF****Space Generative AI with Solar Energy Harvesting****ArXiv CS.AI 50%**

arXiv:2609.01062v1 Announce Type: new Abstract: Satellites are emerging as promising platforms to extend generative \emph{artificial intelligence} (AI) services to remote areas lacking terrestrial infrastructure. However, deploying space generative AI is fundamentally...

<https://arxiv.org/abs/2609.01062>

7 **CONF****Latent Recurrent Thoughts: Recurrent Refinement of Proposed Latents for Reasoning with Frozen LLMs****ArXiv CS.AI 50%**

arXiv:2609.01117v1 Announce Type: new Abstract: Chain-of-thought reasoning unfolds in discrete token space: each step is committed as text, errors propagate, and eliciting good traces presupposes traces to imitate. Reasoning instead in a model's continuous representation space -...

<https://arxiv.org/abs/2609.01117>

8 **CONF****Jailbreaking Text-to-Image Models Through Cracks: Navigating Heterogeneous Safety Filters via Multi-Agent Debate****ArXiv CS.AI 50%**

arXiv:2609.01168v1 Announce Type: new Abstract: Text-to-image (T2I) models remain vulnerable to jailbreak attacks that elicit Not-Safe-For-Work (NSFW) content, despite increasingly being guarded by heterogeneous, multi-layer safety stacks combining text filters, image...

<https://arxiv.org/abs/2609.01168>

9 **CONF****LEAP: Likelihood Elicitation and Aggregation for LLM-based Probabilistic Forecasting****ArXiv CS.AI 50%**

arXiv:2609.01337v1 Announce Type: new Abstract: LLM-based forecasting systems have improved on real-world tasks such as financial markets and sports outcomes, largely through stronger search and tool use. Many systems still ask an LLM to read all collected evidence together and...

<https://arxiv.org/abs/2609.01337>

0 CONF

When Guardrails Look Effective: Construct Validity Failures in LLM Agent Commerce Evaluation

ArXiv CS.AI 50%

arXiv:2609.01519v1 Announce Type: new Abstract: Interactive simulations increasingly evaluate policies in markets populated by language-model agents. Their outputs can look economic---prices, profits, consumer surplus, and welfare---without instantiating the behavior named in...

<https://arxiv.org/abs/2609.01519>

1 CONF

Can LLMs Discover Scientific Laws in Real and Parallel Worlds?

ArXiv CS.AI 50%

arXiv:2609.01552v1 Announce Type: new Abstract: Scientific equation discovery has long been central to scientific progress, proceeding through iterative cycles of hypothesis generation, observational testing, and refinement under scientific constraints. As LLM capabilities...

<https://arxiv.org/abs/2609.01552>

2 CONF

EDRAC: Benchmarking Arabic Dialect Reading Comprehension

ArXiv CS.AI 50%

arXiv:2609.01113v1 Announce Type: cross Abstract: Dialectal Arabic (DA) remains under-resourced compared to Modern Standard Arabic (MSA), particularly for machine reading comprehension (MRC) and question answering (QA). Existing Arabic QA benchmarks primarily focus on formal...

<https://arxiv.org/abs/2609.01113>

3 CONF

Behaviorally Grounded User Profiles from the Wild for Personalized Alignment and Multi-Perspective Reasoning

ArXiv CS.AI 50%

arXiv:2609.00014v1 Announce Type: cross Abstract: Persona-driven techniques increasingly adapt large language models (LLMs) to diverse contexts. However, existing methods predominantly rely on rigid, synthetic personas that flatten individual variation, rely on stereotypes, and...

<https://arxiv.org/abs/2609.00014>

4 CONF

trajectory-judge: What Outcome-Only LLM Judges Miss on Agent Trajectories

ArXiv CS.AI 50%

arXiv:2609.00038v1 Announce Type: cross Abstract: Outcome-only evaluation is the production default for LLM agents: show a judge the request and the final reply and ask whether it was handled well. The metric is structurally blind to an agent that reaches the right answer the...

<https://arxiv.org/abs/2609.00038>

5 CONF

Task-Specific Prompt with Global Context for Multi-Task Graph Pre-Training

ArXiv CS.AI 50%

arXiv:2609.00047v1 Announce Type: cross Abstract: Graph prompt learning is an effective paradigm to adapt pre-trained graph models to downstream tasks in low-resource scenarios. However, existing multi-task graph pre-training frameworks generally use randomly initialized...

<https://arxiv.org/abs/2609.00047>

6 CONF

GUI-CC: Benchmarking Contextual Consistency of GUI World Models as Agent Environments

ArXiv CS.AI 50%

arXiv:2609.00048v1 Announce Type: cross Abstract: GUI world models are increasingly evaluated as one-step next-screen predictors, yet their intended use is often as multi-step environments for GUI agents. This mismatch leaves a key requirement under-tested: generated states must...

<https://arxiv.org/abs/2609.00048>

7 CONF

REAL-Q: E2E LLM Quantization via Dynamic Gradient Descent

ArXiv CS.AI 50%

arXiv:2609.00049v1 Announce Type: cross Abstract: Post-training quantization (PTQ) is essential for deploying large language models (LLMs) under strict resource constraints. State-of-the-art PTQ methods quantize each layer with a single closed-form second-order solver: to remain...

<https://arxiv.org/abs/2609.00049>

8 CONF

Detecting Hidden Behaviors in LLMs via Activation-matched Finetuning

ArXiv CS.AI 50%

arXiv:2609.00351v1 Announce Type: cross Abstract: Large language models can hide hidden behaviors that activate only under narrow conditions, such as backdoor triggers, sleeper-agent deployment cues, sandbagging, or topic-conditioned censorship. Such behaviors are difficult to...

<https://arxiv.org/abs/2609.00351>

9 CONF

Zero-Shot Respiratory Sound Classification through LLM-Augmented Audio-Text Alignment

ArXiv CS.AI 50%

arXiv:2609.00055v1 Announce Type: cross Abstract: Self-supervised respiratory encoders lack semantic grounding in clinical domain needed for zero-shot inference, limiting their utility without task-specific labeled data. We propose a framework that aligns these encoders with...

<https://arxiv.org/abs/2609.00055>

10 CONF

EEG-VID: Task-Guided Latent Predictive Pretraining for EEG Decoding and Assistive Target Selection

ArXiv CS.AI 50%

arXiv:2609.00566v1 Announce Type: cross Abstract: We propose EEG-VID, a task-guided latent predictive pretraining framework for EEG decoding under session and subject shifts. EEG-VID predicts future latent EEG states from recent history using an exponential-moving-average target...

<https://arxiv.org/abs/2609.00566>

11 CONF

DISTAL: Distillation and Self-Supervised Pretraining for Structure-Agnostic Materials Property Prediction

ArXiv CS.AI 50%

arXiv:2609.00059v1 Announce Type: cross Abstract: Materials property prediction remains difficult in low-data settings, where many target properties are supported by only a limited number of labeled samples. Models with the strongest predictive accuracy often depend on crystal...

<https://arxiv.org/abs/2609.00059>

- 2
CONF
ReNFT: Repairing Mode Collapse in Reward Post-Training via Internal Probability-Mass Recalibration
ArXiv CS.AI 50%
arXiv:2609.00061v1 Announce Type: cross Abstract: Reward post-training of diffusion generators inevitably concentrates probability mass on a few reward-favored modes, a mode collapse that erases within-prompt diversity. Existing methods for mitigating collapse rely on external...
<https://arxiv.org/abs/2609.00061>
- 3
CONF
RePro: Proof-Verified Benchmark Rewriting for Reliable Evaluation of LLM Mathematical Problem Solving
ArXiv CS.AI 50%
arXiv:2609.00062v1 Announce Type: cross Abstract: Data contamination undermines the reliable evaluation of large language models (LLMs) on mathematical problem solving. While rewriting-based evaluation mitigates memorization, existing methods lack guarantees of problem validity...
<https://arxiv.org/abs/2609.00062>
- 4
CONF
Attention Sensitivity Is Not Enough: Dissociating Attention-Level and Behavioural In-Context Learning under Fine-Tuning
ArXiv CS.AI 50%
arXiv:2609.00064v1 Announce Type: cross Abstract: In-context learning (ICL) lets large language models adapt to new tasks from demonstrations, and fine-tuning can erode this behaviour. Many preservation diagnostics inspect attention: if attention changes when demonstrations...
<https://arxiv.org/abs/2609.00064>
- 5
CONF
Do Multimodal LLMs See Before They Read? Diagnosing Contextual Sycophancy
ArXiv CS.AI 50%
arXiv:2609.00067v1 Announce Type: cross Abstract: External text can override conflicting image evidence in multimodal large language models, a failure we call multimodal contextual sycophancy. We introduce a 998-case diagnostic that independently varies visual evidence,...
<https://arxiv.org/abs/2609.00067>
- 6
CONF
RW-LoRA: Communication-Efficient Decentralized LoRA Fine-Tuning via Random Walks
ArXiv CS.AI 50%
arXiv:2609.00078v1 Announce Type: cross Abstract: Parameter-efficient fine-tuning methods such as LoRA have become a standard approach for adapting large foundation models. Adopting fine-tuning to distributed settings faces several challenges. Most existing distributed LoRA...
<https://arxiv.org/abs/2609.00078>
- 7
CONF
Retrieval, Scoring, and Decoding Shape Performance and Stability in LLM-based Conversational Recommendation
ArXiv CS.AI 50%
arXiv:2609.00086v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly used as rerankers in conversational recommender systems, yet measured gains depend strongly on the retrieval and inference protocol. On the ReDial conversational movie recommendation...
<https://arxiv.org/abs/2609.00086>

8 **CONF****Assessing Alignment and Stability of Feature Importance Explanations via Weight of Evidence****ArXiv CS.AI 50%**

arXiv:2609.00090v1 Announce Type: cross Abstract: Feature importance Methods (FIMs) are widely used in Explainable AI to interpret model predictions, yet attribution scores alone often provide limited insight into the underlying reasoning process. In this work, we introduce a...

<https://arxiv.org/abs/2609.00090>

9 **CONF****Do General NLP Embeddings Capture Ontological Reasoning?****ArXiv CS.AI 50%**

arXiv:2609.00177v1 Announce Type: cross Abstract: General-purpose NLP embedding models perform well on linguistic tasks, but their ability to capture symbolic ontological structure remains unclear. We introduce AVA, a systematic framework for evaluating whether embeddings...

<https://arxiv.org/abs/2609.00177>

0 **CONF****Assessing Suicide Risk in Arabic Crisis Helpline Calls: A Comparison of Arabic and English Large Language Models****ArXiv CS.AI 50%**

arXiv:2609.00191v1 Announce Type: cross Abstract: Crisis helplines assess suicide risk through structured interviews, a process that is slow and dependent on operator training and workload. Natural language processing could support risk assessment and call prioritization, but...

<https://arxiv.org/abs/2609.00191>

1 **CONF****Provably Efficient Federated Reinforcement Learning with Linear Function Approximation and Logarithmic Communication Cost****ArXiv CS.AI 50%**

arXiv:2609.00193v1 Announce Type: cross Abstract: We study federated online reinforcement learning with linear function approximation. While recent multi-agent reinforcement learning algorithms achieve strong regret guarantees, they typically require sharing raw trajectories....

<https://arxiv.org/abs/2609.00193>

2 **CONF****Distributed Implicit Harm: A Compositional Safety Blind Spot in MLLM-Based Video Moderation****ArXiv CS.AI 50%**

arXiv:2609.00206v1 Announce Type: cross Abstract: Despite their growing use in video moderation, multimodal large language models (MLLMs) exhibit a compositional safety blind spot: videos composed of seemingly benign components can convey harmful meaning when interpreted as a...

<https://arxiv.org/abs/2609.00206>

3 **CONF****Don't Let the Model Write the YAML: Deterministic, Minimal-Diff GitOps Remediation from LLM-Proposed Field Changes****ArXiv CS.AI 50%**

arXiv:2609.00227v1 Announce Type: cross Abstract: LLM agents increasingly diagnose incidents and propose remediations. In a GitOps workflow, applying a fix means editing a version-controlled config file, and the obvious implementation, having the model author the edited file or...

<https://arxiv.org/abs/2609.00227>

- 4
CONF
CoLT-Drive: Counterfactual Long-Tail Benchmarking and Knowledge-Preserving Adaptation for Driving Affordance Prediction
ArXiv CS.AI 50%
arXiv:2609.00242v1 Announce Type: cross Abstract: Long-tail autonomous driving failures are often framed as rare-object recognition errors. We argue that this view is incomplete: the decision-critical question is not only whether a model recognizes an unusual object, but whether...
<https://arxiv.org/abs/2609.00242>
- 5
CONF
CompanionSim: Synthetic Data for Evaluating Anthropomorphism in Human-AI Relationships
ArXiv CS.AI 50%
arXiv:2609.00250v1 Announce Type: cross Abstract: Many people now see AI systems as not just productivity tools but as social companions. Researchers are eager to study the consequences of AI companionship behaviors, such as validation, which evoke trust, empathy, and attachment...
<https://arxiv.org/abs/2609.00250>
- 6
CONF
Delegation Without Trust: An Empirical Gap Analysis of Identity, Authorization, and Runtime Governance in Multi-Agent LLM Systems
ArXiv CS.AI 50%
arXiv:2609.00267v1 Announce Type: cross Abstract: Autonomous LLM agents increasingly act on a user's behalf: they hold credentials, call tools and services, and spawn sub-agents that act further on their behalf. This turns a long-standing distributed-systems question -- who is...
<https://arxiv.org/abs/2609.00267>
- 7
CONF
Cleaner Speech, Weaker Generalization: Revisiting Pitt-Derived Benchmarks for Alzheimer's Disease Detection
ArXiv CS.AI 50%
arXiv:2609.00276v1 Announce Type: cross Abstract: Speech-based Alzheimer's disease (AD) detection increasingly relies on speech-enhanced and curated versions of the Pitt Corpus, where speech enhancement, sample selection, and demographic balancing are often treated as beneficial...
<https://arxiv.org/abs/2609.00276>
- 8
CONF
Geometry-aware Latent Autoregressive Generative Model for PDEs in Complex Domains
ArXiv CS.AI 50%
arXiv:2609.00297v1 Announce Type: cross Abstract: Solving multiphysics partial differential equations (PDEs) remains a major challenge in scientific computing, especially for highly complex μ -scale tortuous geometries critical to energy and chemical engineering. We address...
<https://arxiv.org/abs/2609.00297>
- 9
CONF
A Human-AI Theorem Connecting Spontaneous and Field-Induced Mechanisms of Collective Behavior in One Dimension
ArXiv CS.AI 50%
arXiv:2609.00322v1 Announce Type: cross Abstract: Can an artificial intelligence (AI) generate a scientific hypothesis outside a human collaborator's active hypothesis space (AHS), and can human-AI research be organized to make such breakthroughs more likely? We document such a...
<https://arxiv.org/abs/2609.00322>

00 CONF

Topic Matching in the Wild: Benchmark and Lessons from Real-World ASR Transcripts

ArXiv CS.AI 50%

arXiv:2609.00330v1 Announce Type: cross Abstract: In contact centers, real-time agent-assist tools determine, for each of many predefined topics, whether a live customer utterance is relevant and display a coaching card to the agent when it is. The input is noisy and...

<https://arxiv.org/abs/2609.00330>

01 CONF

Counterfactual Fragility Certificates: Exposing High-Confidence Brittleness under Structured Evidence Failure

ArXiv CS.AI 50%

arXiv:2609.00366v1 Announce Type: cross Abstract: High test accuracy and good aggregate calibration do not show whether an individual prediction is structurally supported by its evidence. In tabular decision systems, failures often occur when a feature family becomes...

<https://arxiv.org/abs/2609.00366>

02 CONF

CordisBench: Can Language Models Reason About Component Lifecycles in Dynamic Agent Harnesses?

ArXiv CS.AI 50%

arXiv:2609.01600v1 Announce Type: cross Abstract: Dynamic agent harnesses let language models change the software that shapes their own execution. This flexibility brings a new reasoning burden: a local plugin change can propagate through dependencies and cleanup. We introduce...

<https://arxiv.org/abs/2609.01600>

03 CONF

Higher Structures in Deep Learning

ArXiv CS.AI 50%

arXiv:2609.00472v1 Announce Type: cross Abstract: We provide an expository introduction on the importance of higher-arity tensor operations to deep learning. Then, we conduct a novel empirical investigation of higher-arity phenomenon in trained neural networks, introduce a...

<https://arxiv.org/abs/2609.00472>

04 CONF

EvoFlint: An Evolutionary Atlas of Multi-Turn LLM Vulnerabilities

ArXiv CS.AI 50%

arXiv:2609.00487v1 Announce Type: cross Abstract: Frontier language models that refuse harmful single-turn prompts often comply when the same intent is reached gradually over many turns, making multi-turn attacks one of the least understood failure modes of large language...

<https://arxiv.org/abs/2609.00487>

05 CONF

RecalibrateGPT: AI Fatigue Resilient Conversational Interfaces

ArXiv CS.AI 50%

arXiv:2609.00506v1 Announce Type: cross Abstract: Large language models are powerful, but their interfaces often devolve into a type $\rightarrow$ read $\rightarrow$ retype loop, creating conversational AI fatigue, cognitive load, and eventual task abandonment. To mitigate this,...

<https://arxiv.org/abs/2609.00506>

06 CONF

The Interlingua Hypothesis: LLMs Translate via a Latent Task-agnostic Feature Space

ArXiv CS.AI 50%

arXiv:2609.00515v1 Announce Type: cross Abstract: Large language models (LLMs) have recently demonstrated improved machine translation performance over strong supervised baselines. This raises questions as to what mechanisms underlie how LLMs perform machine translation between...

<https://arxiv.org/abs/2609.00515>

07 CONF

The Safeguard Worked. Is the LLM System Safer?

ArXiv CS.AI 50%

arXiv:2609.00519v1 Announce Type: cross Abstract: Safeguards in deployed LLM services are evaluated by refusal, attack success, and policy violation rates. Those rates characterize how a control performed on the requests it was tested on. A deployment has to answer a different...

<https://arxiv.org/abs/2609.00519>

08 CONF

EmbodiedSkills: A Unified Framework for Orchestrating, Training, and Deploying VLA Agents

ArXiv CS.AI 50%

arXiv:2609.01281v1 Announce Type: cross Abstract: Vision-language-action (VLA) models map visual observations and language instructions directly to robot actions, but long-horizon tasks require more than action prediction. An agent must coordinate perception, planning,...

<https://arxiv.org/abs/2609.01281>

09 CONF

Predicting Program Exit Code with LLMs and Programming Language Semantics

ArXiv CS.AI 50%

arXiv:2609.00579v1 Announce Type: cross Abstract: Large language models (LLMs) have shown proficiency in various software engineering tasks, such as code generation and translation. However, a key limitation in their performance may be their (lack of) understanding of...

<https://arxiv.org/abs/2609.00579>

10 CONF

Confess What You Know: Forget-Set Misalignment with Model Knowledge in LLM Unlearning

ArXiv CS.AI 50%

arXiv:2609.00605v1 Announce Type: cross Abstract: Machine unlearning for large language models (LLMs) often assumes that a pre-defined forget set matches what the model has memorized, but this frequently breaks in realistic privacy settings where the original training data is...

<https://arxiv.org/abs/2609.00605>

11 CONF

Restrict, Don't Retrain: Inference-Time VLM Guidance for Zero-Shot Aerial Segmentation

ArXiv CS.AI 50%

arXiv:2609.00628v1 Announce Type: cross Abstract: Global welfare often depends on the correct interpretation of aerial and satellite imagery. Acting on such imagery (mapping flooded ground, crop extent, or damaged infrastructure) demands pixel-level segmentation to ensure...

<https://arxiv.org/abs/2609.00628>

12 CONF

Probing Factual Knowledge Transfer with Training Data Interventions

ArXiv CS.AI 50%

arXiv:2609.01341v1 Announce Type: cross Abstract: Do multilingual language models transfer factual knowledge across languages during continued pretraining, or do they mostly recall facts learned directly from the target-language data? To answer this question more reliably, we...

<https://arxiv.org/abs/2609.01341>

13 CONF

EEG-AS: Instance-Level Foundation Model Selection for EEG Foundation Models via Behavior Reconstruction

ArXiv CS.AI 50%

arXiv:2609.00653v1 Announce Type: cross Abstract: Electroencephalography (EEG) is a non-invasive technique for measuring neural activity and has been widely used in neuroscience applications. Recent advances in EEG foundation models have enabled strong performance across diverse...

<https://arxiv.org/abs/2609.00653>

14 CONF

VOIM: Training-Free Open-Vocabulary 3D Instance Mapping for RGB-D and Monocular SLAM

ArXiv CS.AI 50%

arXiv:2609.00775v1 Announce Type: cross Abstract: We present Voxel-Grounded Online Instance Manager (VOIM), a training-free voxel-grounded instance manager that builds open-vocabulary 3D instance maps from RGB-D or from monocular RGB alone, a regime no prior training-free system...

<https://arxiv.org/abs/2609.00775>

15 CONF

A Checklist to assess the energy and carbon impacts of ML/AI applications in Earth System Modeling

ArXiv CS.AI 50%

arXiv:2609.00847v1 Announce Type: cross Abstract: As machine learning and artificial intelligence find their way into nearly every aspect of climate, weather, and Earth system modeling, it is worth pausing to consider what our design decisions imply for the science and for the...

<https://arxiv.org/abs/2609.00847>

16 CONF

Does Fault Localization Beat a Fresh Attempt? A Placebo-Controlled Study of Test-Guided Code Repair

ArXiv CS.AI 50%

arXiv:2609.00854v1 Announce Type: cross Abstract: Fault localization can focus a code model's repair on the statements a failing test implicates, but a targeted edit may succeed merely because it is small, and a second model call may succeed without using the failure at all. We...

<https://arxiv.org/abs/2609.00854>

17 CONF

Vision-Language-Guided Pseudo-Labels for Unsupervised Domain Adaptation in Semantic Segmentation for Waste Sorting

ArXiv CS.AI 50%

arXiv:2609.00898v1 Announce Type: cross Abstract: Obtaining labeled data for semantic segmentation in applied settings (e.g., autonomous driving, industrial waste sorting) is expensive and often infeasible at scale. We present a cross-modal pseudo-labeling pipeline that enables...

<https://arxiv.org/abs/2609.00898>

18 CONF

Context-Grounding Gains Are Mediated by Pre-existing Machinery: Auditing GRPO, SFT, and DPO

ArXiv CS.AI 50%

arXiv:2609.00925v1 Announce Type: cross Abstract: Language models can ignore prompt evidence when it conflicts with memorized knowledge. Post-training can make models follow such evidence more reliably, but it is unclear whether these gains require new machinery or strengthen...

<https://arxiv.org/abs/2609.00925>

19 **CONF****Semi-Supervised Virtual Staining via Morphology Preservation and Histopathological Realism Constraints****ArXiv CS.AI** **50%**

arXiv:2609.00984v1 Announce Type: cross Abstract: Virtual staining aims to computationally generate target-stained histopathological images while reducing the cost and time associated with conventional staining procedures. However, existing methods rely predominantly on strictly...

<https://arxiv.org/abs/2609.00984>

20 **CONF****On the Human and Computer Alignment of Attribute-Based Music Matches****ArXiv CS.AI** **50%**

arXiv:2609.00987v1 Announce Type: cross Abstract: Recent advances in generative AI are raising ethical concerns regarding the originality of generated content and the potential replication of training data, with further implications for transparency, attribution, and...

<https://arxiv.org/abs/2609.00987>

21 **CONF****A Network Science Perspective on Evaluating Deep Graph Generative Models****ArXiv CS.AI** **50%**

arXiv:2609.01015v1 Announce Type: cross Abstract: Traditional network models from network science, such as the Erdos-Renyi and configuration models, generate random networks that reproduce few selected topological properties observed in real-world networks. Deep graph generative...

<https://arxiv.org/abs/2609.01015>

22 **CONF****Causal Evidentiary Governance for High-Risk Machine Learning Systems****ArXiv CS.AI** **50%**

arXiv:2609.01040v1 Announce Type: cross Abstract: Machine learning systems deployed for credit, hiring, and resource distribution are increasingly subject to regulatory oversight from policies such as the EU AI Act and GDPR. Current fairness governance practices rely on...

<https://arxiv.org/abs/2609.01040>

23 **CONF****Towards Provable and Scalable Training of Quantized Neural Networks with Ising Optimization****ArXiv CS.AI** **50%**

arXiv:2506.18240v5 Announce Type: replace-cross Abstract: Training quantized neural networks remains fundamentally challenging due to non-convex loss landscapes and discrete parameter spaces. We introduce an exact Quadratic Constrained Binary Optimization (QCBO) framework with...

<https://arxiv.org/abs/2506.18240>

24 **CONF****From Truncation to Commitment: Persistent Context in Uniform Discrete Diffusion****ArXiv CS.AI** **50%**

arXiv:2609.01043v1 Announce Type: cross Abstract: Uniform-state discrete diffusion models update all tokens in parallel while keeping every position revisable. Even when the commonly used top-\$p\$ rule leaves only one candidate at a position, that choice affects only the current...

<https://arxiv.org/abs/2609.01043>

25 CONF

HiveTraceGuard-Pro: A Compact Generative Guardrail for Prompt Injection, Jailbreaks, and Adversarial Obfuscation

ArXiv CS.AI 50%

arXiv:2609.01046v1 Announce Type: cross Abstract: Production LLMs must handle inputs that attempt to override system instructions, bypass safety policies or elicit harmful responses. A common mitigation is a separate guardrail model. Existing reports, however, provide little...

<https://arxiv.org/abs/2609.01046>

26 CONF

StainPresetNet: Stain Preset Network for Fast Multi-to-Multi Stain Normalization

ArXiv CS.AI 50%

arXiv:2609.01146v1 Announce Type: cross Abstract: Stain normalization reduces color variations caused by variations in staining protocols and imaging conditions, thereby enhancing computer-aided diagnostic system performance. Traditional methods derive mapping relationships from...

<https://arxiv.org/abs/2609.01146>

27 CONF

Towards AI-Assisted Clinical Trial Matching: Practical Considerations, Multicenter Evaluation, and Real-World Deployment

ArXiv CS.AI 50%

arXiv:2609.01202v1 Announce Type: cross Abstract: Clinical trials are essential for advancing cancer care and drug development, but many fail because of insufficient patient enrollment. While there is growing interest in using AI to support patient recruitment, existing systems...

<https://arxiv.org/abs/2609.01202>

28 CONF

Autonomous discovery of new structure-plausibility laws for explainable and rapid crystal diagnosis and screening

ArXiv CS.AI 50%

arXiv:2609.01209v1 Announce Type: cross Abstract: Crystal generators and tool-using agents propose structures faster than density functional theory (DFT) energy and phonon calculations or experiments can assess them. Deciding which candidates merit expensive assessment is...

<https://arxiv.org/abs/2609.01209>

29 CONF

HugAgent: A Human Simulation Benchmark for Individual-Level Reasoning

ArXiv CS.AI 50%

arXiv:2510.15144v4 Announce Type: replace Abstract: Simulating human reasoning in open-ended tasks has long been a central aspiration in AI and cognitive science. While large language models now approximate human responses at scale, they remain tuned to population-level...

<https://arxiv.org/abs/2510.15144>

30 CONF

Position Matters: Feature Inversion Attacks in ViT Split Inference with Token Reduction and Shuffling

ArXiv CS.AI 50%

arXiv:2609.01232v1 Announce Type: cross Abstract: Vision Transformers (ViTs) are increasingly used in split-inference systems, where edge devices transmit intermediate token representations to a remote cloud. In this setting, token reduction lowers computation and communication...

<https://arxiv.org/abs/2609.01232>

31 CONF

From Language to Behavior: Scaling Sequence Transformers for Industrial Recommendation Ranking with Rec-Native Designs

ArXiv CS.AI 50%

arXiv:2609.01240v1 Announce Type: cross Abstract: Scaling Transformers has driven large gains in language modeling, but transplanting this to behavior-sequence modeling in production ranking is challenging: recommendation differs in signal quality, where behavior sequences are...

<https://arxiv.org/abs/2609.01240>

32 CONF

Explore More, Drift Less: Outcome-Only Reinforcement Learning Can Suffice for Long-Horizon Interactive Agents

ArXiv CS.AI 50%

arXiv:2609.01245v1 Announce Type: cross Abstract: Reinforcement learning is a natural way to post-train LLM agents for long-horizon interactive tasks judged only by end-of-task verification, yet a shared belief holds that outcome-only RL soon hits a ceiling on small open models....

<https://arxiv.org/abs/2609.01245>

33 CONF

The Constitutional Coverage Trilemma in AI Governance

ArXiv CS.AI 50%

arXiv:2609.01275v1 Announce Type: cross Abstract: Frontier AI systems function as \emph{constitutional institutions}: each deployed model encodes an implicit ranking among safety, helpfulness, honesty, autonomy, and equity. We ask whether the supply of frontier constitutional...

<https://arxiv.org/abs/2609.01275>

34 CONF

TimeSteer: Inference-Time Speech Scheduling in Joint Audio-Visual Diffusion Models

ArXiv CS.AI 50%

arXiv:2609.01277v1 Announce Type: cross Abstract: Although pretrained joint audio-visual diffusion models offer rich control over \emph{what} to generate, they provide no explicit control over \emph{when} an utterance should occur. To address this, we study \emph{inference-time}...

<https://arxiv.org/abs/2609.01277>

35 CONF

Some Emotions Run Deeper: Layer-wise Probing and Causal Intervention in Large Language Models

ArXiv CS.AI 50%

arXiv:2609.01279v1 Announce Type: cross Abstract: Emotion is expressed in text along a wide spectrum, from surface lexical cues to inferences entangled with content. Most layer-wise analyses of emotion in LLMs use a single corpus, leaving open whether the depth at which emotion...

<https://arxiv.org/abs/2609.01279>

36 CONF

HiLRP: Toward One Trustworthy Explanation for Vision Transformer: Conservation-Valid Attribution via Attention Primitives

ArXiv CS.AI 50%

arXiv:2609.01282v1 Announce Type: cross Abstract: Vision Transformer (ViT) design has become increasingly diverse, with backbones combining convolutional stems, windowed, linear, or multi-axis attention, patch merging, and spatial reduction in various configurations. This...

<https://arxiv.org/abs/2609.01282>

37 CONF

GazeRefine: Expert Gaze as a Test-Time Prompt for Training-Free Medical Image Segmentation

ArXiv CS.AI 50%

arXiv:2609.01310v1 Announce Type: cross Abstract: Medical image segmentation remains difficult to scale because high-performing methods typically rely on dense expert annotations and task-specific training. We introduce GazeRefine, a training-free framework that uses gaze as an...

<https://arxiv.org/abs/2609.01310>

38 CONF

MIDR: Enrichment-Augmented Indexing for Multimodal Document Retrieval

ArXiv CS.AI 50%

arXiv:2609.01316v1 Announce Type: cross Abstract: Retrieval over visually rich documents has a representation problem: important content often lives in tables, charts, figures, and layout relations that plain OCR linearizes, corrupts, or omits. ColPali-family visual retrievers...

<https://arxiv.org/abs/2609.01316>

39 CONF

Bandits in Prod: Hyperparameter Optimization at Inference Time

ArXiv CS.AI 50%

arXiv:2609.01335v1 Announce Type: cross Abstract: Many production systems can assess a configuration only by using it on live requests and observing noisy feedback. Modern agentic systems are a prominent example, with inference-time choices such as model selection, retrieval...

<https://arxiv.org/abs/2609.01335>

40 CONF

CHARM: Character Hallucination for Multicultural Role Play Benchmark

ArXiv CS.AI 50%

arXiv:2609.01352v1 Announce Type: cross Abstract: Role-playing large language models (LLMs) are expected to adopt a character's style while also respecting that character's knowledge boundaries. Prior evaluations detect character hallucination but rarely distinguish whether...

<https://arxiv.org/abs/2609.01352>

41 CONF

Evaluating Multimodal LLMs as Generalist Vision-Language-Action Agents for Drone Control: Commanding, Approaching, Tracking and Searching

ArXiv CS.AI 50%

arXiv:2609.01404v1 Announce Type: cross Abstract: Multimodal Large Language Models (MLLMs) are strong perceivers of images and video. We ask how far that reach extends into acting: dropping an MLLM directly into a drone's control loop, with its entire action space declared...

<https://arxiv.org/abs/2609.01404>

42 CONF

Semantic-Guided Multimodal Preprocessing for Vision Transformer-Based Clear Cell Renal Cell Carcinoma Grading

ArXiv CS.AI 50%

arXiv:2609.01426v1 Announce Type: cross Abstract: Clear cell renal cell carcinoma (CCRCC) grading is essential for treatment planning, yet existing approaches either analyze patch-level images directly or focus solely on nuclei-level classification, without linking to final...

<https://arxiv.org/abs/2609.01426>

43 CONF

Learning Sparse Decision Trees via Transformer Variational Auto-Encoders

ArXiv CS.AI 50%

arXiv:2609.01430v1 Announce Type: cross Abstract: Decision trees are among the most widely used models in machine learning, largely due to their transparent decision logic, making them well-suited for high-stakes decision-making contexts. However, most existing learning...

<https://arxiv.org/abs/2609.01430>

44 CONF

When Safety Routing Breaks: Understanding Alignment Fragility under Benign Fine-Tuning

ArXiv CS.AI 50%

arXiv:2609.01455v1 Announce Type: cross Abstract: Benign fine-tuning severely weakens the safety alignment of large language models (LLMs), so we study why refusal behavior is so fragile. While prior work often attributes this failure to gradient conflict, we propose a...

<https://arxiv.org/abs/2609.01455>

45 CONF

TempCloze: Can Video-LLMs Identify the Missing Middle?

ArXiv CS.AI 50%

arXiv:2609.01515v1 Announce Type: cross Abstract: Temporal reasoning benchmarks for Video-LLMs are often mediated by language, leaving room for linguistic shortcuts from option wording, answer correlations, or language priors. To reduce such shortcuts, we introduce TempCloze, a...

<https://arxiv.org/abs/2609.01515>

46 CONF

Can LLMs Design Video Coding Tools? A Case Study on Planar Mode

ArXiv CS.AI 50%

arXiv:2609.01535v1 Announce Type: cross Abstract: This paper explores whether large language models (LLMs) can design video coding tools, a highly challenging task due to the intricate algorithmic coupling of tool modifications. In particular, we present an empirical case study...

<https://arxiv.org/abs/2609.01535>

47 CONF

Scaling Near-Optimal SFT-RL Annotation Budget Allocation from Small to Large LLMs

ArXiv CS.AI 50%

arXiv:2609.01573v1 Announce Type: cross Abstract: How to divide a fixed annotation budget between supervised fine-tuning (SFT) and reinforcement learning (RL) during LLM post-training remains an open problem. Existing work characterizes only broad trends (e.g., SFT dominates in...

<https://arxiv.org/abs/2609.01573>

48 CONF

Mechanism Design for Alignment and Control

ArXiv CS.AI 50%

arXiv:2609.01595v1 Announce Type: cross Abstract: We develop a framework for mechanism design with AI agents whose alignment (preferences) and capabilities (feasible actions and information) are unknown. We want such agents to act on our behalf so mechanisms must incentivize...

<https://arxiv.org/abs/2609.01595>

49 CONF

Efficient SWE Agent Benchmarking via Trajectory-Aware Evaluation

ArXiv CS.AI 50%

arXiv:2609.01603v1 Announce Type: cross Abstract: Evaluating software engineering agents on realistic benchmarks is costly, since each task may require multi-step code exploration, modification, and test execution. Existing efficient evaluation methods select representative...

<https://arxiv.org/abs/2609.01603>

50 **CONF****GeoGR²:Zero-Shot Geospatial Inference via Geostatistically-Guided Iterative Refinement with LLMs****ArXiv CS.AI** **50%**

arXiv:2508.04080v2 Announce Type: replace Abstract: Standard large language model prompting treats geospatial inference as independent, instance-wise prediction, ignoring the fundamental spatial dependencies that govern geographic reality. Consequently, even advanced models...

<https://arxiv.org/abs/2508.04080>

51 **CONF****Uncovering the Computational Ingredients of Human-Like Representations in LLMs****ArXiv CS.AI** **50%**

arXiv:2510.01030v2 Announce Type: replace Abstract: The human ability to translate diverse perceptual and linguistic inputs into structured behavior has been thought to rest on learning robust representations of concepts. The rapid advancement of transformer-based large language...

<https://arxiv.org/abs/2510.01030>

52 **CONF****Compositional Machine Design as Program Synthesis with LLMs****ArXiv CS.AI** **50%**

arXiv:2510.14980v3 Announce Type: replace Abstract: Large language models (LLMs) have shown strong abilities in writing and revising programs, yet many program-synthesis benchmarks still evaluate programs in symbolic or digital environments. We introduce compositional machine...

<https://arxiv.org/abs/2510.14980>

53 **CONF****Multi-Agent LLM Orchestration Achieves Deterministic, High-Quality Decision Support for Incident Response****ArXiv CS.AI** **50%**

arXiv:2511.15755v3 Announce Type: replace Abstract: Large language models (LLMs) promise to accelerate incident response in production systems, yet single-agent approaches generate vague, unusable recommendations. We present MyAntFarm.ai, a reproducible containerized framework...

<https://arxiv.org/abs/2511.15755>

54 **CONF****LifeAgentBench: Benchmarking LLMs for Long-Horizon, Cross-Dimensional Lifestyle Health Reasoning****ArXiv CS.AI** **50%**

arXiv:2601.13880v2 Announce Type: replace Abstract: Personalized lifestyle health analysis requires long-horizon, multi-dimensional reasoning over heterogeneous lifestyle signals, and recent advances in mobile sensing and large language models (LLMs) make such support...

<https://arxiv.org/abs/2601.13880>

55 **CONF****Ontology-Guided Neuro-Symbolic Inference: Grounding Language Models with Mathematical Domain Knowledge****ArXiv CS.AI** **50%**

arXiv:2602.17826v2 Announce Type: replace Abstract: Language models exhibit fundamental limitations -- hallucination, brittleness, and lack of formal grounding -- that are particularly problematic in high-stakes specialist fields requiring verifiable reasoning. I investigate...

<https://arxiv.org/abs/2602.17826>

56 CONF

Verifiable abstention makes AI leak diagnosis accountable in urban water distribution networks

ArXiv CS.AI 50%

arXiv:2608.18836v2 Announce Type: replace Abstract: Leak localization is usually evaluated as forced-choice prediction, although sparse hydraulic observations may not justify excavation. Here, we quantify a pressure-information limit and use it to recast localization as...

<https://arxiv.org/abs/2608.18836>

57 CONF

RACE: Scalable Statistical Estimation of Functional Consistency in LLM Neurons

ArXiv CS.AI 50%

arXiv:2608.24758v2 Announce Type: replace Abstract: Discovering stable neuron behavior across entire domains remains a challenge in mechanistic interpretability. Existing methods often rely on instance-level point estimates or computationally expensive procedures, which either...

<https://arxiv.org/abs/2608.24758>

58 CONF

AI Alignment through a Game-theoretic Lens: A Survey

ArXiv CS.AI 50%

arXiv:2608.27910v2 Announce Type: replace Abstract: As large language models and increasingly capable AI agents are deployed in high-risk settings, aligning them with complex human values has become a central challenge. Existing alignment methods, while effective in improving...

<https://arxiv.org/abs/2608.27910>

59 CONF

Accelerating Unified Multimodal Models with Core-Expansion Routing and Unified Computation Scheduling

ArXiv CS.AI 50%

arXiv:2608.29291v2 Announce Type: replace Abstract: Unified multimodal models jointly support understanding and generation, but incur substantial redundant computation across tokens, layers, and generation timesteps. Through token-importance probing, we identify an asymmetric...

<https://arxiv.org/abs/2608.29291>

60 CONF

Will the User Ever Know? Covert Indirect Prompt Injection Attacks on Tool-Using LLM Agents

ArXiv CS.AI 50%

arXiv:2608.30362v2 Announce Type: replace Abstract: As LLM agents take real-world actions through tools, indirect prompt injection (IPI) has emerged as a serious threat. The standard metric, Attack Success Rate (ASR), counts whether an injection succeeds but ignores what the...

<https://arxiv.org/abs/2608.30362>

61 CONF

Building Expressive and Tractable Probabilistic Generative Models: A Review

ArXiv CS.AI 50%

arXiv:2402.00759v4 Announce Type: replace-cross Abstract: We present a comprehensive survey of the advancements and techniques in the field of tractable probabilistic generative modeling, primarily focusing on Probabilistic Circuits (PCs). We provide a unified perspective on the...

<https://arxiv.org/abs/2402.00759>

62 CONF

Keep Everyone Happy: Online Fair Division of Numerous Items with Few Copies

ArXiv CS.AI 50%

arXiv:2408.12845v3 Announce Type: replace-cross Abstract: This paper considers a novel variant of the online fair division problem involving multiple agents in which a learner sequentially observes an indivisible item that must be irrevocably allocated to one of the agents to...

<https://arxiv.org/abs/2408.12845>

63 CONF

One-shot Style Transfer LLM log-probabilities for Authorship Attribution and Verification

ArXiv CS.AI 50%

arXiv:2510.13302v4 Announce Type: replace-cross Abstract: Computational stylometry studies writing style through quantitative textual patterns, enabling applications such as authorship attribution, identity linking, and plagiarism detection. Despite the relevance of language...

<https://arxiv.org/abs/2510.13302>

64 CONF

SEBA: Sample-Efficient Black-Box Attacks on Visual Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2511.09681v2 Announce Type: replace-cross Abstract: Visual reinforcement learning has achieved remarkable progress in visual control and robotics, but its vulnerability to adversarial perturbations remains underexplored. Most existing black-box attacks focus on...

<https://arxiv.org/abs/2511.09681>

65 CONF

A Machine Learning-Driven Solution for Denoising Inertial Confinement Fusion Images

ArXiv CS.AI 50%

arXiv:2511.16717v3 Announce Type: replace-cross Abstract: Neutron imaging is essential for diagnosing and optimizing inertial confinement fusion implosions at the National Ignition Facility. Due to the required 10-micrometer resolution, however, neutron image require image...

<https://arxiv.org/abs/2511.16717>

66 CONF

Channel-Adaptive Edge AI: Maximizing Inference Throughput by Adapting Computational Complexity to Channel States

ArXiv CS.AI 50%

arXiv:2603.03146v2 Announce Type: replace-cross Abstract: \emph{Integrated communication and computation} (IC²) has emerged as a new paradigm for enabling efficient edge inference in sixth-generation (6G) networks. However, the design of IC² technologies is hindered by the...

<https://arxiv.org/abs/2603.03146>

67 CONF

Global Attention with Linear Complexity for Exascale Generative Data Assimilation in Earth System Prediction

ArXiv CS.AI 50%

arXiv:2604.16590v2 Announce Type: replace-cross Abstract: Accurate Earth system prediction requires state inference from incomplete observations, but conventional two-stage data assimilation (DA) is computationally prohibitive because repeated PDE-based ensemble forecasts,...

<https://arxiv.org/abs/2604.16590>

68 CONF

The Topological Trouble With Transformers

ArXiv CS.AI 50%

arXiv:2604.17121v5 Announce Type: replace-cross Abstract: Transformers encode structure in sequences via an expanding contextual history. However, their purely feedforward architecture fundamentally limits dynamic state tracking. State tracking -- the iterative updating of...

<https://arxiv.org/abs/2604.17121>

69 CONF

MusTBench: Benchmarking and Advancing Temporal Grounding in Music LLMs

ArXiv CS.AI 50%

arXiv:2605.29300v2 Announce Type: replace-cross Abstract: Recent Large Audio-Language Models (LALMs) have demonstrated promising abilities in understanding musical content. However, whether their responses are grounded in the correct temporal regions of the audio remains...

<https://arxiv.org/abs/2605.29300>

70 CONF

PlanarBench: Evaluating LLM Spatial Reasoning via Planar Graph Drawing

ArXiv CS.AI 50%

arXiv:2606.02010v2 Announce Type: replace-cross Abstract: Existing LLM graph benchmarks typically ask models to answer graph-theoretic questions or compute symbolic solutions rather than construct spatial layouts. Within-task difficulty is also primarily stratified by vertex...

<https://arxiv.org/abs/2606.02010>

71 CONF

Enabling KV Caching of Shared Prefix for Diffusion Language Models

ArXiv CS.AI 50%

arXiv:2606.07571v3 Announce Type: replace-cross Abstract: Key-value (KV) caching for shared prefixes is essential for high-throughput large language model (LLM) serving, but it faces critical challenges in emerging diffusion language models (DLMs). In DLMs, bidirectional...

<https://arxiv.org/abs/2606.07571>

72 CONF

Energy-Based Transformers as Predictors of Reading Difficulty

ArXiv CS.AI 50%

arXiv:2606.23382v2 Announce Type: replace-cross Abstract: Transformer language models have become established tools for modeling human sentence processing, with measures such as surprisal and attention entropy serving as effective predictors of reading difficulty that together...

<https://arxiv.org/abs/2606.23382>

73 CONF

Backspace as a Natural Experiment: An Accelerated Failure Time Model of Selective Post-Error Motor Impairment in Parkinsons Disease

ArXiv CS.AI 50%

arXiv:2607.24796v2 Announce Type: replace-cross Abstract: Parkinson's disease (PD) selectively impairs distinct stages of motor control. Using backspace events as natural error-correction episodes in the public neuroQWERTY MIT-CSXPD dataset (n=57 subjects, 27 PD with UPDRS-III...

<https://arxiv.org/abs/2607.24796>

74 **CONF****Can We Trust In-Distribution Success? Locked Evaluation Reveals Transfer Failure and Sampling-Depth Entanglement in CRISPRi Perturbation Prediction****ArXiv CS.AI 50%**

arXiv:2608.00152v3 Announce Type: replace-cross Abstract: AI evaluation can support the wrong inference when an in-domain benchmark success does not survive distribution shift, or when the benchmark endpoint is entangled with a design factor. We study this problem in CRISPRi...

<https://arxiv.org/abs/2608.00152>

75 **CONF****Debiased Inference for AI-Generated Data without Gold-Standard Labels: Identification via Multiple Imperfect Measurements****ArXiv CS.AI 50%**

arXiv:2608.18294v2 Announce Type: replace-cross Abstract: An increasing number of scholars use AI to measure variables they subsequently include in downstream analyses. Although AI-measured variables are often analyzed as if observed without error, ignoring prediction errors in...

<https://arxiv.org/abs/2608.18294>