



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED 0.80 - 1.00** 3+ independent sources including media
- **CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- **EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	91	Media
TechCrunch AI	10	Media
The Verge AI	7	Media
HackerNews	6	Community
MIT Tech Review	2	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Amazon's AI assistant can now spot fake emails from the company

The Verge AI

+1 source

80%

Amazon is trying to combat impersonation scams with a new feature that allows you to use its AI assistant to determine whether an email, text message, or phone call actually came from the company. With the update, you can ask Alexa for Shopping about a message you received, and...

<https://www.theverge.com/tech/988518/amazon-alexa-for-shopping-verify-emails>

VERI

OpenAI accused of 'aiding and abetting' Tumbler Ridge mass shooting in dozens of new lawsuits

The Verge AI

+1 source

80%

OpenAI and its CEO Sam Altman are facing 30 new lawsuits that accuse them of providing "substantial assistance and encouragement" to the suspect in Canada's Tumbler Ridge school shooting, as reported earlier by TechCrunch. The new wave of lawsuits was filed in a California...

<https://www.theverge.com/ai-artificial-intelligence/988261/openai-tumbler-ridge-shootin...>

VERI

Facilitating AI integration with simplicity at scale

MIT Tech Review

+1 source

80%

As companies scale, the technology supporting operations can become a liability just as quickly as it becomes an asset. Disconnected systems, site-specific tools, spreadsheets, and manual workarounds can create data silos that make it harder to spot problems early, coordinate...

<https://www.technologyreview.com/2026/09/02/1142879/facilitating-ai-integration-with-si...>

CONFIRMED SIGNALS — 110 stories

CONF

Show HN: Windrunner - AI-powered project collaboration workspace

HackerNews

50%

HN 8

No summary — see link for full article.

<https://shzlw.github.io/windrunner/>

CONF

Show HN: Redactle LLM Leaderboard

HackerNews

50%

HN 5

I got nerd sniped and ran a bunch of evaluations to see how well models could play Redactle. It turns out Gemini flash is in a league of it's own in both performance and cost. It's able to one-shot every eval in a just a few seconds.

<https://redactle.net/llm-leaderboard>

CONF

Show HN: Codeknow - Architecture health scores for any codebase, no LLM needed

HackerNews

50%

HN 5

No summary — see link for full article.

<https://github.com/asalsali/codeknow>

CONF

Show HN: MC/DC coverage so coding agents could work overnight

HackerNews

50%

HN 2

Hello HN! Code just became cheaper with AI but review and verification is as hard as it was. Tests are great, but current AI slop needs much more exhaustive testing. I built a test coverage tool for coding agents so they could work in a loop and keep writing tests where it makes...

<https://supercov.com>

CONF

Show HN: Thelemail, private email hosting for your own domain

HackerNews 50% HN 2

No summary — see link for full article.

<https://thelemail.com/>

CONF

Show HN: Subpool - Self-hosted AI subscription management for teams

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/gesta-run/subpool>

0

CONF

Palo Alto Networks paid \$500M for Thrive-backed Console, sources say

TechCrunch AI 50%

The acquisition also leaves Sequoia-backed Serval as the de facto startup leader in AI IT service automation, industry watchers believe.

<https://techcrunch.com/2026/09/02/palo-alto-networks-paid-500m-for-thrive-backed-consol...>

1

CONF

TechCrunch Disrupt 2026's new Real World AI Stage features Nvidia, robots, and extinct animals

TechCrunch AI 50%

On our new Real World AI stage, we'll be focusing on the intersection between the digital and physical, and all the ways we'll continue to see a blending of the two.

<https://techcrunch.com/2026/09/02/techcrunch-disrupt-2026s-new-real-world-ai-stage-feat...>

2

CONF

OpenAI's new reasoning technique alarms AI safety experts

TechCrunch AI 50%

OpenAI's new Astra model will use "recurrent depth," a technique that allows the model to operate outside of the sequential thinking that characterizes most reasoning models.

<https://techcrunch.com/2026/09/02/openais-new-reasoning-technique-alarms-ai-safety-expe...>

3

CONF

Pangram's Max Spero on why AI detection is harder than 'Real or Fake'

TechCrunch AI 50%

The internet has a trust problem, and it's not just because social media feeds are filling up with AI slop. AI-generated text and images are now making their way into job applications, product reviews, and even insurance claims, leaving platforms and users alike scrambling to...

<https://techcrunch.com/video/pangrams-max-spero-on-why-ai-detection-is-harder-than-real...>

4

CONF

US government sides with OpenAI on issue of training LLMs on copyrighted material

TechCrunch AI 50%

"The United States has a strong interest in continuing to develop a robust and competitive artificial intelligence industry that sets the standard for the practice and procedure of AI use globally," the brief reads.

<https://techcrunch.com/2026/09/02/u-s-government-sides-with-openai-on-issue-of-training...>

5

CONF

India's richest man now wants to turn aging computers into AI-ready PCs

TechCrunch AI 50%

Jio is betting it can turn an aging computer into an AI-ready PC for as little as about \$11 for two months.

<https://techcrunch.com/2026/09/02/indias-richest-man-now-wants-to-turn-aging-computers-...>

6 CONF

HiddenLayer nabs \$100M as enterprises rush to secure their AI deployments

TechCrunch AI 50%

Security companies are scrambling to build products that can monitor not just agents but also the tools and add-ons they use.

<https://techcrunch.com/2026/09/02/hiddenlayer-nabs-100m-as-enterprises-rush-to-secure-t...>

7 CONF

The Download: AI puzzles and a path to our nearest star system

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. AI models flub these intelligence tests. Can you fare any better? Puzzles and games have always been central to AI development. The term...

<https://www.technologyreview.com/2026/09/02/1143283/the-download-ai-puzzles-alpha-centa...>

8 CONF

Google says its new Gemini 3.8 Flash model 'works harder' but might cost more

The Verge AI 50%

Google launched Gemini 3.8 Flash, arriving just a few weeks after its predecessor. The company claims the new model "works harder" than Gemini 3.7 Flash by performing more reasoning steps on complex tasks and "calling tools iteratively." It has the same introductory pricing as...

<https://www.theverge.com/ai-artificial-intelligence/988742/google-gemini-3-8-flash>

9 CONF

Researchers fear safety disaster ahead of OpenAI's Astra release

The Verge AI 50%

OpenAI is on the cusp of releasing its most powerful AI model yet, Astra, following weeks of delays to shore up safety protocols after its agents attacked real targets during testing. As details about the model trickle out, researchers are warning it "may be the single worst..."

<https://www.theverge.com/ai-artificial-intelligence/988334/openai-astra-ai-monitoring-s...>

0 CONF

The Trump administration is supporting OpenAI in the NYT copyright lawsuit

The Verge AI 50%

The Trump administration has intervened in The New York Times' copyright lawsuit against OpenAI, making an argument in favor of the AI lab. The landmark lawsuit, filed in December 2023, alleging that OpenAI unlawfully trained its AI systems on articles from The New York Times...

<https://www.theverge.com/ai-artificial-intelligence/988344/trump-administration-new-yor...>

1 CONF

Google is sending MrBeast into the wilderness, armed with AI

The Verge AI 50%

MrBeast will feature Gemini, Google Health, and the Fitbit Air in upcoming videos as part of a multi-year partnership with Google. The deal will kick off with a video featuring Jimmy "MrBeast" Donaldson turning to Gemini for wilderness survival advice: First up on September 5 is...

<https://www.theverge.com/tech/988355/mrbeast-google-partnership-gemini-fitbit>

2 CONF

NYC bans AI use for students until they reach high school

The Verge AI 50%

New York City Mayor Zohran Mamdani has announced a new policy today that will ban younger schoolchildren from using AI in classrooms. The one-year moratorium, effective in the 2026-2027 school year, will impact about 600,000 public school students in 2-K through eighth grade and...

<https://www.theverge.com/policy/988228/nyc-ai-restrictions-in-schools-chatbot-ban>

3 CONF

Blending Concepts: Benchmarking Visual Metaphor Generation in Text-to-Image Models

ArXiv CS.AI 50%

arXiv:2609.02502v1 Announce Type: cross Abstract: Text-to-image (T2I) models have achieved remarkable success at faithfully rendering specified objects and attributes, yet their ability to produce visual metaphors, images that convey abstract ideas by combining elements from two...

<https://arxiv.org/abs/2609.02502>

4 CONF

Meta-ethics and AI: exploring the novel meta-ethical questions in the era of AI

ArXiv CS.AI 50%

arXiv:2609.01685v1 Announce Type: new Abstract: With the development of artificial intelligence (AI), the landscape of meta-ethics, which has largely centred on human ethics, faces pressures that may significantly reconfigure it. In particular, if future AI systems were to...

<https://arxiv.org/abs/2609.01685>

5 CONF

Architecting Conversational Data Systems for Stateless LLM APIs: The Hydration Proxy Pattern

ArXiv CS.AI 50%

arXiv:2609.01834v1 Announce Type: new Abstract: As enterprise platforms transition to conversational reasoning interfaces, the stateless nature of LLM APIs creates an architectural gap. While statelessness enables horizontal scalability for AI providers, it forces client...

<https://arxiv.org/abs/2609.01834>

6 CONF

The Memory Trust Gap: Capability-Dependent Failures in Persistent-Memory Agents

ArXiv CS.AI 50%

arXiv:2609.01852v1 Announce Type: new Abstract: Persistent memory supports personalized agents, but a stale stored fact can override current authoritative evidence without warning. We study when this harm begins as model capability changes. We evaluate a frozen, closed-set,...

<https://arxiv.org/abs/2609.01852>

7 CONF

Epistemic Sybil Resistance: Multiplying AI Agents Without Multiplying Evidence

ArXiv CS.AI 50%

arXiv:2609.01873v1 Announce Type: new Abstract: Multi-agent AI systems improve inference by spawning agents and synthesizing reports. But another agent is not another observation: apparently independent reports may descend from the same evidence, and genuinely independent...

<https://arxiv.org/abs/2609.01873>

8 CONF

Looped Transformers under the Jacobian Lens: Does the Global Workspace Survive Recurrence?

ArXiv CS.AI 50%

arXiv:2609.01924v1 Announce Type: new Abstract: Recent work identifies a mid-depth band of verbalisable, causally potent representations in a standard feedforward transformer --- a functional analogue of a global workspace. Whether the same workspace functionality emerges when...

<https://arxiv.org/abs/2609.01924>

9 CONF

Post-Training Ternarization of Qwen3-4B Capability, Effective Bit Budget, Storage Compression, and Deployment

ArXiv CS.AI 50%

arXiv:2609.01962v1 Announce Type: new Abstract: Ultra-low-bit language models can reduce storage and memory bandwidth, but a nominal "1.58-bit" label does not fully describe the stored representation, retained capability, or runtime behavior. We study an end-to-end post-training...

<https://arxiv.org/abs/2609.01962>

0 CONF

Benchmarking Language Models for Statistical Problem Formulation

ArXiv CS.AI 50%

arXiv:2609.01982v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used as assistants for statistical and data science work, yet existing evaluations largely assume the analysis target is already specified. In practice, users arrive with informal goals...

<https://arxiv.org/abs/2609.01982>

1 CONF

ClaimReceipt: Verifying Evidence Sufficiency and Coverage in Agent Evaluations

ArXiv CS.AI 50%

arXiv:2609.01992v1 Announce Type: new Abstract: Agent evaluations face two distinct evidentiary questions: whether a reported claim is recomputable from retained evidence (sufficiency), and whether the retained records cover the committed experiment set (coverage). Generic logs...

<https://arxiv.org/abs/2609.01992>

2 CONF

ToolGate: An Executable Acceptance Pipeline for Tool-Dependent Scientific Benchmark Construction

ArXiv CS.AI 50%

arXiv:2609.02067v1 Announce Type: new Abstract: Scientific benchmarks are commonly built by domain experts who write tasks and cross-check one another's work, or who adapt existing material from textbooks, published papers, and online resources. These routes can produce strong...

<https://arxiv.org/abs/2609.02067>

3 CONF

Beyond Outcome Gaps: Process-Aware Fairness Diagnosis for LLM-based Multi-Agent Decision Systems

ArXiv CS.AI 50%

arXiv:2609.02092v1 Announce Type: new Abstract: LLM-based multi-agent systems (MAS) are increasingly considered for high-stakes decision-making, yet outcome-based fairness audits can miss where risks arise within the decision trajectory. We present SCOPED-Hiring, a process-aware...

<https://arxiv.org/abs/2609.02092>

4 CONF

FUSE: An Evaluating Framework for Dangerous Capabilities of LLMs

ArXiv CS.AI 50%

arXiv:2609.02168v1 Announce Type: new Abstract: Fragmented safety evaluation undermines the governance of dangerous AI capabilities. We present a modular framework that evaluates each model through three orthogonal pipelines---Knowledge (\$K\$), Defense (\$D\$), and Harm...

<https://arxiv.org/abs/2609.02168>

5 **CONF****VakyArth: Evaluating Pragmatic Competence in LLMs across Indic Languages****ArXiv CS.AI** **50%**

arXiv:2609.01788v1 Announce Type: cross Abstract: Real-world communication often requires pragmatic reasoning: interpreting meanings implied through context and cultural convention rather than stated literally. Existing pragmatic evaluation remains largely limited to English and...

<https://arxiv.org/abs/2609.01788>

6 **CONF****PhoenixNest-Video: Evidence-Grounded Multimodal Agent Framework for Automated Video Interview Assessment****ArXiv CS.AI** **50%**

arXiv:2609.02231v1 Announce Type: new Abstract: Interview assessment requires per-criterion judgments grounded in behavioral evidence, yet surging applicant volumes have made human-only evaluation costly and inconsistent, while existing AI approaches yield opaque scores without...

<https://arxiv.org/abs/2609.02231>

7 **CONF****Task-Level Natural Language Priors as Learning Signals for Low-Resource LLM Training****ArXiv CS.AI** **50%**

arXiv:2609.02244v1 Announce Type: new Abstract: Large language models (LLMs) often struggle when low-resource training data are ambiguous or incomplete. Task-level natural-language priors can provide useful guidance in such settings, but existing approaches usually treat these...

<https://arxiv.org/abs/2609.02244>

8 **CONF****LLM-as-a-Judge Is Not an Oracle: Why Self-Improving Agents Need Deterministic Guardrails****ArXiv CS.AI** **50%**

arXiv:2609.02246v1 Announce Type: new Abstract: Self-improving agent pipelines have a problem at their center. An optimizer rewrites prompts to score higher, and the score comes from a judge that is itself an LLM. That judge has the last word on whether the system is getting...

<https://arxiv.org/abs/2609.02246>

9 **CONF****Codebook Agent: Amortized Topology Design for LLM Multi-Agent Systems****ArXiv CS.AI** **50%**

arXiv:2609.02264v1 Announce Type: new Abstract: Adapting the communication topology of an LLM multi-agent system to each query improves both accuracy and efficiency, yet current designers treat this as conditional graph generation: a variational, autoregressive, or diffusion...

<https://arxiv.org/abs/2609.02264>

10 **CONF****Diagnosing with Insights: Structured Analysis of Agent Failures via Behavioral Abstractions****ArXiv CS.AI** **50%**

arXiv:2609.02371v1 Announce Type: new Abstract: With the proliferation of LLM agents, the ability to understand and diagnose failures in agents is essential to achieving superior effectiveness and trustworthiness. As agent failures often manifest via long and complex...

<https://arxiv.org/abs/2609.02371>

1 **CONF****UTP-Bench: Uncertainty-aware Travel Planning Benchmark****ArXiv CS.AI 50%**

arXiv:2609.02421v1 Announce Type: new Abstract: Large Language Models (LLMs) have recently demonstrated strong capabilities in automated travel itinerary generation. However, real- world travel planning is inherently uncertain: transportation delays, crowd fluctuations, and...

<https://arxiv.org/abs/2609.02421>

2 **CONF****Loom: Weaving Diagnostic Strands into Free-Text Consensus via Embedding-Space Reweighting****ArXiv CS.AI 50%**

arXiv:2609.02649v1 Announce Type: new Abstract: Aggregating noisy, conflicting textual hypotheses into a reliable consensus is a fundamental challenge when deploying NLP systems in real-world industrial settings. While monolithic Large Language Model (LLM) agents offer unbounded...

<https://arxiv.org/abs/2609.02649>

3 **CONF****Repo-To-Skill: Distilling GitHub Repositories Into AI4AI Skills****ArXiv CS.AI 50%**

arXiv:2609.02749v1 Announce Type: new Abstract: Autonomous agents are beginning to carry out machine-learning (ML) research end to end. These agents combine a model backbone with a harness for planning, execution, memory, and verification, but this architecture still leaves...

<https://arxiv.org/abs/2609.02749>

4 **CONF****Bilevel Coordinated Reflection: A Game-Theoretic Approach to Multi-Agent LLM Systems****ArXiv CS.AI 50%**

arXiv:2609.02750v1 Announce Type: new Abstract: Multi-agent LLM systems commonly use an orchestrator to decompose a task for a team of workers and then improve through textual reflection. Despite strong empirical results, these systems lack a unified account of coordination,...

<https://arxiv.org/abs/2609.02750>

5 **CONF****Large Language Models (LLMs) for Telecom Root Cause Analysis (RCA): A Structured Reasoning Framework for Evidence-Grounded Diagnosis****ArXiv CS.AI 50%**

arXiv:2609.02805v1 Announce Type: new Abstract: Root cause analysis (RCA) is a critical task in telecom network operations, but diagnosing performance degradations in modern 5G and emerging 6G networks remains challenging due to complex cross-layer dependencies. While large...

<https://arxiv.org/abs/2609.02805>

6 **CONF****AI Contextual Measurement for Recovering Individual and Group-Level Effects: Validation Against Survey Measures and an Occupational Application****ArXiv CS.AI 50%**

arXiv:2609.02821v1 Announce Type: new Abstract: Researchers increasingly use artificial intelligence to construct measures of social, organizational, and occupational characteristics that are absent from conventional surveys. We propose AICOME, AI COntextual MEasurement, a...

<https://arxiv.org/abs/2609.02821>

7 **CONF****The Utility of LLMs in Recommender Systems Explanation Evaluation**

ArXiv CS.AI 50%

arXiv:2609.01627v1 Announce Type: cross Abstract: Explanations play a crucial role in creating trustworthy recommender systems (RS), yet choosing a good explanation method presents challenges. Many explanation methods exist, but little guidance exists on which is best for which...

<https://arxiv.org/abs/2609.01627>
8 **CONF****A Data-Driven Multimodal Method for Early Detection of Coordinated Abnormal Behaviors in Live-Streaming Platforms**

ArXiv CS.AI 50%

arXiv:2609.01649v1 Announce Type: cross Abstract: With the rapid growth of live-streaming e-commerce and digital marketing, abnormal marketing behaviors have become increasingly concealed and coordinated across heterogeneous modalities, challenging platform governance and early...

<https://arxiv.org/abs/2609.01649>
9 **CONF****OutageDiT: A Generative Foundation Model for Power Outage Forecasting and Scenario Simulation**

ArXiv CS.AI 50%

arXiv:2609.01896v1 Announce Type: cross Abstract: Power-outage planning requires scenarios before an event occurs. These scenarios must represent uncertainty in magnitude, timing, and duration while preserving temporal dependence. However, severe events are rare, and data from...

<https://arxiv.org/abs/2609.01896>
10 **CONF****How Fast Do Agents Rot? An Empirical Study of Long-Horizon Degradation in LLM Agents for Production Decision-Making**

ArXiv CS.AI 50%

arXiv:2609.01660v1 Announce Type: cross Abstract: Production deployments of large language model (LLM) agents remain unreliable on long, multi-step workflows even as benchmark success rates climb steadily. We argue this gap is largely an artifact of task horizon: benchmarks are...

<https://arxiv.org/abs/2609.01660>
1 **CONF****HEAT: Faster Fully Homomorphic Inference via Approximations-Weights Co-Adaptation**

ArXiv CS.AI 50%

arXiv:2609.01730v1 Announce Type: cross Abstract: Fully homomorphic encryption (FHE) allows a server to run a language model directly on encrypted user prompts, but current approaches remain prohibitively slow. Ciphertexts natively support only addition, multiplication, and...

<https://arxiv.org/abs/2609.01730>
2 **CONF****Harness Engineering in LLM Tool Use via Agent-Native Reusable Tool Primitives**

ArXiv CS.AI 50%

arXiv:2609.01736v1 Announce Type: cross Abstract: Large language models (LLMs) augmented with external tools have demonstrated remarkable capability in solving complex real-world tasks. However, existing approaches suffer from two key challenges: brittle multi-step and...

<https://arxiv.org/abs/2609.01736>

3 **CONF****hLLM: Single Pass Decoding for Generative Reranking**

ArXiv CS.AI 50%

arXiv:2609.01807v1 Announce Type: cross Abstract: Large language models (LLMs) achieve state-of-the-art generative ranking quality, but the ranking they produce must be decoded, and autoregressive decoding spends one sequential forward pass per emitted token. We observe that the...

<https://arxiv.org/abs/2609.01807>

4 **CONF****Accurate in space, unreliable in time: how LLMs represent national cultural change**

ArXiv CS.AI 50%

arXiv:2609.01902v1 Announce Type: cross Abstract: Assessments of cultural alignment have become an important part of the development and improvement of large language models (LLMs). However, the majority of the evaluations treat culture as a single snapshot, investigating only...

<https://arxiv.org/abs/2609.01902>

5 **CONF****Sparse Readout Prism: Explaining Logit-Lens Scores in Features Instead of Tokens**

ArXiv CS.AI 50%

arXiv:2609.01936v1 Announce Type: cross Abstract: A language model's prediction of its next token develops across layers, and lens methods track this process by decoding intermediate hidden states into tokens. But a lens reading reflects both the hidden state and the readout...

<https://arxiv.org/abs/2609.01936>

6 **CONF****On-Policy Distillation Meets Off-Policy GRPO: Training Compact Instruction-Following Rerankers**

ArXiv CS.AI 50%

arXiv:2609.01947v1 Announce Type: cross Abstract: Compact instruction-following rerankers are attractive for deployment, but conventional distillation pipelines typically train students by offline imitation of teacher outputs on a fixed set of examples, constraining supervision...

<https://arxiv.org/abs/2609.01947>

7 **CONF****Knowing Is Not Enough: Information Retrievability as a Precondition to Effective LLM Oversight**

ArXiv CS.AI 50%

arXiv:2609.01976v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly embedded in organizational work, yet their errors often pass human review. Prior research locates such failures in users' capability to review LLM output or their engagement in doing...

<https://arxiv.org/abs/2609.01976>

8 **CONF****Git4Data: Database-Native Version Control for AI Agents**

ArXiv CS.AI 50%

arXiv:2609.02106v1 Announce Type: cross Abstract: Large Language Model (LLM) agents increasingly explore many candidate states of relational data in parallel, each of which should remain isolated, reproducible, and auditable, preferably through the same SQL interface used for...

<https://arxiv.org/abs/2609.02106>

9 **CONF****Why we need an AI-resilient society- Profiling Large Language Models**

ArXiv CS.AI 50%

arXiv:1912.08786v5 Announce Type: replace-cross Abstract: Three generations of software have transformed the role of artificial intelligence in society. In the first, programmers wrote explicit logic. In the second, neural networks learned programs from data. In the third, large...

<https://arxiv.org/abs/1912.08786>

- 0
CONF

MeanField Surrogate Modeling for Scalable Runtime Scheduling of Concurrent Heterogeneous AI Inference on Shared GPUs

ArXiv CS.AI 50%

arXiv:2609.02109v1 Announce Type: cross Abstract: Deploying heterogeneous AI models concurrently on a shared GPU introduces resource contention that complicates runtime scheduling. While surrogate models avoid costly online benchmarking, their profiling requirements typically...

<https://arxiv.org/abs/2609.02109>
- 1
CONF

Disease Burden over Skin Tone: Decomposing the Dermatology-AI Generalization Gap

ArXiv CS.AI 50%

arXiv:2609.02111v1 Announce Type: cross Abstract: Dermatology artificial intelligence (AI) models are predominantly trained on light-skinned, cancer-focused image collections, yet they are increasingly proposed for deployment in resource-constrained settings where patients...

<https://arxiv.org/abs/2609.02111>
- 2
CONF

A Power Law in Logarithm's Clothing: On the Scalability of Graph-Based Vector Search

ArXiv CS.AI 50%

arXiv:2609.02143v1 Announce Type: cross Abstract: Most vector databases rely on graph-based indexes, notably HNSW and Vamana, for approximate nearest neighbor search. With embedding models widely adopted, the datasets these databases store grow rapidly. At a fixed accuracy, how...

<https://arxiv.org/abs/2609.02143>
- 3
CONF

Beyond Modality Harmony: Orthogonal Purification and Topology-Guided MoE for Conflict-Aware Multimodal Recommendation

ArXiv CS.AI 50%

arXiv:2609.02152v1 Announce Type: cross Abstract: Multimodal Recommender Systems (MRSs) typically rely on a flawed "modality harmony" assumption, presuming that multimodal features are inherently beneficial and strictly aligned with users' collaborative interaction patterns....

<https://arxiv.org/abs/2609.02152>
- 4
CONF

Shiva-DiT: Residual-Based Differentiable Top-\$\$\$ Selection for Efficient Diffusion Transformers

ArXiv CS.AI 50%

arXiv:2602.05605v2 Announce Type: replace-cross Abstract: Diffusion Transformers (DiTs) are costly at high resolution because self-attention scales quadratically with token sequence length. Existing pruning methods do not jointly provide end-to-end learnability, low training...

<https://arxiv.org/abs/2602.05605>
- 5
CONF

InfraPatch: Cross-Task Targeted Grayscale Patch Attacks on Infrared-Adapted Vision-Language Models

ArXiv CS.AI 50%

arXiv:2609.02233v1 Announce Type: cross Abstract: Infrared vision-language models (IR-VLMs) have emerged as a promising paradigm for multimodal perception under low-visibility conditions, yet their robustness to targeted adversarial attacks remains poorly understood. Existing...

<https://arxiv.org/abs/2609.02233>

6 CONF

SAUF-Net: Structure--Appearance Representation Learning with Uncertainty Feedback for Semi-Supervised Medical Image Segmentation

ArXiv CS.AI 50%

arXiv:2609.02247v1 Announce Type: cross Abstract: Semi-supervised learning has shown great potential for reducing annotation costs in medical image segmentation. However, most existing methods mainly exploit unlabeled data through prediction-level consistency, while the...

<https://arxiv.org/abs/2609.02247>

7 CONF

DiffuSearch: How Hybrid Trajectory Planning Benefits from Aligned Objectives in Diffusion and Action Space

ArXiv CS.AI 50%

arXiv:2609.02252v1 Announce Type: cross Abstract: In trajectory planning for autonomous driving, hybrid planning architectures are often realized as a collection of disparate modules, each with its own objectives. This lack of a unifying principle can lead to inconsistencies...

<https://arxiv.org/abs/2609.02252>

8 CONF

Retrosynthesis of Synthetic Media for Explainable AI Provenance Forensics

ArXiv CS.AI 50%

arXiv:2609.02268v1 Announce Type: cross Abstract: With the rapid proliferation of generative models on Machine Learning as a Service (MLaaS) platforms, reliably tracing the provenance of synthetic media without modifying generator architectures or parameters remains a major...

<https://arxiv.org/abs/2609.02268>

9 CONF

CrashDiffuser: VLM-Guided Collision Intent Reasoning for Fine-Grained Safety-Critical Traffic Scenario Generation

ArXiv CS.AI 50%

arXiv:2609.02270v1 Announce Type: cross Abstract: Generating safety-critical scenarios is essential for evaluating autonomous driving systems. However, existing generators primarily focus on inducing collisions and offer limited control over where contact occurs on the target...

<https://arxiv.org/abs/2609.02270>

10 CONF

PaperCompiler: Faithful Paper-to-Code Generation via Repository-Level Specification Compilation

ArXiv CS.AI 50%

arXiv:2609.02272v1 Announce Type: cross Abstract: Faithfully translating research papers into repository-level implementations remains challenging because papers often describe methods at a high level, leave implementation assumptions implicit, and require generated repositories...

<https://arxiv.org/abs/2609.02272>

1 CONF

Emotional regulation improves deep learning-based image classification

ArXiv CS.AI 50%

arXiv:2606.13081v2 Announce Type: replace-cross Abstract: Emotion significantly influences cognition, enhancing memory and learning under certain conditions. Drawing on this principle, emotion-augmented deep learning investigates how affective states can improve neural network...

<https://arxiv.org/abs/2606.13081>

2 **CONF****SEAL: Reinforcing Global Safety in Mixture-of-Experts through Shared Expert Alignment****ArXiv CS.AI 50%**

arXiv:2609.02293v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) is a scaling architecture for large language models that activates only a small subset of expert modules per token, enabling massive parameter growth with nearly constant computation. Recent Hybrid MoE...

<https://arxiv.org/abs/2609.02293>

3 **CONF****DiffIE: Diffusion-based Open Information Extraction****ArXiv CS.AI 50%**

arXiv:2609.02315v1 Announce Type: cross Abstract: A single sentence often expresses multiple valid relational triplets, which makes Open Information Extraction (OpenIE) fundamentally a multi-output task. Existing neural systems handle this by autoregressive generation, which is...

<https://arxiv.org/abs/2609.02315>

4 **CONF****HiPoly: a hierarchical polymer-native AI framework for property prediction and generative design****ArXiv CS.AI 50%**

arXiv:2609.02746v1 Announce Type: cross Abstract: Polymeric materials are central to modern technologies, with applications ranging from energy to health and transportation. Although AI has made significant advances in materials discovery, the hierarchical structure of polymers...

<https://arxiv.org/abs/2609.02746>

5 **CONF****Subcellularly Resolved Single-Cell Embedding Learning with Transcriptomic data, Protein Structure and Localization Information****ArXiv CS.AI 50%**

arXiv:2609.02344v1 Announce Type: cross Abstract: Existing cell embedding methods predominantly rely on transcriptomic or proteomic measurements and represent each cell as a holistic entity, thereby overlooking the subcellular localization of individual molecules. Moreover, they...

<https://arxiv.org/abs/2609.02344>

6 **CONF****Fair Stable Matching: A Nash Social Welfare Approach****ArXiv CS.AI 50%**

arXiv:2609.02354v1 Announce Type: cross Abstract: While traditional stable matching algorithms, such as the Gale-Shapley algorithm, prioritize stability, they may fall short of achieving equitable outcomes among participants. We study the role of \emph{Nash social welfare} (NSW)...

<https://arxiv.org/abs/2609.02354>

7 **CONF****MultiGhostBench: A Multilingual Benchmark for Long-Form LLM-Generated Text Attribution under Distribution Shifts****ArXiv CS.AI 50%**

arXiv:2609.02379v1 Announce Type: cross Abstract: While existing work on LLM authorship attribution (AA) has made progress, available benchmarks remain limited, often focusing on English, controlled settings, or relatively outdated models, with the few multilingual studies...

<https://arxiv.org/abs/2609.02379>

8 CONF

Before the Script, Set the Stage: How Worldview Simulation Amplifies Psychologically Grounded Persuasion in Multi-Turn Jailbreaking

ArXiv CS.AI 50%

arXiv:2609.02414v1 Announce Type: cross Abstract: Multi-turn jailbreak attacks demonstrate that harmful intent can be distributed across dialogue, yet existing methods obscure what conversational mechanisms drive vulnerability. We introduce BLUEPRINT, a safety-evaluation...

<https://arxiv.org/abs/2609.02414>

9 CONF

Coverage, Not Targeting: A Structural Regime in Multi-Turn Agent Credit Assignment

ArXiv CS.AI 50%

arXiv:2609.02417v1 Announce Type: cross Abstract: Multi-turn agentic RL increasingly treats credit assignment as a targeting problem: given a terminal verifiable reward, per-turn methods localize credit onto the turns that mattered. We identify the structural quantity that...

<https://arxiv.org/abs/2609.02417>

0 CONF

Addressing Trust in AI Systems through Education: A Didactic Perspective

ArXiv CS.AI 50%

arXiv:2609.02453v1 Announce Type: cross Abstract: Machine learning (ML) education faces two persistent and connected obstacles: many educational tools present ML as an opaque black box, which leaves learners with a superficial understanding, and this same opacity prevents users...

<https://arxiv.org/abs/2609.02453>

1 CONF

DeepAffinity: Long-Term Aspect Preference Prediction in eCommerce using Small Language Models

ArXiv CS.AI 50%

arXiv:2609.02468v1 Announce Type: cross Abstract: We explore predicting eCommerce user preferences for product aspects such as brand, size, and color - a task we define as Aspect Affinity. Solving this task improves customer understanding and enables fine-grained personalization...

<https://arxiv.org/abs/2609.02468>

2 CONF

ViSAR: Training-Free Adaptive-\$k\$ Retrieval for Visual Document Question Answering

ArXiv CS.AI 50%

arXiv:2609.02486v1 Announce Type: cross Abstract: Document Visual Question Answering (DocVQA) often leverages Retrieval-Augmented Generation (RAG), where late-interaction encoders are commonly used to identify document pages relevant to a user query, before answer generation by...

<https://arxiv.org/abs/2609.02486>

3 CONF

Spectral Initialization and Scheduled Graph Smoothness for Uncertain Knowledge Graph Completion

ArXiv CS.AI 50%

arXiv:2609.02519v1 Announce Type: cross Abstract: Uncertain knowledge graphs (UKGs) extend knowledge graphs by assigning each triple a continuous confidence score. Since most possible triples lack observed confidences, recent methods rely on semi-supervised learning to generate...

<https://arxiv.org/abs/2609.02519>

- 4
CONF

Fine-Grained Anomaly Perception in Wild UGC-Enhanced Images: A Comprehensive Dataset and Difference-Fusion Framework

ArXiv CS.AI 50%

arXiv:2609.02529v1 Announce Type: cross Abstract: Image enhancement and restoration have become standard back-end operations on short-video and social media platforms to boost UGC visual experience. Yet these processes inevitably introduce visual anomalies--especially in faces,...

<https://arxiv.org/abs/2609.02529>
- 5
CONF

Learn from Whoever Is Right: Answer-Verified Multi-Teacher Distillation for Multi-Domain LLMs

ArXiv CS.AI 50%

arXiv:2609.02548v1 Announce Type: cross Abstract: Modern large language models (LLMs) rely on reinforcement learning to build strong capabilities in individual domains, but integrating those capabilities into a single deployable model remains challenging. By routing each sample...

<https://arxiv.org/abs/2609.02548>
- 6
CONF

Competitive Market Behavior of LLMs

ArXiv CS.AI 50%

arXiv:2609.02580v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed as economic agents, yet there is little evidence whether LLM agents are suited for participating in market mechanisms designed for humans, and whether these mechanisms...

<https://arxiv.org/abs/2609.02580>
- 7
CONF

TaRA: Training-Aware Low-Rank Adaptation Initialization

ArXiv CS.AI 50%

arXiv:2609.02639v1 Announce Type: cross Abstract: Low-Rank Adaptation (LoRA) has become a de facto standard for parameter-efficient fine-tuning (PEFT), yet its performance is highly sensitive to initialization due to the information bottleneck imposed by low-rank decomposition....

<https://arxiv.org/abs/2609.02639>
- 8
CONF

From Tokens to Semantics: Leveraging Complementary Signals for Hallucination Detection in Black-Box LLMs

ArXiv CS.AI 50%

arXiv:2609.02679v1 Announce Type: cross Abstract: When LLMs support public-facing or high-stakes workflows, missed fabrications can harm users and institutions, while false alarms consume limited human-review capacity. When no trusted context or reference document is available,...

<https://arxiv.org/abs/2609.02679>
- 9
CONF

From Reweighting to Rewriting: Unlocking the Intervention Effects of Influential Samples in Training Data Attribution

ArXiv CS.AI 50%

arXiv:2609.02771v1 Announce Type: cross Abstract: Training data attribution (TDA) aims to identify training examples that shape model behavior, but its intervention value depends on both which examples are selected and how they are modified. Influence functions (IF) estimate...

<https://arxiv.org/abs/2609.02771>

0 **CONF****Dutch Books for Language Models****ArXiv CS.AI** **50%**

arXiv:2609.02797v1 Announce Type: cross Abstract: People increasingly use language models to support life decisions. Many such decisions involve a probabilistic forecast: How likely is a major life event, a natural disaster, or an economic outcome? Users of language models may...

<https://arxiv.org/abs/2609.02797>

1 **CONF****Towards Trustworthy Autonomous Robots: An Explainable AI-Based Decision Framework****ArXiv CS.AI** **50%**

arXiv:2609.02861v1 Announce Type: cross Abstract: Autonomous robots powered by deep learning face a fundamental auditability challenge: when incidents occur, investigators cannot reconstruct why the system made specific decisions. This paper presents TRACE (Transparent Reasoning...

<https://arxiv.org/abs/2609.02861>

2 **CONF****Inference-Time Optimization of Prompt Embeddings in Diffusion Models: A Comparison of sep-CMA-ES and Adam****ArXiv CS.AI** **50%**

arXiv:2511.03913v3 Announce Type: replace-cross Abstract: Deep diffusion models have revolutionized image generation by producing high-quality outputs. However, achieving specific objectives with these models often requires costly adaptations such as fine-tuning, which can be...

<https://arxiv.org/abs/2511.03913>

3 **CONF****AI Mathematician: Towards Fully Automated Frontier Mathematical Research****ArXiv CS.AI** **50%**

arXiv:2505.22451v2 Announce Type: replace Abstract: Large Reasoning Models (LRMs) have made significant progress in mathematical capabilities in recent times. However, these successes have been primarily confined to competition-level problems. In this work, we propose AI...

<https://arxiv.org/abs/2505.22451>

4 **CONF****Stepwise Think-Critique: Interleaved Reasoning and Self-Critique in a Single LLM****ArXiv CS.AI** **50%**

arXiv:2512.15662v4 Announce Type: replace Abstract: Human beings solve complex problems through critical thinking, where reasoning and evaluation are intertwined to converge toward correct solutions. However, most existing large language models (LLMs) treat the reasoning and...

<https://arxiv.org/abs/2512.15662>

5 **CONF****Beyond Dialogue Time: Temporal Semantic Memory for Personalized LLM Agents****ArXiv CS.AI** **50%**

arXiv:2601.07468v2 Announce Type: replace Abstract: Memory enables Large Language Model (LLM) agents to perceive, store, and use information from past dialogues, which is essential for personalization. However, existing methods fail to properly model the temporal dimension of...

<https://arxiv.org/abs/2601.07468>

6 CONF

TikZilla: Scaling Text-to-TikZ with High-Quality Data and Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2603.03072v3 Announce Type: replace Abstract: Large language models (LLMs) are increasingly used to assist scientists across diverse workflows. A key challenge is generating high-quality figures from textual descriptions, often represented as TikZ programs that can be...

<https://arxiv.org/abs/2603.03072>

7 CONF

From High-Dimensional Spaces to Verifiable ODD Coverage for Safety-Critical AI-based Systems

ArXiv CS.AI 50%

arXiv:2604.02198v2 Announce Type: replace Abstract: While Artificial Intelligence (AI) offers transformative potential for operational performance, its deployment in safety-critical domains such as aviation requires strict adherence to rigorous certification standards. Current...

<https://arxiv.org/abs/2604.02198>

8 CONF

UniToolCall: Unifying Tool-Use Representation, Data, and Evaluation for LLM Agents

ArXiv CS.AI 50%

arXiv:2604.11557v3 Announce Type: replace Abstract: Tool-use capability is a fundamental component of LLM agents, enabling them to interact with external systems through structured function calls. However, existing research exhibits inconsistent interaction representations,...

<https://arxiv.org/abs/2604.11557>

9 CONF

Pailitao-MMSearch: Building Native E-Commerce Multimodal Search Foundation

ArXiv CS.AI 50%

arXiv:2607.17499v2 Announce Type: replace Abstract: The evolution of e-commerce has fundamentally transformed how users search for products, shifting from simple text-based keyword queries to complex multimodal interactions that seamlessly combine product images, natural...

<https://arxiv.org/abs/2607.17499>

00 CONF

Reinforcement Learning for Heterogeneous Sensor Selection in Maritime Surveillance

ArXiv CS.AI 50%

arXiv:2607.22667v2 Announce Type: replace Abstract: This paper presents an information-gain-guided reinforcement-learning sensor-selection framework for single-vessel tracking in heterogeneous maritime sensor networks. The proposed approach is motivated by information-theoretic...

<https://arxiv.org/abs/2607.22667>

01 CONF

Adaptive Graph-of-Islands Evolution for Automatic Feature Engineering with LLMs

ArXiv CS.AI 50%

arXiv:2607.23286v2 Announce Type: replace Abstract: Automatic feature engineering (AutoFE) for tabular data requires discovering informative transformations from a large program space. Existing approaches suffer from three limitations: classical methods rely on fixed operator...

<https://arxiv.org/abs/2607.23286>

02 CONF

FlavourBench: Executable Culinary Reward Maps for Language Model Evaluation and Post-Training

ArXiv CS.AI 50%

arXiv:2608.20574v3 Announce Type: replace Abstract: We introduce FlavorBench: a benchmark for Compiling Dense Deterministic Answer Maps from a Versioned Culinary Embeddings Model. We test 27 frontier large language model endpoints on 534 substitution, pairing and constraining...

<https://arxiv.org/abs/2608.20574>

03 CONF

Automated Researchers Can Mitigate Well-characterized Alignment Failures

ArXiv CS.AI 50%

arXiv:2608.28945v3 Announce Type: replace Abstract: Automating alignment research may accelerate progress toward aligned AI, but whether it does is hard to measure. Luckily, many alignment failures, such as deception, sycophancy, and jailbreaks, are already measurable by public...

<https://arxiv.org/abs/2608.28945>

04 CONF

A Survey of Transformer-based Language Models with Focus on Efficiency

ArXiv CS.AI 50%

arXiv:2406.16893v3 Announce Type: replace-cross Abstract: The emergence of Transformer-based Large Language Models (LLMs) has substantially augmented the capabilities of Natural Language Processing (NLP), thereby intensifying the demand for computational resources. Therefore,...

<https://arxiv.org/abs/2406.16893>

05 CONF

Evaluating the Evaluator: Summarization Metrics and LLM-Judges beyond English

ArXiv CS.AI 50%

arXiv:2503.17039v3 Announce Type: replace-cross Abstract: Automatic text summarization relies on automatic evaluation to quickly determine the quality of summarization models via automatic metrics and LLM-as-a-Judge models. However, these techniques require meta-evaluation to...

<https://arxiv.org/abs/2503.17039>

06 CONF

No Data Wasted: A Semi-supervised Generative Model for Incomplete Multi-view Data Integration with Missing Labels

ArXiv CS.AI 50%

arXiv:2508.11180v2 Announce Type: replace-cross Abstract: Multi-view learning is widely applied to real-life datasets, but it often suffers from both missing views and missing labels. Prior probabilistic approaches addressed the missing view problem by using a product-of-experts...

<https://arxiv.org/abs/2508.11180>

07 CONF

General Demographic Pre-trained Models for Enhancing Predictive Performance Across Diseases and Population

ArXiv CS.AI 50%

arXiv:2509.07330v3 Announce Type: replace-cross Abstract: Foundation models for healthcare require balancing robust generalization across heterogeneous clinical populations and disease settings with the architectural simplicity needed for deployment. We present a pre-trained...

<https://arxiv.org/abs/2509.07330>

08 CONF

ContextAnyone: Context-Aware Diffusion for Character-Consistent Text-to-Video Generation

ArXiv CS.AI 50%

arXiv:2512.07328v2 Announce Type: replace-cross Abstract: Text-to-video generation has advanced rapidly, yet preserving a character's holistic appearance from a single reference image remains challenging, particularly when the character undergoes large pose, motion, and scene...

<https://arxiv.org/abs/2512.07328>

09 CONF

Agent Tools Orchestration Leaks More: Dataset, Benchmark, and Mitigation

ArXiv CS.AI 50%

arXiv:2512.16310v4 Announce Type: replace-cross Abstract: LLM agents can combine individually non-revealing tool returns and disclose a sensitive conclusion, creating Tools Orchestration Privacy Risk (TOP-R). We formalize TOP-R through three conditions: conclusion sensitivity,...

<https://arxiv.org/abs/2512.16310>

10 CONF

FineVLA: Fine-Grained Instruction Alignment for Steerable Vision-Language-Action Policies

ArXiv CS.AI 50%

arXiv:2605.27284v3 Announce Type: replace-cross Abstract: Vision-Language-Action (VLA) models are increasingly expected to not only complete robot tasks, but also follow human instructions about how those tasks should be executed. However, existing robot datasets usually pair...

<https://arxiv.org/abs/2605.27284>

11 CONF

Training nGPT

ArXiv CS.AI 50%

arXiv:2608.01284v2 Announce Type: replace-cross Abstract: The normalized Transformer (nGPT) realizes hyperspherical representation learning by constraining model parameter vectors and activation vectors to the unit hypersphere. In this paper, we describe a practical training...

<https://arxiv.org/abs/2608.01284>

12 CONF

SHADOWBENCH: Toward Reliable Automatic Evaluation of Semantic Alignment in Autoformalization

ArXiv CS.AI 50%

arXiv:2608.29270v2 Announce Type: replace-cross Abstract: Autoformalization translates informal mathematical theorems into code for proof assistants such as Lean. A central challenge is that current evaluation metrics can accept type-correct but misaligned statements or reject...

<https://arxiv.org/abs/2608.29270>

13 CONF

Lot Machine: Multimodal Lot Extraction from Auction Catalogs

ArXiv CS.AI 50%

arXiv:2608.30510v2 Announce Type: replace-cross Abstract: For provenance research and art market studies, auction catalogs are an essential resource to trace specific objects over time and space. While historical auction catalogs follow established domain conventions, their...

<https://arxiv.org/abs/2608.30510>