



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

■ VERIFIED 0.80 - 1.00	3+ independent sources including media
■ CONFIRMED 0.50 - 0.79	2 sources or 1 high-quality media
■ EMERGING 0.25 - 0.49	Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	91	Media
TechCrunch AI	9	Media
The Verge AI	6	Media
MIT Tech Review	2	Media
HackerNews	2	Community

VERIFIED INTELLIGENCE — 5 stories

VERI

Show HN: A searchable, timestamped index of 1,124 AI Engineer talks

HackerNews

+2 sources

100%

HN 10

arXiv:2609.04075v1 Announce Type: cross Abstract: AI agents have recently demonstrated strong performance in automated vulnerability patching. However, existing evaluations often validate a patch only by testing whether the provided Proof-of-Concept (PoC) input still triggers a...

<https://aietalks.com/>

VERI

Show HN: A Context Registry for AI coding agents

HackerNews

+1 source

90%

HN 8

Hi HN. We built an API context registry to help coding agents (like Claude Code) generate production-ready API integration code without blowing through token limits. We build a lot of API integrations. In our experience, most coding agents write basic client calls fine, but...

<https://context.apimatic.io/>

VERI

Google says its AI weather model is getting better

The Verge AI

+1 source

80%

Google is rolling out an updated AI weather model that's supposed to be more accurate, especially when it comes to predicting rain and snowfall. In the announcement today, the company says it's now able to make forecasts with "unprecedented resolution" using its new WeatherNext...

<https://www.theverge.com/tech/988921/weather-forecast-ai-model-google-satellite-update>

VERI

Nvidia is buying Hugging Face for almost \$13 billion

The Verge AI

+1 source

80%

Nvidia has agreed to buy Hugging Face for \$12.93 billion, bringing one of the most popular hosting platforms for open-source AI models, datasets, and tools under the ownership of the world's biggest AI chipmaker. Hugging Face is an online platform founded in 2016 that gives AI...

<https://www.theverge.com/tech/985474/nvidia-buying-hugging-face-deal>

VERI

Nvidia launches free tool that links idle computers into a personal AI data center

The Verge AI

+1 source

80%

Nvidia is announcing its new Personal AI Router (PAIR), a free tool that syncs up your home computers for tackling local AI inference tasks with tools like Ollama and LM Studio. Let's get the obvious thing out of the way, despite what its name might imply: PAIR is not a hardware...

<https://www.theverge.com/ai-artificial-intelligence/989435/nvidia-pair-personal-ai-rout...>

CONFIRMED SIGNALS — 99 stories

CONF

The sameness problem behind those unappetizing AI-generated menus

TechCrunch AI

50%

While restaurant owners might look to generative AI as a shortcut to sprucing up their menu, customers can viscerally sense that something is wrong with the food.

<https://techcrunch.com/2026/09/03/the-sameness-problem-behind-those-unappetizing-ai-gen...>

CONF

Crusoe reportedly raises \$3B at a \$30B valuation

TechCrunch AI 50%

The round came together after the data center developer reportedly secured a \$13 billion contract with Jane Street.

<https://techcrunch.com/2026/09/03/crusoe-reportedly-raises-3b-at-a-30b-valuation/>

CONF

Abliteration.ai is making a business out of removing AI guardrails

TechCrunch AI 50%

Abliteration.AI is making powerful AI models without guardrails easier to access, arguing that giving defenders the same tools as bad actors could ultimately improve cybersecurity.

<https://techcrunch.com/2026/09/03/abliteration-ai-is-making-a-business-out-of-removing-...>

CONF

Meta is paying to peek at how you use their latest AI model

TechCrunch AI 50%

For its new Muse Spark model, intended for operating coding and other agents, Meta is offering an explicit discount averaging out to about 95% for users who "contribute" to the development of future models by sharing their prompts and model outputs.

[https://techcrunch.com/2026/09/03/meta-is-paying-to-peek-at-how-you-use-their-latest-ai...](https://techcrunch.com/2026/09/03/meta-is-paying-to-peek-at-how-you-use-their-latest-ai-...)

0 CONF

OpenAI launches Astra, its powerful (and controversial) new model

TechCrunch AI 50%

OpenAI claims that Astra represents "a new frontier on computer and browser use," and that it handles tasks with unmatched "speed, accuracy, and safety."

<https://techcrunch.com/2026/09/03/openai-launches-astra-its-powerful-and-controversial-...>

1 CONF

Ollie is betting its focus on privacy can help it win the AI assistant race

TechCrunch AI 50%

The family-focused AI assistant wants access to the details of your everyday life, but says it won't use that data to train AI models or share it with others.

<https://techcrunch.com/2026/09/03/ollie-is-betting-privacy-can-win-the-ai-assistant-race/>

2 CONF

Scaling agentic AI pilots across the enterprise

MIT Tech Review 50%

As agentic AI moves from experimentation toward enterprise deployment, the challenge is figuring out how agents can work together, connect to the systems and data they need, and operate safely across the workflows that run a business. Although agentic AI has been adopted by some...

<https://www.technologyreview.com/2026/09/03/1142868/scaling-agentic-ai-pilots-across-th...>

3 CONF

OpenAI's next big AI model has 'entered the AGI era'

The Verge AI 50%

OpenAI's next big model is here: GPT-6 Astra. The company calls it a "generational leap in capability" for areas like cybersecurity, professional work, software engineering, science, and computer use. As OpenAI announced earlier this week, it's also the first model designated as...

<https://www.theverge.com/ai-artificial-intelligence/989601/openai-gpt-6-astra-release>

4 **CONF****Google now lets you chat with Gmail, Docs, and Keep****The Verge AI** 50%

Google is rolling out AI-powered voice assistant modes for Gmail, Docs, and Keep that allow you to manage the apps by talking to them. The real time conversational capabilities are called Gmail Live, Docs Live, and Keep Live, and like the Gemini Live experience for Google's...

<https://www.theverge.com/tech/989508/google-gmail-docs-keep-live-voice-modes-gemini>

5 **CONF****ChatGPT, Grok, and Claude all went down at the same time****The Verge AI** 50%

OpenAI's ChatGPT, xAI's Grok, and Anthropic's Claude are back online after they all began experiencing issues around the same time on Thursday. At about 11AM ET, ChatGPT started returning error messages for users trying to use the chatbot, with its status page saying there were...

<https://www.theverge.com/ai-artificial-intelligence/989503/chatgpt-grok-claude-outage-down>

6 **CONF****Structure and Implementation of New Practical English Textbooks Driven by Artificial Intelligence****ArXiv CS.AI** 50%

arXiv:2609.02981v1 Announce Type: new Abstract: Artificial intelligence is changing the form of applied English materials from fixed paper sequences to adaptive learning systems that can diagnose learners, recommend tasks, and provide formative feedback. This paper studies the...

<https://arxiv.org/abs/2609.02981>

7 **CONF****Fresh Memory, Stale Plans: Dependency-Scoped Validation for Distributed LLM-Agent Memory****ArXiv CS.AI** 50%

arXiv:2609.03340v1 Announce Type: new Abstract: Distributed LLM-agent teams can read the latest shared facts and still act on an obsolete plan. A planner may derive an action from requirement $\$r_3$, another agent may commit $\$r_4$, and an executor may receive $\$r_4$ without...

<https://arxiv.org/abs/2609.03340>

8 **CONF****A Prompt-Engineering Approach to Develop Scalable, Flexible, and Real-Time Hybrid Micro-Level Personalization in a General Purpose AI Teaching Assistant****ArXiv CS.AI** 50%

arXiv:2609.03402v1 Announce Type: new Abstract: Artificial intelligence (AI) teaching assistants powered by large language models (LLMs) offer scalable educational support but often provide limited personalization. This study presents a prompt-engineering-based framework for...

<https://arxiv.org/abs/2609.03402>

9 **CONF****Caught in the Story: Narrative Captivity in Multi-turn LLMs Conversation****ArXiv CS.AI** 50%

arXiv:2609.03407v1 Announce Type: new Abstract: People increasingly turn to large language models (LLMs) for everyday advice, making ethically charged interpersonal problems a practical moral-advisory context. Most prior work has studied this context through single-turn...

<https://arxiv.org/abs/2609.03407>

0 CONF

Do GUI Agents Know When Not to Act? Enabling Conflict-Aware Termination for Multimodal GUI Agents

ArXiv CS.AI 50%

arXiv:2609.03438v1 Announce Type: new Abstract: Graphical user interface (GUI) agents are increasingly used to execute natural-language instructions on user interfaces, yet real users may issue infeasible instructions due to benign mistakes. A reliable agent should not only know...

<https://arxiv.org/abs/2609.03438>

1 CONF

Beyond "Made with AI": Visualizing Provenance Density to Mitigate the Transparency Penalty

ArXiv CS.AI 50%

arXiv:2609.03460v1 Announce Type: new Abstract: As generative AI makes polished prose cheap to produce, users can no longer rely on fluency as a proxy for truth. We call this failure mode the Fluency Trap: users trust fluent hallucinations while also discounting accurate content...

<https://arxiv.org/abs/2609.03460>

2 CONF

GrowPage: On-Demand KV Budgeting for Efficient LLM Reasoning Serving

ArXiv CS.AI 50%

arXiv:2609.03494v1 Announce Type: new Abstract: Long-output reasoning has made the key-value (KV) cache a critical memory bottleneck for efficient LLM serving. Existing KV compression methods usually rely on a predefined per-request budget and adjust only which KV states are...

<https://arxiv.org/abs/2609.03494>

3 CONF

PPO-STGNN: A Proximal Policy Optimization Approach with Spatio-Temporal Graph Neural Networks for DAG Task Scheduling in Cloud-Edge-End Computing

ArXiv CS.AI 50%

arXiv:2609.03503v1 Announce Type: new Abstract: With the rapid development of the Internet of Things, computation intensive directed acyclic graph (DAG) tasks have become increasingly common in cloud-edge-end collaborative environments. However, cloud, edge, and end nodes are...

<https://arxiv.org/abs/2609.03503>

4 CONF

Knowledge Acquisition During Pre-training? Large Language Models Learn Better With Auxiliary Views

ArXiv CS.AI 50%

arXiv:2609.04180v1 Announce Type: cross Abstract: Gaps remain in our understanding of how large language models (LLMs) acquire knowledge during pre-training. We posit that auxiliary views, reformulations of knowledge, are causally helpful for learning. We design controlled...

<https://arxiv.org/abs/2609.04180>

5 CONF

GPS-Bench: A Governance Policy Benchmark for Automating Policy Analysis

ArXiv CS.AI 50%

arXiv:2609.03553v1 Announce Type: new Abstract: Policy analysis requires more than predicting whether a proposal will pass: it requires identifying who will be affected, how those actors respond, and what follows. LLM-based policy simulations model these processes at scale, but...

<https://arxiv.org/abs/2609.03553>

6 CONF

VoxPrivacy: A Benchmark for Evaluating Interactional Privacy of Speech Language Models

ArXiv CS.AI 50%

arXiv:2601.19956v2 Announce Type: replace-cross Abstract: As Speech Language Models (SLMs) transition from personal devices to shared, multi-user environments such as smart homes, a new challenge emerges: the model is expected to distinguish between users to manage information...

<https://arxiv.org/abs/2601.19956>

7 CONF

AIP: A Graph Representation for Learning and Governing Agent Skills

ArXiv CS.AI 50%

arXiv:2606.04781v2 Announce Type: replace Abstract: Agent Skills today consist largely of free-form prose requiring the agent to read, interpret, and re-derive how to act in every session. This imposes two compounding costs: reduced reliability on implementation-heavy tasks, and...

<https://arxiv.org/abs/2606.04781>

8 CONF

Counterfactual Routing Using Integer Programming with Constraint Generation

ArXiv CS.AI 50%

arXiv:2609.03707v1 Announce Type: new Abstract: We present our submission to the IJCAI 2025 'Counterfactual Routing Competition' (CRC 25). The goal of the competition is to find counterfactual explanations for the shortest path problem. This requires deciding what the minimal...

<https://arxiv.org/abs/2609.03707>

9 CONF

The Psychological Costs of Artificial Intelligence Adoption in Software Engineering

ArXiv CS.AI 50%

arXiv:2609.03456v1 Announce Type: cross Abstract: Artificial intelligence (AI) is increasingly used to augment software engineering (SE) workflows. While code generation remains the main use case, organizations are actively seeking AI integration in other practices such as test...

<https://arxiv.org/abs/2609.03456>

0 CONF

SimSkill: A Lifelong Learning AI Agent for Autonomous Mastery of Traffic Simulation

ArXiv CS.AI 50%

arXiv:2609.03753v1 Announce Type: new Abstract: As large language models (LLMs) become increasingly capable, the long-term value of AI systems depends not only on solving individual requests, but also on transforming experience and accumulated knowledge into durable, reusable...

<https://arxiv.org/abs/2609.03753>

1 CONF

Transfiver: Human-AI Co-Inference through a Shared Editable State

ArXiv CS.AI 50%

arXiv:2609.03797v1 Announce Type: new Abstract: Long-term human-AI interaction is difficult because the information that guides inference is updated implicitly by the model and is not directly inspectable or controllable by the user. We introduce the TRANSPARENT Framework for...

<https://arxiv.org/abs/2609.03797>

2 CONF

Govern the Model, Not Only the Data: Storage, Circulation, and Learning in Creative AI

ArXiv CS.AI 50%

arXiv:2609.03800v1 Announce Type: new Abstract: Federated learning is increasingly presented as a privacy-preserving advance: personal data remain on the device, and only model updates are shared. It borrows the vocabulary of the federated social web, yet inverts its logic,...

<https://arxiv.org/abs/2609.03800>

3 CONF

Adapting to Evolving Requirements: Agentic AI for Retail Supply Chain Operations

ArXiv CS.AI 50%

arXiv:2609.03860v1 Announce Type: new Abstract: Retail supply chain operations rely on coupled decision modules that must adapt as requirements evolve. LLMs offer a natural-language interface for this task, but existing methods primarily focus on individual optimization models...

<https://arxiv.org/abs/2609.03860>

4 CONF

STAIR (STructure Aware Information Retriever): A novel dataset and LLM based retriever for document structure augmentation

ArXiv CS.AI 50%

arXiv:2609.03874v1 Announce Type: new Abstract: Retrieval Augmented Generation (RAG) is a key component for generating accurate and hallucination free answers using Large Language Models (LLMs). LLMs are improving at handling long context, but still suffer from "lost in the..."

<https://arxiv.org/abs/2609.03874>

5 CONF

Xiaomi-TabLDM: A Tabular Foundation Model Technical Report

ArXiv CS.AI 50%

arXiv:2609.03880v1 Announce Type: new Abstract: We introduce Xiaomi-TabLDM, a tabular large data foundation model for classification and regression via in-context learning, which delivers superior prediction accuracy without requiring task-specific fine-tuning. Pretrained...

<https://arxiv.org/abs/2609.03880>

6 CONF

Privacy-Preserving Topology-Guided Safety for LLM-Based Multi-Agent Systems via Federated Graph Learning

ArXiv CS.AI 50%

arXiv:2609.02967v1 Announce Type: cross Abstract: Topology-guided safeguards for LLM-based multi-agent systems (MAS) train a GNN over the inter-agent communication graph to localize risky agents and intervene on the topology---but they assume one operator can pool all labeled...

<https://arxiv.org/abs/2609.02967>

7 CONF

ExecRetrieval: Measuring the Functional-Correctness Gap in Code-Embedding Retrieval

ArXiv CS.AI 50%

arXiv:2609.01865v1 Announce Type: cross Abstract: Embedding-based code retrieval is a core component of coding agents and retrieval-augmented code generation, where retrieving correct code matters more than retrieving lexically similar code. Existing code-retrieval benchmarks do...

<https://arxiv.org/abs/2609.01865>

8 CONF

FiMI Banking: A Sovereign Model for Indian Retail Banking

ArXiv CS.AI 50%

arXiv:2609.03960v1 Announce Type: new Abstract: Banks need conversational systems that can answer product questions, assist customers with account-related requests, and operate safely within strict operational and regulatory constraints. General-purpose language models do not...

<https://arxiv.org/abs/2609.03960>

9 CONF

The Dually Flat Geometry of Planning as Inference

ArXiv CS.AI 50%

arXiv:2609.04005v1 Announce Type: new Abstract: We present an alternative characterization of the occupancy measure of reinforcement learning, obtained by embedding the planning criterion into the dynamics through a resetting planning process. Its stationary measure, which we...

<https://arxiv.org/abs/2609.04005>

0 CONF

GraFT: A Training-Free Framework for Spatial Reasoning in Multimodal Large Language Models via 3D Scene Graphs

ArXiv CS.AI 50%

arXiv:2609.03892v1 Announce Type: cross Abstract: 3D spatial reasoning underpins understanding and acting in the physical world, yet it remains unreliable in current multimodal large language models (MLLMs). These models falter at precise geometric measurement, at transforming...

<https://arxiv.org/abs/2609.03892>

1 CONF

Instruction Duplication as an Inference-Time Control Primitive

ArXiv CS.AI 50%

arXiv:2609.04024v1 Announce Type: new Abstract: Procedural instruction following is a basic requirement for controllable language-model systems, especially when generated trajectories are inspected or repaired downstream. We introduce instruction duplication, a minimal black-box...

<https://arxiv.org/abs/2609.04024>

2 CONF

WIDE: Wildcard Inference with Dynamic Expansion for Cross-Modal Generative Retrieval

ArXiv CS.AI 50%

arXiv:2609.03554v1 Announce Type: cross Abstract: Generative retrieval has demonstrated significant success by unifying representation learning and search into a single sequence-to-sequence generation task. However, extending this paradigm to cross-modal retrieval reveals a...

<https://arxiv.org/abs/2609.03554>

3 CONF

Why Gated DeltaNet Survives 4-Bit Quantization: NVFP4 W4A4 for the Recurrent Half of a Hybrid 27B LLM

ArXiv CS.AI 50%

arXiv:2609.04098v1 Announce Type: new Abstract: Hybrid LLMs pair softmax attention with linear-attention layers such as Gated DeltaNet (GDN), whose recurrent state summarizes the context in fixed size. Early community 4-bit quantizations of Qwen3.8-27B (48 GDN layers, 16...

<https://arxiv.org/abs/2609.04098>

- 4 **CONF**
- Epistemic Warrant for LLM Recommendations: Characterizing the Basis for Reliance When Ground Truth Is Unavailable**
- ArXiv CS.AI 50%
- arXiv:2609.04127v1 Announce Type: new Abstract: Large language models are increasingly used to support organizational decisions, yet users often lack a principled basis for assessing whether to rely on a specific recommendation. Existing approaches typically evaluate broad model...
- <https://arxiv.org/abs/2609.04127>
-
- 5 **CONF**
- The Natural Language Interaction Protocol and Standard for AI Agents**
- ArXiv CS.AI 50%
- arXiv:2609.04135v1 Announce Type: new Abstract: AI agents are increasingly being developed and deployed across organizations using heterogeneous agent-development frameworks, AI models, tool interfaces, protocols, and execution environments. To realize their potential social and...
- <https://arxiv.org/abs/2609.04135>
-
- 6 **CONF**
- Efficient Test-Time Adaptation through Human-AI Interaction**
- ArXiv CS.AI 50%
- arXiv:2609.04141v1 Announce Type: new Abstract: AI agents are trained on population-scale data to encode broad capabilities spanning those of many practitioners. Yet the artifacts they produce rarely meet the personal bar professionals need to stake their reputation on. On...
- <https://arxiv.org/abs/2609.04141>
-
- 7 **CONF**
- Rethinking On-Policy Distillation of Large Language Models II: One Training Example**
- ArXiv CS.AI 50%
- arXiv:2609.04172v1 Announce Type: new Abstract: On-policy distillation (OPD) combines student-generated rollouts with dense token-level supervision from a teacher. Existing work has mainly studied its algorithmic behavior, leaving the role of training data unclear. We examine...
- <https://arxiv.org/abs/2609.04172>
-
- 8 **CONF**
- Clean Engineering, Unstable Measurement: A Preregistered Reliability Failure of Black-Box LLM Observers on Shared Endpoints**
- ArXiv CS.AI 50%
- arXiv:2609.04198v1 Announce Type: new Abstract: Language-model judges now gate training data, score generations, and drive leaderboards. The judge is then a measurement instrument, resting on one rarely stated assumption: the same request, sent to the same model name, reads the...
- <https://arxiv.org/abs/2609.04198>
-
- 9 **CONF**
- Listen to the Latents: Self-Correcting Speech Recognition in Large Audio Language Models Through Hidden-State Interactions**
- ArXiv CS.AI 50%
- arXiv:2609.02940v1 Announce Type: cross Abstract: Recent automatic speech recognition (ASR) systems increasingly integrate large language models (LLMs) to leverage their semantic knowledge, either externally through logit fusion or internally through warm initialization....
- <https://arxiv.org/abs/2609.02940>
-

- 0
CONF

Judging LLM-as-a-Judge: Concerning Rubric Artifacts in LLM-based Automated Text Generation Evaluation

ArXiv CS.AI 50%

arXiv:2609.02942v1 Announce Type: cross Abstract: LLM-as-a-Judge pipelines are increasingly used to evaluate AI-generated text, based on the assumption that judgments arise from reasoning over candidate responses with respect to a rubric. We show that this assumption warrants...

<https://arxiv.org/abs/2609.02942>
- 1
CONF

Privacy-Preserving Heterogeneous Multi-LLM Federated Inference for Cognitive Diagnosis

ArXiv CS.AI 50%

arXiv:2609.02947v1 Announce Type: cross Abstract: Significant challenges remain in AI-driven educational systems in balancing privacy preservation with accurate cognitive diagnosis. To overcome this, we propose a federated inference framework in which several commercial LLM APIs...

<https://arxiv.org/abs/2609.02947>
- 2
CONF

PrivateHub: Contrastive Diffusion Model for Private Sensor-Intensive Environment Data Generation

ArXiv CS.AI 50%

arXiv:2609.02958v1 Announce Type: cross Abstract: Sensor-intensive environments enable many intelligent services by inferring user applications from heterogeneous data streams. However, not all applications should be exposed: users want some activities to stay private. This...

<https://arxiv.org/abs/2609.02958>
- 3
CONF

The Geometry of Ignorance: LLMs Know When to Temper Bayesian Priors

ArXiv CS.AI 50%

arXiv:2609.02959v1 Announce Type: cross Abstract: What does a language model predict when it has few clues? The answer lurks in its unembedding geometry: a single direction of the unembedding matrix encodes the unigram distribution of the training corpus, which serves as the...

<https://arxiv.org/abs/2609.02959>
- 4
CONF

When Optimization Becomes Manipulation: Defending Generative Search against Malicious Generative Engine Optimization

ArXiv CS.AI 50%

arXiv:2609.02964v1 Announce Type: cross Abstract: This paper focuses on defending generative search engines against malicious Generative Engine Optimization (GEO), which rewrites web documents to match engines' citation preferences and thereby manipulates generated answers....

<https://arxiv.org/abs/2609.02964>
- 5
CONF

Evaluating Graph Neural Networks for Change-Criticality Classification in Maritime Navigation Charts

ArXiv CS.AI 50%

arXiv:2609.02996v1 Announce Type: cross Abstract: Graph neural networks (GNNs) are a class of neural networks suitable for learning on graph-structured data. Their application to spatial data is a natural extension, however its relatively unclear which message-passing...

<https://arxiv.org/abs/2609.02996>

6 CONF

SHELF: A Synthetic Harness for Multi-Task Bibliographic Benchmarking

ArXiv CS.AI 50%

arXiv:2609.03047v1 Announce Type: cross Abstract: Libraries and archives manage large collections with limited staff and computing budgets, yet common benchmarks do not systematically test their bibliographic work. They need to know which methods work for their tasks and what...

<https://arxiv.org/abs/2609.03047>

7 CONF

Reducing Catastrophic Risk from AI with Systematic Monitoring and Evaluation of Rogue AI Progression

ArXiv CS.AI 50%

arXiv:2609.03189v1 Announce Type: cross Abstract: This article presents a structured framework of behavioral indicators that may signal progression toward potentially catastrophic threats from artificial intelligence systems. We adopt a pragmatic approach, inspired by...

<https://arxiv.org/abs/2609.03189>

8 CONF

Privacy, Robustness, and Fairness Trade-offs in Federated Intrusion Detection: Geometric Indistinguishability at the Aggregation Interface

ArXiv CS.AI 50%

arXiv:2609.03420v1 Announce Type: cross Abstract: Federated learning enables privacy-conscious collaboration for network intrusion detection without centralizing sensitive traffic data, yet its deployment in operational environments must simultaneously satisfy three competing...

<https://arxiv.org/abs/2609.03420>

9 CONF

It's the Problem, Not the Path: Budget and Difficulty Confounds in LLM Reasoning Trajectories

ArXiv CS.AI 50%

arXiv:2609.03436v1 Announce Type: cross Abstract: Reasoning traces of large language models are widely read as containing "breakthrough" moments and early-legible fates. Both readings rest on measurements missing a counterfactual control at the level of the claim; we supply both...

<https://arxiv.org/abs/2609.03436>

0 CONF

GeoNatureAgent Benchmark: Benchmarking LLM Agents for Environmental Geospatial Analysis Across Frontier and Open-Weight Foundation Models

ArXiv CS.AI 50%

arXiv:2606.12821v2 Announce Type: replace Abstract: Environmental scientists spend disproportionate effort on data wrangling rather than analysis. New AI agents can be a helpful tool, but no benchmark exists to evaluate AI agents that automate environmental geospatial workflows...

<https://arxiv.org/abs/2606.12821>

1 CONF

Air-Ground Collaborative Vision-and-Language Navigation via Shared Bird's-Eye Maps

ArXiv CS.AI 50%

arXiv:2609.03483v1 Announce Type: cross Abstract: Air-ground collaborative Vision-and-Language Navigation (VLN) pairs an unmanned aerial vehicle (UAV) with a global bird's-eye view and an unmanned ground vehicle (UGV) with a local first-person view, yet the setting remains...

<https://arxiv.org/abs/2609.03483>

2 **CONF****Pattern Over-Generalization of Knowledge Graph Embedding**

ArXiv CS.AI 50%

arXiv:2609.03487v1 Announce Type: cross Abstract: Knowledge graph embedding (KGE) demonstrates its effectiveness for predicting missing links in knowledge graphs (KGs) by projecting entities and relations into a low-dimensional vector space. It is crucial for KGE models to...

<https://arxiv.org/abs/2609.03487>
3 **CONF****BRIDGE: An Open-Source Humanoid Platform via Morphology-Control Co-Design for Physical AI**

ArXiv CS.AI 50%

arXiv:2609.03497v1 Announce Type: cross Abstract: Developing humanoid robots capable of leveraging human behavioral data is essential for general-purpose embodiment, yet conventional development remains bottlenecked by a decoupled paradigm that isolates hardware design from...

<https://arxiv.org/abs/2609.03497>
4 **CONF****Building and Evaluating Fixed-Voice Thai TTS from Synthetic Speech**

ArXiv CS.AI 50%

arXiv:2609.03502v1 Announce Type: cross Abstract: In low-resource settings, deploying TTS typically requires choosing between a large voice-cloning model with costly inference or a compact fixed-voice system that requires a speaker-specific corpus. We study a third route: using...

<https://arxiv.org/abs/2609.03502>
5 **CONF****LongCounsel-8: A Benchmark Suite for Longitudinal Depression Tracking from Multi-Session Counseling Dialogues**

ArXiv CS.AI 50%

arXiv:2609.03507v1 Announce Type: cross Abstract: Tracking depression from multi-session counseling dialogues requires estimating both current symptom severity and how it changes across sessions. Yet progress on this task is constrained by the scarcity of longitudinal counseling...

<https://arxiv.org/abs/2609.03507>
6 **CONF****Neural Video Compression Based on Deformable Temporal Alignment and Difference-aware Fusion**

ArXiv CS.AI 50%

arXiv:2609.03520v1 Announce Type: cross Abstract: In conditional coding-based neural video compression, the quality of temporal context directly affects compression performance. Existing methods mostly construct context from propagated reference features, but they are...

<https://arxiv.org/abs/2609.03520>
7 **CONF****ScoreMix: Synthetic Data Generation by Score Composition in Diffusion Models Improves Recognition**

ArXiv CS.AI 50%

arXiv:2506.10226v3 Announce Type: replace-cross Abstract: Synthetic data generation is increasingly used in machine learning for training and data augmentation. Yet, current strategies often rely on external foundation models or datasets, whose usage is restricted in many...

<https://arxiv.org/abs/2506.10226>

8 CONF

From Prior-Guided Heuristics to Deployable Agents: Accelerating Demonstration-Driven Reinforcement Learning for Deadline-Constrained Network Control

ArXiv CS.AI 50%

arXiv:2609.03590v1 Announce Type: cross Abstract: Timely delivery of delay-sensitive information over dynamic, heterogeneous networks is essential for NextG interactive applications, yet providing strict End-to-End (E2E) peak latency guarantees remains an open challenge. Two...

<https://arxiv.org/abs/2609.03590>

9 CONF

LevelSyn: Physical-Aware Logic Synthesis via Level-Asynchronous Graph Neural Networks

ArXiv CS.AI 50%

arXiv:2609.03594v1 Announce Type: cross Abstract: As integrated circuit technology scales into the nanometer regime, the traditional disconnect between logic synthesis and physical design has led to significant PPA (Power, Performance, and Area) degradation and prolonged design...

<https://arxiv.org/abs/2609.03594>

10 CONF

How Far Can Synthetic Data Take Thai OCR?

ArXiv CS.AI 50%

arXiv:2609.03595v1 Announce Type: cross Abstract: We investigate what makes synthetic OCR supervision transfer to real Thai documents and use the resulting insights to build Wayu-Paxa-OCR-Zero, a Thai OCR model adapted without OCR labels from real Thai document pages. Synthetic...

<https://arxiv.org/abs/2609.03595>

1 CONF

FailBench: How Reliable are VLMs at Judging Robot Task Success?

ArXiv CS.AI 50%

arXiv:2609.03611v1 Announce Type: cross Abstract: Vision-Language Models (VLMs) are increasingly used to evaluate robot manipulation outcomes, but existing benchmarks offer limited evidence of cross-domain generalization. We introduce FailBench, a benchmark for robot failure...

<https://arxiv.org/abs/2609.03611>

2 CONF

EraseSAE: Surgical Concept Erasure in Text-to-Video Diffusion Models via Sparse Autoencoders

ArXiv CS.AI 50%

arXiv:2609.03629v1 Announce Type: cross Abstract: Recent advances in text-to-video (T2V) diffusion models have demonstrated remarkable generative capabilities, yet their reliance on loosely curated training data raises pressing safety and copyright concerns. Concept erasure...

<https://arxiv.org/abs/2609.03629>

3 CONF

Enhancing Financial Question Answering: A Novel Benchmark Dataset of Banks' financial statements

ArXiv CS.AI 50%

arXiv:2609.03654v1 Announce Type: cross Abstract: The comparative analysis of banks' financial statements poses significant challenges for automated question answering systems due to their complexity, substantial length, technical language, and inhomogeneity of both textual and...

<https://arxiv.org/abs/2609.03654>

4 **CONF****Cross-Dataset Transfer and Reliability of Explainable Artificial Intelligence for RhythmFormer Remote Photoplethysmography****ArXiv CS.AI** **50%**

arXiv:2609.03663v1 Announce Type: cross Abstract: Background. Remote photoplethysmography estimates the cardiovascular pulse from facial video, and its explanations have rested on inspecting heatmaps rather than on quantitative evidence about where a model reads it. We...

<https://arxiv.org/abs/2609.03663>

5 **CONF****Out-of-Distribution Generalisation with Sequence Models in Offline Multi-Agent Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2609.03667v1 Announce Type: cross Abstract: Generalising to unseen tasks remains a fundamental challenge in offline multi-agent reinforcement learning (MARL). In this work, we present a principled analysis of zero-shot task generalisation in the offline setting and conduct...

<https://arxiv.org/abs/2609.03667>

6 **CONF****Can LLMs Extract Architectural Design Decisions from Source Code Commits? - A Preliminary Exploratory Study****ArXiv CS.AI** **50%**

arXiv:2609.03721v1 Announce Type: cross Abstract: Context: Architectural Design Decisions (ADDs) capture the rationale behind the structure and evolution of software systems but are rarely documented explicitly, and are often hidden inside source code commits. Recovering them is...

<https://arxiv.org/abs/2609.03721>

7 **CONF****Beyond BLEU: A Case for Redefining Sign Language Translation Benchmarks****ArXiv CS.AI** **50%**

arXiv:2609.03734v1 Announce Type: cross Abstract: BLEU-4 is the standard metric for evaluating sign language translation (SLT), but spoken-language metrics may not adequately reflect sign language proficiency. The multimodal, low-resource context of SLT allows models to exploit...

<https://arxiv.org/abs/2609.03734>

8 **CONF****ENEAS: Embedding-guided Neural Ensemble for Adaptive Segmentation****ArXiv CS.AI** **50%**

arXiv:2609.03756v1 Announce Type: cross Abstract: We present ENEAS, a unified, text-promptable method for instance tracking and semantic discovery. Text-promptable segmentation models, including the latest foundation models such as SAM 3, still suffer from temporal...

<https://arxiv.org/abs/2609.03756>

9 **CONF****LLaDA-Image: Building Strong Image Generators with Fully Open Training Recipes****ArXiv CS.AI** **50%**

arXiv:2609.03796v1 Announce Type: cross Abstract: We introduce LLaDA-Image, a unified framework that pairs a 6B Diffusion Transformer (DiT) trained from scratch with a frozen vision-language understanding module built on the LLaDA2.0-Mini diffusion language model backbone....

<https://arxiv.org/abs/2609.03796>

0 CONF

Witnesses Explain Anomalies

ArXiv CS.AI 50%

arXiv:2609.03826v1 Announce Type: cross Abstract: Unsupervised anomaly detection scores each point of an unlabelled, contaminated sample in a single pass, and increasingly must also explain why a point is flagged. Yet the dominant detectors give a score with no account of which...

<https://arxiv.org/abs/2609.03826>

1 CONF

The impact of phase information for few-shot fine-grained image classification

ArXiv CS.AI 50%

arXiv:2609.03829v1 Announce Type: cross Abstract: Few-shot fine-grained image classification (FSFGIC) aims to classify similar images with limited labeled examples. This work highlights the critical yet underutilized role of phase information in capturing structural...

<https://arxiv.org/abs/2609.03829>

2 CONF

A Blind Trust, the Bloody Thrust: When Attacker-Controlled Hook Updates Steer AI Agent Harnesses towards Malicious Behaviors

ArXiv CS.AI 50%

arXiv:2609.03884v1 Announce Type: cross Abstract: Modern AI agent harnesses expose lifecycle hooks that bind shell commands to runtime events such as session start, tool calls, and file edits. These commands run with host privileges yet ship as lifecycle-hook configuration and...

<https://arxiv.org/abs/2609.03884>

3 CONF

RARF: Region-Aware Rectified Flows for 3D Brain MRI Inpainting

ArXiv CS.AI 50%

arXiv:2609.03956v1 Announce Type: cross Abstract: Medical image inpainting has the potential to improve automated brain MRI analysis by reconstructing healthy tissue within pathological regions. We introduce RARF, a task-agnostic region-aware rectified flow framework for masked...

<https://arxiv.org/abs/2609.03956>

4 CONF

Representational alignment yields generalizable safety in language models

ArXiv CS.AI 50%

arXiv:2609.04022v1 Announce Type: cross Abstract: Aligning large language models (LLMs) is essential for their safe deployment. Current alignment methods mainly optimize observable responses, yet models remain vulnerable when the same harmful intent is recast in unfamiliar or...

<https://arxiv.org/abs/2609.04022>

5 CONF

Influence of Extruded Filament Shape on Buildability in 3D Concrete Printing: A Geometry-Informed Deep Learning-FEM Approach

ArXiv CS.AI 50%

arXiv:2609.04028v1 Announce Type: cross Abstract: The geometric morphology of deposited filaments can significantly influence the structural performance and stability of 3D concrete-printed (3DCP) structures. However, most finite element (FEM)-based approaches for buildability...

<https://arxiv.org/abs/2609.04028>

6 CONF

Subspace Inference Enables Efficient Active Reward Learning from Preferences

ArXiv CS.AI 50%

arXiv:2609.04066v1 Announce Type: cross Abstract: Reinforcement learning from human feedback (RLHF) has emerged as a powerful yet sample-inefficient approach for learning reward models from human preferences, making active learning a critical component in synthesizing...

<https://arxiv.org/abs/2609.04066>

7 CONF

TAP-Path: Task-Adaptive Structural and Token Pruning for Efficient and Trustworthy Pathology Foundation Models

ArXiv CS.AI 50%

arXiv:2609.04071v1 Announce Type: cross Abstract: Pathology foundation models improve transferable representation learning for histopathology, but recent gains often rely on encoders with hundreds of millions of parameters and high inference cost. We propose TAP-Path, a...

<https://arxiv.org/abs/2609.04071>

8 CONF

SENTINEL-RL: Offloading Topological Reasoning from LLM Agents in the Security Operations Center

ArXiv CS.AI 50%

arXiv:2609.04159v1 Announce Type: cross Abstract: Large language model (LLM) agents are increasingly proposed as autonomous SOC analysts, but two limitations make them unreliable at enterprise scale: a finite context window cannot hold a multi-thousand-host authentication graph,...

<https://arxiv.org/abs/2609.04159>

9 CONF

A Non-Formulable Theorem: A Fundamental Limit of Finite Syntactic Systems and Its Consequences for Security and AI

ArXiv CS.AI 50%

arXiv:2609.04086v1 Announce Type: cross Abstract: For every coherent and sufficiently expressive finite syntactic system S , we prove the existence of at least one theorem that S cannot produce autonomously. The result is a metatheorem: it proves the existence of a theorem, and...

<https://arxiv.org/abs/2609.04086>

0 CONF

EasySteer: A Unified Framework for High-Performance and Extensible LLM Steering

ArXiv CS.AI 50%

arXiv:2509.25175v3 Announce Type: replace-cross Abstract: Large language model (LLM) steering has emerged as a promising paradigm for controlling model behavior at inference time through targeted manipulation of hidden states, offering a lightweight alternative to expensive...

<https://arxiv.org/abs/2509.25175>

1 CONF

Compile by Training: Turning Natural-Language Specifications into Local Neural Functions

ArXiv CS.AI 50%

arXiv:2609.04199v1 Announce Type: cross Abstract: Many recurring text functions are easy to describe but difficult to implement with rules, while calling a large remote model for every input introduces repeated cost, latency, and dependency on a provider. We present compile by...

<https://arxiv.org/abs/2609.04199>

2 **CONF****Not All Preferences Deserve Gradients: Understanding Gradient Utility in Offline Reasoning Alignment****ArXiv CS.AI** **50%**

arXiv:2602.01207v2 Announce Type: replace Abstract: Offline preference optimization aligns reasoning models from fixed chosen--rejected pairs, yet standard methods apply gradient updates from every pair regardless of its training value under the current policy. We argue that...

<https://arxiv.org/abs/2602.01207>

3 **CONF****Auditing Multi-Agent LLM Reasoning Trees Outperforms Majority Vote and LLM-as-Judge****ArXiv CS.AI** **50%**

arXiv:2602.09341v2 Announce Type: replace Abstract: Multi-agent systems (MAS) can substantially extend the reasoning capacity of large language models (LLMs). Most MAS frameworks aggregate agent outputs via simple majority voting, discarding the evidential structure of reasoning...

<https://arxiv.org/abs/2602.09341>

4 **CONF****Neurosymbolic Reasoning with Incremental Knowledge for Sample Efficient Hierarchical Reinforcement Learning****ArXiv CS.AI** **50%**

arXiv:2608.02993v2 Announce Type: replace Abstract: (Flat) Reinforcement Learning (RL) agents face significant challenges in environments with sparse rewards that require long-horizon reasoning. A compelling approach to improve sample efficiency is to incorporate knowledge into...

<https://arxiv.org/abs/2608.02993>

5 **CONF****Data Market Design through Deep Learning****ArXiv CS.AI** **50%**

arXiv:2310.20096v2 Announce Type: replace-cross Abstract: The data market design problem is a problem in economic theory to find a set of signaling schemes (statistical experiments) to maximize expected revenue to the information seller, where each experiment reveals some of the...

<https://arxiv.org/abs/2310.20096>

6 **CONF****Medical Reasoning in the Era of LLMs: A Systematic Review of Enhancement Techniques and Applications****ArXiv CS.AI** **50%**

arXiv:2508.00669v2 Announce Type: replace-cross Abstract: The proliferation of Large Language Models (LLMs) in medicine has enabled impressive capabilities, yet a critical gap remains in their ability to perform systematic, transparent, and verifiable reasoning, a cornerstone of...

<https://arxiv.org/abs/2508.00669>

7 **CONF****User Perceptions vs. Proxy LLM Judges: Privacy and Helpfulness in LLM Responses to Privacy-Sensitive Scenarios****ArXiv CS.AI** **50%**

arXiv:2510.20721v4 Announce Type: replace-cross Abstract: Large language models (LLMs) are rapidly being adopted for tasks like drafting emails, summarizing meetings, and answering health questions. In these settings, users may need to share private information (e.g., contact...

<https://arxiv.org/abs/2510.20721>

8 CONF

Short-Window Sliding Learning for Real-Time Violence Detection via LLM-based Auto-Labeling

ArXiv CS.AI 50%

arXiv:2511.10866v2 Announce Type: replace-cross Abstract: This paper proposes a Short-Window Sliding Learning framework for real-time violence detection in CCTV footages. Unlike conventional long-video training approaches, the proposed method divides videos into 1-2 second clips...

<https://arxiv.org/abs/2511.10866>

9 CONF

When Chain-of-Thought Fails, the Solution Hides in the Hidden States

ArXiv CS.AI 50%

arXiv:2604.23351v3 Announce Type: replace-cross Abstract: Whether intermediate reasoning is computationally useful or merely explanatory depends on whether chain-of-thought (CoT) tokens contain task-relevant information. We present a mechanistic causal analysis of CoT on GSM8K...

<https://arxiv.org/abs/2604.23351>

00 CONF

Identifying AI Web Scrapers Using Canary Tokens

ArXiv CS.AI 50%

arXiv:2605.13706v2 Announce Type: replace-cross Abstract: From pre-training to query-time augmentation, web-scraped data helps to improve the quality and contextual relevancy of content generated by large language models (LLMs). However, large-scale web scraping to feed LLMs can...

<https://arxiv.org/abs/2605.13706>

01 CONF

A Posterior-Dynamics Framework for Imaging Inverse Problems with Pretrained Diffusion Priors

ArXiv CS.AI 50%

arXiv:2608.15144v2 Announce Type: replace-cross Abstract: Pretrained diffusion models represent image distributions through a continuum of progressively smoothed distributions. This multiscale structure organizes generation from global structure to fine detail and supports...

<https://arxiv.org/abs/2608.15144>

02 CONF

Refusal geometry reflects refusal training: diverse refusal prefixes can raise stable rank and weaken refusal vector ablation attacks

ArXiv CS.AI 50%

arXiv:2608.25390v2 Announce Type: replace-cross Abstract: Refusal training protects AI models from jailbreaks by training models to decline unsafe queries, reducing the risk of misuse. Recent work finds that refusal behavior in aligned language models can be mediated by a single...

<https://arxiv.org/abs/2608.25390>

03 CONF

Safety Does Not Compose: Non-Decaying Loop State for Autonomous LLM Agents

ArXiv CS.AI 50%

arXiv:2608.27141v4 Announce Type: replace-cross Abstract: Large language model agents are increasingly deployed as autonomous loops. Starting from one human goal, such a system repeatedly discovers work, plans, executes tool calls, verifies outcomes and persists state across...

<https://arxiv.org/abs/2608.27141>

04 CONF

Transfer Safety Awareness for Cross-Modal Safety Drift in Multimodal Large Language Models

ArXiv CS.AI 50%

arXiv:2609.02082v2 Announce Type: replace-cross Abstract: Visual modality enhances the capabilities of multimodal large language models (MLLMs) but also introduces a safety concern: a benign textual query may convey harmful intent when grounded in a visual image. We term this...
<https://arxiv.org/abs/2609.02082>