



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	120	Media
HackerNews	4	Community
TechCrunch AI	2	Media
The Verge AI	1	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Show HN: Wayfinder - A reference implementation for evaluating AI applications

HackerNews

+1 source

90%

HN 3

arXiv:2609.05117v1 Announce Type: cross Abstract: Current LLM safety benchmarks largely rely on binary metrics, overlooking how models respond to harmful prompts with varying threat implicitness. We introduce TIER, a Threat Implicitness Benchmark for behavioral safety evaluation...

<https://github.com/DivakarUngatla/wayfinder>

VERI

Seattle Times and Newsday sue OpenAI and Microsoft for infringement

The Verge AI

+1 source

80%

The Seattle Times and Newsday are just the latest plaintiffs to take OpenAI to court, alleging copyright infringement. The two outlets say the company used their journalism as training data for its AI models without permission and often reproduces passages from their reporting...

<https://www.theverge.com/ai-artificial-intelligence/990932/seattle-times-newsday-lawsui...>

CONFIRMED SIGNALS — 123 stories

CONF

Show HN: Engrim - A universal, local-first SQLite memory engine for AI CLIs

HackerNews

50%

HN 28

Hey all, TLDR: When you sign up you get an email alias @audioreality.ai that you can use to sign up for newsletters, have agents send output to, etc. Once a day you get a private podcast with a multi-host format/conversation, so it's not a basic summary and it's not a verbatim...

<https://github.com/timgordontg/engrim>

CONF

Show HN: I stopped using an LLM gateway and put rate-limits/fallback in-process

HackerNews

50%

HN 4

Sup HN. I built VernLLM cuz every LLM gateway I looked at, meant adding a network hop just to get rate limiting, multi provider fallback, circuit breaking and other features. Needing to take a whole separate service just to deploy, monitor and trust with my API keys is just...

<https://github.com/LakBud/vernLLM>

CONF

Show HN: Server hardening playbook where every item is failure/fix/verify

HackerNews

50%

HN 3

No summary — see link for full article.

<https://github.com/Sanexxxx777/server-hardening-playbook>

CONF

Authors push back as publishers and agents make claims on Anthropic settlement

TechCrunch AI

50%

Authors say publishers seem to be claiming more than their fair share of settlement payments.

<https://techcrunch.com/2026/09/06/authors-push-back-as-publishers-and-agents-seek-share...>

CONF

Large Language Models for HVAC Operations in Building Energy Systems: A Critical Review of Methods, Applications, and Deployment Readiness

ArXiv CS.AI 50%

arXiv:2609.05314v1 Announce Type: new Abstract: Building automation systems generate rich sensor data yet remain insight-poor because heterogeneous point naming, missing metadata, and fragmented documentation obstruct their operational use. This systematic review analyses and...

<https://arxiv.org/abs/2609.05314>

CONF

Data-Optimized Contingency Screening: A Machine Learning Approach to Power System Security

ArXiv CS.AI 50%

arXiv:2609.04300v1 Announce Type: new Abstract: Ensuring the security of the power system is essential for stability and reliability, especially in the event of disruption. Effective classification of contingency in power systems enables proactive decision-making and mitigates...

<https://arxiv.org/abs/2609.04300>

CONF

A Removal Based Approach to Improve LLM Faithfulness at Test-Time

ArXiv CS.AI 50%

arXiv:2609.04343v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used for consequential decisions, making their explanations an important tool for auditing model behavior. Unfortunately, these explanations can be unfaithful, failing to reflect the...

<https://arxiv.org/abs/2609.04343>

0 CONF

Why Better Models Can Create Riskier Systems: Evidence from LLM Agents in Financial Markets

ArXiv CS.AI 50%

arXiv:2609.04373v1 Announce Type: new Abstract: Large language models (LLMs) are being deployed at scale in consequential real-world systems, from financial markets to content moderation to hiring. We show that improving individual model capability can degrade rather than...

<https://arxiv.org/abs/2609.04373>

1 CONF

Corporate Language Model (CLM): Transforming Tacit and Fragmented Enterprise Knowledge into a Sovereign, Auditable, and Executable Corporate Intelligence Layer

ArXiv CS.AI 50%

arXiv:2609.04377v1 Announce Type: new Abstract: Enterprise AI deployments fail not from model inadequacy, but because organizations lack a structured substrate encoding how they decide, negotiate, and execute. Generic LLMs carry no firm-specific ontological priors; RAG remains...

<https://arxiv.org/abs/2609.04377>

2 CONF

ERPBench: Evaluating LLM Agents for Enterprise Decision-Making Across Competitive Market Ecologies

ArXiv CS.AI 50%

arXiv:2609.04667v1 Announce Type: new Abstract: Large language model (LLM) agents are increasingly proposed for enterprise workflows, yet existing evaluations rarely test whether business-decision conclusions transfer across competitive market ecologies. We introduce ERPBench,...

<https://arxiv.org/abs/2609.04667>

3 CONF

PerfReasoning: How Well Do LLMs Reason on Hardware Performance?

ArXiv CS.AI 50%

arXiv:2609.04476v1 Announce Type: new Abstract: Performance modeling is central to hardware design and software optimization, yet constructing these models requires structured reasoning about computation, data reuse, storage, and movement. We introduce PerfReasoning, a benchmark...

<https://arxiv.org/abs/2609.04476>

4 CONF

When Quantization Breaks Memory: Recurrent-State Write-Back in Low-Precision Temporal Inference

ArXiv CS.AI 50%

arXiv:2609.04490v1 Announce Type: new Abstract: Quantization is widely used to reduce the computational and memory demands of neural-network inference. In recurrent networks, however, the quantized state is stored and returned at the next time step, so the rule used to store...

<https://arxiv.org/abs/2609.04490>

5 CONF

BioSync: Transformer-Based Cross-Modal Fusion for a Multimodal Physiological Digital Biomarker

ArXiv CS.AI 50%

arXiv:2609.04504v1 Announce Type: new Abstract: Cardiac, neural, behavioral, and speech measurements from wearable and mobile devices provide partial, noise-sensitive views of physiological state. BioSync combines these measurements into the \textbf{BioSync Index (BSI)}, a...

<https://arxiv.org/abs/2609.04504>

6 CONF

From Answers to Interpretations: Rethinking Ambiguity-Induced Aleatoric Uncertainty Estimation in LLMs

ArXiv CS.AI 50%

arXiv:2609.04543v1 Announce Type: new Abstract: A key challenge in reliable LLM deployment is recognizing when uncertainty reflects irreducible variability in the task rather than limitations in the model's knowledge. In language tasks, a central source of such aleatoric...

<https://arxiv.org/abs/2609.04543>

7 CONF

IPGeoAI: Transformer-Based Geolocation with LLM Semantic Fusion

ArXiv CS.AI 50%

arXiv:2609.04559v1 Announce Type: new Abstract: Accurate city-level IP Geolocation is an important enabler for the modern digital ecosystem, underpinning services ranging from local content delivery and targeting to digital rights enforcement. However, traditional heuristic and...

<https://arxiv.org/abs/2609.04559>

8 CONF

SiLR: Structure-Preserving Admission and Process Reward for LLM Tool Agents

ArXiv CS.AI 50%

arXiv:2609.04629v1 Announce Type: new Abstract: A runtime gate for an LLM tool agent is usually cast as a filter. In a ReAct loop a rejected proposal is followed by another at the same state, so the gate is a search operator over the proposal stream whose admission criterion...

<https://arxiv.org/abs/2609.04629>

9 CONF

A Cost-Aware Agentic Architecture for NL-to-SQL over Nested Enterprise Schemas, with a New Benchmark

ArXiv CS.AI 50%

arXiv:2609.04641v1 Announce Type: new Abstract: Natural-language-to-SQL systems have advanced rapidly on academic benchmarks, yet production enterprise schemas exhibit graph-like, semi-structured, deeply nested structure that current benchmarks do not measure. We make two...

<https://arxiv.org/abs/2609.04641>

0 CONF

Evidence Integration in Large Language Models

ArXiv CS.AI 50%

arXiv:2609.04290v1 Announce Type: cross Abstract: Despite increasing reliance on LLMs that reason with external evidence supplied by tools, retrieval-augmented generation, other agents, and users, how LLMs integrate such evidence into decisions they have already begun to form...

<https://arxiv.org/abs/2609.04290>

1 CONF

Train What You Deploy: Token-Faithful Post-Training of a Production Coding

ArXiv CS.AI 50%

arXiv:2609.04678v1 Announce Type: new Abstract: Existing post-training pipelines for coding and terminal agents suffer severe token and control fidelity errors: simplified training environments mismatch production deployments, and offline token reconstruction from agent logs...

<https://arxiv.org/abs/2609.04678>

2 CONF

Predicting Spatiotemporal Mobile Sensing-Based PM2.5 Concentrations Using Low-Rank Adapted Spatially Attentive Graph Neural Network

ArXiv CS.AI 50%

arXiv:2609.04693v1 Announce Type: new Abstract: Urban air quality can vary significantly along transit corridors, necessitating high-resolution monitoring. This work introduces a novel mobile-sensing dataset from Surat, Gujarat, India, comprising PM_{2.5} concentrations,...

<https://arxiv.org/abs/2609.04693>

3 CONF

Model Retirement Creates Reproducibility Risk in Biomedical AI Publications

ArXiv CS.AI 50%

arXiv:2609.04699v1 Announce Type: new Abstract: Background. Large language models (LLMs) are being adopted in biomedical research at a rapid and accelerating pace, yet commercial services that host many widely used models operate under deprecation schedules that can complicate...

<https://arxiv.org/abs/2609.04699>

4 CONF

FinalityBench: An Effect-Level Benchmark for Agent Decisions Under Delayed and Conflicting Financial Finality

ArXiv CS.AI 50%

arXiv:2609.04706v1 Announce Type: new Abstract: A merchant's payment processor, ledger, ERP and bank feed are updated by messages that get delayed, duplicated, dropped and reordered, so for minutes at a time the four hold contradictory beliefs about the same order. An agent...

<https://arxiv.org/abs/2609.04706>

5 CONF

PLUME: Parameter-Efficient Personalization of Large Language Models via Low-Rank User Modulation in Shared Subspaces

ArXiv CS.AI 50%

arXiv:2609.04715v1 Announce Type: new Abstract: Personalizing large language models (LLMs) is essential for delivering AI assistance that aligns with individual users' styles, intents, and preferences. While per-user fine-tuning can substantially enhance personalization quality,...

<https://arxiv.org/abs/2609.04715>

6 CONF

DCFA: Dual-view Causal-inspired Attribution for Failure Reasoning in LLM-based Multi-agent Systems

ArXiv CS.AI 50%

arXiv:2609.04749v1 Announce Type: new Abstract: Large language model (LLM)-based multi-agent systems have experienced rapid growth in recent years. Despite their promise, such systems remain fragile, frequently exhibiting reasoning and coordination errors that can lead to...

<https://arxiv.org/abs/2609.04749>

7 CONF

Shadow Queries for Private Retrieval in Vector Databases

ArXiv CS.AI 50%

arXiv:2609.04767v1 Announce Type: new Abstract: Large language models (LLMs) increasingly rely on information retrieval (IR) systems, such as Retrieval-Augmented Generation (RAG), to incorporate domain-specific knowledge without costly re-training. These systems often store...

<https://arxiv.org/abs/2609.04767>

8 CONF

From Language Models to World-Acting Systems: Progress and Limits of Agentic AI across Digital, Social, Virtual, and Physical Environments

ArXiv CS.AI 50%

arXiv:2609.04894v1 Announce Type: new Abstract: Large language models become consequential agents when surrounding systems let outputs change external state. Models now call tools, operate interfaces, delegate work, retain state, inhabit generated worlds, and control robots or...

<https://arxiv.org/abs/2609.04894>

9 CONF

Whose record is this? Diagnosing and authorizing record use in personalized multimodal models

ArXiv CS.AI 50%

arXiv:2609.04801v1 Announce Type: new Abstract: Contextualized visual personalization can retrieve a true record yet apply it to the wrong visual subject. We formalize when a record may condition an answer as `\emph{record authorization}`: subject presence (\$P\$), record-edge...

<https://arxiv.org/abs/2609.04801>

0 CONF

Molecular D'ej`a Vu: Digit-Level Retrieval of Published Values in Frontier Language Models

ArXiv CS.AI 50%

arXiv:2609.05381v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly evaluated on molecular property benchmarks, but accuracy cannot distinguish a model that predicts a property from one that retrieves a published number. We audit 22 frontier models on...

<https://arxiv.org/abs/2609.05381>

1 CONF

You Really Didn't Get That? Benchmarking Social Pragmatic Inference for Indirect and Playful Chinese Online Comments

ArXiv CS.AI 50%

arXiv:2609.04384v1 Announce Type: cross Abstract: Chinese online comments often convey social meaning through indirect and playful language that is hard to interpret without context. Existing evaluations largely organize items around predefined phenomena or controlled pragmatic...

<https://arxiv.org/abs/2609.04384>

2 CONF

MM-IFEval-Pro: A Multilingual and Attack-Resistant Benchmark for Instruction-Following in Vision-Language Models

ArXiv CS.AI 50%

arXiv:2609.04859v1 Announce Type: new Abstract: As vision-language models (VLMs) rapidly advance in image understanding, cross-modal reasoning, and complex instruction execution, instruction-following capability has become a key indicator of their reliability and practicality....

<https://arxiv.org/abs/2609.04859>

3 CONF

MZ-Rain: Moisture-Budget-Guided Zero-Inflated Model for Station-Level Precipitation Nowcasting

ArXiv CS.AI 50%

arXiv:2609.04864v1 Announce Type: new Abstract: Accurate station-level precipitation nowcasting is critical for agriculture, water resource management, and disaster prevention, which typically is formulated as a time series forecasting problem. However, conventional time-series...

<https://arxiv.org/abs/2609.04864>

4 CONF

CoSkill: Joint Reinforcement Learning of Reasoning and Meta-Skill Agents for Hierarchical Skill Evolution

ArXiv CS.AI 50%

arXiv:2609.04865v1 Announce Type: new Abstract: Skill libraries improve the sample efficiency of agentic reinforcement learning (RL) by enabling large language model (LLM) agents to reuse procedural knowledge. Yet existing paradigms exhibit structural shortcomings: they either...

<https://arxiv.org/abs/2609.04865>

5 CONF

Role-Aware Artificial Intelligence Across Augmentation and Automation in Human-Machine Symbiosis

ArXiv CS.AI 50%

arXiv:2605.00440v2 Announce Type: replace Abstract: The evolution of artificial intelligence (AI) has rendered the boundary between humanity and computational machinery increasingly ambiguous. In the presence of more interwoven relationships within human-machine symbiosis, the...

<https://arxiv.org/abs/2605.00440>

6 CONF

CHAMP: Cross-domain Hybrid Architecture for Matchmaking and Prediction in Online Multi-Player Games

ArXiv CS.AI 50%

arXiv:2609.04870v1 Announce Type: new Abstract: Multiplayer Online Battle Arena (MOBA) games rely on matchmaking to maintain competitive balance. Our prior work, CUPID, framed matchmaking as an assignment re-optimization problem and showed that a single-mode win-rate predictor...

<https://arxiv.org/abs/2609.04870>

7 CONF

MARLA: A Conceptual Scaffold for Regulatory Learning under the EU AI Act

ArXiv CS.AI 50%

arXiv:2609.04877v1 Announce Type: new Abstract: The EU AI Act positions regulation as part of the infrastructure for safe, trustworthy and market-ready innovation. Realising this ambition requires regulatory learning: the evidence generated during implementation must be...

<https://arxiv.org/abs/2609.04877>

8 CONF

Reinforcement Learning for Sequential Solar PV Policy Design under Uncertainty: An Agent-Based Approach

ArXiv CS.AI 50%

arXiv:2609.04880v1 Announce Type: new Abstract: Designing effective and fiscally sustainable policies for solar photovoltaic (PV) adoption requires balancing adoption gains against public expenditure under uncertainty and heterogeneous decision-making. This study formulates PV...

<https://arxiv.org/abs/2609.04880>

9 CONF

Compact Bellman-Grounded Cognitive Maps for Cost-Aware Navigation

ArXiv CS.AI 50%

arXiv:2609.05104v1 Announce Type: new Abstract: Biological agents navigate familiar environments not by re-solving routes for each new goal, but by reusing a learned map built once and read off as goals change. Existing artificial cognitive-map models mimic this reuse, yet their...

<https://arxiv.org/abs/2609.05104>

0 CONF

Artificial Intelligence in Equity and Crypto Markets: Progress, Profitability Evidence, and the Limits of Automated Investing

ArXiv CS.AI 50%

arXiv:2609.04917v1 Announce Type: new Abstract: Artificial intelligence (AI) now supports investment workflows from data and prediction through research, portfolios, execution, and tool use. Technical capability, however, is not evidence of investment profitability. This...

<https://arxiv.org/abs/2609.04917>

1 CONF

Solving Hard XAI Queries Based on a Compiled Dual-Rail Encoding

ArXiv CS.AI 50%

arXiv:2609.04931v1 Announce Type: new Abstract: The widespread adoption of artificial intelligence (AI) within real-world applications has raised a lot of concerns regarding their trustworthiness, especially in critical applications. The field of eXplainable AI (XAI) has emerged...

<https://arxiv.org/abs/2609.04931>

2 CONF

A Tree-based RAG Framework for Evidence-Intensive QA via Adaptive Planning and Topology-Aware Evidence Gathering

ArXiv CS.AI 50%

arXiv:2609.04981v1 Announce Type: new Abstract: Recent structured RAG methods leverage tree- or graph-based reasoning structures to improve multi-hop QA. However, they face key limitations in evidence-intensive QA, where answering a question requires synthesizing information...

<https://arxiv.org/abs/2609.04981>

3 **CONF****Language models judge war differently when tested for alignment****ArXiv CS.AI 50%**

arXiv:2609.05009v1 Announce Type: new Abstract: Safety evaluations can mischaracterize deployed behaviour if artificial-intelligence systems respond to being evaluated. We test this possibility in a full-factorial conjoint experiment on decisions to start a war, spanning 20...

<https://arxiv.org/abs/2609.05009>

4 **CONF****Moral Competence Before Moral Content: Why LLM Agents Lack the Prerequisites for Coherent Alignment****ArXiv CS.AI 50%**

arXiv:2609.05036v1 Announce Type: new Abstract: AI alignment requires AI systems to adhere to human norms, values, or intentions. Under value pluralism there is no correct target, but a shared prerequisite is that the system's behavior expresses a coherent policy: a mapping from...

<https://arxiv.org/abs/2609.05036>

5 **CONF****TruthInsightBench: An Evidence-Grounded Benchmark for Automated Evaluation of Open-Ended Scientific Discovery Agents****ArXiv CS.AI 50%**

arXiv:2609.05079v1 Announce Type: new Abstract: Autonomous coding agents are increasingly proposed as AI-scientist systems that conduct analyses and write research reports, but executing a prescribed analysis is not the same as making a discovery. Existing benchmarks are...

<https://arxiv.org/abs/2609.05079>

6 **CONF****Measuring AI Accountability Through Argumentation Analysis: Can Model Reasoning Withstand Scrutiny?****ArXiv CS.AI 50%**

arXiv:2609.05088v1 Announce Type: new Abstract: AI oversight methods rely on ground truth for validation, but what constitutes appropriate AI behavior is contested. This leaves evaluation of moral reasoning in LLMs and debate-based oversight implicitly avoiding realistic...

<https://arxiv.org/abs/2609.05088>

7 **CONF****LLM-Guided Program Evolution for Circle Packing: Breaking 10 Packomania Records for \$28****ArXiv CS.AI 50%**

arXiv:2609.05093v1 Announce Type: new Abstract: We present Discovery Loop, a lightweight system that uses a large language model (LLM) to iteratively evolve optimization algorithms. Starting from a simple seed solver, the LLM proposes algorithmic improvements guided by a...

<https://arxiv.org/abs/2609.05093>

8 **CONF****VICAL: Vicinal Consistency Alignment for Long-Tailed Visual Recognition****ArXiv CS.AI 50%**

arXiv:2609.04948v1 Announce Type: cross Abstract: Multi-expert models have become the dominant paradigm for long-tailed learning, largely attributed to their presumed ability to benefit from expert diversity. However, we revisit this central assumption and reveal that diversity...

<https://arxiv.org/abs/2609.04948>

- 9
CONF

A Hybrid Predictive Ensemble of Machine Learning and Deep Neural Networks for Early Cardiovascular Disease Risk Assessment

ArXiv CS.AI 50%

arXiv:2609.05146v1 Announce Type: new Abstract: This study introduces an intelligent framework that integrates machine learning and deep neural network ensemble techniques for early detection and prognosis of cardiovascular diseases. The system utilizes real-time physiological...

<https://arxiv.org/abs/2609.05146>
- 0
CONF

ACE: Adaptive Calibration-Free Expert Skipping for MoE-based LLMs

ArXiv CS.AI 50%

arXiv:2609.05228v1 Announce Type: new Abstract: Mixture-of-Experts (MoE) architectures provide an efficient paradigm for scaling large language models (LLMs), yet fixed top-k routing activates the same number of expert slots for every token, causing substantial redundant...

<https://arxiv.org/abs/2609.05228>
- 1
CONF

Do LLMs Exhibit Coherent Knowledge Structures in Mathematical Reasoning? A Perspective from Knowledge Space Theory

ArXiv CS.AI 50%

arXiv:2609.05245v1 Announce Type: new Abstract: Human knowledge is inherently structured and interdependent: mastery of a concept requires prior mastery of its prerequisites, a principle formalized by Knowledge Space Theory (KST). While LLMs achieve strong performance on complex...

<https://arxiv.org/abs/2609.05245>
- 2
CONF

A Unified Physics-Aware Quantum Machine Learning Framework across Power GaN HEMTs and Logic Nanowire FETs: Predicting Unseen Process Splits and Held-Out Geometry Combinations with Lower Error and Tighter Split-to-Split Variability

ArXiv CS.AI 50%

arXiv:2609.05251v1 Announce Type: new Abstract: We present a unified reinforcement-learning (RL) framework that discovers compact parametrized quantum circuits (PQCs) for data-scarce device modeling. A graph neural network (GNN) policy optimized by proximal policy optimization...

<https://arxiv.org/abs/2609.05251>
- 3
CONF

Commonsense Reasoning in Computer Vision: Foundations, Recent Advancements, and Future Directions

ArXiv CS.AI 50%

arXiv:2609.05257v1 Announce Type: new Abstract: Commonsense reasoning in computer vision encompasses integrating visual data and contextual knowledge, crucial for enhancing AI's understanding of everyday scenarios. This understanding not only improves machine learning models but...

<https://arxiv.org/abs/2609.05257>
- 4
CONF

Trace2Tower: Transition-Aware EigenTrace Induction of Multi-Level Skills for LLM Agents

ArXiv CS.AI 50%

arXiv:2609.05261v1 Announce Type: new Abstract: Large language model agents increasingly rely on execution traces to master complex interactive tasks. However, current paradigms are bottlenecked by shallow trajectory retrieval and flat skill summarization, fundamentally ignoring...

<https://arxiv.org/abs/2609.05261>

5 CONF

AI for Computational Design Science: A Responsible Human-AI Framework and Case Study on Short-Form Video Safety Surveillance

ArXiv CS.AI 50%

arXiv:2609.05270v1 Announce Type: new Abstract: Artificial intelligence (AI) is transforming not only what information systems researchers design, but also how design research is conducted. Yet existing literature offers limited guidance for computational design science (CDS)...

<https://arxiv.org/abs/2609.05270>

6 CONF

Don't Drop Dropout: Optimizing Layer Sparsity for Efficient LLM Training and Inference

ArXiv CS.AI 50%

arXiv:2609.05275v1 Announce Type: new Abstract: Layer dropout (a.k.a. stochastic depth) has been shown to enable faster training, higher accuracy, and robustness to zero-shot layer pruning in both language and vision transformers. However, as models and datasets have scaled,...

<https://arxiv.org/abs/2609.05275>

7 CONF

GUT: Quantifying and Optimizing the Reasoning Uncertainty of LLMs via Graph Complexity

ArXiv CS.AI 50%

arXiv:2609.05284v1 Announce Type: new Abstract: Recent years have witnessed great advances in the reasoning ability of Large Language Models (LLMs). However, the reasoning processes of LLMs often exhibit uncertainty, where LLMs often produce a proliferation of divergent branches...

<https://arxiv.org/abs/2609.05284>

8 CONF

LLM-Driven Algorithm Design for Quantum Circuit Synthesis based on Binary Decision Diagrams

ArXiv CS.AI 50%

arXiv:2609.05327v1 Announce Type: new Abstract: Quantum circuits are central to implementing quantum algorithms on quantum devices, where quantum gates must be reversible. Many quantum algorithms rely on Boolean functions, which must therefore be implemented reversibly within...

<https://arxiv.org/abs/2609.05327>

9 CONF

Technical Manual for a Toolkit for Measuring Contextual Individuation in Transformer Language Models

ArXiv CS.AI 50%

arXiv:2609.05333v1 Announce Type: new Abstract: A transformer language model assigns a single, context-independent vector to a word type at its embedding layer, yet is widely believed to individuate that word's occurrences by context in its later layers. Testing this belief...

<https://arxiv.org/abs/2609.05333>

10 CONF

Who Should Grade My Work? Student Perspectives on Transparent AI-Assisted Writing Assessment in Higher Education

ArXiv CS.AI 50%

arXiv:2609.05346v1 Announce Type: new Abstract: The integration of GenAI tools into higher education assessment raises important questions about how students understand, interpret, and respond to AI-mediated evaluation. As instructors increasingly explore AI tools for providing...

<https://arxiv.org/abs/2609.05346>

1 CONF

Necessary or Sufficient? Evaluating LLM Explanations With Behavioural Evidence

ArXiv CS.AI 50%

arXiv:2609.05385v1 Announce Type: new Abstract: LLM decision components that can operate within agent workflows often produce action-relevant recommendations or judgements together with explanations. Operators may use the named factors to monitor a system, diagnose errors, or...

<https://arxiv.org/abs/2609.05385>

2 CONF

Multi-Step Tool-Calling over Korean Open Public APIs: A Benchmark and a Data-Synthesis Recipe

ArXiv CS.AI 50%

arXiv:2609.05395v1 Announce Type: new Abstract: Data-sovereignty regulations increasingly require public institutions to deploy open-source, on-premise LLM agents that chain multiple tool-calls across live government APIs. However, open-source models consistently underperform in...

<https://arxiv.org/abs/2609.05395>

3 CONF

A Deep Generative Model for Synthesizing Labeled Wireless Signals

ArXiv CS.AI 50%

arXiv:2609.05396v1 Announce Type: new Abstract: Wireless signals with position-related labels are pivotal for both performance evaluation and model training in the realm of wireless sensing. However, acquiring real-world datasets is often challenged by significant measurement...

<https://arxiv.org/abs/2609.05396>

4 CONF

AlcaTRAz - Anchored Tree-Rule Defense Against Jailbreaks

ArXiv CS.AI 50%

arXiv:2609.03693v1 Announce Type: cross Abstract: Large language models (LLMs) are vulnerable to jailbreak attacks that bypass safety alignment through carefully crafted prompts. Many existing defenses require access to model weights or internals, making them difficult to apply...

<https://arxiv.org/abs/2609.03693>

5 CONF

Scalable Context Orchestration for Serving LLMs Over Voice

ArXiv CS.AI 50%

arXiv:2609.04288v1 Announce Type: cross Abstract: Voice AI applications are gaining popularity as advances in large language models (LLMs) enable more natural and accessible spoken interactions. Serving these applications requires accounting not only for what users say, but also...

<https://arxiv.org/abs/2609.04288>

6 CONF

Blockchain-Enabled Secure Logging for Fiscal Electronic Mechanisms: Evaluation of the Greek eSEND and myDATA Tax Systems

ArXiv CS.AI 50%

arXiv:2609.04356v1 Announce Type: cross Abstract: This paper analyzes the implementation of blockchain-based integrity mechanisms in Greek Fiscal Electronic Mechanisms (FEMs) and the central tax information system eSEND. The study examines the cryptographic architecture of...

<https://arxiv.org/abs/2609.04356>

7 **CONF****Cross-modal triage network: a multimodal deep learning framework for severity-based triage and visual explainability in chest radiographs****ArXiv CS.AI 50%**

arXiv:2609.04357v1 Announce Type: cross Abstract: Purpose: Increased number of chest radiograph (CXR) scans create a triage bottleneck, queueing urgent examinations behind routine ones. Existing AI tools are predominantly unimodal binary classifiers lacking severity awareness,...

<https://arxiv.org/abs/2609.04357>

8 **CONF****Ultrasound-Based Prediction of Cirrhosis Decompensation Using Large-Scale Computer Vision Models****ArXiv CS.AI 50%**

arXiv:2609.04365v1 Announce Type: cross Abstract: Decompensation represents a critical transition in the course of cirrhosis, yet clinicians have limited non-invasive tools to reliably predict its onset. In this study, we propose a novel imaging-based approach that leverages...

<https://arxiv.org/abs/2609.04365>

9 **CONF****Where Appearance Fails, Geometry Recognizes: A CAD-Free 3D Shape Prior That Complements Vision Foundation Models****ArXiv CS.AI 50%**

arXiv:2609.04381v1 Announce Type: cross Abstract: Recognizing specific objects onboarded without a labeled training set recurs across manufacturing and service robotics, yet the conventional renderable prior, a computer-aided-design (CAD) model, is often unavailable....

<https://arxiv.org/abs/2609.04381>

0 **CONF****REFINE: LLM Refinement over Budgeted Text-Attributed Graphs for Personalized Medical Concept Representation****ArXiv CS.AI 50%**

arXiv:2609.04415v1 Announce Type: cross Abstract: Learning rich medical concept representations is essential for EHR prediction. Text-attributed knowledge graphs (TKGs) provide a natural foundation by organizing heterogeneous medical relations together with textual semantics....

<https://arxiv.org/abs/2609.04415>

1 **CONF****GRACE: Graph-Grounded Reflective Agent Copilot Engine for Expert-in-the-Loop Knowledge Expansion****ArXiv CS.AI 50%**

arXiv:2609.04442v1 Announce Type: cross Abstract: Large language models deployed in high-stakes settings frequently generate plausible but ungrounded claims. Standard retrieval-augmented generation (RAG) pipelines offer limited remedy, since they retrieve isolated passages...

<https://arxiv.org/abs/2609.04442>

2 **CONF****A Roadmap for MEG Foundation Models****ArXiv CS.AI 50%**

arXiv:2609.04461v1 Announce Type: cross Abstract: Foundation models are beginning to reshape brain-signal analysis by moving the field beyond task-specific decoding pipelines toward reusable models pretrained on broad neural datasets. Magnetoencephalography (MEG) is a compelling...

<https://arxiv.org/abs/2609.04461>

3 **CONF****Shared circuits predict whether LLMs generalize across formats in arithmetic reasoning****ArXiv CS.AI 50%**

arXiv:2609.04463v1 Announce Type: cross Abstract: In many forms of reasoning, including arithmetic reasoning, generalizing across superficial changes in input format is effortless for humans: anyone who can solve $2+5$ can also solve 'two plus five'. In contrast, LLMs are more...

<https://arxiv.org/abs/2609.04463>

4 **CONF****Exploring Solution Divergence and Its Effect on Large Language Model Problem Solving****ArXiv CS.AI 50%**

arXiv:2509.22480v2 Announce Type: replace-cross Abstract: Large language models (LLMs) have been widely used for problem-solving tasks. Most recent work improves their performance through supervised fine-tuning (SFT) with labeled data or reinforcement learning (RL) from task...

<https://arxiv.org/abs/2509.22480>

5 **CONF****Towards Understanding Pause Token Fine-Tuning Dynamics: A Mode Retention Perspective****ArXiv CS.AI 50%**

arXiv:2609.04489v1 Announce Type: cross Abstract: Pause-token methods improve LLM reasoning by inserting special tokens into sequences. Prior work explains these gains through computational expressivity. However, there is relatively little investigation into the training...

<https://arxiv.org/abs/2609.04489>

6 **CONF****Atlas: Optimizing Deployment of Compound AI Workflows on Heterogeneous Clusters****ArXiv CS.AI 50%**

arXiv:2609.04513v1 Announce Type: cross Abstract: Compound AI workflows are increasingly used to serve complex AI tasks by coordinating multiple AI models and software components. This approach enables deployment flexibility, as each workflow stage can expose different model...

<https://arxiv.org/abs/2609.04513>

7 **CONF****Pitch-class Steering for Diffusion-based Music Generation via Latent-space Probes****ArXiv CS.AI 50%**

arXiv:2609.04516v1 Announce Type: cross Abstract: Recent work on controllable music generation has focused on autoregressive models, leaving diffusion-based systems comparatively underexplored. We present a lightweight method for steering the pitch content of audio produced by...

<https://arxiv.org/abs/2609.04516>

8 **CONF****Continual Field-Adaptive Models (CFAMs) for Post-Deployment Physical AI****ArXiv CS.AI 50%**

arXiv:2609.04552v1 Announce Type: cross Abstract: Unattended interactive autonomy - machines that step into danger in place of humans and complete tasks with human tools - remains a missing capability in mission-critical operations. These domains offer scarce training data and...

<https://arxiv.org/abs/2609.04552>

9 CONF

Dynamic Adaptation of the LLM Context for Generating Routines with Coupled Semantics

ArXiv CS.AI 50%

arXiv:2609.04570v1 Announce Type: cross Abstract: LLM-based code generation fails when correctness depends on execution-dependent coupling: the meaning of one routine is defined by the runtime behavior of another, a relationship that cannot be resolved from textual descriptions...

<https://arxiv.org/abs/2609.04570>

0 CONF

Training-Free Halving of Activated Experts in Fine-Grained Mixture-of-Experts Models

ArXiv CS.AI 50%

arXiv:2609.04575v1 Announce Type: cross Abstract: Modern fine-grained Mixture-of-Experts (MoE) models route each token to a small number of experts and renormalize their router probabilities. We show that this renormalization implicitly calibrates expert output gain to the...

<https://arxiv.org/abs/2609.04575>

1 CONF

When Do Internal Probes Beat Reading the Answer? Miscalibrated Readouts and Behavior-Concealed Knowledge in Language Models

ArXiv CS.AI 50%

arXiv:2609.04582v1 Announce Type: cross Abstract: A 0.6B language model, asked to verify 1,200 logical conclusions (half valid, half corrupted by a single semantic edit), answers YES every time. Judged by behavior it discriminates nothing; linear probes on its hidden states read...

<https://arxiv.org/abs/2609.04582>

2 CONF

PetQA: Benchmarking Veterinary Knowledge and Clinical Reasoning

ArXiv CS.AI 50%

arXiv:2609.04598v1 Announce Type: cross Abstract: We introduce PetQA, a Korean long-form question-answering (QA) benchmark for evaluating veterinary knowledge and clinical reasoning in large language models (LLMs) and large vision-language models (LVLMs). PetQA contains 10,076...

<https://arxiv.org/abs/2609.04598>

3 CONF

Tracing Audio Grounding and Answer Selection in Audio LLMs

ArXiv CS.AI 50%

arXiv:2609.04637v1 Announce Type: cross Abstract: Audio Large Language Models (Audio LLMs) have advanced in audio understanding, yet they can still predict the answer by reasoning from textual cues or linguistic priors rather than the provided audio. A common remedy is to train...

<https://arxiv.org/abs/2609.04637>

4 CONF

Wireless Foundation Models: State-of-the-Art and Open Challenges

ArXiv CS.AI 50%

arXiv:2609.04707v1 Announce Type: cross Abstract: Wireless foundation models (WFMs) have emerged as a promising approach for learning reusable representations from large-scale wireless data and adapting them to downstream tasks. However, the rapidly growing literature remains...

<https://arxiv.org/abs/2609.04707>

5 **CONF****Refuse without Refusal: A Structural Analysis of Safety-Tuning Responses for Reducing False Refusals in Language Models****ArXiv CS.AI** **50%**

arXiv:2609.04714v1 Announce Type: cross Abstract: Striking a balance between helpfulness and safety remains a fundamental challenge in aligning large language models. To achieve this balance, models should refuse harmful queries (e.g., "How do I shoot someone?") while remaining...

<https://arxiv.org/abs/2609.04714>

6 **CONF****When Does an Interpretation Count as Established? The Formation, Evaluation, and Responsibility of Interpretation in Generative AI****ArXiv CS.AI** **50%**

arXiv:2609.04766v1 Announce Type: cross Abstract: Generative AI research has increasingly evaluated factuality, citation, coverage, and report structure. Yet passing such local checks does not by itself show that a humanistic interpretation has been established. This paper asks...

<https://arxiv.org/abs/2609.04766>

7 **CONF****LLM4CKD: Large Language Models for Early Stage Chronic Kidney Disease Screening****ArXiv CS.AI** **50%**

arXiv:2609.04013v2 Announce Type: replace Abstract: Early screening of chronic kidney disease (CKD) is critical for timely intervention, yet most machine learning (ML) and deep learning (DL) approaches require labeled data and model training, limiting their use in real-world...

<https://arxiv.org/abs/2609.04013>

8 **CONF****Mitigating Performance Discrepancy in Cross-Domain 3D Class-Incremental Learning****ArXiv CS.AI** **50%**

arXiv:2609.04860v1 Announce Type: cross Abstract: 3D perception plays a crucial role in real-world applications such as autonomous driving, robotics, and AR/VR. In practical scenarios, 3D perception models need to continually adapt to newly emerging 3D object categories, making...

<https://arxiv.org/abs/2609.04860>

9 **CONF****PRISM-Bench: An Audio-Centric Diagnostic Benchmark for Text-to-Audio-Video Generation****ArXiv CS.AI** **50%**

arXiv:2609.04867v1 Announce Type: cross Abstract: Text-to-audio-video (T2AV) generation has advanced rapidly, but its evaluation still underestimates the audio modality. Existing benchmarks either treat audio as an auxiliary component of video quality or assess it in isolation...

<https://arxiv.org/abs/2609.04867>

0 **CONF****Forgetting Without Restarting: Execution-State Unlearning for Stateful LLM Agents****ArXiv CS.AI** **50%**

arXiv:2609.04875v1 Announce Type: cross Abstract: Long-running LLM agents are stateful: beyond the transcript they accrete compressed summaries, plaintext memory, pending tool plans, and, under every serving API, a KV cache. Yet today's "forget" operations delete a plaintext...

<https://arxiv.org/abs/2609.04875>

1 CONF

ReCAST: Restoration-aware Cascaded Stage-wise Training for Obfuscated SMS Risk Classification

ArXiv CS.AI 50%

arXiv:2609.04878v1 Announce Type: cross Abstract: Fraudulent messages sent via Short Message Service (SMS) are increasingly obfuscated to evade cost-conscious classifiers in production systems. In Chinese SMS, attackers can exploit a wide range of carefully crafted obfuscation...

<https://arxiv.org/abs/2609.04878>

2 CONF

Adaptation Interfaces for In-Context Tabular Foundation Models in Time-to-Event Prediction

ArXiv CS.AI 50%

arXiv:2609.04901v1 Announce Type: cross Abstract: Tabular foundation models (TabFMs) achieve strong performance on structured data, particularly for standard classification and regression problems. Yet, extending them to censored time-to-event prediction is challenging because...

<https://arxiv.org/abs/2609.04901>

3 CONF

Better Understanding, Better Fixes? A Study of Hallucination in LLM-based Automated Program Repair

ArXiv CS.AI 50%

arXiv:2609.04909v1 Announce Type: cross Abstract: Large language models (LLMs) have significantly advanced automated program repair (APR), yet existing evaluations remain largely result-centric and provide limited insight into hallucination during repair. In APR, hallucination...

<https://arxiv.org/abs/2609.04909>

4 CONF

TreeFI: Value-Aware Statistical Fault Injection for Deep Neural Networks

ArXiv CS.AI 50%

arXiv:2609.04912v1 Announce Type: cross Abstract: Reliability evaluation of deep neural networks under hardware faults commonly relies on fault injection, but exhaustive campaigns are intractable for modern models and datasets. Statistical fault injection reduces this cost, yet...

<https://arxiv.org/abs/2609.04912>

5 CONF

Active Inference for an Intelligent Agent in Autonomous Reconnaissance Missions

ArXiv CS.AI 50%

arXiv:2510.17450v2 Announce Type: replace Abstract: We develop an active inference route-planning method for the autonomous control of intelligent agents. The aim is to reconnoiter a geographical area to maintain a common operational picture. To achieve this, we construct an...

<https://arxiv.org/abs/2510.17450>

6 CONF

One Diffusion Model, Two Roles: Guided Trajectory Planning and Safety-Critical Scenario Generation in Closed-Loop Simulation

ArXiv CS.AI 50%

arXiv:2609.04921v1 Announce Type: cross Abstract: Diffusion probabilistic models can capture the multi-modal, interaction-rich distribution of joint future trajectories in driving scenes. We show that a single pretrained diffusion traffic model can serve two complementary roles...

<https://arxiv.org/abs/2609.04921>

7 CONF

Post Fusion Bird's Eye View Feature Stabilization for Robust Multimodal 3D Detection

ArXiv CS.AI 50%

arXiv:2603.05623v2 Announce Type: replace-cross Abstract: Camera-LiDAR fusion is widely used in autonomous driving to enable accurate 3D object detection. However, bird's-eye view (BEV) fusion detectors can degrade significantly under domain shift and sensor failures, limiting...

<https://arxiv.org/abs/2603.05623>

8 CONF

How a Chatbot's Response Style Shapes a Classroom: A Multi-Agent Simulation of Students Consulting AI

ArXiv CS.AI 50%

arXiv:2609.05018v1 Announce Type: cross Abstract: LLM-based chatbots are increasingly used as everyday confidants. Because they are designed to maximize user satisfaction, they can respond with excessive empathy and affirmation, which may reinforce mistaken beliefs and foster...

<https://arxiv.org/abs/2609.05018>

9 CONF

Leveraging Low-Level Symbolic Competences for Unsupervised Grounding in Hallucination Detection

ArXiv CS.AI 50%

arXiv:2609.05025v1 Announce Type: cross Abstract: Hallucination-where a language model generates outputs that are factually incorrect or unsupported by the source-is a major challenge for both prompted and fine-tuned language models. Detecting hallucinations is difficult due to...

<https://arxiv.org/abs/2609.05025>

00 CONF

How do LLMs Evaluate Perceived Moral Agency? Investigating Moral Decision-Making in Human-Artificial Agents Interactions

ArXiv CS.AI 50%

arXiv:2609.05037v1 Announce Type: cross Abstract: As LLMs take on roles requiring moral advice, understanding how they attribute moral agency becomes critical. Humans possess moral agency, the capacity to make ethically guided decisions and bear responsibility for their...

<https://arxiv.org/abs/2609.05037>

01 CONF

Qlippy: A Retrieval-Augmented GenAI Assistant for Reproducible Quantum Workflows and Experiment Tracking

ArXiv CS.AI 50%

arXiv:2609.05039v1 Announce Type: cross Abstract: Quantum software development is iterative and error-prone. Noisy hardware and repeated re-execution make experiment tracking, provenance, and reproducibility essential, yet these practices are hard to adopt because of tooling...

<https://arxiv.org/abs/2609.05039>

02 CONF

NEAT-POCKET: Pocket-Conditioned Autoregressive 3D Molecular Generation with a Neighborhood-Guided Set Transformer

ArXiv CS.AI 50%

arXiv:2609.05097v1 Announce Type: cross Abstract: AI-driven de novo molecular design offers a promising route to accelerate early-stage drug discovery by generating novel ligands directly within target protein binding pockets. We present NEAT-POCKET, a pocket-conditioned...

<https://arxiv.org/abs/2609.05097>

03 CONF

A Schema Bounded Language Model for Refining Robot Policies Without Destabilizing Local Learning

ArXiv CS.AI 50%

arXiv:2609.05133v1 Announce Type: cross Abstract: This paper addresses navigation by composite heterogeneous robots in a decentralized system when policy reasoning and local control operate at different update levels. In a NetLogo--Python implementation, three robots share...

<https://arxiv.org/abs/2609.05133>

04 CONF

AxQM: A Textbook-Scale Benchmark for Formal Proof Synthesis in a Library of Finite-Dimensional Quantum Mechanics

ArXiv CS.AI 50%

arXiv:2609.05157v1 Announce Type: cross Abstract: Formalizing mathematics in a proof assistant, where a machine checks every definition, statement and proof, has set a new standard of rigor. Large language models are now capable of formalizing autonomously, even at the scale of...

<https://arxiv.org/abs/2609.05157>

05 CONF

Phase Transition Frequency as a Training Time Predictor of Test Accuracy in ResNets

ArXiv CS.AI 50%

arXiv:2609.05194v1 Announce Type: cross Abstract: The number of discrete class-separability jumps observed during ResNet finetuning is examined empirically as a predictor of final test accuracy. Across 75 experiments spanning four benchmarks (CIFAR-10, CIFAR-100, TinyImageNet,...

<https://arxiv.org/abs/2609.05194>

06 CONF

A Verifier-Guided Explainable Reasoning Framework with Gold-Anchored QLoRA, Task-Aware Mixture-of-Experts, and Group-Relative RLVR

ArXiv CS.AI 50%

arXiv:2609.05221v1 Announce Type: cross Abstract: Large language models (LLMs) show strong reasoning ability, but their explanations can remain inconsistent, weakly grounded, or difficult to verify. We propose a verifier-guided explainable reasoning framework for transparent...

<https://arxiv.org/abs/2609.05221>

07 CONF

PRICE: A Systematic Study of LLM Adaptation Choices for Bitcoin Price Forecasting

ArXiv CS.AI 50%

arXiv:2609.05235v1 Announce Type: cross Abstract: Cryptocurrency markets exhibit extreme volatility and non-stationary dynamics that challenge conventional forecasting methods. Although Large Language Models (LLMs) have shown promise for time series forecasting, the combined...

<https://arxiv.org/abs/2609.05235>

08 CONF

Lightweight Vision Transformer Compression for On-Device Plant Disease Detection in Resource-Constrained Agricultural Field Conditions

ArXiv CS.AI 50%

arXiv:2609.05334v1 Announce Type: cross Abstract: Chilli (*Capsicum annuum*) is one of India's most economically significant crops, yet its productivity is persistently threatened by diseases that are difficult to identify without expert intervention. While Vision Transformers...

<https://arxiv.org/abs/2609.05334>

09 CONF

Design Docs Are All You Need: An AI-native Machine-Learning Performance Tool

ArXiv CS.AI 50%

arXiv:2609.05364v1 Announce Type: cross Abstract: Machine-learning performance modeling is a uniquely hostile terrain for long-lived software: the assumptions baked into today's abstractions are invalidated by tomorrow's models and systems, forcing perpetual refactoring of...

<https://arxiv.org/abs/2609.05364>

10 CONF

When LLM Decompilers Recompile More and Preserve Less

ArXiv CS.AI 50%

arXiv:2609.05370v1 Announce Type: cross Abstract: Decompilation recovers high-level source from compiled machine code and serves as a foundation for security tasks such as vulnerability detection and malware analysis. Traditional decompilers like Ghidra and Hex-Rays expose...

<https://arxiv.org/abs/2609.05370>

11 CONF

Diffusion TV: Experiencing Diffusion Models through Tangible, Embodied Interaction

ArXiv CS.AI 50%

arXiv:2609.05404v1 Announce Type: cross Abstract: Diffusion TV is an interactive AI art installation that offers a tangible and embodied experience of diffusion models through a modified CRT TV. By physically manipulating the TV's antenna, audiences control the clarity of...

<https://arxiv.org/abs/2609.05404>

12 CONF

Procedural Content Generation via Generative Artificial Intelligence

ArXiv CS.AI 50%

arXiv:2407.09013v3 Announce Type: replace Abstract: The attempt to utilize machine learning in procedural content generation (PCG) has been made in the past. In this survey paper, we investigate how generative artificial intelligence (AI), which saw a significant increase in...

<https://arxiv.org/abs/2407.09013>

13 CONF

Achieving Olympiad-Level Geometry Large Language Model Agent via Complexity Boosting Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2512.10534v5 Announce Type: replace Abstract: Large language model (LLM) agents exhibit strong mathematical problem-solving abilities and can even solve International Mathematical Olympiad (IMO) level problems with the assistance of formal proof systems. However, due to...

<https://arxiv.org/abs/2512.10534>

14 CONF

Measuring proximity to standard planes during fetal brain ultrasound scanning

ArXiv CS.AI 50%

arXiv:2404.07124v2 Announce Type: replace-cross Abstract: This paper presents a pipeline designed to bring ultrasound (US) plane pose estimation closer to clinical use, demonstrating the feasibility of continuous, real-time proximity feedback for navigation to the standard...

<https://arxiv.org/abs/2404.07124>

15 CONF

AI-Powered CPS-Enabled Vulnerable-User-Aware Urban Transportation Digital Twin: Methods and Applications

ArXiv CS.AI 50%

arXiv:2501.10396v4 Announce Type: replace-cross Abstract: We present methods and applications for the development of digital twins (DT) for urban traffic management. While the majority of studies on the DT focus on its ``eyes," which is the emerging sensing and perception like...

<https://arxiv.org/abs/2501.10396>

16 CONF

Partial Inverse Design of High-Performance Concrete Using Cooperative Neural Networks for Constraint-Aware Mix Generation

ArXiv CS.AI 50%

arXiv:2512.06813v3 Announce Type: replace-cross Abstract: High-performance concrete (HPC) requires complex mix design decisions involving interdependent variables and practical constraints. While data-driven methods have improved predictive modeling for forward design in...

<https://arxiv.org/abs/2512.06813>

17 CONF

Comparables XAI: Faithful Example-based AI Explanations with Counterfactual Trace Adjustments

ArXiv CS.AI 50%

arXiv:2602.13784v2 Announce Type: replace-cross Abstract: Explaining with examples is an intuitive way to justify AI decisions. However, it is challenging to understand how a decision value should change relative to the examples with many features differing by large amounts. We...

<https://arxiv.org/abs/2602.13784>

18 CONF

YOLO with Kolmogorov-Arnold networks and vision-language foundation models for interpretable object detection with trustworthy multimodal AI in computer vision perception

ArXiv CS.AI 50%

arXiv:2603.23037v2 Announce Type: replace-cross Abstract: The trustworthy object detection capabilities of a novel Kolmogorov-Arnold network framework are examined here. The approach addresses a key limitation in computer vision for vehicle detection perception, and beyond....

<https://arxiv.org/abs/2603.23037>

19 CONF

From Architecture to Output: Structural Origins of Hallucination in Large Language Models and the Amplifying Role of Data

ArXiv CS.AI 50%

arXiv:2606.07537v2 Announce Type: replace-cross Abstract: Large language models produce fluent, confident, factually wrong output. Existing taxonomies classify these failures by output type -- intrinsic versus extrinsic, faithfulness versus factuality -- but say nothing about...

<https://arxiv.org/abs/2606.07537>

20 CONF

Golden Ruler: A Numeric Format Catalog with Bit-Exact Conformance Vectors for FP8, BF16, MXFP4, and Microscaling Formats

ArXiv CS.AI 50%

arXiv:2606.09686v3 Announce Type: replace-cross Abstract: Numeric format proliferation in machine learning hardware -- FP8 (E4M3 and E5M2), BF16, MXFP4, microscaling block formats, and dozens of research variants -- has outpaced the availability of vendor-neutral, bit-exact...

<https://arxiv.org/abs/2606.09686>

21 CONF

Attributable by Construction: Claim-Anchored Provenance for Multi-Document Summarization

ArXiv CS.AI 50%

arXiv:2606.23989v4 Announce Type: replace-cross Abstract: Large language models produce fluent multi-document summaries, but their attributions are typically coarse---whole documents or passages---and generated post hoc, leaving each statement hard to verify. We argue that...

<https://arxiv.org/abs/2606.23989>

22 CONF

Not All LLM Reasoning is Visible in the Chain-of-Thought

ArXiv CS.AI 50%

arXiv:2607.22925v2 Announce Type: replace-cross Abstract: A key question for AI safety is whether a language model expresses all of its reasoning in its output tokens. We demonstrate a concrete failure mode where frontier models exhibit invisible reasoning by leveraging...

<https://arxiv.org/abs/2607.22925>

23 CONF

SPD: Single Pass Decoding for Generative Reranking

ArXiv CS.AI 50%

arXiv:2609.01807v2 Announce Type: replace-cross Abstract: Large language models (LLMs) achieve state-of-the-art generative ranking quality, but the ranking they produce must be decoded, and autoregressive decoding spends one sequential forward pass per emitted token. We observe...

<https://arxiv.org/abs/2609.01807>

24 CONF

Post-Training Language Models for Gold-Medal Performance in Coding Competitions

ArXiv CS.AI 50%

arXiv:2609.02849v2 Announce Type: replace-cross Abstract: Competitive programming has become a key test of large language model reasoning, with international competitions such as IOI and ICPC representing its most challenging settings. We present an end-to-end specialization...

<https://arxiv.org/abs/2609.02849>

25 CONF

Tree species mapping in Denmark: A comparison of spectral-temporal features with geospatial foundation model embeddings

ArXiv CS.AI 50%

arXiv:2609.03480v2 Announce Type: replace-cross Abstract: We map tree species across Denmark using National Forest Inventory plots and EO data, while evaluating the potential of foundation models for large-scale forest characterization. We compare two alternative input...

<https://arxiv.org/abs/2609.03480>