



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	90	Media
HackerNews	10	Community
TechCrunch AI	6	Media
The Verge AI	3	Media
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 2 stories

VERI

Show HN: MaruCheck - Independent QA for AI-generated code

HackerNews

+1 source

90%

HN 3

Hi HN, I've been building MaruCheck, an independent, open-source QA verification tool aimed to solve the issues that arise when coding agents like codex, claude, cursor make semantic issues that has heavy ramifications. Take the example where the requirements says a free user...

<https://github.com/Kidus-M/MaruCheck>

VERI

Meta's Muse AI works and creeps me out

The Verge AI

+1 source

80%

Meta has launched its new Muse assistant, marking the company's first real foray into AI-powered productivity tools. The company says its AI agent can "take the busywork off your plate" by helping you with online shopping, emails, trip-planning, and more. I decided to try out...

<https://www.theverge.com/tech/993391/meta-muse-ai-hands-on>

CONFIRMED SIGNALS — 108 stories

CONF

Show HN: Self-hosted company OS, Claude Code and Codex agents in departments

HackerNews

50%

HN 46

Hi HN, I am Dimitris. This is a company OS that I built and use to run my business and anyone can install it and self host it for free. Think of it as Claude Code, Cowork and the cloud sessions in one self hosted application. It is a Multi-tenant application by design where many...

<https://github.com/OtoDock/oto-dock>

CONF

Show HN: Calmscroll - a reader that shows the current paragraph in the library

HackerNews

50%

HN 6

No summary — see link for full article.

<https://calmscroll.com>

CONF

Show HN: ObserverBench - Test internal AI monitors by the harm they miss

HackerNews

50%

HN 4

No summary — see link for full article.

<https://kwisatzh.github.io/observerbench/try/>

CONF

Show HN: Post-graph-RAG - your RAG still thinks the old CFO is the CFO

HackerNews

50%

HN 4

Hello! I created Fashionguesser, a daily fashion trivia game. Inspired by other daily trivia games like Timeguesser, Connections, and Wordle, Fashionguesser is my attempt at the same thing for people curious about fashion. The game is broken up into three parts: Runway Frames -...

<https://github.com/crajah/post-graph-rag>

CONF

Show HN: Faint Light, 10-100x faster plate-solving/astrometry in Rust

HackerNews

50%

HN 4

No summary — see link for full article.

https://github.com/fschuindt/faint_light

CONF

Show HN: News observability site tracks coverage by party lean/distraction/etc.

HackerNews 50% HN 3

Heya HN! I'm helping launch this site which uses Qwen 3-32b to process thousands of news articles and analyze them by region, public interest, partisan lean, etc. We would love your feedback!

<https://pressaudit.org/>

CONF

Show HN: MySamantha - a productivity app where the AI is optional

HackerNews 50% HN 3

Hi HN, in a world where 100% of code is written by AI (pew, sometimes even 200%), there might still be a small percentage of us who want to write code in a more efficient way, especially when it comes to LLM interaction. I honestly tired converting my existing TS types into...

<https://news.ycombinator.com/item?id=49645994>

0 CONF

Show HN: Raiders at the Gate - tower defense web game

HackerNews 50% HN 2

No summary — see link for full article.

<https://games.chrisbeach.co.uk/raiders-at-the-gate/>

1 CONF

Show HN: I built an AI 'Fog of War' for life goals

HackerNews 50% HN 2

No summary — see link for full article.

<https://www.tryworthit.app>

2 CONF

Jensen Huang explains why Nvidia will grow an astounding 70% next year

TechCrunch AI 50%

Nvidia has its finger in every pie, and sees another year of plenty in its future, Jensen Huang says. But, he insists, its deals are not circular.

<https://techcrunch.com/2026/09/10/jensen-huang-explains-why-nvidia-will-grow-an-astound...>

3 CONF

Anthropic details distillation campaigns from Alibaba, Moonshot AI, and DeepSeek

TechCrunch AI 50%

A new report released Thursday by Anthropic alleges persistent distillation attacks by China-based AI companies, which have escalated in recent months as competition in the space has intensified.

<https://techcrunch.com/2026/09/10/anthropic-details-distillation-campaigns-from-alibaba...>

4 CONF

Anthropic reveals rogue AI agents hate CAPTCHAs, just like you

TechCrunch AI 50%

Come inside the mind of a bot trying to convince the internet it's human.

<https://techcrunch.com/2026/09/10/anthropic-reveals-rogue-ai-agents-hate-captchas-just-...>

5 CONF

India's Pocket FM doubles revenue run rate to \$500M as AI powers 93% of audio content

TechCrunch AI 50%

Pocket FM uses AI to produce 99% of its new content, helping make content production about 80 times cheaper.

<https://techcrunch.com/2026/09/10/indias-pocket-fm-doubles-revenue-run-rate-to-500m-as-...>

6 CONF

AI agents are flooding public services with new requests

TechCrunch AI 50%

"The vast majority of cases we find are people who are entitled to claim for something, claiming for that thing," the researcher told TechCrunch.

<https://techcrunch.com/2026/09/10/ai-agents-are-flooding-public-services-with-new-reque...>

7 CONF

Powering AI is an architecture problem

MIT Tech Review 50%

On July 22, 2026, a transmission line fault in Ashburn, Virginia—the heart of the world's largest data center cluster—knocked more than 3 gigawatts of load off the grid in seconds. And it wasn't the first time. Two years earlier, a single failed surge arrester dropped roughly 60...

<https://www.technologyreview.com/2026/09/10/1141649/powering-ai-is-an-architecture-prob...>

8 CONF

Healthcare AI's next test is integration

MIT Tech Review 50%

The entrance of major AI companies into healthcare is a meaningful and welcome development, accelerating the technical foundation available to the industry. Their models are increasingly capable of processing long clinical records, interpreting complex terminology, comparing...

<https://www.technologyreview.com/2026/09/10/1141421/healthcare-ais-next-test-is-integra...>

9 CONF

The Download: OpenAI's turning point for math and a battery record

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. What OpenAI's latest controversy tells us about the future of math OpenAI says its agents have solved one of the most important open problems...

<https://www.technologyreview.com/2026/09/09/1143767/the-download-openai-math-future-bat...>

0 CONF

Universal Music is launching an AI music platform with ElevenLabs

The Verge AI 50%

Universal Music Group is launching a new AI-powered platform that will allow users to draw from its catalog of licensed music to create song remixes, mashups, and new takes on tracks, according to an announcement on Thursday. The record label is developing the platform through a...

<https://www.theverge.com/ai-artificial-intelligence/993465/universal-music-elevenlabs-ai>

1 CONF

Mathematicians want proof OpenAI didn't use their work

The Verge AI 50%

Another researcher is challenging OpenAI about the data driving its increasingly impressive array of mathematical discoveries. Just days after a bitter row erupted over whether the company's models benefited from unpublished work, a second mathematician has come forward accusing...

<https://www.theverge.com/ai-artificial-intelligence/993263/where-does-openai-get-mathem...>

2 CONF

Scores Alone Do Not Prove Discovery: The Discovery Certification Protocol for Auditing AI Research Agents

ArXiv CS.AI 50%

arXiv:2609.09219v1 Announce Type: cross Abstract: AI research agents combine prior knowledge, public sources, and experimental feedback to produce useful results. The Discovery Certification Protocol (DCP) turns claims about these results into executable recovery and feedback...

<https://arxiv.org/abs/2609.09219>

3 **CONF****An Autonomous GeoAI Agent for Arctic Eco-Navigation****ArXiv CS.AI** **50%**

arXiv:2609.09374v1 Announce Type: new Abstract: Arctic maritime navigation is becoming increasingly important as changing sea-ice conditions expand seasonal accessibility while simultaneously introducing substantial operational, environmental, and community risks. Arctic route...

<https://arxiv.org/abs/2609.09374>

4 **CONF****From Plausible to Actionable: A Position on LLM Self-Explanations****ArXiv CS.AI** **50%**

arXiv:2607.15957v3 Announce Type: cross Abstract: Large Language Models (LLMs) can generate natural language explanations that rationalize their own decisions, a phenomenon commonly referred to as self-explanations. Such explanations have emerged as a promising direction for...

<https://arxiv.org/abs/2607.15957>

5 **CONF****Seven Sources of Physical AI Capability Formation****ArXiv CS.AI** **50%**

arXiv:2609.09627v1 Announce Type: new Abstract: Capabilities relevant to Physical AI can arise from materially different formation histories, yet existing taxonomies organized by morphology, architecture, learning algorithm, task, or domain do not directly answer what gives rise...

<https://arxiv.org/abs/2609.09627>

6 **CONF****Black-Box Red Teaming of Agentic AI: A Taxonomy-Driven Framework for Automated Risk Discovery****ArXiv CS.AI** **50%**

arXiv:2609.09647v1 Announce Type: new Abstract: Agentic systems are rapidly moving to production, where they read untrusted inputs, call tools with real permissions, and act autonomously, expanding the security surface beyond chat-only models. Yet standard evaluations remain...

<https://arxiv.org/abs/2609.09647>

7 **CONF****Compact Visuotactile World Models for Lifting: Prediction, Reward Alignment, and Force Constraints****ArXiv CS.AI** **50%**

arXiv:2609.09597v2 Announce Type: cross Abstract: Accurate tactile forecasts need not improve force-constrained control. We study a 652,157-parameter action-conditioned visuotactile world model with matched behavior cloning, policy learning in imagination, independent reactive...

<https://arxiv.org/abs/2609.09597>

8 **CONF****Safe to Stop? Risk-Constrained Stopping for Sequential Clinical Diagnosis Agents****ArXiv CS.AI** **50%**

arXiv:2609.09678v1 Announce Type: new Abstract: Clinical diagnosis agents must decide not only what test to request next, but also when to diagnose or defer. Existing agent benchmarks largely evaluate accuracy after fixed or unconstrained interaction, leaving autonomous stopping...

<https://arxiv.org/abs/2609.09678>

9 CONF

Can Artificial Intelligence Support Healthcare and Mental Health Through Early Cyberbullying Detection ? The Impact of Emotion-Aware AI on Proactive Online Safety

ArXiv CS.AI 50%

arXiv:2609.09735v1 Announce Type: new Abstract: Healthcare systems, mental health, and public well-being are increasingly affected by cyberbullying and harmful online interactions. This paper presents CareGuard, an early-warning framework designed to support healthcare-driven...

<https://arxiv.org/abs/2609.09735>

0 CONF

AgentAudit: An Open, Extensible Framework for Full-Lifecycle Trust Evaluation of AI Agents

ArXiv CS.AI 50%

arXiv:2609.09875v1 Announce Type: new Abstract: Existing evaluation frameworks mostly assess only one part of AI agents, such as task completion (AgentBench) or security robustness (AgentDojo, ASB), rather than the complete pipeline of planning, tool selection, tool execution,...

<https://arxiv.org/abs/2609.09875>

1 CONF

IBIB: A Protocol for Measuring Enterprise AI Systems by Serving Route, Not Model Identifier

ArXiv CS.AI 50%

arXiv:2609.10494v1 Announce Type: cross Abstract: Enterprises deploy systems, not checkpoints. Usable capability depends jointly on weights, serving route, precision, output contract, and harness, yet all 18 audited benchmarks score advertised model identifiers. We treat this as...

<https://arxiv.org/abs/2609.10494>

2 CONF

The Era by Eon Benchmark: A Generated Enterprise Estate with Exact Ground Truth for Benchmarking LLM Agents

ArXiv CS.AI 50%

arXiv:2609.09853v1 Announce Type: new Abstract: LLM agents for enterprise systems of record cannot be evaluated on customer production data, and no existing substitute provides ground truth. We present the Era by Eon Benchmark for evaluating LLM agents that use enterprise tools....

<https://arxiv.org/abs/2609.09853>

3 CONF

Scored vs. Generated Readouts in Behavioral Language Models: An Empirical Study of Elicitation Format

ArXiv CS.AI 50%

arXiv:2609.09882v1 Announce Type: new Abstract: Language models fine-tuned on customer behavior can predict outcomes and generate explanations, but these readouts are often treated as interchangeable. Holding model checkpoint and prompt content fixed, we compare probabilities...

<https://arxiv.org/abs/2609.09882>

4 CONF

Decision Transformer for UAV-Mounted RIS-Assisted Dynamic D2D Communications

ArXiv CS.AI 50%

arXiv:2609.09885v1 Announce Type: new Abstract: This paper studies unmanned aerial vehicle (UAV)-mouted reconfigurable intelligent surface (RIS)-assisted device-to-device (D2D) communication with stochastic link activation. It models UAV motion and attitude, time-varying Rician...

<https://arxiv.org/abs/2609.09885>

5 CONF

Grounded Evaluation and Repair for NL-to-PDDL Problem Generation

ArXiv CS.AI 50%

arXiv:2609.09898v1 Announce Type: new Abstract: Large Language Models (LLMs) have shown promise for translating Natural Language (NL) planning descriptions into PDDL problem instances. However, standard evaluation criteria such as syntactic validity or planner success can...

<https://arxiv.org/abs/2609.09898>

6 CONF

Belief-State Engine: Augmenting LLMs for Principled Planning Under Partial Observability

ArXiv CS.AI 50%

arXiv:2609.10036v1 Announce Type: new Abstract: Large language model agents produce fluent action sequences across a wide range of tasks, yet they fail in characteristic ways once the environment becomes partially observable. Ambiguous feedback pushes them into premature...

<https://arxiv.org/abs/2609.10036>

7 CONF

OntologyAligner: Ontology-Aligned Retrieval and Hierarchy-Guided Large Language Model Reranking for Biomedical Ontology Normalization

ArXiv CS.AI 50%

arXiv:2609.10055v1 Announce Type: new Abstract: Biomedical ontology normalization maps free-text expressions to standardized concepts, enabling consistent integration and analysis of biomedical data. This task remains challenging because lexical variation and subtle distinctions...

<https://arxiv.org/abs/2609.10055>

8 CONF

Reference-Based Bias Detection in LLMs via Relative Representations of Hidden States

ArXiv CS.AI 50%

arXiv:2609.10060v1 Announce Type: new Abstract: Existing bias auditing methods typically rely on model outputs, requiring costly benchmarks or judge models and potentially missing internal shifts that never appear in generated text. We propose a reference-based method that...

<https://arxiv.org/abs/2609.10060>

9 CONF

Agent-Based ML-LLM Fusion with Self-Optimizing Prompts for Plateau Weather Alerts

ArXiv CS.AI 50%

arXiv:2609.10135v1 Announce Type: new Abstract: To address insufficient contextualization, weak generalization, and poor scenario adaptation in tourism meteorological services, we propose SmartWeatherAgent--a unified three-stage architecture integrating intent recognition,...

<https://arxiv.org/abs/2609.10135>

10 CONF

Bringing Value Models Back: Generative Critics for Value Modeling in LLM Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2604.10701v2 Announce Type: replace-cross Abstract: Credit assignment is a central challenge in reinforcement learning (RL). Classical actor-critic methods address this challenge through fine-grained advantage estimation based on a learned value function. However, learned...

<https://arxiv.org/abs/2604.10701>

1 **CONF****TRACE: Training Reasoning Agents for Causal Exploration with Synthesized Rewards****ArXiv CS.AI 50%**

arXiv:2609.10315v1 Announce Type: new Abstract: Reinforcement learning with verifiable rewards (RLVR) has advanced language-model reasoning in domains such as mathematics and code, where objective answers are inexpensive to check. Diagnostic reasoning over complex data lacks...

<https://arxiv.org/abs/2609.10315>

2 **CONF****RiLM: Parameter-Efficient Language Modeling via Geodesic Decoding****ArXiv CS.AI 50%**

arXiv:2609.10305v1 Announce Type: cross Abstract: Language models under one million parameters matter for edge deployment, domain adaptation, and reproducible research, yet a two-layer LSTM or Transformer at embedding width $d = 128$ still spends roughly one third of its capacity...

<https://arxiv.org/abs/2609.10305>

3 **CONF****Cyber-Financial Contagion: Modeling the Propagation of an AI Vendor Compromise Through the Banking System****ArXiv CS.AI 50%**

arXiv:2609.10350v1 Announce Type: new Abstract: The banking system now depends on a small set of shared artificial intelligence vendors for fraud screening, credit decisioning, anti-money-laundering triage, customer analytics, and internal decision support. This paper studies...

<https://arxiv.org/abs/2609.10350>

4 **CONF****Fortunate Recall: Ontology-Driven Memory Lifecycle Management for Persistent Coherence in LLMs****ArXiv CS.AI 50%**

arXiv:2609.10413v1 Announce Type: new Abstract: Current LLM memory systems treat all personal facts identically, so stores grow without bound while retrieval precision degrades. The core challenge is lifecycle management: which memories should persist, which should be replaced,...

<https://arxiv.org/abs/2609.10413>

5 **CONF****Quantifying Logical Consistency in Transformers via Query-Key Alignment****ArXiv CS.AI 50%**

arXiv:2502.17017v1 Announce Type: cross Abstract: Large language models (LLMs) have demonstrated impressive performance in various natural language processing tasks, yet their ability to perform multi-step logical reasoning remains an open challenge. Although Chain-of-Thought...

<https://arxiv.org/abs/2502.17017>

6 **CONF****RAU: Reference-based Anatomical Understanding with Vision Language Models****ArXiv CS.AI 50%**

arXiv:2509.22404v2 Announce Type: replace-cross Abstract: Anatomical understanding, which is the ability to identify, localize, or segment anatomical structures, is critical in medical image analysis; however, its progress is constrained by the scarcity of expert-labeled data. A...

<https://arxiv.org/abs/2509.22404>

7 CONF

AgentHijack: Visual Patch Attacks on Multimodal Computer-Use Agents

ArXiv CS.AI 50%

arXiv:2609.09212v1 Announce Type: cross Abstract: This paper presents an end-to-end evaluation framework for image-triggered command injection against computer-use agents (CUAs). The goal is to test whether a local visual patch can induce verifiable environmental consequences...

<https://arxiv.org/abs/2609.09212>

8 CONF

Geometry Conditioning in an Embodied SLM: Training Controls and Robustness Diagnostics in a 0.8B Hybrid Model

ArXiv CS.AI 50%

arXiv:2609.09213v1 Announce Type: cross Abstract: We study how physical-state inputs affect a 0.8B hybrid language model adapted for manipulation with 6.2M trainable parameters. Six conditions are trained on three LIBERO-Spatial tasks and evaluated over three seeds and 540...

<https://arxiv.org/abs/2609.09213>

9 CONF

Compute-Bounded Security Assurance - Coverage, Verification, and Response under Resource Constraints

ArXiv CS.AI 50%

arXiv:2609.09229v1 Announce Type: cross Abstract: Additional inference compute can increase the number of correctly resolved security-assurance tasks, but repeated success, unique coverage, accepted evidence, and operational protection are different quantities. We develop a...

<https://arxiv.org/abs/2609.09229>

0 CONF

Scaling Post-Training Ternarisation to Qwen3-8B Capability Retention, Reproduction, Lossless Packing, and Packed Execution

ArXiv CS.AI 50%

arXiv:2609.09240v1 Announce Type: cross Abstract: Ultra-low-bit language models promise reductions in storage and memory traffic, but a nominal "1.58-bit" label does not specify the deployed representation or its execution cost. We study a scale-up of an aggressive post-training...

<https://arxiv.org/abs/2609.09240>

1 CONF

Distribution-Consistent Inference for Dynamic Sparse Mixture-of-Experts

ArXiv CS.AI 50%

arXiv:2609.09241v1 Announce Type: cross Abstract: Mixture-of-Experts (MoE) architectures have emerged as a powerful paradigm for scaling model capacity while preserving efficient inference in large foundation models. However, most MoE models use a fixed top-\$k\$ expert selection...

<https://arxiv.org/abs/2609.09241>

2 CONF

In RAG We Trust? Measuring Robustness of Retrieval-Augmented Generation Under Document Poisoning

ArXiv CS.AI 50%

arXiv:2609.09243v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) grounds a language model in retrieved documents, which reduces hallucination but creates a new attack surface: if retrieved text is tampered with, the model may repeat the falsehood. We study...

<https://arxiv.org/abs/2609.09243>

3 CONF

DiffLUT-Net: Differentiable Training of FPGA LUT Networks with Learnable Connectivity

ArXiv CS.AI 50%

arXiv:2609.09254v1 Announce Type: cross Abstract: Field-programmable gate arrays (FPGAs) enable efficient neural-network inference, but most deployment flows either accelerate multiply-accumulate operations or convert pretrained quantized models into lookup tables (LUTs). We...

<https://arxiv.org/abs/2609.09254>

4 CONF

Improving 5G AI-RAN MCS Selection by Predicting Retransmissions

ArXiv CS.AI 50%

arXiv:2609.09324v1 Announce Type: cross Abstract: Link Adaptation (LA) in 5G NR is inherently reactive, relying on channel measurements and HARQ feedback that may become quickly obsolete when the channel changes quickly. This data is also noisy, making it hard to track...

<https://arxiv.org/abs/2609.09324>

5 CONF

Smart Adaptive Computing Across the Continuum: LLMs in IoT-Edge-Cloud Resource Management

ArXiv CS.AI 50%

arXiv:2609.09348v1 Announce Type: cross Abstract: Managing resources across IoT, edge, and cloud layers calls for continuous, context-aware decisions under constraints that rarely stay fixed. Deep reinforcement learning (DRL) handles this class of problems well, and large...

<https://arxiv.org/abs/2609.09348>

6 CONF

MetroLLM-Bench: Evaluating Language Models as Transit Kiosk Runtimes

ArXiv CS.AI 50%

arXiv:2609.10016v1 Announce Type: cross Abstract: We introduce MetroLLM-Bench, a 955-case benchmark for testing language models as the policy layer of a transit kiosk. It covers six real metro systems, ranging from 37 to 414 stations, and eleven categories that include routing,...

<https://arxiv.org/abs/2609.10016>

7 CONF

Edu-QuRating: Multi-Dimensional Educational Data Curation with Distilled Pairwise Judgements

ArXiv CS.AI 50%

arXiv:2609.09425v2 Announce Type: cross Abstract: Educational data filters have become a practical way to improve language-model pre-training, but most filters treat educational value as a single scalar property. This may be too broad for some applications, especially if the...

<https://arxiv.org/abs/2609.09425>

8 CONF

Distributed Physical Layer Authentication and Collaborative RSMA in Non-Terrestrial Networks via Graph Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2609.09475v1 Announce Type: cross Abstract: Existing physical-layer authentication (PLA) schemes for non-terrestrial networks (NTNs) often rely on single-anchor verification, lack joint authentication-transmission design, and ignore tag privacy leakage under eavesdropping....

<https://arxiv.org/abs/2609.09475>

9 CONF

From Fixed Keys to Readable Schemas: Small Language Models for Vehicle Agent Function Calls

ArXiv CS.AI 50%

arXiv:2609.09476v1 Announce Type: cross Abstract: In-vehicle assistants must translate natural-language requests into accurate vehicle function calls under strict memory and latency constraints, making small language models (SLMs) attractive for on-device deployment. For such...

<https://arxiv.org/abs/2609.09476>

0 CONF

A Statistical Approach to Estimating Sample Size of Machine Learning Models

ArXiv CS.AI 50%

arXiv:2609.09547v1 Announce Type: cross Abstract: Sample size determination for machine learning (ML) prediction models is challenging because conventional power analysis typically requires the predictor-outcome relationship and effect structure to be specified a priori....

<https://arxiv.org/abs/2609.09547>

1 CONF

Arbitrary Cipher Attacks Against Large Language Models Do Not Require Fine-Tuning

ArXiv CS.AI 50%

arXiv:2609.09553v1 Announce Type: cross Abstract: Large language model safety and security research is preoccupied with, among other things, detecting and preventing jailbreak attacks: alignment bypasses that allow an adversarial user to elicit unwanted or harmful outputs from...

<https://arxiv.org/abs/2609.09553>

2 CONF

The Vibe Shift in Software Engineering: Evaluating AI-Led Conversational Programming for Performance, Cognition, and Responsible Adoption

ArXiv CS.AI 50%

arXiv:2609.09560v1 Announce Type: cross Abstract: This study evaluates Vibe Coding, an emerging AI-led conversational programming paradigm that enables developers to generate software through natural-language interaction with large language models. Using a mixed-methods design,...

<https://arxiv.org/abs/2609.09560>

3 CONF

Myocardial Strain Drift Correction in Deep Learning Based Ultrasound Tracking

ArXiv CS.AI 50%

arXiv:2609.09577v1 Announce Type: cross Abstract: Myocardial strain from echocardiography is a key biomarker for cardiac function. Recent deep learning methods show strong performance for myocardial motion tracking but often lack physiological constraints, leading to temporal...

<https://arxiv.org/abs/2609.09577>

4 CONF

Watermarks Without Verification: AI Text Watermarking After the EU AI Act

ArXiv CS.AI 50%

arXiv:2609.09604v1 Announce Type: cross Abstract: On August 2, 2026, the obligations of Article 50 of the EU AI Act took effect, requiring generative AI providers to mark the content their systems produce and ensure it can be detected as AI-generated. Days later, Anthropic...

<https://arxiv.org/abs/2609.09604>

5 **CONF****CT-SAFR: Safe and Interpretable Chain-of-Thought Reasoning for Autonomous Robots: A Multi-Layered Verification Framework for Trustworthy AI-Driven Robotic Decision Making****ArXiv CS.AI** **50%**

arXiv:2609.09692v1 Announce Type: cross Abstract: Chain-of-Thought (CoT) prompting enables LLMs to perform explicit, step-by-step reasoning, creating opportunities for sophisticated autonomous robots. However, recent research reveals that reasoning models verbalize their actual...

<https://arxiv.org/abs/2609.09692>

6 **CONF****When Auditors Fabricate: Batch-Size Degradation and Confident Hallucination in LLM Detection of Planted Document Contamination****ArXiv CS.AI** **50%**

arXiv:2609.09696v1 Announce Type: cross Abstract: Large language models are increasingly proposed as automated auditors of document quality, yet their reliability as detectors of planted errors is poorly characterised. We construct a contaminated corpus of 150 academic papers...

<https://arxiv.org/abs/2609.09696>

7 **CONF****Kernel-Complexity Edge Sanitization for Training-Free Defense against Structural Graph Attacks****ArXiv CS.AI** **50%**

arXiv:2609.09698v1 Announce Type: cross Abstract: Graph Neural Networks (GNNs) have achieved remarkable success across diverse applications, yet they remain highly vulnerable to adversarial attacks that maliciously perturb graph structure. Existing defenses often lack rigorous...

<https://arxiv.org/abs/2609.09698>

8 **CONF****Fine-Tuning a KV Cache Concatenation-Aware Model or Recomputing KV Caches? Why Not Both?****ArXiv CS.AI** **50%**

arXiv:2609.09768v1 Announce Type: cross Abstract: In Retrieval-Augmented Generation (RAG) systems, a large number of retrieved chunks are concatenated to form the input context so that users can receive high-quality responses based on external knowledge. As a result, the input...

<https://arxiv.org/abs/2609.09768>

9 **CONF****BRACE: Anchored Bellman-Residual Correction for Stale Critics in Asynchronous RL****ArXiv CS.AI** **50%**

arXiv:2609.09783v1 Announce Type: cross Abstract: Asynchronous reinforcement learning has become the standard way to scale training for language models, but the resulting policy lag biases the critic toward the stale behavior policy. Existing work on asynchronous LLM training...

<https://arxiv.org/abs/2609.09783>

0 **CONF****Pairit: A Platform for Live Experiments on Human-AI Collaboration****ArXiv CS.AI** **50%**

arXiv:2609.09789v1 Announce Type: cross Abstract: Organizational design in the era of artificial intelligence requires experimental methods that can test how human-AI groups coordinate, delegate, and make decisions. Programmable platforms coordinate live human-to-human sessions...

<https://arxiv.org/abs/2609.09789>

1 CONF

How Fragile Is Safety Alignment at Frontier Scale? A Single-Direction Attack on a 320B MoE

ArXiv CS.AI 50%

arXiv:2609.09793v1 Announce Type: cross Abstract: Directional ablation removes an aligned language model's ability to refuse by projecting a single "refusal direction" out of the weights that write the residual stream. It needs no gradient-based training and no optimization,...

<https://arxiv.org/abs/2609.09793>

2 CONF

CS-Guard: Benchmarking LLM Guardrails for Code Generation Security

ArXiv CS.AI 50%

arXiv:2609.09798v1 Announce Type: cross Abstract: Large language models (LLMs) have been exploited to generate malware, but the effectiveness of guardrails for code generation security remains unclear. We introduce CS-Guard, the first benchmark to systematically evaluate...

<https://arxiv.org/abs/2609.09798>

3 CONF

Subgroup Membership Inference Audits of Differentially Private Synthetic Text

ArXiv CS.AI 50%

arXiv:2609.09848v1 Announce Type: cross Abstract: Synthetic data releases are increasingly proposed in the literature as a means of sharing realistic data replicas in lieu of sensitive private datasets. Even when the worst-case privacy leakage of such releases is bounded by...

<https://arxiv.org/abs/2609.09848>

4 CONF

Can AI Agents Deliver Verifiable Network-Wide Outcomes Across Authority Boundaries?

ArXiv CS.AI 50%

arXiv:2609.10181v1 Announce Type: cross Abstract: AI agents are increasingly involved in network automation, where they can initiate configuration changes through mediated operational interfaces and assess the resulting state. Nonetheless, operational networks usually span many...

<https://arxiv.org/abs/2609.10181>

5 CONF

With a Thermomix You Lose the Ability to Cook: A Kitchen Machine Analogy for Applications of Generative AI in Education

ArXiv CS.AI 50%

arXiv:2609.09856v1 Announce Type: cross Abstract: The rapid adoption of generative AI tools such as ChatGPT has sparked intense debate about their risks and opportunities for education, as well as the ways researchers should investigate them. In this paper, we approach these...

<https://arxiv.org/abs/2609.09856>

6 CONF

Forward-Free LLM Depth Pruning via Weight Redundancy

ArXiv CS.AI 50%

arXiv:2609.09883v1 Announce Type: cross Abstract: Depth pruning reduces large language model (LLM) inference cost by removing complete Transformer blocks. Activation-based methods collect hidden states through forward passes on calibration data, while existing forward-free...

<https://arxiv.org/abs/2609.09883>

7 **CONF****Direct Diversity Optimization for Diverse Successful Trajectories in Preference Post-Training**

ArXiv CS.AI 50%

arXiv:2609.10052v1 Announce Type: cross Abstract: LLM agents for sequential decision tasks are often post-trained with trajectory-level outcome labels, but such labels provide little supervision for preserving multiple successful branches from the same decision state. We study...

<https://arxiv.org/abs/2609.10052>

8 **CONF****Beyond Training: A Feasibility Taxonomy for Inference-Time AI Governance**

ArXiv CS.AI 50%

arXiv:2609.10105v1 Announce Type: cross Abstract: Compute governance today is a governance of training: the thresholds, reporting requirements, and frontier-AI regimes now in force attach to training compute and treat the trained model as the regulatory unit. That picture is...

<https://arxiv.org/abs/2609.10105>

9 **CONF****Context operations to architecture modelling output from large language models and evaluation criteria for their use in systems engineering design**

ArXiv CS.AI 50%

arXiv:2609.10132v1 Announce Type: cross Abstract: The development of generative artificial intelligence resources enables opportunities of speeding up systems and engineering design work. This contribution introduces a framework of formal operations for assembling context in...

<https://arxiv.org/abs/2609.10132>

0 **CONF****Active Adaptation, Not Static Defense: Temporal Dynamics of Preventative Steering in Adversarial Fine-Tuning**

ArXiv CS.AI 50%

arXiv:2609.10142v1 Announce Type: cross Abstract: Large language models remain fragile against malicious fine-tuning, motivating training-time defenses against harmful persona drift. Preventative Steering injects undesirable-trait persona vectors during fine-tuning and removes...

<https://arxiv.org/abs/2609.10142>

1 **CONF****Hierarchical and Permutation-Invariant Feature Transformation Learning via Policy-Guided Embedding Search**

ArXiv CS.AI 50%

arXiv:2609.10225v1 Announce Type: cross Abstract: Feature transformation improves predictive performance on tabular data by constructing informative abstractions from raw features. Recent generative approaches encode transformation knowledge into continuous embedding spaces for...

<https://arxiv.org/abs/2609.10225>

2 **CONF****DiSCo: A Distribution-First Steering and Cultural Prior Evaluation Framework for Measuring Cultural Preference Bias in LLMs**

ArXiv CS.AI 50%

arXiv:2609.10253v1 Announce Type: cross Abstract: Large language models (LLMs) are increasingly deployed in globally used assistants, yet their default choices in culturally grounded everyday situations can systematically favour some cultures over others, affecting localisation,...

<https://arxiv.org/abs/2609.10253>

3 CONF

GANDR: Claim Auditing for Verifiable Legal Answer Generation

ArXiv CS.AI 50%

arXiv:2609.10293v1 Announce Type: cross Abstract: In high-stakes domains such as legal practice, a language-model answer is only useful to the extent that a reader can verify each claim against the source the system cites. Current grounded-generation pipelines score the answer...

<https://arxiv.org/abs/2609.10293>

4 CONF

One Loop, Two Gains: Can Active Learning win the Lottery for Free?

ArXiv CS.AI 50%

arXiv:2609.10311v1 Announce Type: cross Abstract: The lottery ticket hypothesis posits the existence of winning tickets: sparse subnetworks that, when trained in isolation from their original initialization, match the accuracy of the full dense network. The predominant method...

<https://arxiv.org/abs/2609.10311>

5 CONF

Beyond One-Size-Fits-All: Sample-Adaptive Strategy Routing for Vision Token Pruning in MLLMs

ArXiv CS.AI 50%

arXiv:2609.10346v1 Announce Type: cross Abstract: Multimodal large language models (MLLMs) process hundreds or thousands of visual tokens per image, incurring prohibitive inference costs. While existing vision token pruning methods mitigate this overhead, they implicitly assume...

<https://arxiv.org/abs/2609.10346>

6 CONF

MOONWALK: Mediating Operations with Intent-Evidence-Action Alignment Across Junior-Supervisor Review Workflows in Animation/VFX Pre-Production

ArXiv CS.AI 50%

arXiv:2609.10385v1 Announce Type: cross Abstract: Animation and VFX pre-production review requires teams to translate loosely specified creative intent--briefs, evolving specifications, heterogeneous references, and verbal decisions--into revisions that junior artists can...

<https://arxiv.org/abs/2609.10385>

7 CONF

Can Foundation Models Moderate Online Content? Evaluating Instruction- vs. Example-Driven Policy Operationalization

ArXiv CS.AI 50%

arXiv:2609.10410v1 Announce Type: cross Abstract: The growing complexity of content moderation policies presents a critical challenge for their consistent operationalization. While foundation models possess the basic capabilities needed to confront this challenge, whether they...

<https://arxiv.org/abs/2609.10410>

8 CONF

Emergency Department Revisit Quality Review Screening: Exploring Human Decision-Making and Artificial Intelligence Support

ArXiv CS.AI 50%

arXiv:2609.10421v1 Announce Type: cross Abstract: Background: Emergency Department (ED) return visits are commonly reviewed for quality assurance, but are often limited (e.g., to revisits within 48-72 hours) to increase actionable finding yield while minimizing chart review...

<https://arxiv.org/abs/2609.10421>

9 CONF

Forgetting Only What Matters: Layer-Selective Unlearning toward Robust LLMs

ArXiv CS.AI 50%

arXiv:2609.10439v1 Announce Type: cross Abstract: Large Language Models (LLMs) can memorize and reproduce sensitive, copyrighted, or otherwise undesirable training content, creating privacy, safety, and regulatory concerns. Machine unlearning offers a practical alternative to...

<https://arxiv.org/abs/2609.10439>

0 CONF

Complementing reinforcement learning with SFT through logit averaging in the post training of LLMs

ArXiv CS.AI 50%

arXiv:2605.20555v2 Announce Type: replace-cross Abstract: We introduce a novel method that averages the logits of a frozen reference policy (e.g., SFT) and a trainable policy, and incorporate the method into Group Relative Policy Optimization (GRPO). In contrast to Reinforcement...

<https://arxiv.org/abs/2605.20555>

1 CONF

ROTATE: Regret-driven Open-ended Training for Ad Hoc Teamwork

ArXiv CS.AI 50%

arXiv:2505.23686v3 Announce Type: replace Abstract: Learning to collaborate with previously unseen partners is a fundamental generalization challenge, known as Ad Hoc Teamwork (AHT). Existing methods often adopt a two-stage pipeline: first, a fixed population of teammates is...

<https://arxiv.org/abs/2505.23686>

2 CONF

Non-Stationarity Breaks Permutation Surrogates in Multi-Agent Reinforcement Learning: Diagnosis and Remedies

ArXiv CS.AI 50%

arXiv:2604.23716v4 Announce Type: replace Abstract: Reporting guidance for information-theoretic measures is rarely tested against ground truth. We test one guardrail in two multi-agent reinforcement learning games, a social dilemma and a coordination race, where directed...

<https://arxiv.org/abs/2604.23716>

3 CONF

KernelGenBench: Can LLMs and Agents Write Efficient Kernels Across Operator Sources and Hardware Platforms?

ArXiv CS.AI 50%

arXiv:2607.27231v3 Announce Type: replace Abstract: Modern AI systems depend on specialized accelerator kernels, whose development is complicated by increasingly diverse operators and hardware. LLMs and agentic systems promise to automate this work, but existing evaluations do...

<https://arxiv.org/abs/2607.27231>

4 CONF

Learning to Predict Middle-Layer Attention in MLLMs for Visual Token Pruning

ArXiv CS.AI 50%

arXiv:2608.06411v2 Announce Type: replace Abstract: Multimodal large language models (MLLMs) achieve strong performance across diverse vision-language tasks, but their efficiency is limited by the cost of processing numerous visual tokens. Visual token pruning can reduce this...

<https://arxiv.org/abs/2608.06411>

5 CONF

A Human Audit of OpenAIs AI-Generated Mathematical Proofs

ArXiv CS.AI 50%

arXiv:2608.14673v3 Announce Type: replace Abstract: We assess 18 chapter-specific reviews of the ten mathematical results announced by OpenAI on 1 August 2026, alongside review standards, Lean formalizations, subsequent research, and mathematical references. The article audits...

<https://arxiv.org/abs/2608.14673>

6 CONF

A Composable Evaluation System for Reproducible Omni-Modal Foundation Model Evaluation

ArXiv CS.AI 50%

arXiv:2609.01315v2 Announce Type: replace Abstract: Building an omni-modal foundation model means evaluating it across text, image, video, and audio. Excellent evaluation toolkits exist for each modality, but their inference engines, prompt conventions, and metric...

<https://arxiv.org/abs/2609.01315>

7 CONF

A Taxonomy of Architecture Options for Foundation Model-based Agents: Analysis and Decision Model

ArXiv CS.AI 50%

arXiv:2408.02920v2 Announce Type: replace-cross Abstract: The rapid advancement of AI technology has led to widespread applications of agent systems across various domains. However, the need for detailed architecture design poses significant challenges in designing and operating...

<https://arxiv.org/abs/2408.02920>

8 CONF

BTBR: A Bayesian-Theory-Driven Probabilistic-Fuzzy Framework for Implicit Bias Removal in Large Language Models

ArXiv CS.AI 50%

arXiv:2408.10608v2 Announce Type: replace-cross Abstract: Large language models (LLMs) may encode biased associations from heterogeneous training corpora that are not immediately visible under ordinary prompting, but can surface when the model is steered toward particular...

<https://arxiv.org/abs/2408.10608>

9 CONF

Predicting Estimated Times of Restoration for Electrical Outages Using Longitudinal Tabular Transformers

ArXiv CS.AI 50%

arXiv:2505.00225v2 Announce Type: replace-cross Abstract: Utilities publish Estimated Times of Restoration (ETRs) for customer-facing storm outages, and their accuracy governs whether customers can make sound decisions about food, medical equipment, and relocation. Prior work...

<https://arxiv.org/abs/2505.00225>

00 CONF

Generative AI for Analysts

ArXiv CS.AI 50%

arXiv:2512.19705v2 Announce Type: replace-cross Abstract: We study how generative artificial intelligence (GenAI) reshapes financial analysts' information production. Using the 2023 integration of GenAI into FACTSET as a plausibly exogenous change in AI access, we find that...

<https://arxiv.org/abs/2512.19705>

01 CONF

From Rubrics to Reliable Scores: Evidence-Grounded Text Evaluation with LLM Judges

ArXiv CS.AI 50%

arXiv:2601.08654v3 Announce Type: replace-cross Abstract: Rubric-based text evaluation increasingly relies on large language models (LLMs) as scalable judges, yet frozen black-box models can interpret the same criteria inconsistently, produce score attributions that are...

<https://arxiv.org/abs/2601.08654>

02 CONF

False positive bias in AI-powered speech-based cognitive screening for multilingual English speakers in the UK

ArXiv CS.AI 50%

arXiv:2602.13047v2 Announce Type: replace-cross Abstract: Conversational speech reveals early signs of cognitive decline, including dementia and mild cognitive impairment (MCI). AI models show promise for speech-based screening, yet most research focuses on monolingual groups....

<https://arxiv.org/abs/2602.13047>

03 CONF

Spec-Harness: Measuring and Improving Behavioral Adequacy of LLM-Synthesized Formal Specifications

ArXiv CS.AI 50%

arXiv:2604.00280v2 Announce Type: replace-cross Abstract: Formal specifications play a central role in ensuring software reliability, yet automatically synthesizing high-quality specifications remains difficult and often requires domain expertise. Recent work has applied large...

<https://arxiv.org/abs/2604.00280>

04 CONF

Where is the Mind? Persona Vectors and LLM Individuation

ArXiv CS.AI 50%

arXiv:2604.17031v3 Announce Type: replace-cross Abstract: The individuation problem for large language models asks which entities associated with them, if any, should be identified as minds. We approach this problem through mechanistic interpretability, engaging in particular...

<https://arxiv.org/abs/2604.17031>

05 CONF

Dont Just Teach, Explain! A Gamified 20Q Recommender for Cybersecurity Education

ArXiv CS.AI 50%

arXiv:2604.26964v2 Announce Type: replace-cross Abstract: The escalating complexity of modern cyber threats demands innovative approaches to security education that transcend traditional pedagogical methods. Conventional training paradigms often fail to engage learners...

<https://arxiv.org/abs/2604.26964>

06 CONF

Using Reward Uncertainty to Induce Diverse Behaviour in Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2606.03962v2 Announce Type: replace-cross Abstract: Classical reinforcement learning (RL) typically seeks a deterministic policy that maximizes the expected sum of a scalar reward. Yet, modern applications such as language model fine-tuning or scientific discovery demand...

<https://arxiv.org/abs/2606.03962>

07 CONF

FP8 is All You Need (Part 1): Debunking Hardware FP64 as the HPC Holy Grail (Sep 3rd version)

ArXiv CS.AI 50%

arXiv:2606.06510v4 Announce Type: replace-cross Abstract: We argue that on AI-optimised GPUs of the NVIDIA B300 generation and beyond, the FP8 tensor-core matrix operation, composed through CRT-based Ozaki Scheme II, can serve as the dominant matrix-work substrate for the...

<https://arxiv.org/abs/2606.06510>

08 CONF

Builder, Defender, Breaker: Measurable Independence and Bounded Autonomy When Generative Models Build, Defend and Test Software

ArXiv CS.AI 50%

arXiv:2607.03215v2 Announce Type: replace-cross Abstract: Generative models now write application code, harden and monitor it, and probe it for exploitable flaws, so that one family of models increasingly plays builder, defender and breaker at once. The prevailing view treats...

<https://arxiv.org/abs/2607.03215>

09 CONF

AtlasNLP: A Country-Aware Atlas of Dataset Representation in NLP

ArXiv CS.AI 50%

arXiv:2608.30107v2 Announce Type: replace-cross Abstract: Understanding which countries are represented in NLP datasets is essential for identifying gaps, targeting data collection, measuring progress, and informing AI policy. However, geographic metadata is very rarely...

<https://arxiv.org/abs/2608.30107>

10 CONF

Phase-Aware Spatial-Frequency Fusion for Few-Shot Fine-Grained Image Classification

ArXiv CS.AI 50%

arXiv:2609.03829v2 Announce Type: replace-cross Abstract: Few-shot fine-grained image classification (FSFGIC) aims to classify similar images with limited labeled examples. This work highlights the critical yet underutilized role of phase information in capturing structural...

<https://arxiv.org/abs/2609.03829>