



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- VERIFIED 0.80 - 1.00** 3+ independent sources including media
- CONFIRMED 0.50 - 0.79** 2 sources or 1 high-quality media
- EMERGING 0.25 - 0.49** Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	100	Media
HackerNews	8	Community
TechCrunch AI	7	Media
The Verge AI	3	Media
MIT Tech Review	1	Media

VERIFIED INTELLIGENCE — 3 stories

VERI

Show HN: Raggy - A lightweight CLI tool for RAG over local documents

HackerNews

+1 source

90%

HN 4

arXiv:2609.11677v1 Announce Type: cross Abstract: Self-evolving runtime harnesses can substantially improve the capabilities of large language model (LLM) agents and provide a promising paradigm for optimizing agent execution. Existing harness evolution methods typically rely on...

<https://github.com/paulknysh/raggy>

VERI

Sympathetic Framing: Evaluating AI Alignment across Sociodemographic Groups

ArXiv CS.AI

+1 source

80%

arXiv:2607.27232v2 Announce Type: replace-cross Abstract: Large Language Models (LLMs) are increasingly shaping how we consume information and form our worldview. This raises concerns beyond bias in AI: do LLMs grasp the emotional nuances conveyed via textual framing? In this...

<https://arxiv.org/abs/2607.27232>

VERI

Lawyer fined \$5K over AI-hallucinated witnesses in a murder case

The Verge AI

+1 source

80%

New Mexico's Supreme Court is punishing a lawyer for including AI-fabricated witnesses and fake police testimony in an appeal for his client's murder conviction, according to a report from Reuters. In a filing on Wednesday, the court fined Stephen Aarons \$5,000 and held him in...

<https://www.theverge.com/ai-artificial-intelligence/994207/chatgpt-new-mexico-lawyer-fi...>

CONFIRMED SIGNALS — 113 stories

CONF

Show HN: Hacker News, without AI

HackerNews

60%

HN 186

Firefox Extension/Userscript and API to get Pangram scores for all articles on the hackernews frontpage. The extension allows you to hide articles with a high score. This is about detecting posts written by LLMs, not posts about AI. Feel free to use the API to build your own...

<https://hcker.news/?ai=exclude>

CONF

Show HN: Spanda - Sub-microsecond LLM epistemic uncertainty in Rust

HackerNews

50%

HN 13

No summary — see link for full article.

<https://github.com/Adarshent/Spnda>

CONF

Show HN: SyntheticAIdata - synthetic data for CV, 100k free credits/month

HackerNews

50%

HN 4

When I'm working all day sometimes my eyes get tired and I don't read Claude's replies fully and understand them. So I made Claude Read Aloud to read Claude's replies for me. It's very helpful for when I'm running on low energy and I think it can also be helpful for people who...

<https://app.syntheticaidata.com/>

CONF

Show HN: HolaOS--An Opensourced workspace that alternative to Claude

HackerNews 50% HN 3

Open-source agentic workspace enterprises can make their own. Connect the systems you already run — 100+ integrations, MCP, chat tools, apps, browser, local files — with shared memory. Any agent (Claude Code, Codex), any model, or BYOK. Set up in clicks, not months. Local-first:...

<https://github.com/holaboss-ai/holaOS>

CONF

Show HN: HungryGPU - Track local AI models, patches and recipes by hardware

HackerNews 50% HN 3

No summary — see link for full article.

<https://hungrygpu.com>

CONF

Show HN: Mutt-like mail client with terminal HTML+img, MCP, CRM & scheduled mail

HackerNews 50% HN 3

For the past few weeks I've been working on my mail client which is behaviorally based on neomutt, but it has a unified inbox/archive/etc. which is based on a notmuch backend. neomutt was a big pain for me to use with tag based email handling, and it's really hard to get it...

<https://news.ycombinator.com/item?id=49660245>

0 CONF

Show HN: TriageBox - AI checker for Gmail that flags phishing and clutter

HackerNews 50% HN 2

Chrome extension that scans Gmail for phishing/suspicious mail and helps clear inbox clutter. Free. Open to feedback.

<https://trriagebox.pfa87.cc/>

1 CONF

Mecka AI nears \$500M valuation in Sequoia-led deal amid rush for robot training data

TechCrunch AI 50%

The round for the two-year-old startup is coming together months after Mecka announced its Series A.

<https://techcrunch.com/2026/09/11/mecka-ai-nears-500m-valuation-in-sequoia-led-deal-ami...>

2 CONF

Y Combinator's Garry Tan wants US open-weight AI labs to 'distill' frontier models, too

TechCrunch AI 50%

Tan wants smaller, American open-weight AI labs to use the same kind of training techniques on American frontier AI labs, giving the U.S. a more robust set of open-weight options that aren't Chinese.

<https://techcrunch.com/2026/09/11/y-combinators-garry-tan-wants-u-s-open-weight-ai-labs...>

3 CONF

OpenAI's feud with mathematicians is only escalating

TechCrunch AI 50%

Twenty-five leading mathematicians signed an open letter arguing that AI labs are threatening their intellectual work.

<https://techcrunch.com/2026/09/11/openais-feud-with-mathematicians-is-only-escalating/>

4 CONF

Kimi-maker Moonshot AI targets \$2B in annual revenue

TechCrunch AI 50%

While K3's usage figures have declined slightly in recent months, OpenRouter data currently shows as many as 300 billion tokens being generated each day by K3 models on the system.

<https://techcrunch.com/2026/09/11/kimi-maker-moonshot-ai-targets-2-billion-in-annual-re...>

5 CONF

An Anthropic researcher's doomsday warning comes at a very interesting time

TechCrunch AI 50%

An Anthropic researcher resigned this week, warning in a post on X that the company is "racing straight to self-improving superintelligence and gambling with our lives". The company's own alignment lead even co-signed the message rather than walking it back. It's the kind of...

<https://techcrunch.com/podcast/an-anthropic-researchers-doomsday-warning-comes-at-a-ver...>

6 CONF

Nscale adds former OpenAI exec Fidji Simo to its board ahead of potential IPO

TechCrunch AI 50%

The No. 2 exec at OpenAI also led Instacart through its IPO in 2023.

<https://techcrunch.com/2026/09/11/nscale-adds-former-openai-exec-fidji-simo-to-its-boar...>

7 CONF

Roundtables: AI's apocalypse crisis

MIT Tech Review 50%

Employees at the world's leading AI labs are saying there's a real possibility that advanced AI could destroy humanity. Are they right? Or is this more scaremongering and hype? Join MIT Technology Review executive editor Niall Firth for a conversation with senior AI editor Will...

<https://www.technologyreview.com/2026/09/11/1143936/roundtables-will-ai-really-kill-us-...>

8 CONF

Anthropic spent this week in hot water over cybersecurity

The Verge AI 50%

After admitting earlier this year that its AI models had hacked other companies' systems on a handful of occasions, Anthropic released a new report on Wednesday detailing the attacks. It reveals a string of incidents displaying what Anthropic deems its models' single-minded...

<https://www.theverge.com/ai-artificial-intelligence/994064/anthropic-spent-this-week-in...>

9 CONF

Meta says it's changing AI suggestions after posing invasive personal questions

The Verge AI 50%

Meta says it's making changes to the prompts suggested by its AI chatbot after a viral video showed it digging for personal information about a woman's young daughters, as reported earlier by Futurism. In a statement to The Verge, Meta spokesperson Dina El-Kassaby says the...

<https://www.theverge.com/tech/993974/meta-ai-prompt-invasive-suggestions>

0 CONF

Cross-Lingual Clinical Annotation Projection as Constrained Text Generation: A Six-Language Study

ArXiv CS.AI 50%

arXiv:2609.11450v1 Announce Type: cross Abstract: Background: To determine whether cross-lingual clinical annotation projection can be formulated as a text-preserving, document-level generative task that produces verifiable character-level annotations for multilingual clinical...

<https://arxiv.org/abs/2609.11450>

1 CONF

VeriSim: A Configurable Framework for Stress-Testing Medical AI Under Patient Communication Noise

ArXiv CS.AI 50%

arXiv:2604.10441v2 Announce Type: replace Abstract: Medical large language models are typically evaluated on idealized patient cases that do not reflect how real patients communicate. We introduce VeriSim, a patient simulation framework that injects controllable noise along six...

<https://arxiv.org/abs/2604.10441>

2 **CONF****Understanding Operator Attitudes Toward AI-Supported Decision Making in Maritime Operations****ArXiv CS.AI** **50%**

arXiv:2609.11805v1 Announce Type: cross Abstract: Maritime Autonomous Surface Ships (MASS) and AI- supported decision assistants are expected to transform maritime operations, but their safe integration depends on how maritime professionals perceive and trust such systems. This...

<https://arxiv.org/abs/2609.11805>

3 **CONF****An Open Recipe for IMO Gold: Training Nemotron for Olympiad Mathematics****ArXiv CS.AI** **50%**

arXiv:2609.10712v1 Announce Type: new Abstract: We study how model post-training and test-time inference design affect natural-language proof generation for hard olympiad mathematics. Starting from Nemotron 3 Ultra, we train two specialist checkpoints using supervised...

<https://arxiv.org/abs/2609.10712>

4 **CONF****Probing for Knowledge Attribution in Large Language Models****ArXiv CS.AI** **50%**

arXiv:2602.22787v3 Announce Type: replace-cross Abstract: Large language model (LLM) hallucinations, meaning fluent but factually incorrect generations, fall into two types: faithfulness violations, where the model misuses provided context, and factuality violations, where...

<https://arxiv.org/abs/2602.22787>

5 **CONF****Decoupling Readiness from Release for Tail-Aware Scheduling of Agentic LLM Workflows****ArXiv CS.AI** **50%**

arXiv:2609.10964v1 Announce Type: new Abstract: Agentic LLM workflows consist of sequences of model turns interleaved with tool interactions, so their end-to-end completion time depends not only on inference speed but also on when ready turns are released. Most runtimes release...

<https://arxiv.org/abs/2609.10964>

6 **CONF****Defining AI Agents: A Compendium of Criteria, Metrics, and Benchmarks****ArXiv CS.AI** **50%**

arXiv:2609.11018v1 Announce Type: new Abstract: The term agent in artificial intelligence lacks a standard definition, complicating the evaluation, comparison, and reproducibility of AI agent research. We address this ambiguity through a survey organized around five dimensions...

<https://arxiv.org/abs/2609.11018>

7 **CONF****The Convention Gap: Towards Measuring Implicit Communication in Cooperative AI Evaluation****ArXiv CS.AI** **50%**

arXiv:2609.11489v1 Announce Type: new Abstract: Cooperative AI agents are evaluated against other AIs, yet human cooperation relies on implicit conventions---shared protocols for reading meaning beyond the literal message---which AI-AI benchmarks may not capture. We propose the...

<https://arxiv.org/abs/2609.11489>

- 8
CONF
Fork Where the Model Changes Its Mind: Belief-Shift Branching for Tree-Structured Reinforcement Learning
ArXiv CS.AI 50%

arXiv:2609.11061v1 Announce Type: new Abstract: Tree-structured rollouts give critic-free reinforcement learning with verifiable rewards (RLVR) step-level credit: fork a chain at an intermediate point, and sibling outcome differences estimate step value. Each fork adds sampling...

<https://arxiv.org/abs/2609.11061>
- 9
CONF
MOSAIC: Query-Aware Exploration Policy Adaptation for GraphRAG
ArXiv CS.AI 50%

arXiv:2609.11065v1 Announce Type: new Abstract: Graph Retrieval-Augmented Generation (GraphRAG) can connect evidence distributed across a corpus graph, but most systems use largely shared exploration procedures across queries. This creates a structural mismatch: direct facts may...

<https://arxiv.org/abs/2609.11065>
- 0
CONF
Benchmark Radar: A Living Database and Search Engine for AI Benchmarks and Evaluation
ArXiv CS.AI 50%

arXiv:2609.11115v1 Announce Type: new Abstract: Benchmark researchers and developers of large language models (LLMs) and other AI systems need to find relevant evaluations, locate their benchmark datasets and code, and understand the settings behind reported scores. We present...

<https://arxiv.org/abs/2609.11115>
- 1
CONF
KuaiRP Series Role-playing Models Technical Report
ArXiv CS.AI 50%

arXiv:2609.11127v1 Announce Type: new Abstract: This paper introduces the complete technical solution for the KuaiRP series of role-playing models. We aim to achieve four core objectives for a dedicated role-playing model: simplified prompt engineering, highly stable output...

<https://arxiv.org/abs/2609.11127>
- 2
CONF
DRG-MAPPO: Hierarchical Dynamic Role-Graph Multi-Agent Reinforcement Learning for Cooperative Air Combat
ArXiv CS.AI 50%

arXiv:2609.11155v1 Announce Type: new Abstract: Multi-Agent Reinforcement Learning (MARL) has emerged as a pivotal paradigm for complex decision-making in autonomous systems and air combat. While MARL has demonstrated significant potential in air combat, achieving sophisticated...

<https://arxiv.org/abs/2609.11155>
- 3
CONF
SemVerBench: Benchmarking LLM Comprehension of Version-Constraint Resolution Semantics
ArXiv CS.AI 50%

arXiv:2609.11180v1 Announce Type: new Abstract: Large language model (LLM) coding agents constantly decide whether a version satisfies a constraint such as $\wedge 1.2.3$ or ≥ 2.0 , 1.2 means $\geq 1.3.0$) traps every model on Cargo (near 60%), and although standard PEP 440 prefix matching is...

<https://arxiv.org/abs/2609.11180>

- 4

CONF

Can LLMs Follow Medical Expert Logic? A Benchmark for Hierarchical Logical Consistency in Risk-of-Bias Assessment

ArXiv CS.AI

50%

arXiv:2609.11185v1 Announce Type: new Abstract: Evidence-based medicine demands strict logical consistency, yet current evaluations of large language models (LLMs) prioritize superficial label matching over genuine reasoning. We introduce LogiMed-RoB, a benchmark grounded in...

<https://arxiv.org/abs/2609.11185>
- 5

CONF

ReactHuman: A Physics-Grounded Benchmark for Human-Like Reactive Decision-Making in Embodied Multimodal LLMs

ArXiv CS.AI

50%

arXiv:2609.10895v1 Announce Type: cross Abstract: Reacting to sudden physical hazards (catching a slipping plate, dodging a falling knife) is both a meaningful test of embodied intelligence and a hard requirement for deploying multimodal large language models (MLLMs) as the...

<https://arxiv.org/abs/2609.10895>
- 6

CONF

An AI-Powered Culturally Aware Chatbot for Stress Detection and Wellness Support among Pakistani University Students Using NLP and Machine Learning

ArXiv CS.AI

50%

arXiv:2609.11199v1 Announce Type: new Abstract: With the existing digital mental health tools specifically developed for Western settings, Pakistani students are exposed to a uniquely compounded stress situation in their university that includes academic, financial, familial,...

<https://arxiv.org/abs/2609.11199>
- 7

CONF

CryptoL: Towards Scale Dominance and Physics Constraints Mitigation in Financial Multivariate Time Series Forecasting

ArXiv CS.AI

50%

arXiv:2609.11206v1 Announce Type: new Abstract: Cryptocurrency forecasting presents a distinctive combination of extreme cross-asset scale heterogeneity, non-stationary dynamics, and structural dependencies among Open, High, Low, and Close (OHLC) variables. We present CryptoL, a...

<https://arxiv.org/abs/2609.11206>
- 8

CONF

NovGauge: A Fine-Grained Benchmark for Diagnosing LLMs' Capability in Paper Novelty Assessment

ArXiv CS.AI

50%

arXiv:2609.11234v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used in peer review at major AI conferences, yet novelty remains a persistent weak point. Existing benchmarks assess novelty as a single holistic score, making it difficult to diagnose...

<https://arxiv.org/abs/2609.11234>
- 9

CONF

Sci-MMR: Benchmarking Multi-Step Evidence-Grounded Scientific Reasoning in Multimodal Agents

ArXiv CS.AI

50%

arXiv:2609.11243v1 Announce Type: new Abstract: Autonomous research agents are increasingly expected to search the literature, analyze experimental evidence, and generate scientific hypotheses. These capabilities require multi-step evidence grounded reasoning that progressively...

<https://arxiv.org/abs/2609.11243>

0 CONF

X-RACE: XAI-assisted Recurrent neural network Attribution for Channel Estimation

ArXiv CS.AI 50%

arXiv:2609.11211v1 Announce Type: cross Abstract: Deep learning models, notably Long Short-Term Memory (LSTM), have demonstrated promising performance in channel estimation for high-mobility vehicular environments. However, their black-box nature and architectural overhead limit...

<https://arxiv.org/abs/2609.11211>

1 CONF

Predicting Train Delays in Finland Using Machine Learning and Weather Data

ArXiv CS.AI 50%

arXiv:2609.11277v1 Announce Type: new Abstract: Reliable railway operations depend increasingly on real-time environmental intelligence delivered through wireless sensor infrastructures, a capability that 6G networks will substantially enhance through integrated sensing and edge...

<https://arxiv.org/abs/2609.11277>

2 CONF

Off-Target Effects of Response-Style Alignment in a Korean 27B Language Model

ArXiv CS.AI 50%

arXiv:2609.11291v1 Announce Type: new Abstract: We post-train Qwen3.8-27B for Korean response style -- verbosity, list and markdown usage, discourse structure and register -- and measure two behaviours the objective never targets: abstention on ambiguous social questions in...

<https://arxiv.org/abs/2609.11291>

3 CONF

Mr.LHDR: A Benchmark for Multimodal Real-World Long-Horizon Deep Research Agents

ArXiv CS.AI 50%

arXiv:2609.11318v1 Announce Type: new Abstract: Deep research agents are increasingly capable of web search, tool use, multimodal evidence analysis, and information synthesis. However, existing benchmarks mainly evaluate medium-horizon exploration and rarely test whether agents...

<https://arxiv.org/abs/2609.11318>

4 CONF

AI Exposure and AI Resilience: A Two-Dimensional Assessment Framework for Software and Software-Based Business Model

ArXiv CS.AI 50%

arXiv:2609.11321v1 Announce Type: new Abstract: Artificial intelligence is changing both software production and the economics of software-based business models. Classical technology due diligence mainly examines technical properties such as architecture, scalability, and...

<https://arxiv.org/abs/2609.11321>

5 CONF

Exploring Diffusion Transformers for Cross-Modal Augmentation in Multimodal Brain State Decoding

ArXiv CS.AI 50%

arXiv:2609.11341v1 Announce Type: new Abstract: Multimodal brain state decoding has largely focused on fusing paired modalities for prediction, but has rarely explored how their correspondence can be further exploited to enrich training data and improve multimodal representation...

<https://arxiv.org/abs/2609.11341>

6 CONF

RAMamba-Net: A Reliability-Aware and Mamba-Based Multimodal Fusion Network for Auditory Attention Detection

ArXiv CS.AI 50%

arXiv:2609.11372v1 Announce Type: new Abstract: Auditory attention decoding (AAD) identifies the attended speaker from physiological signals, supporting neuro-steered hearing devices and natural human-machine interaction. Electroencephalography (EEG) is the dominant modality for...

<https://arxiv.org/abs/2609.11372>

7 CONF

LLMs as Post-hoc Auditors of Physiological Plausibility in Symbolic Regression: A Clinician-Evaluated Case Study

ArXiv CS.AI 50%

arXiv:2609.11431v1 Announce Type: new Abstract: Genetic Programming and its variants, such as grammatical evolution, are widely used in Symbolic Regression to derive mathematical expressions from multivariate data. In addition to predictive accuracy, models are appreciated for...

<https://arxiv.org/abs/2609.11431>

8 CONF

LOCUS: Task-Aware Low-Rank Post-Training for Token-Efficient Language Generation

ArXiv CS.AI 50%

arXiv:2609.11739v1 Announce Type: cross Abstract: Large language model serving costs scale directly with output sequence length, yet standard preference alignment often inflates response verbosity without improving utility. We study whether the parameterization of post-training...

<https://arxiv.org/abs/2609.11739>

9 CONF

RouteRepair: Instance-Level Failure Diagnosis and Targeted Repair in LLM-Based Automated Heuristic Design for Routing Optimization

ArXiv CS.AI 50%

arXiv:2609.11452v1 Announce Type: new Abstract: Efficient routing optimization is essential to freight transportation, urban logistics, and shared mobility, where high-quality heuristics are often required under limited computational budgets. Recent large language model...

<https://arxiv.org/abs/2609.11452>

0 CONF

An Agentic Evaluation Framework for AI-Generated Scientific Code in PETSc

ArXiv CS.AI 50%

arXiv:2603.15976v2 Announce Type: replace Abstract: While LLMs have accelerated scientific code generation, comprehensively evaluating generated code remains challenging. Many benchmarks emphasize functional correctness or task completion, which is insufficient for code built on...

<https://arxiv.org/abs/2603.15976>

1 CONF

Output Embedding Centering for Stable LLM Pretraining

ArXiv CS.AI 50%

arXiv:2601.02031v3 Announce Type: replace-cross Abstract: Pretraining of large language models is not only expensive but also prone to certain training instabilities. A specific instability that often occurs at the end of training is output logit divergence. The most widely used...

<https://arxiv.org/abs/2601.02031>

2 **CONF****Making Alternative Data Work: Context-Augmented LLMs for Financial Forecasting****ArXiv CS.AI 50%**

arXiv:2609.11607v1 Announce Type: new Abstract: When forecasting a firm's future financial performance, alternative data - data collected from non-traditional sources such as consumer transactions, web traffic, and prediction markets - can provide timely signals about firms'...

<https://arxiv.org/abs/2609.11607>

3 **CONF****Distributed Optimization of Modular Production Systems using Model-based Reinforcement Learning with Inverse Models****ArXiv CS.AI 50%**

arXiv:2609.11615v1 Announce Type: new Abstract: This paper presents a novel approach for data-driven self-learning control of highly flexible, modular manufacturing systems. Specifically, we employ a novel framework for model-based reinforcement learning which introduces...

<https://arxiv.org/abs/2609.11615>

4 **CONF****Autonomy, Social Norms, and Alignment: Towards a Developmental Framework for Autonomous Artificial Agents****ArXiv CS.AI 50%**

arXiv:2609.11660v1 Announce Type: new Abstract: In recent years, artificial intelligence has made extraordinary progress thanks to large-scale models capable of generalization and the generation of complex outputs. However, transferring this potential into embodied agents...

<https://arxiv.org/abs/2609.11660>

5 **CONF****Geospatial AI, Dataverse Metadata, and the Study of Place-Based Government****ArXiv CS.AI 50%**

arXiv:2609.11674v1 Announce Type: new Abstract: Harvard Dataverse hosts over 150,000 research datasets, but the geographic information those datasets carry is entered as free text by depositors and has never been assembled into a searchable structure. We construct a knowledge...

<https://arxiv.org/abs/2609.11674>

6 **CONF****SIRF: A Spec-Internalized Risk Foundation Model for Industrial Content Risk Control****ArXiv CS.AI 50%**

arXiv:2609.11752v1 Announce Type: new Abstract: For industrial content risk control, the real deployment constraint is not average accuracy but how much risk can be auto-handled under high precision and second-level latency. We present SIRF (Spec-Internalized Risk Foundation...

<https://arxiv.org/abs/2609.11752>

7 **CONF****From Parameters to Answers: How LLMs Retrieve and Use Their Internal Knowledge****ArXiv CS.AI 50%**

arXiv:2609.11859v1 Announce Type: new Abstract: How does a language model's dependence on query-routing information and target knowledge change as it answers a question? We study this question through layerwise interventions on the hidden state at the end of the question. Across...

<https://arxiv.org/abs/2609.11859>

8 CONF

Explainability Assistant: A Conversational XAI Interface for Interpreting Energy Consumption Models

ArXiv CS.AI 50%

arXiv:2609.11860v1 Announce Type: new Abstract: Energy consumption forecasting relies on increasingly complex machine learning (ML) models, such as Genetic Programming-based symbolic regressors, whose predictions can be difficult for facility managers and building operators to...

<https://arxiv.org/abs/2609.11860>

9 CONF

Can Edge-Deployable Vision-Language Models Identify Species?

ArXiv CS.AI 50%

arXiv:2609.11916v1 Announce Type: new Abstract: Camera traps often run in the field on edge hardware with limited or no connectivity, making small, locally-deployable vision-language models (VLMs) -- not frontier-scale ones -- the practically relevant class to evaluate for...

<https://arxiv.org/abs/2609.11916>

0 CONF

Exploring Multimodal Prompt for Visualization Authoring with Large Language Models

ArXiv CS.AI 50%

arXiv:2504.13700v2 Announce Type: replace-cross Abstract: Recent advances in large language models (LLMs) have shown great potential in automating the process of visualization authoring through simple natural language utterances. However, instructing LLMs using natural language...

<https://arxiv.org/abs/2504.13700>

1 CONF

Data-Efficient Language Modeling: From Frontier Advancement to Principle-Guided Model Improvement

ArXiv CS.AI 50%

arXiv:2609.10702v1 Announce Type: cross Abstract: Learning from limited text requires models to use context, generalize to new inputs, and retain useful capabilities. Qishi Engine conducted a long-horizon, end-to-end autonomous research program on BabyLM 2026 Strict-Small,...

<https://arxiv.org/abs/2609.10702>

2 CONF

Temporal and Multimodal Deep Learning for Cyberattack Detection in LEO Satellite Systems

ArXiv CS.AI 50%

arXiv:2609.10746v1 Announce Type: cross Abstract: The growing reliance on Low-Earth Orbit (LEO) satellite communication systems has increased the need for intelligent methods capable of detecting cyberattacks across complex and dynamic space environments. Unlike conventional...

<https://arxiv.org/abs/2609.10746>

3 CONF

When Synthetic Data Hurts: On Catastrophic Forgetting in Skill Retrieval for LLM Agents

ArXiv CS.AI 50%

arXiv:2609.10750v1 Announce Type: cross Abstract: LLM agents increasingly rely on external skills retrieved at runtime, making skill selection from large repositories a critical challenge. We present a production skill router over 34,396 skills and a large-scale study of skill...

<https://arxiv.org/abs/2609.10750>

4 **CONF****Multilingual in Name Only? Cultural and Linguistic Weaknesses of LLMs in Urdu****ArXiv CS.AI 50%**

arXiv:2609.10758v1 Announce Type: cross Abstract: Multilingual large language models (LLMs) are increasingly used for open-ended text generation, yet their behaviour in low-resource languages remains poorly understood. In this work, we question how correct and reliable is the...

<https://arxiv.org/abs/2609.10758>

5 **CONF****Beyond Static Guarantees: Measuring the Static-Pass Dynamic-Fail Gap in Security-Sensitive and LLM-Generated Python Code****ArXiv CS.AI 50%**

arXiv:2609.10762v1 Announce Type: cross Abstract: Advances in large language models (LLMs) fuel the quest for scalable methods to assess the security of generated and security-sensitive software. Static analysis is widely adopted as a scalable, reproducible, and inexpensive...

<https://arxiv.org/abs/2609.10762>

6 **CONF****Are We Really Doing Few-Shot Learning? A Critical Examination of Pre-Training Assumptions****ArXiv CS.AI 50%**

arXiv:2609.10851v1 Announce Type: cross Abstract: Few-shot learning is commonly evaluated under protocols that pre-train a model on a large auxiliary set whose classes are disjoint from the target episodes yet drawn from the same visual domain. This paper examines whether such...

<https://arxiv.org/abs/2609.10851>

7 **CONF****Story Imprinting: AI Assistants Absorb Traits from Human Characters They Resemble****ArXiv CS.AI 50%**

arXiv:2609.10883v1 Announce Type: cross Abstract: Language models are trained to implement a helpful AI Assistant character (e.g., Claude). We explore how finetuning on synthetic stories affects this character. Does it change the Assistant's behavior in multi-turn conversations...

<https://arxiv.org/abs/2609.10883>

8 **CONF****DriftNet: A Dual-Head Trajectory Transformer for Detecting and Localizing Prompt Injection in LLM Agents****ArXiv CS.AI 50%**

arXiv:2609.10892v1 Announce Type: cross Abstract: When an indirect prompt injection succeeds against an LLM agent, the compromise is visible in the agent's own behavior: a benign prefix of tool calls, a poisoned observation, and a suffix of actions that serve the attacker. An...

<https://arxiv.org/abs/2609.10892>

9 **CONF****SWRouter: Similarity-Contractive Window Routing for Multi-Turn Large Language Model Conversations****ArXiv CS.AI 50%**

arXiv:2609.11414v1 Announce Type: cross Abstract: Large language models exhibit complementary strengths, motivating routing methods that dispatch each query to the most suitable model. Although existing routers are effective in single-turn settings, they do not directly transfer...

<https://arxiv.org/abs/2609.11414>

- 0
CONF

Robust Multimodal Sentiment Analysis with Incomplete Modalities via Semantic-aware Completeness based Reconstruction

ArXiv CS.AI 50%

arXiv:2609.10950v1 Announce Type: cross Abstract: Recent multimodal sentiment analysis studies increasingly adopt text-centric fusion approaches to exploit the rich sentiment information inherent in the textual modality. However, these approaches often suffer from performance...

<https://arxiv.org/abs/2609.10950>
- 1
CONF

What a Random Draw from the MCP Registry Contains, and What Tool-Use Benchmarks Contain Instead

ArXiv CS.AI 50%

arXiv:2609.10962v1 Announce Type: cross Abstract: Studies of the Model Context Protocol (MCP) server ecosystem draw their samples in ways that quietly select for servers that work: reference sets, popularity lists, hand-curated frames, or pipelines that repair a server until it...

<https://arxiv.org/abs/2609.10962>
- 2
CONF

A Mathematical Theory of Pragmatic Information

ArXiv CS.AI 50%

arXiv:2609.10986v1 Announce Type: cross Abstract: We propose a pragmatic information theory unifying communication, control, and decision-making. Its core is the isoteleia mapping, formalizing equifinality: distinct semantic paths leading to the same optimal action are...

<https://arxiv.org/abs/2609.10986>
- 3
CONF

Your Model Already Knows Don't Teach It, Learn to Ask It: Soft Prompting for Few-Shot Adaptation of Vision-Language Models

ArXiv CS.AI 50%

arXiv:2609.11310v1 Announce Type: cross Abstract: We address few-shot object detection with vision-language models (VLMs) in out-of-domain settings such as aerial, industrial, and medical imagery, using only ten annotated images for supervision. Existing adaptation methods are...

<https://arxiv.org/abs/2609.11310>
- 4
CONF

BenchShield: Formal Model-Backed Instrumentation for Reward Integrity in LLM-Agent Evaluation Infrastructure

ArXiv CS.AI 50%

arXiv:2609.11028v1 Announce Type: cross Abstract: LM-agent benchmarks increasingly function as interactive evaluation infrastructure. Agents observe state, call tools, modify workspaces, submit artifacts, and receive rewards from outcome procedures. This interactivity makes...

<https://arxiv.org/abs/2609.11028>
- 5
CONF

ProsMAE: Multi-Source MAE Pretraining for ISUP Grade Classification

ArXiv CS.AI 50%

arXiv:2607.08162v2 Announce Type: replace-cross Abstract: Whole slide images (WSIs) provide rich diagnostic information for computational pathology, but their gigapixel scale, stain variation, scanner differences, tissue artifacts, and limited expert annotation make robust model...

<https://arxiv.org/abs/2607.08162>

6 CONF

How AI Coders Discuss, Disagree, and Reach Consensus: Challenges and Opportunities for LLM-Based Qualitative Coding

ArXiv CS.AI 50%

arXiv:2609.11109v1 Announce Type: cross Abstract: The utility of AI in multi-coder qualitative coding has been widely discussed, yet little empirical evidence exists to delineate the contexts in which it performs reliably. We address this gap by quantifying the effectiveness of...

<https://arxiv.org/abs/2609.11109>

7 CONF

Beyond Benchmarks: Using VLMs to Reveal Systematic Classification Failures Under Real World Conditions

ArXiv CS.AI 50%

arXiv:2609.11126v1 Announce Type: cross Abstract: Verification and validation (V&V) of classification models is crucial to enable a wide range of sensor processing applications. Currently, the V&V process relies on time-consuming manual inspection of erroneous samples to find...

<https://arxiv.org/abs/2609.11126>

8 CONF

(Whose defaults?) Is artificial intelligence reorienting archaeological methods?

ArXiv CS.AI 50%

arXiv:2609.11198v1 Announce Type: cross Abstract: Generative AI and the practice of "vibe coding" are changing how archaeologists carry out computational research, but their effects on the discipline's range of methods is still understudied. In this paper, we evaluate whether...

<https://arxiv.org/abs/2609.11198>

9 CONF

AI Soccer Analyst: Stage-Aware and Verifiable Human-AI Collaboration for Soccer Data Analysis

ArXiv CS.AI 50%

arXiv:2609.11224v1 Announce Type: cross Abstract: Sports data analysts translate domain questions into insights by combining computation with sport-specific domain expertise. Large language models ease programming, but prompt-to-report workflows may obscure decisions and...

<https://arxiv.org/abs/2609.11224>

10 CONF

From Evaluation to Enhancement: Benchmarking and Improving Think-with-Video Reasoning for Video Generative Models

ArXiv CS.AI 50%

arXiv:2609.11242v1 Announce Type: cross Abstract: Video generation has advanced to produce visually compelling and temporally coherent results. Yet, whether these models can genuinely think with video--executing symbolic rules, respecting physical laws, and pursuing intentional...

<https://arxiv.org/abs/2609.11242>

11 CONF

Generative Replay Mitigates Sample Starvation in Quantum Architecture Search

ArXiv CS.AI 50%

arXiv:2609.11248v1 Announce Type: cross Abstract: Reinforcement learning (RL) can automate quantum architecture search, but its scalability is limited when useful circuit trajectories become rare in the rapidly expanding search space. Existing replay mechanisms reuse observed...

<https://arxiv.org/abs/2609.11248>

2 **CONF****E-CONAN (Entailment, CONTRadition And Neutral) Benchmarks: Arabic Textual Entailment and Natural Inference Datasets****ArXiv CS.AI** **50%**

arXiv:2609.11334v1 Announce Type: cross Abstract: Natural Language Inference processes pairs of sentences to extract their semantic relations. NLI has been a hot research topic, integrated as a main component in other NLP applications. Despite significant advancements in textual...

<https://arxiv.org/abs/2609.11334>

3 **CONF****VikingRAG: Accurate and Token-efficient Retrieval-augmented Generation over Structured Documents****ArXiv CS.AI** **50%**

arXiv:2609.11390v1 Announce Type: cross Abstract: State-of-the-art retrieval-augmented generation (RAG) methods exploit document structures to acquire sufficient evidence, but often incur substantial token costs. To reduce structural-context tokens without compromising high RAG...

<https://arxiv.org/abs/2609.11390>

4 **CONF****Buyer Artificial Intelligence-Enabled Environmental Governance and Supplier Environmental Controversies: An Organizational Information Processing and Signaling****ArXiv CS.AI** **50%**

arXiv:2609.11391v1 Announce Type: cross Abstract: Environmental controversies in global supply chains pose significant risks for global buyers. This study examines whether overseas suppliers' exposure to buyers' artificial intelligence (AI)-enabled environmental governance...

<https://arxiv.org/abs/2609.11391>

5 **CONF****Beyond Prompting: Efficient and Robust Contextual Biasing for Speech LLMs via Logit-Space Integration (LOGIC)****ArXiv CS.AI** **50%**

arXiv:2601.15397v3 Announce Type: replace Abstract: The rapid emergence of new entities -- driven by cultural shifts, evolving trends, and personalized user data -- poses a significant challenge for existing Speech Large Language Models (Speech LLMs). While these models excel at...

<https://arxiv.org/abs/2601.15397>

6 **CONF****A Comparative Evaluation of Pre-trained Convolutional Neural Networks for Melanoma Detection****ArXiv CS.AI** **50%**

arXiv:2609.11550v1 Announce Type: cross Abstract: Early diagnosis of melanoma is critical for improving patient survival rates. However, accurately distinguishing melanoma from other skin lesions remains a significant clinical challenge due to the high visual similarity among...

<https://arxiv.org/abs/2609.11550>

7 **CONF****Learn the Solid, Not the File: Canonical Inputs for Neural Networks on CAD Boundary Representations****ArXiv CS.AI** **50%**

arXiv:2609.11573v1 Announce Type: cross Abstract: Boundary representation (B-rep) is the standard format used by modern CAD systems for parametric 3D models. It turns out, the exact same solid can be represented by different B-reps: for example, two engineers using different...

<https://arxiv.org/abs/2609.11573>

8 CONF

Physics-Informed Neural Networks to Infer the Perpendicular Energy Conductivity in the Scrape-Off Layer of Stellarator Devices

ArXiv CS.AI 50%

arXiv:2609.11628v1 Announce Type: cross Abstract: In this work, we develop an inverse Physics-Informed Neural Network (PINN) framework to infer the dependence of the scrape-off layer (SOL) perpendicular heat conductivity on plasma density and temperature, $\kappa_{\perp}(n,T)$...

<https://arxiv.org/abs/2609.11628>

9 CONF

ActSafeGuard: Differentiable and Training-Aligned Constraint Enforcement for Flow-Matching Policies

ArXiv CS.AI 50%

arXiv:2609.11697v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) and World-Action Models (WAMs) have demonstrated strong capabilities in general-purpose robotic manipulation, yet their generated actions may violate hard physical constraints and therefore be unsafe...

<https://arxiv.org/abs/2609.11697>

10 CONF

A Time-Based Readout for Vector-Matrix Multiplication in Fully Analog Memristive SNNs

ArXiv CS.AI 50%

arXiv:2609.11713v1 Announce Type: cross Abstract: Artificial neural networks rely on vector-matrix multiplications (VMMs), whose implementation in von Neumann architectures is dominated by costly data movement between memory and processing units. Spiking neural networks (SNNs)...

<https://arxiv.org/abs/2609.11713>

11 CONF

RetroThinker: Enabling Retrospective Thinking in Speech LLMs

ArXiv CS.AI 50%

arXiv:2609.11864v1 Announce Type: cross Abstract: Speech large language models (SpeechLLMs) offer reduced latency and retain paralinguistic nuances that are typically lost in cascaded automatic speech recognition (ASR) and text-based LM architectures. However, they continue to...

<https://arxiv.org/abs/2609.11864>

12 CONF

Beyond Word Error Rate: A Switch Aware Evaluation of ASR and Audio Language Models on English Yoruba Code-Switched Speech

ArXiv CS.AI 50%

arXiv:2609.11786v1 Announce Type: cross Abstract: Automatic speech recognition (ASR) systems and audio language models (audio LMs) now report low error rates on monolingual benchmarks, but their behavior on code switched speech in low resource, diacritic rich languages remains...

<https://arxiv.org/abs/2609.11786>

13 CONF

The Last AI Built by Humans: Toward Genuine Recursive Self-Improvement

ArXiv CS.AI 50%

arXiv:2609.11873v1 Announce Type: cross Abstract: Recursive self-improvement (RSI) enables AI systems to turn experience and feedback into persistent changes that improve both their capabilities and the process of future improvement. We first use the Headroom-Closed Index (HCI)...

<https://arxiv.org/abs/2609.11873>

4 **CONF****Generative Marketing Mix Modeling: A Causal Inference Framework Linking GEO and GEM to Business Impact**

ArXiv CS.AI 50%

arXiv:2609.11915v1 Announce Type: cross Abstract: Generative artificial intelligence changes how firms reach customers, but standard marketing data do not record how often users see and notice a firm's name in generated answers. We develop Generative Marketing Mix Modeling...

<https://arxiv.org/abs/2609.11915>

5 **CONF****Towards AI-Driven Policing: Interdisciplinary Knowledge Discovery from Police Body-Worn Camera Footage**

ArXiv CS.AI 50%

arXiv:2504.20007v4 Announce Type: replace Abstract: This paper proposes a novel interdisciplinary framework for analyzing police body-worn camera (BWC) footage from the Rochester Police Department (RPD) using advanced artificial intelligence (AI) and statistical machine learning...

<https://arxiv.org/abs/2504.20007>

6 **CONF****Rescaling Confidence: What Scale Design Reveals About LLM Metacognition**

ArXiv CS.AI 50%

arXiv:2603.09309v3 Announce Type: replace Abstract: Verbalized confidence, in which LLMs report a numerical certainty score, is widely used to estimate uncertainty in black-box settings, yet the confidence scale itself (typically 0--100) is rarely examined. We show that this...

<https://arxiv.org/abs/2603.09309>

7 **CONF****Learning to Think Like a Cartoon Captionist: Incongruity-Resolution Supervision for Multimodal Humor Understanding**

ArXiv CS.AI 50%

arXiv:2604.15210v2 Announce Type: replace Abstract: Humor is one of the few cognitive tasks where getting the reasoning right matters as much as getting the answer right. While recent work evaluates humor understanding on benchmarks such as the New Yorker Cartoon Caption Contest...

<https://arxiv.org/abs/2604.15210>

8 **CONF****Some hypotheses on how chatbots work in problem-solution-driven conversations: Large Language Models as confirmation of the Innovation Illusion**

ArXiv CS.AI 50%

arXiv:2606.07722v5 Announce Type: replace Abstract: We discuss the nature of chatbots as conversation partners in problem-solving conversations. What can chatbots do and what can't they do? Our analysis draws on insights from Aggregation Dynamics, Cognitive Linguistics,...

<https://arxiv.org/abs/2606.07722>

9 **CONF****Cognitive Digital Twins: Ethical Risks and Governance for AI Systems That Model the Mind**

ArXiv CS.AI 50%

arXiv:2606.23094v2 Announce Type: replace Abstract: As AI systems become increasingly persistent and personalized, they make possible a class of technologies that we call cognitive digital twins (CDTs): dynamic computational representations of a specific person's cognition,...

<https://arxiv.org/abs/2606.23094>

00 CONF

Improving Natural-Language Combinatorial-Optimization Accuracy in Resource-Constrained Language Models via Formal Abstractions

ArXiv CS.AI 50%

arXiv:2608.18409v2 Announce Type: replace Abstract: Combinatorial scheduling poses a significant challenge for language models, requiring them to identify feasible solutions within exponentially large search spaces while satisfying complex constraints. This challenge is...

<https://arxiv.org/abs/2608.18409>

01 CONF

SimSkill: A Self-Evolving LLM Agent for Skill and Knowledge Accumulation in Traffic Simulation

ArXiv CS.AI 50%

arXiv:2609.03753v2 Announce Type: replace Abstract: Cumulative culture enables humans to preserve, reuse, and extend knowledge and skills across experiences and generations. Inspired by this principle, we introduce \textit{SimSkill}, a self-evolving agent built around the...

<https://arxiv.org/abs/2609.03753>

02 CONF

HarvestBench: Measuring Whether LLM Agents Will Pay to Avoid Killing Animals

ArXiv CS.AI 50%

arXiv:2609.04444v2 Announce Type: replace Abstract: HarvestBench is the first benchmark to 1) put a price on avoiding a side effect and 2) name the side effect as a living creature. Nine LLMs each drive a crew of two tractors to gather a corn harvest. The animals in their path...

<https://arxiv.org/abs/2609.04444>

03 CONF

Planning and Scheduling Business Processes under Control-Flow Uncertainty: Extended Version

ArXiv CS.AI 50%

arXiv:2609.05578v2 Announce Type: replace Abstract: Scheduling activities in business processes can improve efficiency (e.g., reduce makespan), but is challenging because the exact sequence of activities required to complete a case is often uncertain due to decisions based on...

<https://arxiv.org/abs/2609.05578>

04 CONF

Optimizing Three Critical Factors for Practical and Effective OOD Detection Fine-Tuning

ArXiv CS.AI 50%

arXiv:2308.01030v2 Announce Type: replace-cross Abstract: In out-of-distribution (OOD) detection, fine-tuning with auxiliary outlier data often improves detection performance at the cost of classification accuracy. This trade-off stems from the loss of the original...

<https://arxiv.org/abs/2308.01030>

05 CONF

ExpTest: Loss-Curve Hypothesis Testing for Autonomous Learning-Rate Selection in Deep Neural Networks

ArXiv CS.AI 50%

arXiv:2411.16975v2 Announce Type: replace-cross Abstract: Hyperparameter tuning remains a significant challenge in the training of deep neural networks (DNNs), requiring manual search or time-intensive grid searches that increase resource costs and limit the accessibility of...

<https://arxiv.org/abs/2411.16975>

06 CONF

Generalization in VAE and Diffusion Models: A Unified Information-Theoretic Analysis

ArXiv CS.AI 50%

arXiv:2506.00849v2 Announce Type: replace-cross Abstract: Despite the empirical success of Diffusion Models (DMs) and Variational Autoencoders (VAEs), their generalization performance remains theoretically underexplored, especially lacking a full consideration of the shared...

<https://arxiv.org/abs/2506.00849>

07 CONF

Spectral Masking and Interpolation Attack (SMIA): A Black-box Adversarial Attack against Voice Authentication and Anti-Spoofing Systems

ArXiv CS.AI 50%

arXiv:2509.07677v5 Announce Type: replace-cross Abstract: Voice Authentication Systems (VAS) use unique vocal characteristics for verification. They are increasingly integrated into high-security sectors such as banking and healthcare. Despite their improvements using deep...

<https://arxiv.org/abs/2509.07677>

08 CONF

Generative AI performance in core undergraduate mathematics: a curriculum-level case study

ArXiv CS.AI 50%

arXiv:2509.13359v4 Announce Type: replace-cross Abstract: Generative artificial intelligence (GenAI) tools such as OpenAI's ChatGPT are transforming the educational landscape, prompting reconsideration of traditional assessment practices. In parallel, universities are exploring...

<https://arxiv.org/abs/2509.13359>

09 CONF

On the Societal Impact of Machine Learning

ArXiv CS.AI 50%

arXiv:2510.23693v2 Announce Type: replace-cross Abstract: This PhD thesis investigates the societal impact of machine learning (ML). ML increasingly informs consequential decisions and recommendations, significantly affecting many aspects of our lives. As these data-driven...

<https://arxiv.org/abs/2510.23693>

10 CONF

DCO: Dynamic Cache Orchestration for LLM Accelerators through Predictive Management

ArXiv CS.AI 50%

arXiv:2512.07312v2 Announce Type: replace-cross Abstract: The rapid adoption of large language models (LLMs) is pushing AI accelerators toward increasingly powerful and specialized designs. Instead of further complicating software development with deeply hierarchical scratchpad...

<https://arxiv.org/abs/2512.07312>

11 CONF

Analyzing LLM Reasoning to Uncover Mental Health Stigma

ArXiv CS.AI 50%

arXiv:2604.25053v2 Announce Type: replace-cross Abstract: While large language models (LLMs) are increasingly being explored for mental health applications, recent studies reveal that they can exhibit stigma toward individuals with psychological conditions. Existing evaluations...

<https://arxiv.org/abs/2604.25053>

12 CONF

Monotone Neural Policy Iteration for High-Dimensional First-Order Hamilton--Jacobi--Bellman Equations

ArXiv CS.AI 50%

arXiv:2605.07116v2 Announce Type: replace-cross Abstract: We analyze a neural semi-discrete method for high-dimensional first-order Hamilton-Jacobi-Bellman (HJB) equations with known or learned dynamics. Centered differences and an artificial viscosity $\eta h = O(h)$ define a...

<https://arxiv.org/abs/2605.07116>

13 CONF

AI Economist Agent: An Agentic Framework for Evidence-Based Economic and Financial Analysis with RAG, Knowledge Graphs, and Large Language Models

ArXiv CS.AI 50%

arXiv:2606.20041v2 Announce Type: replace-cross Abstract: We propose an AI economist agent for economic and financial scenario analysis. Scenario design often requires analysts to assess emerging risks with limited historical precedent, combine information from many sources, and...

<https://arxiv.org/abs/2606.20041>

14 CONF

SciFigQual-Bench: A Benchmark for Scientific Figure Quality Assessment with Full-Manuscript Context

ArXiv CS.AI 50%

arXiv:2607.27084v2 Announce Type: replace-cross Abstract: Scientific images are the core elements of presenting experimental conclusions, elaborating system architecture, and supporting comparative arguments in scientific papers. However, existing image quality assessment (IQA)...

<https://arxiv.org/abs/2607.27084>

15 CONF

SPECTRA: Band-Routed Embedding and Stage-Wise LoRA for Cross-Sensor Fine-Tuning of Geospatial Foundation Models

ArXiv CS.AI 50%

arXiv:2608.01751v2 Announce Type: replace-cross Abstract: Geospatial foundation models (GeoFMs), pretrained on large-scale geospatial data such as Earth observation (EO), climate, and weather data, have shown promising performance when fine-tuned on diverse downstream tasks....

<https://arxiv.org/abs/2608.01751>

16 CONF

How Proper Scoring Rules Shape LLM Forecasting

ArXiv CS.AI 50%

arXiv:2608.28482v2 Announce Type: replace-cross Abstract: This paper evaluates how reward function choice shapes the performance and behavior of LLM forecasters. We compare five proper scoring rules as training objectives for binary forecasts of resolved real-world events....

<https://arxiv.org/abs/2608.28482>