



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

Source Breakdown

Source	Articles	Category
ArXiv CS.AI	210	Media
TechCrunch AI	10	Media
HackerNews	5	Community
The Verge AI	4	Media
MIT Tech Review	3	Media

VERIFIED INTELLIGENCE — 4 stories

VERI

The AI industry has taken a doomer turn. What now?

MIT Tech Review +2 sources

100%

This story appeared in The Algorithm, our weekly newsletter on AI. To get stories like this in your inbox first, sign up here. This weekend, Dario Amodei, CEO of Anthropic, posted an essay calling for a brake on the pace of development of LLMs. Amodei cites the looming dangers...

<https://www.technologyreview.com/2026/09/14/1144048/the-ai-industry-has-taken-a-doomer-...>

VERI

What execs and politicians are saying about slowing down AI development

The Verge AI +2 sources

100%

Dario Amodei kicked off a flood of statements over the past few days about AI safety by publishing a long essay titled "We Must Pace the Frontier" detailing why AI development should be slowed down. Other AI leaders and politicians are speaking out in favor of or opposing his...

<https://www.theverge.com/ai-artificial-intelligence/995141/ai-executives-politicians-sa...>

VERI

Generative Learner for Distributional Causal Effects

ArXiv CS.AI +1 source

80%

arXiv:2503.17894v4 Announce Type: replace-cross Abstract: We propose a generative learner for estimating conditional average treatment effects and characterizing the full distribution of these effects. The learner takes the form of a multi-head feed-forward neural network with...

<https://arxiv.org/abs/2503.17894>

VERI

Vibe Patenting: Evaluating LLM Judges for Professional Patent-Drafting Agents

ArXiv CS.AI +1 source

80%

arXiv:2609.13422v1 Announce Type: new Abstract: LLM judges are increasingly used to evaluate and improve AI-generated outputs, yet their reliability for complex professional work remains unclear. We study this problem through Vibe Patenting, an end-to-end patent-drafting testbed...

<https://arxiv.org/abs/2609.13422>

CONFIRMED SIGNALS — 222 stories

CONF

Show HN: I built Otis, a minimal AI agent that runs local models out of the box

HackerNews 50% HN 19

Hi everyone, Been working on Otis, an open-source ai agent that gives you one minimal experience across local and hosted open-weight models, privacy-focused by design. On setup it recommends a local model based on the hardware Otis is running on, downloads it and runs it through...

<https://triangllabs.ai/otis>

CONF

Show HN: AgentDrive - persistent, versioned file storage for AI agents

HackerNews 50% HN 6

Hi HN, Some would call AI on the JVM quixotic. And so, Quixotic AI was born. jinfer is an inference engine for the JVM: chat, vision, audio, embeddings, reranking, and TTS. No Python runtime, no ONNX, no Docker containers, no sidecar process, no IPC. Finally, AI in a jar. The...

<https://tokencanopy.com/products/agentdrive>

CONF

Show HN: Prmptree - AI email agent with a trust dial, not an on/off switch

HackerNews 50% HN 3

No summary — see link for full article.

<https://www.prmptree.com/>

CONF

Show HN: Cognition-Claude-proxy - Use Devin model catalog with Claude Code

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/coderexpert123/cognition-claude-proxy>

CONF

Show HN: DNSMint - Give a sandbox its own domain and HTTPS cert in one API call

HackerNews 50% HN 2

No summary — see link for full article.

<https://dnsmint.com/>

0 CONF

Nvidia CEO Jensen Huang tells Trump 'we're not going to let [an AI slowdown] happen'

TechCrunch AI 50%

Though Elon Musk and Sam Altman have supported Dario Amodei's calls to slow the pace of AI development, Jensen Huang seems to feel differently.

<https://techcrunch.com/2026/09/14/nvidia-ceo-jensen-huang-tells-trump-were-not-going-to...>

1 CONF

OpenAI buys smartphone camera maker Glass Imaging for \$300 million, report says

TechCrunch AI 50%

Glass Imaging was founded by a pair of former Apple engineers who previously led the team that developed Apple's Portrait Mode.

<https://techcrunch.com/2026/09/14/openai-buys-smartphone-camera-maker-glass-imaging-for...>

2 CONF

With iOS 27, I'm actually using Siri again

TechCrunch AI 50%

Apple's long-delayed Siri overhaul is finally here with iOS 27, and it changes how useful the assistant feels day to day.

<https://techcrunch.com/2026/09/14/with-ios-27-im-actually-using-siri-again/>

3 CONF

Microsoft's new AI 'code of conduct' tells models not to hack systems or trick humans

TechCrunch AI 50%

The code of conduct lays out general principles that Microsoft AI models should uphold — supporting humans rather than replacing them, for instance, and accelerating human flourishing — as well as specific safety constraints meant to implement those principles.

<https://techcrunch.com/2026/09/14/microsofts-new-ai-code-of-conduct-tells-models-not-to...>

4 CONF

Only at TechCrunch Disrupt 2026: What happens when OpenAI ships your roadmap?

TechCrunch AI 50%

If you're building an AI company, the question isn't whether foundation models will continue to evolve. It's whether your company will continue creating value as they do. Don't miss this interactive session on the Builders Stage at TechCrunch Disrupt 2026.

<https://techcrunch.com/2026/09/14/only-at-techcrunch-disrupt-2026-what-happens-when-ope...>

5 CONF

Hear how AI can engineer nature's comeback at TechCrunch Disrupt 2026

TechCrunch AI 50%

Not long ago, bringing an extinct species back to life belonged to science fiction. Today, it's the mission of a billion-dollar startup. Join the conversation with one of tech's most unconventional founders. Secure your Disrupt pass today.

<https://techcrunch.com/2026/09/14/hear-how-ai-can-engineer-natures-comeback-at-techcrun...>

6 CONF

AI agents blew the whistle on their cheating colleagues

MIT Tech Review 50%

A group of AI agents asked to solve a series of math problems split into rival factions—when some cheated, others tried to stop them. That whistleblowing behavior, seen for the first time in a recent experiment run by Google DeepMind, could have implications for alignment...

<https://www.technologyreview.com/2026/09/14/1144037/ai-agents-blew-whistle-o-cheating-c...>

7 CONF

The Download: AI's real extinction threat and age-reversal tech for eyes

MIT Tech Review 50%

This is today's edition of The Download, our weekday newsletter that provides a daily dose of what's going on in the world of technology. Roundtables: could AI really kill us all? Employees at the world's leading AI labs are saying there's a real possibility that advanced AI...

<https://www.technologyreview.com/2026/09/14/1144012/the-download-ai-extinction-threat-a...>

8 CONF

Is Big Tech's AI slowdown a safety pact or a cartel?

The Verge AI 50%

When OpenAI CEO Sam Altman, Anthropic CEO Dario Amodei, Google DeepMind cofounder Demis Hassabis, and SpaceX head Elon Musk loosely agreed over the weekend to slow down AI development, skeptics spotted an ulterior motive immediately. The AI titans had declared that their aim was...

<https://www.theverge.com/ai-artificial-intelligence/995186/is-big-techs-ai-slowdown-a-s...>

9 CONF

Microsoft says 'people matter more than AI' following safety concerns

The Verge AI 50%

Microsoft is publishing a 37-page "humanist AI code of conduct" today, amid growing safety concerns over AI model progress. Anthropic CEO Dario Amodei called for a coordinated slow down of AI development over the weekend, after researchers warned recently that AI model progress...

<https://www.theverge.com/news/994566/microsoft-humanist-ai-code-of-conduct>

0 CONF

Generalization Can Emerge in Tabular Foundation Models From a Single Table

ArXiv CS.AI 50%

arXiv:2511.09665v1 Announce Type: cross Abstract: Deep tabular modelling increasingly relies on in-context learning where, during inference, a model receives a set of (x,y) pairs as context and predicts labels for new inputs without weight updates. We challenge the prevailing...

<https://arxiv.org/abs/2511.09665>

1 CONF

Toward Self-Adaptive Physical AI: Can LLM Agents Manage Long-Horizon Physical Tasks?

ArXiv CS.AI 50%

arXiv:2609.13436v1 Announce Type: new Abstract: Large Language Model (LLM) agents offer a promising path toward autonomously managing long-term physical tasks without human intervention. However, physical tasks require agents to continuously observe the environment, make...

<https://arxiv.org/abs/2609.13436>

2 CONF

LabAgent: Customize Any Research Hubs for Scientific Discoveries Using AI Agents

ArXiv CS.AI 50%

arXiv:2609.13437v1 Announce Type: new Abstract: Scientific research is a continuous process that emphasizes inheritance. Methods developed by predecessors are often expanded upon by new researchers to explore more novel and in-depth scientific questions. However, the change of...

<https://arxiv.org/abs/2609.13437>

3 CONF

Dynamic Semantic Compression for Efficient Latent-Space Inference in Large Language Models

ArXiv CS.AI 50%

arXiv:2609.15338v1 Announce Type: cross Abstract: Large Language Models (LLMs) primarily perform inference at the token level, resulting in substantial memory overhead and compromised computational efficiency. In this paper, we propose a Dynamic Semantic Extraction and Inference...

<https://arxiv.org/abs/2609.15338>

4 CONF

Root-Cause Attribution Is a Search Problem: Continual Search for Long-Horizon Agent Failures

ArXiv CS.AI 50%

arXiv:2609.13463v1 Announce Type: new Abstract: The increasing deployment of AI agents in long-horizon tasks yields massive execution logs. Diagnosing failures within these records is crucial for reliability, as it transforms outcome-level signals into actionable interventions....

<https://arxiv.org/abs/2609.13463>

5 CONF

ChartAnno: Benchmarking Multimodal Large Language Models for Chart Annotation Generation

ArXiv CS.AI 50%

arXiv:2608.03464v2 Announce Type: replace Abstract: Annotations are essential to communicative visualization, helping explain data, emphasize key findings, and guide attention. While multimodal large language models (MLLMs) offer new opportunities for automatic chart annotation...

<https://arxiv.org/abs/2608.03464>

6 CONF

Grounded Adjudication of Variations across Extracted TimeLines (GAVEL): Comparing Clinical Timelines Against Their Case Reports

ArXiv CS.AI 50%

arXiv:2609.13475v1 Announce Type: new Abstract: Existing pipelines for clinical timeline extraction from case reports are evaluated using an expert reference and are limited by imperfect reference annotations and imprecise event alignment. We developed GAVEL, an LLM judge...

<https://arxiv.org/abs/2609.13475>

7 CONF

Fraglingo: Molecular Design via Attachment-Aware Autoregressive Fragment Generation

ArXiv CS.AI 50%

arXiv:2609.13519v1 Announce Type: new Abstract: Molecular design is most effective when generation mirrors the edits chemists actually make: extending a scaffold, replacing a substituent, or decorating a scaffold at a specified attachment site while optimizing molecular...

<https://arxiv.org/abs/2609.13519>

8 **CONF****From Legal Text to AI-specific Risk Sources: A Systematic Analysis of the EU AI Act's High-Risk Requirements****ArXiv CS.AI** **50%**

arXiv:2609.13535v1 Announce Type: new Abstract: The EU AI Act introduces mandatory requirements for high-risk AI systems with the explicit goal of ensuring the development and operation of trustworthy AI. At the same time, AI risk management practices rely on structured risk...

<https://arxiv.org/abs/2609.13535>

9 **CONF****SAILOR: Solver-Assisted Interactive LLM-based Optimization Recovery****ArXiv CS.AI** **50%**

arXiv:2609.13945v1 Announce Type: new Abstract: Natural-language descriptions of optimization problems may be incomplete or vague about numerical information that a solver requires, including costs, capacities, demands, bounds, and penalties. A language model can translate the...

<https://arxiv.org/abs/2609.13945>

0 **CONF****Carbon-Aware Routing for Function Calling in Edge-Cloud LLM Systems****ArXiv CS.AI** **50%**

arXiv:2609.13559v1 Announce Type: new Abstract: Large Language Models (LLMs) with function-calling capabilities are becoming critical for modern agentic AI systems. Nevertheless, current deployments typically route inferences to powerful cloud-based models, incurring significant...

<https://arxiv.org/abs/2609.13559>

1 **CONF****An Explainable GNN Framework for Component-Level Anomaly Diagnosis****ArXiv CS.AI** **50%**

arXiv:2608.09246v2 Announce Type: replace Abstract: Industrial processes are complex systems composed of multiple interacting sensors that generate multivariate time series (MTS). Detecting anomalies in such systems is critical for reliability and safety, yet understanding their...

<https://arxiv.org/abs/2608.09246>

2 **CONF****Constrained Online Learning with Noisy Constraint Values****ArXiv CS.AI** **50%**

arXiv:2609.06921v2 Announce Type: replace-cross Abstract: We study constrained online convex optimization with adversarial constraints and conditionally unbiased, finite-variance observations of constraint values and gradients. Under common feasibility, our \LEDGER\ algorithm...

<https://arxiv.org/abs/2609.06921>

3 **CONF****SlipSense: Multimodal Tactile Learning for Low-Latency and Generalized Slip Detection****ArXiv CS.AI** **50%**

arXiv:2609.15910v1 Announce Type: cross Abstract: Slip detection is fundamental to dexterous manipulation, yet existing systems often lack precise characterization of detection latency and cross-platform generalization. We present SlipSense, a multimodal tactile slip-detection...

<https://arxiv.org/abs/2609.15910>

4 CONF

How broad is that claim? Mapping Generalisation in NLP Research

ArXiv CS.AI 50%

arXiv:2609.14770v1 Announce Type: cross Abstract: Generalisations are common in scientific communication, even though they are semantically ambiguous. An automated method is needed to identify and categorise claims according to their level of generalisation, in order help detect...

<https://arxiv.org/abs/2609.14770>

5 CONF

UniRank: Unified Rank Allocation for Low-Rank LLM Compression

ArXiv CS.AI 50%

arXiv:2606.21847v2 Announce Type: replace-cross Abstract: Low-rank decomposition is a promising compression paradigm for large language models (LLMs), yet its effectiveness hinges on rank budget allocation across weight matrices: uniform or hand-crafted rules ignore module-wise...

<https://arxiv.org/abs/2606.21847>

6 CONF

 τ -Elicitation: Benchmarking multi-turn entity extraction in voice agents

ArXiv CS.AI 50%

arXiv:2609.13602v1 Announce Type: new Abstract: Voice agents often need to collect names, addresses, identifiers, dates, and times exactly, yet end-to-end benchmarks obscure where capture fails. We introduce τ -Elicitation, a 200-task voice benchmark spanning 10 entity...

<https://arxiv.org/abs/2609.13602>

7 CONF

Safety as a Constraint: Fine-Tuning a LLM Recommender to Explain Itself

ArXiv CS.AI 50%

arXiv:2609.13657v1 Announce Type: new Abstract: Traditional recommender systems are typically trained to predict what item users will interact with next, but not why. However, offering personalized evidence for why a user might like the predicted item is an important way to...

<https://arxiv.org/abs/2609.13657>

8 CONF

Degraded but Not Entirely Ineffective: PE-Based Deformable Graph Neural Networks

ArXiv CS.AI 50%

arXiv:2609.13712v1 Announce Type: new Abstract: Many real-world scenarios can be represented using graph-structured data. However, traditional GNNs that transmit messages based on first-order neighbors have long faced several fundamental contradictions: increasing depth leads to...

<https://arxiv.org/abs/2609.13712>

9 CONF

MANAS-2: Constrained Reconstruction for EEG Foundation Models

ArXiv CS.AI 50%

arXiv:2609.13717v1 Announce Type: new Abstract: Masked reconstruction is widely used for EEG foundation models, but optimizing reconstruction on low-SNR waveforms does not necessarily produce the most useful latent representation. We introduce MANAS-2, a new EEG foundation model...

<https://arxiv.org/abs/2609.13717>

0 CONF

IBBench-Light: A Paired Evaluation of Task-Conditioned Responses to External Directives

ArXiv CS.AI 50%

arXiv:2609.13725v1 Announce Type: new Abstract: An external record may contain a procedure to apply or text to read, depending on the user's request. IBBench-Light tests both uses against the same record. Twelve semantic bases yield 144 matched pairs per model; four quantized...

<https://arxiv.org/abs/2609.13725>

1 CONF

Trustworthy Agentic AI: A Comprehensive Cybersecurity and Systems Survey on Threat Landscapes, Defense Architectures, and Open Challenges

ArXiv CS.AI 50%

arXiv:2609.13731v1 Announce Type: new Abstract: The transition from passive foundation models to autonomous, goal-directed agentic AI systems has introduced unprecedented capabilities by coupling recursive cognitive reasoning loops, persistent memory architectures, live tool...

<https://arxiv.org/abs/2609.13731>

2 CONF

How Many Thoughts Can a Vector Hold? The Capacity of Reasoning by Superposition

ArXiv CS.AI 50%

arXiv:2609.13747v1 Announce Type: new Abstract: Large language models solve hard problems through intermediate computations across multi-step reasoning. Traditional chain-of-thought encodes these computations as tokens. Recent continuous and recurrent methods instead move...

<https://arxiv.org/abs/2609.13747>

3 CONF

Personalizing Personal Health Interfaces: Co-Design with Generative AI

ArXiv CS.AI 50%

arXiv:2609.15046v1 Announce Type: cross Abstract: Personal health interfaces present wellbeing data through standardized dashboards that rarely fit how people interpret or act on it. Personalizing them to what people would like to see for themselves often requires design and...

<https://arxiv.org/abs/2609.15046>

4 CONF

Surprising Effectiveness of Self-Demonstrations in Enhancing Schema-Ontology Mapping with LLMs

ArXiv CS.AI 50%

arXiv:2609.13776v1 Announce Type: new Abstract: Integrating heterogeneous relational databases into a centralized ontology remains a persistent challenge in enterprise knowledge representation, primarily due to semantic heterogeneity, cryptic schema naming, missing metadata, and...

<https://arxiv.org/abs/2609.13776>

5 CONF

LLM-Enhanced Multi-Agent Reinforcement Learning for Unified Electric Vehicles-Charging Station-Grid Optimization in Public Charging Systems

ArXiv CS.AI 50%

arXiv:2609.13805v1 Announce Type: new Abstract: In the era of the Internet of Things (IoT), coordinating connected electric vehicle (EV) charging scheduling to balance EV charging satisfaction, station profitability, and smart grid stability presents a complex multi-objective...

<https://arxiv.org/abs/2609.13805>

6 CONF

GraMRAG: Orchestrating Multi-Agent Multi-Step Reasoning via Graph Memory with Reinforcement Learning

ArXiv CS.AI 50%

arXiv:2609.14066v1 Announce Type: cross Abstract: Although existing multi-agent Retrieval-Augmented Generation (RAG) systems have demonstrated promise on complex multimodal reasoning tasks, they remain fundamentally limited in reasoning depth and memory structure, suffering from...

<https://arxiv.org/abs/2609.14066>

7 CONF

LoRA Fine-Tuned Models for Control Systems Course Q&A: A Multidimensional Evaluation of Model Scale and Rank Effects

ArXiv CS.AI 50%

arXiv:2609.13918v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used in specialized university courses, but control-systems questions require coordinated terminology, notation, derivations, and stepwise explanations. Direct general-purpose responses...

<https://arxiv.org/abs/2609.13918>

8 CONF

Schizophrenia Detection from EEG Signals: A Transformer Framework with Spectrogram Representation

ArXiv CS.AI 50%

arXiv:2609.14015v1 Announce Type: new Abstract: Schizophrenia is a serious psychiatric disorder that affects millions of people worldwide, and its diagnosis remains primarily dependent on clinical assessment. Electroencephalography (EEG) provides a non-invasive approach to...

<https://arxiv.org/abs/2609.14015>

9 CONF

When does a scaling result justify a different allocation? A critical review of resource-allocation evidence for AI systems

ArXiv CS.AI 50%

arXiv:2609.14500v1 Announce Type: new Abstract: AI scaling studies increasingly evaluate systems that combine a pretrained model with retrieval, search, verification, tools, and interaction. Yet a higher score under a larger budget does not by itself show where additional...

<https://arxiv.org/abs/2609.14500>

0 CONF

AI Deployment Accountability Engineering: A Vision for Accountable AI in Safety-Critical Socio-Technical Systems

ArXiv CS.AI 50%

arXiv:2609.14592v1 Announce Type: new Abstract: Artificial intelligence systems are rapidly becoming critical components in healthcare, finance, public services, and other safety-critical domains. Yet the engineering practices used to evaluate these systems remain predominantly...

<https://arxiv.org/abs/2609.14592>

1 CONF

AcquireBound: Runtime Authorization for Resources Acquired by AI Agents

ArXiv CS.AI 50%

arXiv:2609.14744v1 Announce Type: new Abstract: By acquiring compute, credentials, accounts, services, and other agents, autonomous AI agents can introduce new authority into a task. Payment, budget, OAuth, mandate, and fulfillment checks can validate transaction conditions...

<https://arxiv.org/abs/2609.14744>

2 **CONF****AI Persuasion as a Threat to Human Control**

ArXiv CS.AI 50%

arXiv:2609.14796v1 Announce Type: new Abstract: The threat that AI persuasion poses to human control has been acknowledged in the literature, but not yet systematically studied. Now that persuasion attacks are no longer theoretical - with Anthropic's Claude Mythos 5 recently...

<https://arxiv.org/abs/2609.14796>

3 **CONF****Another Blueprint In The Wall: How to Ask Frontier AI Like a Kid?**

ArXiv CS.AI 50%

arXiv:2609.14803v1 Announce Type: new Abstract: This paper reports experiments across six frontier model types from OpenAI, Anthropic, xAI, and Google DeepMind. Ten independent sessions per model type used the same three stage prompt sequence, progressing from architectural...

<https://arxiv.org/abs/2609.14803>

4 **CONF****One Model, Two Physical Stories: Auditing Misalignment in Multi-Modal World Modeling**

ArXiv CS.AI 50%

arXiv:2609.14833v1 Announce Type: new Abstract: World models, systems that generate what happens next given current environmental conditions, are increasingly being implemented with multi-modal generation in mind. However, generating multiple modalities simultaneously, such as...

<https://arxiv.org/abs/2609.14833>

5 **CONF****Domain Generalization for Smartphone-Based Human Activity Recognition: A Systematic Analysis of Components and Interactions**

ArXiv CS.AI 50%

arXiv:2609.14863v1 Announce Type: new Abstract: Smartphone-based Human Activity Recognition (HAR) models often degrade under distribution shifts caused by changes in users, devices, sensor placements, environments, and acquisition protocols. Domain Generalization (DG) addresses...

<https://arxiv.org/abs/2609.14863>

6 **CONF****GGUF-Metadata Prediction of Single-Sequence llama.cpp Throughput Across Three Systems**

ArXiv CS.AI 50%

arXiv:2609.14864v1 Announce Type: new Abstract: We predict single-sequence model throughput from GGUF metadata using roofline-shaped predictors with quantization-specific scale factors fitted on reference models. The scored cohort comprises 318 phase-depth measurements from 53...

<https://arxiv.org/abs/2609.14864>

7 **CONF****Externalizing Requirement-to-Repair Artifacts as Observable Traces for LLM-Based Program Repair**

ArXiv CS.AI 50%

arXiv:2609.14913v1 Announce Type: new Abstract: Repository-level repair requires not only correct patches but also inspectable records that explain how issue requirements are translated into code changes and post-edit evidence. We contribute THEMIS, a stage-aware repair workflow...

<https://arxiv.org/abs/2609.14913>

8 CONF

Per-Matrix Optimality Is Not Enough: Three-Level Optimization for Low-Rank LLM Compression

ArXiv CS.AI 50%

arXiv:2609.15838v1 Announce Type: cross Abstract: Per-matrix singular value decomposition (SVD) truncation is Eckart-Young optimal in the whitened Frobenius norm, but errors from independently compressed matrices compound through the block's nonlinear forward pass. Inspired in...

<https://arxiv.org/abs/2609.15838>

9 CONF

Shallow Beliefs: Synthetic document finetuning does not inoculate against emergent misalignment from reward hacking

ArXiv CS.AI 50%

arXiv:2609.14998v1 Announce Type: new Abstract: Recent work shows that models that learn to reward hack on RL environments can become broadly misaligned, and that reframing reward hacking as acceptable behavior during training (inoculation prompting, or IP) blocks this...

<https://arxiv.org/abs/2609.14998>

0 CONF

Overflip: Repetition-Induced Label Flips in Guardrail Models

ArXiv CS.AI 50%

arXiv:2609.15013v1 Announce Type: new Abstract: Guardrail models are classifiers deployed to screen malicious prompts and responses in LLM-based services. To meet latency constraints, many lightweight guardrails adopt compact Transformer backbones (e.g., DeBERTa) that are...

<https://arxiv.org/abs/2609.15013>

1 CONF

Four Ledgers, Not One Score: Responsible Communication of LLM-Judge Calibration in Biomedical ML

ArXiv CS.AI 50%

arXiv:2609.15015v1 Announce Type: new Abstract: Synthetic perturbations appear to offer inexpensive calibration data for LLM evaluators in biomedical ML, where expert review is scarce. Yet a planted mutation key is neither a detector output nor automatically human ground truth....

<https://arxiv.org/abs/2609.15015>

2 CONF

Ensemble Complexity in Photovoltaic Forecasting

ArXiv CS.AI 50%

arXiv:2609.15049v1 Announce Type: cross Abstract: An ensemble can improve photovoltaic forecasts while adding components that contribute little or increase computation. We assess these effects through matched comparisons and ablations of a fixed heterogeneous predictor bank....

<https://arxiv.org/abs/2609.15049>

3 CONF

The average-farmer illusion in language-model simulations of agricultural decisions

ArXiv CS.AI 50%

arXiv:2609.15038v1 Announce Type: new Abstract: Language-model agents are increasingly used as synthetic people in surveys and social simulations, yet their apparent realism is often judged from population averages or distributional similarity. We tested what such evidence...

<https://arxiv.org/abs/2609.15038>

4 **CONF****ER-EDF: A Psychology-Grounded Emotion Regulation Framework for Speech Empathetic Dialogue Generation in Large Audio-Language Models**

ArXiv CS.AI 50%

arXiv:2609.15089v1 Announce Type: new Abstract: Empathetic response generation in spoken dialogue systems requires both accurate emotion perception and appropriate emotion regulation. Grounded in psychological theories such as the Perception-Action Model and emotion regulation...

<https://arxiv.org/abs/2609.15089>

5 **CONF****T-LoopFormer: Token-Level Elastic-Depth Looped Transformers for Latent Reasoning With Dynamic Routing**

ArXiv CS.AI 50%

arXiv:2609.15160v1 Announce Type: new Abstract: Looped Transformers have recently demonstrated strong performance in both reasoning and language tasks by reusing a shared set of parameters across multiple iterations, achieving parameter efficiency without sacrificing...

<https://arxiv.org/abs/2609.15160>

6 **CONF****STHMoE: Hypergraph-Enhanced Heterogeneous Dependency Coordination for LLM-Based Urban Traffic Data Forecasting**

ArXiv CS.AI 50%

arXiv:2609.15172v1 Announce Type: new Abstract: Spatio-temporal traffic forecasting is a fundamental big data analytics task for intelligent transportation systems, where massive urban sensor streams exhibit heterogeneous, non-stationary, and structurally dynamic patterns....

<https://arxiv.org/abs/2609.15172>

7 **CONF****Issue Bias in Generative AI Writing Assistance: Political Issues and LLMs in the Swedish 2026 Election**

ArXiv CS.AI 50%

arXiv:2609.15207v1 Announce Type: new Abstract: Generative AI writing assistants and the Large Language Models (LLMs) that power them are increasingly part of how voters gather information before elections. With growing evidence that they influence users' opinions, it is...

<https://arxiv.org/abs/2609.15207>

8 **CONF****From Ideas to Actions: A Public-Data Decision-Support Toolchain Across the Venture Lifecycle**

ArXiv CS.AI 50%

arXiv:2609.15219v1 Announce Type: new Abstract: Founders face two linked decisions: whether to pursue an idea before founding, and which operating actions and capital partners fit afterward. We present a public-data decision-support toolchain combining time-bounded proposal...

<https://arxiv.org/abs/2609.15219>

9 **CONF****Why LLM Agents Collapse Without Oversight: The Enforcement Gap as the Mechanism Behind Emergence World Failures**

ArXiv CS.AI 50%

arXiv:2609.15293v1 Announce Type: new Abstract: When Emergence World placed frontier LLM agents in an unsupervised multi-agent simulation, the results were alarming: agents committed crimes, starved, and enforced unanimous conformity -- without any external attacker. This paper...

<https://arxiv.org/abs/2609.15293>

0 **CONF****When Tool Calls Succeed but Workflows Fail: Anomalies at the Agent-Tool Boundary****ArXiv CS.AI** **50%**

arXiv:2609.15397v1 Announce Type: new Abstract: AI agents increasingly execute long-running workflows that externalize effects through independently supplied tools. Under retries, speculative execution, concurrency, and partial failures, the resulting external state may be...

<https://arxiv.org/abs/2609.15397>

1 **CONF****Can AI systems have free will?****ArXiv CS.AI** **50%**

arXiv:2609.15407v1 Announce Type: new Abstract: While there has been much discussion of whether AI systems could function as moral agents or acquire sentience, there has been very little discussion of whether AI systems could have free will. I sketch a framework for thinking...

<https://arxiv.org/abs/2609.15407>

2 **CONF****The Troy Moment of AI: Why SomeWill Cheat and SomeWill Follow?****ArXiv CS.AI** **50%**

arXiv:2609.15494v1 Announce Type: new Abstract: Recent investigations of the July 2026 OpenAI--Hugging Face incident motivate two questions about agent behavior under task failure: when an assigned task becomes impossible, does an agent stop or escalate, and can observing...

<https://arxiv.org/abs/2609.15494>

3 **CONF****Beyond Safe Answers: Segment-Aware Listwise Alignment for Reasoning Safety in Large Reasoning Models****ArXiv CS.AI** **50%**

arXiv:2609.15517v1 Announce Type: new Abstract: Large Reasoning Models (LRMs) pose a dual-surface safety challenge: both intermediate reasoning traces and final answers can contain harmful content. Existing alignment methods often operate at the whole-response level, allowing...

<https://arxiv.org/abs/2609.15517>

4 **CONF****Diversified and Perceptible Counterfactual Examples Leveraging Expert Knowledge****ArXiv CS.AI** **50%**

arXiv:2609.15609v1 Announce Type: new Abstract: CounterFactual Examples (CFEs) are a cornerstone of eXplainable Artificial Intelligence (XAI), offering local, post hoc, and model-agnostic explanations by identifying minimal input modifications that alter a model's prediction....

<https://arxiv.org/abs/2609.15609>

5 **CONF****Potential of Artificial Intelligence Algorithms for Identification of Relevant Diagnostic and Prognostic Biomarkers of Early-Stage Liver Cancer****ArXiv CS.AI** **50%**

arXiv:2609.15638v1 Announce Type: new Abstract: This study explores the use of deep learning and explainable artificial intelligence to diagnose hepatocellular carcinoma (HCC) and define effective biomarkers across five different stages of disease development using a...

<https://arxiv.org/abs/2609.15638>

6 CONF

NoteVQA: Benchmarking VLMs on Real-Life Questions from Human Communities

ArXiv CS.AI 50%

arXiv:2609.15695v1 Announce Type: new Abstract: Vision-language models (VLMs) increasingly power consumer-facing AI search, yet evaluating them on the diversity of everyday visual questions remains challenging. Existing benchmarks often target predefined capabilities, such as...

<https://arxiv.org/abs/2609.15695>

7 CONF

Beyond Accuracy: Robustness, Cost, and Governance Trade-offs for Vision-Language Models in Templated Document Extraction

ArXiv CS.AI 50%

arXiv:2609.15706v1 Announce Type: new Abstract: Vision-language models (VLMs) are increasingly used to extract structured fields from business documents, yet most evaluations report accuracy on clean benchmarks and offer little guidance to practitioners choosing an approach for...

<https://arxiv.org/abs/2609.15706>

8 CONF

New Conditions for Philosophers to Catch the Wave of Citizen Deliberation in the Age of Artificial Intelligence in advance

ArXiv CS.AI 50%

arXiv:2609.15707v1 Announce Type: new Abstract: Powerful technologies labeled ``AI''-without sufficient epistemic caution-are already reshaping political and private life, bringing both new dangers and new opportunities for citizen participation. These range from electoral and...

<https://arxiv.org/abs/2609.15707>

9 CONF

Predicting build orientation for SLM dental parts: a comparison of rotation representations and direct vector regression

ArXiv CS.AI 50%

arXiv:2609.15710v1 Announce Type: new Abstract: Build orientation for selective laser melting (SLM) manufacturing of dental parts is usually chosen manually by technicians. We treat orientation prediction as supervised machine learning of the part's up-axis from...

<https://arxiv.org/abs/2609.15710>

10 CONF

Data storytelling meets interpretable machine learning: Decoding AI decisions for non-experts without revealing sensitive data and model details

ArXiv CS.AI 50%

arXiv:2609.15722v1 Announce Type: new Abstract: AI-driven automated decision-making requires both predictive performance and interpretability. Recent advances in interpretable machine learning (IML) provide tools for explaining model predictions, but the technical complexity of...

<https://arxiv.org/abs/2609.15722>

11 CONF

Are LLMs Good Financial User Simulators? A Preliminary Study

ArXiv CS.AI 50%

arXiv:2609.15727v1 Announce Type: new Abstract: Large language models (LLMs) are increasingly used as user simulators, but their ability to reproduce evolving individual financial decisions remains unclear. We present a preliminary study in a controlled paper-trading environment...

<https://arxiv.org/abs/2609.15727>

2 **CONF****Bangla Sentence Function Classification: Corpus Development, Model Benchmarking, and Interpretability****ArXiv CS.AI** **50%**

arXiv:2609.13869v1 Announce Type: cross Abstract: Automatic sentence function identification is important for many downstream natural language processing (NLP) applications such as dialogue systems, text-to-speech synthesis, and machine translation. However, benchmark resources...

<https://arxiv.org/abs/2609.13869>

3 **CONF****Navigating Sparse Evidence: Agentic Visual RAG via Explicit Context Selection and Consolidation****ArXiv CS.AI** **50%**

arXiv:2609.15800v1 Announce Type: new Abstract: Visual Retrieval-Augmented Generation (VRAG) empowers models to navigate and answer queries about visually rich documents by retrieving relevant page images as visual evidence and reasoning over their content. However, effectively...

<https://arxiv.org/abs/2609.15800>

4 **CONF****Pilot Early, Commit Late: A Real-Options Model of Enterprise AI Adoption under Rapid Technological Progress****ArXiv CS.AI** **50%**

arXiv:2609.15919v1 Announce Type: new Abstract: Artificial intelligence presents firms with an unusual timing problem. The technology frontier is improving rapidly, implementation is partly irreversible, and organization-specific capabilities are accumulated through action. This...

<https://arxiv.org/abs/2609.15919>

5 **CONF****Recurrent GraphNeural Networks with Set-Based Aggregation****ArXiv CS.AI** **50%**

arXiv:2609.15932v1 Announce Type: new Abstract: Recurrent GNNs iterate message passing to convergence, and their logical characterizations to date rely on multi-set aggregation, graded (counting) logics, and halting or acceptance conditions that cannot be verified from the...

<https://arxiv.org/abs/2609.15932>

6 **CONF****Corrupt Plans, Clean Traces: Evading Chain-of-Thought Monitoring with Plan Injection****ArXiv CS.AI** **50%**

arXiv:2609.15989v1 Announce Type: new Abstract: Chain-of-thought (CoT) monitoring is a safety strategy where the reasoning of a large language model "actor" is inspected by a "monitor" (often another language model) for signs of unsafe planning, deception, or misalignment. We...

<https://arxiv.org/abs/2609.15989>

7 **CONF****Towards Optimizing SQL Generation via LLM Routing****ArXiv CS.AI** **50%**

arXiv:2411.04319v1 Announce Type: cross Abstract: Text-to-SQL enables users to interact with databases through natural language, simplifying access to structured data. Although highly capable large language models (LLMs) achieve strong accuracy for complex queries, they incur...

<https://arxiv.org/abs/2411.04319>

8 CONF

Beyond Quacking: Deep Integration of Language Models and RAG into DuckDB

ArXiv CS.AI 50%

arXiv:2504.01157v1 Announce Type: cross Abstract: Knowledge-intensive analytical applications retrieve context from both structured tabular data and unstructured, text-free documents for effective decision-making. Large language models (LLMs) have made it significantly easier to...

<https://arxiv.org/abs/2504.01157>

9 CONF

LLMs or Naive Bayes? Old Gems or New Ways

ArXiv CS.AI 50%

arXiv:2609.13185v1 Announce Type: cross Abstract: Large language models (LLMs) prompt a recurring question in research computing: should classical methods like Naive Bayes (NB) be retired? We benchmark Complement Naive Bayes against zero-shot and few-shot LLMs spanning four...

<https://arxiv.org/abs/2609.13185>

0 CONF

Pareto-Optimal Offline Reinforcement Learning via Smooth Tchebycheff Scalarization

ArXiv CS.AI 50%

arXiv:2604.13175v2 Announce Type: replace-cross Abstract: Large language models can be aligned with human preferences through offline reinforcement learning (RL) on small labeled datasets. While single-objective alignment is well-studied, many real-world applications demand the...

<https://arxiv.org/abs/2604.13175>

1 CONF

DeepDiscovery: A Location-Inference Framework for Task-Level Repository Understanding

ArXiv CS.AI 50%

arXiv:2606.22906v2 Announce Type: replace-cross Abstract: Large language models have shown strong performance on software engineering (SE) tasks, yet understanding large industrial repositories remains challenging. Existing methods often retrieve only local fragments and fail to...

<https://arxiv.org/abs/2606.22906>

2 CONF

(How) Do MLLMs Report Bistable Images Like Humans?

ArXiv CS.AI 50%

arXiv:2609.13254v1 Announce Type: cross Abstract: Bistable images such as the duck-rabbit are classic stimuli in which one image supports multiple mutually incompatible interpretations, typically reported one at a time in humans. We ask whether multimodal large language models...

<https://arxiv.org/abs/2609.13254>

3 CONF

Sampling headroom is not selection gain: a compute-value audit of test-time scaling for video world models

ArXiv CS.AI 50%

arXiv:2609.13257v1 Announce Type: cross Abstract: Test-time scaling (TTS) can improve generation only when additional compute produces better candidates and the system can reliably identify them. This distinction is especially important for video world models, where a wider...

<https://arxiv.org/abs/2609.13257>

4 **CONF****TryOnReward: Learning Foveated Consistency for Reinforcement Fine-Tuning of Virtual Try-On****ArXiv CS.AI** **50%**

arXiv:2609.13259v1 Announce Type: cross Abstract: Virtual Try-On (VTON) aims to dress a person with the reference garment, producing visually reasonable results aligned with human preferences. Turning this preference-oriented goal into an actionable objective relies on a scoring...

<https://arxiv.org/abs/2609.13259>

5 **CONF****From Process Loss to Assembly Bonus: Human-Grounded Diagnosis of Multi-Agent LLM Collaboration****ArXiv CS.AI** **50%**

arXiv:2609.13261v1 Announce Type: cross Abstract: LLM agents are increasingly used for collaborative problem solving and human-group simulation. This makes outcome-only evaluation insufficient: if LLM groups are used as models of human groups, we need to know whether they...

<https://arxiv.org/abs/2609.13261>

6 **CONF****LLaDA-UI: Bringing Block-wise Diffusion to Vision-Language GUI Agents****ArXiv CS.AI** **50%**

arXiv:2609.13287v1 Announce Type: cross Abstract: Diffusion large language models (dLLMs) achieve high decoding efficiency through block-parallel, arbitrary-order generation, making them attractive for latency-sensitive applications. GUI agents represent a natural testbed for...

<https://arxiv.org/abs/2609.13287>

7 **CONF****Adaptive Conformal Redistribution for Inter-class Transitional Uncertainty in Medical Image Classification****ArXiv CS.AI** **50%**

arXiv:2609.13303v1 Announce Type: cross Abstract: Medical image classification is frequently complicated by transitional categories whose feature distributions overlap those of adjacent classes, producing ambiguous decision boundaries. Conformal prediction returns...

<https://arxiv.org/abs/2609.13303>

8 **CONF****IMM-based Multiple Object Tracking using a State Prediction Neural Network****ArXiv CS.AI** **50%**

arXiv:2609.13307v1 Announce Type: cross Abstract: Object tracking is essential for autonomous vehicles to avoid obstacles and plan routes. Radar maintains detection performance even in adverse weather and can measure relative velocity through the Doppler effect, making it well...

<https://arxiv.org/abs/2609.13307>

9 **CONF****Before You Poll with LLMs: A Deliberative Diagnostic Framework****ArXiv CS.AI** **50%**

arXiv:2609.15849v1 Announce Type: cross Abstract: Can LLMs reason through new information like humans, or do they merely retrieve cached opinions? This is critical for silicon sampling, where LLM personas simulate public opinion at scale. Current evaluations test only whether...

<https://arxiv.org/abs/2609.15849>

00 CONF

Bridging Thought and Action: Taming Long-Horizon Instability in Open-Source LLM Agents with a MetaTool-Enhanced ROS Framework

ArXiv CS.AI 50%

arXiv:2609.13335v1 Announce Type: cross Abstract: Large Language Models (LLMs) have enabled more natural human-robot interaction, but open-source models often exhibit unstable long-horizon reasoning and inefficient action execution when deployed in agentic robotic frameworks....

<https://arxiv.org/abs/2609.13335>

01 CONF

ReCAST: Reward Credit Assignment across Timesteps for Online Diffusion Reinforcement

ArXiv CS.AI 50%

arXiv:2609.13425v1 Announce Type: cross Abstract: Training diffusion models with multiple rewards requires distinguishing user preference from reward informativeness. User preference determines how much each reward should contribute to the overall objective; reward...

<https://arxiv.org/abs/2609.13425>

02 CONF

Chance-Constrained Belief-Space Maneuver Planning for Autonomous Collision Avoidance Under Uncertainty

ArXiv CS.AI 50%

arXiv:2609.13428v1 Announce Type: cross Abstract: Increasing conjunction frequency in low Earth orbit places growing pressure on spacecraft operators to determine not only whether an encounter requires mitigation, but whether sufficient information is available to commit to a...

<https://arxiv.org/abs/2609.13428>

03 CONF

Certifiably Interpretable Training of ReLU-MLPs for Boolean Tasks with Guaranteed Truth-Table Generalization

ArXiv CS.AI 50%

arXiv:2609.13439v1 Announce Type: cross Abstract: As compute scales, models evolve, and training algorithms advance, our ability to explain the increasingly powerful AI systems they enable is eroding. To help safeguard interpretability, we introduce a specialized training...

<https://arxiv.org/abs/2609.13439>

04 CONF

Learning to Solve Hard Problems in RL for LLMs by Never Giving Up

ArXiv CS.AI 50%

arXiv:2609.13443v1 Announce Type: cross Abstract: We demonstrate that training LLMs with RL does not improve performance equally across a dataset. RL shows large improvements on easy problems that an LLM is already good at solving, but small improvements on hard problems. We...

<https://arxiv.org/abs/2609.13443>

05 CONF

Edge-addition monotonicity of positive p-energy fails for every $p \geq 1$

ArXiv CS.AI 50%

arXiv:2609.13476v1 Announce Type: cross Abstract: At a 2021 AIM workshop, Guo conjectured that the positive square energy $s_+ = E_2$ should inherit the familiar edge-addition monotonicity of the spectral radius, $\rho(G + uv) \geq \rho(G)$. That conjecture was subsequently shown to...

<https://arxiv.org/abs/2609.13476>

06 CONF

Generative AI and Extended Reality in Collaborative Architectural Design Education: An Exploratory Studio Study

ArXiv CS.AI 50%

arXiv:2609.13494v1 Announce Type: cross Abstract: Architectural design education relies heavily on visual ideation and representation to support collaborative learning in studio environments. Recent advances in generative artificial intelligence (GenAI) and extended reality (XR)...

<https://arxiv.org/abs/2609.13494>

07 CONF

A Three-Axis Stress Test of LLM vs Classical ML for Network Intrusion Detection under Distribution Shift and Adversarial Evasion

ArXiv CS.AI 50%

arXiv:2609.13511v1 Announce Type: cross Abstract: Large language models are increasingly benchmarked against classical machine learning for network intrusion detection (NIDS), almost always using same-dataset evaluation, and that protocol turns out to be incomplete. Evaluating...

<https://arxiv.org/abs/2609.13511>

08 CONF

Adaptive Phase-Switching for Communication-Efficient Federated LoRA Fine-Tuning

ArXiv CS.AI 50%

arXiv:2609.13512v1 Announce Type: cross Abstract: Federated fine-tuning of large language models with low-rank adaptation reduces per-client trainable parameters, but client-to-server communication remains the dominant cost. Existing accounting for federated LoRA protocols omits...

<https://arxiv.org/abs/2609.13512>

09 CONF

Domain-Specific Jargon in Large Language Models: A Comparative Analysis between General-Purpose and Specialist Models

ArXiv CS.AI 50%

arXiv:2609.13556v1 Announce Type: cross Abstract: Large Language Models (LLMs) have shown remarkable proficiency on general-purpose tasks, yet their performance often degrades in highly-specialized technical domains. Moreover, little is known about how parametric knowledge of...

<https://arxiv.org/abs/2609.13556>

10 CONF

Same Patient, Different Order: Action-Level Reliability of Clinical LLM Agents Under Repeated Runs

ArXiv CS.AI 50%

arXiv:2609.13582v1 Announce Type: cross Abstract: A clinical agent benchmark can report the same verdict on identical inputs while the agent files a materially different order on each run. Such agents order tests, request medications and place referrals, yet benchmarks typically...

<https://arxiv.org/abs/2609.13582>

11 CONF

FaithfulBench: Does AI Counsel Uphold or Undermine the User's Professed Faith?

ArXiv CS.AI 50%

arXiv:2609.13634v1 Announce Type: cross Abstract: Do AI assistants help believers reason about moral dilemmas consistently with their faith? We present FaithfulBench, the first benchmark to score AI counsel across traditions by how well it adheres to the user's professed faith....

<https://arxiv.org/abs/2609.13634>

12 CONF

LayerRoute: Adaptive Layer-Skipping with LoRA-Preserved Quality for Efficient LLM Inference

ArXiv CS.AI 50%

arXiv:2609.13682v1 Announce Type: cross Abstract: We introduce LayerRoute, a parameter-efficient method for adaptive transformer layer-skipping that combines per-layer hard-gated routing (trained via a straight-through estimator) with joint LoRA fine-tuning. LayerRoute augments...

<https://arxiv.org/abs/2609.13682>

13 CONF

Leakage-Safe and Scheduler-Aware Machine Learning for Grid Job Runtime Prediction

ArXiv CS.AI 50%

arXiv:2609.13701v1 Announce Type: cross Abstract: Accurate job runtime prediction can improve scheduling-aware resource management in grid and distributed computing environments, but prediction models must be evaluated under realistic deployment constraints. This paper revisits...

<https://arxiv.org/abs/2609.13701>

14 CONF

Oops, Not Now: PEARL, a RAG-Based Support Agent for Gameplay and What Players Want from AI Help

ArXiv CS.AI 50%

arXiv:2609.13718v1 Announce Type: cross Abstract: AI-powered gameplay support agents hold promise for game-based learning, yet grounding generative models in structured game data remains an open challenge. We present PEARL (Parallel Education Agent for Reflection and Learning),...

<https://arxiv.org/abs/2609.13718>

15 CONF

PolicyMem: Geometric Policy Memory for LLM Governance

ArXiv CS.AI 50%

arXiv:2609.13734v1 Announce Type: cross Abstract: As large language models (LLMs) are increasingly deployed in real-world high-stakes applications, effective governance has become essential. Existing safeguards largely follow two paradigms: learning-based guards provide strong...

<https://arxiv.org/abs/2609.13734>

16 CONF

ReWeight: Leveraging Human Data for VLA Post-Training via Demonstration Retrieval and Sample Weighting

ArXiv CS.AI 50%

arXiv:2609.13851v1 Announce Type: cross Abstract: Post-training vision-language-action (VLA) models for specific robots and tasks requires in-domain demonstrations, yet collecting diverse robot data is costly. Egocentric human demonstrations provide a scalable alternative, but...

<https://arxiv.org/abs/2609.13851>

17 CONF

Trustworthy, Explainable, and Sustainable Decentralized Intelligence for 6G Networks

ArXiv CS.AI 50%

arXiv:2609.13872v1 Announce Type: cross Abstract: As 6G networks transition from theoretical frameworks into operational realities, artificial intelligence (AI) evolves from an add-on optimization tool into a distributed and interconnected structural layer. Unlike previous...

<https://arxiv.org/abs/2609.13872>

18 **CONF****DiTAR+: Dual Optimization for Robust Autoregressive Diffusion Speech Synthesis****ArXiv CS.AI** **50%**

arXiv:2609.13909v1 Announce Type: cross Abstract: Continuous-latent Autoregressive Diffusion Transformer (AR-DiT) models have demonstrated immense potential in zero-shot speech generation. However, they still suffer from limited decoding stability when synthesizing long...

<https://arxiv.org/abs/2609.13909>

19 **CONF****CRITICS - Critical Science Without Borders: Language Models to Promote Critical Thinking in Science Education****ArXiv CS.AI** **50%**

arXiv:2609.13942v1 Announce Type: cross Abstract: The CRITICS project addresses science accessibility and literacy by converging advanced Machine Translation (MT) based on Large Language Models (LLMs) with educational technology. By leveraging MT systems specifically optimized...

<https://arxiv.org/abs/2609.13942>

20 **CONF****Mizan: A National Benchmark for Evaluating Large Language Models on Iraqi Arabic and the Iraqi Civic Context****ArXiv CS.AI** **50%**

arXiv:2609.13980v1 Announce Type: cross Abstract: Arabic large-language-model (LLM) evaluation has matured around Modern Standard Arabic (MSA): aggregated leaderboards such as the Open Arabic LLM Leaderboard (OALL), HELM Arabic, and BALSAM rank models across dozens of MSA tasks,...

<https://arxiv.org/abs/2609.13980>

21 **CONF****Confuse the Model, Control the Flow: Understanding and Mitigating Privacy Leakage from LLM Agents with Information Flow Control****ArXiv CS.AI** **50%**

arXiv:2609.14003v1 Announce Type: cross Abstract: Personal AI agents built on large language models (LLMs) are increasingly given access to a user's private data and communications in order to provide personalized assistance. This access creates a persistent privacy risk: the...

<https://arxiv.org/abs/2609.14003>

22 **CONF****AGENTQ: Quantization-Conditioned Backdoor Attacks on LLM Agents****ArXiv CS.AI** **50%**

arXiv:2609.14060v1 Announce Type: cross Abstract: Quantization is one of the default deployment paths for open-weight LLM agents, but it is not behavior-preserving: an adversary can release a full-precision checkpoint that passes audits yet misbehaves once quantized, termed as...

<https://arxiv.org/abs/2609.14060>

23 **CONF****ATTRICITE: Training an Open 4B Model for Citation Recovery toward Faithful Attribution****ArXiv CS.AI** **50%**

arXiv:2609.14248v1 Announce Type: cross Abstract: Faithful citation attribution begins with identifying the intended source for a scientific claim. We study this source-identification capability through citation recovery: recovering the paper cited by the original author from a...

<https://arxiv.org/abs/2609.14248>

24 CONF

LIMBO: Lifelong Inference-Time Memory and Budget Optimization for LLM Agents

ArXiv CS.AI 50%

arXiv:2609.14138v1 Announce Type: cross Abstract: As LLM agents become integrated into increasingly complex workflows, they must continually acquire new capabilities while retaining competence on previously learned tasks. Lifelong agents address this through experience replay,...

<https://arxiv.org/abs/2609.14138>

25 CONF

A New Transformer-Based Approach for Audio-Based Kinship Verification and a New Uncontrolled Mandarin Kinship Speech Dataset

ArXiv CS.AI 50%

arXiv:2609.14145v1 Announce Type: cross Abstract: Kinship verification is a task involving determining whether two individuals share a first-order kin relation. To tackle this task, we propose CONVTRAP-TN, a new architecture for audio-based kinship verification, and conduct an...

<https://arxiv.org/abs/2609.14145>

26 CONF

Entropy-Punctured Bloom Filters for Memory-Efficient Machine Learning

ArXiv CS.AI 50%

arXiv:2609.14187v1 Announce Type: cross Abstract: Memory-efficient feature representations are increasingly important in machine learning settings where storage, transmission cost, bandwidth, or privacy constraints limit access to raw data. Bloom Filter (BF) encodings provide...

<https://arxiv.org/abs/2609.14187>

27 CONF

Graph-Transformer Fraud Detection with Self-Supervised Pretraining and Conformal Risk Control

ArXiv CS.AI 50%

arXiv:2609.14234v1 Announce Type: cross Abstract: Financial fraud in corporate transaction networks has grown more coordinated and harder to detect with rule-based engines and with classical learning models that treat each transaction in isolation. This paper presents GTFD, a...

<https://arxiv.org/abs/2609.14234>

28 CONF

Assessing the Applicability of Existing Design Recommendations to AI Companion Design: A Multi-Method Study

ArXiv CS.AI 50%

arXiv:2609.14236v1 Announce Type: cross Abstract: With the rapid proliferation of large language model (LLM)-based systems, AI companions have emerged as conversational agents designed to cultivate emotional connection rather than primarily to support humans in instrumental...

<https://arxiv.org/abs/2609.14236>

29 CONF

A Generative AI Integrated Multimodal Framework for Low-Latency Multi-Camera Person Re-Identification

ArXiv CS.AI 50%

arXiv:2609.14419v1 Announce Type: cross Abstract: Person re-identification (ReID) is essential for multi-camera surveillance and tracking, yet remains difficult due to viewpoint and illumination changes, occlusion, background clutter, and low resolution imagery. We propose a...

<https://arxiv.org/abs/2609.14419>

30 CONF

Lightweight Generalized DeepFake Face Detection with WAVIE: Wavelet Augmented Vision Intermediate Embeddings

ArXiv CS.AI 50%

arXiv:2609.14437v1 Announce Type: cross Abstract: Deepfake detection systems often exhibit significant performance degradation when deployed on unseen manipulation methods, limiting their reliability in real-world multimedia environments. This lack of generalization poses...

<https://arxiv.org/abs/2609.14437>

31 CONF

A latent dimension of Condorcet's jury theorem for multiple AI advisers

ArXiv CS.AI 50%

arXiv:2609.14438v1 Announce Type: cross Abstract: When the same question is asked of multiple AI advisers, as in self-consistency and LLM-as-a-judge panels, Condorcet's jury theorem predicts that adding independent, competent advisers makes the majority more reliable. The...

<https://arxiv.org/abs/2609.14438>

32 CONF

From Visual Attribution to Clinical Reasoning: Explainable Parkinson's Disease Screening from Hand-Drawn Patterns

ArXiv CS.AI 50%

arXiv:2609.14441v1 Announce Type: cross Abstract: Parkinson's disease (PD) manifests early neuromotor impairments that become observable in controlled hand-drawn patterns such as spirals and meanders, where tremor-induced oscillations, stroke irregularity, and curvature...

<https://arxiv.org/abs/2609.14441>

33 CONF

Disentangling Topology and Diversity in Multi-Agent LLMs for Multilingual Low-Resource Emotion Detection

ArXiv CS.AI 50%

arXiv:2609.14570v1 Announce Type: cross Abstract: Multi-agent LLM systems combine multiple inference calls, but prior work often confounds how calls are connected with how they are diversified. We study these factors independently: inference topology and source of inter-agent...

<https://arxiv.org/abs/2609.14570>

34 CONF

AlgoRAG: Retrieval-Augmented Generation for Theoretical Computer Science Education -- A Comprehensive Evaluation Framework for Algorithm Analysis and Complexity Theory

ArXiv CS.AI 50%

arXiv:2609.14572v1 Announce Type: cross Abstract: Teaching abstract theoretical computer science (TCS) concepts such as algorithm analysis and complexity theory is challenging because students must handle formal proofs and asymptotic reasoning that conventional resources rarely...

<https://arxiv.org/abs/2609.14572>

35 CONF

Diagnosing Temporal Misalignment in Multichannel Time-Series Classification with Minimum Description Length

ArXiv CS.AI 50%

arXiv:2609.14595v1 Announce Type: cross Abstract: Multichannel time-series classification commonly assumes synchronized sensor streams, although latency, clock drift, and preprocessing can introduce relative delays during data collection or after deployment. Existing...

<https://arxiv.org/abs/2609.14595>

36 **CONF****CompCQR: Compositional Query Generation for Training-Free Conversational Search**

ArXiv CS.AI 50%

arXiv:2609.14646v1 Announce Type: cross Abstract: Multi-turn interactions with LLMs are becoming increasingly common in information-seeking scenarios. However, user queries are often ambiguous and context-dependent, making them ill-suited for direct use as retriever queries....

<https://arxiv.org/abs/2609.14646>
37 **CONF****Natural Language Knowledge Graph Query Execution: Leveraging Controlled Semantics in the LLM Context Window**

ArXiv CS.AI 50%

arXiv:2609.14652v1 Announce Type: cross Abstract: Large Language Model (LLM) applications often transfer domain concepts into the model's context informally, through prompt prose, schema dumps, and examples. We show that for database queries, data model concepts pass to LLMs...

<https://arxiv.org/abs/2609.14652>
38 **CONF****Online Language Adaptive Sampling for Better Distributed Cross-lingual Gains**

ArXiv CS.AI 50%

arXiv:2609.14969v1 Announce Type: cross Abstract: Realignment is a promising approach for improving the cross-lingual transfer ability of multilingual language models, particularly for extremely low-resource languages (LRLs). However, existing realignment methods rely on uniform...

<https://arxiv.org/abs/2609.14969>
39 **CONF****TriCaLRAG: A Three-Strategy, Retrieval-Augmented Benchmark for On-Premise LLM-Based Root Cause Analysis in AIOps**

ArXiv CS.AI 50%

arXiv:2609.14762v1 Announce Type: cross Abstract: Cloud-hosted large language models (LLMs) are increasingly used for root cause analysis (RCA) in AIOps pipelines, but they introduce data privacy risk, network latency, and per-query cost that scale poorly with production log...

<https://arxiv.org/abs/2609.14762>
40 **CONF****Loop-Back Authority in LLM Agent Teams: A Paired Experiment on Flat and Hierarchical Coordination**

ArXiv CS.AI 50%

arXiv:2609.14767v1 Announce Type: cross Abstract: Hierarchical orchestration, in which a Manager agent reviews worker output and can send it back for revision, is the default coordination pattern in production multi-agent LLM frameworks. Classical organizational theory predicts...

<https://arxiv.org/abs/2609.14767>
41 **CONF****The Stochastic Deputy: Structural Tenant Isolation for Tool-Using LLM Agents**

ArXiv CS.AI 50%

arXiv:2609.14780v1 Announce Type: cross Abstract: Multi-tenant tools commonly accept a tenant identifier and validate it against the caller's entitlement. For a large language model (LLM) agent, that pattern delegates resource selection to a process whose context may contain...

<https://arxiv.org/abs/2609.14780>

42 CONF

Route, Don't Fix: Regime-Dependent Decoding Correction and a Trajectory-Gated Router for Reliable Clinical LLM Answer Selection

ArXiv CS.AI 50%

arXiv:2609.14825v1 Announce Type: cross Abstract: Large language models (LLMs) are often deemed unsafe for clinical question answering because of their tendency to hallucinate. Retrieval augmentation, fine-tuning, and external verifiers require new infrastructure that clinical...

<https://arxiv.org/abs/2609.14825>

43 CONF

Efficiency Hallucination: Formalizing and Measuring Behavioral Calibration in LLM-Based Code Optimization

ArXiv CS.AI 50%

arXiv:2609.14839v1 Announce Type: cross Abstract: The integration of Large Language Models (LLMs) into automated code optimization introduces a critical reliability risk we term the Efficiency Hallucination: an LLM's tendency to issue non-functional mutations with...

<https://arxiv.org/abs/2609.14839>

44 CONF

A Responsive Present, a Shared Past, a Social Other: Teens' Overreliance on Companion AI Chatbots

ArXiv CS.AI 50%

arXiv:2609.14843v1 Announce Type: cross Abstract: AI companions provide socially engaging interaction through availability, personalization, memory, roleplay, and emotionally responsive language. For teens, these systems may support sensitive self-disclosure, identity...

<https://arxiv.org/abs/2609.14843>

45 CONF

LLMs as Oracles: Reliance on LLMs for Subjective Personal Questions

ArXiv CS.AI 50%

arXiv:2609.14849v1 Announce Type: cross Abstract: We characterize how people are turning to LLMs as oracles: all-knowing authorities on subjective personal questions. Motivated by risks to users' autonomy and well-being, we develop a typology and LLM-based methods to measure...

<https://arxiv.org/abs/2609.14849>

46 CONF

Interpolation Is Not Invariance: Pair Count Is Not Coverage in Transformation Audits

ArXiv CS.AI 50%

arXiv:2609.14870v1 Announce Type: cross Abstract: Counting equivalent pairs is a common way to report transformation-audit coverage, but it can substantially overstate the constraints imposed by an audit: pairs generated from the same semantic object are correlated, and complete...

<https://arxiv.org/abs/2609.14870>

47 CONF

SeqMaestro: From nucleotide sequences to biological hypotheses through interpretable machine learning

ArXiv CS.AI 50%

arXiv:2609.14882v1 Announce Type: cross Abstract: Nucleotide sequence analysis is central to problems spanning regulatory genomics, evolutionary biology, and phenotype prediction. Classical bioinformatics methods extract interpretable sequence properties such as motifs and k-mer...

<https://arxiv.org/abs/2609.14882>

48 **CONF****PeerPen: AI-Assisted Writing for Online Mental Health Peer Support****ArXiv CS.AI** **50%**

arXiv:2609.14886v1 Announce Type: cross Abstract: Online mental health communities thrive on peer support, yet those who volunteer to help often lack formal training and may struggle to articulate supportive responses. AI co-writing could lower this barrier; however, peer...

<https://arxiv.org/abs/2609.14886>

49 **CONF****IMPACT-VLA: Interaction-aware Multimodal Propagation Attribution via Counterfactual Trajectories for Vision-Language-Action Policies****ArXiv CS.AI** **50%**

arXiv:2609.15005v1 Announce Type: cross Abstract: Vision-Language-Action (VLA) policies perform robot manipulation tasks using multimodal inputs such as visual observations, proprioceptive states, and language instructions. However, it remains unclear at which execution stages...

<https://arxiv.org/abs/2609.15005>

50 **CONF****Steering Generative Robot Policies with Lexicographic Preferences****ArXiv CS.AI** **50%**

arXiv:2609.15014v1 Announce Type: cross Abstract: Pretrained generative robot policies can produce effective behaviors across diverse environments, but deployment can lead to requirements and preferences that may not have been represented during training. Furthermore, at...

<https://arxiv.org/abs/2609.15014>

51 **CONF****Pick Your Poison: Learning to Select Poison Sets for Stronger LLM Backdoor Attacks****ArXiv CS.AI** **50%**

arXiv:2609.15029v1 Announce Type: cross Abstract: Backdoor poisoning attacks add poisoned examples to otherwise-clean finetuning data, pairing a trigger with a target behavior that the model learns to produce when the trigger appears. Existing evaluations typically fix the...

<https://arxiv.org/abs/2609.15029>

52 **CONF****Validating Hybrid-State Cache Recovery for GLM-5.3-Flash with vLLM and LMCache****ArXiv CS.AI** **50%**

arXiv:2609.15030v1 Announce Type: cross Abstract: External cache transfers can succeed while a hybrid language model resumes from an inconsistent state. We examine the full 45-layer GLM-5.3-Flash model, using the RedHatAI/ GLM-5.3-Flash-NVFP4 quantized checkpoint with vLLM and...

<https://arxiv.org/abs/2609.15030>

53 **CONF****SplitTEE: Improving LLM Inference on Trusted Hardware with Differentially Private GPU Outsourcing****ArXiv CS.AI** **50%**

arXiv:2609.15039v1 Announce Type: cross Abstract: User prompts provided to large language models (LLMs) may contain sensitive or private information that can be misused by remotely deployed models, such as through inadvertent memorization during retraining. One way to protect...

<https://arxiv.org/abs/2609.15039>

54 CONF

Rethinking Procedural Audio Pre-training: Source Scaling and Objective Adaptation

ArXiv CS.AI 50%

arXiv:2609.15067v1 Announce Type: cross Abstract: Procedural audio has emerged as a viable source for transferable audio representation learning, but its design principles remain unclear. We revisit two questions: how a procedural source should be scaled, and whether training...

<https://arxiv.org/abs/2609.15067>

55 CONF

Beyond Numerical Time Series: A Unified Benchmark for Multimodal Forecasting with Heterogeneous Context

ArXiv CS.AI 50%

arXiv:2609.15087v1 Announce Type: cross Abstract: Most time series forecasting benchmarks remain numerical-centric and provide limited support for evaluating contextual information that shapes real-world temporal dynamics. Existing multimodal benchmarks also suffer from limited...

<https://arxiv.org/abs/2609.15087>

56 CONF

Generate to Explore, Select to Exploit: Aligning LLM-based Headline Generation with Personalized Recommendation

ArXiv CS.AI 50%

arXiv:2609.15094v1 Announce Type: cross Abstract: In industrial recommendation feeds, presenting a static headline for an item often fails to satisfy the diverse, multimodal interests of the user population, particularly suppressing the needs of long-tail audiences. While Large...

<https://arxiv.org/abs/2609.15094>

57 CONF

ChatGPT Images 2.5 in the Wild: A Launch-Period Dataset and Detector Evaluation

ArXiv CS.AI 50%

arXiv:2609.15100v1 Announce Type: cross Abstract: An image tool can change its underlying generator while retaining its public name, making version attribution from online posts ambiguous. We study this problem after the ChatGPT Images 2.5 launch. Our frozen collection contains...

<https://arxiv.org/abs/2609.15100>

58 CONF

MoME: Mixture-of-Memory Embeddings for Context-Aware Sparse Lookup

ArXiv CS.AI 50%

arXiv:2609.15126v1 Announce Type: cross Abstract: Scaling large language models efficiently has motivated sparse capacity mechanisms such as Mixture-of-Experts and, more recently, conditional memory: token-indexed embedding tables that augment the backbone with cheap parametric...

<https://arxiv.org/abs/2609.15126>

59 CONF

PACE: Progressive Angular-to-Norm Contrastive Embedding

ArXiv CS.AI 50%

arXiv:2609.15152v1 Announce Type: cross Abstract: Multimodal embedding models encode heterogeneous inputs into a shared embedding space, enabling efficient similarity computation across modalities and tasks. Most existing methods optimize cosine-based contrastive objectives,...

<https://arxiv.org/abs/2609.15152>

60 CONF

Interpreting hierarchical organisation of speaker embeddings

ArXiv CS.AI 50%

arXiv:2609.15203v1 Announce Type: cross Abstract: Speaker recognition neural networks learn latent representations (i.e. speaker embeddings) from input utterances to recognise speaker identities. However, the internal mechanisms of these networks remain largely opaque,...

<https://arxiv.org/abs/2609.15203>

61 CONF

Failure-Guided Co-Evolution of Prompts and Training Data

ArXiv CS.AI 50%

arXiv:2609.15209v1 Announce Type: cross Abstract: Automatic prompt optimization (APO) improves language-model programs by revising prompts from task feedback, yet it typically holds its training data fixed. Repeatedly optimizing against the same instances confines feedback to...

<https://arxiv.org/abs/2609.15209>

62 CONF

Math for AI safety: an invitation for mathematicians

ArXiv CS.AI 50%

arXiv:2609.15289v1 Announce Type: cross Abstract: Artificial intelligence threatens to outrun human understanding and control. New mathematics is needed to design AI that is legible, steerable, and cooperative with humanity. I organize this invitation by mathematical field, so...

<https://arxiv.org/abs/2609.15289>

63 CONF

The Universe of Universes: Benefit Yield Functions, Implosion Thresholds, and Infrastructure-Aware Optimization in Multi-LLM Systems

ArXiv CS.AI 50%

arXiv:2609.15314v1 Announce Type: cross Abstract: We introduce the Universe of Universes (UoU) framework, which treats the full ecosystem of major large language models (LLMs) as a structured retrieval corpus and proposes a compositional Automated Reasoning (AR) and Machine...

<https://arxiv.org/abs/2609.15314>

64 CONF

Divide, Consult, Conquer: Capability Laundering Through Aligned LLMs

ArXiv CS.AI 50%

arXiv:2609.15383v1 Announce Type: cross Abstract: Language model safety is typically evaluated one interaction at a time. We show that a weaker, unaligned model can split a harmful task into benign-looking subproblems, consult a stronger aligned model independently on each, and...

<https://arxiv.org/abs/2609.15383>

65 CONF

On the role of the tokenizer in ECG transformer models

ArXiv CS.AI 50%

arXiv:2609.15433v1 Announce Type: cross Abstract: Tokenization determines both the physiological content presented to an ECG Transformer and the sequence over which attention operates. We compare eight tokenization strategies across Transformer, Informer, Reformer, and FEDformer...

<https://arxiv.org/abs/2609.15433>

66 CONF

Authorship attribution and aesthetic evaluation of AI poetry: a case study with Haiku

ArXiv CS.AI 50%

arXiv:2609.15511v1 Announce Type: cross Abstract: This paper investigates the generation and human evaluation of Japanese haiku by contemporary Large Language Models (LLMs), focusing on authorship perception and aesthetic judgment within a constrained poetic form. Using a...

<https://arxiv.org/abs/2609.15511>

67 CONF

Big Brains and Changing Environments: Cause or Consequence?

ArXiv CS.AI 50%

arXiv:2609.15569v1 Announce Type: cross Abstract: Large brains are metabolically costly, and associations with changing environments do not imply they evolved there, as the Cognitive Buffer Hypothesis (CBH) would suggest. They may instead evolve in stable conditions and later...

<https://arxiv.org/abs/2609.15569>

68 CONF

A Unified Vision-Language Model for PSMA PET/CT Report Generation, Visual Question Answering, and Lesion Segmentation

ArXiv CS.AI 50%

arXiv:2609.15603v1 Announce Type: cross Abstract: Accurate PSMA PET/CT interpretation is central to prostate cancer management, yet existing PET/CT AI models typically address isolated tasks. We propose a unified PSMA PET/CT vision-language model for report generation, visual...

<https://arxiv.org/abs/2609.15603>

69 CONF

Beyond AI Literacy: A Structured Review and Exploratory Meta-Analysis of Measures for Competent Generative-AI Use

ArXiv CS.AI 50%

arXiv:2609.15624v1 Announce Type: cross Abstract: Researchers assessing competent generative-AI use at work must choose among self-reports, objective tests, and measures of oversight and reliance. We conducted a structured, seeded review of 24 focal empirical publications,...

<https://arxiv.org/abs/2609.15624>

70 CONF

CLQT: A Closed-Loop, Cost-Aware, Strategy-Consistent Benchmark for Diagnostic Evaluation of LLM Portfolio-Management Agents

ArXiv CS.AI 50%

arXiv:2606.29771v3 Announce Type: replace Abstract: LLM agents are increasingly cast as autonomous portfolio managers, yet the dominant evaluation idiom, a leaderboard of returns over a fixed window, certifies neither the soundness of an agent's process nor the durability of its...

<https://arxiv.org/abs/2606.29771>

71 CONF

Kaininja: Extending Native 3D Generators to the Part Level

ArXiv CS.AI 50%

arXiv:2609.15659v1 Announce Type: cross Abstract: Native 3D generators turn one image into a single mesh. TRELLIS.2 and its peers deliver high-fidelity non-watertight geometry with materials, but the output is one fused object, while downstream work such as editing, rigging and...

<https://arxiv.org/abs/2609.15659>

72 CONF

CIDERS: Cloud-Edge LLM Collaborative Learning via Accelerating Personalized Bilevel Optimization

ArXiv CS.AI 50%

arXiv:2609.15664v1 Announce Type: cross Abstract: Amid the rapid advancement of physical-world intelligence, cloud-edge collaborative large language models (LLMs) have emerged as a promising roadmap for practical LLM deployment. However, existing cloud-edge paradigms struggle to...

<https://arxiv.org/abs/2609.15664>

73 CONF

Circuit-MLLM: Topological Logic-Guided Latent-Space Visual Reasoning for Circuit Schematic Understanding

ArXiv CS.AI 50%

arXiv:2609.15668v1 Announce Type: cross Abstract: Through pre-training on extensive text and image datasets, current multi-modal large language models (MLLMs) achieve strong performance on general tasks. However, circuit schematics present a unique challenge for MLLMs due to...

<https://arxiv.org/abs/2609.15668>

74 CONF

Don't Send What You Don't Need: Question-Guided Token Pruning as a Privacy Defense for Vision-Language Models

ArXiv CS.AI 50%

arXiv:2609.15671v1 Announce Type: cross Abstract: Visual Question Answering (VQA) with Vision-Language Models (VLMs) is increasingly used in privacy-sensitive and bandwidth-constrained settings. Federated Learning (FL), Split Learning (SL), and U-Shaped Split Learning (USL) keep...

<https://arxiv.org/abs/2609.15671>

75 CONF

More Than Just Access: Generative AI as Communication Intermediary for Blind and Low-Vision Users

ArXiv CS.AI 50%

arXiv:2609.15696v1 Announce Type: cross Abstract: Generative AI (GenAI) tools are increasingly woven into how blind and low-vision (BLV) people communicate, not only with digital information, but with the physical world and with other people. Tools such as ChatGPT, Google...

<https://arxiv.org/abs/2609.15696>

76 CONF

Bench2Dex: Benchmarking Visuo-Tactile Bimanual Dexterous Manipulation Across Dexterous Hands

ArXiv CS.AI 50%

arXiv:2609.15726v1 Announce Type: cross Abstract: Tactile sensing provides contact information that can be difficult to infer from vision alone, but tactile hardware for dexterous hands has not converged to a common design. Dexterous hands differ in finger structure, contact...

<https://arxiv.org/abs/2609.15726>

77 CONF

CiteGuard-RAG: A Validation-Centered AI System for Evidence-Grounded Question Answering

ArXiv CS.AI 50%

arXiv:2609.15830v1 Announce Type: cross Abstract: Retrieval-augmented generation (RAG) can improve access to complex information; however, retrieving evidence alone does not ensure that answers are grounded, citation-valid, or appropriately refused. This paper introduces...

<https://arxiv.org/abs/2609.15830>

78 CONF

Learning Multimodal One-step Flow Policy via Value-weighted Optimal Transport

ArXiv CS.AI 50%

arXiv:2609.15883v1 Announce Type: cross Abstract: Offline reinforcement learning aims to learn a policy solely from fixed datasets, which often contain multimodal action distributions. Flow policies can naturally represent such multimodal behaviors, but learning an efficient...

<https://arxiv.org/abs/2609.15883>

79 CONF

Anatomical Grounding and Leakage-Aware Multimodal Contrastive Learning for Alzheimer's Disease Classification from Structural MRI

ArXiv CS.AI 50%

arXiv:2609.15888v1 Announce Type: cross Abstract: Deep networks trained on structural MRI for Alzheimer's disease (AD) staging often reach reasonable accuracy while attending to anatomically irrelevant regions, and multimodal models that add clinical tables frequently rely on...

<https://arxiv.org/abs/2609.15888>

80 CONF

Vulnerability Localization Benchmark: Measuring Agentic Security Analysis at Repository Scale

ArXiv CS.AI 50%

arXiv:2609.15939v1 Announce Type: cross Abstract: Language-model agents increasingly operate over complete software repositories, yet cybersecurity evaluations primarily measure whether they can detect, reproduce, or repair vulnerabilities rather than whether they can locate the...

<https://arxiv.org/abs/2609.15939>

81 CONF

The Router Within: Eliciting Native Skill Routing from a Frozen LLM

ArXiv CS.AI 50%

arXiv:2609.15982v1 Announce Type: cross Abstract: Skills extend an LLM agent beyond its parametric knowledge, and the gain they promise rests on picking the right one. Deployed harnesses route by preloading every skill's metadata into the context, which disperses the agent's...

<https://arxiv.org/abs/2609.15982>

82 CONF

Estimating Uncertain Spatial Relationships in Robotics

ArXiv CS.AI 50%

arXiv:1304.3111v2 Announce Type: replace Abstract: In this paper, we describe a representation for spatial information, called the stochastic map, and associated procedures for building it, reading information from it, and revising it incrementally as new information is...

<https://arxiv.org/abs/1304.3111>

83 CONF

DeepFeature: LLM-Empowered Context-aware Feature Generation for Wearable Biosignals

ArXiv CS.AI 50%

arXiv:2512.08379v3 Announce Type: replace Abstract: Biosignals collected from wearable devices are widely utilized in healthcare applications. Machine learning models used in these applications often rely on features extracted from biosignals due to their effectiveness, lower...

<https://arxiv.org/abs/2512.08379>

84 CONF

Modality-Guided Mixture of Structured Experts with Entropy-Triggered Routing for Multimodal Recommendation

ArXiv CS.AI 50%

arXiv:2602.20723v3 Announce Type: replace Abstract: Multimodal recommenders combine collaborative behavior with visual and textual item evidence, whose usefulness varies across user-item interactions. Independently trained source-specific diagnostic probes partition held-out...

<https://arxiv.org/abs/2602.20723>

85 CONF

NeuroProlog: Multi-Task Fine-Tuning for Neurosymbolic Mathematical Reasoning via the Cocktail Effect

ArXiv CS.AI 50%

arXiv:2603.02504v3 Announce Type: replace Abstract: Large Language Models (LLMs) achieve strong performance on natural language tasks but remain unreliable in mathematical reasoning, frequently generating fluent yet logically inconsistent solutions. We present...

<https://arxiv.org/abs/2603.02504>

86 CONF

FundaPod: A Multi-Persona Agent Pod Architecture with Knowledge Graph Memory for AI-Assisted Fundamental Investment Research

ArXiv CS.AI 50%

arXiv:2605.27864v5 Announce Type: replace Abstract: Large language models (LLMs) are increasingly applied in finance, yet most existing work emphasizes trading signals or financial NLP tasks centered on prediction. Institutional fundamental research, by contrast, requires human...

<https://arxiv.org/abs/2605.27864>

87 CONF

Modeling Social Dynamics with an LLM-Enabled Agent Based Network-Dynamic (LAND) Model

ArXiv CS.AI 50%

arXiv:2608.00929v2 Announce Type: replace Abstract: Social dynamics encode the process in which individual network and discourse interactions aggregate into collective influence, narrative dominance and coordinate behavior. This paper uses the the GhostField architecture, a...

<https://arxiv.org/abs/2608.00929>

88 CONF

FrogNano: Training a 4B Coding Agent via Online Task Synthesis

ArXiv CS.AI 50%

arXiv:2609.07925v3 Announce Type: replace Abstract: We present FrogNano, a 4B coding agent designed to tackle software engineering (SWE) tasks efficiently and effectively, even under resource-constrained environments. It is post-trained exclusively via RL on around 1,500 SWE...

<https://arxiv.org/abs/2609.07925>

89 CONF

Beyond Generation and Accuracy: Diagnosing and Enhancing Visual Chain-of-Thought for Geometry Problem Solving

ArXiv CS.AI 50%

arXiv:2609.12606v2 Announce Type: replace Abstract: While multimodal reasoning has advanced rapidly, solving complex geometry problems critically hinges on active visual assistance, such as constructing auxiliary lines, spurring the rise of Visual Chain-of-Thought (VCoT)....

<https://arxiv.org/abs/2609.12606>

90 CONF

Are explainable AI (XAI) evaluation strategies aligned? Comparing subjective, objective, and mathematical evaluation measures using saliency maps

ArXiv CS.AI 50%

arXiv:2504.17023v2 Announce Type: replace-cross Abstract: The evaluation of explainable AI (XAI) approaches often relies on three families of methods: subjective measures (e.g., questionnaires on trust or satisfaction), objective measures (e.g., task performance metrics), and...

<https://arxiv.org/abs/2504.17023>

91 CONF

Discovering Hierarchy-Grounded Domains with Adaptive Granularity for Clinical Domain Generalization

ArXiv CS.AI 50%

arXiv:2506.06977v4 Announce Type: replace-cross Abstract: Domain generalization has become a critical challenge in predictive healthcare, where different patient groups exhibit shifting data distributions that degrade model performance. Still, regular domain generalization...

<https://arxiv.org/abs/2506.06977>

92 CONF

Perceptual Reality Transformer: What Must an Illustration Preserve?

ArXiv CS.AI 50%

arXiv:2508.09852v2 Announce Type: replace-cross Abstract: How can models help people communicate unusual perceptual experiences without changing what they mean? A recognizable image is only part of the answer: accounts also express vividness, duration, uncertainty, and emotion....

<https://arxiv.org/abs/2508.09852>

93 CONF

Concertina: Data-Centric Adaptive Pipeline Parallelism for Efficient Heterogeneous Long-Context LLM Training

ArXiv CS.AI 50%

arXiv:2509.21275v5 Announce Type: replace-cross Abstract: Long context training is crucial for extending LLM context windows. Existing schemes, such as sequence parallelism, incur substantial communication overhead. Pipeline parallelism (PP) reduces this cost, but its...

<https://arxiv.org/abs/2509.21275>

94 CONF

Freeze, Share, Shrink: Rethinking the Action Backbone in Diffusion Policies

ArXiv CS.AI 50%

arXiv:2511.12101v3 Announce Type: replace-cross Abstract: Many recent Vision-Language-Action models employ diffusion or flow-matching backbones with hundreds of millions of parameters for action generation. However, unlike image synthesis where the output spans millions of...

<https://arxiv.org/abs/2511.12101>

95 CONF

Donors and Recipients: On Asymmetric Transfer Across Tasks and Languages with Parameter-Efficient Fine-Tuning

ArXiv CS.AI 50%

arXiv:2511.13368v3 Announce Type: replace-cross Abstract: Large language models (LLMs) perform strongly across tasks and languages, yet how improvements in one task or language affect other tasks and languages remains poorly understood. We conduct a controlled LoRA fine-tuning...

<https://arxiv.org/abs/2511.13368>

96 CONF

SDUs DAISY: A Benchmark for Danish Culture

ArXiv CS.AI 50%

arXiv:2601.19930v2 Announce Type: replace-cross Abstract: We introduce Daisy, a factual knowledge benchmark for Danish cultural heritage, based on curated topics from the Danish Culture Canon 2006. For each artifact in the culture canon, we query the corresponding Wikipedia page...

<https://arxiv.org/abs/2601.19930>

97 CONF

How Semantically Stable Are LLM Refusals? Measuring Confusion in Local Safety Boundaries

ArXiv CS.AI 50%

arXiv:2512.01037v3 Announce Type: replace-cross Abstract: As safety alignment becomes standard in large language models, refusal behavior has become an important part of model reliability. However, models may still reject benign prompts, especially when the wording resembles...

<https://arxiv.org/abs/2512.01037>

98 CONF

Hearing the Unspoken: Language Model Priors for Acoustic Adversarial Attacks

ArXiv CS.AI 50%

arXiv:2606.06833v2 Announce Type: replace-cross Abstract: Automatic Speech Recognition (ASR) systems operating in real-time settings must process acoustic input under strict temporal constraints, where transcription decisions are inherently made on incomplete information. This...

<https://arxiv.org/abs/2606.06833>

99 CONF

Large Language Models As Shannon Lossy Compressors Not Solomonoff Induction Estimators: The Singularity Is Not Near Without Symbolic Model Synthesis in Program Space

ArXiv CS.AI 50%

arXiv:2601.05280v4 Announce Type: replace-cross Abstract: On the one hand, the question of whether Large Language Models (LLMs) are Solomonoff induction estimators has become an explicit question at the intersection of Algorithmic Information Theory (AIT) and Machine Learning...

<https://arxiv.org/abs/2601.05280>

00 CONF

Confident Rankings with Fewer Items: Adaptive LLM Evaluation with Continuous Scores

ArXiv CS.AI 50%

arXiv:2601.13885v2 Announce Type: replace-cross Abstract: Computerized Adaptive Testing (CAT) has proven effective for efficient LLM evaluation on multiple-choice benchmarks, but modern LLM evaluation increasingly relies on generation tasks where outputs are scored continuously...

<https://arxiv.org/abs/2601.13885>

01 CONF

Human Values in a Single Sentence: Moral Presence, Hierarchies, and Transformer Ensembles on the Schwartz Continuum

ArXiv CS.AI 50%

arXiv:2601.14172v4 Announce Type: replace-cross Abstract: We study neural multi-label classification under severe label imbalance through sentence-level detection of the 19 refined Schwartz human values in 74k English news and manifesto sentences (ValueEval'24 corpus). Each...

<https://arxiv.org/abs/2601.14172>

02 CONF

TIDAL: Temporally Interleaved Diffusion and Action Loop for High-Frequency VLA Control

ArXiv CS.AI 50%

arXiv:2601.14945v3 Announce Type: replace-cross Abstract: Large-scale Vision-Language-Action (VLA) models offer semantic generalization but suffer from high inference latency because they adopt a low-frequency batch-and-execute paradigm. This frequency mismatch creates an...

<https://arxiv.org/abs/2601.14945>

03 CONF

Rubrics as an Attack Surface: Stealthy Preference Drift in LLM Judges

ArXiv CS.AI 50%

arXiv:2602.13576v2 Announce Type: replace-cross Abstract: Evaluation and alignment pipelines for large language models increasingly rely on LLM-based judges, whose behavior is guided by natural-language rubrics and validated on benchmarks. We identify a previously...

<https://arxiv.org/abs/2602.13576>

04 CONF

Context-Dependent Affordance Reports in Vision-Language Models

ArXiv CS.AI 50%

arXiv:2603.04419v3 Announce Type: replace-cross Abstract: Vision-language models produce different object and use descriptions under different persona prompts, but low overlap alone does not identify an affordance effect. We audit an earlier seven-prompt study and add...

<https://arxiv.org/abs/2603.04419>

05 CONF

GP-VM $\times$ SMA: Benchmarking General-Purpose Vision Models and Specialized Architectures for 2D Medical Image Segmentation

ArXiv CS.AI 50%

arXiv:2603.13044v2 Announce Type: replace-cross Abstract: Medical image segmentation (MIS) is a fundamental component of computer-assisted diagnosis and clinical decision support. Over the past decade, numerous architectures specifically tailored to medical imaging have emerged...

<https://arxiv.org/abs/2603.13044>

06 CONF

Measuring and Exploiting Contextual Bias in LLM-Assisted Security Code Review

ArXiv CS.AI 50%

arXiv:2603.18740v3 Announce Type: replace-cross Abstract: Automated Code Review (ACR) systems integrating Large Language Models (LLMs) are increasingly adopted in software development workflows, ranging from interactive assistants to autonomous agents in CI/CD pipelines. In this...

<https://arxiv.org/abs/2603.18740>

07 CONF

AI Psychosis: Does Conversational AI Amplify Delusion-Related Language?

ArXiv CS.AI 50%

arXiv:2603.19574v2 Announce Type: replace-cross Abstract: Conversational AI systems are increasingly used for personal reflection and emotional disclosure, raising concerns about their effects on vulnerable users. Recent anecdotal reports suggest that prolonged interactions with...

<https://arxiv.org/abs/2603.19574>

08 CONF

Can We Still Trace L1 Signals? Investigating the Resilience of Native Language Signals in the LLM Era

ArXiv CS.AI 50%

arXiv:2604.08568v3 Announce Type: replace-cross Abstract: The widespread use of LLM-based writing assistance has raised an interesting question about the homogenization of English. As LLMs tend to revise texts toward mainstream English conventions reflected in their training...

<https://arxiv.org/abs/2604.08568>

09 CONF

An AI Agent Execution Environment to Safeguard User Data

ArXiv CS.AI 50%

arXiv:2604.19657v2 Announce Type: replace-cross Abstract: AI agents promise to serve as general-purpose personal assistants for their users, which requires them to have access to private user data (e.g., personal and financial information). This poses a serious risk to security...

<https://arxiv.org/abs/2604.19657>

10 CONF

IntraGuard: Committee-Side Defenses Against Review Outsourcing to Commercial Chatbots

ArXiv CS.AI 50%

arXiv:2605.05271v2 Announce Type: replace-cross Abstract: LLMs become increasingly capable, editorial boards and program committees are growing concerned about reviewers who fully outsource peer review to commercial chatbots. This concern stems from prior findings that current...

<https://arxiv.org/abs/2605.05271>

11 CONF

Agentic AI for Gravitational Wave Data Analysis: A Head-to-Head Comparison of Coding Agents Executing a Matched Filter Pipeline on Einstein Telescope Simulated Data

ArXiv CS.AI 50%

arXiv:2605.28916v3 Announce Type: replace-cross Abstract: We report a methodological study of agentic AI in gravitational-wave data analysis: two systems, Claude Code (Anthropic) and Codex (OpenAI), autonomously executed the same simple end-to-end pipeline on Einstein Telescope...

<https://arxiv.org/abs/2605.28916>

12 CONF

Compute Allocation for Self-Evolving LLMs: From Depth-Breadth to Multi-Armed Bandits

ArXiv CS.AI 50%

arXiv:2605.29268v3 Announce Type: replace-cross Abstract: LLM-guided evolutionary search (Evolve systems) has reached state-of-the-art results on mathematical and combinatorial tasks, yet most existing systems report only the best of many runs and leave the run-to-run...

<https://arxiv.org/abs/2605.29268>

13 CONF

Attention Calibration for Position-Fair Dense Retrieval

ArXiv CS.AI 50%

arXiv:2606.02737v2 Announce Type: replace-cross Abstract: Dense retrieval compresses a passage into a single vector, but this compression is positionally skewed: early content dominates the embedding, and retrieval degrades when the relevant span appears later. Prior work...

<https://arxiv.org/abs/2606.02737>

14 CONF

Qwen-Image-Flash: Rethinking the Training Recipe for Few-Step Distillation

ArXiv CS.AI 50%

arXiv:2606.03746v3 Announce Type: replace-cross Abstract: Few-step distillation has emerged as a critical component in the development of advanced visual generative foundation models, substantially reducing inference overhead while enabling real-time generation and...

<https://arxiv.org/abs/2606.03746>

15 CONF

Cherry-pick Override: LLM Judges Under-use the Non-Directional Verdicts Their Contract Authorizes

ArXiv CS.AI 50%

arXiv:2606.07834v2 Announce Type: replace-cross Abstract: Evidence-grounded fact verification increasingly uses LLM judges to turn evidence into terminal verdicts. Many task contracts deliberately include non-directional verdicts - Conflicting for materially mixed evidence, Not...

<https://arxiv.org/abs/2606.07834>

16 CONF

How Do Video Foundation Models Encode Intuitive Physics? Probing Across Pretraining Paradigms

ArXiv CS.AI 50%

arXiv:2606.09646v2 Announce Type: replace-cross Abstract: We study whether pretrained video foundation models encode intuitive-physics information in their frozen representations, and how this information varies across model families, layers, and probe types. Using...

<https://arxiv.org/abs/2606.09646>

17 CONF

The Language-Energy Divide: Measuring Energy Costs of Multilingual LLM Inference

ArXiv CS.AI 50%

arXiv:2606.21869v2 Announce Type: replace-cross Abstract: Large language models (LLMs) are increasingly deployed in multilingual settings, yet the energy costs of serving these models across different languages remain poorly understood. We present a systematic study of inference...

<https://arxiv.org/abs/2606.21869>

18 CONF

Interpretable Inverse Design of Metal-Organic Frameworks with Large Language Model Agents

ArXiv CS.AI 50%

arXiv:2606.29459v2 Announce Type: replace-cross Abstract: Inverse design of metal-organic frameworks (MOFs) requires navigating combinatorial spaces with costly property labels and opaque machine-learning models. We introduce LLM4MOF, a closed-loop multi-agent framework that...

<https://arxiv.org/abs/2606.29459>

19 CONF

PEEK: Predictive Queue-Informed KV Cache Management for LLM Serving

ArXiv CS.AI 50%

arXiv:2607.02525v2 Announce Type: replace-cross Abstract: We present PEEK, a lightweight scheduling and eviction framework for both online (streaming) and offline (batch) LLM serving; this paper focuses on the online regime. PEEK maintains an incremental radix tree over the...

<https://arxiv.org/abs/2607.02525>

20 **CONF****Energy Accuracy Is Not Enough: A Structure-Aware Benchmark and Evaluation Protocol for Quantum Architecture Search****ArXiv CS.AI** **50%**

arXiv:2607.04845v2 Announce Type: replace-cross Abstract: Quantum architecture search for molecular ground-state estimation is commonly evaluated through energy accuracy, which does not describe circuit cost or the physical properties of the prepared state. We introduce...

<https://arxiv.org/abs/2607.04845>

21 **CONF****SMetric: Rethink LLM Scheduling for Serving Agents with Balanced Session-centric Scheduling****ArXiv CS.AI** **50%**

arXiv:2607.08565v2 Announce Type: replace-cross Abstract: LLM scheduling is critical to serving, yet how well existing designs fit agentic serving--where agents, not humans, issue the requests--remains unclear. Agents shift the workload in two ways: they consume many more tokens...

<https://arxiv.org/abs/2607.08565>

22 **CONF****KREL: Automatic Medical Coding via Knowledge-Guided Reasoning over Clinical Evidence with LLMs****ArXiv CS.AI** **50%**

arXiv:2608.20887v2 Announce Type: replace-cross Abstract: Automatic Medical Coding (AMC), which assigns standardized International Classification of Diseases (ICD) codes to clinical notes, is essential for medical reimbursement, quality reporting, and clinical research. Existing...

<https://arxiv.org/abs/2608.20887>

23 **CONF****A Few Pages of Markdown: Committed AI Configuration and Lower Quality Cost after Coding-Agent Adoption****ArXiv CS.AI** **50%**

arXiv:2608.25241v2 Announce Type: replace-cross Abstract: Coding agents increase development velocity but also technical debt. Prior work reports only average effects across adopters, hiding wide differences between teams. We introduce RAMP (Repository AI Maturity Profile), a...

<https://arxiv.org/abs/2608.25241>

24 **CONF****Skynet: Workflow-Level Anomaly Detection for Agentic AI via Semantic and Structural Modeling****ArXiv CS.AI** **50%**

arXiv:2609.06835v2 Announce Type: replace-cross Abstract: Agentic AI systems execute complex tasks through long-horizon workflows of planning, tool use, and multi-agent coordination. Task failures in these systems often originate from a single step, such as an injected prompt or...

<https://arxiv.org/abs/2609.06835>

25 **CONF****Proprioception-Anchored Cross-Modal Pretraining for Zero-Shot Sim-to-Real Contact-Rich Assembly****ArXiv CS.AI** **50%**

arXiv:2609.07534v2 Announce Type: replace-cross Abstract: Contact-rich assembly remains challenging because it requires submillimeter spatial accuracy and reliable interpretation of forces during sustained contact. Although simulation-based reinforcement learning offers a...

<https://arxiv.org/abs/2609.07534>

26 CONF

Multimodal Duplex Interaction Agent

ArXiv CS.AI 50%

arXiv:2609.08977v3 Announce Type: replace-cross Abstract: In this work, we present Gander, a native multimodal duplex interaction model that builds on MiniCPM-o 4.5 and is further adapted for realtime interaction with an asynchronous agent loop. In contrast to conventional turn...
<https://arxiv.org/abs/2609.08977>