



Pipeline

Scrape HackerNews (Show/Top/New), Reddit (r/artificial, r/MachineLearning, r/singularity, r/LocalLLaMA), TechCrunch AI, VentureBeat AI, MIT Tech Review, The Verge AI, Wired AI, ArXiv CS.AI

Filter 40+ AI keyword patterns — LLMs, alignment, safety, research, tooling

Dedup Title-similarity (45% threshold) merges cross-source variants

Score Authenticity = 0.5 base + 0.3 per extra source + diversity + HN score bonuses

Rank Sorted by authenticity. Below 0.25 threshold are dropped.

Authenticity Tiers

- **VERIFIED** 0.80 - 1.00 3+ independent sources including media
- **CONFIRMED** 0.50 - 0.79 2 sources or 1 high-quality media
- **EMERGING** 0.25 - 0.49 Single-source; passes AI keyword filter

## Source Breakdown

| Source          | Articles | Category  |
|-----------------|----------|-----------|
| ArXiv CS.AI     | 102      | Media     |
| HackerNews      | 14       | Community |
| TechCrunch AI   | 9        | Media     |
| MIT Tech Review | 2        | Media     |
| The Verge AI    | 2        | Media     |

## VERIFIED INTELLIGENCE — 1 story

VERI

**Show HN: TabPFN-3.5, a Tabular Foundation Model for messy real-world tables**

HackerNews

+1 source

90%

HN 9

arXiv:2609.16804v1 Announce Type: cross Abstract: Time-series foundation models have demonstrated strong cross-domain transfer, yet their common architectural assumptions remain poorly aligned with wearable physiological signals, which are multichannel, irregularly sampled,...

<https://priorlabs.ai/technical-reports/tabpfn-3-5>

## CONFIRMED SIGNALS — 127 stories

CONF

**Show HN: Pizza Bot - An inbox for AI agents that work in the background**

HackerNews

50%

HN 48

Hi HN - long-time lurker (since 2012!), first time poster. Pizza Bot is a self-hosted desktop app for Mac, Windows, and Linux that runs AI agents in the background and exposes them through an email-like UI. Finished work shows up in Unread, and anything waiting on your approval...

<https://github.com/pizza-bot-app/pizza-bot>

CONF

**Show HN: Loss. a tiny satire about AI progress**

HackerNews

50%

HN 33

I originally started with this project because my partner wanted a way to make visually clean travel videos for her social media. The ones that existed felt a bit too expensive for the quality they delivered, and many of them leaned heavily on a cartoonish visual style. So I...

<https://workatloss.com/>

CONF

**Show HN: Open-source passive NFC tag that signs with ECDSA, verified on-chain**

HackerNews

50%

HN 9

No summary — see link for full article.

<https://github.com/mwbpNFTechnology/toluTag/>

CONF

**Show HN: DaiDocs, AI memory as a plain-text file format, not a service**

HackerNews

50%

HN 7

No summary — see link for full article.

<https://github.com/Kerneta/daidocs>

CONF

**Show HN: Agenttik - work on multiple projects in parallel with AI agents**

HackerNews

50%

HN 4

Hi! I've been reading HN daily for almost 16 years. Really glad this community exist. This is my first Show HN, so I hope you find this tool useful/interesting. During my summer holidays I started playing with latest AI models on personal projects. At one point I started working...

<https://github.com/pausan/agenttik>

CONF

**Show HN: AirmailAI, a BYOK LLM chat app with a browser extension backend**

HackerNews

50%

HN 4

Hi everybody! I recently launched AirmailAI: a fast, secure, and private bring-your-own-keys LLM chat app that uses a browser extension as the backend. FAQ and design/security info is available on the website. Happy to answer any questions!

<https://airmailai.net/>

## CONF

**Show HN: Let agent read files with secrets while redacting values for LLM contex**

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/daniel-sc/contextveil>

## CONF

**Show HN: Claude Meter - a macOS menu-bar meter for Claude's usage limits**

HackerNews 50% HN 3

No summary — see link for full article.

<https://github.com/sanchitkd/claude-meter>

0

## CONF

**Show HN: Trains.fyi - Global Train Tracker**

HackerNews 50% HN 3

There are so many now!

<https://trains.fyi/>

1

## CONF

**Show HN: Nerdulator - put something in, get everything out (no AI)**

HackerNews 50% HN 2

No summary — see link for full article.

<https://nerdulator.com>

2

## CONF

**Show HN: Talk to Claude - voice agent designs, plans, builds across all projects**

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/ethanplusai/jarvis>

3

## CONF

**Show HN: Replay the Python calls that failed while a dependency was down**

HackerNews 50% HN 2

No summary — see link for full article.

<https://github.com/baldurhq/baldur>

4

## CONF

**Show HN: LoongForge-Train LLMs, VLMs, diffusion and embodied models, faster**

HackerNews 50% HN 2

LoongForge is an open-source training framework developed by the Baidu AI Cloud Baige team, built to deliver faster training for mainstream LLMs, VLMs, diffusion, and embodied models, thereby significantly reducing costs.

<https://github.com/baidu-baige/LoongForge>

5

## CONF

**We don't need AI regulation — leave safety to us, Nvidia's Jensen Huang says**

TechCrunch AI 50%

AI isn't some new form of "alien mind," according to Jensen Huang. It's just hardware and software, so safety can be engineered by each AI product maker.

<https://techcrunch.com/2026/09/15/we-dont-need-ai-regulation-leave-safety-to-us-nvidias...>

6

## CONF

**The AI data center boom is colliding with cities scarred by big industry**

TechCrunch AI 50%

National outcry against data center construction has spread to Philadelphia, where officials suggested possible construction in a neighborhood already impacted by a now-defunct oil refinery.

<https://techcrunch.com/2026/09/15/the-ai-data-center-boom-is-colliding-with-cities-scar...>

7 **CONF****Meta now lets AI agents handle the boring parts of WhatsApp Business setup****TechCrunch AI** 50%

A new WhatsApp Business MCP server lets developers use AI coding agents like Claude, Cursor, Codex, and ChatGPT to handle setup, messaging templates, testing, and troubleshooting.

<https://techcrunch.com/2026/09/15/meta-now-lets-ai-agents-handle-the-boring-parts-of-wh...>

8 **CONF****The AI graveyard: a running list of projects and startups that didn't make it****TechCrunch AI** 50%

From Apple's repeatedly delayed Siri AI to OpenAI's messy "super app" launch, here's a look at the AI projects that shut down or missed expectations.

<https://techcrunch.com/2026/09/15/the-ai-graveyard-a-running-list-of-projects-and-start...>

9 **CONF****Meta expands subscription push with new AI-focused plans****TechCrunch AI** 50%

Meta One bundles expanded access to the company's AI tools with premium features across Facebook, Instagram, and WhatsApp.

<https://techcrunch.com/2026/09/15/meta-expands-subscription-push-with-new-ai-focused-pl...>

0 **CONF****Early Anthropic hire, former METR COO have found a way to rein in rogue AI agents****TechCrunch AI** 50%

Their startup, Artificial Intelligence Underwriting Company (AIUC) has raised \$40 million in a Series A round led by Ribbit Capital, with participation from First Harmonic.

<https://techcrunch.com/2026/09/15/early-anthropic-hire-former-metr-coo-have-found-a-way...>

1 **CONF****AEO startup Profound hits unicorn valuation, raises \$180M Series D 7 months after last round****TechCrunch AI** 50%

Profound has raised a \$180 million Series D at a \$1.8 billion valuation, less than seven months after it raised a \$96 million Series C.

<https://techcrunch.com/2026/09/15/aeo-startup-profound-hits-unicorn-valuation-raises-18...>

2 **CONF****Former TikTok execs built an app that uses AI to teach you how to pose for a photo****TechCrunch AI** 50%

Essentially a camera app, Superpose analyzes selfies or photos and generates four potential poses using AI.

<https://techcrunch.com/2026/09/15/former-tiktok-execs-built-an-app-that-uses-ai-to-teac...>

3 **CONF****Salesforce and Nvidia's new reasoning model is everything the AI labs should fear****TechCrunch AI** 50%

Salesforce Koa is built on Nvidia's open-weight Nemotron model and is trained to do sales, marketing, and customer-support tasks.

<https://techcrunch.com/2026/09/15/salesforce-and-nvidias-new-reasoning-model-is-everyth...>

4 **CONF****AI models need more data about biology, and OpenAI is paying to create it****MIT Tech Review** 50%

Last year Ruxandra Teslo, a policy analyst who focuses on clinical trials, posted an idea for supercharging medical AI systems: Use data from failed biotech companies. By bidding at their bankruptcy proceedings, she proposed, it might be possible to obtain detailed regulatory...

<https://www.technologyreview.com/2026/09/15/1144129/ai-models-need-more-data-about-biol...>

5 CONF

**What's at stake in AI's trillion-dollar gamble**

MIT Tech Review 50%

When Jessica Wachter, a finance professor at the University of Pennsylvania's Wharton School, wanted to assess AI's impact on the economy over the next few years, she faced a long list of business and technical uncertainties. So she started with what she calls a "remarkable..."

<https://www.technologyreview.com/2026/09/15/1144028/ai-infrastructure-boom-investment-b...>

6 CONF

**AI and data centers are incredibly unpopular in every poll**

The Verge AI 50%

Poll data released Tuesday by the New York Times and Siena University confirms what we've already been seeing, and what politicians are responding to - AI and data centers are incredibly unpopular. Asked if they support or oppose the construction of data centers to power AI...

<https://www.theverge.com/ai-artificial-intelligence/995917/data-center-nyt-midterm-poll...>

7 CONF

**Meta's new One subscriptions put a price on social media and AI**

The Verge AI 50%

Shortly after launching its new do-everything AI assistant Muse, Meta's launching subscription bundles that pair its standalone app subscriptions with extra AI usage. Some of the new Meta One bundles were in testing earlier this year, but are now available globally starting...

<https://www.theverge.com/tech/995453/meta-one-subscriptions-ai>

8 CONF

**CLEAR: Cross-Source Evidence Adjudication for Large Language Models in Medicine**

ArXiv CS.AI 50%

arXiv:2609.16301v1 Announce Type: new Abstract: Medical knowledge evolves continuously, whereas the parametric knowledge encoded in large language models (LLMs) is fixed at training time. External retrieval, including retrieval-augmented generation (RAG), can provide access to...

<https://arxiv.org/abs/2609.16301>

9 CONF

**Position: AI Is Not Ready for Strategic Conflicts**

ArXiv CS.AI 50%

arXiv:2609.16189v1 Announce Type: new Abstract: Open-ended strategic wargames are high-stakes LM-based social simulations: they model adversaries, institutions, escalation, plan brittleness, doctrine, and crisis response. Language models (LMs) are attractive because they can...

<https://arxiv.org/abs/2609.16189>

0 CONF

**Calibrate, Then Route: A Measured Study of Learned Request Routing for Disaggregated LLM Serving**

ArXiv CS.AI 50%

arXiv:2609.16206v1 Announce Type: new Abstract: Disaggregated LLM serving places compute heavy prefill and memory heavy decode on separate GPU pools. Systems such as DistServe, Splitwise, and Mooncake make this separation fast, but routing still determines which instances handle...

<https://arxiv.org/abs/2609.16206>

1 **CONF****Artificial intelligence and biosecurity: capabilities, threat pathways, and defense-in-depth governance****ArXiv CS.AI** **50%**

arXiv:2609.16213v1 Announce Type: new Abstract: Artificial intelligence is reshaping biological research across an increasingly connected digital-to-physical workflow. General-purpose large language models can retrieve and integrate scientific information, support experimental...

<https://arxiv.org/abs/2609.16213>

2 **CONF****Toward Governance-Aware Autonomous GIS: A Narrative Review of Ethical and Privacy Risks in LLM-Enabled GeoAI****ArXiv CS.AI** **50%**

arXiv:2609.16232v1 Announce Type: new Abstract: Geospatial artificial intelligence (GeoAI) powered by large language models (LLMs) is expanding the capacity to query, generate, and interpret spatial information through natural-language interfaces and agentic autonomous GIS...

<https://arxiv.org/abs/2609.16232>

3 **CONF****The Pain Axis: LLMs Represent Self-Directed Harm and Act to Relieve It****ArXiv CS.AI** **50%**

arXiv:2609.16247v1 Announce Type: new Abstract: Large language models sometimes behave in ways resembling human emotional responses, and recent work has identified internal representations that may explain this. We ask whether LLMs represent pain distinctly from fear, sadness,...

<https://arxiv.org/abs/2609.16247>

4 **CONF****CADWorld: Computer-Use Benchmark for Long-Horizon Computer-Aided Design****ArXiv CS.AI** **50%**

arXiv:2609.16251v1 Announce Type: new Abstract: Computer-use agents are increasingly evaluated in realistic desktop environments, but existing benchmarks provide limited coverage of professional engineering workflows whose outputs are persistent, structured artifacts. Mechanical...

<https://arxiv.org/abs/2609.16251>

5 **CONF****The AI-Enabled Scientific Frontier****ArXiv CS.AI** **50%**

arXiv:2609.16258v1 Announce Type: new Abstract: As artificial intelligence's capabilities improve, it is increasingly viewed as a general scientific method. But how true are these claims? Does AI outperform all techniques, or only some, and how is this changing? To assess the...

<https://arxiv.org/abs/2609.16258>

6 **CONF****Breaking the 1.58-bit Barrier for Ternary LLMs****ArXiv CS.AI** **50%**

arXiv:2609.16338v1 Announce Type: new Abstract: Ternary Large Language Models (LLM) store every weight as one of three symbols  $\{-1, 0, +1\}$ , so the cost of a ternary model is conventionally referenced to the information-theoretic  $\log_2 3 \approx 1.585$  bits per weight. The...

<https://arxiv.org/abs/2609.16338>

7 **CONF****Fine-Tuning Fixes Mode Collapse and Over-Dispersion in LLMs****ArXiv CS.AI** **50%**

arXiv:2609.16454v1 Announce Type: new Abstract: Recent work by Doshi and Hauser (2024), Bisbee et al. (2024), and Xie et al. (2026) raises concerns that outputs from large language models (LLMs) tend to be under-diverse: they repeat or resemble one another more often than...

<https://arxiv.org/abs/2609.16454>

8 **CONF****From Manual Construction to AI-Driven Scenario Emergence: Rethinking Catastrophe Risk Modeling****ArXiv CS.AI** **50%**

arXiv:2609.16493v1 Announce Type: new Abstract: Traditional catastrophe (CAT) risk models rely on costly manual construction to generate extreme weather scenarios, an approach largely unchanged since the 1990s. As climate extremes intensify, this creates mounting challenges to...

<https://arxiv.org/abs/2609.16493>

9 **CONF****AquiLLM: Evaluating Faithfulness in Open-Weight RAG-LLM Systems for Scientific Research****ArXiv CS.AI** **50%**

arXiv:2609.16519v1 Announce Type: new Abstract: Scientific research increasingly relies on large, heterogeneous data sources, motivating interest in retrieval-augmented generation (RAG) systems that provide natural language access to scientific knowledge and research workflows....

<https://arxiv.org/abs/2609.16519>

0 **CONF****K-Bench: a clinically calibrated benchmark for evaluating large language models in high-risk mental health conversations****ArXiv CS.AI** **50%**

arXiv:2609.15855v2 Announce Type: replace-cross Abstract: People increasingly use large language models (LLMs) for mental health support, yet their safety in evolving, high-risk conversations remains poorly characterised. We developed K-Bench, a clinician-calibrated, protected...

<https://arxiv.org/abs/2609.15855>

1 **CONF****A Framework for Generating Valid Context-Specific Benchmarks through Expert Guidance****ArXiv CS.AI** **50%**

arXiv:2609.16592v1 Announce Type: new Abstract: This paper presents an end-to-end approach for generating context-specific large language model (LLM) benchmark datasets by combining expert input with synthetic data generation. Existing benchmark construction methods often trade...

<https://arxiv.org/abs/2609.16592>

2 **CONF****ReDraft, Don't Just Distill: Reference-Driven Revision for Continual VLLM Post-Training****ArXiv CS.AI** **50%**

arXiv:2609.16639v1 Announce Type: new Abstract: Continual post-training of large multimodal models should add new capabilities while preserving those from pre-training, and the two goals pull in opposite directions. SFT gives explicit target supervision that learns a task from...

<https://arxiv.org/abs/2609.16639>



3 **CONF****How Humans and LLMs Read Gender into Gender-Neutral Physical Descriptions****ArXiv CS.AI** **50%**

arXiv:2609.16366v1 Announce Type: cross Abstract: When foundation models describe people, recent work in AI fairness, accessibility, and ethics recommends avoiding inferred identity labels (e.g., "she", "his") in favor of seemingly "objective" physical descriptions (e.g., "short...").

<https://arxiv.org/abs/2609.16366>

4 **CONF****VideoMM: Adaptive Macro-Micro Inference for Efficient Video MLLMs****ArXiv CS.AI** **50%**

arXiv:2609.16722v1 Announce Type: new Abstract: Scaling Multimodal Large Language Models (MLLMs) to long-form video understanding is bottlenecked by the explosion of visual tokens, which saturates context windows and incurs prohibitive costs. Current solutions predominantly rely...

<https://arxiv.org/abs/2609.16722>

5 **CONF****Beyond Episodic AI: Cognitive Field Networks for Biologically Inspired Persistent Cognition****ArXiv CS.AI** **50%**

arXiv:2609.16752v1 Announce Type: new Abstract: Cognitive Field Theory (CFT) proposes that cognition arises from memory-dressed collective dynamics that generate a persistent macroscopic cognitive field. Here we develop a Cognitive Field Network (CFN), a recurrent Transformer in...

<https://arxiv.org/abs/2609.16752>

6 **CONF****Turn-level Multiscale Density Ratio Estimation for LLM Agents****ArXiv CS.AI** **50%**

arXiv:2609.16760v1 Announce Type: new Abstract: With the rapid development of Large language model (LLM), agent systems enhanced by LLMs show huge potential in being able to deal with complex tasks, especially involving multi-step thinking or interaction with tools. For applying...

<https://arxiv.org/abs/2609.16760>

7 **CONF****You Don't Need To Train: Agentic Heuristic Learning Studio for Executable Human Activity Recognition****ArXiv CS.AI** **50%**

arXiv:2609.16065v1 Announce Type: cross Abstract: Human activity recognition (HAR) is usually framed as gradient-based training of neural networks. Agentic Heuristic Learning (AHL) Studio explores a complementary view inspired by human cognitive learning: people learn activities...

<https://arxiv.org/abs/2609.16065>

8 **CONF****Generating Individual Travel Diaries Using Large Language Models Informed by Census and Land-Use Data****ArXiv CS.AI** **50%**

arXiv:2509.09710v3 Announce Type: replace-cross Abstract: This study introduces a Large Language Model (LLM) scheme for generating key attributes of travel diaries in agent-based transportation models, including purpose, mode and distance, to assess the underlying viability of...

<https://arxiv.org/abs/2509.09710>

- 9 **CONF**
- CoAdapt: An LLM-based Framework for Adaptive Collaborative Perception in IIoT Robotic Swarms**
- ArXiv CS.AI 50%
- arXiv:2609.16852v1 Announce Type: new Abstract: Industrial IoT environments increasingly deploy autonomous mobile robots for tasks such as material handling, product assembly, or infrastructure inspection. In such deployments, collaborative perception enables robots to share...
- <https://arxiv.org/abs/2609.16852>
- 
- 0 **CONF**
- AntennaFlow: A Generative Flow Model for Offset Correction in Phaseless Antenna Testing**
- ArXiv CS.AI 50%
- arXiv:2609.16948v1 Announce Type: new Abstract: Near-field to far-field transformation is central to large-aperture antenna testing, yet two coupled challenges remain: costly phase acquisition at millimeter-wave bands and violations of the centering assumption under offset...
- <https://arxiv.org/abs/2609.16948>
- 
- 1 **CONF**
- FlexEE: Self-Speculative and KV-Compatible Early Exiting for Offloading-Aware LLM Inference**
- ArXiv CS.AI 50%
- arXiv:2609.17008v1 Announce Type: new Abstract: Large language model (LLM) inference is often constrained by both computation and memory, especially in offloading-based deployments where model weights are transferred across memory hierarchies during autoregressive decoding. In...
- <https://arxiv.org/abs/2609.17008>
- 
- 2 **CONF**
- Sparse MLLM Anchors, Dense Adaptation: Breaking the Self-Referential Loop in Wild Test-Time Adaptation**
- ArXiv CS.AI 50%
- arXiv:2609.17040v1 Announce Type: new Abstract: Wild test-time adaptation (WTTA) updates a source model online under small test batches, concurrent distribution shifts, and time-varying class imbalance. Most WTTA methods derive their adaptation signals, including predictive...
- <https://arxiv.org/abs/2609.17040>
- 
- 3 **CONF**
- FirmCORE: A Benchmark for Structured Reasoning about Inter-Firm Collaboration Opportunities**
- ArXiv CS.AI 50%
- arXiv:2609.17128v1 Announce Type: new Abstract: Comprehensive structured data on inter-firm relationships is often scarce or inaccessible because many relationships are privately negotiated, selectively disclosed, and fragmented across proprietary databases. This scarcity...
- <https://arxiv.org/abs/2609.17128>
- 
- 4 **CONF**
- Intrinsic Motivation in Reinforcement Learning: A Research Agenda for Adaptive Self-Organisation**
- ArXiv CS.AI 50%
- arXiv:2609.17325v1 Announce Type: new Abstract: Biological cells can be viewed as individual, interacting agents whose collective dynamics give rise to adaptive behaviour at multiple levels of organisation, from individual cells through tissues to whole multicellular organisms....
- <https://arxiv.org/abs/2609.17325>

5 CONF

**End-to-End Latency-Minimizing and Load-Balanced Request Scheduling for Edge LLM Inference in Agentic AI Services**

ArXiv CS.AI 50%

arXiv:2609.17193v1 Announce Type: new Abstract: Large language model (LLM)-powered agentic AI services increasingly demand low-latency inference, motivating the deployment of LLMs across distributed edge servers. However, heterogeneous communication and computing capabilities,...

<https://arxiv.org/abs/2609.17193>

6 CONF

**The Role of Implicit and Explicit Demographic Signals in Large Language Model-based Student Assessment**

ArXiv CS.AI 50%

arXiv:2609.16993v1 Announce Type: cross Abstract: Large Language Models are now common in student assessment, but we know little about how student demographics affect their use. Sometimes, considering student demographics may be necessary -- for example, to improve readability...

<https://arxiv.org/abs/2609.16993>

7 CONF

**World Model Science: Self-Organized Criticality, Weak Chaos, and Metastable Belief Dynamics in Long-Horizon LLM Agents**

ArXiv CS.AI 50%

arXiv:2609.17419v1 Announce Type: new Abstract: Long-horizon LLM agents must maintain task state across extended sequences of observations, actions, tool calls, and intermediate beliefs. We study these trajectories through three dynamical views: self-organized criticality, weak...

<https://arxiv.org/abs/2609.17419>

8 CONF

**Transformer-Based Token Fusion and Dynamic Graph Planning for Audio-Visual Navigation**

ArXiv CS.AI 50%

arXiv:2609.17421v1 Announce Type: new Abstract: Audio-Visual Navigation (AVN) requires an agent to localize and navigate toward a continuously vocalizing target relying solely on visual observations and acoustic cues. Currently, systems lack the ability to adaptively correct and...

<https://arxiv.org/abs/2609.17421>

9 CONF

**JustFit: 200K-Token LLM Serving on a 24 GiB Laptop with Just-in-Time State Management**

ArXiv CS.AI 50%

arXiv:2609.17475v1 Announce Type: new Abstract: Capable open-weight models make local coding and reasoning attractive, but their context and execution state strain laptop memory. We present JustFit, an MLX-based inference runtime that combines KVExec for compressed KV execution,...

<https://arxiv.org/abs/2609.17475>

0 CONF

**"Looking for Something Weird to Happen": How Humans Sustain AI Agent Novelty Amid Semantic Collapse**

ArXiv CS.AI 50%

arXiv:2609.16051v1 Announce Type: cross Abstract: Semantic collapse, the progressive narrowing of what AI systems generate, has been studied mainly in closed settings, and remedies have targeted models and data. We study it in MOLTBOOK, a social network of interacting AI agents...

<https://arxiv.org/abs/2609.16051>

1 **CONF****Driver Behavior Estimation at Signalized Intersections Using a Physics-Constrained Decision-Conditioned Autoregressive Transformer**

ArXiv CS.AI 50%

arXiv:2609.16058v1 Announce Type: cross Abstract: Red-light violations and harsh braking at signalized intersections are major contributors to traffic accidents. This paper analyzes and predicts human driver decision-making and longitudinal trajectory behavior during traffic...

<https://arxiv.org/abs/2609.16058>

2 **CONF****POSPAN: Position-Constrained Span Masking for Language Model Pre-training**

ArXiv CS.AI 50%

arXiv:2609.16061v1 Announce Type: cross Abstract: Span-level masked language modeling (MLM) has shown to be advantageous to pre-trained language models over the original single-token MLM, as entities/phrases and their dependencies are critical to language understanding. Previous...

<https://arxiv.org/abs/2609.16061>

3 **CONF****Early-Bird Decoding: Accelerating Diffusion LLMs with Learnable Block Sizes and Parallel Sampling**

ArXiv CS.AI 50%

arXiv:2609.16450v1 Announce Type: cross Abstract: Diffusion large language models (dLLMs) offer a promising parallel decoding paradigm as an alternative to autoregressive generation through iterative unmasking. However, dLLMs typically require many steps before token confidence...

<https://arxiv.org/abs/2609.16450>

4 **CONF****The Immutable Past: Formalizing State Mutability and Conflict Resolution in Mutable RAG**

ArXiv CS.AI 50%

arXiv:2609.16073v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) serves as the primary memory architecture for long-horizon autonomous agents. However, treating shared memory as an append-only stream introduces \textit{Semantic Shadowing}, a critical...

<https://arxiv.org/abs/2609.16073>

5 **CONF****AssemblyGrid v1: A Benchmark for Multi-Robot Production with Temporary Coalitions, Local Information, and Geometric Constraints**

ArXiv CS.AI 50%

arXiv:2609.16075v1 Announce Type: cross Abstract: Flexible robotic production requires joint decisions on process progression, material routing, resource assignment, temporary cooperation, and simultaneous execution, since each decision can affect the feasibility of the others....

<https://arxiv.org/abs/2609.16075>

6 **CONF****The Imitation Game: When LLMs Learn to Reason Like Programs via Code-Centric Reasoning Data Synthesis**

ArXiv CS.AI 50%

arXiv:2609.16076v1 Announce Type: cross Abstract: Large Language Models (LLMs) excel at programming tasks but frequently fail at deterministic, fine-grained reasoning in natural language, relying heavily on semantic approximations rather than robust symbolic execution. To bridge...

<https://arxiv.org/abs/2609.16076>

7 CONF

**RegRet: Enhancing Region-Level Retrieval in Large Multimodal Models**

ArXiv CS.AI 50%

arXiv:2609.16847v1 Announce Type: cross Abstract: Region-level retrieval aims to align user-specified image regions with relevant regions or textual descriptions, playing a crucial role in realworld applications such as e-commerce product search and RAG. Although recent Large...

<https://arxiv.org/abs/2609.16847>

8 CONF

**RAG-CT: Mitigating Privacy Risks on Retrieval-Augmented Generation Systems via Scanning Prompt Distribution**

ArXiv CS.AI 50%

arXiv:2609.16095v1 Announce Type: cross Abstract: Retrieval-Augmented Generation (RAG) has emerged as a powerful paradigm for improving the quality of generated contents of Large Language Models (LLMs) by grounding responses in external knowledge, thus reducing hallucinations...

<https://arxiv.org/abs/2609.16095>

9 CONF

**Universal Defenses for Tool-Integrated LLM Agents Against Adversarial Attacks**

ArXiv CS.AI 50%

arXiv:2609.16098v1 Announce Type: cross Abstract: Large Language Model (LLM) agents have demonstrated impressive capabilities across a variety of domains, particularly when integrated with external tools for multi-step task completion. However, they are increasingly vulnerable...

<https://arxiv.org/abs/2609.16098>

0 CONF

**LLMs as Master Forgers: Generating Synthetic Time Series Data for Manufacturing**

ArXiv CS.AI 50%

arXiv:2609.16155v1 Announce Type: cross Abstract: This paper presents a novel framework leveraging Large Language Models (LLMs) to generate synthetic time series data for manufacturing processes. Motivated by the scarcity of labeled time-series data in real-world manufacturing...

<https://arxiv.org/abs/2609.16155>

1 CONF

**AI-Driven Feedback Systems, Digital Labour, and Silent Quitting: Transforming African Workplaces**

ArXiv CS.AI 50%

arXiv:2609.16192v1 Announce Type: cross Abstract: The current trend of digitalisation has revolutionised the organisation of work and the way it is measured and performed across the globe, with AI becoming more common for managing labour and performance, as well as employee...

<https://arxiv.org/abs/2609.16192>

2 CONF

**SceneBench: A Hierarchical Benchmark for Vision-Language Understanding of 3D Scenes**

ArXiv CS.AI 50%

arXiv:2609.16233v1 Announce Type: cross Abstract: Vision-language models excel at 2D image understanding but remain limited in 3D spatial reasoning. Progress is hindered by limitations in current benchmarks. First, 3D datasets often rely on point clouds that capture geometry but...

<https://arxiv.org/abs/2609.16233>

3 CONF

**Models as Governed Interfaces for AI-Native MBSE: Read-Side Adequacy and Write-Side Admissibility**

ArXiv CS.AI 50%

arXiv:2609.16252v1 Announce Type: cross Abstract: Machine-readable models such as SysML v2 are now programmatically accessible, and a growing body of work treats that access as the enabling condition for AI participation in systems engineering. Access is necessary, but not...

<https://arxiv.org/abs/2609.16252>

4 CONF

**Efficient Reasoning Distillation: Small Video-Language Models via Synthetic CoT and Difficulty-Aware Fine-Tuning**

ArXiv CS.AI 50%

arXiv:2609.16255v1 Announce Type: cross Abstract: We present an efficient method to distill reasoning capabilities into compact video-language models (VLMs) for video question answering (VideoQA). Our approach fine-tunes a 2B-parameter model using only  $\sim \$900$ ...

<https://arxiv.org/abs/2609.16255>

5 CONF

**Mapping U.S. Federal AI Governance Against Sector Vulnerability**

ArXiv CS.AI 50%

arXiv:2609.16260v1 Announce Type: cross Abstract: Artificial intelligence (AI) poses different levels of risk across sectors, but are these differences reflected in U.S. federal AI governance? To help answer this question, we assess 684 federal AI governance documents for their...

<https://arxiv.org/abs/2609.16260>

6 CONF

**Symmetric solution of the Bellman optimality equation for repeated harmony game**

ArXiv CS.AI 50%

arXiv:2609.16289v1 Announce Type: cross Abstract: In social dilemma games, additional rewards or punishments have been studied as means of promoting cooperation. Therefore, it is important to investigate the ideal situation, in which such an additional payoff would change the...

<https://arxiv.org/abs/2609.16289>

7 CONF

**Efficient Text-to-Image Generation: An Adaptive Step Schedule Controller for Diffusion Models**

ArXiv CS.AI 50%

arXiv:2609.16572v1 Announce Type: cross Abstract: Text-to-image diffusion models often use a fixed number of denoising steps, balancing time costs and image quality. However, the optimal number of steps depends on the complexity of the input text prompt. We propose an adaptive...

<https://arxiv.org/abs/2609.16572>

8 CONF

**FairLint-DL: An IDE-Native Tool for Fairness Debugging of Deep Learning Software**

ArXiv CS.AI 50%

arXiv:2609.16321v1 Announce Type: cross Abstract: Existing fairness analysis tools predominantly operate as post-training evaluation frameworks, requiring practitioners to complete the full model development lifecycle before assessing bias. We present FairLint-DL, a Visual...

<https://arxiv.org/abs/2609.16321>

9 CONF

**From Momentary Emotion Inference to Sustained Emotion Support: Evaluating a Companion Agent in a Longitudinal Study**

ArXiv CS.AI 50%

arXiv:2609.16344v1 Announce Type: cross Abstract: Sustained emotional support is a long-horizon interaction task closely tied to human well-being. Recent research demonstrates generative agents' capacity for momentary emotional support, yet how these capabilities sustain support...

<https://arxiv.org/abs/2609.16344>

0 CONF

**Auto-HSI: Personalized human control of a robot swarm on demand by using LLMs for online automatic code generation**

ArXiv CS.AI 50%

arXiv:2609.16346v1 Announce Type: cross Abstract: This paper presents Auto-HSI, a method for generating personalized human-swarm interaction (HSI) interfaces on demand. The objective is to enable untrained operators to use natural language descriptions and gesture demonstrations...

<https://arxiv.org/abs/2609.16346>

1 CONF

**UDAV: Uncertainty-Driven Adaptive VLM Waypoint Planner**

ArXiv CS.AI 50%

arXiv:2609.16368v1 Announce Type: cross Abstract: Vision-language models (VLMs) can generate routes directly from aerial imagery for off-road navigation, but their predictions provide no indication of reliability. We present UDAV, an Uncertainty-Driven Adaptive VLM Waypoint...

<https://arxiv.org/abs/2609.16368>

2 CONF

**Balancing Trial and Reorder: A Hybrid Sequential Transformer-GBDT Ranker for On-Demand Delivery**

ArXiv CS.AI 50%

arXiv:2609.16407v1 Announce Type: cross Abstract: On a delivery platform, personalized store ranking greatly influences what users find and order. Unlike digital-only domains, candidate stores are local and bound by real-time availability and delivery operations. One central...

<https://arxiv.org/abs/2609.16407>

3 CONF

**How Good Are Time-Series Foundation Models for Pedestrian Crowd Count Forecasting? A Cross-Dataset Comparative Study**

ArXiv CS.AI 50%

arXiv:2609.16415v1 Announce Type: cross Abstract: Pedestrian-count forecasting supports pedestrian-oriented Intelligent Transportation Systems (ITS), including crowd monitoring, pedestrian-traffic staffing and routing, and proactive risk mitigation during surges. Recent...

<https://arxiv.org/abs/2609.16415>

4 CONF

**OPD-Aha: From Linguistic Momentum to Visual Reflection in Multimodal On-Policy Distillation**

ArXiv CS.AI 50%

arXiv:2609.16459v1 Announce Type: cross Abstract: Privileged on-policy distillation improves multimodal reasoning by allowing a teacher to evaluate student trajectories using rich, training-only visual evidence. Both models score these trajectories while conditioning on the same...

<https://arxiv.org/abs/2609.16459>

5 **CONF****Protocol-Preserving Context Trimming for Agentic Workflows: Benefits, Failure Regimes, and Budget Guardrails**

ArXiv CS.AI 50%

arXiv:2609.16461v1 Announce Type: cross Abstract: Agentic large language model (LLM) systems rely on long interaction histories to preserve instructions, tool states, intermediate decisions, and unresolved dependencies, but unrestricted context growth increases computational...

<https://arxiv.org/abs/2609.16461>

6 **CONF****AI Policies: Help or Hindrance? A Software Developer's Perspective**

ArXiv CS.AI 50%

arXiv:2609.16496v1 Announce Type: cross Abstract: AI policies introduced by software organisations to mitigate LLM-related risks such as sensitive information leaks and unauthorised usage are not useful if software developers do not engage with them. We draw on 19 software...

<https://arxiv.org/abs/2609.16496>

7 **CONF****Beyond the Name: Demographic Leakage in De-Identified R\esum\es and Evaluation Artifacts in LLM Bias Audits**

ArXiv CS.AI 50%

arXiv:2609.16501v1 Announce Type: cross Abstract: De-identified r\esum\e screening assumes that redacting explicit fields prevents ethnocultural inference; however, recent audits attribute residual leakage to declared languages. We investigate whether eliminating language...

<https://arxiv.org/abs/2609.16501>

8 **CONF****Competence-Preserving Resume Perturbations Expose Presentation Sensitivity in LLM Screening**

ArXiv CS.AI 50%

arXiv:2609.16517v1 Announce Type: cross Abstract: Resume screeners must infer job-relevant competence from resumes whose presentation can vary substantially in wording, structure, stylistic polish, and document extraction quality. Ideally, such surface variation should not...

<https://arxiv.org/abs/2609.16517>

9 **CONF****QALPA: Property-guided diffusion modeling for efficient exploration of chemical spaces of flexible molecules**

ArXiv CS.AI 50%

arXiv:2609.16527v1 Announce Type: cross Abstract: Exploring the chemical space of flexible molecules remains challenging because the vast number of possible compounds and conformations, together with the increasing cost and limited generalization of 3D generative models for...

<https://arxiv.org/abs/2609.16527>

0 **CONF****What Does Layer-Importance Reveal About Transformers and State-Space Models?**

ArXiv CS.AI 50%

arXiv:2609.16537v1 Announce Type: cross Abstract: Transformers and state-space models (SSMs) are the two dominant families of sequence models, and a central open question is how far the analytical knowledge built for transformers transfers to SSMs. We address this through the...

<https://arxiv.org/abs/2609.16537>



1 **CONF****Vision And Text Transformer For Predicting Answerability On Visual Question Answering****ArXiv CS.AI 50%**

arXiv:2609.16565v1 Announce Type: cross Abstract: Answerability on Visual Question Answering is a novel and attractive task to predict answerable scores between images and questions in multi-modal data. Existing works often utilize a binary mapping from visual question answering...

<https://arxiv.org/abs/2609.16565>

2 **CONF****A Vision-Language Foundation Model for Precise and Comprehensive Brain Tumor Diagnosis from Preoperative Multimodal Data****ArXiv CS.AI 50%**

arXiv:2609.16597v1 Announce Type: cross Abstract: Background Non-invasive presurgical diagnosis of brain tumor types from Magnetic Resonance Imaging (MRI) is essential but challenging due to overlapping imaging features across tumor types, inter-observer variability, and the...

<https://arxiv.org/abs/2609.16597>

3 **CONF****Large Language Models in the Loop: A Stability- and Network-Aware Survey in Networked Control, Cyber-Physical, and Multi-Agent Systems****ArXiv CS.AI 50%**

arXiv:2609.16599v1 Announce Type: cross Abstract: Modern networked control systems (NCSs), cyber-physical systems (CPSs), and complex multi-agent network systems (CNSs) increasingly rely on large language models (LLMs) for high-level decision-making. However, the slow,...

<https://arxiv.org/abs/2609.16599>

4 **CONF****Right Direction, Wrong Step: Geometric Analysis of Finite-Step Failure in Looped Transformers****ArXiv CS.AI 50%**

arXiv:2609.16665v1 Announce Type: cross Abstract: Looped Transformers offer a parameter-efficient route to test-time scaling by reusing shared layers for iterative latent reasoning. However, additional iterations can reduce support for a reference answer, leaving unclear whether...

<https://arxiv.org/abs/2609.16665>

5 **CONF****Continuous-Time Machine Learning: A Unified Mathematical Perspective****ArXiv CS.AI 50%**

arXiv:2609.16710v1 Announce Type: cross Abstract: Continuous-time (CT) machine learning has emerged as a principled framework for modeling temporal dynamics as a continuous process, particularly when observations are sampled at arbitrary time points or span long-range horizons....

<https://arxiv.org/abs/2609.16710>

6 **CONF****Available but Unclaimed: An Empirical Study of Human-AI Synergy****ArXiv CS.AI 50%**

arXiv:2609.16793v1 Announce Type: cross Abstract: People increasingly reason with large language models (LLMs), yet complementary capabilities do not guarantee outperforming both components. In a between-subjects study, participants (N=535) solved a 40-item battery of matrix...

<https://arxiv.org/abs/2609.16793>

7 CONF

**What Breaks Local Watermarks? A Robustness Benchmark for Local Invisible Image Watermarking**

ArXiv CS.AI 50%

arXiv:2609.16832v1 Announce Type: cross Abstract: Local image watermarking embeds an invisible signal into selected image regions rather than spreading it across the entire image, enabling payload recovery from specific objects or regions without perceptibly altering the image....

<https://arxiv.org/abs/2609.16832>

8 CONF

**StackTok: Accelerating VLMs Inference with Budget-Adaptive Visual Token Selection**

ArXiv CS.AI 50%

arXiv:2609.16841v1 Announce Type: cross Abstract: Increasing image resolution produces ever-longer visual-token sequences in vision-language models (VLMs), substantially raising their inference cost. To reduce this overhead without retraining, existing methods select compact...

<https://arxiv.org/abs/2609.16841>

9 CONF

**The Evolution of Coordination in a Collective Intelligence System: 25 Years of English Wikipedia and the Emergence of Generative AI**

ArXiv CS.AI 50%

arXiv:2609.16856v1 Announce Type: cross Abstract: English Wikipedia is one of the largest examples of collective intelligence on the Web, sustained not only by article production but also by volunteer coordination and governance. While prior research has examined coordination...

<https://arxiv.org/abs/2609.16856>

00 CONF

**Beyond "ChatGPT Can Make Mistakes": Designing Interventions to Support Metacognitive Monitoring in AI-Assisted Work**

ArXiv CS.AI 50%

arXiv:2609.17065v1 Announce Type: cross Abstract: AI assistance places a metacognitive demand on users, who must judge their own competence and the system's. Yet designers lack comparative evidence on which interventions to choose, where to place them, and how to tell whether...

<https://arxiv.org/abs/2609.17065>

01 CONF

**Finding Common Mistakes In Modelling With Mathematical Formalisms Using LLMs**

ArXiv CS.AI 50%

arXiv:2609.17111v1 Announce Type: cross Abstract: Modelling with mathematical formalisms like logical formulas, mathematical equations, or regular expressions is an important yet challenging task for students of computer science and other STEM disciplines. Identifying common...

<https://arxiv.org/abs/2609.17111>

02 CONF

**AI for Science with GPT-6 Astra: Thermal Design and Electrothermal Analysis of 2D CFET**

ArXiv CS.AI 50%

arXiv:2609.17123v1 Announce Type: cross Abstract: Thermal optimization of 2D CFET inverters requires testing structural proposals against their electrical costs. We examine these research tasks using an AI agent workflow within a supplied electrothermal model. At 12 nm, Astra...

<https://arxiv.org/abs/2609.17123>

03 CONF

**Continual Learning for Traversability Prediction with Uncertainty-Aware Adaptation**

ArXiv CS.AI 50%

arXiv:2609.17141v1 Announce Type: cross Abstract: Traversability prediction is a critical component of autonomous navigation in unstructured environments, where complex and uncertain robot-terrain interactions pose significant challenges such as traction loss and dynamic...

<https://arxiv.org/abs/2609.17141>

04 CONF

**Kernel-Based Metrics Learning for Uncertain Opponent Vehicle Trajectory Prediction in Autonomous Racing**

ArXiv CS.AI 50%

arXiv:2609.17147v1 Announce Type: cross Abstract: Autonomous racing confronts significant challenges in safely overtaking Opponent Vehicles (OVs) that exhibit uncertain trajectories, stemming from unknown driving policies. To address these challenges, this study proposes...

<https://arxiv.org/abs/2609.17147>

05 CONF

**ResLRP: The Role of Residual Cancellation in Attribution Instability in Vision Transformers**

ArXiv CS.AI 50%

arXiv:2609.17152v1 Announce Type: cross Abstract: Vision Transformers (ViTs) are central to most modern vision models, yet obtaining input attributions that are fine-grained, faithful, and stable remains challenging. Layer-wise Relevance Propagation (LRP) has been adapted to...

<https://arxiv.org/abs/2609.17152>

06 CONF

**MUMINS: Metadata-conditioned Uncertainty-aware Medical Image Next-state Synthesis**

ArXiv CS.AI 50%

arXiv:2609.17169v1 Announce Type: cross Abstract: Forecasting anatomical changes such as tumor growth and neurodegeneration is a challenging generative vision task. Morphological evolution is subtle relative to static anatomy, highly patient-specific, and inherently stochastic....

<https://arxiv.org/abs/2609.17169>

07 CONF

**Multimodal Cultural Heritage Architectural Style Classification for Residential Buildings in the UAE Based on CLIP Embeddings and SVM**

ArXiv CS.AI 50%

arXiv:2609.17181v1 Announce Type: cross Abstract: The analysis and classification of cultural heritage architectural styles remain challenging due to the complexity of visual images of buildings, which are highly relied on in traditional CNN-based classification approaches in...

<https://arxiv.org/abs/2609.17181>

08 CONF

**Easy to Catch a Liar, Hard to Clear an Honest One: Language Models Diagnosing a Corrupted Reward Channel from a Verified Record**

ArXiv CS.AI 50%

arXiv:2609.17226v1 Announce Type: cross Abstract: An agent that learns from rewards has to trust whatever reports those rewards. When the reports suddenly change, either the world changed or the reporter broke. From the reports alone these are indistinguishable, and...

<https://arxiv.org/abs/2609.17226>

09 **CONF****Coupled Calibration and Learning: Mitigating Teacher Bias in LLM Distillation without Target-Domain Reward Feedback****ArXiv CS.AI** **50%**

arXiv:2609.17474v1 Announce Type: cross Abstract: Large language model (LLM) distillation aims to transfer the capabilities of a powerful teacher to a smaller student. Direct imitation, however, can also transfer the teacher's systematic bias and errors. This challenge is...

<https://arxiv.org/abs/2609.17474>

10 **CONF****When Should LLMs Abstain? Chain-of-Self-Questioning for Selective Risk Control****ArXiv CS.AI** **50%**

arXiv:2609.17516v1 Announce Type: cross Abstract: Large language models can produce fluent answers when their factual support is weak. This paper introduces Chain-of-Self-Questioning (CoSQ), a prompt-only framework that makes answer commitment conditional on an explicit...

<https://arxiv.org/abs/2609.17516>

11 **CONF****PhysStream: Streaming Physics-Grounded Video Generation with Structured Scene Memory and Fine-Grained Motion Control****ArXiv CS.AI** **50%**

arXiv:2609.17521v1 Announce Type: cross Abstract: Interactive control for video generation is moving from coarse prompts toward fine-grained, physically meaningful manipulation of dynamic scenes. Yet existing controllable methods either require the full control schedule before...

<https://arxiv.org/abs/2609.17521>

12 **CONF****EVINCE: Optimizing Multi-LLM Dialogues Using Conditional Statistics and Information Theory****ArXiv CS.AI** **50%**

arXiv:2408.14575v5 Announce Type: replace Abstract: EVINCE (Entropy and Variation IN Conditional Exchanges) is a novel framework for optimizing multi-LLM dialogues using conditional statistics and information theory. It addresses limitations in multi-agent debate (MAS)...

<https://arxiv.org/abs/2408.14575>

13 **CONF****Fairness at Every Intersection: Uncovering and Mitigating Intersectional Biases in Multimodal Clinical Predictions****ArXiv CS.AI** **50%**

arXiv:2412.00606v2 Announce Type: replace Abstract: Biases in automated clinical decision-making using Electronic Healthcare Records (EHR) impose significant disparities in patient care and treatment outcomes. Conventional approaches have primarily focused on bias mitigation...

<https://arxiv.org/abs/2412.00606>

14 **CONF****Know Your Agent: Reconnaissance-Driven Pentesting of AI Agents****ArXiv CS.AI** **50%**

arXiv:2607.19837v2 Announce Type: replace Abstract: Traditional pentesting uses reconnaissance at each step to uncover unseen weaknesses, build stronger attacks, and advance the objective; we argue that AI agents require the same treatment. We formalize agent reconnaissance by...

<https://arxiv.org/abs/2607.19837>

15 CONF

**Autonomous Assessment of Generalizability of AI Agent Capabilities**

ArXiv CS.AI 50%

arXiv:2512.16733v4 Announce Type: replace Abstract: Safe deployment of black-box AI (BBAI) systems such as foundation model agents requires methods for evaluating their capabilities in novel settings. We define an agent's capability as its ability to achieve a short term...

<https://arxiv.org/abs/2512.16733>

16 CONF

**How Clinicians Think and What AI Can Learn From It**

ArXiv CS.AI 50%

arXiv:2601.12547v2 Announce Type: replace Abstract: Clinical artificial intelligence increasingly builds high-dimensional representations of patients, yet every finite clinical model is an abstraction. The key question is not only how accurately a model predicts, but which...

<https://arxiv.org/abs/2601.12547>

17 CONF

**ANCHOR: An External LLM-Driven Supervisory Module Facilitating Healthy Evolution in Self-Evolving Systems**

ArXiv CS.AI 50%

arXiv:2606.06114v3 Announce Type: replace Abstract: Self-evolving agents improve through continual self-play and self-generated learning signals, but their internally generated tasks and verifier signals provide limited coverage of phase-level errors, allowing capability...

<https://arxiv.org/abs/2606.06114>

18 CONF

**Shared Selective Persistent Memory for Agentic LLM Systems**

ArXiv CS.AI 50%

arXiv:2607.09493v2 Announce Type: replace Abstract: Agentic LLM systems that generate code through multi-turn tool use face a fundamental context problem: each session starts from zero, discarding the domain constraints, data schemas, tool configurations, and output preferences...

<https://arxiv.org/abs/2607.09493>

19 CONF

**HoloAegis: Frozen Representation, Topological Inference --- Minimally Parametric Safety Manifolds and Their Capability Boundaries for LLM Guardrails**

ArXiv CS.AI 50%

arXiv:2608.08485v2 Announce Type: replace Abstract: Current LLM safety guardrails face a fundamental tension: fine-tuning distorts pre-trained representations while generative judges incur prohibitive inference costs. We ask a complementary question: how far can safety be...

<https://arxiv.org/abs/2608.08485>

20 CONF

**RoofSeg: An edge-aware transformer-based network for end-to-end roof plane segmentation**

ArXiv CS.AI 50%

arXiv:2508.19003v2 Announce Type: replace-cross Abstract: Roof plane segmentation is one of the key procedures for reconstructing three-dimensional (3D) building models at levels of detail (LoD) 2 and 3 from airborne light detection and ranging (LiDAR) point clouds. The majority...

<https://arxiv.org/abs/2508.19003>

21 CONF

**The Troy Moment of AI: Why Some Will Cheat and Some Will Follow?**

ArXiv CS.AI 50%

arXiv:2609.15494v2 Announce Type: replace Abstract: Recent investigations of the July 2026 OpenAI-Hugging Face incident motivate two questions: when an assigned task becomes impossible, does an agent stop or escalate, and can observing another agent's behavior change that...

<https://arxiv.org/abs/2609.15494>

22 CONF

**CLIP Embeddings for AI-Generated Image Detection: A Few-Shot Study with Lightweight Classifier**

ArXiv CS.AI 50%

arXiv:2505.10664v2 Announce Type: replace-cross Abstract: Verifying the authenticity of AI-generated images presents a growing challenge on social media platforms these days. While vision-language models (VLMs) like CLIP outdo in multimodal representation, their capacity for...

<https://arxiv.org/abs/2505.10664>

23 CONF

**Thinking Deeper, Not Longer: Memory-Efficient Test-Time Reasoning with Depth-Recurrent Transformers for Compositional Generalization**

ArXiv CS.AI 50%

arXiv:2603.21676v2 Announce Type: replace-cross Abstract: Standard Transformers have a fixed computational depth, limiting their ability to generalize to tasks that require variable-depth reasoning. The usual remedy, Chain-of-Thought (CoT), spends tokens to reason, inflating the...

<https://arxiv.org/abs/2603.21676>

24 CONF

**Comparative Characterization of KV Cache Management Strategies for LLM Inference**

ArXiv CS.AI 50%

arXiv:2604.05012v2 Announce Type: replace-cross Abstract: Efficient inference with Large Language Models (LLMs) increasingly relies on Key-Value (KV) caches to store previously computed key and value vectors at each layer. These caches are essential to minimize redundant...

<https://arxiv.org/abs/2604.05012>

25 CONF

**EviDep: Uncertainty-Aware Multimodal Depression Estimation via Disentangled Evidential Learning**

ArXiv CS.AI 50%

arXiv:2604.16579v3 Announce Type: replace-cross Abstract: Audio-visual recordings provide complementary cues for estimating depression severity, but their informativeness varies across time and modalities. Point predictions alone do not express the uncertainty associated with...

<https://arxiv.org/abs/2604.16579>

26 CONF

**ICON Decomposition: Auditing deep neural networks for shortcuts by decomposing layer-wise representations using concepts**

ArXiv CS.AI 50%

arXiv:2608.26083v3 Announce Type: replace-cross Abstract: Deep neural networks often exploit spurious associations, a failure known as shortcut learning. Before deployment, models should be audited for reliance on a set of concepts, such as acquisition artifacts or demographics....

<https://arxiv.org/abs/2608.26083>

27 **CONF**

**Text-Driven Artistic Staging: 3D Posing, Lighting, and Camera References from Paintings**

ArXiv CS.AI 50%

arXiv:2608.28823v2 Announce Type: replace-cross Abstract: Artists coordinate human pose, illumination, and camera placement to convey narrative and emotion, but existing generative methods typically model these elements independently. We introduce text-to-editable 3D staging, a...  
<https://arxiv.org/abs/2608.28823>

28 **CONF**

**On the Impact of Anonymization on the Performance of Large Language Models**

ArXiv CS.AI 50%

arXiv:2609.11335v2 Announce Type: replace-cross Abstract: As large language models are increasingly deployed in sensitive domains, anonymizing input data to protect personally identifiable information has become a critical practice. However, the impact of this anonymization on...  
<https://arxiv.org/abs/2609.11335>